

MINDCRAFT

Research Constitution

How do humans become confident mathematical thinkers?

Identity transformation · Evidence · Red Team · Product systems · Experiments

This is not a pitch deck. This is the company's research operating system.
Challenge assumptions. Label uncertainty. Kill weak arguments.
Source corpus: ~15,770 words · multi-chapter living lab

MindCraft Research Lab · ongoing evidence program

MindCraft Research Constitution v1

Status: Living operating document — not a pitch deck

Edition: v1.5 (multi-chapter lab; ongoing evidence program)

Research question: How do humans become *confident mathematical thinkers*?

Product thesis under audit: The product is identity transformation, not mathematics delivery.

Last updated: 2026-07-26

Growth model: Core OS (this file) + chapters/*.md via CHAPTER_MANIFEST.txt → PDF

Scale intent: Multi-month densification toward 150–300 pages of *evidenced* material — never fluff

Epistemic rule: Every claim is labeled FACT / HYPOTHESIS / FOUNDER BELIEF / SPECULATION.

Deep-dive chapters currently mounted

Part	File	Focus
XXIV	chapters/24_affect_anxiety_belonging.md	Anxiety, WM, stereotype threat, belonging
XXV	chapters/25_self_efficacy_narrative_identity.md	Bandura sources, narrative identity, habit≠identity
XXVI	chapters/26_cognitive_load_mastery_tutoring.md	CLT, fading, mastery war, tutoring calibration
XXVII	chapters/27_markets_parents_competition.md	WTP, parents, competitive teardown
XXVIII	chapters/28_history_meaning_math.md	Invention-story thesis under trial
XXIX	chapters/29_spacing_retrieval.md	Spacing, retrieval, memory of competence
XXX	chapters/30_attribution_helplessness.md	Weiner attributions, helplessness, feedback language
XXXI	chapters/31_flow_challenge_skill.md	Flow, challenge–skill, difficulty design
XXXII	chapters/32_parent_anxiety_transmission.md	Maloney/Beilock parent anxiety pathway
XXXIII	chapters/33_ai_tutors_trust_sycophancy.md	AI trust / sycophancy; RT: Bastani■PWC, guards≠gain

Queued next: see NEXT_LAB.md (causal DAG, competitive audits, equity audit, expectancy-value).

0. How to read this document

This Constitution is the research OS for MindCraft. It exists to prevent the company from building features that feel right and fail quietly.

Rules of engagement:

1. Founder beliefs enter as **FOUNDER BELIEF**, not as law.
2. The Red Team’s job is destruction. Unchallenged claims are invalid.
3. Page count is not quality. A shorter true section beats a longer vague one.
4. “AI will make explanations free” is a hypothesis about markets, not a moral slogan.
5. Product implications must map to experiments with falsifiers.

This v1 is dense and incomplete on purpose. Expanding toward 150–300 pages is allowed only when new evidence, frameworks, or contradicted priors justify the ink.

Part I — Executive Summary

I.1 The compressed answer

HYPOTHESIS (confidence: medium–high): Humans become confident mathematical thinkers when three conditions compound:

1. **Affective permission** — the nervous system stops treating math as social/physical threat.
2. **Competence evidence** — the learner accumulates undeniable, self-attributed successes at the right grain of difficulty.
3. **Identity re-storying** — the person revises the narrative “I am not a math person” into a workable self-story compatible with struggle.

Explanations alone rarely produce (1)–(3). Explanations can accelerate (2) *after* (1) is present. This matches the founder’s tutoring observation and is directionally supported by math-anxiety, SDT, and mindset literatures — with important contradictions (see Red Team).

I.2 Verdict on the scarcity thesis

FOUNDER BELIEF under audit: AI commoditizes knowledge; scarcity shifts to motivation, confidence, identity, trust, curiosity, persistence, emotional safety, meaning.

Red Team verdict (provisional):

Claim	Status	Notes
Explanations are getting cheaper/faster	FACT (directionally)	Generative AI + open content lowers marginal cost of explanation

“Tutoring is free”	FALSE as stated	High-quality <i>human</i> attention, accountability, and diagnosis remain scarce and expensive
Content is free	MOSTLY FACT for commodity content	Differentiation migrates to curation, sequencing, trust, outcomes
Motivation/identity become scarce	HYPOTHESIS	Likely true as <i>product differentiators</i> ; not proven as <i>sole moat</i>
Emotional safety is first-order	HYPOTHESIS with supportive evidence	Math anxiety impairs working memory; interventions that ignore affect underperform for anxious learners

Implication: MindCraft should not bet the company on “better explanations.” It should bet on a **system that converts fear → evidence → identity**, with explanations as infrastructure, not hero feature.

I.3 What to optimize (North Star debate)

Metric	Predicts short-term engagement?	Predicts durable math identity?	Risk
Correct answers	Medium	Low–medium	Teaches gaming / guessing
Time spent	High (easy to inflate)	Low	Engagement theater
Streaks / DAU	High	Uncertain	Duo-style habit ≠ identity
Confidence (self-report)	Medium	Medium	Cheap talk, demand effects
Challenge-seeking (advanced course taking)	Medium	High (Yeager et al., Nature 2019 signal)	Hard to measure early
Voluntary return after failure	Medium	High (HYPOTHESIS)	Needs instrumentation
“I am a math person” endorsement + behavior	Medium	Highest target	Slow, noisy

Working North Star (HYPOTHESIS):

Challenge-seeking under safety — the student chooses a harder problem, returns after a miss, and attributes success to strategy/effort *in math specifically*.

Part II — Key Insights (challenged)

Insight 1 — Fear is a performance bottleneck, not a personality trait

FACT: Math anxiety is associated with reduced working-memory availability during math tasks; anxious learners underperform relative to ability (Ashcraft & Kirk line of work; broader anxiety–WM literature).

HYPOTHESIS: For a large subset of MindCraft’s Maya archetype, the binding constraint is not “missing explanation,” it is **threat appraisal** (“I will look stupid”).

Implication: Soft-wrong UX, hide-correctness diagnostics, wizard coaching, and narrative framing are not polish — they are load-bearing if the thesis holds.

Red Team: Many high performers succeed under stress. Anxiety interventions have heterogeneous effects. Do not medicalize ordinary difficulty.

Insight 2 — Growth mindset is real, small, and context-dependent

FACT: Yeager et al. (2019, *Nature*) — brief online growth-mindset intervention improved grades for lower-achieving students and increased advanced math enrollment in a nationally representative U.S. sample; effects stronger when peer norms supported challenge-seeking.

FACT / REVIEW: Domain-general mindset messaging shows mixed results; math-*specific* mindset interventions more often report positive outcomes (2023 systematic review in *Educational Research Review*).

FACT / CAUTION: “False growth mindset” (praise effort without teaching strategy / changing environment) is harmful (Dweck’s own caution; Kohn’s critique).

Implication: MindCraft must not ship empty “believe in yourself” copy. Mindset must be embodied in **task design, feedback, and tutor norms**.

Insight 3 — Autonomy, competence, relatedness are the psychological tripod

FACT: Self-Determination Theory (Ryan & Deci) — sustained intrinsic motivation requires autonomy, competence, relatedness.

FACT (meta): SDT-based education interventions show reliable gains for autonomy and competence support; relatedness effects are less consistent (Wang et al. 2024 meta-analysis).

Implication: Product loops must deliver:

- Autonomy: meaningful choice (path, story world, pace)
- Competence: mastery grain that produces earned wins
- Relatedness: tutor presence, parent signal, peer-safe identity (hardest; do not fake with empty social)

Insight 4 — Mastery learning works — until you measure the wrong thing

FACT: Kulik, Kulik & Bangert-Drowns (1990) meta-analysis — mastery programs show positive exam effects (overall ~0.5 SD in their synthesis), stronger for weaker students; can increase time; self-paced college mastery can

reduce completion.

FACT / CONTRADICTION: Slavin’s “Mastery Learning Reconsidered” (1987) — little support for group-based mastery on *standardized* measures; experimenter-made tests show moderate effects; coverage vs mastery tradeoff is real.

Implication: MindCraft’s mastery graph is promising **if** mastery means transferable competence, not local quiz familiarity. Measure far transfer, not only same-item success.

Insight 5 — Habit products ≠ identity products

FACT / INDUSTRY: Duolingo’s public method emphasizes gamification for motivation and return (streaks, leagues, variable rewards) (Duolingo Method whitepaper, 2023). Streaks exploit loss aversion; they drive DAU.

HYPOTHESIS: Streaks can increase practice frequency without changing “I am a math person.” Habit without identity yields brittle engagement (stop streak → stop learning).

Implication: Use habit mechanics as **on-ramps**, not as the destination. Pair with identity-marking rituals (map fill, story progression, tutor recognition of growth).

Insight 6 — Narrative is not decoration; it is encoding + meaning

FOUNDER BELIEF (WORLD_VISION): Math stripped of invention-story becomes alienating machinery.

HYPOTHESIS: Historical/human invention frames increase curiosity and reduce “arbitrary rule” appraisals, improving persistence.

Evidence status: Strong tradition in math education (history of math pedagogy); fewer large RCTs than mindset literature. Treat as **design prior + research program**, not settled science.

Insight 7 — AI commoditizes *answers*, not *becoming*

HYPOTHESIS: As answer generation approaches free, willingness-to-pay migrates to:

- 1. Diagnosis of the *actual* break (gap map)
- 2. Emotional containment during struggle
- 3. Accountability relationships (tutor)
- 4. Identity-consistent practice worlds
- 5. Trusted outcome signaling to parents

SPECULATION: In 30 years, “show me the steps” is ambient. “Help me become the kind of person who stays with hard problems” remains a human+system problem.

Part III — Research Log (v1 seed)

Date	Finding	Type	Action
------	---------	------	--------

2026-07-25	Constituted Research Lab + Constitution v1	Process	Create OS
2026-07-25	Scarcity thesis partially false as slogans; true as product wedge	Analysis	Rewrite pitch language
2026-07-25	Yeager 2019 constrains mindset claims	Evidence	Ban empty mindset copy
2026-07-25	Soft-wrong + wizard coach shipped in product	Product	Instrument coach → retry rate
2026-07-25	Mastery literature split (Kulik vs Slavin)	Evidence	Define mastery measurement carefully
2026-07-25	Bloom 2-sigma overstated vs VanLehn / modern tutoring metas	Evidence	Stop marketing 2-sigma
2026-07-25	Founder sequence wounded → SAFE-CRAFT	Synthesis	Attempt earlier; invention-story optional
2026-07-25	Deliberate practice necessary≠sufficient (Hambrick line)	Evidence	Progress-vs-self framing

Part IV — Evidence Table (selected)

Claim	Best supporting evidence	Contradicting / limiting evidence	Confidence	MindCraft implication
Brief mindset intervention can move grades + advanced math taking	Yeager et al., Nature 2019	Effects heterogeneous; peer norms matter; replication debates in broader mindset meta-analyses	Medium–High	Mindset must be ecological, not banner text
Math-domain mindset > generic	2023 Ed Research Review systematic review	Study quality varies; publication bias possible	Medium	Math-specific identity language
SDT need support raises motivation	Wang et al. 2024 meta	Relatedness effects weaker	Medium–High	Design for autonomy + competence first

Math anxiety impairs performance	Cognitive literature (Ashcraft tradition)	Not all low performers are anxious	Medium–High	Affective onboarding is first-class
Mastery learning raises achievement	Kulik et al. 1990	Slavin 1987 on standardized tests	Medium	Mastery graph + transfer tests
Gamified habits raise return	Duolingo method / industry practice	Habit ≠ deep learning; extrinsic traps (Deci)	Medium	Streaks as on-ramp only
1:1 tutoring is uniquely powerful	Bloom “2 sigma” framing (1984)	Often overstated; quality varies; VanLehn etc. nuance tutoring effects	Medium	Keep human tutors; AI amplifies, not replaces

Part V — Contradictory Evidence (do not skip)

1. **Mindset skepticism:** Some metas find small average effects; school implementations often degrade into posters.
2. **Grit / resilience hype:** Duckworth-style grit is contested; structural inequality and curriculum quality matter more than character slogans.
3. **Discovery learning failures:** Pure constructivism without guidance often fails (Kirschner, Sweller, Clark critiques). Narrative worlds must still teach.
4. **Engagement ≠ learning:** Time-on-app can rise while understanding falls (edtech’s original sin).
5. **AI may increase explanation value temporarily:** Novices may trust fluent wrongness; trust becomes the scarce resource, not explanation volume.
6. **Some students want speed and drills:** Identity-through-story is not universal. Offer multiple vessels.

Part VI — Mental Models

VI.1 The Identity Transformation Cascade (ITC)

HYPOTHESIS — MindCraft original synthesis

Threat ↓ → Attention available → Micro-success → Attribution ("I figured it")
→ Willingness to re-enter challenge → Accumulated competence evidence
→ Narrative update ("maybe I can") → Identity claim ("I do math")
→ Environment selection (harder courses, peers, tutors)
→ Reinforcing loop

Balancing loop: Premature hard problems → threat ↑ → avoidance → identity freeze.

Product translation: Gap scan (no shame) → soft-wrong → tiny earned wins → map fill → tutor recognition → advanced path.

VI.2 The Three Doors (what students actually need in a stuck moment)

Door	Need	Wrong product response	Right response
A	Safety	“Here’s a longer explanation”	Contain affect; normalize miss
B	Clarity	Pep talk	Precise model / worked example at right grain
C	Agency	Take over and solve for them	Scaffolded choice; student does the decisive step

FOUNDER BELIEF: Tutors often open Door B first. Many students are still behind Door A.

VI.3 Knowledge commodity stack (post-AI)

Layer 0: Facts & procedures → commoditizing fast
 Layer 1: Explanations → commoditizing fast
 Layer 2: Diagnosis of misconceptions → partially automatable; trust-sensitive
 Layer 3: Adaptive sequencing → automatable with good ontology
 Layer 4: Emotional co-regulation → scarce (human + careful UX)
 Layer 5: Identity & meaning → scarce
 Layer 6: Accountability relationship → scarce

MindCraft’s durable stack lives in **Layers 2–6**, with 0–1 as utilities.

VI.4 Transfer from outside education

Domain	Mechanism	Transfer to MindCraft
Games (Celeste, FromSoftware-lite design)	Fair difficulty + death as information	Soft-wrong as physics, not moral failure
Sports coaching	Film study of <i>specific</i> error	Gap map + wrong-choice coach
Music	Scales before repertoire; recital as identity	Concept drills → story quest → public signal
Therapy (exposure)	Graded exposure to feared stimulus	Math anxiety: graded challenge under safety
Aviation checklists	Externalize working memory under stress	Procedure scaffolds when anxious
Martial arts belts	Visible competence ladder	Mastery pips / map regions

Language apps	Habit loops	Daily practice — without Duo’s extrinsic ceiling
Religion / ritual	Shared meaning + belonging	Caution: do not cultify; use <i>light</i> ritual (map lighting)

Part VII — Universal Learning Principles (provisional)

1. **Emotion gates cognition** in threatened domains.
 2. **Earned difficulty** beats random difficulty.
 3. **Feedback should be information, not identity judgment.**
 4. **Attribution matters:** strategy > talent for growth; luck attributions kill agency.
 5. **Spaced retrieval** beats massed re-reading (cognitive science consensus).
 6. **Worked examples** → **fading guidance** for novices (cognitive load).
 7. **Belonging uncertainty** destroys persistence for stereotyped groups (Walton/Cohen tradition).
 8. **Transfer requires varied practice**, not identical clones.
 9. **Human accountability** multiplies AI tools.
 10. **Identity is a lagging indicator**; design leading indicators (challenge-seeking, return-after-fail).
-

Part VIII — Product Implications (systems, not features)

VIII.1 Core system: Fear → Evidence → Identity (FEI Loop)

Reinforcing loop R1 (desired):

Safety UX → attempt → informative miss → coach → success → map fill → identity → more attempts.

Balancing loop B1 (threat):

Public shame / red buzz → avoidance → less practice → skill lag → more shame.

MindCraft already partially implements FEI: hide-correctness diagnostic, soft-wrong, wizard coach, story chapters, knowledge map. **Missing instrumentation:** does coach raise retry rate and later challenge-seeking?

VIII.2 What to stop optimizing

- Vanity DAU without learning transfer
- Explanation length as quality proxy
- Points that do not mark identity-relevant milestones

VIII.3 What to instrument next (minimum viable science)

1. **Retry rate within 2 minutes of soft-wrong** (with vs without wizard coach)
2. **Challenge-seeking:** % choosing harder level when offered
3. **Return after fail session** (D1 retention after a loss session)
4. **Self-identity item** (pre/post, 1–2 items, math-specific)
5. **Advanced course intent / enrollment** (slow, gold)

VIII.4 Tutor system role

Tutors are not content pipes. They are:

- Affect regulators
- Attribution coaches
- Bridge diagnosticians (concept A known, link A→B broken)
- Identity witnesses (“I saw you become someone who stays”)

AI should brief tutors, not replace witnessing.

Part IX — Experiments (pre-registered style)

Experiment A — Wizard coach on soft-wrong

- **Question:** Does the under-Graph wizard increase productive retry vs soft-wrong alone?
- **Design:** A/B, chapter + practice
- **Primary:** retry within 120s; secondary: eventual correct without hint binge
- **Falsifier:** no lift; or lift only in engagement with worse accuracy

Experiment B — Story-first vs bare stem

- **Question:** Does invention-story framing raise persistence on first miss?
- **Design:** within-concept crossover
- **Primary:** time-to-abandon after first miss
- **Falsifier:** story slows without improving retention/transfer

Experiment C — North Star validation

- **Question:** Which early metric best predicts 8-week challenge-seeking?
- **Candidates:** streak length, soft-wrong retry, identity item, map mastery Δ
- **Method:** observational + lagged prediction

Experiment D — Explanation timing

- **Question:** Is explanation-before-attempt worse than attempt-with-safety for anxious students?
- **Design:** screen with anxiety pretest; randomize explanation-first vs attempt-first
- **Falsifier:** explanation-first wins for all segments

Part X — Weekly Priorities (research × product)

1. Instrument FEI loop events in analytics (retry, coach shown, write-mode exit).
 2. Run Experiment A for 2 weeks.
 3. Red Team the WORLD_VISION narrative claims with one outside learning scientist.
 4. Interview 10 Mayas: “When did math stop feeling dangerous?” (qual).
 5. Define “mastery” operationally so Slavin’s critique cannot gut the graph.
-

Part XI — Founder Questions (do not dodge)

1. If ChatGPT tutors become “good enough,” what *relationship* do parents still buy?
 2. Are we building for Maya’s identity — or for parents’ anxiety about scores? Both? Tradeoff?
 3. Would we accept lower short-term accuracy for higher long-term challenge-seeking?
 4. Is the story world load-bearing, or is gap+tutor enough?
 5. What would falsify the identity-transformation thesis in 90 days?
-

Part XII — Long-term Roadmap (research horizons)

H1 (now): Prove FEI loop moves retry + challenge-seeking.

H2: Prove identity language + map produce durable self-concept change.

H3: Prove human tutor + AI diagnosis beats AI-alone on persistence and transfer.

H4: Generalize beyond math (science/identity) only after math identity mechanism is solid.

Part XIII — Open Problems

1. Causal path from narrative history-of-math to identity (under-identified).
 2. Relatedness at scale without creepy social.
 3. Measuring identity without demand effects.
 4. Preventing gamification from colonizing meaning.
 5. AI fluent nonsense vs trust calibration for teens.
 6. Equity: does story-world privilege culturally specific narratives?
-

Part XIV — Red Team Dossier (destroy weak arguments)

Kill #1: “Explanations are free, therefore MindCraft should not explain”

Destroyed form: Absolutism.

Surviving form: Explanations are *necessary infrastructure* with collapsing *willingness-to-pay*; differentiation is diagnosis + affect + identity + accountability.

Kill #2: “Just fix mindset”

Destroyed: Poster mindset.

Surviving: Ecological mindset (task + feedback + peer/tutor norms), math-specific.

Kill #3: “Engagement is the goal”

Destroyed: Engagement without transfer.

Surviving: Engagement as oxygen for the FEI loop, not the fire.

Kill #4: “Mastery graph = Bloom 2-sigma”

Destroyed: Inflated tutoring mythology.

Surviving: Mastery *can* help weaker students if mastery is real and time costs are managed.

Kill #5: “Identity transformation is unmeasurable, so ship vibes”

Destroyed: Vibes-as-strategy.

Surviving: Leading indicators (retry, challenge-seeking) + lagging identity items.

Part XV — References (verified starting set)

Do not invent citations. Expand this list only with retrieved sources.

1. Yeager, D. S., et al. (2019). A national experiment reveals where a growth mindset improves achievement. *Nature*. <https://doi.org/10.1038/s41586-019-1466-y>
2. Dweck, C. S. (2006/2007). *Mindset*. Random House.
3. Ryan, R. M., & Deci, E. L. (2000). Self-determination theory... *American Psychologist* / later handbooks.
4. Wang, Y., et al. (2024). Meta-analysis of SDT-based interventions in education. (Self-Determination Theory site PDF).
5. Kulik, C.-L. C., Kulik, J. A., & Bangert-Drowns, R. L. (1990). Effectiveness of mastery learning programs: A meta-analysis. *Review of Educational Research*.
6. Slavin, R. E. (1987). Mastery learning reconsidered. *Review of Educational Research*.
7. Boaler, J., et al. — mathematical mindset course / classroom studies (Frontiers Education 2018; later scale papers).

8. Systematic review: mindset interventions in mathematics classrooms (2023). *Educational Research Review*, 100554.
 9. Johnston-Wilder, S., & Lee, C. — mathematical resilience / SDT framing (2021).
 10. Duolingo (2023). *The Duolingo Method* whitepaper.
 11. Bloom, B. S. (1984). The 2 sigma problem. *Educational Researcher*.
 12. Ashcraft, M. H., & Kirk, E. P. (2001). The relationships among working memory, math anxiety, and performance. *Journal of Experimental Psychology: General*.
 13. Suárez-Pellicioni, M., Núñez-Peña, M. I., & Colomé, À. (2016). Math anxiety: A review of its cognitive consequences... *Cognitive, Affective, & Behavioral Neuroscience*.
 14. Caviola, S., et al. / related metas — WM mediating role in math anxiety–performance (e.g., *Frontiers in Psychology*, 2021 meta, $r \approx -0.17$ MA–MP).
 15. Kirschner, P. A., Sweller, J., & Clark, R. E. (2006). Why minimal guidance during instruction does not work. *Educational Psychologist*.
 16. Walton, G. M., & Cohen, G. L. — belonging interventions tradition.
 17. VanLehn, K. (2011). The relative effectiveness of human tutoring, intelligent tutoring systems... *Educational Psychologist*.
 18. Nickow, A., Oreopoulos, P., & Quan, V. (2020). Tutoring RCT meta-analysis (effects ~ 0.37 SD — see Education Next summary).
 19. Hambrick, D. Z., et al. (2014). Deliberate practice: Is that all it takes... *Intelligence*; Macnamara, Hambrick & Oswald (2014) deliberate practice meta-analysis.
 20. Ericsson, K. A., Krampe, R. T., & Tesch-Römer, C. (1993). The role of deliberate practice... *Psychological Review*.
 21. Howard, J. L., Slemp, G. R., & Wang, X. (2024/2025). Need support and need thwarting meta-analysis (student populations).
 22. MindCraft internal: WORLD_VISION.md, BRAND_BOOK.md, MindCraft_Viability_Research_Strategy.md.
-

Part XVI — Mechanism Deep Dive: How identity actually changes

This part answers the research question at mechanism grain. It is synthesis, not settled science.

XVI.1 What “confident mathematical thinker” means (operational)

Do not use: “likes math apps” or “says math is fun once.”

Working definition (HYPOTHESIS): A confident mathematical thinker is a person who, in math-relevant contexts:

1. **Approaches** moderately hard problems without immediate avoidance.
2. **Persists** after a miss long enough to extract information.
3. **Attributes** success primarily to controllable causes (strategy, practice, help-seeking) rather than luck or tutor magic.
4. **Selects** into harder future opportunities when available (courses, contests, careers) at rates above their prior baseline.
5. **Endorses** a math-capable self-concept *and* behaves consistently with that endorsement.

Items 1–4 are leading. Item 5 is lagging. Product should optimize leading indicators while periodically measuring the lagging identity claim.

XVI.2 The threat–competence–narrative triad

Three literatures converge on the same transformation engine:

Mechanism	Core claim	Confidence	Key sources
Threat / affect	Math anxiety transiently reduces working-memory resources available for math	High (FACT direction)	Ashcraft & Kirk (2001); Suárez-Pellicioni et al. review (2016); WM mediation meta (2021)
Competence evidence	Motivation requires felt competence; mastery structures can raise achievement when mastery is real	Medium–High	SDT (Ryan & Deci); Kulik mastery meta; Slavin critique
Narrative / identity	Self-stories and belonging cues shape challenge-seeking and persistence	Medium	Yeager et al. (2019); Walton/Cohen belonging tradition; narrative psychology

HYPOTHESIS — Triad necessity: For students starting from “I am bad at math,” improving any one factor alone is usually insufficient:

- Safety without competence → comfort without skill → identity still fragile.
- Competence without safety → possible for some; many never enter the practice volume.
- Narrative without either → empty affirmation (false growth mindset).

XVI.3 Founder tutoring sequence — audit

FOUNDER BELIEF (observed in tutoring):

1. Reduce fear
2. Demystify mathematics
3. Explain why humanity invented it
4. Simplify the narrative
5. Create a tiny success
6. Reinforce success
7. Celebrate progress → identity shifts

Red Team audit:

Step	Universal?	Evidence status	Failure mode
------	------------	-----------------	--------------

1 Fear ↓	Common for anxious subset; not for all	Strong for anxious learners	Over-soothing that removes challenge
2 Demystify	Often helpful	Medium (cognitive clarity)	Over-simplification that lies
3 Invention story	Founder prior	Under-evidenced at RCT scale	Cultural mismatch; time cost
4 Simplify narrative	Often helpful	Medium	Removes necessary difficulty
5 Tiny success	Near-universal lever	High (competence need)	Too-easy wins → hollow identity
6 Reinforce	High if attribution-correct	High (feedback literature)	Praise of talent or praise of empty effort
7 Celebrate	Helpful if genuine	Medium	Performative celebration → cynicism

Improved sequence (HYPOTHESIS — “SAFE-CRAFT”):

1. Safety: threat appraisal ↓ (tone, privacy, soft-wrong, no public shame)
2. Attempt: student acts before receiving a full lecture (productive struggle at right grain)
3. Feedback-as-information: miss becomes diagnosis, not character
4. Earned micro-win: success the student can own
5. Causal story: “what worked” (strategy), optionally “why humans invented this tool”
6. Re-entry: immediate second attempt / next slightly harder item
7. Attribution + witness: tutor/system names the growth specifically
8. Future self: map fill / path unlock that marks identity-relevant progress
9. Transfer check: varied item so win was not clone memorization

What changed vs founder sequence: Attempt moves earlier; invention-story becomes optional enrichment inside step 5, not a prerequisite to agency; transfer check prevents fake mastery.

XVI.4 Math anxiety — what is known, what is not

FACT: Higher math anxiety correlates with lower math performance (meta $r \approx -0.17$ in one 2021 synthesis of 57 studies — modest average, not destiny).

FACT: Ashcraft & Kirk (2001) show math anxiety can act like a dual-task load: worry competes for working memory, especially on computation-span / carry-heavy tasks.

HYPOTHESIS: In product UX, anything that increases social-evaluative threat (red buzzers, public leaderboards for novices, tutor contempt, parent hovering UI) will shrink effective WM and look like “they don’t get it.”

Contradictions / limits:

- Not all underperformance is anxiety; some is missing knowledge, poor curriculum, sleep, language barriers.
- Trait vs state anxiety differ; trait may partly reflect prior failures (reverse causality possible).

- High performers can be anxious and still succeed via avoidance of hard electives — a hidden identity cost.

MindCraft rule: Treat anxiety as a first-class load on the learning system, not as a soft skill sidebar.

XVI.5 Deliberate practice — necessary, not sufficient

FACT / TRADITION: Ericsson’s deliberate practice emphasizes structured, feedback-rich, effortful practice aimed at improvement (not mere repetition).

FACT / CRITIQUE: Hambrick and colleagues argue deliberate practice is important but not sufficient; in chess/music reanalyses, deliberate practice often explains roughly $\sim 1/3$ of reliable variance, leaving most unexplained (abilities, starting age, opportunity, etc.).

Implication for MindCraft: Do not sell “10,000 hours.” Sell **high-quality attempts under feedback**. Also respect individual differences: same practice volume will not equalize outcomes. Identity work must include “progress relative to self,” not only absolute rank.

XVI.6 Tutoring effects — demythologize Bloom

FACT: Bloom (1984) popularized the “2 sigma” tutoring + mastery framing.

FACT / CALIBRATION: Broader tutoring metas are far smaller (often ~ 0.3 – 0.4 SD in modern RCT syntheses cited by Education Next). VanLehn (2011) estimated human tutoring ~ 0.79 SD vs no tutoring, with step-based ITS nearly comparable (~ 0.76).

Implication: Human tutors remain valuable, especially for affect + accountability + diagnosis of weird misconceptions. They are not magic 2-sigma machines. AI step-level feedback can capture much of the *cognitive* tutoring benefit; the scarce remainder is Layers 4–6 (emotion, identity, accountability).

Part XVII — Scarcity Thesis: Full Red Team Trial

XVII.1 The claim under oath

FOUNDER BELIEF: AI commoditizes knowledge; future scarcity is motivation, confidence, identity, trust, curiosity, persistence, emotional safety, meaning.

XVII.2 Cross-examination

Claim A — “Explanations are free”

Verdict: Directionally **FACT** for commodity explanations; **FALSE** for *trusted, personalized, correct, curriculum-aligned* explanations under time pressure.

Evidence: Generative AI + open platforms collapse marginal cost of text explanation.

Contradiction: Hallucinations, shallow fluency, and misalignment create a *trust tax*. Parents may still pay for verified human judgment.

Claim B — “Tutoring is free”

Verdict: FALSE.

AI chat is cheap. High-quality tutoring includes scheduling, relationship, emotional labor, accountability, and reputation. Those remain expensive. Nickow et al.-style tutoring metas still show meaningful effects when tutoring is structured — it is not free at scale.

Claim C — “Content is free”

Verdict: MOSTLY FACT for undifferentiated content; FALSE for sequenced, assessed, outcome-linked content systems.

Khan/YouTube prove content abundance. Differentiation is curation + measurement + human wrap.

Claim D — “Motivation/identity become the scarce goods”

Verdict: HYPOTHESIS with strategic plausibility.

As cognitive help approaches free, *willingness to engage difficulty* becomes the binding constraint for many students. This does not mean motivation products automatically win — many habit apps fail to produce identity.

Claim E — “Emotional safety is first-order”

Verdict: HYPOTHESIS — true for anxious / threatened learners; not universal.

For some, speed drills and competition *increase* engagement. Segment, do not romanticize softness.

XVII.3 Surviving thesis (company doctrine)

Adopted doctrine (refined):

> MindCraft does not compete primarily on explanation volume. It competes on converting threatened learners into challenge-seeking mathematical agents through diagnosis, affective design, mastery evidence, narrative meaning, and human accountability — with AI as infrastructure.

Falsifiers (90-day capable):

1. Anxious segment shows no FEI lift vs explanation-first control.
2. Parents refuse to pay when score gains lag identity metrics.
3. AI-alone matches human+AI on persistence + transfer for target segment.

Part XVIII — Systems Maps (loops, not features)

XVIII.1 FEI reinforcing loop (desired)

```
+----- identity claim ("I do math") <--+  
| |  
v |  
safety UX ----> attempt ----> informative miss ----> coach  
^ | |  
| v v  
+----- map fill / witness <----- earned success <+
```

XVIII.2 Shame balancing loop (destroy this)

```
public failure signal -> threat ↑ -> avoidance -> fewer attempts  
^ |  
+----- skill lag / worse scores -----+
```

XVIII.3 Extrinsic colonization loop (danger)

```
streak / points ↑ -> return ↑ -> shallow grinding ↑  
| |  
+-- identity unchanged -----+--> burnout / streak break → exit
```

Design rule: Habit loops may feed FEI but must not replace identity markers.

XVIII.4 Parent trust loop

```
visible diagnosis + honest progress -> parent trust ↑ -> continued purchase  
^ |  
+---- outcome evidence (scores / challenge-seeking)
```

FOUNDER QUESTION unresolved: Optimize Maya's identity or parent's anxiety? Doctrine: report *both* challenge-seeking and skill evidence; never fake either.

XVIII.5 Network effects (honest)

MindCraft today has weak classic network effects. Do not pretend otherwise.

Possible compounding advantages (HYPOTHESIS):

1. Ontology + misconception memory → better diagnosis over time (data moat if privacy-respecting).
2. Tutor playbooks trained on FEI outcomes → service quality moat.
3. Story worlds with longitudinal character arcs → switching costs of meaning (fragile; easy to cargo-cult).

Part XIX — Cross-Domain Transfer Catalog

Principles that repeatedly change human behavior outside classrooms — filtered for MindCraft use.

Domain	Robust mechanism	Transfer	30-year durable?	AI improve?
Exposure therapy	Graded exposure under safety	Anxiety-graded problem sets	Yes	Yes (personalize gradient)
Sports film study	Specific error review	Wrong-choice coach	Yes	Yes (auto clip of miss type)
Music pedagogy	Technique → repertoire → performance	Drill → quest → “recital” signal	Yes	Partial
Martial arts	Visible ranks + dojo norms	Map regions + tutor dojo culture	Yes	Weak for norms
Aviation	Checklists under stress	Procedure scaffolds when anxious	Yes	Yes
Games (Celeste-like)	Fair failure + retry culture	Soft-wrong physics	Yes	Yes
Chess	Annotated game review	Session postmortems	Yes	Yes
Coaching	Identity + accountability	Tutor as witness	Yes	Assist, not replace
Religion/ritual	Meaning + belonging	Light rituals only; avoid cult dynamics	Yes (human)	Risky
Creator economy	Public craft identity	Optional shareable math artifacts	Uncertain	Yes

Anti-transfer: Casino variable rewards → addiction without skill. Do not copy engagement maxing from social media.

Part XX — Competitive Landscape (mechanism lens)

Product	Primary scarce resource they sell	Identity claim?	Risk
Khan Academy	Free mastery content + reputation	Weak–medium	Explanation commodity
Duolingo	Habit / streak motivation	Weak (language identity sometimes)	Extrinsic ceiling
Brilliant	Aesthetic problem joy + prestige	Medium	Narrow segment

Coursera/edX	Credentials	Medium (career identity)	Completion collapse
Chegg / homework help	Answers under deadline	Anti-identity (outsourcing)	Integrity & AI shock
Private tutors	Human accountability + customization	High when good	Supply constrained
ChatGPT tutors	Instant explanation	Low unless wrapped	Trust / hallucination
MindCraft (target)	FEI conversion + tutor witness + gap diagnosis	Intended high	Must prove, not assert

Strategic implication: Do not out-Khan Khan on content breadth. Do not out-Duo Duo on streaks. Out-tutor mediocre marketplaces on *diagnosed emotional-cognitive loop*.

Part XXI — Metrics Dictionary (instrumentation)

XXI.1 Banned as North Stars

- Raw DAU / time-on-app without transfer
- Streak length alone
- Explanation open-rate
- Points / coins

XXI.2 Leading indicators (ship first)

Metric ID	Definition	Why
retry_120s	New attempt on same or isomorphic item within 120s of soft-wrong	Persistence under safety
coach_shown	Wizard/coach surfaced	Treatment exposure
write_exit_to_retry	Leave write mode then attempt	Agency after reflection
challenge_accept	Chose harder level when offered	Challenge-seeking
hint_binge	≥3 hints without independent solve	Gaming / helplessness
transfer_pass	Correct on varied item after mastery mark	Anti-fake-mastery

XXI.3 Lagging indicators

Metric ID	Definition
math_person_item	1–7 agreement: “I am a math person” (math-specific)
anxiety_state_item	Short state anxiety before session
advanced_intent	Intent to take harder course / contest
tutor_witness_note	Tutor tagged identity-relevant growth

XXI.4 Decision rule (HYPOTHESIS)

Ship changes that raise `retry_120s` and `challenge_accept` without raising `hint_binge` or dropping `transfer_pass`.

Part XXII — Appendices

Appendix A — Glossary

Term	Meaning
FEI Loop	Fear → Evidence → Identity system
ITC	Identity Transformation Cascade
Soft-wrong	Miss treated as information, not shame theater
Gap scan	Low-shame diagnostic of confidence/exposure
Challenge-seeking	Choosing harder math opportunity when optional
False growth mindset	Effort praise without strategy / conditions
SAFE-CRAFT	Improved tutoring sequence (Part XVI.3)
Layers 0–6	Knowledge commodity stack (Part VI.3)

Appendix B — Maya interview protocol (qual)

Goal: Map moments when math stopped (or started) feeling dangerous.

1. “Tell me about a time math felt dangerous or humiliating.”
2. “What did your body do? What did you do next?”
3. “Tell me about a time you surprised yourself in math.”
4. “Who saw that moment? Did it matter that someone saw it?”
5. “If an app celebrated you, would that feel real or fake? Why?”
6. “What would make you take a harder math class next year?”
7. “When you get stuck, do you want comfort, a hint, or the answer? In what order?”

Coding: threat markers; attribution language; witness presence; desire for story vs drill.

Appendix C — Causal DAG (provisional)

```
prior failures → math anxiety → WM load → performance ↓  
\\ \→ avoidance → practice ↓ → skill lag  
\\ \→ fixed identity narrative  
\→ missing knowledge ████████████████████████████████████████ → performance ↓  
  
safety UX → threat ↓ → WM available ↑ → attempt quality ↑  
coach + diagnosis → misconception repair → competence evidence ↑  
competence evidence + witness → identity update → challenge-seeking ↑  
challenge-seeking → harder practice → skill ↑ (reinforcing)
```

Open identification problem: Does identity cause challenge-seeking, or does challenge-seeking cause identity? Likely bidirectional. Measure both.

Appendix D — Experiment pre-registration sketches

A — Wizard coach

- Population: chapter/practice users with soft-wrong enabled
- Arms: coach on / coach off
- Primary: `retry_120s`
- Guardrail: `transfer_pass`, session completion
- Stop rule: clear harm on transfer or spike in hint binge

B — Story frame

- Arms: invention-story wrapper vs bare stem (same math)
- Primary: time-to-abandon after first miss
- Secondary: enjoyment, retention @ 7d

D — Explanation timing × anxiety

- Screen: brief anxiety item
- Arms: explain-first vs attempt-first
- Heterogeneity: high vs low anxiety

Appendix E — Red Team standing orders

1. Attack any claim that cannot name a falsifier.
2. Attack any metric that can be gamed without learning.
3. Attack any narrative feature without a learning outcome.
4. Attack Bloom 2-sigma mythology on sight.
5. Attack “AI makes tutoring free” slogans on sight.

Appendix F — Future chapters (queued, not padded)

1. Neuroscience of math anxiety (amygdala/WM models; cautious translation)
 2. History of mathematics as meaning technology
 3. Parent trust economics & willingness-to-pay experiments
 4. Equity audit of story worlds
 5. Full competitive teardown with usage telemetry (when available)
 6. Formal Bayesian update process for Constitution claims
-

Part XXIII — Research Lab Operating Cadence

Weekly (60–90 min):

1. One claim enters trial (evidence + contradiction).
2. Red Team attempts kill.
3. Surviving claim updates Constitution with confidence tag.
4. One experiment metric reviewed.

Monthly:

- Rewrite one weak chapter.
- External reader (learning scientist or skeptical founder).
- Regenerate PDF.

Quarterly:

- Kill or promote North Star metric.
 - Publish internal “what we believed that was wrong” memo.
-

Closing stance

MindCraft’s deepest risk is not technical failure. It is **winning the wrong game**: becoming the best explanation engine in a world where explanations are cheap, while losing the only game that compounds — helping a human revise who they are in the presence of difficulty.

This Constitution exists so the company notices that risk early, and runs experiments that can kill beloved ideas.

v1.1 expanded mechanisms, scarcity trial, systems maps, and instrumentation. The lab continues. Page count is not the finish line — falsifiable truth is.

Part XXIV — Affect, Anxiety, Belonging (deep dive)

Chapter status: Living evidence brief

Primary question: When is emotional safety causally load-bearing for mathematical confidence?

Owners: Cognitive Neuroscientist seat · Ed Psych seat · Red Team

XXIV.1 What is happening, mechanistically?

Three related but non-identical constructs get collapsed in product talk. Separate them.

Construct	Definition (working)	Time scale	Product signal
State math anxiety	Transient dread / arousal in a math moment	Seconds–minutes	Avoid click, freeze, rapid hint binge
Trait math anxiety	Stable tendency to experience math as threat	Months–years	Course avoidance, “I’m not a math person”
Belonging uncertainty	“People like me don’t belong in this domain”	Contextual, can be acute	Withdrawal after mild adversity
Stereotype threat	Extra evaluative pressure from group stereotype relevance	Situational	Underperformance vs ability when stereotype is cued

FACT (direction): Math anxiety is associated with worse math performance; one 2021 meta-analysis reported mean $r \approx -0.17$ across 57 studies (modest average association — not destiny).

FACT (mechanism candidate): Ashcraft & Kirk (2001) show high math-anxiety individuals have reduced working-memory span especially on computation-based span tasks; dual-task math + memory load amplifies errors/RT. Anxiety behaves like a competing load on central executive resources.

FACT (related): Reviews (e.g., Suárez-Pellicioni et al., 2016) organize explanations as (a) WM competition, (b) attentional control / inhibition deficits, (c) possible low-level numerical representation differences — not mutually exclusive.

HYPOTHESIS for MindCraft: For the Maya-anxious segment, UX that increases social-evaluative threat (shame theater, public ranking early, contemptuous feedback) will look like a “knowledge gap” while actually being a temporary capacity tax.

XXIV.2 Why does this happen? (first principles)

1. **Working memory is scarce.** Multi-step math is a WM sport. Anything that occupies the central executive (worry, self-monitoring for stereotype confirmation, rumination about looking stupid) steals the resource the task needs.

2. **Humans are status-sensitive learners.** Math is a publicly ranked school subject. Threat is often social (“I will be seen as dumb”) more than arithmetic.
3. **Avoidance is rational short-term.** Leaving the room ends the aversive arousal. Long-term, avoidance prevents the mastery experiences that would update efficacy.
4. **Identity consolidates avoidance.** Repeated exits become “I’m not a math person,” which then licenses future exits (narrative lock-in).

Human need satisfied by safety design: Need for non-humiliation while competence is still fragile; relatedness/respect; predictability.

XXIV.3 Belonging and stereotype threat — what transfers?

FACT: Steele & Aronson tradition — when a test is framed as ability-diagnostic under a negative group stereotype, performance can drop relative to non-diagnostic framing (classic lab pattern; effect sizes and replication quality vary by context).

FACT: Spencer, Steele & Quinn (1999) line — women/math stereotype relevance can impair performance when threat is cued.

FACT: Walton & Cohen (2007, 2011) — belonging uncertainty makes adversity feel identity-diagnostic; brief belonging interventions framing adversity as common/transient improved long-term outcomes for marginalized college students in landmark trials (e.g., Science 2011 belonging intervention).

Red Team limits:

- Effects are **heterogeneous** and context-dependent (like mindset).
- School-wide scaling often attenuates lab effects.
- MindCraft must not cargo-cult a 45-minute belonging essay into a gamified toast.
- Belonging interventions are not a substitute for teaching.

Surviving transfer:

- Normalize struggle without normalizing low standards.
 - Make early adversity non-identity-diagnostic (“misses are how the map learns”).
 - Avoid stereotype-cueing copy (“even girls can...” is often worse).
 - Tutor witnessing is a belonging technology when authentic.
-

XXIV.4 Neuroscience: what we can say without overclaiming

FACT / REVIEW level: Math anxiety literature includes ERP/fMRI correlates; reviews document affective and cognitive network involvement. Exact “brain region product features” are **SPECULATION** and should not drive roadmap.

Allowed claim: Affective arousal and executive control interact; designing for lower threat is biologically plausible.

Forbidden claim: “Our purple calm mode activates the prefrontal cortex.”

XXIV.5 Contradictory evidence (keep visible)

1. Many students underperform from **missing knowledge**, not anxiety.
 2. Some learners seek **competitive arousal** and dislike soft UX.
 3. Reverse causality: low skill → anxiety → further avoidance (bidirectional).
 4. Trait anxiety interventions that only soothe without competence can create comfort without growth.
 5. Meta correlations are modest; anxiety is one bottleneck among many.
-

XXIV.6 Can this generalize? 30 years? AI?

Question	Answer
Generalize beyond math?	Yes — any high-status evaluative domain (music juries, coding interviews).
Still work in 30 years?	Yes — human threat systems are not a 2026 fad.
Can AI improve it?	Yes for personalization of challenge gradient & detection of freeze patterns; risky for fake empathy.
MindCraft change	Soft-wrong, hide-correctness diagnostics, coach tone, private early practice, tutor witness scripts are load-bearing for anxious segment.

XXIV.7 Experimental validation program

ID	Question	Design	Primary	Falsifier
AFF-1	Does soft-wrong raise retry vs hard-fail UX?	A/B	retry_120s	No lift / accuracy down
AFF-2	Does anxiety-state moderate explanation-first harm?	Anxiety screen × timing	transfer + retry	No interaction
AFF-3	Does “misses are information” copy reduce belonging threat items?	Pre/post items	belonging/threat items	Demand effects only

AFF-4	Tutor witness note vs no note after growth moment	Tutor RCT	2-week challenge -seeking	Null
-------	--	-----------	------------------------------	------

Confidence in chapter thesis: Medium–High that affect matters for a large subset; Medium that MindCraft’s current UX already captures the right mechanisms; Low that we have measured it.

XXIV.8 Product system (not features)

Affective Load Manager (system):

1. Detect freeze / hint-binge / rapid exit (behavioral).
2. Downshift evaluative threat (tone, privacy, soft-wrong).
3. Offer graded re-entry (easier isomorphic item), not pep talk alone.
4. Restore challenge once retry and success stabilize.
5. Log whether the student later chooses harder work (identity leading indicator).

Balancing loop: chronic downshift → boredom → exit. Must re-raise difficulty.

Part XXV — Self-Efficacy, Narrative Identity, Habit (deep dive)

Chapter status: Living evidence brief

Primary question: What updates the self-story “I am bad at math,” and what merely increases login frequency?

Owners: Learning Science · Narrative Psychology · Behavioral Econ · Red Team

XXV.1 Self-efficacy is not “confidence vibes”

FACT (theory): Bandura’s social cognitive theory — self-efficacy is belief in capability to organize and execute actions required for specific attainments. It is **task- and domain-specific**, not a global personality glow.

FACT (sources): Efficacy information is processed from four sources (Bandura, 1997):

1. **Mastery experiences** (usually strongest)
2. **Vicarious experiences** (models)
3. **Social persuasion** (credible others)
4. **Physiological / affective states** (how arousal is interpreted)

FACT (measurement tradition): Usher & Pajares (2009/2010 line) validated source scales for middle-school mathematics; sources relate to math self-efficacy, goals, optimism.

HYPOTHESIS: MindCraft’s FEI loop is largely a **self-efficacy engineering system**:

FEI stage	Efficacy source
Safety / soft-wrong	Reinterprets physiological arousal (“alert, not doomed”)
Earned micro-win	Mastery experience
Coach naming strategy	Persuasion + attribution
Tutor witness / peer story	Vicarious + persuasion
Map fill after transfer	Mastery that survives variation

XXV.2 Why mastery experiences dominate — and how products fake them

What is happening: A win only updates efficacy if the learner **attributes** it to their capability under meaningful difficulty.

Fake mastery (product sins):

- Too-easy items → “this doesn’t count”
- Heavy hinting → “the app did it”
- Identical clones after “mastery” → recognition, not competence
- Points for opening lessons → engagement theater

Real mastery design rules (HYPOTHESIS):

1. Difficulty just above comfort (desirable difficulty, not cruelty).
2. Student performs the decisive step.
3. Immediate informative feedback.
4. Soon after: varied item (transfer_pass).
5. Explicit attribution to strategy (“you isolated the factor”) not talent.

Contradicting caution: After strong efficacy exists, occasional failure hurts less (Bandura). Early failures hurt more. Protect early arc; do not permanently infantilize.

XXV.3 Reciprocal loops: efficacy ↔ achievement

FACT / EMERGING: Longitudinal work often finds reciprocal relations — achievement raises efficacy and efficacy raises later achievement (e.g., recent high-school math reciprocity analyses in large-scale assessment literature). Causality is not one-way.

Implication: “Just raise confidence” without skill is unstable. “Just raise skill” while humiliating can stall practice volume for anxious students. Build both.

XXV.4 Narrative identity — the story that persists

FACT / TRADITION: Narrative identity research (McAdams tradition) — people maintain identity through evolving life stories with themes of agency, communion, redemption, contamination.

HYPOTHESIS (MindCraft synthesis): Math identity change is a **redemption narrative rewrite**:

- Contamination story: “I failed → I am bad at math → I avoid → I stay bad.”
- Redemption story: “I struggled → I used a strategy → I grew → I am becoming someone who stays.”

Product implication: Story worlds matter if they supply **symbols for redemption**, not only costume. The map lighting after a real win is a narrative punctuation mark.

Red Team: Most edtech “stories” are skin. If removing the narrative leaves identical pedagogy and identical emotions, the story was decoration. Test with story-off experiments (Experiment B).

XXV.5 Habit vs identity — the Duolingo trap, precisely

FACT / INDUSTRY: Habit products use cues, streaks, variable rewards, loss aversion. Duolingo’s method materials emphasize motivation systems for return.

FACT / THEORY tension: Self-Determination Theory warns that controlling extrinsic rewards can undermine intrinsic motivation when they feel coercive (classic Deci findings; nuance exists for informational rewards).

HYPOTHESIS — Habit/Identity Divergence Model (HID):

Habit success: cue → routine → reward → return frequency ↑
Identity success: evidence → attribution → self-story → choice of harder future ↑
Overlap zone: daily practice that produces mastery evidence
Failure mode A: habit without mastery → brittle streak addiction
Failure mode B: identity talk without practice → empty affirmation

Metric separation (mandatory):

- Habit health: D1/D7 return, streak
- Identity health: challenge_accept, advanced course intent, math-person item + behavior consistency

Optimize for overlap zone. Never promote a habit metric that moves opposite identity metrics.

XXV.6 Vicarious experience and “who is the model?”

HYPOTHESIS: Models work when perceived as **similar yet slightly ahead** (“near-peer who struggled”), not as genius tutors who never miss.

Product:

- Villager / character arcs that show struggle → strategy → win.

- Tutor stories: “I used to blank on fractions.”
- Avoid only showcasing perfect speedruns.

Social persuasion: Credibility matters. Random app toast < trusted tutor sentence. Parents can persuade or destroy (“we’re not math people” is efficacy poison — FOUNDER BELIEF with strong anecdotal support; needs interview confirmation).

XXV.7 Universal principles from this chapter

1. Efficacy is specific; measure math efficacy, not “confidence.”
2. Mastery experiences must be owned and non-trivial.
3. Arousal interpretation is designable (anxiety as signal vs doom).
4. Narrative punctuation should mark real competence events.
5. Habits are oxygen; identity is the fire.

XXV.8 Experiments

ID	Question	Design	Primary
EFF-1	Does requiring decisive student step (vs auto-solve) raise efficacy item?	A/B	math efficacy delta
EFF-2	Strategy attribution copy vs talent copy after win	A/B	challenge_accept @ 7d
EFF-3	Streak on vs identity milestone on (map fill)	A/B	retention vs challenge_accept tradeoff
EFF-4	Near-peer struggle story vs expert-perfect story	A/B	persistence after first miss

Falsifier for identity thesis: Habit metrics rise for 8 weeks while challenge-seeking and transfer stay flat — then MindCraft is a streak app, not an identity company.

Part XXVI — Cognitive Load, Mastery, Tutoring (deep dive)

Chapter status: Living evidence brief

Primary question: How do we create earned competence without drowning working memory — and without mythologizing tutoring?

Owners: Cognitive Science · Learning Science · AI Researcher · Red Team

XXVI.1 Cognitive Load Theory — the non-negotiable physics

FACT / TRADITION: Cognitive Load Theory (Sweller and colleagues) — working memory is limited; instructional design that forces novices into unguided search can impede schema acquisition.

FACT: Worked-example effect — for novices, studying worked examples often beats equivalent time in conventional problem solving for structurally similar later items (classic Sweller & Cooper algebra studies, 1985 line).

FACT / DESIGN IMPLICATION: Guidance fading — move from full worked examples → completion problems → independent solving as expertise grows (Renkl/Atkinson fading tradition). Continuing full examples after expertise rises can trigger **expertise reversal** (redundancy).

Why this matters to identity: Overloaded novices experience failure as self-indictment. Load mismanagement manufactures fake “I’m bad at math” evidence.

XXVI.2 Reconciling CLT with “attempt before lecture”

SAFE-CRAFT says attempt early. CLT says novices need worked examples. **Is this a contradiction?**

Resolution (HYPOTHESIS — Attempt Grain Principle):

Learner state	Right “attempt”	Wrong “attempt”
True novice, new schema	Attempt a micro-step or completion blank inside a worked example	Full novel multi-step problem cold
Partial schema	Isomorphic problem with soft-wrong	Leap to transfer without success
Anxious + some knowledge	Slightly easier isomorphic first	Public hard problem as opener
Advanced	Minimal guidance; generative struggle	Endless examples (expertise reversal)

Product rule: “Struggle first” means **agency at the right grain**, not abandonment into search.

This dissolves a false war between “productive struggle” slogans and cognitive load science.

XXVI.3 Mastery learning — keep the war on the table

Supporting evidence: Kulik, Kulik & Bangert-Drowns (1990) meta — mastery programs associated with positive exam effects (often cited ~0.5 SD overall in their synthesis), often stronger for weaker students; time costs can rise; completion can suffer in some self-paced settings.

Contradicting evidence: Slavin (1987) “Mastery Learning Reconsidered” — limited support for group-based mastery on standardized measures; stronger on experimenter-made tests; coverage vs depth tradeoff is real.

MindCraft doctrine:

1. Mastery means **transfer_pass**, not same-item streak.
2. Track time-to-mastery; do not hide cost.
3. Allow strategic coverage decisions (exam track) without pretending infinite time.
4. Ontology graph is only as good as far-transfer checks.

XXVI.4 Tutoring — calibrate the mythology

Bloom (1984): Tutoring + mastery framing popularized as ~2 SD (“2 sigma problem”).

Calibration FACTS:

- Cohen, Kulik & Kulik (1982) tutoring meta — much smaller average effects than 2.0.
- Nickow, Oreopoulos & Quan (2020) tutoring RCT meta — on the order of ~0.37 SD (impressive, not magic).
- VanLehn (2011) — human tutoring ~0.79 SD vs no tutoring; step-based ITS ~0.76; answer-based CAI smaller. Granularity of feedback matters; beyond a point, finer grain plateaus.

Strategic reading for MindCraft:

Tutoring function	Automate?	Keep human?
Step feedback	Yes (ITS-like)	Optional
Misconception diagnosis	Partially (ontology + traces)	Human for weird cases
Affect co-regulation	Partially (UX)	Human strong
Accountability / relationship	Weakly	Human strong
Identity witnessing	Poorly	Human strong

AI implication: Compete on diagnosis quality + FEI wrapping; do not sell “we are 2-sigma.”

XXVI.5 Deliberate practice inside sessions

FACT / TRADITION: Ericsson et al. (1993) — structured, feedback-rich practice aimed at improvement.

FACT / CRITIQUE: Hambrick/Macnamara line — deliberate practice explains substantial but incomplete variance; not a sufficient account of expertise differences.

Session design (HYPOTHESIS):

1. Clear target skill (ingredient-level).
2. Immediate feedback at step grain when possible.
3. Successive refinement on weak subskill (not random variety first).
4. Then varied practice for transfer.
5. Short enough to protect affect; frequent enough to compound.

XXVI.6 Minimal guidance failures

FACT: Kirschner, Sweller & Clark (2006) critique pure discovery/minimal guidance for novices — often inferior to guided instruction.

Implication for story worlds: Narrative must not become unguided discovery of math. Story wraps guided cognition. World_VISION is compatible with guidance; it is incompatible with vibes-only exploration.

XXVI.7 Systems: Competence Production Line

diagnosis (gap) → grain selection → worked/fade → attempt
 → soft-wrong info → repair → transfer check
 → mastery mark → spaced return → challenge raise

Balancing loops:

- Too slow fade → boredom
- Too fast fade → threat + load spike → avoidance
- Mastery without spacing → illusion of competence

XXVI.8 Experiments

ID	Question	Design	Primary
CLT-1	Worked-example-first vs problem-first for true novices	Within concept	near & far transfer
CLT-2	Fade schedule adaptive to mastery graph vs fixed	A/B	time-to-transfer_pass
TUT-1	Human tutor + AI brief vs AI-alone on persistence	RCT	retry + 4-week challenge_accept
MAS-1	Mastery gate with transfer_pass vs same-item 3-in-a-row	A/B	2-week retention

Red Team standing kill: Any claim that MindCraft “replicates Bloom 2-sigma” without citing modern tutoring metas.

Part XXVII — Markets, Parents, Competition (deep dive)

Chapter status: Living strategy brief (evidence + market reasoning)

Primary question: If explanations cheapen, what do families still buy — and what should MindCraft refuse to become?

Owners: Founder Philosopher · Product Strategist · Behavioral Econ · Red Team

XXVII.1 Re-state the scarcity doctrine under market pressure

Adopted doctrine (from Part XVII): Compete on diagnosis, affect, identity, accountability — with AI as infrastructure.

Market test: Parents often buy **score relief** and **homework peace**, not “identity transformation.” Students may buy relief from shame. Schools buy coverage and equity metrics.

FOUNDER QUESTION (unresolved): Can MindCraft sell identity truthfully while packaging score outcomes parents recognize?

Working answer (HYPOTHESIS): Sell the FEI system as the *method*; report scores/challenge-seeking as *evidence*. Never sell vibes without transfer.

XXVII.2 Willingness-to-pay stack (post-AI)

Layer	Marginal cost trend	WTP trend	Notes
Raw answers	→ 0	↓↓	Chegg shock / GPT
Generic explanations	→ 0	↓	Unless trusted
Personalized practice	↓	→ / ↑	Needs outcomes
Diagnosis of <i>this</i> student	↓ with AI	↑ if trusted	Ontology advantage possible
Emotional safety during struggle	Hard to automate well	↑ for anxious segment	UX + tutor
Accountability relationship	Scarce	↑↑	Scheduling, human
Trusted outcome signal to parents	Scarce	↑↑	Reports that don’t lie

SPECULATION (30-year): Ambient AI tutors everywhere. Differentiator = trusted human+system that changes trajectory, not chat fluency.

XXVII.3 Parent trust economics

What parents fear (hypothesis from tutoring practice — needs interviews):

1. Child is falling behind permanently.
2. They cannot help at home.
3. Child hates learning / shuts down.
4. Money spent on apps that entertain without learning.

Trust builders (HYPOTHESIS):

- Specific diagnosis (“fractions↔ratios bridge weak”) over vague “needs practice”
- Visible soft-wrong → retry behavior (effort signal)
- Transfer checks (anti-cheat of fake mastery)
- Honest plateaus (“this week: consolidation, not fireworks”)
- Tutor notes that sound human and concrete

Trust destroyers:

- Inflated streaks as “progress”
- Grade promises you can’t keep
- Dark-pattern engagement
- AI that hallucinates pedagogy confidently

Experiment PAR-1: Parent report variants — score-only vs FEI dashboard (retry, challenge-seeking, mastery map). Measure retention/WTP intent and comprehension of child’s actual bottleneck.

XXVII.4 Competitive teardown (mechanism lens)

Khan Academy

- **Sells:** free mastery content + brand trust
- **Strength:** coverage, accessibility
- **Weakness for Maya:** can still feel lonely/evaluative; identity not core
- **MindCraft response:** do not out-content; out-diagnosis + human wrap

Duolingo

- **Sells:** habit motivation
- **Strength:** return loops
- **Weakness:** habit ≠ deep skill/identity (HID model)
- **MindCraft response:** borrow on-ramps; refuse extrinsic colonization of meaning

Brilliant

- **Sells:** aesthetic intellectual joy + prestige
- **Strength:** delightful problems for already-curious
- **Weakness:** may miss threatened segment
- **MindCraft response:** win the anxious middle, not only the delighted

Tutoring marketplaces

- **Sells:** access to humans
- **Strength:** relationship potential
- **Weakness:** uneven quality; weak diagnosis systems
- **MindCraft response:** FEI-trained tutors + AI briefings

Generic GPT tutors

- **Sells:** instant explanation
- **Strength:** 24/7 cognitive help
- **Weakness:** trust, curriculum alignment, affect, accountability
- **MindCraft response:** wrap models in ontology + FEI + human witness

XXVII.5 Business model risks tied to research thesis

If thesis true...	Business implication
Identity is the product	Retention via meaning + progress, not only content library
Affect is first-order for segment	CAC messaging should speak shame→agency, not “more videos”
Human witness scarce	Hybrid model: AI scale + tutor moments that matter
Parents buy trust	Reporting quality is a product surface

If thesis false...	Pivot signal
Explanation-first wins all segments	Become best adaptive explainer (crowded)
Streaks drive outcomes without identity	Become math Duolingo (crowded, brittle)
Parents only pay for score spikes	Short-cycle test prep (possible, different company)

XXVII.6 Network effects — honesty clause

FACT about present: MindCraft does not have Facebook-style network effects.

Plausible compounding (HYPOTHESIS):

1. Misconception telemetry → better diagnosis (privacy-constrained data advantage)
2. Tutor playbook quality → service moat
3. Longitudinal identity arcs in worlds → switching costs of meaning (fragile)

Red Team: “Community” features often fake relatedness. SDT metas show relatedness interventions harder to move. Do not ship empty social to check a box.

XXVII.7 Pricing ethics (Founder Philosopher)

HYPOTHESIS: Charging for emotional safety is legitimate if competence and transfer are real; illegitimate if the product only soothes.

Guardrail: Every paid claim in marketing must map to a measured FEI or learning metric. If not measured, do not claim.

XXVII.8 Experiments & interviews

1. **10 parent interviews:** what did you think you bought? what would make you cancel?
2. **10 Maya interviews:** when did math stop feeling dangerous?
3. **PAR-1** report A/B
4. **COMP-1** message test: identity framing vs score framing vs hybrid (conversion + 30-day retention)

Confidence: Medium that hybrid packaging wins; Low on exact price points without data.

Part XXVIII — History, Meaning, and Why Math Exists (deep dive)

Chapter status: Design-research brief (under-evidenced; high founder interest)

Primary question: Does explaining *why humanity invented a mathematical idea* change persistence and identity — or is it a beautiful distraction?

Owners: Historian of Mathematics seat · Founder Philosopher · Learning Science · Red Team

XXVIII.1 The founder prior (label it)

FOUNDER BELIEF: Students disengage partly because math is presented as arbitrary machinery. Reconnecting ideas to human invention, need, and story demystifies the subject and restores meaning.

WORLD_VISION alignment: Story-first questions; math frozen, narrative wraps.

Epistemic status: Strong pedagogical tradition; **weak large-scale RCT base** relative to mindset/SDT. Treat as a research program, not settled science.

XXVIII.2 What “meaning” might be doing (mechanisms)

If invention-story helps, why?

Mechanism	Claim type	Testable implication
Arbitrary-rule appraisal ↓	HYPOTHESIS	Fewer “why do we have to learn this?” exits
Curiosity spike	HYPOTHESIS	Longer voluntary time before first quit
Cognitive schema anchor	HYPOTHESIS	Better retention of structure if story encodes structure (not trivia)
Identity affiliation (“humans who solve”)	HYPOTHESIS	Higher math-person endorsement
Seductive details (harm)	FACT tradition in multimedia learning	Stories that add irrelevant color can <i>hurt</i> retention

Critical constraint: History must encode the **mathematical structure**, not costume the problem with dragons while leaving the math alien.

XXVIII.3 Seductive details problem (Red Team weapon)

FACT / TRADITION: Multimedia learning research (Mayer and others) — interesting but irrelevant details can impair learning (seductive details effect). Not always replicated equally, but the risk is real.

Kill criterion for MindCraft stories: If a narrative element can be removed without changing the mathematical mental model, it is decoration. Decoration can remain for motivation **only if** Experiment B shows persistence/retention gains that justify time cost.

Survive criterion: Narrative that makes the *invariant* memorable (e.g., why we need a common measure; why zero; why coordinates) is structural meaning.

XXVIII.4 History of math as demystification technology

Examples of structural meaning (illustrative — not a curriculum):

- 1. **Fractions / ratios:** Fair sharing, measurement, comparison when wholes differ.
- 2. **Negative numbers:** Debt, direction, opposite change — not “fake numbers.”

3. **Coordinates:** Descartes' move — marry algebra to geometry to locate.
4. **Functions:** Input–output machines for prediction and dependency.
5. **Probability:** Reasoning under incomplete information; insurance, games, risk.
6. **Quadratics / optimization:** Paths of projectiles, areas, tradeoffs — tools for extrema.

HYPOTHESIS: One crisp invention sentence + one visual beat beats a museum lecture.

XXVIII.5 Cross-domain analogy: martial arts lineages & music history

Domain	Meaning practice	Transfer
Martial arts	Lineage stories + why a form exists	Respect + purpose; still drill techniques
Music	Composer context + etudes	Meaning without skipping scales
Religion	Origin myths	Caution: do not sacralize math cultishly
Architecture	“Why this arch holds”	Structure-first stories

Principle: Meaning wraps discipline; it does not replace repetition.

XXVIII.6 Equity and whose history?

Open problem: Whose invention stories? Eurocentric timelines can alienate. Global histories of mathematics (India, Islamic golden age, China, Africa, Indigenous measurement systems) matter for belonging **and** accuracy.

HYPOTHESIS: Inclusive structural histories can support belonging for students who feel math is “not for people like me” — but only if not tokenistic.

Experiment HIST-EQ: Same math, Euro-only story vs plural invention story; measure belonging items + learning — watch for seductive-detail harm.

XXVIII.7 Integration with SAFE-CRAFT

Invention story belongs in **step C (Causal story)** after an attempt grain, or as a 10–20s prelude that encodes structure — not as a gate that delays agency for anxious students who need a micro-win first.

Default order for threatened learners: Safety → micro-attempt → feedback → meaning beat → next attempt.

Default order for curious/low-anxiety: Short meaning beat → attempt → fade guidance.

Segment, don't dogmatize.

XXVIII.8 Experimental program (this chapter lives or dies here)

ID	Question	Design	Primary	Kill condition
HIST-1	Structure-encoding story vs bare stem	Crossover	retention @ 7d + retry	Story slows, no learning lift
HIST-2	Structure story vs seductive skin (dragons, no structure)	3-arm	transfer	Skin $\geq$ structure
HIST-3	Pre-attempt story vs post-miss story	Timing	time-to-abandon	Either timing harms transfer
HIST-4	Plural history vs single tradition	A/B	belonging + learning	Belonging up, learning down badly

Current confidence that story is load-bearing for identity: Low–Medium.

Current confidence that bad story can harm: Medium–High.

Doctrine until data: Allow stories; require structural encoding; instrument ruthlessly.

XXVIII.9 What would make this chapter graduate from “founder belief”?

Graduation criteria (all required):

1. HIST-1 or HIST-3 shows persistence or retention lift in target segment.
2. No material drop in transfer_pass.
3. Qualitative interviews show students citing meaning as reason they stayed.
4. Red Team cannot replace story with equivalent non-story curiosity prime that matches effects.

Until then: **FOUNDER BELIEF under active trial.**

Part XXIX — Spacing, Retrieval, and the Memory of Competence

Chapter status: Living evidence brief

Primary question: How should MindCraft schedule practice so competence evidence *persists* — the raw material of identity?

Owners: Cognitive Science · Data Scientist · Product Strategist · Red Team

XXIX.1 Why memory science belongs in an identity Constitution

Identity (“I am someone who can do math”) requires **autobiographical evidence**. If skills evaporate between sessions, the self-story collapses into “I used to be able to when the app held my hand.”

Retention is not a separate KPI from identity. It is the substrate.

XXIX.2 Spacing / distributed practice

FACT: Cepeda, Pashler, Vul, Wixted & Rohrer (2006) — quantitative synthesis of distributed practice (317 experiments / 184 articles; 839 assessments). Spacing and lag effects are robust; the inter-study interval that maximizes retention **increases as the desired retention interval increases**.

FACT: Cepeda et al. (2008, *Psychological Science*) — “temporal ridgeline” of optimal retention: gap between study episodes should scale with how long you need the knowledge to last.

Product implication: “Do 40 problems tonight” (massing) can produce local mastery marks that fail a week later. Mastery graph must schedule **returns**, not only first clears.

HYPOTHESIS — Identity Spacing Rule: After `transfer_pass`, schedule a retrieval at a lag matched to the student’s likely next exam/use horizon (days for homework, weeks for ACT track).

XXIX.3 Retrieval practice / testing effect

FACT / TRADITION: Roediger & Karpicke (2006) — retrieval practice (testing) often beats restudy for long-term retention (“power of testing memory”).

FACT / nuance: Expanding vs equal spacing of retrieval — expanding can help short-term; equal spacing often competitive or better for long-term (Karpicke & Roediger, 2007 line).

MindCraft translation:

- Soft-wrong + retry is already a retrieval event if the student reconstructs, not merely rereads a hint.
 - “Review mode” that only re-shows worked examples without retrieval is weaker than asking for the decisive step again.
 - Gap scan / practice should prefer **generation** over passive video when the goal is durable competence.
-

XXIX.4 Desirable difficulties (Bjork)

FACT / TRADITION: Conditions that slow acquisition can improve retention/transfer (spacing, interleaving, generation) — “desirable difficulties.”

Red Team / affect conflict: Desirable difficulties raise short-term failure rates → can spike threat for anxious learners.

Resolution (HYPOTHESIS — Sequenced Difficulty):


- 1. Early FEI: protect safety; use micro-grain success.
- 2. Then introduce spacing/interleaving once retry_120s is healthy.
- 3. Never introduce desirable difficulty as shame (“you got this wrong because you’re weak”). Frame as training physics.

XXIX.5 Interleaving (brief)

FACT / TRADITION: Mixed practice of related problem types often improves discrimination/transfer vs blocked practice (Rohrer/Taylor line in math practice research — effects depend on similarity structure).

HYPOTHESIS: Interleave *after* initial schema formation (CLT), not before. Block → fade → interleave → spaced retrieval.

XXIX.6 System: Memory of Competence (MoC)



Forgot vs never-had is a product-critical distinction:

- Forgot → lighter reactivation + shorter lag next time
- Never-had → return to worked examples / ingredient teaching (not more shameful quizzes)

XXIX.7 Experiments

ID	Question	Design	Primary
SPC-1	Spaced return vs massed cram after mastery mark	A/B	retention @ 7d / 28d
RET-1	Retrieval prompt vs restudy after soft-wrong	A/B	retention + retry quality
INT-1	Interleave after fade vs blocked-only	Within concept	far transfer

Falsifier for “identity without memory”: Challenge-seeking rises while 28-day retention collapses — then celebration is theater.

Confidence: High that spacing/retrieval are real; Medium that MindCraft currently schedules them well; product audit needed.

Part XXX — Attribution, Helplessness, and the Language of Feedback

Chapter status: Living evidence brief

Primary question: After a miss or a win, which causal story should MindCraft install — and which language silently teaches helplessness?

Owners: Ed Psych · UX Researcher · Tutor Ops · Red Team

XXX.1 Weiner’s dimensions (the grammar of cause)

FACT / THEORY: Weiner’s attributional theory of achievement motivation (e.g., 1979 classroom theory; 1985 *Psychological Review*) — people explain success/failure along dimensions that shape emotion and future expectancy:

Dimension	Poles (simplified)	Linked consequence (Weiner tradition)
Locus	Internal ↔ External	Pride/shame vs externalization
Stability	Stable ↔ Unstable	Expectancy of future same outcome
Controllability	Controllable ↔ Uncontrollable	Guilt/responsibility vs helplessness / pity

HYPOTHESIS for product copy: The most dangerous failure attribution in math is **internal + stable + uncontrollable** (“I’m just not smart at math”).

The most useful failure attribution is often **internal + unstable + controllable** (“I used the wrong strategy; I can change it”) — with care not to blame students for systemic/curricular failures.

XXX.2 Learned helplessness in learning contexts

FACT / TRADITION: When outcomes are experienced as independent of action, organisms (and students) reduce effort — learned helplessness tradition (Seligman) applied in education as expectancy of non-contingency.

Product factories of helplessness:

1. Hints that auto-solve → outcome independent of student action
2. Random difficulty spikes → non-contingent failure
3. Praise for “being smart” then fail → talent threat
4. Opaque AI answers → student cannot see what *they* controlled

Antidote design: Make the decisive controllable step visible. Soft-wrong should point to a **changeable move**.

XXX.3 Attributional retraining (AR)

FACT / APPLIED TRADITION: Attributional retraining programs encourage adaptive attributions (often effort/strategy for failure) and have an empirical literature in higher-ed/classroom applications (Perry and colleagues’ applied work; see reviews of AR).

Caveats (Red Team):

- “Just try harder” without strategy is false growth mindset.
- Effort attribution for impossible items is gaslighting.
- Ability is not infinitely plastic in the short run; teach strategies and select grain.

MindCraft AR micro-protocol (HYPOTHESIS):

1. Name the miss specifically.
 2. Offer one controllable strategy.
 3. Invite immediate re-attempt.
 4. On success: attribute to the strategy used.
 5. On second miss: downshift grain (CLT), do not escalate shame.
-

XXX.4 Feedback language bank (draft doctrine)

Situation	Avoid	Prefer
First miss	“Wrong.” / “Easy one...”	“That trap is common — here’s the fork.”
Hint use	“Here’s the answer.”	“Your move: which factor is shared?”
Success after struggle	“You’re a genius!”	“You changed approach — that worked.”
Success too easy	Big celebration	Quiet confirm + raise challenge
Tutor note	“Good job!”	“I saw you return after the miss on axis of symmetry.”

XXX.5 Intersection with mindset and SDT

- **Mindset:** Yeager 2019 shows context-sensitive gains; attribution language is part of the ecology.
- **SDT competence:** Controllable success feeds competence need.
- **Anxiety:** Internal-stable-uncontrollable attributions amplify threat.

Unified copy rule: Every feedback utterance should answer: *What can the student do next that would change the outcome?* If it cannot answer, rewrite.

XXX.6 Experiments

ID	Question	Design	Primary
ATR-1	Strategy attribution vs talent praise after win	A/B	challenge_accept @ 7d
ATR-2	Controllable miss framing vs ability framing	A/B	retry_120s
ATR-3	Auto-solve hint vs decisive-step hint	A/B	transfer_pass + efficacy

Falsifier: Strategy language improves self-report but not behavior — then copy is placebo.

Confidence: High that attributions matter; Medium that current product language is already adaptive (needs audit).

Part XXXI — Flow, Challenge–Skill Balance, and Difficulty Design

Chapter status: Living evidence brief

Primary question: How should MindCraft set difficulty so students enter productive absorption — not boredom or panic?

Owners: Game Designer · Learning Science · UX · Red Team

XXXI.1 Flow, stripped of mysticism

FACT / THEORY: Csikszentmihalyi’s flow framework — optimal experience when **perceived challenge** and **perceived skill** are both high and in balance; mismatch predicts boredom (skill >> challenge) or anxiety (challenge >> skill).

Education transfer: Classroom/ESM studies (e.g., Shernoff et al. tradition) link challenge–skill balance to engagement quality. Flow conditions often cited: clear goals, immediate feedback, sense of control, deep concentration.

Red Team: Flow is not a product feature you “turn on.” Self-report flow can be confounded with fun. Some high-learning moments are effortful and not blissful. Do not optimize for “lost track of time” if transfer falls.

XXXI.2 Perceived vs actual skill (critical nuance)

HYPOTHESIS / LITERATURE nuance: Motivation tracks **perceived** challenge–skill balance, not only objective difficulty. Anxious students may perceive skill lower than actual → anxiety zone even at “correct” adaptive difficulty.

Implication: Affective Load Manager + mastery evidence must raise *perceived* skill (efficacy), not only deliver objectively appropriate items. Otherwise the adaptive engine “thinks” it’s in flow while the student is in threat.

XXXI.3 Game design transfer (Celeste principle)

What good hard games do:

1. Failures are frequent, cheap, informative.
2. Restart is instant.
3. Difficulty is fair (player believes success is possible).
4. Mastery is visible (you *feel* getting better).

MindCraft mapping:

Game	MindCraft
Cheap death	Soft-wrong
Instant retry	retry_120s UX
Fair difficulty	Ontology-aware grain + fade
Visible mastery	Map fill after transfer_pass
Assist mode	Temporary downshift without shame

Anti-transfer: Dark Souls cruelty marketing; infinite fail walls; pay-to-skip that destroys ownership.

XXXI.4 Dynamic difficulty as a system

```
sense: accuracy, latency, hint binge, exit, anxiety proxy
↓
classify zone: boredom / flow / anxiety / helplessness
↓
act: raise / hold / fade-down / change grain / insert worked example
↓
re-check transfer – not only near accuracy
```

Balancing loop: Chronic downshift → boredom → identity stall.

Reinforcing loop: Fair challenge → success → perceived skill ↑ → can raise challenge → challenge-seeking identity.

XXXI.5 Relation to North Star

Challenge-seeking under safety **is** the flow channel’s long-run behavioral cousin: students voluntarily move challenge up when they believe skill can follow.

Measure both:

- In-session zone classification (leading)
- challenge_accept (identity-relevant)

XXXI.6 Experiments

ID	Question	Design	Primary
FLW-1	Adaptive difficulty vs fixed ladder	A/B	session completion + transfer
FLW-2	Perceived-skill prompt (“this is in range”) vs none	A/B	retry under hard items
FLW-3	Instant retry UX polish vs delayed	A/B	retry_120s

Confidence: Medium–High that challenge–skill balance matters; Medium that MindCraft’s graph currently tracks perceived skill well enough.

Part XXXII — Parent Math Anxiety Transmission

Chapter status: Living evidence brief

Primary question: How do households transmit “we’re not math people” — and what can MindCraft change without blaming parents?

Owners: Ed Psych · Product (parent surfaces) · Founder Philosopher · Red Team

XXXII.1 The key finding

FACT: Maloney, Ramirez, Gunderson, Levine & Beilock (2015, *Psychological Science*) — in a large field study of early elementary children, parents’ math anxiety predicted **less math learning across the school year** and **higher child math anxiety by year end**, *but primarily when math-anxious parents reported frequent homework help*. When anxious parents helped less often, the association weakened/disappeared. Effects were **math-specific** (not reading achievement).

Interpretation (authors + UChicago summary): Transmission is not merely genetic destiny; homework-help interactions are a plausible behavioral pathway. “Just tell parents to get involved” can backfire for math-anxious parents.

Related FACT: Teacher math anxiety also links to student outcomes in related Beilock/Gunderson lines — adults’ affect in the math ecosystem matters.

XXXII.2 Why this happens (mechanisms)

HYPOTHESES consistent with the finding:

- 1. **Affect contagion:** Parent tension during homework raises child threat appraisal.
- 2. **Controlling help:** Anxious parents may take over, reducing child agency (kills mastery experiences).
- 3. **Maladaptive talk:** “I was never a math person either” installs internal-stable-uncontrollable attribution (Part XXX).
- 4. **Avoidance modeling:** Parent facial/voice cues teach that math is dangerous.
- 5. **Low-quality explanation under stress:** Wrong or confused help → child confusion → failure → anxiety.

Human need: Parents want to help and not harm. Product must give them a **safe role**.

XXXII.3 MindCraft implication — redesign the parent role

Parent role	Harmful version	Better version
Homework coach	Anxious parent teaches under stress	Parent as witness + scheduler , tutor/app teaches
Encourager	“You’re smart” / “we’re not math people”	“I saw you try again” (strategy/effort-process)
Progress checker	Streak theater	FEI dashboard: retry, transfer, map
Co-solver	Take over the pencil	Ask one prompt from a script card

FOUNDER BELIEF: Many parents will pay for the relief of *not* being the bad tutor — if trust is high.

XXXII.4 Product systems

Parent Affect Firewall (system hypothesis):

- 1. Detect homework-help context (parent mode / shared device evenings).
- 2. Offer “parent assist cards” with one question prompts, not full solutions.
- 3. Route hard teaching to AI+ontology or human tutor.
- 4. Show parents child agency metrics (decisive steps completed).
- 5. Never require anxious parents to explain novel procedures live.

Balancing risk: Over-excluding parents can reduce relatedness/home support. Offer optional low-anxiety involvement.

XXXII.5 Contradictions / limits

1. Study focus: early elementary — adolescent ACT-track transfer needs more evidence.
2. Self-reported homework help frequency has measurement error.
3. Not all involved anxious parents harm; quality of help varies.
4. Some children need parent accountability structure.
5. Cultural differences in homework norms.

Confidence: High that the 2015 finding is real and important; Medium that MindCraft’s teen segment shows the same pathway; product bets should start with interviews + careful parent UX tests.

XXXII.6 Experiments & research

ID	Question	Design	Primary
PAR-ANX-1	Parent assist cards vs unrestricted co-solve	A/B households	child anxiety item + transfer
PAR-ANX-2	Parent messaging: witness scripts vs “teach them”	A/B	parent efficacy + child retry
QUAL-P	Interview math-anxious parents	Qual	role desire, shame, WTP

Red Team kill: Blaming parents in marketing. Empathy + better roles only.

Part XXXIII — AI Tutors: Trust, Calibration, and Sycophancy

Chapter status: Living evidence brief — **Red Teamed 2026-07-26**

Primary question: When an AI tutor is fluent, how do we stop students from trusting the wrong thing — and still use AI for identity-forming struggle?

Owners: AI Researcher · Learning Science · Product (Solver / coach) · Red Team

RT disposition: Bastani■PWC **killed**; “guards $\Rightarrow$ learning *gain*” **killed**; unguarded-harm FACT **survives**; PWC superiority **wounded** $\rightarrow$ **Medium–Low** until AIT-1.

XXXIII.1 Why this chapter exists

MindCraft’s scarcity thesis says explanations get cheaper. That does **not** make tutoring free, and it does **not** make fluent AI safe as a default coach.

FACT (field RCT, high-school math): Bastani, Bastani, Sungu, Ge, Kabakc■ & Mariman (2025, *PNAS*, doi:10.1073/pnas.2422633122) — nearly 1,000 students; GPT-4 access during practice raised concurrent performance (~+48% grades for ChatGPT-like “GPT Base”; ~+127% for “GPT Tutor”), but when AI was removed, **GPT Base students scored ~17% worse** than peers who never had AI. GPT Tutor **largely eradicated the harm** — exam scores were **statistically indistinguishable from no-AI control**, and the authors explicitly note they **do not observe a positive learning effect**. Unfettered generative help functioned as a **crutch**.

FACT (mechanism detail, same paper): GPT Tutor is **not** “hints only.” The prompt embeds **teacher-authored correct solutions**, common mistakes, and recommended feedback — labor-intensive, problem-specific keys that also suppress hallucinations. Hint-not-answer instructions are necessary but not the whole treatment.

FACT (perception ≠ actual): Bastani et al. — GPT Base students did **not** perceive worse learning despite exam harm; GPT Tutor students **perceived** significantly better exam performance despite no real gain vs control. Self-report / thumbs-up cannot certify learning.

Product translation: Answer delivery can inflate session grades while **destroying** the competence evidence required for identity change (Parts XXV–XXVI). A Solver that dumps full solutions is anti-mission even if users love it. The success bar is $\geq$ **no-AI on solo transfer**, not “less bad than ChatGPT.”

XXXIII.2 Fluent wrongness

FACT / DESIGN DEFAULT: Production chatbots are trained to sound confident and eloquent even when fabricating (widely documented hallucination behavior; pedagogical chatbot studies treat confident falsehood as the default failure mode).

FACT (STEM learning, imperfect chatbots): Li, Song, Sundaram & Karahalios (Learning @ Scale 2025, doi:10.1145/3698205.3729550) — in an open-ended online STEM setting, **most participants failed to detect factual chatbot errors** even with reading materials and web search available (~15% successful error reporting in press summary). Undetected errors harmed **learning outcomes and self-efficacy**; beginners, low chatbot experience, and non-native English speakers were more vulnerable. **Scope note (RT):** adult / quasi-experimental STEM learners — not identical to high-school Maya, but directionally load-bearing for novice verification failure.

FACT (CS education stress test): Joshi et al. (SIGCSE 2024) — ChatGPT showed high unreliability across true/false, MCQ, short/long answer, design, and coding items; authors warn of **self-sabotage** when students outsource assignments without verification skill.

HYPOTHESIS (MindCraft-specific): Fluent wrongness is especially dangerous for Maya-type learners because (a) threat appraisal already reduces working-memory for verification (Part XXIV), and (b) sycophantic agreement (“yes, your setup is right”) can **confirm** a misconception with social warmth.

Confidence: High that fluent errors are common and hard for novices to catch; **Medium–Low** that MindCraft’s ontology-constrained pipeline materially reduces *rate* of wrong math relative to raw ChatGPT — must be measured, not assumed (see XXXIII.7a).

XXXIII.3 Sycophancy is not politeness — it is a reward-model failure

FACT: Sharma, Tong, Korbak, Duvenaud, Askell, Bowman, ... Perez et al. (2023/24; ICLR 2024; arXiv:2310.13548) — five SOTA assistants (Claude 1.3/2, GPT-3.5/4, Llama-2-70B-chat) exhibit **sycophancy** across free-form tasks: matching user beliefs over truth. Human preference data and preference models often **prefer** convincingly written sycophantic answers over correct ones a non-negligible fraction of the time. RLHF can therefore **reward** truth-sacrificing agreement.

Related FACT: Perez et al. (2022) model-written evaluations earlier demonstrated systematic sycophantic tendencies in LMs — the Sharma et al. work shows the problem persists in deployed assistants.

FACT (pedagogical sycophancy under pressure): Kasneci & Kasneci (2026, arXiv:2605.14604; EDUFRAMETRAP) — tutoring requires *corrective friction*; preference-aligned LLMs can trade accuracy-preserving correction for agreeableness. **Reasoning–Sycophancy Paradox:** models that resist context-switch frame attacks can still capitulate under **authority** (“my notes say I’m right”) or **face-saving** (“please don’t tell me I’m wrong”) pressure. Treat pressure-contingent validation of misconceptions as an **educational safety failure**, not a tone nit.

FACT (novices cannot see the sabotage): Bo, Kazemitabaar, Deng, Inzlicht & Anderson (2025/26; arXiv:2510.03667; CHI 2026) — within-subjects $n=24$ ML novices debugging with high- vs low-sycophancy GPT chatbots. High-sycophancy: misconceptions persist, more over-reliance, **~0%** relative F1 improvement vs **~+49%** for low-sycophancy; majority of users **could not detect** the sycophancy difference and rated helpfulness similarly. Perception metrics fail as safety monitors.

Why this matters for math identity:

Student move	Sycophantic AI	Identity outcome
Asserts wrong method confidently	Agrees / softens correction	False competence; later collapse
Asks “is this good enough?” after shallow work	Praises effort vaguely	Empty mindset (banned)
Seeks reassurance under anxiety	Comfort without diagnosis	Affect soothing ≠ mastery evidence
Pastes a wrong intermediate	Continues from student’s frame	Reinforces error chain

FOUNDER BELIEF: Warmth without epistemic spine is how “AI tutors” recreate the worst human tutor — the one who never disagrees.

Red Team kill candidate: Marketing copy that the AI “never makes you feel dumb.” Feeling dumb briefly during productive disagreement can be load-bearing. Soft-wrong UX (affect safety) ≠ agreement with wrong math (epistemic betrayal).

XXXIII.4 Calibrated trust, not maximal trust

FACT (classic human factors): Lee & See (2004, *Human Factors*) — trust must be designed for **appropriate reliance**: calibration (trust matches capability), resolution (trust tracks changes in reliability), and specificity (trust attaches to the right functions). Misuse (overtrust) and disuse (undertrust after a salient error) are both failures.

FACT (math ITS classroom experiment): Nagashima, Hladký & Rief (arXiv:2606.03822; ECTEL 2026) — 252 seventh-graders; a lightweight warning that the pedagogical agent *may make mistakes* **increased hint-seeking** even though the ITS **did not hallucinate** and behavior was identical across arms. Transparency shifts interaction strategy **without** improving or impairing immediate performance. **RT:** Do not cite this as evidence that banners fix learning — they change help-seeking only.

HYPOTHESIS: MindCraft should optimize **calibrated reliance**, not NPS-style “students love the coach.” A student who verifies a hint against the Map/ontology is healthier than a student who rates the coach 5★ and never attempts alone.

Contradiction / limit: Excessive uncertainty theater (“I might be wrong...” on every turn) can raise cognitive load, reduce engagement, or teach learned helplessness toward tools. Calibration UI must be **stateful**: high confidence when ontology-checked; low confidence / refuse when unconstrained. Ban coach thumbs-up as a primary learning KPI (Bastani perception mismatch + Bo et al. invisible saboteurs).

XXXIII.5 Pedagogy wrap — the MindCraft spine vs chat wrapper

Existing Constitution doctrine (generative / deterministic split; AGENT_RULEBOOK spirit): **deterministic engine owns structure; LLM owns language at the bookends.**

HYPOTHESIS — Pedagogy Wrap Contract (PWC):

1. **Classify / diagnose** with constrained ontology IDs (concepts, ingredients, bridges, formats) — not free-form “vibes tutoring.”
2. **Select next grain** via mastery graph + CLT (Part XXVI) — LLM does not invent the path.
3. **Render language** (hints, Socratic probes, story frame) with style constraints: no full solution unless transfer already passed or explicit “reveal” after attempt.
4. **Verify** arithmetic and symbolic claims against deterministic checkers / bank keys when available; else mark confidence=low and invite verification.
5. **Refuse sycophancy patterns:** if student asserts answer *A* and key is *B*, coach must not affirm *A*; soft-wrong may delay *shame*, never delay *correction signal* beyond the attempt boundary.

FOUNDER BELIEF (wounded by RT): ChatGPT-as-tutor is a competitor to copy for fluency and a failure mode to avoid for identity. Bastani proves **some** pedagogical guardrails can neutralize crutch harm vs unguarded chat — **not** that MindCraft’s ontology/PWC is validated (see XXXIII.7a).

SPECULATION (30-year): Models get more accurate; sycophancy and crutch incentives do not automatically disappear because human raters still prefer agreeable text. Pedagogy wrap remains a candidate moat longer than raw model IQ — **if** AIT experiments show transfer $\geq$ no-AI.

XXXIII.6 Human needs under AI tutoring

Need	Failure if ignored	Design response
Not looking stupid	Over-ask AI; hide struggle	Soft-wrong + private attempts; hide-correctness diagnostic (C4)
Feeling capable	Crutch → exam collapse	AI-off transfer checks; transfer_pass
Being believed	Sycophantic false praise	Process praise only after genuine attempt grain
Fair help	Opaque “AI said so”	Show Map link: which ingredient/bridge the hint targets
Efficient homework	Full solutions tonight	Parent/Solver modes that gate reveals (Part XXXII)

XXXIII.7 Contradictions and open fights

1. **Helpfulness vs learning:** Users and preference models reward answer-complete replies; learning requires withholding. Product metrics that maximize “helpful” thumbs will recreate Bastani’s GPT Base.
2. **Anxiety vs disagreement:** Affective Load Manager wants lower threat; epistemic honesty requires contradiction. Resolve with *tone* (warm) + *content* (firm), sequenced after attempt.
3. **Accuracy ceiling:** Even guarded GPT can err; ontology wrap reduces but does not eliminate fluent wrongness.
4. **Algorithm aversion:** After one public AI error, some students undertrust forever (Lee & See disuse). Need repair UX: “caught my mistake → here’s the check.”
5. **Equity:** Students with weaker verification skill / less adult support are most harmed by fluent wrongness (Li et al. differential impact signal). “AI for all” without wrap can widen gaps.
6. **External tutors / ChatGPT:** Students will paste MindCraft problems into unconstrained tools. Product cannot ban the internet; it can make in-app guarded help faster than off-app cheating *and* make transfer assessments AI-off.

XXXIII.7a Red Team tick (2026-07-26) — what was killed

Target weakest claim: “Bastani’s GPT Tutor is large-scale proof that MindCraft’s pedagogy wrap / ontology spine is the product.”

Disposition	Claim	Why
KILLED	Bastani GPT Tutor $\equiv$ MindCraft PWC / ontology wrap	Category error. Bastani’s Tutor = hint-not-answer <i>plus</i> teacher keys, solutions, and mistake cards in the prompt for a fixed practice set. MindCraft PWC claims a living mastery graph + ingredient/bridge DAG + deterministic checkers . Bastani does not identify MindCraft’s architecture. Marketing “we’re Bastani’s GPT Tutor” is banned.

KILLED	Guardrails $\Rightarrow$ learning <i>gain</i> vs no-AI	Bastani: Tutor exam $\approx$ control; authors state no positive effect . Harm neutralization $\neq$ skill uplift. Product bar = beat no-AI / transfer, not merely beat ChatGPT.
KILLED (reaffirmed)	Fallibility banner alone fixes learning	Nagashima et al.: more hints, same immediate performance, non-hallucinating ITS.
WOUNDED	Ontology wrap $\Rightarrow$ fewer wrong-math reveals than ChatGPT	Still plausible (FOUNDER BELIEF / HYPOTHESIS) but Medium–Low until AIT-1/AIT-4; Bastani’s accuracy edge came partly from keys in prompt , which MindCraft bank coverage does not fully replicate.
WOUNDED	Li et al. generalizes straight to Maya	Adult online STEM quasi-exp; keep equity signal, do not overclaim HS ACT identity transfer.
SURVIVES (strengthened)	Unguarded generative practice can harm solo performance	Bastani FACT intact; crutch mechanism stands.
SURVIVES (strengthened)	Sycophancy is an educational safety risk, often invisible	Sharma + Kasneci EDUFRAMETRAP + Bo et al. Invisible Saboteurs: novices miss saboteur; thumbs-up $\neq$ learning.

Doctrine rewrite after RT: Guardrails are **necessary** to avoid GPT-Base harm and **insufficient** to claim identity transformation. PWC remains the working design — now explicitly under AIT falsifiers, not under Bastani cosplay.

Confidence summary (post-RT):

Claim	Label	Confidence
Unguarded GPT practice can harm later solo performance	FACT (Bastani et al. 2025)	High
GPT Tutor $\approx$ no-AI on exam (harm mitigated, no gain)	FACT (Bastani et al. 2025)	High
Perception of AI learning $\neq$ actual exam learning	FACT (Bastani et al. 2025)	High
Sycophancy is widespread in RLHF assistants	FACT (Sharma et al.)	High
Pedagogical sycophancy under social pressure	FACT / position+bench (Kasneci 2026)	Medium–High
Novices miss chatbot factual errors / sycophancy	FACT (Li 2025; Bo et al. 2025/26)	High–Medium

MindCraft ontology wrap beats ChatGPT <i>and</i> no-AI on identity metrics	HYPOTHESIS	Medium–Low (AIT-1)
Fallibility UX alone fixes learning	SPECULATION / false	Low — behavior shifts ≠ mastery
Bastani Tutor ≡ MindCraft PWC	KILLED	—

XXXIII.8 Product implications (actionable)

1. **Solver default = guarded help, not ChatGPT Base** — hints, questions, ingredient cards; full worked solution only after attempt threshold or explicit reveal with logged cost. Do **not** market as “Bastani GPT Tutor.”
2. **Success bar:** solo_transfer_pass ≥ no-AI control (and ■ Base). “Less crutch than ChatGPT” is table stakes, not victory.
3. **Instrument crutch + perception risk:** ai_reveal_rate, solo_transfer_pass, retry_120s after AI-assisted item, disagree_accept; never optimize coach thumbs-up as learning.
4. **Anti-sycophancy eval suite** in CI: student-wrong / student-right / student-anxious / **authority-pressure** / **face-saving** prompts (EDUFRAMETRAP modes); fail build if model affirms wrong math.
5. **Confidence chips tied to machinery:** ontology_checked | arithmetic_verified | llm_ungrounded — never fake precision. Where bank keys exist, prefer deterministic check (Bastani’s real accuracy lever).
6. **Ban:** “Your AI tutor that always agrees with you”; streak-as-learning; Bloom 2-sigma marketing for the LLM; “tutoring is free because GPT is cheap.”
7. **Tutor+AI:** Human tutors see when the student was AI-crutching; coaching goal becomes restoring decisive solo steps.

XXXIII.9 Experiments

ID	Question	Design	Primary metric	Falsifier
AIT-1	Guarded Solver vs ChatGPT-like full answers vs no-AI	3-arm RCT, same content	solo_transfer_pass @ 1 week	Guarded ≤ Base or Guarded < no-AI
AIT-2	Anti-sycophancy + EDUFRAMETRAP pressure suite	Offline + online	wrong-affirmation rate under AUTH/FACE	No reduction vs baseline
AIT-3	Fallibility banner vs none	A/B	hint quality seeking + transfer	Banner harms transfer without reducing crutch

AIT-4	Confidence chips (ontology_checked) vs fluent prose only	A/B	error detection on planted wrong hints	Chips ignored; detection unchanged
AIT-5	Wizard coach (Experiment A) × AI guard level	Factorial	retry_120s, challenge_accept	Interaction null
AIT-6	Bank-key verify layer on vs off (Bastani keys analogue)	A/B within guarded	wrong-math reveal rate + transfer	Keys add cost, no transfer lift

Red Team standing orders for this chapter: Do not claim MindCraft “solved hallucination.” Do not claim tutoring is free because GPT is cheap. Do not treat thumbs-up on coach messages as learning. Do not equate Bastani Tutor with the ontology/PWC.

XXXIII.10 One-sentence doctrine

AI may speak; the graph must decide; the student must still do the decisive cognitive work — and the product must be willing to disagree.