

MINDCRAFT

Research Constitution

How do humans become confident mathematical thinkers?

Identity transformation · Evidence · Red Team · Product systems · Experiments

This is not a pitch deck. This is the company's research operating system.
Challenge assumptions. Label uncertainty. Kill weak arguments.
Source corpus: ~338,128 words · multi-chapter living lab

MindCraft Research Lab · ongoing evidence program

MindCraft Research Constitution v1

Status: Living operating document — not a pitch deck

Edition: v1.19 (Synthesizer consolidation of Parts CXLIII–CL into surviving doctrine)

Research question: How do humans become *confident mathematical thinkers*?

Product thesis under audit: The product is identity transformation, not mathematics delivery.

Last updated: 2026-08-15 (Synthesizer v1.19 — CXLIII–CL company law)

Growth model: Core OS (this file) + chapters/*.md via CHAPTER_MANIFEST.txt → PDF

Scale intent: Multi-month densification toward 150–300 pages of *evidenced* material — never fluff

Epistemic rule: Every claim is labeled FACT / HYPOTHESIS / FOUNDER BELIEF / SPECULATION.

Deep-dive chapters currently mounted

Part	File	Focus
XXIV	chapters/24_affect_anxiety_belonging.md	Anxiety, WM, stereotype threat, belonging
XXV	chapters/25_self_efficacy_narrative_identity.md	Bandura sources, narrative identity, habit≠identity
XXVI	chapters/26_cognitive_load_mastery_tutoring.md	CLT, fading, mastery war, tutoring calibration
XXVII	chapters/27_markets_parents_competition.md	WTP, parents, competitive teardown
XXVIII	chapters/28_history_meaning_math.md	Invention-story thesis under trial
XXIX	chapters/29_spacing_retrieval.md	Spacing, retrieval, memory of competence
XXX	chapters/30_attribution_helplessness.md	Weiner attributions, helplessness, feedback language
XXXI	chapters/31_flow_challenge_skill.md	Flow, challenge–skill, difficulty design
XXXII	chapters/32_parent_anxiety_transmission.md	Maloney/Beilock parent anxiety pathway
XXXIII	chapters/33_ai_tutors_trust_sycophancy.md	AI trust / sycophancy; RT: Bastani■PWC, guards≠gain
XXXIV	chapters/34_formal_causal_dag_identification.md	FEI causal DAG; L0–L4 claim ladder; confounders
XXXV	chapters/35_competitive_session_audits.md	Khan/Duo/Brilliant/ChatGPT session mechanism audits
XXXVI	chapters/36_equity_audit_story_worlds.md	Equity audit of story worlds; whose history; tokenism kill

XXXVII	chapters/37_expectancy_value_eccles.md	Eccles SEVT; utility/cost/attainment → choice; EVT experiments
XXXVIII	chapters/38_goal_orientation.md	Mastery vs performance goal structures; appearance ban; GO experiments
XXXIX	chapters/39_interleaving_vs_blocking.md	Interleave vs block; strategy selection; IL experiments
XL	chapters/40_self_explanation_prompts.md	Self-explanation prompts; Chi/Renkl → coach UX; SE experiments
XLI	chapters/41_desirable_difficulties_x_anxiety.md	Desirable difficulties × anxiety; SAFE-DD; Bjork vs Ashcraft
XLII	chapters/42_social_comparison_leaderboards.md	Social comparison & leaderboards; SAFE-COMPARE; ranks help/hurt
XLIII	chapters/43_habit_formation_science.md	Habit formation (Wood/Clear); SAFE-HABIT; cue≠colonization
XLIV	chapters/44_intrinsic_motivation_killers.md	Intrinsic motivation killers (Deci); SAFE-REWARD; controlling≠informational
XLV	chapters/45_mathematical_resilience.md	Mathematical resilience (JW/Lee); SAFE-RESILIENCE; grit theater kill
XLVI	chapters/46_teacher_tutor_expectancy.md	Tutor expectancy (Pygmalion/Golem); SAFE-EXPECTANCY; trait-label ban
XLVII	chapters/47_sleep_stress_learning.md	Sleep/stress/learning; SAFE-TIMING; all-nighter & sleep-app kill
XLVIII	chapters/48_transfer_variation_theory.md	Transfer & variation theory; SAFE-TRANSFER; story≠far-transfer kill
XLIX	chapters/49_misconceptions_as_productive.md	Misconceptions as productive; SAFE-MISCON; soft-wrong / diagnostic wrongs
L	chapters/50_confidence_calibration.md	Confidence calibration; SAFE-CALIB; over/underconfidence; C4 hide-correctness

LI	chapters/51_deliberate_practice_tutoring.md	Deliberate practice in tutoring; SAFE-DP; Ericsson vs Hambrick; 45-min spine
LII	chapters/52_community_of_practice.md	Community of practice; SAFE-CoP; Lave/Wenger LPP; careful school transfer
LIII	chapters/53_religion_ritual_light_touch.md	Religion/ritual light-touch; SAFE-RITUAL; meaning without cult dynamics
LIV	chapters/54_military_aviation_brief_debrief.md	Military/aviation brief-debrief; SAFE-AAR; after-action reviews for sessions
LV	chapters/55_sports_film_study_pedagogy.md	Sports film-study pedagogy; SAFE-FILM; error clips → coach cards
LVI	chapters/56_music_pedagogy_ladders.md	Music pedagogy ladders; SAFE-MUSIC; scales → repertoire → recital
LVII	chapters/57_chess_annotation_metacognition.md	Chess annotation & metacognition; SAFE-ANNOT; postmortem UX
LVIII	chapters/58_therapy_graded_exposure.md	Therapy graded exposure; SAFE-EXPOSE; math anxiety ladders
LIX	chapters/59_parent_wtp_experiments.md	Parent WTP experiments; SAFE-WTP; CBC/conjoint + message tests
LX	chapters/60_trust_after_ai_hallucination.md	Trust after AI hallucination; SAFE-REPAIR; escalate-to-human
LXI	chapters/61_measurement_of_math_identity.md	Measurement of math identity; SAFE-IDMEASURE; scales + demand effects
LXII	chapters/62_longitudinal_identity_change.md	Longitudinal identity change; SAFE-LONGID; 8/26/52 week horizons
LXIII	chapters/63_act_exam_pressure.md	ACT/exam pressure; SAFE-EXAM; learn affect vs exam affect
LXIV	chapters/64_multilingual_ell_math_identity.md	Multilingual/ELL math identity; SAFE-ELL; language vs math load

LXV	chapters/65_gender_math_stereotypes_update.md	Gender & math stereotypes update; SAFE-GENDER; parity vs climate
LXVI	chapters/66_socioeconomic_constraint_on_grit.md	Socioeconomic constraint on grit; SAFE-STRUCTURE; structure > slogans
LXVII	chapters/67_ontology_as_diagnosis_moat.md	Ontology as diagnosis moat; SAFE-ONTOLOGY; state×graph×FEI×transfer
LXVIII	chapters/68_human_in_the_loop_tutor_ops.md	Human-in-the-loop tutor ops; SAFE-HITL; playbooks/QA/FEI training
LXIX	chapters/69_privacy_student_affect_data.md	Privacy & student affect data; SAFE-PRIVACY; ERT ban; check-in ethics
LXX	chapters/70_constitution_meta_bayesian_lab.md	Constitution meta; SAFE-LABMETA; falsify/check/demote lab process
LXXI	chapters/71_peer_nearpeer_expert_tutoring.md	Peer vs near-peer vs expert tutoring; SAFE-TUTORGRAIN; hire bar
LXXII	chapters/72_generated_question_validity_key_risk.md	Generated question validity & key risk; SAFE-GENQ; verify-before-ship
LXXIII	chapters/73_district_procurement_privacy_gtm.md	District procurement & privacy GTM; SAFE-PROCURE; trust packet / NDPA
LXXIV	chapters/74_spaced_retrieval_schedules_product_ux.md	Spaced retrieval schedules in product UX; SAFE-SCHED; ISI×RI / equal vs expanding
LXXV	chapters/75_tutor_workforce_pipeline_quality_drift.md	Tutor workforce pipeline & quality drift; SAFE-WORKFORCE; fidelity over tenure
LXXVI	chapters/76_forgetting_curves_product_honesty.md	Forgetting curves as product honesty; SAFE-FORGET; aged evidence / delayed probes
LXXVII	chapters/77_adaptive_spacing_algorithms_vs_fixed_calendars.md	Adaptive spacing vs fixed calendars; SAFE-ADAPT; calendar-first / banded personalization

LXXVIII	chapters/78_cram_products_vs_durable_identity.md	Cram products vs durable-identity positioning; SAFE-DURABLE; dual-rail GTM / no point guarantees
LXXIX	chapters/79_feedback_timing_immediate_vs_delayed.md	Feedback timing immediate vs delayed; SAFE-FBTIME; mode-conditional clocks / focus>timing
LXXX	chapters/80_bridge_gap_product_narrative.md	Bridge-gap product narrative; SAFE-BRIDGE; connection failures > isolated nodes
LXXXI	chapters/81_cold_start_without_fake_mastery.md	Cold-start without fake mastery; SAFE-COLD; humble prior + labeled seed + probes
LXXXII	chapters/82_format_axis_product_narrative.md	Format-axis product narrative; SAFE-FORMAT; FormatId gaps / conversion evidence
LXXXIII	chapters/83_parent_dashboard_honesty_ux.md	Parent dashboard honesty UX; SAFE-PDASH; proof age / MoC without shame
LXXXIV	chapters/84_tutor_talk_ratio_instrumentation.md	Tutor talk-ratio instrumentation; SAFE-TALK; construction > airtime
LXXXV	chapters/85_story_world_cognitive_load_vs_identity.md	Story-world cognitive load vs identity payoff; SAFE-STORYLOAD; CLT-budgeted wrap
LXXXVI	chapters/86_competence_evidence_social_proof.md	Competence evidence as social proof; SAFE-PROOF; witnessable solo MoC
LXXXVII	chapters/87_bank_coverage_honesty_marketing.md	Bank coverage honesty in marketing; SAFE-COVER; matrix > item-count
LXXXVIII	chapters/88_worked_example_fading_solver_ux.md	Worked-example fading in Solver UX; SAFE-FADE; completion → solo
LXXXIX	chapters/89_help_seeking_vs_help_abuse.md	Help-seeking vs help abuse; SAFE-HELP; instrumental > executive
XC	chapters/90_product_analytics_fei_north_stars.md	Product analytics for FEI North Stars; SAFE-INSTRUMENT; ship the four

XCI	chapters/91_hint_economy_contingent_scaffolding.md	Hint economy & contingent scaffolding; SAFE-HINT; soft→hard priced peeks
XCII	chapters/92_explanation_length_germane_load.md	Explanation length vs germane load; SAFE-EXPLAIN; short coach vs monologue
XCIII	chapters/93_retrieval_failure_modes_practice_ux.md	Retrieval failure modes; SAFE-RETRIEVE; TOT vs blank vs struggle bounds
XCIV	chapters/94_productive_failure_vs_guided_success.md	Productive failure vs guided success; SAFE-PF; generate→consolidate fidelity
XCV	chapters/95_cognitive_apprenticeship_tutor_playbooks.md	Cognitive apprenticeship in tutor playbooks; SAFE-APPRENTICE; model→coach→fade
XCVI	chapters/96_attention_residue_device_distraction.md	Attention residue & device distraction; SAFE-ATTN; attempt-window protection
XCVII	chapters/97_dual_coding_diagram_formatid_load.md	Dual coding & diagram FormatId load; SAFE-DUAL; contiguous structural figures
XCVIII	chapters/98_metacognitive_monitoring_gap_scan.md	Metacognitive monitoring in gap-scan; SAFE-MONITOR; named judgments → control
XCIX	chapters/99_opportunity_cost_tutor_minutes_vs_ai.md	Tutor minutes vs AI scaffold; SAFE-ROI; FEI per tutor-minute / Map minute budgets
C	chapters/100_worked_example_vs_problem_solving_timing.md	Worked-example vs PS timing across expertise; SAFE-EXPTIME; grain-local adaptive fade
CI	chapters/101_self_regulated_learning_cycles_practice_ux.md	SRL cycles (Zimmerman) in Practice UX; SAFE-SRL; phase→control micro-prompts
CII	chapters/102_generative_learning_activities_practice.md	Generative learning (Fiorella/Mayer) in Practice; SAFE-GENERATE; SOI-forcing micro-GLAs
CIII	chapters/103_error_climate_psychological_safety_tutoring.md	Error climate & psychological safety in tutoring; SAFE-ERRCLIMATE; stay-diagnose under FEI

CIV	chapters/104_homework_help_vs_practice_identity.md	Homework help vs practice identity; SAFE-HWHELP; dual-rail Solver/Practice
CV	chapters/105_gesture_embodiment_math_practice.md	Gesture & embodiment in math practice; SAFE-GESTURE; strategy-tied optional motion
CVI	chapters/106_peer_explanation_quality_knowledge_building.md	Peer explanation quality; SAFE-PEERX; knowledge-building vs telling
CVII	chapters/107_curiosity_prediction_before_reveal.md	Curiosity & prediction before reveal; SAFE-PREDICT; commit→surprise→encode
CVIII	chapters/108_analogical_reasoning_bridge_teaching.md	Analogical reasoning & bridge teaching; SAFE-ANALOG; structure-map joins
CIX	chapters/109_writing_to_learn_math_practice.md	Writing-to-learn in math practice; SAFE-WRITE; Notes/protocol chips
CX	chapters/110_spatial_skills_math_identity.md	Spatial skills & math identity; SAFE-SPATIAL; FormatId hops
CXI	chapters/111_formative_assessment_cycles_tutor_sessions.md	Formative assessment cycles; SAFE-FA; exit tickets vs FEI
CXII	chapters/112_feedback_literacy_coach_uptake.md	Feedback literacy / coach uptake; SAFE-FBLIT; delivery≠uptake
CXIII	chapters/113_collaborative_problem_solving_vs_solo_fei.md	Collaborative PS vs solo FEI; SAFE-COLLAB; group≠solo mastery
CXIV	chapters/114_desirable_difficulty_scheduling_under_exam_deadlines.md	DD scheduling under exam deadlines; SAFE-DDSCHED; hold mix / ease load
CXV	chapters/115_tutor_identity_belonging_near_peer_role_models.md	Tutor identity & belonging; SAFE-ROLEMODEL; fidelity-first near-peer
CXVI	chapters/116_offline_practice_fidelity_paper_whiteboard_capture.md	Offline practice fidelity; SAFE-OFFLINE; attempt→capture→confirm
CXVII	chapters/117_worked_example_fading_under_high_math_anxiety.md	Fade × anxiety; SAFE-FADEX; protect then fade
CXVIII	chapters/118_parent_message_framing_after_soft_wrong_streaks.md	Parent soft-wrong framing; SAFE-SOFTMSG; grain+CTA not nag

CXIX	chapters/119_classroom_orchestration_map_live_tutoring.md	Map live-session orchestration; SAFE-MAPORCH; brief→slice→adjust
CXX	chapters/120_tutor_burnout_early_warning_without_wellness_theater.md	Tutor burnout early-warning; SAFE-BURNWARN; fidelity/load→redesign
CXXI	chapters/121_explanation_quality_rubrics_ai_coach_wraps.md	AI coach wrap QA rubrics; SAFE-EXPLAINQA; principle-short gates
CXXII	chapters/122_partial_credit_soft_wrong_taxonomy.md	Soft-wrong mark taxonomy; SAFE-TAXON; type before points
CXXIII	chapters/123_exam_week_tutor_load_shedding_playbooks.md	Exam-week tutor load shedding; SAFE-LOADSHED; acuity shed under crunch
CXXIV	chapters/124_distractor_design_diagnostic_soft_wrong_mc.md	Distractor design for soft-wrong MC; SAFE-DISTRACTOR; misconception-mapped options
CXXV	chapters/125_confidence_tagged_misses_gap_scan_recovery.md	Confidence-tagged gap-scan miss recovery; SAFE-CONFMISS; CRI fork without Identity Score™
CXXVI	chapters/126_tutor_micro_debrief_fidelity_under_time_caps.md	Tutor micro-debrief fidelity under time caps; SAFE-MICRODEB; capped AAR without lecture cosplay
CXXVII	chapters/127_parent_crunch_week_message_tests.md	Parent crunch-week message tests; SAFE-CRUNCHMSG; shed honesty vs 24/7 CBC
CXXVIII	chapters/128_format_trap_distractors_diagram_graph_items.md	Format-trap distractors on diagram/graph items; SAFE-FORMTRAP; visual traps without decoration theater
CXXIX	chapters/129_session_open_brief_fidelity_under_late_joins.md	Session-open brief fidelity under late joins; SAFE-MICROBRIEF; capped shared-model open
CXXX	chapters/130_dual_rail_solver_cta_copy_under_crunch.md	Dual-rail Solver CTA copy under crunch; SAFE-CRUNCHCTA; help tonight vs Practice gym without Chegg hero
CXXXI	chapters/131_figure_key_verify_loops_generated_format_id.md	Figure key-verify loops for generated FormatId stems; SAFE-FIGKEY; structural gates without fluency SVG pass

CXXXII	chapters/132_async_prebrief_cards_late_tutor_wait.md	Async pre-brief cards while waiting for late tutor; SAFE-ASYNCBRIEF; student-owned Set without idle lobby
CXXXIII	chapters/133_one_tap_solver_practice_transfer_grain.md	One-tap Solver→Practice transfer grain; SAFE-XFERGRAIN; handoff without dump-and-quit
CXXXIV	chapters/134_hide_correctness_eligibility_formatid.md	Hide-correctness FormatId eligibility gates; SAFE-HCELIG; C4 stricter than Practice browse
CXXXV	chapters/135_known_finite_eta_honesty_late_tutor_lobby.md	Known-finite ETA honesty in late-tutor lobby; SAFE-KNOWNETA; finite band without soon/shame
CXXXVI	chapters/136_soft_require_vs_hard_gate_follow_on_help_avoidance.md	Soft-require vs hard-gate follow-on; SAFE-SOFTGATE; soft-require default under help-avoidance
CXXXVII	chapters/137_reveal_practice_vs_c4_split_pools_sealed_formatid.md	Reveal-Practice vs C4 split pools; SAFE-SPLITPOOL; graded use without silent Map writes
CXXXVIII	chapters/138_formatid_gap_scan_composition_sealed_stock_scarcity.md	FormatId gap-scan composition under sealed-stock scarcity; SAFE-SCANCOMP; honest holes vs pad theater
CXXXIX	chapters/139_parent_trust_copy_visual_diagnostic_holes.md	Parent trust copy when visual diagnostic holes remain; SAFE-HOLETRUST; labeled holes without shame/pad
CXL	chapters/140_promote_to_c4_tutor_qa_workflow.md	Promote-to-C4 tutor QA for Practice-only FormatId; SAFE-PROMOTE; ticketed checklist without play-count flip
CXLI	chapters/141_exposure_cool_down_after_reveal_same_stem_c4.md	Exposure cool-down after reveal before same-stem C4; SAFE-COOLDOWN; anti-answer-hunt without pool starvation
CXLII	chapters/142_student_facing_hole_copy_after_gap_scan.md	Student-facing hole copy after gap-scan; SAFE-STUDHOLE; no shame, no fake greens; parity with HOLETRUST

CXLIII	chapters/143_promote_backlog_sla_parent_hole_eta.md	Promote backlog SLA vs parent hole-ETA honesty; SAFE-FILLETA; capacity-grounded verify bands
CXLIV	chapters/144_second_rater_thresholds_figured_c4.md	Second-rater thresholds for figured C4 admits; SAFE-DUALRATE; risk-triggered dual without cost theater
CXLV	chapters/145_isomorphic_stem_family_hashing_cooldown.md	Isomorphic stem-family hashing for cool-down; SAFE-FAMHASH; sealed family unit without over-block
CXLVI	chapters/146_gap_scan_exit_emotion_without_anxiety_score.md	Gap-scan exit emotion without Anxiety Score™; SAFE-EXITEMO; destake copy after holes
CXLVII	chapters/147_dual_rater_capacity_planning_filleeta.md	Dual-rater capacity planning vs FILLETA band honesty; SAFE-CAPPLAN; dual minutes in fill estimator
CXLVIII	chapters/148_template_trust_inheritance_first_of_cell_dual.md	Template-trust inheritance after first-of-cell dual; SAFE-TMPTRUST; inherit intensity not seals
CXLIX	chapters/149_post_promote_repair_stratified_dashboards.md	Post-promote repair stratified dashboards; SAFE-REPAIRSTRAT; path-true IRR without Kappa % NS
CL	chapters/150_dual_capacity_hire_vs_contract_sme_panels.md	Dual-capacity hire vs contract SME panels; SAFE-SURGEANEL; make-or-buy dual under exam spikes
CLI	chapters/151_template_trust_revoke_sample_dual_cadence.md	Template-trust revoke & sample-dual cadence; SAFE-REVOKE; switching-rule tighten/revoke under repair spikes
CLII	chapters/152_path_conditioned_dual_blame_selection_vs_quality.md	Path-conditioned dual blame tests; SAFE-PATHBLAME; selection vs quality without Dual Shame Score™

CLIII	chapters/153_contract_sme_panel_seal_onboarding_parallel_run.md	Contract SME panel seal onboarding & parallel-run gates; SAFE-PANELRAMP; ramp before C4 authority
CLIV	chapters/154_post_spike_dual_demobilization_fidelity_cliff.md	Post-spike dual demobilization without fidelity cliff; SAFE-DEMOB; handover + capture, no invent-work

Queued next: see NEXT_LAB.md (Part CLV gold-set path×stratum probes, then 156+). Company law through CL (v1.19); CLI–CLIV provisional pending synthesizer.

Synthesizer note (v1.19): Eight researcher chapters (CXLIII–CL) landed after v1.18 as provisional I.4 densifiers — capacity-grounded hole-fill ETA bands, risk-triggered second-rater thresholds, sealed stem-family hashing for cool-down, destake gap-scan exit without Anxiety Score™, dual minutes as first-class FILLETA capacity, conditional template-trust inheritance after blueprint dual, path-stratified post-promote repair dashboards without Kappa % NS, and staff-led hire/contract SME panels for exam-week dual surge. This edition promotes each stack to company law (removes provisional tags), formalizes Red Team kills #97–#104, confirms FILLETA/DUALRATE/FAMHASH/EXITEMO/CAPPLAN/TMPTRUST/REPAIRSTRAT/SURGEANAL experiment families in Part IX, and refreshes metrics/glossary/competitive implication. Deep-dive files remain authoritative for citations; the OS keeps only *surviving* product rules.

0. How to read this document

This Constitution is the research OS for MindCraft. It exists to prevent the company from building features that feel right and fail quietly.

Rules of engagement:

1. Founder beliefs enter as **FOUNDER BELIEF**, not as law.
2. The Red Team’s job is destruction. Unchallenged claims are invalid.
3. Page count is not quality. A shorter true section beats a longer vague one.
4. “AI will make explanations free” is a hypothesis about markets, not a moral slogan.
5. Product implications must map to experiments with falsifiers.

This v1 is dense and incomplete on purpose. Expanding toward 150–300 pages is allowed only when new evidence, frameworks, or contradicted priors justify the ink.

Part I — Executive Summary

I.1 The compressed answer

HYPOTHESIS (confidence: medium–high): Humans become confident mathematical thinkers when three conditions compound:

1. **Affective permission** — the nervous system stops treating math as social/physical threat.
2. **Competence evidence** — the learner accumulates undeniable, self-attributed successes at the right grain of difficulty — including *strategy selection* and *self-accounted why*, not only blocked accuracy.
3. **Identity re-storying** — the person revises the narrative “I am not a math person” into a workable self-story compatible with struggle — via expectancy × value × cost (Eccles), not posters.

Explanations alone rarely produce (1)–(3). Fluent AI explanations can *accelerate wrong confidence* (Bastani unguarded GPT; Part XXXIII). Explanations help (2) only *after* (1) and when the student still does the decisive generative work (Parts XXVI, XL).

I.2 Verdict on the scarcity thesis

FOUNDER BELIEF under audit: AI commoditizes knowledge; scarcity shifts to motivation, confidence, identity, trust, curiosity, persistence, emotional safety, meaning.

Red Team verdict (updated v1.19):

Claim	Status	Notes
Explanations are getting cheaper/faster	FACT (directionally)	Generative AI + open content lowers marginal cost of explanation
“Tutoring is free”	FALSE as stated	Attention, accountability, diagnosis, calibrated trust remain scarce
Content is free	MOSTLY FACT for commodity content	Differentiation = curation, sequencing, trust, outcomes
Motivation/identity become scarce	HYPOTHESIS	Product differentiator; not proven sole moat — instrument $E \times V \times \text{Cost}$ separately (XXXVII)
Emotional safety is first-order	HYPOTHESIS with supportive evidence	Ashcraft WM tax; segment, don’t romanticize softness
Unguarded AI chat improves solo learning	KILLED	Bastani: concurrent $\uparrow$ then solo exam $\downarrow$; guards neutralize harm $\neq$ prove gain
Bastani “GPT Tutor” $\equiv$ MindCraft ontology wrap	KILLED	Category error (XXXIII RT)
Streaks / blocked fluency = mastery	KILLED as North Star	HID + XXXVIII/XXXIX; measure delayed mix + transfer
Harder-always / anxiety-blind “desirable difficulty”	KILLED	SAFE-DD (XLI); equip→destake→dose before Bjork dose

Ranks / XP / grit meters as identity engines	KILLED	SAFE-COMPARE / REWARD / RESILIENCE (XLII–XLV)
Story immersion $\equiv$ far transfer; mistake posters $\equiv$ learning	KILLED	SAFE-TRANSFER / MISCON (XLVIII–XLIX)
Raise confidence / Belief Score™ / CA→grades	KILLED	SAFE-CALIB (L); Foster: appropriate confidence, not inflate
Hours / forums / ritual→score / domain cosplay	KILLED	SAFE-DP / CoP / RITUAL (LI–LIII); borrow method not costume
Tip-flood / engine-dump / flooding $\equiv$ pedagogy	KILLED	SAFE-AAR / FILM / MUSIC / ANNOT / EXPOSE (LIV–LVIII)
Likert $\equiv$ WTP / survey\$ $\equiv$ revenue / ACT-point packages	KILLED	SAFE-WTP (LIX); CBC + bias controls; message $\neq$ package
Never-hallucinates / apology $\equiv$ repair / Identity Score™ / 8-week math person	KILLED	SAFE-REPAIR / IDMEASURE / LONGID (LX–LXII)
Day-one ACT flood / EL $\equiv$ ability / pink track / grit-as-character	KILLED	SAFE-EXAM / ELL / GENDER / STRUCTURE (LXIII–LXVI)
Graph file / RAG chat $\equiv$ diagnosis moat	KILLED	SAFE-ONTOLOGY (LXVII); moat = state $\times$ ontology $\times$ FEI $\times$ solo transfer
Warm human / hours / got-it $\equiv$ tutor quality	KILLED	SAFE-HITL (LXVIII); Map-briefed prompt $>$ pour + FEI QA
Emotion AI / Anxiety Score™ / FERPA $\equiv$ ethics	KILLED	SAFE-PRIVACY (LXIX); ERT ban; compliance floor $\neq$ doctrine
Page-count science / Bayesian theater / cite-wash	KILLED	SAFE-LABMETA (LXX); synthesizer must demote
Ivy/expert résumé / peer Discord $\equiv$ tutoring quality	KILLED	SAFE-TUTORGRAIN (LXXI); trained near-peer + hire-by-fidelity
LLM items $\equiv$ shipped bank / item-count hero	KILLED	SAFE-GENQ (LXXII); verify-before-ship; drop-rate gate
FERPA badge / click-wrap / biometric school SKU $\equiv$ GTM	KILLED	SAFE-PROCURE (LXXIII); privacy review = sales gate
Expanding SRS / perfect-interval AI / streak-as-spacing	KILLED	SAFE-SCHED (LXXIV); horizon-matched ISI $\times$ RI; equal-ish lags
Tutor headcount / tenure / wellness theater $\equiv$ quality	KILLED	SAFE-WORKFORCE (LXXV); fidelity@tenure + coach-capacity hire gate

Mastery fireworks / Ebbinghaus Score™ ≡ durable competence	KILLED	SAFE-FORGET (LXXXVI); aged evidence + delayed probes
SM-2/FSRS / black-box perfect-interval AI ≡ math identity	KILLED	SAFE-ADAPT (LXXVII); calendar-first; banded inspectable adapt
Cram / bootcamp / +points packages ≡ durable identity	KILLED	SAFE-DURABLE (LXXVIII); dual-rail GTM; no point guarantees
Instant-feedback AI / always-delay ≡ learning science	KILLED	SAFE-FBTIME (LXXIX); mode-conditional clocks; focus>timing
Bridge-count / re-explain ≡ connection mastery	KILLED	SAFE-BRIDGE (LXXX); remediate the join; measure join transfer
Day-one mastery fireworks / confidence ≡ mastery	KILLED	SAFE-COLD (LXXXI); humble prior + labeled seed + probes
Format-count / every-format / Format Personality™	KILLED	SAFE-FORMAT (LXXXII); conversion evidence; structural tags only
Tonight-% / streak / shame-rank parent portal ≡ trust	KILLED	SAFE-PDASH (LXXXIII); proof-age hero; autonomy>control
Talk-% / silence theater / got-it ≡ Socratic quality	KILLED	SAFE-TALK (LXXXIV); construction > airtime; prompt density
Immersion / lore / Story Engagement ≡ learning	KILLED	SAFE-STORYLOAD (LXXXV); CLT-budgeted structure wrap
User-count / star-wall / vague praise ≡ learning proof	KILLED	SAFE-PROOF (LXXXVI); dated solo competence artifacts
Item-count / complete-ACT / unsynced GENQ ≡ coverage	KILLED	SAFE-COVER (LXXXVII); matrix + gaps + use-tier
Always-full-worked / never-fade / solution-first ≡ learning	KILLED	SAFE-FADE (LXXXVIII); completion ladder + solo proof
Unlimited-hints / help-NPS / avoidance-as-grit ≡ learning	KILLED	SAFE-HELP (LXXXIX); instrumental contingent help + solo
DAU/streak/XP / FEI Score™ / thumbs-up ≡ learning NS	KILLED	SAFE-INSTRUMENT (XC); FEI four + XXI.4 co-gates
Unlimited free hard peeks / Hint Score™ / XP-for-hints ≡ learning	KILLED	SAFE-HINT (XCI); soft-first contingent + construction cost

Longer $\Rightarrow$ better / Explanation Score TM / monologue $\Rightarrow$ SE $\equiv$ learning	KILLED	SAFE-EXPLAIN (XCII); short principle-first + SE-before-wrap
All-struggle $\Rightarrow$ productive / Struggle Score TM / blank-time / never-stuck $\equiv$ learning	KILLED	SAFE-RETRIEVE (XCIII); mode-contingent stall UX
Every-miss $\Rightarrow$ PF / fail-first-always / PF Score TM / discovery-without-consolidation $\equiv$ learning	KILLED	SAFE-PF (XCIV); fidelity-gated generate $\rightarrow$ consolidate
Guild cosplay / modeling $\Rightarrow$ lecture / never-fade / Apprenticeship Score TM $\equiv$ learning	KILLED	SAFE-APPRENTICE (XCV); phase-fidelity model $\rightarrow$ coach $\rightarrow$ fade
Focus Score TM / phone-shame / mid-attempt DAU pushes / multitasking-as-talent $\equiv$ learning	KILLED	SAFE-ATTN (XCVI); attempt-window protection + tech breaks
Dual Coding Score TM / always-add-picture / decoration $\Rightarrow$ dual-coding / visual-learner meshing / animation-default	KILLED	SAFE-DUAL (XCVII); contiguous structural figures + format-hop proof
Monitoring Score TM / metacognition-without-control / easy-only $\Rightarrow$ calibration / one-scan calibration	KILLED	SAFE-MONITOR (XCVIII); named judgments $\rightarrow$ miss-class control
AI-replaces-tutors / tutoring-is-free / unlimited-human / hours NS / Session ROI Score TM	KILLED	SAFE-ROI (XCIX); Map minute budgets + FEI per tutor-minute
Universal E0 $\rightarrow$ E3 destiny / fixed $\Rightarrow$ adaptive / confidence $\Rightarrow$ fade stage / Example Timing Score TM	KILLED	SAFE-EXPTIME (C); grain-local adaptive example $\leftrightarrow$ solve timing
SRL Score TM / reflection-streak / survey $\Rightarrow$ event / streak $\Rightarrow$ SRL / forethought-without-control	KILLED	SAFE-SRL (CI); phase $\rightarrow$ control cycles + FEI proof
Generative Learning Score TM / AI $\Rightarrow$ generation / teach-for-XP / eight-menu GLA bomb	KILLED	SAFE-GENERATE (CII); SOI-forcing sparse GLAs + transfer proof
Psychological Safety Score TM / poster safe-space / Bermuda-triangle / praise-fog $\Rightarrow$ climate	KILLED	SAFE-ERRCLIMATE (CIII); stay-diagnose + adaptive action under FEI

Homework Completion Score™ / tonight-% / Chegg-dump / assisted⇒ready / purity-ban help	KILLED	SAFE-HWHELP (CIV); dual-rail instrumental help + solo proof
Embodiment Score™ / styles meshing / always-gesture / action⇒gesture / pose-camera AI	KILLED	SAFE-GESTURE (CV); strategy-tied optional motion + transfer proof
Peer Explanation Score™ / teach-for-XP / telling⇒quality / Discord⇒building / AI⇒peer teaching	KILLED	SAFE-PEERX (CVI); building-gated teach-back + solo proof
Curiosity Score™ / cliffhanger-without-commit / postdiction / instant-reveal entertainment	KILLED	SAFE-PREDICT (CVII); commit→surprise→encode repair
Analogy Score™ / Mapping Minutes / mention⇒pedagogy / surface-story⇒structure / AI likeness dump	KILLED	SAFE-ANALOG (CVIII); co-present structure-map + student alignment
Writing Score™ / Journal Minutes / freewrite⇒learning / longer⇒deeper / AI rewrite⇒WTL	KILLED	SAFE-WRITE (CIX); short prompted cognitive+metacognitive protocols + FEI
Spatial Score™ / Mental Rotation Minutes / brain-training⇒identity / boys-better spatial / Spatial Personality™	KILLED	SAFE-SPATIAL (CX); FormatId-embedded hops + conditional transfer
Formative Score™ / Assessment Minutes / exit-ticket⇒learning / 0.4–0.7 AfL ads	KILLED	SAFE-FA (CXI); evidence→adjust cycles + FEI
Feedback Literacy Score™ / Feedback Minutes / delivery⇒uptake / thumbs⇒used / longer⇒literate	KILLED	SAFE-FBLIT (CXII); uptake loop + FEI
Collaboration Score™ / Group Minutes / Discord⇒CSCL / group-correct⇒ready	KILLED	SAFE-COLLAB (CXIII); solo gate + accountability
Desirable Difficulty Score™ / Challenge Minutes / always-ease / always-max / cram⇒ready	KILLED	SAFE-DDSCHED (CXIV); horizon hold/ease + stakes≠DD
Representation Score™ / Diversity Minutes / costume⇒belonging / match⇒quality / empty likeness praise	KILLED	SAFE-ROLEMODEL (CXV); fidelity-first + competence-visible near-peers

Capture Score™ / Photo Minutes / OCR≡mastery / always-camera / photo→answer hero	KILLED	SAFE-OFFLINE (CXVI); attempt→capture→confirm + no Photomath pipe
Forever-E0 anxious track / day-one fade flood / Anxiety·Calm Score™ fade gates	KILLED	SAFE-FADEX (CXVII); protect-then-fade; no forever dump
Miss Streak / soft-wrong shame push / live wrong-answer stalk / empty mindset parent letter	KILLED	SAFE-SOFTMSG (CXVIII); grain+failure-enhancing+one CTA
Always-on student Map / Map Minutes·Map Score™ / mid-attempt full-graph dump	KILLED	SAFE-MAPORCH (CXIX); brief→grain-slice→adjust
Burnout Score™ / Wellness Minutes / wellness-primary under caseload spike	KILLED	SAFE-BURNWARN (CXX); fidelity/load→org redesign
Fluency≡pedagogy / Explanation·Rubric Score™ / LLM-judge-alone ship / soft-pass wrong math	KILLED	SAFE-EXPLAINQA (CXXI); principle-short gates + hard-fail math
Soft-wrong≡half-credit / Partial Credit Score™ / all-wrongs-equal / gradebook cosplay	KILLED	SAFE-TAXON (CXXII); type before points
Hours-booked crunch NS / Load Shed Score™ / infinite-availability / pour≡exam support	KILLED	SAFE-LOADSHED (CXXIII); acuity shed under crunch
Distractor Score™ / always-four≡better / beauty traps / unmapped “diagnostic” MC	KILLED	SAFE-DISTRACTOR (CXXIV); misconception-mapped options
Identity Score™ / Confidence Personality™ / equal-weight recovery / inflate-confidence NS	KILLED	SAFE-CONFMISS (CXXV); CRI fork + hypercorrection priority
Skip-as-default / lecture-as-debrief / Microdebrief Score™ / +25% ACT debrief ads	KILLED	SAFE-MICRODEB (CXXVI); capped student-first micro-AAR with traces
24/7 default / hours-hero crunch / fear-parent emails / Availability Score™ / silent shed	KILLED	SAFE-CRUNCHMSG (CXXVII); shed-honesty + thinking-minutes CBC
FormatTrap Score™ / seductive stem art≡FormatId / beauty visual traps / unmapped visual “diagnostic”	KILLED	SAFE-FORMTRAP (CXXVIII); trap-mapped FormatId distractors

Brief Score™ / “no time for brief” / pour-open / late-join shame / I-PASS cosplay ACT ads	KILLED	SAFE-MICROBRIEF (CXXIX); capped shared-model open + student synthesis
Chegg-hero crunch CTA / fear-countdown dump / Dual-CTA Score™ / assisted⇒ready	KILLED	SAFE-CRUNCHCTA (CXXX); dual labeled rails + follow-on Practice
Fluency SVG⇒keyed figure / judge-alone figure ship / Figure Key Score™ / unverified FormatId diagnostic	KILLED	SAFE-FIGKEY (CXXXI); structural figure verify + use gating
Idle spinner lobby / pour-open wait / Wait Minutes·Lobby Score™ / entertainment-primary fill	KILLED	SAFE-ASYNCBRIEF (CXXXII); session-related Set card + verify handoff
Dump-and-quit Solver close / maze handoff / Transfer Score™ / assisted⇒ready	KILLED	SAFE-XFERGRAIN (CXXXIII); one-tap edge-matched Practice grain + solo gate
Tagged FormatId⇒C4 / fluency eligibility / coverage-padding / HideCorrectness Score™	KILLED	SAFE-HCELIG (CXXXIV); seals-first C4 eligibility; honest holes
Soon-default / optimistic ETA slip / apology-primary / ETA Score™ / shame feed	KILLED	SAFE-KNOWNETA (CXXXV); known-finite ETA + refresh honesty
Universal hard-gate after every hint / Gate Score™ / hint-tax shame	KILLED	SAFE-SOFTGATE (CXXXVI); soft-require default; conditional hard-gate + avoidance falsifiers
Unified sealed bank ⇒ Practice+C4 / silent dual-write / auto-promote / SplitPool Score™	KILLED	SAFE-SPLITPOOL (CXXXVII); graded pools; ticketed promote
Symbolic/Practice pad into C4 / equal-quota theater / Composition Score™ / diagnose-every-format ads	KILLED	SAFE-SCANCOMP (CXXXVIII); blueprint + sealed stock; labeled holes
Hide-holes for conversion / child-deficit shame / complete-visual ads / Hole Trust Score™	KILLED	SAFE-HOLETRUST (CXXXIX); label hole + one fill CTA
Play-count/thumbs auto-promote / silent hcelig_pass / Promote Score™	KILLED	SAFE-PROMOTE (CXL); ticketed checklist; demote faster than promote

Immediate same-stem C4 after reveal / forever·global cool-down / Cool-Down Score™ / pad under cool-down	KILLED	SAFE-COOLDOWN (CXLI); per-student cool-down + isomorphic substitutes
Parent-only hole honesty / fake greens / ability-shame / empty mindset poster / Student Hole Score™	KILLED	SAFE-STUDHOLE (CXLII); parity honesty + one Practice CTA
Coming-soon / play-count fill bars / date soft-pass / Filleta Score™	KILLED	SAFE-FILLETA (CXLIII); capacity-grounded verify bands
Always-dual / never-dual figured C4 / Dual Rater Score™ / mammography-cosplay guarantees	KILLED	SAFE-DUALRATE (CXLIV); risk-triggered dual; seals before second rater
Item_id-only cool-down / concept-wide ban / embedding-only Family Hash Score™	KILLED	SAFE-FAMHASH (CXLV); sealed family compromise set
Anxiety·Calm·Exit Emotion Score™ / consoling ability comfort / ego Completeness % / emotion-AI exit	KILLED	SAFE-EXITEMO (CXLVI); destake exit + one Practice CTA
Dual-omitted FILLETA / ETA-skipped dual / always-dual without FTE / Dual Rater·Capacity Score™	KILLED	SAFE-CAPPLAN (CXLVII); dual minutes in fill estimator
Forever-dual clones / silent inherit / Template Trust Score™ / FERF≡C4-ready	KILLED	SAFE-TMPTRUST (CXLVIII); inherit intensity not seals; sample+revoke
Kappa % NS / pooled repair vanity / Landis-band marketing / Dual Rater Score™ IRR	KILLED	SAFE-REPAIRSTRAT (CXLIX); path-stratified repair first
Dual Rater Score™ / Surge Hire % / Contract Minutes / unsealed freelancer dual / hospital-surge cosplay	KILLED	SAFE-SURGEANAL (CL); sealed make-or-buy dual portfolio

Implication: Do not bet on “better explanations,” graph-file cosplay, warm-tutor theater, emotion cameras, page-count science, Ivy hire theater, unverified AI banks, FERPA-badge GTM, expanding-SRS mystique, headcount vanity, mastery fireworks, FSRS brand, cram packages, instant-feedback theater, bridge-count ads, day-one greens, format-count cosplay, parent surveillance portals, Talk Ratio theater, immersive lore-as-pedagogy, star-wall / user-count learning proof, item-count / complete-ACT coverage theater, always-full-worked / unlimited-solutions Solver theater, unlimited-hints / never-stuck answer-dump theater, DAU/streak/XP / FEI Score™ dashboard theater, unlimited free hard peeks / Hint Score™ theater, longer≡better / unlimited thorough AI essay theater, struggle-theater / blank-time / always-wait / instant-dump stall UX, fail-first / every-miss-as-PF / discovery-without-consolidation

cosplay, guild / master-tutor / modeling-as-lecture / never-fade apprenticeship costume, Focus Score™ / phone-shame / mid-attempt DAU pushes, Dual Coding Score™ / always-add-picture / visual-learner meshing, Monitoring Score™ / one-scan calibration, Session ROI Score™ / tutoring-is-free, Example Timing Score™ / universal E0→E3, SRL Score™ / reflection-streak, Generative Learning Score™ / AI≡generation, Psychological Safety Score™ / Bermuda-triangle, Homework Completion Score™ / Chegg-dump, Embodiment Score™ / pose-camera, Peer Explanation Score™ / Discord≡building, Curiosity Score™ / cliffhanger-without-commit, Analogy Score™ / surface-story≡structure, Writing Score™ / Journal Minutes, Spatial Score™ / Mental Rotation Minutes, Formative Score™ / 0.4–0.7 AfL ads, Feedback Literacy Score™ / delivery≡uptake, Collaboration Score™ / Discord≡CSCL / group-correct≡ready, Desirable Difficulty Score™ / Challenge Minutes / always-ease-before-ACT / always-max-DD-until-bell / cram-week≡ready calendar theater, or Capture Score™ / Photo Minutes / OCR≡mastery / Photomath-pipe notebook AI, forever-E0 / Anxiety·Calm Score™ fade gates, Miss Streak / soft-wrong shame push parent theater, always-on Map / Map Minutes mid-attempt dumps, Burnout Score™ / wellness-primary under caseload, fluency≡pedagogy / LLM-judge-alone wrap ship, Partial Credit Score™ / all-wrongs-equal gradebook cosplay, hours-booked crunch / Load Shed Score™ / pour≡exam support, Distractor Score™ / always-four / beauty-trap MC, or Identity Score™ / Confidence Personality™ / equal-weight recovery / inflate-confidence theater, Microdebrief Score™ / skip-as-default lecture-close, 24/7·Availability Score™ / fear-parent crunch emails, FormatTrap Score™ / beauty visual traps, Brief Score™ / pour-open / late-join shame, Chegg-hero Dual-CTA Score™, Figure Key Score™ / fluency SVG pass, Wait Minutes·Lobby Score™ / entertainment lobby, or Transfer Score™ / dump-and-quit / maze handoff, HideCorrectness Score™ / tagged≡C4 eligibility, ETA Score™ / soon-wait theater, Gate Score™ / universal hard-gate, SplitPool Score™ / silent dual-use, Composition Score™ / pad-scan theater, Hole Trust Score™ / complete-visual ads under holes, Promote Score™ / play-count flip, Cool-Down Score™ / immediate same-stem C4, or Student Hole Score™ / fake greens / parent-only honesty, FilleETA Score™ / play-count fill bars / date soft-pass, Dual Rater Score™ / always-dual-never-dual / mammography cosplay, Family Hash Score™ / item_id-only·concept-wide cool-down, Exit Emotion Score™ / emotion-AI exit comfort, Capacity Score™ / ETA-skipped dual, Template Trust Score™ / silent inherit / forever-dual clones, Kappa % / pooled IRR vanity, or Surge Hire % / unsealed freelancer dual / hospital-surge cosplay — subordinate UX under FEI + surviving doctrine. Bet on **FEI + pedagogy wrap + SAFE-* stack**: fear→evidence→identity, with inspectable diagnosis and surviving doctrine densifiers — not calendar theater or Score™ costumes.

I.3 What to optimize (North Star debate)

Metric	Predicts short-term engagement?	Predicts durable math identity?	Risk
Correct answers (esp. blocked)	Medium	Low–medium	Fluency illusion; gaming
Time spent / chat tokens	High	Low	Engagement / AI-monologue theater
Streaks / DAU	High	Uncertain	Habit ≠ identity (HID)
Confidence (self-report)	Medium	Medium	Cheap talk; overtrust after fluent AI
Challenge-seeking (challenge_accept + why)	Medium	High	Must code mastery vs appearance/normative motive (XXXVIII)
Voluntary return after failure (retry_120s)	Medium	High	Needs coach content discipline (SE prompts)

Delayed mixed / transfer_pass / solo_transfer_pass	Medium	High	Anti-fake-mastery; anti-AI-crutch
“I am a math person” + behavior	Medium	Highest target	Slow, noisy; L3+ claim only

Working North Star (HYPOTHESIS):

Challenge-seeking under safety with transfer — the student chooses a harder problem for a *mastery* reason, returns after a miss, can pick the strategy on a mixed delayed set, and attributes success to strategy — without needing the model to finish the thought.

I.4 Surviving doctrine stack (Synthesizer merge through CL)

Duplicate frameworks collapsed. Deep dives own citations; this table is **company law until killed**.

Layer	Surviving rule	Source parts	Commercial use
FEI / SAFE-CRAFT	Safety → attempt → informative miss → earned win → attribution → transfer	XVI, VIII	Session demo in marketing; not feature matrix
Claim ladder L0–L4	Do not sell identity (L3) from session A/Bs (L1)	XXXIV	Parent/copy hygiene; pre-reg estimands
PWC (pedagogy wrap constraint)	Guarded AI + ontology spine; success = solo_transfer_pass ≥ no-AI	XXXIII, XXXV	Ban “we’re Bastani’s Tutor”; ban unguarded Solver hero
E×V×Cost	Instrument expectancy, value, cost separately; student-authored utility > staff utility dump	XXXVII	Kill single “motivation” KPI
TARGET / mastery climate	Default-ban appearance UX + streak-as-North-Star; soft-wrong = mastery evaluation structure	XXXVIII	CTA copy audit; appearance climate kill
SAFE-DD	Equip → destake → dose → frame → measure; desirable difficulty only after WM/threat capacity	XLI	Ban hardness-as-virtue; gate IL/SE by anxiety segment

SAFE-COMPARE	Criterion feedback default; peer ranks opt-in; no forced infinite named leaderboards for Maya	XLII	Leaderboard / league kill as growth engine
SAFE-HABIT	Cue practice not the counter; miss ≠ identity failure; habit metrics subordinate to FEI	XLIII	Streak colonization kill; if-then onboarding OK
SAFE-REWARD	No expected tokens for mere engagement; informational feedback > inducement; transfer-gated unlocks	XLIV	XP-as-love / coin-for-open kill
SAFE-RESILIENCE	Growth-zone challenge + support recruitment; ban grit NS, danger-zone-as-DD, Resilience Score™	XLV	Grit-poster / character-meter kill
SAFE-EXPECTANCY	Task-only tutor briefs; ban trait/weakness labels; CIOF mediation; cold-start scripts	XLVI	Pygmalion marketing kill; tutor ethics
SAFE-TIMING	Encode → sleep → transfer claim; never reward midnight volume; stress/restriction-aware load; pressure ≠ DD	XLVII	All-nighter / sleep-app brand kill
SAFE-TRANSFER	Name the hop; VT vary-on-purpose; hug exam + bridge abstraction; story ≠ transfer; blocked ≠ ready	XLVIII	Narrative-theater / blocked-accuracy vanity kill
SAFE-MISCON	Elicit → classify → destake → springboard → route; soft-wrong ≠ empty kindness; confidence-tiered feedback	XLIX	Mistake-poster / red-X-as-rigor / equal-wrong graph kill

SAFE-CALIB	Item-level confidence → miss class → tiered feedback; raise <i>appropriate</i> confidence, not Belief Score™	L	Hide-correctness diagnostic; ban CA→grades marketing
SAFE-DP	Diagnose → isolate → goal → stretch → attempt → method feedback → revise → prove → stop	LI	Ban 10k-hours / minutes / coverage theater as DP
SAFE-CoP	Name the practice; peripheral = real work; old-timer required; trajectory > community tab	LII	Ban Discord/streak ≡ LPP; “join the practice” not “join the community”
SAFE-RITUAL	Light, optional, practice-pointing marks; doctrinal default; anti-totalism; no ritual→score	LIII	Ban streak-as-liturgy; family/order milieu copy
SAFE-AAR	Brief → attempt with traces → four-question student-first close → one improve-how	LIV	Ban bootcamp brand; +25% ACT from debrief metas
SAFE-FILM	Objective clip + ≤3 cues + learner-timed review + self↻model; next-attempt commitment	LV	Ban tip-flood film room; sports-academy cosplay
SAFE-MUSIC	Scales (goalful ingredients) → chosen repertoire → transfer-as-recital proof	LVI	Ban virtuoso cosplay; hours KPI; bare-pass theater
SAFE-ANNOT	Annotate-before-reveal ; principle tags; forced critical restudy; classify→drill; annotate wins	LVII	Ban GM cosplay; engine-dump≡coaching; rating NS
SAFE-EXPOSE	Situation hierarchy; destaked approach; expectancy violation; fade safety behaviors; multi-context	LVIII	Ban therapy branding; flooding; Calm Score™; one-session cure

SAFE-WTP	Parent CBC with price+neither; ship-mapped attributes; message≠package; cash pilot before stickers	LIX	Ban Likert≡WTP; survey\$≡revenue; ACT-point SKUs
SAFE-REPAIR	Detect → classify → own explanatorily → Map-checked correct → destaked re-attempt → triggered escalate	LX	Ban never-hallucinates; rote apology≡repair; infinite free escalate
SAFE-IDMEASURE	Validated Cribbs-factor battery + FEI behavioral co-primary + demand-aware admin; narrative probe light	LXI	Ban Identity Score™; two-item≡FEI; splash Likert; empty “math person” praise
SAFE-LONGID	Stage claims: 8w process / 26w triangulated movement / 52w persistence; fade-aware packages	LXII	Ban math-person-in-8-weeks; end-of-program Likert≡durable
SAFE-EXAM	Dual rail MA+TA; destake-to-learn / dose-to-prove; readiness = delayed mixed under stakes	LXIII	Ban day-one ACT flood; blocked≡exam-ready; ACT guarantees; Calm Score™ exam NS
SAFE-ELL	Separate language vs math load; linguistic modify without math dilute; format triangulation before verdict	LXIV	Ban math-is-language-free; forever procedure ELL track; translate≡identity
SAFE-GENDER	Parity hygiene; cue hygiene; belonging via practice recognition; no ST→score / pink cosplay	LXV	Ban innate-gap premise; threat-theater ACT; girl-SKU / pink track
SAFE-STRUCTURE	Structure/scaffolds before grit slogans; bandwidth hygiene; behavior≠trait; no character-blame equity	LXVI	Ban grit-NS / Grit Score™ / poor-kids-need-grit / streak-as-character

SAFE-ONTOLOGY	Diagnosis-before-dialo gue; inspectable state; LLM bookends; moat = state×ontology×FEI×tr ansfer not graph file	LXVII	Ban node-count moat / RAG≡KT / fluency-as-diagnosis / ITS-2σ / smarter-than-ChatGPT
SAFE-HITL	Map-briefed tutors; prompt>pour; answer evidence not “got it?”; complementarity; coaching cycles; FEI QA; packetized escalate	LXVIII	Ban warm≡quality / hours NS / Discord≡QA / Ivy≡playbook / explain-first / humans-skip-Map
SAFE-PRIVACY	Self-authored short-TTL affect → pedagogy modifiers only; ERT ban; contextual integrity; refusal-safe	LXIX	Ban face/voice emotion AI / Anxiety Score™ / FERPA≡ethics / mood gates / sell-affect
SAFE-LABMETA	Label → check → kill/wound/survive → demote copy; small FEI hard core; synthesizer must delete; page count vanity	LXX	Ban page-count NS / Bayesian theater / cite-wash / belief-by-repetition / science-backed ads without falsifiers
SAFE-TUTORGRAI N	Default trained near-peer; structure>pairing; no 2σ; Map>credential diagnosis; suppress knowledge-telling; expert=escalate	LXXI	Ban Ivy/PhD≡FEI / Discord-peer≡tutoring / explainers hire bar / expertise≡diagnosis / tutor-learning NS
SAFE-GENQ	Verify-before-ship; key fails hard-fail; Kane IUA for bank uses; drop-rate gate; coverage≠NS; LLM≠classical AIG without constraints	LXXII	Ban LLM-items≡shipped / item-count hero / unverified diagnostic keys / fluency≡keyed correctness / scale-at-30%-drop
SAFE-PROCURE	Privacy review = GTM gate; NDPA-ready DPA; Exhibit A honesty; free=paid approval; no bio / no pupil-data ads	LXXIII	Ban FERPA-badge≡ready / click-wrap school path / biometric school SKU / marketing-first without packet

SAFE-SCHED	Horizon-matched ISI×RI; delayed first return; equal-ish lags for long RI; generation Returns; exam dual rail; no streak-as-spacing	LXXIV	Ban expanding-SRS hero / perfect-interval AI / cram≡ready / overlearn-tonight / restudy≡review / shuffle≡spacing
SAFE-WORKFORCE	Fidelity-over-tenure; continuous coaching; coach-capacity gates hiring; burnout→load redesign; named drift channels; retention≠NS	LXXV	Ban headcount/hours/r etention-% NS / tenure≡quality / onboarding≡coaching / wellness-theater burnout / Burnout Score™ / Ivy-bench pipeline
SAFE-FORGET	Age evidence; JOLs≠retention; delayed probes as truth; returns as competence protection; no shame-curve hero	LXXVI	Ban mastery-fireworks ≡durable / Ebbinghaus Score™ / streak-as-memory / confidence≡retention / cured-forgetting ads / retention% guarantees
SAFE-ADAPT	Calendar-first; adapt only inside RI bands with due+why; selection-among-due > opaque rewrite; flashcard lineage≠FEI	LXXVII	Ban SM-2/FSRS/Anki hero / black-box perfect-interval AI / flashcard-MAE≡ACT / engagement≡learning / silent ease-hell
SAFE-DURABLE	Dual-rail GTM: durable learn + late prove; cram-intercept not cosplay; legible MoC for WTP; no point guarantees	LXXVIII	Ban cram≡ready / ACT-point packages / overlearn-tonight / shame-parent GTM / streak-as-durable
SAFE-FBTIME	Mode-conditional timing: diagnostic hide / learn micro-delay+SE / prove KR; focus>timing; prior-knowledge switch	LXXIX	Ban instant-feedback AI hero / always-delay≡science / feedback≡learning / green-check NS / ego-toast≡FI
SAFE-BRIDGE	Connection-first diagnosis when endpoints non-naive; remediate relation with co-present supports; measure join transfer	LXXX	Ban bridge-count moat / re-explain≡bridge / coverage≡connection / we-connect-everything / shame-edge Map

SAFE-COLD	Humble prior → labeled seed → hide-correctness probes; Map chrome = evidence age; honest mapping ≠ fake knowing	LXXXI	Ban day-one mastery fireworks / confidence≡mastery / skip-to-green / prior-as-biography / placement-belt cold start
SAFE-FORMAT	Format as first-class evidence; conversion as object; support before stack; structural FormatId only	LXXXII	Ban format-count moat / every-format ads / decoration≡coverage / concept≡format-flexibl e / Format Personality™
SAFE-PDASH	Beliefs before chrome; proof-age hero; one improvement CTA; autonomy>control; ranks/affect off	LXXXIII	Ban tonight-%/streak parent NS / fireworks portal / shame ranks / live wrong-answer stalk / gradebook cosplay
SAFE-TALK	Construction > airtime; prompt density + constructive share + wait bands; pour-streak QA; human=AI schema	LXXXIV	Ban Talk Ratio Score™ / talk-% NS / 50/50 dogma / silence≡Socratic / got-it≡construction / AI monologue≡dialogue
SAFE-STORYLOAD	CLT budgets narrative; structure-encode or cut; coherence > immersion; novices get less lore; skip-story; story≠transfer	LXXXV	Ban immersion/lore NS / Story Engagement Score™ / seductive trivia≡pedagogy / forced lore walls / fandom-skin hero
SAFE-PROOF	Competence > crowd; enactive dated solo artifacts; provincial similarity; student-owned thin portfolio; two-sided honesty	LXXXVI	Ban user-count/star-wall learning proof / vague praise / Identity Score™ splash / +points case studies / forced public portfolio NS
SAFE-COVER	Matrix > total; gaps as features; use-tier labels; examTag honesty; GENQ never counted unverified; prior substantiation via audit	LXXXVII	Ban item-count hero / complete-ACT puffery / GCSE-as-ACT / unsynced GENQ counts / ontology-nodes≡bank

SAFE-FADE	Stage not costume; completion bridge; backward fade first; expertise-aware up/down; attempt grain; proof = solo transfer	LXXXVIII	Ban always-full-worked / never-fade / solution-first / view $\equiv$ mastery / Fade Score™ / unlimited-solutions hero
SAFE-HELP	Instrumental > executive; abuse $\cup$ avoidance fail; contingent stage; SE before bottom-out; metacognition $\neq$ score; dignity telemetry	LXXXIX	Ban unlimited-hints / help-NPS NS / avoidance-as-grit / Help Score™ / parent hint-stalk / never-stuck dump
SAFE-INSTRUMENT	Ship the four FEI events; NSM stack + stage OMTM; XXI.4 co-gates; anti-Goodhart; anti-gaming companions; Kane IUA for uses	XC	Ban DAU/streak/XP NS / FEI Score™ / single-KPI tutor pay / thumbs-up $\equiv$ learning / identity ads without IUA
SAFE-HINT	Soft before hard; contingency up/down; peek cost=construction; fade $\times$ hint one machine; no XP/Hint Score™; inspectable stage	XCI	Ban unlimited free hard peeks / Hint Score™ / XP-for-hints / shame-timer cost / black-box perfect-hint AI / fade+free dump
SAFE-EXPLAIN	Short principle-first before monologue; germane $\neq$ more text; SE before wrap; length fades with expertise; coherence>charm; no token NS	XCII	Ban longer $\equiv$ better / Explanation Score™ / AI monologue $\equiv$ SE / unlimited thorough AI essays / fixed long wrap / seductive coach fluff
SAFE-RETRIEVE	Name stall mode; attempt before heavy help; TOT $\rightarrow$ cue; freeze $\rightarrow$ bounded soft; progress $\neq$ blank-time; no Struggle Score™	XCIII	Ban all-struggle $\equiv$ productive / blank-time NS / always-wait / instant dump / never-stuck hero / errors-always-scar

SAFE-PF	PF = generate→consolidate on student RSMs; fidelity+prior gates; guided default when thin/exam/freeze; no PF Score™	XCIV	Ban every-miss=PF / fail-first-always / gener ation-without-consolid ation / discovery-as-identity / ACT-from-invent
SAFE-APPRENTICE	Model control/heuristics → coach → scaffold→fade; articulate; shared ca_phase; proof = solo+articulation	XCV	Ban guild cosplay / modeling=lecture / never-fade / talk-%=CA / Apprenticeship Score™ / CA ads without phases
SAFE-ATTN	Protect mid-attempt windows; residue-aware park/resume; opt-in focus + tech breaks; presence hygiene without shame	XCVI	Ban Focus Score™ / phone-shame / mid-attempt DAU pushes / multitasking-as-talent / guaranteed ACT from focus
SAFE-DUAL	Structural dual coding; contiguous labels; coherence weed; static annotated default; conversion stays the object	XCVII	Ban Dual Coding Score™ / always-add-picture / decoration=learning / styles meshing / animation-default / split-legend density
SAFE-MONITOR	Named judgments; C4 destake; confidence×outcome classes; control wiring mandatory; densifies SAFE-CALIB/COLD	XCVIII	Ban Monitoring Score™ / metacognitio n-without-control / easy-only=calibration / one-scan calibration / raise-metacognition identity
SAFE-ROI	Tutor minutes scarce; Map minute budgets; constrained AI for reps; humans for joins/desta ke/repair/witness; FEI per tutor-minute	XCIX	Ban AI-replaces-tutors / tutoring-is-free / unlimited-human / hours NS / Session ROI Score™ / VanLehn-as-ChatGPT
SAFE-EXPTIME	Grain-local example↔solve timing; completion bridge; adaptive>fixed; rapid rechecks; densifies SAFE-FADE	C	Ban universal E0→E3 destiny / fixed=adaptive / confidence=stage / Example Timing Score™ / always-struggle-first

SAFE-SRL	Phase→control cycles; process goals; attribution→plan writeback; microanalysis is>inventories; densifies MONITOR/AAR	CI	Ban SRL Score™ / reflection-streak NS / survey≡event / streak≡SRL / forethought-without-control / ability-shame reflection
SAFE-GENERATE	SOI-forcing sparse GLAs; student produces; teach-back act; guided FormatId draw; densifies EXPLAIN/DUAL/SRL/SE	CII	Ban Generative Learning Score™ / AI≡generation / teach-for-XP / eight-menu bomb / verbatim summary theater
SAFE-ERRCLIMATE	Error climate > slogan; stay-diagnose-reattempt; destake learn / keep FEI; densifies MISCON/HELP/TALK/AAR	CIII	Ban Safety Score™ / poster safe-space / Bermuda-triangle / shame reel / praise-fog≡climate
SAFE-HWHELP	Dual-rail Help vs Practice; instrumental unstuck; assisted≠solo; help recruits proof; densifies HELP/FADE/ROI	CIV	Ban Completion Score™ / Chegg-dump / assisted≡ready / Bastani Base hero / purity-ban help
SAFE-GESTURE	Strategy-tied optional gesture/draw; gesture≠action for transfer; respect non-gesturers; densifies DUAL/GENERATE/FORMAT	CV	Ban Embodiment Score™ / styles meshing / always-gesture / pose-camera / AR body≡FEI
SAFE-PEERX	Building over telling; inferential prompts; apply after explain; capture act; grain+QA; densifies GENERATE/EXPLAIN/HELP	CVI	Ban Peer Explanation Score™ / teach-for-XP / Discord≡building / AI≡peer teaching
SAFE-PREDICT	Commit before reveal; curiosity+surprise for expectancy violations; encode repair; mode gates; densifies FBTIME/MONITOR/MISCON	CVII	Ban Curiosity Score™ / cliffhanger-without-commit / postdiction / guess-for-XP / curiosity ACT ads

SAFE-ANALOG	Structure-map joins; co-present compare; student alignment; supports required; densifies BRIDGE/TRANSFER/DUAL	CVIII	Ban Analogy Score™ / mention≡pedagogy / surface-story≡structure / AI likeness dump / always-analogy
SAFE-WRITE	Short prompted WTL; cognitive+metacognitive protocols; Notes dual-rail; AI cues not authors; densifies GENERATE/SE/SRL	CIX	Ban Writing Score™ / Journal Minutes / freewrite≡learning / longer≡deeper / AI≡WTL / always-journal
SAFE-SPATIAL	Spatial via FormatId hops; conditional transfer; destake spatial anxiety; densifies FORMAT/DUAL/GESTURE/GENDER	CX	Ban Spatial Score™ / Mental Rotation Minutes / brain-training≡identity / boys-better / Spatial Personality™
SAFE-FA	Formative cycles: evidence→adjust→close gap; exit tickets only if wired; densifies HITL/AAR/FBTIME/SRL	CXI	Ban Formative Score™ / Assessment Minutes / ticket≡learning / 0.4–0.7 ads
SAFE-FBLIT	Feedback literacy = sense+use; uptake loop; proactive recipience; densifies FBTIME/FA/EXPLAIN/HINT	CXII	Ban Feedback Literacy Score™ / Feedback Minutes / delivery≡uptake / thumbs≡used
SAFE-COLLAB	Structured collab + individual accountability + solo FEI gate; densifies PEERX/CoP/PROOF	CXIII	Ban Collaboration Score™ / Group Minutes / Discord≡CSCL / group-correct≡ready
SAFE-DDSCHED	Horizon-matched hold/ease of DD; stakes≠DD; Protect+probe final 72h; densifies DD/EXAM/SCHED/CRAM	CXIV	Ban Desirable Difficulty Score™ / Challenge Minutes / always-ease / always-max / cram≡ready
SAFE-ROLEMODE L	Fidelity-first near-peer; competence-visible attainable bios; optional preference match; densifies GENERATOR/ELL/TUTORGRAIN	CXV	Ban Representation Score™ / Diversity Minutes / costume≡belonging / match≡quality

SAFE-OFFLINE	Offline-first attempt; post-commit capture; verify before Map; densifies ATTN/HELP /HWHELP/PROOF/G ENQ	CXVI	Ban Capture Score™ / Photo Minutes / OCR≡mastery / photo→answer hero
SAFE-FADEX	Protect-then-fade under HMA; longer E0/E1 hold without forever track; peek binge = safety behavior; densifies FADE/EXPOSE/DD	CXVII	Ban forever-E0 / day-one fade flood / Anxiety·Calm Score™ fade gates / Fade Score™
SAFE-SOFTMSG	Grain + failure-enhancing frame + one CTA; autonomy>control; densifies PDASH/FBL IT/ERRCLIMATE/MI SCON	CXVIII	Ban Miss Streak / live wrong-answer stalk / empty mindset posters / Parent Correctness Score™
SAFE-MAPORCH	Brief→grain-slice→adj ust; tutor-private default; no mid-attempt full-graph; densifies HI TL/FA/ATTN/BRIDG E	CXIX	Ban always-on student Map / Map Minutes·Map Score™ / dual-screen default / ontology tour≡tutoring
SAFE-BURNWARN	Fidelity/load early-warn → org redesign; JD-R; densifies WORKFORC E/PRIVACY/ROI	CXX	Ban Burnout Score™ / Wellness Minutes / mandatory mood gate / wellness-primary under caseload
SAFE-EXPLAINQA	Principle-short rubric gates; hard-fail math/key; densifies EX PLAIN/GENQ/REPAI R	CXXI	Ban fluency≡pedagogy / Explanation·Rubric Score™ / LLM-judge-alone / soft-pass wrong math
SAFE-TAXON	Typed marks drive route+graph; soft-wrong wraps type; densifies MISCON/FB LIT/FBTIME	CXXII	Ban soft-wrong≡half-credit / Partial Credit Score™ / all-wrongs-equal / gradebook cosplay

SAFE-LOADSHED	Named exam-week pause list; acuity FEI minutes; org redesign first; densifies BURN WARN/DDSCHEd/W ORKFORCE/EXAM/ROI	CXXIII	Ban hours-booked crunch NS / Load Shed Score™ / infinite-availability / pour≡exam support
SAFE-DISTRACTOR	Misconception-mapped functioning distractors; function>fashion; densifies TAXON/MISCON/GENQ	CXXIV	Ban Distractor Score™ / always-four≡better / beauty traps / unmapped “diagnostic” MC
SAFE-CONFMISS	Confidence×type recovery dose; CRI fork; hypercorrection priority; densifies TAXON/CALIB/MONITOR/COLD	CXXV	Ban Identity Score™ / Confidence Personality™ / equal-weight recovery / inflate-confidence NS
SAFE-MICRODEB	Capped student-first micro-AAR with traces; protect under LOADSHED;densifies AAR/LOADSHED/HITL/TALK/FBLIT	CXXVI	Ban skip-as-default / lecture-as-debrief / Microdebrief Score™ / +25% ACT debrief ads / vibes-only close
SAFE-CRUNCHMSG	Shed-honesty + protected thinking-minutes CBC/messages; densifies LOADSHED/WTP/SOFTMSG/PDASH/EXAM	CXXVII	Ban 24/7 default / hours-hero / fear-parent emails / silent shed / Availability Score™
SAFE-FORMTRAP	Named format-trap visual distractors; coherent structural stems; densifies DISTRACTOR/FORMAT/DUAL/GENQ	CXXVIII	Ban FormatTrap Score™ / seductive stems / beauty visual traps / unmapped visual “diagnostic”
SAFE-MICROBRIEF	Capped shared-model open (edge/success/trap/role + synthesis); densifies AAR/MICRODEB/LOADSHED/HITL/EXPECTANCY	CXXIX	Ban “no time for brief” / pour-open / Brief Score™ / late-join shame / aviation·OR·I-PASS ACT cosplay

SAFE-CRUNCHCTA	Dual labeled Solver/Practice CTAs under crunch; instrumental verbs; fiduciary defaults; follow-on Practice; densifies HWHELP/C RUNCHMSG/HELP/D URABLE/SOFTMSG	CXXX	Ban Chegg-hero CTA / fear-countdown dump / Dual-CTA Score™ / silent Practice collapse / assisted⇒ready
SAFE-FIGKEY	Structural figure key-verify for FormatId generation; deterministic spine; use gating; densifies GEN Q/FORMTRAP/DUAL /COVER/EXPLAINQ A	CXXXI	Ban fluency SVG pass / judge-alone figure ship / Figure Key Score™ / unverified FormatId diagnostic / FormatId-count hero
SAFE-ASYNCBRIEF	Capped async Set card on late/early wait; session-related occupied wait; if-then chip; live verify-not-replace; densifies MICROBRIEF/SRL/ATTN/LOADS HED/EXPECTANCY	CXXXII	Ban idle spinner / pour-open wait / Wait Minutes·Lobby Score™ / entertainment-primary / forced marathon Practice-before-join / late-shame feed
SAFE-XFERGRAIN	One-tap edge-matched Practice grain after instrumental Solver; short default; solo gate; SCHED return; densifies HWHELP/C RUNCHCTA/SCHED/ HELP/FADE/PROOF	CXXXIII	Ban dump-and-quit / maze handoff / Transfer Score™ / assisted⇒ready / marathon-after-hint / X P-for-help-without-solo
SAFE-HCELIG	Seals-first FormatId eligibility for C4; hcelig_pass ≠ tagged-in-bank; honest holes > pad; densifies FIGKEY/GENQ/CALIB/COLD/FORMAT/C OVER/FORMTRAP	CXXXIV	Ban tagged⇒C4 / fluency eligibility / coverage-padding / HideCorrectness Score™

SAFE-KNOWNETA	Known-finite late-tutor ETA with refresh honesty; occupied wait under ASYNCBRIEF; densifies ASYNCBRIEF/ MICROBRIEF/LOA DSHED/EXPECTANCY	CXXXV	Ban soon-default / optimistic slip / apology-primary / ETA Score™ / tutor-shame feed
SAFE-SOFTGATE	Soft-require default after instrumental help; hard-gate only with avoidance falsifiers; densifies XFERGRAIN/ HELP/CRUNCHCTA	CXXXVI	Ban universal hard-gate / Gate Score™ / Forced Follow-On % NS / hint-tax shame
SAFE-SPLITPOOL	Graded Practice-reveal vs C4 pools; no silent Map dual-write; ticketed promote; densifies HCELIG/RE PAIR/GENQ/COVER	CXXXVII	Ban unified bank≡diagnose / silent dual-write / auto-promote / SplitPool Score™
SAFE-SCANCOMP	Blueprint composition from sealed C4 stock; omit+label holes; densifies HCELIG/SPLIT POOL/COVER/FOR MAT	CXXXVIII	Ban symbolic pad / equal-quota theater / Composition Score™ / diagnose-every-format ads
SAFE-HOLETRUST	Parent hole honesty: label + seal reason + one fill CTA; densifies SCANCOMP/PDASH/ COVER/WTP/SOFTMSG	CXXXIX	Ban hide-holes / child-deficit shame / complete-visual ads / Hole Trust Score™
SAFE-PROMOTE	Ticketed Practice→C4 QA checklist + seals re-verify; demote faster than promote; densifies SPLITPOOL/REPAIR/ FIGKEY/HCELIG/COVER	CXL	Ban play-count/thumbs auto-promote / silent hcelig_pass / Promote Score™
SAFE-COOLDOWN	Per-student cool-down after reveal before same-family C4; isomorphic substitutes then holes; densifies SPLITPOOL/SCANCOMP/ CALIB/COLD	CXLI	Ban immediate same-stem C4 / forever·global lock default / Cool-Down Score™ / pad under cool-down

SAFE-STUDHOLE	Student hole parity with HOLETRUST; product-hole frame; one Practice CTA; no fake greens; densifies HOLETRUST/COLD/CALIB/SOFTMSG/FB LIT	CXLII	Ban parent-only honesty / fake greens / ability shame / empty mindset poster / Student Hole Score™
SAFE-FILLETA	Capacity-grounded finite fill bands; no play-count progress; no date soft-pass seals; densifies PROMOTE/KNOWNETA/HOLETRUST/STUDHOLE	CXLIII	Ban coming-soon / play-count fill bars / Filleta Score™ / complete-visual-by-date
SAFE-DUALRATE	Risk-triggered second rater; seals before dual; disagreement→reject/repair; densifies PROMOTE/FIGKEY/HITL/FILLETA	CXLIV	Ban always-dual / never-dual / Dual Rater Score™ / mammography-cosplay ACT ads
SAFE-FAMHASH	Sealed isomorphic family as cool-down compromise unit; substitutes outside family; densifies COOLDOWN/GENQ/FIGKEY/SCANCOMP	CXLV	Ban item_id-only / concept-wide ban / embedding-only Family Hash Score™
SAFE-EXITEMO	Destake gap-scan exit copy; task/process + one CTA; no mood telemetry; densifies STUDHOLE/PRIVACY/EXPOSE/SOFTMSG	CXLVI	Ban Anxiety·Calm·Exit Emotion Score™ / consoling ability comfort / ego Completeness % / emotion-AI exit
SAFE-CAPPLAN	Dual minutes as first-class FILLETA capacity line; headroom; label dual scarcity; densifies DUALRATE/FILLETA/WORKFORCE	CXLVII	Ban dual-omitted ETA / ETA-skipped dual / always-dual without FTE / Dual Rater·Capacity Score™
SAFE-TMPTRUST	Blueprint dual once; inherit single-rater under sealed sameness + sample audit + revoke; densifies DUALRATE/FAMHASH/FIGKEY/CAPPLAN	CXLVIII	Ban forever-dual clones / silent inherit / Template Trust Score™ / FERF≡C4-ready

SAFE-REPAIRSTRAT	Path-stratified post-promote repair + contingency drill-downs; κ aide not NS; densifies DUALRATE/PROMOTE/INSTRUMENT/TMPTRUST	CXLIX	Ban Kappa % NS / pooled repair vanity / Landis-band marketing / Dual Shame boards
SAFE-SURGEPANEL	Make-or-buy dual portfolio with seals, parallel-run, demobilization honesty; densifies CAPPLAN/LOADSHED/WORKFORCE/DUALRATE	CL	Ban Surge Hire % / unsealed freelancer dual / always-hire·contract dogma / hospital-surge cosplay
SAFE-PANELRAMP <i>(provisional)</i>	Seal literacy → shadow dual → graduate before contract C4; densifies SURGEPANEL/FIGKEY/DUALRATE/CAPPLAN	CLIII	Ban Parallel Score™ / day-one C4 / Loom≡ready / SOW≡scaled dual
SAFE-DEMOB <i>(provisional)</i>	Planned post-spike dual wind-down with handover, capture, FILLETA demob honesty; densifies SURGEPANEL/WORKFORCE/REPAIRSTRAT	CLIV	Ban Demobilization Score™ / invent-work / cliff-cut / forever-panel latch
Block → near-miss interleave → spaced mix	Blocking = acquisition scaffold; delayed mixed accuracy = readiness signal	XXIX, XXXIX, LXXIV	Ban “shuffle = science”; ban blocked-accuracy vanity
Student-generated why	Faded examples + structured principle/misconception prompts <i>before</i> AI wrap	XXVI, XL	Ban AI-monologue≡SE; ban explain-own-wrong-first default
Competitive wedge	Do not out-content Khan, out-streak Duo, out-delight Brilliant, or out-fluency ChatGPT	XXXV, XX	Sell recoverable struggle + competence evidence + solo transfer
Equity of worlds	Story wrap is identity technology; tokenism and stereotype-cueing copy are kills	XXXVI, XXVIII	HIST-EQ; belonging without “even you can”

Merged / demoted (do not treat as separate products): “AI tutor,” “mastery path,” “growth mindset,” “engagement,” “grit,” “habit streak,” “celebrate mistakes,” “community of practice,” “deliberate practice brand,” “film study,” “recital,” “grandmaster review,” “exposure therapy,” “girl STEM brand,” “WTP survey,” “Identity Score™,” “8-week math person,” “character equity,” “knowledge-graph moat,” “emotion AI empathy,” “science-backed page count,” “Ivy tutor SKU,” “AI wrote N questions,” “FERPA-compliant badge,” “perfect-interval SRS,” “tutor headcount / retention-%,” “mastery fireworks / personal Ebbinghaus,” “FSRS/Anki brain optimizer,” “cram/bootcamp/+points packages,” “instant feedback AI / always-delay science,” “bridge-count / we connect everything,” “day-one personalized mastery / placement belt,” “every format / visual learning / Format Personality™,” “parent fireworks portal / Family Leaderboard / live wrong-answer stalk,” “immersive story world / lore engagement / Story Engagement Score™,” “user-count / star-wall / vague confidence testimonials as learning proof,” “item-count / complete ACT bank / AI fills all holes coverage theater,” “always-full-worked / unlimited solutions / never-fade Solver,” “unlimited hints / never stuck / Help Score™,” “DAU/streak/XP North Star / FEI Score™ / thumbs-up=learning,” “unlimited free hard peeks / Hint Score™ / XP-for-hints / perfect-hint AI,” “longer=better / Explanation Score™ / unlimited thorough AI essays / monologue=SE,” “maximize productive struggle minutes / Struggle Score™ / blank-time grit / always-wait / never-stuck dump,” “fail-first always / every miss is productive failure / PF Score™ / discovery-without-consolidation,” “guild / master tutor / cognitive apprenticeship™ costume / Apprenticeship Score™ / modeling-minutes / never-fade coach,” “Focus Score™ / deep-work theater / phone-shame / mid-attempt streak pings / multitasking digital-native talent,” “Dual Coding Score™ / always-add-picture / decoration-as-dual-coding / visual-learner meshing / animation-first multimedia / split-legend density theater,” and “Monitoring Score™ / metacognition-without-control / easy-only calibration / one-scan calibration theater,” and “AI-replaces-tutors / tutoring-is-free / unlimited-human / hours-booked / Session ROI Score™ theater,” and “SRL Score™ / reflection-streak / survey-SRL=event / streak=self-regulation theater,” and “universal E0→E3 destiny / Example Timing Score™ / fixed=adaptive fade costume,” and “Generative Learning Score™ / AI=generation / teach-for-XP / eight-strategy GLA bomb,” and “Psychological Safety Score™ / poster safe-space / Bermuda-triangle / praise-fog climate,” and “Homework Completion Score™ / tonight-% / Chegg-dump / assisted=ready / purity-ban help,” and “Embodiment Score™ / kinesthetic styles / always-gesture / pose-camera AI,” and “Peer Explanation Score™ / teach-for-XP / Discord=knowledge-building / AI=peer teaching,” and “Curiosity Score™ / cliffhanger-without-commit / postdiction / guess-for-XP / instant-reveal entertainment,” and “Analogy Score™ / Mapping Minutes / mention=pedagogy / surface-story=structure / AI likeness dump,” and “Writing Score™ / Journal Minutes / freewrite=learning / longer=deeper / AI=WTL,” and “Spatial Score™ / Mental Rotation Minutes / brain-training=identity / Spatial Personality™,” and “Formative Score™ / Assessment Minutes / exit-ticket=learning / 0.4–0.7 AfL ads,” and “Feedback Literacy Score™ / Feedback Minutes / delivery=uptake / thumbs=used,” and “Collaboration Score™ / Group Minutes / Discord=CSCL / group-correct=ready,” and “Desirable Difficulty Score™ / Challenge Minutes / always-ease-before-ACT / always-max-DD / cram-week=ready,” and “Representation Score™ / Diversity Minutes / costume=belonging / match=quality,” and “Capture Score™ / Photo Minutes / OCR=mastery / photo→answer / Photomath-pipe,” and “forever-E0 anxious track / Anxiety·Calm Score™ fade gates / Fade Score™,” and “Miss Streak / soft-wrong shame push / live wrong-answer stalk / empty mindset parent letter,” and “always-on Map / Map Minutes·Map Score™ / mid-attempt full-graph,” and “Burnout Score™ / Wellness Minutes / wellness-primary under caseload,” and “fluency=pedagogy / Explanation·Rubric Score™ / LLM-judge-alone wrap ship,” and “Partial Credit Score™ / soft-wrong=half-credit / all-wrongs-equal gradebook,” and “hours-booked crunch / Load Shed Score™ / pour=exam support,” and “Distractor Score™ / always-four / beauty-trap MC,” and “Identity Score™ / Confidence Personality™ / equal-weight recovery / inflate-confidence,” and “Microdebrief Score™ / Debrief Minutes / skip-as-default lecture-close,” and “24/7 availability default / hours-hero crunch / fear-parent emails / Availability Score™ / silent shed,” and “FormatTrap Score™ / Visual Distractor Minutes / seductive stems / beauty visual traps / unmapped visual diagnostic,” and “Brief Score™ / Brief Minutes / Opening Streak / no-time-for-brief / pour-open / late-join shame / I-PASS cosplay ACT ads,” and “Chegg-hero crunch CTA / finish-tonight primary / fear-countdown dump / Dual-CTA Score™ / Help Tonight Minutes / silent Practice collapse / assisted=ready,” and

“fluency SVG≡keyed figure / LLM-as-judge-alone figure ship / Figure Key Score™ / Visual QA Minutes / unverified FormatId diagnostic / FormatId-count or AI-drew-N-diagrams hero / soft-pass wrong geometry,” and “Wait Minutes / Lobby Score™ / Prebrief Streak / idle spinner / pour-open wait / entertainment-primary fill,” and “dump-and-quit Solver close / maze handoff / Transfer Score™ / Handoff Minutes / Follow-on Streak / assisted≡ready / marathon-after-hint / XP-for-help-without-solo,” and “universal hard-gate after every hint / Gate Score™ / Forced Follow-On % / hint-tax shame / ignore help-avoidance when Forced % rises,” and “hide-format_not_probed for conversion / child-deficit visual-hole shame / complete-visual ads under holes / Hole Trust Score™ / Coverage Satisfaction / Honesty Minutes / fear-scarcity hole WTP,” and “play-count/thumbs auto-promote / silent hcelig_pass / fluency soft-pass promote / Promote Score™ / QA Minutes / Promotion Throughput %,” and “immediate same-stem C4 after reveal / forever-global cool-down lock / Cool-Down Score™ / Exposure Minutes / Lock Duration % / pad under cool-down,” and “parent-only hole honesty / fake greens for not_probed / ability-shame hole frames / empty growth-mindset poster without CTA / Student Hole Score™ / Completeness % / Honesty Minutes / absence≡mastery,” and “coming-soon / play-count fill bars / date soft-pass / Filleta Score™,” and “always-dual / never-dual / Dual Rater Score™ / mammography cosplay,” and “item_id-only cool-down / concept-wide ban / Family Hash Score™,” and “Anxiety·Calm·Exit Emotion Score™ / consoling ability comfort / emotion-AI exit,” and “dual-omitted FILLETA / ETA-skipped dual / Capacity Score™,” and “forever-dual clones / silent inherit / Template Trust Score™,” and “Kappa % NS / pooled repair vanity / Landis-band marketing,” and “Surge Hire % / Contract Minutes / unsealed freelancer dual / hospital-surge cosplay” are *not* independent North Stars — they are subordinate UX under FEI + the rows above (borrow *method*, never *costume*).

Next research bottleneck: Part CLV **Gold-set planted probes for path×stratum monitors**. Company law through SAFE-SURGE PANEL (CL); **SAFE-REVOKE (CLI) + SAFE-PATHBLAME (CLII) + SAFE-PANELRAMP (CLIII) + SAFE-DEMOB (CLIV) provisional** densifiers pending synthesizer. Experiment families .../SURGE PANEL/REVOKE/PATHBLAME/PANELRAMP/DEMOB gate claims above L1. Researcher count since v1.19 synthesizer = 4.

Part II — Key Insights (challenged)

Insight 1 — Fear is a performance bottleneck, not a personality trait

FACT: Math anxiety is associated with reduced working-memory availability during math tasks; anxious learners underperform relative to ability (Ashcraft & Kirk line of work; broader anxiety–WM literature).

HYPOTHESIS: For a large subset of MindCraft’s Maya archetype, the binding constraint is not “missing explanation,” it is **threat appraisal** (“I will look stupid”).

Implication: Soft-wrong UX, hide-correctness diagnostics, wizard coaching, and narrative framing are not polish — they are load-bearing if the thesis holds.

Red Team: Many high performers succeed under stress. Anxiety interventions have heterogeneous effects. Do not medicalize ordinary difficulty.

Insight 2 — Growth mindset is real, small, and context-dependent

FACT: Yeager et al. (2019, *Nature*) — brief online growth-mindset intervention improved grades for lower-achieving students and increased advanced math enrollment in a nationally representative U.S. sample; effects stronger when peer norms supported challenge-seeking.

FACT / REVIEW: Domain-general mindset messaging shows mixed results; math-specific mindset interventions more often report positive outcomes (2023 systematic review in *Educational Research Review*).

FACT / CAUTION: “False growth mindset” (praise effort without teaching strategy / changing environment) is harmful (Dweck’s own caution; Kohn’s critique).

Implication: MindCraft must not ship empty “believe in yourself” copy. Mindset must be embodied in **task design, feedback, and tutor norms**.

Insight 3 — Autonomy, competence, relatedness are the psychological tripod

FACT: Self-Determination Theory (Ryan & Deci) — sustained intrinsic motivation requires autonomy, competence, relatedness.

FACT (meta): SDT-based education interventions show reliable gains for autonomy and competence support; relatedness effects are less consistent (Wang et al. 2024 meta-analysis).

Implication: Product loops must deliver:

- Autonomy: meaningful choice (path, story world, pace)
- Competence: mastery grain that produces earned wins
- Relatedness: tutor presence, parent signal, peer-safe identity (hardest; do not fake with empty social)

Insight 4 — Mastery learning works — until you measure the wrong thing

FACT: Kulik, Kulik & Bangert-Drowns (1990) meta-analysis — mastery programs show positive exam effects (overall ~0.5 SD in their synthesis), stronger for weaker students; can increase time; self-paced college mastery can reduce completion.

FACT / CONTRADICTION: Slavin’s “Mastery Learning Reconsidered” (1987) — little support for group-based mastery on *standardized* measures; experimenter-made tests show moderate effects; coverage vs mastery tradeoff is real.

Implication: MindCraft’s mastery graph is promising **if** mastery means transferable competence, not local quiz familiarity. Measure far transfer, not only same-item success.

Insight 5 — Habit products ≠ identity products

FACT / INDUSTRY: Duolingo’s public method emphasizes gamification for motivation and return (streaks, leagues, variable rewards) (Duolingo Method whitepaper, 2023). Streaks exploit loss aversion; they drive DAU.

HYPOTHESIS: Streaks can increase practice frequency without changing “I am a math person.” Habit without identity yields brittle engagement (stop streak → stop learning).

Implication: Use habit mechanics as **on-ramps**, not as the destination. Pair with identity-marking rituals (map fill, story progression, tutor recognition of growth).

Insight 6 — Narrative is not decoration; it is encoding + meaning

FOUNDER BELIEF (WORLD_VISION): Math stripped of invention-story becomes alienating machinery.

HYPOTHESIS: Historical/human invention frames increase curiosity and reduce “arbitrary rule” appraisals, improving persistence.

Evidence status: Strong tradition in math education (history of math pedagogy); fewer large RCTs than mindset literature. Treat as **design prior + research program**, not settled science.

Insight 7 — AI commoditizes *answers*, not *becoming*

HYPOTHESIS: As answer generation approaches free, willingness-to-pay migrates to:

- 1. Diagnosis of the *actual* break (gap map)
- 2. Emotional containment during struggle
- 3. Accountability relationships (tutor)
- 4. Identity-consistent practice worlds
- 5. Trusted outcome signaling to parents

SPECULATION: In 30 years, “show me the steps” is ambient. “Help me become the kind of person who stays with hard problems” remains a human+system problem.

Part III — Research Log (v1 seed)

Date	Finding	Type	Action
2026-08-16	SAFE-DEMOB: Post-spike dual demobilization without fidelity cliff; kill Demobilization Score™ / invent-work / cliff-cut / forever-panel latch	Evidence	DEMOB-1...5; next = CLV gold-set path×stratum probes
2026-08-16	SAFE-PANELRAMP: Contract SME panel seal onboarding & parallel-run gates; kill Parallel Score™ / day-one C4 / Loom≡ready / SOW≡scaled dual	Evidence	PANELRAMP-1...5; next = CLIV post-spike demobilization

2026-08-15	SAFE-TMPTRUST: Template-trust inheritance after first-of-cell dual; kill forever-dual clones / silent inherit / Template Trust Score™ / FERF≡C4-ready	Evidence	TMPTRUST-1...5; next = CXLIX post-promote repair dashboards
2026-08-15	SAFE-REPAIRSTRAT: Post-promote repair stratified dashboards; kill Kappa % NS / pooled repair vanity / Landis-band marketing	Evidence	REPAIRSTRAT-1...5; next = CL dual-capacity hire panels
2026-08-15	SAFE-PATHBLAME: Path-conditioned dual blame tests (selection vs quality); kill Dual Shame Score™ / unconditioned dual rankings / Propensity Theater™ / case-mix cosplay	Evidence	PATHBLAME-1...5; next = CLIII contract SME panel onboarding
2026-08-15	SAFE-REVOKE: Template-trust revoke & sample-dual cadence under repair spikes; kill forever-trust latch / frozen sample-% / Template Trust Score™ / AQL≡ACT / CUSUM cosplay	Evidence	REVOKE-1...5; next = CLII path-conditioned dual blame
2026-08-15	SAFE-SURGEPANEL: Dual-capacity hire vs contract SME panels under exam spikes; kill Surge Hire % / unsealed freelancer dual / hospital-surge cosplay	Evidence	SURGEPANEL-1...5; next = CLI template-trust revoke cadence

2026-08-15	SAFE-CAPPLAN: Dual-rater capacity planning vs FILLETA band honesty; kill dual-omitted ETA / ETA-skipped dual / always-dual without FTE / Dual Rater Score™	Evidence	CAPPLAN-1...5; next = CXLVIII template-trust inheritance
2026-08-14	SAFE-EXITEMO: Gap-scan exit emotion without Anxiety Score™; kill Anxiety/Calm/Exit Emotion Score™ / consoling comfort / ego Completeness % / emotion-AI exit	Evidence	EXITEMO-1...5; next = CXLVII dual-rater capacity × FILLETA
2026-08-14	SAFE-FAMHASH: Isomorphic stem-family hashing for cool-down; kill item_id-only / concept-wide ban / embedding-only Family Hash Score™	Evidence	FAMHASH-1...5; next = CXLVI gap-scan exit emotion
2026-08-14	SAFE-DUALRATE: Second-rater thresholds for figured C4 admits; kill always-dual / never-dual / Dual Rater Score™ / mammography-cosplay guarantees	Evidence	DUALRATE-1...5; next = CXLV isomorphic cool-down hashing
2026-08-14	SAFE-FILLETA: Promote backlog SLA vs parent hole-ETA honesty; kill coming-soon / play-count fill bars / date soft-pass / FilleTA Score™	Evidence	FILLETA-1...5; next = CXLIV second-rater thresholds
2026-08-15	Synthesizer v1.19: merge CXLIII–CL SAFE-* into I.4; kills #97–#104	Synthesis	Surviving commercial law; next = CLI template-trust revoke

2026-08-14	Synthesizer v1.18: merge CXXXIV–CXLII SAFE-* into I.4; kills #88–#96	Synthesis	Surviving commercial law; next = CXLIII promote backlog SLA
2026-08-14	SAFE-STUDHOLE: Student-facing hole copy after gap-scan; kill parent-only honesty / fake greens / ability shame / empty mindset poster / Student Hole Score™	Evidence	STUDHOLE-1...5; next = CXLIII promote backlog SLA
2026-08-13	SAFE-COOLDOWN: Exposure cool-down after reveal before same-stem C4; kill immediate same-stem C4 / Cool-Down Score™ / pad under cool-down	Evidence	COOLDOWN-1...5; next = CXLII student-facing hole copy
2026-08-13	SAFE-PROMOTE: Promote-to-C4 tutor QA for Practice-only FormatId; kill play-count/thumbs auto-promote / silent hcelig_pass / Promote Score™	Evidence	PROMOTE-1...5; next = CXLI exposure cool-down
2026-08-13	SAFE-HOLETRUST: Parent trust copy when visual diagnostic holes remain; kill hide-holes / shame frames / complete-visual ads / Hole Trust Score™	Evidence	HOLETRUST-1...5; next = CXL promote-to-C4 tutor QA
2026-08-13	SAFE-SCANCOMP: FormatId gap-scan composition under sealed-stock scarcity; kill symbolic pad / equal-quota theater / Composition Score™ / all-formats ads under thin C4	Evidence	SCANCOMP-1...5; next = CXXXIX parent trust copy when visual holes remain

2026-08-13	SAFE-SPLITPOOL: Reveal-Practice vs C4 split pools for sealed FormatId; kill unified dual-use / silent Map writes / auto-promote / SplitPool Score™	Evidence	SPLITPOOL-1...5; next = CXXXVIII FormatId gap-scan composition
2026-08-13	SAFE-SOFTGATE: soft-require vs hard-gate follow-on under help-avoidance; kill universal hard-gate / Gate Score™ / hint-tax shame	Evidence	SOFTGATE-1...5; next = CXXXVII Reveal-Practice vs C4 split pools
2026-08-13	SAFE-KNOWNETA: known-finite ETA honesty in late-tutor lobby; kill soon-default / optimistic slip / apology-primary / ETA Score™ / shame feed	Evidence	KNOWNETA-1...5; next = CXXXVI soft-require vs hard-gate transfer
2026-08-13	SAFE-HCELIG: hide-correctness FormatId eligibility gates; kill tagged=C4 / fluency eligibility / coverage-padding / HideCorrectness Score™	Evidence	HCELIG-1...5; next = CXXXV known-finite ETA lobby
2026-08-12	Synthesizer v1.17: merge CXXVI–CXXXIII SAFE-* into I.4; kills #80–#87	Synthesis	Surviving commercial law; next = CXXXIV hide-correctness FormatId eligibility
2026-08-12	SAFE-XFERGRAIN: one-tap Solver→Practice transfer grain; kill dump-and-quit / maze handoff / Transfer Score™ / assisted⇒ready / marathon-after-hint	Evidence	XFERGRAIN-1...5; next = CXXXIV hide-correctness FormatId eligibility

2026-08-12	SAFE-ASYNCBRIEF: async pre-brief cards on late/early wait; kill idle spinner / pour-open wait / Wait Minutes·Lobby Score™ / entertainment-primary / Practice-before-join marathon	Evidence	ASYNCBRIEF-1...5; next = CXXXIII one-tap Solver→Practice transfer
2026-08-12	SAFE-FIGKEY: figure key-verify loops for generated FormatId stems; kill fluency SVG≡keyed / judge-alone ship / Figure Key Score™ / unverified FormatId diagnostic	Evidence	FIGKEY-1...5; next = CXXXII async pre-brief cards
2026-08-12	SAFE-CRUNCHCTA: dual-rail Solver CTA copy under crunch; kill Chegg-hero CTA / fear-countdown dump / Dual-CTA Score™ / silent Practice collapse / assisted≡ready	Evidence	CRUNCHCTA-1...5; next = CXXXI figure key-verify loops
2026-08-12	SAFE-MICROBRIEF: session-open brief fidelity under late joins; kill no-time-for-brief / pour-open / Brief Score™ / late-join shame / I-PASS cosplay ACT ads	Evidence	MICROBRIEF-1...5; next = CXXX dual-rail Solver CTA under crunch
2026-08-12	SAFE-FORMTRAP: format-trap distractors on diagram/graph items; kill FormatTrap Score™ / seductive stems / beauty visual traps / unmapped visual diagnostic	Evidence	FORMTRAP-1...5; next = CXXIX session-open brief fidelity

2026-08-11	SAFE-CRUNCHMSG: parent crunch-week message tests; kill 24/7 default / hours-hero / fear-parent emails / Availability Score™	Evidence	CRUNCHMSG-1...5; next = CXXVIII format-trap distractors
2026-08-11	SAFE-MICRODEB: tutor micro-debrief fidelity under time caps; kill skip-as-default / lecture-as-debrief / Microdebrief Score™ / +25% ACT debrief ads	Evidence	MICRODEB-1...5; next = CXXVII parent crunch-week message tests
2026-08-11	Synthesizer v1.16: merge CXVII–CXXV SAFE-* into I.4; kills #71–#79	Synthesis	Surviving commercial law; next = CXXVI tutor micro-debrief fidelity
2026-08-11	SAFE-CONFMISS: confidence-tagged gap-scan miss recovery; kill Identity Score™ / equal-weight recovery / inflate-confidence NS	Evidence	CONFMISS-1...5; next = CXXVI tutor micro-debrief fidelity
2026-08-11	SAFE-TAXON: typed soft-wrong mark taxonomy; kill Partial Credit Score™ / all-wrongs-equal / gradebook cosplay	Evidence	TAXON-1...5; next = CXXIII exam-week load shedding
2026-08-11	SAFE-DISTRACTOR: misconception-mapped MC distractors; kill Distractor Score™ / always-four / beauty traps / unmapped “diagnostic”	Evidence	DISTRACTOR-1...5; next = CXXV confidence-tagged gap-scan misses
2026-08-11	SAFE-LOADSHED: exam-week tutor load shedding; kill hours-booked crunch NS / infinite-availability / pour≡exam support	Evidence	LOADSHED-1...5; next = CXXIV distractor design

2026-08-11	SAFE-EXPLAINQA: principle-short QA gates for AI coach wraps; kill fluency $\equiv$ pedagogy / Explanation Score™	Evidence	EXPLAINQA-1...5; next = CXXII soft-wrong taxonomy
2026-08-10	SAFE-MAPORCH: brief $\rightarrow$ grain-slice $\rightarrow$ adjust Map timing; kill always-on / Map Minutes / mid-attempt full-graph	Evidence	MAPORCH-1...5; next = CXX tutor burnout early-warning
2026-08-10	SAFE-SOFTMSG: grain+failure-enhancing+one CTA after soft-wrong; kill Miss Streak / live stalk / empty mindset posters	Evidence	SOFTMSG-1...5; next = CXIX Map mid-session orchestration
2026-08-10	SAFE-FADEX: protect-then-fade under HMA; kill forever-E0 / flood / Anxiety·Calm Score™ fade gates	Evidence	FADEX-1...5; next = CXVIII parent soft-wrong framing
2026-08-10	Synthesizer v1.15: merge CVIII–CXVI SAFE-* into I.4; kills #62–#70	Synthesis	Surviving commercial law; next = CXVII fade $\times$ anxiety
2026-08-10	SAFE-OFFLINE: attempt $\rightarrow$ capture $\rightarrow$ confirm; kill Capture Score™ / OCR $\equiv$ mastery / photo $\rightarrow$ answer hero	Evidence	OFFLINE-1...5; next = CXVII fade $\times$ anxiety
2026-08-10	SAFE-ROLEMODEL: fidelity-first near-peer belonging; kill Representation Score™ / costume $\equiv$ belonging / match $\equiv$ quality	Evidence	ROLEMODEL-1...5; next = CXVI offline practice fidelity
2026-08-10	SAFE-DDSCHED: horizon-matched DD hold/ease; kill Desirable Difficulty Score™ / always-ease / always-max / cram $\equiv$ ready	Evidence	DDSCHED-1...5; next = CXV tutor identity & belonging

2026-08-09	SAFE-COLLAB: structured collab + solo FEI gate; kill Collaboration Score™ / Discord=CSCL / group-correct=ready	Evidence	COLLAB-1...5; next = CXIV DD under exam deadlines
2026-08-09	SAFE-FBLIT: feedback literacy = sense+use; kill Feedback Literacy Score™ / delivery=uptake / thumbs=used	Evidence	FBLIT-1...5; next = CXIII collaborative vs solo FEI
2026-08-09	SAFE-WRITE: short prompted WTL protocols; kill Writing Score™ / Journal Minutes / freewrite=learning / AI=WTL	Evidence	WRITE-1...5; next = CX spatial skills
2026-08-09	SAFE-SPATIAL: FormatId-embedded spatial hops; kill Spatial Score™ / Mental Rotation Minutes / brain-training=identity	Evidence	SPATIAL-1...5; next = CXI formative assessment
2026-08-09	SAFE-FA: formative cycles not AfL badge; kill Formative Score™ / 0.4–0.7 legend ads	Evidence	FA-1...5; next = CXII feedback literacy
2026-08-09	SAFE-ANALOG: structure-map joins; kill Analogy Score™ / mention=pedagogy / surface-story=structure / AI likeness dump	Evidence	ANALOG-1...5; next = CIX writing-to-learn
2026-08-09	Synthesizer v1.14: merge C–CVII SAFE-* into I.4; kills #54–#61	Synthesis	Surviving commercial law; next = CVIII analogical bridges
2026-08-09	SAFE-PREDICT: commit before reveal; kill Curiosity Score™ / cliffhanger-without-co mmit / postdiction / guess-for-XP	Evidence	PREDICT-1...5; next = CVIII analogical bridges (Synthesizer eligible)

2026-08-08	SAFE-PEERX: building-gated teach-back; kill Peer Explanation Score™ / teach-for-XP / Discord≡building / AI≡peer teaching	Evidence	PEERX-1...5; next = CVII curiosity & prediction
2026-08-08	SAFE-GESTURE: strategy-tied optional gesture/draw; kill Embodiment Score™ / styles meshing / always-gesture / pose-camera	Evidence	GESTURE-1...5; next = CVI peer explanation quality
2026-08-08	SAFE-ERRCLIMATE: stay-diagnose error climate in tutoring; kill Safety Score™ / Bermuda-triangle / praise-fog	Evidence	ERRCLIM-1...5; next = CIV homework vs practice identity
2026-08-08	SAFE-HWHELP: dual-rail Solver/Practice; kill Completion Score™ / Chegg-dump / assisted≡ready	Evidence	HWHELP-1...5; next = CV gesture & embodiment
2026-08-08	SAFE-GENERATE: SOI-forcing sparse GLAs in Practice; kill Generative Learning Score™ / AI≡generation / teach-for-XP	Evidence	GENACT-1...5; next = CIII error climate
2026-08-08	SAFE-SRL: phase→control Zimmerman cycles in Practice; kill SRL Score™ / streak≡SRL / survey≡event	Evidence	SRL-1...5; next = CII generative learning
2026-08-08	SAFE-EXPTIME: grain-local example↔solve timing; kill universal E0→E3 destiny / fixed≡adaptive / Timing Score™	Evidence	EXPTIME-1...5; next = CI SRL cycles

2026-08-07	Synthesizer v1.13: merge XCII–XCIX SAFE-* into I.4; kills #46–#53	Synthesis	Surviving commercial law; next = C worked-example timing
2026-08-07	SAFE-ROI: Map minute budgets; kill AI-replaces-tutors / tutoring-is-free / hours NS / Session ROI Score™	Evidence	ROI-1...5; next = C worked-example timing
2026-08-07	SAFE-ATTN: protect attempt windows; kill Focus Score™ / phone-shame / mid-attempt DAU pushes	Evidence	ATTN-1...5; next = XCVII dual coding
2026-08-07	SAFE-DUAL: contiguous structural figures; kill Dual Coding Score™ / styles meshing / always-add-picture	Evidence	DUAL-1...5; next = XCVIII metacog monitoring
2026-08-07	SAFE-MONITOR: named judgments → control; kill Monitoring Score™ / one-scan calibration	Evidence	MONITOR-1...5; next = XCIX tutor vs AI ROI
2026-08-07	SAFE-APPRENTICE: model→coach→fade + articulate; kill guild cosplay / modeling≡lecture	Evidence	APPRENTICE-1...5; next = XCVI attention residue
2026-08-07	SAFE-PF: fidelity-gated generate→consolidate; kill every-miss≡PF / fail-first-always	Evidence	PF-1...5; next = XCV apprenticeship
2026-08-06	Synthesizer v1.12: merge LXXXIV–XCI SAFE-* into I.4; kills #38–#45	Synthesis	Surviving commercial law; next = XCII explanation length
2026-08-05	Synthesizer v1.11: merge LXXV–LXXXIII SAFE-* into I.4; kills #29–#37	Synthesis	Surviving commercial law; next = LXXXIV tutor talk-ratio

2026-08-04	Synthesizer v1.10: merge LXVII–LXXIV SAFE-* into I.4; kills #21–#28	Synthesis	Surviving commercial law; next = LXXV tutor workforce
2026-08-02	Synthesizer v1.9: merge LIX–LXVI SAFE-* into I.4; kills #18–#20	Synthesis	Surviving commercial law; next = LXVII ontology moat
2026-08-01	Synthesizer v1.8: merge L–LVIII SAFE-* into I.4; kills #15–#17	Synthesis	Surviving commercial law; next = LIX parent WTP
2026-08-01	SAFE-EXPOSE: graded approach; kill therapy cosplay / flooding / Calm Score™	Evidence	EXP-O1...5; expectancy violation
2026-08-01	SAFE-ANNOT / MUSIC / FILM: postmortem + ladder + sparse clips	Evidence	ANNOT/MUSIC/FIL M-1...5
2026-07-31	Synthesizer v1.7: merge XLI–XLIX SAFE-* into I.4; kills #10–#14	Synthesis	Surviving commercial law; next = L confidence calibration
2026-07-31	Misconceptions productive: SAFE-MISCON; soft-wrong ≠ empty kindness	Evidence	MIS-1...4; springboard + classify
2026-07-30	Transfer/VT: story ≠ far transfer; blocked ≠ ready	Evidence	TR-1...4; SAFE-TRANSFER
2026-07-30	Sleep/stress: encode→sleep→claim; pressure ≠ DD	Evidence	TIM-1...4; SAFE-TIMING
2026-07-29	Synthesizer v1.6: merge XXXIII–XL into I.4 doctrine stack	Synthesis	Surviving commercial law; next was XLI DD×anxiety
2026-07-29	SE prompts: AI monologue ■ SE; structure before wrap	Evidence	SE-1...4; densify Exp A/D
2026-07-29	Interleaving: block→n ear-miss→spaced mix; kill shuffle=science	Evidence	IL-1...4; ban blocked vanity

2026-07-29	Goal climate: appearance/streak NS killed; TARGET audit	Evidence	GO-1...4
2026-07-29	Eccles E×V×Cost: kill single motivation KPI	Evidence	EVT-1...4
2026-07-25	Constituted Research Lab + Constitution v1	Process	Create OS
2026-07-25	Scarcity thesis partially false as slogans; true as product wedge	Analysis	Rewrite pitch language
2026-07-25	Yeager 2019 constrains mindset claims	Evidence	Ban empty mindset copy
2026-07-25	Soft-wrong + wizard coach shipped in product	Product	Instrument coach → retry rate
2026-07-25	Mastery literature split (Kulik vs Slavin)	Evidence	Define mastery measurement carefully
2026-07-25	Bloom 2-sigma overstated vs VanLehn / modern tutoring metas	Evidence	Stop marketing 2-sigma
2026-07-25	Founder sequence wounded → SAFE-CRAFT	Synthesis	Attempt earlier; invention-story optional
2026-07-25	Deliberate practice necessary≠sufficient (Hambrick line)	Evidence	Progress-vs-self framing

Part IV — Evidence Table (selected)

Claim	Best supporting evidence	Contradicting / limiting evidence	Confidence	MindCraft implication
Brief mindset intervention can move grades + advanced math taking	Yeager et al., Nature 2019	Effects heterogeneous; peer norms matter; replication debates in broader mindset meta-analyses	Medium–High	Mindset must be ecological, not banner text

Math-domain mindset > generic	2023 Ed Research Review systematic review	Study quality varies; publication bias possible	Medium	Math-specific identity language
SDT need support raises motivation	Wang et al. 2024 meta	Relatedness effects weaker	Medium–High	Design for autonomy + competence first
Math anxiety impairs performance	Cognitive literature (Ashcraft tradition)	Not all low performers are anxious	Medium–High	Affective onboarding is first-class
Mastery learning raises achievement	Kulik et al. 1990	Slavin 1987 on standardized tests	Medium	Mastery graph + transfer tests
Gamified habits raise return	Duolingo method / industry practice	Habit ≠ deep learning; extrinsic traps (Deci)	Medium	Streaks as on-ramp only
1:1 tutoring is uniquely powerful	Bloom “2 sigma” framing (1984)	Often overstated; quality varies; VanLehn etc. nuance tutoring effects	Medium	Keep human tutors; AI amplifies, not replaces

Part V — Contradictory Evidence (do not skip)

1. **Mindset skepticism:** Some metas find small average effects; school implementations often degrade into posters.
2. **Grit / resilience hype:** Duckworth-style grit is contested; structural inequality and curriculum quality matter more than character slogans.
3. **Discovery learning failures:** Pure constructivism without guidance often fails (Kirschner, Sweller, Clark critiques). Narrative worlds must still teach.
4. **Engagement ≠ learning:** Time-on-app can rise while understanding falls (edtech’s original sin).
5. **AI may increase explanation value temporarily:** Novices may trust fluent wrongness; trust becomes the scarce resource, not explanation volume.
6. **Some students want speed and drills:** Identity-through-story is not universal. Offer multiple vessels.

Part VI — Mental Models

VI.1 The Identity Transformation Cascade (ITC)

HYPOTHESIS — MindCraft original synthesis

Threat ↓ → Attention available → Micro-success → Attribution ("I figured it")
→ Willingness to re-enter challenge → Accumulated competence evidence
→ Narrative update ("maybe I can") → Identity claim ("I do math")
→ Environment selection (harder courses, peers, tutors)
→ Reinforcing loop

Balancing loop: Premature hard problems → threat ↑ → avoidance → identity freeze.

Product translation: Gap scan (no shame) → soft-wrong → tiny earned wins → map fill → tutor recognition → advanced path.

VI.2 The Three Doors (what students actually need in a stuck moment)

Door	Need	Wrong product response	Right response
A	Safety	“Here’s a longer explanation”	Contain affect; normalize miss
B	Clarity	Pep talk	Precise model / worked example at right grain
C	Agency	Take over and solve for them	Scaffolded choice; student does the decisive step

FOUNDER BELIEF: Tutors often open Door B first. Many students are still behind Door A.

VI.3 Knowledge commodity stack (post-AI)

Layer 0: Facts & procedures → commoditizing fast
Layer 1: Explanations → commoditizing fast
Layer 2: Diagnosis of misconceptions → partially automatable; trust-sensitive
Layer 3: Adaptive sequencing → automatable with good ontology
Layer 4: Emotional co-regulation → scarce (human + careful UX)
Layer 5: Identity & meaning → scarce
Layer 6: Accountability relationship → scarce

MindCraft’s durable stack lives in **Layers 2–6**, with 0–1 as utilities.

VI.4 Transfer from outside education

Domain	Mechanism	Transfer to MindCraft (surviving law)
Games (Celeste, FromSoftware-lite design)	Fair difficulty + death as information	Soft-wrong as physics (SAFE-MISCON / FEI)

Sports coaching	Film study of <i>specific</i> error	SAFE-FILM: ≤3-cue attempt clips + self-model
Music	Scales → repertoire → recital	SAFE-MUSIC: ingredient drills → chosen pieces → transfer proof
Therapy (exposure)	Graded approach + expectancy violation	SAFE-EXPOSE: situation hierarchy under destaked stakes
Aviation / military	Brief → attempt → structured debrief	SAFE-AAR: four-question student-first close
Chess annotation	Postmortem before engine; principle tags	SAFE-ANNOT: annotate→classify→drill; wins too
Martial arts belts	Visible competence ladder	Mastery pips / map regions (criterion, not league)
Language apps	Habit loops	SAFE-HABIT cue practice — without Duo's extrinsic ceiling
Religion / ritual	Shared meaning + belonging	SAFE-RITUAL: light optional marks; anti-totalism
Craft apprenticeship	Legitimate peripheral participation	SAFE-CoP: real peripheral work + old-timer; not Discord KPIs

Part VII — Universal Learning Principles (provisional)

1. **Emotion gates cognition** in threatened domains.
 2. **Earned difficulty** beats random difficulty.
 3. **Feedback should be information, not identity judgment.**
 4. **Attribution matters:** strategy > talent for growth; luck attributions kill agency.
 5. **Spaced retrieval** beats massed re-reading (cognitive science consensus).
 6. **Worked examples** → **fading guidance** for novices (cognitive load).
 7. **Belonging uncertainty** destroys persistence for stereotyped groups (Walton/Cohen tradition).
 8. **Transfer requires varied practice**, not identical clones.
 9. **Human accountability** multiplies AI tools.
 10. **Identity is a lagging indicator;** design leading indicators (challenge-seeking, return-after-fail).
-

Part VIII — Product Implications (systems, not features)

VIII.1 Core system: Fear → Evidence → Identity (FEI Loop)

Reinforcing loop R1 (desired):

Safety UX → attempt → informative miss → coach → success → map fill → identity → more attempts.

Balancing loop B1 (threat):

Public shame / red buzz → avoidance → less practice → skill lag → more shame.

MindCraft already partially implements FEI: hide-correctness diagnostic, soft-wrong, wizard coach, story chapters, knowledge map. **Missing instrumentation:** does coach raise retry rate and later challenge-seeking?

VIII.2 What to stop optimizing

- Vanity DAU without learning transfer
- Explanation length as quality proxy
- Points that do not mark identity-relevant milestones

VIII.3 What to instrument next (minimum viable science)

1. **Retry rate within 2 minutes of soft-wrong** (with vs without wizard coach)
2. **Challenge-seeking:** % choosing harder level when offered
3. **Return after fail session** (D1 retention after a loss session)
4. **Self-identity item** (pre/post, 1–2 items, math-specific)
5. **Advanced course intent / enrollment** (slow, gold)

VIII.4 Tutor system role

Tutors are not content pipes. They are:

- Affect regulators
- Attribution coaches
- Bridge diagnosticians (concept A known, link A→B broken)
- Identity witnesses (“I saw you become someone who stays”)

AI should brief tutors, not replace witnessing.

Part IX — Experiments (pre-registered style)

Core A–D remain. Chapter ticks densified them into families — prefer the densified IDs when shipping.

Experiment A — Wizard coach on soft-wrong

- **Question:** Does the under-Graph wizard increase productive retry vs soft-wrong alone?
- **Design:** A/B, chapter + practice
- **Primary:** retry within 120s; secondary: eventual correct without hint binge

- **Falsifier:** no lift; or lift only in engagement with worse accuracy
- **Densified content (v1.6):** Coach turns must be SE-constrained (principle / misconception contrast), not pep talk — see SE-1...4 / SE-2 as operational pairs with Experiment D.

Experiment B — Story-first vs bare stem

- **Question:** Does invention-story framing raise persistence on first miss?
- **Design:** within-concept crossover
- **Primary:** time-to-abandon after first miss
- **Falsifier:** story slows without improving retention/transfer
- **Equity densify:** HIST-EQ / Part XXXVI — whose story, tokenism kill.

Experiment C — North Star validation

- **Question:** Which early metric best predicts 8-week challenge-seeking?
- **Candidates:** streak length, soft-wrong retry, identity item, map mastery Δ , delayed mix accuracy
- **Method:** observational + lagged prediction
- **v1.6 prior:** Expect streaks to lose; retry_120s + transfer_pass + mastery-motive challenge_accept to win.

Experiment D — Explanation timing

- **Question:** Is explanation-before-attempt worse than attempt-with-safety for anxious students?
- **Design:** screen with anxiety pretest; randomize explanation-first vs attempt-first
- **Falsifier:** explanation-first wins for all segments
- **Densified:** Prefer SE-2/SE-3 (contrastive / structured generation) over monologue-first.

Experiment families mounted in chapters (do not orphan)

Family	Identifies (max claim without follow-up)	Parts
AIT-*	Guarded AI → solo transfer / anti-sycophancy	XXXIII
DAG-*	Estimand hygiene; covariate vs mediator	XXXIV
CSA-*	Session preference vs competitors; paste under guards	XXXV
HIST-EQ / equity probes	Story belonging without stereotype cue	XXXVI
EVT-*	Expectancy / value / cost levers	XXXVII
GO-*	Mastery vs performance CTA climate	XXXVIII
IL-*	Schedule → discrimination / delayed mix	XXXIX
SE-*	Prompt regime → transfer / misconception recurrence	XL

DD-*	Equip/destake/dose × anxiety → retry + delayed mix	XLI
SC-*	Criterion vs peer rank climate → affect + challenge	XLII
HAB-*	Cue/if-then vs streak colonization → return without identity harm	XLIII
IM-*	Informational vs engagement-loot rewards → free-choice + transfer	XLIV
RES-*	Growth-zone + support recruitment vs grit theater	XLV
EXP-*	Tutor brief / CIOF / AI expectancy → talk ratio + solo transfer	XLVI
TIM-*	Practice timing × sleep/stress → delayed transfer	XLVII
TR-*	VT / hug / bridge sequence → format-hop + solo transfer	XLVIII
MIS-*	Soft-wrong springboard / mis-mapped distractors / confidence tier	XLIX
CAL-*	Confidence-tiered feedback / hide-correctness / underconf unlock	L
DP-*	Session spine / edge targeting / method feedback / effort cap	LI
CoP-*	Named practice / peripheral work / old-timer / trajectory metrics	LII
RIT-*	Light ritual vs silent / recovery marker / anti-streak liturgy	LIII
AAR-*	Brief + four-question close / student-first / trace-backed debrief	LIV
FILM-*	Sparse clip cues / self-controlled review / self↪model pairing	LV
MUSIC-*	Ladder UX / repertoire choice / identified technique / transfer recital	LVI
ANNOT-*	Annotate-before-reveal / forced restudy / classify→drill / annotate wins	LVII

EXP-O-*	Situation hierarchy / expectancy check / fade safety / multi-context	LVIII
WTP-*	Parent CBC / FEI report vs score / message arms / cash pilot	LIX
REPAIR-*	Post-hallucination repair / reattempt / packetized escalate	LX
IDM-*	Validated identity battery / demand controls / FEI co-primary	LXI
LONG-*	8/26/52w horizons / fade checks / staged identity claims	LXII
EXAM-*	Destake-learn / dose-prove ladders / delayed mixed under stakes	LXIII
ELL-*	Language vs math load / linguistic modify / format triangulation	LXIV
GEND-*	Cue hygiene / practice recognition / no ST→score / no pink track	LXV
STRUCT-*	Bandwidth hygiene / structure>slogans / access honesty / no grit-NS	LXVI
ONTO-*	Ontology-constrained next vs chat / bridge route / Map brief / state vs summary	LXVII
HITL-*	Map-briefed tutor / prompt-first playbook / coached QA / tutor+Map vs AI-alone	LXVIII
PRIV-*	Check-in TTL / soften-on-stress / task-only brief / ERT preference / refusal-safe	LXIX
LABMETA-*	Pre-reg gate / Red Team bite / copy lint / synthesizer demote / external skeptic	LXX
GRAIN-*	Trained near-peer / Map vs expert-no-Map / peer-pair / hire-by-fidelity / parent CBC grain	LXXI
GENQ-*	Verify-on vs off / generated vs seed / diagnostic seal / prompt-harden drop rate / parent CBC verified-bank	LXXII

PROCURE-*	Trust-packet-first / NDPA vs custom / LEA-approved free pilot / district vs parent message / no-bio attribute	LXXIII
SCHED-*	Horizon equal-lag vs massed / expanding vs equal / delayed first return / generation vs restudy / ACT-date scheduler	LXXIV
WORK-*	Continuous coaching vs onboard-only / coach-capacity hire cap / caseload vs wellness / Map-brief hard-require / parent CBC fidelity	LXXV
FORGET-*	Aged-node vs fireworks / honest post-mission vs confetti / parent proof-age CBC / RI-salience vs fluency JOL / clear-state expiry	LXXVI
ADAPT-*	Fixed calendar vs banded modulation / due-set priority vs FIFO / inspectable vs opaque smart review / FSRs-urgency priority-only / stress-capped pull-forward	LXXVII
CRAM-*	Durable vs cram-pack CBC / cram-intercept UX / dual-rail label / score-lift vs proof-age ads / massed week vs horizon Returns	LXXVIII
FB-*	Micro-delay+SE vs instant KR / tiered timing / diagnostic hide / prove-rail KR / process vs ego toast	LXXIX
BRIDGE-*	Bridge-card vs endpoint re-teach / Map edge CTA / detector↔tutor / co-present vs sequential / parent CBC connection copy	LXXX
COLD-*	Seed±probes / honest Map chrome vs fireworks / uniform-easy QA / contradicting-reason on easy / parent CBC honest-map	LXXXI

FORMAT-*	Format-gap CTA / co-present conversion / structural vs decorative / supported multi-format / parent CBC format-flex	LXXXII
PDASH-*	Proof-age strip vs fireworks / improvement digest / absolute vs rank / autonomy vs nag / parent CBC MoC portal	LXXXIII
TALK-*	Prompt-first vs explain-first / coach dashboard vs clip QA / SE gate / AI prompt-cap / parent CBC construction	LXXXIV
STORYLOAD-*	Structure wrap vs plain / lore budget / skip-story / mastery-gated reveal / parent CBC clear-story vs immersion	LXXXV
PROOF-*	Competence-card vs star-wall landing / provincial case / two-sided MoC vs +points / opt-in share / parent CBC dated solo	LXXXVI
COVER-*	Honest matrix vs item-count hero / parent CBC holes+seal / district audit packet / in-app gap callout / audit cadence	LXXXVII
FADE-*	Backward completion ladder vs always-full / forward vs backward / fade-up after solos / SE-on-blank / parent CBC fade story	LXXXVIII
FADEX-*	Longer E0/E1 hold for high-MA vs aggressive fade / protective WE on bridge / incorrect mid-rung / destake vs timed fade / parent CBC protect-then-fade	CXVII
SOFTMSG-*	Grain+failure-enhancing+CTA vs miss-shame / autonomy vs controlling tip / anxious-parent pressure-reduce / digest+MoC / parent CBC honest-grain	CXVIII
MAPORCH-*	Grain-slice vs full-graph / tutor-private vs dual-screen / brief+between-attempt vs always-on / bridge slice vs red gallery / CBC timed Map	CXIX

HELP-*	Contingent instrumental vs unlimited bottom-out / SE-before-hint / help-invite vs freeze / meta-msg vs gate / parent CBC next-step help	LXXXIX
INSTR-*	Lagged FEI vs streak/DAU prediction / tutor KPI co-gates / event fire audit / motive-coded challenge CTA / parent CBC FEI report	XC
HINT-*	Soft-first contingent vs free hard / SE-cost vs dwell / fade×hint couple / peek budget / parent CBC next-step hints	XCI
EXPLAIN-*	Principle-short vs monologue / SE-before-wrap / length×fade / soft-wrong short mark / parent CBC clear-short	XCII
RETRIEVE-*	TOT micro-cue vs instant hint / freeze bound soft vs long wait / attempt-before-help / progress-aware dwell / parent CBC attempt-then-coach	XCIII
PF-*	High-fidelity PF vs DI / student-RSM C&C vs generic lecture / thin-prior gate / freeze-exit PF / parent CBC invent-then-compare	XCIV
APPRENTICE-*	Playbook model→coach→fade+articulate vs explain-pour / control-talk model vs steps-only / within-session fade / reciprocal beat / parent CBC take-the-wheel	XCV
ATTN-*	Mid-attempt push suppression / opt-in focus mode / tech-break between items / phone-salience tip / parent CBC attempt-window vs streak / Maya residue phenomenology	XCVI
DUAL-*	Contiguous annotated figure vs split legend / structural vs decorative / static vs animation default / coach+figure length / parent CBC format-hop	XCVII

MONITOR-*	Named judgment+miss-class vs easy-only / C4 destake / control-wired routing / one-scan vs revisit / parent CBC monitoring honesty	XCVIII
ROI-*	Map minute budget vs open hours / AI-reps+human-join vs AI-alone / FEI-per-tutor-minute dashboard / escalate precision / parent CBC scarce-human honesty	XCIX
EXPTIME-*	Fixed vs adaptive fade / concept vs concept×format initial stage / confidence→E vs probe→E / always-E0 vs always-E3 vs adaptive / parent CBC step-down help	C
SRL-*	Forethought chips vs autopilot / attribution→plan writeback / microanalysis vs journal / process vs ability attribution UI / parent CBC plans-the-attempt	CI
GENACT-*	Teach-fictive-peer vs solution-view / guided draw vs free vs provided / SOI summarize vs XP / sparse vs eight-menu / parent CBC generates-explanation	CII
ERRCLIM-*	Stay-diagnose vs redirect / soft-wrong+Map vs praise-only / parent CBC diagnosed-mistakes / error-episode QA vs warmth NPS / public vs private error visibility	CIII
HWHELP-*	Dual-rail Help vs Practice / assisted vs solo gate / costly dump vs purity ban / parent CBC tonight-help honesty	CIV
GESTURE-*	Strategy-tied gesture vs none / observe vs produce vs draw / FormatId-triggered / tutor model→fade / parent CBC optional motion	CV

PEERX-*	Building teach-back vs telling / inferential Q vs summary / explain→apply / HITL building code / near-peer matched vs open room	CVI
PREDICT-*	Soft-wrong predict→delay→repair vs instant dump / predict vs postdiction / predict vs example / demo predict vs watch / always vs soft-wrong-only gate	CVII
ANALOG-*	Co-present structure-map vs sequential re-teach / high-support cues vs bare compare / student map vs AI likeness / story-skin vs structure / always vs severity-gated	CVIII
WRITE-*	Short prompted protocol vs none vs long freewrite / cognitive+meta vs meta-only / student draft vs AI polish / written vs oral / weekly chips vs diary streak	CIX
SPATIAL-*	FormatId-embedded hop vs isolated rotation gym / scaffolded vs flood / destake vs ability-shame / boys-better costume A/B	CX
FA-*	Wired exit vs XP ticket / evidence→adjust vs got-it theater / FEI vs Formative Score™ / HITL cycle QA	CXI
FBLIT-*	Uptake gate vs essay+thumbs / decodeable vs jargon / affect-managed vs ego / compare-chip vs reveal-only	CXII
COLLAB-*	Scripted+solo-gate vs solo-only / open room vs roles / shared unlock vs solo unlock / accountability vs team XP	CXIII
DDSCHEd-*	Hold-mix vs ease-blocked last 14d / Protect+probe 72h vs scored flood / horizon ISI vs cram / dual-rail vs single scored calendar	CXIV

ROLEMODEL-*	Competence-visible bio vs demographic-only / preference match vs random-qualified / attainable vs superstar / belonging brief / parent likeness copy	CXV
OFFLINE-*	Offline-first vs always-camera / post-commit capture vs mid-solve phone / OCR-confirm vs auto-green / photo→coach vs photo→answer / parent CBC paper-proof	CXVI
BURNWARN-*	Fidelity/load early-warn vs wellness-primary / caseload redesign vs mood gate / coach-capacity hire / parent CBC fidelity honesty	CXX
EXPLAINQA-*	Principle-short rubric gate vs fluency / human audit vs LLM-judge-alone / soft-wrong short mark / parent CBC clear-short	CXXI
TAXON-*	Typed soft-wrong vs binary / type-conditioned deltas / step-join mark / parent categorical vs % / auto-type+audit	CXXII
LOADSHED-*	Acuity queue+pause list vs hours-booked / org redesign vs wellness-primary / dual-rail crunch / parent CBC shed honesty	CXXIII
DISTRACTOR-*	Mapped 3-option vs unmapped 4 / retire dead options / GENQ distractor-seal / theory-named soft-wrong / parent CBC grain-from-option	CXXIV
CONFMISS-*	Class-conditioned recovery vs equal deck / high-conf springboard priority / easy-vs-probe collision / parent CBC grain vs Personality™ / conf×type severity	CXXV
MICRODEB-*	Capped micro-AAR vs skip / student-first vs lecture-close / traces vs vibes / LOADSHED protect / parent CBC brief-close	CXXVI

CRUNCHMSG-*	Shed-honesty vs 24/7 CBC / thinking-minutes vs hours-hero / fear vs fiduciary / silent shed vs named pause / parent CBC crunch copy	CXXVII
FORMTRAP-*	Trap-mapped visual distractors vs beauty stems / conversion route vs topic-only / GENQ trap seal / parent CBC format-miss honesty	CXXVIII
MICROBRIEF-*	Capped shared-model open vs pour / late-join brief vs skip / student synthesis / AAR Set pair / parent CBC goals-before-problems	CXXIX
CRUNCHCTA-*	Dual labeled rails vs Chegg-hero / follow-on Practice vs dump-quit / fiduciary default / fear-countdown vs calm dual / parent CBC help+gym	CXXX
FIGKEY-*	Structural figure verify vs fluency SVG / use-gate vs unverified diagnostic / judge-alone reject / stem^key constraints / parent CBC keyed figures	CXXXI
ASYNCBRIEF-*	Occupied Set wait vs spinner / if-then chip / verify-not-replace / no entertainment-primary / parent CBC waiting-as-prep	CXXXII
XFERGRAIN-*	One-tap same-edge grain vs dump-quit / short vs maze / solo gate / SCHED return / parent CBC help-then-prove	CXXXIII
HCELIG-*	HCELIG-pass-only C4 vs mixed tagged / Practice-vs-C4 tier / honest holes vs pad / graphicacy label / parent CBC sealed visuals	CXXXIV
KNOWNETA-*	Finite ETA vs soon / optimistic slip vs refresh / apology-only vs duration / shame feed vs neutral / parent CBC wait honesty	CXXXV
SOFTGATE-*	Soft-require vs hard-gate / avoidance falsifier / autonomy copy / short hard-gate only / parent CBC no-hint-tax	CXXXVI

SPLITPOOL-*	Split pools vs unified dual-use / silent Map block / ticketed promote / coverage tables by use / parent CBC Practice≠C4	CXXXVII
SCANCOMP-*	Sealed-only compose vs pad / omit+label vs quota theater / depth-where-stock / exposure caps / parent CBC holes-as-features	CXXXVIII
HOLETRUST-*	Label+CTA vs hide / shame vs product-hole / complete-visual ads vs honesty / fear WTP vs fiduciary / parent CBC hole copy	CXXXIX
PROMOTE-*	Ticket+checklist vs play-count / demote-on-repair / second-rater triggers / fluency soft-pass ban / parent CBC fill path	CXL
COOLDOWN-*	Cool-down bar vs immediate same-stem / isomorphic substitutes / holes under empty cell / duration band / parent CBC memory≠diagnosis	CXLI
STUDHOLE-*	Student parity vs parent-only / fake greens vs chips / ability shame vs product-hole / mindset poster vs CTA / student CBC hole copy	CXLII
FILLETA-*	Capacity fill band vs soon / conservative vs optimistic slip / Practice CTA / ticket-state vs play-count bar / parent CBC verify honesty	CXLIII
DUALRATE-*	Triggered dual vs always/never / consensus vs either-admit / T1 vs T1∧T2 / IRR sample cadence / parent CBC dual honesty vs hospital cosplay	CXLIV
FAMHASH-*	Family vs item_id / narrow vs concept-wide / embed-auto vs SME seal / R2 on-off / parent CBC same-family honesty	CXLV
EXITEMO-*	Destake exit vs Anxiety Score™ / ability comfort vs strategy CTA / Completeness % / emotion-AI exit / parent-student CBC	CXLVI

CAPPLAN-*	Dual-aware FILLETA vs fantasy / util headroom / named dual constraint / always-dual@fixed FTE / parent CBC dual-honest band	CXLVII
TMPTRUST-*	Inherit+sample vs forever-dual / silent inherit / T5-only vs T1 clear / sample % / parent CBC blueprint honesty	CXLVIII
REPAIRSTRAT-*	Path-stratified repair vs pooled κ / contingency drill-down / inherit discord revoke / windowed drift / parent CBC no Kappa % ads	CXLIX
SURGEPANEL-*	Hire vs contract dual / parallel-run gate / dual_source repair / FILLETA surge label / demobilization without invent-work	CL
PANELRAMP-*	Seal literacy / shadow dual / graduate gate / FILLETA contract_ramp / DRIFT sample	CLIII
DEMOB-*	Planned handover / invent-work ban / capture gate / FILLETA demob / post-exit repair window	CLIV
REVOKE-*	Revoke-on-repair vs soft-investigate / switching sample vs frozen % / Shewhart vs CUSUM trip / re-qualify paths / parent CBC pause honesty	CLI

Part X — Weekly Priorities (research × product)

1. Instrument FEI loop events in analytics (retry, coach shown, write-mode exit).
 2. Run Experiment A for 2 weeks.
 3. Red Team the WORLD_VISION narrative claims with one outside learning scientist.
 4. Interview 10 Mayas: “When did math stop feeling dangerous?” (qual).
 5. Define “mastery” operationally so Slavin’s critique cannot gut the graph.
-

Part XI — Founder Questions (do not dodge)

1. If ChatGPT tutors become “good enough,” what *relationship* do parents still buy?

2. Are we building for Maya's identity — or for parents' anxiety about scores? Both? Tradeoff?
 3. Would we accept lower short-term accuracy for higher long-term challenge-seeking?
 4. Is the story world load-bearing, or is gap+tutor enough?
 5. What would falsify the identity-transformation thesis in 90 days?
-

Part XII — Long-term Roadmap (research horizons)

H1 (now): Prove FEI loop moves retry + challenge-seeking.

H2: Prove identity language + map produce durable self-concept change.

H3: Prove human tutor + AI diagnosis beats AI-alone on persistence and transfer.

H4: Generalize beyond math (science/identity) only after math identity mechanism is solid.

Part XIII — Open Problems

1. Causal path from narrative history-of-math to identity (under-identified).
 2. Relatedness at scale without creepy social.
 3. Measuring identity without demand effects.
 4. Preventing gamification from colonizing meaning.
 5. AI fluent nonsense vs trust calibration for teens.
 6. Equity: does story-world privilege culturally specific narratives?
-

Part XIV — Red Team Dossier (destroy weak arguments)

Kill #1: “Explanations are free, therefore MindCraft should not explain”

Destroyed form: Absolutism.

Surviving form: Explanations are *necessary infrastructure* with collapsing *willingness-to-pay*; differentiation is diagnosis + affect + identity + accountability.

Kill #2: “Just fix mindset”

Destroyed: Poster mindset.

Surviving: Ecological mindset (task + feedback + peer/tutor norms), math-specific — TARGET structures (XXXVIII).

Kill #3: “Engagement is the goal”

Destroyed: Engagement without transfer.

Surviving: Engagement as oxygen for the FEI loop, not the fire.

Kill #4: “Mastery graph = Bloom 2-sigma”

Destroyed: Inflated tutoring mythology.

Surviving: Mastery *can* help weaker students if mastery is real (transfer_pass) and time costs are managed.

Kill #5: “Identity transformation is unmeasurable, so ship vibes”

Destroyed: Vibes-as-strategy.

Surviving: Leading indicators (retry, challenge-seeking, delayed mix) + lagging identity items; L0–L4 claim ladder (XXXIV).

Kill #6 (v1.6): “AI tutor / Bastani GPT Tutor is our product”

Destroyed: Unguarded chat-as-learning; category-error marketing; guards-as-gain.

Surviving: Pedagogy wrap + solo_transfer_pass $\geq$ no-AI bar (XXXIII RT).

Kill #7 (v1.6): “Blocked accuracy / streaks / shuffle prove learning”

Destroyed: Fluency illusion; streak North Star; random shuffle $\equiv$ interleaving science.

Surviving: Short block $\rightarrow$ near-miss interleave $\rightarrow$ spaced mixed check (XXXIX); streaks as on-ramp only (HID).

Kill #8 (v1.6): “Any explain box / AI monologue = self-explanation science”

Destroyed: Untargeted justify-your-wrong-guess; chat length as depth.

Surviving: Structured principle/misconception SE on correct or known-incorrect content before AI wrap (XL).

Kill #9 (v1.6): “One motivation dial”

Destroyed: Collapsing expectancy and value; staff “math is useful” dumps as EVT.

Surviving: $E \times V \times \text{Cost}$ triad with student-authored utility (XXXVII).

Kill #10 (v1.7): “Harder is always better / desirable difficulty ignores anxiety”

Destroyed: Anxiety-blind interleave/SE/pressure-as-DD; hardness theater as rigor brand.

Surviving: SAFE-DD equip $\rightarrow$ destake $\rightarrow$ dose $\rightarrow$ frame $\rightarrow$ measure (XLI); gate IL/SE by segment.

Kill #11 (v1.7): “Leaderboards / streaks / XP create math identity”

Destroyed: Appearance climates, streak colonization, engagement-loot as love-of-math.

Surviving: SAFE-COMPARE + SAFE-HABIT + SAFE-REWARD — criterion feedback, cue practice, informational unlocks only (XLII–XLIV).

Kill #12 (v1.7): “Grit / Resilience Score™ / stay in the danger zone”

Destroyed: Character meters; danger-zone-as-desirable-difficulty; grit posters as product.

Surviving: SAFE-RESILIENCE growth-zone + support recruitment; measure recovery/transfer not MRS theater (XLV).

Kill #13 (v1.7): “Pygmalion labels / warm AI that finishes the problem”

Destroyed: Trait/weakness tutor briefs; warmth-without-output-demand as expectancy science.

Surviving: SAFE-EXPECTANCY task briefs + CIOF; AI must clear PWC/solo_transfer_pass (XLVI + XXXIII).

Kill #14 (v1.7): “Story / sleep tips / mistake celebration = transfer”

Destroyed: Narrative immersion as automatic far transfer; sleep-app brand; empty mistake posters; red-X-as-rigor; equal graph updates for all wrongs.

Surviving: SAFE-TRANSFER + SAFE-TIMING + SAFE-MISCON — vary on purpose, encode→sleep→claim, elicit→classify→springboard→route (XLVII–XLIX).

Kill #15 (v1.8): “Raise confidence / Belief Score™ / CA raises grades”

Destroyed: Inflate-confidence brand; confidence meters as identity; confidence assessment alone as attainment lever; equal updates for guess vs high-conf misconception.

Surviving: SAFE-CALIB — item-level elicit → miss class → tiered feedback; appropriate confidence; C4 hide-correctness as diagnostic hygiene (L).

Kill #16 (v1.8): “Hours / forums / ritual→score / domain costume = identity”

Destroyed: 10k-hours and minutes KPIs; Discord/streak ≡ CoP; unretracted ritual→performance marketing; bootcamp/conservatory/GM/therapy branding as belonging.

Surviving: SAFE-DP + SAFE-CoP + SAFE-RITUAL — diagnose→prove spine; named practice with real peripheral work; light optional marks pointing at recover/prove (LI–LIII).

Kill #17 (v1.8): “Tip-flood / engine-dump / flooding ≡ high-performance pedagogy”

Destroyed: Uncapped film-room tips; Stockfish/LLM monologue as annotation; day-one timed mixed ACT as “exposure”; calm-to-zero / Calm Score™; loss-only shame reels; bare-pass recital theater.

Surviving: SAFE-AAR + SAFE-FILM + SAFE-MUSIC + SAFE-ANNOT + SAFE-EXPOSE — brief–debrief, sparse clips, skills→pieces→proof, annotate→drill, graded approach with expectancy violation (LIV–LVIII).

Kill #18 (v1.9): “Likert≡WTP / survey dollars≡revenue / ACT-point packages”

Destroyed: Parent survey willingness as price truth; engagement-streak or unlimited-AI as default parent value props; guaranteed ACT points in tiers; message arms treated as package economics.

Surviving: SAFE-WTP — CBC with price+neither, ship-mapped attributes (diagnosis, FEI report, human minutes, AI policy), cheap-talk/cash pilots, message≠package (LIX).

Kill #19 (v1.9): “Never-hallucinates / apology≡repair / Identity Score™ / 8-week math person”

Destroyed: Invincible-AI brand; rote or empathy apology as trust fix; thumbs-up≡repaired; permanent disuse brand; Math Identity Score™; two-item HSLS pre/post≡FEI; splash-screen identity Likert; empty AI “math person” praise; end-of-program Likert≡durable identity; math-person-in-8-weeks.

Surviving: SAFE-REPAIR + SAFE-IDMEASURE + SAFE-LONGID — detect→own→Map-correct→reattempt→triggered escalate; validated factors + FEI behavioral co-primary; staged 8/26/52w claims with fade checks (LX–LXII).

Kill #20 (v1.9): “Day-one ACT flood / EL≡ability / pink track / grit-as-character”

Destroyed: Timed ACT as onboarding; blocked accuracy as exam-ready; Calm Score™ exam NS; math-is-language-free; forever procedure-only ELL track; translate≡identity; innate boys-better premise; ST-removal→ACT; pink/for-her easy track; grit-as-NS / Grit Score™ / poor-kids-need-grit / streak-as-character / scarcity-dropout≡moral-failure.

Surviving: SAFE-EXAM + SAFE-ELL + SAFE-GENDER + SAFE-STRUCTURE — destake-to-learn/dose-to-prove; separate language vs math load; parity + cue hygiene + practice recognition; structure/scaffolds before slogans (LXIII–LXVI).

Kill #21 (v1.10): “Graph file / RAG chat ≡ diagnosis moat”

Destroyed: Node-count or ontology JSON as moat; RAG or chat-memory as knowledge tracing; fluency/tokens as diagnosis; Bloom ITS-2σ; “smarter than ChatGPT”; ALEKS-cosplay without FEI transfer proof.

Surviving: SAFE-ONTOLOGY — diagnosis-before-dialogue; inspectable longitudinal state; LLM bookends under PWC; moat = state×ontology×FEI×solo transfer (LXVII).

Kill #22 (v1.10): “Warm human / hours / got-it ≡ tutor quality”

Destroyed: Warm presence or Ivy credentials as quality; hours-booked NS; “Do you understand?” / smile / stars as mastery or QA; Discord≡tutor QA; explain-first brand; AI-replaces-tutors or humans-skip-Map; onboarding deck≡coaching.

Surviving: SAFE-HITL — Map-briefed tutors; prompt>pour; answer evidence; complementarity; coaching cycles with fidelity; FEI QA; packetized escalate (LXVIII).

Kill #23 (v1.10): “Emotion AI / Anxiety Score™ / FERPA≡ethics”

Destroyed: Face/voice/wearable emotion recognition as empathy or personalization; Anxiety Score™ / continuous mood NS or parent SKU; FERPA-compliant as ethics proof; mandatory mood gates; emotion-trait tutor briefs; sell/enrich affect; empathy-camera brand.

Surviving: SAFE-PRIVACY — self-authored short-TTL check-in → pedagogy modifiers only; ERT ban; contextual integrity; refusal-safe; compliance is a floor not a doctrine (LXIX).

Kill #24 (v1.10): “Page-count science / Bayesian theater / cite-wash”

Destroyed: Chapter/page count as lab NS; “Bayesian update” without model check or experiment; DOI lists that change no product rule; FOUNDER BELIEF promoted to FACT by repetition; synthesizer that never demotes;

“science-backed Constitution” ads without falsifiers; fake precision confidence numbers.

Surviving: SAFE-LABMETA — label → Red Team/check → kill/wound/survive → claim-ladder demotion of copy; small FEI hard core; hostile prior on single-study overclaim; synthesizer must delete (LXX).

Kill #25 (v1.10): “Ivy/expert résumé / peer Discord ≡ tutoring quality”

Destroyed: PhD/Ivy/expert credentials as FEI proof or default SKU; Bloom 2σ human-tutor marketing; unstructured peer Discord ≡ product tutoring; “great explainers” hire bar; expertise ≡ accurate diagnosis; tutor-learning as student North Star.

Surviving: SAFE-TUTORGRAIN — trained near-peer default; structure>pairing; Map carries diagnosis; suppress knowledge-telling; expert=escalate/coach-of-tutors; hire by fidelity drills (LXXI).

Kill #26 (v1.10): “LLM items ≡ shipped bank / item-count hero / unverified diagnostic keys”

Destroyed: LLM-generated items as ready-to-ship without independent key verify; item-count / “AI wrote N questions” as hero metric or ACT/identity proof; unverified keys as C4 diagnostic or mastery ground truth; fluency of stem ≡ keyed correctness; clone floods ≡ transfer; scaling generation at ~30% verify drop.

Surviving: SAFE-GENQ — verify-before-ship; key fails hard-fail; Kane IUA for bank uses; drop-rate gate before scale; coverage≠NS; seal C4/mastery writes to verified keys (LXXII).

Kill #27 (v1.10): “FERPA badge / click-wrap / biometric school SKU ≡ district GTM”

Destroyed: FERPA-compliant badge as district-ready or ethics; teacher click-wrap / free viral classroom as primary school path; biometric/face-voice empathy as school differentiator; marketing-first RFP without DPA+Exhibit A; parent WTP copy as LEA trust story; targeted ads / sell-affect on pupil records.

Surviving: SAFE-PROCURE — privacy review = GTM gate; NDPA-ready contracting; Exhibit A honesty; free=paid approval rigor; no biometrics; no pupil-data ads; deletion on exit (LXXIII).

Kill #28 (v1.10): “Expanding SRS / perfect-interval AI / streak-as-spacing ≡ durable math”

Destroyed: Expanding-retrieval hero marketing; “AI found your perfect interval”; streak/daily-open as the spacing engine; cram-week ≡ ACT-ready; overlearn-tonight as retention; tip-restudy as Review; shuffle ≡ spaced retrieval; schedule-guaranteed ACT points.

Surviving: SAFE-SCHED — horizon-matched ISI×RI; delayed first return; equal-ish lags for long RI; generation Returns on Practice/Map; exam dual rail (spaced learn / timed prove); anxiety-aware lag stretch; instrument SCHED-* before SRS brand claims (LXXIV).

Kill #29 (v1.11): “Tutor headcount / tenure / wellness theater ≡ quality”

Destroyed: Tutor headcount / hours / retention-% as North Stars; tenure months ≡ FEI quality; onboarding deck ≡ continuous coaching; wellness-module theater as burnout fix; Burnout Score™; marketplace-volume / Ivy-bench as attrition insurance.

Surviving: SAFE-WORKFORCE — fidelity-over-tenure; continuous coaching; coach-capacity gates hiring; burnout→load redesign; named drift channels; retention without fidelity = vanity (LXXV).

Kill #30 (v1.11): “Mastery fireworks / Ebbinghaus Score™ ≡ durable competence”

Destroyed: Tonight clear / mastery fireworks as durable competence; personal Ebbinghaus / Forgetting Score™ as UI hero; streak-as-memory; confidence=retention; cured-forgetting ads; guaranteed retention% from decay chrome.

Surviving: SAFE-FORGET — age the evidence; JOLs≠retention; delayed probes as truth surface; frame returns as competence protection; no shame-curve hero (LXXVI).

Kill #31 (v1.11): “SM-2 / FSRS / black-box perfect-interval AI ≡ math identity”

Destroyed: SM-2/Anki/FSRS as hero brand; silent ease-hell ISI rewrite; flashcard MAE / recall-prob wins as ACT or solo_transfer proof; scheduler DAU/engagement as learning success; “perfect interval AI” copy.

Surviving: SAFE-ADAPT — calendar-first (SAFE-SCHED); adapt only inside declared RI bands with visible due+why; prefer selection-among-due over opaque rewrite; flashcard lineage ≠ FEI endpoints; instrument ADAPT-* before any SRS-optimizer marketing (LXXVII).

Kill #32 (v1.11): “Cram / bootcamp / +points packages ≡ durable math identity”

Destroyed: Cram-week ≡ ACT-ready or math-person; guaranteed ACT/SAT point SKUs; overlearn-tonight as retention method; score-only default parent brand; shame-parent GTM; streak/tonight-accuracy as durable proof; unlabeled “final mastery blitz.”

Surviving: SAFE-DURABLE — dual-rail GTM (durable learn + late prove); cram-intercept not cram-cosplay; legible MoC (proof age / returns) for WTP; no point guarantees; instrument CRAM-* before anti-cram attack ads (LXXVIII).

Kill #33 (v1.11): “Instant feedback AI / always-delay science ≡ learning”

Destroyed: Instant-feedback AI as hero claim; always-delay=desirable-difficulty science; feedback=learning (ignore Kluger harm rate); green-check / latency NPS as North Star; ego-toast=formative feedback; soft-wrong delay as Calm Score™.

Surviving: SAFE-FBTIME — mode-conditional timing (diagnostic hide / learn micro-delay+SE / prove KR labeled); focus>timing; prior-knowledge switch; instrument FB-* before timing brand ads (LXXIX).

Kill #34 (v1.11): “Bridge-count / re-explain ≡ connection mastery”

Destroyed: Bridge-/edge-count as moat or quality; chat or tutor re-explain of endpoint ≡ bridging the join; coverage of both topics ≡ connection; absolute “we connect everything”; missing-prereq conflated with bridge gap; shame-edge Map theater; analogy mention without cognitive supports ≡ pedagogy.

Surviving: SAFE-BRIDGE — connection-first narrative when endpoints non-naive; remediate the relation with co-present supports; measure join solo_transfer_pass; instrument BRIDGE-* before connection-hero ads (LXXX).

Kill #35 (v1.11): “Day-one mastery fireworks / confidence ≡ mastery”

Destroyed: Onboarding green Map / mastery confetti from empty or seed-only state; self-rated easy ≡ mastery or ACT readiness; skip-scan writing optimistic mastery; population prior sold as student biography; placement-belt / Identity Score™ cold start; diagnostic-completion % as North Star; chat fluency as calibrated prior.

Surviving: SAFE-COLD — humble prior → labeled seed → hide-correctness probes; Map chrome matches evidence age; sell honest mapping not fake knowing; instrument COLD-* before day-one personalization-hero ads (LXXXI).

Kill #36 (v1.11): “Format-count / every-format / decoration ≡ format mastery”

Destroyed: FormatId vocabulary size as moat; absolute “we teach every format” / visual-learning-for-all ads; decorative story art $\equiv$ FormatId coverage; single-vessel concept greens $\equiv$ format-flexible readiness; format shuffle $\equiv$ science without conversion; Format Personality™ / learning-style quiz as North Star.

Surviving: SAFE-FORMAT — format as first-class evidence; conversion as learning object with supports; structural tags only; severity in worstWeakness; instrument FORMAT-* before format-hero ads (LXXXII).

Kill #37 (v1.11): “Tonight-% / streak / shame-rank parent portal $\equiv$ trust”

Destroyed: Tonight accuracy / streak / XP as parent North Star; mastery fireworks on parent home; classmate shame ranks as default; live wrong-answer surveillance; Anxiety Score™ on parent view; gradebook/SIS cosplay as differentiation; controlling homework-nag loop; guaranteed points from parent widgets.

Surviving: SAFE-PDASH — belief correction via proof age + return agenda; one improvement CTA; autonomy support > control; ranks/affect off; differentiate on MoC not SIS; instrument PDASH-* before Family Leaderboard / live-progress campaigns (LXXXIII).

Kill #38 (v1.12): “Talk-% / silence / got-it $\equiv$ Socratic quality”

Destroyed: Talk Ratio Score™ or raw tutor/student talk-% as North Star; 50/50 airtime dogma; silence theater as “Socratic”; “got it?” as constructive evidence; AI monologue sold as dialogue via token counts; emotion-AI scoring of talk; Bloom-from-student-talk marketing.

Surviving: SAFE-TALK — construction > airtime; prompt density + student constructive share + wait bands + pour-streak QA; same schema for human and AI; instrument TALK-* before Socratic/talk-balance campaigns (LXXXIV).

Kill #39 (v1.12): “Immersion / lore / Story Engagement $\equiv$ learning”

Destroyed: Immersion or lore-depth as North Star / Story Engagement Score™; “story always helps learning”; seductive franchise trivia as pedagogy; forced pre-attempt lore walls; surface fandom personalization as hero; day-one encyclopedia before first attempt; story immersion $\equiv$ far transfer.

Surviving: SAFE-STORYLOAD — CLT budgets narrative; structure-encode or cut; coherence > immersion; novices get less lore; skip-story; language $\neq$ math; story $\neq$ transfer; instrument STORYLOAD-* before immersive-world campaigns (LXXXV).

Kill #40 (v1.12): “User-count / star-wall / vague praise $\equiv$ learning proof”

Destroyed: User-count / DAU / star-average as primary learning social proof; vague confidence testimonials without a named hop; Identity Score™ splash as proof; +points / guaranteed ACT case studies; Family Leaderboard as social proof; AI “math person” praise as competence; forced public portfolio / viral share as North Star.

Surviving: SAFE-PROOF — competence > crowd; enactive dated solo artifacts; provincial similarity; student-owned thin portfolio; two-sided honesty; opt-in dignity; instrument PROOF-* before star/user-count learning ads (LXXXVI).

Kill #41 (v1.12): “Item-count / complete-ACT / unsynced GENQ $\equiv$ coverage”

Destroyed: Item-count / “N questions” as marketing hero or North Star; “complete ACT bank / every concept” without concept $\times$ level honesty; GCSE/other-examTag sold as ACT inventory; counting unverified or unsynced generated items; coverage % $\equiv$ exam readiness or identity; ontology node count $\equiv$ playable bank; stale-audit coverage ads; “AI fills any hole instantly” as substantiation.

Surviving: SAFE-COVER — matrix > total; gaps as features; use-tier labels (Kane IUA); examTag honesty; GENQ never counted unverified; prior substantiation via regenerated audit; instrument COVER-* before complete-bank / item-count campaigns (LXXXVII).

Kill #42 (v1.12): “Always-full-worked / never-fade / solution-first $\equiv$ learning”

Destroyed: Always-full-worked or solution-first Solver as learning North Star / brand; never-fade step scaffolding; viewing steps $\equiv$ mastery; expertise-blind full dumps for near-fluent learners; unlimited perfect solutions as differentiator vs ChatGPT; Fade Score™ or guidance-% vanity NS; forward-only fade dogma without local test.

Surviving: SAFE-FADE — stage not costume; completion bridge; backward fade first; expertise-aware up/down; attempt grain before reveal; proof = solo transfer; instrument FADE-* before unlimited-solutions / always-show-every-step campaigns (LXXXVIII).

Kill #43 (v1.12): “Unlimited-hints / help-NPS / avoidance-as-grit $\equiv$ learning”

Destroyed: Unlimited hints / always-open bottom-out as learning North Star; hint volume or help NPS as pedagogy proof; help avoidance as grit; Help Tutor metacognition as guaranteed domain/ACT gains; gaming detectors as parent shame dashboards; “AI always knows when you need help” black-box brand; Help Score™ vanity NS; permanent tutor dependency as “support.”

Surviving: SAFE-HELP — instrumental > executive; abuse and avoidance both fail; contingent stage not unlimited menu; SE/attempt before bottom-out; metacognition $\neq$ score magic; dignity telemetry; proof = solo transfer; instrument HELP-* before never-stuck / unlimited-hints campaigns (LXXXIX).

Kill #44 (v1.12): “DAU/streak/XP / FEI Score™ / thumbs-up $\equiv$ learning North Star”

Destroyed: DAU / streak / XP / time-on-app as learning North Stars; blocked-session accuracy as readiness NS; FEI Score™ composite sold as proof; single-event tutor pay or shame ranks on `retry_120s` alone; coach thumbs-up / help NPS without `solo_transfer_pass`; science-backed identity ads from unvalidated dashboards; celebrating FEI lifts when `hint_binge` / `ai_reveal_rate` rise.

Surviving: SAFE-INSTRUMENT — ship the four FEI events; NSM stack + stage OMTM; XXI.4 co-primary gates; anti-Goodhart / anti-gaming companions; Kane IUA for marketing uses; dignity telemetry; instrument INSTR-* / Experiment C before identity-from-dashboard campaigns (XC).

Kill #45 (v1.12): “Unlimited free hard peeks / Hint Score™ / XP-for-hints $\equiv$ learning”

Destroyed: Unlimited free hard / bottom-out peeks as learning North Star or kindness brand; Hint Score™ or raw hint-click engagement; XP/streak rewards for peeking; shame timers or parent live peek-stalk as “cost”; black-box “AI perfect hint” mystique; fade costume with a free hard submenu underneath.

Surviving: SAFE-HINT — soft before hard; contingency up/down; peek cost = construction; fade $\times$ hint one machine; separate ITS soft/hard from Saye & Brush hard/soft scaffolds; inspectable stage; instrument HINT-* before unlimited-free-hints campaigns (XCI).

Kill #46 (v1.13): “Longer $\equiv$ better / Explanation Score™ / monologue $\equiv$ SE $\equiv$ learning”

Destroyed: Longer coach/Solver essays as quality or care; token count / thoroughness NPS / Explanation Score™ as learning KPIs; AI monologue labeled as self-explanation; unlimited thorough AI explanation brand; fixed multi-paragraph wrap across expertise stages; seductive coach fluff as pedagogy.

Surviving: SAFE-EXPLAIN — short principle-first before monologue; germane ≠ more text; SE before wrap; length fades with expertise/fade×hint; coherence > charm; no token North Star; instrument EXPLAIN-* before unlimited-thorough-AI campaigns (XCII).

Kill #47 (v1.13): “All-struggle≡productive / Struggle Score™ / blank-time / never-stuck ≡ learning”

Destroyed: Treating every stall as productive struggle or desirable difficulty; blank-seconds / dwell-XP / Struggle Score™ as learning North Stars; always-wait silence theater; instant hard dump on first hesitation; never-stuck answer-delivery brand; “errors always scar memory” fear marketing that skips attempts.

Surviving: SAFE-RETRIEVE — name stall mode; attempt before heavy help; TOT → wait + micro-cue; freeze → bounded soft scaffold; progress marks productive struggle (not clock alone); unsuccessful retrieval + feedback can potentiate; instrument RETRIEVE-* before struggle-theater / never-stuck campaigns (XCIII).

Kill #48 (v1.13): “Every-miss≡PF / fail-first-always / PF Score™ / discovery-without-consolidation ≡ learning”

Destroyed: Labeling every unsuccessful attempt as productive failure; fail-first / invent-always as brand or North Star; generation without Phase-2 consolidation that builds on student RSMs; PF Score™ / generation-minutes KPIs; pure discovery-as-identity; anxiety-forced endless generation cosplay; guaranteed ACT points from invent-first.

Surviving: SAFE-PF — designed generate→consolidate-on-student-RSMs; fidelity + prior-knowledge gates; guided/I-PS default when priors thin, procedural/exam chrome, or freeze (SAFE-RETRIEVE exit); proof = conceptual + transfer; instrument PF-* before invent-first campaigns (XCIV).

Kill #49 (v1.13): “Guild cosplay / modeling≡lecture / never-fade / Apprenticeship Score™ ≡ learning”

Destroyed: Guild / master / journeyman costume as pedagogy proof; modeling as lecture or AI monologue without student attempt; never-fade permanent coaching as kindness; talk-% alone as cognitive apprenticeship; Apprenticeship Score™ / modeling-minutes NS; Discord≡apprenticeship; reciprocal-teaching costume without turn-taking; CA marketing without inspectable phase logs.

Surviving: SAFE-APPRENTICE — model control/heuristics → coach → scaffold→fade; articulation required; shared ca_phase for HITL+AI; proof = solo transfer + articulation quality; instrument APPRENTICE-* before guild / apprenticeship-costume campaigns (XCV).

Kill #50 (v1.13): “Focus Score™ / phone-shame / mid-attempt DAU pushes / multitasking-as-talent ≡ learning”

Destroyed: Focus Score™ / Deep Work Score™ as North Star; phone-shame or character-moralizing as pedagogy; “digital natives multitask fine” as free pass; mid-attempt streak/DAU pushes as engagement science; phone presence as proof of modern learning; focus-kit ads promising guaranteed ACT points; covert gaze/mic attention surveillance.

Surviving: SAFE-ATTN — protect mid-attempt windows; residue-aware park/resume; opt-in focus + tech breaks; presence hygiene without shame; interrupt telemetry ≠ productive struggle; instrument ATTN-* before focus-costume campaigns (XCVI).

Kill #51 (v1.13): “Dual Coding Score™ / always-add-picture / styles meshing / animation-default ≡ learning”

Destroyed: Dual Coding Score™ / Visual Engagement Score™ as North Star; always-add-a-picture as dual-coding science; decorative chrome ≡ FormatId learning; visual-learner meshing tracks; animation/video as default modernity over static annotated figures; stacked narration+on-screen text+busy figure as thoroughness; split legends as information density; unverified generated figures as diagnostic truth; guaranteed ACT points from multimedia dual-coding ads.

Surviving: SAFE-DUAL — structural dual coding; contiguous labels; coherence weeding; static annotated default; conversion remains the object (SAFE-FORMAT); instrument DUAL-* before dual-coding / visual-learner campaigns (XCVII).

Kill #52 (v1.13): “Monitoring Score™ / metacognition-without-control / one-scan calibration ≡ learning”

Destroyed: Monitoring Score™ / Metacognition % as North Star; “we teach metacognition” ads without judgment→control wiring; global easy/kinda/hard alone as calibrated monitoring science; immediate post-item surety equated to delayed JOL accuracy; confidence slider without miss-class routing; raise-metacognition / raise-confidence as identity transformation; one gap-scan as calibration cure; guaranteed ACT points from metacognitive gap-scan packaging.

Surviving: SAFE-MONITOR — named judgment types; C4 destake; confidence×outcome miss classes; control wiring mandatory; densifies SAFE-CALIB / SAFE-COLD; instrument MONITOR-* before metacognition brand campaigns (XCVIII).

Kill #53 (v1.13): “AI-replaces-tutors / tutoring-is-free / hours-booked / Session ROI Score™ ≡ learning”

Destroyed: AI replaces tutors / tutoring-is-free slogans; unlimited human entitlement SKUs; hours booked / session count as North Star; Session ROI Score™ gamified ops vanity; VanLehn ITS≈human as ChatGPT-dump license; Bastani practice-with-AI scores as learning proof; school pooled tutoring ES as thin marketplace ACT guarantee.

Surviving: SAFE-ROI — tutor minutes as scarce inventory; Map minute budgets; constrained AI for reps; humans for joins/destake/repair/witness; FEI per tutor-minute; instrument ROI-* before AI-replaces-tutors / unlimited-human campaigns (XCIX).

Kill #54 (v1.14): “Universal E0→E3 / Example Timing Score™ / fixed≡adaptive ≡ learning”

Destroyed: Universal fixed E0→E3 destiny as pedagogy; treating fixed fade schedules as adaptive science; gap-scan confidence alone as fade-stage truth; Example Timing Score™ vanity NS; always-example or always-struggle-first brand; black-box perfect-fade AI.

Surviving: SAFE-EXPTIME — grain-local example↔solve timing; completion bridge; adaptive>fixed when evidence exists; rapid rechecks beat Likert destiny; densifies SAFE-FADE; instrument EXPTIME-* before universal E0→E3 / Timing Score™ campaigns (C).

Kill #55 (v1.14): “SRL Score™ / reflection-streak / survey≡event / streak≡SRL ≡ learning”

Destroyed: SRL Score™ / reflection-streak as North Star; trait survey SRL equated to event microanalysis; streak≡self-regulation; forethought without control writeback; ability-shame reflection; metacognition-module ACT

guarantees.

Surviving: SAFE-SRL — Zimmerman forethought→performance→reflection with phase→control writeback; microanalysis>trait surveys; process attributions; densifies SAFE-MONITOR/AAR/ATTN/HELP; instrument SRL-* before SRL Score™ / streak≡SRL / metacognition-module ads (CI).

Kill #56 (v1.14): “Generative Learning Score™ / AI≡generation / teach-for-XP ≡ learning”

Destroyed: Generative Learning Score™ as NS; AI monologue labeled as student generation; verbatim summary theater; unguided doodle≡dual-coding; teach-for-XP; eight-strategy menu bombs; ACT guarantees from GLA branding.

Surviving: SAFE-GENERATE — SOI-forcing sparse generative asks; student produces (coach does not substitute); teach-back requires the explanation act; guided FormatId draw/map; densifies SAFE-EXPLAIN/DUAL/SRL/SE; instrument GENACT-* before Generative Learning Score™ / AI≡generation / eight-strategy ACT ads (CII).

Kill #57 (v1.14): “Psychological Safety Score™ / poster safe-space / Bermuda-triangle ≡ learning”

Destroyed: Psychological Safety Score™ / Comfort % as NS; poster-only safe space; Bermuda-triangle redirect away from diagnosis; shame reel; praise-fog AI≡climate; destake≡no-standards; ACT guarantees from climate branding.

Surviving: SAFE-ERRCLIMATE — error climate > safety slogan; stay-diagnose-scaffold-reattempt; destake learning episodes without deleting FEI standards; adaptive action metrics over comfort surveys; densifies SAFE-MISCON/HELP/TALK/AAR/EXAM/PDASH; instrument ERRCLIM-* before Psychological Safety Score™ / poster safe-space / Bermuda-triangle / praise-fog ACT ads (CIII).

Kill #58 (v1.14): “Homework Completion Score™ / tonight-% / Chegg-dump / assisted≡ready ≡ learning”

Destroyed: Homework Completion Score™ / tonight-% as NS; “we finish your worksheet” brand; assisted≡solo-ready; Chegg-clone dump hero; Bastani Base Solver ads; purity ban on all homework help; parent homework stalk; grade guarantees from Solver.

Surviving: SAFE-HWHELP — dual-rail Help vs Practice; instrumental unstuck under deadline; assisted ≠ unassisted; costly dump; help recruits solo proof; densifies SAFE-HELP/FADE/ROI/DURABLE/PROOF; instrument HWHELP-* before Homework Completion Score™ / Chegg-clone / Bastani Base / purity-ban ACT ads (CIV).

Kill #59 (v1.14): “Embodiment Score™ / styles meshing / always-gesture / pose-camera ≡ learning”

Destroyed: Embodiment Score™ / Gesture Score™ / Motion Minutes as NS; kinesthetic-style meshing; always-add-gesture dogma; manipulative≡gesture≡abstraction; forced produce-only; pose-camera embodiment AI; AR body immersion≡FEI; ACT guarantees from gesture kits.

Surviving: SAFE-GESTURE — strategy-tied optional gesture/finger-trace/draw; gesture ≠ object-action for transfer; respect non-gesturers; gesture+language; FormatId-triggered; tutor model→fade; densifies SAFE-DUAL/GENERATE/FORMAT/APPRENTICE; instrument GESTURE-* before Embodiment Score™ / styles meshing / pose-camera / AR-body ACT ads (CV).

Kill #60 (v1.14): “Peer Explanation Score™ / teach-for-XP / Discord≡building / AI≡peer teaching ≡ learning”

Destroyed: Peer Explanation Score™ / Teach Minutes as NS; teach-for-XP; knowledge-telling≡quality; Discord≡knowledge-building; AI monologue≡peer teaching; unlock-without-solo after explain; peer-tutoring ACT guarantees.

Surviving: SAFE-PEERX — knowledge-building over telling; inferential prompts; apply after explain; capture student explanation act; near-peer grain+QA; densifies SAFE-GENERATE/EXPLAIN/HELP/TUTORGRAIN/APPRENTICE; instrument PEERX-* before Peer Explanation Score™ / teach-for-XP / Discord≡building / AI-monologue≡peer teaching ACT ads (CVI).

Kill #61 (v1.14): “Curiosity Score™ / cliffhanger-without-commit / postdiction / guess-for-XP ≡ learning”

Destroyed: Curiosity Score™ / Engagement-from-curiosity as NS; cliffhanger without prediction commit; postdiction theater; instant-reveal entertainment; guess-for-XP; always-predict dogma; curiosity ACT guarantees.

Surviving: SAFE-PREDICT — commit before reveal; prediction-induced curiosity and surprise for expectancy-violating joins; encode the repair; demo/figure predict; mode gates; densifies SAFE-FBTIME/MONITOR/RETRIEVE/MISCON/PF; instrument PREDICT-* before Curiosity Score™ / cliffhanger-without-commit / postdiction / guess-for-XP / curiosity ACT ads (CVII).

Kill #62 (v1.15): “Analogy Score™ / Mapping Minutes / mention≡pedagogy / surface-story≡structure ≡ learning”

Destroyed: Analogy Score™ / Mapping Minutes as NS; analogy-mention≡pedagogy; surface-story≡structure; AI likeness dump without student mapping; always-analogy dogma; HOT Score™; Discord≡supported compare; analogy ACT guarantees.

Surviving: SAFE-ANALOG — structure-map joins; co-present compare; student alignment; cognitive supports required; name where analogy breaks; densifies SAFE-BRIDGE/TRANSFER/DUAL/GESTURE/PREDICT; instrument ANALOG-* before Analogy Score™ / mention≡pedagogy / surface-story≡structure / AI likeness dump / always-analogy / analogy ACT ads (CVIII).

Kill #63 (v1.15): “Writing Score™ / Journal Minutes / freewrite≡learning / AI≡WTL ≡ learning”

Destroyed: Writing Score™ / Journal Minutes as NS; freewrite diary≡learning; longer-essay≡deeper; AI rewrite≡WTL; Notes scrapbook without FEI; always-journal dogma; writing ACT guarantees.

Surviving: SAFE-WRITE — short prompted cognitive+metacognitive writing-to-learn; Notes dual-rail with FEI; AI may cue not author; densifies SAFE-GENERATE/SE/EXPLAIN/SRL/PEERX/PROOF; instrument WRITE-* before Writing Score™ / Journal Minutes / freewrite≡learning / longer≡deeper / AI≡WTL / always-journal / writing ACT ads (CIX).

Kill #64 (v1.15): “Spatial Score™ / Mental Rotation Minutes / brain-training≡identity ≡ learning”

Destroyed: Spatial Score™ / Mental Rotation Minutes as NS; isolated brain-training≡math identity; innate or boys-better spatial costume; always-add-3D; Spatial Personality™; unscaffolded diagram flood; spatial ACT guarantees.

Surviving: SAFE-SPATIAL — FormatId-embedded spatial competence; conditional transfer; destake spatial anxiety; densifies SAFE-FORMAT/DUAL/GESTURE/TRANSFER/GENDER; instrument SPATIAL-* before Spatial Score™ / Mental Rotation Minutes / brain-training=identity / boys-better / Spatial Personality™ / spatial ACT ads (CX).

Kill #65 (v1.15): “Formative Score™ / Assessment Minutes / exit-ticket=learning / 0.4–0.7 AfL ads = learning”

Destroyed: Formative Score™ / Assessment Minutes as NS; exit-ticket=learning; Black–Wiliam 0.4–0.7 product ads; more-quizzes=formative; gradebook cosplay; “got it?”=closed cycle; AfL ACT guarantees.

Surviving: SAFE-FA — formative-as-process; Sadler gap-closing; no 0.4–0.7 product legend; exit tickets only if wired to next move; FEI over Formative Score™; densifies SAFE-HITL/AAR/FBTIME/MONITOR/SRL/INSTRUMENT; instrument FA-* before AfL marketing (CXI).

Kill #66 (v1.15): “Feedback Literacy Score™ / Feedback Minutes / delivery=uptake = learning”

Destroyed: Feedback Literacy Score™ / Feedback Minutes as NS; delivery=uptake; thumbs-up=used; longer=more literate; illiteracy-shame without design fix; always-more-feedback dogma; feedback-literacy ACT guarantees.

Surviving: SAFE-FBLIT — feedback literacy as sense+use; proactive recipience; uptake loop UX; densifies SAFE-FBTIME/FA/EXPLAIN/HINT/HELP/SRL/MONITOR/ERRCLIMATE; instrument FBLIT-* before Feedback Literacy Score™ / delivery=uptake / thumbs=used / Feedback Minutes ads (CXII).

Kill #67 (v1.15): “Collaboration Score™ / Group Minutes / Discord=CSCL / group-correct=ready = learning”

Destroyed: Collaboration Score™ / Group Minutes as NS; Discord=CSCL learning; group-correct=ready; free-rider-blind unlocks; always-collab dogma; collab ACT guarantees.

Surviving: SAFE-COLLAB — structured collab with group goals + individual accountability; mandatory solo FEI gate; densifies SAFE-PEERX/CoP/PROOF/HELP/INSTRUMENT/ERRCLIMATE; instrument COLLAB-* before Collaboration Score™ / Group Minutes / Discord=CSCL / group-correct=ready / always-collab / collab ACT ads (CXIII).

Kill #68 (v1.15): “Desirable Difficulty Score™ / Challenge Minutes / always-ease / always-max / cram=ready = learning”

Destroyed: Desirable Difficulty Score™ / Challenge Minutes as NS; always-ease-before-ACT; blocked-comfort=ready; always-max-DD-until-bell; last-48h scored flood-as-DD; cram-week=ready; schedule ACT-point guarantees.

Surviving: SAFE-DDSCHED — horizon-matched hold/ease of desirable difficulties (Build / Join / Protect+probe); stakes≠DD; ease load not readiness standards in final ~72h; densifies SAFE-DD/EXAM/SCHED/DURABLE/CRAM/RETRIEVE/EXPOSE; instrument DDSCHED-* before Desirable Difficulty Score™ / Challenge Minutes / always-ease / always-max / cram=ready / schedule ACT guarantees (CXIV).

Kill #69 (v1.15): “Representation Score™ / Diversity Minutes / costume=belonging / match=quality = learning”

Destroyed: Representation Score™ / Diversity Minutes as NS; token STEM costume⇒belonging; demographic match⇒quality; empty likeness praise⇒recognition; Belonging Score™; matching ACT guarantees; soft standards for matched pairs.

Surviving: SAFE-ROLEMODEL — fidelity before likeness; competence-visible attainable near-peers; optional preference match; recognition = practice evidence; densifies SAFE-GENDER/ELL/TUTORGRAIN/HITL/EXPECTANCY/CoP/IDMEASURE; instrument ROLEMODEL-* before Representation Score™ / Diversity Minutes / costume⇒belonging / match⇒quality / empty likeness praise / Belonging Score™ / matching ACT guarantees (CXV).

Kill #70 (v1.15): “Capture Score™ / Photo Minutes / OCR⇒mastery / photo→answer ≡ learning”

Destroyed: Capture Score™ / Photo Minutes as NS; OCR⇒mastery; always-on camera; photo→answer Practice hero; Capture ACT guarantees; phone-companion focus theater.

Surviving: SAFE-OFFLINE — offline-first attempt grain; capture after commit; residue-aware brief ritual; verify OCR before Map write; no photo→answer dump; densifies SAFE-ATTN/HELP/HWHELP/GESTURE/PROOF/GENQ/FORMAT; instrument OFFLINE-* before Capture Score™ / Photo Minutes / OCR⇒mastery / always-camera / photo→answer hero / Capture ACT guarantees (CXVI).

Kill #71 (v1.16): “Forever-E0 / day-one fade flood / Anxiety·Calm Score™ fade gates ≡ learning”

Destroyed: Forever-E0 anxious track; day-one aggressive fade flood for HMA; Anxiety Score™ or Calm Score™ fade gates; Fade Score™ NS; permanent peek “support”; soft standards for anxious personas; protect-without-fade ChatGPT dump; ACT guarantees from anxiety-aware scaffolding.

Surviving: SAFE-FADEX — protective worked study under anxiety WM load; longer E0/E1 hold without permanent persona track; destake; treat peek binge as safety behavior to fade; proof = solo transfer + approach; densifies SAFE-FADE/EXPOSE/DD; instrument FADEX-* before forever-E0 / flood / Calm Score™ fade campaigns (CXVII).

Kill #72 (v1.16): “Miss Streak / soft-wrong shame push / empty mindset parent letter ≡ learning”

Destroyed: Miss Streak / soft-wrong shame push; live wrong-answer stalk; Parent Correctness Score™; empty “mistakes are good!” without grain+CTA; controlling “sit until correct” default; growth-mindset parent letter⇒FEI; ACT guarantees from parent messaging.

Surviving: SAFE-SOFTMSG — grain + failure-enhancing frame + one improvement CTA; autonomy>control; densifies SAFE-PDASH/FBLIT/ERRCLIMATE/MISCON; instrument SOFTMSG-* before Miss Streak / shame-push / empty mindset parent campaigns (CXVIII).

Kill #73 (v1.16): “Always-on Map / Map Minutes / mid-attempt full-graph ≡ learning”

Destroyed: Always-on student Map; Map Minutes or Map Score™ NS; mid-attempt full-graph dump; Map-open⇒cycle-closed; dual-screen Map+problem default; ontology tour⇒tutoring; ACT guarantees from Map airtime.

Surviving: SAFE-MAPORCH — brief→grain-slice→adjust; tutor-private default; between-attempt opens; densifies SAFE-HITL/FA/ATTN/BRIDGE/ONTOLOGY; instrument MAPORCH-* before always-on Map / Map Minutes /

mid-attempt dump campaigns (CXIX).

Kill #74 (v1.16): “Burnout Score™ / Wellness Minutes / wellness-primary under caseload ≡ quality”

Destroyed: Burnout Score™ / Wellness Minutes NS; mandatory mood gate; tutor emotion-AI; wellness-module-as-primary when caseload is the demand; retention-% without fidelity; sell tutor-affect dashboards; ACT guarantees from wellness packaging.

Surviving: SAFE-BURNWARN — fidelity/load early-warn → org redesign (JD-R); densifies SAFE-WORKFORCE/PRIVACY/ROI; instrument BURNWARN-* before Burnout Score™ / wellness-primary campaigns (CXX).

Kill #75 (v1.16): “Fluency≡pedagogy / Explanation·Rubric Score™ / LLM-judge-alone ≡ learning”

Destroyed: Fluency/human-likeness≡pedagogy; thumbs-up thoroughness NS; Explanation Score™ or Rubric Score™; LLM-judge-alone ship authority; soft-pass wrong math because prose “helped”; reveal-answer default as guidance; auto-wrap-QA replacing GENQ key verify; ACT guarantees from rubric-scored AI coach.

Surviving: SAFE-EXPLAINQA — principle-short rubric gates; hard-fail math/key; densifies SAFE-EXPLAIN/GENQ/REPAIR; instrument EXPLAINQA-* before fluency≡pedagogy / LLM-judge-alone campaigns (CXXI).

Kill #76 (v1.16): “Soft-wrong≡half-credit / Partial Credit Score™ / all-wrongs-equal ≡ learning”

Destroyed: Soft-wrong≡half-credit; Partial Credit Score™ or Mark Minutes NS; all-wrongs-equal mastery deltas; gradebook cosplay (% banners, letter grades, live point stalk); ordered PCM steps≡identity; empty celebrate-mistakes without mark→route; ACT guarantees from partial-credit packaging.

Surviving: SAFE-TAXON — typed marks drive routing and type-conditioned graph updates; soft-wrong wraps type; densifies SAFE-MISCON/FBLIT/FBTIME; instrument TAXON-* before Partial Credit Score™ / gradebook cosplay campaigns (CXXII).

Kill #77 (v1.16): “Hours-booked crunch / Load Shed Score™ / pour≡exam support ≡ quality”

Destroyed: Hours-booked / session-volume crunch NS; Load Shed Score™ or Crunch Minutes; infinite-availability exam brand; wellness-primary under caseload spike; hire-surge without coach bandwidth; pour-mode≡exam support; hospital-cosplay ads; soft standards for crunch personas; ACT guarantees from crunch playbooks.

Surviving: SAFE-LOADSHED — named exam-week pause list; acuity FEI minutes; org redesign first; densifies SAFE-BURNWARN/DDSCHED/WORKFORCE/EXAM/ROI/HWHELP; instrument LOADSHED-* before hours-booked crunch / infinite-availability campaigns (CXXIII).

Kill #78 (v1.16): “Distractor Score™ / always-four / beauty traps / unmapped diagnostic MC ≡ learning”

Destroyed: Distractor Score™ / Option Minutes; always-four≡better; beauty traps·absurd fillers·NOTA-AOTA default; unmapped wrongs≡diagnostic; LLM distractors without mis_ seal; option-count or item-count NS; ACT guarantees from distractor packaging.

Surviving: SAFE-DISTRACTOR — misconception-mapped functioning distractors; function over fashion; soft-wrong from chosen option; densifies SAFE-TAXON/MISCON/GENQ; instrument DISTRACTOR-* before Distractor Score™ / always-four / beauty-trap campaigns (CXXIV).

Kill #79 (v1.16): “Identity Score™ / Confidence Personality™ / equal-weight recovery / inflate-confidence ≡ learning”

Destroyed: Identity Score™ / Confidence Personality™ from hard/kinda/easy; equal-weight recovery for all misses; easy≡mastered; inflate-confidence as recovery NS; Belief Score™; CA-without-route; ACT guarantees from confidence-tagged recovery packaging.

Surviving: SAFE-CONFMISS — confidence×type recovery dose; CRI fork; hypercorrection priority for sure misses; densifies SAFE-TAXON/CALIB/MONITOR/COLD; instrument CONFMISS-* before Identity Score™ / equal-weight recovery / inflate-confidence campaigns (CXXV).

Kill #80 (v1.17): “Skip-as-default / lecture-as-debrief / Microdebrief Score™ / +25% ACT from debriefs ≡ learning”

Destroyed: “No time for debrief” default; lecture-as-micro-debrief; Microdebrief Score™·Debrief Minutes·Reflection Streak NS; +25% ACT from borrowed debrief metas; military-cosplay micro-AAR; vibes-only close without traces; therapy-length mandate under crunch; ACT guarantees from micro-debrief packaging.

Surviving: SAFE-MICRODEB — capped student-first micro-AAR with objective traces; protect under LOADSHED; densifies SAFE-AAR/LOADSHED/HITL/TALK/FBLIT; instrument MICRODEB-* before skip-as-default / lecture-close / Microdebrief Score™ campaigns (CXXVI).

Kill #81 (v1.17): “24/7 availability / hours-hero / fear-parent emails / Availability Score™ ≡ care”

Destroyed: 24/7·infinite-availability default parent value prop; hours-booked crunch hero; fear-appeal parent emails as default; silent shed≡care; Likert-WTP for “thinking minutes”; Availability Score™·Crunch Message Score™·Open-Late Minutes NS; ACT guarantees from crunch messaging.

Surviving: SAFE-CRUNCHMSG — shed-honesty + protected thinking-minutes CBC/messages; densifies SAFE-LOADSHED/WTP/SOFTMSG/PDASH/EXAM; instrument CRUNCHMSG-* before 24/7 / hours-hero / fear-parent campaigns (CXXVII).

Kill #82 (v1.17): “FormatTrap Score™ / seductive stems / beauty visual traps / unmapped visual diagnostic ≡ learning”

Destroyed: FormatTrap Score™ / Visual Distractor Minutes; seductive stem art≡FormatId; beauty-trap option pictures; unmapped visual wrongs≡diagnostic; decoration≡dual-coding distractors; LLM figure options without trap seal; visual miss≡topic-only; ACT guarantees from visual-distractor packaging.

Surviving: SAFE-FORMTRAP — named format-trap visual distractors; coherent structural stems; densifies SAFE-DISTRACTOR/FORMAT/DUAL/GENQ; instrument FORMTRAP-* before FormatTrap Score™ / beauty visual trap campaigns (CXXVIII).

Kill #83 (v1.17): “No time for brief / pour-open / Brief Score™ / late-join shame / I-PASS cosplay ≡ learning”

Destroyed: “No time for brief/goals” default; pour-open lecture-as-brief; Brief Score™·Brief Minutes·Opening Streak NS; late-join shame open; aviation·OR·I-PASS cosplay ACT ads; vibes-only open without named edge; soft standards via “we connected”; ACT guarantees from micro-brief packaging.

Surviving: SAFE-MICROBRIEF — capped shared-model open (edge/success/trap/role + student synthesis); densifies SAFE-AAR/MICRODEB/LOADSHED/HITL/EXPECTANCY; instrument MICROBRIEF-* before no-time-for-brief / pour-open / Brief Score™ campaigns (CXXIX).

Kill #84 (v1.17): “Chegg-hero crunch CTA / Dual-CTA Score™ / assisted⇒ready ⇒ learning”

Destroyed: Single Chegg-hero Solver CTA under crunch; “finish tonight’s homework” primary brand; fear-countdown→dump CTAs; Dual-CTA Score™·Help Tonight Minutes·Crunch Click NS; silent Practice-rail collapse in exam week; assisted⇒ready from Solver CTA; ACT guarantees from dual-rail CTA packaging.

Surviving: SAFE-CRUNCHCTA — dual labeled Solver/Practice CTAs under crunch; instrumental verbs; fiduciary defaults; follow-on Practice; densifies SAFE-HWHELP/CRUNCHMSG/HELP/DURABLE/SOFTMSG; instrument CRUNCHCTA-* before Chegg-hero / Dual-CTA Score™ / fear-dump campaigns (CXXX).

Kill #85 (v1.17): “Fluency SVG⇒keyed figure / judge-alone figure ship / Figure Key Score™ ⇒ learning”

Destroyed: Fluency SVG / pretty-render⇒keyed figure; LLM-as-judge-alone figure ship; Figure Key Score™·Visual QA Minutes; unverified FormatId figures as diagnostic ground truth; FormatId-count or “AI drew N diagrams” hero; soft-pass wrong geometry because traps were mapped; ACT guarantees from figure-generation packaging.

Surviving: SAFE-FIGKEY — structural figure key-verify for FormatId generation; deterministic spine; use gating; densifies SAFE-GENQ/FORMTRAP/DUAL/COVER/EXPLAINQA; instrument FIGKEY-* before fluency SVG / judge-alone / Figure Key Score™ campaigns (CXXXI).

Kill #86 (v1.17): “Idle spinner / pour-open wait / Wait Minutes·Lobby Score™ / entertainment lobby ⇒ learning”

Destroyed: Idle spinner-only late lobby; pour-open AI wait dump; Wait Minutes·Lobby Score™·Prebrief Streak NS; entertainment-primary fill; forced marathon Practice-before-join; late-tutor shame feed; ACT guarantees from async pre-brief packaging.

Surviving: SAFE-ASYNCBRIEF — capped async Set card on late/early wait; session-related occupied wait; if-then chip; live verify-not-replace; densifies SAFE-MICROBRIEF/SRL/ATTN/LOADSHED/EXPECTANCY; instrument ASYNCBRIEF-* before idle spinner / Lobby Score™ / pour-open wait campaigns (CXXXII).

Kill #87 (v1.17): “Dump-and-quit / maze handoff / Transfer Score™ / assisted⇒ready ⇒ learning”

Destroyed: Dump-and-quit Solver close; multi-step maze handoff; Transfer Score™·Handoff Minutes·Follow-on Streak NS; assisted⇒ready; marathon-after-soft-hint; XP-for-help-without-solo; ACT guarantees from transfer-grain packaging.

Surviving: SAFE-XFERGRAIN — one-tap edge-matched Practice grain after instrumental Solver; short default; solo gate; SCHED return; densifies SAFE-HWHELP/CRUNCHCTA/SCHED/HELP/FADE/PROOF; instrument XFERGRAIN-* before dump-and-quit / Transfer Score™ / maze handoff campaigns (CXXXIII).

Kill #88 (v1.18): “Tagged FormatId≡C4 / fluency eligibility / coverage-padding / HideCorrectness Score™ ≡ learning”

Destroyed: Tagged-FormatId≡C4-eligible; fluency-render eligibility; coverage-padding unsealed visuals into gap-scan; HideCorrectness Score™·Eligibility Minutes·FormatId-count-in-scan NS; miss-on-broken-figure≡format-gap without seals; ACT guarantees from visual diagnostic packaging.

Surviving: SAFE-HCELIG — seals-first FormatId eligibility for C4; hcelig_pass distinct from in-bank; honest holes > pad; densifies SAFE-FIGKEY/GENQ/CALIB/COLD/FORMAT/COVER/FORMTRAP; instrument HCELIG-* before tagged≡C4 / fluency eligibility / HideCorrectness Score™ campaigns (CXXXIV).

Kill #89 (v1.18): “Soon-default / optimistic ETA slip / apology-primary / ETA Score™ ≡ care”

Destroyed: “Soon”/“almost” default late lobby; optimistic ETA slip loops; apology-primary without duration; ETA Score™·Wait Satisfaction·Lobby Minutes NS; student-facing tutor-shame feed; fake-precision countdown; ACT guarantees from wait-UX packaging.

Surviving: SAFE-KNOWNETA — known-finite ETA with refresh honesty; densifies SAFE-ASYNCBRIEF/MICROBRIEF/LOADSHED/EXPECTANCY; instrument KNOWNETA-* before soon-default / ETA Score™ / shame-feed campaigns (CXXXV).

Kill #90 (v1.18): “Universal hard-gate / Gate Score™ / hint-tax shame ≡ learning”

Destroyed: Universal hard-gate after every hint; optional-only dump chip as sole close; Gate Score™·Forced Follow-On %·Soft-Require Compliance Minutes NS; “can’t leave until Practice” hint-tax shame; ignore help-avoidance when Forced % rises; ACT guarantees from soft-vs-hard gate packaging.

Surviving: SAFE-SOFTGATE — soft-require default; conditional short hard-gate with avoidance falsifiers; densifies SAFE-XFERGRAIN/HELP/CRUNCHCTA; instrument SOFTGATE-* before Gate Score™ / universal hard-gate campaigns (CXXXVI).

Kill #91 (v1.18): “Unified bank ≡ Practice+C4 / silent dual-write / SplitPool Score™ ≡ learning”

Destroyed: Unified sealed bank ≡ Practice+C4; silent dual-write Practice→Map; auto-promote from play counts; SplitPool Score™·Dual-Use Minutes·Pool-Separation % NS; padding C4 with Practice-only visuals; ACT guarantees from one-bank packaging.

Surviving: SAFE-SPLITPOOL — graded Practice-reveal vs C4 pools; ticketed promote; densifies SAFE-HCELIG/REPAIR/GENQ/COVER; instrument SPLITPOOL-* before unified dual-use / SplitPool Score™ campaigns (CXXXVII).

Kill #92 (v1.18): “Pad-scan / equal-quota theater / Composition Score™ ≡ learning”

Destroyed: Symbolic/Practice pad into C4 FormatId blueprint cells; equal-quota theater under N≈1 sealed stock; silent underrepresentation sold as format-flexible; Composition Score™·Format Balance %·Scan Completeness Minutes NS; “diagnose every format” ads under thin C4 stock; absence≡mastery; ACT guarantees from balanced-scan packaging.

Surviving: SAFE-SCANCOMP — blueprint composition from sealed stock with labeled holes; densifies SAFE-HCELIG/SPLITPOOL/COVER/FORMAT; instrument SCANCOMP-* before pad / Composition Score™ campaigns (CXXXVIII).

Kill #93 (v1.18): “Hide-holes / complete-visual ads / Hole Trust Score™ ≡ trust”

Destroyed: Hide-format_not_probed for conversion; child-deficit visual-hole shame; “complete visual diagnostic” ads while holes remain; Hole Trust Score™·Coverage Satisfaction·Honesty Minutes NS; absence≡format mastery; fear-scarcity hole WTP; ACT guarantees from hole packaging.

Surviving: SAFE-HOLETRUST — label hole + seal reason + one fill CTA; densifies SAFE-SCANCOMP/PDASH/COVER/WTP/SOFTMSG; instrument HOLETRUST-* before hide-holes / Hole Trust Score™ campaigns (CXXXIX).

Kill #94 (v1.18): “Play-count auto-promote / Promote Score™ ≡ bank quality”

Destroyed: Play-count / thumbs-up auto-promote to C4; silent hcelig_pass flip; fluency soft-pass at promote; Promote Score™·QA Minutes·Promotion Throughput % NS; padding C4 while critical REPAIR open; ACT or complete-visual guarantees from promote packaging.

Surviving: SAFE-PROMOTE — ticketed checklist + seal re-verify; demote faster than promote; densifies SAFE-SPLITPOOL/REPAIR/FIGKEY/HCELIQ/COVER; instrument PROMOTE-* before play-count / Promote Score™ campaigns (CXL).

Kill #95 (v1.18): “Immediate same-stem C4 / Cool-Down Score™ ≡ diagnosis”

Destroyed: Immediate same-stem C4 after reveal; forever or global cool-down lock as default; Cool-Down Score™·Exposure Minutes·Lock Duration % NS; pad under cool-down; isomorphic-clone leak as “fresh”; ACT or complete-visual guarantees from cool-down packaging.

Surviving: SAFE-COOLDOWN — per-student cool-down after reveal before same-family C4; isomorphic substitutes then holes; densifies SAFE-SPLITPOOL/SCANCOMP/CALIB/COLD; instrument COOLDOWN-* before immediate same-stem / Cool-Down Score™ campaigns (CXLI).

Kill #96 (v1.18): “Parent-only hole honesty / fake greens / Student Hole Score™ ≡ identity”

Destroyed: Parent-only hole honesty; fake greens for not_probed; ability-shame hole frames; empty growth-mindset poster without CTA; Student Hole Score™·Completeness %·Honesty Minutes NS; absence≡mastery; ACT or identity guarantees from student hole packaging.

Surviving: SAFE-STUDHOLE — student parity honesty; product-hole frame; one Practice CTA; no fake greens; densifies SAFE-HOLETRUST/COLD/CALIB/SOFTMSG/FBLIT; instrument STUDHOLE-* before parent-only / fake-green / Student Hole Score™ campaigns (CXLII).

Kill #97 (v1.19): “Coming-soon / play-count fill bars / Filleta Score™ ≡ trust”

Destroyed: Coming-soon / almost-ready default hole-fill; play-count fill bars; optimistic slip loops; date-driven soft-pass promote; Filleta Score™·Backlog Minutes·Honesty Clock NS; ACT or complete-visual-by-date guarantees from fill packaging.

Surviving: SAFE-FILLETA — capacity-grounded finite fill bands; no play-count progress; no date soft-pass seals; densifies SAFE-PROMOTE/KNOWNETA/HOLETRUST/STUDHOLE; instrument FILLETA-* before coming-soon / Filleta Score™ campaigns (CXLIII).

Kill #98 (v1.19): “Always-dual / never-dual / Dual Rater Score™ ≡ quality”

Destroyed: Always-dual costume; never-dual figured C4; Dual Rater Score™·Agreement Minutes·Kappa % NS; seal-waiving second rater; mammography-cosplay ACT ads; ETA-skipped mandatory dual; Discord-emoji≡second rater.

Surviving: SAFE-DUALRATE — risk-triggered second rater; seals before dual; disagreement→reject/repair; densifies SAFE-PROMOTE/FIGKEY/HITL/FILLETA; instrument DUALRATE-* before Dual Rater Score™ / always-dual campaigns (CXLIV).

Kill #99 (v1.19): “Item_id-only / concept-wide / Family Hash Score™ ≡ cool-down”

Destroyed: Item_id-only cool-down under GENQ clones; concept×FormatId cool-down as default family; embedding-only family without SME seal; Family Hash Score™·Clone Minutes·Similarity % NS; pad when family bar empties a cell; ACT or complete-visual guarantees from hashing packaging.

Surviving: SAFE-FAMHASH — sealed isomorphic family as cool-down compromise unit; substitutes outside family; densifies SAFE-COOLDOWN/GENQ/FIGKEY/SCANCOMP; instrument FAMHASH-* before item_id-only / Family Hash Score™ campaigns (CXLV).

Kill #100 (v1.19): “Anxiety·Calm·Exit Emotion Score™ / consoling comfort ≡ care”

Destroyed: Anxiety Score™ / Calm Score™ / Exit Emotion Score™ / Mood Minutes NS; consoling “not everyone can...” ability talk; ego Completeness % beside hole honesty; mandatory mood gate; emotion-AI exit empathy; fake greens as affect management; calm-to-zero or ACT guarantees from gap-scan exit packaging.

Surviving: SAFE-EXITEMO — destake exit copy; task/process + one Practice CTA; no mood telemetry; densifies SAFE-STUDHOLE/PRIVACY/EXPOSE/SOFTMSG; instrument EXITEMO-* before Exit Emotion Score™ / emotion-AI exit campaigns (CXLVI).

Kill #101 (v1.19): “Dual-omitted FILLETA / ETA-skipped dual / Capacity Score™ ≡ honesty”

Destroyed: Dual-omitted FILLETA; ETA-skipped dual; always-dual without dual FTE; 100% dual-util plans; Dual Rater Score™·Capacity Score™·Reviewer Minutes NS; hospital-cosplay dual-capacity ads; complete-visual-by-date or ACT guarantees from dual capacity packaging.

Surviving: SAFE-CAPPLAN — dual minutes as first-class FILLETA capacity line; headroom; label dual scarcity; densifies SAFE-DUALRATE/FILLETA/WORKFORCE; instrument CAPPLAN-* before ETA-skipped dual / Capacity Score™ campaigns (CXLVII).

Kill #102 (v1.19): “Forever-dual clones / silent inherit / Template Trust Score™ ≡ quality”

Destroyed: Forever-dual on sealed isomorphic clones after blueprint dual; silent inherit without template+family seals; Template Trust Score™·Inherit %·Family Calibration Score™ NS; first dual ≡ forever trust without sample audit; embedding-only inheritance; FERF≡C4-ready ads; ACT or complete-visual guarantees from template-inheritance packaging.

Surviving: SAFE-TMPTRUST — blueprint dual once; inherit single-rater under sealed sameness + sample audit + revoke; densifies SAFE-DUALRATE/FAMHASH/FIGKEY/CAPPLAN; instrument TMPTRUST-* before forever-dual / Template Trust Score™ campaigns (CXLVIII).

Kill #103 (v1.19): “Kappa % NS / pooled repair vanity / Landis-band marketing ≡ integrity”

Destroyed: Kappa % / Dual Rater Score™ / Agreement Minutes NS; pooled-only post-promote repair vanity; Landis–Koch “almost perfect” marketing band; unconditioned dual-shame boards; mammography-cosplay IRR heroes; ACT or complete-visual guarantees from IRR packaging.

Surviving: SAFE-REPAIRSTRAT — path-stratified repair + contingency drill-downs; κ as aide not NS; densifies SAFE-DUALRATE/PROMOTE/INSTRUMENT/TMPTRUST; instrument REPAIRSTRAT-* before Kappa % / Dual Rater Score™ IRR campaigns (CXLIX).

Kill #104 (v1.19): “Surge Hire % / unsealed freelancer dual / hospital-surge cosplay $\equiv$ capacity”

Destroyed: Dual Rater Score™ / Surge Hire % / Contract Minutes NS; hire-surge theater; unsealed Discord-freelancer $\equiv$ second rater; always-hire or always-contract dogma; hospital-surge-cosplay dual ads; complete-visual-by-exam-week or ACT guarantees from surge packaging.

Surviving: SAFE-SURGEPANEL — make-or-buy dual portfolio with seals, parallel-run, demobilization honesty; densifies SAFE-CAPPLAN/LOADSHED/WORKFORCE/DUALRATE; instrument SURGEPANEL-* before Surge Hire % / unsealed freelancer dual campaigns (CL).

Part XV — References (verified starting set)

Do not invent citations. Expand this list only with retrieved sources.

1. Yeager, D. S., et al. (2019). A national experiment reveals where a growth mindset improves achievement. *Nature*. <https://doi.org/10.1038/s41586-019-1466-y>
2. Dweck, C. S. (2006/2007). *Mindset*. Random House.
3. Ryan, R. M., & Deci, E. L. (2000). Self-determination theory... *American Psychologist* / later handbooks.
4. Wang, Y., et al. (2024). Meta-analysis of SDT-based interventions in education. (Self-Determination Theory site PDF).
5. Kulik, C.-L. C., Kulik, J. A., & Bangert-Drowns, R. L. (1990). Effectiveness of mastery learning programs: A meta-analysis. *Review of Educational Research*.
6. Slavin, R. E. (1987). Mastery learning reconsidered. *Review of Educational Research*.
7. Boaler, J., et al. — mathematical mindset course / classroom studies (Frontiers Education 2018; later scale papers).
8. Systematic review: mindset interventions in mathematics classrooms (2023). *Educational Research Review*, 100554.
9. Johnston-Wilder, S., & Lee, C. — mathematical resilience / SDT framing (2021).
10. Duolingo (2023). *The Duolingo Method* whitepaper.
11. Bloom, B. S. (1984). The 2 sigma problem. *Educational Researcher*.
12. Ashcraft, M. H., & Kirk, E. P. (2001). The relationships among working memory, math anxiety, and performance. *Journal of Experimental Psychology: General*.
13. Suárez-Pellicioni, M., Núñez-Peña, M. I., & Colomé, À. (2016). Math anxiety: A review of its cognitive consequences... *Cognitive, Affective, & Behavioral Neuroscience*.
14. Caviola, S., et al. / related metas — WM mediating role in math anxiety–performance (e.g., *Frontiers in Psychology*, 2021 meta, $r \approx -0.17$ MA–MP).

15. Kirschner, P. A., Sweller, J., & Clark, R. E. (2006). Why minimal guidance during instruction does not work. *Educational Psychologist*.
16. Walton, G. M., & Cohen, G. L. — belonging interventions tradition.
17. VanLehn, K. (2011). The relative effectiveness of human tutoring, intelligent tutoring systems... *Educational Psychologist*.
18. Nickow, A., Oreopoulos, P., & Quan, V. (2020). Tutoring RCT meta-analysis (effects ~0.37 SD — see Education Next summary).
19. Hambrick, D. Z., et al. (2014). Deliberate practice: Is that all it takes... *Intelligence*; Macnamara, Hambrick & Oswald (2014) deliberate practice meta-analysis.
20. Ericsson, K. A., Krampe, R. T., & Tesch-Römer, C. (1993). The role of deliberate practice... *Psychological Review*.
21. Howard, J. L., Slemp, G. R., & Wang, X. (2024/2025). Need support and need thwarting meta-analysis (student populations).
22. MindCraft internal: WORLD_VISION.md, BRAND_BOOK.md, MindCraft_Viability_Research_Strategy.md.

Part XVI — Mechanism Deep Dive: How identity actually changes

This part answers the research question at mechanism grain. It is synthesis, not settled science.

XVI.1 What “confident mathematical thinker” means (operational)

Do not use: “likes math apps” or “says math is fun once.”

Working definition (HYPOTHESIS): A confident mathematical thinker is a person who, in math-relevant contexts:

1. **Approaches** moderately hard problems without immediate avoidance.
2. **Persists** after a miss long enough to extract information.
3. **Attributes** success primarily to controllable causes (strategy, practice, help-seeking) rather than luck or tutor magic.
4. **Selects** into harder future opportunities when available (courses, contests, careers) at rates above their prior baseline.
5. **Endorses** a math-capable self-concept *and* behaves consistently with that endorsement.

Items 1–4 are leading. Item 5 is lagging. Product should optimize leading indicators while periodically measuring the lagging identity claim.

XVI.2 The threat–competence–narrative triad

Three literatures converge on the same transformation engine:

Mechanism	Core claim	Confidence	Key sources
-----------	------------	------------	-------------

Threat / affect	Math anxiety transiently reduces working-memory resources available for math	High (FACT direction)	Ashcraft & Kirk (2001); Suárez-Pellicioni et al. review (2016); WM mediation meta (2021)
Competence evidence	Motivation requires felt competence; mastery structures can raise achievement when mastery is real	Medium–High	SDT (Ryan & Deci); Kulik mastery meta; Slavin critique
Narrative / identity	Self-stories and belonging cues shape challenge-seeking and persistence	Medium	Yeager et al. (2019); Walton/Cohen belonging tradition; narrative psychology

HYPOTHESIS — Triad necessity: For students starting from “I am bad at math,” improving any one factor alone is usually insufficient:

- Safety without competence → comfort without skill → identity still fragile.
- Competence without safety → possible for some; many never enter the practice volume.
- Narrative without either → empty affirmation (false growth mindset).

XVI.3 Founder tutoring sequence — audit

FOUNDER BELIEF (observed in tutoring):

1. Reduce fear
2. Demystify mathematics
3. Explain why humanity invented it
4. Simplify the narrative
5. Create a tiny success
6. Reinforce success
7. Celebrate progress → identity shifts

Red Team audit:

Step	Universal?	Evidence status	Failure mode
1 Fear ↓	Common for anxious subset; not for all	Strong for anxious learners	Over-soothing that removes challenge
2 Demystify	Often helpful	Medium (cognitive clarity)	Over-simplification that lies
3 Invention story	Founder prior	Under-evidenced at RCT scale	Cultural mismatch; time cost
4 Simplify narrative	Often helpful	Medium	Removes necessary difficulty

5 Tiny success	Near-universal lever	High (competence need)	Too-easy wins → hollow identity
6 Reinforce	High if attribution-correct	High (feedback literature)	Praise of talent or praise of empty effort
7 Celebrate	Helpful if genuine	Medium	Performative celebration → cynicism

Improved sequence (HYPOTHESIS — “SAFE-CRAFT”):

1. Safety: threat appraisal ↓ (tone, privacy, soft-wrong, no public shame)
2. Attempt: student acts before receiving a full lecture (productive struggle at right grain)
3. Feedback-as-information: miss becomes diagnosis, not character
4. Earned micro-win: success the student can own
5. Causal story: “what worked” (strategy), optionally “why humans invented this tool”
6. Re-entry: immediate second attempt / next slightly harder item
7. Attribution + witness: tutor/system names the growth specifically
8. Future self: map fill / path unlock that marks identity-relevant progress
9. Transfer check: varied item so win was not clone memorization

What changed vs founder sequence: Attempt moves earlier; invention-story becomes optional enrichment inside step 5, not a prerequisite to agency; transfer check prevents fake mastery.

XVI.4 Math anxiety — what is known, what is not

FACT: Higher math anxiety correlates with lower math performance (meta $r \approx -0.17$ in one 2021 synthesis of 57 studies — modest average, not destiny).

FACT: Ashcraft & Kirk (2001) show math anxiety can act like a dual-task load: worry competes for working memory, especially on computation-span / carry-heavy tasks.

HYPOTHESIS: In product UX, anything that increases social-evaluative threat (red buzzers, public leaderboards for novices, tutor contempt, parent hovering UI) will shrink effective WM and look like “they don’t get it.”

Contradictions / limits:

- Not all underperformance is anxiety; some is missing knowledge, poor curriculum, sleep, language barriers.
- Trait vs state anxiety differ; trait may partly reflect prior failures (reverse causality possible).
- High performers can be anxious and still succeed via avoidance of hard electives — a hidden identity cost.

MindCraft rule: Treat anxiety as a first-class load on the learning system, not as a soft skill sidebar.

XVI.5 Deliberate practice — necessary, not sufficient

FACT / TRADITION: Ericsson’s deliberate practice emphasizes structured, feedback-rich, effortful practice aimed at improvement (not mere repetition).

FACT / CRITIQUE: Hambrick and colleagues argue deliberate practice is important but not sufficient; in chess/music reanalyses, deliberate practice often explains roughly ~1/3 of reliable variance, leaving most unexplained (abilities, starting age, opportunity, etc.).

Implication for MindCraft: Do not sell “10,000 hours.” Sell **high-quality attempts under feedback**. Also respect individual differences; same practice volume will not equalize outcomes. Identity work must include “progress relative to self,” not only absolute rank.

XVI.6 Tutoring effects — demythologize Bloom

FACT: Bloom (1984) popularized the “2 sigma” tutoring + mastery framing.

FACT / CALIBRATION: Broader tutoring metas are far smaller (often ~0.3–0.4 SD in modern RCT syntheses cited by Education Next). VanLehn (2011) estimated human tutoring ~0.79 SD vs no tutoring, with step-based ITS nearly comparable (~0.76).

Implication: Human tutors remain valuable, especially for affect + accountability + diagnosis of weird misconceptions. They are not magic 2-sigma machines. AI step-level feedback can capture much of the *cognitive* tutoring benefit; the scarce remainder is Layers 4–6 (emotion, identity, accountability).

Part XVII — Scarcity Thesis: Full Red Team Trial

XVII.1 The claim under oath

FOUNDER BELIEF: AI commoditizes knowledge; future scarcity is motivation, confidence, identity, trust, curiosity, persistence, emotional safety, meaning.

XVII.2 Cross-examination

Claim A — “Explanations are free”

Verdict: Directionally **FACT** for commodity explanations; **FALSE** for *trusted, personalized, correct, curriculum-aligned* explanations under time pressure.

Evidence: Generative AI + open platforms collapse marginal cost of text explanation.

Contradiction: Hallucinations, shallow fluency, and misalignment create a *trust tax*. Parents may still pay for verified human judgment.

Claim B — “Tutoring is free”

Verdict: FALSE.

AI chat is cheap. High-quality tutoring includes scheduling, relationship, emotional labor, accountability, and reputation. Those remain expensive. Nickow et al.–style tutoring metas still show meaningful effects when tutoring is

structured — it is not free at scale.

Claim C — “Content is free”

Verdict: **MOSTLY FACT** for undifferentiated content; **FALSE** for sequenced, assessed, outcome-linked content systems.

Khan/YouTube prove content abundance. Differentiation is curation + measurement + human wrap.

Claim D — “Motivation/identity become the scarce goods”

Verdict: **HYPOTHESIS** with strategic plausibility.

As cognitive help approaches free, *willingness to engage difficulty* becomes the binding constraint for many students. This does not mean motivation products automatically win — many habit apps fail to produce identity.

Claim E — “Emotional safety is first-order”

Verdict: **HYPOTHESIS** — true for anxious / threatened learners; not universal.

For some, speed drills and competition *increase* engagement. Segment, do not romanticize softness.

XVII.3 Surviving thesis (company doctrine)

Adopted doctrine (refined v1.7):

> MindCraft does not compete primarily on explanation volume, AI tutor warmth, streak/XP theater, grit meters, or story immersion. It competes on converting threatened learners into challenge-seeking mathematical agents through diagnosis, affective design, mastery-climate evaluation, sequenced desirable difficulties (SAFE-DD), student-generated why, variation-designed transfer, productive-error routing, narrative meaning under equity constraints, and human accountability — with AI as *guarded* infrastructure that must clear `solo_transfer_pass` ≥ no-AI.

Falsifiers (90-day capable):

1. Anxious segment shows no FEI lift vs explanation-first control.
2. Parents refuse to pay when score gains lag identity metrics.
3. AI-alone matches human+AI on persistence + transfer for target segment.
4. Guarded Solver raises concurrent scores but loses `solo_transfer_pass` vs no-AI (PWC fail).
5. Interleave/SE regimes raise delayed mix but destroy `retry_120s` in high-anxiety segment without sequenced SAFE-DD (XLI; pending DD-1/DD-2).
6. Soft-wrong without springboard/route matches key-only on retry + misconception recurrence (MIS-1 fail).
7. VT/hug sequences fail to beat random surface skins on next-day format-hop transfer (TR-1 fail).

Part XVIII — Systems Maps (loops, not features)

XVIII.1 FEI reinforcing loop (desired)

```
+----- identity claim ("I do math") <--+  
| |  
v |  
safety UX ----> attempt ----> informative miss ----> coach  
^ | |  
| v v  
+----- map fill / witness <----- earned success <+
```

XVIII.2 Shame balancing loop (destroy this)

```
public failure signal -> threat ↑ -> avoidance -> fewer attempts  
^ |  
+----- skill lag / worse scores -----+
```

XVIII.3 Extrinsic colonization loop (danger)

```
streak / points ↑ -> return ↑ -> shallow grinding ↑  
| |  
+-- identity unchanged -----+--> burnout / streak break → exit
```

Design rule: Habit loops may feed FEI but must not replace identity markers.

XVIII.4 Parent trust loop

```
visible diagnosis + honest progress -> parent trust ↑ -> continued purchase  
^ |  
+---- outcome evidence (scores / challenge-seeking)
```

FOUNDER QUESTION unresolved: Optimize Maya's identity or parent's anxiety? Doctrine: report *both* challenge-seeking and skill evidence; never fake either.

XVIII.5 Network effects (honest)

MindCraft today has weak classic network effects. Do not pretend otherwise.

Possible compounding advantages (HYPOTHESIS):

1. Ontology + misconception memory → better diagnosis over time (data moat if privacy-respecting).
2. Tutor playbooks trained on FEI outcomes → service quality moat.
3. Story worlds with longitudinal character arcs → switching costs of meaning (fragile; easy to cargo-cult).

Part XIX — Cross-Domain Transfer Catalog

Principles that repeatedly change human behavior outside classrooms — filtered for MindCraft use.

Domain	Robust mechanism	Transfer	30-year durable?	AI improve?
Exposure therapy	Graded exposure under safety	Anxiety-graded problem sets	Yes	Yes (personalize gradient)
Sports film study	Specific error review	Wrong-choice coach	Yes	Yes (auto clip of miss type)
Music pedagogy	Technique → repertoire → performance	Drill → quest → “recital” signal	Yes	Partial
Martial arts	Visible ranks + dojo norms	Map regions + tutor dojo culture	Yes	Weak for norms
Aviation	Checklists under stress	Procedure scaffolds when anxious	Yes	Yes
Games (Celeste-like)	Fair failure + retry culture	Soft-wrong physics	Yes	Yes
Chess	Annotated game review	Session postmortems	Yes	Yes
Coaching	Identity + accountability	Tutor as witness	Yes	Assist, not replace
Religion/ritual	Meaning + belonging	Light rituals only; avoid cult dynamics	Yes (human)	Risky
Creator economy	Public craft identity	Optional shareable math artifacts	Uncertain	Yes

Anti-transfer: Casino variable rewards → addiction without skill. Do not copy engagement maxing from social media.

Part XX — Competitive Landscape (mechanism lens)

Product	Primary scarce resource they sell	Identity claim?	Risk
Khan Academy	Free mastery content + reputation	Weak–medium	Explanation commodity; “mastery” without transfer_pass
Duolingo	Habit / streak motivation	Weak (language identity sometimes)	Extrinsic ceiling; appearance of consistency

Brilliant	Aesthetic problem joy + prestige	Medium	Narrow segment; not Maya threat profile
Coursera/edX	Credentials	Medium (career identity)	Completion collapse
Chegg / homework help	Answers under deadline	Anti-identity (outsourcing)	Integrity & AI shock
Private tutors	Human accountability + customization	High when good	Supply constrained
ChatGPT tutors	Instant explanation	Low unless wrapped	Trust / hallucination; Bastani Base harm
MindCraft (target)	FEI conversion + tutor witness + gap diagnosis + solo transfer	Intended high	Must prove, not assert

Strategic implication (v1.19 densified): Do not out-Khan Khan on content breadth or item-count theater (SAFE-COVER). Do not out-ChatGPT on always-full-worked / unlimited-solutions Solver (SAFE-FADE). Do not out-Duo Duo on streaks/leagues/XP or expanding-SRS mystique. Do not out-Brilliant Brilliant on puzzle delight. Do not out-ChatGPT on fluency or graph-file cosplay. Do not out-Anki/FSRS on black-box interval theater. Do not out-Kaplan/bootcamp on massed cram or +points packages (SAFE-DURABLE). Do not out-instant-feedback AI or always-delay cosplay (SAFE-FBTIME). Do not out-node-weakness theater when the failure is a join (SAFE-BRIDGE). Do not out-day-one mastery fireworks or placement-belt cold starts (SAFE-COLD). Do not out-format-count / every-format / visual-learning-style quizzes (SAFE-FORMAT). Do not out-Dual-Coding-Score / always-add-picture / animation-first multimedia theater (SAFE-DUAL). Do not out-Explanation-Score / longer=better monologue / monologue=SE theater (SAFE-EXPLAIN). Do not out-Struggle-Score / blank-time / always-wait / never-stuck stall theater (SAFE-RETRIEVE). Do not out-fail-first / every-miss-as-PF / discovery-without-consolidation cosplay (SAFE-PF). Do not out-guild / modeling-as-lecture / Apprenticeship Score™ costume (SAFE-APPRENTICE). Do not out-Focus-Score / phone-shame / mid-attempt DAU-push theater (SAFE-ATTN). Do not out-Monitoring-Score / metacognition-without-control / one-scan calibration theater (SAFE-MONITOR). Do not out-AI-replaces-tutors / tutoring-is-free / unlimited-human / hours-booked / Session ROI Score™ theater (SAFE-ROI). Do not out-Example-Timing-Score / universal E0→E3 destiny / fixed=adaptive fade costume (SAFE-EXPTIME). Do not out-SRL-Score / reflection-streak / survey=event theater (SAFE-SRL). Do not out-Generative-Learning-Score / AI=generation / teach-for-XP / eight-menu GLA bomb (SAFE-GENERATE). Do not out-Psychological-Safety-Score / poster safe-space / Bermuda-triangle / praise-fog theater (SAFE-ERRCLIMATE). Do not out-Homework-Completion-Score / tonight-% / Chegg-dump / assisted=ready theater (SAFE-HWHELP). Do not out-Embodiment-Score / styles meshing / always-gesture / pose-camera theater (SAFE-GESTURE). Do not out-Peer-Explanation-Score / teach-for-XP / Discord=building / AI=peer teaching theater (SAFE-PEERX). Do not out-Curiosity-Score / cliffhanger-without-commit / postdiction / guess-for-XP theater (SAFE-PREDICT). Do not out-Analogy-Score / Mapping Minutes / surface-story=structure / AI likeness dump theater (SAFE-ANALOG). Do not out-Writing-Score / Journal Minutes / freewrite / AI=WTL theater (SAFE-WRITE). Do not out-Spatial-Score / Mental Rotation Minutes / brain-training=identity theater (SAFE-SPATIAL). Do not out-Formative-Score / Assessment Minutes / 0.4–0.7 AfL legend theater (SAFE-FA). Do not out-Feedback-Literacy-Score / delivery=uptake / thumbs=used theater (SAFE-FBLIT). Do not out-Collaboration-Score / Group Minutes / Discord=CSCL / group-correct=ready theater (SAFE-COLLAB). Do not out-Desirable-Difficulty-Score / Challenge Minutes / always-ease / always-max / cram=ready calendar theater (SAFE-DDSCHED). Do not out-Representation-Score / Diversity Minutes / costume=belonging / match=quality theater (SAFE-ROLEMODEL). Do not out-Capture-Score / Photo Minutes / OCR=mastery / Photomath-pipe notebook AI (SAFE-OFFLINE). Do not out-forever-E0 /

Anxiety·Calm Score™ fade gates (SAFE-FADEX). Do not out-Miss Streak / soft-wrong shame-push parent theater (SAFE-SOFTMSG). Do not out-always-on Map / Map Minutes mid-attempt dumps (SAFE-MAPORCH). Do not out-Burnout Score™ / wellness-primary under caseload (SAFE-BURNWARN). Do not out-fluency=pedagogy / LLM-judge-alone wrap ship (SAFE-EXPLAINQA). Do not out-Partial Credit Score™ / all-wrongs-equal gradebook cosplay (SAFE-TAXON). Do not out-hours-booked crunch / Load Shed Score™ / pour=exam support (SAFE-LOADSHED). Do not out-Distractor Score™ / always-four / beauty-trap MC (SAFE-DISTRACTOR). Do not out-Identity Score™ / Confidence Personality™ / equal-weight recovery / inflate-confidence theater (SAFE-CONFMISS). Do not out-parent fireworks portals, Family Leaderboards, or live wrong-answer stalk (SAFE-PDASH). Do not out-Talk Ratio Score / silence theater / Socratic-by-talk-% ads (SAFE-TALK). Do not out-immersive lore / Story Engagement Score / franchise trivia-as-pedagogy (SAFE-STORYLOAD). Do not out-user-count / star-wall / vague-confidence testimonials as learning proof (SAFE-PROOF). Do not out-DAU/streak/XP / FEI Score™ dashboard theater as learning proof (SAFE-INSTRUMENT). Do not out-tutor-headcount vanity or wellness-theater fidelity (SAFE-WORKFORCE). Do not out-grit character apps, therapy Calm Score™, sports academies, conservatories, grandmaster cosplay, pink STEM SKUs, survey-priced ACT guarantees, Ivy-tutor theater, emotion cameras, or FERPA-badge GTM. Do not out-Microdebrief Score™ / skip-as-default lecture-close (SAFE-MICRODEB). Do not out-24/7 Availability Score™ / fear-parent crunch emails (SAFE-CRUNCHMSG). Do not out-FormatTrap Score™ / beauty visual traps (SAFE-FORMTRAP). Do not out-Brief Score™ / pour-open / late-join shame (SAFE-MICROBRIEF). Do not out-Chegg-hero Dual-CTA Score™ under crunch (SAFE-CRUNCHCTA). Do not out-Figure Key Score™ / fluency SVG pass (SAFE-FIGKEY). Do not out-Wait Minutes·Lobby Score™ / entertainment lobby (SAFE-ASYNCBRIEF). Do not out-Transfer Score™ / dump-and-quit / maze handoff (SAFE-XFERGRAIN). Do not out-HideCorrectness Score™ / tagged=C4 / fluency eligibility / coverage-padding (SAFE-HCELIG). Do not out-ETA Score™ / soon-wait / optimistic slip / tutor-shame feed (SAFE-KNOWNETA). Do not out-Gate Score™ / universal hard-gate / hint-tax shame (SAFE-SOFTGATE). Do not out-SplitPool Score™ / silent dual-use / auto-promote (SAFE-SPLITPOOL). Do not out-Composition Score™ / pad-scan / diagnose-every-format theater (SAFE-SCANCOMP). Do not out-Hole Trust Score™ / hide-holes / complete-visual ads under holes (SAFE-HOLETRUST). Do not out-Promote Score™ / play-count flip / silent hcelig_pass (SAFE-PROMOTE). Do not out-Cool-Down Score™ / immediate same-stem C4 / pad under cool-down (SAFE-COOLDOWN). Do not out-Student Hole Score™ / fake greens / parent-only honesty / ability-shame hole frames (SAFE-STUDHOLE). Do not out-FilleTA Score™ / play-count fill bars / date soft-pass (SAFE-FILLETA). Do not out-Dual Rater Score™ / always-dual-never-dual / mammography cosplay (SAFE-DUALRATE). Do not out-Family Hash Score™ / item_id-only·concept-wide cool-down (SAFE-FAMHASH). Do not out-Exit Emotion Score™ / emotion-AI exit comfort (SAFE-EXITEMO). Do not out-Capacity Score™ / ETA-skipped dual (SAFE-CAPPLAN). Do not out-Template Trust Score™ / silent inherit / forever-dual clones (SAFE-TMPTRUST). Do not out-Kappa % / pooled IRR vanity (SAFE-REPAIRSTRAT). Do not out-Surge Hire % / unsealed freelancer dual / hospital-surge cosplay (SAFE-SURGEPANEL). Out-compete on the **session + ops + honesty stack**: inspectable diagnosis (SAFE-ONTOLOGY), honest cold-start mapping (SAFE-COLD), format-gap conversion evidence (SAFE-FORMAT), contiguous structural dual-channel figures (SAFE-DUAL), parent MoC proof-age surfaces (SAFE-PDASH), witnessable dated solo competence artifacts (SAFE-PROOF), blueprint-honest coverage matrices with gaps named (SAFE-COVER), faded guidance that earns solo carry (SAFE-FADE), instrumental contingent help not unlimited-hint dumps (SAFE-HELP), FEI North Star events under co-gates not vanity OMTMs (SAFE-INSTRUMENT), contingent soft→hard peeks priced in construction not free dumps (SAFE-HINT), length-contingent principle coaches with SE-before-wrap (SAFE-EXPLAIN), mode-contingent retrieval (TOT cue vs freeze scaffold) without Struggle Score™ (SAFE-RETRIEVE), fidelity-gated productive-failure generate→consolidate when priors allow (SAFE-PF), model→coach→fade apprenticeship phases with articulation (SAFE-APPRENTICE), attempt windows protected from residue and mid-attempt pushes (SAFE-ATTN), connection-first Map when endpoints are green (SAFE-BRIDGE), Map-briefed prompt>pour humans with minute budgets (SAFE-HITL / TUTORGRAIN / WORKFORCE / ROI), grain-local example↔solve timing not universal E0→E3 (SAFE-EXPTIME), phase→control

SRL micro-cycles not reflection streaks (SAFE-SRL), SOI-forcing sparse generative asks not AI≡generation (SAFE-GENERATE), stay-diagnose error climate under FEI not poster safety (SAFE-ERRCLIMATE), dual-rail Help vs Practice with solo proof (SAFE-HWHELP), strategy-tied optional gesture/draw not embodiment cosplay (SAFE-GESTURE), building-gated teach-back with solo proof (SAFE-PEERX), predict-commit→surprise→encode not curiosity theater (SAFE-PREDICT), supported structure-map joins not analogy costume (SAFE-ANALOG), short prompted WTL protocols not journal minutes (SAFE-WRITE), FormatId-embedded spatial hops not rotation gyms (SAFE-SPATIAL), evidence→adjust formative cycles not AfL badges (SAFE-FA), feedback uptake loops not delivery≡literacy (SAFE-FBLIT), structured collab with solo FEI gate not Discord≡CSCL (SAFE-COLLAB), horizon-matched DD hold/ease not always-ease/max calendar theater (SAFE-DDSCHED), fidelity-first competence-visible near-peers not Representation Score™ (SAFE-ROLEMODEL), offline-first attempt→capture→confirm not Photomath pipe (SAFE-OFFLINE), protect-then-fade under HMA not forever-E0 (SAFE-FADEX), grain+CTA parent soft-wrong honesty not Miss Streak (SAFE-SOFTMSG), brief→slice→adjust Map not always-on dumps (SAFE-MAPORCH), fidelity/load early-warn→redesign not Burnout Score™ (SAFE-BURNWARN), principle-short wrap QA not fluency theater (SAFE-EXPLAINQA), typed soft-wrong marks not Partial Credit Score™ (SAFE-TAXON), exam-week acuity shed not hours-booked crunch (SAFE-LOADSHED), misconception-mapped distractors not beauty traps (SAFE-DISTRACTOR), confidence×type recovery not Identity Score™ (SAFE-CONFMIS), capped micro-debrief fidelity (SAFE-MICRODEB), shed-honesty crunch parent copy (SAFE-CRUNCHMSG), trap-mapped FormatId distractors (SAFE-FORMTRAP), capped shared-model opens (SAFE-MICROBRIEF), dual-rail crunch CTAs (SAFE-CRUNCHCTA), structural figure key-verify (SAFE-FIGKEY), async Set pre-briefs on wait (SAFE-ASYNCBRIEF), one-tap Solver→Practice transfer grain (SAFE-XFERGRAIN), seals-first FormatId C4 eligibility (SAFE-HCELIG), known-finite late ETAs (SAFE-KNOWNETA), soft-require follow-on with avoidance falsifiers (SAFE-SOFTGATE), graded Practice/C4 pools (SAFE-SPLITPOOL), scarcity-honest scan composition (SAFE-SCANCOMP), parent hole-trust copy (SAFE-HOLETRUST), ticketed promote-to-C4 (SAFE-PROMOTE), reveal→C4 cool-down (SAFE-COOLDOWN), student hole parity copy (SAFE-STUDHOLE), capacity-grounded fill bands (SAFE-FILLETA), risk-triggered dual review (SAFE-DUALRATE), sealed family cool-down hashing (SAFE-FAMHASH), destake gap-scan exits (SAFE-EXITEMO), dual-aware capacity planning (SAFE-CAPPLAN), conditional template-trust inheritance (SAFE-TMPTRUST), path-stratified repair dashboards (SAFE-REPAIRSTRAT), sealed make-or-buy dual surge panels (SAFE-SURGEANEL), verify-before-ship banks (SAFE-GENQ), privacy-bound affect + district trust packets (SAFE-PRIVACY / PROCURE), horizon-matched returns (SAFE-SCHED), calendar-first banded adaptivity (SAFE-ADAPT), time-honest aged evidence (SAFE-FORGET), dual-rail durable GTM (SAFE-DURABLE), mode-conditional feedback clocks (SAFE-FBTIME), SAFE-DP spine under SAFE-DD, productive-error + calibration (SAFE-MISCON / CALIB), recoverable AI truth (SAFE-REPAIR), validated identity on long clocks (SAFE-IDMEASURE / LONGID), structure honesty (SAFE-ELL / GENDER / STRUCTURE), parent CBC (SAFE-WTP), and transfer when help is gone — then *say that* in marketing (CSA-2). Lab process itself obeys SAFE-LABMETA (falsify before “science-backed”).

Part XXI — Metrics Dictionary (instrumentation)

XXI.1 Banned as North Stars

- Raw DAU / time-on-app without transfer
- Streak length alone

- Explanation open-rate / AI chat token count
- Points / coins / XP velocity
- Blocked-session accuracy alone
- Coach thumbs-up / “felt helpful” without solo_transfer_pass
- Named leaderboard rank / league tier
- Resilience Score™ / grit meter / MRS / Grit Score™ as in-app KPI
- Midnight minutes practiced / all-nighter volume
- Belief Score™ / confidence meters / confidence streaks
- Community tab visits / Discord volume / games played / rating
- Calm Score™ / anxiety-must-hit-zero meters
- Tip-recall / film-room engagement without method-change + transfer
- Math Identity Score™ / two-item “math person” alone / splash identity Likert
- Parent survey WTP dollars / guaranteed ACT-point package tiers
- NAEP EL-gap closure / gender STEM-gap closure as product claims
- Ontology node-count / “powered by our knowledge graph” without FEI outcome
- Hours-booked / smile-star tutor QA without solo_transfer_pass
- Anxiety Score™ / face-voice emotion engagement meters
- Chapter/page count / “science-backed Constitution” without falsifiers
- Item-count / “AI wrote N questions” without verify rate
- Complete-ACT / every-concept coverage ads without concept×level matrix + gaps
- FERPA-compliant badge / teacher click-wrap as district readiness
- Expanding-SRS / “perfect interval AI” / streak-as-spacing engines
- Tutor headcount / retention-% without fidelity@tenure
- Mastery fireworks / personal Ebbinghaus / Forgetting Score™ as durable-competence claims
- SM-2/FSRS/Anki hero / black-box perfect-interval AI as learning NS
- Cram-week / ACT-point package tiers as durable-identity claims
- Instant-feedback latency / always-delay as learning NS
- Bridge-count / “we connect everything” without join transfer
- Day-one mastery greens / placement-belt cold start
- Format-count / Format Personality™ / learning-style quiz NS
- Parent tonight-% / Family Leaderboard / live wrong-answer stalk
- Talk Ratio Score™ / raw talk-% / 50/50 airtime dogma as Socratic NS
- Story Engagement Score™ / immersion / lore-depth as learning NS
- Fade Score™ / always-full-worked / unlimited-solutions as learning NS
- Help Score™ / unlimited-hints / never-stuck dump as learning NS
- Hint Score™ / XP-for-hints / unlimited free hard peeks as learning NS
- FEI Score™ / single-event tutor commissions on retry_120s alone
- Coach thumbs-up / help NPS without solo_transfer_pass
- Explanation Score™ / token-count quality / longer⇒better monologue as learning NS
- Struggle Score™ / blank-time / dwell-XP as productive-struggle NS
- PF Score™ / fail-first-always / generation-minutes without consolidation

- Apprenticeship Score™ / modeling-minutes / guild cosplay as pedagogy NS
- Focus Score™ / Deep Work Score™ / mid-attempt DAU pushes as learning NS
- Dual Coding Score™ / Visual Engagement Score™ / always-add-picture as learning NS
- Monitoring Score™ / Metacognition % / one-scan calibration as learning NS
- Session ROI Score™ / hours-booked / AI-replaces-tutors as learning NS
- Example Timing Score™ / universal E0→E3 destiny / fixed≡adaptive fade as learning NS
- SRL Score™ / reflection-streak / survey-SRL≡event as learning NS
- Generative Learning Score™ / teach-for-XP / AI≡generation as learning NS
- Psychological Safety Score™ / Comfort % / poster safe-space as learning NS
- Homework Completion Score™ / tonight-% / Chegg-dump / assisted≡ready as learning NS
- Embodiment Score™ / Gesture Score™ / Motion Minutes / pose-camera as learning NS
- Peer Explanation Score™ / Teach Minutes / Discord≡building as learning NS
- Curiosity Score™ / cliffhanger-without-commit / postdiction / guess-for-XP as learning NS
- Analogy Score™ / Mapping Minutes / surface-story≡structure as learning NS
- Writing Score™ / Journal Minutes / freewrite≡learning as learning NS
- Spatial Score™ / Mental Rotation Minutes / brain-training≡identity as learning NS
- Formative Score™ / Assessment Minutes / 0.4–0.7 AfL ads as learning NS
- Feedback Literacy Score™ / Feedback Minutes / delivery≡uptake as learning NS
- Collaboration Score™ / Group Minutes / Discord≡CSCL / group-correct≡ready as learning NS
- Desirable Difficulty Score™ / Challenge Minutes / always-ease / always-max / cram≡ready as learning NS
- Representation Score™ / Diversity Minutes / costume≡belonging / match≡quality as learning NS
- Capture Score™ / Photo Minutes / OCR≡mastery / photo→answer as learning NS
- Forever-E0 anxious track / Anxiety·Calm Score™ fade gates / Fade Score™ as learning NS
- Miss Streak / soft-wrong shame push / Parent Correctness Score™ / empty mindset parent letter as learning NS
- Always-on Map / Map Minutes·Map Score™ / mid-attempt full-graph as learning NS
- Burnout Score™ / Wellness Minutes / wellness-primary under caseload as learning NS
- Fluency≡pedagogy / Explanation·Rubric Score™ / LLM-judge-alone wrap ship as learning NS
- Soft-wrong≡half-credit / Partial Credit Score™ / all-wrongs-equal / gradebook cosplay as learning NS
- Hours-booked crunch / Load Shed Score™ / infinite-availability / pour≡exam support as learning NS
- Distractor Score™ / Option Minutes / always-four≡better / beauty traps as learning NS
- Identity Score™ / Confidence Personality™ / equal-weight recovery / inflate-confidence as learning NS
- Microdebrief Score™ / Debrief Minutes / skip-as-default lecture-close / +25% ACT debrief ads as learning NS
- 24/7 availability default / Availability Score™ / hours-hero crunch / fear-parent emails / silent shed as care NS
- FormatTrap Score™ / Visual Distractor Minutes / seductive stems / beauty visual traps as learning NS
- Brief Score™ / Brief Minutes / Opening Streak / no-time-for-brief / pour-open / late-join shame as learning NS
- Dual-CTA Score™ / Help Tonight Minutes / Chegg-hero crunch CTA / fear-countdown dump as learning NS
- Figure Key Score™ / Visual QA Minutes / fluency SVG≡keyed / unverified FormatId diagnostic as learning NS
- Wait Minutes / Lobby Score™ / Prebrief Streak / idle spinner / entertainment-primary lobby as learning NS
- Transfer Score™ / Handoff Minutes / Follow-on Streak / dump-and-quit / maze handoff / assisted≡ready as learning NS
- HideCorrectness Score™ / Eligibility Minutes / tagged≡C4 / coverage-padding unsealed visuals as learning NS
- ETA Score™ / Wait Satisfaction / soon-default / optimistic slip / tutor-shame feed as care NS

- Gate Score™ / Forced Follow-On % / universal hard-gate / hint-tax shame as learning NS
- SplitPool Score™ / Dual-Use Minutes / silent dual-write / auto-promote as learning NS
- Composition Score™ / Format Balance % / pad-scan / diagnose-every-format ads as learning NS
- Hole Trust Score™ / Coverage Satisfaction / hide-holes / complete-visual ads under holes as trust NS
- Promote Score™ / QA Minutes / play-count auto-promote / silent hcelig_pass as bank-quality NS
- Cool-Down Score™ / Exposure Minutes / immediate same-stem C4 / pad under cool-down as diagnosis NS
- Student Hole Score™ / Completeness % / fake greens / parent-only hole honesty as identity NS
- Filleta Score™ / Backlog Minutes / play-count fill bars / date soft-pass / coming-soon hole clocks as trust NS
- Dual Rater Score™ / Agreement Minutes / Kappa % / always-dual·never-dual / mammography cosplay as quality NS
- Family Hash Score™ / Clone Minutes / item_id-only·concept-wide·embedding-only cool-down as learning NS
- Exit Emotion Score™ / Anxiety·Calm Score™ exit / Mood Minutes / emotion-AI exit empathy as care NS
- Capacity Score™ / Reviewer Minutes / dual-omitted FILLETA / ETA-skipped dual as honesty NS
- Template Trust Score™ / Inherit % / silent inherit / forever-dual clones / FERF≡C4-ready as quality NS
- Kappa % NS / pooled repair vanity / Landis-band marketing / Dual Shame boards as integrity NS
- Surge Hire % / Contract Minutes / unsealed freelancer dual / hospital-surge cosplay as capacity NS

XXI.2 Leading indicators (ship first)

Metric ID	Definition	Why
retry_120s	New attempt on same or isomorphic item within 120s of soft-wrong	Persistence under safety
coach_shown	Wizard/coach surfaced	Treatment exposure
write_exit_to_retry	Leave write mode then attempt	Agency after reflection
challenge_accept	Chose harder level when offered	Challenge-seeking (instrument motive)
hint_binge	≥3 hints without independent solve	Gaming / helplessness / failed exposure fade
transfer_pass	Correct on varied item after mastery mark	Anti-fake-mastery
solo_transfer_pass	Transfer with AI/Solver closed or denied	Anti-crutch (XXXIII)
ai_reveal_rate	Full-answer / unguarded reveal frequency	Crutch exposure
strategy_class_error	Wrong <i>procedure family</i> on mixed set	Interleaving discrimination (XXXIX)
se_principle_hit	Correct structured principle/ingredient pick when prompted	SE quality proxy (XL)

help_recruit_then_solo	Opened card/tutor/peer help then solved without full reveal	SAFE-RESILIENCE support path (XLV)
misconception_route_hit	Soft-wrong routed to mapped mis_ / ingredient failure_mode	SAFE-MISCON classification (XLIX)
format_hop_pass	Correct after declared format/context hop (VT check)	SAFE-TRANSFER discernment (XLVIII)
confidence_miss_tier	High vs low confidence on miss (when elicited)	SAFE-CALIB / hypercorrection (L)
annotate_before_reveal	Student principle/why logged before key/coach/model	SAFE-ANNOT / SAFE-AAR (LVII / LIV)
film_cue_count	Coaching points attached to attempt clip (cap ≤3)	SAFE-FILM less-is-more (LV)
avoided_format_approach	Attempt on previously avoided format/stakes step	SAFE-EXPOSE approach rate (LVIII)
repair_reattempt_pass	Correct destaked re-attempt after owned AI miss	SAFE-REPAIR calibration (LX)
escalation_resolve_48h	Packetized human handoff closed with fix	SAFE-REPAIR escalate SLA (LX)
improve_how_specificity	Debrief yields concrete next-attempt tweak	SAFE-AAR close quality (LIV)
stakes_ladder_step	Declared pressure step completed after destaked equip	SAFE-EXAM dose-to-prove (LXIII)
lang_vs_math_miss_flag	Miss tagged language-parse vs concept (when elicited)	SAFE-ELL load separation (LXIV)
identity_battery_ok	Validated multi-factor identity admin completed (not splash)	SAFE-IDMEASURE demand-aware (LXI)
bandwidth_tax_skip	Incomplete mission under self-reported scarcity/stress (not grit)	SAFE-STRUCTURE hygiene (LXVI)
map_brief_open	Tutor/session opened Map diagnosis before first explain	SAFE-HITL / ONTOLOGY brief fidelity (LXVIII / LXVII)
key_verify_pass	Generated item passed independent key check before bank write	SAFE-GENQ ship gate (LXXII)
return_isi_days	Days between concept exposure and scheduled retrieval return	SAFE-SCHED horizon lag (LXXIV)
trust_packet_ready	DPA + Exhibit A + data map + deletion/breach posted for LEA	SAFE-PROCURE GTM gate (LXXIII)
tutor_fidelity_30d / tutor_fidelity_90d	FEI rubric sample pass rate at tenure horizons	SAFE-WORKFORCE drift gate (LXXV)
coach_util	Coach sample bandwidth vs active tutor count	SAFE-WORKFORCE hire-cap (LXXV)

last_proof_age_days	Days since last delayed/generation proof on a concept	SAFE-FORGET aged evidence (LXXVI)
retention_probe_7d / retention_probe_28d	Pass rate on scheduled delayed retrieval probes	SAFE-FORGET / MoC honesty (LXXVI)
return_due_reason	Shown rationale chip for next return (horizon vs evidence pull)	SAFE-ADAPT inspectability (LXXVII)
adapt_band_isi_days	Actual ISI vs horizon baseline (within declared band)	SAFE-ADAPT modulation audit (LXXVII)
cram_intercept_shown	User requested massed week and saw RI tradeoff + denser Returns offer	SAFE-DURABLE cram-intercept (LXXVIII)
prove_rail_labeled	Exam-week intensity sessions tagged prove (not mastery)	SAFE-DURABLE dual-rail honesty (LXXVIII)
fb_microdelay_ms	Forced SE/attempt gate duration before learn-rail reveal	SAFE-FBTIME learn gate (LXXIX)
soft_wrong_se_before_reveal	Soft-wrong path required student why before coach	SAFE-FBTIME / SAFE-SE (LXXIX)
bridge_gap_mission_start	Practice launched from bridge-gap CTA (not node-only)	SAFE-BRIDGE Map/Practice (LXXX)
join_solo_transfer_pass	Solo pass on delayed item requiring named bridge	SAFE-BRIDGE join transfer (LXXX)
seed_probe_disagree	High-confidence seed vs hide-correctness probe miss on same concept	SAFE-COLD calibration (LXXXI)
cold_map_honesty	First-session Map showed untouched/seed labels (not mastery greens)	SAFE-COLD chrome (LXXXI)
format_gap_mission_start	Practice launched from format-gap CTA (not topic-only)	SAFE-FORMAT Map/Practice (LXXXII)
format_hop_solo_transfer_pass	Solo pass on delayed item in previously weak FormatId	SAFE-FORMAT conversion (LXXXII)
parent_proof_age_view	Parent opened proof-age strip / MoC surface	SAFE-PDASH honesty (LXXXIII)
parent_return_cta_start	Student return/mission started after parent CTA or digest	SAFE-PDASH action loop (LXXXIII)
parent_softmsg_open	Parent opened soft-wrong grain digest (not miss-streak hero)	SAFE-SOFTMSG honesty (CXVIII)
parent_intrusion_tip_shown	Autonomy / anti-control tip shown after soft-wrong run	SAFE-SOFTMSG autonomy>control (CXVIII)

map_slice_adjust	Mid-session grain-slice Map open that changed next prompt/mission	SAFE-MAPORCH adjust (CXIX)
map_student_dual_open	Student-facing Map open concurrent with active attempt (demote / gate)	SAFE-MAPORCH anti-dual (CXIX)
student_constructive_share	Why/SE/soft-wrong turns / student turns (sampled session)	SAFE-TALK construction fidelity (LXXXIV)
prompt_density	Eliciting prompts / tutor turns	SAFE-TALK prompt>pour (LXXXIV)
wait_ms_p50	Median wait time 1 and 2 (ms)	SAFE-TALK wait coaching (LXXXIV)
pour_streak_max	Longest consecutive tutor explain turns	SAFE-TALK monologue detect (LXXXIV)
story_time_to_first_attempt_s	Seconds from module open to first math attempt	SAFE-STORYLOAD pre-attempt budget (LXXXV)
story_wrap_noninferior	Transfer/retention within margin vs plain stem	SAFE-STORYLOAD learning gate (LXXXV)
proof_card_share_opt_in	Opt-in shares of dated solo-transfer proof cards	SAFE-PROOF share dignity (LXXXVI)
proof_card_named_hop	Share/view events that name concept/bridge/FormatId + proof type	SAFE-PROOF specificity (LXXXVI)
coverage_matrix_view	Parent/LEA/admin opened dated concept×level coverage matrix	SAFE-COVER honesty surface (LXXXVII)
coverage_claim_mismatch	Marketed coverage vs playable audit disagree (ops incident)	SAFE-COVER substantiation (LXXXVII)
fade_stage	Solver/help guidance stage E0–E3 at attempt start	SAFE-FADE ladder (LXXXVIII)
exptime_grain	Concept×format×bridge grain used to set/recheck fade stage	SAFE-EXPTIME (C)
srl_phase_event	Forethought / performance / reflection micro-event with control writeback	SAFE-SRL (CI)
gla_soi_event	Sparse generative ask (summarize/teach/draw/map) with SOI quality gate	SAFE-GENERATE (CII)
error_episode_fidelity	Tutor/coach stayed, diagnosed, scaffolded, required re-attempt after miss	SAFE-ERRCLIMATE (CIII)
fade_step_attempt	Student-generated completion/SE before next worked unlock	SAFE-FADE attempt grain (LXXXVIII)

help_deliberate_ms	Time on non-bottom hint level before next action	SAFE-HELP deliberate help (LXXXIX)
help_executive_race	≥2 hint-level advances under short dwell before bottom-out	SAFE-HELP abuse detect (LXXXIX)
challenge_motive	Enum on challenge_accept: mastery / appearance / normative / unknown	SAFE-INSTRUMENT motive gate (XC)
fei_event_fire_ok	Soft-wrong path emitted retry/coach/binge companions without dark funnel	SAFE-INSTRUMENT completeness (XC)
hint_level_open	Soft/principle vs hard/bottom-out level shown at request	SAFE-HINT ladder grain (XCI)
hard_peek_cost_ok	Hard level advanced only after SE/blank/dwell/stage gate	SAFE-HINT construction cost (XCI)
coach_wrap_tokens	Coach/Solver wrap token count before student SE	SAFE-EXPLAIN length discipline (XCII)
se_before_wrap	Student principle/why logged before full coach wrap	SAFE-EXPLAIN SE-before-wrap (XCII)
stall_mode	Enum: tot / blank_freeze / progress / avoidance / unknown	SAFE-RETRIEVE mode classify (XCIII)
tot_cue_then_retry	Micro-cue after TOT wait then independent retry	SAFE-RETRIEVE TOT path (XCIII)
pf_phase	Mission phase: generate / consolidate / guided	SAFE-PF sequence fidelity (XCIV)
pf_rsm_captured	Student RSM attempt logged before Phase-2 C&C	SAFE-PF consolidation input (XCIV)
ca_phase	Shared model/coach/scaffold/fade/articulate/reflect label	SAFE-APPRENTICE phase log (XCV)
ca_articulate_ok	Student articulation before bottom-out reveal	SAFE-APPRENTICE articulate gate (XCV)
attempt_interrupted	Focus lost / push / tab-away during active attempt	SAFE-ATTN residue (XCVI)
focus_mode_opt_in	Session used opt-in focus / push-suppress mode	SAFE-ATTN hygiene (XCVI)
figure_contiguous_ok	Stem+figure labels spatially contiguous (audit flag)	SAFE-DUAL contiguity (XCVII)
decorative_figure_reject	Decorative figure blocked from diagnostic use	SAFE-DUAL coherence (XCVII)
judgment_type	Gap-scan judgment enum: ease / RCJ / delayed_JOL	SAFE-MONITOR named judgment (XCVIII)
monitor_control_wired	Confidence×outcome miss class drove a routing action	SAFE-MONITOR control wiring (XCVIII)

tutor_minute_budget_ok	Human session stayed within Map minute budget	SAFE-ROI scarce inventory (XCIX)
fei_per_tutor_minute	FEI events / billed tutor minutes (ops; not vanity NS)	SAFE-ROI allocation proof (XCIX)
help_rail_mode	Session labeled Help (instrumental unstuck) vs Practice (FEI gym)	SAFE-HWHELP dual-rail (CIV)
assisted_then_solo_gate	After Help, solo transfer required before ready claim	SAFE-HWHELP assisted≠solo (CIV)
gesture_strategy_opt_in	Optional strategy-tied gesture/trace/draw used (not forced)	SAFE-GESTURE optional motion (CV)
peerx_building_code	Teach-back coded building vs telling (HITL/AI rubric)	SAFE-PEERX quality gate (CVI)
predict_commit_before_reveal	Prediction logged before feedback/key/demo reveal	SAFE-PREDICT commit wick (CVII)
predict_encode_repair	After miss/surprise, repair encode before next item	SAFE-PREDICT encode loop (CVII)
analog_map_act	Student produced source↔target alignment on a join	SAFE-ANALOG mapping act (CVIII)
analog_copresent_ok	Source kept visible during target comparison	SAFE-ANALOG support (CVIII)
write_protocol_ok	Short cognitive+metacognitive Notes/protocol chip completed with content	SAFE-WRITE protocol (CIX)
write_student_authored	Explanation draft authored by student (not AI-substituted)	SAFE-WRITE authorship (CIX)
spatial_format_hop_ok	Delayed diagram/graph/number_line hop after remediation	SAFE-SPATIAL format hop (CX)
spatial_embed_not_arcade	Spatial practice tied to concept/FormatId (not isolated rotation gym)	SAFE-SPATIAL embed (CX)
fa_cycle_closed	In-session evidence led to a named next attempt (Sadler act)	SAFE-FA cycle (CXI)
exit_ticket_wired	Wrap item changed Map/mission (not XP-only)	SAFE-FA exit wire (CXI)
feedback_uptake_act	After coach/tutor feedback, student performed named next attempt	SAFE-FBLIT uptake (CXII)
feedback_decode_ok	Soft-wrong/coach at concept/ingredient/FormatId grain (not jargon essay)	SAFE-FBLIT decode (CXII)

solo_after_collab	Solo transfer attempt required after collaborative episode before unlock	SAFE-COLLAB solo gate (CXIII)
collab_accountability_ok	Group goal + individual check present (not shared-green only)	SAFE-COLLAB Slavin hinge (CXIII)
ddsched_band_ok	Practice band matches days-to-exam (Build / Join / Protect+probe)	SAFE-DDSCHEM horizon (CXIV)
ddsched_stakes_split	Train sets destaked; prove probes labeled (stakes≠DD)	SAFE-DDSCHEM stakes split (CXIV)
rolemodel_fidelity_first	Tutor passed GRAIN fidelity before any preference match applied	SAFE-ROLEMODEL hire bar (CXV)
rolemodel_competence_visible	Tutor bio/cue shows attainable near-peer competence (not costume-only)	SAFE-ROLEMODEL Marx signal (CXV)
offline_attempt_before_capture	Student committed attempt offline before capture unlock	SAFE-OFFLINE attempt grain (CXVI)
capture_ocr_confirmed	OCR/Map write confirmed by student or tutor before mastery event	SAFE-OFFLINE verify (CXVI)
photo_solve_side_door	Capture path used as answer-app / full-solution dump	SAFE-OFFLINE anti-Photomath (CXVI)
fadex_hold_ok	High-MA learner held at protective E0/E1 then advanced without forever track	SAFE-FADEX protect-then-fade (CXVII)
fadex_peek_binge	Consecutive full peeks treated as safety behavior to fade	SAFE-FADEX exposure join (CXVII)
burnwarn_fidelity_alert	Fidelity/load early-warn fired before exit (not wellness-primary)	SAFE-BURNWARN early-warn (CXX)
burnwarn_redesign_act	Caseload/org redesign action logged after alert	SAFE-BURNWARN JD-R (CXX)
explainqa_gate_pass	Coach wrap passed principle-short + math/key hard-fail before ship	SAFE-EXPLAINQA gate (CXXI)
explainqa_fluency_reject	Fluent wrap rejected for math/key fail or monologue costume	SAFE-EXPLAINQA anti-fluency (CXXI)
taxon_mark_type	Miss assigned typed mark (slip/bug/lack/blank/assist)	SAFE-TAXON type (CXXII)
taxon_route_hit	Typed mark drove Map/route action (not binary red-X only)	SAFE-TAXON route (CXXII)

loadshed_pause_fired	Named exam-week elective paused per playbook	SAFE-LOADSHED pause list (CXXIII)
loadshed_acuity_ok	Remaining HITL minutes prioritized high-acuity FEI under crunch	SAFE-LOADSHED acuity (CXXIII)
distractor_mapped_ok	Chosen MC distractor carries named mis_/trap seal	SAFE-DISTRACTOR map (CXXIV)
distractor_dead_retired	Non-functioning distractor retired after item analysis	SAFE-DISTRACTOR function (CXXIV)
microdeb_complete	Close hit four-beat micro-AAR within time cap (student spoke first)	SAFE-MICRODEB complete (CXXVI)
microdeb_trace_ok	≥ 1 objective attempt/media trace cited in close	SAFE-MICRODEB traces (CXXVI)
crunchmsg_shed_named	Parent crunch message named pause/shed honestly (not silent or 24/7)	SAFE-CRUNCHMSG honesty (CXXVII)
crunchmsg_minutes_frame	Copy framed protected thinking minutes (not hours-booked hero)	SAFE-CRUNCHMSG framing (CXXVII)
formtrap_seal_ok	Visual MC distractor carries named format_trap_id (or mis_)	SAFE-FORMTRAP seal (CXXVIII)
formtrap_route_hit	Format-trap pick drove conversion/FormatId route (not topic-only)	SAFE-FORMTRAP route (CXXVIII)
brief_complete	Late-join/on-time open hit four-beat micro-brief (or full brief)	SAFE-MICROBRIEF complete (CXXIX)
synthesis_ok	Student restated edge/success before first problem	SAFE-MICROBRIEF receiver synthesis (CXXIX)
rail_chosen	Student/parent selected help-tonight vs Practice-gym CTA	SAFE-CRUNCHCTA rail (CXXX)
followon_start	Practice grain started after Solver instrumental unstick	SAFE-CRUNCHCTA follow-on (CXXX)
figkey_pass	Generated FormatId figure passed structural/constraint verify with stem^key	SAFE-FIGKEY pass (CXXXI)
figkey_drop_reason	Coded fail (stem mismatch / geometry / label / judge-alone reject)	SAFE-FIGKEY drop (CXXXI)
asyncbrief_complete	Late/early lobby Set card submitted with four beats within cap	SAFE-ASYNCBRIEF complete (CXXXII)

plan_chip_ok	If-then plan chip present and session-related	SAFE-ASYNCBRIEF implementation intention (CXXXII)
xfer_near_attempt	Post-Solver Practice grain started on same edge (one-tap)	SAFE-XFERGRAIN near attempt (CXXXIII)
xfer_near_pass	Near transfer item passed with Solver/AI closed or denied	SAFE-XFERGRAIN near pass (CXXXIII)
xfer_return_due	SAFE-SCHED return scheduled after transfer grain	SAFE-XFERGRAIN × SCHED (CXXXIII)
hcelig_pass	FormatId item passed seals for C4 hide-correctness pool	SAFE-HCELIG pass (CXXXIV)
hcelig_block_reason	Coded block (figkey/genq/formtrap/graphicacy) when tagged≠eligible	SAFE-HCELIG block (CXXXIV)
eta_finite_shown	Late lobby showed known-finite ETA band (not soon)	SAFE-KNOWNETA finite (CXXXV)
eta_refresh_honest	ETA refresh without optimistic under-promise	SAFE-KNOWNETA refresh (CXXXV)
softgate_default	Post-help close used soft-require (Close available)	SAFE-SOFTGATE default (CXXXVI)
hardgate_trigger_ok	Hard-gate only under named trigger + avoidance falsifier clear	SAFE-SOFTGATE conditional (CXXXVI)
help_avoid_after_gate	Solver open/help recruit drop after gate escalation	SAFE-SOFTGATE avoidance falsifier (CXXXVI)
pool_practice_only	Item served from Practice-reveal pool (not C4)	SAFE-SPLITPOOL Practice (CXXXVII)
pool_c4_hcelig	Item served from C4 HCELIG pool	SAFE-SPLITPOOL C4 (CXXXVII)
silent_dual_write_block	Practice-only outcome blocked from C4 Map mastery write	SAFE-SPLITPOOL integrity (CXXXVII)
scancomp_hole_flag	Gap-scan emitted format_not_probed for empty sealed cell	SAFE-SCANCOMP hole (CXXXVIII)
scancomp_pad_reject	Selector rejected symbolic/Practice pad into C4 cell	SAFE-SCANCOMP anti-pad (CXXXVIII)
holetrust_labeled	Parent summary listed probed vs not_probed + seal reason	SAFE-HOLETRUST label (CXXXIX)
holetrust_cta_one	Exactly one fill CTA paired with hole (Practice or promote path)	SAFE-HOLETRUST CTA (CXXXIX)
promote_ticket_open	Practice→C4 promote ticket created with checklist	SAFE-PROMOTE ticket (CXL)

promote_pass	Ticket closed with seal re-verify + reviewer id → hcelig_pass	SAFE-PROMOTE pass (CXL)
promote_demote	C4 item demoted on critical repair faster than promote cycle	SAFE-PROMOTE demote (CXL)
cooldown_until	Per-student stem-family cool-down bar active after reveal	SAFE-COOLDOWN bar (CXL I)
cooldown_subst_ok	Isomorphic out-of-family substitute used (or hole) under cool-down	SAFE-COOLDOWN substitute (CXL I)
studhole_parity	Student hole facts match parent HOLETRUST (no dual script)	SAFE-STUDHOLE parity (CXL II)
studhole_no_fake_green	Unprobed cells not rendered as mastered/flexible	SAFE-STUDHOLE anti-pad (CXL II)
studhole_cta_one	Student summary had exactly one Practice/open-more CTA	SAFE-STUDHOLE CTA (CXL II)
confmiss_tier_route	Confidence×type class chose recovery dose (not equal deck)	SAFE-CONFMISS CRI fork (CXXV)
confmiss_highconf_priority	High-conf miss prioritized for attended destaked springboard	SAFE-CONFMISS hypercorrection (CXXV)

XXI.3 Lagging indicators

Metric ID	Definition
math_person_item	1–7 agreement: “I am a math person” (math-specific; never sole endpoint)
cribbs_interest / cribbs_recognition / cribbs_competence	Factor scores from validated battery (LXI)
anxiety_state_item	Short state anxiety before session
advanced_intent	Intent to take harder course / contest
tutor_witness_note	Tutor tagged identity-relevant growth
identity_fade_26w / identity_fade_52w	Triangulated identity movement vs baseline at horizons (LXII)

XXI.4 Decision rule (HYPOTHESIS)

Ship changes that raise `retry_120s` and `mastery-motive challenge_accept` without raising `hint_binge` / `ai_reveal_rate`, and without dropping `transfer_pass` or `solo_transfer_pass`. Prefer delayed mixed accuracy over blocked session accuracy.

Part XXII — Appendices

Appendix A — Glossary

Term	Meaning
FEI Loop	Fear → Evidence → Identity system
ITC	Identity Transformation Cascade
Soft-wrong	Miss treated as information, not shame theater
Gap scan	Low-shame diagnostic of confidence/exposure
Challenge-seeking	Choosing harder math opportunity when optional
False growth mindset	Effort praise without strategy / conditions
SAFE-CRAFT	Improved tutoring sequence (Part XVI.3)
Layers 0–6	Knowledge commodity stack (Part VI.3)
HID	Habit/Identity Divergence — streaks can raise return without identity
MoC	Memory of Competence — spaced retrieval after transfer_pass
PWC	Pedagogy Wrap Constraint — guarded AI + student generation + solo transfer bar
Claim ladder L0–L4	Max defensible causal/marketing claim given design (Part XXXIV)
$E \times V \times \text{Cost}$	Eccles expectancy $\times$ subjective task value $\times$ cost triad
TARGET	Classroom/goal-structure dimensions cueing mastery vs performance
SAFE-DD	Equip → destake → dose → frame → measure (Part XLI)
SAFE-COMPARE	Criterion-default comparison policy (Part XLII)
SAFE-HABIT	Cue practice without streak colonization (Part XLIII)
SAFE-REWARD	Informational / non-controlling reward policy (Part XLIV)
SAFE-RESILIENCE	Growth-zone + support recruitment, not grit meters (Part XLV)
SAFE-EXPECTANCY	Task briefs + CIOF; ban trait labels (Part XLVI)
SAFE-TIMING	Encode → sleep → delayed transfer claim (Part XLVII)
SAFE-TRANSFER	Variation/hug/bridge transfer design (Part XLVIII)

SAFE-MISCON	Productive-error elicit → classify → springboard → route (Part XLIX)
SAFE-CALIB	Item-level confidence → miss class → tiered feedback (Part L)
SAFE-DP	Diagnose → isolate → attempt → method feedback → prove spine (Part LI)
SAFE-CoP	Named practice; peripheral = real work; trajectory > tab (Part LII)
SAFE-RITUAL	Light optional practice-pointing marks; anti-totalism (Part LIII)
SAFE-AAR	Brief → traces → four-question student-first close (Part LIV)
SAFE-FILM	Sparse attempt clips + self-model + next commitment (Part LV)
SAFE-MUSIC	Scales → repertoire → transfer-as-recital ladder (Part LVI)
SAFE-ANNOT	Annotate-before-reveal → classify → drill; annotate wins (Part LVII)
SAFE-EXPOSE	Situation hierarchy + expectancy violation + fade safety (Part LVIII)
SAFE-WTP	Parent CBC + bias controls + message≠package (Part LIX)
SAFE-REPAIR	Detect→own→correct→reattempt→triggered escalate (Part LX)
SAFE-IDMEASURE	Validated identity factors + FEI behavior + demand-aware admin (Part LXI)
SAFE-LONGID	Stage 8w process / 26w triangulated identity / 52w persistence; fade-aware (Part LXII)
SAFE-EXAM	Destake-to-learn / dose-to-prove; delayed mixed under stakes; no ACT guarantees (Part LXIII)
SAFE-ELL	Language load ≠ math load; modify incidental English not math; no deficit track (Part LXIV)
SAFE-GENDER	Parity hygiene; cue hygiene; practice recognition; no ST→score / pink cosplay (Part LXV)
SAFE-STRUCTURE	Structure > slogans; bandwidth hygiene; no grit-NS / character-blame equity (Part LXVI)
SAFE-ONTOLOGY	Diagnosis-before-dialogue; inspectable state; LLM bookends; moat≠graph file (Part LXVII)
SAFE-HITL	Map-briefed tutors; prompt>pour; coaching cycles; FEI QA; packetized escalate (Part LXVIII)

SAFE-PRIVACY	Self-authored short-TTL affect → pedagogy only; ERT ban; contextual integrity (Part LXIX)
SAFE-LABMETA	Label→check→kill/demote; small FEI core; page-count vanity kill (Part LXX)
SAFE-TUTORGRAIN	Trained near-peer default; structure>pairing; hire by fidelity; expert=escalate (Part LXXI)
SAFE-GENQ	Verify-before-ship; key hard-fail; Kane IUA; drop-rate gate; coverage≠NS (Part LXXII)
SAFE-PROCURE	Privacy review = GTM gate; NDPA-ready; Exhibit A honesty; no bio / no pupil ads (Part LXXIII)
SAFE-SCHED	Horizon-matched ISI×RI; delayed first return; equal-ish lags; generation Returns; no streak-as-spacing (Part LXXIV)
SAFE-WORKFORCE	Fidelity-over-tenure; continuous coaching; coach-capacity hire gate; burnout→load redesign; no headcount NS (Part LXXV)
SAFE-FORGET	Age evidence; JOLs≠retention; delayed probes as truth; returns as competence protection; no Ebbinghaus Score™ (Part LXXVI)
SAFE-ADAPT	Calendar-first; banded inspectable personalization; selection>opaque rewrite; flashcard lineage≠FEI (Part LXXVII)
SAFE-DURABLE	Dual-rail GTM; cram-intercept not cosplay; legible MoC for WTP; no point guarantees (Part LXXVIII)
SAFE-FBTIME	Mode-conditional timing; focus>timing; prior-knowledge switch; no instant-feedback hero (Part LXXIX)
SAFE-BRIDGE	Connection-first when endpoints non-naive; remediate relation; measure join transfer; no bridge-count moat (Part LXXX)
SAFE-COLD	Humble prior → labeled seed → probes; Map = evidence age; no day-one mastery fireworks (Part LXXXI)
SAFE-FORMAT	Format as first-class evidence; conversion as object; support before stack; no format-count / style-quiz NS (Part LXXXII)
SAFE-PDASH	Beliefs before chrome; proof-age hero; one improvement CTA; autonomy>control; ranks/affect off (Part LXXXIII)
SAFE-TALK	Construction > airtime; prompt density + constructive share + wait bands; no Talk Ratio Score™ (Part LXXXIV)

SAFE-STORYLOAD	CLT-budgeted narrative wrap; structure-encode or cut; coherence > immersion; no Story Engagement Score™ (Part LXXXV)
SAFE-PROOF	Competence > crowd; dated solo artifacts; provincial similarity; thin portfolio; no star-wall / +points proof (Part LXXXVI)
SAFE-COVER	Matrix > total; gaps as features; use-tier labels; examTag honesty; no item-count / complete-ACT / unsynced GENQ (Part LXXXVII)
SAFE-FADE	Stage not costume; completion bridge; backward fade; expertise-aware; attempt grain; no unlimited-solutions / never-fade hero (Part LXXXVIII)
SAFE-HELP	Instrumental > executive; abuse∪avoidance fail; contingent stage; SE before bottom-out; no unlimited-hints / Help Score™ (Part LXXXIX)
SAFE-INSTRUMENT	Ship the four; NSM+OMTM; XXI.4 co-gates; anti-Goodhart; Kane IUA; no DAU/streak/FEI Score™ NS (Part XC)
SAFE-HINT	Soft before hard; contingency; construction-priced peeks; fade×hint; no Hint Score™ / free hard dump (Part XCI)
SAFE-EXPLAIN	Short principle-first; germane≠more text; SE before wrap; length fades; no Explanation Score™ / monologue≡SE (Part XCII)
SAFE-RETRIEVE	Mode-contingent stalls; TOT→cue; freeze→soft bound; attempt before help; no Struggle Score™ / never-stuck (Part XCIII)
SAFE-PF	Fidelity-gated generate→consolidate on student RSMs; guided default when thin/exam/freeze; no PF Score™ / every-miss≡PF (Part XCIV)
SAFE-APPRENTICE	Model→coach→scaffold→fade + articulate; shared ca_phase; no guild cosplay / modeling≡lecture / Apprenticeship Score™ (Part XCV)
SAFE-ATTN	Protect attempt windows; residue park/resume; opt-in focus + tech breaks; no Focus Score™ / phone-shame / mid-attempt DAU pushes (Part XCVI)
SAFE-DUAL	Structural dual coding; contiguous labels; coherence weed; static annotated default; no Dual Coding Score™ / styles meshing / always-add-picture (Part XCVII)

SAFE-MONITOR	Named judgments → miss-class control; C4 destake; densifies SAFE-CALIB/COLD; no Monitoring Score™ / one-scan calibration (Part XCVIII)
SAFE-ROI	Map minute budgets; constrained AI for reps; humans for joins; FEI per tutor-minute; no AI-replaces-tutors / hours NS (Part XCIX)
SAFE-EXPTIME	Grain-local example↔solve timing; adaptive>fixed; no universal E0→E3 / Timing Score™ (Part C)
SAFE-SRL	Phase→control SRL cycles; microanalysis>surveys; no SRL Score™ / streak≡SRL (Part CI)
SAFE-GENERATE	SOI-forcing sparse GLAs; no Generative Learning Score™ / AI≡generation / teach-for-XP (Part CII)
SAFE-ERRCLIMATE	Stay-diagnose error climate; no Safety Score™ / Bermuda-triangle / praise-fog (Part CIII)
SAFE-HWHELP	Dual-rail Help vs Practice; no Completion Score™ / Chegg-dump / assisted≡ready (Part CIV)
SAFE-GESTURE	Strategy-tied optional gesture/draw; no Embodiment Score™ / styles meshing / pose-camera (Part CV)
SAFE-PEERX	Building-gated teach-back; no Peer Explanation Score™ / teach-for-XP / Discord≡building (Part CVI)
SAFE-PREDICT	Commit before reveal; encode repair; no Curiosity Score™ / cliffhanger-without-commit / postdiction (Part CVII)
SAFE-ANALOG	Structure-map joins; co-present compare; no Analogy Score™ / mention≡pedagogy / surface-story≡structure (Part CVIII)
SAFE-WRITE	Short prompted WTL protocols; no Writing Score™ / Journal Minutes / freewrite≡learning / AI≡WTL (Part CIX)
SAFE-SPATIAL	FormatId-embedded spatial hops; no Spatial Score™ / Mental Rotation Minutes / brain-training≡identity (Part CX)
SAFE-FA	Evidence→adjust cycles; no Formative Score™ / Assessment Minutes / 0.4–0.7 AfL ads (Part CXI)
SAFE-FBLIT	Sense+use uptake loops; no Feedback Literacy Score™ / Feedback Minutes / delivery≡uptake / thumbs≡used (Part CXII)
SAFE-COLLAB	Structured collab + solo FEI gate; no Collaboration Score™ / Group Minutes / Discord≡CSCL / group-correct≡ready (Part CXIII)

SAFE-DDSCHED	Horizon-matched DD hold/ease; stakes≠DD; no Desirable Difficulty Score™ / Challenge Minutes / always-ease / always-max / cram⇒ready (Part CXIV)
SAFE-ROLEMODEL	Fidelity-first near-peer belonging; competence-visible attainable bios; no Representation Score™ / costume⇒belonging / match⇒quality (Part CXV)
SAFE-OFFLINE	Attempt→capture→confirm; no Capture Score™ / Photo Minutes / OCR⇒mastery / photo→answer hero (Part CXVI)
SAFE-FADEX	Protect-then-fade under HMA; no forever-E0 / flood / Anxiety·Calm Score™ fade gates (Part CXVII)
SAFE-SOFTMSG	Grain+failure-enhancing+one CTA; no Miss Streak / shame-push / empty mindset parent letter (Part CXVIII)
SAFE-MAPORCH	Brief→grain-slice→adjust; no always-on Map / Map Minutes / mid-attempt full-graph (Part CXIX)
SAFE-BURNWARN	Fidelity/load early-warn→redesign; no Burnout Score™ / wellness-primary under caseload (Part CXX)
SAFE-EXPLAINQA	Principle-short wrap QA gates; no fluency⇒pedagogy / LLM-judge-alone / soft-pass wrong math (Part CXXI)
SAFE-TAXON	Typed soft-wrong marks; no Partial Credit Score™ / all-wrongs-equal / gradebook cosplay (Part CXXII)
SAFE-LOADSHED	Exam-week acuity shed; no hours-booked crunch / Load Shed Score™ / pour⇒exam support (Part CXXIII)
SAFE-DISTRACTOR	Misconception-mapped functioning distractors; soft-wrong from chosen option; no Distractor Score™ / always-four / beauty traps (Part CXXIV)
SAFE-CONFMISS	Confidence×type gap-scan miss recovery; CRI fork + hypercorrection priority; no Identity Score™ / equal-weight recovery / inflate-confidence NS (Part CXXV)
SAFE-MICRODEB	Capped student-first micro-AAR with traces; no Microdebrief Score™ / skip-as-default / lecture-as-debrief (Part CXXVI)
SAFE-CRUNCHMSG	Shed-honesty crunch parent CBC/messages; no 24/7 default / hours-hero / fear-parent emails / Availability Score™ (Part CXXVII)
SAFE-FORMTRAP	Named format-trap visual distractors; coherent structural stems; no FormatTrap Score™ / seductive stems / beauty visual traps (Part CXXVIII)

SAFE-MICROBRIEF	Capped shared-model open + student synthesis; no Brief Score™ / no-time-for-brief / pour-open / late-join shame (Part CXXIX)
SAFE-CRUNCHCTA	Dual labeled Solver/Practice CTAs under crunch; instrumental verbs + follow-on Practice; no Chegg-hero / Dual-CTA Score™ / fear-dump (Part CXXX)
SAFE-FIGKEY	Structural figure key-verify for FormatId generation; deterministic spine + use gating; no fluency SVG / judge-alone / Figure Key Score™ (Part CXXXI)
SAFE-ASYNCBRIEF	Capped async Set card on late/early wait; session-related occupied wait + verify handoff; no Wait Minutes / Lobby Score™ / pour-open wait (Part CXXXII)
SAFE-XFERGRAIN	One-tap edge-matched Practice grain after instrumental Solver; solo gate + SCHED return; no dump-and-quit / Transfer Score™ / maze handoff (Part CXXXIII)
SAFE-HCELIG	Seals-first FormatId C4 eligibility; honest holes > pad; no HideCorrectness Score™ / tagged≡C4 (Part CXXXIV)
SAFE-KNOWNETA	Known-finite late-tutor ETA + refresh honesty; no ETA Score™ / soon-default / shame feed (Part CXXXV)
SAFE-SOFTGATE	Soft-require default; conditional hard-gate + avoidance falsifiers; no Gate Score™ / hint-tax shame (Part CXXXVI)
SAFE-SPLITPOOL	Graded Practice vs C4 pools; ticketed promote; no SplitPool Score™ / silent dual-write (Part CXXXVII)
SAFE-SCANCOMP	Sealed-stock scan composition with labeled holes; no Composition Score™ / pad theater (Part CXXXVIII)
SAFE-HOLETRUST	Parent hole honesty + one fill CTA; no Hole Trust Score™ / hide-holes / complete-visual ads (Part CXXXIX)
SAFE-PROMOTE	Ticketed Practice→C4 QA; demote faster than promote; no Promote Score™ / play-count flip (Part CXL)
SAFE-COOLDOWN	Per-student reveal→C4 cool-down + substitutes/holes; no Cool-Down Score™ / immediate same-stem C4 (Part CXLI)

SAFE-STUDHOLE	Student hole parity + product-hole frame + one CTA; no Student Hole Score™ / fake greens (Part CXLII)
SAFE-FILLETA	Capacity-grounded finite fill bands; no Filleta Score™ / play-count bars / date soft-pass (Part CXLIII)
SAFE-DUALRATE	Risk-triggered second rater; seals before dual; no Dual Rater Score™ / always-dual·never-dual (Part CXLIV)
SAFE-FAMHASH	Sealed isomorphic family cool-down unit; no Family Hash Score™ / item_id-only·concept-wide ban (Part CXLV)
SAFE-EXITEMO	Destake gap-scan exit + one CTA; no Exit Emotion Score™ / emotion-AI exit comfort (Part CXLVI)
SAFE-CAPPLAN	Dual minutes in FILLETA estimator; no Capacity Score™ / ETA-skipped dual (Part CXLVII)
SAFE-TMPTRUST	Inherit intensity not seals; sample+revoke; no Template Trust Score™ / silent inherit (Part CXLVIII)
SAFE-REPAIRSTRAT	Path-stratified repair dashboards; κ aide not NS; no Kappa % / pooled IRR vanity (Part CXLIX)
SAFE-SURGE PANEL	Sealed make-or-buy dual surge portfolio; no Surge Hire % / unsealed freelancer dual (Part CL)
SAFE-REVOKE	Revocable template-trust + switching sample cadence; no Template Trust Score™ / forever-trust latch (Part CLI; provisional)
SAFE-PATHBLAME	Path×stratum blame before dual/single quality verdicts; no Dual Shame Score™ / unconditioned rankings (Part CLII; provisional)
SAFE-PANELRAMP	Seal literacy + shadow dual before contract C4 authority; no Parallel Score™ / day-one C4 (Part CLIII; provisional)
SAFE-DEMOB	Planned post-spike dual wind-down; no Demobilization Score™ / invent-work / cliff-cut (Part CLIV; provisional)

Appendix B — Maya interview protocol (qual)

Goal: Map moments when math stopped (or started) feeling dangerous.

1. “Tell me about a time math felt dangerous or humiliating.”
2. “What did your body do? What did you do next?”
3. “Tell me about a time you surprised yourself in math.”

2. Attack any metric that can be gamed without learning.
3. Attack any narrative feature without a learning outcome.
4. Attack Bloom 2-sigma mythology on sight.
5. Attack “AI makes tutoring free” slogans on sight.

Appendix F — Future chapters (queued, not padded)

1. Neuroscience of math anxiety (amygdala/WM models; cautious translation)
2. History of mathematics as meaning technology — under trial via XXVIII + HIST
3. Parent trust economics & willingness-to-pay — **DONE** Part LIX / SAFE-WTP (builds EVT-4 / XXVII)
- 3b. Trust after AI hallucination / repair + escalate — **DONE** Part LX / SAFE-REPAIR (builds XXXIII)
- 3c. Measurement of math identity — **DONE** Part LXI / SAFE-IDMEASURE
- 3d. Longitudinal identity change models — **DONE** Part LXII / SAFE-LONGID
- 3e. ACT/exam pressure special case — **DONE** Part LXIII / SAFE-EXAM
- 3f. Multilingual / ELL math identity — **DONE** Part LXIV / SAFE-ELL
- 3g. Gender & math stereotypes update — **DONE** Part LXV / SAFE-GENDER
- 3h. Socioeconomic constraint on “grit” — **DONE** Part LXVI / SAFE-STRUCTURE
- 3i. Ontology as diagnosis moat — **DONE** Part LXVII / SAFE-ONTOLOGY
- 3j. Human-in-the-loop tutor ops — **DONE** Part LXVIII / SAFE-HITL
- 3k. Privacy & student affect data — **DONE** Part LXIX / SAFE-PRIVACY
- 3l. Constitution meta / Bayesian lab process — **DONE** Part LXX / SAFE-LABMETA
- 3m. Peer vs near-peer vs expert tutoring — **DONE** Part LXXI / SAFE-TUTORGRAIN
- 3n. Generated question validity & key risk — **DONE** Part LXXII / SAFE-GENQ
- 3o. District procurement & privacy review as GTM — **DONE** Part LXXIII / SAFE-PROCURE
- 3p. Spaced retrieval schedules in product UX — **DONE** Part LXXIV / SAFE-SCHED
- 3q. Tutor workforce pipeline & quality drift — **DONE** Part LXXV / SAFE-WORKFORCE
- 3r. Forgetting curves as product honesty — **DONE** Part LXXVI / SAFE-FORGET
- 3s. Adaptive spacing algorithms vs fixed calendars — **DONE** Part LXXVII / SAFE-ADAPT
- 3t. Cram products vs durable-identity positioning — **DONE** Part LXXVIII / SAFE-DURABLE
- 3u. Feedback timing immediate vs delayed — **DONE** Part LXXIX / SAFE-FBTIME
- 3v. Bridge-gap product narrative — **DONE** Part LXXX / SAFE-BRIDGE
- 3w. Cold-start without fake mastery — **DONE** Part LXXXI / SAFE-COLD
- 3x. Format-axis product narrative — **DONE** Part LXXXII / SAFE-FORMAT
- 3y. Parent dashboard honesty UX — **DONE** Part LXXXIII / SAFE-PDASH
- 3z. Tutor talk-ratio instrumentation — **DONE** Part LXXXIV / SAFE-TALK

3aa. Story-world cognitive load vs identity payoff — **DONE** Part LXXXV / SAFE-STORYLOAD

3ab. Competence evidence as social proof — **DONE** Part LXXXVI / SAFE-PROOF

3ac. Bank coverage honesty in marketing — **DONE** Part LXXXVII / SAFE-COVER

3ad. Worked-example fading in Solver UX — **DONE** Part LXXXVIII / SAFE-FADE

3ae. Help-seeking vs help abuse — **DONE** Part LXXXIX / SAFE-HELP

3af. Product analytics for FEI North Stars — **DONE** Part XC / SAFE-INSTRUMENT

3ag. Hint economy & contingent scaffolding — **DONE** Part XCI / SAFE-HINT

3ah. Explanation length vs germane load — **DONE** Part XCII / SAFE-EXPLAIN

3ai. Retrieval failure modes in practice UX — **DONE** Part XCIII / SAFE-RETRIEVE

3aj. Productive failure vs guided success sequencing — **DONE** Part XCIV / SAFE-PF

3ak. Cognitive apprenticeship in tutor playbooks — **DONE** Part XCV / SAFE-APPRENTICE

3al. Attention residue & device distraction in practice — **DONE** Part XCVI / SAFE-ATTN

3am. Dual coding & diagram FormatId load — **DONE** Part XCVII / SAFE-DUAL

3an. Metacognitive monitoring in gap-scan — **DONE** Part XCVIII / SAFE-MONITOR

3ao. Opportunity cost of tutor minutes vs AI scaffold — **DONE** Part XCIX / SAFE-ROI

3ap. Worked-example vs problem-solving timing across expertise — **DONE** Part C / SAFE-EXPTIME

3aq. Self-regulated learning cycles (Zimmerman) in Practice UX — **DONE** Part CI / SAFE-SRL

3ar. Generative learning activities (Fiorella/Mayer) in Practice — **DONE** Part CII / SAFE-GENERATE

3as. Error climate & psychological safety in tutoring — **DONE** Part CIII / SAFE-ERRCLIMATE

3at. Homework help vs practice identity conflict — **DONE** Part CIV / SAFE-HWHELP

3au. Gesture & embodiment in math practice — **DONE** Part CV / SAFE-GESTURE

3av. Peer explanation quality (knowledge-building vs telling) — **DONE** Part CVI / SAFE-PEERX

3aw. Curiosity & prediction before reveal — **DONE** Part CVII / SAFE-PREDICT

3ax. Analogical reasoning & bridge teaching — **DONE** Part CVIII / SAFE-ANALOG

3ay. Writing-to-learn in math practice — **DONE** Part CIX / SAFE-WRITE

3az. Spatial skills & math identity — **DONE** Part CX / SAFE-SPATIAL

3ba. Formative assessment cycles in tutor sessions — **DONE** Part CXI / SAFE-FA

3bb. Feedback literacy (coach feedback uptake) — **DONE** Part CXII / SAFE-FBLIT

3bc. Collaborative problem-solving vs solo FEI — **DONE** Part CXIII / SAFE-COLLAB

3bd. Desirable difficulty scheduling under exam deadlines — **DONE** Part CXIV / SAFE-DDSCHED

3be. Tutor identity & belonging (near-peer role models) — **DONE** Part CXV / SAFE-ROLEMODEL

3bf. Offline practice fidelity (paper/whiteboard capture) — **DONE** Part CXVI / SAFE-OFFLINE

3bg. **Synthesizer v1.15** — merge CVIII–CXVI into I.4 company law — **DONE**

3bh. Worked-example fading under high math anxiety — **DONE** Part CXVII / SAFE-FADEX

3bi. Parent message framing after soft-wrong streaks — **DONE** Part CXVIII / SAFE-SOFTMSG

3bj. Classroom orchestration of Map during live tutoring — **DONE** Part CXIX / SAFE-MAPORCH

3bk. Tutor burnout early-warning without wellness theater — **DONE** Part CXX / SAFE-BURNWARN

3bl. Explanation quality rubrics for AI coach wraps — **DONE** Part CXXI / SAFE-EXPLAINQA

3bm. Partial-credit & soft-wrong taxonomy productization — **DONE** Part CXXII / SAFE-TAXON

3bn. Exam-week tutor load shedding playbooks — **DONE** Part CXXIII / SAFE-LOADSHED

3bo. Distractor design for diagnostic soft-wrong MC — **DONE** Part CXXIV / SAFE-DISTRACTOR

3bp. Confidence-tagged misses in gap-scan recovery — **DONE** Part CXXV / SAFE-CONFMISS

3bq. **Synthesizer v1.16** — merge CXVII–CXXV into I.4 company law — **DONE**

3br. Tutor micro-debrief fidelity under time caps — **DONE** Part CXXVI / SAFE-MICRODEB

3bs. Parent crunch-week message tests (availability vs thinking minutes) — **DONE** Part CXXVII / SAFE-CRUNCHMSG

3bt. Format-trap distractors on diagram/graph items — **DONE** Part CXXVIII / SAFE-FORMTRAP

3bu. Session-open brief fidelity under late joins — **DONE** Part CXXIX / SAFE-MICROBRIEF

3bv. Dual-rail Solver CTA copy under crunch — **DONE** Part CXXX / SAFE-CRUNCHCTA

3bw. Figure key-verify loops for generated FormatId stems — **DONE** Part CXXXI / SAFE-FIGKEY

3bx. Async pre-brief cards while waiting for late tutor — **DONE** Part CXXXII / SAFE-ASYNCBRIEF

3by. One-tap Solver→Practice transfer grain after instrumental help — **DONE** Part CXXXIII / SAFE-XFERGRAIN

3bz. **Synthesizer v1.17** — merge CXXVI–CXXXIII into I.4 company law — **DONE**

3ca. Hide-correctness eligibility gates for FormatId items — **DONE** Part CXXXIV / SAFE-HCELIG

3cb. Known-finite ETA honesty in late-tutor lobby copy — **DONE** Part CXXXV / SAFE-KNOWNETA

3cc. Soft-require vs hard-gate follow-on under help-avoidance risk — **DONE** Part CXXXVI / SAFE-SOFTGATE

3cd. Reveal-Practice vs C4 split pools for sealed FormatId — **DONE** Part CXXXVII / SAFE-SPLITPOOL

3ce. FormatId gap-scan composition under sealed-stock scarcity — **DONE** Part CXXXVIII / SAFE-SCANCOMP

3cf. Parent trust copy when visual diagnostic holes remain — **DONE** Part CXXXIX / SAFE-HOLETRUST

3cg. Promote-to-C4 tutor QA workflow for Practice-only FormatId — **DONE** Part CXL / SAFE-PROMOTE

3ch. Exposure cool-down after reveal before same-stem C4 — **DONE** Part CXLI / SAFE-COOLDOWN

3ci. Student-facing hole copy after gap-scan — **DONE** Part CXLII / SAFE-STUDHOLE

3ci2. **Synthesizer v1.18** — merge CXXXIV–CXLII into I.4 company law — **DONE**

3ci3. Promote backlog SLA vs parent hole-ETA honesty — **DONE** Part CXLIII / SAFE-FILLETA

3ci4. Second-rater thresholds for figured C4 admits — **DONE** Part CXLIV / SAFE-DUALRATE

3ci5. Isomorphic stem-family hashing for cool-down — **DONE** Part CXLV / SAFE-FAMHASH

3cj. Gap-scan exit emotion without Anxiety Score™ — **DONE** Part CXLVI / SAFE-EXITEMO

- 3ck. Dual-rater capacity planning vs FILLETA band honesty — **DONE** Part CXLVII / SAFE-CAPPLAN
- 3cl. Template-trust inheritance after first-of-cell dual admit — **DONE** Part CXLVIII / SAFE-TMPTRUST
- 3cm. Post-promote repair stratified dashboards (single vs dual) — **DONE** Part CXLIX / SAFE-REPAIRSTRAT
- 3cn. Dual-capacity hire vs contract SME panels under exam-week spikes — **DONE** Part CL / SAFE-SURGEPANEL
- 3cn2. **Synthesizer v1.19** — merge CXLIII–CL into I.4 company law — **DONE**
- 3co. Template-trust revoke & sample-dual cadence — **DONE** Part CLI / SAFE-REVOKE (id 151)
- 3cp. Path-conditioned dual blame tests — **DONE** Part CLII / SAFE-PATHBLAME (id 152)
- 3cq. Contract SME panel seal onboarding & parallel-run gates — **DONE** Part CLIII / SAFE-PANELRAMP (id 153)
- 3cr. Post-spike dual demobilization without fidelity cliff — **DONE** Part CLIV / SAFE-DEMOB (id 154)
- 3cs. Gold-set planted probes for path×stratum monitors — **NEXT** id 155
- 3ct. Generator vs rater poison separation on inherit repair — queued id 156
- 3cu. LEA packet path-mix honesty without Dual Shame theater — queued id 157
- 4. Equity audit of story worlds — **DONE** (Part XXXVI)
- 5. Competitive teardown — **DONE** session audits (Part XXXV); usage telemetry still open
- 6. Formal Bayesian update process for Constitution claims — **DONE** Part LXX / SAFE-LABMETA
- 7. **DONE (2026-07-30)**: Desirable difficulties × anxiety (Part XLI) — SAFE-DD stack
- 8. **DONE (2026-07-30)**: Social comparison & leaderboards (Part XLII) — SAFE-COMPARE stack
- 9. **DONE (2026-07-30)**: Habit formation science without identity colonization (Part XLIII) — SAFE-HABIT stack
- 10. **DONE (2026-07-30)**: Intrinsic motivation killers / Deci caveats (Part XLIV) — SAFE-REWARD stack
- 11. **DONE (2026-07-30)**: Mathematical resilience (Part XLV / id 45) — SAFE-RESILIENCE stack
- 12. **DONE (2026-07-30)**: Teacher/tutor expectancy effects (Part XLVI / id 46) — SAFE-EXPECTANCY stack
- 13. **DONE (2026-07-30)**: Sleep, stress, and learning (Part XLVII / id 47) — SAFE-TIMING stack
- 14. **DONE (2026-07-30)**: Transfer & variation theory (Part XLVIII / id 48) — SAFE-TRANSFER stack
- 15. **DONE (2026-07-31)**: Misconceptions as productive (Part XLIX / id 49) — SAFE-MISCON stack
- 16. **Synthesizer v1.7 (2026-07-31)**: Merged XLI–XLIX into I.4 / XIV / IX / XXI — no new chapter
- 17. **DONE (2026-07-31)**: Confidence calibration (Part L / id 50) — SAFE-CALIB stack
- 18. **DONE (2026-07-31)**: Deliberate practice in tutoring sessions (Part LI / id 51) — SAFE-DP stack
- 19. **DONE (2026-07-31)**: Community of practice / Lave–Wenger (Part LII / id 52) — SAFE-CoP stack
- 20. **DONE (2026-07-31)**: Religion/ritual light-touch (Part LIII / id 53) — SAFE-RITUAL stack
- 21. **DONE (2026-07-31)**: Military/aviation brief-debrief (Part LIV / id 54) — SAFE-AAR stack
- 22. **DONE (2026-07-31)**: Sports film-study pedagogy (Part LV / id 55) — SAFE-FILM stack
- 23. **DONE (2026-08-01)**: Music pedagogy ladders (Part LVI / id 56) — SAFE-MUSIC stack
- 24. **DONE (2026-08-01)**: Chess annotation & metacognition (Part LVII / id 57) — SAFE-ANNOT stack
- 25. **DONE (2026-08-01)**: Therapy graded exposure (Part LVIII / id 58) — SAFE-EXPOSE stack
- 26. **Synthesizer v1.8 (2026-08-01)**: Merged L–LVIII into I.4 / XIV / IX / XXI — no new chapter
- 27. **DONE (2026-08-01)**: Parent WTP experiments design (Part LIX / id 59) — SAFE-WTP stack
- 28. **DONE (2026-08-01)**: Trust after AI hallucination (Part LX / id 60) — SAFE-REPAIR stack
- 29. **DONE (2026-08-01)**: Measurement of math identity (Part LXI / id 61) — SAFE-IDMEASURE stack

30. **DONE (2026-08-02):** Longitudinal identity change models (Part LXII / id 62) — SAFE-LONGID stack
31. **DONE (2026-08-02):** ACT/exam pressure special case (Part LXIII / id 63) — SAFE-EXAM stack
32. **DONE (2026-08-02):** Multilingual / ELL math identity (Part LXIV / id 64) — SAFE-ELL stack
33. **DONE (2026-08-02):** Gender & math stereotypes update (Part LXV / id 65) — SAFE-GENDER stack
34. **DONE (2026-08-02):** Socioeconomic constraint on “grit” (Part LXVI / id 66) — SAFE-STRUCTURE stack
35. **Synthesizer v1.9 (2026-08-02):** Merged LIX–LXVI into I.4 / XIV / IX / XXI — no new chapter
36. **DONE (2026-08-02):** Ontology as diagnosis moat (Part LXVII / id 67) — SAFE-ONTOLOGY stack
37. **DONE (2026-08-03):** Human-in-the-loop tutor ops (Part LXVIII / id 68) — SAFE-HITL stack
38. **DONE (2026-08-03):** Privacy & student affect data (Part LXIX / id 69) — SAFE-PRIVACY stack
39. **DONE (2026-08-03):** Constitution meta / Bayesian lab process (Part LXX / id 70) — SAFE-LABMETA stack
40. **DONE (2026-08-03):** Peer vs near-peer vs expert tutoring (Part LXXI / id 71) — SAFE-TUTORGRAIN stack
41. **DONE (2026-08-03):** Generated question validity & key risk (Part LXXII / id 72) — SAFE-GENQ stack
42. **DONE (2026-08-03):** District procurement & privacy review as GTM (Part LXXIII / id 73) — SAFE-PROCURE stack
43. **DONE (2026-08-04):** Spaced retrieval schedules in product UX (Part LXXIV / id 74) — SAFE-SCHED stack
44. **Synthesizer v1.10 (2026-08-04):** Merged LXVII–LXXIV into I.4 / XIV / IX / XXI — no new chapter
45. **DONE (2026-08-04):** Tutor workforce pipeline & quality drift (Part LXXV / id 75) — SAFE-WORKFORCE stack
46. **DONE (2026-08-04):** Forgetting curves as product honesty (Part LXXVI / id 76) — SAFE-FORGET stack
47. **DONE (2026-08-04):** Adaptive spacing algorithms vs fixed calendars (Part LXXVII / id 77) — SAFE-ADAPT stack
48. **DONE (2026-08-04):** Cram products vs durable-identity positioning (Part LXXVIII / id 78) — SAFE-DURABLE stack
49. **DONE (2026-08-04):** Feedback timing immediate vs delayed (Part LXXIX / id 79) — SAFE-FBTIME stack
50. **DONE (2026-08-04):** Bridge-gap product narrative (Part LXXX / id 80) — SAFE-BRIDGE stack
51. **DONE (2026-08-05):** Cold-start without fake mastery (Part LXXXI / id 81) — SAFE-COLD stack
52. **DONE (2026-08-05):** Format-axis product narrative (Part LXXXII / id 82) — SAFE-FORMAT stack
53. **DONE (2026-08-05):** Parent dashboard honesty UX (Part LXXXIII / id 83) — SAFE-PDASH stack
54. **Synthesizer v1.11 (2026-08-05):** Merged LXXV–LXXXIII into I.4 / XIV / IX / XXI — no new chapter
55. **DONE (2026-08-05):** Tutor talk-ratio instrumentation (Part LXXXIV / id 84) — SAFE-TALK stack
56. **DONE (2026-08-05):** Story-world cognitive load vs identity payoff (Part LXXXV / id 85) — SAFE-STORYLOAD stack
57. **DONE (2026-08-05):** Competence evidence as social proof (Part LXXXVI / id 86) — SAFE-PROOF stack
58. **DONE (2026-08-05):** Bank coverage honesty in marketing (Part LXXXVII / id 87) — SAFE-COVER stack
59. **DONE (2026-08-06):** Worked-example fading in Solver UX (Part LXXXVIII / id 88) — SAFE-FADE stack
60. **DONE (2026-08-06):** Help-seeking vs help abuse (Part LXXXIX / id 89) — SAFE-HELP stack
61. **DONE (2026-08-06):** Product analytics for FEI North Stars (Part XC / id 90) — SAFE-INSTRUMENT stack
62. **DONE (2026-08-06):** Hint economy & contingent scaffolding (Part XCI / id 91) — SAFE-HINT stack
63. **Synthesizer v1.12 (2026-08-06):** Merged LXXXIV–XCI into I.4 / XIV / IX / XXI — no new chapter
64. **DONE (2026-08-06):** Explanation length vs germane load (Part XCII / id 92) — SAFE-EXPLAIN stack
65. **DONE (2026-08-06):** Retrieval failure modes in practice UX (Part XCIII / id 93) — SAFE-RETRIEVE stack
66. **DONE (2026-08-07):** Productive failure vs guided success sequencing (Part XCIV / id 94) — SAFE-PF stack

67. **DONE (2026-08-07):** Cognitive apprenticeship in tutor playbooks (Part XCV / id 95) — SAFE-APPRENTICE stack
68. **DONE (2026-08-07):** Attention residue & device distraction in practice (Part XCVI / id 96) — SAFE-ATTN stack
69. **DONE (2026-08-07):** Dual coding & diagram FormatId load (Part XCVII / id 97) — SAFE-DUAL stack
70. **DONE (2026-08-07):** Metacognitive monitoring in gap-scan (Part XCVIII / id 98) — SAFE-MONITOR stack
71. **DONE (2026-08-07):** Opportunity cost of tutor minutes vs AI scaffold (Part XCIX / id 99) — SAFE-ROI stack
72. **Synthesizer v1.13 (2026-08-07):** Merged XCII–XCIX into I.4 / XIV / IX / XXI — no new chapter
73. **DONE (2026-08-08):** Worked-example vs problem-solving timing across expertise (Part C / id 100) — SAFE-EXPTIME stack
74. **DONE (2026-08-08):** Self-regulated learning cycles in Practice UX (Part CI / id 101) — SAFE-SRL stack
75. **DONE (2026-08-08):** Generative learning activities in Practice (Part CII / id 102) — SAFE-GENERATE stack
76. **DONE (2026-08-08):** Error climate & psychological safety in tutoring (Part CIII / id 103) — SAFE-ERRCLIMATE stack
77. **DONE (2026-08-08):** Homework help vs practice identity conflict (Part CIV / id 104) — SAFE-HWHELP stack
78. **DONE (2026-08-08):** Gesture & embodiment in math practice (Part CV / id 105) — SAFE-GESTURE stack
79. **DONE (2026-08-08):** Peer explanation quality (Part CVI / id 106) — SAFE-PEERX stack
80. **DONE (2026-08-09):** Curiosity & prediction before reveal (Part CVII / id 107) — SAFE-PREDICT stack
81. **Synthesizer v1.14 (2026-08-09):** Merged C–CVII into I.4 / XIV / IX / XXI — no new chapter
82. **DONE (2026-08-09):** Analogical reasoning & bridge teaching (Part CVIII / id 108) — SAFE-ANALOG stack
83. **DONE (2026-08-09):** Writing-to-learn in math practice (Part CIX / id 109) — SAFE-WRITE stack
84. **DONE (2026-08-09):** Spatial skills & math identity (Part CX / id 110) — SAFE-SPATIAL stack
85. **DONE (2026-08-09):** Formative assessment cycles in tutor sessions (Part CXI / id 111) — SAFE-FA stack
86. **DONE (2026-08-09):** Feedback literacy / coach uptake (Part CXII / id 112) — SAFE-FBLIT stack
87. **DONE (2026-08-09):** Collaborative problem-solving vs solo FEI (Part CXIII / id 113) — SAFE-COLLAB stack
88. **DONE (2026-08-10):** Desirable difficulty scheduling under exam deadlines (Part CXIV / id 114) — SAFE-DDSCHED stack
89. **DONE (2026-08-10):** Tutor identity & belonging / near-peer role models (Part CXV / id 115) — SAFE-ROLEMODEL stack
90. **DONE (2026-08-10):** Offline practice fidelity / paper-whiteboard capture (Part CXVI / id 116) — SAFE-OFFLINE stack
91. **Synthesizer v1.15 (2026-08-10):** Merged CVIII–CXVI into I.4 / XIV / IX / XXI — no new chapter
92. **DONE (2026-08-10):** Worked-example fading under high math anxiety (Part CXVII / id 117) — SAFE-FADEX stack
93. **DONE (2026-08-10):** Parent message framing after soft-wrong streaks (Part CXVIII / id 118) — SAFE-SOFTMSG stack
94. **DONE (2026-08-10):** Classroom orchestration of Map during live tutoring (Part CXIX / id 119) — SAFE-MAPORCH stack
95. **DONE (2026-08-10):** Tutor burnout early-warning without wellness theater (Part CXX / id 120) — SAFE-BURNWARN stack
96. **DONE (2026-08-11):** Explanation quality rubrics for AI coach wraps (Part CXXI / id 121) — SAFE-EXPLAINQA stack

97. **DONE (2026-08-11):** Partial-credit & soft-wrong taxonomy (Part CXXII / id 122) — SAFE-TAXON stack
98. **DONE (2026-08-11):** Exam-week tutor load shedding playbooks (Part CXXIII / id 123) — SAFE-LOADSHED stack
99. **DONE (2026-08-11):** Distractor design for diagnostic soft-wrong MC (Part CXXIV / id 124) — SAFE-DISTRACTOR stack
100. **DONE (2026-08-11):** Confidence-tagged misses in gap-scan recovery (Part CXXV / id 125) — SAFE-CONFMISS stack
101. **DONE (2026-08-11):** Tutor micro-debrief fidelity under time caps (Part CXXVI / id 126) — SAFE-MICRODEB stack
102. **DONE (2026-08-11):** Parent crunch-week message tests (Part CXXVII / id 127) — SAFE-CRUNCHMSG stack
103. **DONE (2026-08-12):** Format-trap distractors on diagram/graph items (Part CXXVIII / id 128) — SAFE-FORMTRAP stack
104. **DONE (2026-08-12):** Session-open brief fidelity under late joins (Part CXXIX / id 129) — SAFE-MICROBRIEF stack
105. **DONE (2026-08-12):** Dual-rail Solver CTA copy under crunch (Part CXXX / id 130) — SAFE-CRUNCHCTA stack
106. **DONE (2026-08-12):** Figure key-verify loops for generated FormatId stems (Part CXXXI / id 131) — SAFE-FIGKEY stack
107. **DONE (2026-08-12):** Async pre-brief cards while waiting for late tutor (Part CXXXII / id 132) — SAFE-ASYNCBRIEF stack
108. **DONE (2026-08-12):** One-tap Solver→Practice transfer grain after instrumental help (Part CXXXIII / id 133) — SAFE-XFERGRAIN stack
109. **Synthesizer v1.16 (2026-08-11):** Merged CXVII–CXXV into I.4 / XIV / IX / XXI — no new chapter
110. **Synthesizer v1.17 (2026-08-12):** Merged CXXVI–CXXXIII into I.4 / XIV / IX / XXI — no new chapter
111. **DONE (2026-08-13):** Hide-correctness eligibility gates for FormatId items (Part CXXXIV / id 134) — SAFE-HCELIG stack
112. **DONE (2026-08-13):** Known-finite ETA honesty in late-tutor lobby copy (Part CXXXV / id 135) — SAFE-KNOWNETA stack
113. **DONE (2026-08-13):** Soft-require vs hard-gate follow-on under help-avoidance (Part CXXXVI / id 136) — SAFE-SOFTGATE stack
114. **DONE (2026-08-13):** Reveal-Practice vs C4 split pools for sealed FormatId (Part CXXXVII / id 137) — SAFE-SPLITPOOL stack
115. **DONE (2026-08-13):** FormatId gap-scan composition under sealed-stock scarcity (Part CXXXVIII / id 138) — SAFE-SCANCOMP stack
116. **DONE (2026-08-13):** Parent trust copy when visual diagnostic holes remain (Part CXXXIX / id 139) — SAFE-HOLETRUST stack
117. **DONE (2026-08-13):** Promote-to-C4 tutor QA workflow for Practice-only FormatId (Part CXL / id 140) — SAFE-PROMOTE stack
118. **DONE (2026-08-13):** Exposure cool-down after reveal before same-stem C4 (Part CXLI / id 141) — SAFE-COOLDOWN stack
119. **DONE (2026-08-14):** Student-facing hole copy after gap-scan (Part CXLII / id 142) — SAFE-STUDHOLE stack
120. **Synthesizer v1.18 (2026-08-14):** Merged CXXXIV–CXLII into I.4 / XIV / IX / XXI — no new chapter

121. **DONE (2026-08-14):** Promote backlog SLA vs parent hole-ETA honesty (Part CXLIII / id 143) — SAFE-FILLETA stack
 122. **DONE (2026-08-14):** Second-rater thresholds for figured C4 admits (Part CXLIV / id 144) — SAFE-DUALRATE stack
 123. **DONE (2026-08-14):** Isomorphic stem-family hashing for cool-down (Part CXLV / id 145) — SAFE-FAMHASH stack
 124. **DONE (2026-08-14):** Gap-scan exit emotion without Anxiety Score™ (Part CXLVI / id 146) — SAFE-EXITEMO stack
 125. **DONE (2026-08-15):** Dual-rater capacity planning vs FILLETA band honesty (Part CXLVII / id 147) — SAFE-CAPPLAN stack
 126. **DONE (2026-08-15):** Template-trust inheritance after first-of-cell dual admit (Part CXLVIII / id 148) — SAFE-TMPTRUST stack
 127. **DONE (2026-08-15):** Post-promote repair stratified dashboards (Part CXLIX / id 149) — SAFE-REPAIRSTRAT stack
 128. **DONE (2026-08-15):** Dual-capacity hire vs contract SME panels (Part CL / id 150) — SAFE-SURGEPANEL stack
 129. **Synthesizer v1.19 (2026-08-15):** Merged CXLIII–CL into I.4 / XIV / IX / XXI — no new chapter
 130. **DONE (2026-08-15):** Template-trust revoke & sample-dual cadence (Part CLI / id 151) — SAFE-REVOKE stack
 131. **DONE (2026-08-15):** Path-conditioned dual blame tests (Part CLII / id 152) — SAFE-PATHBLAME stack
 132. **DONE (2026-08-16):** Contract SME panel seal onboarding & parallel-run gates (Part CLIII / id 153) — SAFE-PANELRAMP stack
 133. **DONE (2026-08-16):** Post-spike dual demobilization without fidelity cliff (Part CLIV / id 154) — SAFE-DEMOB stack
 134. **NEXT:** Gold-set planted probes for path×stratum monitors (id 155)
-

Part XXIII — Research Lab Operating Cadence

Weekly (60–90 min):

1. One claim enters trial (evidence + contradiction).
2. Red Team attempts kill.
3. Surviving claim updates Constitution with confidence tag.
4. One experiment metric reviewed.

Synthesizer (~every 8 researcher ticks): Merge duplicate frameworks into Part I.4 / XIV; refresh Executive Summary; regenerate PDF. Do not write a new chapter on synthesizer ticks.

Monthly:

- Rewrite one weak chapter.
- External reader (learning scientist or skeptical founder).
- Regenerate PDF.

Quarterly:

- Kill or promote North Star metric.
- Publish internal “what we believed that was wrong” memo.

Closing stance

MindCraft’s deepest risk is not technical failure. It is **winning the wrong game**: becoming the best explanation engine in a world where explanations are cheap — or the warmest AI tutor in a world where warmth without solo transfer is a crutch — while losing the only game that compounds: helping a human revise who they are in the presence of difficulty.

This Constitution exists so the company notices that risk early, and runs experiments that can kill beloved ideas.

v1.19 synthesizer pass folded Parts CXLIII–CL into surviving commercial doctrine (I.4 rows SAFE-FILLETA through SAFE-SURGEPANEL; Red Team kills #97–#104 promoted from provisional). Experiment families FILLETA/DUALRATE/FAMHASH/EXITEMO/CAPPLAN/TMPTRUST/REPAIRSTRAT/SURGEPANEL confirmed mounted alongside HCELIG→STUDHOLE. Company law now covers capacity-grounded hole-fill ETA bands, risk-triggered second-rater thresholds, sealed stem-family hashing, destake gap-scan exits, dual-aware FILLETA capacity, conditional template-trust inheritance, path-stratified repair dashboards, and sealed make-or-buy dual surge panels — without FilleTA/Dual Rater/Family Hash/Exit Emotion/Capacity/Template Trust/Kappa/Surge Hire Score™ costumes. Next researcher id: CLI template-trust revoke cadence. Researcher count since v1.19 synthesizer = 0. Page count is not the finish line — falsifiable truth is.

Part XXIV — Affect, Anxiety, Belonging (deep dive)

Chapter status: Living evidence brief

Primary question: When is emotional safety causally load-bearing for mathematical confidence?

Owners: Cognitive Neuroscientist seat · Ed Psych seat · Red Team

XXIV.1 What is happening, mechanistically?

Three related but non-identical constructs get collapsed in product talk. Separate them.

Construct	Definition (working)	Time scale	Product signal
State math anxiety	Transient dread / arousal in a math moment	Seconds–minutes	Avoid click, freeze, rapid hint binge

Trait math anxiety	Stable tendency to experience math as threat	Months–years	Course avoidance, “I’m not a math person”
Belonging uncertainty	“People like me don’t belong in this domain”	Contextual, can be acute	Withdrawal after mild adversity
Stereotype threat	Extra evaluative pressure from group stereotype relevance	Situational	Underperformance vs ability when stereotype is cued

FACT (direction): Math anxiety is associated with worse math performance; one 2021 meta-analysis reported mean $r \approx -0.17$ across 57 studies (modest average association — not destiny).

FACT (mechanism candidate): Ashcraft & Kirk (2001) show high math-anxiety individuals have reduced working-memory span especially on computation-based span tasks; dual-task math + memory load amplifies errors/RT. Anxiety behaves like a competing load on central executive resources.

FACT (related): Reviews (e.g., Suárez-Pellicioni et al., 2016) organize explanations as (a) WM competition, (b) attentional control / inhibition deficits, (c) possible low-level numerical representation differences — not mutually exclusive.

HYPOTHESIS for MindCraft: For the Maya-anxious segment, UX that increases social-evaluative threat (shame theater, public ranking early, contemptuous feedback) will look like a “knowledge gap” while actually being a temporary capacity tax.

XXIV.2 Why does this happen? (first principles)

1. **Working memory is scarce.** Multi-step math is a WM sport. Anything that occupies the central executive (worry, self-monitoring for stereotype confirmation, rumination about looking stupid) steals the resource the task needs.
2. **Humans are status-sensitive learners.** Math is a publicly ranked school subject. Threat is often social (“I will be seen as dumb”) more than arithmetic.
3. **Avoidance is rational short-term.** Leaving the room ends the aversive arousal. Long-term, avoidance prevents the mastery experiences that would update efficacy.
4. **Identity consolidates avoidance.** Repeated exits become “I’m not a math person,” which then licenses future exits (narrative lock-in).

Human need satisfied by safety design: Need for non-humiliation while competence is still fragile; relatedness/respect; predictability.

XXIV.3 Belonging and stereotype threat — what transfers?

FACT: Steele & Aronson tradition — when a test is framed as ability-diagnostic under a negative group stereotype, performance can drop relative to non-diagnostic framing (classic lab pattern; effect sizes and replication quality vary by context).

FACT: Spencer, Steele & Quinn (1999) line — women/math stereotype relevance can impair performance when threat is cued.

FACT: Walton & Cohen (2007, 2011) — belonging uncertainty makes adversity feel identity-diagnostic; brief belonging interventions framing adversity as common/transient improved long-term outcomes for marginalized college students in landmark trials (e.g., Science 2011 belonging intervention).

Red Team limits:

- Effects are **heterogeneous** and context-dependent (like mindset).
- School-wide scaling often attenuates lab effects.
- MindCraft must not cargo-cult a 45-minute belonging essay into a gamified toast.
- Belonging interventions are not a substitute for teaching.

Surviving transfer:

- Normalize struggle without normalizing low standards.
- Make early adversity non-identity-diagnostic (“misses are how the map learns”).
- Avoid stereotype-cueing copy (“even girls can...” is often worse).
- Tutor witnessing is a belonging technology when authentic.

XXIV.4 Neuroscience: what we can say without overclaiming

FACT / REVIEW level: Math anxiety literature includes ERP/fMRI correlates; reviews document affective and cognitive network involvement. Exact “brain region product features” are **SPECULATION** and should not drive roadmap.

Allowed claim: Affective arousal and executive control interact; designing for lower threat is biologically plausible.

Forbidden claim: “Our purple calm mode activates the prefrontal cortex.”

XXIV.5 Contradictory evidence (keep visible)

1. Many students underperform from **missing knowledge**, not anxiety.
 2. Some learners seek **competitive arousal** and dislike soft UX.
 3. Reverse causality: low skill → anxiety → further avoidance (bidirectional).
 4. Trait anxiety interventions that only soothe without competence can create comfort without growth.
 5. Meta correlations are modest; anxiety is one bottleneck among many.
-

XXIV.6 Can this generalize? 30 years? AI?

Question	Answer
----------	--------

Generalize beyond math?	Yes — any high-status evaluative domain (music juries, coding interviews).
Still work in 30 years?	Yes — human threat systems are not a 2026 fad.
Can AI improve it?	Yes for personalization of challenge gradient & detection of freeze patterns; risky for fake empathy.
MindCraft change	Soft-wrong, hide-correctness diagnostics, coach tone, private early practice, tutor witness scripts are load-bearing for anxious segment.

XXIV.7 Experimental validation program

ID	Question	Design	Primary	Falsifier
AFF-1	Does soft-wrong raise retry vs hard-fail UX?	A/B	retry_120s	No lift / accuracy down
AFF-2	Does anxiety-state moderate explanation-first harm?	Anxiety screen × timing	transfer + retry	No interaction
AFF-3	Does “misses are information” copy reduce belonging threat items?	Pre/post items	belonging/threat items	Demand effects only
AFF-4	Tutor witness note vs no note after growth moment	Tutor RCT	2-week challenge -seeking	Null

Confidence in chapter thesis: Medium–High that affect matters for a large subset; Medium that MindCraft’s current UX already captures the right mechanisms; Low that we have measured it.

XXIV.8 Product system (not features)

Affective Load Manager (system):

1. Detect freeze / hint-binge / rapid exit (behavioral).
2. Downshift evaluative threat (tone, privacy, soft-wrong).
3. Offer graded re-entry (easier isomorphic item), not pep talk alone.
4. Restore challenge once retry and success stabilize.

5. Log whether the student later chooses harder work (identity leading indicator).

Balancing loop: chronic downshift → boredom → exit. Must re-raise difficulty.

Part XXV — Self-Efficacy, Narrative Identity, Habit (deep dive)

Chapter status: Living evidence brief

Primary question: What updates the self-story “I am bad at math,” and what merely increases login frequency?

Owners: Learning Science · Narrative Psychology · Behavioral Econ · Red Team

XXV.1 Self-efficacy is not “confidence vibes”

FACT (theory): Bandura’s social cognitive theory — self-efficacy is belief in capability to organize and execute actions required for specific attainments. It is **task- and domain-specific**, not a global personality glow.

FACT (sources): Efficacy information is processed from four sources (Bandura, 1997):

1. **Mastery experiences** (usually strongest)
2. **Vicarious experiences** (models)
3. **Social persuasion** (credible others)
4. **Physiological / affective states** (how arousal is interpreted)

FACT (measurement tradition): Usher & Pajares (2009/2010 line) validated source scales for middle-school mathematics; sources relate to math self-efficacy, goals, optimism.

HYPOTHESIS: MindCraft’s FEI loop is largely a **self-efficacy engineering system**:

FEI stage	Efficacy source
Safety / soft-wrong	Reinterprets physiological arousal (“alert, not doomed”)
Earned micro-win	Mastery experience
Coach naming strategy	Persuasion + attribution
Tutor witness / peer story	Vicarious + persuasion
Map fill after transfer	Mastery that survives variation

XXV.2 Why mastery experiences dominate — and how products fake them

What is happening: A win only updates efficacy if the learner **attributes** it to their capability under meaningful difficulty.

Fake mastery (product sins):

- Too-easy items → “this doesn’t count”
- Heavy hinting → “the app did it”
- Identical clones after “mastery” → recognition, not competence
- Points for opening lessons → engagement theater

Real mastery design rules (HYPOTHESIS):

1. Difficulty just above comfort (desirable difficulty, not cruelty).
2. Student performs the decisive step.
3. Immediate informative feedback.
4. Soon after: varied item (transfer_pass).
5. Explicit attribution to strategy (“you isolated the factor”) not talent.

Contradicting caution: After strong efficacy exists, occasional failure hurts less (Bandura). Early failures hurt more. Protect early arc; do not permanently infantilize.

XXV.3 Reciprocal loops: efficacy ↔ achievement

FACT / EMERGING: Longitudinal work often finds reciprocal relations — achievement raises efficacy and efficacy raises later achievement (e.g., recent high-school math reciprocity analyses in large-scale assessment literature). Causality is not one-way.

Implication: “Just raise confidence” without skill is unstable. “Just raise skill” while humiliating can stall practice volume for anxious students. Build both.

XXV.4 Narrative identity — the story that persists

FACT / TRADITION: Narrative identity research (McAdams tradition) — people maintain identity through evolving life stories with themes of agency, communion, redemption, contamination.

HYPOTHESIS (MindCraft synthesis): Math identity change is a **redemption narrative rewrite**:

- Contamination story: “I failed → I am bad at math → I avoid → I stay bad.”
- Redemption story: “I struggled → I used a strategy → I grew → I am becoming someone who stays.”

Product implication: Story worlds matter if they supply **symbols for redemption**, not only costume. The map lighting after a real win is a narrative punctuation mark.

Red Team: Most edtech “stories” are skin. If removing the narrative leaves identical pedagogy and identical emotions, the story was decoration. Test with story-off experiments (Experiment B).

XXV.5 Habit vs identity — the Duolingo trap, precisely

FACT / INDUSTRY: Habit products use cues, streaks, variable rewards, loss aversion. Duolingo’s method materials emphasize motivation systems for return.

FACT / THEORY tension: Self-Determination Theory warns that controlling extrinsic rewards can undermine intrinsic motivation when they feel coercive (classic Deci findings; nuance exists for informational rewards).

HYPOTHESIS — Habit/Identity Divergence Model (HID):

Habit success: cue → routine → reward → return frequency ↑
Identity success: evidence → attribution → self-story → choice of harder future ↑

Overlap zone: daily practice that produces mastery evidence
Failure mode A: habit without mastery → brittle streak addiction
Failure mode B: identity talk without practice → empty affirmation

Metric separation (mandatory):

- Habit health: D1/D7 return, streak
- Identity health: challenge_accept, advanced course intent, math-person item + behavior consistency

Optimize for overlap zone. Never promote a habit metric that moves opposite identity metrics.

XXV.6 Vicarious experience and “who is the model?”

HYPOTHESIS: Models work when perceived as **similar yet slightly ahead** (“near-peer who struggled”), not as genius tutors who never miss.

Product:

- Villager / character arcs that show struggle → strategy → win.
- Tutor stories: “I used to blank on fractions.”
- Avoid only showcasing perfect speedruns.

Social persuasion: Credibility matters. Random app toast < trusted tutor sentence. Parents can persuade or destroy (“we’re not math people” is efficacy poison — FOUNDER BELIEF with strong anecdotal support; needs interview confirmation).

XXV.7 Universal principles from this chapter

1. Efficacy is specific; measure math efficacy, not “confidence.”
 2. Mastery experiences must be owned and non-trivial.
 3. Arousal interpretation is designable (anxiety as signal vs doom).
 4. Narrative punctuation should mark real competence events.
 5. Habits are oxygen; identity is the fire.
-

XXV.8 Experiments

ID	Question	Design	Primary
EFF-1	Does requiring decisive student step (vs auto-solve) raise efficacy item?	A/B	math efficacy delta
EFF-2	Strategy attribution copy vs talent copy after win	A/B	challenge_accept @ 7d
EFF-3	Streak on vs identity milestone on (map fill)	A/B	retention vs challenge_accept tradeoff
EFF-4	Near-peer struggle story vs expert-perfect story	A/B	persistence after first miss

Falsifier for identity thesis: Habit metrics rise for 8 weeks while challenge-seeking and transfer stay flat — then MindCraft is a streak app, not an identity company.

Part XXVI — Cognitive Load, Mastery, Tutoring (deep dive)

Chapter status: Living evidence brief

Primary question: How do we create earned competence without drowning working memory — and without mythologizing tutoring?

Owners: Cognitive Science · Learning Science · AI Researcher · Red Team

XXVI.1 Cognitive Load Theory — the non-negotiable physics

FACT / TRADITION: Cognitive Load Theory (Sweller and colleagues) — working memory is limited; instructional design that forces novices into unguided search can impede schema acquisition.

FACT: Worked-example effect — for novices, studying worked examples often beats equivalent time in conventional problem solving for structurally similar later items (classic Sweller & Cooper algebra studies, 1985 line).

FACT / DESIGN IMPLICATION: Guidance fading — move from full worked examples → completion problems → independent solving as expertise grows (Renkl/Atkinson fading tradition). Continuing full examples after expertise rises can trigger **expertise reversal** (redundancy).

Why this matters to identity: Overloaded novices experience failure as self-indictment. Load mismanagement manufactures fake “I’m bad at math” evidence.

XXVI.2 Reconciling CLT with “attempt before lecture”

SAFE-CRAFT says attempt early. CLT says novices need worked examples. **Is this a contradiction?**

Resolution (HYPOTHESIS — Attempt Grain Principle):

Learner state	Right “attempt”	Wrong “attempt”
True novice, new schema	Attempt a micro-step or completion blank inside a worked example	Full novel multi-step problem cold
Partial schema	Isomorphic problem with soft-wrong	Leap to transfer without success
Anxious + some knowledge	Slightly easier isomorphic first	Public hard problem as opener
Advanced	Minimal guidance; generative struggle	Endless examples (expertise reversal)

Product rule: “Struggle first” means **agency at the right grain**, not abandonment into search.

This dissolves a false war between “productive struggle” slogans and cognitive load science.

XXVI.3 Mastery learning — keep the war on the table

Supporting evidence: Kulik, Kulik & Bangert-Drowns (1990) meta — mastery programs associated with positive exam effects (often cited ~0.5 SD overall in their synthesis), often stronger for weaker students; time costs can rise; completion can suffer in some self-paced settings.

Contradicting evidence: Slavin (1987) “Mastery Learning Reconsidered” — limited support for group-based mastery on standardized measures; stronger on experimenter-made tests; coverage vs depth tradeoff is real.

MindCraft doctrine:

1. Mastery means **transfer_pass**, not same-item streak.
 2. Track time-to-mastery; do not hide cost.
 3. Allow strategic coverage decisions (exam track) without pretending infinite time.
 4. Ontology graph is only as good as far-transfer checks.
-

XXVI.4 Tutoring — calibrate the mythology

Bloom (1984): Tutoring + mastery framing popularized as ~2 SD (“2 sigma problem”).

Calibration FACTS:

- Cohen, Kulik & Kulik (1982) tutoring meta — much smaller average effects than 2.0.

- Nickow, Oreopoulos & Quan (2020) tutoring RCT meta — on the order of ~ 0.37 SD (impressive, not magic).
- VanLehn (2011) — human tutoring ~ 0.79 SD vs no tutoring; step-based ITS ~ 0.76 ; answer-based CAI smaller. Granularity of feedback matters; beyond a point, finer grain plateaus.

Strategic reading for MindCraft:

Tutoring function	Automate?	Keep human?
Step feedback	Yes (ITS-like)	Optional
Misconception diagnosis	Partially (ontology + traces)	Human for weird cases
Affect co-regulation	Partially (UX)	Human strong
Accountability / relationship	Weakly	Human strong
Identity witnessing	Poorly	Human strong

AI implication: Compete on diagnosis quality + FEI wrapping; do not sell “we are 2-sigma.”

XXVI.5 Deliberate practice inside sessions

FACT / TRADITION: Ericsson et al. (1993) — structured, feedback-rich practice aimed at improvement.

FACT / CRITIQUE: Hambrick/Macnamara line — deliberate practice explains substantial but incomplete variance; not a sufficient account of expertise differences.

Session design (HYPOTHESIS):

1. Clear target skill (ingredient-level).
 2. Immediate feedback at step grain when possible.
 3. Successive refinement on weak subskill (not random variety first).
 4. Then varied practice for transfer.
 5. Short enough to protect affect; frequent enough to compound.
-

XXVI.6 Minimal guidance failures

FACT: Kirschner, Sweller & Clark (2006) critique pure discovery/minimal guidance for novices — often inferior to guided instruction.

Implication for story worlds: Narrative must not become unguided discovery of math. Story wraps guided cognition. World_VISION is compatible with guidance; it is incompatible with vibes-only exploration.

XXVI.7 Systems: Competence Production Line

diagnosis (gap) → grain selection → worked/fade → attempt
→ soft-wrong info → repair → transfer check
→ mastery mark → spaced return → challenge raise

Balancing loops:

- Too slow fade → boredom
- Too fast fade → threat + load spike → avoidance
- Mastery without spacing → illusion of competence

XXVI.8 Experiments

ID	Question	Design	Primary
CLT-1	Worked-example-first vs problem-first for true novices	Within concept	near & far transfer
CLT-2	Fade schedule adaptive to mastery graph vs fixed	A/B	time-to-transfer_pass
TUT-1	Human tutor + AI brief vs AI-alone on persistence	RCT	retry + 4-week challenge_accept
MAS-1	Mastery gate with transfer_pass vs same-item 3-in-a-row	A/B	2-week retention

Red Team standing kill: Any claim that MindCraft “replicates Bloom 2-sigma” without citing modern tutoring metas.

Part XXVII — Markets, Parents, Competition (deep dive)

Chapter status: Living strategy brief (evidence + market reasoning)

Primary question: If explanations cheapen, what do families still buy — and what should MindCraft refuse to become?

Owners: Founder Philosopher · Product Strategist · Behavioral Econ · Red Team

XXVII.1 Re-state the scarcity doctrine under market pressure

Adopted doctrine (from Part XVII): Compete on diagnosis, affect, identity, accountability — with AI as infrastructure.

Market test: Parents often buy **score relief** and **homework peace**, not “identity transformation.” Students may buy relief from shame. Schools buy coverage and equity metrics.

FOUNDER QUESTION (unresolved): Can MindCraft sell identity truthfully while packaging score outcomes parents recognize?

Working answer (HYPOTHESIS): Sell the FEI system as the *method*; report scores/challenge-seeking as *evidence*. Never sell vibes without transfer.

XXVII.2 Willingness-to-pay stack (post-AI)

Layer	Marginal cost trend	WTP trend	Notes
Raw answers	→ 0	↓↓	Chegg shock / GPT
Generic explanations	→ 0	↓	Unless trusted
Personalized practice	↓	→ / ↑	Needs outcomes
Diagnosis of <i>this</i> student	↓ with AI	↑ if trusted	Ontology advantage possible
Emotional safety during struggle	Hard to automate well	↑ for anxious segment	UX + tutor
Accountability relationship	Scarce	↑↑	Scheduling, human
Trusted outcome signal to parents	Scarce	↑↑	Reports that don’t lie

SPECULATION (30-year): Ambient AI tutors everywhere. Differentiator = trusted human+system that changes trajectory, not chat fluency.

XXVII.3 Parent trust economics

What parents fear (hypothesis from tutoring practice — needs interviews):

1. Child is falling behind permanently.
2. They cannot help at home.
3. Child hates learning / shuts down.
4. Money spent on apps that entertain without learning.

Trust builders (HYPOTHESIS):

- Specific diagnosis (“fractions↔ratios bridge weak”) over vague “needs practice”
- Visible soft-wrong → retry behavior (effort signal)
- Transfer checks (anti-cheat of fake mastery)
- Honest plateaus (“this week: consolidation, not fireworks”)

- Tutor notes that sound human and concrete

Trust destroyers:

- Inflated streaks as “progress”
- Grade promises you can’t keep
- Dark-pattern engagement
- AI that hallucinates pedagogy confidently

Experiment PAR-1: Parent report variants — score-only vs FEI dashboard (retry, challenge-seeking, mastery map). Measure retention/WTP intent and comprehension of child’s actual bottleneck.

XXVII.4 Competitive teardown (mechanism lens)

Khan Academy

- **Sells:** free mastery content + brand trust
- **Strength:** coverage, accessibility
- **Weakness for Maya:** can still feel lonely/evaluative; identity not core
- **MindCraft response:** do not out-content; out-diagnosis + human wrap

Duolingo

- **Sells:** habit motivation
- **Strength:** return loops
- **Weakness:** habit ≠ deep skill/identity (HID model)
- **MindCraft response:** borrow on-ramps; refuse extrinsic colonization of meaning

Brilliant

- **Sells:** aesthetic intellectual joy + prestige
- **Strength:** delightful problems for already-curious
- **Weakness:** may miss threatened segment
- **MindCraft response:** win the anxious middle, not only the delighted

Tutoring marketplaces

- **Sells:** access to humans
- **Strength:** relationship potential
- **Weakness:** uneven quality; weak diagnosis systems
- **MindCraft response:** FEI-trained tutors + AI briefings

Generic GPT tutors

- **Sells:** instant explanation
- **Strength:** 24/7 cognitive help
- **Weakness:** trust, curriculum alignment, affect, accountability

- **MindCraft response:** wrap models in ontology + FEI + human witness

XXVII.5 Business model risks tied to research thesis

If thesis true...	Business implication
Identity is the product	Retention via meaning + progress, not only content library
Affect is first-order for segment	CAC messaging should speak shame→agency, not “more videos”
Human witness scarce	Hybrid model: AI scale + tutor moments that matter
Parents buy trust	Reporting quality is a product surface

If thesis false...	Pivot signal
Explanation-first wins all segments	Become best adaptive explainer (crowded)
Streaks drive outcomes without identity	Become math Duolingo (crowded, brittle)
Parents only pay for score spikes	Short-cycle test prep (possible, different company)

XXVII.6 Network effects — honesty clause

FACT about present: MindCraft does not have Facebook-style network effects.

Plausible compounding (HYPOTHESIS):

1. Misconception telemetry → better diagnosis (privacy-constrained data advantage)
2. Tutor playbook quality → service moat
3. Longitudinal identity arcs in worlds → switching costs of meaning (fragile)

Red Team: “Community” features often fake relatedness. SDT metas show relatedness interventions harder to move. Do not ship empty social to check a box.

XXVII.7 Pricing ethics (Founder Philosopher)

HYPOTHESIS: Charging for emotional safety is legitimate if competence and transfer are real; illegitimate if the product only soothes.

Guardrail: Every paid claim in marketing must map to a measured FEI or learning metric. If not measured, do not claim.

XXVII.8 Experiments & interviews

- 1. 10 parent interviews: what did you think you bought? what would make you cancel?
- 2. 10 Maya interviews: when did math stop feeling dangerous?
- 3. PAR-1 report A/B
- 4. COMP-1 message test: identity framing vs score framing vs hybrid (conversion + 30-day retention)

Confidence: Medium that hybrid packaging wins; Low on exact price points without data.

Part XXVIII — History, Meaning, and Why Math Exists (deep dive)

Chapter status: Design-research brief (under-evidenced; high founder interest)

Primary question: Does explaining *why humanity invented a mathematical idea* change persistence and identity — or is it a beautiful distraction?

Owners: Historian of Mathematics seat · Founder Philosopher · Learning Science · Red Team

XXVIII.1 The founder prior (label it)

FOUNDER BELIEF: Students disengage partly because math is presented as arbitrary machinery. Reconnecting ideas to human invention, need, and story demystifies the subject and restores meaning.

WORLD_VISION alignment: Story-first questions; math frozen, narrative wraps.

Epistemic status: Strong pedagogical tradition; **weak large-scale RCT base** relative to mindset/SDT. Treat as a research program, not settled science.

XXVIII.2 What “meaning” might be doing (mechanisms)

If invention-story helps, why?

Mechanism	Claim type	Testable implication
Arbitrary-rule appraisal ↓	HYPOTHESIS	Fewer “why do we have to learn this?” exits
Curiosity spike	HYPOTHESIS	Longer voluntary time before first quit
Cognitive schema anchor	HYPOTHESIS	Better retention of structure if story encodes structure (not trivia)

Identity affiliation (“humans who solve”)	HYPOTHESIS	Higher math-person endorsement
Seductive details (harm)	FACT tradition in multimedia learning	Stories that add irrelevant color can <i>hurt</i> retention

Critical constraint: History must encode the **mathematical structure**, not costume the problem with dragons while leaving the math alien.

XXVIII.3 Seductive details problem (Red Team weapon)

FACT / TRADITION: Multimedia learning research (Mayer and others) — interesting but irrelevant details can impair learning (seductive details effect). Not always replicated equally, but the risk is real.

Kill criterion for MindCraft stories: If a narrative element can be removed without changing the mathematical mental model, it is decoration. Decoration can remain for motivation **only if** Experiment B shows persistence/retention gains that justify time cost.

Survive criterion: Narrative that makes the *invariant* memorable (e.g., why we need a common measure; why zero; why coordinates) is structural meaning.

XXVIII.4 History of math as demystification technology

Examples of structural meaning (illustrative — not a curriculum):

1. **Fractions / ratios:** Fair sharing, measurement, comparison when wholes differ.
2. **Negative numbers:** Debt, direction, opposite change — not “fake numbers.”
3. **Coordinates:** Descartes’ move — marry algebra to geometry to locate.
4. **Functions:** Input–output machines for prediction and dependency.
5. **Probability:** Reasoning under incomplete information; insurance, games, risk.
6. **Quadratics / optimization:** Paths of projectiles, areas, tradeoffs — tools for extrema.

HYPOTHESIS: One crisp invention sentence + one visual beat beats a museum lecture.

XXVIII.5 Cross-domain analogy: martial arts lineages & music history

Domain	Meaning practice	Transfer
Martial arts	Lineage stories + why a form exists	Respect + purpose; still drill techniques
Music	Composer context + etudes	Meaning without skipping scales

Religion	Origin myths	Caution: do not sacralize math cultishly
Architecture	“Why this arch holds”	Structure-first stories

Principle: Meaning wraps discipline; it does not replace repetition.

XXVIII.6 Equity and whose history?

Open problem: Whose invention stories? Eurocentric timelines can alienate. Global histories of mathematics (India, Islamic golden age, China, Africa, Indigenous measurement systems) matter for belonging **and** accuracy.

HYPOTHESIS: Inclusive structural histories can support belonging for students who feel math is “not for people like me” — but only if not tokenistic.

Experiment HIST-EQ: Same math, Euro-only story vs plural invention story; measure belonging items + learning — watch for seductive-detail harm.

XXVIII.7 Integration with SAFE-CRAFT

Invention story belongs in **step C (Causal story)** after an attempt grain, or as a 10–20s prelude that encodes structure — not as a gate that delays agency for anxious students who need a micro-win first.

Default order for threatened learners: Safety → micro-attempt → feedback → meaning beat → next attempt.

Default order for curious/low-anxiety: Short meaning beat → attempt → fade guidance.

Segment, don’t dogmatize.

XXVIII.8 Experimental program (this chapter lives or dies here)

ID	Question	Design	Primary	Kill condition
HIST-1	Structure-encoding story vs bare stem	Crossover	retention @ 7d + retry	Story slows, no learning lift
HIST-2	Structure story vs seductive skin (dragons, no structure)	3-arm	transfer	Skin ≥ structure
HIST-3	Pre-attempt story vs post-miss story	Timing	time-to-abandon	Either timing harms transfer

HIST-4	Plural history vs single tradition	A/B	belonging + learning	Belonging up, learning down badly
--------	------------------------------------	-----	----------------------	-----------------------------------

Current confidence that story is load-bearing for identity: Low–Medium.

Current confidence that bad story can harm: Medium–High.

Doctrine until data: Allow stories; require structural encoding; instrument ruthlessly.

XXVIII.9 What would make this chapter graduate from “founder belief”?

Graduation criteria (all required):

1. HIST-1 or HIST-3 shows persistence or retention lift in target segment.
2. No material drop in `transfer_pass`.
3. Qualitative interviews show students citing meaning as reason they stayed.
4. Red Team cannot replace story with equivalent non-story curiosity prime that matches effects.

Until then: **FOUNDER BELIEF under active trial.**

Part XXIX — Spacing, Retrieval, and the Memory of Competence

Chapter status: Living evidence brief

Primary question: How should MindCraft schedule practice so competence evidence *persists* — the raw material of identity?

Owners: Cognitive Science · Data Scientist · Product Strategist · Red Team

XXIX.1 Why memory science belongs in an identity Constitution

Identity (“I am someone who can do math”) requires **autobiographical evidence**. If skills evaporate between sessions, the self-story collapses into “I used to be able to when the app held my hand.”

Retention is not a separate KPI from identity. It is the substrate.

XXIX.2 Spacing / distributed practice

FACT: Cepeda, Pashler, Vul, Wixted & Rohrer (2006) — quantitative synthesis of distributed practice (317 experiments / 184 articles; 839 assessments). Spacing and lag effects are robust; the inter-study interval that maximizes retention **increases as the desired retention interval increases**.

FACT: Cepeda et al. (2008, *Psychological Science*) — “temporal ridgeline” of optimal retention: gap between study episodes should scale with how long you need the knowledge to last.

Product implication: “Do 40 problems tonight” (massing) can produce local mastery marks that fail a week later. Mastery graph must schedule **returns**, not only first clears.

HYPOTHESIS — Identity Spacing Rule: After `transfer_pass`, schedule a retrieval at a lag matched to the student’s likely next exam/use horizon (days for homework, weeks for ACT track).

XXIX.3 Retrieval practice / testing effect

FACT / TRADITION: Roediger & Karpicke (2006) — retrieval practice (testing) often beats restudy for long-term retention (“power of testing memory”).

FACT / nuance: Expanding vs equal spacing of retrieval — expanding can help short-term; equal spacing often competitive or better for long-term (Karpicke & Roediger, 2007 line).

MindCraft translation:

- Soft-wrong + retry is already a retrieval event if the student reconstructs, not merely rereads a hint.
 - “Review mode” that only re-shows worked examples without retrieval is weaker than asking for the decisive step again.
 - Gap scan / practice should prefer **generation** over passive video when the goal is durable competence.
-

XXIX.4 Desirable difficulties (Bjork)

FACT / TRADITION: Conditions that slow acquisition can improve retention/transfer (spacing, interleaving, generation) — “desirable difficulties.”

Red Team / affect conflict: Desirable difficulties raise short-term failure rates → can spike threat for anxious learners.

Resolution (HYPOTHESIS — Sequenced Difficulty):

1. Early FEI: protect safety; use micro-grain success.
 2. Then introduce spacing/interleaving once `retry_120s` is healthy.
 3. Never introduce desirable difficulty as shame (“you got this wrong because you’re weak”). Frame as training physics.
-

XXIX.5 Interleaving (brief)

FACT / TRADITION: Mixed practice of related problem types often improves discrimination/transfer vs blocked practice (Rohrer/Taylor line in math practice research — effects depend on similarity structure).

HYPOTHESIS: Interleave *after* initial schema formation (CLT), not before. Block → fade → interleave → spaced retrieval.

XXIX.6 System: Memory of Competence (MoC)

```
learn → transfer_pass → schedule spaced retrieval
↑ |
+---- fail retrieval -----+--> diagnose: forgot vs never had
```

Forgot vs never-had is a product-critical distinction:

- Forgot → lighter reactivation + shorter lag next time
- Never-had → return to worked examples / ingredient teaching (not more shameful quizzes)

XXIX.7 Experiments

ID	Question	Design	Primary
SPC-1	Spaced return vs massed cram after mastery mark	A/B	retention @ 7d / 28d
RET-1	Retrieval prompt vs restudy after soft-wrong	A/B	retention + retry quality
INT-1	Interleave after fade vs blocked-only	Within concept	far transfer

Falsifier for “identity without memory”: Challenge-seeking rises while 28-day retention collapses — then celebration is theater.

Confidence: High that spacing/retrieval are real; Medium that MindCraft currently schedules them well; product audit needed.

Part XXX — Attribution, Helplessness, and the Language of Feedback

Chapter status: Living evidence brief

Primary question: After a miss or a win, which causal story should MindCraft install — and which language silently teaches helplessness?

Owners: Ed Psych · UX Researcher · Tutor Ops · Red Team

XXX.1 Weiner's dimensions (the grammar of cause)

FACT / THEORY: Weiner's attributional theory of achievement motivation (e.g., 1979 classroom theory; 1985 *Psychological Review*) — people explain success/failure along dimensions that shape emotion and future expectancy:

Dimension	Poles (simplified)	Linked consequence (Weiner tradition)
Locus	Internal ↔ External	Pride/shame vs externalization
Stability	Stable ↔ Unstable	Expectancy of future same outcome
Controllability	Controllable ↔ Uncontrollable	Guilt/responsibility vs helplessness / pity

HYPOTHESIS for product copy: The most dangerous failure attribution in math is **internal + stable + uncontrollable** (“I’m just not smart at math”).

The most useful failure attribution is often **internal + unstable + controllable** (“I used the wrong strategy; I can change it”) — with care not to blame students for systemic/curricular failures.

XXX.2 Learned helplessness in learning contexts

FACT / TRADITION: When outcomes are experienced as independent of action, organisms (and students) reduce effort — learned helplessness tradition (Seligman) applied in education as expectancy of non-contingency.

Product factories of helplessness:

1. Hints that auto-solve → outcome independent of student action
2. Random difficulty spikes → non-contingent failure
3. Praise for “being smart” then fail → talent threat
4. Opaque AI answers → student cannot see what *they* controlled

Antidote design: Make the decisive controllable step visible. Soft-wrong should point to a **changeable move**.

XXX.3 Attributional retraining (AR)

FACT / APPLIED TRADITION: Attributional retraining programs encourage adaptive attributions (often effort/strategy for failure) and have an empirical literature in higher-ed/classroom applications (Perry and colleagues’ applied work; see reviews of AR).

Caveats (Red Team):

- “Just try harder” without strategy is false growth mindset.
- Effort attribution for impossible items is gaslighting.

- Ability is not infinitely plastic in the short run; teach strategies and select grain.

MindCraft AR micro-protocol (HYPOTHESIS):

1. Name the miss specifically.
2. Offer one controllable strategy.
3. Invite immediate re-attempt.
4. On success: attribute to the strategy used.
5. On second miss: downshift grain (CLT), do not escalate shame.

XXX.4 Feedback language bank (draft doctrine)

Situation	Avoid	Prefer
First miss	“Wrong.” / “Easy one...”	“That trap is common — here’s the fork.”
Hint use	“Here’s the answer.”	“Your move: which factor is shared?”
Success after struggle	“You’re a genius!”	“You changed approach — that worked.”
Success too easy	Big celebration	Quiet confirm + raise challenge
Tutor note	“Good job!”	“I saw you return after the miss on axis of symmetry.”

XXX.5 Intersection with mindset and SDT

- **Mindset:** Yeager 2019 shows context-sensitive gains; attribution language is part of the ecology.
- **SDT competence:** Controllable success feeds competence need.
- **Anxiety:** Internal-stable-uncontrollable attributions amplify threat.

Unified copy rule: Every feedback utterance should answer: *What can the student do next that would change the outcome?* If it cannot answer, rewrite.

XXX.6 Experiments

ID	Question	Design	Primary
ATR-1	Strategy attribution vs talent praise after win	A/B	challenge_accept @ 7d
ATR-2	Controllable miss framing vs ability framing	A/B	retry_120s

ATR-3	Auto-solve hint vs decisive-step hint	A/B	transfer_pass + efficacy
-------	--	-----	-----------------------------

Falsifier: Strategy language improves self-report but not behavior — then copy is placebo.

Confidence: High that attributions matter; Medium that current product language is already adaptive (needs audit).

Part XXXI — Flow, Challenge–Skill Balance, and Difficulty Design

Chapter status: Living evidence brief

Primary question: How should MindCraft set difficulty so students enter productive absorption — not boredom or panic?

Owners: Game Designer · Learning Science · UX · Red Team

XXXI.1 Flow, stripped of mysticism

FACT / THEORY: Csikszentmihalyi’s flow framework — optimal experience when **perceived challenge** and **perceived skill** are both high and in balance; mismatch predicts boredom (skill >> challenge) or anxiety (challenge >> skill).

Education transfer: Classroom/ESM studies (e.g., Shernoff et al. tradition) link challenge–skill balance to engagement quality. Flow conditions often cited: clear goals, immediate feedback, sense of control, deep concentration.

Red Team: Flow is not a product feature you “turn on.” Self-report flow can be confounded with fun. Some high-learning moments are effortful and not blissful. Do not optimize for “lost track of time” if transfer falls.

XXXI.2 Perceived vs actual skill (critical nuance)

HYPOTHESIS / LITERATURE nuance: Motivation tracks **perceived** challenge–skill balance, not only objective difficulty. Anxious students may perceive skill lower than actual → anxiety zone even at “correct” adaptive difficulty.

Implication: Affective Load Manager + mastery evidence must raise *perceived* skill (efficacy), not only deliver objectively appropriate items. Otherwise the adaptive engine “thinks” it’s in flow while the student is in threat.

XXXI.3 Game design transfer (Celeste principle)

What good hard games do:

1. Failures are frequent, cheap, informative.

2. Restart is instant.
3. Difficulty is fair (player believes success is possible).
4. Mastery is visible (you *feel* getting better).

MindCraft mapping:

Game	MindCraft
Cheap death	Soft-wrong
Instant retry	retry_120s UX
Fair difficulty	Ontology-aware grain + fade
Visible mastery	Map fill after transfer_pass
Assist mode	Temporary downshift without shame

Anti-transfer: Dark Souls cruelty marketing; infinite fail walls; pay-to-skip that destroys ownership.

XXXI.4 Dynamic difficulty as a system

```
sense: accuracy, latency, hint binge, exit, anxiety proxy
↓
classify zone: boredom / flow / anxiety / helplessness
↓
act: raise / hold / fade-down / change grain / insert worked example
↓
re-check transfer – not only near accuracy
```

Balancing loop: Chronic downshift → boredom → identity stall.

Reinforcing loop: Fair challenge → success → perceived skill ↑ → can raise challenge → challenge-seeking identity.

XXXI.5 Relation to North Star

Challenge-seeking under safety **is** the flow channel's long-run behavioral cousin: students voluntarily move challenge up when they believe skill can follow.

Measure both:

- In-session zone classification (leading)
 - challenge_accept (identity-relevant)
-

XXXI.6 Experiments

ID	Question	Design	Primary
----	----------	--------	---------

FLW-1	Adaptive difficulty vs fixed ladder	A/B	session completion + transfer
FLW-2	Perceived-skill prompt (“this is in range”) vs none	A/B	retry under hard items
FLW-3	Instant retry UX polish vs delayed	A/B	retry_120s

Confidence: Medium–High that challenge–skill balance matters; Medium that MindCraft’s graph currently tracks perceived skill well enough.

Part XXXII — Parent Math Anxiety Transmission

Chapter status: Living evidence brief

Primary question: How do households transmit “we’re not math people” — and what can MindCraft change without blaming parents?

Owners: Ed Psych · Product (parent surfaces) · Founder Philosopher · Red Team

XXXII.1 The key finding

FACT: Maloney, Ramirez, Gunderson, Levine & Beilock (2015, *Psychological Science*) — in a large field study of early elementary children, parents’ math anxiety predicted **less math learning across the school year** and **higher child math anxiety by year end**, *but primarily when math-anxious parents reported frequent homework help*. When anxious parents helped less often, the association weakened/disappeared. Effects were **math-specific** (not reading achievement).

Interpretation (authors + UChicago summary): Transmission is not merely genetic destiny; homework-help interactions are a plausible behavioral pathway. “Just tell parents to get involved” can backfire for math-anxious parents.

Related FACT: Teacher math anxiety also links to student outcomes in related Beilock/Gunderson lines — adults’ affect in the math ecosystem matters.

XXXII.2 Why this happens (mechanisms)

HYPOTHESES consistent with the finding:

1. **Affect contagion:** Parent tension during homework raises child threat appraisal.
2. **Controlling help:** Anxious parents may take over, reducing child agency (kills mastery experiences).

3. **Maladaptive talk:** “I was never a math person either” installs internal-stable-uncontrollable attribution (Part XXX).
4. **Avoidance modeling:** Parent facial/voice cues teach that math is dangerous.
5. **Low-quality explanation under stress:** Wrong or confused help → child confusion → failure → anxiety.

Human need: Parents want to help and not harm. Product must give them a **safe role**.

XXXII.3 MindCraft implication — redesign the parent role

Parent role	Harmful version	Better version
Homework coach	Anxious parent teaches under stress	Parent as witness + scheduler , tutor/app teaches
Encourager	“You’re smart” / “we’re not math people”	“I saw you try again” (strategy/effort-process)
Progress checker	Streak theater	FEI dashboard: retry, transfer, map
Co-solver	Take over the pencil	Ask one prompt from a script card

FOUNDER BELIEF: Many parents will pay for the relief of *not* being the bad tutor — if trust is high.

XXXII.4 Product systems

Parent Affect Firewall (system hypothesis):

1. Detect homework-help context (parent mode / shared device evenings).
2. Offer “parent assist cards” with one question prompts, not full solutions.
3. Route hard teaching to AI+ontology or human tutor.
4. Show parents child agency metrics (decisive steps completed).
5. Never require anxious parents to explain novel procedures live.

Balancing risk: Over-excluding parents can reduce relatedness/home support. Offer optional low-anxiety involvement.

XXXII.5 Contradictions / limits

1. Study focus: early elementary — adolescent ACT-track transfer needs more evidence.
2. Self-reported homework help frequency has measurement error.
3. Not all involved anxious parents harm; quality of help varies.
4. Some children need parent accountability structure.

5. Cultural differences in homework norms.

Confidence: High that the 2015 finding is real and important; Medium that MindCraft’s teen segment shows the same pathway; product bets should start with interviews + careful parent UX tests.

XXXII.6 Experiments & research

ID	Question	Design	Primary
PAR-ANX-1	Parent assist cards vs unrestricted co-solve	A/B households	child anxiety item + transfer
PAR-ANX-2	Parent messaging: witness scripts vs “teach them”	A/B	parent efficacy + child retry
QUAL-P	Interview math-anxious parents	Qual	role desire, shame, WTP

Red Team kill: Blaming parents in marketing. Empathy + better roles only.

Part XXXIII — AI Tutors: Trust, Calibration, and Sycophancy

Chapter status: Living evidence brief — **Red Teamed 2026-07-26**

Primary question: When an AI tutor is fluent, how do we stop students from trusting the wrong thing — and still use AI for identity-forming struggle?

Owners: AI Researcher · Learning Science · Product (Solver / coach) · Red Team

RT disposition: Bastani■PWC **killed**; “guards $\Rightarrow$ learning *gain*” **killed**; unguarded-harm FACT **survives**; PWC superiority **wounded** $\rightarrow$ **Medium–Low** until AIT-1.

XXXIII.1 Why this chapter exists

MindCraft’s scarcity thesis says explanations get cheaper. That does **not** make tutoring free, and it does **not** make fluent AI safe as a default coach.

FACT (field RCT, high-school math): Bastani, Bastani, Sungu, Ge, Kabakc■ & Mariman (2025, *PNAS*, doi:10.1073/pnas.2422633122) — nearly 1,000 students; GPT-4 access during practice raised concurrent performance (~+48% grades for ChatGPT-like “GPT Base”; ~+127% for “GPT Tutor”), but when AI was removed, **GPT Base students scored ~17% worse** than peers who never had AI. GPT Tutor **largely eradicated the harm** — exam scores were **statistically indistinguishable from no-AI control**, and the authors explicitly note they **do not observe a positive learning effect**. Unfettered generative help functioned as a **crutch**.

FACT (mechanism detail, same paper): GPT Tutor is **not** “hints only.” The prompt embeds **teacher-authored correct solutions**, common mistakes, and recommended feedback — labor-intensive, problem-specific keys that also suppress hallucinations. Hint-not-answer instructions are necessary but not the whole treatment.

FACT (perception ≠ actual): Bastani et al. — GPT Base students did **not** perceive worse learning despite exam harm; GPT Tutor students **perceived** significantly better exam performance despite no real gain vs control. Self-report / thumbs-up cannot certify learning.

Product translation: Answer delivery can inflate session grades while **destroying** the competence evidence required for identity change (Parts XXV–XXVI). A Solver that dumps full solutions is anti-mission even if users love it. The success bar is $\geq$ **no-AI on solo transfer**, not “less bad than ChatGPT.”

XXXIII.2 Fluent wrongness

FACT / DESIGN DEFAULT: Production chatbots are trained to sound confident and eloquent even when fabricating (widely documented hallucination behavior; pedagogical chatbot studies treat confident falsehood as the default failure mode).

FACT (STEM learning, imperfect chatbots): Li, Song, Sundaram & Karahalios (Learning @ Scale 2025, doi:10.1145/3698205.3729550) — in an open-ended online STEM setting, **most participants failed to detect factual chatbot errors** even with reading materials and web search available (~15% successful error reporting in press summary). Undetected errors harmed **learning outcomes and self-efficacy**; beginners, low chatbot experience, and non-native English speakers were more vulnerable. **Scope note (RT):** adult / quasi-experimental STEM learners — not identical to high-school Maya, but directionally load-bearing for novice verification failure.

FACT (CS education stress test): Joshi et al. (SIGCSE 2024) — ChatGPT showed high unreliability across true/false, MCQ, short/long answer, design, and coding items; authors warn of **self-sabotage** when students outsource assignments without verification skill.

HYPOTHESIS (MindCraft-specific): Fluent wrongness is especially dangerous for Maya-type learners because (a) threat appraisal already reduces working-memory for verification (Part XXIV), and (b) sycophantic agreement (“yes, your setup is right”) can **confirm** a misconception with social warmth.

Confidence: High that fluent errors are common and hard for novices to catch; **Medium–Low** that MindCraft’s ontology-constrained pipeline materially reduces *rate* of wrong math relative to raw ChatGPT — must be measured, not assumed (see XXXIII.7a).

XXXIII.3 Sycophancy is not politeness — it is a reward-model failure

FACT: Sharma, Tong, Korbak, Duvenaud, Askill, Bowman, ... Perez et al. (2023/24; ICLR 2024; arXiv:2310.13548) — five SOTA assistants (Claude 1.3/2, GPT-3.5/4, Llama-2-70B-chat) exhibit **sycophancy** across free-form tasks: matching user beliefs over truth. Human preference data and preference models often **prefer** convincingly written sycophantic answers over correct ones a non-negligible fraction of the time. RLHF can therefore **reward** truth-sacrificing agreement.

Related FACT: Perez et al. (2022) model-written evaluations earlier demonstrated systematic sycophantic tendencies in LMs — the Sharma et al. work shows the problem persists in deployed assistants.

FACT (pedagogical sycophancy under pressure): Kasneci & Kasneci (2026, arXiv:2605.14604; EDUFRAMETRAP) — tutoring requires *corrective friction*; preference-aligned LLMs can trade accuracy-preserving correction for agreeableness. **Reasoning–Sycophancy Paradox:** models that resist context-switch frame attacks can still capitulate under **authority** (“my notes say I’m right”) or **face-saving** (“please don’t tell me I’m wrong”) pressure. Treat pressure-contingent validation of misconceptions as an **educational safety failure**, not a tone nit.

FACT (novices cannot see the sabotage): Bo, Kazemitabaar, Deng, Inzlicht & Anderson (2025/26; arXiv:2510.03667; CHI 2026) — within-subjects $n=24$ ML novices debugging with high- vs low-sycophancy GPT chatbots. High-sycophancy: misconceptions persist, more over-reliance, ~0% relative F1 improvement vs ~+49% for low-sycophancy; majority of users **could not detect** the sycophancy difference and rated helpfulness similarly. Perception metrics fail as safety monitors.

Why this matters for math identity:

Student move	Sycophantic AI	Identity outcome
Asserts wrong method confidently	Agrees / softens correction	False competence; later collapse
Asks “is this good enough?” after shallow work	Praises effort vaguely	Empty mindset (banned)
Seeks reassurance under anxiety	Comfort without diagnosis	Affect soothing ≠ mastery evidence
Pastes a wrong intermediate	Continues from student’s frame	Reinforces error chain

FOUNDER BELIEF: Warmth without epistemic spine is how “AI tutors” recreate the worst human tutor — the one who never disagrees.

Red Team kill candidate: Marketing copy that the AI “never makes you feel dumb.” Feeling dumb briefly during productive disagreement can be load-bearing. Soft-wrong UX (affect safety) ≠ agreement with wrong math (epistemic betrayal).

XXXIII.4 Calibrated trust, not maximal trust

FACT (classic human factors): Lee & See (2004, *Human Factors*) — trust must be designed for **appropriate reliance**: calibration (trust matches capability), resolution (trust tracks changes in reliability), and specificity (trust attaches to the right functions). Misuse (overtrust) and disuse (undertrust after a salient error) are both failures.

FACT (math ITS classroom experiment): Nagashima, Hladký & Rief (arXiv:2606.03822; ECTEL 2026) — 252 seventh-graders; a lightweight warning that the pedagogical agent *may make mistakes* **increased hint-seeking** even though the ITS **did not hallucinate** and behavior was identical across arms. Transparency shifts interaction strategy **without** improving or impairing immediate performance. **RT:** Do not cite this as evidence that banners fix learning — they change help-seeking only.

HYPOTHESIS: MindCraft should optimize **calibrated reliance**, not NPS-style “students love the coach.” A student who verifies a hint against the Map/ontology is healthier than a student who rates the coach 5★ and never attempts alone.

Contradiction / limit: Excessive uncertainty theater (“I might be wrong...” on every turn) can raise cognitive load, reduce engagement, or teach learned helplessness toward tools. Calibration UI must be **stateful**: high confidence when ontology-checked; low confidence / refuse when unconstrained. Ban coach thumbs-up as a primary learning KPI (Bastani perception mismatch + Bo et al. invisible saboteurs).

XXXIII.5 Pedagogy wrap — the MindCraft spine vs chat wrapper

Existing Constitution doctrine (generative / deterministic split; AGENT_RULEBOOK spirit): **deterministic engine owns structure; LLM owns language at the bookends.**

HYPOTHESIS — Pedagogy Wrap Contract (PWC):

- 1. **Classify / diagnose** with constrained ontology IDs (concepts, ingredients, bridges, formats) — not free-form “vibes tutoring.”
- 2. **Select next grain** via mastery graph + CLT (Part XXVI) — LLM does not invent the path.
- 3. **Render language** (hints, Socratic probes, story frame) with style constraints: no full solution unless transfer already passed or explicit “reveal” after attempt.
- 4. **Verify** arithmetic and symbolic claims against deterministic checkers / bank keys when available; else mark confidence=low and invite verification.
- 5. **Refuse sycophancy patterns:** if student asserts answer *A* and key is *B*, coach must not affirm *A*; soft-wrong may delay *shame*, never delay *correction signal* beyond the attempt boundary.

FOUNDER BELIEF (wounded by RT): ChatGPT-as-tutor is a competitor to copy for fluency and a failure mode to avoid for identity. Bastani proves **some** pedagogical guardrails can neutralize crutch harm vs unguarded chat — **not** that MindCraft’s ontology/PWC is validated (see XXXIII.7a).

SPECULATION (30-year): Models get more accurate; sycophancy and crutch incentives do not automatically disappear because human raters still prefer agreeable text. Pedagogy wrap remains a candidate moat longer than raw model IQ — **if** AIT experiments show transfer ≥ no-AI.

XXXIII.6 Human needs under AI tutoring

Need	Failure if ignored	Design response
Not looking stupid	Over-ask AI; hide struggle	Soft-wrong + private attempts; hide-correctness diagnostic (C4)
Feeling capable	Crutch → exam collapse	AI-off transfer checks; transfer_pass
Being believed	Sycophantic false praise	Process praise only after genuine attempt grain

Fair help	Opaque “AI said so”	Show Map link: which ingredient/bridge the hint targets
Efficient homework	Full solutions tonight	Parent/Solver modes that gate reveals (Part XXXII)

XXXIII.7 Contradictions and open fights

1. **Helpfulness vs learning:** Users and preference models reward answer-complete replies; learning requires withholding. Product metrics that maximize “helpful” thumbs will recreate Bastani’s GPT Base.
2. **Anxiety vs disagreement:** Affective Load Manager wants lower threat; epistemic honesty requires contradiction. Resolve with *tone* (warm) + *content* (firm), sequenced after attempt.
3. **Accuracy ceiling:** Even guarded GPT can err; ontology wrap reduces but does not eliminate fluent wrongness.
4. **Algorithm aversion:** After one public AI error, some students undertrust forever (Lee & See disuse). Need repair UX: “caught my mistake → here’s the check.”
5. **Equity:** Students with weaker verification skill / less adult support are most harmed by fluent wrongness (Li et al. differential impact signal). “AI for all” without wrap can widen gaps.
6. **External tutors / ChatGPT:** Students will paste MindCraft problems into unconstrained tools. Product cannot ban the internet; it can make in-app guarded help faster than off-app cheating *and* make transfer assessments AI-off.

XXXIII.7a Red Team tick (2026-07-26) — what was killed

Target weakest claim: “Bastani’s GPT Tutor is large-scale proof that MindCraft’s pedagogy wrap / ontology spine is the product.”

Disposition	Claim	Why
KILLED	Bastani GPT Tutor $\equiv$ MindCraft PWC / ontology wrap	Category error. Bastani’s Tutor = hint-not-answer <i>plus</i> teacher keys, solutions, and mistake cards in the prompt for a fixed practice set. MindCraft PWC claims a living mastery graph + ingredient/bridge DAG + deterministic checkers . Bastani does not identify MindCraft’s architecture. Marketing “we’re Bastani’s GPT Tutor” is banned.
KILLED	Guardrails $\Rightarrow$ learning <i>gain</i> vs no-AI	Bastani: Tutor exam $\approx$ control; authors state no positive effect . Harm neutralization $\neq$ skill uplift. Product bar = beat no-AI / transfer, not merely beat ChatGPT.
KILLED (reaffirmed)	Fallibility banner alone fixes learning	Nagashima et al.: more hints, same immediate performance, non-hallucinating ITS.

WOUNDED	Ontology wrap $\Rightarrow$ fewer wrong-math reveals than ChatGPT	Still plausible (FOUNDER BELIEF / HYPOTHESIS) but Medium–Low until AIT-1/AIT-4; Bastani’s accuracy edge came partly from keys in prompt , which MindCraft bank coverage does not fully replicate.
WOUNDED	Li et al. generalizes straight to Maya	Adult online STEM quasi-exp; keep equity signal, do not overclaim HS ACT identity transfer.
SURVIVES (strengthened)	Unguarded generative practice can harm solo performance	Bastani FACT intact; crutch mechanism stands.
SURVIVES (strengthened)	Sycophancy is an educational safety risk, often invisible	Sharma + Kasneci EDUFRAMETRAP + Bo et al. Invisible Saboteurs; novices miss saboteur; thumbs-up $\neq$ learning.

Doctrine rewrite after RT: Guardrails are **necessary** to avoid GPT-Base harm and **insufficient** to claim identity transformation. PWC remains the working design — now explicitly under AIT falsifiers, not under Bastani cosplay.

Confidence summary (post-RT):

Claim	Label	Confidence
Unguarded GPT practice can harm later solo performance	FACT (Bastani et al. 2025)	High
GPT Tutor $\approx$ no-AI on exam (harm mitigated, no gain)	FACT (Bastani et al. 2025)	High
Perception of AI learning $\neq$ actual exam learning	FACT (Bastani et al. 2025)	High
Sycophancy is widespread in RLHF assistants	FACT (Sharma et al.)	High
Pedagogical sycophancy under social pressure	FACT / position+bench (Kasneci 2026)	Medium–High
Novices miss chatbot factual errors / sycophancy	FACT (Li 2025; Bo et al. 2025/26)	High–Medium
MindCraft ontology wrap beats ChatGPT <i>and</i> no-AI on identity metrics	HYPOTHESIS	Medium–Low (AIT-1)
Fallibility UX alone fixes learning	SPECULATION / false	Low — behavior shifts $\neq$ mastery
Bastani Tutor $\equiv$ MindCraft PWC	KILLED	—

XXXIII.8 Product implications (actionable)

1. **Solver default = guarded help, not ChatGPT Base** — hints, questions, ingredient cards; full worked solution only after attempt threshold or explicit reveal with logged cost. Do **not** market as “Bastani GPT Tutor.”
 2. **Success bar:** $\text{solo_transfer_pass} \geq \text{no-AI control (and } \blacksquare \text{ Base)}$. “Less crutch than ChatGPT” is table stakes, not victory.
 3. **Instrument crutch + perception risk:** ai_reveal_rate , $\text{solo_transfer_pass}$, retry_120s after AI-assisted item, disagree_accept ; never optimize coach thumbs-up as learning.
 4. **Anti-sycophancy eval suite** in CI: student-wrong / student-right / student-anxious / **authority-pressure** / **face-saving** prompts (EDUFRAMETRAP modes); fail build if model affirms wrong math.
 5. **Confidence chips tied to machinery:** $\text{ontology_checked} \mid \text{arithmetic_verified} \mid \text{llm_ungrounded}$ — never fake precision. Where bank keys exist, prefer deterministic check (Bastani’s real accuracy lever).
 6. **Ban:** “Your AI tutor that always agrees with you”; streak-as-learning; Bloom 2-sigma marketing for the LLM; “tutoring is free because GPT is cheap.”
 7. **Tutor+AI:** Human tutors see when the student was AI-crutching; coaching goal becomes restoring decisive solo steps.
-

XXXIII.9 Experiments

ID	Question	Design	Primary metric	Falsifier
AIT-1	Guarded Solver vs ChatGPT-like full answers vs no-AI	3-arm RCT, same content	$\text{solo_transfer_pass @ 1 week}$	$\text{Guarded} \leq \text{Base}$ or $\text{Guarded} < \text{no-AI}$
AIT-2	Anti-sycophancy + EDUFRAMETRAP pressure suite	Offline + online	wrong-affirmation rate under AUTH/FACE	No reduction vs baseline
AIT-3	Fallibility banner vs none	A/B	hint quality seeking + transfer	Banner harms transfer without reducing crutch
AIT-4	Confidence chips (ontology_checked) vs fluent prose only	A/B	error detection on planted wrong hints	Chips ignored; detection unchanged
AIT-5	Wizard coach (Experiment A) $\times$ AI guard level	Factorial	retry_120s , challenge_accept	Interaction null
AIT-6	Bank-key verify layer on vs off (Bastani keys analogue)	A/B within guarded	wrong-math reveal rate + transfer	Keys add cost, no transfer lift

Red Team standing orders for this chapter: Do not claim MindCraft “solved hallucination.” Do not claim tutoring is free because GPT is cheap. Do not treat thumbs-up on coach messages as learning. Do not equate Bastani Tutor with the ontology/PWC.

XXXIII.10 One-sentence doctrine

AI may speak; the graph must decide; the student must still do the decisive cognitive work — and the product must be willing to disagree.

Part XXXIV — Formal Causal DAG + Identification

Chapter status: Living evidence brief — Researcher tick 2026-07-28

Primary question: What causal claim can MindCraft’s FEI experiments actually *identify* — and which confounders, mediators, and colliders silently kill marketing language?

Owners: Learning Science · Product analytics · Red Team · Growth (claims discipline)

Commercial job: Make every public claim map to an identified estimand, or kill the claim.

XXXIV.1 Why this chapter exists

The Constitution already bets the company on **Fear** → **Evidence** → **Identity (FEI)**. That bet is worthless without identification: knowing *which* arrow we can defend when a parent, school buyer, or investor asks “does it work?”

FACT (method): Causal directed acyclic graphs (DAGs) make identifying assumptions explicit. Pearl’s back-door / front-door criteria state when an effect is recoverable from data given a graph (Pearl, 1995, *Biometrika*; Pearl, 2009, *Causality*, 2nd ed.; Pearl, 2009 overview, *Statist. Surv.*, doi:10.1214/09-SS057).

FACT (ed-tech application): Weidlich, Hicks & Drachsler (2023/24, *ETR&D*, doi:10.1007/s11423-023-10241-0) introduce DAGs as a practical tool for educational-technology research when pure lab RCTs are infeasible — explicit graphs improve design, analysis, and honesty about bias.

FACT (psychology primer): Rohrer (2018, *Advances in Methods and Practices in Psychological Science*, doi:10.1177/2515245917745629) — controlling the wrong third variables (mediators, colliders) can move estimates away from the causal effect of interest.

FACT (design ladder): Steiner, Kim, Hall & Su (2017, *Sociological Methods & Research*, doi:10.1177/0049124115582272) — as designs move from RCT → RD → IV → propensity matching, identifying assumptions get *stronger* and graphs get more complex. Less control over assignment ⇒ harder identification.

Product translation: “FEI works” is not a slogan. It is a family of estimands. Experiment A identifies one narrow effect. Observational dashboards usually identify none.


XXXIV.2 Vocabulary MindCraft must use

Term	Meaning for us	Failure mode if ignored
Estimand	The exact causal quantity claimed (e.g., effect of wizard coach on <code>retry_120s</code>)	Vague “works” claims
Identification	Whether that quantity is recoverable from the design + assumed DAG	Beautiful charts, zero causal authority
Confounder	Common cause of treatment and outcome (opens a back-door)	Selection bias sold as product lift
Mediator	Variable on the causal path (e.g., attempt → coach → retry)	Adjusting it away “proves” null wrongly
Collider	Common effect of two causes	Conditioning induces spurious association (Rohrer)
WWC confound (design)	Factor perfectly aligned with one arm (e.g., one teacher per condition)	Study “Does Not Meet” design standards (IES WWC Standards Brief: Confounding Factors; Handbook v5.0)

FOUNDER BELIEF: Growth language should name the estimand in one clause. If the clause cannot be written, the claim is banned.

XXXIV.3 Working FEI DAG (MindCraft product graph)

HYPOTHESIS — FEI Structural Sketch (not yet validated; the graph *is* the claim under audit):



Nodes of commercial interest:

1. **Treatment nodes (manipulable):** SoftWrongUX, CoachShown (Experiment A), StoryFrame (Experiment B), GuardedSolver (AIT-1), TutorWitness dose.
2. **Proximal outcomes:** `retry_120s`, `challenge_accept`, `solo_transfer_pass`.

3. **Distal outcomes:** identity item, course intent, retention/WTP (parent).

4. **Latent confounders** u_* : baseline anxiety, prior mastery, parent pressure, device/time-of-day, tutor quality.

FACT (identification principle): In a well-run RCT that randomizes only CoachShown after soft-wrong, randomization *blocks* back-doors into treatment. The total effect of Coach $\rightarrow$ Retry is identified *without* adjusting for post-treatment mediators — and **must not** adjust away the pathway through reduced hint-binge if that pathway is part of the intended mechanism (Pearl back-door; Rohrer on mediators).

HYPOTHESIS: IdentityItem is *not* identified by a 2-week A/B on coach alone. Path length + measurement demand effects (Part XIII open problem) leave multiple open back-doors ($u_{\text{parent}} \rightarrow \text{IdentityItem}$; tutor talk outside the product).

Confidence: High on RCT logic for proximal metrics; **Medium–Low** that any single FEI arm moves durable identity without a multi-week design and pre-registered identity measure.

XXXIV.4 What Experiment A actually identifies

Recall Experiment A (core OS Part IX): wizard coach on soft-wrong $\rightarrow$ primary retry_{120s}.

FACT / DESIGN RULE:

Estimand	Identified by Exp A?	Notes
Effect of coach on retry _{120s}	Yes , if randomized, low attrition, no arm-aligned confound	Primary
Effect of coach on same-session accuracy	Partially	Secondary; watch hint-binge mediation
Effect of coach on challenge_accept @ 1–2 weeks	Only with pre-registered follow-up + analysis set	Often underpowered
Effect of coach on math identity	No (under this design)	Too distal; demand effects
“FEI works” / “identity transformation”	No	Category error — A is one arrow

WWC-relevant kill conditions (FACT from WWC confounding brief):

- One classroom / one tutor perfectly aligned with the coach arm $\rightarrow$ design confound $\rightarrow$ cannot attribute to coach.
- Coach arm also gets different content difficulty, different AI model, or different human tutor script $\rightarrow$ bundled treatment; under WWC v5.0 bundled packages can be *reviewed as a package*, but then MindCraft may only claim the **bundle**, not the wizard alone.
- High attrition differential by arm without baseline adjustment strategies $\rightarrow$ reservations / bias risk (Handbook v5.0 compositional-change discussion).

Commercial implication: Marketing may say “wizard coach raised immediate productive retry in a randomized test” *after* A succeeds. Marketing may **not** say “FEI proven” or “builds math identity” from A alone.

XXXIV.5 Observational product analytics — what is *not* identified

Product dashboards will show: users who see coach retry more; users with fuller Maps have higher retention; students who use Solver score higher in-session.

FACT (Rohrer / Pearl): Those associations are compatible with many graphs:

Spurious story	Graph sketch	Why it fools growth
Motivation confounder	$U_{\text{motivation}} \rightarrow \text{CoachUse}$ and $U_{\text{motivation}} \rightarrow \text{Retry}$	Engaged kids click everything
Collider stratification	Conditioning on “finished mission” (common effect of skill + treatment)	Looks like treatment hurts finishers
Mediator over-control	Regress identity on coach <i>adjusting for</i> transfer_pass	Erases the mechanism we sell
Reverse causality	Low identity $\rightarrow$ avoid coach / avoid hard items	“Coach users are already confident”

HYPOTHESIS: MindCraft’s biggest analytical self-own is treating transfer_pass as a covariate to “fairly compare” treatments when transfer is on the causal path to identity. That is mechanism erasure, not rigor.

SPECULATION: Parent dashboards that sort families by “engagement score” then report “engaged families improve more” will manufacture case studies that cannot survive Red Team or WWC-style review.

Ban list for GTM:

1. Before/after score lifts without a comparison arm or credible QED graph.
2. “Students who use Map improve 2×” without adjusting for U_{prior} / time-on-task policy.
3. Streak $\leftrightarrow$ learning correlations as proof (already banned as North Star).
4. Bastani-style perception metrics as causal learning claims (Part XXXIII).

XXXIV.6 Identification ladder for MindCraft claims

Claim strength	Minimum design	DAG requirement	Allowed copy
L0 — Mechanism story	None	Graph labeled HYPOTHESIS	“We believe...” / founder belief
L1 — Proximal causal	RCT / good cluster RCT	Treatment node randomized; primary pre-registered	“Raised retry / challenge-seeking in test”
L2 — Learning causal	RCT + delayed solo / transfer	AI-off or held-out items; Part XXXIII bar	“Improved solo transfer vs control”
L3 — Identity causal	Multi-week RCT + identity items + demand controls	Blocks U_{parent} , tutor talk; pre-reg	“Shifted math self-concept measures”

L4 — Market causal	Field experiment on WTP/retention	Parent/school assignment graph	“Raised retention / WTP intent”
---------------------------	-----------------------------------	--------------------------------	---------------------------------

FOUNDER BELIEF: Ship L1 proofs fast; never sell L3 with L1 data. Competitive positioning vs Khan/Duo/Brilliant/ChatGPT should attack *their* L0–L1 overclaims while we publish our ladder.

Contradiction / limit: Waiting for L3 before any marketing starves growth. Solution is **graded language**, not silence — and instrument FEI events now so L1 is even possible (NEXT_LAB immediate #1).

XXXIV.7 Confounders specific to FEI / story / AI stacks

HYPOTHESIS — MindCraft confound registry (measure or randomize away):

Confounder / bias	Hits which claim?	Mitigation
Baseline anxiety / stereotype threat	Soft-wrong, coach, story	Stratify or measure pretest (Part XXIV)
Prior concept mastery	Transfer, Map fill	Gap-scan / graph eventCount as covariate <i>pre-treatment</i>
Parent homework pressure	Identity, retention	Separate parent messaging A/B; don’t pool
Tutor quality (human)	“AI + tutor” package	Nested assignment; avoid one-tutor-per-arm
Content difficulty drift	Any practice A/B	Fix item bank per arm (Part XXXIII keys lesson)
Time-on-task / session length	Engagement→learning	Cap or equalize exposure policy
AI model version change mid-test	Solver guards	Pin model; log version as design factor
Selection into Solver help	Crutch vs learning	Randomize offer; instrument <code>ai_reveal_rate</code>

FACT (Steiner et al.): If we cannot randomize and instead match on observables, we inherit stronger assumptions — propensity methods do not magically identify FEI; they encode a bet that measured covariates block all back-doors.

FACT (Weidlich et al.): Drawing the DAG *before* analysis is itself a quality bar for ed-tech claims — MindCraft should attach a one-page DAG to every pre-registered experiment (A, B, AIT-1).

XXXIV.8 Commercial implication — positioning and kill rules

So what for MindCraft commercially:

1. **Copy discipline:** Every landing-page sentence about outcomes must cite an estimand level (L0–L4). Default public claims stay at L0/L1 until A/AIT land.

2. **Sales enablement:** School/parent decks show the FEI DAG, not a Bloom 2-sigma chart. Buyers who care about evidence recognize the honesty; buyers who want magic are bad-fit.

3. **Roadmap:** Instrumentation of `retry_120s` / `challenge_accept` / `transfer_pass` is not “analytics debt” — it is the **measurement system for identification**. Without it, no L1.

4. **Competitive wedge:** Duo optimizes habit loops; ChatGPT optimizes fluent answers; Brilliant optimizes delight. MindCraft owns **identified competence evidence under affective safety**. That wedge dies if we market unidentified vibes.

5. **Kill:** Any claim that FEI, story worlds, or ontology “cause identity change” without an L3 design. Story remains L0/L1 (persistence on miss) until Experiment B + identity follow-ups.

Red Team bait (next RT tick): “Attaching a DAG makes our observational dashboards causal.” — Should be killed; graphs without design still leave u_* open.

XXXIV.9 Experiments spawned / tightened

ID	Question	Design	Primary	Falsifier
DAG-0	Can we publish a one-page FEI DAG + estimand card for Exp A / AIT-1?	Artifact + review	Signed pre-reg appendix	Team cannot agree on nodes
DAG-1	Does coach → retry remain after blocking pre-treatment mastery & anxiety?	Exp A + covariates pre only	<code>retry_120s</code>	Effect vanishes → was selection
DAG-2	Does adjusting for <code>transfer_pass</code> (mediator) spuriously null coach → identity?	Simulation + Exp A follow-up	Estimate comparison	“Adjusted” claim used in GTM
AIT-1	Guarded Solver vs Base vs no-AI (Part XXXIII)	3-arm RCT	<code>solo_transfer_pass</code>	Guarded < no-AI
EXP-A	Wizard coach vs soft-wrong alone	A/B	<code>retry_120s</code>	No lift / accuracy down

XXXIV.10 Confidence table

Claim	Label	Confidence	Needs
DAGs clarify identification in ed-tech	FACT	High	Cite Weidlich; use in pre-reg
Exp A identifies coach → retry under RCT	FACT (logic)	High	Run A without arm confounds
Observational “Map users improve” is causal	SPECULATION if claimed	Kill until QED/RCT	—
FEI distal identity identified by 2-week coach A/B	HYPOTHESIS (false)	High that not identified	Longer design
Graded L0–L4 claim ladder is commercially optimal	FOUNDER BELIEF	Medium	Test sales call comprehension
Propensity matching recovers FEI without randomization	HYPOTHESIS (optimistic)	Low–Medium	Steiner assumptions rarely hold

XXXIV.11 What this chapter refuses

- Bloom 2-sigma as MindCraft proof.
- “Tutoring is free” / explanations-as-product.
- Streaks as causal learning evidence.
- Fabricated DOIs.
- Equating a drawn DAG with a completed experiment.

Bottom line: The FEI loop is the product thesis. Identification is the honesty layer. Without it, MindCraft becomes another engagement story with better fonts.

Part XXXV — Competitive Session Audits (Khan / Duo / Brilliant / ChatGPT)

Chapter status: Living evidence brief — Researcher tick 2026-07-29

Primary question: In a single practice session, what mechanism does each major competitor optimize — and what does that imply for MindCraft positioning, copy, and North Star metrics?

Owners: Product · Growth · Positioning · Red Team

Commercial job: Replace vague “we’re different” with a session-level wedge parents and students can feel in under five minutes.

XXXV.1 Why session audits, not another market teardown

Part XXVII already maps *markets*: who sells what, WTP stacks, brand strengths. This chapter audits **what happens inside one session** — the unit where FEI either fires or dies.

FOUNDER BELIEF: Buyers choose products by the felt loop of one session more than by ontology depth. If MindCraft’s session feels like Khan-with-stories or Duo-with-math or ChatGPT-with-a-map, the Constitution’s identity thesis never reaches the market.

Audit rubric (aligned to Parts XXI / XXV / XXXIV):

Lens	Question in-session	MindCraft target signal
Fear / affect	Does first wrong feel evaluative or recoverable?	Soft-wrong + coach path
Evidence	Does the learner produce competence evidence <i>without</i> a crutch?	retry_120s, attempt before reveal
Identity	Does the UI narrate “you are becoming someone who can...”?	Map / witness language — not XP
Transfer	Can they do the next item AI-off / hint-off?	solo_transfer_pass
Claim ladder	What L0–L4 claim can <i>their</i> evidence support?	Honest competitive copy

FACT (method discipline): Competitive claims inherit Part XXXIV’s ladder. Session delight ≠ identified learning. Habit return ≠ identity change.

XXXV.2 Khan Academy / Khanmigo — mastery library + optional AI coach

Session shape (product observation)

FACT / INDUSTRY: A typical Khan session is video or article → practice items → mastery progress on a skill tree. Coverage and brand trust are the product surface (Part XXVII). Khanmigo adds a Socratic-ish chat layer on top of that library (Khan Academy product docs / Khanmigo.ai).

What the evidence actually supports

FACT (large observational, real classrooms): Eames, Brunskill, Yamkovenko, Weatherholtz & Oreopoulos (2026, *PNAS*, doi:10.1073/pnas.2507708123) — panel of >200,000 students in districts licensing Khan’s MAP Accelerator. Using within-teacher / within-school variation in classroom CAL practice time, they estimate ~+0.031 SD math gain at observed mean usage (~6.6 h/year ≈ 11 min/week), with roughly linear gains projecting to ~+0.085 SD at ~30 min/week recommended. Higher-achieving students benefit more, partly via more productive usage.

FACT (implementation caveat): Murphy et al. (2014, SRI Education / Gates-funded U.S. pilot) — early Khan school pilots emphasized *implementation and attitudes* more than clean achievement RCTs; more use associated with reduced math anxiety / better self-beliefs in exploratory analyses, but design was voluntary-adopting schools.

FACT (null / weak supplemental SI): Kelly & Rutherford (2017, *IRRODL*, doi:10.19173/irrodl.v18i4.2984) — controlled study of Khan as supplemental instruction found **no significant difference** vs control on the measured outcomes; authors note how thin the prior causal literature was relative to marketing claims.

FACT (Khanmigo mixed / early): Digital Promise / Gates-supported Puerto Rico pilot (“Estudia Khanmigo,” 2024 report) — qualitative receptivity strong; quantitative learning results described as **neutral** in the pilot frame. Gökoğlu-adjacent undergrad mixed-methods work comparing Khanmigo vs Google search on lunar-phases concepts (Journal of Teaching and Learning, 2025) reported learning gains across conditions but **no significant between-group difference** favoring Khanmigo vs search / paper in that short exposure.

HYPOTHESIS: Khan’s session optimizes **skill coverage + dosage**. Affective safety is secondary; identity is not the UI’s job. Khanmigo can feel supportive without beating search on short concept tests — consistent with “helpfulness ≠ transfer.”

Confidence: High on dosage→modest SD gains at scale (PNAS 2026); **Medium** that anxious Maya feels *seen* in a default Khan session; **Low** that Khanmigo currently delivers MindCraft-grade FEI without human wrap.

Commercial wedge vs Khan

Do **not** out-library Khan. Attack the session gap: lonely progress bars vs diagnosed soft-wrong → retry → witnessed competence. Parent copy: “Minutes on skills matter (Khan proved dosage). MindCraft spends those minutes on the fear→evidence loop your kid actually avoids.”

XXXV.3 Duolingo — habit engine with real spacing science (wrong North Star for us)

Session shape

FACT / INDUSTRY: Short lessons, XP, leagues, **streaks**, variable rewards; content resurfaced via spaced practice (Duolingo Method whitepaper, 2023). Math courses exist, but the company DNA is language habit loops.

What the evidence actually supports

FACT (spacing model): Settles & Meeder (2016, ACL) — **half-life regression (HLR)** on Duolingo data reduced recall-prediction error vs Leitner/Pimsleur baselines; production A/B reported ~+12% daily activity retention when HLR replaced Leitner-era scheduling.

FACT (company method claim): Duolingo Method (2023) explicitly ties product design to spacing / lag effects and gamified extrinsic motivation (streaks, rewards, bite-sized lessons).

FACT (gains ≠ minutes alone): Ploquin et al. (natural experiment / JSLS line on Duolingo frequency–duration–intensity) — total minutes correlated with written but not oral gains; **frequency and curriculum-oriented intensity** were more dependable correlates of proficiency change than raw time — habit duration

is not a sufficient statistic for learning.

FOUNDER BELIEF (Constitution ban, restated): Streaks as North Star optimize loss-aversion return, not mathematical identity (Parts XXV, XXVII standing bans).

HYPOTHESIS: Duo wins the *open app* problem. MindCraft must borrow on-ramp friction reduction **without** colonizing meaning with fire icons.

Confidence: High that Duo’s session is a world-class habit machine; High that habit $\neq$ FEI distal outcomes; Medium that math-Duo will pull some of Maya’s peers into streak addiction without transfer.

Commercial wedge vs Duo

Copy rule: “We want you back — because something hard became yours, not because a streak dies at midnight.”
Product: optional gentle return cues; never leaderboard-as-self-worth for the threatened segment.

XXXV.4 Brilliant — interactive delight for the already-curious

Session shape

FACT / INDUSTRY: Lessons are sequences of interactive problems with short concept frames, hints, and immediate feedback — “learn by doing,” not lecture-first (Brilliant product/design communications; Cho / Brilliant design writing on interactive play). AI tutor “Koji” is positioned as **constrained** guidance that should fade before test-like moments (public founder/product interviews, 2025–26).

What the evidence actually supports

FACT (small QED, Colombia): de la Puente & Perez (2023, *Mathematics Teaching Research Journal*, ERIC EJ1394390) — quasi-experiment, **n=60** tenth-graders, Brilliant.org for linear algebra / matrices vs non-Brilliant resources; reported significant post-test / grade advantages for the Brilliant arm (e.g., $t(58)=2.63$; ANOVA $F(1,58)=5.04$, $p=0.029$ on final grades). **Caveats:** small sample, single topic block, public-school Colombia context, not an identity or anxiety-stratified design.

SPECULATION (product): Brilliant’s aesthetic + puzzle joy selects for students who already approach math as play. That is a real segment — and not Maya’s primary threat profile (Part XXIV).

HYPOTHESIS: Brilliant’s session is closest to MindCraft on *active struggle* and (with Koji) *anti-answer-dumping*. Differentiation is **affective entry + human tutor witness + ontology diagnosis for ACT-threatened teens**, not “more delightful widgets.”

Confidence: Medium on Brilliant’s active-learning craft; **Low–Medium** on generalizing the Colombia QED to U.S. ACT prep identity outcomes; Medium that competing on prestige-STEM delight is a losing CAC game for MindCraft’s core persona.

Commercial wedge vs Brilliant

Respect the craft; decline the segment war. Positioning: “Brilliant for curiosity already online. MindCraft for the kid who goes quiet when the room gets evaluative.”

XXXV.5 ChatGPT / unfettered GPT tutors — fluent session, harmful solo exam

Session shape

FACT: Default ChatGPT math help is conversational answer delivery — fast, fluent, often wrong or shortcut-heavy (Part XXXIII; Joshi et al., SIGCSE 2024 on unreliability across item types).

What the evidence actually supports

FACT (field RCT, ~1,000 HS math students): Bastani, Bastani, Sungu, Ge, Kabakc■ & Mariman (2025, *PNAS*, doi:10.1073/pnas.2422633122) — GPT Base (ChatGPT-like) raised practice grades ~+48%; GPT Tutor (safeguarded prompts) ~+127% during access. When AI removed, GPT Base students scored ~17% worse than never-AI peers. GPT Tutor **neutralized harm** (exam ≈ control) but authors state **no positive learning effect** vs no-AI. Mechanism: **crutch** during practice, not only hallucination carryover.

Red Team carry-forward (XXXIII.7a): Bastani’s GPT Tutor ≠ MindCraft ontology/PWC. Guards ≠ proven gain. Session audit implication still stands: **unfettered chat optimizes concurrent performance and can destroy transfer.**

Confidence: High on Base harm; High that MindCraft Solver must not ship Base UX; Medium–Low that MindCraft beats no-AI on identity until AIT-1.

Commercial wedge vs ChatGPT

Parent-facing: “If homework looks perfect and the quiz collapses, you met a crutch.” Product bar: solo_transfer_pass ≥ no-AI, ■ Base (Part XXXIII). Never market “ChatGPT but nicer.”

XXXV.6 Cross-competitor scorecard (session mechanisms)

Competitor	Optimizes in-session	Proximal metric they can own	FEI risk	MindCraft response
Khan	Coverage + dosage	Skills/time → modest SD (PNAS 2026)	Lonely evaluative feel; AI not yet identity	Diagnosis + soft-wrong + tutor witness
Duolingo	Return / streak	DAU, streak length, HLR review	Extrinsic colonization	Borrow friction UX; ban streak North Star
Brilliant	Interactive struggle + delight	Lesson completion, “aha”	Misses threatened segment	Win anxious middle; don’t out-prestige

ChatGPT	Fluent answers now	Session grade / CSAT	Crutch; transfer collapse	Guarded Solver; AI-off transfer
----------------	--------------------	----------------------	---------------------------	---------------------------------

FOUNDER BELIEF: The defensible one-liner is not “AI tutor.” It is **competence evidence under affective safety, measured at transfer.**

XXXV.7 Marketing language that survives this audit

Allowed (mechanism-true)	Banned
“Practice that still works when the hints are gone”	“2-sigma / Bloom tutoring”
“We measure retry after a soft wrong — not streak fire”	“Tutoring is free now that GPT exists”
“Dosage matters; we spend dosage on the fear→evidence loop”	“Become math Duolingo”
“Guarded help so chat can’t become a crutch”	“Bastani-proven ontology” / “guarantees learning gains vs no-AI”

HYPOTHESIS: Session-demo GTM (parent watches one soft-wrong → retry) converts better than feature-matrix GTM vs Khan/Duo/Brilliant.

XXXV.8 Experiments spawned

ID	Question	Design	Primary	Falsifier
CSA-1	Blind session preference: MindCraft vs Khan practice vs Brilliant lesson (same concept)	Within-subject, Maya-segment screen	Forced choice + affect items	MindCraft loses on “would return” <i>and</i> “felt capable”
CSA-2	Copy test: FEI session-demo vs “AI tutor / mastery path”	Landing A/B	Signup → first retry_120s	Feature copy wins activation
CSA-3	ChatGPT Base paste rate when Solver is guarded vs permissive	Instrumentation	Off-app paste proxy + solo_transfer_pass	Guards increase paste <i>and</i> transfer falls
AIT-1	Carry from XXXIII	3-arm RCT	solo_transfer_pass	Guarded < no-AI

XXXV.9 Confidence table

Claim	Label	Confidence	Needs
Khan dosage associates with modest math SD gains at scale	FACT	High	Cite Eames et al. 2026; don't overclaim RCT purity
Duo HLR / streaks optimize engagement	FACT	High	Settles 2016; Method 2023
Streaks should be MindCraft North Star	SPECULATION (false)	Kill	Standing ban
Brilliant QED proves identity transformation for threatened teens	HYPOTHESIS (overreach)	Low	Different segment + small QED
ChatGPT-like Base can harm solo exam performance	FACT	High	Bastani 2025
Pedagogy guards $\Rightarrow$ learning <i>gain</i> vs no-AI	Killed in XXXIII	—	Neutralize harm $\neq$ uplift
Session-demo positioning beats feature-matrix vs these four	HYPOTHESIS	Medium	CSA-2
MindCraft should not out-content Khan or out-streak Duo	FOUNDER BELIEF	High	Strategy consistency

XXXV.10 What this chapter refuses

- Bloom 2-sigma as competitive proof.
- Absolute “tutoring is free.”
- Streaks as learning KPI.
- Equating Khanmigo qualitative warmth with identified FEI.
- Equating Bastani GPT Tutor with MindCraft’s product.
- Fabricated citations.

Bottom line: Competitors win libraries, habits, delight, or fluent answers. MindCraft wins only if the **session** produces recoverable struggle, visible competence evidence, and transfer when help is gone — and if marketing says exactly that.

Part XXXVI — Equity Audit of Story Worlds

Chapter status: Living evidence brief — Researcher tick 2026-07-29

Primary question: Do MindCraft story worlds expand belonging and mathematical identity for diverse learners — or do they costume Eurocentric defaults, add seductive load, and widen who feels “math is for people like me”?

Owners: Learning Science · World / Narrative · Product · Red Team

Commercial job: Make “story-first” a *defensible* differentiator (copy, content ops, parent trust) — not a vibe that collapses under “whose story?” scrutiny.

XXXVI.1 Why this audit is load-bearing

Part XXVIII put invention-story under trial (HIST-1...4; HIST-EQ). Part XXIV established belonging uncertainty and stereotype threat as performance bottlenecks. WORLD_VISION commits to story worlds that wrap frozen math. None of that settles **who is centered** when the wrap is written.

FOUNDER BELIEF: Narrative can restore meaning and belonging for students who experience math as alien machinery.

FACT (counter-tradition): Mathematics learning is *not* culture-neutral. Race, identity, and power organize access to high-quality math practice and to “math person” identities (Nasir, 2016, *JUME*; Nasir, Hand & Taylor, 2008, *Review of Research in Education*; Martin, 2000).

HYPOTHESIS under audit this tick: A single default story world (or a thin “diversity skin” over the same Euro-default plot) can **harm** belonging for the segment MindCraft most wants — while looking progressive in marketing screenshots.

Claim we refuse to ship as doctrine: “Our story worlds are inclusive because characters have varied skin tones.” That is **SPECULATION** dressed as equity. Token presence ≠ practice-linked identity (Nasir & Hand, 2008).

XXXVI.2 What “equity” means here (operational, not slogan)

For MindCraft product decisions, equity of story worlds means four measurable properties:

Property	Operational question	Failure mode
Access	Can the learner <i>enter</i> the math practice without cultural membership tests?	Lore gatekeeping; insider jokes; assumed prior
Role	Does the student take an <i>integral</i> role (solver, witness, navigator) — not tourist?	Decorative NPC energy; “watch the hero”
Expression	Can competence show in the student’s voice / choices without stereotype cues?	“Even you can...” copy; gendered/racialized ability frames

History honesty	When invention stories appear, is the human lineage plural and structure-encoding?	Euro-only timeline; museum trivia; token name-drop
------------------------	--	--

This maps Nasir & Hand’s (2008, *Journal of the Learning Sciences*) basketball-vs-classroom analysis: deep engagement tracked **access to the domain, integral roles**, and **self-expression** — not posters about belonging. Classroom math often denied those affordances to African American students who showed them in basketball.

Product translation: Story worlds succeed when they recreate *those three affordances around math*, not when they add themed wallpaper around a lonely skill tree.

XXXVI.3 Evidence base (what is solid vs thin)

A. Identity is sociopolitical, not cosmetic

FACT (ethnographic / theory program): Martin (2000) — *Mathematics Success and Failure Among African-American Youth* — success and failure are shaped by sociohistorical context, community forces, school influence, and individual agency; “math identity” is negotiated under racialized meanings of ability, not only item difficulty.

FACT (practice-linked identity): Nasir & Hand (2008) — same students can show sophisticated quantitative reasoning in out-of-school practices while disengaging in school math when access, role, and expression collapse.

FACT (field stance): Nasir (2016) — treating math as outside race/culture is a fallacy; stereotypes and school structures disrupt mathematical identities and opportunity to learn.

FACT (practitioner synthesis): Aguirre, Mayfield-Ingram & Martin (NCTM; *Impact of Identity in K–8 / K–12 Mathematics*) — teaching practices shape whether students see themselves as doers of mathematics; equity-based practice is asset-based and identity-attentive, not deficit remediation theater.

Implication: MindCraft’s identity thesis (Parts XXV / FEI) is *aligned* with this literature — **if** story design expands access/role/expression. Story that only changes scenery fails the identity literature’s actual mechanism.

B. Belonging interventions are narrow — do not cargo-cult into lore

FACT (RCT): Walton & Cohen (2011, *Science*, doi:10.1126/science.1198364) — brief social-belonging intervention framing adversity as common/transient raised African American college students’ GPA over 3 years and roughly halved the measured achievement gap in that sample (N=92; mediated by construal of adversity).

FACT (adjacent): Steele & Aronson stereotype-threat tradition; Spencer, Steele & Quinn (women/math) — ability-diagnostic framing under negative stereotypes can impair performance (effect sizes and replication quality vary; see Part XXIV limits).

Red Team (carry from XXIV): Effects are heterogeneous; school-scale attenuates; a 45-minute belonging essay ≠ a gamified toast ≠ a fantasy prologue. Belonging is not a substitute for teaching.

HYPOTHESIS: Story worlds help belonging when they make early adversity **non-identity-diagnostic** (“misses teach the map”) *and* refuse stereotype-cueing copy — not when they lecture “you belong here” in lore text.

C. Whose history? Plural roots are historical **FACT**; classroom transfer is **HYPOTHESIS**

FACT (historiography): Joseph (1991/2011, *The Crest of the Peacock*) documents non-European mathematical lineages (Egypt, Mesopotamia, India, China, Islamic world, and more) against Eurocentric “Greeks-only greatness” narratives.

FACT (program): D’Ambrosio’s ethnomathematics programme (e.g., 1985–1990 line; ethno + mathema + tics) treats mathematical knowing as culturally situated practice, not a single civilizational monopoly.

FACT (critical pedagogy tradition): Gutstein (2006, *Reading and Writing the World with Mathematics*) — math can be used to read/write sociopolitical conditions; equity is not only “access to algebra” but agency. **Caveat for MindCraft:** Full Freirean classroom projects are not the product’s unit of delivery; borrow the *honesty about power*, not a forced activism skin on ACT prep items.

HYPOTHESIS (from XXVIII HIST-EQ / HIST-4): Plural invention stories that encode structure can lift belonging without crushing retention — **if** they are not seductive trivia.

D. Seductive details and “diversity costume” share a kill path

FACT: Harp & Mayer (1998, *Journal of Educational Psychology*, doi:10.1037/0022-0663.90.3.414) — interesting but irrelevant details can damage learning via cognitive interest mechanisms.

FACT (meta-analytic): Sundararajan & Adesope (2020, *Educational Psychology Review*, doi:10.1007/s10648-020-09522-4) — including seductive details tends to hinder learning; moderators exist (format, pacing, subject), but the mean direction is caution, not license.

Kill criterion (FOUNDER BELIEF → testable): If a cultural reference, character trait, or “heritage beat” can be removed without changing the mathematical mental model **and** without changing who can take an integral role, it is decoration. Decoration that only signal brand virtue are equity theater **and** CLT risk.

Survive criterion: Cultural/historical content that (a) encodes the invariant, (b) expands access/role/expression for under-affirmed students, and (c) does not cue stereotypes.

XXXVI.4 Failure modes specific to MindCraft story worlds

Failure mode	What it looks like in product	Likely harm	Epistemic label
Euro-default adventure	Castles, Greco-Roman “genius,” white default heroes	Quiet alienation; “not for people like me”	HYPOTHESIS (strong prior from identity lit)
Token palette swap	Same plot; recolored sprites	Cynical read; no identity mechanism	HYPOTHESIS

Heritage trivia dump	Long museum lore before first attempt	Seductive load; delayed micro-win for anxious Maya	FACT risk (CLT / seductive details)
Stereotype-friendly tropes	“Street-smart” vs “book-smart”; gendered magic roles	Cue threat / essentialism	HYPOTHESIS (aligned with threat lit)
Forced activism costume	Every item a justice pamphlet	Parent trust fracture; math time stolen; off-mission	SPECULATION commercially; Gutstein ≠ ACT session
Single-world monoculture	One lore for all students forever	No cultural fit; no choice autonomy (SDT)	HYPOTHESIS

FOUNDER BELIEF (constrained): Plural *structural* histories + multiple entry aesthetics beat one mega-franchise lore — but only with HIST/EQ experiments and content-ops cost discipline.

XXXVI.5 Audit rubric (ship this as content ops, not a vibes review)

Before a story module graduates from draft:

1. **Structure test:** Name the mathematical invariant the narrative makes memorable. If blank → rewrite or cut.
 2. **Role test:** In the first 60 seconds, does the student *do* math as an agent, or watch lore? Threatened learners need Safety → micro-attempt before long meaning (XXVIII.7).
 3. **Belonging-copy test:** Ban ability-diagnostic group framing (“girls struggle with...”, “unlike other kids...”). Prefer adversity-as-common-and-transient framing (Walton & Cohen mechanism — adapted carefully; do not fake an RCT in UI).
 4. **History plurality test:** If an invention beat exists, is at least one non-Euro lineage available in the corpus for that concept family over time (not every single card)? Token one-off “Al-Khwarizmi cameo” without structure = fail.
 5. **Seductive-details test:** Remove the cultural beat; if learning items and role affordances unchanged, cut or move post-success.
 6. **Stereotype scan:** Two reviewers (ideally including someone outside the default writer demographic) flag tropes; unresolved flags block ship.
 7. **Claim ladder (Part XXXIV):** Marketing may say “we design stories so more students can take an integral role in the math” (L0/L1 product intent). It may **not** say “our lore closes achievement gaps” without L3 evidence.
-

XXXVI.6 Commercial implications (copy, positioning, roadmap)

Marketing language that survives

- **Allowed:** “Math isn’t a private club with one history. Stories should make the *idea* human — and give every learner a real role in solving it.”
- **Allowed:** “We refuse decoration that wastes working memory. Narrative earns its place when it encodes the math or restores a sense that struggle isn’t proof you don’t belong.”

- **Banned:** “Culturally magical world closes the gap.”
- **Banned:** “Diverse characters = equitable outcomes.”
- **Banned:** Absolute tutoring-is-free / 2-sigma / streaks-as-identity (standing bans).

Competitive stance

Khan’s library and Duo’s streaks are largely **culture-thin** at the session surface (Part XXXV). MindCraft *chooses* narrative risk. That is only a wedge if the equity audit is visible in product craft — otherwise competitors can dismiss story as gimmick **and** critics can dismiss it as exclusionary. Honest position: “We accept narrative risk; we instrument belonging + learning; we kill modules that fail the rubric.”

Growth / parent trust

Parents in threatened segments buy **dignity + results**. Equity theater without `retry_120s` / transfer reads as politics instead of help. Lead with FEI evidence; treat plural history as accuracy and belonging hygiene, not a culture-war banner unless the parent segment asks for it.

Product bets

1. **Content ops checklist** above as merge gate for story modules.
2. **At least two aesthetic/cultural entry frames** for core onboarding arcs (same math, different wrap) — autonomy without fragmenting ontology.
3. **HIST-EQ / EQ-1...3** before scaling “world lore” as brand centerpiece.
4. Tutor witness remains the primary belonging technology for high-threat students (XXIV); story is amplifier, not substitute.

XXXVI.7 Experiments spawned

ID	Question	Design	Primary	Kill condition
EQ-1	Euro-default story vs plural structure-encoding story (same items)	A/B or crossover	belonging/threat items + <code>retry_120s</code> + <code>retention@7d</code>	Belonging flat; or learning drop > pre-reg threshold
EQ-2	Token palette-swap vs role/access redesign (same art budget)	3-arm	integral-role self-report + attempt latency	Palette $\geq$ role redesign on belonging
EQ-3	Pre-attempt heritage lore vs post-success meaning beat	Timing	time-to-first-attempt; transfer	Pre-lore delays attempt with no belonging lift

HIST-EQ	(from XXVIII) Plural vs single-tradition invention	A/B	belonging + learning	Belonging up, transfer down badly
EQ-QUAL	10 Maya interviews: “who is this world for?”	Protocol B + equity probes	theme codes	Students read world as “not for me” at >threshold

Pre-reg discipline (XXXIV): EQ-* identify story→belonging/persistence. They do **not** identify story→GPA gap closure. Do not sell L3 from L1.

XXXVI.8 Confidence table

Claim	Label	Confidence
Math learning is entangled with race/culture/identity	FACT (research consensus in equity math ed)	High
Access / role / expression predict engagement better than posters	FACT (Nasir & Hand mechanisms) + HYPOTHESIS for digital worlds	Medium–High
Default Euro adventure can alienate	HYPOTHESIS	Medium
Token diversity without role redesign fails	HYPOTHESIS	Medium–High
Seductive cultural trivia can hurt learning	FACT (seductive-details tradition + meta-analysis direction)	Medium–High
Plural invention stories lift belonging <i>and</i> learning	HYPOTHESIS (HIST-EQ unrun)	Low–Medium
Story can replace tutor witnessing for threat	SPECULATION / likely false	Low (against)
Equity marketing without FEI metrics is brand risk	FOUNDER BELIEF	Medium–High

XXXVI.9 What this chapter kills

1. **Kill:** “Inclusive skins” as the equity strategy.
2. **Kill:** Shipping long pre-attempt lore as belonging intervention.
3. **Kill:** Marketing claims that story worlds close racial/gender achievement gaps without identified designs.

4. **Wound:** Single-franchise world as forever default — survivable only as temporary scaffold with EQ-1/2 on the roadmap.

5. **Survive (constrained):** Story-first WORLD_VISION — **if** narrative encodes structure, expands integral roles, passes stereotype scan, and stays on the L0–L2 claim ladder until experiments graduate it.

Doctrine until data: Plural, structure-encoding, role-rich worlds under audit > monoculture lore > token palette swaps > heritage trivia dumps.

Part XXXVII — Expectancy-Value (Eccles): Utility, Cost, Attainment → Course Choice

Chapter status: Living evidence brief — Researcher tick 2026-07-29

Primary question: Why do students *choose* (or abandon) math pathways — and which MindCraft surfaces move expectancy vs value vs cost?

Owners: Learning Science · Product · Growth / Positioning · Red Team

Commercial job: Replace vague “make math meaningful” copy with a **defensible E×V map**: which product moves which lever, what parents hear, what we refuse to claim.

XXXVII.1 Why this chapter exists

Parts XXV (self-efficacy / narrative identity) and XXIV (anxiety / belonging) explain *can I* and *am I safe*. They do not fully explain *will I keep choosing math* when Instagram, sports, and sleep compete.

FACT (program lineage): Eccles-Parsons et al. (1983) built an achievement-choice model to explain gendered STEM participation: performance and choice are most proximally driven by **expectancies for success** and **subjective task values** (STV), themselves shaped by socializers, culture, and prior experience (*Expectancies, values, and academic behaviors*, in Spence, *Achievement and Achievement Motives*).

FACT (rename / situating): Eccles & Wigfield (2020, *Contemporary Educational Psychology*, doi:10.1016/j.cedpsych.2020.101859) rename the framework **Situated Expectancy-Value Theory (SEVT)** — expectancies and values are *situation-bound hierarchies*, not a single global “motivation score.”

FOUNDER BELIEF under audit: MindCraft’s identity thesis only converts to commercial retention / course-persistence if product reduces **cost** and raises **attainment + utility** while expectancy (FEI evidence) rises — not if we only make prettier explanations.

Claim we refuse as doctrine: “Story worlds automatically raise utility value.” That is **SPECULATION**. Utility is a *student-authored* link to goals (Hulleman tradition), not a brand aesthetic.

XXXVII.2 The model in product language (not poster language)

Construct	Plain definition	MindCraft surface (hypothesis)	Failure if ignored
Expectancy	“Can I succeed <i>on the next relevant tasks</i> ?”	Gap-scan → level gate; soft-wrong → retry; tutor witness; FEI micro-wins	High value + low expectancy → avoidance / crutch (Bastani Base risk, Part XXXIII)
Intrinsic value	Enjoyment / interest in the activity itself	Story-module engagement; challenge-skill (XXXI)	Fun without transfer = Duo-streak trap (XXXV)
Attainment value	Importance to <i>who I am</i> / identity	Identity re-storying (XXV); integral role in worlds (XXXVI)	Scoreboard status ≠ “math person”
Utility value	Usefulness for current/future goals	ACT roadmap honesty; parent-facing “why this concept”; student-written relevance prompts	Pure test-prep cynicism; or fake career lore
Cost	Effort, opportunity, emotional/psychological price	Anxiety load (XXIV); session length; shame of failure; time vs other activities	High cost kills choice even when expectancy/value look “ok”

FACT (classic split): Meece, Wigfield & Eccles (1990, *Journal of Educational Psychology*, doi:10.1037/0022-0663.82.1.60) — in young adolescents (N=250), **performance expectancies** predicted subsequent math grades; **value perceptions** predicted **course enrollment intentions**; math anxiety was predicted by ability/expectancy/value perceptions but did not independently drive grades/enrollment in their SEM.

Product translation: Instrument **both** outcome classes. `retry_120s` / `transfer` ≈ expectancy path. `Intent-to-continue` / `challenge_accept` / “take next ACT practice block” ≈ value path. Do not collapse both into a streak KPI (standing ban).

XXXVII.3 Subjective task value — four components, four commercial mistakes

A. Intrinsic value ≠ engagement theater

FACT: Intrinsic value = enjoyment of the task (Eccles-Parsons et al., 1983; Wigfield & Eccles, 2000, *CEP*, doi:10.1006/ceps.1999.1015).

HYPOTHESIS: Story can raise intrinsic value *if* it encodes structure and preserves integral roles (XXVIII / XXXVI). Seductive details that raise “interest” but crush learning (Harp & Mayer; Sundararajan & Adesope — carry) are **fake intrinsic** for MindCraft’s purpose.

Kill: Marketing that equates “delightful world” with durable math choice without FEI + value instruments.

B. Attainment value is the identity wedge

FACT (theory refinement): Attainment value moved from “importance of doing well” toward linkage with personal/social identities (Eccles, 2005/2009 line; summarized in Eccles & Wigfield, 2024, *Educational Psychology Review*, doi:10.1007/s10648-024-09888-9).

FOUNDER BELIEF: “I am becoming someone who can think mathematically” is attainment value. Streaks are not.

Commercial implication: Parent and teen copy should sound like *identity competence*, not *gamified attendance*. Allowed: “Evidence you can trust yourself on the next hard problem.” Banned: “Don’t break your streak — that means you’re a math person.”

C. Utility value — the only mass-market EVT intervention with strong field RCTs

FACT (RCT): Hulleman & Harackiewicz (2009, *Science*, doi:10.1126/science.1177067) — ninth-grade science students randomly assigned to write how material connects to their lives (or a known other) vs summary control; **relevance writing raised interest and grades especially for students with low success expectations.**

FACT (lab + classroom): Hulleman, Godes, Hendricks & Harackiewicz (2010, *JEP*, doi:10.1037/a0019506) — utility writing raised perceived utility and interest, again strongest for lower performers/expectancies; utility perceptions mediated interest effects.

HYPOTHESIS for MindCraft: A 2–4 minute **student-authored** utility prompt *after* a micro-win (not before first attempt under high threat) beats brand-authored “math is useful for STEM careers” copy — especially for Maya with low expectancy.

Red Team wound: Effects are domain/context sensitive; science classroom ≠ ACT app session; dosage and timing matter (SEVT situating). Do not paste “Hulleman proved story worlds work.”

Kill: Shipping utility as *staff-written* career propaganda dumped pre-attempt (CLT + seductive + autonomy threat).

D. Cost — historically neglected, commercially decisive

FACT: Cost was theorized early (effort, opportunity, psychological cost of failure) but empirically under-measured for decades (Wigfield & Cambria reviews; Eccles & Wigfield, 2020 SEVT paper stresses cost elaboration).

FACT (measurement): Flake, Barron, Hulleman, McCoach & Welsh (2015, *CEP*, doi:10.1016/j.cedpsych.2015.03.002) — validated multi-dimensional cost (task effort, outside effort, loss of valued alternatives, emotional cost); cost is separable from expectancy/value and relates to student outcomes.

HYPOTHESIS: For anxious Maya, **emotional cost** and **effort cost** dominate early sessions; for busy juniors, **opportunity cost** (vs sleep/jobs/sports) dominates weekly return. Product that only raises “fun” while leaving shame-of-wrongness intact fails SEVT choice.

Product translation: Soft-wrong / hide-correctness diagnostic (C4) / tutor tone (Part XXX) are **cost interventions**, not just pedagogy niceties. Session design that demands 40 minutes of lore before first attempt raises opportunity + effort cost without expectancy gain (tie to EQ-3).

XXXVII.4 Hierarchies and situations (why one “motivation score” is a product bug)

FACT: SEVT emphasizes *hierarchies* of expectancies and values across competing activities and situations (Eccles & Wigfield, 2020; Wigfield & Eccles developmental reviews). Choosing math means math beats alternatives *this week*, not that motivation is globally high.

HYPOTHESIS: Duo wins short-horizon intrinsic/utility-for-streak; Khan wins low-cost access library; ChatGPT Base raises short-term utility (“get the answer”) while crushing expectancy-for-unaided performance (Part XXXIII/XXXV). MindCraft’s wedge is **attainment + expectancy evidence with managed cost** — not out-streaking Duo or out-libraring Khan.

SPECULATION (commercial): Parents purchase when they believe we move the *child’s hierarchy* toward math without catastrophic emotional cost — “they’ll open it again without a fight” is a cost/expectancy story sold as calm, not as dopamine.

XXXVII.5 Socializers — parents and tutors as EVT machinery

FACT (model structure): Socializers’ beliefs and behaviors shape children’s interpretations, affective memories, and thus expectancies/values (Eccles-Parsons et al., 1983; SEVT socialization arm).

Carry: Part XXXII (parent anxiety transmission) — parent math anxiety can raise the child’s **emotional cost** and lower expectancy via homework help patterns (Maloney/Beilock line).

HYPOTHESIS: Parent dashboard copy that celebrates *process evidence* (retry, transfer) raises parent-transmitted attainment/utility messages; copy that only shows streak/percentile can raise performance pressure → cost.

Product implication: Parent WTP experiments (Part XXVII queue) should manipulate EVT framing: “identity/competence evidence” vs “minutes practiced” vs “AI explanations unlimited.” Pre-register which frame shifts purchase intent without raising reported homework conflict.

XXXVII.6 Mapping MindCraft features → EVT levers (audit table)

Feature / bet	Primary lever	Secondary	Claim ladder max without new data
Gap-scan + level gating	Expectancy calibration	Cost ↓ (fewer impossible first items)	L1 product intent
Soft-wrong → retry_120s	Expectancy evidence	Emotional cost ↓	L1–L2 with FEI instrumentation
Story worlds (structure-encoding)	Intrinsic + attainment	Utility if student links	L0–L1; EQ/HIST required
Student utility writing (post-win)	Utility	Intrinsic (situational interest)	L1; EVT-1 experiment

ACT exam-mode roadmap	Utility (instrumental)	Expectancy if honest sequencing	L1 honesty; not “score guarantee”
Tutor human sessions	Expectancy + cost	Attainment (witnessed identity)	L1; do not claim 2-sigma
Solver without guards	Short-term utility (answer)	Expectancy ↓ (crutch)	Kill as hero feature
Streaks / leaderboards	Pseudo-attainment / extrinsic	Cost ↑ under threat	Standing ban as North Star

XXXVII.7 Commercial implications (copy, positioning, growth)

Marketing language that survives

- **Allowed:** “Expectancy without value is empty drills. Value without expectancy is inspirational avoidance. We design for both — and we treat the emotional price of being wrong as a first-class design constraint.”
- **Allowed:** “Utility is something the student *writes*, not a slogan we paste.”
- **Allowed (parent):** “Fewer fights at the kitchen table” = cost reduction claim — keep qualitative until measured.
- **Banned:** “Our world makes math useful” (untested utility).
- **Banned:** “Fun = they’ll choose Precalculus.”
- **Banned:** Absolute tutoring-is-free / 2-sigma / streaks-as-identity.

Positioning vs competitors

Competitor	EVT profile (audit, not insult)	MindCraft contrast
Khan	Low cost access; mixed expectancy; thin attainment	Session FEI + identity evidence
Duo	Intrinsic/habit; cost managed via brevity; attainment via streak fiction	Transfer + real expectancy
Brilliant	Intrinsic puzzle; utility unclear for ACT Maya	Exam-honest utility + anxiety cost
ChatGPT Base	Max short-term utility; expectancy poison	Guarded Solver + unaided transfer

Growth / North Star hygiene

Optimize a **pair**: expectancy evidence (retry_120s, transfer_pass) **and** a value/choice proxy (challenge_accept, weekly voluntary return, intent-to-take-next-block). If value proxy rises while transfer falls, you are Duo-pilling the product — kill the feature.

XXXVII.8 Experiments spawned

ID	Question	Design	Primary	Kill condition
EVT-1	Post-win student utility writing vs staff utility blurb vs none	3-arm A/B after micro-win	situational interest + 7d return + transfer@7d	Utility arms lose transfer or raise time-to-next-attempt badly
EVT-2	ACT-utility framing vs identity-attainment framing in session CTA	2-arm	challenge_accept ; parent/teen preference	Attainment wins preference but tanks attempt rate in anxious segment (segmented kill)
EVT-3	Cost questionnaire (Flake-style short form) pre/post first week	Observational → then intervene on top cost	which cost dimension predicts churn	If emotional cost predicts churn but product only ships “more fun,” roadmap invalid
EVT-4	Parent copy: competence-evidence vs minutes vs AI-unlimited	WTP / conflict items	purchase intent; reported homework conflict	Competence frame raises conflict without WTP lift
EVT-QUAL	10 Maya interviews: “why would you take another math course?”	Appendix B + EVT probes	coded expectancy /value/cost	Students cannot articulate any non-grade reason → attainment/utility product gap

Pre-reg (XXXIV): EVT-* identify levers on **in-app choice and short-horizon learning**. They do **not** identify long-term course enrollment or career STEM entry without follow-up designs. Do not sell “we increase AP Calculus enrollment” from EVT-1.

XXXVII.9 Confidence table

Claim	Label	Confidence
Choice $\approx$ f(expectancy, STV) in Eccles tradition	FACT (large corpus)	High
Expectancy→performance, value→enrollment intentions (adolescent math)	FACT (Meece et al., 1990 SEM)	Medium–High (sample/era limits)
STV has intrinsic / attainment / utility / cost structure	FACT (theory + measurement work)	High

Utility-value writing can lift interest/grades for low-expectancy students	FACT (Hulleman & Harackiewicz, 2009)	Medium–High (context limits)
Cost is multi-dimensional and outcome-relevant	FACT (Flake et al., 2015)	Medium–High
Story worlds raise utility by themselves	SPECULATION	Low
MindCraft should instrument E and V separately	FOUNDER BELIEF / HYPOTHESIS	Medium–High
Soft-wrong + tutor tone primarily cut emotional cost	HYPOTHESIS	Medium
Streaks raise attainment value	SPECULATION / likely false for identity	Low (against)

XXXVII.10 What this chapter kills

1. **Kill:** Single “motivation” KPI as if E and V were the same dial.
2. **Kill:** Staff-authored “math is useful” dumps as a substitute for student-authored utility.
3. **Kill:** Treating story delight as proof of course-choice impact.
4. **Wound:** Pure ACT-utility positioning — survives for juniors near test date; wounds long-horizon identity brand if it becomes the only message.
5. **Survive (constrained):** Identity-transformation thesis — **if** product explicitly raises expectancy *and* attainment/utility *and* designs against cost, with EVT-1...4 on the roadmap.

Doctrine until data: Design and market the **E×V×Cost triad**. Expectancy evidence without value is a drill app. Value talk without expectancy is a poster. Ignoring cost is how anxious students churn while NPS looks fine among the already-confident.

Part XXXVIII — Goal Orientation: Mastery vs Performance Framing

Chapter status: Living evidence brief — Researcher tick 2026-07-29

Primary question: When MindCraft frames goals as *get better at this* vs *look smart / beat others / don’t look dumb*, which motivational pattern do we induce — and what must product + copy refuse?

Owners: Learning Science · Product · Growth / Positioning · Red Team

Commercial job: Replace empty “mastery learning” branding with a **goal-structure audit**: which surfaces cue mastery-approach, which cue appearance/performance, which cue avoidance — and which we kill even if they juice short-term engagement.

XXXVIII.1 Why this chapter exists

Parts XXV (self-efficacy), XXX (attribution / helplessness), XXXI (challenge–skill), and XXXVII (expectancy-value) answer *can I, why did I fail, is difficulty calibrated, and will I choose math*. They do not fully answer: **what success definition is the product teaching?**

A product that says “identity transformation” while UI, CTAs, and parent dashboards scream *percentile / streak / don’t mess up* is running a **performance-avoidance climate** with mastery marketing on top.

FOUNDER BELIEF under audit: MindCraft’s North Star metrics (retry_120s, challenge_accept, transfer_pass) only stay coherent if the **goal structure** of sessions is mastery-approach dominant. If recognition and evaluation structures are performance/appearance, those metrics will fight the UI.

Claim we refuse as doctrine: “Any competitive element raises motivation.” That collapses normative, appearance, and avoidance goals — meta-analyses show they are not the same animal.

XXXVIII.2 The constructs (product language, not poster language)

Goal	Plain definition	Typical student cue	MindCraft failure mode if dominant
Mastery-approach	Improve competence vs self/task standard	“Understand this; try a harder variant”	None if paired with transfer; risk = slow coverage if time unmanaged
Mastery-avoidance	Avoid failing to learn / avoid imperfect understanding	“Don’t leave gaps; don’t misunderstand”	Perfectionism, over-checking, refusal to attempt under uncertainty
Performance-approach (normative)	Outperform peers	Rank, leaderboard, “top of class”	Can raise grades for some; often raises threat for Maya; identity ≠ rank
Performance-approach (appearance)	Demonstrate talent / look smart	Public correctness, “prove you’re smart”	Meta-analytically worse than normative; shame + shallow strategy
Performance-avoidance	Avoid looking incompetent	Hide wrongness; skip hard items; Solver crutch	Helpless pattern; FEI collapse; Part XXX / XXXIII risk

FACT (classic dichotomy + classroom structures): Ames (1992, *Journal of Educational Psychology*, doi:10.1037/0022-0663.84.3.261) — classroom **task, evaluation/recognition, and authority** structures make mastery vs performance goals salient and elicit qualitatively different motivational patterns; interventions must treat structures as interdependent, not one “growth mindset poster.”

FACT (student perceptions → strategies): Ames & Archer (1988, *JEP*, doi:10.1037/0022-0663.80.3.260) — among 176 secondary students, perceived **mastery** classroom goals related to more effective strategies, preference for challenge, positive attitude, and effort→success beliefs; perceived **performance** salience related to ability focus, negative ability evaluation, and attributing failure to low ability.

FACT (implicit theories → goals → patterns): Dweck & Leggett (1988, *Psychological Review*, doi:10.1037/0033-295X.95.2.256) — entity vs incremental theories orient students toward performance vs learning goals, which set up helpless vs mastery-oriented response patterns after difficulty.

Product translation: Soft-wrong + retry is not only an expectancy/cost intervention (XXXVII) — it is a **goal-structure** intervention. Public scoreboards and “don’t break the streak” are evaluation/recognition structures that cue performance/appearance.

XXXVIII.3 Approach–avoidance split (why “performance goals” is too coarse)

FACT (trichotomous model): Elliot & Church (1997, *JPSP*, doi:10.1037/0022-3514.72.1.218) — mastery, performance-approach, and performance-avoidance are separable; mastery facilitated intrinsic motivation; performance-approach enhanced graded performance; **performance-avoidance harmed both** intrinsic motivation and grades (college classroom path model).

FACT (2×2 framework): Elliot & McGregor (2001, *JPSP*, doi:10.1037/0022-3514.80.3.501) — adds **mastery-avoidance**; four goals show distinct antecedents/outcomes; mastery-avoidance sits between mastery-approach (more adaptive) and performance-avoidance (more maladaptive) on several indicators.

HYPOTHESIS for MindCraft: Anxious Maya often defaults to **performance-avoidance** (“don’t look dumb”) or **mastery-avoidance** (“I must understand perfectly before I try”). Product that only bans leaderboards but still frames every miss as identity threat will not fix the pattern.

Kill: Marketing that equates “mastery mode” with “zero mistakes allowed on the path.” That is mastery-avoidance cosplay.

XXXVIII.4 Performance-approach is not one knob — measurement kills bad slogans

FACT (meta-analysis on measures): Hulleman, Schrager, Bodmann & Harackiewicz (2010, *Psychological Bulletin*, doi:10.1037/a0018947) — across 243 correlational studies ($N \approx 91k$), “performance-approach” scales that emphasize **normative outperformance** correlate *positively* with performance outcomes ($r \approx .14$), while scales emphasizing **appearance / looking talented** correlate *negatively* ($r \approx -.14$). Same label, different constructs.

FACT (extended outcomes meta): Senko & Dawson (2017, *Journal of Educational Psychology*, doi:10.1037/edu0000160) — the normative vs appearance split generalizes beyond grades to competence perceptions and self-regulation; age confounds with measure type were addressed.

FACT (theory status review): Senko, Hulleman & Harackiewicz (2011, *Educational Psychologist*, doi:10.1080/00461520.2011.538646) — achievement goal theory at a crossroads: multiple controversies (including when performance-approach helps) require precise operationalization before practice claims.

Commercial implication: If MindCraft ever ships social comparison, it must pre-register whether the cue is **normative skill evidence** (dangerous for threatened learners; still not identity) or **appearance theater** (ban). Default = **neither**. Do not cite “performance goals sometimes help grades” as license for public shame UX.

SPECULATION (positioning): “Beat your friends” copy will read as Duo/Brilliant hybrid; “prove you’re smart” reads as SAT-cram Instagram. Neither is the identity-transformation wedge.

XXXVIII.5 Classroom / product goal structures (TARGET), not vibes

FACT (schoolwide / TARGET lineage): Maehr & Midgley (1991, *Educational Psychologist*, doi:10.1080/00461520.1991.9653140) — motivation change requires altering school/classroom goal structures, not only exhorting students; dimensions map to task design, authority, recognition, grouping, evaluation, time (TARGET family; Ames, 1992).

FACT (review): Meece, Anderman & Anderman (2006, *Annual Review of Psychology*, doi:10.1146/annurev.psych.56.091103.070258) — classroom goal structures predict personal goals and achievement-related patterns; mastery structures are the design target for adaptive motivation.

FACT (longitudinal TARGET → mastery goals): Lüftenegger et al. (2014, *Educational Psychology*, doi:10.1080/01443410.2013.814189) — TARGET as a latent structure positively predicted students’ mastery goal orientations (N=1680 secondary; longitudinal).

MindCraft TARGET audit (hypothesis map):

TARGET dim	Mastery-leaning product move	Performance/appearance trap
Task	Varied challenge matched to skill (XXXI); format variety with transfer	Same-item grind that feels like “prove speed”
Authority	Student choice of next concept <i>after</i> honest diagnosis	Forced path that only optimizes for looking ahead of peers
Recognition	Process evidence: retry, transfer, bridge repair	Streak flames, public %ile, “genius” badges
Grouping	Tutor as witness; optional co-op without rank	Ranked classrooms / public wrongness
Evaluation	Soft-wrong diagnostic (C4); hide correctness when graph-corrupting	Instant public right/wrong as identity verdict
Time	Short sessions that still allow struggle → retry	Timed pressure that cues avoidance

FOUNDER BELIEF: PawHub “Practice the weak spot” is mastery-approach if copy is “repair this gap” and **not** “you’re behind everyone.” Gap language can flip to performance-avoidance under shame framing.

XXXVIII.6 Interaction with expectancy, anxiety, and AI crutches

FACT (goal pursuit × expectancy): Senko & Hulleman (2013, *JEP*, doi:10.1037/a0031136) — for mastery-approach and performance-approach, the harder a goal appears to attain, the less students pursue it (with costs to interest/achievement); mastery-approach goals were generally judged easier to attain than performance-approach; situational cues (material complexity) shift mastery-approach expectancies.

Carry: Part XXIV (anxiety) + XXXIII (AI overtrust) — under high threat, performance-avoidance + Solver-without-guards is the cheap escape from a hard-looking mastery or performance goal.

HYPOTHESIS: Level gating after gap-scan (easy→L3, kinda→L2, hard→L1) is a **goal-attainment-expectancy** design: it makes mastery-approach goals feel attainable. Removing gates to “challenge everyone” without expectancy supports will increase avoidance and Solver dependence.

Product implication: CTAs should make the next mastery goal *feel attainable* (Senko & Hulleman) without lying about transfer. Honest sequencing ≠ performance theater.

XXXVIII.7 Competitive product audit (goal climate lens)

Competitor	Dominant goal climate (audit)	What MindCraft must not copy
Duolingo	Habit + streak recognition; appearance of consistency	Streak-as-identity / streak-as-North-Star (standing ban)
Khan	Mastery language + energy points / badges mixed	“Mastery” label without transfer_pass
Brilliant	Puzzle competence; mild normative “smart people do this”	Appearance-smart branding without anxiety cost design
ChatGPT Base	Outcome utility (“get it right now”) → avoidance of struggle	Solver as hero; unaided transfer optional

HYPOTHESIS: MindCraft’s wedge is **mastery-approach structures with expectancy-managed tasks and anti-appearance evaluation** — not out-ranking Duo and not out-explaining Khan.

XXXVIII.8 Commercial implications (copy, product, growth)

Marketing language that survives

- **Allowed:** “We design for improvement goals — understand, repair, transfer — not for looking smart in public.”
- **Allowed:** “Wrong answers are evidence for the next attempt, not a verdict on who you are.” (ties XXX / soft-wrong)
- **Allowed (parent):** “We measure whether they take the next hard problem, not whether they kept a streak.”
- **Banned:** “Become a top 1% math student” as hero claim (appearance/normative mashup; undefended).

- **Banned:** “Never break your streak — that’s how mastery works.”
- **Banned:** Empty growth-mindset posters without TARGET structure changes (Dweck mis-applied as sticker).
- **Banned:** 2-sigma / tutoring-is-free absolutes.

Feature claim ladder (pre-reg hygiene, Part XXXIV)

Feature	Goal climate claim max without GO-* data
Soft-wrong → retry_120s	L1 intent: mastery-approach evaluation structure
challenge_accept CTA	L1: approach (need to verify mastery vs performance wording)
transfer_pass	L1–L2 anti-fake-mastery (competence standard ≠ appearance)
Leaderboards / public ranks	L0 default kill for Maya segment
Streak North Star	Kill (standing ban)
“Mastery graph” marketing	L1 only if transfer instrumented; else label is cosmetics

Growth / North Star hygiene

challenge_accept is **not automatically mastery-approach**. Instrument **why** they accept: repair gap (mastery) vs beat peer (normative) vs avoid looking behind (avoidance). If acceptance rises because of shame CTAs, kill the copy even if the metric looks green.

XXXVIII.9 Experiments spawned

ID	Question	Design	Primary	Kill condition
GO-1	Mastery-approach CTA (“repair this gap”) vs normative (“catch up to peers”) vs appearance (“show you’re smart”)	3-arm on Practice launch	challenge_accept ; retry_120s; 7d return; anxiety item	Normative/appearance win short accept but lose retry/transfer or raise anxiety
GO-2	Recognition: process badge (retry/transfer) vs streak badge vs none	3-arm	transfer@7d; identity item; self-reported goal	Streak badge wins engagement, loses transfer / raises avoidance

GO-3	Evaluation: soft-wrong first vs instant public correctness	2-arm on early session	next-attempt latency; avoidance skips; Solver opens	Instant correctness raises “finish” rate but kills retry_120s
GO-4	Goal-climate pulse (short PALS-like maste ry/performance classroom items adapted to app) weekly	Observational → then redesign recognition	which climate predicts churn / Solver dependence	If product self-reports “mastery brand” but climate score is performance → roadmap invalid
GO-QUAL	10 Maya interviews: “What would ‘doing well’ here mean?”	Appendix B + goal probes	coded mastery / normative / appearance / avoidance	Majority answer only “not look dumb” → evaluat ion/recognition redesign mandatory

Pre-reg (XXXIV): GO-* identify **in-app goal climate** → **short-horizon FEI behaviors**. They do **not** identify long-term GPA or STEM major from CTA wording alone. Do not sell “we create mastery-oriented students for life” from GO-1.

XXXVIII.10 Confidence table

Claim	Label	Confidence
Mastery vs performance goal structures shape strategies/attitudes	FACT (Ames & Archer, 1988; Ames, 1992)	High
Implicit theories orient goals → helpless vs mastery patterns	FACT (Dweck & Leggett, 1988)	High
Performance-avoidance is broadly maladaptive for motivation/grades	FACT (Elliot & Church, 1997 lineage)	High
2×2 adds mastery-avoidance as distinct	FACT (Elliot & McGregor, 2001)	High
“Performance-approach” effects depend on normative vs appearance ops	FACT (Hulleman et al., 2010; Senko & Dawson, 2017)	High
TARGET-like structures can raise mastery orientations	FACT (Maehr & Midgley; Lüftenegger et al., 2014)	Medium–High (school ≠ app)
Soft-wrong + process recognition = mastery evaluation structure	HYPOTHESIS	Medium–High

Any leaderboard helps MindCraft’s Maya segment	SPECULATION / likely false	Low (against)
Streaks instantiate mastery goals	SPECULATION / false for identity	Low (against)
MindCraft should ban appearance-goal UX by default	FOUNDER BELIEF / HYPOTHESIS	Medium–High

XXXVIII.11 What this chapter kills

1. **Kill:** Treating “mastery” as a brand adjective without evaluation/recognition structures that cue improvement standards.
2. **Kill:** Streak / public rank / “look smart” as North Star or hero copy (appearance + standing streak ban).
3. **Kill:** Collapsing all performance goals into “competition is motivating.”
4. **Wound:** Normative comparison as a growth lever — survives only as a pre-registered, segmented experiment (GO-1), never as default Maya UX.
5. **Survive (constrained):** Identity-transformation thesis — **if** product TARGET structures are mastery-approach dominant, avoidance/appearance cues are designed out, and GO-1...4/QUAL sit on the roadmap.

Doctrine until data: Market and build **mastery-approach goal structures** (task challenge + process recognition + non-shaming evaluation + attainable next goals). Performance-avoidance and appearance goals are product bugs dressed as engagement. Normative comparison is not the default growth loop.

Part XXXIX — Interleaving vs Blocking: When to Mix Problem Types

Chapter status: Living evidence brief — Researcher tick 2026-07-29

Primary question: When should MindCraft mix problem types (interleave) vs group same-type drills (block) — and what must product + copy refuse to claim?

Owners: Learning Science · Product (Practice sequencing) · Growth / Positioning · Red Team

Commercial job: Turn “we shuffle questions” from a cute engineering detail into a **defensible practice-design claim:** strategy *selection*, not just strategy *execution* — without selling interleaving as always-on magic or as engagement theater.

XXXIX.1 Why this chapter exists

Part XXIX (spacing / retrieval) already flags interleaving as a desirable difficulty and hypothesizes **Block** → **fade** → **interleave** → **spaced retrieval**. Part XXVI (CLT) warns that early overload kills learning. Part XXXI (challenge–skill) and XXXVIII (goal climate) warn that harder-feeling practice can cue avoidance if framed as shame.

This chapter densifies the **mix vs block** decision for Practice sessions, ingredient cards, and ACT-style mixed review — because textbooks and most edtech default to **blocked practice** (a dozen slope problems after the slope lesson), and students experience that fluency as “I get it” while exams demand **choose the strategy from the problem itself**.

FOUNDER BELIEF under audit: Identity transformation requires competence evidence that survives *transfer*. Blocked fluency is often **illusory competence** — high session accuracy, weak delayed discrimination.

Claim we refuse as doctrine: “Always interleave everything from the first item.” That ignores schema formation, CLT, and similarity moderators — and can spike threat for anxious Maya before `retry_120s` is healthy (XXIX.4).

XXXIX.2 Constructs (product language)

Schedule	Plain definition	What the student practices	Failure mode if misused
Blocked	Consecutive items share the same strategy / concept	<i>Execute</i> a known procedure	Fluency illusion; weak strategy choice on exams
Interleaved	Consecutive items require different strategies (no safe “same as last”)	<i>Select</i> then execute	Early overload; frustration; avoidance if unscaffolded
Hybrid / faded	Short block for schema → rising mix	Execute → select	Best candidate for MindCraft default

Operational definition for MindCraft (HYPOTHESIS): Interleaving = consecutive Practice items do **not** share the same primary strategy class (concept + dominant ingredient / archetype family). Mere format shuffle (word → diagram) with the same strategy is **not** full interleaving — it is representation variety (useful; different claim).

XXXIX.3 Mechanism: why mixture changes the pedagogical demand

FACT (math laboratory / classroom line): Rohrer & Taylor (2007, *Instructional Science*, doi:10.1007/s11251-007-9015-8) — Experiment 2: after learning multiple problem types, **mixed** vs **blocked** practice; one-week delayed test vastly favored mixed practice. Experiment 1 separately showed spacing > massing for a single type. Shuffled textbooks alter **both** spacing and mixing; they are related but not identical levers.

FACT (strategy choice vs execution): Rohrer, Dedrick & Stershic (2015, *Journal of Educational Psychology*, doi:10.1037/edu0000001) — 126 seventh-graders, same problems over ~3 months, interleaved vs blocked arrangement; after review, unannounced tests favored interleaving at 1 day ($d = 0.42$) and 30 days ($d = 0.79$). Core mechanism claim: blocked practice lets students know the strategy *before reading the problem*; interleaved practice forces strategy choice from problem features — the skill exams actually require.

FACT (not only look-alike problems): Rohrer, Dedrick & Burgess (2014, *Psychonomic Bulletin & Review*, doi:10.3758/s13423-014-0588-3) — interleaved benefit is **not** limited to superficially similar problem kinds;

reordering changes pedagogical demand (choose vs execute) even when types do not look alike.

FACT (classroom-scale RCT): Rohrer, Dedrick, Hartwig & Cheung (2020, *JEP*, doi:10.1037/edu0000367; Online First 2019) — preregistered cluster RCT, 54 seventh-grade classes, mostly interleaved vs mostly blocked assignments over months; after interleaved review for both, one-month unannounced test: 61% vs 38%, $d = 0.83$. Teachers implemented without training; anonymous pre-results survey supported feasibility. Caveats remain (materials, population), but this is the strongest classroom efficacy signal in the math interleaving line.

Product translation: MindCraft’s transfer_pass and mixed ACT practice are **closer to the interleaved demand** than a single-concept blocked grind. A session that is 12 items of linear_equations after “Learn linear equations” is textbook-default blocking — fine as *early schema practice*, dangerous as the *only* schedule if we market exam readiness.

XXXIX.4 Meta-analytic moderators (do not overclaim “interleaving always wins”)

FACT (inductive learning meta-analysis): Brunmair & Richter (2019, *Psychological Bulletin*, doi:10.1037/bul0000209) — 59 studies / 238 effect sizes / 158 samples; overall interleaving vs blocking Hedges’ $g = 0.42$. **Material matters:** paintings/visual high ($g \approx 0.67$); **mathematical tasks small but positive ($g \approx 0.34$)**; words favor **blocking** ($g \approx -0.39$); expository texts ambiguous. Moderators: stronger interleaving when **between-category similarity is high, within-category similarity is low**, and material is more complex — consistent with attentional bias / sequential attention accounts (Carvalho & Goldstone lineage).

FACT (attentional bias framework): Carvalho & Goldstone (2015, *Frontiers in Psychology*, doi:10.3389/fpsyg.2015.00505) — study sequence biases what is encoded: interleaving emphasizes **between-category differences**; blocking emphasizes **within-category similarities**. Prediction: when categories are hard to tell apart, interleave; when the job is to notice what members of one category share, block can help.

HYPOTHESIS for MindCraft: Interleave **confusable** strategy families (e.g., independent vs dependent probability; similar ACT archetypes that share surface cues but need different ingredients). Block when introducing a *new* ingredient representation until the student can execute with worked-example fade (XXVI) — then mix against near-miss foils.

Kill: Marketing that says “our AI interleaves everything like elite tutors” without specifying **confusable sets** and **when blocking is intentional**.

XXXIX.5 Fluency illusion, affect, and why students prefer blocking

FACT / TRADITION (desirable difficulties): Conditions that slow acquisition (spacing, interleaving, generation) often improve later retention/transfer (Bjork lineage; see XXIX.4). Interleaving feels harder *during* practice; blocked practice feels fluent.

HYPOTHESIS (Koriat/Bjork-style monitoring, applied): Students and parents misread blocked fluency as mastery. Parent dashboards that celebrate “12/12 on slope today” without a mixed delayed check sell **performance theater** (XXXVIII) dressed as learning.

Carry from XXIV / XXXVII / XXXVIII: Raising interleaving dose raises short-term error rate → for anxious Maya, that can cue performance-avoidance unless (a) soft-wrong / mastery evaluation, (b) expectancy-managed difficulty, (c) explicit frame: “this mix trains choosing, not proving you’re dumb.”

Product implication: Copy for interleaved sessions: “**Practice choosing the right move**” — not “hard mode for elites” and not streak pressure.

XXXIX.6 Sequencing doctrine (ties XXVI + XXIX)

FOUNDER BELIEF / HYPOTHESIS (sequenced schedule):

1. **Acquire (block-lite):** short same-strategy set + worked examples / cards until execution is stable enough that errors are mostly slips, not total confusion.
2. **Discriminate (interleave):** mix with **near-miss** types that share surface features (bridge gaps, format gaps, ACT archetype neighbors).
3. **Retain (space):** redistribute the same mix across days (XXIX) — interleaving *guarantees* some spacing of each type, but lag still needs explicit scheduling.
4. **Transfer check:** unannounced mixed set → transfer_pass; failure routes to diagnosis (forgot vs never-had; wrong strategy class vs execution bug).

Contradiction to hold: Some math/category studies report null or blocking advantages under specific similarity structures (noted in Brunmair & Richter’s discussion of mixed math findings; e.g., certain representation-interleaving designs). MindCraft must treat **interleaving axis** (strategy class) as primary for exam readiness, and treat **representation mix** as a separate experiment track (format axis / C2–C3).

XXXIX.7 Competitive / product audit (schedule lens)

Surface	Typical schedule	What MindCraft should do
Textbook / workbook	Heavy blocking after lesson	Do not clone as “mission complete”
Khan-like mastery	Often skill-blocked until “mastered”	Mastery label without mixed delayed check = fluency risk
Duo-like	Habit + short items; mix varies by skill tree	Habit ≠ discrimination training
Brilliant	Puzzle variety; not always systematic near-miss mix	Variety ≠ controlled interleaving
ChatGPT Solver	Single-problem depth	Zero schedule pedagogy; cannot substitute mix training
MindCraft Practice today	Often concept-scoped sessions + bank shuffle	Shuffle within one concept ≠ full strategy interleaving; need cross-type mix modes

HYPOTHESIS: PawHub “Practice the weak spot” correctly targets a gap, but a pure blocked weak-spot grind can raise local accuracy while leaving **strategy selection** untrained. Pair weak-spot blocks with a **discrimination coda**: 3–5 interleaved near-misses including the weak type.

XXXIX.8 Commercial implications (copy, product, growth)

Marketing language that survives

- **Allowed:** “Exams don’t announce the chapter. We train choosing the strategy — not only repeating the last one.”
- **Allowed:** “Blocked practice feels easy; mixed practice builds the skill tests actually measure.”
- **Allowed (parent):** “We check whether they can pick the right approach on a mixed set days later — not just a perfect streak on one topic.”
- **Banned:** “Interleaving always doubles scores” / absolute effect-size marketing (Rohrer classroom $d \approx 0.83$ is **one** RCT context; Brunmair math $g \approx 0.34$ is the broader inductive-math signal — do not collapse).
- **Banned:** “Always mix from day one.”
- **Banned:** Selling shuffle/random as science without near-miss design.
- **Banned:** 2-sigma / tutoring-is-free / streak-as-mastery.

Feature claim ladder (Part XXXIV hygiene)

Feature	Claim max without IL-* data
Within-concept item shuffle	L0–L1 “variety”; not full interleaving claim
Cross-strategy interleaved coda after weak-spot block	L1 intent: discrimination training
Delayed mixed transfer_pass	L1–L2 anti-fluency-illusion
“Interleaving improves ACT” hero claim	L0 until IL-1/IL-2 pre-registered

Growth / North Star hygiene

Session accuracy under **blocked** schedules will look better than under interleaved ones. **Do not** optimize North Star on blocked accuracy. Prefer: mixed delayed accuracy, strategy-selection error rate (wrong procedure class), transfer_pass, and retry_120s after mix-induced misses.

XXXIX.9 Experiments spawned

ID	Question	Design	Primary	Kill condition
----	----------	--------	---------	----------------

IL-1	Blocked-only weak-spot session vs block + interleaved near-miss coda	2-arm; same total items	Delayed (48h) mixed accuracy; strategy-class error rate	Coda hurts delayed mix and raises avoidance without transfer gain
IL-2	Early interleave (from item 1) vs fade-in interleave after short block	2-arm on new concept	Acquisition errors; retry_120s; 7d mixed test	Early-all-mix wins delayed score but kills return/anxiety for Maya segment
IL-3	Near-miss interleave (confusable archetypes) vs far mix (unrelated easy types)	2-arm	Discrimination errors on near-miss test	Far mix equals near-miss on discrimination → “any shuffle” doctrine dies
IL-4	Parent/student judgment: blocked fluency vs mixed score — which do they trust?	Survey + optional choice	Preference for blocked; willingness-to-pay framing	If parents only buy blocked dashboards, copy/education mandatory before UX ships mix-heavy
IL-QUAL	10 Maya interviews: “When did practice feel like the real test?”	Appendix B	coded choose-vs-execute moments	Nobody describes strategy choice → product isn’t creating the demand state

Pre-reg (XXXIV): IL-* identify **schedule** → **discrimination** / **delayed mix**. They do **not** identify long-term STEM identity from one coda. Do not sell “interleaving creates math people.”

Supersedes / densifies: XXIX INT-1 (interleave after fade vs blocked-only) → prefer IL-1/IL-2 as the operational pair.

XXXIX.10 Confidence table

Claim	Label	Confidence
Mixed/interleaved math practice often beats blocked on delayed tests	FACT (Rohrer & Taylor, 2007; Rohrer et al., 2015, 2020)	High (math classroom/lab line)
Mechanism includes forced strategy selection (not only spacing)	FACT / strong interpretation in Rohrer line	High

Benefit can appear for non-look-alike problem kinds	FACT (Rohrer et al., 2014)	High
Overall inductive interleaving effect moderated by material & similarity	FACT (Brunmair & Richter, 2019)	High
Math subset effect smaller than visual category learning	FACT ($g \approx 0.34$ vs higher for paintings)	High
Blocking can help when within-category similarity learning is the job	FACT (attentional bias framework) / HYPOTHESIS for early schema	Medium–High
Always-interleave-from-item-1 is optimal for anxious Maya	SPECULATION / likely false	Low (against)
Within-concept shuffle $\equiv$ interleaving science claim	SPECULATION / false	Low (against)
Default MindCraft path = short block $\rightarrow$ near-miss interleave $\rightarrow$ spaced mix check	FOUNDER BELIEF / HYPOTHESIS	Medium–High

XXXIX.11 What this chapter kills

1. **Kill:** Treating blocked session accuracy as proof of exam readiness or identity change.
2. **Kill:** “We randomly shuffle, therefore we interleave” as a marketing claim.
3. **Kill:** Absolute always-interleave / always-double-scores claims.
4. **Wound:** Pure weak-spot blocked grind as the only Practice mode — survives as *acquisition*, not as the full loop.
5. **Survive (constrained):** Desirable-difficulty interleaving for **strategy discrimination** — after schema foothold, with near-miss design, affect-safe framing, and delayed mixed transfer_pass.

Doctrine until data: Market and build **choose-the-strategy training** via faded interleaving of confusable types. Blocking is a scaffold, not the product. Fluency without mixed delayed evidence is a vanity metric.

Part XL — Self-Explanation Prompts: Chi/Renkl $\rightarrow$ Coach UX

Chapter status: Living evidence brief — Researcher tick 2026-07-29

Primary question: When do prompts that force the student to *generate* an explanation improve math learning — and what must MindCraft’s coach / cards / Solver refuse to do with “explain yourself” UX?

Owners: Learning Science · Product (coach, ingredient cards, Solver) · Growth / Positioning · Red Team

Commercial job: Turn “we explain everything” and “AI tutor talks at you” into a **defensible generation claim:** students build transferable understanding by explaining *why steps work* (and why common wrong moves fail) — without selling open-ended “justify your guess” as science, and without confusing fluent AI explanation to the student

with self-explanation *by* the student.

XL.1 Why this chapter exists

Part XXVI (CLT / worked examples / fading) already puts examples and completion problems in the acquisition path. Part XXIX / XXXIX demand retrieval and strategy *selection*. Part XXXIII warns that fluent AI explanations can create overtrust. Experiment A (wizard coach) and soft-wrong already assume that *after* an error, something productive happens.

This chapter densifies the **prompt layer**: what the product asks the student to *produce* — principle for a step, why a misconception is wrong, anticipation of the next operator — vs what it merely *shows*.

FOUNDER BELIEF under audit: Identity as a mathematical thinker requires evidence that the student can *account for* a move, not only select A/B/C/D. Passive delivery without generation is Layer-1 scarcity relief, not identity transformation.

Claim we refuse as doctrine: “Any explain box after a wrong answer is self-explanation science.” Untargeted prompts, explain-your-wrong-guess-without-feedback, and AI monologues are different mechanisms — some help, some hurt, some are theater.

XL.2 Constructs (product language)

Construct	Plain definition	Product surface	Failure mode
Self-explanation	Learner generates inferences beyond the given text/example <i>for themselves</i>	Prompted coach turns; card reflection; step-principle pick	Empty “why?” with no scaffold; AI does the explaining for them
Instructional explanation	Expert/AI/tutor explains <i>to</i> the learner	Solver narration; worked example text	Can reduce germane effort if overused (expertise / passivity)
Worked-example study	Study full solution before / instead of cold solving	Ingredient cards; fade path	Without SE prompts → shallow copy (Chi “poor” pattern)
Principle / glossary citation	Name the rule that justifies a step	Aleven-style menu; ingredient id pick	Menu-click without comprehension if too shallow
Misconception contrast	Explain why a common wrong path fails	Soft-wrong + “why not B”; bank distractors	Asking only “defend your wrong answer” before feedback

Operational definition for MindCraft (HYPOTHESIS): A self-explanation prompt is load-bearing only if (1) the to-be-explained content is tagged **correct or known-incorrect**, (2) the response format scaffolds principle / operator–goal / misconception contrast (not free essay only), and (3) success is measured by **transfer / delayed mix**,

not by longer time-on-task or chat length.

XL.3 Classic effect: spontaneous quality, then elicitation

FACT (good vs poor example study): Chi, Bassok, Lewis, Reimann & Glaser (1989, *Cognitive Science*, doi:10.1207/s15516709cog1302_1) — talk-aloud while studying mechanics worked examples. “Good” students generated many explanations that refined conditions for actions and linked steps to principles; “poor” students under-explained, monitored poorly, and relied heavily on examples at solve time. Self-explanation quality predicted understanding and example-independence.

FACT (elicitation works): Chi, de Leeuw, Chiu & LaVancher (1994, *Cognitive Science*, doi:10.1016/0364-0213(94)90016-7) — eighth-graders prompted to self-explain after each line of a circulatory-system text outperformed controls who reread; high explainers built more accurate mental models. Prompts can *create* the behavior that spontaneous good students show.

FACT (quality ≠ mere time): Renkl (1997, *Cognitive Science*, doi:10.1016/S0364-0213(99)80017-2) — probability worked examples; time-on-task controlled. Learning gains still predicted by qualitative SE features: **principle-based** explanations, **operator–goal** explication, **anticipative** reasoning. Two effective profiles: anticipative reasoners and principle-based explainers.

Product translation: Coach UX should elicit *those* moves — “What rule licenses this step?” / “What is this step trying to achieve?” / “What happens next?” — not “Write anything about your feelings about math” and not “Hear the AI explain for 90 seconds.”

XL.4 Scaffolded prompts + fading (coach design spine)

FACT (fading + principle prompts): Atkinson, Renkl & Merrill (2003, *Journal of Educational Psychology*, doi:10.1037/0022-0663.95.4.774) — successively fading worked-out steps helps near transfer vs example–problem pairs, but far transfer is unreliable from fading alone. Combining fading with prompts to identify the **underlying principle** for each worked step produced medium-to-large effects on **near and far transfer without extra time on task**.

FACT (Cognitive Tutor classroom): Alevan & Koedinger (2002, *Cognitive Science*, doi:10.1016/S0364-0213(02)00061-7) — geometry Cognitive Tutor; students who explained steps by citing geometric rules (glossary) learned with greater understanding and better transfer than problem-solving-only peers. Interpretation: better-integrated declarative knowledge, less shallow procedural knowledge. Scales in classroom software — not only lab talk-alouds.

HYPOTHESIS for MindCraft: Ingredient card_templates + bridge cards are the worked-example substrate; coach prompts should (a) attach principle/ingredient labels to faded steps, (b) require a brief generation or structured pick *before* revealing the next full AI paragraph, (c) fade prompts as mastery rises (expertise reversal risk if we nag forever).

Tie to SAFE-CRAFT (XXVI): Micro-attempt / completion blank inside an example *plus* a principle prompt beats cold multi-step solve *or* passive read-only cards.

XL.5 Math meta-analysis: what we may claim (and must not)

FACT (math SE meta-analysis): Rittle-Johnson, Loehr & Durkin (2017, *ZDM*, doi:10.1007/s11858-017-0834-z) — prompted self-explanation in mathematics: small-to-moderate immediate gains — procedural knowledge $ES \approx .28$, conceptual $ES \approx .33$, procedural transfer $ES \approx .46$. Effects **stronger when high-quality explanations are scaffolded** (training or structured responses). Time-on-task control did **not** erase the benefit (moderator). **Caveats they stress:** evidence for reliable classroom (esp. primary/secondary non-ITS) effects and for **retention over delay** is much thinner; delayed procedural/conceptual effects often weak/null while delayed transfer $ES \approx .32$ in the subset that measured it; substantial heterogeneity (some null/negative studies).

Instructional recommendations from that paper (treat as design doctrine, not slogans):

1. Scaffold high-quality explanations (structure or train).
2. Design prompts so they do **not** steal attention from other critical content.
3. Prompt learners to explain **correct** information.
4. Prompt why **common misconceptions** are incorrect.

Kill: Marketing “self-explanation always raises grades” / absolute classroom-retention claims from lab immediate ES alone. Claim ladder: L1 “structured why-this-step prompts aim at transfer”; L2+ only after SE-* experiments with delayed mixed transfer_pass.

XL.6 Constraints: when prompts backfire

FACT (constraint review): Rittle-Johnson & Loehr (2017, *Psychonomic Bulletin & Review*, doi:10.3758/s13423-016-1079-5) — four constraints: (1) SE favors principle-guided domains and transfer, can hurt detail/exception learning; (2) explaining **one’s own** (often wrong) ideas can entrench them — better to explain known-correct / known-incorrect content; (3) prompts focus attention — mis-aimed prompts tax the wrong content; (4) alternatives (instructional explanation, retrieval, unfamiliar problems) can match or beat SE in some conditions.

HYPOTHESIS for soft-wrong: Order matters. After a wrong MC choice: (A) brief correct-principle or contrastive “why B fails” prompt with scaffold → help; (B) “Explain why you picked C” *before* any correctness signal → risk of justifying the misconception (constraint 2). Prefer: soft-wrong frame → show contrastive target → generate/select why the foil fails → then optional deeper AI wrap (XXXIII: calibration > fluency).

Contradiction to hold: Open generation is the romantic “thinker” move; structured glossary/ingredient picks (Aleven) often deliver the measurable classroom gain. MindCraft should not shame structured picks as “not real thinking” — they are the scalable SE scaffold.

XL.7 Competitive / product audit (explanation lens)

Surface	Typical explanation pattern	MindCraft stance
Khan / video	Instructional explanation; optional reflection	Do not clone lecture-first as identity product

Duo	Rare deep SE; habit + bits	Habit ≠ principle generation
Brilliant	Puzzle insight; light justify	Insight ≠ scaffolded SE curriculum
ChatGPT Solver	High-fluency instructional explanation	Not self-explanation; overtrust risk (XXXIII)
Cognitive Tutor lineage	Step + rule citation	Closest evidence-backed software pattern
MindCraft today	Cards + Solver + coach drafts	Need student generation gates before AI walls of text; wire misconception bank into contrastive prompts

FOUNDER BELIEF: Solver that answers *for* Maya without a generation gate competes on Layer-1 (answers get cheap) and loses the identity thesis. Coach that forces a principle pick / misconception contrast competes on Layer 2–6.

XL.8 Commercial implications (copy, product, growth)

Marketing language that survives

- **Allowed:** “We don’t just show the steps — we ask you to say *why* the step works.”
- **Allowed:** “Wrong answers become contrast: why that trap fails, not a shame spiral.”
- **Allowed (parent):** “We train explaining the rule, not only picking the letter — that’s what transfers.”
- **Banned:** “AI explains everything so tutoring is free.”
- **Banned:** “Any chat box = self-explanation science.”
- **Banned:** Absolute meta-analytic grade claims; 2-sigma; streak-as-understanding.
- **Banned:** Forcing students to justify unchecked wrong guesses as the default pedagogy.

Feature claim ladder (Part XXXIV hygiene)

Feature	Claim max without SE-* data
AI worked-example narration	L0 instructional help; not SE claim
Structured principle / ingredient pick on faded step	L1 SE scaffold (Aleven/Atkinson-aligned intent)
Contrastive “why this distractor fails” after soft-wrong	L1 misconception-SE intent
Open essay “explain your answer” every item	L0 until quality+affect validated; risk of load + avoidance
“SE raises ACT / identity” hero	L0 until SE-1/SE-2 delayed transfer

Growth / North Star hygiene

Chat turns, explanation word count, and Solver dwell time are **vanity** if uncoupled from transfer_pass and strategy-class errors. Prefer: rate of principle-consistent explanations (rubric or structured correct pick), misconception-contrast success, delayed mixed accuracy, retry_120s after contrastive miss — not tokens generated by the model.

XL.9 Experiments spawned

ID	Question	Design	Primary	Kill condition
SE-1	Principle-prompt on faded card steps vs same cards with AI narration only	2-arm; matched time budget	48h near + far / mixed transfer; shallow-copy errors	Narration-only $\geq$ prompted on transfer and lower load/avoidance
SE-2	Soft-wrong $\rightarrow$ contrastive “why foil fails” vs soft-wrong $\rightarrow$ explain-own-choice first	2-arm	Misconception recurrence; retry_120s; delayed item isomorphic	Explain-own-first wins transfer $\rightarrow$ constraint-2 doctrine wounded
SE-3	Structured ingredient/glossary pick vs open text SE vs no SE	3-arm	Transfer; time; completion rate (Maya anxiety segment)	Open text wins transfer but kills return $\rightarrow$ ship structured default
SE-4	Prompt every step vs adaptive fade of SE prompts with mastery	2-arm	Transfer; prompt skip/nag irritation	Always-on prompts equal fade on transfer but raise churn $\rightarrow$ fade mandatory
SE-QUAL	10 Maya interviews: “When did an explanation feel like <i>yours</i> vs the app talking?”	Appendix B	coded ownership / passivity moments	Nobody distinguishes $\rightarrow$ UX isn’t creating generation

Pre-reg (XXXIV): SE-* identify **prompt regime** $\rightarrow$ **transfer** / **misconception recurrence**. They do **not** identify long-term identity from one session of explain boxes. Densifies Experiment A: coach content should be SE-constrained, not generic pep talk.

Supersedes / densifies: XXVI worked-example fading notes; Experiment D (explanation timing $\times$ anxiety) $\rightarrow$ prefer SE-2/SE-3 as operational pairs; XXXIII “pedagogy wrap” $\rightarrow$ wrap = student generation + calibration, not longer AI prose.

XL.10 Confidence table

Claim	Label	Confidence
Higher-quality spontaneous SE while studying examples predicts better learning	FACT (Chi et al., 1989)	High
Prompting SE can improve understanding vs reread / no prompt	FACT (Chi et al., 1994; math meta)	High
Qualitative SE features predict gains beyond time-on-task	FACT (Renkl, 1997)	High
Fading + principle prompts aid near & far transfer without extra time	FACT (Atkinson et al., 2003)	High
Step explanation via rule citation helps Cognitive Tutor geometry transfer	FACT (Aleven & Koedinger, 2002)	High
Math prompted-SE immediate effects small–moderate; scaffold helps	FACT (Rittle-Johnson et al., 2017 ZDM)	High
Classroom + delayed retention evidence thinner than immediate lab/ITS	FACT (same meta)	High
Explaining own wrong ideas can entrench; prefer correct/incorrect targets	FACT (constraint review)	High
AI instructional explanation ≡ self-explanation benefit	SPECULATION / false	Low (against)
Default MindCraft path = faded examples + structured principle/misconception prompts before AI wrap	FOUNDER BELIEF / HYPOTHESIS	Medium–High

XL.11 What this chapter kills

1. **Kill:** Equating Solver/AI monologue with self-explanation pedagogy.
2. **Kill:** “Explain why you chose that” *before* correctness/contrast as the default after every wrong answer.
3. **Kill:** Absolute “SE always raises scores / always works in class / always sticks for months” marketing.
4. **Wound:** Open-ended explain boxes as the only SE UX — survives as optional depth, not as the scalable default.
5. **Survive (constrained):** Scaffolded prompts on **correct steps and known misconceptions**, paired with worked-example fading, measured by transfer — the coach spine for identity-relevant competence evidence.

Doctrine until data: Build and market **student-generated why** (structured first) on faded examples and contrastive wrongs. AI explains *after* or *around* that generation — never as a substitute for it. Fluency of the model is not mastery of the student.

Part XLI — Desirable Difficulties × Anxiety: When Struggle Helps vs When It Becomes Threat

Chapter status: Living evidence brief — Researcher tick 2026-07-30

Primary question: How should MindCraft reconcile Bjork-style desirable difficulties (spacing, interleaving, retrieval, generation) with Ashcraft-style math-anxiety / working-memory threat — so product + copy do not sell either “grind harder” or “never make it hard”?

Owners: Learning Science · Affect / Anxiety · Product (Practice sequencing + evaluation climate) · Growth / Positioning · Red Team

Commercial job: Resolve the I.4 bottleneck: ship faded desirable difficulties **without** converting anxious Maya’s session into shame theater — and market the distinction between *productive struggle* and *threat* in language parents and students can trust.

XLI.1 Why this chapter exists

Parts XXIX (spacing/retrieval), XXXIX (interleaving), and XL (self-explanation) all lean on **desirable difficulties**: conditions that slow short-term performance but improve durable access and transfer. Part XXIV (affect/anxiety) and Ashcraft’s line show that math anxiety **already** loads working memory like a dual task. Part XXXI (challenge–skill) and XXXVIII (goal climate) warn that harder-feeling practice can cue avoidance under performance threat.

v1.6 left an explicit bottleneck: “*Part XLI must resolve desirable difficulties × anxiety so interleaving/SE/spacing do not ship as shame.*”

FOUNDER BELIEF under audit: Identity transformation requires competence evidence that survives delay and mix — which *requires* some desirable difficulty — but for threatened learners, unsequenced difficulty is not “elite training”; it is **threat + WM collapse + exit**.

Claim we refuse as doctrine: “Struggle is always good” / “always crank desirable difficulties” *and* its twin “anxious students should only get easy fluency drills forever.” Both are commercially tempting; both fail Maya.

XLI.2 Constructs (product language)

Construct	Plain definition	Product signal	Failure mode if misused
-----------	------------------	----------------	-------------------------

Desirable difficulty (DD)	Challenge that forces encoding/retrieval processes that support later access	Mixed review, spaced retries, generation before reveal	Feels worse today → students & parents quit if misframed
Undesirable difficulty	Challenge the learner cannot productively meet (knowledge, cues, or affect load)	Early full interleave; open SE under panic; high-stakes quiz theater	Errors without learning; avoidance; “I’m not a math person”
Math anxiety load	Worry/rumination consuming WM resources needed for multi-step math	Freeze, hint binge, rapid exit after first hard miss	Interpreted as “lazy” or “needs more struggle”
Stakes / evaluation climate	Whether the hard moment is framed as judgment vs training	Soft-wrong, mastery climate, no streak-as-worth	High stakes turn retrieval into threat (Hinze & Rapp)
Equipped success	Prior knowledge + cues sufficient to <i>respond</i> to the difficulty	Block-lite → fade → mix; worked examples; structured SE	DD without equipment = cruelty dressed as science

Operational definition (HYPOTHESIS): A MindCraft difficulty is *desirable* only when (a) the student has enough schema/cues to respond with non-chance success or productive partial success, **and** (b) evaluation climate keeps the moment in mastery/training frame, **and** (c) affect load is not already saturating WM. Else it is undesirable — even if the *same* schedule would be desirable for a calm, prepared peer.

XLI.3 Bjork line: difficulty that builds durable access

FACT (origin): Bjork (1994, in Metcalfe & Shimamura, *Metacognition: Knowing about knowing*, MIT Press, pp. 185–205) — conditions that slow acquisition (spacing, variation, testing/generation as learning events, contextual interference) can improve long-term retention and transfer relative to conditions that make performance look strong during practice.

FACT (desirable ≠ any difficulty): Bjork & Bjork (2011/2014 reviews; JARMAC 2020 comment, doi:10.1016/j.jarmac.2020.09.003) — many difficulties are undesirable forever; a difficulty is desirable only when the learner can respond successfully with the encoding/retrieval processes the difficulty is meant to trigger. *If the learner is not equipped, the difficulty becomes undesirable.*

FACT (learning vs performance illusion): Instructors and students systematically misread current fluency as learning (Bjork lineage; see also XXIX / XXXIX). Blocked / massed / reread conditions feel productive; spaced / interleaved / tested conditions feel worse *and often teach more*.

Product translation: MindCraft cannot sell “we make practice feel easy” as proof of learning, nor “we make practice feel hard” as proof of rigor. The commercial claim must be **durable choosing and retrieving under exam-like demand** — measured by delayed mix / transfer_pass / solo_transfer_pass — not by session comfort or session pain.

XLI.4 Ashcraft line: anxiety as a competing load

FACT (WM disruption): Ashcraft & Kirk (2001, *JEP: General*, doi:10.1037/0096-3445.130.2.224) — high math-anxiety individuals show reduced WM spans (especially computation-based spans); dual-task math + memory load amplifies RT/errors. Anxiety acts as a **transitory disruption of working memory**, not merely as a stable “low ability” trait.

FACT (review / education bridge): Ashcraft & Krause (2007, *Psychonomic Bulletin & Review*, doi:10.3758/BF03194059) — math beyond simple retrieval depends on WM; high math anxiety functions like a resource-demanding secondary task (preoccupation with fear). Global consequence: avoidance of elective math, majors, and careers.

FACT (carry from XXIV): Association of math anxiety with performance is real but modest on average (meta $r \approx -0.17$); mechanism candidates include WM competition + avoidance. Anxiety is **one bottleneck among many**, not destiny.

Conflict statement (HYPOTHESIS): Desirable difficulties often *increase* momentary retrieval/selection demand — exactly the demand anxiety already taxes. Unbuffered DD therefore risks converting a would-be desirable difficulty into an undesirable one for HMA students: same schedule, different learner state → different pedagogy.

XLI.5 Evidence that stakes and anxiety change whether “hard practice” helps

FACT (pressure can kill retrieval benefits): Hinze & Rapp (2014, *Applied Cognitive Psychology*, doi:10.1002/acp.3032) — high-stakes vs low-stakes quizzes: quiz performance similar, but **final test better after low-stakes**; only low-stakes quizzes beat reread. Authors propose **retrieval disruption**: pressure can impair the memory-formation benefits of retrieval even when online quiz scores look fine.

FACT (classroom retrieval often lowers reported test anxiety): Agarwal, D’Antonio, Roediger, McDermott & McDaniel (2014, *JARMAC*, doi:10.1016/j.jarmac.2014.07.002) — survey of 1,408 middle/high-school students in classrooms using retrieval practice; 72% reported feeling **less nervous** for unit tests/exams when retrieval practice was used; 92% said it helped learning. Self-report, not RCT — but directionally opposite to “tests always raise anxiety.”

FACT (meta: practice tests tend to reduce test anxiety): Yang, Li, Zhao, Luo & Shanks (2023, *Educational Psychology Review*, doi:10.1007/s10648-023-09801-w) — 24 studies / 25 effects / $N \approx 3,374$; practice tests reduce test anxiety with Hedges’ $g \approx -0.52$; easy practice tests more mitigating than difficult ones. **Caveat:** mostly classroom/quiz anxiety outcomes — not a free pass to ship brutal mixed ACT sims as day-one “practice tests.”

FACT (anxiety can reshape what is learned, not only performance): Fioriti, Pizzie, Evans, Green & Lyons (2025, *JEP:LMC*, doi:10.1037/xlm0001453) — adults learning unfamiliar multiplication across days: high math anxiety associated with **impaired procedural learning** on unrepeated problems but **accelerated retrieval learning** on heavily repeated problems — consistent with effort allocation toward declarative retrieval routes and away from procedural efficiency gains.

HYPOTHESIS for MindCraft: Anxious Maya may *prefer* and even accelerate on repeated retrieval of the same items (looks like mastery, feels safer) while under-building procedural flexibility and strategy selection — exactly the exam skills interleaving/SE target. Product that only rewards repeated-item fluency can **serve anxiety’s preference** while

failing transfer.

XLI.6 Resolution doctrine (not a slogan)

FOUNDER BELIEF / HYPOTHESIS — SAFE-DD stack (ties FEI / mastery climate / IL / SE):

1. **Equip before you strain.** Short block-lite + worked-example / card fade until execution errors are mostly slips (XXVI, XXXIX). DD without schema is undesirable.
2. **Lower stakes while raising retrieval.** Soft-wrong, mastery climate, no public rank, no streak-as-worth (XXXVIII; Hinze & Rapp). Retrieval as *training*, not judgment.
3. **Dose difficulty to remaining WM.** Prefer structured SE picks over open essays under high affect (XL SE-3); prefer faded interleave over day-one full mix (IL-2); prefer spaced short sets over marathon mixed threat.
4. **Frame the feeling.** Copy: “This mix trains *choosing* — feeling slower is the workout, not proof you’re behind.” Never: “hard mode for elites” / “grind through the pain.”
5. **Measure the right win.** Session accuracy under block is not the North Star. Prefer *retry_120s*, delayed mix / *transfer_pass*, *solo_transfer_pass*, and anxiety-segment retention — not streaks or hint-free speed alone.
6. **Do not abandon DD for the anxious segment.** Permanent easy fluency is identity colonization by avoidance (XXV). Graded exposure to equipped difficulty is the path — therapy-adjacent but not a therapy claim (queue id 58).

Kill the false dichotomy: Bjork vs Ashcraft is not “science wars.” Bjork already says unequipped difficulty is undesirable; Ashcraft explains *why anxiety unequips* (WM). The product job is **re-equip + destake + sequence**, then apply DD.

XLI.7 Competitive / marketing implications

Competitor pattern	What it often does under anxiety	MindCraft wedge
Duolingo-like streaks	Turns miss into identity failure; pressures retrieval	Ban streak-as-learning; train without shame timers
Blocked textbook / many apps	Fluency comfort; weak strategy choice	Honest: early block OK; claim exam-ready only after mix
Brutal “exam sim” day one	High stakes + hard mix → Hinze-style disruption + exit	Soft sim → rise stakes after <i>retry_120s</i> healthy
ChatGPT tutor monologue	Soothes without generation; may reduce threat short-term	Student-generated <i>why before</i> AI wrap (XL); calm ≠ mastery
“Growth mindset” posters	Empty struggle praise without equipment	Struggle + scaffolds + mastery climate + transfer metrics

Marketing language that survives

- **Allowed:** “Hard practice only after you’re equipped — then we mix so exams don’t surprise you.”
- **Allowed:** “Feeling slower on mixed sets is the training signal, not a verdict.”

- **Allowed (parent):** “We separate *training retrieval* from *graded judgment* so practice builds memory instead of panic.”
- **Banned:** “Struggle always equals learning.”
- **Banned:** “We never make anxious students struggle.”
- **Banned:** Streak / blocked fluency / AI comfort as proof of readiness.
- **Banned:** Bloom 2-sigma; absolute tutoring-is-free.

Feature claim ladder (Part XXXIV hygiene)

Feature	Claim max without DD×ANX data
Soft-wrong + mastery copy on hard items	L1 affect-safe practice climate
Block → fade → interleave schedule	L1 sequenced DD (intent aligned XXXIX)
Low-stakes spaced retrieval before “exam mode”	L1 Hinze-aligned retrieval hygiene
“Our struggle pedagogy raises ACT for anxious students”	L0 until DD-1/DD-2
Always-on max mix / open SE under threat	L0 — likely undesirable difficulty

XL.I.8 Experiments spawned

ID	Question	Design	Primary	Kill condition
DD-1	Faded DD (block→mix + soft-wrong) vs early full mix same items	2-arm; stratify high vs low math-anxiety pretest	7d transfer_pass; retry_120s; exit after first miss	Early-mix wins transfer and retry in HMA segment → fade doctrine wounded
DD-2	Low-stakes retrieval sets vs same items framed as scored “exam drill”	2-arm (Hinze-style stakes)	Delayed retention; state anxiety; return D+1	High-stakes ≥ low-stakes on delayed + no anxiety cost → stakes doctrine overstated
DD-3	Structured SE under mix vs open SE under mix for HMA segment	2-arm (densifies SE-3)	Completion; transfer; hint binge	Open SE wins without churn → structured-default less necessary
DD-4	Repeated-item fluency path vs varied procedural path (Fioriti-inspired)	2-arm within HMA	Procedural transfer / strategy-class errors vs item-repeat accuracy	Repeat-fluency equals varied on transfer → anxiety-shaped preference harmless

DD-QUAL	10 Maya interviews: “When did hard practice feel like training vs attack?”	Appendix B	coded stake/equipment moments	Nobody separates → UX framing invisible
---------	---	------------	-------------------------------	---

Pre-reg (XXXIV): DD-* identify **schedule** × **stakes** × **anxiety** → **transfer/return**. They do **not** identify long-term identity from one hard session. Densifies Experiment D and IL-2/SE-3 as the anxiety-gated versions of desirable difficulty.

Supersedes / densifies: XXIX.4 “never introduce DD as shame”; XXXIX IL-2; XL SE-3; I.4 bottleneck → now has an operational SAFE-DD stack pending DD-1/2.

XLI.9 Confidence table

Claim	Label	Confidence
Spacing/interleaving/retrieval/generation can improve long-term learning vs fluent short-term conditions	FACT (Bjork lineage)	High
Difficulties are desirable only if the learner can respond successfully	FACT (Bjork & Bjork)	High
Math anxiety disrupts WM / acts like dual-task load	FACT (Ashcraft & Kirk, 2001; Ashcraft & Krause, 2007)	High
High stakes can erase retrieval-practice benefits on delayed tests	FACT (Hinze & Rapp, 2014)	High
Classroom retrieval programs often associate with <i>lower</i> reported exam nerves	FACT (Agarwal et al., 2014; self-report)	Medium–High
Practice testing meta-analytically reduces test anxiety (esp. easier quizzes)	FACT (Yang et al., 2023)	High
HMA may accelerate repeated-item retrieval learning while impairing procedural learning	FACT (Fioriti et al., 2025)	Medium–High (adult lab; arithmetic)
SAFE-DD (equip → destake → dose → frame → measure) is the right MindCraft default	FOUNDER BELIEF / HYPOTHESIS	Medium–High

“Struggle always” or “never struggle anxious kids” is good marketing	SPECULATION / false	Low (against)
--	---------------------	---------------

XL.I.10 What this chapter kills

1. **Kill:** “Desirable difficulties always help — turn difficulty up for everyone.” (Ignores Bjork’s own *equipped* clause + Ashcraft WM.)
2. **Kill:** “Anxious students should only get easy, blocked fluency forever.” (Serves avoidance; fails transfer; Fioriti-shaped trap.)
3. **Kill:** High-stakes “exam mode” as the default vehicle for retrieval practice. (Hinze & Rapp.)
4. **Wound:** Self-report that “quizzes reduce anxiety” as license for brutal day-one mixed sims — survives as *low-stakes, well-sequenced* retrieval, not as cruelty.
5. **Survive (constrained):** Faded, low-stakes, mastery-framed desirable difficulties that produce delayed mix evidence — the commercial spine for “exam-ready without panic theater.”

Doctrine until data: Ship **SAFE-DD:** equip → lower stakes → dose to WM → frame as training → measure delayed transfer and `retry_120s` in the anxiety segment. Market productive struggle only when those guards are visible. Never sell pain or comfort as mastery.

Part XLII — Social Comparison & Leaderboards: When Ranks Help vs Hurt Novices

Chapter status: Living evidence brief — Researcher tick 2026-07-30 (UTC hour 3)

Primary question: When do social-comparison surfaces (especially leaderboards) motivate learning versus threaten identity, and what should MindCraft default for Maya — the anxious / novice / identity-fragile segment?

Owners: Motivation · Affect / Anxiety · Product (Practice / Community UX) · Growth / Positioning · Red Team

Commercial job: Kill “ranks = engagement” as a North Star; harden the I.4 TARGET rule (default-ban appearance UX) into a shippable comparison policy competitors can be audited against.

XLII.1 Why this chapter exists

Part XXXVIII (goal orientation) already bans **appearance** goal structures and streak-as-North-Star. Part XXIV / XLI show social-evaluative threat taxes working memory and can turn desirable difficulty into threat. Competitive products (Duolingo-like ranks, classroom Top-N boards, public classroom streaks) still sell comparison as “motivation.”

FOUNDER BELIEF under audit: Identity transformation needs competence evidence — not a public ladder that tells bottom-half Maya she is behind *people* rather than behind a *skill criterion*.

Claim we refuse as doctrine: “Leaderboards motivate learners” (full stop) — and its twin “if we anonymize, any infinite rank is fine.”

XLII.2 Constructs (product language)

Construct	Plain definition	Product signal	Failure mode if misused
Social comparison	Evaluating self via others’ standing (Festinger)	Any peer rank, percentile, “class average,” public tutor callout	Turns mastery climate into appearance climate
Upward / downward / lateral	Compare to better / worse / similar others	Top-N board; “you beat 80%”; neighbor band	Extreme upward gap → give-up; downward → complacent ego boost without learning
Absolute vs relative board	All players ranked vs local neighbors / Top-N only	Infinite classroom board vs “near you” strip	Absolute + real names = max face threat for low ranks
Criterion / mastery feedback	Progress vs a task standard, not vs peers	Path progress, concept mastery bars, personal bests	Still can be gamed if criterion ≠ transfer
Normative / appearance feedback	Standing vs others; looking smart/fast	Public leaderboard, streak leagues	Identity threat; WM load; exit after first public miss

Operational definition (HYPOTHESIS): A MindCraft comparison surface is *safe* only when (a) the default is **self/criterion**, (b) peer comparison is **opt-in**, (c) comparison targets are **attainable** (local band, not unreachable Top-1), (d) metrics ranked are **effort/process-compatible** with mastery climate (not speed-only or streak-only), and (e) low standing is never **publicly named** for the anxiety / novice segment.

XLII.3 Festinger → competition: the comparison engine

FACT (origin): Festinger (1954, *Human Relations*, “A Theory of Social Comparison Processes”) — when objective standards are ambiguous, people evaluate abilities/opinions by comparing with others; ability comparison carries a unidirectional drive upward (improve toward better standards).

FACT (selection + reaction meta): Gerber, Wheeler & Suls (2018, *Psychological Bulletin*, doi:10.1037/bul0000127) — across 60+ years: when choosing up vs down, strong preference for **upward** comparison under no threat; **contrast** (self-evaluation moves *away* from the target) dominates reactions — ability estimates and affect often decline after upward contact. People often choose upward comparisons *even though* the typical reaction is self-deflating contrast.

FACT (competition model): Garcia, Tor & Schiff (2013, *Perspectives on Psychological Science*, doi:10.1177/1745691613504114) — social comparison is a major source of competitive attitudes/behavior; situational

amplifiers include **proximity to a standard** (near #1 vs far away) and **N** (few vs many competitors). Ranking UIs are engineered comparison machines.

Product translation: Shipping a leaderboard is not a neutral “engagement widget.” It is a deliberate injection of comparison → contrast risk → competitive framing. For Maya, that often means **appearance climate** (XXXVIII), not mastery.

XLII.4 Leaderboards can move behavior — under narrow conditions

FACT (HE experiment): Landers & Landers (2014, *Simulation & Gaming*, doi:10.1177/1046878114563662) — random assignment to a course wiki **leaderboard** (not tied to grades) increased project interactions (~+29.6) vs control; time-on-task mediated academic performance. Supports: leaderboards *can* raise engagement via goal/conflict attributes **in that setting**.

Caveats (FACT / scope): Undergraduate / course project; not math-anxious high-school Maya; not public infinite classroom shame; achievement metric was interaction volume, not durable transfer. Landers’ theory of gamified learning (2014, doi:10.1177/1046878114563660) treats game attributes as mediators/moderators of *instructional* quality — gamification does not replace pedagogy.

FACT (HE systematic review): Li, Liang, Fryer & Shum (2024, *Journal of Computer Assisted Learning*, doi:10.1111/jcal.13077) — 20 articles / 22 studies / 29 interventions (2014–2023), **all higher education**; ~half ≤1 hour. Leaderboards *can* benefit motivation/engagement/performance, but **effects hinge on design**. Authors recommend (among others): absolute boards for motivation *with* anonymity especially for low ranks; mix performance + engagement criteria; caution that HE short-duration findings do **not** generalize cleanly to K–12.

HYPOTHESIS: Landers-style gains are most likely when learners already have task competence + low identity threat + voluntary/low-stakes ranking. MindCraft’s default user is the opposite segment on day one.

XLII.5 Math-specific: anxiety × infinite ranks

FACT (primary maths game, large N): Almo, Rocha, Brennan & Dondio (2024, *International Journal of Serious Games*, doi:10.17083/ijsg.v11i4.794) — N=1,389 Irish primary students, 6-week *Seven Spells* programme with an **infinite leaderboard**. Leaderboard enjoyment predicted by rank position **and** maths anxiety: **maths-anxious players disliked the leaderboard more than non-anxious peers even after controlling for position**. Preference for playing against others also predicted enjoyment. Small correlation between liking the board and liking the game.

Implication (HYPOTHESIS): For anxious learners, public infinite rank is not a free engagement boost — it is a trait-moderated threat surface. “They’re low because they’re bad at the game” does not exhaust the mechanism; anxiety predicts dislike **net of rank**.

FACT (negative feedback climate): Fong, Patall, Vasquez & Stautberg (2019, *Educational Psychology Review*, doi:10.1007/s10648-018-9446-6) — meta-analysis (78 studies): negative feedback reduces intrinsic motivation vs positive feedback; harm lessens when feedback includes instructional how-to-improve detail and uses **criterion-based** standards (vs purely normative). Public low rank is often pure normative negative feedback without instructional content.

Bridge to XLI / XXIV: Social-evaluative threat (public miss, public bottom third) is exactly the stakes Hinze & Rapp warn can disrupt retrieval benefits — and Ashcraft-style WM load. Infinite boards amplify that climate at scale.

XLII.6 Resolution doctrine — SAFE-COMPARE

FOUNDER BELIEF / HYPOTHESIS — SAFE-COMPARE stack (ties TARGET / SAFE-DD / FEI):

1. **Default = criterion, not peers.** Show progress vs path / mastery / personal best. Peers off by default.
2. **Opt-in competition only.** Never force infinite boards onto gap-scan novices or high-anxiety segment.
3. **If peers: local + attainable.** Prefer neighbor bands / similar-ability cohorts over Top-1 theater (Garcia proximity + Gerber contrast).
4. **Anonymize low standing.** No named public bottom ranks (Li et al. Rec. 2; Bai et al. cited therein).
5. **Rank what mastery climate allows.** Prefer process-compatible signals (retries completed, transfer attempts, concepts unlocked) over speed leagues / streak leagues. Never rank “hint-free seconds” as identity.
6. **Instructional wrap on any relative miss.** If relative feedback ships, attach *how to improve next* (Fong moderators) — not only “you’re #47.”
7. **Measure the right win.** `retry_120s`, `challenge_accept` (mastery-coded), `transfer_pass` / `solo_transfer_pass` — not `leaderboard_views` or rank delta as learning.

Kill the false dichotomy: “Competition is toxic” vs “competition is always motivating.” Competition can raise time-on-task for some prepared learners (Landers). For Maya default, **unbuffered infinite peer rank is appearance UX** — banned until SC experiments say otherwise.

XLII.7 Competitive / marketing implications

Competitor pattern	Comparison mechanism	MindCraft wedge
Duolingo-like leagues / ranks	Normative weekly standing; streak adjacent	Mastery path progress; ban streak/rank as learning proof
Classroom Top-N boards	Extreme upward targets; face threat	Criterion bars; optional local bands
Khan / practice dashboards	Often self-paced; weaker public rank	Keep self/criterion; don’t “gamify catch-up” with shame ranks
Brilliant prestige signals	Ability-appearance among peers	Sell <i>solo transfer</i> and <i>challenge-seeking why</i> , not percentile flex
ChatGPT tutor alone	No peer board — but social comparison still via school/parents	Own the <i>safe</i> comparison story: progress vs yesterday’s self

Marketing language that survives

- **Allowed:** “Progress against *your* path — not a public ladder.”
- **Allowed:** “Competition is optional; mastery climate is the default.”
- **Allowed (parent):** “We don’t put anxious kids on a named scoreboard to ‘motivate’ them.”
- **Banned:** “Leaderboards keep students hooked” as a learning claim.
- **Banned:** Rank / league / streak as North Star or as proof of identity change.
- **Banned:** Bloom 2-sigma; absolute tutoring-is-free; empty mindset posters.

Feature claim ladder (Part XXXIV hygiene)

Feature	Claim max without SC data
Criterion progress / personal bests default	L1 mastery-climate alignment (XXXVIII)
Opt-in anonymized local band	L1 SAFE-COMPARE intent
“Our leaderboard raises ACT / identity for anxious Maya”	L0 until SC-1/SC-2
Forced infinite named classroom board	L0 — appearance UX; likely harm for HMA

XLII.8 Experiments spawned

ID	Question	Design	Primary	Kill condition
SC-1	Criterion-only progress vs infinite named leaderboard (same practice items)	2-arm; stratify high vs low math-anxiety	retry_120s; exit after public miss; 7d transfer_pass	Infinite board wins transfer and retry in HMA → default-ban wounded
SC-2	Opt-in local anonymized band vs forced absolute board	2–3 arm	Enjoyment; state anxiety; return D+1; challenge_accept motive codes	Forced absolute ≥ opt-in local on retention without anxiety cost → anonymity/opt-in overstated
SC-3	Rank process metrics (retries, mix attempts) vs rank speed/streak	2-arm within volunteers	Strategy-class errors; transfer_pass vs session speed	Speed rank equals process rank on transfer → metric choice moot
SC-4	Criterion negative feedback + how-to-improve vs rank-only “you’re #N”	2-arm (Fong-inspired)	Intrinsic motivation proxy; retry	Rank-only ≥ instructional criterion → Fong moderator fails in product

SC-QUAL	10 Maya interviews: “When did seeing others’ scores help vs freeze you?”	Appendix B	coded comparison threat moments	Nobody distinguishes → copy invisible
---------	---	------------	---------------------------------	---------------------------------------

Pre-reg (XXXIV): SC-* identify **comparison surface** × **anxiety** → **return/transfer**. They do **not** identify that “healthy competition builds character” from one league season.

Densifies: XXXVIII appearance ban; I.4 TARGET “leaderboard kill for Maya default”; XXIV social-evaluative threat hypothesis.

XLII.9 Confidence table

Claim	Label	Confidence
People use others to evaluate abilities when standards are ambiguous	FACT (Festinger, 1954)	High
Upward comparison often preferred; contrast often dominates self-eval/affect	FACT (Gerber et al., 2018)	High
Ranking contexts amplify competitive social comparison	FACT (Garcia et al., 2013)	High
Leaderboards can increase time-on-task / mediate performance in some HE settings	FACT (Landers & Landers, 2014)	High (scope-limited)
Leaderboard effects in education are design-dependent; HE evidence ≠ K–12	FACT (Li et al., 2024)	High
Maths-anxious players dislike infinite boards more, net of rank	FACT (Almo et al., 2024)	High
Normative negative feedback tends to hurt intrinsic motivation vs positive; criterion + how-to softens	FACT (Fong et al., 2019)	High
SAFE-COMPARE default (criterion, opt-in, local, anonymize low, process metrics) is right for MindCraft	FOUNDER BELIEF / HYPOTHESIS	Medium–High
“Leaderboards motivate everyone / raise learning for Maya”	SPECULATION / false as universal	Low (against)

XLII.10 What this chapter kills

1. **Kill:** Universal “leaderboards motivate students” as product doctrine or marketing. (Design- and trait-dependent; HE ≠ Maya.)
2. **Kill:** Infinite named public ranks as default engagement for novices / high math-anxiety. (Almo et al.; TARGET / appearance.)
3. **Kill:** Rank delta / league standing / streak leagues as North Star or as proof of identity transformation.
4. **Wound:** Absolute-board motivation gains from HE reviews as license for K–12 public boards — survives only as *design hypothesis*, not as Maya default.
5. **Survive (constrained):** Opt-in, anonymized, attainable, criterion-wrapped comparison for learners who already seek competition — never as the onboarding spine.

Doctrine until data: Ship **SAFE-COMPARE:** criterion default → peer opt-in → local/attainable → anonymize low standing → rank mastery-compatible process → instructional wrap → measure transfer/retry, not rank views. Market “progress without public ladders.” Never sell ranks as belonging.

Part XLIII — Habit Formation Science (Wood/Clear): Cue Design Without Identity Colonization

Chapter status: Living evidence brief — Researcher tick 2026-07-30 (UTC hour 6 ≡ Red Team slot; ch43 never written → prefer Researcher per rotation)

Primary question: How should MindCraft use habit science (context cues, repetition, automaticity) to raise *return to practice* without making streaks / daily check-ins the student’s identity — or mistaking habit for learning?

Owners: Motivation · Product (Practice / notifications / recognition) · Growth / Positioning · Red Team

Commercial job: Harden Insight 5 / HID into a shippable **SAFE-HABIT** stack competitors (especially Duo-like streak products) can be audited against; kill “habit = mastery” and “21 days” folklore in copy and KPIs.

XLIII.1 Why this chapter exists

Parts XXV (habit ≠ identity), XXXVIII (appearance / streak North Star ban), XLI (SAFE-DD), and XLII (SAFE-COMPARE) already treat streaks and ranks as threat surfaces when they become *worth*. Market gravity still pulls edtech toward Duolingo-shaped habit loops: push cue → micro-session → streak reward → loss aversion.

FOUNDER BELIEF under audit: Identity transformation needs *reliable return to difficulty* — so cue design is load-bearing — but the habit that forms must be “open practice and try the next hard thing,” not “protect a number that proves I am consistent.”

Claims we refuse as doctrine:

1. “Streaks = learning” / streak length as North Star.

2. “Habits take 21 days” as science.
3. Treating popular packaging (*Atomic Habits*) as peer-reviewed habit theory without Wood/Gardner/Lally grounding.
4. Habit loops that **colonize identity** — the student becomes “a streak-keeper” rather than “someone who can recover from wrong.”

XLIII.2 Constructs (product language)

Construct	Plain definition	Product signal	Failure mode if misused
Habit (scientific)	Learned context→response association in procedural memory; cue activates response with limited goal mediation (Wood)	Stable time/place/preceding-action cue for practice open	Shipping “motivation” when the real lever is <i>cue stability</i>
Automaticity	Behavior feels efficient / less deliberative	Self-report automaticity; shorter time-to-first-problem	Confusing app-open automaticity with <i>math skill</i> automaticity
Context cue	Physical setting, time, people, prior action in a sequence	After-school phone dock; tutor session end → 10-min solo; calendar block	Push spam as fake context; unstable cues never habitize
Implementation intention	If-then plan: when Y, do X (Gollwitzer)	Onboarding “After dinner, I open Minecraft Map”	Vague goals (“practice more”) with no situational trigger
Identity colonization (product sense)	Habit metric becomes who the student <i>is</i> / fears losing	Streak as self-worth; freeze culture; XP-as-self	HID failure — return without identity change; break → exit
SAFE-HABIT	Cue + small start + miss-forgiveness + learning metrics above habit metrics	See XLIII.6	Habit theater that still ranks streak / DAU as proof

Operational definition (HYPOTHESIS): A MindCraft habit surface is *safe* only when (a) cues target **opening equipped practice**, not protecting a counter, (b) miss recovery is cheap (Lally: one miss ≠ wipe), (c) recognition prefers process/identity markers over streak length, and (d) North Stars remain `retry_120s` / `transfer_pass` / `solo_transfer_pass` / identity items — never streak days.

XLIII.3 Wood: habits are context–response, not vibes

FACT (review): Wood & Rünger (2016, *Annual Review of Psychology*, doi:10.1146/annurev-psych-122414-033417) — habits are learned associations between recurring **context cues** and responses, strengthened gradually via associative / reward learning; once formed, perception of the cue activates the response representation with limited need for executive control. Defining features of habit automaticity include cue activation and **relative insensitivity to short-term goal/outcome changes**.

FACT (habit–goal interface): Wood & Neal (2007, *Psychological Review*, “A New Look at Habits and the Habit–Goal Interface”) — goals help *start* repetition and exposure to cues; once habits form, contexts can trigger responses without a mediating goal. Habit change therefore often requires changing cues / environments, not only “trying harder.”

FACT (systems view): Wood, Mazar & Neal (2021, *Perspectives on Psychological Science*, habits vs goals as separate but interacting systems) — habitual responses can activate regardless of current goals; people may still inhibit them, but cue-driven activation is real. Product translation: a streak notification is a **Pavlovian/context cue** engineered by the app — not proof the student wants math.

Product translation: If MindCraft wants return, design **stable contexts** (same after-school window; end-of-tutor-session stack; device home-screen tile) more than pep talks. If MindCraft wants identity, do **not** make the strongest cue “don’t break the number.”

XLIII.4 Clear / Fogg packaging vs the evidence spine

FACT (popularization, not primary science): Clear (*Atomic Habits*, 2018) packages cue / craving / response / reward loops and “habit stacking” (“After [current habit], I will [new habit]”) for a mass audience. Fogg (*Tiny Habits*) popularized tiny start + anchor habits. Neither replaces Wood’s procedural-memory account.

FACT (what stacking actually is): Habit stacking’s evidence spine is largely **implementation intentions** — Gollwitzer & Sheeran (2006, *Advances in Experimental Social Psychology*, doi:10.1016/S0065-2601(06)38002-1): across 94 tests, if-then plans show a medium-to-large effect on goal attainment ($d \approx 0.65$), via heightened cue accessibility and automated initiation.

HYPOTHESIS: Clear is usable *copy scaffolding* for onboarding (“After tutoring, open Map for one mission”) if MindCraft cites the Gollwitzer mechanism and measures initiation → equipped practice, not Clear’s compound-growth metaphors as science.

Kill: Marketing that says “built on Atomic Habits science” as if the book were a methods paper. Prefer: “cue-based practice plans (implementation intentions; habit research).”

XLIII.5 How long, how fragile, what to measure

FACT (timeline myth): Lally, van Jaarsveld, Potts & Wardle (2010, *European Journal of Social Psychology*, doi:10.1002/ejsp.674) — daily repetition in a consistent context; automaticity rises on an asymptotic curve; among good-fit participants, median time to ~95% of asymptote $\approx$ **66 days** (range **18–254**). Missing **one** opportunity did not materially derail formation. ~Half of motivated participants still failed to perform consistently enough to reach habit status in the study window.

Kill: “21 days to form a habit” (Maltz folklore, not this literature). Also kill product copy that implies a week-long streak *is* a habit.

FACT (definition hygiene): Gardner (2015, *Health Psychology Review*, doi:10.1080/17437199.2013.876238) — habit as a *process* whereby a stimulus generates an impulse to act via learned S–R association; impulse need not equal behavior. Most health-habit studies are correlational / self-report-heavy — confidence ceiling for edtech transfer claims.

FACT (measurement trap): Verplanken & Orbell (2003, *Journal of Applied Social Psychology*, doi:10.1111/j.1559-1816.2003.tb01951.x) — SRHI includes history of repetition, automaticity, **and expressing identity**. Self-report habit strength can smuggle identity language into a frequency metric.

Product implication (HYPOTHESIS): Do not treat SRHI-like “this is typically me” items as proof of *math identity transformation* (Part XXV / id 61 queue). Separate: (1) practice-open automaticity, (2) math self-efficacy / narrative identity scales, (3) transfer metrics.

XLIII.6 Duolingo-shaped products: habit works — and colonizes

FACT (industry, self-report): Duolingo publicly designs streaks using habit research framing; company blog claims learners who reach short streaks are far more likely to continue (e.g., 7-day streak correlational lifts — *not* causal RCTs published as independent learning science). Streak freezes / separation of streak vs daily goal are retention experiments on **engagement**, not transfer.

FACT (misuse qualitative): Hadi Mogavi et al. (2022, *Learning @ Scale*, doi:10.1145/3491140.3528274) — forum analysis + interviews on Duolingo: **gamification misuse** when users fixate on XP/streaks/leaderboards and get distracted from learning; reasons include competitiveness, playfulness overlearning, dark nudges, compulsion, herding. Ramifications touch learning efficiency, well-being, and ethics.

Bridge to XLII / XXXVIII: Streak leagues + infinite ranks are the same family as appearance climates — consistency theater. Wood predicts cues will keep firing even when the *learning* goal cooled; Hadi Mogavi shows the product can train the wrong response (protect streak / farm XP).

FOUNDER BELIEF: MindCraft may borrow **cue stability** + **small start** + **miss forgiveness**. MindCraft must not borrow **streak-as-self** or **loss-aversion-as-pedagogy**.

XLIII.7 Resolution doctrine — SAFE-HABIT

FOUNDER BELIEF / HYPOTHESIS — SAFE-HABIT stack (ties HID / TARGET / SAFE-DD / SAFE-COMPARE):

1. **Cue the practice, not the counter.** Default cues: calendar/time, tutor-session end stack, story-world return point — not “your streak dies at midnight” as the primary cue.
2. **Implementation intention onboarding.** One written if-then in gap-scan aftermath or first session (“After [stable cue], I open one mission”). Measure initiation, not slogan recall.

3. **Tiny equipped start.** First action $\leq$ friction of one problem with SAFE-DD equipment — not a multi-screen habit tutorial.
4. **Miss $\neq$ identity failure.** One miss does not reset belonging (Lally). Prefer gentle resume copy over shame / “you broke it.” Streak freezes are a *retention patch*, not a learning feature — if shipped, never as North Star.
5. **Habit metrics subordinate.** Track D1/D7 return, cue-honoring rate, time-to-first-problem as *diagnostics*. Learning North Stars remain retry / transfer / solo transfer / challenge_accept (mastery-coded).
6. **Identity markers elsewhere.** Map fill, story unlocks, tutor recognition of *recovery*, parent “they tried the hard one again” — not streak length as “who they are.”
7. **No identity colonization.** Ban copy that equates worth with unbroken days. Ban ranking streaks (XLII). Ban streak leagues as proof of transformation.

Kill the false dichotomy: “We refuse all habit design” vs “we must out-Duo Duo.” Refuse *colonizing* habit design; keep *cue science*.

XLIII.8 Competitive / marketing implications

Competitor pattern	Habit mechanism	MindCraft wedge
Duolingo streaks / freezes	Loss aversion + daily cue; retention-first	Cue practice return; miss-forgiving; never streak-as-learning
Duo leagues + XP	Habit + appearance compound	SAFE-COMPARE + SAFE-HABIT; sell solo transfer
Khan session streaks / energy	Mild habit / progress bars	Keep criterion progress; don’t add shame timers
Brilliant daily puzzles	Delight + return habit	Return <i>into</i> recoverable struggle, not only delight
ChatGPT “study every day” tips	Advice without cue architecture	Own <i>if-then</i> + <i>equipped mission</i> as the plan

Marketing language that survives

- **Allowed:** “Same time, same cue — open one hard problem and recover.”
- **Allowed:** “Habits get you to the desk; identity is what you do when it’s wrong.”
- **Allowed (parent):** “We build a practice cue without making a broken streak mean failure.”
- **Banned:** “Streaks create math identity.”
- **Banned:** “21-day habit challenge = mastery.”
- **Banned:** Streak / DAU as North Star or as proof of transformation.
- **Banned:** Bloom 2-sigma; tutoring-is-free; empty mindset posters.

Feature claim ladder (Part XXXIV hygiene)

Feature	Claim max without HAB data
---------	----------------------------

If-then practice cue + tiny equipped start	L1 initiation design (Gollwitzer-aligned)
Miss-forgiving resume without streak shame	L1 HID hygiene
“Our streak raises ACT / math identity”	L0 until HAB-1/HAB-2
Forced streak + public streak rank as default	L0 — colonization UX

XLIII.9 Experiments spawned

ID	Question	Design	Primary	Kill condition
HAB-1	If-then cue plan vs streak-primary onboarding (same missions)	2-arm; new users	D7 return; time-to-first-problem; 14d transfer_pass	Streak arm wins transfer and identity item → on-ramp-only doctrine wounded
HAB-2	Miss-shame copy (“you broke your streak”) vs miss-forgiving resume	2-arm after forced miss	Return D+1; state anxiety; retry_120s	Shame ≥ forgive on return without anxiety cost → forgive overstated
HAB-3	Stable calendar cue vs push-only reminders	2-arm	Cue-honoring rate; 4-week automaticity proxy	Push-only equals calendar on return → context-stability claim weak
HAB-4	Recognition: process/recovery badge vs streak badge vs none	3-arm (ties GO-2)	challenge_accept motives; identity item; streak obsession items	Streak badge wins engagement and transfer → HID weaker than feared
HAB-QUAL	10 Maya interviews: “What does your streak (elsewhere) make you feel when you miss?”	Appendix B	coded colonization vs helpful cue	Nobody reports threat → copy over-indexed

Pre-reg (XXXIV): HAB-* identify **cue/recognition** × **return/transfer**. They do **not** identify that “daily habit builds mathematicians” from DAU alone.

Densifies: Insight 5 / HID; I.4 TARGET streak ban; XXXVIII appearance; XLII streak-league kill.

XLIII.10 Confidence table

Claim	Label	Confidence
Habits are context–response associations; cue-triggered; relatively goal-insensitive once strong	FACT (Wood & Rünger, 2016; Wood & Neal, 2007)	High
Goals help form habits via repetition/exposure; changing cues aids habit change	FACT (Wood line)	High
Implementation intentions raise goal attainment (medium–large)	FACT (Gollwitzer & Sheeran, 2006)	High
Automaticity plateaus on the order of weeks–months; median ~66d in Lally sample; high variance; one miss often OK	FACT (Lally et al., 2010)	High
“21 days” is not the Lally result	FACT	High
Clear/Fogg are popular packaging; stacking ≈ if-then planning	HYPOTHESIS / FACT (mechanism mapping)	Medium–High
Gamification can distract from learning (Duo case)	FACT (Hadi Mogavi et al., 2022)	High (qualitative)
Streaks can raise return without identity / transfer	HYPOTHESIS (HID; industry)	Medium–High
SAFE-HABIT is right default for MindCraft	FOUNDER BELIEF / HYPOTHESIS	Medium–High
“Streaks = learning / identity”	SPECULATION / false as universal	Low (against)

XLIII.11 What this chapter kills

1. **Kill:** Streak length / DAU / “habit streak” as North Star or as proof of identity transformation. (HID; standing ban densified with Wood + Duo misuse.)
2. **Kill:** “21-day habit” marketing as science. (Lally variance + asymptote.)
3. **Kill:** Equating *Atomic Habits* branding with peer-reviewed habit doctrine. (Use Wood/Gardner/Gollwitzer; Clear as optional UX copy only.)
4. **Kill:** Shame-on-miss and public streak ranks as default engagement for Maya. (Colonization; XLII.)
5. **Wound:** “Any daily reminder is habit science” — survives only if tied to **stable context + repeated equipped response**, not push spam.
6. **Survive (constrained):** Cue design, tiny starts, implementation intentions, and miss forgiveness as **on-ramps** to FEI / transfer — never as the product thesis.

Doctrine until data: Ship **SAFE-HABIT**: cue practice (not the counter) → if-then onboarding → tiny equipped start → miss ≠ identity failure → habit metrics subordinate → identity markers elsewhere → ban streak colonization. Market “we get you back to the hard problem — without making a broken streak mean you’re not a math person.”

Part XLIV — Intrinsic Motivation Killers (Deci): Controlling Rewards and the Caveats

Chapter status: Living evidence brief — Researcher tick 2026-07-30 (UTC hour 9; hour%6≠0; 3 researcher entries since synthesizer v1.6 → Researcher, not Synthesizer)

Primary question: When do MindCraft rewards, XP, badges, streaks, and “do this for the prize” UX *undermine* the intrinsic interest and autonomy we claim to build — and when are they informational rather than controlling?

Owners: Motivation · Product (Practice / recognition / parent copy) · Growth / Positioning · Red Team

Commercial job: Turn the vague “extrinsic bad / Deci says so” slogan into a shippable **SAFE-REWARD** stack competitors can be audited against; kill both “all rewards kill love of math” absolutism and “gamify everything” as identity strategy.

XLIV.1 Why this chapter exists

Parts XXV (habit ≠ identity), XXXVIII (appearance climates), XLII (SAFE-COMPARE), and XLIII (SAFE-HABIT) already treat streaks, ranks, and consistency theater as threat surfaces. Market gravity still pulls edtech toward *expected tangible / token rewards for doing the task*: XP for opening the app, coins for finishing any problem, badges for login streaks, parent dashboards that score “minutes practiced.”

FOUNDER BELIEF under audit: Identity transformation needs *felt autonomy* and *competence information* — not the absence of all rewards. The killer is **controlling functional significance** (pressure to behave for the inducement), not the existence of feedback or unexpected acknowledgment.

Claims we refuse as doctrine:

1. “All rewards kill intrinsic motivation” (Cameron–Pierce counter-literature exists; CET is differentiated).
 2. “Gamification is always fine if engagement rises” (engagement ≠ free-choice interest; Hadi Mogavi / Duo misuse).
 3. “Verbal praise always helps kids” (Deci: controlling praise / “you should keep it up” can undermine).
 4. Treating SDT as a poster (“autonomy / competence / relatedness”) without contingency design.
-

XLIV.2 Constructs (product language)

Construct	Plain definition	Product signal	Failure mode if misused
-----------	------------------	----------------	-------------------------

Intrinsic motivation	Doing the activity for inherent interest / satisfaction	Free-choice return to hard problems; curiosity asks; challenge_accept for mastery reasons	Confusing “opens app for XP” with interest
Extrinsic motivation (spectrum)	Behavior for separable outcome; can be more/less autonomous (external → introjected → identified → integrated)	Parent-pressure vs “I chose ACT path because...”	Treating all extrinsic as poison or all as fine
Controlling vs informational	Event felt as pressure to behave vs competence feedback (CET)	“Finish 10 for a coin” vs “you used the bridge correctly”	Same badge can be either — design of <i>how</i> it is framed
Engagement-contingent reward	Expected reward for <i>doing</i> the task (no quality bar)	XP for any attempt; login coins	Classic undermine path (Deci 1999)
Completion-contingent	Expected reward for finishing	“Complete mission → loot” regardless of strategy quality	Similar undermine; weak competence signal
Performance-contingent	Reward for meeting a standard	Badge for transfer pass; tutor note for recovery	Can inform competence <i>or</i> control — climate decides
SAFE-REWARD	Prefer informational feedback; avoid expected tangible for mere engagement; autonomy-supportive framing	See XLIV.7	Reward theater that still trains external PLOC

Operational definition (HYPOTHESIS): A MindCraft reward surface is *safe* only when (a) it primarily conveys **competence / strategy information**, (b) expected tangible / token rewards are **not** the reason to open practice, (c) framing is autonomy-supportive (choice, rationale, non-pressuring language), and (d) North Stars remain `retry_120s` / `transfer_pass` / `solo_transfer_pass` / `identity` items — never XP velocity or badge count.

XLIV.3 SDT spine (needs, not posters)

FACT (overview): Ryan & Deci (2000, *American Psychologist*, 55(1), 68–78, doi:10.1037/0003-066X.55.1.68) — Self-Determination Theory posits innate needs for **competence, autonomy, and relatedness**; social contexts that support these needs facilitate intrinsic motivation and internalization; contexts that thwart them yield alienation and depleted motivation.

FACT (CET): Cognitive Evaluation Theory (Deci & Ryan line; summarized in Deci, Koestner & Ryan, 1999) — rewards and external events have **functional significance**: controlling aspects shift perceived locus of causality

outward (undermine intrinsic motivation); informational aspects affirm competence (can enhance). Competence feedback without autonomy does not reliably enhance intrinsic motivation.

Product translation: MindCraft’s brand voice (“you’re becoming someone who recovers”) must not be paired with UX that says “do this or lose the prize.” Relatedness (tutor, story world, parent) is not a license for controlling rewards.

XLIV.4 The classic finding — and the differentiated meta

FACT (origin): Deci (1971, *Journal of Personality and Social Psychology*, 18(1), 105–115, doi:10.1037/h0030644) — expected extrinsic rewards for an interesting activity can reduce subsequent free-choice engagement (the “undermining” / overjustification pattern that launched the literature).

FACT (Deci–Koestner–Ryan 1999 meta): Deci, Koestner & Ryan (1999, *Psychological Bulletin*, 125(6), 627–668, doi:10.1037/0033-2909.125.6.627) — 128 studies. As predicted by CET:

- **Engagement-contingent** rewards undermine free-choice intrinsic motivation ($d \approx -0.40$).
- **Completion-contingent** ($d \approx -0.36$) and **performance-contingent** ($d \approx -0.28$) also undermine free-choice behavior on average.
- Engagement- and completion-contingent also hurt self-reported interest ($d \approx -0.15$ / -0.17).
- **Unexpected** tangible rewards: no average undermine on free-choice.
- **Positive feedback / verbal rewards:** enhance free-choice ($d \approx 0.33$) and interest ($d \approx 0.31$) on average.
- Tangible rewards tended to be **more detrimental for children** than college students; verbal rewards less enhancing for children than for college students.

FACT (education restatement): Deci, Koestner & Ryan (2001, *Review of Educational Research*, 71(1), 1–27, doi:10.3102/00346543071001001) — argue Cameron & Pierce (1994) understated undermining; conclude tangible rewards have a substantial undermining effect when CET contingencies are respected; defend keeping CET for educational practice.

Product implication (HYPOTHESIS): Expected XP/coins for *mere engagement* (open app, any answer, any completion) is the highest-risk MindCraft pattern for Maya-aged users. Informational soft-wrong + strategy feedback is the low-risk pattern. Performance-contingent badges *can* work if they mark real competence (transfer) and are not administered as pressure.

XLIV.5 Caveats — Cameron/Pierce, climate, and “not all rewards”

FACT (counter-meta): Cameron & Pierce (1994, *Review of Educational Research*, 64(3), 363–423, doi:10.3102/00346543064003363) — across 96 between-groups experiments, *overall* reward did not decrease intrinsic motivation; verbal praise increased it; the clearest negative effect was expected tangible rewards for *simply doing* a task (free-choice time after reward removal). They argued negative effects are circumscribed and avoidable.

FACT (debate): Ryan & Deci (1996, *RER*), Lepper et al. (1996), Kohn (1996) vs Cameron & Pierce (1996, *RER*) — methodological dispute over aggregation, moderators, and what “overall” means. **Not settled as “rewards are harmless.”**

FOUNDER BELIEF / resolution for MindCraft: Do **not** cite “Deci proved all gamification is evil.” Do **not** cite Cameron to greenlight engagement-contingent loot. Cite the **contingency**: expected tangible for mere doing is the kill zone; informational verbal/competence feedback is the build zone; performance-contingent is climate-sensitive.

FACT (controlling praise): Ryan (1982) and Ryan, Mims & Koestner (1983) — positive feedback administered with controlling language (e.g., “you *should* keep up the good work”) reduces intrinsic motivation relative to informational administration. Deci 1999 summarizes interpersonal climate as a moderator of performance-contingent rewards.

Kill: “We praise a lot, so SDT is satisfied.” Praise can be controlling.

FACT (choice as autonomy support): Patall, Cooper & Robinson (2008, *Psychological Bulletin*, 134(2), 270–300, doi:10.1037/0033-2909.134.2.270) — meta of 41 studies: providing choice enhances intrinsic motivation, effort, task performance, and perceived competence. Effects stronger for instructionally *irrelevant* choices, when ~2–4 successive choices are offered, when rewards are **not** stacked after the choice manipulation, and vs highly controlling controls; stronger for children than adults.

Product implication: Meaningful micro-choice (which story frame, which representation after mastery structure is set) can support autonomy — but **choice + immediately dangling reward** weakens the choice benefit. Do not “choose your path → earn 50 coins.”

XLIV.6 Competitive patterns under CET audit

Competitor pattern	CET reading	MindCraft wedge
Duolingo XP / streaks / leagues	Expected engagement- and completion-contingent tokens + appearance climate (XLII/XLIII)	Informational recovery feedback; no XP-as-reason-to-open; SAFE-HABIT cues
Khan energy / points-ish progress	Mild completion signals; mostly mastery bars	Keep criterion progress; avoid coinization of minutes
Brilliant delight unlocks	Intrinsic interest scaffold; some completion loot	Keep delight <i>inside</i> the problem; don’t pay for opening
ChatGPT “I’ll quiz you if you want”	Low token reward; high fluency risk (XXXIII)	Autonomy of asking ≠ competence evidence; measure solo transfer
Parent dashboards: minutes / day streak	Parent-facing controlling contingency spills to student	Report <code>retry_120s</code> , transfer, <code>challenge_accept</code> — not “minutes bought with coins”

Bridge: XLIII says cue the practice, not the counter. XLIV says: if you must reward, reward **informative competence**, not engagement.

XLIV.7 Resolution doctrine — SAFE-REWARD

FOUNDER BELIEF / HYPOTHESIS — SAFE-REWARD stack (ties TARGET / SAFE-COMPARE / SAFE-HABIT / FEI):

1. **No expected tangible / token for mere engagement.** Ban default XP-for-open, coins-for-any-answer, loot-for-login. (Deci 1999 engagement-contingent.)
2. **Feedback > inducement.** Soft-wrong, strategy labels, bridge-gap naming, and “you recovered” tutor notes are the primary reinforcer. Prefer unexpected acknowledgment over announced prizes.
3. **If performance-contingent, mark real competence.** Badges/unlocks tied to `transfer_pass` / `solo_transfer_pass` / `mastery-coded challenge_accept` — not to streak days or raw volume. Administer without “you should” pressure.
4. **Autonomy-supportive framing.** Offer rationale; allow meaningful micro-choice (Patall); never “do this or lose the reward.”
5. **Don’t stack choice with dangling rewards.** Choice then immediate coin undermines the autonomy gain.
6. **Child-risk premium.** Maya-aged users are in the higher tangible-harm band of Deci 1999 — stricter default than adult SaaS gamification norms.
7. **Metrics hygiene.** Track XP/badge counts as *diagnostics of system behavior*, never as North Stars. Learning NS remain `retry` / `transfer` / `solo transfer` / `identity`.
8. **Parent copy:** Sell “she chose the harder one again” and competence evidence — not “she earned 400 XP.”

Kill the false dichotomy: “No rewards ever” vs “gamify everything.” Refuse *controlling engagement loot*; keep *informational competence signals* and story unlocks that are earned by transfer, not by attendance.

XLIV.8 Marketing language that survives

- **Allowed:** “We don’t pay kids to open the app — we show them what they figured out.”
- **Allowed:** “Unlocks follow real transfer, not login streaks.”
- **Allowed (parent):** “Progress you can trust: recovery and transfer — not points theater.”
- **Banned:** “Rewards create love of math.”
- **Banned:** “XP builds identity.”
- **Banned:** Absolute “all extrinsic motivation is bad” / “tutoring is free so motivation is free.”
- **Banned:** Bloom 2-sigma; empty mindset posters; streak/XP as North Star.

Feature claim ladder (Part XXXIV hygiene)

Feature	Claim max without IM data
Informational soft-wrong + strategy feedback (no engagement loot)	L1 CET-aligned hygiene
Unexpected tutor/parent acknowledgment after recovery	L1 (unexpected tangible / social)
“Our XP system raises intrinsic math interest”	L0 until IM-1/IM-2
Expected coins for any completion as default	L0 — engagement-contingent kill zone

XLIV.9 Experiments spawned

ID	Question	Design	Primary	Kill condition
IM-1	Engagement-loot (XP for any attempt) vs no-loot informational feedback (same missions)	2-arm; new users	Free-choice analog; voluntary hard-problem rate post-session; 14d interest item; transfer_pass	Loot arm wins free-choice and transfer → engagement-loot ban overstated
IM-2	Performance badge on transfer_pass (informational copy) vs same badge with controlling copy (“you should keep earning”)	2-arm	challenge_accept motives; interest; return	Controlling ≥ informational on interest → climate claim weak
IM-3	Micro-choice without reward vs choice + dangling coins	2-arm (Patall moderator)	Intrinsic motivation proxy; effort; autonomy need satisfaction	Choice+coins ≥ choice-only → “don’t stack” overstated
IM-4	Unexpected recovery acknowledgment vs announced prize for recovery	2-arm	Free-choice return; attribution coding	Announced ≥ unexpected on free-choice without control feelings → unexpected preference weaker
IM-QUAL	10 Maya interviews: “When a learning app gives points, what does that make you feel about the math?”	Appendix B	coded controlling vs informational	Nobody reports control → copy over-indexed

Pre-reg (XXXIV): IM-* identify **reward contingency** × **free-choice/interest/transfer**. They do **not** identify “gamification builds mathematicians” from XP velocity.

Densifies: TARGET appearance ban; SAFE-HABIT; SAFE-COMPARE; Insight on extrinsic traps; standing streak NS ban.

XLIV.10 Confidence table

Claim	Label	Confidence
Autonomy, competence, relatedness support intrinsic motivation / wellness	FACT (Ryan & Deci, 2000)	High
Expected engagement-/completion-contingent tangible rewards undermine free-choice IM on average	FACT (Deci, Koestner & Ryan, 1999)	High
Unexpected tangible rewards do not show average undermine	FACT (Deci et al., 1999)	High
Positive feedback enhances IM on average; controlling administration can reverse	FACT (Deci et al., 1999; Ryan line)	High
Tangible rewards more detrimental for children than college samples	FACT (Deci et al., 1999)	High
Cameron & Pierce find overall null with circumscribed negatives	FACT (1994); interpretation contested	Medium (dispute)
Choice enhances IM / effort / competence perceptions	FACT (Patall et al., 2008)	High
Stacking rewards after choice weakens choice benefit	FACT (Patall moderator)	Medium–High
SAFE-REWARD is right default for MindCraft	FOUNDER BELIEF / HYPOTHESIS	Medium–High
“All rewards kill intrinsic motivation” / “XP creates math love”	SPECULATION / false as universals	Low (against)

XLIV.11 What this chapter kills

1. **Kill:** Expected engagement-contingent XP/coins/loot as default motivation architecture for Maya. (Deci 1999 $d \approx -0.40$ free-choice.)
2. **Kill:** Marketing that “rewards create love of math” or “XP = identity.” (External PLOC; HID.)
3. **Kill:** Absolute slogan “all extrinsic rewards are bad” used to ban informational feedback and earned transfer unlocks. (CET differentiation; Cameron caveat acknowledged.)
4. **Kill:** Controlling praise / “you should keep earning” as the voice of competence support.
5. **Wound:** “Performance badges are always safe” — survives only with informational climate + real competence criterion.

6. **Survive (constrained):** Unexpected acknowledgment, strategy-informative feedback, autonomy-supportive micro-choice, and transfer-gated unlocks under **SAFE-REWARD**.

Doctrine until data: Ship **SAFE-REWARD**: no expected tokens for mere engagement → feedback over inducement → performance marks only for real transfer/recovery → autonomy framing → don't stack choice with dangling coins → child-risk premium → XP never North Star. Market “we don't pay kids to practice — we make the competence visible.”

Part XLV — Mathematical Resilience: Growth Under Adversity Without Grit Theater

Chapter status: Living evidence brief — Researcher tick 2026-07-30 (UTC hour 12 ≡ Red Team slot, but ch45 never written → prefer Researcher per rotation; 4 researcher entries since synthesizer v1.6 → not yet Synthesizer)

Primary question: What is *mathematical resilience* as an actionable construct (Johnston-Wilder & Lee), how does it differ from grit posters and empty “bounce back” copy, and what product stack keeps Maya in the growth zone without dumping her into the danger zone?

Owners: Motivation / Affect · Product (Practice / coach / tutor briefing) · Positioning · Red Team

Commercial job: Replace “build grit” / “resilience mindset” marketing with a shippable **SAFE-RESILIENCE** stack competitors can be audited against — growth-zone design, support recruitment, struggle-as-normal, value without coercion — tied to FEI / SAFE-DD / SAFE-REWARD.

XLV.1 Why this chapter exists

Parts XXIV (anxiety), XXV (self-efficacy / habit≠identity), XXX (attribution / helplessness), XXXI (flow / challenge-skill), XLI (SAFE-DD), and XLIV (SAFE-REWARD) already treat threat, recovery, and controlling incentives. Market and school language still collapses everything into **grit** or generic **resilience** slogans: “just push through,” “bounce back harder,” streak-as-character, leaderboard-as-grit-proof.

FOUNDER BELIEF under audit: Identity transformation under difficulty needs *domain-specific* resilience — knowing struggle is normal in math, valuing the work, knowing how to work at it, and **recruiting support** before anxiety flips the lid — not a personality transplant labeled grit.

Claims we refuse as doctrine:

1. Duckworth-style grit as MindCraft's North Star or primary product claim (Credé et al. wound the construct).
 2. “Desirable difficulty” = keep students in the danger/anxiety zone (conflicts SAFE-DD / Ashcraft).
 3. Empty resilience posters without support architecture (tutor cards, peers, ontology scaffolds).
 4. Treating mathematical resilience as synonymous with growth-mindset wall art.
-

XLV.2 Constructs (product language)

Construct	Plain definition	Product signal	Failure mode if misused
Developmental resilience (Masten)	Ordinary adaptive systems enabling competence despite adversity — not rare “super kids”	Design that protects basic systems (safety, belonging, competence feedback)	Selling “extraordinary character” as the product
Mathematical resilience (JW/Lee)	Positive stance that lets learners engage despite affective barriers: growth belief, value, knowing how to work at math, knowing how to recruit support	Stay in growth zone; ask for help; recover after stuck; continue study path	Posterizing the four aspects without UX
Growth Zone Model	Comfort (easy / little growth) → Growth (challenge + support) → Danger/Anxiety (threat / thinking collapses)	Challenge calibrated + support on tap; exit danger fast	Framing red-zone panic as “grit training”
Grit (Duckworth)	Perseverance + consistency of interest toward long-term goals	Tempting NS copy	Near-conscientiousness ; weak intervention leverage (Credé)
SAFE-RESILIENCE	Equip → name struggle → recruit support → destake danger → measure recovery/transfer — never grit KPI	See XLV.7	Resilience theater that still shames help-seeking

Operational definition (HYPOTHESIS): A MindCraft session builds mathematical resilience only when the student (a) experiences struggle as *expected* not as personal defect, (b) has a clear path to **support** (cards, tutor, peer, worked scaffold) before threat spikes, (c) returns to challenge after recovery (retry_120s / challenge_accept), and (d) can transfer without the model finishing the thought — not when streak or “grit score” rises.

XLV.3 Ordinary magic — resilience is not a unicorn trait

FACT: Masten (2001, *American Psychologist*, 56(3), 227–238, doi:10.1037/0003-066X.56.3.227) — resilience in development is often **ordinary magic**: competence despite adversity typically arises from normative adaptational systems (attachment, learning systems, self-regulation, community supports). Greatest threats are those that *compromise* these protective systems. Policy implication: protect and restore ordinary systems rather than hunt rare resilient personalities.

Product translation: MindCraft should not market “we forge resilient geniuses.” Market “we keep the ordinary protective systems online while math gets hard” — safety, relatedness (tutor/world), competence information,

recoverable miss. That is Masten-aligned and FEI-aligned.

Kill: “Resilience is a rare character trait we select for.”

XLV.4 Mathematical resilience as a pragmatic construct

FACT (origin framing): Johnston-Wilder & Lee (2010) introduced **mathematical resilience** as a pragmatic positive construct opposite avoidance, anxiety, and helplessness — enabling learners to engage with mathematics despite affective barriers (cited as foundational in later scale and handbook work; e.g. Lee & Johnston-Wilder construct summaries; Johnston-Wilder & Lee handbook chapter, doi:10.4324/9781315740713-34).

FACT (four aspects — construct literature): Lee & Johnston-Wilder summarize mathematical resilience as four pragmatic aspects:

1. **Growth** — belief that mathematical capability can develop with effort and appropriate support (growth/incremental theory of learning).
2. **Value** — mathematics has personal / world value and the learner is valued in the learning community.
3. **Knowing how to work at mathematics** — struggle is expected; strategies exist for stuckness.
4. **Knowing how to recruit support** — peers, adults, resources; stuck is not permanent solitude.

FACT (Growth Zone Model): Johnston-Wilder et al. (2013 onward; e.g. pre-service teacher applications, Warwick WRAP) — three zones:

- **Comfort (green):** routine; little growth.
- **Growth (amber/yellow):** challenge requiring support, persistence, agency; productive struggle.
- **Danger / anxiety (red):** challenge perceived beyond capability even with support; survival response; repeated exposure → persistent anxiety / avoidance.

HYPOTHESIS (product): MindCraft’s “soft-wrong + cards + tutor” stack is literally a **support-recruitment machine** for staying in the growth zone. Competitors that only escalate difficulty (or shame help) shove Maya into red.

Bridge to SAFE-DD (XLI): Desirable difficulties belong in the *growth* zone with equipment and destaking — not in the danger zone. “Struggle” ≠ “panic.”

XLV.5 Measurement — MRS is attitudes, not magic scores

FACT: Kooken, Welsh, McCoach, Johnston-Wilder & Lee (2016, *Measurement and Evaluation in Counseling and Development*, 49(3), 217–242, doi:10.1177/0748175615596782) — developed and validated the **Mathematical Resilience Scale (MRS)** via EFA/CFA across three samples. Three correlated factors: **Value**, **Struggle**, and **Growth**. Positioned as a way to gauge likelihood of participation and persistence in mathematics.

Product implication (HYPOTHESIS): MRS-style items (value / struggle-as-normal / growth belief) can inform parent research and longitudinal identity panels — but they are **not** a North Star and not a gamified in-app grit meter. Self-report attitudes ≠ transfer. Pair with `retry_120s`, `challenge_accept`, `transfer_pass`, help-recruitment events (opened card / asked tutor / peer).

Kill: Shipping a “Resilience Score™” badge as proof of identity change (L0 → L3 leap; Part XXXIV hygiene).

XLV.6 Grit theater — what we refuse to sell

FACT (origin): Duckworth, Peterson, Matthews & Kelly (2007, *Journal of Personality and Social Psychology*, 92(6), 1087–1101, doi:10.1037/0022-3514.92.6.1087) — grit as perseverance and passion for long-term goals; predictive claims in selective samples fueled popular education discourse.

FACT (meta wound): Credé, Tynan & Harms (2017, *Journal of Personality and Social Psychology*, 113(3), 492–511, doi:10.1037/pspp0000102) — 584 effects, 88 samples, $N \approx 66,807$: higher-order grit structure not confirmed; grit **very strongly** overlaps conscientiousness; overall grit–performance/retention relations **modest**; perseverance-of-effort facet carries more criterion validity than consistency-of-interest; **interventions to enhance grit may have only weak effects**; construct validity questioned; primary utility may lie in perseverance facet alone.

FOUNDER BELIEF / resolution: Perseverance *after equipped struggle* remains commercially meaningful (retry_{120s}). Packaging that as Duckworth “grit” invites overclaim, personality fatalism, and equity harm (blame the child’s character when curriculum/support failed). Prefer **mathematical resilience** language: growth zone + support recruitment + struggle-as-normal.

Standing Constitution already warned (core OS): grit/resilience hype is contested; structural inequality and curriculum quality matter more than character slogans. This chapter densifies that kill with Credé + JW/Lee specificity.

Kill: “MindCraft builds grit.” **Survive (constrained):** “MindCraft keeps you in the growth zone — challenge with support, recovery without shame.”

XLV.7 Competitive audit — who shoves students into the red zone?

Pattern	Zone reading	MindCraft wedge
Duolingo streak shame / league pressure	Danger + appearance climate (XLII/XLIII); “grit” via fear of break	Miss ≠ identity failure; cue practice not counter; no grit KPI
Khan mastery lock without affect design	Can strand in stuck without emotional exit tools	Soft-wrong + ingredient cards as support recruitment
Brilliant hard delight	Often growth-zone when curiosity high; weak for high-anxiety Maya	SAFE-DD equip/destake before delight escalation
Unguarded ChatGPT “just try harder / here’s the answer”	Either skips struggle (comfort via answer) or fluent wrong confidence	PWC + solo transfer; support = pedagogy wrap not answer dump
School “resilience assembly”	Poster without recruit-support architecture	Operationalize support buttons in-session

XLV.8 Resolution doctrine — SAFE-RESILIENCE

FOUNDER BELIEF / HYPOTHESIS — SAFE-RESILIENCE stack (ties FEI / SAFE-DD / TARGET / SAFE-REWARD / EVT value):

1. **Name the zones in product language.** Comfort / growth / danger — coach and tutor briefs use this, not “be gritty.”
2. **Equip before escalate.** SAFE-DD first: representations, faded examples, destaked attempts — then raise challenge.
3. **Struggle is expected, not character failure.** Attribution copy (XXX): strategy and effort-with-strategy; never “you’re not resilient enough.”
4. **Support recruitment is a first-class action.** Cards, ask-tutor, peer, worked scaffold — log as positive resilience behavior, not as “cheating the streak.”
5. **Exit the danger zone fast.** Anxiety spike → reduce stakes, shorten item, switch representation, offer support — do not double dose “desirable difficulty.”
6. **Value without coercion.** Eccles utility/attainment (XXXVII) and story worlds — not “suffer because grit.”
7. **Measure recovery and transfer, not grit scores.** `retry_120s`, `help-then-solo` sequences, `challenge_accept`, `transfer_pass` / `solo_transfer_pass`; optional MRS-like survey offline.
8. **Parent copy:** Sell “she got stuck, recruited help, and came back” — not “she’s a gritty kid now.”

Kill the false dichotomy: “Coddle forever in comfort” vs “forge character in the red zone.” Refuse both. Ship growth-zone challenge with support on tap.

XLV.9 Marketing language that survives

- **Allowed:** “Hard math with a way out — support when you’re stuck, challenge when you’re ready.”
- **Allowed:** “Stuck isn’t failure. Stuck is when you recruit the next tool.”
- **Allowed (parent):** “We train recovery and transfer — not grit slogans.”
- **Banned:** “We build grit.” / “Resilience score.” / “Push through the panic.”
- **Banned:** Bloom 2-sigma; empty mindset posters; streaks as proof of resilience; absolute tutoring-is-free.

Feature claim ladder (Part XXXIV hygiene)

Feature	Claim max without RES data
Growth-zone framing + support CTAs in coach	L1 hygiene aligned with JW/Lee
Soft-wrong + card recruitment → retry	L1 mechanism claim
“Raises mathematical resilience (MRS)” as product outcome	L0 until RES-1
“Builds grit / character”	L0 — banned packaging

XLV.10 Experiments spawned

ID	Question	Design	Primary	Kill condition
RES-1	Growth-zone coach copy + explicit support CTA vs grit/perseverance copy (same difficulty)	2-arm; anxiety-stratified	retry_120s; help-recruit rate; post affect; transfer_pass	Grit copy $\geq$ growth-zone on retry+transfer without raising danger affect $\rightarrow$ zone framing overstated
RES-2	Support-first after miss (card/tutor offer) vs no offer (same soft-wrong)	2-arm	Time-to-retry; danger-zone self-report; solo transfer later	No-offer $\geq$ support-first on retry without harm $\rightarrow$ recruitment less load-bearing
RES-3	Escalate difficulty after equip vs escalate without equip (SAFE-DD $\times$ RES)	2 $\times$ 2	Challenge_accept ; WM/anxiety proxy; transfer	Unequipped escalate wins transfer in high-anxiety $\rightarrow$ SAFE-DD sequencing weak
RES-4	MRS Value/Struggle/Growth Δ over 4 weeks of SAFE-RESILIENCE UX vs control	Survey + behavior	MRS factors + behavioral NS	MRS rises, behavior flat $\rightarrow$ attitude theater
RES-QUAL	10 Maya interviews: “When you’re stuck in math, what does ‘being resilient’ mean?”	Appendix B	coded grit-vs-support meanings	Everyone already means support recruitment $\rightarrow$ copy shift lower priority

Pre-reg (XXXIV): RES-* identify **zone framing** $\times$ **support recruitment** $\times$ **recovery/transfer**. They do **not** identify “we create resilient mathematicians” from MRS alone.

Densifies: SAFE-DD; FEI; TARGET; standing grit-hype ban; help-seeking as virtue not streak cheat.

XLV.11 Confidence table

Claim	Label	Confidence
Developmental resilience often arises from ordinary protective systems	FACT (Masten, 2001)	High

Mathematical resilience = growth + value + how-to-work + recruit support	FACT (JW/Lee construct line)	High (construct); Medium (intervention effect sizes sparse)
Growth Zone Model: comfort / growth / danger	FACT (JW et al. model literature)	High as pedagogy model; Medium as RCT-proven UX
MRS measures Value, Struggle, Growth attitudes	FACT (Kookken et al., 2016)	High
Grit higher-order structure weak; overlaps conscientiousness; modest prediction; weak intervention leverage	FACT (Credé et al., 2017)	High
Danger-zone exposure as “training” harms anxious learners	HYPOTHESIS (Ashcraft line + SAFE-DD)	Medium–High
SAFE-RESILIENCE is right default for MindCraft	FOUNDER BELIEF / HYPOTHESIS	Medium–High
“Grit posters raise transfer” / “Resilience Score™ = identity”	SPECULATION / false as doctrine	Low (against)

XLV.12 What this chapter kills

1. **Kill:** Grit / “build character” as MindCraft marketing North Star. (Credé 2017; Constitution grit warning.)
2. **Kill:** Treating panic / danger-zone as desirable difficulty. (Growth Zone Model + SAFE-DD.)
3. **Kill:** Resilience as rare trait branding (“extraordinary kids”). (Masten ordinary magic.)
4. **Kill:** In-app Resilience Score™ as proof of identity transformation. (MRS ≠ transfer; claim ladder.)
5. **Wound:** Growth-mindset language alone — survives only inside growth + value + struggle know-how + **support recruitment**, not as wall art.
6. **Survive (constrained):** Growth-zone challenge with support on tap; struggle-as-normal; recovery metrics; parent story of “stuck → help → return.”

Doctrine until data: Ship **SAFE-RESILIENCE**: name zones → equip before escalate → struggle expected → support recruitment first-class → exit danger fast → value without coercion → measure recovery/transfer not grit. Market “hard math with a way out.”

Part XLVI — Teacher/Tutor Expectancy: Pygmalion Without Magic, Golem Without Labels

Chapter status: Living evidence brief — Researcher tick 2026-07-30 (UTC hour 15; hour%6≠0 → Researcher; 5 researcher entries since synthesizer v1.6 → not yet Synthesizer)

Primary question: What do teacher/tutor expectancy effects actually justify for MindCraft — and what briefing ethics keep college tutors from installing a Golem on Maya?

Owners: Tutor ops / Product (briefing cards) · Equity · Positioning · Red Team

Commercial job: Replace “we believe in every student” poster copy and “Pygmalion IQ” folklore with a shippable **SAFE-EXPECTANCY** stack — evidence-based briefs, class-level high challenge + support, ban stigmatizing labels, measure mediation behaviors not vibes.

XLVI.1 Why this chapter exists

Parts XXV (self-efficacy), XXX (attribution), XXXVI (equity), XXXVII (EVT expectancy), and XLV (support recruitment) already treat *student-side* “can I?” and belonging. They do not fully answer: **what beliefs do human tutors carry into the session, how do those beliefs leak into climate/input/output/feedback, and how should MindCraft’s pre-session briefing be designed so gap-scan weakness does not become a self-fulfilling low label?**

FOUNDER BELIEF under audit: A knowledge-graph + tutor hybrid only compounds identity if the *human* in the loop treats Maya as capable of the next hard step with support — not as a fixed “weak student” whose low mastery scores license colder climate, thinner input, and fewer chances to respond.

Claims we refuse as doctrine:

1. Marketing the classic Pygmalion IQ story as MindCraft’s differentiator (“our tutors raise IQ by believing harder”).
 2. Briefing tutors with trait labels (“low ability,” “slow,” “not STEM material”) drawn from gap-scan or weakness ranking.
 3. Treating teacher-expectancy literature as proof that positive vibes alone move transfer.
 4. Ignoring Golem / differential-treatment risks for stigmatized or “weakness-tagged” students.
-

XLVI.2 Constructs (product language)

Construct	Plain definition	Product signal	Failure mode if misused
Induced expectancy (Pygmalion paradigm)	Authority given false “bloomers” info → treats them differently → outcomes shift	Early-session / cold-start briefing window	Selling IQ magic; ignoring replication limits
Accuracy vs prophecy	Expectations often predict because they track real prior performance	Gap-scan / graph as <i>task</i> evidence, not character	Confusing prediction with causation
Golem effect	Low expectations → colder / more dogmatic treatment → worse performance	Ban negative trait labels in tutor cards	“Honest” weakness dump that licenses withdrawal

Four-factor mediation	Climate, input, output opportunity, feedback carry expectancy into behavior	Audit tutor/coach on those four channels	Soft climate without hard input/output
High-expectation teacher (class-level)	High challenge + high support for <i>all</i> , not only nominated stars	Tutor playbook: everyone gets hard Qs + scaffolds	Favorites / “bloomers only” tracks
SAFE-EXPECTANCY	Evidence brief → high challenge-with-support → no trait labels → mediate via CIOF → measure behaviors	See XLVI.7	Belief posters; Golem briefs

Operational definition (HYPOTHESIS): A MindCraft tutor session is expectancy-*safe* only when (a) the brief states **next playable tasks and scaffolds**, not trait judgments, (b) the tutor offers **hard-enough** work with support on tap (XLV growth zone), (c) climate stays warm under struggle, and (d) the student gets **output opportunity** (speak/solve before tutor finishes) — not when the tutor “believes harder” while finishing the math.

XLVI.3 Origin story — real, contested, not marketing copy

FACT: Rosenthal & Jacobson (1968, *Pygmalion in the Classroom*; also *Urban Review* summary, doi:10.1007/BF02322211) — teachers told that randomly selected elementary students were poised for intellectual blooming; those students later showed larger IQ gains on average in the reported analyses, especially in early grades. The study launched classroom expectancy research and intense methodological critique (e.g. Thorndike and later replication debates).

FACT (meta, induction credibility): Raudenbush (1984, *Journal of Educational Psychology*, 76(1), 85–97, doi:10.1037/0022-0663.76.1.85) — synthesis of 18 Pygmalion-style experiments: expectancy effects on pupil IQ are larger when teachers have had **little prior contact** with pupils at induction (≤ 2 weeks); effects shrink when teachers already know students well. Grade pattern: stronger early-grade effects in the corpus; design features of induction matter.

Product translation: MindCraft’s **first tutor sessions and cold-start briefs** are the high-leverage expectancy window — analogous to early-year induction. Once a tutor has rich behavioral evidence, “random bloomers” folklore is the wrong metaphor; **accurate, task-level** expectations dominate.

Kill: “Pygmalion proves we raise IQ by believing.” Sell mediation design in the first hours of contact, not IQ alchemy.

XLVI.4 What 35 years actually say (Jussim & Harber)

FACT (review): Jussim & Harber (2005, *Personality and Social Psychology Review*, 9(2), 131–155, doi:10.1207/s15327957pspr0902_3) — after ~35 years:

1. Self-fulfilling prophecies in classrooms **do occur**, but effects are **typically small**, do not accumulate greatly across perceivers/time, and may dissipate more than snowball.

2. **Powerful** effects may selectively hit students from **stigmatized** social groups.
3. Whether expectancy effects reliably change *intelligence*, and whether they do more harm than good in general, remains unclear.
4. Teacher expectations often predict outcomes **because they are accurate**, not only because they are self-fulfilling.

HYPOTHESIS for MindCraft: Gap-scan / mastery graphs will *accurately* predict struggle. Accuracy is not permission to install a Golem. The commercial risk is **translating accurate weakness into low-challenge, low-warmth treatment** — exactly the path that can still produce small-but-real prophecy costs, amplified for already-stereotyped students (ties Part XXXVI equity).

Wound: “Expectancy effects are huge and ubiquitous.” Survives only as: small average effects; context/stigma moderators; accuracy dominates prediction.

XLVI.5 Mediation — climate, input, output, feedback

FACT (behavioral mediation): Brophy & Good (1970, *Journal of Educational Psychology*, doi:10.1037/h0029908) — observational study: teachers demanded better performance and praised it more for high-expectation students; accepted poor performance more readily from lows and were less likely to praise good performance from lows when it occurred.

FACT (four-factor meta): Harris & Rosenthal (1985, *Psychological Bulletin*, 97(3), 363–386, doi:10.1037/0033-2909.97.3.363) — 31 meta-analyses of mediating behaviors; Rosenthal’s four-factor frame (**climate, feedback, input, output/opportunity**) organizes how expectancies become treatment. Expectancy→behavior and behavior→outcome links support multiple mediators (e.g. warmer climate, longer/richer interaction, questioning). Later summaries often note **feedback** as the weaker of the four average links, pushing affect/effort emphasis.

Product checklist (HYPOTHESIS — CIOF audit):

Channel	Tutor/coach do	Anti-pattern
Climate	Warmth under struggle; curiosity about wrong paths	Cold distance after “weak” brief
Input	Richer explanations, multiple representations (cards)	Thin “just try” for lows; lectures only for highs
Output	Wait time; let Maya finish; SE prompts (Part XL)	Tutor finishes every thought (Solver crutch)
Feedback	Criterion + strategy (XXX); soft-wrong	Praise only highs; criticism without how-to

Kill: Climate-only positivity (“you’ve got this!”) without input/output redesign — that is poster Pygmalion.

XLVI.6 Golem risk — negative expectancy is the ops hazard

FACT: Babad, Inbar & Rosenthal (1982, *Journal of Educational Psychology*, 74(4), 459–474, doi:10.1037/0022-0663.74.4.459) — distinguished positive expectancy outcomes (“Galatea”) from **negative** ones (“Golem”). Teachers high in susceptibility to biasing information treated low-potential nominees more negatively (notably dogmatic behavior); unbiased teachers treated groups more equitably. Differential negative treatment showed up in teacher behavior **and** student task performance.

FACT (student-perceived differentiation): Brattesani, Weinstein & Marshall (1984, *Journal of Educational Psychology*, 76(2), 236–247, doi:10.1037/0022-0663.76.2.236) — in classrooms where students perceive **high differential treatment**, teacher expectations contribute more to achievement than in low-PDT classrooms. Students can see the track.

Product translation: A tutor card that says “Maya is weak on linear equations / low confidence / hard gap-scan” is biasing information. High-bias tutors will leak Golem through climate and dogmatism. Low-PDT design = same high standards + differentiated *scaffolds*, not differentiated *respect*.

Kill: “Honest labels help tutors calibrate.” Honest **task evidence** helps; honest **trait verdicts** harm.

XLVI.7 High-expectation teachers — class-level, not favorites

FACT: Rubie-Davies (2007, *British Journal of Educational Psychology*, doi:10.1348/000709906X101601) — contrasts high- vs low-expectation *teachers* (expectations for the class as a whole). High-expectation teachers spent more time teaching, oriented lessons, linked prior knowledge, explained carefully, asked mixed open/closed questions including **high-level questions of all students**, and managed more positively; lows were more dismissive of incorrect answers and relied on closed questioning.

HYPOTHESIS: MindCraft’s competitive wedge vs “AI that flatters then finishes” and vs “tutor who only drills the labeled weak topic gently” is **Rubie-Davies-shaped ops**: every student gets high-level questioning + scaffolds; incorrect answers get learning support, not dismissal.

Bridge to SAFE-RESILIENCE / SAFE-DD: High expectation ≠ danger zone. High expectation = growth-zone challenge with support recruitment — not grit theater and not thin input for “lows.”

XLVI.8 Competitive audit (expectancy lens)

Competitor pattern	Expectancy mechanism	MindCraft counter
Duo	Habit/streak as “you’re the type who shows up”; little academic expectancy mediation	Transfer + CIOF; no streak-as-belief
Khan	Self-serve library; weak human expectancy channel	Human tutor brief + graph without Golem labels
Brilliant	Prestige challenge; can cue performance-approach for novices	Mastery structures + expectancy-managed gates

ChatGPT Base	Sycophantic high-warmth climate + answer dump (fake high expectancy, zero output demand)	Warm climate plus forced solo output / PWC
Unbriefed college tutor	Forms expectancies from accent, race, “seems lost,” or raw weakness dump	SAFE-EXPECTANCY brief + playbook

XLVI.9 Resolution doctrine — SAFE-EXPECTANCY

FOUNDER BELIEF / HYPOTHESIS — SAFE-EXPECTANCY stack:

- 1. Brief tasks, not traits.** Tutor card: current concept, recommended level, formats struggled, successful scaffolds last time, next playable question IDs — **never** “low ability / not a math person.”
- 2. Protect the cold-start window.** First sessions: explicit high-challenge-with-support script (Raudenbush induction logic).
- 3. Class-level high expectation.** Rubie-Davies: hard questions for everyone; wrong answers → teach, don’t dismiss.
- 4. Mediate via CIOF.** Climate warmth + rich input (cards) + output opportunity (wait / SE) + strategy feedback — audit these, not “belief intensity.”
- 5. Accuracy without Golem.** Graph weakness drives *scaffold choice*, not *respect withdrawal*.
- 6. Stigma premium.** Extra review when briefs could activate stereotype (XXXVI); ban demographic proxies in tutor UI.
- 7. AI coach inherits the same rules.** No “you’re behind your peers” cold open; no finishing her thought under “support.”
- 8. Measure mediation + FEI, not IQ folklore.** Tutor wait-time, student talk ratio, help-then-solo, `retry_120s`, `transfer_pass` / `solo_transfer_pass`.

Kill the false dichotomy: “Ignore data, just believe” vs “dump the weakness label so tutors are realistic.” Refuse both. Ship **evidence-based high challenge with equitable treatment**.

XLVI.10 Marketing language that survives

- **Allowed:** “Tutors see what to teach next — not who she is forever.”
- **Allowed:** “High challenge for everyone, support on tap — no tracks of respect.”
- **Allowed (parent):** “We brief tutors on the next skill, not on a fixed ability story.”
- **Banned:** “Pygmalion IQ gains.” / “Believe and they bloom.” / “We tell tutors who’s weak so they know.”
- **Banned:** Bloom 2-sigma; empty mindset posters; streaks as expectancy proof; tutoring-is-free.

Feature claim ladder (Part XXXIV hygiene)

Feature	Claim max without EXP data
---------	----------------------------

Trait-label ban + task-only tutor brief	L1 hygiene aligned with Golem literature
CIOF playbook / wait-time coaching	L1 mechanism claim
“Raises IQ via tutor expectancy”	L0 — banned
“Eliminates achievement gaps via Pygmalion”	L0 until EXP-* + equity design (Jussim caveats)

XLVI.11 Experiments spawned

ID	Question	Design	Primary	Kill condition
EXP-1	Task-only brief vs trait/weakness-label brief (same content otherwise)	2-arm tutors; record session	Student talk ratio; wait-time; post affect; transfer_pass	Label brief ≥ task-only on transfer without colder climate → Golem risk overstated for our tutors
EXP-2	High-expectation playbook (hard Q + scaffold) vs “remediate gently only” after same gap-scan	2-arm	challenge_accept; solo transfer; help-recruit	Gentle-only wins transfer+return in high-anxiety → Rubie-Davies transfer fails for Maya
EXP-3	Cold-start script (first 2 sessions) vs unstructured tutor open	2-arm new matches	FEI events week 1; retention to session 3	Unstructured ≥ scripted → induction window less load-bearing
EXP-4	AI coach: warm+finish vs warm+output-demand (SE/wait)	2-arm	solo_transfer_pass; PWC	Warm+finish ≥ warm+output on solo transfer → climate-only expectancy myth revived
EXP-QUAL	10 tutor + 10 Maya: “What did the brief make you assume?”	Interview	coded trait vs task inferences	Tutors already ignore labels → UI priority shifts

Pre-reg (XXXIV): EXP-* identify **brief content** × **CIOF mediation** × **transfer**. They do **not** identify “we create bloomers” or IQ change.

Densifies: Equity (XXXVI); EVT expectancy; SAFE-RESILIENCE support; FEI; tutor ops queue item 68.

XLVI.12 Confidence table

Claim	Label	Confidence
Rosenthal & Jacobson launched expectancy research; original IQ story is contested	FACT	High
Induced expectancy effects larger when teacher knows student less (Raudenbush, 1984)	FACT	High
Average self-fulfilling effects small; accuracy often drives prediction; stigma can amplify (Jussim & Harber, 2005)	FACT	High
Differential treatment observable; highs get more demand/praise (Brophy & Good, 1970)	FACT	High
Four-factor mediation frame useful (Harris & Rosenthal, 1985)	FACT	High (framework); Medium (which channel dominates in tutoring apps)
Golem / high-bias negative treatment real (Babad et al., 1982)	FACT	High
High-expectation teachers differ in practice for whole class (Rubie-Davies, 2007)	FACT	High (observational/descriptive line); Medium (causal ops transfer)
SAFE-EXPECTANCY is right default for MindCraft tutor briefing	FOUNDER BELIEF / HYPOTHESIS	Medium–High
“Pygmalion IQ” / belief posters as product claim	SPECULATION / false as doctrine	Low (against)

XLVI.13 What this chapter kills

1. **Kill:** Marketing classic Pygmalion IQ blooming as MindCraft proof. (Raudenbush limits; Jussim accuracy; claim ladder.)
2. **Kill:** Tutor briefs that install trait/weakness identity labels. (Babad Golem; Weinstein PDT.)
3. **Kill:** Climate-only positivity without input/output demand. (Harris & Rosenthal mediation; Bastani/PWC line.)
4. **Kill:** “Just be realistic — tell tutors she’s weak” as ops hygiene.
5. **Wound:** Expectancy as a large, cumulative social-problem engine — survives only as small average effects with stigma/PDT moderators.
6. **Survive (constrained):** Cold-start high-challenge-with-support; task-only briefs; CIOF audit; Rubie-Davies whole-class high expectation; measure mediation + transfer.

Doctrine until data: Ship **SAFE-EXPECTANCY**: brief the next skill, not the fixed self; high challenge for everyone with scaffolds; warm climate under struggle; force student output; never weaponize gap-scan as a Golem tag. Market “tutors see what to teach next — not who she is forever.”

Part XLVII — Sleep, Stress, and Learning: Timing Practice Without All-Nighter Theater

Chapter status: Living evidence brief — Researcher tick 2026-07-30 (UTC hour 18 ≡ Red Team slot, but ch47 never written → prefer Researcher per rotation; 6 researcher entries since synthesizer v1.6 → not yet Synthesizer)

Primary question: How should MindCraft time practice and exam-week product behavior given sleep-dependent consolidation and stress/pressure drain on working memory — without becoming a sleep-coaching app or glorifying late-night grind?

Owners: Product (session timing / exam modes) · Parent messaging · Affect pipeline · Red Team

Commercial job: Ship a **SAFE-TIMING** stack — overnight spacing as consolidation infrastructure, exam-week load softening, ban streak/all-nighter heroism, measure delayed transfer not midnight minutes.

XLVII.1 Why this chapter exists

Parts XXIX (spacing/retrieval), XXIV/XLI (anxiety × desirable difficulty), XXXI (flow/challenge–skill), and XLV (growth vs danger zone) already treat *within-session* difficulty and affect. They do not fully answer: **when** should Maya practice relative to sleep, and what should the product do in high-stakes stress weeks when encoding, attention, and WM are compromised?

FOUNDER BELIEF under audit: Identity transformation requires recoverable struggle *and* biological conditions that let struggle stick. A product that pushes “one more mission” at 1 a.m. during exam week can look engaged while sabotaging consolidation and installing danger-zone affect.

Claims we refuse as doctrine:

- 1. Marketing MindCraft as a sleep-hygiene / wellness brand (“we fix teen sleep”).
- 2. Celebrating all-nighters, midnight streaks, or “grind hours” as learning proof.
- 3. Treating acute exam pressure as desirable difficulty.
- 4. Absolute claims that “sleep always consolidates math” independent of encoding quality, task type, and sleep opportunity.

XLVII.2 Constructs (product language)

Construct	Plain definition	Product signal	Failure mode if misused
-----------	------------------	----------------	-------------------------

Sleep-dependent consolidation	Offline reprocessing after encoding that stabilizes/reshapes memory	Space practice across nights; delay transfer checks	Sell “sleep = magic” without encoding first
Encoding under restriction	Multi-night short sleep impairs <i>new</i> learning capacity	Soften new-concept loads after chronic short sleep flags	Blame student “laziness” for restriction biology
Pressure / choke	High-stakes appraisal consumes WM needed for hard math	Destake practice; hide peer ranks in exam week	Treat choke as low ability
Math-anxiety dual-task	Worry acts like a secondary load on WM	SAFE-DD before hard items	Add timers/leaderboards under anxiety
SAFE-TIMING	Time practice for consolidation + protect WM under stress	See XLVII.8	Sleep-app feature creep; grind theater

Operational definition (HYPOTHESIS): A MindCraft practice plan is *timing-safe* when (a) hard encoding sessions are followed by a real sleep opportunity before high-stakes transfer claims, (b) multi-night sleep debt / high stress triggers **load softening** (fewer new concepts, more retrieval of known, support on tap), and (c) engagement metrics never reward post-midnight volume as mastery.

XLVII.3 Sleep consolidates — and evolves — memory

FACT (systems review): Diekelmann & Born (2010, *Nature Reviews Neuroscience*, 11, 114–126, doi:10.1038/nrn2762) — sleep optimizes consolidation depending on learning conditions and sleep timing; slow-wave sleep (SWS) and REM support complementary system- vs synaptic-consolidation roles; consolidation can change representations quantitatively *and* qualitatively.

FACT (broader offline processing): Stickgold (2005, *Nature*, 437, 1272–1278, doi:10.1038/nature04286) — converging evidence that offline memory reprocessing during sleep is a major component of how memories are formed and shaped (not a single all-or-none event).

FACT (integrative/schema angle): Walker & Stickgold (2010, *Nature Reviews Neuroscience*, 11, 218, doi:10.1038/nrn2762-c1) — beyond veridical preservation, sleep (often REM-linked in their framing) supports unitization, assimilation into remote networks, and abstraction of rules/schemas — relevant to MindCraft’s transfer goal, not only verbatim item memory.

FACT (mechanism exemplar): Gais & Born (2004, *PNAS*, 101(7), 2140–2144, doi:10.1073/pnas.0305404101) — elevating acetylcholine during SWS-rich sleep blocked declarative word-pair consolidation while sparing a nondeclarative skill task — evidence that the *neurochemical milieu* of SWS matters for declarative consolidation, not “rest” generically.

Product translation: Spacing missions across nights (Part XXIX) is not only spacing theory — it is a **consolidation window**. Claiming “exam-ready” from a same-day blocked accuracy binge without overnight delay is scientifically lazy (ties XXXIX readiness signal).

Kill: “Sleep replaces practice.” Consolidation operates on what was encoded. No encoding → nothing to consolidate.

XLVII.4 Adolescent sleep restriction is causal for learning damage

Maya is a high-schooler. Population sleep debt is not a vibe; experimental restriction shows classroom-relevant harm.

FACT (simulated classroom): Beebe, Rose & Amin (2010, *Journal of Adolescent Health*, 47(5), 523–525, doi:10.1016/j.jadohealth.2010.03.005) — within-subject multi-night restriction vs healthy duration: sleep-restricted adolescents showed poorer learning, more inattentive behaviors, and lower arousal in a simulated classroom (medium-to-large effects; small *n*, preliminary but experimental).

FACT (review): Beebe (2011, *Pediatric Clinics of North America*, doi:10.1016/j.pcl.2011.03.002) — synthesizes cognitive, behavioral, and functional consequences of inadequate sleep in children/adolescents; chronic restriction is common and not benign.

FACT (long-term factual retention): Cousins, Wong & Chee (2019, *Journal of Adolescent Health*, 65(4), 549–557, doi:10.1016/j.jadohealth.2019.04.030) — adolescents (15–18) with 5 h sleep opportunity for 5 nights vs 9 h controls: after four restriction nights, learning detailed factual material showed **more forgetting** at 30 min (~26%), Day 3 (~34%), and Day 42 (~65%) relative to controls. Vigilance also fell; memory impairment did not simply track vigilance correlation.

HYPOTHESIS for MindCraft: A “productive” ACT cram week of late missions on short sleep can **look** like engagement while destroying the encoding that gap-scan and practice are supposed to build. Parent-visible “minutes practiced” without sleep context is a vanity metric.

Wound: “Weekend catch-up fixes the week.” Literature flags incomplete restitution and cumulative cost (see Beebe 2019-line summaries); do not market weekend binge as full repair.

XLVII.5 Stress and pressure — WM theft, not character failure

Sleep loss and evaluative stress often co-occur in exam weeks. Separate the mechanisms.

FACT (math anxiety × WM): Ashcraft & Krause (2007, *Psychonomic Bulletin & Review*, 14, 243–248, doi:10.3758/BF03194059) — math performance depends on WM for beyond-retrieval arithmetic; high math anxiety behaves like a dual-task (worry consumes capacity); anxious students also avoid math pathways. Complements Ashcraft & Kirk (2001) line already in Part XXIV.

FACT (choking under pressure): Beilock & Carr (2005, *Psychological Science*, 16(2), 101–105, doi:10.1111/j.0956-7976.2005.00789.x) — performance pressure harmed **high**-WM individuals on the most WM-demanding math problems — pressure consumes the capacity their superior strategies rely on. Not “dumb under stress”; often “capable under load.”

Bridge to SAFE-DD / FEI: Exam-week timers, public ranks (XLII), and “prove you’re ready tonight” framing are pressure injectors. Desirable difficulty assumes encoding capacity; pressure + anxiety can erase the dose.

HYPOTHESIS: Affective check-in stress > 0.7 already softens target_mastery in /recommend — that is the correct direction. Extend it: under high stress / self-reported short sleep, prefer retrieval + support recruitment over new-concept flooding and appearance UX.

Kill: “Pressure makes diamonds” as practice design. Pressure makes WM poorer for the problems that need it most.

XLVII.6 Competitive audit (timing lens)

Competitor pattern	Timing / stress mechanism	MindCraft counter
Duo	Streaks + late-night nagging; habit over consolidation	SAFE-HABIT + SAFE-TIMING: miss≠identity; no midnight streak heroism
Khan	Infinite library; learner can cram to exhaustion	Spaced missions + delayed transfer as readiness
Brilliant	Delight / prestige sessions; can invite performance pressure	Mastery climate; destake under stress
ChatGPT Base	Always-on 2 a.m. answers; no consolidation pedagogy	Human/AI wrap that schedules <i>struggle then sleep</i> , not answer dumps
Tutor shops	Pack hours before exam Saturday	Brief tutors: encode early in week; retrieval near exam; never shame sleep

XLVII.7 Exam weeks — special product mode, not harder grind

SPECULATION → **product hypothesis:** Default “more missions” during ACT week maximizes anxiety and restriction. Safer mode:

1. **Front-load encoding** earlier in the week when sleep opportunity is better.
2. **Shift to retrieval + mixed review** 48–72 h pre-exam (XXIX/XXXIX).
3. **Lower appearance cues** (no leaderboards, no streak threats).
4. **Keep growth-zone items** but exit danger fast (XLV); prefer scaffolds over new ontology leaps.
5. **Parent copy:** sleep is part of the study plan — not “rest is for quitters.”

FOUNDER BELIEF: Selling calm competence under exam load beats selling heroic sleep debt.

XLVII.8 Resolution doctrine — SAFE-TIMING

FOUNDER BELIEF / HYPOTHESIS — **SAFE-TIMING stack:**

1. **Encode** → **sleep** → **transfer claim**. Treat overnight delay as part of readiness (with XXIX spacing), not optional wellness.
2. **Never reward midnight volume**. Ban streak/XP incentives that spike after local quiet hours; engagement ≠ consolidation.
3. **Restriction-aware load**. If multi-night short sleep or high stress is signaled (check-in, parent note, exam-week flag): fewer new concepts, more retrieval, CIOF support (XLVI), SAFE-DD destake.
4. **Pressure** ≠ **DD**. Timers, ranks, and “prove it tonight” are kill switches under high anxiety/stress.
5. **Exam-week mode**. Soften novelty; maximize recoverable retrieval; protect sleep opportunity messaging.
6. **Do not become a sleep clinic**. Nudge timing and load — do not diagnose disorders or sell melatonin theater.
7. **Tutor briefs include timing**. “Practice early evening when possible; protect sleep before transfer checks” — task ops, not character judgment.
8. **Measure delayed** transfer_pass / solo_transfer_pass, not minutes-after-midnight.

Kill the false dichotomy: “Ignore biology, just grind” vs “pivot the company into sleep coaching.” Refuse both. Ship **timing-aware learning load**.

XLVII.9 Marketing language that survives

- **Allowed:** “Hard practice sticks better when she sleeps on it.”
- **Allowed:** “We don’t measure mastery at 1 a.m.”
- **Allowed (parent):** “Exam week ≠ all-nighter week — we shift to review and protect sleep.”
- **Banned:** “Our app fixes teen sleep.” / “Grind while they sleep.” / Streak-as-study-ethic.
- **Banned:** Bloom 2-sigma; empty mindset; tutoring-is-free; pressure-as-pedagogy.

Feature claim ladder (Part XXXIV hygiene)

Feature	Claim max without TIM-* data
Overnight spacing before transfer checks	L1 mechanism aligned with consolidation reviews
Exam-week load softening + destake	L1 product hygiene under Ashcraft/Beilock
“Sleep mode raises ACT scores”	L0–L1 only with pre-reg; not L3 identity
Sleep-clinic / wellness brand claims	L0 — banned

XLVII.10 Experiments spawned

ID	Question	Design	Primary	Kill condition
----	----------	--------	---------	----------------

TIM-1	Same content: practice ending ≥ 1 h before habitual bedtime vs late-night block; transfer next evening	2-arm	next-day / 48h transfer_pass	Late $\geq$ early $\rightarrow$ consolidation timing overstated for our items
TIM-2	Exam-week mode (retrieval+ destake) vs “more new missions” in last 5 days pre-mock	2-arm	mock score; affect; sleep opportunity self-report	More-new wins mock+return $\rightarrow$ softening wrong default
TIM-3	High-stress check-in $\rightarrow$ auto soft path vs ignore stress	2-arm	retry_120s; WM-demand item accuracy; return	Ignore $\geq$ soft on transfer without affect cost $\rightarrow$ modifier useless
TIM-4	Parent message: “sleep is study” vs “push one more hour”	message A/B	next-night session timing; parent WTP/trust	Push-hour increases volume <i>and</i> next-day transfer $\rightarrow$ grind narrative survives
TIM-QUAL	10 Maya exam-week diaries: when practice helped vs harmed	interview	coded sleep debt $\times$ danger zone $\times$ cram	Students already self-time well $\rightarrow$ UI priority drops

Pre-reg (XXXIV): TIM-* identify **timing $\times$ load $\times$ delayed transfer**. They do **not** identify clinical sleep treatment or ACT score miracles from a bedtime tip.

Densifies: XXIX spacing; XLI SAFE-DD; XLII compare ban; XLIII habit without colonization; XLV danger exit; affect check-in; Part LXIII exam-pressure queue.

XLVII.11 Confidence table

Claim	Label	Confidence
Sleep supports declarative/procedural consolidation with stage-sensitive mechanisms (Diekelmann & Born, 2010; Stickgold, 2005)	FACT	High

Sleep can support schema/abstraction beyond verbatim storage (Walker & Stickgold, 2010)	FACT	Medium–High (active research debate on REM vs SWS mapping)
Low ACh during SWS matters for some declarative consolidation (Gais & Born, 2004)	FACT	High (specific manipulation); Medium (product generality)
Multi-night adolescent restriction impairs classroom learning/attention (Beebe et al., 2010; Beebe, 2011)	FACT	High (direction); Medium (effect sizes / sample limits)
Restriction impairs long-term factual retention in teens (Cousins et al., 2019)	FACT	High
Math anxiety consumes WM (Ashcraft & Krause, 2007)	FACT	High
Pressure can choke high-WM math performance (Beilock & Carr, 2005)	FACT	High
SAFE-TIMING is right default for MindCraft scheduling/exam mode	FOUNDER BELIEF / HYPOTHESIS	Medium–High
MindCraft-as-sleep-app is a good growth wedge	SPECULATION / false as doctrine	Low (against)

XLVII.12 What this chapter kills

1. **Kill:** All-nighter / midnight-streak heroism as learning culture. (Cousins retention; Beebe classroom; SAFE-HABIT.)
2. **Kill:** Treating exam pressure and timers as desirable difficulty. (Beilock choke; Ashcraft dual-task; SAFE-DD.)
3. **Kill:** Same-day cram accuracy as exam-ready proof without overnight/delayed transfer. (Diekelmann/Born; XXIX/XXXIX.)
4. **Kill:** Pivoting brand into clinical sleep wellness.
5. **Wound:** “Sleep always helps every math skill equally.” Survives only as: condition- and task-dependent consolidation; encoding first.
6. **Survive (constrained):** Encode→sleep→claim; restriction/stress-aware load; exam-week retrieval mode; parent “sleep is study”; measure delayed transfer.

Doctrine until data: Ship **SAFE-TIMING:** practice hard enough to encode, sleep enough to consolidate, soften novelty under debt and pressure, never sell grind-at-1-a.m. as identity. Market “hard practice sticks better when she sleeps on it.”

Part XLVIII — Transfer & Variation Theory: Far Transfer Design in the Ontology

Chapter status: Living evidence brief — Researcher tick 2026-07-30 (UTC hour 21; hour%6≠0 → Researcher; 7 researcher entries since synthesizer v1.6 → not yet Synthesizer; this tick becomes the 8th)

Primary question: How should MindCraft design question sequences and “transfer” claims so Maya discerns *critical mathematical structure* — without selling story-world immersion or blocked accuracy as far transfer?

Owners: Product (questionBank / ontology / pathfinder) · Content ops · Marketing claim ladder · Red Team

Commercial job: Ship a **SAFE-TRANSFER** stack — deliberate variation/invariance for discernment, hugged near transfer toward ACT formats, bridged high-road abstraction only where designed, measure delayed/format/solo transfer — kill “narrative = transfer for free.”

XLVIII.1 Why this chapter exists

Parts XXIX (spacing/retrieval), XXXIX (interleaving), XL (self-explanation), XLI (SAFE-DD), and XLVII (encode→sleep→claim) already constrain *when* and *how hard* practice should feel. They do not fully answer: **what pattern of examples** makes a concept transferable, and what marketing language about “transfer” survives evidence?

MindCraft’s ontology already encodes ingredients, bridges, combinations, formats, and archetypes. The risk is product language that treats any format change, story skin, or mixed set as “far transfer” — while competitors sell fluency theater as readiness.

FOUNDER BELIEF under audit: Identity as a mathematical thinker requires *flexible discernment* of structure across contexts — not one beautiful tutorial problem solved once. That flexibility is **designed**, not an automatic side-effect of lore.

Claims we refuse as doctrine:

- 1. Story worlds automatically produce far transfer.
- 2. Blocked same-day accuracy ≡ exam-ready / transferable skill.
- 3. “We teach for transfer” without specifying content×context dimensions (Barnett & Ceci).
- 4. Random interleaving or random story skins as a substitute for *principled* variation of critical aspects.

XLVIII.2 Constructs (product language)

Construct	Plain definition	Product signal	Failure mode if misused
Near / far transfer	Application of learning across some distance in content and/or context	Label every transfer claim with <i>which</i> dimensions moved	Vague “far transfer” marketing

Low-road / high-road	Automatic similarity trigger vs mindful abstraction	Hugging (ACT-like items) vs bridging (explicit compare/abstract prompts)	Expect far transfer from low-road drill alone
Object of learning	What is to be discerned (content + critical aspects)	Ingredient / bridge as the intended discernment target	Teach “one thing” without opening critical dimensions
Variation / invariance	What changes vs what stays fixed across a sequence	Example sets that separate critical aspects	Vary everything at once → noise; vary nothing → surface fixation
SAFE-TRANSFER	Transfer claims + sequences that survive evidence	See XLVIII.8	Story-skin theater; blocked-accuracy readiness

Operational definition (HYPOTHESIS): A MindCraft mission is transfer-*safe* when (a) the sequence makes at least one critical aspect *discernable* via planned contrast/generalization/separation/fusion, (b) “transfer” checks change a declared Barnett–Ceci dimension (format, surface story, temporal delay, or solo vs scaffolded), and (c) marketing never claims far transfer from immersion alone.

XLVIII.3 Far transfer is rare without design — name the dimensions

FACT (taxonomy): Barnett & Ceci (2002, *Psychological Bulletin*, 128(4), 612–637, doi:10.1037/0033-2909.128.4.612) — a century of “does far transfer occur?” debates compared apples to oranges; they propose a nine-dimension content×context framework. Estimating a *single* far-transfer effect size is misguided; evidence is substantial under some conditions and untested under others.

FACT (prosecution): Detterman (1993, in Detterman & Sternberg, Eds., *Transfer on Trial*, Ablex) — argues spontaneous far/general transfer is rare; apparent transfer often reflects cues/hints or other mechanisms; humans more often learn specific facts and patterns than free-floating abstractions.

FACT (analogical noticing failure): Gick & Holyoak (1980, *Cognitive Psychology*, 12, 306–355, doi:10.1016/0010-0285(80)90013-4) — subjects who read an analogous military “fortress” story often failed to spontaneously apply the convergence principle to Duncker’s radiation problem; **hints** dramatically raised solution rates. Structural isomorphism ≠ noticed transfer.

Product translation: MindCraft must not market “she’ll just see it on the ACT” after a kitchen-world linear-equations arc. Spontaneous far transfer is the *exception*. Design either **hugs** the exam context or **bridges** with explicit abstraction (compare two instances; name the invariant).

Kill: “Our stories create far transfer.” Stories can raise utility/attainment (XXXVII) and belonging (XXXVI); they do not substitute for noticing structure.

XLVIII.4 Mechanisms that *do* teach for transfer — low road, high road, hug, bridge

FACT (mechanisms): Salomon & Perkins (1989, *Educational Psychologist*, 24(2), 113–142, doi:10.1207/s15326985ep2402_1) — **low-road** transfer: extensive, varied practice → automatic triggering in similar stimulus conditions; **high-road** transfer: mindful abstraction, forward- or backward-reaching search for connections. Findings of transfer/nontransfer reflect whether those conditions were met.

FACT (instructional strategies): Perkins & Salomon (1992, “Transfer of Learning,” *International Encyclopedia of Education*, 2nd ed.; building on Perkins & Salomon, 1988, *Educational Leadership*, 46(1), 22–32) — **hugging** (practice that approximates target performance) vs **bridging** (prompt abstraction, metacognition, planned connections). Education gets transfer when designed for it — not by hoping.

HYPOTHESIS for MindCraft:

- **Hug:** format-tagged ACT/GCSE-like items, timed retrieval near exam, archetype-matched instances (Layer 2/3) after ingredient mastery.
- **Bridge:** coach prompts that force comparison (“what stayed the same?”), self-explanation of the invariant (XL), and bridge-gap cards when two concepts are known but unconnected.
- **Low-road alone** (Duo-like drill skins) will not deliver ACT problem-type novelty. **High-road alone** (chatty abstraction without varied instances) will not stick.

Wound: Absolute “transfer can’t be taught” (hard Detterman reading). Survives only as: spontaneous far transfer is rare; *cued and designed* transfer is the product job.

XLVIII.5 Variation theory — discernment requires planned difference

FACT (foundations): Marton & Booth (1997, *Learning and Awareness*, Lawrence Erlbaum) — learning is experiencing an object of learning in a new way; discerning critical aspects requires experiencing patterns of **variation and invariance**.

FACT (math classroom design): Kullberg, Runesson Kempe & Marton (2017, *ZDM Mathematics Education*, 49, 559–569, doi:10.1007/s11858-017-0858-4) — variation theory as a design principle changes *what is made possible to learn*; teachers using VT systematically open critical aspects (e.g., relating concepts rather than teaching one-at-a-time isolation); example sets matter, and teaching must **draw attention** to the patterns (presence of variation ≠ experienced variation).

FACT (dimensions of possible variation): Watson & Mason (line; e.g., Watson & Mason, 2005/2006 math-education elaborations of Marton) — to learn a mathematical concept is to become aware of which aspects can vary while remaining an example, and over what **range of permissible change**. Teacher and student may stress different dimensions.

FACT (Chinese bianshi tradition): Gu, Huang & Marton (2004, in Fan et al., Eds., *How Chinese Learn Mathematics*, World Scientific, pp. 309–347, doi:10.1142/9789812562241_0012) — Chinese “teaching with variation” systematically juxtaposes conceptual and procedural variants so students deepen structure; indigenous practice, not a Western bolt-on.

FACT (curriculum variation problems): Sun (2011a, *Educational Studies in Mathematics*, 76(1), 65–85, doi:10.1007/s10649-010-9263-4) — Chinese textbook “variation problems” organize connectedness (one problem,

multiple changes/solutions) to highlight invariants; contrasts with curricula that treat problems as isolated context decorations.

Product translation for the ontology:

1. Pick the **object of learning** (ingredient or bridge), not “finish five questions.”
2. Hold the critical structure **invariant**; vary one dimension at a time (numbers → representation → story skin → exam archetype).
3. After separation, **fuse** (vary jointly) only when WM/SAFE-DD allows (XXVI/XLI).
4. Prompt noticing (“what changed / what stayed the same?”) — ties XL self-explanation.
5. Interleaving (XXXIX) is related but not identical: IL trains discrimination among *categories*; VT trains discernment of *critical aspects within* an object of learning. Use both deliberately.

Kill: “More random story skins = better transfer.” Random surface change without invariant focus is entertainment variance, not Martonian variation.

XLVIII.6 Competitive audit (transfer / variation lens)

Competitor pattern	Transfer mechanism	MindCraft counter
Khan	Huge item pool; learner may block by topic; transfer left to chance	VT sequences + delayed mixed transfer checks
Duo	Low-road habit; skin-deep variety; streaks as proof	SAFE-HABIT + SAFE-TRANSFER: variety serves discernment, not XP
Brilliant	Elegant interactive instances; may under-specify exam hug	Hug ACT formats after insight; measure transfer_pass
ChatGPT Base	Fluent explanations; weak structured example contrast; sycophancy risk (XXXIII)	Deterministic card/path spine; SE + contrast prompts; no monologue⇒mastery
Tutor shops	Ad-hoc worksheets; expert tutors often vary well by instinct	Codify variation patterns in bank + tutor briefs (CIOF-compatible, XLVI)

XLVIII.7 Ontology / bank design implications (MindCraft-specific)

HYPOTHESIS — variation layers already latent in the stack:

- **Ingredient card styles** (geometric / algebraic / procedural) = representation variation with concept held.
- **Format axis** (word_problem, diagram, symbolic_expression, ...) = context/content-adjacent transfer dimensions — prefer matched then force a declared hop.
- **Bridges** = high-road objects of learning (connection itself is the critical aspect).
- **Archetypes (Layer 2)** = exam-pattern hug targets after near mastery.

- **Combinations** = fusion stage (multiple ingredients co-activated) — do not dump at cold start.

SPECULATION → **product rule**: Generation and bank curation should tag each practice *bundle* with: `invariant_aspect`, `varied_dimension`, `transfer_check_dimension`. Untagged bundles default to “drill,” not “transfer mission.”

FOUNDER BELIEF: Pathfinder recommendations that only maximize weakness severity without sequencing variation will grind the same surface until boredom — and still fail the ACT’s novel wrap.

XLVIII.8 Resolution doctrine — SAFE-TRANSFER

FOUNDER BELIEF / HYPOTHESIS — SAFE-TRANSFER stack:

1. **Name the hop.** Every transfer claim states which Barnett–Ceci-style dimension moved (format, story surface, delay, scaffold removal, domain).
2. **Design variation, don’t sprinkle it.** Sequences: contrast → separate critical aspect → generalize range → fuse when ready.
3. **Hug the exam; bridge the abstraction.** ACT-like items after structure is discernable; explicit compare/SE prompts for high-road.
4. **Never equate story immersion with transfer.** Narrative is motivation/belonging infrastructure (WORLD_VISION / XXXVI–XXXVII), not a transfer engine by itself.
5. **Blocked accuracy ≠ transferable skill.** Require delayed / mixed / format-shifted / solo checks (`transfer_pass`, `solo_transfer_pass`).
6. **Draw attention to the pattern.** Coach/tutor line: “What stayed the same?” is load-bearing UX, not flavor text.
7. **Respect load and anxiety.** Fusion + far hops only under SAFE-DD; exam week prefers hugged retrieval (XLVII).
8. **Measure what you sell.** Instrument transfer hops; kill missions that raise same-surface accuracy without hop gains.

Kill the false dichotomy: “Transfer is a myth, only drill” vs “Our world magically creates genius generalizers.” Refuse both. Ship **designed discernment**.

XLVIII.9 Marketing language that survives

- **Allowed**: “We change one thing at a time until the structure is obvious — then we change the wrapping.”
- **Allowed**: “Ready means she can do it in a new format tomorrow, not five identical items tonight.”
- **Allowed (parent)**: “Practice looks varied on purpose — same math backbone, different disguises.”
- **Banned**: “Immersive stories create far transfer.” / “AI explanations transfer automatically.” / Blocked-streak-as-mastery.
- **Banned**: Bloom 2-sigma; empty mindset; tutoring-is-free; vague far-transfer heroics.

Feature claim ladder (Part XXXIV hygiene)

Feature	Claim max without TR-* data
---------	-----------------------------

VT-sequenced bundles + noticing prompts	L1 mechanism aligned with Marton/Kullberg/Gu
Format-hop / delayed transfer checks as readiness	L1 product metric hygiene (Barnett–Ceci + Gick/Holyoak)
“Raises ACT by teaching far transfer”	L0–L1 only with pre-reg estimand; not L3 identity
Story-world $\equiv$ far transfer	L0 — banned

XLVIII.10 Experiments spawned

ID	Question	Design	Primary	Kill condition
TR-1	Same concept: VT sequence (one varied dimension) vs random surface skins	2-arm	next-day format-hop transfer_pass	Random $\geq$ VT $\rightarrow$ principled variation overstated
TR-2	Bridge object: explicit compare/SE vs monologue explanation then items	2-arm	bridge-gap item accuracy; solo_transfer_pass	Monologue $\geq$ compare $\rightarrow$ XL/bridge UX useless
TR-3	Hug-then-hop: ACT-archetype items after VT vs archetype-first cold	2-arm	archetype accuracy; affect; time	Cold $\geq$ VT-prep $\rightarrow$ hug sequencing wrong
TR-4	Transfer check: delayed mixed vs same-day blocked as “mission complete” gate	2-arm	48h transfer; return	Same-day gate predicts 48h as well $\rightarrow$ delay rule overstated
TR-QUAL	10 Maya think-alouds on “what changed?” prompts	interview	coded discernment of invariant	Students already name invariants $\rightarrow$ prompt priority drops

Pre-reg (XXXIV): TR-* identify **sequence design** $\times$ **declared transfer hop**. They do **not** identify “stories cause identity” or ACT score miracles from a single varied set.

Densifies: XXIX delayed readiness; XXXIX IL discrimination; XL SE; XLI SAFE-DD; format-weakness workstream; Layer 2/3 archetype bank; Part XLIX misconceptions queue.

XLVIII.11 Confidence table

Claim	Label	Confidence
Far-transfer debates need multi-dimension taxonomy (Barnett & Ceci, 2002)	FACT	High
Spontaneous far/general transfer is rare; cues matter (Detterman, 1993; Gick & Holyoak, 1980)	FACT	High
Low-road vs high-road + hug/bridge are usable design levers (Salomon & Perkins, 1989; Perkins & Salomon, 1992/1988)	FACT	High
Discernment requires variation/invariance patterns (Marton & Booth, 1997; Kullberg et al., 2017)	FACT	High
Chinese variation / bianshi is a systematic task-design tradition (Gu et al., 2004; Sun, 2011a)	FACT	High
SAFE-TRANSFER is right default for MindCraft bank/path/copy	FOUNDER BELIEF / HYPOTHESIS	Medium–High
Story immersion alone produces far transfer	SPECULATION / false as doctrine	Low (against)

XLVIII.12 What this chapter kills

1. **Kill:** Marketing story worlds as automatic far-transfer engines. (Gick/Holyoak noticing; Detterman rarity.)
2. **Kill:** Blocked same-surface accuracy as transferable / exam-ready proof. (Barnett–Ceci; XXIX/XXXIX readiness.)
3. **Kill:** Random cosmetic variety as “variation theory.” (Kullberg: what is made possible to learn; attention to patterns.)
4. **Kill:** AI monologue explanations as sufficient for transfer. (Hint/schema need; XL SE; XXXIII trust.)
5. **Wound:** “Transfer cannot be taught.” Survives only as: spontaneous far transfer is rare; designed hug/bridge/VT can raise hop success.
6. **Survive (constrained):** Named transfer hops; VT sequences; hug ACT + bridge abstraction; noticing prompts; delayed/format/solo metrics.

Doctrine until data: Ship **SAFE-TRANSFER**: vary on purpose so the backbone becomes obvious, hug the exam when claiming readiness, bridge with explicit abstraction, never sell lore or identical-item streaks as transfer. Market “same math backbone, different disguises — then prove it in a new wrapping tomorrow.”

Part XLIX — Misconceptions as Productive: Diagnostic Wrong Answers & Soft-Wrong Science

Chapter status: Living evidence brief — Researcher tick 2026-07-31 (UTC hour 00 ≡ Red Team slot, but ch49 never written → prefer Researcher per rotation; 8 researcher entries since synthesizer v1.6 → Synthesizer deferred to next eligible non-RT hour)

Primary question: How should MindCraft treat wrong answers — as shame tokens to erase, or as *diagnostic resources* that drive soft-wrong coaching, misconception routing, and honest mastery updates?

Owners: Product (Diagnostic / practice soft-wrong / Layer-4 misconception memory) · Engine (ingredient/misconception links) · Marketing claim ladder · Red Team

Commercial job: Ship a **SAFE-MISCON** stack — elicit → classify (lack vs confident misconception) → springboard without public shame → route remediation → measure retry + recalibration — kill “instant red-X pedagogy” and “all wrongs are equal.”

XLIX.1 Why this chapter exists

MindCraft already ships pieces of this thesis: soft-wrong UX, hide-correctness diagnostic (C4), Eedi misconception bank, Layer-1 canonical_misconception_family fields, and Experiment A (wizard coach on soft-wrong). What the Constitution lacked was an evidence spine that says *why* those features are load-bearing — and which competitor postures they kill.

Parts XXX (attribution), XXXVIII (mastery climate), XL (self-explanation), XLI (SAFE-DD), and XLVIII (SAFE-TRANSFER) constrain how misses feel and how sequences vary. They do not fully answer: **what is a productive wrong answer**, and when does revealing correctness help vs corrupt learning and the graph?

FOUNDER BELIEF under audit: Identity as a mathematical thinker grows when a miss becomes *information about structure* — not a grade, not a personality verdict, and not a fluent AI monologue that skips the student’s idea.

Claims we refuse as doctrine:

- 1. Misconceptions are unitary “bugs” that instruction must confront-and-replace wholesale.
 - 2. Instant correctness feedback = better learning for every miss.
 - 3. Soft-wrong tone alone equals diagnostic teaching.
 - 4. All incorrect selections update mastery the same way (guess ≡ confident wrong ≡ slip).
 - 5. Marketing “we celebrate mistakes” without routing, classification, or retry instrumentation.
-

XLIX.2 Constructs (product language)

Construct	Plain definition	Product signal	Failure mode if misused
-----------	------------------	----------------	-------------------------

Misconception (classic)	Stable, theory-like wrong idea that interferes with learning	Named <code>misconception_id</code> ; distractors that diagnose it	Treat every wrong as a deep theory
Knowledge-in-pieces / resources	Fragmentary, context-sensitive knowledge elements; flawed <i>and</i> productive	Route to refine/reorganize, not only “delete”	Romanticize every wrong as deep wisdom
Diagnostic wrong	Incorrect response that reveals <i>which</i> resource/misconception fired	Option → misconception map; confidence tier	Beauty distractors with no diagnostic map
Soft-wrong	Low-shame evaluation structure: miss is safe, informative, retry-inviting	Tone + no public rank + coach invite	Soft tone with no next move = empty kindness
Hide-correctness (C4)	Serve real items; record outcome; do not reveal key during diagnostic	Gap-scan / FEI probes	Revealing a wrong key as “gotcha” corrupts trust + can seed wrong keys into memory theater
SAFE-MISCON	Productive-error doctrine that survives evidence	See XLIX.8	Mistake-celebration posters; red-X gamification

Operational definition (HYPOTHESIS): A MindCraft miss is *productive* when (a) the response maps to a known misconception family or resource pattern, (b) the UX destakes identity threat (soft-wrong / mastery climate), (c) the next move uses the wrong as a springboard (compare, conflict, SE, bridge card) rather than only stating the key, and (d) telemetry separates lack-of-knowledge, slip, and high-confidence misconception for graph updates.

XLIX.3 Misconceptions reconceived — flawed *and* productive

FACT (constructivist critique): Smith, diSessa & Roschelle (1993/1994, *Journal of the Learning Sciences*, 3(2), 115–163, doi:10.1207/s15327809jls0302_1) — acknowledge empirical findings of misconceptions research but reject the accompanying view that students hold strongly held flawed ideas that instruction must *confront and replace*. That view overemphasizes discontinuity with experts and conflicts with constructivism: advanced knowledge is built from prior understandings. They sketch learning as refinement and reorganization within a complex-systems view; misconceptions are **both flawed and productive**.

FACT (knowledge in pieces): diSessa (1993, *Cognition and Instruction*, 10(2–3), 105–225, doi:10.1207/s1532690xc1002&3_2) — intuitive knowledge as weakly organized elements (p-prims / sense of mechanism), not a coherent wrong theory ready for wholesale replacement. Expertise densifies and reorganizes the system.

FACT (instructional lens): Hammer (1996, *Journal of the Learning Sciences*, 5(2), 97–127, doi:10.1207/s15327809jls0502_1) — “misconceptions” vs “p-prims” change what a teacher sees as the instructional task (replace a theory vs refine activation conditions of resources).

Product translation: MindCraft’s Layer-1 `failure_mode` / `canonical_misconception_family` and Eedi `mis_` IDs should drive **refinement routes** (when does this resource work? when does it fail?) — not only a shameful “you have Misconception X” badge. Coach copy: “That move works for this case — here’s where it breaks” beats “Wrong. Correct answer is B.”

Kill: Product language that treats Maya as carrying a disease-like misconception to be purged. Survive: diagnostic naming for *engine routing*, student-facing language as productive refinement.

XLIX.4 Errors as springboards — inquiry, not erasure

FACT (teaching experiment): Borasi (1994, *Journal for Research in Mathematics Education*, 25(2), 166–208, doi:10.5951/jresmetheduc.25.2.0166) — secondary students can capitalize on errors as **springboards for inquiry**; variations of the strategy exist; the point is stimulating mathematical inquiry from the error, not only teacher-side diagnosis for remediation.

FACT (error analysis tradition): Radatz (1979, *Journal for Research in Mathematics Education*, 10(3), 163–172, doi:10.5951/jresmetheduc.10.3.0163) — systematic error analysis has long diagnostic value for teachers; Borasi’s contribution is turning that value toward *student* inquiry.

FACT (diagnostic teaching / conflict): Bell (1993a, *Educational Studies in Mathematics*, 24, 5–34; 1993b, *Educational Studies in Mathematics*, 24, 115–137) and Swan’s diagnostic-teaching line (e.g., Swan, 2001, “Dealing with misconceptions in mathematics,” in *Issues in Mathematics Teaching*, Routledge) — long-term learning improves when lessons **expose** existing thinking, generate cognitive conflict across methods/answers, and resolve via reflective discussion — vs curricula that avoid analyzing errors or front-load explanations then graded practice.

HYPOTHESIS for MindCraft: Soft-wrong + wizard coach should implement a *compressed* springboard cycle: expose idea → contrast with invariant/conflict item → student-generated why (XL) → optional AI wrap under PWC — not Duo-style instant red/green with explanation dump.

Wound: Classroom conflict discussion does not transfer 1:1 to solo app UX. Survive as: design *micro-conflict* (two worked paths; “which claim survives?”) rather than claiming Socratic seminar in a phone UI.

XLIX.5 Soft-wrong science — safety is necessary, not sufficient

FACT (affect × attempt): Prior Constitution densification (XXIV anxiety/WM; XXXVIII mastery climate; XLI SAFE-DD destake) — threat appraisal suppresses attempt quality; evaluation structure that makes misses identity-safe raises willingness to engage difficulty.

FOUNDER BELIEF: Soft-wrong is mastery *evaluation structure*, not brand voice. Without diagnostic mapping + next move, soft-wrong is empty kindness.

HYPOTHESIS: Soft-wrong’s commercial wedge vs Khan/Duo is not “nicer feedback” — it is **recoverable struggle with competence evidence** (retry_120s, challenge_accept, solo_transfer_pass). Competitors can copy warmer copy; they rarely ship miss → misconception route → bridge/ingredient card → delayed transfer check.

Contradicting pressure: Parents often want instant correctness theater (“show the answer”). Claim ladder: L0–L1 session metrics for soft-wrong; do not sell L3 identity from tone alone (XXXIV).

Kill: Soft-wrong as marketing slogan without Experiment A instrumentation. Survive: soft-wrong as FEI stage (informative miss) gated by SAFE-MISCON routing.

XLIX.6 Confidence, calibration, and hide-correctness

FACT (hypercorrection): Butterfield & Metcalfe (2001, *Journal of Experimental Psychology: Learning, Memory, and Cognition*, 27(6), 1491–1494, doi:10.1037//0278-7393.27.6.1491) — contrary to interference predictions, **high-confidence errors** were *more* likely to be corrected after feedback than low-confidence errors (hypercorrection).

FACT (mechanism sketch): Butterfield & Metcalfe (2006, *Memory & Cognition* / follow-on attentional work; see also Metcalfe & Finn line) — corrective feedback after surprising high-confidence errors captures more attention; later work notes latent knowledge / “knew it all along” contributions. Product implication: confidence on wrongs is diagnostically precious, not just a UI flourish.

FACT (authentic math diagnostics online): Foster, Woodhead, Barton & Clark-Wilson (2022 advance of 2021 online-first; *Educational Studies in Mathematics*, doi:10.1007/s10649-021-10084-7) — large-scale diagnostic quiz data: higher confidence on an incorrect response associated with higher later correction probability on analogous items (hypercorrection signal in authentic math learning), after controls discussed in-paper; confidence assessment can prompt reflection.

FACT (calibration hygiene): Confidence–accuracy work in diagnostics (e.g., overconfidence on alternative conceptions; gender patterns in under/overconfidence across studies) — poor calibration is common; separating *misinformed* (confident wrong) from *uninformed* (low confidence / blank) changes instruction.

Product translation — C4 hide-correctness: During gap-scan / identity-sensitive diagnostics, revealing keys mid-flow can (a) spike threat, (b) turn the session into answer-hunting, (c) risk encoding wrong “gotcha” narratives. Record outcomes + confidence; defer key reveal to a designed springboard moment (or never in pure diagnostic). **Do not** market hide-correctness as “we never teach the answer” — that is false and banned under tutoring-is-free style slogans.

HYPOTHESIS: Practice mode should often *use* hypercorrection: after a high-confidence miss, prioritize attended corrective feedback + isomorphic retest; after low-confidence miss, prioritize scaffold/hug before confrontation (SAFE-DD equip).

Kill: Treating every wrong as equal graph evidence. Survive: confidence-tiered update policy (aligns with Layer-4 evidence_update_policy intent: don’t over-update from one item; separate can-do from can-recognize).

XLIX.7 What this means for the bank, ontology, and Solver

Bank / C5: Distractors should be designed misconceptions, not random wrongs. Eedi ingestion already mints `mis_ IDs` — enrichment linking misconceptions → ingredients remains the annotation bottleneck (CLAUDE.md). Without links, soft-wrong cannot route.

Ontology: Prefer nested ingredient `failure_mode` + Layer-2 `common_misconception_ids` as the join surface for coach cards. Variation sequences (XLVIII) that juxtapose the misconception-triggering surface vs the invariant are VT done right.

Solver / AI: Under PWC (XXXIII), AI wrap after student springboard — not “fluent correct answer” that ignores the selected distractor. Sycophancy risk: agreeing with a wrong path. Guard: name the student’s move, state where it applies, show the break.

Competitive wedge (XXXV): Duo celebrates streaks and instant binary feedback; ChatGPT delivers fluent keys; Khan explains. Minecraft sells **diagnostic wrong** → **recoverable path** with evidence the parent can see (retry, transfer) without leaderboard shame (SAFE-COMPARE).

XLIX.8 SAFE-MISCON stack (ship rule)

- 1. **Elicit** — Real items; prefer misconception-tagged distractors.
- 2. **Classify** — Lack / slip / high-confidence misconception (confidence + item family).
- 3. **Destake** — Soft-wrong mastery climate; no public shame; no streak-as-punishment for miss.
- 4. **Springboard** — Micro-conflict, compare, SE prompt, or Borasi-style inquiry move *using the wrong*.
- 5. **Route** — Ingredient / bridge / VT contrast card from misconception map.
- 6. **Reveal with purpose** — Feedback timing by confidence tier; hide-correctness in diagnostic; attended feedback after high-confidence miss.
- 7. **Measure** — `retry_120s`, recalibration Δ , misconception recurrence, delayed isomorphic accuracy — not “mistakes celebrated” vanity.

Marketing language that survives: “Wrong answers that teach,” “find the pattern in the miss,” “confidence that matches skill.”

Marketing language that dies: “We never say you’re wrong,” “mistakes are magic,” instant red-X as rigor, Soft-wrong™ as identity transformation (L3 from tone).

XLIX.9 Claim ladder

Claim	Max ladder without new data
Soft-wrong raises attempt/retry vs harsh red-X	L1 (Experiment A)
Misconception-tagged distractors improve routing vs random wrongs	L1
Confidence-tiered feedback beats uniform instant key	L1

Hide-correctness diagnostic yields cleaner weakness ranking	L1 (C1/C4 product bet)
Productive-error UX → identity as mathematical thinker	L3 — ban until transfer + identity items

XLIX.10 Experiments spawned

ID	Question	Design	Primary	Kill condition
MIS-1	Soft-wrong + springboard coach vs soft-wrong + key-only	2-arm	retry_120s; next-item accuracy	Key-only ≥ springboard → inquiry UX overstated
MIS-2	Misconception-mapped distractors vs random wrongs	2-arm	route precision; recurrence of same mis-ID	Random ≥ mapped → bank annotation ROI fails
MIS-3	Confidence-tiered feedback (hypercorrection-aware) vs uniform instant	2-arm	high-conf miss correction @ 48h	Uniform ≥ tiered → tiering complexity unjustified
MIS-4	Hide-correctness diagnostic vs reveal-as-you-go for gap-scan	2-arm	completion; affect; later worstweakness stability	Reveal ≥ hide on weakness quality + affect → C4 overstated
MIS-QUAL	10 Maya think-alouds after soft-wrong	interview	coded: shame vs curiosity; names own idea?	Students already springboard unaided → coach priority drops

Pre-reg (XXXIV): MIS-1...4 identify **error-handling design × classification**. They do **not** identify “celebrating mistakes causes identity.”

Densifies: Experiment A; C4 diagnostic; Layer-4 misconception memory; XL SE; XLVIII VT conflict pairs; Part L confidence calibration queue.

XLIX.11 Confidence table

Claim	Label	Confidence
-------	-------	------------

Misconceptions are both flawed and productive resources (Smith et al., 1993/94)	FACT	High
Knowledge-in-pieces / p-prims reframes instructional task (diSessa, 1993; Hammer, 1996)	FACT	High
Errors can be springboards for inquiry (Borasi, 1994)	FACT	High
Diagnostic teaching / conflict discussion beats error-avoidant exposition for long-term learning (Bell/Swan line)	FACT	Medium–High (classroom evidence; app transfer open)
High-confidence errors are hypercorrected after feedback (Butterfield & Metcalfe, 2001; Foster et al. math diagnostics)	FACT	High
SAFE-MISCON is right default for MindCraft soft-wrong / diagnostic / bank	FOUNDER BELIEF / HYPOTHESIS	Medium–High
Soft tone alone transforms identity	SPECULATION / false as doctrine	Low (against)

XLIX.12 What this chapter kills

1. **Kill:** Misconceptions-as-disease / confront-and-replace-only pedagogy as company doctrine. (Smith/diSessa/Roschelle; Hammer.)
2. **Kill:** Instant red-X + answer dump as the default “rigorous” product. (Borasi springboard; Bell/Swan diagnostic teaching.)
3. **Kill:** Soft-wrong as empty brand kindness without diagnostic route + retry metrics.
4. **Kill:** Equal graph updates for guess, slip, and confident misconception. (Hypercorrection / calibration literature.)
5. **Kill:** Hide-correctness marketed as “we never teach answers.”
6. **Wound:** Classroom conflict discussion $\equiv$ mobile UX. Survive as designed micro-conflict + SE, not seminar cosplay.
7. **Survive (constrained):** Misconception-mapped distractors; confidence-tiered feedback; soft-wrong as FEI informative miss; springboard-before-monologue; MIS-1...4.

Doctrine until data: Ship **SAFE-MISCON:** treat wrong answers as diagnostic fuel, destake identity threat, springboard from the student’s idea, route to structure, measure retry and recalibration — never sell mistake-celebration posters or red-X theater as learning science.

Part L — Confidence Calibration: Over/Underconfidence, Hide-Correctness, and Honest Self-Knowledge

Chapter status: Living evidence brief — Researcher tick 2026-07-31 (UTC hour 06 ≡ Red Team slot, but ch50 never written → prefer Researcher per rotation; 0 researcher entries since synthesizer v1.7)

Primary question: How should MindCraft treat student *confidence* — as a vibe to inflate, a score to gamify, or a *calibrated signal* that drives feedback, graph updates, and identity-safe diagnostics?

Owners: Product (Diagnostic C4 / soft-wrong / confidence UX) · Engine (evidence_update_policy / confidence tiers) · Marketing claim ladder · Red Team

Commercial job: Ship a **SAFE-CALIB** stack — elicit confidence → classify miss type → destake → recalibrate without inflation theater → tier feedback → measure confidence_miss_tier — kill “raise confidence” as North Star and “how sure?” without routing.

L.1 Why this chapter exists

The research question of this Constitution is how humans become *confident mathematical thinkers*. That phrase is dangerous if “confident” means louder, happier, or more certain. The commercial and scientific load-bearing meaning is **calibrated competence**: knowing when you know, knowing when you do not, and updating both skill and self-appraisal under informative misses.

Parts XXV (self-efficacy), XXX (attribution), XXXIII (AI overtrust), XLIX (SAFE-MISCON / hypercorrection), and the v1.7 metric confidence_miss_tier already point here. What was missing was a chapter that (a) defines calibration rigorously, (b) separates over- from underconfidence product responses, (c) ties C4 hide-correctness to metacognitive hygiene, and (d) refuses confidence theater as marketing.

FOUNDER BELIEF under audit: Identity as a mathematical thinker includes *honest self-knowledge* — not perpetual certainty, and not chronic self-doubt dressed as humility.

Claims we refuse as doctrine:

- 1. Raise confidence = improve learning (always).
- 2. Confidence meters / “Belief Score™” as identity North Stars.
- 3. Asking “how sure?” without changing feedback, routing, or graph updates.
- 4. Hide-correctness = never teaching answers.
- 5. Dunning–Kruger as a meme that justifies shaming low performers.
- 6. Instant correctness reveal as the only path to calibration.

L.2 Constructs (product language)

Construct	Plain definition	Product signal	Failure mode if misused
-----------	------------------	----------------	-------------------------

Calibration	Match between stated confidence and empirical accuracy across items	Confidence × correctness curves; bias & absolute accuracy indices	One-item “I feel sure” = calibrated
Overconfidence	Confidence systematically exceeds accuracy	High-conf wrongs; skipped review; false mastery	Shame the student; ignore hypercorrection opportunity
Underconfidence	Accuracy exceeds confidence	Correct + low conf; stuck on secure content	Force “believe harder” posters; skip challenge progression
Confidence of response	Certainty that <i>this</i> answer is correct (item-level)	0–10 or low/med/high next to answer	Global “math confidence” survey only
Hide-correctness (C4)	Record outcome; defer or withhold key reveal in diagnostic	Gap-scan / FEI probes	Never-teach-answers slogan; corrupt key theater
SAFE-CALIB	Calibration doctrine that survives evidence	See L.8	Confidence gamification; inflation theater

Operational definition (HYPOTHESIS): MindCraft treats calibration as a *session and graph* property when (a) item-level confidence is elicited on diagnostic and soft-wrong paths, (b) high-conf misses trigger attended corrective / springboard sequences, (c) low-conf corrects unlock challenge progression rather than endless drill, (d) mastery updates weight miss type (guess ≠ confident misconception), and (e) marketing claims stay at L0–L1 until CAL experiments report.

L.3 Calibration science — what “well calibrated” means

FACT (definition): Lichtenstein & Fischhoff (1977, *Organizational Behavior and Human Performance*, 20, 159–183; classic line) and Fischhoff, Slovic & Lichtenstein (1977, *Journal of Experimental Psychology: Human Perception and Performance*, 3(4), 552–564, doi:10.1037/0096-1523.3.4.552) — a judge is well calibrated if, over the long run, for propositions assigned probability p , the proportion true $\approx p$. Extreme certainty is routinely *inappropriate*: people are wrong too often when they feel certain.

FACT (bias pattern): Across general-knowledge and many educational tasks, the dominant bias is **overconfidence**, especially when accuracy is near chance; as accuracy rises, overconfidence often shrinks and underconfidence can appear at high accuracy (Lichtenstein & Fischhoff, 1977 summary pattern; replicated in educational monitoring work).

FACT (measurement plurality): Lingel, Lenhart & Schneider (2019, *ZDM*, metacognitive monitoring in mathematics — ERIC EJ1222655) — absolute accuracy, relative accuracy (e.g. Gamma), and diagnostic sensitivity/specificity are *not interchangeable*. Method choices change conclusions. Seventh-graders showed **pervasive overconfidence**; retrospective judgments and some scale formats improved monitoring precision vs prospective.

Product translation: MindCraft should log item-level confidence + correctness and report **bias** (over/under) and **discrimination** (does higher confidence predict correctness?) separately. Do not ship a single “Calibration %” vanity

metric parents can misunderstand.

Kill: One glamorous calibration number as growth KPI. Survive: confidence_miss_tier + bias curves for product learning, not for public leaderboards (SAFE-COMPARE).

L.4 Overconfidence, underconfidence, and the dual commercial problem

FACT (school math overconfidence): García, Rodríguez, González-Castro, González-Pienda & Torrance (2016, *Metacognition and Learning*; ERIC EJ1105998) — elementary students solving math problems were poorly calibrated and tended toward **overconfidence**; inaccurate calibrators differed in reported metacognitive process use.

FACT (ability × miscalibration): Kruger & Dunning (1999, *Journal of Personality and Social Psychology*, 77(6), 1121–1134, doi:10.1037/0022-3514.77.6.1121) — bottom-quartile performers on humor/grammar/logic tasks grossly overestimated ability; incompetence can impair the metacognitive skill needed to recognize incompetence. Improving skill can improve self-appraisal.

Wound (popular DK): The meme version (“stupid people are always overconfident”) overclaims, ignores regression-to-the-mean critiques, and invites product cruelty. Use as **FACT** about *metacognitive dual burden in low skill*, not as a brand insult toward Maya.

FACT (underconfidence is also costly): Foster’s CA program (Foster, 2016, *Educational Studies in Mathematics*, 91(2), 271–288, doi:10.1007/s10649-015-9660-9; Foster, 2016, *Mathematics Teaching*, 251, 11–13; Centre for Mathematical Cognition summary) — underconfidence can reduce satisfaction, block progression to harder material, and trap students in repetitive secure practice; overconfidence can embed uncorrected errors and complacency. Pedagogical aim is **appropriate** confidence, not maximal confidence.

HYPOTHESIS for MindCraft: Overconfident Maya needs destaked confrontation + springboard (SAFE-MISCON + hypercorrection). Underconfident Maya needs earned-win evidence and challenge unlocks (FEI), not pep talks. One “confidence booster” feature cannot serve both.

Kill: “We build confidence” as undifferentiated slogan. Survive: “confidence that matches skill” / “know what you know.”

L.5 Confidence assessment in classrooms — promise and null on attainment

FACT (CA mechanism): Confidence assessment asks students to rate certainty alongside each answer; scoring that rewards correct high-confidence and penalizes incorrect high-confidence (and also punishes under-reporting when correct) incentivizes *truthful* ratings (Foster, 2016; Gardner-Medwin line in HE). Aim includes discouraging guess-and-hope.

FACT (student acceptability): Foster (2016 ESM) — secondary students can make usable confidence judgments on procedures; generally positive about CA in short trials.

FACT (null on summative attainment — load-bearing contradiction): Foster (2021/2022, *International Journal of Science and Mathematics Education*, 20, 1411–1429, doi:10.1007/s10763-021-10207-9) — quasi-experimental CA across four secondary schools (N = 475), meta-analytic Cohen’s $d \approx -0.02$ [95% CI $-0.22, 0.19$], Bayes factor favoring

null on attainment gains vs controls. Conclusion: incorporating CA into low-stakes formative assessments does **not** appear detrimental — and also does **not** show a clear positive attainment effect from CA alone.

Commercial implication (critical): Shipping a “how sure?” widget is cheap and marketable. Evidence says **CA alone is not an attainment engine**. MindCraft’s wedge is not novelty CA — it is CA **wired** to soft-wrong routing, misconception maps, confidence-tiered feedback, and graph policy. Competitors can copy a confidence slider; they rarely ship the spine.

HYPOTHESIS: MindCraft CA improves *process metrics* (retry, misconception route hit, delayed isomorphic correction after high-conf miss) before it moves ACT scores. Claim ladder: process L1 first; attainment L2+ only after CAL + MIS experiments.

Kill: Marketing “confidence assessment raises grades.” Survive: CA as diagnostic fuel + metacognitive prompt, non-inferior on attainment per Foster 2021, valuable when coupled to FEI/SAFE-MISCON.

L.6 Hypercorrection, contradicting reasons, and feedback design

FACT (hypercorrection): Butterfield & Metcalfe (2001, *JEP:LMC*, 27(6), 1491–1494, doi:10.1037//0278-7393.27.6.1491) — high-confidence errors more often corrected after feedback than low-confidence errors.

FACT (math diagnostics): Foster, Woodhead, Barton & Clark-Wilson (*Educational Studies in Mathematics*, doi:10.1007/s10649-021-10084-7) — authentic online math diagnostic data show hypercorrection-consistent patterns; confidence assessment can prompt reflection.

FACT (debiasing via contradicting reasons): Koriat, Lichtenstein & Fischhoff (1980, *Journal of Experimental Psychology: Human Learning and Memory*, 6(2), 107–118) — listing reasons *against* the chosen answer improved appropriateness of confidence; supporting reasons alone did not. Confidence tracks strength of supporting evidence; people neglect disconfirming evidence unless prompted.

Product translation:

1. After high-conf miss → attended feedback + micro-conflict / “why might this be wrong?” (SE + Koriat) before fluent AI monologue (PWC / XXXIII).
2. After low-conf miss → scaffold/hug first (SAFE-DD), then lighter confrontation.
3. After low-conf correct → show competence evidence; invite slightly harder item (underconfidence unlock).
4. C4 hide-correctness during gap-scan: elicit confidence + answer; **defer** key reveal so the session stays diagnostic, not answer-hunting — then use purposeful reveal in practice springboards (XLIX).

Kill: Uniform instant red/green for every miss. Survive: confidence-tiered reveal + contradicting-reason prompt (aligns MIS-3).

L.7 Hide-correctness, AI trust, and graph integrity

FOUNDER BELIEF: C4 exists because revealing keys mid-diagnostic can spike threat, turn FEI into answer theater, and (if keys are wrong) poison trust. Recording outcomes without reveal preserves graph honesty for weakness ranking

(C1).

HYPOTHESIS: Calibration *improves* when students later reconcile confidence with delayed feedback in a destaked springboard — not when every item is a public gotcha.

Link to XXXIII: Fluent AI tutors create **overtrust** (calibration failure toward the *system*). Student self-appraisal and AI-skepticism are sibling problems.

Layer-4 alignment: Confidence tiers operationalize `evidence_update_policy` — high-conf wrong → stronger misconception evidence; low-conf wrong → weaker / “uninformed” path.

Kill: Equal mastery deltas for all wrongs; “we never teach the answer.” Survive: purposeful reveal; confidence-weighted updates; AI wrap after student generation.

L.8 SAFE-CALIB stack (ship rule)

1. **Elicit** — Item-level confidence on diagnostic and on soft-wrong / challenge items (not only global surveys).
2. **Classify** — Correct/incorrect × low/med/high → miss classes (guess, slip, confident misconception, underconfident correct).
3. **Destake** — Soft-wrong / mastery climate; never public-shame high-conf wrongs; no streak punishment for honest low confidence.
4. **Recalibrate** — Show mismatch privately (“sure + wrong” or “unsure + right”) as information, not identity verdict.
5. **Tier feedback** — Hypercorrection path for high-conf miss; hug-then-dose for low-conf miss; unlock challenge after underconfident correct.
6. **Prompt against** — One contradicting reason / “where does this break?” before answer dump (Koriat).
7. **Measure** — `confidence_miss_tier`, calibration bias Δ , high-conf miss correction @ delay, `underconf→challenge_accept` — not Belief Score™ or confidence streaks.

Marketing language that survives: “Know what you know,” “confidence that matches skill,” “wrong answers that teach,” “honest diagnostics.”

Marketing language that dies: “We raise confidence,” “believe in yourself” as product, Confidence League ranks, “never say you’re wrong,” CA-raises-grades claims without data.

L.9 Competitive wedge (brief)

Duo (binary + streaks), Khan (explain-first), Brilliant (scaffold dopamine), and ChatGPT tutors (fluent certainty → overtrust) all fail *calibration* differently. MindCraft’s L1-safe line: not louder confidence — **earned, calibrated** confidence that survives a format hop (solo transfer + confidence on unaided items + PWC).

L.10 Claim ladder

Claim	Max ladder without new data
Item-level CA is usable with teens and non-inferior on attainment vs BAU	L1 (Foster 2016; 2021 null)
High-conf misses deserve different feedback than low-conf misses	L1 (hypercorrection literature)
Contradicting-reason prompts improve confidence appropriateness	L1 (Koriat et al., 1980)
C4 hide-correctness improves weakness ranking quality + affect vs reveal-as-you-go	L1 product bet (MIS-4 / CAL-4)
Confidence-tiered graph updates beat equal-wrong updates	L1
CA alone raises summative math attainment	Banned pending contrary RCT (Foster 2021 against)
Calibrated confidence → identity as mathematical thinker	L3 — ban until transfer + identity instruments

L.11 Experiments spawned

ID	Question	Design	Primary	Kill condition
CAL-1	Confidence-tiered feedback vs uniform instant key	2-arm	high-conf miss correction @ 48h; retry_120s	Uniform $\geq$ tiered → tiering unjustified
CAL-2	Contradicting-reason prompt vs explain-correct-only after miss	2-arm	calibration bias Δ ; next isomorphic accuracy	Explain-only $\geq$ prompt → Koriat UX overstated for app
CAL-3	Underconfident-correct → challenge unlock vs more secure drill	2-arm	challenge_accept; later accuracy	Drill $\geq$ unlock on growth → underconf policy wrong
CAL-4	Hide-correctness diagnostic vs reveal-as-you-go (shared with MIS-4)	2-arm	completion; affect; worstweakness stability	Reveal wins on weakness + affect → C4 overstated
CAL-5	Confidence-weighted mastery updates vs equal-wrong updates	2-arm offline / sim + online	prediction of next-session accuracy	Equal $\geq$ weighted → engine complexity unjustified

CAL-QUAL	10 Maya interviews: meaning of “sure” / shame after high-conf miss	interview	coded: threat vs curiosity; wants reveal when?	Students refuse confidence UX → redesign elicit
----------	--	-----------	--	---

Pre-reg (XXXIV): CAL-* identify **confidence** × **feedback** / **update design**. They do **not** identify “raising confidence causes identity.”

Densifies: MIS-3/4; C4; Layer-4 evidence policy; XXV self-efficacy; XXXIII AI overtrust; metric confidence_miss_tier.

L.12 Confidence table

Claim	Label	Confidence
Calibration = long-run match of confidence to accuracy (Fischhoff/Lichtenstein line)	FACT	High
Overconfidence is common; extreme certainty often unwarranted	FACT	High
School math students often poorly calibrated / overconfident (García et al.; Lingel et al.)	FACT	High
DK dual burden for low skill is real; meme uses are overstated	FACT / WOUND	Medium–High
CA usable + non-inferior attainment; not a proven attainment booster alone (Foster 2016; 2021)	FACT	High
Hypercorrection after high-conf errors (Butterfield & Metcalfe; Foster et al.)	FACT	High
Contradicting reasons improve confidence appropriateness (Koriat et al., 1980)	FACT	High
SAFE-CALIB is right default for MindCraft diagnostic / soft-wrong / graph	FOUNDER BELIEF / HYPOTHESIS	Medium–High
Raise-confidence marketing transforms identity	SPECULATION / false as doctrine	Low (against)

L.13 What this chapter kills

1. **Kill:** “Raise confidence” as undifferentiated North Star or brand promise. (Foster: appropriate confidence; under- and over- both hurt.)
2. **Kill:** Confidence meters, Belief Score™, or confidence streaks as identity engines. (SAFE-COMPARE / SAFE-HABIT alignment.)
3. **Kill:** “How sure?” UI without tiered feedback, routing, or update policy. (Foster 2021 null on CA-alone attainment.)
4. **Kill:** CA / confidence assessment marketed as raising grades without new evidence.
5. **Kill:** Equal graph updates for guess vs high-confidence misconception.
6. **Kill:** Hide-correctness as “we never teach answers.”
7. **Kill:** Dunning–Kruger shaming as brand voice.
8. **Wound:** Classroom CA $\equiv$ app CA; HE medical CA $\equiv$ HS math. Survive as: elicit + truthful incentive design + wired spine, not worksheet cosplay.
9. **Survive (constrained):** Item-level confidence; SAFE-CALIB stack; hypercorrection-aware feedback; underconfidence unlocks; CAL-1...5; coupling to SAFE-MISCON.

Doctrine until data: Ship **SAFE-CALIB**: treat confidence as a diagnostic signal for honest self-knowledge and differential pedagogy — never as a vibe to inflate, a league to rank, or a substitute for FEI evidence.

Part LI — Deliberate Practice in Tutoring Sessions: Ericsson vs Hambrick, and the 45-Minute Design

Chapter status: Living evidence brief — Researcher tick 2026-07-31 (UTC hour 09; hour%6≠0; 1 researcher entry since synthesizer v1.7 → Researcher)

Primary question: What should a MindCraft tutoring / practice session *do* if “practice” is not volume, not streak minutes, and not answer delivery — and how far can Ericsson’s deliberate-practice (DP) frame travel before Macnamara–Hambrick evidence kills overclaim?

Owners: Product (session architecture / tutor briefing) · Engine (weakness targeting / feedback loops) · Marketing claim ladder · Red Team

Commercial job: Ship a **SAFE-DP** session stack — diagnose → isolate edge skill → goal → effortful attempt → informative feedback → revise → brief rest → solo check — kill 10,000-hours theater, grind-volume North Stars, and “we do deliberate practice” as brand without a coach-designed edge task.

LI.1 Why this chapter exists

MindCraft sells sessions: college tutors, practice missions, Solver, Map. The Constitution’s research question is identity transformation under difficulty. Deliberate practice is the popular science answer to “how experts get good.” It is also a marketing trap — competitors and edtech decks cite Ericsson while shipping naive drill (more of the same),

gamified volume (streaks/XP), or fluent AI monologue (pseudo-coaching).

Parts XXVI (CLT / fading), XXIX (spacing/retrieval), XXXI (challenge–skill), XXXIX (interleave), XL (self-explanation), XLI (SAFE-DD), and L (SAFE-CALIB) already constrain *what kind* of difficulty and feedback are safe. What was missing was a chapter that (a) defines DP vs purposeful vs naive practice, (b) faces the Macnamara–Hambrick variance numbers honestly, (c) translates math-education DP (Lehtinen line) into a **45-minute** tutor/product design, and (d) refuses talent-denial and grind-as-virtue as doctrine.

FOUNDER BELIEF under audit: A high-quality MindCraft session is closer to *coach-designed practice at the edge of ability with feedback* than to homework help, content library, or engagement loops.

Claims we refuse as doctrine:

1. 10,000 hours / “practice makes perfect” as causal story.
2. Accumulated minutes / streaks as proof of DP.
3. DP explains nearly all expert variance (popular overread).
4. Any hard drill = DP; AI chat without diagnosis = DP coach; talent-denial as brand.

LI.2 Constructs (product language)

Construct	Plain definition	Product signal	Failure mode if misused
Naive practice	Repeating what you can already do; play / homework slog	High accuracy, low struggle, no edge target	“Completed 40 problems” vanity
Purposeful practice	Goal + focus + feedback, possibly self-directed (Ericsson & Pool, <i>Peak</i> , 2016)	Student sets a micro-goal; app gives feedback	Duo-like streaks dressed as purpose
Deliberate practice (strict)	Coach/teacher designs tasks beyond current reliable performance; immediate informative feedback; refinement; builds mental representations (Ericsson et al., 1993; Ericsson & Pool, 2016)	Tutor/engine picks weakness; task just beyond; feedback revises method	Labeling all structured homework “DP”
Structured practice (meta-analytic)	Broad “practice designed to improve” used in Macnamara et al. (2014) — often looser than Ericsson’s DP	Hours logged in studies	Equating study-hours with DP

Mental representation	Organized internal model that lets experts monitor and improve (Ericsson line; Lehtinen et al., 2017)	Transfer / format hop / self-monitor	Fluency without representation
SAFE-DP	DP doctrine that survives Hambrick wound + anxiety/CLT gates	See LI.8	Grind theater; 10k-hours copy

Operational definition (HYPOTHESIS): A MindCraft block counts as *deliberate-practice-shaped* when (a) target is a diagnosed edge (concept, ingredient, bridge, format, or miss class), (b) task is beyond reliable solo performance but solvable with brief support, (c) feedback changes the *next attempt's method*, not only the answer key, (d) student generates at least one why / check before AI wrap (PWC / XL), and (e) session ends with a short unaided probe (solo_transfer_pass or isomorphic check) — not with “minutes practiced.”

LI.3 Ericsson’s core claim — what is actually said

FACT (foundational paper): Ericsson, Krampe & Tesch-Römer (1993, *Psychological Review*, 100(3), 363–406, doi:10.1037/0033-295X.100.3.363) — expert performance is strongly associated with accumulated *deliberate practice*: highly structured activities designed to improve specific aspects of performance, typically with teacher/coach guidance, clear goals, feedback, and effortful concentration. Mere experience / play / mindless repetition do not substitute.

FACT (effort constraint): Same paper and later summaries — effective daily DP duration is limited; early practice often ~1 hour or less; laboratory extended-practice protocols commonly ~1 hr sessions; pushing past optimal duration yields diminishing returns and fatigue. Rest is part of the training system, not laziness (aligns SAFE-TIMING).

FACT (Peak distinction): Ericsson & Pool (2016, *Peak*) — **purposeful practice** = well-defined goal, focus, feedback, get-out-of-comfort-zone; **deliberate practice** = purposeful practice *plus* a field with known training methods and a teacher who designs/adapts exercises. Popular “10,000 hours” readings flatten this distinction and often cite Gladwell’s popularization rather than Ericsson’s constraints.

Product translation: MindCraft’s tutor + ontology spine can *aspire* to the coach role (diagnose → design edge task → feedback). Solo app loops without diagnosis are at best purposeful practice — and often naive practice with a nicer UI.

Kill: “10,000 hours of MindCraft.” Survive: “coach-designed practice at your edge, in a session you can sustain.”

LI.4 The Hambrick / Macnamara wound — variance is not destiny theater

FACT (meta-analysis): Macnamara, Hambrick & Oswald (2014, *Psychological Science*, 25(8), 1608–1618, doi:10.1177/0956797614535810) — across domains, deliberate-practice measures explained roughly **26%** of variance in games, **21%** music, **18%** sports, **4%** education, and **<1%** professions. Conclusion: DP is important, but **not as important as popular arguments claim**.

FACT (sports follow-up): Macnamara, Moreau & Hambrick (2016, *Perspectives on Psychological Science*, 11(3), 333–350, doi:10.1177/1745691616635591) — in sports, substantial unexplained variance remains; elite vs sub-elite comparisons weaken simple “more DP → elite” stories.

FACT (reply chain): Ericsson (2016, *Perspectives on Psychological Science*, doi:10.1177/1745691616635600) argues meta-analyses sum *heterogeneous practice hours* and dilute the construct; Macnamara & Hambrick (2016, doi:10.1177/1745691616635614) reply that Ericsson’s criteria shift and that existing evidence does not support “largely accounted for by DP.”

FACT (definition rescue attempt): Ericsson & Harwell (2019, *Frontiers in Psychology*, doi:10.3389/fpsyg.2019.02396) — re-emphasize original DP criteria; claim Macnamara et al. often measured **structured practice**, not strict DP; reanalyses with stricter inclusion raise explained variance. Independent methodologists still warn against treating any reanalysis as settling talent, opportunity, and starting-skill confounds.

Wound for MindCraft marketing: Education’s ~4% figure is a Red Team grenade against “science says deliberate practice is everything.” Even if Ericsson is right that many education studies measured weak proxies, MindCraft cannot sell DP as a near-sufficient cause of identity or ACT gains.

HYPOTHESIS: MindCraft should treat DP as a **session design grammar** (how to spend 45 minutes), not as a **variance-explained brand claim**. Commercial strength is *individualized edge practice with feedback* — the part both camps agree matters more than mindless hours — without promising that practice volume explains Maya’s future.

Kill: “Experts are made of deliberate practice — so we made you an expert factory.” Survive: “Most homework isn’t practice that changes you; we design the next hard thing on purpose.”

LI.5 Mathematics education — drill vs deliberate practice

FACT (math DP review): Lehtinen, Hannula-Sormunen, McMullen & Gruber (2017, *ZDM Mathematics Education*, 49, 625–636, doi:10.1007/s11858-017-0856-6) — school math “practice” is often **drill-and-practice** that automatizes procedures but can yield *inert* routines rather than adaptive number knowledge; DP in expertise research involves thinking, problem solving, reflection, and dynamically retargeted training — not mechanical fluency alone. They summarize Ericsson (2016) features: coach-known skills, beyond-current-ability effort, specific goals, full attention, feedback + modification, building mental representations, improving prior skills with correct fundamentals.

FACT (study time ≠ improvement): Plant, Ericsson, Hill & Asberg (2005, *Contemporary Educational Psychology*, 30(1), 96–116) — in higher education, amount of study time did not significantly predict GPA once quality/concentration/goals entered; improvement tied to concentrated, goal-directed learning (cited in Lehtinen et al. as caution against raw hours).

HYPOTHESIS for MindCraft: Ontology ingredients + bridges are the *right decomposition* for DP-style targeting (isolate the failing subprocess, then rebuild), but only if sessions refuse “do 20 similar ACT items” as the default. Blocked accuracy is acquisition theater (XXXIX); DP needs the edge miss class (SAFE-MISCON / SAFE-CALIB), then varied probes (SAFE-TRANSFER).

Link to CLT / anxiety: Beyond-ability tasks without destake/equip violate SAFE-DD. DP’s “near-maximal effort” is **not** permission for threat flooding. For Maya, the coach’s first job is often making the edge task *informative*, not merely hard.

Kill: Back-to-basics grind marketed as Ericsson. Survive: adaptive practice that builds representations — Lehtinen’s contrast as competitive copy vs worksheet mills.

LI.6 What a 45-minute MindCraft session should look like

FOUNDER BELIEF: Human tutor sessions (~45–60 min) and digital missions should share one DP-shaped spine so marketing demos, tutor briefs, and app UX do not diverge into “warm homework help” vs “streak game.”

Proposed spine (HYPOTHESIS — SAFE-DP):

Minute block	Move	Maps to
0–5	Affect + goal: name <i>one</i> edge target (weakness / bridge / format / miss class)	FEI safety; $E \times V$; SAFE-EXPECTANCY task-only language
5–12	Quick diagnose / warm isomorphic (hide-correctness OK if FEI)	C4; SAFE-CALIB elicit if used
12–28	Edge attempts: 1–3 items beyond reliable solo; student generation before explain	DP effort; PWC; SE (XL); SAFE-DD dose
28–36	Feedback that revises method: contradicting reason / springboard / ingredient card — not answer dump	Koriat; SAFE-MISCON; ingredient runtime
36–42	Near transfer / format hop / interleaved foil	SAFE-TRANSFER; XXXIX
42–45	Solo check + attribution (“what changed in your method?”) + next cue	solo_transfer_pass; SAFE-HABIT if-then; no streak sermon

Effort constraint product rule: Prefer one sharp DP cycle over stuffing eight topics. Marathon “grind nights” violate SAFE-TIMING and Ericsson’s own duration cautions.

Tutor ethics: Brief tutors on *task* targets from the graph — never trait labels (SAFE-EXPECTANCY). “Today we strengthen the functions→rate bridge” beats “she’s weak and unmotivated.”

AI role: Engine proposes the edge target and card order; LLM wraps language *after* student attempt (PWC). Unguarded Solver-as-session is not DP.

LI.7 Competitive wedge (brief)

Khan often explain-then-block; Duo sells streak volume; Brilliant scaffolds delight; ChatGPT tutors monologue without diagnosis; traditional tutoring delivers answers. MindCraft’s L1-safe line: not “more practice” —

better-designed practice (one edge, method-changing feedback, solo proof). Minutes, XP, and fluent help are not DP.

LI.8 SAFE-DP stack (ship rule)

1. **Diagnose** — Pick one edge from graph / gap / tutor brief (concept, ingredient, bridge, format, or confidence-miss class).
2. **Isolate** — Decompose to the failing subprocess (ontology), not a random worksheet.
3. **Goal** — State a performance micro-goal for the block (“solve rate problems when the story hides the function”).
4. **Stretch** — Tasks beyond reliable solo, inside SAFE-DD capacity (equip/destake first if anxious).
5. **Attempt** — Student generates; no AI monologue-first (PWC / XL).
6. **Feedback** — Informative, method-revising; confidence-tiered when elicited (SAFE-CALIB / MISCON).
7. **Revise** — Immediate retry or isomorphic variant that uses the new method.
8. **Prove** — Short unaided / format-hop check; log `solo_transfer_pass`, not minutes.
9. **Stop** — Respect effort constraint; schedule sleep/spacing (SAFE-TIMING / XXIX).
10. **Measure** — Edge-target hit rate, method-change coding, transfer — never 10k-hours progress bars.

Marketing language that survives: “Practice at your edge,” “coach-designed sessions,” “feedback that changes your next try,” “show what you can do alone.”

Marketing language that dies: “10,000 hours,” “grind mindset,” “practice makes perfect,” “XP = mastery,” “our AI is your deliberate-practice coach” without diagnosis spine, talent-denial (“anyone can be elite if they DP”).

LI.9 Claim ladder

Claim	Max ladder without new data
Strict DP ≠ mere study hours / drill fluency	L1 (Ericsson 1993; Lehtinen 2017; Plant et al. 2005)
Popular “DP explains almost all expertise” overclaims; education variance small in Macnamara et al.	L1 (2014 meta; reply chain)
Coach/teacher-designed edge tasks + feedback beat naive repetition for skill change	L1 (expertise literature consensus on quality > raw hours)
MindCraft 45-min SAFE-DP spine improves solo transfer vs homework-help control	L1 product bet — needs DP experiments
DP volume alone creates identity as mathematical thinker	Banned (L3 without instruments; Hambrick wound)
Talent/opportunity irrelevant	Banned as brand

LI.10 Experiments spawned

ID	Question	Design	Primary	Kill condition
DP-1	SAFE-DP spine vs homework-help / explain-first session (same tutor time)	2-arm	solo_transfer_pass; method-change code	Help $\geq$ DP spine $\rightarrow$ spine overstated
DP-2	One edge target vs multi-topic coverage in 45 min	2-arm	delayed isomorphic accuracy; cognitive load / affect	Coverage $\geq$ focus on learning $\rightarrow$ “one edge” wrong for ACT cram segment
DP-3	Engine-proposed ingredient target vs tutor-chosen topic without graph	2-arm	edge-hit; transfer	Tutor-only $\geq$ engine $\rightarrow$ graph targeting weak
DP-4	Feedback revises method (SE + springboard) vs answer-key-only after miss	2-arm	next-attempt success; retry_120s	Key-only $\geq$ method feedback $\rightarrow$ DP feedback claim weak
DP-5	Session cap (~45–60 focused min) vs encouraged grind (+30 min similar items)	2-arm	delayed accuracy; sleep/affect	Grind $\geq$ cap $\rightarrow$ effort-constraint messaging wrong for this population
DP-QUAL	10 Maya + 5 tutor interviews: “real practice” vs help	interview	coded: edge / threat / agency	Spine feels like cruelty $\rightarrow$ redesign destake

Pre-reg (XXXIV): DP-* identify **session architecture / targeting / feedback quality** — not “DP causes identity” or “explains elite careers.” Densifies FEI, PWC, SAFE-DD/TRANSFER/MISCON/CALIB/TIMING; prefer transfer metrics over minutes.

LI.11 Confidence table

Claim	Label	Confidence
Ericsson et al. (1993): DP = coach-guided, feedback-rich, effortful practice beyond current performance	FACT	High

Practice duration constrained; rest matters; purposeful ≠ deliberate without coach (<i>Peak</i>)	FACT	High
Macnamara et al. (2014): modest DP variance; ~4% education; definition dispute unsettled	FACT	High
Lehtinen et al.: math drill ≠ DP; Plant et al.: hours ≠ achievement when quality ignored	FACT	High
SAFE-DP is right default session grammar for MindCraft	FOUNDER BELIEF / HYPOTHESIS	Medium–High
10k-hours / talent-denial / grind NS transform learners	SPECULATION / false as doctrine	Low (against)

LI.12 What this chapter kills

1. **Kill:** 10,000-hours / “practice makes perfect” as causal brand.
2. **Kill:** Streak minutes, problem counts, or XP as DP proof.
3. **Kill:** “DP explains expertise — therefore our app creates experts” (education ~4% wound).
4. **Kill:** Hard drill ≡ DP; AI monologue ≡ DP coach; talent-denial; multi-topic coverage theater as default rigor.
5. **Wound:** Violin-lab DP ≠ anxious ACT novice tutoring — survive as quality grammar + SAFE-DD, not music cosplay.
6. **Survive:** SAFE-DP spine; ontology edge targeting; method-revising feedback; effort caps; DP-1...5; “better-designed practice” copy.

Doctrine until data: Ship **SAFE-DP:** one diagnosed edge, effortful attempt, feedback that changes the next try, short solo proof, stop before grind — never sell hours, streaks, or talent-denial as the learning engine.

Part LII — Community of Practice (Lave/Wenger): Legitimate Peripheral Participation — Careful Transfer

Chapter status: Living evidence brief — Researcher tick 2026-07-31 (UTC hour 12 ≡ Red Team slot, but ch52 never written → prefer Researcher per rotation; 2 researcher entries since synthesizer v1.7 → Researcher)

Primary question: What does *community of practice* / *legitimate peripheral participation* (LPP) actually claim — and how far can MindCraft transfer those ideas into product, tutoring, story worlds, and “Community” UX before the transfer becomes cosplay?

Owners: Product (social / tutor / world design) · Marketing claim ladder · Equity (whose community?) · Red Team

Commercial job: Ship a **SAFE-CoP** stack — participation trajectories with real old-timers and real practices, not Discord-as-community, leaderboard-as-belonging, or “we’re a CoP” brand — while killing apprenticeship theater that

ignores Boylan/Fuller wounds.

LII.1 Why this chapter exists

The Constitution’s thesis is identity transformation under difficulty. Lave & Wenger’s situated-learning line is the classic social account of how newcomers become members of a practice — midwives, tailors, quartermasters — not how students complete problem sets. Edtech decks love the phrase *community of practice* because it sounds like belonging, mentors, and identity without specifying *what practice* students are joining.

MindCraft already has social-adjacent surfaces: college tutors, story worlds, Community nav, parent/tutor roles, Map as shared ontology. Without this chapter, those surfaces drift toward (a) Duo-style leagues as belonging, (b) forum/Discord theater as CoP, or (c) AI chat as “apprenticeship” with no old-timer practice. Parts XXV (identity), XXXVI (equity of worlds), XLII (SAFE-COMPARE), XLVI (tutor expectancy), and LI (SAFE-DP coach role) constrain pieces; what was missing was an honest transfer of LPP into product law.

FOUNDER BELIEF under audit: Becoming a confident mathematical thinker is partly *joining a practice* (ways of posing, checking, arguing, recovering from wrong), not only acquiring items — but school/ACT prep is not a tailor shop, and naming the product a CoP does not make LPP true.

Claims we refuse as doctrine:

1. “MindCraft is a community of practice” as marketing without naming the practice and the peripheral→central trajectory.
 2. Forums, streaks, or chat volume as proof of LPP.
 3. AI monologue as old-timer apprenticeship.
 4. School/ACT classrooms (or solo app sessions) as interchangeable with ethnographic apprenticeship sites.
-

LII.2 Constructs (product language)

Construct	Plain definition	Product signal	Failure mode if misused
Community of practice (CoP)	People engaged over time in a joint enterprise with shared repertoire (Wenger, 1998)	Tutors + peers + world characters sharing <i>how math is done here</i>	Calling any group chat a CoP
Legitimate peripheral participation (LPP)	Newcomers participate in real practice at the edge; access and legitimacy grow toward fuller roles (Lave & Wenger, 1991)	Student contributes partial reasoning / checks; tutor/old-timer holds the full craft	Watching a video ≠ peripheral participation

Old-timer / newcomer	Asymmetric competence + responsibility in the practice	College tutor, veteran Maya peer, coach spine	Peer leaderboard as fake expertise hierarchy
Repertoire	Shared ways of talking, tools, stories, criteria of “good work”	Soft-wrong language; Map ingredients; story check-moves	Lore dump without practice criteria
Acquisition metaphor	Learning as gaining knowledge/skills as possessions (Sfard, 1998)	Mastery graph, item bank, transfer probes	Only metaphor → identity ignored
Participation metaphor	Learning as becoming a member of a practice (Sfard, 1998)	“How we recover from wrong here”	Only metaphor → contentless belonging
SAFE-CoP	CoP doctrine that survives Boylan/Fuller transfer wounds + Sfard pluralism	See LII.8	Apprenticeship cosplay; Discord theater

Operational definition (HYPOTHESIS): A MindCraft interaction counts as *LPP-shaped* when (a) the student does a real piece of the practice (attempt, explain, check, revise — not only consume), (b) an old-timer (human tutor or SE-constrained coach) models criteria of competence, (c) legitimacy is task/role based (you may try this edge), not trait based (you are a weak student — SAFE-EXPECTANCY), and (d) progress is visible as *changing participation* (more solo responsibility, wider repertoire) — not as XP.

LII.3 Lave & Wenger (1991) — what LPP actually says

FACT (foundational book): Lave & Wenger (1991, *Situated Learning: Legitimate Peripheral Participation*, Cambridge University Press, ISBN 9780521423748; doi:10.1017/CBO9780511815355) — learning viewed as situated activity has as its central defining characteristic **legitimate peripheral participation**: learners participate in communities of practitioners and move toward fuller participation in the sociocultural practices of the community. LPP names relations among newcomers and old-timers, activities, identities, artifacts, and communities of knowledge and practice — not a classroom seating chart.

FACT (cases): Their analytic cases are midwives, Yucatec midwives’ apprentices, Vai and Gola tailors, navy quartermasters, butchers, and Alcoholics Anonymous — craft and mutual-aid settings with durable joint enterprises, not standardized exam prep.

FACT (learning subsumed): They argue that becoming a participant *includes* skill learning; identity and membership are not add-ons after content delivery.

Product translation: MindCraft’s research question (“confident mathematical thinkers”) rhymes with *identity in practice*. The danger is importing the *label* while shipping acquisition-only UX (explain → pick A–D → streak).

Kill: “Situated learning = our story world backdrop.” Survive: “Participation means doing pieces of mathematical practice with growing responsibility.”

LII.4 Wenger (1998) — meaning, practice, community, identity

FACT: Wenger (1998, *Communities of Practice: Learning, Meaning, and Identity*, Cambridge University Press; doi:10.1017/CBO9780511803932) — social theory of learning with four interlocking components: **meaning** (experiencing the world as meaningful), **practice** (shared historical/social resources for mutual engagement), **community** (social configurations where competence is recognized), **identity** (learning changes who we are). CoPs are informal groupings that form as people pursue shared enterprises over time — primary unit of analysis is neither lone individual nor formal institution alone.

HYPOTHESIS for MindCraft: The product must specify the *enterprise* (e.g., recovering from productive wrongs on ACT-path concepts; telling-and-checking in a story world) and the *repertoire* (soft-wrong, SE prompts, Map language). Without enterprise + repertoire, “Community” nav is a social tab, not a CoP.

Wound: Corporate/edtech CoP literature often flattens Wenger into “knowledge-sharing groups.” MindCraft must not sell community features as identity transformation (claim ladder L3) from engagement metrics (L0/L1).

LII.5 The transfer wound — school math is not a tailor shop

FACT (school mathematics critique): Boylan (2010, *Teaching and Teacher Education*, 26(1), 61–70, doi:10.1016/j.tate.2009.08.005) — LPP was developed from informal apprenticeship contexts; applicability to school classrooms is **problematic**, especially where teacher-centred, decontextualised procedural practices dominate “usual” school mathematics. He critiques LPP transfer on social practice, learning relationships, power, agency, and identity; proposes **ecologies of participation** to capture multidimensional participation (imposed/voluntary, engagement/disengagement, identity risk/maintain, marginal/peripheral, etc.). He notes that if LPP is not found in usual school math, treating it as a *universal* feature of situated learning is undermined — while still valuing Lave & Wenger as a contrastive model and intervention heuristic.

FACT (workplace reassessment): Fuller, Hodkinson, Hodkinson & Unwin (2005, *British Educational Research Journal*, 31(1), 49–68, doi:10.1080/0141192052000310029) — LPP remains insightful for workplace learning research, but has **significant limitations** in complex institutional settings (manufacturing, secondary schools): diverse patterns of participation; institutional environments configure opportunities and barriers that ethnographic craft cases under-specify.

FACT (metaphor pluralism): Sfard (1998, *Educational Researcher*, 27(2), 4–13, doi:10.3102/0013189X027002004) — acquisition and participation metaphors each entail desirable and harmful practices; **choosing just one** risks theoretical distortion and bad pedagogy. MindCraft’s graph/mastery engine is acquisition-shaped; identity thesis is participation-shaped — doctrine needs both, not a CoP slogan that deletes assessment.

FACT (math classroom situated work — careful ally): Boaler (1999 AERA paper / related classroom studies; see also Boaler & Greeno identity/agency line summarized in math-education CoP literature) — classroom communities’ constraints and affordances shape what students learn to *do*; formalized classrooms attune students to practices that may not transfer. Useful as evidence that *participation forms matter*, not as proof that MindCraft already is an apprenticeship CoP.

HYPOTHESIS: Treat LPP as a **design aspiration for tutor+practice sessions** (peripheral contribution → fuller solo responsibility) and a **diagnostic for fake community**, not as a claim that the app is a CoP.

Kill: Universal “students learn by LPP here.” Survive: “Usual school math often blocks LPP; we design for legitimate attempts at the edge with old-timer criteria.”

LII.6 What “practice” means for MindCraft (not vibes)

Align with Constitution practices already named:

1. **Diagnose** → **attempt** → **informative miss** → **revise** → **solo check** (FEI / SAFE-DP / SAFE-MISCON).
2. **Student-generated why before AI wrap** (PWC / XL).
3. **Named transfer hops** (SAFE-TRANSFER).
4. **Calibrated confidence talk** (SAFE-CALIB) — saying “I’m sure” as part of the repertoire, not a Belief Score™.
5. **Story worlds as wrappers of frozen math** (equity + meaning) — characters can model *peripheral roles* (watcher → helper → solver) without claiming ethnographic authenticity.

SPECULATION (to test): Peer “Community” features help only when they circulate the repertoire above (how we check; how we destake wrongs). Social chat about school stress without practice criteria may raise belonging feelings without changing mathematical participation (wound for growth loops).

LII.7 Competitive positioning

Competitor pattern	CoP-shaped risk	MindCraft wedge
Duo leagues / streaks	Belonging as rank theater (SAFE-COMPARE)	Criterion competence, not named infinite ranks
Khan comments / help forums	Peripheral <i>talk about</i> problems ≠ LPP in craft	Tutor/old-timer + attempt-first
Brilliant delight social	Shared puzzles without recovery repertoire	Soft-wrong + solo transfer as membership proof
ChatGPT “study buddy”	Fluency without legitimacy criteria	PWC + ontology spine; AI ≠ old-timer by default
Discord “study community”	CoP brand, acquisition optional	Enterprise + repertoire required for CoP copy

Commercial implication: Do not out-community Discord. Sell *becoming someone who participates differently in hard math* — with evidence (solo_transfer_pass, method-change after miss), not member counts.

LII.8 SAFE-CoP stack (product law)

1. **Name the practice** — joint enterprise in copy and UX (“how we recover and prove on this edge”), not “join our community.”
2. **Peripheral must be real work** — attempt, partial solution, check, SE prompt; watching/scrolling is not LPP.
3. **Old-timer required for CoP claims** — human tutor or SE-constrained coach with criteria; unguarded AI chat ≠ apprenticeship.
4. **Legitimacy without trait labels** — access to roles by task readiness; ban Golem scripts (SAFE-EXPECTANCY).
5. **Trajectory > tab** — measure changing participation (solo share up; help faded) not Community page visits.
6. **Plural metaphors** — keep acquisition instruments (graph, bank); never replace them with belonging KPIs (Sfard).
7. **Equity of membership** — whose history/stories define the repertoire (Part XXXVI); token “community” without belonging audit is a kill.
8. **Claim ladder** — CoP/LPP language max L1 product bet until CoP-* experiments; ban L3 “identity via community features” without instruments.

LII.9 Claim ladder

Claim	Max ladder without new data
LPP describes apprenticeship-style learning in craft/mutual-aid cases	L1 (Lave & Wenger, 1991)
Usual school math often fails LPP conditions; transfer is non-trivial	L1 (Boylan, 2010; Fuller et al., 2005)
Acquisition + participation metaphors both needed; one alone distorts	L1 (Sfard, 1998)
MindCraft tutor sessions can implement LPP-shaped trajectories	L1 product bet — needs CoP experiments
Forums / streaks / AI chat = CoP / LPP	Banned
Shipping Community features creates mathematical identity	Banned L3 without instruments

LII.10 Experiments spawned

ID	Question	Design	Primary	Kill condition
CoP-1	LPP-shaped session (peripheral contribution + old-timer criteria + fade to solo) vs explain-first homework help	2-arm	solo_transfer_pass; coded participation moves	Help ≥ LPP spine → CoP framing overstated

CoP-2	Named repertoire onboarding (“how we check here”) vs feature tour	2-arm	next-session method language; retry_120s	Tour ≥ repertoire → practice-naming useless
CoP-3	Human tutor old-timer vs unguarded AI buddy (same time)	2-arm	solo transfer; overtrust flags (XXXIII)	AI ≥ tutor on solo → old-timer claim weak for this segment
CoP-4	Community chat with repertoire prompts vs open social chat	2-arm	practice-talk rate; delayed accuracy	Open ≥ prompted → repertoire constraint wrong
CoP-5	Story peripheral roles (watcher→helper→solver) vs flat problem list	2-arm	challenge_accept ; belonging items; transfer	Flat ≥ roles on learning → role trajectory theater
CoP-QUAL	10 Maya + 5 tutor: “who is in your math people?”	interview	coded: old-timer / repertoire / legitimacy	“Community” = friends only, no practice → social tab demote

Pre-reg (XXXIV): CoP-* identify **participation design** / **old-timer necessity** / **fake-community kill** — not “CoP causes identity” or “LPP is universal in apps.”

LII.11 Confidence table

Claim	Label	Confidence
Lave & Wenger (1991): LPP as central descriptor of situated learning in apprenticeship cases	FACT	High
Wenger (1998): CoP as meaning –practice–community–identity	FACT	High
Boylan (2010): LPP transfer to usual school math is problematic; ecologies of participation	FACT	High
Fuller et al. (2005): LPP limited in complex institutions	FACT	High
Sfard (1998): do not choose only acquisition or only participation	FACT	High

SAFE-CoP is right default for MindCraft social/tutor design	FOUNDER BELIEF / HYPOTHESIS	Medium
Community features alone transform math identity	SPECULATION / false as doctrine	Low (against)

LII.12 What this chapter kills

1. **Kill:** “We are a community of practice” as brand without enterprise + repertoire + old-timer.
2. **Kill:** Forums, Discord, streaks, leagues, or chat volume as LPP proof.
3. **Kill:** AI monologue $\equiv$ apprenticeship / old-timer.
4. **Kill:** Treating school/ACT/app sessions as ethnographically equivalent to Lave & Wenger’s craft cases.
5. **Wound:** Participation-only pedagogy that deletes mastery/transfer instruments (Sfard).
6. **Survive:** SAFE-CoP; LPP-shaped tutor spine; peripheral = real work; trajectory metrics; CoP-1...5; “join the practice of recovering and proving” copy — not “join the community.”

Doctrine until data: Ship **SAFE-CoP** — name the practice, require real peripheral work and old-timer criteria, measure participation trajectories under acquisition+participation pluralism — never sell social tabs or CoP labels as identity engines.

Part LIII — Religion / Ritual Light-Touch Design: Meaning Without Cult Dynamics

Chapter status: Living evidence brief — Researcher tick 2026-07-31 (UTC hour 15; hour%6≠0; 3 researcher entries since synthesizer v1.7 → Researcher)

Primary question: How can MindCraft design *light-touch rituals* that mark transitions, regulate affect, and carry meaning — without drifting into cult dynamics, sacred-science branding, or high-arousal initiation theater?

Owners: Product (session / world / onboarding arcs) · Brand & copy · Equity / trust · Red Team

Commercial job: Ship a **SAFE-RITUAL** stack — brief, opt-in, secular, reversible meaning-markers tied to real practice — while killing “join the order,” streak liturgies, confession milieus, and unretracted ritual→grade overclaims.

LIII.1 Why this chapter exists

The Constitution’s thesis is identity transformation under difficulty. Humans mark identity shifts with *rituals*: greetings before hard work, closing circles after miss→revise, naming a first solo transfer. Edtech and consumer apps already smuggle ritual language — streaks as devotion, leagues as congregations, onboarding as initiation, founder myth as sacred science — usually without naming the load-bearing psychology or the failure modes.

MindCraft has natural ritual-shaped moments: gap-scan completion, first soft-wrong recovery, tutor session open/close, story-world threshold crossings, Map unlocks. Parts XXIV (affect), XXV (identity), XXXI (flow), XLIII (SAFE-HABIT), XLIV (SAFE-REWARD), XLV (resilience), LII (SAFE-CoP) constrain pieces. What was missing was an honest doctrine that (a) uses ritual science for *light* meaning, (b) borrows transition structure without religion cosplay, and (c) installs a cult-dynamics kill list so “meaning” never becomes totalism.

FOUNDER BELIEF under audit: Becoming a confident mathematical thinker needs *marked moments* (not only item completion) — but the company must never sell belonging-through-devotion, milieus that punish exit, or doctrine that overrides the student’s evidence.

Claims we refuse as doctrine:

1. Rituals (or “sacred practice”) as proven grade engines from unretracted lab theater.
2. Streak / league / XP liturgy as identity North Star.
3. High-arousal initiation, confession milieus, or “you’re not one of us until...” gates.
4. Founder/product doctrine as sacred science that cannot be falsified (contradicts the lab itself).

LIII.2 Constructs (product language)

Construct	Plain definition	Product signal	Failure mode if misused
Ritual (psych)	Predefined, symbolic, often repetitive sequence that is <i>more than</i> its instrumental steps (Hobson et al., 2018)	Brief open/close; naming a miss class; candle-free “we begin”	Calling every UX animation a ritual
Routine / habit	Repeated behavior cue→response without required symbolic meaning (SAFE-HABIT)	If-then practice cue	Relabeling habit as sacred
Rite of passage	Transition marked by separation → liminal → reaggregation (van Gennep; Turner)	Gap-scan → liminal “not yet rated” → diagnosticCompleted + first path	Initiation hazing; permanent limbo
Imagistic mode	Infrequent, high-arousal rituals → vivid episodic memory; small exclusive groups (Whitehouse, 2004)	Rare world “boss” moments	Trauma theater; faction bonding over learning

Doctrinal mode	Frequent, low-arousal repetition → semantic/procedural memory; larger inclusive scale (Whitehouse, 2004)	Soft-wrong language; Map repertoire; session grammar	Indoctrination scripts; empty liturgy
Totalism / cult dynamics	Thought-reform pattern: milieu control, confession cult, sacred science, loaded language, doctrine-over-person, exit costs (Lifton)	Anti-patterns in copy, retention, community	“Family,” purity tests, unfalsifiable brand
SAFE-RITUAL	Light-touch ritual doctrine that survives cult kills + retraction wounds	See LIII.8	Sacred streak; initiation cosplay

Operational definition (HYPOTHESIS): A MindCraft moment counts as *SAFE-RITUAL* when (a) it is brief and optional or easily skippable, (b) symbolism points at *practice repertoire* (recover, prove, calibrate) not brand devotion, (c) exit/skip has no status punishment, (d) no confession or purity demand, (e) claims stay on the ladder (affect/transition aid ≠ causal identity engine).

LIII.3 Psychology of ritual — regulatory functions (not magic grades)

FACT (integrative review): Hobson, Schroeder, Risen, Xygalatas & Inzlicht (2018, *Personality and Social Psychology Review*, 22(3), 260–284, doi:10.1177/1088868317734944) — organize empirical ritual work around three regulatory functions: **(a) emotion regulation, (b) performance goal states, (c) social connection**. Mechanisms include bottom-up features (repetition, formality, synchrony) and top-down meaning (belief that the sequence *is* a ritual / carries symbolic significance).

FACT (grief / control): Norton & Gino (2014, *Journal of Experimental Psychology: General*, 143(1), 266–272, doi:10.1037/a0031772) — rituals after losses (loved ones, relationships, lotteries) can reduce grief feelings and increase felt control relative to no-ritual controls. Useful as evidence that symbolic sequences can regulate emotion and agency — **not** as a math-learning RCT.

FACT (consumption / meaning frame): Vohs, Wang, Gino & Norton (2013, *Psychological Science*, 24(9), 1714–1721, doi:10.1177/0956797613478949) — ritualized actions can enhance consumption experience via greater involvement/savoring. Ally for “mark the moment” design; kill if used to justify gamified consumption of XP as learning.

WOUND (retraction — do not market): Brooks, Schroeder, Risen, Gino, Galinsky, Norton & Schweitzer (2016, *Organizational Behavior and Human Decision Processes*, 137, 71–85, doi:10.1016/j.obhdp.2016.07.004) — claimed pre-performance rituals reduce anxiety and improve performance (including high-stakes math-like tasks). **This article has been retracted.** MindCraft must **not** cite it as surviving evidence or sell “rituals raise exam scores.” Treat as a Red Team kill on ritual→grade overclaim.

HYPOTHESIS for MindCraft: Prefer Hobson’s *emotion regulation* and *goal-state* functions for session openers/closers and soft-wrong recovery markers; treat performance/grade claims as **banned** until unretracted, pre-registered evidence on MindCraft metrics (retry_120s, solo_transfer_pass, affect check-ins).

Kill: “Science proves rituals improve math performance.” Survive: “Light symbolic sequences can regulate affect and mark transitions; grade claims need new evidence.”

LIII.4 Modes of religiosity — why high-arousal initiation is the wrong default

FACT: Whitehouse (2004, *Modes of Religiosity: A Cognitive Theory of Religious Transmission*, AltaMira Press / Rowman & Littlefield, ISBN 978-0-7591-0615-4) — religions tend to coalesce around **imagistic** vs **doctrinal** poles. Imagistic: infrequent, high-arousal rituals → flashbulb-like episodic memory; small, exclusive, often intense group bonds. Doctrinal: frequent, low-arousal, routinized teaching/ritual → semantic and procedural memory; larger, more inclusive, centrally scaffolded communities.

FACT (extreme ritual ally — careful): Xygalatas, Mitkidis, Fischer, Reddish, Skewes, Geertz et al. (2013, *Psychological Science*, 24(8), 1602–1605, doi:10.1177/0956797612472910) — extreme rituals can increase prosociality in field settings. This is **not** a product recommendation for teens in ACT prep; it shows imagistic arousal can bond groups — which is exactly the risk surface for cult-adjacent design.

HYPOTHESIS: MindCraft’s default should be **doctrinal-light**: frequent, low-arousal repertoire (session grammar, soft-wrong language, Map names) — not imagistic initiation (ordeal unlocks, shame fires, exclusive survivor factions). Rare story “threshold” moments may use mild imagistic *flavor* only if destaked, optional, and never required for belonging (SAFE-DD / equity).

Product translation: Gap-scan and first recovery are transitions, not hazing. World chapters can feel special without pain-as-proof.

Kill: Bootcamp/ordeal branding; “you haven’t really suffered until...” Survive: Frequent low-arousal practice markers + rare optional celebration.

LIII.5 Rites of passage — structure without cult costume

FACT: van Gennep (1909; English trans. 1960, *The Rites of Passage*, University of Chicago Press) — life transitions commonly share three phases: **separation**, **margin/limen**, **aggregation** (reincorporation).

FACT: Turner (1969, *The Ritual Process*; see also “Liminality and Communitas”) — elaborates the liminal phase as *betwixt and between*: ambiguous status, stripped markers, intense potential for *communitas* (anti-structure bonding). Powerful analytically; dangerous if product designers chase liminality as engagement hack (keep students permanently “in between,” dependent on the brand for status).

HYPOTHESIS for MindCraft transitions:

Phase	Product move	Guardrail
-------	--------------	-----------

Separation	Leave “I don’t know where I stand” → enter gap-scan / session	No shame script; task language only (SAFE-EXPECTANCY)
Liminal	Hide-correctness diagnostic; soft-wrong; “not yet aggregated”	Time-bounded; no status mockery; skippable ceremony
Aggregation	diagnosticCompleted; first path; named repertoire (“how we recover here”)	Grant competence access, not trait identity; no purity test

Link to SAFE-CoP: Aggregation should admit the student to *practice roles* (peripheral → fuller), not to a brand congregation.

Kill: Permanent liminality as retention. Survive: Clear exits into rights/obligations of the practice (you may attempt this edge solo).

LIII.6 Cult dynamics — the kill list (Lifton as anti-pattern checklist)

FACT (classic totalitarian criteria): Lifton (1961, *Thought Reform and the Psychology of Totalism*; later editions restate cult criteria) — thought-reform / totalistic environments characteristically include patterns such as: **milieu control, mystical manipulation, demand for purity, cult of confession, sacred science, loading the language, doctrine over person, dispensing of existence** (who counts as real/valid). Cult-like groups often combine charismatic leader-focus, thought-reform patterns, and exploitation of idealism.

FOUNDER BELIEF: MindCraft’s research lab and product ethics require the *opposite* of sacred science: every doctrine chapter is falsifiable; Red Team is a feature; students may leave without metaphysical damnation.

Product anti-patterns (SPECULATION → treat as design law until wounded):

Lifton pattern	Edtech failure mode	MindCraft ban
Milieu control	App as only legitimate math space; shame offline help	Allow tutors, school, parents; no “only here” purity
Mystical manipulation	Streak miracles; “the algorithm chose you” destiny copy	Transparent recommend reasons; no fate talk
Demand for purity	Perfect streaks; zero-miss identity	Soft-wrong; SAFE-MISCON; miss as data
Cult of confession	Forced public failure theater / oversharing	Private by default; opt-in sharing
Sacred science	Unfalsifiable brand claims; ban critique	Claim ladder; lab kills; no 2-sigma marketing
Loading the language	In-group jargon that replaces thinking	Repertoire words OK; ban thought-stopping slogans
Doctrine over person	Override student evidence with brand story	Graph + outcomes over vibes; SAFE-CALIB

Dispensing of existence	Rank/league as who “counts”	SAFE-COMPARE; no exile UX
-------------------------	-----------------------------	---------------------------

Kill: “We’re a movement / family / order.” Survive: “We’re a practice with optional meaning-markers and exit dignity.”

LIII.7 Competitive audit — ritual theater in the market

Duo streaks/leagues = doctrinal liturgy + mild imagistic competition (habit colonization; cue habits without sacred streak NS). Khan mastery energy = completion ritual (keep instruments; add recovery markers). Brilliant “aha” = imagistic spark (celebrate then probe transfer). GPT chat = pseudo-confession / endless liminality (SE-constrained close). Studygram hustle = purity + confession (ban grind-as-virtue).

Commercial implication: Position as **meaning without membership taxes** — mark *mathematical practice*, not brand devotion. Copy: “mark the recovery,” “open like a craft,” “close with what changed” — never “join the faithful.”

LIII.8 SAFE-RITUAL stack (product law)

1. **Light and short** — seconds to a minute; never a sermon that displaces attempts (SAFE-DP time budget).
2. **Optional / skippable** — no status loss for skip; no streak-as-morality.
3. **Point at practice** — symbolize recover / prove / calibrate / transfer — not founder, brand, or league.
4. **Doctrinal default** — frequent low-arousal repertoire; rare celebration only, never ordeal.
5. **Bounded liminality** — transitions resolve into clear rights (you may try X); no permanent betwixt for retention.
6. **Anti-totalism** — ban milieu monopoly, confession theater, sacred science, purity, exit shame.
7. **Affect ≠ grade claim** — may aim at calm / agency / belonging-to-practice; **no** unretracted ritual→score marketing.
8. **Equity** — no borrowed sacred forms that mock students’ real religions; secular by default (Part XXXVI).
9. **Claim ladder** — ritual language max L1 product bet until RIT-* experiments; ban L3 identity-via-ceremony without instruments.

LIII.9 Claim ladder

Claim	Max ladder without new data
Rituals can regulate emotion / goal states / connection (mechanism review)	L1 (Hobson et al., 2018)
Symbolic sequences can increase felt control after loss	L1 (Norton & Gino, 2014)
Imagistic vs doctrinal modes describe divergent ritual ecologies	L1 (Whitehouse, 2004)

Rites of passage often share separation–liminal–aggregation	L1 (van Gennep; Turner)
Lifton patterns are useful anti-design checks for edtech	L1 design heuristic (not a clinical diagnosis of competitors)
MindCraft light rituals improve <code>retry_120s</code> / affect / solo transfer	L1 product bet — needs RIT experiments
Rituals raise ACT scores / prove identity transformation	Banned without unretracted evidence
Streak liturgy / ordeal initiation / “family” milieus as identity engines	Banned

LIII.10 Experiments spawned

ID	Question	Design	Primary	Kill condition
RIT-1	30s session open ritual (name edge + breath/pose optional) vs silent start	2-arm	pre-task affect; <code>challenge_accept</code>	Silent $\geq$ ritual on start+transfer $\rightarrow$ opener useless
RIT-2	Soft-wrong recovery marker (“name the miss class”) vs correct/incorrect only	2-arm	<code>retry_120s</code> ; method-change code	Binary $\geq$ marker $\rightarrow$ recovery ritual theater
RIT-3	Gap-scan aggregation ceremony (brief) vs auto-redirect no mark	2-arm	next-session return; belonging -to-practice item	Auto $\geq$ ceremony $\rightarrow$ skip ceremony permanently
RIT-4	Optional celebration on first <code>solo_transfer_pass</code> vs none	2-arm	delayed transfer; skip rate	Forced $>$ optional engagement but worse affect $\rightarrow$ force banned
RIT-5	Streak-framed cue vs practice-repertoire cue (same schedule)	2-arm	4-week return; identity items; anxiety	Streak $\geq$ repertoire on learning $\rightarrow$ still ban streak as NS; copy only
RIT-QUAL	10 Maya: which moments felt meaningful vs creepy/cultish?	interview	coded: light / pressure / exit fear	“Family/order” language $\rightarrow$ kill that copy

Pre-reg (XXXIV): RIT-* identify **transition/affect design and anti-totalism** — not “ritual causes identity” or “rituals raise scores.”

LIII.11 Confidence table

Claim	Label	Confidence
Hobson et al. (2018): three regulatory functions of ritual	FACT	High
Norton & Gino (2014): rituals can ease grief / raise felt control	FACT	High
Brooks et al. (2016) ritual→performance paper	FACT it was published; retracted — do not use as evidence	High (against use)
Whitehouse (2004): imagistic vs doctrinal modes	FACT	High
van Gennep / Turner: passage / liminality structure	FACT	High
Lifton totalism criteria as edtech anti-patterns	HYPOTHESIS / design heuristic	Medium–High
SAFE-RITUAL is right default for MindCraft meaning UX	FOUNDER BELIEF / HYPOTHESIS	Medium
Ceremony alone transforms math identity	SPECULATION / false as doctrine	Low (against)

LIII.12 What this chapter kills

1. **Kill:** Ritual→exam-score marketing (especially via retracted Brooks et al., 2016).
2. **Kill:** Streak / league / XP as sacred liturgy or North Star.
3. **Kill:** High-arousal initiation, confession milieus, purity tests, exit shame.
4. **Kill:** “Movement / family / order” copy that loads Lifton patterns.
5. **Wound:** Permanent liminality-as-retention; imagistic faction bonding.
6. **Survive:** SAFE-RITUAL; doctrinal-light repertoire; bounded passages; anti-totalism checklist; RIT-1...5; “mark the recovery” copy — not “join the faithful.”

Doctrine until data: Ship **SAFE-RITUAL** — brief, skippable, practice-pointing meaning-markers under an explicit cult-dynamics ban — never sell devotion, ordeal, or unretracted ritual-performance claims as the identity engine.

Part LIV — Military / Aviation Brief–Debrief: After-Action Reviews for Math Sessions

Chapter status: Living evidence brief — Researcher tick 2026-07-31 (UTC hour 18 ≡ Red Team slot, but ch54 never written → prefer Researcher per rotation; 4 researcher entries since synthesizer v1.7 → Researcher)

Primary question: What transferable structure do military after-action reviews and aviation brief/debrief cycles offer for MindCraft tutoring and practice sessions — without militaristic cosplay, blame culture, or overclaiming effect sizes as ACT marketing?

Owners: Product (session spine / Notes close) · Tutor ops · Coach UX · Brand & copy · Red Team

Commercial job: Ship a **SAFE-AAR** stack — brief → attempt → structured debrief with objective traces — while killing bootcamp branding, lecture-as-debrief, trauma-debrief confusion, and “debriefs improve performance 25%” as an undifferentiated growth claim.

LIV.1 Why this chapter exists

Parts LI (SAFE-DP) and LIII (SAFE-RITUAL) already constrain *time budgets* and *light meaning markers*. What was missing is the **operational cycle** high-reliability domains use to convert experience into the next attempt: a short **brief** that sets intent and roles, an event under observation, and a **debrief / after-action review (AAR)** that compares expected vs actual performance without collapsing into critique-as-humiliation.

MindCraft already has the raw materials: session open (affect check-in), attempt traces, soft-wrong codes, outcomes posted to the graph, Notes/session summaries, and coach cards. Competitors either skip structured reflection (answer-delivery chat), fake it with streak liturgy (Duo), or replace it with instructor monologue (many tutoring products and AI wrap text). The commercial wedge is not “we’re like the military.” It is: **every hard attempt leaves a learning artifact the student helps interpret.**

FOUNDER BELIEF under audit: Identity transformation under difficulty requires *reviewed attempts*, not only more problems — but review must be non-punitive, student-participatory, and tied to the next edge (SAFE-DP), not to brand devotion (SAFE-RITUAL) or trait shame (SAFE-EXPECTANCY).

Claims we refuse as doctrine:

1. Militaristic / “warrior math” branding as belonging technology.
 2. AAR ≡ critical-incident stress debriefing (trauma processing) — different animal.
 3. Marketing “+25% performance” from debrief meta-analyses as if it were an ACT score guarantee.
 4. Instructor/AI lecture that *calls itself* a debrief while blocking student self-analysis.
-

LIV.2 Constructs (product language)

Construct	Plain definition	Product signal	Failure mode if misused
-----------	------------------	----------------	-------------------------

Brief	Short pre-event alignment on intent, success criteria, roles, likely traps (CRM / FAA markers)	Session open: target edge, level, “what counts as success,” optional coach role	10-min sermon; fate talk; trait labels
After-action review (AAR)	Systematic technique turning a recent event into learning via feedback, reflection, discussion (Keiser & Arthur)	Post-mission 4-question close; tutor end-of-session AAR	Critique lecture; blame; vibes-only
Debrief (training)	Structured reflection after simulation/performance; often facilitator-guided	Coach/tutor facilitation; self-led when traces exist	Endless chat liminality; AI monologue
Objective performance media	Records of what happened (video, logs, instruments) — not memory alone	Attempt list, soft-wrong codes, timing, outcomes JSON	Relying on “how did it feel?” only
Facilitation (vs lecture)	Elicit crew/student self-analysis; deepen participation (Dismukes et al.)	Advocacy–inquiry; student names frame first	Tutor/AI tells the moral of the story
Good judgment stance	Share evaluative judgment + curiosity about learner’s frames (Rudolph et al.)	Soft-wrong + “what were you seeing?”	Fake nonjudgmental praise; or blame
Alignment	Match participants ↔ focus ↔ measurement (team vs individual)	Solo practice AAR on solo metrics; tutor+student AAR on shared plan	Team vibes measured by individual XP
SAFE-AAR	Brief–attempt–debrief doctrine that survives militarism / blame / overclaim kills	See LIV.8	Bootcamp cosplay; 25%-as-marketing

Operational definition (HYPOTHESIS): A MindCraft close counts as *SAFE-AAR* when (a) a pre-brief named the intended edge and success criterion, (b) the debrief answers expected / actual / sustain / improve using **at least one objective trace**, (c) the student produces ≥ 1 self-analytic statement before any AI wrap, (d) the output is a next-attempt commitment (not a grade sermon), (e) no trait labels, no punishment for miss, no militaristic belonging tax.

LIV.3 Army AAR doctrine — four questions, not a roast

FACT (doctrine lineage): U.S. Army training circular **TC 25-20** (*A Leader's Guide to After Action Reviews*, 30 Sep 1993) and subsequent training doctrine institutionalize the AAR as a professional discussion focused on training objectives — not a humiliation ritual. Contemporary restatements (e.g., Center for Army Lessons Learned handbooks; USAID AAR guidance adapted from TC 25-20) center four questions:

1. What was **expected** to happen?
2. What **actually** happened?
3. What went **well**, and why?
4. What can be **improved**, and how?

FACT (process intent): Formal Army AAR guidance emphasizes open-ended questioning so participants discover gaps vs standards; facilitators must avoid critique-as-lecture (FM 25-101 Appendix G / CALL lineage). Focus: **results vs objectives**, sustain strengths, actionable fixes.

HYPOTHESIS for MindCraft: Map onto a 60–120s practice close and a 3–5 min tutor close — brief names expected edge; traces show actual; student names one sustain; one improve-how (SAFE-HABIT if-then or method change). Kill “AAR = public roast.” Survive: standards-referenced discovery.

LIV.4 Meta-analytic evidence — real gains, careful transfer

FACT: Tannenbaum & Cerasoli (2013, *Human Factors*, 55(1), 231–245, doi:10.1177/0018720812448394) — meta-analysis of 46 samples ($N = 2,136$). Debriefs / AARs improve effectiveness vs control by about **25% on average** ($d \approx 0.67$). Effects similar for teams and individuals, simulated and real settings, medical and nonmedical samples. **Alignment** (participants, intent, measurement level) bolsters effects; facilitated and structured debriefs trend stronger (facilitation estimate based on fewer nonfacilitated studies).

FACT: Keiser & Arthur (2021, *Journal of Applied Psychology*, 106(7), 1007–1032; advance online 2020, doi:10.1037/apl0000821) — bare-bones meta-analysis of 61 studies (107 *ds*; 915 teams and 3,499 individuals). Overall AAR effect $d \approx 0.79$ across multiple training criteria — larger than Tannenbaum & Cerasoli's task-performance-focused estimate. Two characteristics consistently help: **(a) alignment** to individual vs team, **(b) objective performance review media**. Facilitation interacts: self-led + team-aligned and self-led + objective media among stronger combinations; structure matters more in military than healthcare contexts.

HYPOTHESIS (transfer wound): These samples are largely military, healthcare simulation, and organizational training — **not** ACT prep for anxious Maya. Treat $d \approx 0.67$ – 0.79 as evidence that structured post-event reflection is a high-leverage training design pattern — **not** license to claim “MindCraft AARs raise SAT/ACT by 25%.”

Commercial claim ladder:

- L1: Debriefs/AARs improve training outcomes in reviewed domains (cite metas).
- L1 product bet: SAFE-AAR improves MindCraft instruments (retry_120s, method-change codes, solo_transfer_pass).
- **Banned:** Undifferentiated “science says +25%” on landing pages without domain caveat and MindCraft pre-reg.

Kill: 2-sigma-style overclaim via borrowed meta d . Survive: “Structured brief–debrief with traces beats vibes-only close.”

LIV.5 Aviation — brief as error management; debrief as facilitated self-analysis

FACT (CRM briefing as performance marker): FAA Advisory Circular **AC 120-51E** (*Crew Resource Management Training*) lists effective **briefings** among CRM behavioral markers: thorough, address coordination/planning/problems, establish open communication, set expectations for SOP deviations and automation duties. Briefings are error-management tools — they align the crew *before* the event, not after the crash.

FACT (CRM evolution): Helmreich, Merritt & Wilhelm (1999, *International Journal of Aviation Psychology*, 9(1), 19–32; see also Helmreich line of NASA/UT work) — CRM evolved toward error management: avoid, trap, mitigate. Briefing strategies and situation awareness modules are framed as countermeasures, not soft skills theater.

FACT (LOFT debrief facilitation): Dismukes, Jobe & McDonnell (1997, NASA TM-110442, *LOFT Debriefings: An Analysis of Instructor Techniques and Crew Participation*) — analyzed 36 LOFT debriefings at five U.S. airlines; developed the Debriefing Assessment Battery. **Facilitation** (vs instructor lecture) increases depth of crew participation and self-analysis of CRM performance; instructor skill varied dramatically; crews were responsive but rarely fully led their own debriefs. Companion: McDonnell, Jobe & Dismukes (1997, NASA TM-112192, *Facilitating LOS Debriefings: A Training Manual*).

HYPOTHESIS for MindCraft:

- **Brief** = gap-scan / mission card / tutor open: name the edge, success criterion, known misconception trap, role of coach (hints vs silence).
- **Debrief** = student-led analysis with tutor/coach as facilitator; AI may *prompt* advocacy–inquiry, must not deliver the sermon first (aligns with Part XL SE constraints and Part XXXIII AI trust).

Product translation: Tutor briefing ethics from SAFE-EXPECTANCY (task-only language) + CRM brief markers (coordination, traps, automation = “when does the coach speak?”).

Kill: Captain-monologue onboarding. Survive: short interactive brief + facilitated close.

LIV.6 Healthcare simulation — good judgment, not fake nonjudgment

FACT: Rudolph, Simon, Dufresne & Raemer (2006, *Simulation in Healthcare*, 1(1), 49–55, doi:10.1097/01266021-200600110-00006; PubMed 19088574) — “There’s no such thing as ‘nonjudgmental’ debriefing.” **Debriefing with good judgment:** (1) frames drive actions → results; (2) stance of genuine curiosity (mistakes as puzzles); (3) **advocacy–inquiry** pairs (state observation + judgment, then ask what the learner’s frame was). Helps instructors share evaluative judgment without destroying trust.

Link to SAFE-MISCON / SAFE-CALIB: Soft-wrong is not empty kindness; it is destaked judgment plus frame inquiry. Hide-correctness diagnostic is a special brief (no outcome reveal) — its “debrief” must not smuggle keys that corrupt the graph (C4).

HYPOTHESIS: MindCraft coach cards should default to advocacy–inquiry templates (“I notice you treated the inequality like an equation when you multiplied by -1 — what were you assuming?”) rather than praise-only or scold-only.

Kill: “Nonjudgmental forever” (hides standards) and “brutal honesty” (hides curiosity). Survive: good-judgment debrief.

LIV.7 Distinctions and kill list (scope control)

Do not confuse training AARs with **critical-incident stress debriefing** (trauma processing — different ethics). Ban militaristic belonging (“warrior,” bootcamp shame), streak-as-AAR, AI monologue labeled as debrief, vibes-only hot-wash when traces exist, public leaderboard “reviews” (SAFE-COMPARE), and “+25% smarter” ads from borrowed meta *d*.

Wound: Hot-wash alone under-captures learning — useful as phase 1, insufficient as the whole product. MindCraft default: immediate 60s close + optional next-day Notes deepening.

LIV.8 SAFE-AAR stack (product law)

1. **Brief before attempt** — name edge, success criterion, trap, coach role (CRM brief markers).
 2. **Attempt under observation** — log objective media (outcomes, codes, timing).
 3. **Four-question close** — expected / actual / sustain / improve (TC 25-20 lineage).
 4. **Student speaks first** — ≥ 1 self-analytic statement before tutor/AI wrap (Dismukes facilitation).
 5. **Good judgment** — advocacy-inquiry; standards visible; curiosity about frames (Rudolph).
 6. **Align level** — solo metrics for solo practice; shared plan metrics for tutor sessions (Tannenbaum; Keiser).
 7. **Next-attempt output** — one concrete improve-how tied to SAFE-DP time budget.
 8. **Non-punitive climate** — miss = data; no trait labels (SAFE-EXPECTANCY); no militaristic shame.
 9. **Claim ladder** — meta *d* stays L1 domain evidence; MindCraft ACT claims need AAR-* experiments.
 10. **Anti-cosplay** — borrow the *cycle*, not the costume; secular practice language (SAFE-RITUAL).
-

LIV.9 Competitive audit — who actually closes the loop?

Khan tells correctness without student-led expected/actual. Duo substitutes streak liturgy for AAR. Brilliant celebrates aha but rarely closes sustain/improve. ChatGPT tutors lecture after chatty “goals.” Human tutoring quality varies with facilitator skill (Dismukes variance). MindCraft already logs graph outcomes — formalize SAFE-AAR in practice close + Notes.

Commercial implication: Position as **reviewed struggle** — brief the edge, attempt with traces, close expected/actual — never “military-grade” or borrowed “+25%” miracles. Parent copy: expected vs happened → what changes next. Student copy: “name what you thought would happen.”

LIV.10 Claim ladder

Claim	Max ladder without new data
-------	-----------------------------

AARs/debriefs improve training outcomes in meta-analyzed domains ($d \approx 0.67\text{--}0.79$)	L1 (Tannenbaum & Cerasoli, 2013; Keiser & Arthur, 2021)
Four-question AAR structure is institutionalized Army training practice	L1 (TC 25-20 lineage / CALL restatements)
Facilitation deepens self-analysis vs lecture in LOFT debriefs	L1 (Dismukes et al., 1997)
CRM briefings are performance markers for coordination / error management	L1 (FAA AC 120-51E; Helmreich line)
Good-judgment advocacy–inquiry is a validated simulation debrief stance	L1 method (Rudolph et al., 2006)
SAFE-AAR improves MindCraft <code>retry_120s</code> / method-change / solo transfer	L1 product bet — needs AAR experiments
“+25%” / meta d as ACT or identity guarantee	Banned
Bootcamp / warrior branding as belonging engine	Banned
AAR $\equiv$ trauma CISD	Banned conflation

LIV.11 Experiments spawned

ID	Question	Design	Primary	Kill condition
AAR-1	4-question practice close vs “nice job / try again” toast	2-arm	<code>retry_120s</code> ; method-change code rate	Toast $\geq$ AAR on retry+transfer $\rightarrow$ close theater
AAR-2	Brief (edge+success+trap) vs no brief, same items	2-arm	<code>challenge_accept</code> ; first-attempt process quality	No-brief $\geq$ brief $\rightarrow$ brief bloat; shorten or cut
AAR-3	Objective-trace AAR vs memory-only AAR	2-arm	accuracy of self-report; next-session transfer	Memory $\geq$ trace $\rightarrow$ still keep traces for graph; copy only
AAR-4	Student-first facilitation vs AI-monologue wrap	2-arm	SE quality rubric; <code>solo_transfer_pass</code>	Monologue $\geq$ student-first $\rightarrow$ contradicts XL; kill monologue default
AAR-5	Tutor trained on advocacy–inquiry vs business-as-usual close	2-arm cluster	session notes quality; student affect; 2-week transfer	BAU $\geq$ trained $\rightarrow$ training content fail

AAR-QUAL	10 Maya + 5 tutors: which closes felt useful vs military/shamey?	interview	coded: useful / cosplay / blame	“Bootcamp” language → kill that copy
----------	--	-----------	---------------------------------	--------------------------------------

Pre-reg (XXXIV): AAR-* identify **session architecture** (brief–debrief design) — not “military methods raise ACT” or identity-from-discipline myths.

LIV.12 Confidence table

Claim	Label	Confidence
Tannenbaum & Cerasoli (2013): debriefs $d \approx 0.67$	FACT	High
Keiser & Arthur (2021): AAR $d \approx 0.79$; alignment + objective media matter	FACT	High
TC 25-20 / Army AAR four-question core	FACT	High
Dismukes et al. (1997): facilitation → deeper crew self-analysis	FACT	High
Rudolph et al. (2006): good-judgment / advocacy–inquiry	FACT (method paper)	High
FAA CRM briefing markers	FACT (advisory doctrine)	High
Meta effects transfer cleanly to ACT prep for anxious teens	HYPOTHESIS / wounded	Low–Medium
SAFE-AAR is right default session spine for MindCraft	FOUNDER BELIEF / HYPOTHESIS	Medium
Military cosplay improves belonging or scores	SPECULATION / false as doctrine	Low (against)

LIV.13 What this chapter kills

1. **Kill:** “Military-grade” / bootcamp / warrior branding as identity or efficacy claim.
2. **Kill:** Marketing meta-analytic d / “+25%” as undifferentiated ACT or IQ promise.
3. **Kill:** Conflating training AARs with trauma CISM.
4. **Kill:** Lecture or AI monologue that labels itself a debrief while blocking student self-analysis.
5. **Kill:** Vibes-only closes without objective traces when traces exist.

6. **Wound:** Hot-wash-only products; facilitator skill variance (train tutors; constrain AI).

7. **Survive:** SAFE-AAR; four-question close; CRM-style brief; facilitation + good judgment; AAR-1...5; copy as **reviewed struggle** — not costume cosplay.

Doctrine until data: Ship **SAFE-AAR** — brief the edge, attempt with traces, close with expected/actual/sustain/improve under non-punitive facilitation — and never sell discipline theater or borrowed percentage miracles as the identity engine.

Part LV — Sports Film-Study Pedagogy: Error Clips → Coach Cards

Chapter status: Living evidence brief — Researcher tick 2026-07-31 (UTC hour 21; hour%6≠0; 5 researcher entries since synthesizer v1.7 → Researcher)

Primary question: What transferable design does sports film-study / video-based feedback (VBF) offer for MindCraft practice review — turning attempt traces into sparse, student-controlled “clips” and coach cards — without sports-cosplay branding, feedback floods, or highlight-reel shame?

Owners: Product (Solver / Notes / coach cards) · Tutor ops · Coach UX · Brand & copy · Red Team

Commercial job: Ship a **SAFE-FILM** stack — objective attempt clips + less-is-more cueing + learner-timed review + self-model pairing — while killing “we’re like a sports academy” marketing, monologue film rooms, and flood-of-tips as pedagogy.

LV.1 Why this chapter exists

Part LIV (SAFE-AAR) established the *cycle*: brief → attempt with traces → structured close. Sports film study is the **media layer** of that cycle — high-performance domains do not rely on memory alone; they rewatch selected moments, cue attention, and commit the next play. MindCraft already logs the analogue of game film: attempt lists, soft-wrong codes (mis_ IDs), timing, format tags, outcomes JSON. Competitors either never replay process (answer-key chat), replay only correctness (Khan-style), or flood the learner with AI commentary that no one remembers a week later.

FOUNDER BELIEF under audit: Identity as a mathematical thinker grows when the learner can *see their own decision frames* on a specific attempt — not when a brand dresses practice as a locker room. Film-study transfer is about **selective objective media + facilitation**, not jersey cosplay.

Claims we refuse as doctrine:

1. Sports-academy / “elite athlete” branding as belonging or efficacy technology.
 2. Coach monologue over 30 tips per “film room” (recall collapses).
 3. Highlight-reel of only successes *or* only failures as identity theater.
 4. Treating declarative recall of tips as if it were transfer to next-attempt performance.
 5. Equating any video UI with film-study pedagogy.
-

LV.2 Constructs (product language)

Construct	Plain definition	Product signal	Failure mode if misused
Film study / VBF	Structured review of performance via recorded moments + guided attention	Replay attempt; soft-wrong clip; coach card on <i>that</i> step	Passive highlight reel; shame montage
Performance analysis (PA)	Evaluate/describe performance with media (Pearson et al.)	Graph outcomes + process codes as “tape”	Stats without learner agency
Self-observation	Learner watches <i>own</i> prior attempt	“Replay your steps” before AI wrap	Narcissistic doom-scroll of errors
Skilled / expert model	Observation of a competent demonstration	Worked example / correct pathway card paired to same archetype	Model-only → “I can’t do that” gap
Self-modeling / feedforward	Edited view of self succeeding at a near-future challenge (Dowrick)	Show student’s own prior <i>successful</i> transfer on related edge	Fake montages that invent competence
Attentional cueing	Direct gaze/attention to the decision that mattered	Coach card highlights inequality sign flip / bridge gap	Cue everything = cue nothing
Self-controlled feedback	Learner chooses <i>when</i> to request video/KP review	“Review this attempt?” on demand vs forced after every item	Autonomy theater with no cue quality
SAFE-FILM	Film-study doctrine that survives sports cosplay / flood / shame kills	See LV.8	Academy branding; tip dumps

Operational definition (HYPOTHESIS): A MindCraft review counts as *SAFE-FILM* when (a) the clip is an **objective attempt trace** (not vibes), (b) ≤ 3 attentional cues / coaching points are attached, (c) the learner can request or skip replay (self-control), (d) when a miss is shown, a **paired skilled model or student’s own prior success** on a related edge is available (self-model), (e) the output is a next-attempt commitment (SAFE-AAR improve-how), (f) no public shame reel, no sports-cosplay belonging tax.

LV.3 Applied models of observation — who / what / when / how

FACT: Ste-Marie, Law, Rymal, Jenny, Hall & McCullagh (2012, *International Review of Sport and Exercise Psychology*, 5(2), 145–176, doi:10.1080/1750984X.2012.665076) advance an **Applied Model for the Use of Observation (AMUO)**: practitioners should assess observer and task characteristics, then vary **who** is observed, **what** is observed (and instructional features), **when** observation occurs, and **how** information is delivered — not treat “show

video” as a single intervention.

FACT: Ste-Marie et al. (2020, *Research Quarterly for Exercise and Sport*, doi:10.1080/02701367.2019.1693489) revisit AMUO across 2011–2018 applied-task articles; observation remains effective across a wider task base, but **guidelines still require intentional Who/What/When/How design** — effectiveness is not automatic from “video exists.”

HYPOTHESIS for MindCraft: Map AMUO onto practice:

- **Who:** self (attempt replay) ± skilled model (worked pathway for same archetype) ± peer only when SAFE-COMPARE allows.
- **What:** the decision frame that produced the soft-wrong — not the whole worksheet.
- **When:** after attempt, before next isomorphic item; optional delayed Notes deepen (aligns SAFE-AAR hot-wash + later deepen).
- **How:** sparse cueing + student-first naming (facilitation), not AI monologue.

Kill: “We added a replay button” as film-study. Survive: designed observation with AMUO knobs.

LV.4 Video feedback as performance analysis — learning ≠ tip retention

FACT: Pearson, Webb, Milligan & Dicks (2023/2025 print, *International Review of Sport and Exercise Psychology*, 18(1), 417–439, doi:10.1080/1750984X.2023.2235700) — integrative review of video feedback as a facet of PA (16 studies from 5,838 screened; OSF registration noted). Delivery mode (coach-led classroom vs self-led), environment, and **scheduling** shape learning; scheduling often aids **declarative knowledge retention** more clearly than transfer to on-field performance. Collaborative coach–athlete practices may enhance engagement. Authors call for better frameworks and experimental variety — PA is common; *learning-optimized* PA is under-specified.

Wound (transfer): Elite sport PA classrooms ≠ Maya’s 12-minute practice session. Treat Pearson et al. as a warning against equating “we reviewed the tape” with “performance transferred.”

Commercial implication: MindCraft must instrument **next-attempt behavior** (retry_120s, method-change codes, solo_transfer_pass) — not tip-quiz scores after coach cards.

LV.5 Feedback floods kill recall — less is more

FACT: Mason, Farrow & Hattie (2020, *International Journal of Sports Science & Coaching*; doi:10.1177/1747954120951080) — exploratory study of one-to-one post-match video feedback in an AFL club (6 coach–player dyads). Coaches delivered on the order of ~**30 feedback messages** per meeting. One week later, players recalled about **50% of summarised themes** but only about **6% of fine-grained feedback idea units**. Authors encourage a “**less is more**” approach.

HYPOTHESIS: AI tutors and “helpful” coach UIs recreate the AFL flood at machine speed — worse for anxious novices (SAFE-DD × anxiety; CLT). Cap coach-card points per clip at **1–3**. Prefer one bridge/misconception handle over a laundry list.

Product translation: Soft-wrong → single primary mis_ / ingredient cue → one advocacy–inquiry prompt (Rudolph/SAFE-AAR) → one next-attempt tweak. Kill “here are eight tips from your film.”

LV.6 Self-controlled video feedback — autonomy of *when*, not absence of coaching

FACT: van der Meer, van den Hoven, van der Kamp & Savelsbergh (2024, *Research Quarterly for Exercise and Sport*, 95(2), 537–545, doi:10.1080/02701367.2023.2275801; online 2023) — intermediate tennis players; self-controlled video feedback (request) vs yoked externally controlled schedule; coach provided video with attentional cueing and transitional tactical statements. Self-controlled group showed **larger tactical performance gains** vs pretest at posttest and **one-week retention**. Self-efficacy did not differ by group; gains were not predicted by self-efficacy / self-regulative questionnaire scores in this sample.

FACT (supporting line): Aiken, Fairbrother & Post (2012, *Frontiers in Psychology*, 3:338, doi:10.3389/fpsyg.2012.00338) — self-controlled video feedback benefits basketball set-shot learning relative to yoked schedules (motor-learning line often cited in later VBF work).

Wound: Meta-analytic debate exists on the strength of self-controlled feedback advantages (van der Meer et al. note McKay et al., 2022 vs Jimenez-Diaz et al., 2021). Do **not** market “science proves self-paced video raises ACT.” Treat as design prior: **learner-timed review + quality cueing** beats forced tip dumps *or* no review.

HYPOTHESIS for MindCraft: After a miss (or optional after a hard success), offer “Review this attempt?” rather than auto-playing a 90s AI lecture. Forced review every item → fatigue and controlling climate (SAFE-REWARD / SAFE-DD).

LV.7 Self-model pairing and feedforward — identity-safe media

FACT (pairing): Robertson, St. Germain & Ste-Marie (2018, *Journal of Motor Learning and Development*, 6(1), 18–34, doi:10.1123/jmld.2016-0027) — gymnasts learning two skills in a within-subject design: self-observation **paired with a skilled model** outperformed self-observation alone later in acquisition; error-identification **response sensitivity** was also higher for paired skills.

FACT (self-modeling / feedforward): Dowrick (1999, *Applied and Preventive Psychology*, 8(1), 23–39, doi:10.1016/S0962-1849(99)80009-2) reviews self-modeling; Dowrick (2012, *WIREs Cognitive Science*, 3(2), 215–230, doi:10.1002/wcs.1156) theorizes “learning from the future” via reconfigured component skills. Dowrick, Kim-Rupnow & Power (2006, *Journal of Special Education*, 39(4), 194–207, doi:10.1177/00224669060390040101) — video **feedforward** plus tutoring improved reading fluency; in 9/10 cases improvement rate was greatest during feedforward (small *N* — transfer carefully).

HYPOTHESIS: A miss clip should not stand alone. Pair with (1) a skilled worked pathway for the same archetype and/or (2) the student’s own prior success on a related edge. Feedforward spirit without fake editing: **show competence already owned, then the gap**.

Kill: Failure-only shame reels; success-only dopamine montages that hide productive misconceptions (SAFE-MISCON). Survive: paired clip design.

LV.8 SAFE-FILM stack (product law)

1. **Clip = objective attempt media** — steps, codes, timing; not “how did it feel?” alone (extends Keiser objective-media finding via LIV).
 2. **Less is more** — ≤ 3 cues per clip; prefer one primary soft-wrong handle (Mason et al.).
 3. **AMUO design** — decide who/what/when/how before shipping replay UI (Ste-Marie 2012/2020).
 4. **Self-controlled timing** — learner can request review; default is offer, not force-flood (van der Meer et al.; Aiken et al.).
 5. **Self-model pairing** — miss clip + skilled pathway and/or own prior success (Robertson/Ste-Marie line; Dowrick feedforward spirit).
 6. **Student names the frame first** — then advocacy–inquiry coach card (SAFE-AAR / Rudolph).
 7. **Transfer metric, not tip quiz** — instrument next attempt (Pearson wound).
 8. **Anti-cosplay** — borrow the *media pedagogy*, not locker-room brand (SAFE-RITUAL / SAFE-AAR).
 9. **No public error theater** — film study is private or tutor-dyad unless SAFE-COMPARE + consent (SAFE-CoP caution).
 10. **Claim ladder** — sport VBF evidence stays L1 domain evidence; MindCraft ACT/identity claims need FILM-* experiments.
-

LV.9 Competitive audit — who runs film study?

ChatGPT-style tutors rarely replay *your* attempt structure and risk tip floods. Khan shows correctness with weak process pairing. Duo streak liturgy $\neq$ film study. Brilliant ahas rarely close a self-model loop. Human “sports-style” tutoring inherits facilitator variance (Pearson). MindCraft already has attempt traces + soft-wrong + cards + Notes — **formalize SAFE-FILM**.

Commercial implication: Position as **reviewed decisions** — cue the one hinge, try the next play — never “train like the pros.” Parent copy: replay *how* they thought. Student copy: “pick a clip to review.”

LV.10 Claim ladder

Claim	Max ladder without new data
Observation interventions should vary who/what/when/how (AMUO)	L1 (Ste-Marie et al., 2012; 2020)
VBF-as-PA often improves declarative retention more clearly than transfer; design/delivery matter	L1 (Pearson et al., 2023)
Dense coach tip floods $\rightarrow$ poor week-later fine-grained recall (~6%); less-is-more recommended	L1 exploratory (Mason, Farrow & Hattie, 2020)
Self-controlled video feedback can outperform yoked schedules on tactical learning (tennis)	L1 single study (van der Meer et al., 2024)

Self-skilled-model pairing can beat self-observation alone (gymnastics)	L1 (Robertson, St. Germain & Ste-Marie, 2018)
Video feedforward can accelerate skill acquisition in some educational samples	L1 small-N (Dowrick et al., 2006); theory (Dowrick, 2012)
SAFE-FILM improves MindCraft retry_120s / method-change / solo_transfer_pass	L1 product bet — needs FILM experiments
“Sports academy methods raise ACT / identity” undifferentiated	Banned
Highlight-shame reels / tip floods as pedagogy	Banned

LV.11 Experiments spawned

ID	Question	Design	Primary	Kill condition
FILM-1	≤3-cue clip card vs uncapped AI tip flood after miss	2-arm	tip recall @24h; retry_120s; method-change	Flood ≥ sparse on transfer → revisit cue quality only
FILM-2	Self-controlled “Review?” vs forced replay every miss	2-arm	completion; affect; next-item transfer	Forced ≥ choice on transfer <i>and</i> better affect → keep forced for novices only
FILM-3	Miss-only clip vs miss+skilled model vs miss+own prior success	3-arm	error ID; solo_transfer_pass	Miss-only ≥ paired → pairing theater
FILM-4	AMUO-scripted coach card (who/what/when/how) vs generic praise/explain	2-arm	SE quality; transfer	Generic ≥ AMUO → card content fail
FILM-5	Film close + SAFE-AAR four questions vs film without AAR questions	2-arm	improve-how specificity; week-later transfer	No-AAR ≥ AAR → film without close still ok; keep AAR optional
FILM-QUAL	10 Maya: which clip UX felt useful vs shamey/sportsy?	interview	coded: useful / cosplay / flood / shame	“Athlete” language → kill that copy

Pre-reg (XXXIV): FILM-* identify **review media architecture** — not “sports methods raise scores” or identity-from-athleticism myths.

LV.12 Confidence table

Claim	Label	Confidence
Ste-Marie et al. AMUO who/what/when/how	FACT	High
Pearson et al. VBF-as-PA: delivery/scheduling; transfer caution	FACT	High
Mason et al.: ~30 tips; ~6% fine-grained week recall; less-is-more	FACT (exploratory, small <i>N</i>)	Medium–High
van der Meer et al.: self-controlled VBF > yoked on tennis tactics	FACT (single study)	Medium
Robertson, St. Germain & Ste-Marie: self∪model > self alone	FACT	Medium–High
Dowrick feedforward reading gains	FACT (small <i>N</i>)	Medium
Sport VBF transfers cleanly to ACT prep for anxious teens	HYPOTHESIS / wounded	Low–Medium
SAFE-FILM is right default review layer for MindCraft	FOUNDER BELIEF / HYPOTHESIS	Medium
Sports-cosplay branding improves belonging or scores	SPECULATION / false as doctrine	Low (against)

LV.13 What this chapter kills

1. **Kill:** Sports-academy / pro-athlete / locker-room branding as efficacy or belonging claim.
2. **Kill:** Feedback floods (30 tips / uncapped AI wrap) labeled as film study.
3. **Kill:** Equating tip recall or “we watched the replay” with transfer / exam readiness.
4. **Kill:** Failure-only shame reels and success-only dopamine montages as identity engines.
5. **Kill:** Replay UI without AMUO who/what/when/how design.
6. **Wound:** Classroom PA and elite dyads under-transfer to short teen sessions; facilitator variance.
7. **Survive:** SAFE-FILM; less-is-more clips; self-controlled review timing; self∪model pairing; feedforward spirit via own prior success; FILM-1...5; copy as **reviewed decisions** — not sports cosplay.

Doctrine until data: Ship **SAFE-FILM** — objective attempt clips, sparse cues, learner-timed review, paired models, transfer metrics — stacked on SAFE-AAR’s brief–debrief spine, and never sell athlete mythology or tip floods as the transformation engine.

Part LVI — Music Pedagogy Ladders: Scales → Repertoire → Recital Identity

Chapter status: Living evidence brief — Researcher tick 2026-08-01 (UTC hour 00 ≡ Red Team slot, but ch56 never written → prefer Researcher per rotation; 6 researcher entries since synthesizer v1.7 → Researcher)

Primary question: What transferable design does instrumental music pedagogy offer for MindCraft’s learning ladder — technique drills → meaningful repertoire → public competence signal — without conservatory cosplay, minutes-practiced vanity, or recital theater as identity?

Owners: Product (Practice path / story quests / exam signal) · Coach UX · Brand & copy · Red Team

Commercial job: Ship a **SAFE-MUSIC** stack — scales-as-ingredient drills, repertoire-as-story/exam pieces with learner choice, recital-as-transfer proof not stage cosplay — while killing “practice like a virtuoso,” hours-logged as mastery, and pass-floor exam identity.

LVI.1 Why this chapter exists

The core OS already sketches a music analogy (drill → quest → “recital” signal). Parts LI (SAFE-DP) and LV (SAFE-FILM) borrowed *session architecture* and *review media* from high-performance domains. Music pedagogy is the clearest **ladder** domain: beginners do not start with concertos; they build technique, then play repertoire they care about, then face graded exams or recitals that mark identity — *if* practice quality, motivation, and self-belief cohere.

MindCraft’s product already has the three rungs in embryo: ingredient/concept drills (scales), story-wrapped and exam-tagged questions (repertoire), ACT/diagnostic public signals (recital). Competitors collapse the ladder: Duo gamifies minutes; Khan floods explanations; ChatGPT skips technique for fluent answers; Brilliant jumps to ahas without graded public proof.

FOUNDER BELIEF under audit: Mathematical identity grows when learners climb a **named ladder** — controlled technique → chosen meaningful pieces → evidence they can perform under evaluation — not when a brand dresses homework as a conservatory.

Claims we refuse as doctrine:

1. Conservatory / “virtuoso / prodigy” branding as belonging or efficacy.
 2. Practice **minutes** or streaks as achievement (quantity without formal practice quality).
 3. Scales-forever without repertoire *or* repertoire without technique scaffolds.
 4. Recital/stage theater as identity without solo transfer evidence.
 5. Treating a bare “pass” exam outcome as motivational success.
-

LVI.2 Constructs (product language)

Construct	Music meaning	MindCraft analogue	Failure mode if misused
Scales / technique	Isolated skill building with immediate feedback	Ingredient drills; format reps; bridge micro-skills	Endless drill without meaning → boredom / controlling motivation
Repertoire	Whole pieces worth playing; often student-chosen	Story quests; archetype instances; exam-tagged sets	Pretty narrative with no technique → theater
Recital / grade exam	Public or adjudicated performance identity marker	Diagnostic; timed mixed set; solo_transfer_pass	Stage cosplay; pass-floor as “done”
Formal practice	Goal-directed time with SRL + deliberate strategies	SAFE-DP session spine with goals + feedback	Naive play-through of worksheets
Ineffective strategies	Error-chasing without analysis; mindless repetition	Immediate correction spam; tip floods; blocked vanity accuracy	Hallam “ineffective strategies” factor
SAFE-MUSIC	Ladder doctrine that survives cosplay / quantity / theater kills	See LVI.8	Conservatory brand; minutes KPI

Operational definition (HYPOTHESIS): A MindCraft path counts as *SAFE-MUSIC* when (a) technique drills are **goal-tagged ingredients** with feedback (not infinite XP), (b) repertoire includes **learner-chosen or interest-aligned** problem sets when feasible, (c) public signals require **mixed / transfer** performance not blocked accuracy, (d) minutes practiced never outrank formal-practice quality metrics, (e) branding names the ladder (skills → pieces → proof) without virtuoso mythology.

LVI.3 Deliberate practice’s home domain — music, not Gladwell hours

FACT: Ericsson, Krampe & Tesch-Römer (1993, *Psychological Review*, 100(3), 363–406, doi:10.1037/0033-295X.100.3.363) grounded deliberate practice in the training of instrumental musicians: effortful, teacher-guided activities with goals, immediate informative feedback, and repetition after reflection — distinct from playful performance or mere hours logged. Music academy students’ structured solitary practice was the empirical prototype.

Wound (already in SAFE-DP / Part LI): Hambrick-line critiques and domain limits mean “10,000 hours” and undifferentiated practice time are **banned** marketing. Music supplies the *definition*, not a guarantee that ACT scores rise from branded “deliberate practice.”

Product implication: Scales = short, goalful ingredient reps with feedback; never sell “log more hours this week” as the transformation engine (aligns Bonneville-Roussy below; SAFE-HABIT / SAFE-REWARD).

LVI.4 Strategy quality rises with expertise — not all practice develops

FACT: Hallam, Rinta, Varvarigou, Creech, Papageorgi, Gomes & Lanipekun (2012, *Psychology of Music*, 40(5), 652–680, doi:10.1177/0305735612443868) surveyed **N = 3,325** young instrumentalists across nine expertise/grade levels. Factor analysis yielded seven practice-related factors, including: systematic practice strategies; organization of practice; use of recordings/metronome for feedback; analytic strategies; **ineffective strategies**; concentration; immediate correction of errors. Statistically significant linear relationships with grade level appeared for systematic strategies, recordings/metronome use, ineffective strategies, and immediate error correction — but **not** for organization, analytic strategies, or concentration. Educational implication: expertise tracks *which* strategies learners adopt, not merely “they practiced.”

HYPOTHESIS for MindCraft: Instrument strategy_quality proxies (sectioning hard items, requesting SAFE-FILM clip, naming soft-wrong before retry) — not minutes_in_app. Decline of enjoyment of practice with advancement in related Hallam work is a wound: advanced ladders need repertoire meaning, not more grind UI.

Kill: Progress bars powered only by time-on-task. Survive: strategy factors as product instrumentation.

LVI.5 Quantity is not enough — formal practice mediates time → achievement

FACT: Bonneville-Roussy & Bouffard (2015, *Psychology of Music*, 43(5), 686–704, doi:10.1177/0305735614534910; online 2014) propose an integrative framework: **formal practice** = goal-directed, focused practice that includes both self-regulation and deliberate practice strategies. In a 4-month prospective SEM study, their integrative model predicted musical achievement better than traditional time-only approaches; practice time predicted achievement **when associated with formal practice**, not as raw quantity.

Commercial implication: Parent dashboards that lead with “hours practiced” recreate the failed music predictor. Lead with formal-practice signals: goals set, feedback used, transfer passed. Copy: “focused reps on the weak bridge,” never “she practiced 4 hours.”

Aligns: SAFE-DP (goals/feedback/rep), SAFE-HABIT (cue ≠ colonization), streak ban as North Star.

LVI.6 Repertoire choice — interest upgrades practice behavior

FACT: Renwick & McPherson (2002, *British Journal of Music Education*, 19(2), 173–181, doi:10.1017/S0265051702000256) — case study of a beginning clarinetist (“Clarissa”). When practising **self-selected** repertoire vs teacher-assigned pieces, she engaged strategies typical of more advanced stages (silent fingering, silent thinking, singing), spent more time, and persevered through difficulty. Authors link choice/interest to higher-level cognitive engagement (consistent with broader motivation findings cited therein).

FACT / FOUNDER-adjacent: Evans (2015) applies Self-Determination Theory to music education: scales are often **not** intrinsically enjoyable; motivation quality depends on whether learners internalize the *value* of technique (identified regulation) vs pure external control (teacher said so / praise). External control of scale drills risks the Deci overjustification pattern already killed in SAFE-REWARD (Part XLIV).

HYPOTHESIS: MindCraft “repertoire” should mix (1) teacher/engine-assigned technique (scales) framed with **why this ingredient unlocks the next piece**, and (2) **choice among story arcs / problem sets** when mastery allows — choice upgrades strategy quality, not just delight metrics.

Kill: Forced syllabus-only grind with no choice forever. **Kill:** choice-only playground with no technique ladder. **Survive:** scales with identified value + repertoire choice windows.

LVI.7 Recital / grade exams — identity markers that can demotivate

FACT: Hallam, Papageorgi, Varvarigou & Creech (2020 print / 2019 online, *Psychology of Music*, doi:10.1177/0305735618816168) — **N = 2,131** musicians aged 6–19 across expertise and exam outcomes (fail → highly commended). Those awarded **merely a pass** tended to undertake the **least** practice and responded least positively on organization, recordings/metronome use, analytic strategies, social life, and self-belief in musical ability. Those who **failed** were most likely to adopt ineffective strategies and less likely to enjoy performing, playing, lessons, and practice. Motivation models in the paper foreground musical **identity**, self-belief, goals, attributions, and supportive environment — not exam theater alone.

Wound: Graded exams are identity technology *and* stress technology (parent/teacher reputation dynamics noted in related literature the authors cite). A MindCraft “recital” that is only a score reveal can create pass-floor demotivation or fail-shame — opposite of FEI.

HYPOTHESIS: Public signals must be (a) criterion-referenced transfer proofs (solo_transfer_pass, mixed formats), (b) mastery-climate framed (SAFE TARGET / Part XXXVIII), (c) never sold as “your recital proves you’re a genius,” (d) paired with SAFE-AAR debrief + SAFE-FILM clip on the weak bar, not curtain-call XP.

Kill: Recital cosplay / standing ovation UI as identity. **Kill:** celebrating bare pass as transformative. **Survive:** adjudicated *competence evidence* with recovery path.

LVI.8 SAFE-MUSIC stack (product law)

1. **Name the three rungs** — technique (scales) → repertoire (pieces) → proof (recital/transfer).
2. **Scales = goalful ingredient drills** with immediate feedback (Ericsson definition; SAFE-DP).
3. **Repertoire includes choice** when possible — interest upgrades strategy quality (Renwick & McPherson).
4. **Frame technique value** (identified regulation) — why this drill unlocks the piece (Evans / SDT); avoid pure controlling rewards (SAFE-REWARD).
5. **Formal practice > minutes** — instrument strategy quality; never North-Star hours (Bonneville-Roussy & Bouffard; Hallam 2012).
6. **Public signal = mixed/transfer proof** — not blocked accuracy or stage theater (Hallam 2019 wound + SAFE-TRANSFER).
7. **Pass-floor ≠ win** — design recovery for bare-pass and fail states; protect self-belief (Hallam exam outcomes).
8. **Anti-cosplay** — borrow the *ladder*, not conservatory/prodigy brand (SAFE-RITUAL / SAFE-AAR pattern).
9. **Stack with prior SAFE-*** — DP session spine, FILM review of the hard bar, AAR close, CALIB confidence tiers.

10. **Claim ladder** — music evidence stays L1 domain evidence; MindCraft ACT/identity claims need MUSIC-* experiments.

LVI.9 Competitive audit — who runs the ladder?

Competitor	Scales	Repertoire	Recital	Ladder failure
Khan	Skill practice	Weak story meaning	Unit mastery quizzes	Explanation flood; weak identity proof
Duo	Micro-drills	Limited choice depth	Leagues/streaks as “recital”	Quantity/streak liturgy (banned NS)
Brilliant	Insight puzzles	Delight-as-repertoire	Weak graded public ACT proof	Aha without adjudicated transfer
ChatGPT tutor	Skips scales	Answers as repertoire	No solo proof	Fluent wrongness; no ladder
Human music-style tutor	Often strong	Often assigned-only	Real juries	Variance; stress; syllabus boredom

Commercial implication: Position MindCraft as **skills** → **pieces** → **proof** — “drill the hinge, play the piece, show you can transfer” — never “train like Juilliard” or “practice minutes that compound.” Parent copy: quality of practice and transfer proof. Student copy: “unlock the next piece” / “show what you can play cold.”

LVI.10 Claim ladder

Claim	Max ladder without new data
Deliberate practice prototype is music training with goals, feedback, repetition	L1 (Ericsson et al., 1993)
Strategy factors shift with grade level in large youth sample; time alone insufficient framing	L1 (Hallam et al., 2012)
Practice time predicts achievement when tied to formal (SRL+DP) practice	L1 (Bonneville-Roussy & Bouffard, 2015)
Student-selected repertoire can upgrade practice strategies vs assigned pieces	L1 case study (Renwick & McPherson, 2002)
Scales often need internalized value, not pure intrinsic fun	L1 theoretical/applied SDT (Evans, 2015)
Bare-pass exam outcomes associate with weaker practice/motivation profiles	L1 observational (Hallam et al., 2019/2020)

SAFE-MUSIC improves MindCraft transfer / identity metrics vs minutes-led UX	L1 product bet — needs MUSIC experiments
“Conservatory methods raise ACT / identity” undifferentiated	Banned
Hours practiced / streak as mastery or recital	Banned

LVI.11 Experiments spawned

ID	Question	Design	Primary	Kill condition
MUSIC-1	Named ladder UX (scales→piece→proof) vs flat question feed	2-arm	completion; solo _transfer_pass; perceived progress	Flat ≥ ladder on transfer → ladder is copy theater
MUSIC-2	Choice among 2–3 repertoire sets vs assigned-only	2-arm	strategy proxies; time-on-hard-items; affect	Assigned ≥ choice on transfer <i>and</i> affect → choice optional
MUSIC-3	Technique framed with “unlocks piece X” vs controlling “do your scales”	2-arm	persistence; identified reasons; next-piece start rate	Controlling ≥ identified → SDT frame fail
MUSIC-4	Parent dashboard: formal-practice metrics vs minutes-first	2-arm	parent WTP/comprehension; student pressure talk	Minutes-first better WTP without pressure → still wound educationally; don’t ship
MUSIC-5	Recital = mixed transfer set + AAR vs score-reveal only	2-arm	self-belief; retry; week-later transfer	Reveal-only ≥ AAR on belief <i>and</i> transfer → simplify
MUSIC-QUAL	10 Maya: does “scales/pieces/recital” language feel clarifying or elitist?	interview	coded: clarifying / cosplay / pressure	Conservatory words → kill that copy

Pre-reg (XXXIV): MUSIC-* identify **ladder architecture** — not “music methods raise scores” or prodigy-identity myths.

LVI.12 Confidence table

Claim	Label	Confidence
Ericsson et al.: DP rooted in music training structure	FACT	High
Hallam et al. 2012: large-N strategy factors × grade	FACT	High
Bonneville-Roussy & Bouffard: formal practice mediates time→achievement	FACT	High
Renwick & McPherson: choice upgrades strategies (case)	FACT (small <i>N</i>)	Medium
Evans SDT: scales need internalized value	FACT / applied theory	Medium–High
Hallam et al. 2019: pass-floor weakest practice/motivation profile	FACT (observational)	Medium–High
Music ladder transfers cleanly to ACT prep for anxious teens	HYPOTHESIS / wounded	Low–Medium
SAFE-MUSIC is right default path metaphor for MindCraft	FOUNDER BELIEF / HYPOTHESIS	Medium
Conservatory cosplay / hours KPI improve belonging or scores	SPECULATION / false as doctrine	Low (against)

LVI.13 What this chapter kills

1. **Kill:** Conservatory / virtuoso / prodigy branding as efficacy or belonging claim.
2. **Kill:** Minutes practiced, streaks, or hours-logged as North Star mastery (quantity without formal practice).
3. **Kill:** Scales-only forever without repertoire meaning *or* repertoire theater without technique.
4. **Kill:** Recital/stage cosplay and bare-pass celebration as identity transformation.
5. **Kill:** Equating “we practiced a lot” with exam readiness or musical/math identity.
6. **Wound:** Graded public performance can demotivate (pass-floor, fail-shame, parent pressure); transfer carefully.
7. **Survive:** SAFE-MUSIC ladder; formal-practice metrics; choiceful repertoire; identified technique value; transfer-as-recital; MUSIC-1...5; copy as **skills** → **pieces** → **proof**.

Doctrine until data: Ship **SAFE-MUSIC** — goalful scales, meaningful chosen pieces, adjudicated transfer proof — stacked on SAFE-DP / SAFE-FILM / SAFE-AAR, and never sell conservatory mythology or practice minutes as the transformation engine.

Part LVII — Chess Annotation & Metacognition: Postmortem UX Without Grandmaster Cosplay

Chapter status: Living evidence brief — Researcher tick 2026-08-01 (UTC hour 03; hour%6≠0; 7 researcher entries since synthesizer v1.7 → Researcher)

Primary question: What transferable design does chess *annotation* / *postmortem* offer for MindCraft’s after-attempt metacognition — student-generated explanation of critical positions, error classification into next drills, analysis of wins as well as losses — without grandmaster cosplay, engine-eval floods, or rating-as-identity?

Owners: Product (Practice debrief / Solver review / Notes) · Coach UX · ML soft-wrong → drill loop · Brand & copy · Red Team

Commercial job: Ship a **SAFE-ANNOT** stack — annotate before reveal, self-explain critical moves, classify → next drill, analyze successes not only failures — while killing “think like a GM,” engine dump ≡ coaching, and games-played/rating as learning proof.

LVII.1 Why this chapter exists

Parts LIV–LVI borrowed *session close* (SAFE-AAR), *review media* (SAFE-FILM), and *ladder architecture* (SAFE-MUSIC) from adjacent high-performance domains. Chess is the classic cognitive-science laboratory for **move choice under uncertainty** and for **post-game annotation** as deliberate learning — the closest analogue to MindCraft’s “attempt → review → next weak bridge” loop.

Competitors already fake the chess metaphor poorly: ChatGPT tutors monologue the “best line”; puzzle apps spam tactics without annotation; Duo/Khan celebrate streaks and accuracy without a student-authored postmortem. MindCraft’s soft-wrong science (SAFE-MISCON), hide-correctness diagnostic (SAFE-CALIB), and coach self-explanation doctrine (Part XL / SAFE-SE) need a **postmortem UX law** that is evidence-shaped, not chess-club branding.

FOUNDER BELIEF under audit: Mathematical identity grows when learners *annotate their own critical attempts* — name what they saw, why they chose, what principle failed — then convert that into one next drill, not when an engine or LLM floods them with better answers.

Claims we refuse as doctrine:

1. Grandmaster / “think like a master” branding as efficacy or belonging.
 2. Engine evaluation bars / tip floods labeled as “annotation.”
 3. Games played, rating, streaks, or puzzle counts as North Star mastery.
 4. Loss-only shame review (never annotate wins / successful strategies).
 5. Novice JOLs / vague “how well did you understand?” without forced restudy of critical moves.
 6. AI postmortem as automatically closing skill gaps (self-selection wound).
-

LVII.2 Constructs (product language)

Construct	Chess meaning	MindCraft analogue	Failure mode if misused
Critical position	Branch point where evaluation swings	Soft-wrong item; bridge fail; timed blunder	Annotating every move → exhaustion
Annotation / postmortem	Written/verbal account of thought at critical nodes	Student-first debrief after attempt/session	Engine dump before student speech
Candidate moves	Alternatives considered before choice	Alternative strategies / formats considered	Fake multiple-choice theater
Self-explanation	Why this move / principle applies	Chi/Renkl-style why-prompts (Part XL)	“Explain” that restates the tip
Engine / AI feedback	Stockfish-style eval after human thought	Coach card / model solution <i>after</i> annotate	Reveal-first collapses generation
Template / chunk	Pattern recognition structures in LTM	Ingredient / archetype recognition	“Memorize more patterns” vanity
SAFE-ANNOT	Postmortem doctrine that survives cosplay / flood / rating kills	See LVII.9	Chess branding; eval bars as pedagogy

Operational definition (HYPOTHESIS): A MindCraft review counts as *SAFE-ANNOT* when (a) the learner generates explanation of **1–3 critical attempts before** correctness/model reveal, (b) annotation names a **principle or misconception family** (not vibes), (c) the loop ends in a **linked next drill** (motif → ingredient set), (d) wins/successes are eligible for annotation (not only failures), (e) engine/LLM text never substitutes for student generation, (f) branding borrows the *method*, not grandmaster mythology.

LVII.3 de Groot — verbal protocols are the annotation prototype

FACT: Adriaan de Groot’s classic work (*Het denken van den Schaker*, 1946; English *Thought and Choice in Chess*, 1965/1978, Mouton / Amsterdam University Press lineage) collected **thinking-aloud protocols** from players of varying strength (including masters and world champions) who were asked to choose a move in unfamiliar tournament positions *as if in a real game*, verbalizing plans, calculations, and considerations. The object was the **choice-of-move problem**, not whole-game biography.

FACT (structure): de Groot described phases of chess thought commonly summarized as orientation → exploration → investigation → proof (see secondary summaries in the chess-programming / cognitive literature citing de Groot). Expertise showed up less as “search farther on every branch” in the original macrostructure story, and more as **better initial perception and selection of what to investigate**.

Wound (citation hygiene): Bilali█, McLeod & Gobet (2008, *Cognitive Science* line / “Sixty years of citing de Groot”) document that textbooks often **misreport** de Groot as “experts and novices search identically.” Later replications find additional search-structure differences. Product implication: do **not** market “masters don’t calculate more — they just see” as a slogan that excuses skipping hard work *or* as a claim that MindCraft’s annotation UI replicates grandmaster perception.

Product implication: Annotation UX = structured verbalization at critical nodes (what did you notice → candidates → why this choice → what would falsify it), not a vibes journal. Aligns SAFE-AAR question spine and Part XL self-explanation prompts.

LVII.4 Chunks and templates — recognition scaffolds search (not “genius”)

FACT: Chase & Simon (1973) chunking theory — experts recall meaningful game positions far better than novices; advantage collapses on random boards — grounded expertise in **pattern memory**, not mystical IQ (classic result; see also Gobet & Simon extensions).

FACT: Gobet & Simon’s **template theory** (e.g. Gobet & Simon, 1996; Gobet, 1998 *Cognition* comparison of expert-memory theories) proposes that recurring chunks evolve into richer templates with slots — unifying low-level chunks with higher-level schematic knowledge. Gobet (2005) notes deliberate practice matters because unsupervised “just play” can acquire **irrelevant** chunks that fail to improve (or even hinder) performance.

Commercial implication: MindCraft “annotation” should cue **ingredient/archetype recognition** (“which pattern is this?”) before deep calculation monologues. Kill copy that says students must “think like Kasparov.” Survive: pattern labels tied to ontology IDs after student attempt.

Kill: Genius branding. Kill: unstructured blitz-volume as chunk building. Survive: deliberate attention to the patterns that matter (SAFE-DP).

LVII.5 Self-explanation beats prediction-only for chess novices

FACT: de Bruin, Rikers & Schmidt (2007, *Contemporary Educational Psychology*, 32(2), 188–205, doi:10.1016/j.cedpsych.2006.01.001) — novices learning a King+Rook vs King endgame. Conditions: observe computer moves; predict next move; **predict + self-explain**. Self-explanation condition showed better principled understanding (predictions applying endgame principles) and more often checkmated in the test phase. Prediction alone did **not** beat observation.

HYPOTHESIS for MindCraft: Postmortem prompts must force **why/principle language**, not mere “what would you play next?” prediction clicks. This is the chess-domain replication of Chi/Renkl SE doctrine already in Part XL — use it to harden coach UX after attempts.

Kill: Silent engine scrubbing. Kill: “tap the best move” prediction theater without explanation. Survive: annotate-why before reveal.

LVII.6 Novice metacognition is fragile — JOLs can fail; forced restudy helps

FACT: de Bruin, Rikers & Schmidt (2005, *Applied Cognitive Psychology*, 19(2), 167–181, doi:10.1002/acp.1109) — novice chess endgame learning with judgments of learning (JOLs) and move selection for restudy. Forced selection of moves for restudy improved learning vs free selection (even controlling for number restudied). Providing JOLs showed

better *self-regulatory behavior* on some measures but **no or negative** performance benefit vs no-JOL — authors argue JOLs may place high ineffective load on novice working memory.

FACT: de Bruin, Rikers & Schmidt (2007, *European Journal of Cognitive Psychology* / related metacomprehension skill-acquisition paper, doi:10.1080/09541440701326204) — experienced chess players showed higher metacomprehension accuracy and self-regulation than novices when learning an endgame; novice absolute accuracy near zero. Expertise affects monitoring quality in skill domains.

Product implication: Do not ship bare confidence sliders as “metacognition” for Maya-novices (also SAFE-CALIB wound). Prefer **forced critical-move restudy** + principle annotation. Gate fancy monitoring UI by expertise / after scaffolding.

Kill: “Rate your understanding 1–5” as the whole postmortem. Survive: select 1–3 critical items → explain → restudy linked drill.

LVII.7 Analyzing games beats mere volume — and wins matter

FACT (observational, large-N): Yiannakoulis (2026, *Simulation & Gaming*, doi:10.1177/10468781261443352) — ~2M online blitz games / ~2.1k active Lichess players (1600–1800 band). Computer-assisted game analysis positively associated with rating improvement; **analyzing wins** showed a clearer association than analyzing losses in main models; total games played only weakly associated. Authors stress measurement noise (analysis flag is blunt; biases toward null) and exploratory observational status — not an RCT.

HYPOTHESIS: MindCraft should prompt annotation on **successful hard items** and efficient solutions (Dowrick-adjacent feedforward / SAFE-FILM success clips), not only on failures — “study your brilliancies” as well as soft-wrongs. Loss-only review risks ostrich/shame dynamics the paper discusses via adjacent feedback literature.

Wound: Associational evidence ≠ “annotation raises ACT.” Claim ladder stays L1 for chess rating deltas; MindCraft transfer needs ANNOT experiments.

Kill: “Play more puzzles” as improvement doctrine. Kill: failure-only shame reels. Survive: reflective analysis of critical successes and failures.

LVII.8 AI / engine feedback — self-selection and fluency wounds

FACT / high-relevance wound: Riedl & Bogert (2024/2026 update, arXiv:2409.18660) — 5+ years, ~52k individuals on an online chess platform with optional AI game analysis. Motivated and higher-skilled players **self-select** into AI feedback and use it more productively; apparent learning gains can be an **illusion of AI effectiveness** once endogenous motivation is accounted for; AI access can **widen skill gaps**; centralized AI feedback can reduce intellectual diversity (supported via platform natural experiments in the paper).

Aligns: Part XXXIII AI trust/sycophancy — fluent eval bars are not pedagogy; overtrust risk. SAFE-FILM less-is-more; SAFE-AAR facilitation not lecture.

Product implication: Optional “engine” / model solution is a **second pass** after student annotation. Do not market auto-analysis as closing equity gaps. Instrument who seeks postmortem help (selection bias).

Kill: “AI reviews every game so everyone improves equally.” Kill: reveal-first Stockfish cosplay in Solver. Survive: human-first annotation → sparse model contrast → one drill.

LVII.9 SAFE-ANNOT stack (product law)

1. **Student annotates first** — 1–3 critical attempts before correctness / coach / model reveal.
 2. **Prompt for principles** — why/candidates/what would change your mind (de Groot structure + de Bruin SE).
 3. **Forced restudy selection** for novices — don’t rely on raw JOLs (de Bruin 2005).
 4. **Classify → next drill** — motif/misconception → ingredient set (close the loop; not insight theater).
 5. **Annotate wins too** — successful strategies and hard-earned corrects (Yiannakoulis).
 6. **Sparse AI second pass** — contrast, don’t replace; watch self-selection (Riedl & Bogert).
 7. **Pattern labels after attempt** — chunks/templates as ontology tags, not genius myths (Chase/Simon; Gobet).
 8. **Anti-cosplay** — borrow postmortem method, not “grandmaster academy” brand (SAFE-RITUAL / SAFE-MUSIC pattern).
 9. **Stack** — SAFE-AAR close, SAFE-FILM clip, SAFE-MISCON soft-wrong, SAFE-SE prompts, SAFE-CALIB tiers.
 10. **Claim ladder** — chess evidence = L1 domain; MindCraft ACT/identity claims need ANNOT-* experiments.
-

LVII.10 Competitive audit — who runs a real postmortem?

Competitor	Annotation?	Student generation first?	Loop to next drill?	Failure mode
Khan	Hints/explanations	Often reveal/explain-first	Weak systematic soft-wrong→drill	Explanation flood
Duo	Minimal	No	Streak/XP loop	Quantity liturgy
Brilliant	Aha rationales	Partial	Weak exam transfer proof	Insight without postmortem habit
ChatGPT tutor	Fluent postmortem monologue	Usually model-first	Rare durable drill link	Sycophancy; fluent wrongness
Chess.com / Lichess	Engine analysis UX	Optional; often engine-first	Puzzle motifs if user closes loop	Eval bars; self-selection
Human tutor	Variable verbal postmortem	Often strong	Variable	Unscalable; shame risk

Commercial implication: Position MindCraft as **annotate** → **classify** → **drill** — “replay the critical move in your own words, then train the hinge” — never “AI grandmaster reviews your session.” Parent copy: reflection quality and linked practice, not games played. Student copy: “name the mistake / name the win, then one next rep.”

LVII.11 Claim ladder

Claim	Max ladder without new data
de Groot: verbal protocols reveal structured move-choice thought	L1 (classic methods)
Chase/Simon + Gobet/Simon: expertise tied to chunks/templates, not mysticism	L1
Self-explanation + prediction > prediction/observation for chess-novice principles	L1 (de Bruin et al., 2007 CEP)
Forced restudy selection helps novices; JOLs can fail/overload	L1 (de Bruin et al., 2005 ACP)
Metacomprehension accuracy rises with chess expertise in skill learning	L1 (de Bruin et al., 2007 EJCP)
Game analysis associated with rating gains; play volume weak; wins analysis salient	L1 observational (Yiannakoulias, 2026)
AI feedback gains partly selection; can widen gaps / reduce diversity	L1 observational + natural experiments (Riedl & Bogert)
SAFE-ANNOT improves MindCraft transfer / identity vs reveal-first UX	L1 product bet — needs ANNOT experiments
“Grandmaster methods” / engine dumps raise ACT undifferentiated	Banned
Rating / streak / games-played as mastery North Star	Banned

LVII.12 Experiments spawned

ID	Question	Design	Primary	Kill condition
ANNOT-1	Annotate-before-reveal vs reveal-first coach	2-arm	retry_120s; principle-coded annotation quality; week-later transfer	Reveal-first $\geq$ annotate on transfer $\rightarrow$ annotation is friction
ANNOT-2	SE-why prompts vs prediction-only taps	2-arm	principled language rate; next-drill completion	Prediction $\geq$ SE on transfer $\rightarrow$ simplify
ANNOT-3	Forced 2 critical restudies vs free JOL-only	2-arm	restudy adherence; transfer	JOL-only $\geq$ forced on transfer <i>and</i> affect $\rightarrow$ revisit

ANNOT-4	Annotate wins+losses vs losses-only	2-arm	return rate; self-belief; strategy reuse	Losses-only $\geq$ mixed on learning without shame cost $\rightarrow$ still prefer mixed ethically
ANNOT-5	AI second-pass on vs off after annotation	2-arm	overtrust incidents; transfer; time-to-drill	AI-on hurts calibration or transfer $\rightarrow$ default off for novices
ANNOT-QUAL	10 Maya: does “postmortem / annotate” language feel clarifying or chess-elitist?	interview	coded: clarifying / cosplay / shame	Chess words $\rightarrow$ plain “replay the hard step”

Pre-reg (XXXIV): ANNOT-* identify **postmortem architecture** — not “chess training raises ACT” or grandmaster-identity myths.

LVII.13 Confidence table

Claim	Label	Confidence
de Groot protocols established structured verbal move-choice analysis	FACT	High
Chunk/template accounts dominate expert chess memory explanations	FACT	High
de Bruin 2007: SE helps chess novices discover principles	FACT	High
de Bruin 2005: forced restudy > free; JOLs shaky for novices	FACT	High
Expertise improves metacomprehension in chess skill learning	FACT	Medium–High
Yiannakoulias: analysis $\leftrightarrow$ rating; play volume weak (observational)	FACT (assoc.)	Medium
Riedl & Bogert: AI feedback self-selection / gap / diversity wounds	FACT (platform)	Medium–High

Chess postmortem transfers cleanly to anxious ACT teens' identity	HYPOTHESIS / wounded	Low–Medium
SAFE-ANNOT is right default review metaphor for MindCraft	FOUNDER BELIEF / HYPOTHESIS	Medium
Grandmaster cosplay / engine-first dumps improve belonging or scores	SPECULATION / false as doctrine	Low (against)

LVII.14 What this chapter kills

1. **Kill:** Grandmaster / “think like a master” branding as efficacy or belonging.
2. **Kill:** Engine-eval floods and LLM monologues labeled as annotation or coaching.
3. **Kill:** Games played, rating, streaks, or puzzle counts as North Star mastery.
4. **Kill:** Loss-only shame review; never studying successful critical decisions.
5. **Kill:** Bare JOLs / confidence sliders as sufficient metacognition for novices.
6. **Kill:** “AI analyzes every attempt so gaps close equally” (selection + diversity wound).
7. **Wound:** Textbook oversimplification of de Groot search findings; don’t sloganize perception.
8. **Survive:** SAFE-ANNOT; annotate-before-reveal; SE principles; forced critical restudy; classify→drill; annotate wins; sparse AI second pass; ANNOT-1...5.

Doctrine until data: Ship **SAFE-ANNOT** — student-generated postmortem on critical attempts, principle-tagged, linked to the next drill, stacked on SAFE-AAR / SAFE-FILM / SAFE-SE / SAFE-MISCON — and never sell chess mythology, engine bars, or volume metrics as the transformation engine.

Part LVIII — Therapy: Graded Exposure (Math Anxiety Ladders)

Chapter status: Living evidence brief — Researcher tick 2026-08-01

Primary question: How should MindCraft borrow *graded exposure* / *inhibitory learning* from anxiety treatment — without therapy cosplay, flooding, or “calm until zero” theater — so threatened learners approach math instead of avoiding it?

Owners: Affect / Anxiety · Learning Science · Product (Practice sequencing + stakes) · Growth / Positioning · Red Team

Commercial job: Turn math-anxiety avoidance into a **designed approach ladder** that parents can trust and students can finish — then market approach + expectancy violation, not “we cured anxiety” or day-one ACT sims as exposure therapy.

LVIII.1 Why this chapter exists

Parts XXIV (affect/anxiety), XLI (SAFE-DD), XLV (SAFE-RESILIENCE), and L (SAFE-CALIB) already say: anxiety taxes working memory, unbuffered desirable difficulty becomes threat, grit posters fail, and “raise confidence” is not a North Star. What they do **not** yet specify is the **approach schedule** when the threat *is the math itself* — timed items, word problems, graphs, exams, public board energy.

Competitors fake exposure poorly: Khan dumps harder items after a wrong answer; Duo streaks punish skip; Brilliant opens with delightful puzzles that never feel like the threat Maya avoids; ChatGPT tutors soothe (“you’ve got this”) while the student never approaches the feared format. Therapy apps and school worksheets sometimes cosplay CBT with breathing gifs and no expectancy violation.

FOUNDER BELIEF under audit: Identity transformation for high-math-anxiety (HMA) Maya requires **approaching feared math situations under destaked conditions until threat expectancies are violated** — not forever-easy fluency, not flooding, and not clinical cosplay.

Claims we refuse as doctrine:

1. Day-one full ACT / mixed timed exams labeled “exposure therapy.”
2. Habituation-only success (“stay until anxiety = 0”).
3. Clinical CBT / therapist branding without clinician oversight.
4. Permanent safety behaviors (hint binge, skip, parent rescue) as “support.”
5. Strict linear hierarchy forever (never vary order/context).
6. Anxiety meters / “Calm Score™” as North Star.
7. Exposure → ACT +2 σ / Bloom marketing.

LVIII.2 Constructs (product language)

Construct	Therapy meaning	MindCraft analogue	Failure mode if misused
Fear hierarchy	Ranked approach steps from least → most feared	Format/stakes ladder (untimed soft → timed mix)	Hierarchy as forever easy-mode
Graded exposure	Approach feared CS while learning safety	Planned approach to avoided item types	Random harder = fake exposure
Expectancy violation	Feared outcome does not occur (or is tolerable)	Soft-wrong + survive + retry without catastrophe	Calm breathing without approach
Safety behavior	Acts that prevent learning (escape, ritual)	Instant full reveal; skip; streak protect	“Support” that blocks learning
Inhibitory learning	New CS—no US memory competes with fear memory	Multiple contexts/formats after approach	One safe session → exam panic
SAFE-EXPOSE	Exposure doctrine that survives flooding / cosplay / meter kills	See LVIII.9	Therapy branding; flood sims

Operational definition (HYPOTHESIS): A MindCraft session counts as *SAFE-EXPOSE* when (a) the learner **approaches** a previously avoided or high-SUDS math situation, (b) stakes are **destaked** enough that WM can engage (SAFE-DD), (c) a named **threat expectancy** is checked (“I’ll freeze / blank / be stupid”), (d) outcome **violates or revises** that expectancy without catastrophe framing, (e) **safety behaviors are faded** (hints → cards → solo), (f) later practice **varies context/format** so learning is not locked to one cozy UI, (g) product never claims to *treat* clinical anxiety disorders.

LVIII.3 Why avoidance is the commercial enemy

FACT (construct + avoidance): Hembree (1990, *Journal for Research in Mathematics Education*, 21(1), 33–46, doi:10.2307/749455 / 10.5951/jresmetheduc.21.1.0033) — meta-analysis of 151 studies: mathematics anxiety relates to poorer achievement, poorer attitudes, and is **bound directly to avoidance** of math; valid treatments that reduce anxiety tend to accompany improved performance.

FACT (WM disruption): Ashcraft & Kirk (2001, *JEP: General*, doi:10.1037/0096-3445.130.2.224) — HMA individuals show reduced WM spans under math demand; anxiety acts as a **transitory dual-task load**, amplifying errors especially on carry/complex items (see also Ashcraft & Krause, 2007, *Psychonomic Bulletin & Review*).

Commercial implication: The product problem is not “Maya needs nicer explanations.” It is **avoidance of the situations that would create competence evidence**. Marketing that promises comfort without approach sells temporary relief and loses the identity loop (FEI). Marketing that floods her with timed mixed ACT on day one sells rigor theater and churn.

Kill: “More content solves anxiety.” Kill: “Just make it fun.” Survive: designed approach to avoided formats under recoverable stakes.

LVIII.4 Classical graded exposure and math desensitization lineage

FACT (lineage): Systematic desensitization (Wolpe tradition) pairs gradual hierarchy ascent with incompatible responses (often relaxation). Education/counseling literature adapted this to math anxiety via hierarchy construction + gradual approach (see secondary reviews; e.g. IntechOpen synthesis citing Wolpe 1958 and school applications — cite cautiously as review, not primary RCT).

FACT (math-anxiety treatment comparison): Zettle (2003, *The Psychological Record*, 53, 197–215, doi:10.1007/BF03395440) — Acceptance and Commitment Therapy (ACT) vs systematic desensitization for mathematics anxiety; both approaches target anxiety-related avoidance, with different mechanisms (acceptance/values vs desensitization). Product takeaway: **approach + new relationship to anxious thoughts** matters more than branding one school of therapy.

FACT (case / protocol shape): Open-access childhood MA CBT case work (e.g. Coping Cat–inspired hierarchy + graded exposure; PMC8283868) shows the clinical shape — least→most feared math situations. **Wound:** N=1 ≠ product efficacy claim.

HYPOTHESIS for MindCraft: Borrow *hierarchy + graded approach + destake*; do **not** ship therapist roleplay or “we treat anxiety disorders.”

Kill: Clinical cosplay. Survive: hierarchy of *math situations* (formats, timing, stakes), not feelings posters.

LVIII.5 Inhibitory learning upgrades habituation theater

FACT (mechanism shift): Craske, Treanor, Conway, Zbozinek & Vervliet (2014, *Behaviour Research and Therapy*, 58, 10–23, doi:10.1016/j.brat.2014.04.006; PMC4114726) — exposure optimized via **inhibitory learning**: goal is new inhibitory CS–no US associations and expectancy violation, **not** necessarily within-session fear reduction to zero. Strategies include expectancy violation, removal of safety signals, variability, multiple contexts, retrieval cues, affect labeling (among others).

FACT (clinical overview): Craske inhibitory-regulation line (doi:10.1159/000381574) — anxious individuals often show inhibitory learning/regulation deficits; strengthen inhibitory retrieval, not just “wait until calm.” Variability-in-exposure reviews (e.g. PMC6884337) argue varying hierarchy order/intensity can aid expectancy violation and cross-context retrieval — school-product transfer remains empirical.

Product implication: Success metric is **approach + violated expectancy + later approach in new context** (challenge_accept, retry_120s, format hops), **not** “anxiety slider hit zero.” Soft-wrong + survive is expectancy violation done right. Breathing-only without approach is cosplay.

Kill: Habituation-as-North-Star. Kill: “Stay in the hard problem until you feel calm.” Survive: name the feared outcome → approach → check what happened → vary next.

LVIII.6 Extinction is fragile — renewal kills one-session miracles

FACT: Bouton (2002, *Biological Psychiatry*, 52(10), 976–986, doi:10.1016/S0006-3223(02)01546-9; PubMed 12437938) — extinction does **not** erase the original fear association; it creates new, **context-dependent** learning. Return of fear via **renewal** (ABA / ABC patterns) is expected when context changes after extinction.

FACT (learning & memory bridge): Bouton (2004, *Learning & Memory*, 11, 485–494) — contexts set the occasion for CS–US vs CS–no US; mere removal from the extinction context can renew responding.

Commercial implication: One cozy MindCraft session where Maya survived word problems does **not** license “anxiety cured” copy — exam hall / parental watching / timed ACT is a **new context**. Product must plan **multi-context approach** (home app → timed → mixed → low-stakes quiz climate) and retrieve inhibitory learning with cues (SAFE-AAR “what was I afraid would happen?”).

Kill: One-and-done exposure marketing. Survive: spaced approach across formats/stakes (ties SAFE-TIMING / SAFE-TRANSFER).

LVIII.7 Brief cognitive tools that are *not* exposure — and must not replace it

FACT: Park, Ramirez & Beilock (2014, *Journal of Experimental Psychology: Applied*, doi:10.1037/xap0000013) — short **expressive writing** before a math exam reduced the HMA–LMA performance gap (N=80); among HMAs,

anxiety/cause/insight language related positively to math performance. Related classroom writing finding: Ramirez & Beilock (2011, *Science*, doi:10.1126/science.1199427).

HYPOTHESIS: Writing / affect labeling can free WM *before* approach (Craske lists affect labeling among optimizers) — but writing **without** subsequent approach is incomplete for avoidance. Product: optional 60–90s worry-dump *then* hierarchy step — not a journaling app substitute for practice.

Kill: Mindfulness theater as the whole product. Survive: brief unload → approach → expectancy check.

LVIII.8 Crosswalk to SAFE-DD / FEI / soft-wrong

Exposure principle	MindCraft already-adjacent law	Product move
Equip before strain	SAFE-DD equip	Hierarchy starts below panic; cards available then faded
Destake while approaching	Soft-wrong; mastery climate	No public rank; no streak-as-worth during expose steps
Expectancy violation	FEI fear→evidence	Prompt: “What did you expect?” → “What happened?”
Remove safety signals	Hint fade / solo_transfer	Cap hint depth; require attempt before full reveal
Variability / multi-context	SAFE-TRANSFER / VT	After approach succeeds, hop format/timing
Do not flood	SAFE-RESILIENCE growth-zone	Ban day-one full timed mix for HMA segment

FOUNDER BELIEF: SAFE-EXPOSE is the **affective schedule** for SAFE-DD — difficulty dosing without approach design is incomplete for Maya who never opens the hard door.

LVIII.9 SAFE-EXPOSE stack (surviving product rules)

1. **Build a situation hierarchy, not a vibes hierarchy.** Steps = concrete math situations (untimed symbolic → word → diagram → timed soft → short mix). Prefer student-ranked SUDS-lite once, then product defaults.
2. **Approach before polish.** The session’s job is approach + expectancy check, not explanation volume.
3. **Destake the approach.** Soft-wrong, mastery climate, no league, no streak identity (SAFE-COMPARE / HABIT / REWARD).
4. **Name expectancy → violate → label.** Short prompt triad; optional expressive writing only as WM unload.
5. **Fade safety behaviors.** Hints → structured cards → annotate (SAFE-ANNOT) → solo transfer. Permanent hint binge = failed exposure.
6. **Vary after success.** Do not lock learning to one cozy context (Bouton renewal).
7. **Measure approach, not calm meters.** Primary: challenge_accept, retry_120s, avoided-format approach rate, later transfer_pass. Anxiety self-report is covariate, not North Star.

8. **Anti-cosplay.** No “therapist mode,” no diagnosis, no clinical claims. Plain language: “practice the scary step under safe rules.”

LVIII.10 Competitive audit — who actually runs exposure?

Competitor	Hierarchy?	Destaked approach?	Expectancy check?	Failure mode
Khan	Weak (skill tree ≠ fear ladder)	Partial	Rare	Harder after miss ≈ flood/shame
Duo	No	No (streak threat)	No	Avoidance via skip; streak colonization
Brilliant	Puzzle delight ≠ feared exam format	Low stakes yes	Rare	Fun ≠ approach to avoided threat
ChatGPT tutor	Ad hoc	Emotional soothe	Often sycophantic reassurance	Soothe without approach; fluent wrongness
Human tutor	Often strong	Variable	Variable	Unscalable; trait-label risk (SAFE-EXPECTANCY)

Commercial implication: Position as **approach ladders for feared math situations** — scary step under safe rules, then prove in a new format — never “AI therapy” and never day-one elite timed drills.

LVIII.11 Claim ladder

Claim	Max ladder without new data
Math anxiety tied to avoidance; treatments that work often improve performance	L1 (Hembree, 1990)
Anxiety disrupts WM during math	L1 (Ashcraft & Kirk, 2001)
Exposure benefits from inhibitory learning / expectancy violation > calm-to-zero	L1 clinical translation (Craske et al., 2014+)
Extinction is context-bound; renewal expected	L1 (Bouton, 2002/2004)
Expressive writing can shrink HMA performance gap pre-test	L1 (Park/Ramirez/Beilock)
Zettle: SD and ACT both target math-anxiety avoidance pathways	L1 clinical trial lineage

SAFE-EXPOSE raises MindCraft challenge_accept / transfer for HMA segment	L1 product bet — needs EXP-O experiments
Flooding / day-one ACT sims / Calm Score™ / therapy cosplay	Banned as doctrine

LVIII.12 Experiments spawned

ID	Question	Design	Primary	Kill condition
EXP-O1	Hierarchy+destake vs random harder after miss	2-arm HMA segment	challenge_acce pt; exit; week transfer	Random $\geq$ hierarchy on transfer without churn cost $\rightarrow$ revisit
EXP-O2	Expectancy-viola tion prompts on vs off	2-arm	expectancy revision rate; retry_120s	Prompts add friction without learning $\rightarrow$ shorten/kill
EXP-O3	Safety-behavior fade vs permanent hints	2-arm	solo_transfer_ pass; hint depth	Permanent hints $\geq$ fade on transfer $\rightarrow$ still prefer fade ethically; redesign fade
EXP-O4	Single-context approach vs multi-format renewal plan	2-arm	format-hop pass after delay	Single $\geq$ multi $\rightarrow$ still keep multi for exam ecology
EXP-O5	Worry-write then approach vs approach-only	2-arm HMA	accuracy under load; completion	Write hurts or delays without gain $\rightarrow$ optional/off
EXP-O-QUAL	10 Maya: “exposure ladder” language vs “practice the scary step”	interview	coded: clarifying / clinical / shame	Clinical words $\rightarrow$ plain language only

Pre-reg (XXXIV): EXP-O* identify **approach architecture for anxiety** — not “we treat anxiety disorders” or ACT score guarantees.

LVIII.13 Confidence table

Claim	Label	Confidence
Anxiety–avoidance–performance cluster is real at population level	FACT	High
WM dual-task account of online math-anxiety cost	FACT	High
Inhibitory learning reframes exposure away from habituation-only	FACT (clinical science)	High
Bouton renewal / context-bound extinction	FACT	High
Expressive writing helps some HMA test performance	FACT	Medium–High
Case/protocol SD for child MA generalizes to ACT teens in apps	HYPOTHESIS / wounded	Low–Medium
SAFE-EXPOSE is right default for MindCraft HMA segment	FOUNDER BELIEF / HYPOTHESIS	Medium
Therapy branding / flooding / calm meters improve belonging or scores	SPECULATION / false as doctrine	Low (against)

LVIII.14 What this chapter kills

1. **Kill:** Day-one full timed mixed ACT / flooding labeled as exposure therapy.
2. **Kill:** Habituation-only success (“anxiety must hit zero”).
3. **Kill:** Clinical CBT / therapist / diagnosis cosplay in product or ads.
4. **Kill:** Permanent hint binge, skip, or parent-rescue as “support.”
5. **Kill:** Anxiety meters / Calm Score™ / “we cured math anxiety” as North Star.
6. **Kill:** One cozy session ≡ exam-ready (renewal wound).
7. **Kill:** Bloom 2-sigma / exposure→score miracle marketing.
8. **Wound:** N=1 clinical cases and review chapters as efficacy proof.
9. **Survive:** SAFE-EXPOSE; hierarchy of situations; destaked approach; expectancy violation; fade safety behaviors; multi-context renewal plan; EXP-O1...5.

Doctrine until data: Ship **SAFE-EXPOSE** — graded approach to feared math situations under recoverable stakes, with expectancy checks and faded safety behaviors, stacked on SAFE-DD / FEI / SAFE-CALIB — and never sell therapy cosplay, flooding, or calm meters as the transformation engine.

Part LIX — Parent WTP Experiments Design

Chapter status: Living evidence + experiment-design brief — Researcher tick 2026-08-01

Primary question: How should MindCraft *measure* what parents will pay for — and which message frames move purchase intent — without confusing stated preference for cash, or selling score theater / vibe identity?

Owners: Growth / Positioning · Behavioral Econ · Product (parent surfaces) · Red Team

Commercial job: Ship a **pre-registered WTP / conjoint / message stack** that tells founders which attributes to build and which claims to kill — before pricing pages and parent dashboards ossify around the wrong North Star.

LIX.1 Why this chapter exists

Part XXVII mapped the post-AI WTP stack (answers → explanations → diagnosis → affect → accountability → trusted outcome signal) and spawned PAR-1 / COMP-1. Part XXXII flagged parent anxiety transmission; Part XXXVII's EVT-4 already asked competence-evidence vs minutes vs AI-unlimited. What the Constitution still lacked is the **methods contract**: how to run conjoint / discrete-choice and message tests so results can raise claim-ladder level (Part XXXIV L4 market causal) instead of producing inspirational survey quotes.

FOUNDER BELIEF under audit: Families will pay for *trusted diagnosis + transfer evidence + human accountability* more than for unlimited AI chat — **if** we prove it with trade-off designs, not founder intuition.

Claims we refuse as doctrine:

1. “Parents only buy score spikes” as settled fact without trade-off data.
 2. “Identity transformation” as a purchasable SKU without FEI evidence parents can parse.
 3. Likert “how much would you pay?” as serious WTP.
 4. Stated WTP ≡ revenue (hypothetical bias ignored).
 5. Streak / minutes / “AI unlimited” as the parent value prop by default.
 6. One national average price as the pricing strategy.
-

LIX.2 Constructs (product language)

Construct	Research meaning	MindCraft analogue	Failure mode if misused
WTP	Max price at which a package is chosen over alternatives / outside option	Monthly plan parents select in CBC	Treat survey WTP as sticker price
Attribute	Package feature with levels	Diagnosis depth; tutor hours; report type; price	Too many attributes → noise
CBC / DCE	Choice among multi-attribute profiles	Parent chooses Plan A vs B vs neither	Rank-all-features surveys
Part-worth	Utility of a level	Relative value of FEI report vs score-only	Over-interpret small segments

Message test	Framing → intent / comprehension	Landing / email / dashboard copy	Confuse click with retention
SAFE-WTP	Experiment doctrine that survives vibe pricing	See LIX.10	Sell identity without transfer proof

Operational definition (HYPOTHESIS): A MindCraft parent-pricing study counts as *SAFE-WTP* when (a) it uses **choice-based** trade-offs (or incentive-aligned follow-ups), not open “name a price,” (b) attributes map to **shippable product surfaces**, (c) a **price** or clear outside option is present, (d) **hypothetical-bias mitigations** are pre-registered, (e) primary outcomes separate **purchase intent**, **comprehension of child’s bottleneck**, and **homework-conflict risk**, (f) results update packaging — not vanity decks.

LIX.3 Market facts — demand exists; composition is uneven

FACT (US tutoring industry geography): Kim, Goodman & West (2024, CESifo Working Paper No. 11251 / IZA DP 17178) — U.S. private tutoring *centers* more than tripled ~1997–2022 (~3,000 → ~10,000); locations concentrate in high-income, high-parent-education areas; >half in top income quintile; residual association with Asian American population share conditional on income/education. School-market structure (private/charter/magnet) is only marginally predictive.

FACT (household spend snapshot, cited therein): Using Consumer Expenditure Survey calculations, authors report ~6–7% of U.S. families with children 6–17 paid for tutoring in the past year (2022), averaging about **\$437 in months with a purchase**; top 10% / 5% of monthly spend near **\$1,200 / \$2,000**. (Treat as CES-derived descriptive, not MindCraft price ceiling.)

FACT (shadow-education framing): Bray (2010) and Bray & Lykins (2012, ADB) document private supplementary tutoring (“shadow education”) as a global, fee-based parallel to schooling — useful for competitive context, not a US price guide.

Commercial implication: There is a real paid market adjacent to MindCraft — but it is **skewed**. Pricing and positioning that assume median-America “all parents pay \$X for identity” will miss both the premium competitive segment and equity constraints. Segment conjoint (income × anxiety × exam stakes) is mandatory.

Kill: One national WTP number as strategy. Survive: segment-aware CBC + equity-aware packaging.

LIX.4 Why Likert WTP fails — use stated choice

FACT (method foundation): Louviere, Hensher & Swait (2000, *Stated Choice Methods*, Cambridge) — stated preference via controlled choice experiments recovers relative attribute utilities under random-utility assumptions; preferable to unstructured preference talk when products are multi-attribute.

FACT (RUM lineage): McFadden (1974) random-utility discrete choice — choices reveal trade-offs; with a monetary attribute, researchers can express part-worths as WTP for levels.

FACT (tutoring-adjacent DCE): Gambi & De Witte (2024, *Higher Education*, doi:10.1007/s10734-024-01358-z) — DCE + information experiment on HE students’ tutoring preferences/consumption; high stated WTP coexists with low private uptake; SES demand vs ability-to-pay tension; behavioral frictions (status quo, image) matter. **Wound for K-12 parents:** student HE sample ≠ Maya’s parent — method transfers, coefficients do not. Broader HE DCE literature (UK/EU/Polish university-attribute studies building on Louviere et al.) shows the same tool recovers monetary trade-offs for staff expertise, flexibility, and mode.

Product implication: Stop asking “Would you pay for an identity-transforming math world?” Ask: *Between these two monthly packages, which do you choose — or neither?*

LIX.5 Hypothetical bias — stated ≠ paid

FACT: Murphy, Allen, Stevens & Weatherhead (2005, *Environmental and Resource Economics*, 30(3), 313–325) — meta-analysis of hypothetical vs actual WTP; median hypothetical/actual ratio ~1.35 with severe positive skew; choice-based elicitation associated with lower bias than some contingent-valuation formats (model-sensitive).

FACT (mitigation lineage): Cummings & Taylor (1999, *American Economic Review*) — “cheap talk” scripts can reduce hypothetical overstatement (effects heterogeneous later). Schmidt & Bijmolt (2020, *JAMS*, doi:10.1007/s11747-019-00666-6) — consumer-goods meta: prefer indirect/choice over direct “what would you pay?”; incentive-aligned mechanisms help when feasible.

HYPOTHESIS for MindCraft: Parent CBC will still inflate absolute dollars; **rank-order of attributes** (FEI report vs streak dashboard vs AI chat hours vs human tutor minutes) is the decision-useful output. Validate top packages with a small **deposit** / **waitlist** / **paid pilot** before national pricing.

Kill: Survey WTP as P&L input. Survive: relative part-worths → packaging; cash tests → prices.

LIX.6 Attribute set for MindCraft parent CBC (shippable)

Design rule (FOUNDER BELIEF): Every attribute level must map to a screen or ops cost we can actually ship in 90 days.

Attribute	Levels (draft)	Ties to doctrine
Diagnosis depth	Vague “needs practice” / concept weakness list / bridge+format gap map	Ontology advantage; XXVII trust
Evidence report	Score/minutes only / FEI pack (retry_120s, challenge-seeking, transfer) / FEI + honest plateau note	PAR-1; anti-streak
Human accountability	None / async tutor note weekly / 2× live 25-min / 4× live	Scarcity layer

AI help policy	Unlimited chat / guarded pedagogy wrap / attempt-first then cards	XXXIII; SE constraints
Practice architecture	Streak gamification / mastery climate soft-wrong / SAFE-EXPOSE ladder language	XLII–XLIV; LVIII
Price	e.g. \$49 / \$99 / \$179 / \$299 per month (calibrate to segment)	Must include realistic span
Outside option	“Neither — keep free Khan/GPT/status quo”	Critical for WTP

HYPOTHESIS: Diagnosis depth + FEI report + modest human accountability will dominate unlimited AI and streak theater **in the anxious-middle segment**; score-chaser segment may overweight timed-test sims — measure, don’t assume.

Red Team: If FEI report wins stated choice but parents cannot *explain* the child’s bottleneck after viewing it, comprehension failed — packaging theater.

LIX.7 Message tests are not conjoint (run both)

Conjoint answers *what package*; messages answer *which story opens the wallet without raising homework conflict*.

Pre-registered message arms (builds EVT-4 / COMP-1):

Arm	Parent-facing frame	Risk
M1 Score relief	“Raise ACT math”	Overclaim; retention crash if flat
M2 Homework peace	“End nightly fights”	Soothing without competence
M3 Diagnosis trust	“See the exact bridge that’s stuck”	Jargon; need plain language
M4 Identity / confidence	“Become someone who doesn’t freeze”	Belief Score™ / raise-confidence kill (v1.7–v1.8)
M5 Hybrid	Diagnosis + transfer check + honest plateau	Longer copy; may win retention

Primary outcomes: click → trial start; 14-day activation; parent comprehension quiz (name bottleneck); reported homework conflict; **not** vanity “brand love.”

Kill: Identity-only ads that cannot point to a measurable FEI metric. Survive: hybrid that sells method, shows evidence.

LIX.8 Qualitative gate before fielding CBC

FACT / process: Part XXVII and QUAL-P (XXXII) already require ~10 parent interviews. CBC attributes invented in a founder Slack will mis-specify levels.

Protocol (HYPOTHESIS): Semi-structured interviews → code “what I thought I bought / what would make me cancel / words I trust / words that feel scammy” → freeze attribute glossary → then CBC. Do not reverse the order.

Segment quotas (SPECULATION until interviews): (1) HMA-Maya parents (anxiety/avoidance), (2) score-chaser / selective-college, (3) homework-peace exhausted, (4) price-sensitive remedial. Run latent-class or split models; do not average them into one “parent.”

LIX.9 Equity and pricing ethics

FACT (concentration): Kim et al. (2024) — private center supply concentrates where income/education are high. Packaging only for that segment reproduces shadow-education inequality.

FOUNDER BELIEF: Charging for diagnosis+accountability is legitimate if transfer evidence is real (XXVII.7). Illegitimate if the paid tier is streak cosmetics + GPT wrap.

Guardrails:

- Paid claims ↔ measured FEI / learning metrics.
- No “guaranteed score jump” in any CBC level or ad arm.
- Consider a **honest free tier** that still shows diagnosis (trust acquisition) vs freemium that hides the bottleneck (dark pattern).
- Report segment WTP separately; do not market top-quintile WTP as universal.

LIX.10 SAFE-WTP stack (surviving product rules)

1. **Choice over chatter.** CBC/DCE or incentive-aligned choice; ban standalone “name your price.”
2. **Ship-mapped attributes only.** No fictional features.
3. **Always include price + neither.** Otherwise WTP is fantasy.
4. **Bias protocol.** Cheap talk / certainty / consequentiality / small cash validation — pre-register which.
5. **Three outcomes.** Intent ≠ comprehension ≠ conflict. Optimize packaging on all three.
6. **Segment, don’t average.** Latent class or pre-defined segments.
7. **Messages ≠ packages.** Run COMP/EVT message tests alongside CBC.
8. **Anti-theater.** FEI report that parents cannot interpret loses even if chosen.
9. **Anti-overclaim.** No 2σ, no guaranteed ACT points, no “AI tutor replaces human” as paid promise.
10. **Update doctrine with data.** Part-worths that kill founder favorites get logged as kills.

LIX.11 Claim ladder

Claim	Max ladder without new data
-------	-----------------------------

US private tutoring centers grew sharply; high-income concentration	L1 (Kim/Goodman/West 2024)
Shadow education is a real global fee market	L1 (Bray line)
Stated choice recovers relative utilities better than unstructured WTP talk	L1 (Louviere/Hensher/Swait; McFadden)
Hypothetical WTP inflates; choice formats help but do not erase bias	L1 (Murphy et al. 2005; Cummings & Taylor)
HE tutoring DCE: high WTP, low uptake, SES frictions	L1 adjacent (Gambi & De Witte 2024) — transfer wound to K-12 parents
MindCraft parents prefer FEI+diagnosis over unlimited AI / streaks	L0–L1 product bet — needs WTP-* experiments
Survey dollars = revenue forecast	Banned

LIX.12 Experiments spawned

ID	Question	Design	Primary	Kill condition
WTP-1	FEI report vs score/minutes vs streak dashboard in CBC	CBC, price + neither	part-worths; segment classes	Streak/minutes dominate FEI <i>and</i> later retention → revisit packaging (still educational wound)
WTP-2	Diagnosis depth (vague / concept / bridge+format) × price	CBC	WTP for bridge+format; comprehension	Depth raises intent but kills comprehension → simplify language
WTP-3	Human minutes vs AI-unlimited vs guarded wrap	CBC	part-worth human; churn proxy intent	AI-unlimited wins on intent <i>and</i> 30-day solo transfer → thesis wounded
WTP-4	Message arms M1–M5	RCT landing/email	trial start; conflict; comprehension	Identity arm wins clicks + raises conflict without retention → kill identity-only
WTP-5	Cheap-talk / certainty calibration on vs off	2× embedded in CBC	calibration gap	No gap → still keep cash pilot

WTP-6	Deposit / paid pilot of top CBC winner vs #2	Field	pay-through; 30-day retention	Stated winner loses cash → trust revealed preference
WTP-QUAL	10 parents: cancel triggers + trusted words	Interview	codebook → attribute freeze	Founder glossary ≠ parent glossary → rewrite levels

Pre-reg (XXXIV): WTP-* identify **parent package preference and message effects** — not “we proved identity sells” or national price optimality from N=200 survey.

Ties: EVT-4, PAR-1, COMP-1, MUSIC-4, TIM-4 become arms or covariates inside this family — do not fork duplicate surveys.

LIX.13 Confidence table

Claim	Label	Confidence
Paid tutoring demand exists and is income-skewed in the US center market	FACT	High
CES spend snapshots describe <i>purchasers</i> , not all parents’ WTP	FACT	High
CBC/DCE is the right class of tool for multi-attribute tutoring packages	FACT (methods)	High
Absolute stated dollars will be biased upward	FACT / HYPOTHESIS	High
FEI+diagnosis beats unlimited AI for MindCraft’s anxious-middle segment	HYPOTHESIS	Medium
Hybrid message (diagnosis + transfer + honest plateau) wins retention vs score-only	HYPOTHESIS	Medium
Identity-only parent ads are net-positive	SPECULATION / against doctrine until data	Low (against)
One price point fits all segments	SPECULATION / false as strategy	Low (against)

LIX.14 What this chapter kills

1. **Kill:** Likert / open-ended “how much would you pay?” as pricing research.

2. **Kill:** Treating survey WTP as revenue or national sticker price.
3. **Kill:** “Parents only buy scores” or “parents buy identity” as settled without CBC.
4. **Kill:** Streak / minutes / AI-unlimited as default parent value props without losing on part-worths.
5. **Kill:** Guaranteed ACT points / 2σ / Bloom in any package level or ad.
6. **Kill:** Averaging anxious-middle and score-chaser into one “parent persona.”
7. **Wound:** Importing HE or non-US tutoring conjoint coefficients as US ACT-parent truth.
8. **Survive:** SAFE-WTP; CBC with price+neither; ship-mapped attributes; bias protocol; message≠package; WTP-1...6 + QUAL; cash validation before scale.

Doctrine until data: Run **SAFE-WTP** — choice-based parent trade-offs on diagnosis, FEI evidence, human accountability, AI policy, and price — then package what wins *and* what parents can explain, not what flatters the identity thesis.

Part LX — Trust After AI Hallucination

Chapter status: Living evidence + repair-protocol brief — Researcher tick 2026-08-01

Primary question: After MindCraft’s AI is *caught* being fluently wrong (or quietly wrong and later exposed), how do we repair *calibrated* trust — and when must we escalate to a human — without apology theater, never-wrong marketing, or post-error abandonment?

Owners: AI / Product (Solver · coach) · Learning Science · Human Factors · Parent trust · Red Team

Commercial job: Ship a **SAFE-REPAIR** doctrine that turns inevitable model errors into a *positioning asset* (honest plateau + human accountability) rather than a churn event or a “AI that never fails” lie.

LX.1 Why this chapter exists

Part XXXIII owned *prevention and calibration*: pedagogy wrap, anti-sycophancy, Bastani unguarded-harm FACT, Lee & See appropriate reliance. Part L (SAFE-CALIB) owned *student* confidence after misses. What the Constitution still lacked is the **post-violation playbook**: what the product does in the *minutes after* a hallucination is detected — by checker, by student flag, by tutor review, or by exam-collapse discovery.

FOUNDER BELIEF under audit: Families and students will tolerate imperfect AI *if* the system (a) admits competence failure fast, (b) shows a concrete repair path tied to the Map/ontology, and (c) escalates to a human when stakes or repeated failure demand it — and will *punish* products that either hide errors or drown users in soft apologies without correction.

Claims we refuse as doctrine:

1. “Our AI doesn’t hallucinate” as marketing (false category; Bastani notes GPT Base hallucinations as mechanism risk).
2. Rote “Sorry!” as trust repair.
3. Empathy-only recovery that never corrects the math.
4. Permanent undertrust (“never trust the coach”) as the safety story.

5. Unlimited human escalation as the default UX (destroys unit economics and tutors).
6. Student thumbs-up / “was this helpful?” as proof the repair worked.

LX.2 Constructs (product language)

Construct	Research meaning	MindCraft analogue	Failure mode if misused
Hallucination / fluent wrongness	Confident false content	Wrong hint, invented step, affirming wrong setup	Treat only wild fabrications; ignore soft arithmetic slips
Trust violation	Event that breaks reliance calibration	Caught wrong coach move; silent wrong→exam miss	Every minor UI glitch as “trust crisis”
Competence vs integrity failure	Ability miss vs values/honesty miss (Kim et al.)	Bad algebra vs hiding known wrongness	Apologize for everything the same way
Trust repair	Actions that restore <i>appropriate</i> reliance	Admit → correct → verify → optional escalate	Maximize liking, not calibration
Misuse / disuse	Over-reliance / under-reliance (Parasuraman & Riley)	Copying Solver; abandoning coach after one miss	Chase either extreme as “engagement”
Escalation	Hand-off to human when AI insufficient	Tutor queue / session brief with error packet	Escalate everything or nothing
SAFE-REPAIR	Surviving post-error doctrine	See LX.10	Apology cosplay; never-wrong brand

Operational definition (HYPOTHESIS): A MindCraft post-error flow counts as *SAFE-REPAIR* when (a) the system **detects or accepts** the error signal, (b) labels it as a **competence** miss (not student shame), (c) issues an **explanatory** repair with a **deterministic correction path** (bank key / checker / Map ingredient), (d) invites a **solo re-attempt** before full solution dump, (e) offers **escalation** only under explicit triggers, and (f) measures repair by **re-engagement + verified next attempt**, not apology NPS.

LX.3 Prior art already in the Constitution (do not re-litigate)

FACT (high-school math field RCT): Bastani, Bastani, Sungu, Ge, Kaback & Mariman (2025, *PNAS*, doi:10.1073/pnas.2422633122) — unguarded GPT access inflated practice then **harmed** solo exam (~–17% vs never-AI); guarded “GPT Tutor” largely neutralized harm but did **not** show positive learning vs control; Base hallucinations and crutch copying are named mechanisms. Perception ≠ learning.

FACT (novice verification failure): Li, Song, Sundaram & Karahalios (L@S 2025, doi:10.1145/3698205.3729550) — most learners failed to detect chatbot factual errors; undetected errors harmed outcomes and self-efficacy.

FACT (appropriate reliance): Lee & See (2004, *Human Factors*, doi:10.1518/hfes.46.1.50_30392) — design for calibration, resolution, specificity; misuse and disuse are both failures.

FACT (use/misuse/disuse): Parasuraman & Riley (1997, *Human Factors*) — operators over-rely, under-rely, or abuse automation; trust is a primary mediator.

FACT (trust layers): Hoff & Bashir (2015, *Human Factors*, doi:10.1177/0018720814547570) — dispositional, situational, and learned trust; post-error recovery is largely a **learned-trust** problem.

Product implication: Prevention (PWC / ontology / checkers) remains primary. This chapter owns what happens when prevention fails in the wild — which it will.

LX.4 Trust repair is a designed capability, not a tone setting

FACT (HMI agenda): de Visser, Pak & Shaw (2018, *Ergonomics*, doi:10.1080/00140139.2018.1457701; “From ‘automation’ to ‘autonomy’”) — as systems become more autonomous, **trust repair** after violations becomes a first-class design requirement; competence vs integrity framing and context/risk matter; human–human repair findings transfer only partially.

FACT (human–human baseline): Kim, Ferrin, Cooper & Dirks (2004, *Journal of Applied Psychology*, doi:10.1037/0021-9010.89.1.104) — apology repairs **competence** violations better than denial; for **integrity** violations, denial of culpability can outperform apology (when evidence supports innocence). Follow-on attribution work (Kim, Dirks, Cooper & Ferrin, 2006, *OBHDP*, doi:10.1016/j.obhdp.2005.07.002) shows internal vs external blame interacts with violation type.

HYPOTHESIS for MindCraft math errors: Nearly all Solver/coach hallucinations are **competence** failures. Correct response class = **apology + ownership + concrete fix**, not “the curriculum is wrong” externalization, and not integrity-theater (“we would never...”). If the product *knew* the key and still affirmed the student’s wrong answer (sycophancy), treat that as closer to **integrity** — repair requires admitting the agreement failure, not only the math slip.

Wound (field realism): Liu, Du, Ma, Zhang & Yan (2024, *IEEE Trans. Human-Machine Systems*, doi:10.1109/THMS.2024.3434680) — after automated-driving failure in a test-track study (N=257), verbal apology+explanation+promise **did not restore damaged trust**; human-voice repair only partially mitigated other attitudes, mainly among less experienced drivers. **Transfer wound:** ADS safety failure ≠ wrong algebra hint — but the kill candidate survives: **do not assume apology scripts restore trust scores**.

Kill: Apology copy as the repair KPI. Survive: behavioral re-calibration after a verified correct path.

LX.5 What users prefer after LLM mistakes (preference ≠ recovery)

FACT (preregistered preference study): Zhang et al. (2025/26; *CHI / ACM*, doi:10.1145/3793679; arXiv:2507.02745) — N=162 Prolific; pairwise preferences among **rote**, **explanatory**, and **empathic** apologies after bias, fabrication (hallucination), and factual-error scenarios. **Explanatory** apologies generally preferred; **empathic** preferred in bias contexts; for hallucinations, users saw seriousness but showed **no clear apology-type preference** — uncertainty, not a settled recipe.

FACT (critical review): Harland et al. (2025, *Artificial Intelligence Review*, doi:10.1007/s10462-025-11305-8) — AI apology literature (2020–2023 synthesis) stresses that apology is only one element; identification, explanation, and **system behavior change** are required for meaningful repair. Purely discursive apology without corrective action is ethically and practically thin.

HYPOTHESIS: For Maya after a wrong coach step, the winning stack is **explanatory ownership + Map-linked correction + destaked re-attempt**, not empathic padding (“that must feel frustrating”) alone. Empathy may reduce shame (Part XXIV) but does not restore calibration.

Contradiction / limit: Some GenAI service-recovery work finds gratitude/appreciation or humorous styles can raise *willingness to forgive* in consumer settings (e.g. Systems 2024 mixed methods on language style × gratitude/apology — domain = service UX, not math learning). Treat forgive-intent as **engagement theater** unless paired with verified next-attempt success.

Commercial implication: Marketing may say “when we’re wrong, we show our work and hand you a human.” Product must actually do the second half.

LX.6 Escalation to human — scarce, triggered, packetized

FACT (hybrid classroom deployment): Kazemitabaar et al. (2025, arXiv:2510.14457) — programming course N=82; 673 AI hints; students rated 146 (22%) unhelpful; of those, only **16 (11%)** escalated to instructors. Instructor replies on escalated cases were incorrect/insufficient ~half the time. Escalation is **underused** by students and **not automatically high-quality** for humans.

FACT (transparency shifts strategy, not accuracy): Nagashima, Hladký & Rief (arXiv:2606.03822; ECTEL 2026) — fallibility warning increased hint-seeking without changing identical ITS behavior or immediate performance (already in XXXIII).

HYPOTHESIS — MindCraft escalation triggers (product contract):

1. **Checker conflict:** model claim ≠ bank key / symbolic checker → auto-flag; do not wait for student detective work.
2. **Repeat competence miss:** ≥2 verified wrong coach moves in one mission on same concept/ingredient → offer human.
3. **Affect spike + stuck:** high self-reported threat / long idle after soft-wrong → soft escalate (tutor ping), not more AI monologue.
4. **Integrity-class event:** sycophantic affirm of wrong answer against known key → mandatory human-visible audit packet.
5. **Parent/exam stakes mode:** optional “human-reviewed session” SKU (ties SAFE-WTP human-minutes attribute) — not free infinite chat.

Escalation packet (HYPOTHESIS): Human receives: problem ID, student attempt trace, AI utterance, contradiction evidence, ontology IDs, student affect flag — not a blank “help?” thread. Aligns with SAFE-AAR / SAFE-ANNOT (trace before coach dump).

Wound: Human-in-the-loop is not magic (Kazemitabaar instructor error rate). Escalation must preserve **tutor attention hygiene** — packetized, not “AI failed, you debug the LLM live.”

Press-adjacent note (SPECULATION until primary paper cited in-repo): Hybrid AI-draft + human-approve tutoring reports (e.g. Eedi/LearnLM coverage in education press) suggest low hallucination rates *under* human gatekeeping — useful as competitive *direction*, not as MindCraft proof. Prefer primary methods papers before marketing “near-zero hallucinations.”

LX.7 Identity stakes — why math trust violations cut deeper

HYPOTHESIS: For Maya, a fluent wrong coach move is not only a UX bug; it reactivates the identity narrative “even the helper confirms I’m lost” *or* the opposite overtrust “the AI said I’m right” (Parts XXV, XXX, L). Repair must separate:

- **System competence failure** (external, unstable, specific — Weiner-friendly)
- **Student strategy miss** (controllable after correct feedback)

Never: “you’re just not a math person” and never: “trust me completely next time.”

FOUNDER BELIEF: Soft-wrong already destakes *student* error. Post-hallucination UX must destake *AI* error without teaching helplessness toward tools (Lee & See disuse).

LX.8 Competitive positioning

Competitor move	Failure mode	MindCraft counter (HYPOTHESIS)
ChatGPT raw	Hallucinate + apologize vaguely + continue	Checker + Map correction + solo re-attempt
“AI tutor never wrong” ads	Integrity violation when caught	Honest fallibility + repair SLA
Unlimited chat, no human	Disuse after salient miss; or silent misuse	Triggered human minutes as paid scarce resource
Thumbs-up quality loop	Bastani/Bo: perception fails	repair_reattempt_pass, escalation_resolve_48h

Copy that survives Red Team: “We catch wrong steps, show the fix, and bring a tutor when it matters.”

Copy that dies: “Our AI never makes mistakes” / “Sorry you’re frustrated!” without a corrected path.

LX.9 SAFE-REPAIR stack (surviving product rules)

1. **Detect** — prefer deterministic contradiction over student-reported vibes.
2. **Classify** — competence vs integrity-class (sycophancy-against-key).
3. **Own** — short explanatory apology; no blame-the-student; no vague empathy wall.
4. **Correct** — show Map/ingredient + checked next step; ban full solution until re-attempt window unless escalate.
5. **Re-attempt** — destaked item; measure success.

6. **Calibrate UI** — temporarily lower coach confidence badge on that skill grain (SAFE-CALIB cousin).
7. **Escalate** — only on triggers; send packet; SLA visible to parent/student.
8. **Learn** — log to evaluation set; do not only “prompt harder.”
9. **Anti-cosplay** — no therapy script, no airline-captain apology theater, no never-wrong brand.

LX.10 Experiments (REPAIR family)

ID	Contrast	Primary outcome	Falsifier
REPAIR-1	Explanatory+correct path vs rote apology vs empathy-only after injected wrong hint	repair_reattempt_pss @10m; trust calibration item	Empathy-only wins reattempt → revisit stack
REPAIR-2	Auto-detect+own vs silent correct-and-continue	next-mission return; overtrust score	Silent fix retains more misuse → keep detect
REPAIR-3	Escalation offer after 1 vs 2 vs 3 verified AI misses	escalate rate; human resolve quality; cost/session	Instant escalate floods tutors without lift
REPAIR-4	Packetized tutor handoff vs blank “talk to tutor”	time-to-correct; tutor CSAT; reattempt	Packet no better → packet UX fail
REPAIR-5	Fallibility banner always-on vs post-error-only vs none	hint abuse; solo transfer	Always-on destroys use without calibration gain
REPAIR-6	Parent-visible “AI miss + repair” FEI note vs hidden	parent trust / WTP message arm	Hidden wins retention → wound transparency thesis
REPAIR-QUAL	10 Maya + 5 tutors: what “sorry” means after wrong hint	codebook → copy freeze	Founder apology ≠ user need

Pre-reg (XXXIV): REPAIR-* identify **post-error calibration and escalation efficiency** — not “apology increases ACT scores” or “human tutors eliminate hallucinations.”

Ties: AIT-, CAL-, ANNOT-, AAR-, WTP human-minutes, EXP-O affect gates.

LX.11 Confidence table

Claim	Label	Confidence
-------	-------	------------

Unguarded generative help can harm solo math learning; hallucinations contribute	FACT	High
Novices often fail to detect fluent errors	FACT	High
Appropriate reliance (not max trust) is the right design target	FACT (HF)	High
Apology scripts alone restore post-failure trust	SPECULATION / often false	Low (against) — Liu et al. wound
Explanatory repairs preferred over rote for many LLM errors	FACT (preference)	Medium–High
Preference among apology types for <i>hallucination</i> is unsettled	FACT	High
Competence-framed ownership + corrective action is right class for math slips	HYPOTHESIS	Medium–High
Triggered, packetized human escalation beats unlimited chat <i>and</i> no-human	HYPOTHESIS	Medium
Parent-visible repair notes raise WTP for hybrid SKUs	HYPOTHESIS	Medium
“Never hallucinates” is a viable brand claim	SPECULATION / false	Low (against)

LX.12 What this chapter kills

1. **Kill:** “Our AI doesn’t / won’t hallucinate” marketing.
2. **Kill:** Rote or empathy-only apology as the repair product.
3. **Kill:** Silent overwrite of wrong AI output without student-visible ownership when detection exists.
4. **Kill:** Thumbs-up / satisfaction as proof of trust repair.
5. **Kill:** Infinite free human escalation as default.
6. **Kill:** Permanent “never trust the AI” as brand (disuse).
7. **Wound:** Importing ADS trust-repair nulls or consumer GenAI “forgiveness” as ACT learning proof.
8. **Survive:** SAFE-REPAIR; competence ownership + checked correction + re-attempt; triggered packetized escalation; REPAIR-1...6 + QUAL.

Doctrine until data: After a hallucination, MindCraft sells **recoverable truth** — detect, own, correct, prove, escalate when scarce human attention is worth it — not warmth, not invincibility.

Part LXI — Measurement of Math Identity

Chapter status: Living evidence + measurement brief — Researcher tick 2026-08-01

Primary question: How should MindCraft *measure* mathematics identity — without mistaking a Likert bump for transformation, without inventing a proprietary Belief Score™, and without demand-effect theater?

Owners: Learning Science · Product analytics · Diagnostic / FEI · Parent messaging · Red Team

Commercial job: Ship a **SAFE-IDMEASURE** doctrine: validated subscale items + behavioral FEI corroboration + narrative probes — so marketing, experiments, and parent copy can claim *identity movement* only when measurement survives demand effects.

LXI.1 Why this chapter exists

The Constitution’s research question is identity transformation. Parts XXV (narrative identity / habit≠identity), XXXVII (expectancy-value), L (SAFE-CALIB), and LIX (SAFE-WTP) already ban empty confidence and survey-dollars-as-revenue. What was still missing is the **instrument layer**: which published scales are usable for Maya-aged students, what they actually factor into, how demand and social desirability corrupt them, and how MindCraft should triangulate without building a vanity Identity Meter™.

FOUNDER BELIEF under audit: Identity change is real and commercially load-bearing — but only if we measure it with instruments that have validity evidence *and* refuse to treat a post-session “I am a math person” checkbox as product success.

Claims we refuse as doctrine:

- 1. A proprietary “Math Identity Score™” as North Star.
- 2. Two HSLS-style items alone as causal proof of FEI.
- 3. Pre→post Likert lift after identity-themed copy as marketing evidence.
- 4. Competence/self-efficacy survey ≡ mathematics identity.
- 5. Ignoring recognition (others’ appraisals) because it is hard to productize.
- 6. Treating qualitative storytelling as “unscientific” *or* Likert as “the only science.”

LXI.2 Constructs (product language)

Construct	Research meaning	MindCraft analogue	Failure mode if misused
Mathematics identity	How a learner sees self <i>in relation to math</i> (thickening self-perception)	“Math person” / story world belonging without cosplay	Brand slogan without instrument
Interest	Desire/curiosity to engage math	Challenge accept; story-module return	Fun NPS ≡ identity

Recognition	Self + perceived others as “math person”	Tutor/parent brief language; peer opt-in	Forced public praise / leaderboard recognition
Competence / performance beliefs	Beliefs about understanding / doing math	Mastery graph; calibrated confidence	Self-efficacy survey alone as identity
Narrative identity	Reifying, endorsable, significant stories (Sfard & Prusak)	Notes / AAR “what changed about me as a learner”	Essay theater without transfer
Demand / SDR	Response bias toward expected / socially approved answers	Post-mission identity survey after identity ads	Survey lift sold as FEI
SAFE-IDMEASURE	Surviving measurement doctrine	See LXI.9	Identity Score™ productization

Operational definition (HYPOTHESIS): MindCraft claims *identity movement* only when (a) a **validated multi-item subscale** (interest + recognition + competence/performance) moves with preregistered effect size, (b) at least one **behavioral FEI** marker co-moves (challenge_accept, solo_transfer_pass, or destaked retry), and (c) a light **narrative probe** shows actual→designated story shift — not when any single channel moves alone.

LXI.3 What the field actually measures (not slogans)

FACT (SEM on college calculus, large N): Cribbs, Hazari, Sonnert & Sadler (2015, *Child Development*, 86(4), 1048–1062, doi:10.1111/cdev.12363) — FICSMath; more than 9,000 U.S. college calculus students. SEM of mathematics identity: competence/performance beliefs act largely *indirectly* on “math person” identity via **interest** and **external recognition**. Competence alone is **not** sufficient for identity development.

Product kill: Shipping only “I feel smarter at math” / self-efficacy meters as identity proof. Survive: design and messaging that create *interest* and *recognition*, not competence theater alone.

FACT (K–12 instrument validation): Cribbs & Utley (2023/2024, *Mathematics Education Research Journal*, 36, 767–789, doi:10.1007/s13394-023-00474-w) — N=1559 students grades 5–12 in two U.S. districts; adapted undergraduate items; expert content review + readability (Flesch–Kincaid ~grade 3); EFA/CFA support a **16-item** scale with three factors: **interest, recognition, competence/performance** (competence and performance again load together, echoing Cass et al., 2011). Fit indices in CFA: CFI≈0.97, RMSEA≈0.063, SRMR≈0.028. Negatively correlated with math anxiety (mAMAS) as concurrent validity.

Commercial implication: For Maya (high-school), the closest *validated* efficient instrument is this lineage — not an invented MindCraft rubric. Prefer licensed/adapted item sets with permission over DIY “identity KPIs.”

FACT (two-item national survey tradition): HSLS:09 and follow-ups include math-identity items of the form *I see myself as a math person / others see me as a math person* (Ingels et al., NCES HSLS documentation; Cronbach α reported ~0.84–0.88 in base/follow-up). Miller & Wang (2019, *Journal of Youth and Adolescence*, 48, 2038–2050, doi:10.1007/s10964-019-01115-x) use classroom-practice → motivational beliefs → mathematics identity mediation (N=525 sixth-graders); identity operationalized in that literature stream often collapses to short “math person”

self/other items.

Wound: Cribbs & Utley (2023) explicitly note that several large studies using NAEP/HSLs-style items **did not report full validity evidence** for the measure as used. Two items are useful for **screening and national comparison**, weak as sole **product experiment** endpoints.

FACT (science-identity ancestor): Carlone & Johnson (2007, *Journal of Research in Science Teaching*, 44, 1187–1218, doi:10.1002/tea.20237) — recognition, competence, performance as analytic lens for science identity (women of color). Hazari et al. added **interest** when moving to physics/math identity measurement traditions that Cribbs inherits.

FACT (narrative operationalization): Sfard & Prusak (2005, *Educational Researcher*, 34(4), 14–22, doi:10.3102/0013189X034004014) — identity as collections of **reifying, significant, endorsable** stories; learning as closing the gap between **actual** and **designated** identity narratives. Not a Likert substitute — a complementary analytic.

LXI.4 Demand effects, social desirability, and identity advertising

FACT (response styles): Kreitchmann, Abad, Ponsoda, Nieto & Morillo (2019, *Frontiers in Psychology*, 10:2309, doi:10.3389/fpsyg.2019.02309) — social-desirability and acquiescence biases distort Likert scores; ignoring them worsens validity; modeling SDR or forced-choice designs can help; each strategy has tradeoffs (reliability vs criterion validity).

HYPOTHESIS (MindCraft-specific demand): After a session that *markets* identity (“become a math person”), immediately asking “I am a math person” maximizes **demand characteristics** — students infer the desired answer. Risk rises when coaches, tutors, or parents use the same phrase as praise (Part XLVI expectancy ethics).

HYPOTHESIS: Separating **measurement windows** from **identity marketing copy** (e.g., measure at week boundaries, not mission-complete splash) reduces artifactual lift.

Kill: Pre/post “math person” Likert on the same screen as identity-themed celebration UI, sold as FEI evidence.

Survive: Blinded or temporally separated administration; behavioral corroboration; optional forced-choice / IRT-aware scoring later if scale volume justifies it.

LXI.5 What must not be confused with identity

Nearby construct	Why it is not identity	Constitution home
Self-efficacy	Can rise without “math person” story or recognition	XXV, Bandura
Calibration / confidence	Match of certainty to accuracy — honesty, not belonging	L / SAFE-CALIB
Mastery graph score	Performance evidence; necessary but not sufficient (Cribbs 2015)	Engine / XXI

Streak / habit	Cue-driven return ≠ identity thickening	XLIII / SAFE-HABIT
Anxiety reduction	Correlated (Cribbs & Utley) but distinct	XXIV, LVIII
Parent WTP	Purchasing ≠ student identity	LIX / SAFE-WTP

FOUNDER BELIEF (constrained): FEI still holds — fear→evidence→identity — but **identity** here means the Cribbs-factor bundle + narrative shift, not a mood.

LXI.6 Equity and recognition measurement (do not flatten)

FACT / tradition: Recognition is not only teacher praise; Martin and equity-oriented identity work stress how race, gender, and institutional positioning shape who gets to be seen as mathematical (see Part XXXVI equity audit; Carlone & Johnson recognition as gate).

HYPOTHESIS: Product “recognition” that is only AI cheerleading (“Great job, math person!”) will inflate Likert recognition subscales without changing peer/teacher/parent stories — a **false recognition** pathway.

Product rule: Prefer recognition that is (a) **earned** (tied to specific transferable performance), (b) **multi-source when possible** (tutor brief, parent note, optional peer — never forced leaderboard), (c) **non-trait** (task language per SAFE-EXPECTANCY).

LXI.7 Recommended MindCraft measurement stack

Tier A — Primary quantitative (product experiments): Adapt Cribbs & Utley (2023) **16-item / 3-factor** instrument (or a preregistered short form of the same factors) for 5–12; administer at baseline, 4 weeks, 8 weeks — **not** on every mission complete.

Tier B — Behavioral corroboration (always co-primary): challenge_accept, solo_transfer_pass, retry_120s after informative miss, destaked re-attempt after soft-wrong. Identity claim fails if Likert rises and these stay flat (XXV falsifier).

Tier C — Narrative light probe (qual, Appendix B adjacent): 2–3 prompts: actual story (“Right now, what kind of math learner are you?”), designated story (“Who are you becoming?”), recognition source (“Who has treated you as capable at math lately — and about what?”). Code for actual→designated gap closure (Sfard & Prusak), contamination→redemption themes (McAdams tradition in XXV).

Tier D — Screening only: HSLS-style two items for funnel diagnostics and parent dashboard *context* — never as sole RCT endpoint.

Explicit ban: No composite “Identity Score™” marketed to parents. Report **factor means + behavioral FEI**, or say nothing.

LXI.8 Experiments (IDM family)

ID	Contrast	Primary endpoints	Kill / survive rule
IDM-1	Cribbs 16-item at weeks 0/4/8 vs HSLS 2-item only	Factor α ; sensitivity to FEI intervention	2-item moves, 16-item flat → kill 2-item-as-proof
IDM-2	Measure at week boundary vs mission-complete splash	Δ interest/recognition; demand proxy	Splash-only lift → demand wound
IDM-3	Interest+recognition-focused UX vs competence-meter UX	Identity factors; challenge_accept	Competence UX lifts efficacy not identity → confirms Cribbs SEM
IDM-4	Earned tutor/parent recognition note vs AI “math person” praise	Recognition subscale; transfer	AI praise lifts Likert only → kill false recognition
IDM-5	Likert-only claim vs Likert+behavioral co-primary in experiment report	Pre-reg analysis plan adherence	Likert-alone “win” rejected
IDM-QUAL	10 Maya interviews: what “math person” means; demand cues	Codebook → item wording freeze	Founder phrase ≠ student phrase

Pre-reg (XXXIV): IDM-* identify **measurement validity and demand**, not “MindCraft raises identity by X% forever.”

Ties: EFF-, EVT-, CAL-*, WTP message arms (identity vs score packaging), HIST recognition equity.

LXI.9 SAFE-IDMEASURE stack (surviving doctrine)

1. **Use validated factors** — interest, recognition, competence/performance — not inventing a brand score.
2. **Triangulate** — subscale + FEI behavior + light narrative; one channel alone is insufficient.
3. **Separate measure from marketing** — no identity survey on identity celebration screens.
4. **Competence ≠ identity** — graph and efficacy gains without interest/recognition are incomplete.
5. **Recognition must be earned and multi-source-capable** — ban empty AI “math person” praise as the recognition product.
6. **Report factors, not a trademarked composite** — parent copy may say “students see themselves differently *and* take harder problems” only when both move.
7. **Guard demand/SDR** — preregister; consider timing, neutral framing, later forced-choice if needed.
8. **Longitudinal ownership** — single-session Δ is SPECULATION; 4–8 week windows are the minimum serious bar (feeds Part LXII).

LXI.10 Commercial implications (copy, product, positioning)

Marketing language that survives: “We track whether students *and* the people around them start treating them as someone who does hard math — and whether their practice behavior matches that story.”

Marketing language that dies: “+32% Math Identity Score™”; “Become a math person in one session”; “Our AI recognizes you as a genius.”

Positioning vs Khan / Duo / Brilliant / GPT tutors: Competitors can show accuracy, streaks, or fluency. MindCraft’s differentiated *claim* is identity — which **raises the measurement bar**, not lowers it. If we cannot measure better than a streak app’s mood survey, we do not get to use identity as the brand wedge.

Growth / instrumentation: Add IDM battery to research panel / experiment flags; do not ship Identity Meter on the default student HUD.

Parent WTP link (LIX): Identity packaging in conjoint is legitimate *only* if IDM-5-style evidence exists; otherwise sell FEI session proof and transfer, not identity poetry.

LXI.11 Confidence table

Claim	Label	Confidence
Competence/performance beliefs are insufficient alone for math identity; interest + recognition matter	FACT (SEM)	High
Validated 5–12 multi-item instrument exists (Cribbs & Utley 2023)	FACT	High
HSLs two-item forms are widely used but thin as sole experiment endpoints	FACT / HYPOTHESIS	High / Medium–High
Immediate post-mission identity Likert is demand-contaminated	HYPOTHESIS	Medium–High
AI praise alone produces false recognition subscale lift	HYPOTHESIS	Medium
Narrative probes add validity beyond Likert	HYPOTHESIS (theory-backed)	Medium
Proprietary Identity Score™ improves learning or WTP	SPECULATION / refuse	Low (against)
Single-session identity transformation is measurable and marketable	SPECULATION / refuse	Low (against)

LXI.12 What this chapter kills

1. **Kill:** Math Identity Score™ / Belief Score™-style composites as North Star.
2. **Kill:** Two-item “math person” pre/post as causal FEI proof.
3. **Kill:** Competence or confidence survey $\equiv$ identity.
4. **Kill:** Identity survey glued to identity marketing splash screens.
5. **Kill:** Empty AI “you’re a math person” as recognition product.
6. **Kill:** Single-session identity transformation claims.
7. **Wound:** Importing undergraduate-only instruments without 5–12 validity evidence.
8. **Survive:** SAFE-IDMEASURE; Cribbs-factor triangulation; IDM-1...5 + QUAL; behavioral co-primary.

Doctrine until data: MindCraft may *aim* at identity — but may **claim** identity movement only when validated interest/recognition/competence–performance subscales move **with** FEI behavior (and preferably narrative gap closure), under demand-aware administration.

Part LXII — Longitudinal Identity Change Models

Chapter status: Living evidence + horizon brief — Researcher tick 2026-08-02

Primary question: What does *honest* math-identity change look like at **8 / 26 / 52 weeks** — and which commercial claims survive fade-out, stability, and dimensional-comparison evidence?

Owners: Learning Science · Product analytics · Growth / retention · Parent messaging · Red Team

Commercial job: Ship a **SAFE-LONGID** doctrine: staged claim horizons so MindCraft can sell *direction and persistence*, not “math person in eight weeks” theater, and so subscription / package design matches how identity actually moves.

LXII.1 Why this chapter exists

Part LXI (SAFE-IDMEASURE) fixed *how* to measure identity. This chapter fixes *when* a movement counts as change vs noise, demand, or temporary lift. Founders and growth teams instinctively want a short horizon (“show identity by week 8 for renewals”). Developmental evidence is slower, stickier, and often **differentiating away from math** through high school — not smoothly ascending toward “math person.”

FOUNDER BELIEF under audit: Identity transformation is the product thesis — but durable identity change is a **multi-horizon** process. Selling L3 identity claims on L1 session windows (Part XXXIV claim ladder) is already banned; selling them on an **8-week Likert bump without persistence** is the next failure mode.

Claims we refuse as doctrine:

1. Eight-week subscale lift alone $\equiv$ transformation (marketable or scientific).
2. Linear “always becoming more of a math person” growth curves for Maya.
3. One-semester tutoring $\equiv$ occupational / STEM-pipeline identity.
4. End-of-program effects that ignore fade-out after support ends.

5. Stability evidence as proof that identity *cannot* change (fatalism).
6. Multi-year school-department reforms as templates for a 45-minute app session.

LXII.2 Constructs (product language)

Construct	Research meaning	MindCraft analogue	Failure mode if misused
Horizon	Follow-up window for an estimand	8w / 26w / 52w experiment clocks	Selling 52w outcomes on 8w data
Stability	Rank-order / mean identity often sticky	Baseline IDM battery	“Nothing moves → kill product” or “must move weekly”
Differentiation	Math vs verbal self-concepts diverge over grades	Challenge accept in <i>math</i> not generic hustle	Generic confidence copy
Persistence vs fade-out	Treatment effects that hold vs shrink after end	Retention after plan pause / summer	Celebrate end-of-program only
Process signal	Early FEI / interest / approach markers	challenge_accept, destaked retry	Process ≡ identity
SAFE-LONGID	Surviving longitudinal claim doctrine	See LXII.9	Forever-identity packaging

Operational definition (HYPOTHESIS): MindCraft may claim *identity movement at horizon H* only when SAFE-IDMEASURE triangulation holds **at H**, and for $H \geq 26$ weeks also shows **non-fade** relative to an earlier post-test (or preregistered maintenance criterion). Eight weeks may support *process / direction* claims, not L3 identity brand claims.

LXII.3 Developmental facts (not slogans)

FACT (dimensional comparison meta-analysis): Wan, Lauermaann, Bailey & Eccles (2021, *Psychological Bulletin*, 147(9), 867–889, doi:10.1037/bul0000340) — 142 samples, ~211k participants, 426 effect sizes. Correlations between math- and language-related motivational beliefs *decline* across grades (from roughly $r \approx 0.32$ in Grades 1–4 toward near-zero / slight negative in Grades 9–12). Students increasingly treat domains as **either/or**.

Commercial implication: High-school Maya is in the window where “I’m a words person, not a math person” consolidates. Product copy that treats identity as a generic confidence dial fights the developmental grain.

FACT (person-type trajectories across K–12): Wan, Tian, Guo & Lauermaann (2026, *Child Development*, doi:10.1093/chidev/aacag069) — four U.S. longitudinal samples ($Ns \approx 971$ –5,681; ages ~6–18). Annual classification into math person / language person / both-high / neither-high. **Both-high shrinks**; language-person increases; **math-person rises through middle school then plateaus in high school**; girls’ shift toward language is steeper. In samples with adult follow-up, high-school math-person status associates with more math-intensive occupations.

Product kill: Marketing that implies continuous year-on-year math-identity ascent through HS. Survive: **sustain and deepen** math identity in the HS plateau window — especially for students at risk of language-only sorting.

FACT (expectancy–value over years, not weeks): Simpkins, Davis-Kean & Eccles (2006, *Developmental Psychology*, 42(1), 70–83, doi:10.1037/0012-1649.42.1.70) — N=227; 5th-grade out-of-school math/science activity → 6th/10th-grade expectancies–values → high-school course taking beyond grades. Belief→choice chains run on **multi-year** clocks.

FACT (related classroom→belief→choice chain): Simpkins, Davis-Kean & Eccles / follow-on EVT longitudinal work (e.g., classroom experiences in middle school predicting later courses and aspirations; see also Wang lineage cited in LXI via Miller & Wang 2019) — identity-adjacent motivational beliefs are **path-dependent**, not session-local.

FACT (attainment value / “science person” as persistence predictor): Andersen & Ward (2014, *Science Education*, 98(2), 216–242, doi:10.1002/sce.21092) — HSLs:09 high-ability ninth-graders; science attainment value (“science person”-type identification) predicts STEM persistence plans across Black, Hispanic, and White subgroups (with utility operating differently by group). Snapshot is grade 9, but the *outcome* is planned persistence — a longitudinal *intent*.

Wound / transfer caution: HSLs items are thin (LXI); use as directional evidence that **identity-like attainment value predicts plans**, not as MindCraft’s instrument.

LXII.4 Stability is real — and so is change under the right ecology

FACT (college sample stability): Cribbs, Tassell, Hazari, Sadler & Sonnert (2022, *RIPEM*, 12(2), 68–91, doi:10.37001/ripem.v12i2.2923) — N=131 undergraduates; Wilcoxon signed-rank on global math-identity self-perceptions over ~three college years → **non-significant change** (identity reported as stable). Qualitative themes: formative experiences often involve *performing* and *teaching/sharing*; many believe one *can* become a math person while still harboring innate-ability beliefs.

Interpretation for MindCraft (HYPOTHESIS): Stability under business-as-usual schooling is the null. A tutoring product that expects weekly identity Likert jumps is fighting a sticky construct. Change requires **repeated recognition + interest + competence evidence** (Cribbs SEM 2015; LXI), not one delightful onboarding.

FACT (multi-year equitable departments): Boaler & Staples (2008, *Teachers College Record*, 110(3), 608–645, doi:10.1177/016146810811000302) — Railside longitudinal mixed-methods (~700 students across three CA high schools). Heterogeneous, reform-oriented department associated with larger achievement gains and narrowing of ethnic gaps by roughly the **end of year two** on several measures; students also reported more enjoyment / progression. Horn (2008, *Mathematical Thinking and Learning*, 10(3), 201–239, doi:10.1080/10986060802216177) — turnaround student identities emerge across **departmental worlds over years**, not a single classroom moment.

Kill: Citing Railside/Horn as proof that an edtech session architecture produces identity in weeks.

Survive: Identity ecology is **relational and sustained**; MindCraft’s analogue is tutor+parent recognition + FEI practice cadence over months — not a splash screen.

FACT (short belief interventions ≠ identity doctrine): Boaler, Dieckmann, Pérez-Núñez, Sun & Williams (2018, *Frontiers in Education*, 3:26, doi:10.3389/educ.2018.00026) — teacher-randomized MOOC (~six ~15-min sessions, Oct–Dec) shifted math *beliefs*/engagement and standardized scores vs control on a ~semester research clock. Useful as proof that **beliefs can move in months with math-specific messages** — but Constitution already bans empty mindset

posters; do **not** rebrand MOOC belief lifts as Cribbs-factor identity or FEI.

LXII.5 Fade-out: why 8-week wins die without maintenance

FACT (fade-out is normal): Bailey, Duncan, Cunha, Foorman & Yeager (2020, *Psychological Science in the Public Interest*, 21(2), 55–110, doi:10.1177/1529100620915848) — fade-out of educational intervention effects is **widespread**, often coexists with some persistence, and is substantive (not only measurement artifact). Persistence depends on skill type, institutional constraints, and **complementarities** with later environments.

FACT (earlier geometric-decline pattern): Bailey, Duncan, Odgers & Yu (2017, *Journal of Research on Educational Effectiveness*, 10(1), 7–39, doi:10.1080/19345747.2016.1232459) — post-treatment impacts commonly shrink within 1–2 years; early skill gains are poor guarantees of long-run advantage without supporting environments.

HYPOTHESIS (MindCraft-specific): An 8-week FEI + interest lift that is followed by **disengagement, exam-only cramming, or AI-crutch Solver use** will fade on identity factors by 26–52 weeks — even if week-8 IDM looks green.

Product implication: Retention design is not a growth vanity metric; it is **identity infrastructure**. Pause/summer plans need light maintenance cues (SAFE-HABIT subordinate to FEI), not streak guilt.

LXII.6 Horizon model for MindCraft (8 / 26 / 52)

Horizon	What can honestly move	Allowed commercial speech	Forbidden speech
~8 weeks	Process: challenge_accept, destaked retry, soft-wrong return; possible early interest/recognition subscale <i>direction</i> under SAFE-IDMEASURE	“Students start approaching harder problems and sticking after misses”	“Became a math person”; “identity transformed”; Identity Score™
~26 weeks	Triangulated factor movement + FEI co-primary; narrative actual→designated shift; first fade check vs week 8	“Over a semester, self-view <i>and</i> challenge behavior moved together” (only if true)	Guaranteed ACT points; occupational STEM destiny
~52 weeks	Persistence of week-26 gains; course/intent signals where observable; maintenance under reduced support	Annual renewals tied to <i>sustained</i> approach + transfer, not mood	“Forever math person after one year with us” without follow-up

FOUNDER BELIEF (constrained): Package tiers and parent CBC (Part LIX) should map to horizons: diagnostic + 8w process proof; semester identity-direction report; annual persistence review — never sell the annual claim with only an 8w dashboard.

SPECULATION (explicit): Exact week cutoffs are product conventions informed by school semesters and subscription cycles, not natural laws. Preregister windows; do not retrofit.

LXII.7 Analytic models (what “longitudinal” means in practice)

MindCraft does not need a 10-year panel to be rigorous. Minimum serious designs:

- 1. **Repeated measures / growth:** IDM factors at 0/8/26/52; random slopes; report **individual heterogeneity** (not only means).
- 2. **Latent transition / person-types (inspired by Wan 2026):** math / language / both / neither-style states — watch **both-high erosion** and math-person **plateau** risk.
- 3. **Joint process–identity models:** FEI events as time-varying covariates of interest/recognition (tests FEI pathway).
- 4. **Fade contrast:** end-of-active-support vs +8–12 weeks post-pause.
- 5. **Demand control:** never collect IDM on identity-celebration UI (LXI).

Pre-reg (XXXIV): LONG-* experiments identify **horizon-valid estimands**, not “MindCraft creates math people by $\Delta=X$.”

LXII.8 Experiments (LONG family)

ID	Contrast	Primary endpoints	Kill / survive rule
LONG-1	8w IDM+FEI vs waitlist / active control	Interest/recognition Δ ; challenge_accept	Factor $\uparrow$ FEI flat $\rightarrow$ kill identity claim; FEI $\uparrow$ factor flat $\rightarrow$ process-only speech
LONG-2	26w continuation vs 8w-then-pause	Factors at 26; fade from 8	Pause erases 8w lift $\rightarrow$ kill “8w locks identity”
LONG-3	Recognition-rich tutor/parent notes cadence vs AI praise cadence (extends IDM-4)	Recognition subscale; transfer; 26w hold	AI-only recognition fades faster $\rightarrow$ confirms false recognition
LONG-4	Maintenance micro-dose (1–2 FEI sessions/mo) vs full stop after 26w	52w factors + FEI	No maintenance $\rightarrow$ fade; informs annual package

LONG-5	Person-type transition coding (Wan-style) exploratory	Transition probs; gender split	Both-high collapse without math sustain → HS copy priority
LONG-QUAL	10 Maya + 10 parent interviews at 0/8/26: story change vs score obsession	Narrative codebook	Parent wants 8w identity certificate → WTP honesty constraint

Ties: IDM-, EVT-, FEI/SAFE-CRAFT, WTP package horizons, REPAIR trust persistence, HIST equity of who gets sustained recognition.

LXII.9 SAFE-LONGID stack (surviving doctrine)

1. **Stage claims by horizon** — 8w = process/direction; 26w = triangulated identity movement; 52w = persistence/maintenance.
 2. **Expect stickiness** — null is stable identity under weak ecology (Cribbs 2022); design for repeated FEI + earned recognition.
 3. **Respect differentiation** — HS students sort math vs language; sustain math identity, don't sell generic grit.
 4. **Plan for fade-out** — end-of-program green lights are incomplete; measure post-pause.
 5. **Do not launder mindset MOOCs or multi-year school reforms as app-session identity proof.**
 6. **Align packages to clocks** — CBC attributes and renewals match measurable horizons (LIX).
 7. **Heterogeneity first** — mean lifts can hide non-responders; report who moves.
 8. **Keep SAFE-IDMEASURE** — same instruments, longer clocks; no new Identity Score™.
-

LXII.10 Commercial implications (copy, product, positioning, growth)

Marketing language that survives: “In the first weeks we look for *approach under safety*. Over a semester we look for whether students’ *self-story and challenge behavior* move together. Over a year we look for whether that holds when life gets busy.”

Marketing language that dies: “Math person in 8 weeks”; “Guaranteed identity transformation this semester”; “Join and permanently become STEM”; linear forever-up charts on landing pages.

Positioning vs Khan / Duo / Brilliant / GPT tutors: Competitors optimize session accuracy, streaks, or fluency — short clocks. MindCraft’s wedge is **durable self-as-math-doer**. That only works if we **own the long clock** in measurement and packaging. If we market identity on Duo’s time horizon, we lose the wedge and inherit their vanity metrics.

Growth: Instrument cohort retention *as* identity maintenance; annual plans should include LONG-4-style maintenance, not only denser content. Churn after week 6 is not merely CAC failure — it is **hypothesis failure** for L3 claims.

Vision / future: The company roadmap should treat 26–52 week panels as first-class research infrastructure (research panel flag), not a nice-to-have after feature shipping.

LXII.11 Confidence table

Claim	Label	Confidence
Math/verbal motivational beliefs differentiate across school years	FACT (meta)	High
Math-person rates plateau in HS; both-high shrinks; language-person rises	FACT (multi-sample longitudinal)	High
Global math identity can be stable over years under status-quo college experience	FACT (small N)	Medium–High
Multi-year equitable school departments can shift achievement/identity ecology	FACT (case / mixed methods)	Medium–High (external validity to edtech low)
Educational intervention effects commonly fade without supporting environments	FACT (reviews)	High
8w triangulated lift without 26w hold is insufficient for L3 identity marketing	HYPOTHESIS / doctrine	Medium–High
Staged 8/26/52 product clocks are the right commercial mapping	FOUNDER BELIEF / HYPOTHESIS	Medium
MindCraft will produce 52w identity persistence with maintenance micro-dosing	SPECULATION	Low–Medium (test via LONG-4)

LXII.12 What this chapter kills

1. **Kill:** “Math person in 8 weeks” / single-horizon identity transformation marketing.
2. **Kill:** Treating end-of-program identity Likert as durable without fade checks.
3. **Kill:** Linear always-rising math-identity curves for high school.
4. **Kill:** Equating short mindset/belief MOOC effects with Cribbs-factor identity + FEI.
5. **Kill:** Citing Railside/Horn multi-year worlds as proof of session-scale app identity.
6. **Kill:** Annual “forever STEM identity” packaging without 52w persistence evidence.
7. **Wound:** Using HSLS two-item snapshots as longitudinal product proof (re-anchor LXI).
8. **Survive:** SAFE-LONGID staged horizons; differentiation-aware HS positioning; fade-aware retention; LONG-1...5 + QUAL.

Doctrine until data: Identity is the North Star *target*, not an 8-week badge. Ship process truth early; claim identity movement only on triangulated **semester+** clocks; claim persistence only when fade-out is measured and survived.

Part LXIII — ACT / Exam Pressure Special Case

Chapter status: Living evidence + affect-separation brief — Researcher tick 2026-08-02

Primary question: How is **high-stakes exam affect** (ACT/SAT/state tests) different from **learning-session affect** — and what product, copy, and metric rules follow so MindCraft does not confuse destaked FEI practice with exam readiness (or flood Maya with day-one timed sims)?

Owners: Learning Science · Practice / diagnostic UX · Parent messaging · Growth · Red Team

Commercial job: Ship a **SAFE-EXAM** doctrine: separate *learn* clocks from *perform under pressure* clocks; sell honest prep architecture, not guaranteed points or calm-app theater.

LXIII.1 Why this chapter exists

MindCraft already owns math anxiety (XXIV), desirable difficulties under threat (XLI / SAFE-DD), sleep/stress timing (XLVII), and graded approach ladders (LVIII / SAFE-EXPOSE). Parents still buy **exam outcomes**. Tutors still schedule **timed mixed sets**. Product still risks one of two failures:

- 1. **Flooding:** treat every practice minute as a mini-ACT (performance climate; WM choke; churn).
- 2. **Under-transfer:** optimize cozy FEI accuracy and market “exam-ready” without ever dosing pressure.

FOUNDER BELIEF under audit: Identity transformation happens in *learning affect* (safety → recoverable miss → earned win). College-gatekeeping happens in *exam affect* (stakes, clock, evaluation threat). Conflating them either destroys learning or lies about readiness.

Claims we refuse as doctrine:

- 1. Day-one full timed ACT / mixed flood labeled “exposure therapy” or “rigor.”
 - 2. Low-stakes blocked accuracy ≡ ACT-ready.
 - 3. Guaranteed ACT point lifts / 2σ / Bloom in any package (already SAFE-WTP / standing ban).
 - 4. Calm Score™ / “zero anxiety before test day” as North Star (SAFE-EXPOSE).
 - 5. Expressive-writing or one pre-test ritual as substitute for content + approach ladders.
 - 6. Stereotype-threat lab effects as unlimited marketing for “our UI closes the ACT gap.”
 - 7. Cortisol / stress-bias evidence as “tests are invalid so skip prep.”
-

LXIII.2 Constructs (product language)

Construct	Research meaning	MindCraft analogue	Failure mode if misused
-----------	------------------	--------------------	-------------------------

Learning affect	Destaked / mastery-climate emotion during acquisition	FEI session; soft-wrong; hide-correctness diagnostic	Treat as identical to test day
Exam / evaluation affect	Worry + physiological arousal under evaluative stakes	Timed mixed; mock ACT week; parent-visible score	Flood every session
Math anxiety (MA)	Domain-specific threat of math	Concept avoidance; Solver over-reliance	Collapse into “nerves”
Test anxiety (TA)	Evaluation anxiety across subjects	Clock panic; blanking; rank fear	Only calming apps
Pressure dose	Intentional stakes/time/mix intensity	Laddered timed sets after equip	Hardness-as-virtue
Readiness signal	Delayed mixed under mild→moderate pressure	transfer_pass / timed mix after SAFE-DD	Blocked green checks
SAFE-EXAM	Surviving exam-pressure doctrine	See LXIII.9	Score guarantees; flood/cozy extremes

Operational definition (HYPOTHESIS): A MindCraft path is *exam-honest* when (a) acquisition stays under SAFE-DD / SAFE-EXPOSE destake rules, (b) pressure is **dosed after equip**, (c) readiness claims require **delayed mixed** performance under declared stakes (not blocked accuracy), (d) parent copy separates “building the skill” from “proving under the clock,” (e) no package promises fixed ACT points.

LXIII.3 Exam pressure is real — and not the same as math anxiety

FACT (test anxiety meta): Hembree (1988, *Review of Educational Research*, 58(1), 47–77, doi:10.3102/00346543058001047) — 562 studies: test anxiety relates to poorer performance; inversely to self-esteem; treatments that reduce TA often accompany better test/GPA outcomes. Classic baseline: evaluation threat is a performance construct, not “motivation fluff.”

FACT (math anxiety meta): Hembree (1990, *JRME*, 21(1), 33–46, doi:10.5951/jresmetheduc.21.1.0033) — 151 studies: MA tied to achievement, attitudes, and **avoidance**. Prep products that only soothe without approach leave the avoidance loop intact (LVIII).

FACT (MA vs TA on math performance): Caviola, Toffalini, Giofrè, Mercader Ruiz, Szűcs & Mammarella (2022, *Educational Psychology Review*, 34, 363–399, doi:10.1007/s10648-021-09618-5) — 177 studies / ~906k participants. MA–math achievement overall $r \approx -0.30$; TA–math achievement $r \approx -0.23$. Both matter; MA is the stronger domain hitch. WM mediation of anxiety→math is present but **weak**.

Commercial implication: Parent copy that says only “we reduce test nerves” undersells the math-specific threat. Copy that says only “we fix concepts” undersells evaluation affect on test week. Product needs **both rails**: concept FEI *and* laddered exam affect — sequenced, not mashed.

FACT (WM under math stress): Ashcraft & Kirk (2001, *JEP: General*, doi:10.1037/0096-3445.130.2.224); Beilock (2008, *Current Directions in Psychological Science*, 17(5), 339–343, doi:10.1111/j.1467-8721.2008.00602.x); Beilock & Carr (2005, *Psychological Science*, 16(2), 101–105) — pressure/worry consumes WM; high-WM performers can choke when strategies rely on WM. Exam pressure is a **cognitive-resource** problem, not only a vibes problem.

LXIII.4 High-stakes weeks change bodies — scores are not pure learning meters

FACT (field cortisol × high-stakes tests): Heissel, Adam, Doleac, Figlio & Meer (2021, *Education Finance and Policy*, 16(2), 183–208, doi:10.1162/edfp_a_00306; NBER w25305) — low-income New Orleans students grades 3–8; saliva cortisol in testing weeks vs baseline weeks. Homeroom cortisol ~18% higher before high-stakes tests (boys ~35%). Large cortisol *increases or decreases* associated with ~0.40 SD lower-than-expected scores vs other academic performance — suggestive **stress bias** in score interpretation (correlational on the score link; still product-relevant).

Wound: Sample is K–8 high-stakes *school* tests, not ACT sitters. Mechanism transfers as caution; coefficients do not become “ACT –80 points for everyone.”

Product implication: Do not treat a student’s worst timed mock as pure “doesn’t know content.” Instrument **condition** (sleep, stakes frame, prior flooding) when diagnosing. Do not market “tests are unfair so skip practice” — parents buy access; MindCraft sells **honest preparation under known bias**, not nihilism.

LXIII.5 Pressure during *learning* can kill the testing effect

FACT (high-stakes quizzes blunt retrieval benefits): Hinze & Rapp (2014, *Applied Cognitive Psychology*, 28(4), 597–606, doi:10.1002/acp.3032) — high-stakes vs low-stakes quizzes: initial quiz accuracy similar, but **final retention better after low-stakes**; only low-stakes beat reread controls. Performance pressure during practice can disrupt encoding benefits of retrieval (retrieval-disruption account).

HYPOTHESIS (MindCraft): Running every practice card as a scored, parent-visible, timed ACT item will look “serious” and **underproduce learning** relative to destaked retrieval + later pressure doses.

Tie to SAFE-DD / SAFE-TIMING: Desirable difficulty ≠ maximum stakes. Pressure is a **transfer/context** variable after equip, not the acquisition default.

LXIII.6 Brief worry-offload helps some — it is not the product

FACT (expressive writing before exams): Ramirez & Beilock (2011, *Science*, 331(6014), 211–213, doi:10.1126/science.1199427) — brief pre-exam expressive writing about worries improved exam performance, especially for habitually test-anxious students (lab + classroom field experiments).

Kill: Rebranding MindCraft as a journaling app that “boosts ACT.” Writing is a **pre-performance micro-tool**, optional in SAFE-EXPOSE / SAFE-AAR brief — never a substitute for ontology-backed practice, FEI, or pressure ladders.

Survive: Optional 2–3 minute worry dump before *timed* modules for high-TA segment; measure whether timed accuracy rises without reducing later solo_transfer_pass.

LXIII.7 Stereotype threat: real lab phenomenon, fragile marketing claim

FACT (women’s math under threat): Spencer, Steele & Quinn (1999, *Journal of Experimental Social Psychology*, 35(1), 4–28, doi:10.1006/jesp.1998.1373) — describing a difficult math test as gender-diagnostic can depress women’s performance; reducing threat can eliminate the gap in those designs.

FACT / wound (overclaim): Stoet & Geary (2012, *Review of General Psychology*, doi:10.1037/a0026617) — many purported replications are weak or confounded; enthusiasm that stereotype threat *explains* the gender math gap outruns the clean evidence.

Commercial rule: Design against stereotype-cueing copy and tokenism (XXXVI). Do **not** sell “our story world removes stereotype threat → +X ACT.” Equity is a design constraint and HIST target, not a score guarantee.

LXIII.8 Two clocks: learn vs prove

Phase	Affect goal	Stakes / time	Allowed metrics	Forbidden claim
Acquire	Learning affect; FEI	Destaked; untimed or soft clock	challenge_accept, destaked retry, soft-wrong return	“Exam ready”
Consolidate	Mild evaluative spice	Short timed slices; criterion feedback	Near-mix accuracy; calibration	Leaderboard as NS
Prove (prep)	Exam affect dose	Declared timed mixed / mock	Delayed mixed under pressure; error taxonomy	Guaranteed composite points
Prove (real)	External gate	Official ACT/SAT	Post-hoc outcome study only	Causal brand claim without pre-reg

FOUNDER BELIEF (constrained): Gap-scan and hide-correctness diagnostic (C4 / SAFE-CALIB) stay in **Acquire** affect — they measure knowledge/confidence without turning onboarding into humiliation theater. Timed ACT sims belong **after** equip and approach ladders (SAFE-EXPOSE), not as day-one brand ritual.

SPECULATION: Optimal first full timed mixed set sits after ~2–4 weeks of concept FEI for a typical Maya — test via EXAM-1; do not hardcode without data.

LXIII.9 SAFE-EXAM stack (surviving doctrine)

1. **Separate affects** — learning affect $\neq$ exam affect; never optimize one metric for both.
2. **Destake to learn; dose to prove** — acquisition defaults low-stakes (Hinze & Rapp); pressure is scheduled.
3. **Ladder pressure** — time $\rightarrow$ mix $\rightarrow$ parent-visible mock; gate with SAFE-DD / SAFE-EXPOSE.
4. **Readiness $\neq$ blocked green** — claim exam-prep progress only with delayed mixed under declared stakes.
5. **Dual rail diagnosis** — miss under cozy $\neq$ miss under clock; log condition tags.
6. **Micro-tools subordinate** — expressive writing / brief AAR optional; never the hero feature.
7. **Anti-guarantee** — no ACT point promises; no 2σ ; package “prep architecture,” not outcomes.
8. **Anti-nihilism** — stress-bias evidence informs interpretation and equity, not “skip the exam.”
9. **Stereotype-safe copy** — no threat-cueing; no threat-removal score claims.
10. **Parent speech hygiene** — “building under safety” vs “proving under the clock” in WTP messages (LIX).

LXIII.10 Experiments (EXAM family)

ID	Contrast	Primary endpoints	Kill / survive rule
EXAM-1	Early timed-flood vs destake-then-dose	4w delayed mixed; churn; challenge_accept	Flood $\uparrow$ early mock, $\downarrow$ delayed / $\uparrow$ churn $\rightarrow$ kill flood-default
EXAM-2	Parent-visible daily score vs private criterion	Learning retention; anxiety item; retention	Visible score $\uparrow$ TA, $\downarrow$ Hinze-style retention $\rightarrow$ kill public scoreboard default
EXAM-3	Optional expressive writing before timed module	Timed accuracy; transfer next session	Helps HTA only $\rightarrow$ segment; helps none $\rightarrow$ demote
EXAM-4	Pressure dose after SAFE-EXPOSE ladder vs same dose without ladder	Timed mix; approach; SUDS/expectancy	Ladder wins $\rightarrow$ keep; equal $\rightarrow$ simplify
EXAM-5	Readiness CTA: blocked accuracy vs delayed mixed under mild time	Parent trust; false-ready rate	Blocked CTA predicts mock fail $\rightarrow$ kill blocked-as-ready
EXAM-QUAL	10 Maya + 10 parent: “ready” language meanings	Codebook: score vs approach vs calm	Parents equate calm $\equiv$ ready $\rightarrow$ copy correction

Ties: SAFE-DD, SAFE-EXPOSE, SAFE-CALIB, SAFE-TRANSFER, SAFE-WTP, SAFE-TIMING, LONG horizons for post-exam identity (don’t claim identity from one ACT sit).

LXIII.11 Commercial implications (copy, product, positioning, growth)

Marketing language that survives: “We build skill under conditions where working memory can work — then we practice *proving* under the clock, on purpose.” / “Low-stakes struggle first; timed mixed when the ladder says so.” / “We won’t sell you a point guarantee; we will show you readiness signals that include pressure.”

Marketing language that dies: “ACT +4 guaranteed”; “Day-one real exam conditions”; “Become calm and the score follows”; “Our AI removes test anxiety”; “Practice accuracy means test-day ready.”

Positioning vs Khan / Duo / Brilliant / GPT tutors: Khan/Duo optimize volume and streaks under mostly low evaluative threat; ChatGPT optimizes fluent help (crutch risk); Brilliant optimizes delight. MindCraft’s wedge is **recoverable struggle** → **competence evidence** → **solo transfer**, then an explicit **pressure ladder** competitors either skip (too cozy) or flood (too harsh). Own the sequencing story.

Growth: Instrument `timed_mix_attempted`, `pressure_dose_level`, `mock_vs_expected`, and `false-ready` (blocked-green then mock-fail). Do not grow on raw timed-item counts. Exam-season campaigns sell **architecture** (ladder + diagnosis), not miracle composites.

Vision: Future ACT track is a first-class *affect curriculum* parallel to the concept ontology — pressure as a designed dimension, not an accidental byproduct of “serious mode.”

LXIII.12 Confidence table

Claim	Label	Confidence
TA and MA both negatively relate to math achievement (MA stronger in large meta)	FACT	High
Evaluative pressure can consume WM and induce choke	FACT	High
High-stakes weeks can shift cortisol; large swings associate with worse-than-expected scores	FACT (field; score link correlational)	Medium–High
High-stakes practice quizzes can erase retrieval-learning benefits despite equal quiz scores	FACT (lab)	High
Expressive writing can help some high-TA students on exams	FACT	Medium–High
Stereotype threat can depress lab math performance; gap-explanation overclaim is unsafe	FACT + wound	Medium
Destake-then-dose beats flood-default for MindCraft retention + delayed mixed	HYPOTHESIS	Medium–High

SAFE-EXAM sequencing is the right commercial mapping for ACT parents	FOUNDER BELIEF / HYPOTHESIS	Medium
Exact week of first full mock is universal	SPECULATION	Low (EXAM-1)

LXIII.13 What this chapter kills

1. **Kill:** Day-one timed ACT flood as default pedagogy or “exposure.”
2. **Kill:** Blocked / cozy accuracy marketed as exam-ready.
3. **Kill:** Guaranteed ACT points / 2σ / Bloom in exam-season copy.
4. **Kill:** Calm-to-zero / Calm Score™ as exam North Star.
5. **Kill:** Journaling / expressive writing as the product.
6. **Kill:** “We eliminate stereotype threat → score lift” marketing.
7. **Kill:** Using stress-bias evidence to sell anti-prep nihilism.
8. **Wound:** Treating K–8 cortisol coefficients as ACT effect sizes.
9. **Survive:** Dual-rail MA+TA framing; destake-to-learn / dose-to-prove; EXAM-1...5; parent speech split; pressure as designed curriculum dimension.

Doctrine until data: Exam pressure is a **special case of affect**, not a synonym for math anxiety and not a synonym for desirable difficulty. Ship learning under safety; schedule proving under the clock; sell the architecture, never the composite.

Part LXIV — Multilingual / ELL Math Identity

Chapter status: Living evidence + language–math separation brief — Researcher tick 2026-08-02

Primary question: How do **language load** and **math load** come apart for English learners / multilingual students — and what product, copy, and metric rules stop MindCraft from misdiagnosing identity, dumbing down math, or selling translation theater?

Owners: Learning Science · Question bank / formats · Story worlds (equity) · Diagnostic UX · Parent messaging · Red Team

Commercial job: Ship a **SAFE-ELL** doctrine: measure and scaffold linguistic access without collapsing “math person” into English fluency; sell equitable *access to mathematical practice*, not simplified forever tracks or “we translate so identity follows.”

LXIV.1 Why this chapter exists

MindCraft already owns equity of story worlds (XXXVI), identity measurement (LXI), exam pressure (LXIII), and format-tagged practice. A large and growing share of U.S. secondary students are classified as English learners (ELs) or are multilingual without the label. Product still risks three failures:

1. **Confound:** treat wrong answers on dense English word problems as proof of weak *math* identity / mastery.
2. **Deficit curriculum:** strip context and discourse until only procedures remain — then call it “ELL support.”
3. **Access theater:** auto-translate, glossaries, or bilingual badges marketed as identity transformation.

FOUNDER BELIEF under audit: Mathematical identity is participation in *mathematical practices* (reasoning, representing, arguing) — not conversational English polish. Language can gate that participation without defining the student’s mathematical self.

Claims we refuse as doctrine:

1. “Math is language-free, so EL gaps are ability gaps.”
2. Forever-simplified / procedure-only tracks as equity.
3. Vocabulary flashcards or translation alone as math-identity product.
4. Auto-translate / dual-language UI as proof MindCraft “closes the EL gap.”
5. Homogeneous “ELL” persona in copy (one level, one home language, one need).
6. Using EL status as a soft excuse to skip challenge, FEI, or pressure ladders.
7. Marketing “story worlds fix belonging for bilingual kids” without HIST / ELL experiments.

LXIV.2 Constructs (product language)

Construct	Research meaning	MindCraft analogue	Failure mode if misused
Language load	Linguistic complexity demanded to parse/produce the task	Dense word_problem stems; idioms; passive voice	Call it “math hardness”
Math load	Conceptual / procedural demand of the mathematics	Ingredient DAG; format; mix	Ignore language confound
BICS vs CALP	Conversational fluency vs academic language	Chatty UI vs problem register	Assume chatty English ⇒ ready for ACT stems
Language as resource	Home language / multimodal resources support math	Optional L1 notes; diagrams; gestures analogue in UI	Token “hola” branding
Language as deficiency	Treat EL status as lack	Remedial forever track	Identity colonization as “catch-up English”
Differential item functioning (DIF)	Items unfairly harder for ELs at equal math proficiency	Wordy ACT-style items	Bank blindness

SAFE-ELL	Surviving multilingual doctrine	See LXIV.9	Translate-theater; deficit tracks
----------	---------------------------------	------------	-----------------------------------

Operational definition (HYPOTHESIS): A MindCraft path is *language-honest* when (a) diagnostics and outcomes can separate language-parse fails from math-concept fails, (b) scaffolds reduce *unnecessary* linguistic complexity without removing mathematical discourse opportunities, (c) story wrap does not add stereotype-cueing or opaque cultural idioms as load, (d) identity claims never equate English fluency with being a “math person,” (e) challenge and transfer standards stay intact (SAFE-TRANSFER / SAFE-EXAM).

LXIV.3 The EL–math score gap is real — and partly linguistic access

FACT (NAEP EL gaps): NCES Condition of Education — Mathematics Performance (using 2022 NAEP): grade 4 EL average 216 vs non-EL 239 (**23-point** gap); grade 8 EL 241 vs non-EL 277 (**36-point** gap). Gaps are persistent across assessment years since disaggregation began (1996 for grades 4/8). Grade 12 (2019): EL 111 vs non-EL 152 (**41-point** gap). Source: NCES COE indicator CNC / Digest table 222.12.

Wound: NAEP gaps confound language classification, opportunity to learn, SES, and math proficiency. They are **market and equity facts**, not proof that ELs “aren’t math people,” and not MindCraft effect sizes.

Commercial implication: Parent and school copy can name the access problem honestly. Do not sell “we close the 36-point NAEP gap.” Sell architecture that reduces *construct-irrelevant* language load while keeping math demand.

LXIV.4 Linguistic complexity changes who the item measures

FACT (linguistic modification helps): Abedi & Lord (2001, *Applied Measurement in Education*, 14(3), 219–234, doi:10.1207/S15324818AME1403_2) — 1,174 grade-8 students; NAEP math items vs parallel **linguistically modified** items. ELs scored lower than proficient English speakers overall; modification produced small but significant gains, with larger benefits for ELs, low-SES, and lower-track math students. Preference interviews favored revised wording.

FACT (DIF on word problems): Martiniello (2008, *Harvard Educational Review*, 78(2), 333–368, doi:10.17763/haer.78.2.70783570r1111t32) — linguistic complexity of MCAS grade-4 math word problems associated with DIF against ELs relative to non-ELs of **comparable math proficiency**; think-alouds show reading-comprehension traps (syntax, unfamiliar vocabulary, dense noun phrases) that are not the math construct.

HYPOTHESIS (MindCraft): A wrong answer on a high-language-load *word_problem* after a correct *symbolic_expression* / *diagram* item on the same ingredient is more often a **language-gate** than a concept hole — and treating it as identity evidence (“I’m bad at math”) is product malpractice.

Product implication: Prefer format triangulation (already in bank formats) before mastery or identity updates. Flag high linguistic-complexity stems in authoring. Offer a **language-simplified parallel** for practice (same math, lower incidental English) without making it the only track.

LXIV.5 Conversation fluency ≠ academic math language

FACT (BICS/CALP distinction): Cummins (1979, *Working Papers on Bilingualism*, No. 19, 121–129) introduced **BICS** (basic interpersonal communicative skills) vs **CALP** (cognitive/academic language proficiency) to warn that conversational L2 fluency is not academic register command. Restated in Cummins (2008, in Hornberger, *Encyclopedia of Language and Education*, doi:10.1007/978-0-387-30424-3_36).

Commercial kill: Onboarding that hears fluent chat English and assumes ACT word-problem readiness. Conversely: withholding hard math because a student is still building CALP. Support **register access** (glosses of academic terms, multimodal stems) while keeping cognitive demand.

LXIV.6 Sociocultural turn: participation and resources, not vocab lists

FACT (situated bilingual math): Moschkovich (2002, *Mathematical Thinking and Learning*, 4(2–3), 189–212, doi:10.1207/s15327833mtl04023_5) — moves beyond vocabulary-acquisition views toward participation in mathematical practices; competence includes multimodal and bilingual resources, not only formal English lexicon.

FACT (equitable practice principles): Moschkovich (2013, *Journal of Urban Mathematics Education*, 6(1), doi:10.21423/jume-v6i1a204) — principles/guidelines: treat language as **resource not deficiency**; emphasize academic achievement (not only English learning); support discourse about math ideas (not low-level language drills as the main diet); draw on multiple classroom resources (objects, drawings, graphs, gestures) and home languages.

FACT (field synthesis): de Araujo, Roberts, Willey & Zahner (2018, *Review of Educational Research*, 88(6), 879–919, doi:10.3102/0034654318798093) — review of 75 empirical K–12 EL×math studies (2000–2015) through a sociocultural lens; documents persistent marginalization patterns and the need for cumulative, not only case-study, research.

FACT (task selection wound): de Araujo (2017, *Curriculum Inquiry*, doi:10.1080/03626784.2017.1368351) — secondary teachers accommodating ELLs often selected **repetitive, procedure-focused, context-stripped** tasks, driven by beliefs that ELLs as a homogeneous group cannot handle richer work.

Kill: MindCraft “ELL mode” that only serves blocked procedures and never story/discourse/transfer. That is deficit curriculum dressed as care — and it starves identity (LXI recognition + competence evidence).

Survive: Multimodal + optional L1 support + full mathematical practice (FEI, soft-wrong, transfer).

LXIV.7 Story worlds and identity: equity constraint, not EL miracle

Tie to XXXVI / SAFE-IDMEASURE: Story wrap can support belonging or add idiom load and tokenism. Bilingual characters without mathematical agency are costume. Empty AI praise in simplified English (“you are a math person”) remains banned (LXI).

HYPOTHESIS: For multilingual Maya, recognition that counts is **being heard as a mathematical thinker** (even with imperfect English or code-switching notes), not celebrating English progress as math progress.

SPECULATION: Optional student L1 annotation on Solver traces may raise challenge_accept for mid-proficiency ELs without hurting delayed transfer — test in ELL-3; do not ship as brand without data.

LXIV.8 Two loads: diagnose separately

Signal	Likely load	Product move	Forbidden interpretation
Fails wordy stem; passes symbolic/diagram twin	Language	Offer simplified parallel; gloss; do not slash mastery hard	“Not a math person”
Fails all formats on ingredient	Math	FEI / SAFE-MISCON route	“Just a language issue — skip concept”
Fluent chat; fails academic stems	CALP gap	Register scaffolds; keep math demand	Assume ready for ACT flood
Avoids word problems only	Threat + language	SAFE-EXPOSE ladder on language-load situations	Forever avoid word problems
Parent wants “English tutoring branded as math”	Market pull	Separate message: math practice with language access	Rebrand as ESL app

FOUNDER BELIEF (constrained): Hide-correctness diagnostic (C4) must not silently treat language-gated misses as calibrated underconfidence in *math*. Prefer format-balanced probes for EL-classified or self-identified multilingual students.

LXIV.9 SAFE-ELL stack (surviving doctrine)

1. **Separate loads** — language load $\neq$ math load; instrument both when possible.
 2. **Language as resource** — optional L1 / multimodal supports; never deficiency branding.
 3. **Modify incidental English, not math** — Abedi-style linguistic simplification for practice parallels; keep discourse opportunities (Moschkovich).
 4. **No deficit track** — ban procedure-only forever “ELL path” (de Araujo task-selection wound).
 5. **Format triangulation before identity/mastery verdicts** — especially on word_problem.
 6. **Heterogeneity** — never one ELL persona; ask language history / preference lightly.
 7. **Anti-theater** — translation/glossary $\neq$ identity; $\neq$ NAEP-gap closure claim.
 8. **Anti-excuse** — EL status does not waive FEI, transfer, or SAFE-EXAM dosing — sequence scaffolds, don’t skip.
 9. **Story equity** — no idiom-heavy unexplained slang as default; no token bilingual costume (XXXVI).
 10. **Parent speech hygiene** — “access to show what they know” vs “easier math forever.”
-

LXIV.10 Experiments (ELL family)

ID	Contrast	Primary endpoints	Kill / survive rule
ELL-1	High-language stem vs linguistically modified twin (same math)	Accuracy delta; time; self-report “stuck on words”	Large EL-only delta → keep twins; no delta → deprioritize authoring cost
ELL-2	Format-balanced diagnostic vs word-problem-heavy diagnostic	Mastery misclassify rate; identity item demand	Word-heavy ↑ false weak → require balance for multilingual
ELL-3	Optional L1 note / gloss on vs off	challenge_accept; delayed solo_transfer_pass	Helps approach, hurts transfer → redesign; helps both → ship optional
ELL-4	Procedure-only “support” track vs full multimodal practice	4w transfer; belonging/recognition items	Procedure track ↓ transfer / identity → kill deficit default
ELL-5	Auto-translate banner vs linguistic simplify + diagram	Trust; accuracy; reattempt	Translate-only wins UX but loses transfer → kill as sole solution
ELL-QUAL	10 multilingual Maya (+ parents): “math person” vs “English hard”	Codebook: load attribution	Students blame self-as-dumb for word traps → copy + UI fix

Ties: SAFE-TRANSFER, SAFE-MISCON, SAFE-IDMEASURE, SAFE-EXAM (ACT stems are high CALP), SAFE-EXPOSE (language-load as situation tier), XXXVI HIST-EQ.

LXIV.11 Commercial implications (copy, product, positioning, growth)

Marketing language that survives: “We separate *reading the problem* from *doing the math* — so language never masquerades as destiny.” / “Support for multilingual learners means access to full mathematical practice, not a forever-easy track.” / “Same math, clearer language when needed — then prove it in more than one format.”

Marketing language that dies: “Math is universal, language doesn’t matter”; “ELL mode = easier problems”; “We translate, so the gap closes”; “Fluent English means ACT-ready”; “Bilingual avatar = belonging solved.”

Positioning vs Khan / Duo / Brilliant / GPT tutors: Khan/Duo often add language demand without diagnosing it; GPT tutors can drown ELs in fluent English monologue (SE risk, XL) or over-simplify into procedures; Brilliant’s delight still assumes English UI literacy. MindCraft’s wedge: **format + linguistic-access scaffolding inside identity-safe struggle** — recoverable misses that are honestly classified.

Growth: Instrument language_simplify_used, format_mismatch_flag, self_report_word_stuck, and EL/multilingual opt-in (privacy-sensitive). Do not grow on “ELL engagement minutes.” School sales: lead with **fairer**

evidence of what students know, not miracle gap closure.

Vision: Question bank and story modules treat linguistic complexity as a first-class authoring dimension beside format and concept — parallel to pressure as a designed affect dimension (LXIII).

LXIV.12 Confidence table

Claim	Label	Confidence
Persistent NAEP EL vs non-EL math score gaps	FACT	High
Linguistic modification can improve math-item performance for ELs / lower-track students	FACT	High
Linguistic complexity can create DIF against ELs at comparable math proficiency	FACT	High
Conversational fluency $\neq$ academic language proficiency (BICS/CALP)	FACT (framework; widely used)	High
Equitable EL math instruction emphasizes resource view + discourse, not vocab-only	FACT (design principles / review)	High
Teachers often over-accommodate with procedure-stripped tasks	FACT (small qualitative)	Medium
Format triangulation reduces false “weak math” labels for multilingual users	HYPOTHESIS	Medium–High
SAFE-ELL is the right commercial mapping for MindCraft	FOUNDER BELIEF / HYPOTHESIS	Medium
Optional L1 notes raise approach without hurting transfer	SPECULATION	Low (ELL-3)

LXIV.13 What this chapter kills

1. **Kill:** Math-is-language-free $\rightarrow$ EL gap = ability/identity verdict.
2. **Kill:** Forever procedure-only / context-stripped “ELL support” tracks.
3. **Kill:** Vocabulary drills or auto-translate as the math-identity product.
4. **Kill:** “We close the NAEP EL gap” / guaranteed equity lift marketing.
5. **Kill:** Homogeneous ELL persona and stereotype-cueing bilingual costume.

6. **Kill:** Using EL status to waive challenge, FEI, transfer, or exam dosing.
7. **Wound:** Treating NAEP gaps as pure language or pure math — they are entangled.
8. **Survive:** Load separation; linguistic modify without math dilute; format triangulation; Moschkovich resource/discourse stance; ELL-1...5; parent speech split.

Doctrine until data: Multilingual math identity work is **access to mathematical practice under honest load diagnosis** — not English tutoring in a math skin, and not easier math forever. Scaffold the language gate; keep the math climb.

Part LXV — Gender & Math Stereotypes Update

Chapter status: Living evidence + stereotype/identity brief — Researcher tick 2026-08-02

Primary question: What does *current* evidence (not 1990s posters) say about gender, math performance, stereotypes, and identity — and what copy, product, and metric rules stop MindCraft from selling innate-gap myths, threat-removal score miracles, or pink-STEM cosplay as transformation?

Owners: Learning Science · Story worlds (equity) · Identity measurement · Parent messaging · Diagnostic / exam UX · Red Team

Commercial job: Ship a **SAFE-GENDER** doctrine: treat mean performance parity and stereotype/interest gaps as separable facts; design belonging without essentialism; never sell “we fixed stereotype threat so scores rose.”

LXV.1 Why this chapter exists

MindCraft already owns belonging/anxiety (XXIV), expectancy-value (XXXVII), equity of story worlds (XXXVI), identity measurement (LXI), exam pressure (LXIII), and ELL load separation (LXIV). Gender stereotypes remain a marketing landmine and a product design constraint:

1. **Ability myth:** parents and culture still narrate “boys are better at math” even when mean school performance is near parity.
2. **Threat theater:** early stereotype-threat (ST) findings get oversold as a turnkey score fix.
3. **Cosplay equity:** pink branding, “girl power” badges, or token women-in-STEM stickers sold as identity transformation.

FOUNDER BELIEF under audit: Mathematical identity for Maya is earned participation + recognition + competence evidence (LXI) — not gendered packaging. Stereotypes can gate *interest, belonging, and course choice* even when average test scores look equal.

Claims we refuse as doctrine:

1. Innate male math superiority as product premise or parent copy.
2. “Stereotype threat explains the gender STEM gap” as a settled marketing claim.
3. ST-removal / affirmation as guaranteed ACT points.
4. Girl-washed UI or “for her” tracks as the identity product.
5. Homogeneous “girls in STEM” persona (ignores race, class, ELL, interest).

6. Using gender to soften challenge, FEI, or transfer standards.
7. “We close the gender STEM gap” without measured mechanisms and fade checks.

LXV.2 Constructs (product language)

Construct	Research meaning	MindCraft analogue	Failure mode if misused
Mean performance gap	Average score difference by gender	Bank accuracy; ACT practice	Inflate tiny ds into destiny
Variance / tail gap	Male excess at extremes (debated)	High-end challenge path	Summers-style innate marketing
Stereotype (explicit)	Stated “math is for boys” beliefs	Parent/student surveys	Treat as only signal
Stereotype (implicit)	IAT-style gender–science associations	Not productized	Fake IAT gamification
Stereotype threat	Cue → underperformance under evaluative pressure	Exam-pressure UX (LXIII)	Miracle intervention brand
Cultural field stereotypes	Who belongs in CS/eng culture	Story characters; tutor tone	Token costume without practice
Relative strength / interest	Intra-individual math vs verbal preference	Course / path choice	Essentialize preferences as biology
SAFE-GENDER	Surviving gender×math doctrine	See LXV.9	Pink STEM; ST→score; innate gap

Operational definition (HYPOTHESIS): A MindCraft experience is *gender-honest* when (a) it does not encode ability stereotypes in copy or character casting, (b) it separates performance evidence from interest/identity evidence, (c) it designs belonging via recognition of mathematical practice (not gendered aesthetics), (d) any threat/belonging intervention is pre-registered and not sold as score insurance, (e) challenge and transfer standards stay intact across gender.

LXV.3 Mean math performance: parity is the update

FACT (gender similarities — broad): Hyde (2005, *American Psychologist*, 60(6), 581–592, doi:10.1037/0003-066X.60.6.581) reviewed 46 meta-analyses and advanced the **gender similarities hypothesis**: males and females are similar on most psychological variables; overstated difference claims carry costs.

FACT (school math scores near zero d): Hyde, Lindberg, Linn, Ellis & Williams (2008, *Science*, 321(5888), 494–495, doi:10.1126/science.1160364) analyzed U.S. state assessments covering >7 **million** students; weighted mean gender effect size $\approx d = 0.0065$ (grades 2–11) — consistent with no meaningful mean difference. NAEP hard / complex

items at grade 12 averaged about $d \approx 0.07$ (trivial).

FACT (meta update): Lindberg, Hyde, Petersen & Linn (2010, *Psychological Bulletin*, 136(6), 1123–1135, doi:10.1037/a0021276) — overall gender difference in math performance ≈ 0 ; small male advantages appear in some high-school / high-performing slices (e.g., reported ds on the order of ~ 0.2 – 0.4 in those slices — still not “boys own math”).

Commercial implication: Parent copy can say **school math scores no longer support a boys-are-better story**. Do not sell MindCraft as closing a huge mean gap that largely is not there. Sell the remaining problem: **interest, confidence, recognition, course choice, and stereotype climate**.

Wound: Parity on means does not erase STEM major / career segregation, anxiety gaps, or high-tail debates. Parity $\neq$ identity solved.

LXV.4 Affect and attitudes still gendered — update the 1990s carefully

FACT (attitudes/affect meta — classic still cited): Hyde, Fennema, Ryan, Frost & Hopp (1990, *Psychology of Women Quarterly*, 14(3), 299–324, doi:10.1111/j.1471-6402.1990.tb00022.x) found gender differences in math attitudes/affect often larger than in performance — stereotypes, confidence, and anxiety channels.

HYPOTHESIS (MindCraft, tying XXXVII / LXI): For many Mayas, the binding constraint is **expectancy + attainment value + recognition**, not a missing lecture. Stereotype climate can tax those without moving mean item accuracy much in low-stakes practice.

Kill: “Girls underperform on our bank, so we need easier girl mode.” Prefer diagnose anxiety, format, language load, and mastery climate first.

LXV.5 Implicit stereotypes predict gaps — not a MindCraft IAT product

FACT (nation-level IAT $\leftrightarrow$ TIMSS gaps): Nosek et al. (2009, *PNAS*, 106(26), 10593–10597, doi:10.1073/pnas.0809921106) — $\sim 70\%$ of $>500k$ gender–science IATs across 34 countries showed science=male associations; **nation-level implicit stereotypes predicted** 8th-grade TIMSS science/math sex differences (science: $r \approx 0.60$); explicit self-report stereotypes did not add predictive validity.

Wound: Ecological / correlational; IAT samples are volunteer web-takers; does not identify MindCraft causal paths. Do **not** ship a “stereotype IAT score” as identity metric.

Survive: Cultural stereotype climate is real enough to design against in story casting, tutor language, and parent speech — without psychologizing students via IAT gamification.

LXV.6 Stereotype threat: origin real, school-age literature wounded

FACT (classic college ST): Spencer, Steele & Quinn (1999, *Journal of Experimental Social Psychology*, 35(1), 4–28, doi:10.1006/jesp.1998.1373) — women underperformed men when a math test was described as showing gender differences; performance equalized under gender-fair framing (highly identified college samples).

FACT (children/adolescents meta — small + bias risk): Flore & Wicherts (2015, *Journal of School Psychology*, 53(1), 25–44, doi:10.1016/j.jsp.2014.10.002) — 47 effect sizes; mean effect $\approx g = -0.22$; **none of tested moderators significant**; multiple signs of **publication bias**; called for large pre-registered replications.

FACT (large null registered report): Flore, Mulder & Wicherts (2018, *Comprehensive Results in Social Psychology*, 3(1), doi:10.1080/23743603.2018.1559647) — Dutch high schools, $N = 2,064$, ages ~13–14; **no overall ST effect** on girls’ math scores; no moderated ST effects for domain/gender ID, math anxiety, or difficulty in that sample.

FACT (U.S. school-age nulls): Ganley, Mingle, Ryan, Ryan, Vasilyeva & Perry (2013, *Developmental Psychology*, 49(10), 1886–1897, doi:10.1037/a0031412) — three studies, $N = 931$; **no evidence** school-age girls’ math performance was impacted by ST activation methods from implicit to explicit; discuss publication-bias risk in the child ST literature.

Doyle & Voyer (2016, *Learning and Individual Differences*, 47, 103–116, doi:10.1016/j.lindif.2015.12.018) — meta-analysis of stereotype *manipulations* on math/spatial tests; useful as additional ST literature, not as MindCraft score insurance.

Commercial kill (already foreshadowed in LXIII standing bans): Do not market “we remove stereotype threat → +ACT points.” Destake-to-learn / dose-to-prove remains the exam doctrine; ST cues are one pressure channel among many, not a magic lever.

Survive: Avoid **gratuitous stereotype cues** in product (gendered ability jokes, “even for girls,” male-only genius tropes). That is hygiene, not an efficacy claim.

LXV.7 Cultural stereotypes as gatekeepers (interest $\neq$ ability)

FACT (CS/eng culture stereotypes): Cheryan, Master & Meltzoff (2015, *Frontiers in Psychology*, 6:49, doi:10.3389/fpsyg.2015.00049) — stereotypes of computer science/engineering as male, isolated, machinery-focused, “brilliance”-coded gatekeep girls’ interest; **diversifying representations of people/work/environments** can raise belonging and interest without necessarily changing ability beliefs overnight.

Tie to XXXVI / SAFE-IDMEASURE: Story worlds and tutor casting are identity technology. Token “add a girl character” without mathematical agency is costume. Empty AI praise (“you’re a math girl!”) is banned recognition theater (LXI).

HYPOTHESIS: For MindCraft, the highest-leverage gender-equity move is **recognition of real mathematical practice** (annotations, soft-wrong springboards, transfer proof) inside non-stereotyped worlds — not a separate pink product SKU.

LXV.8 Gender-equality paradox: segregation $\neq$ mean inability

FACT (PISA / STEM graduation pattern): Stoet & Geary (2018, *Psychological Science*, 29(4), 581–593, doi:10.1177/0956797617741719) — PISA science/math/reading $N = 472,242$; girls similar or better in science in 2/3 of

countries; yet sex differences in *relative* academic strengths and STEM degree pursuit **increase with national gender equality** (the “gender-equality paradox”).

FACT (stereotype mediation account): Breda, Jouini, Napp & Thebault (2020, *PNAS*, 117(49), 31063–31069, doi:10.1073/pnas.2008704117) — ~300k 15-year-olds, 64 countries; stereotype that “**math is not for girls**” is *stronger* in more egalitarian/developed countries and can **account for** the paradox linking development to math-intensive field segregation (vs pure “preferences free to express” story alone).

Wound: Cross-country observational; contested commentary exists (e.g., Richardson et al. 2020 on Stoet & Geary). Do not import as U.S. ACT causal law.

Commercial implication: U.S. parent markets are high-equality contexts where **horizontal gendering of interests** can be strong even when scores look equal. Selling “girls can’t do math so buy us” is false; selling “interest and stereotype climate still steer courses” is honest. Neither justifies pink cosplay nor challenge-softening.

LXV.9 SAFE-GENDER stack (surviving doctrine)

1. **Parity hygiene** — do not premise product on a large mean ability gap; cite similarities when talking scores.
2. **Separate channels** — performance ≠ interest ≠ identity ≠ stereotype climate; instrument separately (EVT + IDM).
3. **No innate-gap marketing** — ban Summers-style or “boys brains” copy forever.
4. **No ST→score guarantee** — Flore/Ganley wounds; ST is not ACT insurance (ties SAFE-EXAM).
5. **Cue hygiene** — no gendered ability jokes, “even you,” male-genius tropes in stems/stories/AI.
6. **Belonging via practice** — Cheryan-style diversify *who does math here*; recognition of work > aesthetics.
7. **Anti-cosplay** — no girl-only easy track; no pink SKU as equity.
8. **Anti-excuse** — gender does not waive FEI, transfer, or exam dosing — sequence support, don’t skip.
9. **Intersectionality light** — gender × race × ELL × class; one “STEM girl” persona is a kill (ties XXXVI / ELL).
10. **Parent speech** — “scores are close; confidence and course choice still get gendered” vs “we’ll make her a mathematician in 8 weeks.”

LXV.10 Experiments (GEND family)

ID	Contrast	Primary endpoints	Kill / survive rule
GEND-1	Stereotype-cued copy vs cue-clean copy (same math)	Accuracy; challenge_accept; belonging/recognition items	Cue hurts approach → enforce hygiene; no accuracy delta → still keep hygiene as ethics
GEND-2	Gendered “for her” skin vs neutral craft UI	4w transfer; identity items; reattempt	Skin ↑ vanity engagement, ↓ transfer → kill cosplay

GEND-3	Diverse mathematician agency in story vs token cameo	Belonging; challenge_accept; HIST-EQ codes	Token ≤ baseline → require agency casting
GEND-4	Recognition-of-practice feedback vs “girl power” praise	IDM factors; solo_transfer_pass	Praise-only ↑ Likert, flat transfer → kill praise track
GEND-5	Destaked practice → later gendered-stakes probe (exam ladder)	Accuracy under stakes; skip rate	Destake helps approach without transfer → redesign; both → keep SAFE-EXAM sequencing
GEND-QUAL	10 Maya interviews: “when did ‘math person’ feel gendered?”	Codebook: cue sources (peers/parents/media/UI)	Feeds copy bans + story brief

Ties: SAFE-EXAM, SAFE-IDMEASURE, SAFE-LONGID, SAFE-EXPECTANCY, XXXVI HIST-EQ, XXXVII EVT.

LXV.11 Commercial implications (copy, product, positioning, growth)

Marketing language that survives: “Average school math scores don’t justify ‘boys are better’ — what still gets gendered is confidence, belonging, and who keeps going.” / “We design so stereotypes never become the curriculum.” / “Recognition for real mathematical work — not girl-washed packaging.”

Marketing language that dies: “Fix her math brain”; “stereotype-threat vaccine → +ACT”; “built for girls” easy mode; “we close the STEM gender gap”; male-genius lore; empty “math girl” badges.

Positioning vs Khan / Duo / Brilliant / GPT tutors: Competitors can accidentally amplify stereotypes (male-coded “genius grind,” pink gamification, or fluent AI that mirrors parent stereotypes). MindCraft’s wedge: **identity-safe struggle + competence evidence** under cue hygiene — same FEI spine for every student, with story equity that diversifies *agency*, not difficulty.

Growth: Instrument stereotype_cue_flag (authoring), belonging/recognition items (IDM), course-intent proxies — not “% female DAU” as learning NS. School/parent sales: honesty about interest/identity channels; never innate-gap fear marketing.

Vision: Gender enters MindCraft as **climate and recognition design**, co-equal with language load (ELL) and exam affect (EXAM) — first-class authoring constraints, not a seasonal campaign.

LXV.12 Confidence table

Claim	Label	Confidence
-------	-------	------------

U.S. mean K–12 math gender gaps are trivial in large assessments	FACT	High
Overall meta gender math performance difference ≈ 0 (with small HS/high-tail caveats)	FACT	High
Math attitudes/affect often show larger gender differences than performance	FACT (classic meta; still directionally relevant)	Medium–High
Nation-level implicit gender–science stereotypes correlate with TIMSS gaps	FACT	Medium–High (correlational)
School-age ST average effect small and publication-bias-vulnerable; large RR null exists	FACT	High
Cultural STEM stereotypes gate interest/belonging and are malleable via representation	FACT / HYPOTHESIS (lab→product transfer)	Medium–High
Gender-equality paradox pattern in PISA/STEM degrees	FACT (pattern); mechanisms contested	Medium
Essentialist math-gender stereotypes help explain that paradox	HYPOTHESIS (Breda et al. account)	Medium
SAFE-GENDER is the right commercial mapping for MindCraft	FOUNDER BELIEF / HYPOTHESIS	Medium
Cue-clean + practice recognition beats pink cosplay on transfer	HYPOTHESIS	Medium (GEND-2/4)

LXV.13 What this chapter kills

1. **Kill:** Innate male math superiority as marketing or product premise.
2. **Kill:** Large mean ability-gap narrative contradicted by Hyde/Lindberg line.
3. **Kill:** Stereotype-threat removal as ACT/score guarantee (Flore meta bias + Flore RR + Ganley nulls).
4. **Kill:** Pink / “for her” / girl-power packaging as identity transformation.
5. **Kill:** Token STEM women without mathematical agency; empty “math girl” AI praise.
6. **Kill:** “We close the gender STEM gap” / 8-week mathematician claims.
7. **Kill:** Softening FEI/transfer because the student is a girl.
8. **Wound:** Treating preferences in high-equality countries as pure biology (Stoet/Geary vs Breda tension).
9. **Survive:** Mean parity hygiene; cue hygiene; Cheryan belonging-via-representation; EVT/identity channel separation; GEND-1...5; parent speech that names confidence/choice without innate myths.

Doctrine until data: Gender×math work at MindCraft is **stereotype-climate hygiene + earned recognition inside equitable practice** — not ability mythology, not ST theater, and not pink cosplay.

Part LXVI — Socioeconomic Constraint on “Grit”

Chapter status: Living evidence + equity/positioning brief — Researcher tick 2026-08-02

Primary question: When schools, apps, and marketers sell *grit* as the lever for math success, what does evidence say about construct validity — and how do socioeconomic structure, cognitive bandwidth, and opportunity-to-learn constrain (or fake) “perseverance” as a product story?

Owners: Equity · Parent messaging / WTP · Motivation (SAFE-RESILIENCE) · Product (session load, access) · Red Team

Commercial job: Ship a **SAFE-STRUCTURE** doctrine: structure and scaffolds before slogans; never sell low-SES Maya’s stuckness as a character defect; never ship a Grit Score™ as North Star.

LXVI.1 Why this chapter exists

Part XLV already kills grit *theater* as a resilience brand. Part LIX documents tutoring markets skewed by income. Parts LXIV–LXV separate language load and gender climate from ability myths. What remains unfinished is the **class channel**: parents and schools still hear “build grit” as the fix for achievement gaps that Reardon-scale income gradients and poverty-related cognitive load make structural.

FOUNDER BELIEF under audit: MindCraft can help identity under difficulty only if it treats SES as a **constraint on available bandwidth, time, tutoring dollars, and sense of control** — not as proof that some kids lack passion-and-perseverance DNA.

Claims we refuse as doctrine:

1. Grit as MindCraft’s primary product claim or North Star metric.
 2. “Poor kids just need more grit” / character lectures as equity.
 3. Grit Score™, streak-as-grit, or leaderboard-as-character as learning proof.
 4. Ignoring tutoring-market SES skew when pricing or packaging (ties SAFE-WTP).
 5. Treating low session completion under scarcity as moral failure.
 6. Importing Alan-style classroom grit RCTs as U.S. ACT guarantees.
 7. Softening FEI/transfer standards “because SES” instead of removing cognitive taxes and adding support.
-

LXVI.2 Constructs (product language)

Construct	Research meaning	MindCraft analogue	Failure mode if misused
-----------	------------------	--------------------	-------------------------

Grit (Duckworth)	Perseverance + consistency of interest for long-term goals	Tempting brand copy	Near-conscientiousness ; slogan equity
Perseverance of effort (PE)	Facet with more criterion validity	retry_120s, challenge_accept	Relabel PE as total grit
Consistency of interest (CI)	Facet weaker for academics	Path stickiness	Shame topic switching
Conscientiousness	Big Five trait overlapping grit	Study habits UX	Jangle fallacy product
Sense of control	Belief one can shape outcomes (SES-linked)	Destaked agency; clear next step	“Just believe harder”
Cognitive bandwidth / scarcity load	Poverty-related concerns consume WM/attention	Reduce UI tax; short sessions	Blame “lazy focus”
Income achievement gap	90–10 family-income score gaps	Market + equity fact	Sell grit as gap closer
SAFE-STRUCTURE	Structure > slogans doctrine	See LXVI.9	Character lectures; Grit Score™

Operational definition (HYPOTHESIS): A MindCraft experience is *structure-honest* when (a) it never diagnoses low engagement under scarcity as low grit, (b) it minimizes cognitive taxes (forms, clutter, timed floods) especially for bandwidth-taxed households, (c) it measures approach/recovery/transfer behavior not trait self-report, (d) pricing/access experiments acknowledge ability-to-pay (SAFE-WTP), and (e) challenge standards stay intact while support density rises.

LXVI.3 What grit claimed — and what the meta wound

FACT (origin construct): Duckworth, Peterson, Matthews & Kelly (2007, *Journal of Personality and Social Psychology*, 92(6), 1087–1101, doi:10.1037/0022-3514.92.6.1087) defined grit as perseverance and passion for long-term goals; across adult attainment, Ivy GPA, West Point retention, and Spelling Bee samples, grit accounted for ~4% of variance in success outcomes on average, correlated highly with conscientiousness, and showed some incremental prediction beyond IQ/conscientiousness in those samples.

FACT (short scale): Duckworth & Quinn (2009, *Journal of Personality Assessment*, 91(2), 166–174, doi:10.1080/00223890802634290) — Grit-S (8 items; PE + CI facets).

FACT (meta wound): Credé, Tynan & Harms (2017, *Journal of Personality and Social Psychology*, 113(3), 492–511, doi:10.1037/pspp0000102) — **584** effect sizes, **88** samples, **N = 66,807**. Higher-order grit structure not confirmed; grit only **moderately** related to performance/retention; grit **very strongly** related to conscientiousness; **PE ■ CI** for criterion validity; interventions to “enhance grit” likely **weak** on performance; construct validity questioned; utility may lie mainly in the perseverance facet.

FACT (policy critique): Credé (2018, *Educational Researcher*, 47(9), doi:10.3102/0013189X18801322) — no empirical reason to fuse passion+perseverance into one construct; little support that grit is a *particularly* strong

educational predictor or that it is readily intervention-responsive as marketed.

Commercial kill: Do not sell MindCraft as a grit factory. Do not ship Grit Score™. If any perseverance language survives, it must mean **designed recovery after soft-wrong** (SAFE-RESILIENCE / SAFE-DD), not trait branding.

Survive: Domain-specific persistence behaviors (retry_120s, help-recruitment, transfer) remain product signals — subordinate to FEI, not a personality product line.

LXVI.4 SES is not a grit deficit — structure shows up in scores and markets

FACT (income achievement gradient): Reardon (2011), in Duncan & Murnane (Eds.), *Whither Opportunity?* (Russell Sage) — the U.S. **90–10 income achievement gap** grew roughly **30–40%** for cohorts born ~2001 vs ~25 years earlier; income gaps became large relative to historical black–white gaps; gap is already large at school entry. (Later work debates trend magnitude; the *level* of SES–achievement association remains a central market fact.)

FACT (paid tutoring skew): Kim, Goodman & West (2024, CESifo WP 11251 / IZA DP 17178) — U.S. tutoring-center expansion and high-income concentration (also summarized in Part LIX). Bray (2010) / Bray & Lykins (2012, ADB) — shadow education is stratified.

Commercial implication: Parents who can buy “more grit” narratives often already buy **structure** (tutors, quiet time, test prep). Positioning that shames lower-resource families for “not pushing hard enough” is false and brand-toxic. Equity copy must name **access, time, and bandwidth**, not character.

Wound: Income gaps confound OTL, school quality, health, stress, language, and prior knowledge. They are not MindCraft effect sizes and not proof that grit interventions close Reardon gaps.

LXVI.5 Poverty taxes cognition — “not trying” is often load

FACT (scarcity → cognitive performance): Mani, Mullainathan, Shafir & Zhao (2013, *Science*, 341(6149), 976–980, doi:10.1126/science.1238041) — financial concerns experimentally reduce cognitive performance among lower-income mall participants (not well-off); the same Indian sugarcane farmers perform worse **pre-harvest** (scarce) than **post-harvest** (less scarce), not explained by nutrition/work-time alone; authors calibrate lab effects on the order of **~13 IQ points** / comparable to a night of sleep loss on Raven’s-style tasks. Claim is about **context of poverty**, not fixed traits of “poor people.”

HYPOTHESIS (MindCraft): Exam-week fees, noisy shared devices, after-job homework slots, and parent bill stress are **dual-task taxes** on math WM — same family as anxiety load (Ashcraft) and exam pressure (SAFE-EXAM). Product that adds cognitive taxes (dense UI, day-one timed ACT flood, endless forms) will look like “low grit” in analytics while actually measuring scarcity.

Product rules: Prefer short sharp DP cycles (SAFE-DP); offline-friendly / resume drafts; low-chrome practice; destake-to-learn before stakes; never moralize incomplete missions.

Kill: “They bounced off because they lack grit.” Prefer diagnose load, anxiety, language, format, and access first.

LXVI.6 Sociology of grit: SES → sense of control → perseverance (not pure will)

FACT (SES–control–grit path): Kwon (2021, *Socius*, 7, doi:10.1177/23780231211005216) — U.S. and South Korea survey data; higher SES partially translates into grit (especially **perseverance of effort** in Korea) **via stronger sense of control**; education sometimes **negatively** associates with consistency-of-interest, suggesting stick-to-it-iveness can be a *supplementary* resource when other resources are scarce — not a simple “rich = gritty” cartoon.

HYPOTHESIS: Sense of control in MindCraft is engineered by **clear diagnosis** → **next ingredient** → **recoverable miss**, not by posters. Low-control contexts need *more* scaffold clarity, not louder character talk.

Tie to EVT (XXXVII): Expectancy collapses when structure fails. Selling passion without expectancy is inspirational avoidance.

LXVI.7 Can grit be taught? Real RCT — narrow transfer, not ACT insurance

FACT (classroom intervention RCT): Alan, Boneva & Ertac (2019, *Quarterly Journal of Economics*, 134(3), 1121–1162, doi:10.1093/qje/qjz006) — randomized educational intervention in elementary samples; treated students more likely to exert effort to build task-specific ability; follow-up ~**2.5 years** later ≈ **0.2 SD** on a standardized math test.

Wound / kill for marketing: This is **not** a license to claim MindCraft “installs grit → +ACT points.” Sample/setting ≠ U.S. high-stakes ACT; effect is modest; Credé line still wounds trait-level marketing and easy malleability claims. Borrow the *design idea* (normalize failure → try again with strategy) already owned by SAFE-RESILIENCE / soft-wrong — do not import the brand name as North Star.

SPECULATION: MindCraft’s closest analogue is **attempt** → **soft-wrong** → **support recruitment** → **retry**, which trains perseverance *behavior* without claiming personality change.

LXVI.8 Intersection with identity, ELL, gender, and exam pressure

SES rarely travels alone. Language load (SAFE-ELL), stereotype climate (SAFE-GENDER), and high-stakes affect (SAFE-EXAM) compound bandwidth taxes. Homogeneous “gritty hustle” personas erase that stack.

FOUNDER BELIEF: Recognition of mathematical practice (SAFE-IDMEASURE) must be available without requiring paid co-curricular polish that only high-income students get. Otherwise identity becomes another SES luxury good.

LXVI.9 SAFE-STRUCTURE stack (surviving doctrine)

1. **Structure before slogans** — scaffolds, diagnosis, cards, tutor briefs beat grit posters.
2. **No Grit Score™ / streak-as-character** — already banned as NS; reaffirm under SES equity.
3. **Bandwidth hygiene** — minimize cognitive taxes; short sessions; resume without shame.

4. **Behavior $\neq$ trait** — instrument `retry_120s`, `help-recruit`, `transfer` — not Duckworth Likert as product KPI.
5. **PE not CI worship** — allow healthy topic/path revision; don't shame interest shifts.
6. **Access honesty** — pricing/scholarship experiments acknowledge ability-to-pay (SAFE-WTP); don't equate unpaid with ungritty.
7. **No character-blame equity** — never explain SES gaps as missing passion.
8. **No ACT-from-grit guarantee** — Alan $\neq$ exam insurance; SAFE-EXAM dosing stays.
9. **Support density under scarcity** — more on-ramps, not softer math standards.
10. **Parent speech** — “we remove friction and build recoverable struggle” vs “we build grit.”

LXVI.10 Experiments (STRUCT family)

ID	Contrast	Primary endpoints	Kill / survive rule
STRUCT-1	Low-chrome / resume-friendly practice vs high-tax clutter UI	Session completion; <code>retry_120s</code> ; <code>transfer</code>	Clutter $\downarrow$ completion more for self-reported scarcity $\rightarrow$ enforce bandwidth hygiene
STRUCT-2	Soft-wrong + support recruitment vs grit-poster interstitial (same items)	<code>challenge_accept</code> ; <code>help-recruit</code> ; delayed accuracy	Poster $\geq$ scaffold on <code>transfer</code> $\rightarrow$ unexpected; else kill poster track
STRUCT-3	Grit-S / PE Likert vs behavioral perseverance suite as predictors	Incremental R^2 on <code>transfer_pass</code>	Likert adds nothing $\rightarrow$ ban trait KPI
STRUCT-4	Scholarship / access messaging vs “build character” messaging (parent CBC arm)	WTP; brand trust; offense codes	Character-blame $\downarrow$ trust in lower-income segment $\rightarrow$ kill
STRUCT-5	Exam-week destake ladder vs day-one timed flood under self-reported stress/scarcity	Skip rate; accuracy under later stakes	Flood harms scarce/stressed more $\rightarrow$ SAFE-EXAM+STRUCTURE
STRUCT-QUAL	10 Maya + 10 parent interviews: time, devices, work, quiet, money stress	Codebook of structural barriers	Feeds copy bans + session defaults

Ties: SAFE-RESILIENCE, SAFE-WTP, SAFE-EXAM, SAFE-DD, SAFE-DP, SAFE-ELL, SAFE-GENDER, SAFE-IDMEASURE.

LXVI.11 Commercial implications (copy, product, positioning, growth)

Marketing language that survives: “Struggle isn’t a character test — it’s a design problem.” / “We build recoverable practice and clear next steps, not grit lectures.” / “When life is loud, the product should get quieter — short sessions, saved progress, support on tap.”

Marketing language that dies: “We forge grit”; “poor performers lack passion”; Grit Score™; streak-as-moral-proof; “just push through” as equity; guaranteed ACT from character curriculum.

Positioning vs Khan / Duo / Brilliant / GPT tutors: Duo-style streaks and hustle aesthetics can read as character judgment when life constraints dominate. ChatGPT fluency doesn’t remove bandwidth taxes. MindCraft’s wedge: **ontology diagnosis + soft-wrong support + transfer proof** that respects structural load — structure as product, not sermon.

Growth: Segment analytics by access proxies (device class, session time-of-day, scholarship flag) without building a shady “SES score.” Instrument structural barriers in QUAL; never optimize shame. School/district sales: lead with OTL and scaffolding, not grit PD.

Vision: Class enters MindCraft as **bandwidth, access, and sense-of-control design** — co-equal with ELL load and gender climate — so identity transformation is not a luxury good.

LXVI.12 Confidence table

Claim	Label	Confidence
Duckworth grit predicts some success outcomes (~4% avg variance in origin paper)	FACT	High
Credé meta: strong conscientiousness overlap; weak higher-order grit; PE > CI; weak intervention expectation	FACT	High
Credé 2018: grit not a uniquely strong / easily trained edu lever as marketed	FACT	High
U.S. income achievement gaps are large and historically salient (Reardon line)	FACT	High (trend debates Medium)
Tutoring spend / centers concentrate among higher-income families	FACT	High
Poverty-related concerns can impair cognitive performance (~sleep-loss / ~13 IQ calibration in Mani et al.)	FACT	High (context transfer to HS homework Medium)
SES → sense of control → grit facets (Kwon)	FACT (observational, cross-cultural)	Medium–High

Alan et al. classroom grit intervention → ~0.2 SD math at 2.5y	FACT	High (external validity to MindCraft/ACT Low–Medium)
SAFE-STRUCTURE is the right commercial mapping	FOUNDER BELIEF / HYPOTHESIS	Medium
Bandwidth hygiene + soft-wrong beats grit posters on transfer under scarcity	HYPOTHESIS	Medium (STRUCT-1/2)

LXVI.13 What this chapter kills

1. **Kill:** Grit as MindCraft North Star, primary ad claim, or Grit Score™.
2. **Kill:** “Low-SES students lack grit” / character-blame equity.
3. **Kill:** Streaks or leaderboards as proof of character.
4. **Kill:** Importing Alan RCT as ACT-point insurance.
5. **Kill:** Treating scarcity-driven dropout as moral failure.
6. **Kill:** Softening math standards instead of removing cognitive taxes.
7. **Wound:** Treating all perseverance language as forbidden — behavioral recovery metrics still survive under SAFE-RESILIENCE.
8. **Survive:** Credé construct hygiene; Mani bandwidth humility; Reardon/WTP structural honesty; Kwon control pathway; STRUCT-1...5; structure-before-slogans parent speech.

Doctrine until data: Socioeconomic constraint means MindCraft sells **scaffolds, bandwidth hygiene, and recoverable struggle** — never grit sermons as the fix for structural gaps.

Part LXVII — Ontology as Diagnosis Moat

Chapter status: Living evidence + positioning brief — Researcher tick 2026-08-02

Primary question: When does a structured knowledge representation (concept/ingredient graph + persistent student state) beat fluent chat tutoring for *diagnosis* — and what, exactly, can MindCraft claim as a commercial moat without cosplaying ALEKS or inventing “AI tutor superiority”?

Owners: Engine (ontology / pathfinder) · Product (Map / Solver / Notes) · Positioning vs Khan/Duo/Brilliant/GPT · Red Team

Commercial job: Ship a **SAFE-ONTOLOGY** doctrine: diagnosis is a stateful graph problem first; language models are bookends; the moat is longitudinal, privacy-respecting student state joined to a canonical ontology under FEI — not a JSON file and not chat fluency.

LXVII.1 Why this chapter exists

Parts XXXIII / XXXV / LX already kill unguarded Solver-as-hero, Bastani-Tutor≡MindCraft category errors, and apology≡repair. What remains unfinished is the **engine wedge**: founders believe the standardized ontology + mastery graph is the product moat. Competitors can wrap GPT tomorrow. Parents hear “AI tutor” as interchangeable. Sales needs a defensible sentence for *why not just ChatGPT*.

FOUNDER BELIEF under audit: Knowledge graphs beat chat when the job is *which gap, in what order, with what evidence* — not when the job is *sound fluent*.

Claims we refuse as doctrine:

1. “We have a knowledge graph” ≡ durable competitive moat by itself.
2. RAG-wrapped ChatGPT ≡ MindCraft diagnosis (category error with Bastani-Tutor≡ontology).
3. Chat fluency / token volume as proof of diagnosis quality.
4. Ontology coverage theater (node count as learning).
5. Blocked accuracy on graph-sequenced drills ≡ exam-ready / identity change.
6. Absolute “ITS beat humans / 2-sigma” Bloom marketing.
7. Selling the graph without transfer / solo proof (PWC still binds).

LXVII.2 Constructs (product language)

Construct	Research meaning	MindCraft analogue	Failure mode if misused
Student model	Latent mastery / knowledge state over skills	Concept + ingredient mastery; bridge gaps	Treat chat memory as student model
Knowledge tracing (BKT lineage)	Bayesian update of skill mastery from attempts	/record-outcomes, decay, strength scoring	Black-box “AI knows you” without inspectable state
Knowledge space / learning space	Combinatorial feasible states; outer fringe	Pathfinder trim; “ready next”	ALEKS cosplay; assessment-only brand
Model tracing / cognitive model	Step-level comparison to expert production rules	Ingredient cards + soft-wrong classify	Claim step-model without traces
Ontology / concept graph	Directed prereqs, bridges, combinations	Layer-1 ontology; Map	Node-count vanity; stale alias maps
Diagnosis	Infer <i>what is broken</i> and <i>what to teach next</i>	Gap-scan → /recommend → worstWeakness	Fluency-as-diagnosis
Pedagogy wrap (PWC)	Guards + spine; solo transfer ≥ no-AI	Solver under Map	Unguarded chat hero
SAFE-ONTOLOGY	Graph-first diagnosis doctrine	See LXVII.9	Graph file as moat slogan

Operational definition (HYPOTHESIS): A MindCraft experience is *diagnosis-honest* when (a) every practice/session outcome updates an inspectable concept/ingredient/bridge state, (b) next-action is constrained by ontology edges not free chat wander, (c) LLM text cannot invent mastery or hide wrongness, (d) parent/tutor can see *why this next*, and (e) success is still `solo_transfer_pass` / delayed mix — not chat satisfaction.

LXVII.3 What structured tutors actually earned (and did not)

FACT (relative tutoring effectiveness): VanLehn (2011, *Educational Psychologist*, 46(4), 197–221, doi:10.1080/00461520.2011.611369) — human tutoring effect $\approx d = 0.79$ vs no tutoring (far below Bloom’s mythic 2σ); step-based ITS $\approx d = 0.76$. Interaction *grain* matters more than “AI” branding; answer-only CAI sits weaker than step-based tutors.

FACT (K–12 math ITS meta): Steenbergen-Hu & Cooper (2013, *Journal of Educational Psychology*, 105(4), 970–987, doi:10.1037/a0032447) — 26 reports / 34 samples; overall ITS vs classroom often **near zero to small positive** ($g \approx 0.01$ – 0.09); shorter interventions looked stronger than year-long; general-population samples outperformed low-achiever samples on average. **Wound:** “We have ITS architecture” is not an ACT guarantee and not equity magic.

FACT (college ITS meta): Steenbergen-Hu & Cooper (2014, *Journal of Educational Psychology*, 106(2), 331–347, doi:10.1037/a0034752) — moderate positive overall ($g \approx 0.32$ – 0.37); ITS beat many non-human alternatives; **worse than human tutoring** where compared (negative mean in that thin subset). Pedagogy and teachers still matter.

Commercial implication: Sell *inspectable diagnosis* + *recoverable struggle*, not “our AI is as good as a tutor” or Bloom 2σ . VanLehn’s anti- 2σ fact stays standing ban.

LXVII.4 Persistent skill models — the job chat still fails

FACT (knowledge tracing): Corbett & Anderson (1994/1995, *User Modeling and User-Adapted Interaction*, 4, 253–278, doi:10.1007/BF01099821) — Bayesian *knowledge tracing* over production-rule skills in an ACT-based programming tutor; mastery-driven exercise selection; model predicts posttest from component mastery estimates.

FACT (cognitive tutors as method, not costume): Anderson, Corbett, Koedinger & Pelletier (1995, *Journal of the Learning Sciences*, 4(2), 167–207, doi:10.1207/s15327809jls0402_2) — decade of ACT-R Cognitive Tutor lessons: explicit cognitive model, model-tracing assistance, knowledge-tracing mastery sequencing. Achievement gains tied to *component mastery design*, not chatbot persona.

FACT (knowledge spaces): Doignon & Falmagne (1985, *International Journal of Man-Machine Studies*, 23(2), 175–196, doi:10.1016/S0020-7373(85)80031-6) introduced knowledge-space combinatorics; Falmagne, Koppen, Villano, Doignon & Johannesen (1990, *Psychological Review*, 97(2), 201–224, doi:10.1037/0033-295X.97.2.201) — build/test/search knowledge spaces; practical “what student can do / ready to learn” lists. ALEKS commercializes this lineage (Falmagne & Doignon *Learning Spaces*, 2011; Matayoshi et al. JMP 2021 practical KST/ALEKS perspective) — millions of users, assessment→fringe learning.

HYPOTHESIS (chat wound): Stateless or weakly stateful LLM tutors optimize for *plausible next utterance*. They do not reliably maintain a falsifiable mastery vector across weeks, bridge gaps, or format axes (Shetye 2024 /

conversational-tutor reviews in Part XXXIII line; systematic KT×LLM reviews note diagnosis as open problem). Bastani et al. (2025, *PNAS*, doi:10.1073/pnas.2422633122) shows unguarded GPT Base ↑ practice then ↓ solo exam (~17%) — fluency without learning architecture.

Kill: “ChatGPT remembers our conversation” ≡ knowledge tracing. Memory ≠ mastery model.

Survive: Hybrid design already in MindCraft doctrine — deterministic spine (ontology, outcomes, recommend) with LLM as language bookends (classify/render), not the adjudicator of truth.

LXVII.5 When graphs beat chat (and when they do not)

HYPOTHESIS — graphs win when:

1. **Prerequisite / bridge failure** — student “knows” A and B but fails $A \rightarrow B$ (MindCraft bridge-gap severity). Chat often re-explains A.
2. **Longitudinal planning** — exam-mode path, decay, spacing need durable state (SAFE-TIMING / SAFE-EXAM).
3. **Comparable severity ranking** — C1 worstWeakness across concept + format gaps needs shared scales, not vibes.
4. **Tutor/parent accountability** — “why this next” must be auditable (SAFE-EXPECTANCY task briefs).
5. **Anti-hallucination ground** — Map-checked correct after repair (SAFE-REPAIR).

HYPOTHESIS — chat wins (or ties) when:

1. Open-ended *wording* help after diagnosis is already known.
2. Affective micro-coaching language (still SE-constrained; Parts XL / A).
3. Rapid authoring of surface variants *after* essence/keys verified — not as the diagnostic authority.

Wound (ALEKS already exists): Knowledge-space assessment is not a greenfield moat. Differentiation cannot be “we have states.” Differentiation is **ontology grain (ingredients/bridges) × FEI identity loop × human tutor witnessing × transfer proof** under PWC.

SPECULATION: A pure chat competitor can fake Map UI screenshots; they cannot fake longitudinal eventCount, bridge evidence, and solo_transfer_pass without building the spine — *if* MindCraft instruments and publishes those metrics honestly.

LXVII.6 Ontology is infrastructure, not the North Star

FOUNDER BELIEF refined: The ontology is the *spine* that makes diagnosis, cards, and recommendations non-hallucinated. It is not the emotional product Maya buys and not the parent headline.

FACT-adjacent commercial pattern: Steenbergen-Hu metas + VanLehn show structured tutors can help *and* can underwhelm vs classroom/human baselines depending on implementation. Selling “42 concepts” or “179 ingredients” as the hero metric invites coverage theater (same failure mode as hours-practiced / node-count).

Product rules:

- Map shows *evidence-backed* weakness and next route — not a trophy wall of nodes.
- Gap-scan seeds state; practice appends; never let Solver overwrite mastery from fluent dialogue alone.
- Soft-wrong + misconception tags update ingredient/bridge state (SAFE-MISCON).
- Generation fills coverage holes; generation does not replace ontology joins (C5).

Kill: Homepage that leads with “powered by our knowledge graph” without a human-need sentence (fear → evidence → identity).

LXVII.7 Positioning vs Khan / Duo / Brilliant / GPT

Competitor pattern	What they optimize	MindCraft counter under SAFE-ONTOLOGY
Khan	Content breadth + mastery dashes	Diagnosis of <i>your</i> bridge/format gap + recoverable miss — not another video library
Duo	Streak / habit loops	Habit subordinate to FEI; graph state ≠ streak (SAFE-HABIT)
Brilliant	Delightful interactive puzzles	Delight without longitudinal diagnosis is entertainment; we keep Map memory
ChatGPT tutors	Fluent explanation / homework finish	Fluency is cheap; inspectable next-skill + solo transfer is scarce (PWC)

Marketing language that survives: “We don’t just explain the problem — we keep a living map of what broke and what unlocks next.” / “Chat can talk; diagnosis needs memory you can trust.” / “Your tutor sees the same map the AI must obey.”

Marketing language that dies: “Our AI is smarter than ChatGPT”; “knowledge graph = guaranteed ACT”; “2-sigma Cognitive Tutor”; “we’re ALEKS but cooler”; node-count bragging; Solver as the product hero.

LXVII.8 Moat economics (what compounds)

HYPOTHESIS (data moat): Privacy-respecting longitudinal events (outcomes, soft-wrongs, bridge evidence, format profile) joined to a *shared* canonical ID contract compound; prompt libraries and chat logs commoditize faster.

HYPOTHESIS (workflow moat): Tutor ops that brief from Map (Part LXVIII queued) + student Map literacy create switching costs chat tabs lack — *if* QA and FEI training are real (not Discord theater; SAFE-CoP).

Wound: Without instrumentation, “ontology moat” is founder poetry. Without privacy ethics (queue id 69), event density becomes liability.

Claim ladder: L1 = A/B on recommend quality / time-to-useful-next; L2 = tutor agreement with system diagnosis; L3+ identity still needs SAFE-IDMEASURE / LONGID — graph Δ alone ≠ “became a math person.”

LXVII.9 SAFE-ONTOLOGY stack (surviving doctrine)

1. **Diagnosis before dialogue** — state update and next-action from ontology/outcomes; chat cannot set mastery.
 2. **Inspectable student model** — concept/ingredient/bridge/format fields a tutor can audit.
 3. **LLM = bookends** — classify/render under PWC; never adjudicate ground truth alone.
 4. **Bridges & formats first-class** — not only node mastery (C1 severity).
 5. **No graph-file moat slogan** — moat = state $\times$ ontology $\times$ FEI $\times$ transfer proof.
 6. **No RAG $\equiv$ MindCraft** — retrieval $\neq$ knowledge tracing.
 7. **No fluency-as-diagnosis / token KPIs** — ban chat satisfaction as learning proof.
 8. **No Bloom/ITS-2 σ marketing** — VanLehn standing ban.
 9. **Coverage fills gaps; coverage $\neq$ product** — generation serves bank holes under verify.
 10. **Parent/tutor speech** — “living map of gaps” not “smarter chatbot.”
-

LXVII.10 Experiments (ONTO family)

ID	Contrast	Primary endpoints	Kill / survive rule
ONTO-1	Ontology-constrained next (/recommend) vs free chat “what should I do next?” (same student state hidden from chat)	Time-to-attempt; topic relevance coded; transfer@7d	Chat $\geq$ ontology on transfer $\rightarrow$ wound moat claim
ONTO-2	Bridge-gap card route vs re-explain prerequisite concept only	Bridge item accuracy; severity Δ	Re-explain wins $\rightarrow$ bridge detector theater
ONTO-3	Tutor shown Map diagnosis vs tutor chat-only brief	Tutor–system agreement; session FEI events	Map no lift $\rightarrow$ briefing UX fail not ontology fail
ONTO-4	Outcomes-appended state vs chat-memory “profile” summary only	Delayed mix; calibration of next difficulty	Summary $\approx$ state $\rightarrow$ simplify; else keep event spine
ONTO-5	Solver under Map guards vs unguarded Solver (replication of PWC/AIT)	solo_transfer_pass	Unguarded $\geq$ guarded $\rightarrow$ revisit PWC
ONTO-QUAL	10 parent + 10 tutor: “why not ChatGPT?” codebook	Message resonance; confusion codes	Feeds SAFE-WTP CBC attribute “living diagnosis map”

Ties: PWC, SAFE-REPAIR, SAFE-MISCON, SAFE-TRANSFER, SAFE-CALIB, SAFE-DP, SAFE-WTP.

LXVII.11 Commercial implications (copy, product, growth, vision)

Copy: Lead with the *job* — find the real gap, practice the unlock, prove it alone — not with graph jargon. Use “map,” “gap,” “next unlock,” “your tutor sees the same map.”

Product: Keep PawHub / Practice launches driven by /recommend + playable questions; never let Solver become the homepage. Instrument diagnosis quality (ONTO-1...4) before scaling ontology marketing.

Growth: School/district pitch = shared diagnostic language across tutors + audit trail; consumer pitch = “ChatGPT finishes homework; we build a map you still own when the chat ends.”

Vision: Thirty-year bet stays identity transformation — ontology is the **memory substrate** that makes earned competence evidence cumulative rather than session-amnesiac.

LXVII.12 Confidence table

Claim	Label	Confidence
Human tutoring $\approx d=0.79$; step ITS $\approx d=0.76$ vs no tutoring (VanLehn 2011); Bloom 2σ not baseline	FACT	High
K–12 math ITS meta often small vs classroom (Steenbergen-Hu & Cooper 2013)	FACT	High
College ITS moderate positive; not reliably $\geq$ human tutors (2014 meta)	FACT	High
BKT / Cognitive Tutor lineage: explicit skill model + mastery sequencing predicts/posttests	FACT	High
Knowledge spaces / ALEKS: combinatorial states $\rightarrow$ “can do / ready to learn”	FACT	High
Unguarded GPT can harm solo exam performance (Bastani et al. 2025 PNAS)	FACT	High
LLM chat alone lacks reliable longitudinal diagnostic student models	HYPOTHESIS (strongly supported)	Medium–High

MindCraft moat = state×ontology×FEI×transfer, not graph file or chat	FOUNDER BELIEF / HYPOTHESIS	Medium
Ontology-constrained next beats free chat on delayed transfer (ONTO-1)	HYPOTHESIS	Medium
SAFE-ONTOLOGY is right commercial mapping	FOUNDER BELIEF	Medium

LXVII.13 What this chapter kills

1. **Kill:** Knowledge-graph file / node count as the moat slogan.
2. **Kill:** RAG or chat memory $\equiv$ knowledge tracing / MindCraft diagnosis.
3. **Kill:** Fluency, thumbs-up, or Solver tokens as diagnosis quality.
4. **Kill:** Bloom 2 σ / “ITS = human tutor” absolute marketing.
5. **Kill:** ALEKS-cosplay assessment brand without FEI identity loop.
6. **Kill:** “Smarter than ChatGPT” comparative AI brag.
7. **Wound:** Treating chat language help as useless — it survives as *bookend* under PWC.
8. **Survive:** VanLehn grain lesson; Corbett/Anderson student models; KST fringe idea; Bastani guardrails; ONTO-1...5; diagnosis-before-dialogue parent speech.

Doctrine until data: Ontology beats chat when the scarce job is **trusted, longitudinal diagnosis under FEI** — and only then. Chat remains language infrastructure, never the adjudicator of what Maya knows.

Part LXVIII — Human-in-the-Loop Tutor Ops

Chapter status: Living evidence + ops/product brief — Researcher tick 2026-08-03

Primary question: What does *quality* human tutoring look like as an operating system — playbooks, QA, FEI training, escalation — when MindCraft already has Map diagnosis and AI bookends?

Owners: Tutor ops · Product (tutor brief / session tools) · Engine (Map as shared truth) · Positioning vs “AI replaces tutors” / “warm humans = quality”

Commercial job: Ship a **SAFE-HITL** doctrine: humans are not a vibe layer on chat; they are trained operators who brief from inspectable state, prompt more than lecture, refuse polite false mastery, escalate when AI fails, and are coached at a cadence that survives scale — or the tutor brand is theater.

LXVIII.1 Why this chapter exists

Part LXVII argues diagnosis is a stateful graph problem. Part XLVI bans trait-label expectancy. Part LX kills apology⇒repair and requires triggered human escalate. Part LI wants deliberate-practice session spines. Part LIV wants brief/debrief. None of that ships unless **tutor operations** exist: who briefs how, what “good session” means, how QA works, how new tutors learn FEI without Discord cosplay.

FOUNDER BELIEF under audit: College tutors + Map + AI bookends beat AI-alone *and* beat warm untrained humans — but only if ops encode pedagogy, not “hire nice people.”

Claims we refuse as doctrine:

1. Warm human presence ≡ tutoring quality.
2. Session hours / booked minutes as quality KPI.
3. “Do you understand?” / student smile as mastery check.
4. Slack/Discord chat ≡ CoP or tutor QA (SAFE-CoP already wounded forum⇒LPP).
5. AI coach monologue as tutor training.
6. Empathy scripts without Map task brief (SAFE-EXPECTANCY).
7. Infinite free human escalation (SAFE-REPAIR).
8. “Our tutors are Ivy” as substitute for playbook + transfer proof.

LXVIII.2 Constructs (ops language)

Construct	Research meaning	MindCraft analogue	Failure mode if misused
Naturalistic tutoring	Typical unskilled peer/near-peer dialogue	Default college tutor without playbook	Assume “1:1” automatically deep
Scaffold / prompt	Elicit student construction vs deliver explanation	Coach cards; SE prompts; AAR student-first	Lecture-first “help”
Comprehension-gauging Q	“Got it?” checks	End-of-explain polls	False mastery theater
Politeness tax	Face-saving softens needed correction	Soft-wrong avoided; wrong keys unchallenged	Rapport > diagnosis
Teacher–AI complementarity	Humans for what software cannot; shared orchestration	Tutor + Map + Solver under PWC	Glasses cosplay / dashboard vanity
Coaching (PD)	Sustained, observation-linked development	Tutor QA cycles; FEI rubrics	One-off onboarding deck
SAFE-HITL	Ops doctrine for human loop	See LXVIII.9	Human-as-marketing-prop

Operational definition (HYPOTHESIS): A MindCraft tutor session is *HITL-honest* when (a) pre-brief is Map/task-only (no trait labels), (b) talk ratio favors student construction over tutor monologue, (c) soft-wrongs are

elicited not papered over, (d) AI outputs are challenged against Map ground truth, (e) close uses short AAR tied to next practice, and (f) QA samples sessions against a FEI rubric — not star ratings alone.

LXVIII.3 What naturalistic tutors actually do (and miss)

FACT (collaborative dialogue, not expert strategy): Graesser, Person & Magliano (1995, *Applied Cognitive Psychology*, 9(6), 495–522, doi:10.1002/acp.2350090604) — naturalistic tutoring with *normal unskilled* tutors (grad→undergrad research methods; HS→7th-grade algebra) is dominated by collaborative problem solving, question answering, and explanation on specific examples. Sophisticated strategies common in ITS aspiration (deep error diagnosis systems, ideal Socratic method, affect tactics) were largely **absent or underdeveloped**.

FACT (question quality > question count): Graesser & Person (1994, *American Educational Research Journal*, 31(1), 104–137, doi:10.3102/00028312031001104) — student questions ~240× more frequent in tutoring than classroom; after experience, **quality** of student questions correlated with achievement; **frequency** did not. Tutors need training to improve question-asking — volume is not the metric.

FACT (“Do you understand?” misleads): Person, Graesser, Magliano & Kreuz (1994, *Learning and Individual Differences*, 6(2), 205–229, doi:10.1016/1041-6080(94)90010-8) — quality of student *answers* is the reliable signal of understanding; student questions are weak; answers to tutor comprehension-gauging questions are **very misleading**.

FACT (politeness can inhibit pedagogy): Person, Kreuz, Zwaan & Graesser (1995, *Cognition and Instruction*, 13(2), 161–188, doi:10.1207/s1532690xc1302_1) — Gricean rules and Brown & Levinson politeness strategies appear throughout tutoring dialogue; they can support rapport **and** inhibit effective tutoring via softened/indirect feedback — especially in less constrained domains.

Commercial implication: Marketing “real tutors” without ops is selling the *setting* Graesser already showed works *somewhat* even when unskilled — while missing the failure modes that destroy FEI (false mastery checks, polite non-correction, lecture-heavy help). Ops must train against those defaults.

LXVIII.4 Prompting beats pouring — training content, not personality

FACT (interactive scaffolding ≈ didactic when prompts replace explanations): Chi, Siler, Jeong, Yamauchi & Hausmann (2001, *Cognitive Science*, 25(4), 471–533, doi:10.1207/s15516709cog2504_1) — when tutors were **suppressed from giving explanations and feedback** and pushed to prompt/scaffold, students learned **as effectively** as with traditional didactic tutoring; gains tied to deeper/more scaffolding episodes and greater student control (with reading-ability limits).

HYPOTHESIS (MindCraft mapping): Tutor training should default to *prompt* → *student try* → *soft-wrong diagnose* → *Map-aligned next*, not *explain until nod*. AI coach cards that monologue violate Chi’s interactive lesson and Part XL SE doctrine.

Cross-link: SAFE-DP (45-min spine), SAFE-SE / Part XL, SAFE-MISCON (productive wrongs), SAFE-AAR (student-first close).

Kill: “Our tutors explain better than AI” as the brand. Explanation is cheap (Parts XXXIII/XXXV). Scarce skill = **eliciting construction under diagnosis**.

LXVIII.5 Human–AI complementarity is an ops design, not a slogan

FACT (teacher–AI co-orchestration can lift learning): Holstein, McLaren & Aleven (2020, *AI Magazine*, 41(2), doi:10.1002/aaai.12058) — field evidence that students learn more when teachers and AI tutors work *together* with real-time analytics support (Lumilo line) than in AI-supported classrooms without that teacher orchestration aid. Design premise: AI personalizes content/pace; humans handle situations software is ill-suited for; conflicts arise if AI and human plans are uncoordinated.

FACT (co-design of orchestration tools): Holstein, McLaren & Aleven (2019, *Journal of Learning Analytics*, 6(2), 27–52, doi:10.18608/jla.2019.62.3) — iterative co-design of real-time teacher awareness tools for AI-enhanced K–12; analytics must match practitioner needs, not dump model internals.

HYPOTHESIS (1:1 transfer): MindCraft’s analogue is not smart glasses — it is **pre-session Map brief + live soft-wrong/bridge alerts + Solver escalate packet (SAFE-REPAIR)**. Tutor without Map is Holstein’s “AI classroom without orchestration support.” Map without trained tutor is dashboard vanity.

Wound: Classroom orchestration ≠ remote 1:1 college tutor. Do not market Lumilo effect sizes as ACT guarantees. Borrow *complementarity principle*, not costume.

Kill: “AI replaces tutors” *and* “humans don’t need the graph.” HITL means shared truth + divided labor.

LXVIII.6 Coaching tutors beats hiring theater — and scale is the wound

FACT (coaching meta): Kraft, Blazar & Hogan (2018, *Review of Educational Research*, 88(4), 547–588, doi:10.3102/0034654318759268) — 60 causal-design studies; pooled coaching effects ≈ **0.49 SD** on instructional practice and ≈ **0.18 SD** on student achievement. Much evidence from U.S. literacy coaching in early grades. **Scale wound:** larger effectiveness trials show only a **fraction** of small efficacy-trial effects — coaching is hard to take to scale while keeping quality.

HYPOTHESIS: MindCraft tutor ops that treat onboarding as a PDF + one shadow session will reproduce the scale wound. Surviving design = small-cohort coaching cycles with session clips / traces (SAFE-FILM sparseness; SAFE-AAR), FEI rubric, and Map-brief fidelity checks — not “hire 200 tutors” vanity.

Cross-link SAFE-EXPECTANCY: Briefs are task/concept/next-move only; ban “she’s weak / anxious type” labels that become Golem scripts.

Kill: Headcount as product quality. Kill star-rating-only QA (smile ≠ transfer).

LXVIII.7 Playbook skeleton (product-facing)

HYPOTHESIS — minimum viable tutor playbook (pre→during→post):

1. **Pre (≤3 min):** Open Map diagnosis — target concept/bridge/format, recommended level, last soft-wrongs; task brief only (SAFE-EXPECTANCY).
2. **Open:** One destaked warm item or prior miss — not day-one ACT flood (SAFE-EXAM).
3. **Core:** Prompt-heavy cycle; student writes/tries before tutor speech; treat wrong answers as diagnostic (SAFE-MISCON); no permanent hint leash (SAFE-EXPOSE).
4. **AI rule:** Solver/coach text is draft; Map-checked ground truth; on hallucination run SAFE-REPAIR escalate packet — tutor does not improvise integrity denial.
5. **Close (SAFE-AAR light):** What was the goal? What evidence of progress? What broke? What is the next solo attempt?
6. **QA sample:** Rubric on talk ratio, Map fidelity, soft-wrong use, no trait labels, escalate correctness — not “energy.”

SPECULATION: Parents will pay for *auditable* tutor quality (session brief they can see + solo transfer trend) more than for “elite tutor” adjectives — test under SAFE-WTP CBC attribute “tutor follows your Map.”

LXVIII.8 Positioning

Competitor pattern	What they optimize	MindCraft HITL counter
ChatGPT / AI tutors	Fluent homework finish	Human + Map when AI fails; solo transfer still NS
Marketplace tutors	Hours / credentials	Playbook + Map brief + FEI QA
Khan / content	Video + practice	Diagnosis-informed human time, not another explainer
Duo streaks	Habit	Tutor does not become streak police (SAFE-HABIT)

Marketing that survives: “Your tutor works from the same living map the practice engine uses.” / “We train tutors to make you think — not to nod along.” / “When the AI is wrong, a human takes the packet — with a repair path, not a shrug.”

Marketing that dies: “Real humans, zero AI”; “AI so good you don’t need tutors”; “Ivy League tutors”; hours-booked vanity; Discord community as quality assurance.

LXVIII.9 SAFE-HITL stack (surviving doctrine)

1. **Map-briefed humans** — every session opens from inspectable diagnosis (SAFE-ONTOLOGY).
2. **Prompt > pour** — Chi-interactive default; lecture is the exception.
3. **Answer evidence > “got it?”** — ban comprehension-gauging as mastery.
4. **Polite ≠ correct** — train direct-but-kind correction; soft-wrong welcome.
5. **Complementarity** — AI for pace/draft language; human for diagnosis conflict, affect, escalate.

6. **Coaching cycles** — sustained QA; resist scale-without-fidelity (Kraft wound).
7. **FEI training** — tutors scored on recoverable struggle + next solo plan, not vibes.
8. **Escalate with packets** — SAFE-REPAIR; no infinite free human chat.
9. **No trait briefs** — SAFE-EXPECTANCY.
10. **No Discord=QA / hours=quality / Ivy=playbook.**

LXVIII.10 Experiments (HITL family)

ID	Contrast	Primary endpoints	Kill / survive rule
HITL-1	Map-briefed tutor vs chat-summary brief only	Tutor–system agreement; FEI event rate; student talk ratio	Chat brief $\approx$ Map $\rightarrow$ briefing UX fail
HITL-2	Prompt-first playbook vs explain-first default	Soft-wrong rate; <code>retry_120s</code> ; delayed mix	Explain-first wins transfer $\rightarrow$ revisit Chi mapping
HITL-3	Coached tutors (biweekly clip QA) vs onboarding-only	Rubric fidelity @30d; <code>solo_transfer_pass</code>	No lift $\rightarrow$ coaching theater (Kraft scale wound)
HITL-4	Tutor+Map+guarded Solver vs AI-alone practice	Persistence; transfer; escalate appropriateness	AI-alone $\geq$ HITL on transfer $\rightarrow$ human ops not earning keep
HITL-5	Ban “got it?” + require worked evidence vs free probes	False-mastery incidents coded	Free probes no worse $\rightarrow$ training cost unjustified
HITL-QUAL	10 tutor + 10 parent interviews on brief/QA trust	Message codebook	Feeds WTP attribute “Map-aligned tutor”

Ties: ONTO-3, EXP-, REPAIR-, AAR-, DP-, SE / Part XL, SAFE-EXPECTANCY.

LXVIII.11 Commercial implications

Copy: Sell *trained operators on a shared map*, not “humans vs robots.” Lead with what the tutor *does differently* in the first three minutes (opens your gap, not a generic worksheet).

Product: Tutor console = Map brief + soft-wrong feed + escalate packet + 4-question close. Student-facing: same Map the tutor sees (trust). Instrument HITL-1...4 before scaling marketplace headcount.

Growth: B2B/school pitch = consistent tutor quality via playbook + audit trail; consumer pitch = “not a random tutor hour — a session aimed at your next unlock.”

Vision: Thirty-year identity bet needs *witnesses* who can recognize competence honestly (Part LXI recognition factor) without colonizing with trait labels — HITL is how recognition stays human and falsifiable.

LXVIII.12 Confidence table

Claim	Label	Confidence
Naturalistic unskilled tutoring is collaborative Q&A/examples; deep ideal strategies rare (Graesser et al. 1995)	FACT	High
Student question <i>quality</i> (not count) links to achievement after experience (Graesser & Person 1994)	FACT	High
“Do you understand?” answers mislead; answer quality is better signal (Person et al. 1994)	FACT	High
Politeness strategies can inhibit effective tutoring (Person et al. 1995)	FACT	High
Prompt/scaffold-suppressed-explanation tutoring $\approx$ didactic on learning (Chi et al. 2001)	FACT	High
Teacher–AI co-orchestration with real-time awareness can improve learning vs AI-classroom without it (Holstein line)	FACT	Medium–High (classroom context)
Coaching improves instruction (~ 0.49 SD) and achievement (~ 0.18 SD); scale shrinks effects (Kraft et al. 2018)	FACT	High
Map-briefed + coached MindCraft tutors beat unbriefed warmth on FEI/transfer	HYPOTHESIS	Medium
SAFE-HITL is right commercial mapping	FOUNDER BELIEF	Medium

LXVIII.13 What this chapter kills

1. **Kill:** Warm human presence / Ivy credentials as quality proof.

2. **Kill:** Session hours or booking volume as North Star.
3. **Kill:** “Got it?” / smile / star-rating as mastery or QA.
4. **Kill:** Discord/Slack community $\equiv$ tutor training or QA.
5. **Kill:** Explain-first tutor brand (“we explain better than ChatGPT”).
6. **Kill:** AI-replaces-tutors *and* humans-don’t-need-Map dual fantasies.
7. **Kill:** One-off onboarding without coaching cycles (ignore Kraft scale wound).
8. **Wound:** Overclaiming Lumilo classroom effect sizes as MindCraft ACT insurance.
9. **Survive:** Graesser naturalistic baseline + training targets; Chi prompt default; Holstein complementarity; Kraft coaching-with-fidelity; HITL-1...5; Map-briefed FEI ops.

Doctrine until data: Human tutors are a **trainable operating system** on top of diagnosis — not a marketing prop and not an unexamined 1:1 myth. Without playbooks, QA, and FEI training, “human-in-the-loop” is just another expensive chatbot with a face.

Part LXIX — Privacy & Student Affect Data

Chapter status: Living evidence + product/ethics brief — Researcher tick 2026-08-03

Primary question: What affect data may MindCraft collect, infer, store, and act on — especially anxiety/stress telemetry — without turning identity transformation into emotional surveillance?

Owners: Product (check-in / recommend modifiers) · Engine (affective_state reads) · Legal/trust · Parent marketing · Tutor ops (brief hygiene)

Commercial job: Ship a **SAFE-PRIVACY** doctrine: voluntary, purpose-bound, short-lived, non-biometric affect *signals* that soften pedagogy — never face/voice emotion AI, never Anxiety Score™, never “FERPA compliant” as ethics proof, never trait labels for tutors or ads.

LXIX.1 Why this chapter exists

Parts XXIV / XLI / XLVII / LVIII / LXIII already treat affect as load-bearing for learning. The live stack already uses a **pre-session text check-in** (short free text $\rightarrow$ structured fields $\rightarrow$ `affective_state/{student_id}/latest`) so `/recommend` can soften `target_mastery` when stress is high and inject self-reported struggles. That is a product bet: affect informs *dose and destake*, not a mood dashboard brand.

The commercial temptation is the opposite: webcam “engagement,” voice stress, continuous Anxiety Score™, tutor heatmaps of “frustrated,” parent dashboards of raw emotion, or selling “empathic AI” against Khan/Duo. Emotion AI edtech markets that pitch. The research and child-rights literature say that pitch is scientifically fragile and ethically hot.

FOUNDER BELIEF under audit: MindCraft can use *self-authored* affect for FEI safety without becoming an emotional surveillance company — *if* collection is narrow, reversible, and never treated as ground-truth emotion.

Claims we refuse as doctrine:

1. Facial / vocal emotion AI $\equiv$ personalization or empathy.
2. Anxiety Score™ / continuous mood meter as North Star or parent report.
3. “FERPA compliant” / school-official exception $\equiv$ ethical clearance.
4. Opt-out checkboxes as meaningful consent for minors under compulsion.
5. Inferred affect as durable trait in tutor briefs (SAFE-EXPECTANCY collision).
6. Affect data as ad targeting, insurer, or third-party enrichment feedstock.
7. Soften-challenge forever because a model “detected stress.”

LXIX.2 Constructs (privacy language for product)

Construct	Research meaning	MindCraft analogue	Failure mode if misused
Affective computing	Systems that sense/respond to affect (Picard)	Soften recommend; expose ladders	Webcam “empathy” cosplay
Emotion recognition tech (ERT)	Claim to read emotion from face/voice/biometrics	Default ban	Barrett wound + McStay/Katirai harms
Self-report check-in	Student-authored state in words/scales	Pre-session 2–3 sentences $\rightarrow$ structured fields	Coercion; trait storage; no expiry
Contextual integrity	Privacy = appropriate flows for context norms	Tutoring context $\neq$ workplace emotion monitoring	“Parents already know stress” excuse
Purpose limitation	Use only for stated education/pedagogy purpose	Soften dose / inject struggle / destake	Scope creep to marketing/scoring
Emotional surveillance	Continuous monitoring that shapes behavior under watch	Live face/voice dashboards	Panopticon anxiety; performative calm
SAFE-PRIVACY	Doctrine for affect data	See LXIX.9	Privacy theater via policy PDF

Operational definition (HYPOTHESIS): An affect data practice is *MindCraft-honest* when (a) the student (or parent for young minors) knowingly authors the signal, (b) purpose is limited to session pedagogy modifiers + optional tutor task briefing, (c) retention is short and deletable, (d) no biometric ERT, (e) no durable emotion trait on the student card, and (f) refusal to check in still yields a usable (less customized) path.

LXIX.3 Affective computing $\neq$ license to watch faces

FACT (field founding + early ethics): Picard (1997, *Affective Computing*, MIT Press) framed computing that relates to, arises from, or deliberately influences emotion — and devoted explicit attention to moral/social risks (privacy, manipulation, deceptive machines). Picard (2003, *International Journal of Human-Computer Studies*, 59(1–2), 55–64, doi:10.1016/S1071-5819(03)00052-1) restates the privacy criticism head-on: emotions are “ultimately personal and private”; detection/recognition/manipulation can be the “ultimate breach” if done without acceptable norms — while also arguing ethical HCI uses exist when transparent and user-benefiting.

Commercial implication: Citing Picard to sell classroom cameras is selective quotation. The founding literature already flags privacy as first-class. MindCraft’s lawful heir is **user-controlled self-report** → **respectful response**, not covert sensing.

LXIX.4 Emotion-from-face science does not support product ERT

FACT (expressions are not reliable emotion meters): Barrett, Adolphs, Marsella, Martinez & Pollak (2019, *Psychological Science in the Public Interest*, 20(1), 1–68, doi:10.1177/1529100619832930) — the common view that a person’s emotional state can be readily inferred from facial movements is **not** supported as a strong, general rule. Expressions vary across cultures, situations, and people; similar face configurations map to multiple categories; scowls often communicate something *other* than the expected emotion. They urge caution for AI, education, and security consumers of emotion research.

FACT (edtech emotional AI risk): McStay (2020, *Learning, Media and Technology*, 45(3), 270–283, doi:10.1080/17439884.2020.1686016) — evaluates facial-coding emotional AI marketed into education/SEL quantification; finds serious method risks (material effects on students) plus ethical/legal clash of private EdTech interests with public education goods; concludes **significant risks** in classroom deployment.

FACT (ethics review of ERT): Katirai (2023, *AI and Ethics*, doi:10.1007/s43681-023-00307-3) — structured review (43 articles): biased/unfair outcomes from faulty premises; sensitivity of emotion data; harm risk in consequential settings including **education**; guidelines stress defined scope, fairness, privacy — and the review finds issues may be “potentially insurmountable” even as commercial ERT expands.

FACT (soft biometrics / weak privacy consensus): McStay (2020, *Big Data & Society*, doi:10.1177/2053951720904386) — stakeholder interviews + UK survey: only a *weak* consensus on privacy for emotional AI / soft biometrics; limited window to set stronger norms before deployment normalizes.

Kill: Webcam engagement meters, “we can see when you’re stuck,” voice-stress ACT coach, any vendor demo that maps smile→mastery.

Wound: Text sentiment models over check-in prose are *not* Barrett-safe either — they are inferences. Prefer structured student choices (“stressed / okay / ready”) plus free text stored as evidence of *self-report*, not as ground-truth emotion labels sold upstream.

LXIX.5 Student privacy law is procedural — ethics still required

FACT (FERPA’s limits in big-data edtech): Zeide (2016, *Drexel Law Review*, 8, 339–480; SSRN:2821837) — FERPA’s FIPPs-flavored access-control model **delegates** most privacy decisions to institutions via broad exceptions

(notably the school-official exception); transparency/oversight/accountability are thin; notice-and-consent is a poor fit under compulsory schooling and credential pressure.

FACT (purpose-limitation theater): Zeide (2017, *University of Miami Law Review*, 71(2), 494–533) — “education purpose” limitations in FERPA/state reforms are often **procedural** (on behalf of / authorized by schools), not substantive student-interest protections; data-driven tools change what student information is and how decisions are made; ethical discourse must go beyond checkbox purpose.

FACT (contextual integrity): Nissenbaum (2010, *Privacy in Context*, Stanford University Press; see also 2004, *Washington Law Review*, 79, 119) — privacy violations are inappropriate information *flows* relative to context norms (actors, attributes, transmission principles), not merely “data left the device.”

HYPOTHESIS (MindCraft mapping): A stress check-in that only the recommend engine and (optionally) a Map-briefed tutor see for *this session*, with short TTL, can respect tutoring-context norms. Piping the same signal into parent weekly “Anxiety Score,” marketing cohorts, or tutor trait cards **breaks** contextual integrity even if a lawyer stamps FERPA.

Kill: Marketing “bank-grade / FERPA compliant” as proof we may collect facial emotion or forever-retain affect. Compliance ≠ appropriateness.

LXIX.6 What MindCraft should collect (and refuse)

FOUNDER BELIEF → **product rule (provisional SAFE-PRIVACY):**

Allowed (with hygiene)	Banned by default
Voluntary pre-session text / short Likert stress	Webcam, mic emotion, wearables for “engagement”
Explicit struggle tags student selects	Auto-inferred personality / “fragile” traits
Session-scoped modifiers (soften mastery target; destake)	Permanent Anxiety Score™ on profile / parent SKU
Aggregate, de-identified product analytics	Sale/enrichment of affect to third parties
Tutor brief: <i>task</i> implications only (“high stress → destake first item”)	Tutor brief: emotion labels as person descriptors
Student delete / skip check-in without product exile	Dark-pattern mandatory mood gate every open

Cross-links: SAFE-TIMING (stress softens load); SAFE-DD / SAFE-EXPOSE (destake before dose); SAFE-EXPECTANCY (no trait labels); SAFE-EXAM (learn vs prove rails); SAFE-HITL (tutors see Map tasks, not psych dossiers); SAFE-STRUCTURE (scarce bandwidth ≠ moral failure).

HYPOTHESIS: Soften-on-high-stress survives only if we measure that softened paths still produce `retry_120s` and delayed mix — not merely calmer NPS. Permanent under-challenge from stale stress flags is a privacy *and* learning failure.

LXIX.7 Commercial positioning

Vs Emotion-AI EdTech: Do not compete on “empathic cameras.” Compete on **honest self-report** → **inspectable pedagogy change** (Map + FEI). That is a trust wedge for parents who fear surveillance schools and for districts whose procurement now asks biometric questions.

Vs ChatGPT tutors: Generic chat may *perform* empathy in text while logging prompts indefinitely. MindCraft can claim *narrower* affect handling: purpose-bound fields, short retention, no biometric ERT — but only if true in the privacy policy *and* the schema.

Parent copy (SAFE-WTP compatible): “We ask how today’s session feels — so practice doesn’t pile on when stress is high. We don’t watch faces or sell mood data.” Do not promise clinical mental-health care.

Growth: Treat privacy review packets (data map, retention, ERT ban, delete path) as a **sales asset** for school pilots — Zeide implies schools outsource tools; buyers need substantive answers, not FERPA slogans.

LXIX.8 Experiments (PRIV family)

ID	Contrast	Primary endpoints	Kill / survive rule
PRIV-1	Check-in on (session TTL) vs no check-in	retry_120s; delayed mix; skip rate	Check-in harms transfer or >X% feel coerced → redesign/remove
PRIV-2	Soften-on-high-stress vs ignore stress	WM-demand item accuracy; return; transfer	Ignore ≥ soften on learning without affect cost → modifier useless
PRIV-3	Task-only tutor affect note vs emotion-label brief	Student trust items; expectancy language in tutor talk	Label brief ↑ trait talk / ↓ challenge → ban labels (EXPECTANCY)
PRIV-4	24h TTL vs 30d retain affect fields	Modifier usefulness; parent/student deletion requests	Long retain no lift → shorten; creep detected → kill retain
PRIV-5	Explicit ERT mock (face) vs text check-in preference	Forced choice; WTP attribute in CBC	Face preferred <i>and</i> lifts FEI → reopen only with ethics board (unlikely)
PRIV-QUAL	10 Maya + 10 parent interviews on affect data comfort	Codebook: surveillance vs care	Feeds WTP “no webcam emotion” attribute

Ties: TIM-3, EXP-O-, EXAM-, HITL-1, WTP-*, SAFE-EXPECTANCY.

LXIX.9 SAFE-PRIVACY doctrine (provisional)

1. **Self-authored first.** Prefer student (or parent-mediated) report over inferred biometrics.
 2. **Purpose-bound.** Affect → pedagogy modifiers and task briefs only; never ads, scores-as-SKU, or trait archives.
 3. **Short TTL + delete.** Session-to-days, not forever; visible delete/skip.
 4. **ERT ban.** No face/voice/wearable emotion recognition in product roadmap without extraordinary evidence *and* rights review (Barrett + McStay bar is high).
 5. **Contextual integrity.** Flows must match tutoring norms; parent summaries = actions taken (“we destaked”), not raw emotion dumps.
 6. **Refusal-safe.** Skip check-in still gets a full product.
 7. **Compliance ≠ ethics.** FERPA/COPPA/state student-privacy laws are floors; Nissenbaum appropriateness is the product bar.
 8. **Measure learning, not calm.** Soften flags judged by FEI metrics, not Anxiety Score™ decline.
-

LXIX.10 Confidence table

Claim	Label	Confidence
Picard founding work includes privacy/manipulation concerns (1997; 2003 IJHCS)	FACT	High
Facial movements are unreliable general emotion meters (Barrett et al. 2019)	FACT	High
Classroom facial emotional AI carries significant method + ethics risks (McStay 2020 LMT)	FACT	High
ERT ethics literature flags bias, sensitive data, education harms (Katirai 2023)	FACT	High
FERPA/school-official + purpose limits are weak substantive privacy for edtech (Zeide 2016/2017)	FACT	High
Contextual integrity is the right evaluation frame for affect flows (Nissenbaum)	FACT (framework)	High
Text check-in soften → better FEI than no affect use	HYPOTHESIS	Medium
SAFE-PRIVACY is right commercial mapping for MindCraft	FOUNDER BELIEF	Medium–High

LXIX.11 What this chapter kills

1. **Kill:** Facial/vocal/wearable emotion AI as MindCraft personalization.
2. **Kill:** Anxiety Score™ / continuous mood dashboard as NS or parent SKU.
3. **Kill:** “FERPA compliant” as substitute for substantive affect ethics.
4. **Kill:** Opt-out theater / mandatory mood gates for minors.
5. **Kill:** Emotion labels as durable tutor-brief traits.
6. **Kill:** Affect data sale, ad cohorts, or insurer-style enrichment.
7. **Kill:** Empathy-camera marketing vs Khan/Duo.
8. **Wound:** Treating LLM sentiment over journals as “measured anxiety.”
9. **Survive:** Voluntary short-TTL self-report → soften/destake; PRIV-1...5; Barrett/McStay/Zeide/Nissenbaum stack; ERT ban as trust wedge.

Doctrine until data: Affect is for **permission and dosing**, not for watching. If the company needs a camera to know a student is scared, it has already lost the FEI plot — and the parent trust market.

Part LXX — Constitution Meta: How the Lab Updates Beliefs

Chapter status: Living evidence + process brief — Researcher tick 2026-08-03

Primary question: How should MindCraft revise Constitution claims so the company gets *truer* product law — not a longer PDF, not a shrine of founder priors, and not “science-backed” marketing from unfalsified prose?

Owners: Research Lab · Founders · Product/analytics (experiment falsifiers) · Brand (claim ladder hygiene)

Commercial job: Ship a **SAFE-LABMETA** doctrine: label → check → kill/wound/survive → demote marketing until FEI evidence; Bayesian *practice* means model-check and revise, not vibes-in-Bayes-clothing.

LXX.1 Why this chapter exists

The Constitution is already large. Without a belief-update OS, volume becomes cargo cult: agents append chapters, confidence tables fill with “Medium,” and marketing quietly promotes FOUNDER BELIEF as FACT because the PDF looks thick.

FOUNDER BELIEF under audit: A living research Constitution can be a *competitive* asset — sharper copy, killed bad metrics, defensible trust packets — *if and only if* it behaves like a research programme with falsifiers, not a manifesto that only grows.

Claims we refuse as doctrine:

1. Chapter count / page count / “150–300 pages” as progress North Star.
2. “Bayesian” branding without prior, likelihood, or model check — vibes posterior.
3. Cite-washing: DOI lists that do not change a product rule or kill a slogan.
4. Repetition of FOUNDER BELIEF until it is treated as FACT.

5. Synthesizer as marketing refresh that preserves weak claims.
6. “Science-backed” / “research Constitution” in ads without named falsifiers and claim ladder L0–L4.
7. Fake numeric confidence (e.g. “87% sure”) without evidence protocol.

LXX.2 Constructs (lab language for company law)

Construct	Research meaning	MindCraft analogue	Failure mode if misused
Falsifiability	Scientific claims risk refutation by observation (Popper)	Every SAFE-* row needs a kill condition	Unfalsifiable identity slogans
Research programme	Hard core + protective belt of auxiliaries (Lakatos)	FEI hard core; SAFE-* as adjustable belt	Hardening every slogan as core
Progressive vs degenerating	Novel predictions confirmed vs endless ad-hoc patches	Experiments that bite product	Adding caveats forever, never killing
Bayesian practice	Build → infer → <i>check model against data</i> → revise (Gelman/Shalizi)	Tick: write → Red Team → experiment → synthesizer	Prior worship / no checks
Researcher degrees of freedom	Flexible analysis inflates false positives (Simmons et al.)	Post-hoc metric shopping after A/B	“It worked on engagement”
Pre-registration / disclosure	Timestamp plans; reduce publishability bias (Nosek et al.)	Part IX families; Appendix D sketches	Secret endpoint swaps
Replication reality	Many published effects shrink or vanish (OSC; Ioannidis frame)	Distrust one study → ACT guarantee	Overclaim from single paper

Operational definition (HYPOTHESIS): A Constitution update is *honest* when (a) claims keep FACT/HYPOTHESIS/FOUNDER BELIEF/SPECULATION labels, (b) Red Team can name what would kill the claim, (c) commercial copy cannot climb the claim ladder without matching evidence, and (d) synthesizer *removes* as often as it merges.

LXX.3 Falsifiability is the lab’s anti-marketing vaccine

FACT (demarcation / method): Popper (1959 English ed. of *Logik der Forschung*; Routledge reprints) argues scientific theories are distinguished by *falsifiability*: universal claims are never proven by accumulation of confirming instances, but can be contradicted by reproducible effects. Single stray anomalies do not auto-refute; *reproducible* refuting effects do.

Commercial implication: MindCraft marketing that cannot state a disconfirming user outcome (“if solo_transfer_pass does not beat unguarded chat, we retire Solver-as-hero”) is not “science-backed” — it is branding with footnotes.

Kill: “We transform identity” without behavioral co-primaries (SAFE-IDMEASURE / LONGID).

Survive: Soft-wrong, Map diagnosis, PWC — each already carries falsifiers in Part IX families.

LXX.4 Treat the Constitution as a Lakatosian programme — not a Kuhn shrine

FACT (research programmes): Lakatos (1970, in Lakatos & Musgrave, *Criticism and the Growth of Knowledge*; developed in *The Methodology of Scientific Research Programmes*) models science as sequences of theories with a **hard core** protected by a **protective belt** of auxiliaries. The negative heuristic forbids directing every anomaly at the core; the positive heuristic guides how to revise the belt. Programmes are judged progressive vs degenerating by whether they yield novel confirmed content rather than endless patchwork.

HYPOTHESIS (MindCraft mapping):

- **Hard core (methodologically sticky):** Fear → Evidence → Identity (FEI); identity ≠ answer delivery; explanations-are-cheap market hypothesis; claim ladder L0–L4.
- **Protective belt (must take hits):** Individual SAFE-* product rules, UX defaults, copy bans, experiment families.
- **Degenerating smell:** New chapters that only add caveats around a failing metric (e.g. streaks) without killing the metric.

FOUNDER BELIEF: FEI may stay core while many SAFE-* rows die. That is health, not betrayal.

Kill: Treating every Red Team kill as optional flavor text while shipping the banned feature anyway.

Wound: Over-hardening FEI itself — if transfer never moves and threat never drops, the core must face a crisis review (quarterly “what we believed that was wrong” memo in Part XXIII).

LXX.5 “Bayesian” without model checking is cosplay

FACT (practice ≠ inductivist myth): Gelman & Shalizi (2013, *British Journal of Mathematical and Statistical Psychology*, 66(1), 8–38, doi:10.1111/j.2044-8317.2011.02037.x) argue that successful Bayesian *data analysis* fits hypothetico-deductive practice: models are built, posteriors computed, then **checked** (e.g. posterior predictive checks) and **revised**. Model checking sits outside naive “Bayes = automatic rationality / pure induction.” Priors are part of the model, not sacred beliefs immune to data.

Commercial implication: Lab agents must not write “we Bayesian-updated to High confidence” after citing three friendly papers. The analogue of a posterior predictive check is: **did the experiment family falsify the product rule?** If experiments are unrun, confidence stays capped (typically Medium for HYPOTHESIS; FOUNDER BELIEF never becomes FACT by eloquence).

Lab update rule (HYPOTHESIS → doctrine candidate):

1. **Prior:** label + stated confidence + commercial temptation.

2. **Likelihood:** citations with contradicting evidence mandatory.
3. **Check:** Red Team tick + named falsifier in Part IX.
4. **Revise:** kill / wound / survive; synthesizer promotes only survivors to I.4.
5. **Demote copy:** marketing may use only surviving I.4 rows at the claim-ladder level earned.

Kill: Bayesian Score™ / “posterior confidence” widgets in investor decks.

Survive: Epistemic labels + experiment falsifiers + synthesizer deletions.

LXX.6 Why published ed/psych findings are a hostile prior

FACT (positive predictive value frame): Ioannidis (2005, *PLoS Medicine*, 2(8), e124, doi:10.1371/journal.pmed.0020124) argues many published research findings are likely false when studies are small, effects small, designs flexible, bias high, or pre-study odds poor — claimed findings can track prevailing bias.

FACT (replication estimate): Open Science Collaboration (2015, *Science*, 349(6251), aac4716, doi:10.1126/science.aac4716) — 100 psychology replications: original significant rate ~97%; replication significant rate **36%**; replication effect sizes roughly **half** originals on average.

FACT (flexibility): Simmons, Nelson & Simonsohn (2011, *Psychological Science*, 22(11), 1359–1366, doi:10.1177/0956797611417632) show undisclosed flexibility in collection/analysis (“researcher degrees of freedom”) can make almost any hypothesis appear significant; propose disclosure requirements (and relatedly motivate pre-registration culture).

FACT (incentive repair): Nosek, Spies & Motyl (2012, *Perspectives on Psychological Science*, 7(6), 615–631, doi:10.1177/1745691612459058) — incentives favor novelty/publishability over truth; propose restructuring practices toward transparency (including open workflow tools such as OSF) so accuracy can compete with publishability.

Commercial implication: MindCraft must not turn a single encouraging tutoring or mindset paper into ACT-point packaging. External literature sets **priors and ban lists**; company truth requires **MindCraft experiments** (FEI endpoints), not literature cosplay.

Kill: “Meta-analysis says tutoring works → our app raises ACT 3 points.”

Survive: Literature → mechanism hypothesis → pre-registered family → ship only what survives.

LXX.7 How a tick is supposed to change belief (not page count)

Map the existing cadence (skill + Part XXIII) onto Gelman/Shalizi + Lakatos:

Tick	Epistemic job	Degenerate failure
Researcher	Add labeled claims + citations + commercial implication + experiment IDs	Cite-wash; no falsifier; fluff length

Red Team	Attempt kill of weakest claim; update confidence/contradictions	Soft applause; no log of kill/wound
Synthesizer	Merge survivors into I.4; <i>delete</i> duplicates/weak args; refresh Exec Summary	Accumulate rows; never demote

HYPOTHESIS (stopping rule): If ≥ 8 researcher ticks land without synthesizer, provisional SAFE- rows *are not** company law for ads — they are candidates. Marketing that quotes provisional doctrine as settled FACT violates claim ladder.

SPECULATION: Unattended agents drift toward agreeable synthesis; Red Team hours (UTC % 6 == 0) exist to counter that drift — unless the next chapter was never written (prefer Researcher so the queue does not starve).

LXX.8 Commercial positioning

Vs competitors’ “science-backed” pages: Do not out-footnote Khan/Duo/Brilliant. *Out-kill:* publish (internally, then selectively externally) what metrics and slogans you retired. Trust is a progressive research programme signal.

Parent/district copy: “We pre-register what ‘better’ means — retry, transfer, solo prove — and we drop features that only juice streaks.” Never: “Our 200-page Constitution proves identity transformation.”

Growth: Experiment instrumentation (Part XXI) is GTM infrastructure: schools and serious parents buy *inspectable* claims. SAFE-LABMETA makes SAFE-WTP and district privacy packets coherent — same honesty stack.

Investor hygiene: Deck claims must cite I.4 survivors + claim ladder level, not chapter count since last board meeting.

LXX.9 Experiments (LABMETA family)

ID	Contrast	Primary endpoints	Kill / survive rule
LABMETA-1	Pre-registered FEI endpoints vs post-hoc engagement primary	Pre-reg adherence; false-positive feature ships	Post-hoc primary wins shipping → process fail; restore pre-reg gate
LABMETA-2	Red Team mandatory before I.4 promote vs promote-on-write	Kill rate; later reverse-kills	Zero kills over N ticks → RT theater; change RT incentives
LABMETA-3	Marketing copy lint vs I.4+ladder vs unrestricted “science-backed”	Claim-ladder violations caught; parent trust items	Lint catches nothing while ads overclaim → strengthen lint

LABMETA-4	Synthesizer deletion quota (must demote ≥ 1 weak row) vs merge-only	Rows removed; Exec Summary length	Merge-only $\rightarrow$ synthesizer fail; enforce demote
LABMETA-5	External skeptical reader quarterly vs internal-only	Claims wounded/killed by outsider	No external wounds ever $\rightarrow$ capture; require outsider
LABMETA-QUAL	5 staff write “belief we retired” memos	Codebook of real vs performative kills	Feeds quarterly wrong-belief memo

Ties: DAG- (*estimands*), Part XXI metrics, claim ladder L0–L4, all SAFE- falsifiers.

LXX.10 SAFE-LABMETA doctrine (provisional)

1. **Labels are load-bearing.** FACT/HYPOTHESIS/FOUNDER BELIEF/SPECULATION stay visible; repetition does not promote labels.
2. **Falsifier or it does not ship as law.** No SAFE-* without kill condition and experiment family hook.
3. **Check, then update.** Citations without Red Team / model-check analogue are bibliography, not belief change (Gelman & Shalizi).
4. **Hard core is small.** FEI + claim ladder + cheap-explanations market bet; everything else is belt.
5. **Progressive or cut.** Auxiliaries that only patch failing North Stars (streaks, calm meters, costume pedagogy) get killed, not footnoted.
6. **Hostile prior on single papers.** OSC/Ioannidis/Simmons frame: external lit informs; MindCraft FEI data decides shipping.
7. **Disclosure over cleverness.** Pre-reg endpoints; no silent metric swaps (Nosek; Simmons).
8. **Synthesizer must delete.** Merge without demotion is document inflation.
9. **Page count is vanity.** Truth density and killed claims are the lab KPIs.
10. **Marketing $\leq$ evidence.** “Science-backed” only for I.4 survivors at earned ladder level.

LXX.11 Confidence table

Claim	Label	Confidence
Popperian falsifiability as demarcation/method ideal (1959 LoSD)	FACT (philosophy)	High
Lakatos hard core / protective belt / progressive vs degenerating (1970)	FACT (framework)	High
Bayesian practice centers model checking + revision (Gelman & Shalizi 2013)	FACT	High

Many published findings have poor PPV / replication attrition (Ioannidis 2005; OSC 2015)	FACT	High
Analysis flexibility inflates false positives; disclosure helps (Simmons et al. 2011)	FACT	High
Incentives skew publishability; transparency reforms help (Nosek et al. 2012)	FACT	High
SAFE-LABMETA maps correctly onto MindCraft commercial risk	FOUNDER BELIEF	Medium–High
Enforcing deletion quotas + copy lint raises parent trust	HYPOTHESIS	Medium

LXX.12 What this chapter kills

1. **Kill:** Page/chapter count as lab or company North Star.
2. **Kill:** “Bayesian update” theater without checks, experiments, or label discipline.
3. **Kill:** Cite-washing that never changes I.4 or ads.
4. **Kill:** Promoting FOUNDER BELIEF → FACT by repetition or PDF girth.
5. **Kill:** Synthesizer-as-brochure that never demotes weak doctrine.
6. **Kill:** “Science-backed Constitution” marketing without falsifiers + claim ladder.
7. **Kill:** Fake precision confidence numbers as persuasion.
8. **Wound:** Treating FEI hard core as unchallengeable forever without quarterly crisis review.
9. **Survive:** Label → Red Team → experiment → synthesizer demote/promote; Gelman/Shalizi check ethos; Lakatos belt discipline; hostile prior on single-study overclaim; LABMETA-1...5.

Doctrine until data: The lab’s product is **revised belief that constrains shipping and copy** — not pages. If the Constitution cannot kill a beloved feature, it is a mood board with DOIs.

Part LXXI — Peer vs Near-Peer vs Expert Tutoring

Chapter status: Living evidence + hire/ops brief — Researcher tick 2026-08-03

Primary question: At what *grain of tutor expertise* does MindCraft get learning and identity gains — same-age peer, near-peer (college→HS), or expert (teacher/pro) — and what hire bar survives evidence rather than Ivy theater?

Owners: Tutor ops · Hiring · Product (tutor matching) · Positioning vs marketplace “elite tutors” / peer-only Discord help

Commercial job: Ship a **SAFE-TUTORGRAIN** doctrine: MindCraft’s default human is a *trained near-peer operator on the Map* — not “any peer is fine,” not “PhD/Ivy $\equiv$ better FEI,” and not Bloom 2σ expert cosplay. Hire for trainable Map fidelity + prompt>pour; measure grain with transfer, not résumé.

LXXI.1 Why this chapter exists

Part LXVIII (SAFE-HITL) specifies *how* tutors operate once hired. This chapter specifies *who* to hire and what “expertise” is allowed to mean in copy. MindCraft’s production model is college tutors with high-school students — **near-peer** / **cross-age** in Topping’s typology — while competitors sell either peer Discord vibes or “expert” marketplace credentials. The wrong grain kills unit economics *or* pedagogy: untrained same-age peers default to knowledge-telling; credentialed experts still fail diagnosis and can over-pour; untrained near-peers reproduce Graesser’s naturalistic baseline without FEI.

FOUNDER BELIEF under audit: Near-peer + Map + HITL playbook beats (a) expert-without-Map lecture and (b) peer-without-structure chat — on solo transfer and recoverable struggle — at a cost parents will pay.

Claims we refuse as doctrine:

1. Expert / Ivy / teacher-credential $\equiv$ better FEI outcomes without Map+playbook.
 2. Any peer help (Discord, classmate) $\equiv$ MindCraft tutoring.
 3. Bloom “2-sigma human tutors” as marketing (standing ban; VanLehn wound).
 4. Domain-expert monologue $\equiv$ tutoring quality (Chi interactive kill already).
 5. Tutor learning as a substitute for *student* transfer (Roscoe/Chi bias).
 6. Hire bar = GPA / school brand without fidelity QA.
-

LXXI.2 Constructs (ops language)

Construct	Research meaning	MindCraft analogue	Failure mode if misused
Same-age peer	Tutor $\approx$ tutee grade/age; often reciprocal	Optional student $\leftrightarrow$ student practice (not default paid tutor)	Unstructured Discord $\equiv$ product
Near-peer / cross-age	Tutor older/further along same trajectory	College tutor $\leftrightarrow$ HS Maya	“Big sibling vibes” without playbook
Expert tutor	Teacher / experienced / high domain+pedagogy	Rare escalate / coach-of-tutors	Ivy/PhD SKU as quality
Knowledge-telling	Deliver summary; low monitoring/elaboration	Explain-first college tutor default	Fluency without FEI
Knowledge-building	Monitor, integrate, elaborate (tutor <i>and</i> tutee)	Prompt>pour; soft-wrong springboard	Treat as automatic from “being smart”

Expertise grain	How much credential/experience changes outcomes	Hire bar + matching policy	Resume theater
SAFE-TUTORGRAIN	Doctrine for who tutors	See LXXI.9	Peer-chaos or expert-cosplay

Operational definition (HYPOTHESIS): A MindCraft human tutor is *grain-honest* when they are (a) near-peer by default (college→HS), (b) screened for Map-brief fidelity and prompt-first moves, not school brand, (c) coached under SAFE-HITL, and (d) escalated to expert grain only for packetized hard cases (SAFE-REPAIR / clinical-affect out of scope) — never marketed as “experts on every call.”

LXXI.3 Peer tutoring works — and structure is the moderator

FACT (classic school meta): Cohen, Kulik & Kulik (1982, *American Educational Research Journal*, 19(2), 237–248, doi:10.3102/00028312019002237) — 65 tutoring-program evaluations; tutored students outperform controls on exams and show more positive subject attitudes; *tutors* also gain understanding/attitudes; little/no self-esteem effect. Peer/near-peer school tutoring is not a null intervention.

FACT (updated peer meta + moderators): Leung (2015, *Journal of Educational Psychology*, 107(2), 558–579, doi:10.1037/a0037698) — comprehensive peer-tutoring meta; positive impact on achievement; moderators include role structure (same-age reciprocal vs nonreciprocal vs cross-age), educational level, and program parameters grounded in role theory / interdependent contingencies — i.e. **design beats “put two kids together.”**

FACT (HE typology): Topping (1996, *Higher Education*, 32(3), 321–345, doi:10.1007/BF00138870) — peer tutoring in further/higher education has many formats; cross-year small-group, PSI, and supplemental models are among those with substantial evidence and wider dissemination merit. MindCraft’s college→HS pairing is closer to **cross-age** / **cross-year** than to same-year reciprocal undergrad SI — do not cite Topping as “Discord study groups work.”

Commercial implication: Peer/near-peer is *evidence-eligible* as a delivery grain. Unstructured peer chat is not. Sell structured near-peer sessions, not “community help.”

LXXI.4 Expert humans are not 2σ magic — and diagnosis is weak

FACT (human tutoring ES vs folklore): VanLehn (2011, *Educational Psychologist*, 46(4), 197–221, doi:10.1080/00461520.2011.611369) — review finds human tutoring effect size $\approx d = 0.79$ vs no-tutoring same content — far below Bloom’s folklore ~ 2.0 ; step-based ITS $\approx d = 0.76$, nearly matching human. Interaction-plateau: finer-than-step grain does not reliably beat step grain.

FACT (tutors mis-model students): Chi, Siler & Jeong (2004, *Cognition and Instruction*, 22(3), 363–386, doi:10.1207/s1532690xci2203_4) — novice-but-domain-knowledgeable tutors monitor student understanding poorly; miss most student-initiated knowledge deficits; everyday tutors are usually novices (Fitz-Gibbon cited therein) yet still produce gains (ties Cohen et al.).

Cross-link: Graesser naturalistic tutors lack deep ideal strategies (Part LXVIII); Putnam-style curricular scripts appear even in experienced tutors (cited in tutoring/ITS literature). Expertise ≠ accurate student model.

Kill: “Our expert tutors close the 2-sigma gap.” Kill resume=diagnosis. Wound: treating VanLehn’s human≈ITS parity as “AI replaces humans” — MindCraft’s bet is FEI + Map complementarity (SAFE-HITL / SAFE-ONTOLOGY), not ITS cosplay.

LXXI.5 Peer tutors default to knowledge-telling

FACT (tutor learning review): Roscoe & Chi (2007, *Review of Educational Research*, 77(4), 534–574, doi:10.3102/0034654307309920) — peer tutors *can* learn from tutoring via reflective knowledge-building (monitor, integrate, elaborate), but show a pervasive **knowledge-telling bias** — summarizing/delivering with little monitoring — even when trained; tutor-learning potential is often unrealized.

FACT (process study): Roscoe & Chi (2008, *Instructional Science*, 36(4), 321–350, doi:10.1007/s11251-007-9034-5) — untrained peer tutors learn most when explaining/responding incorporates reflective knowledge-building; knowledge-telling dominates; inferential tutee questions more often elicit elaborative tutor responses.

HYPOTHESIS (MindCraft): College near-peers without HITL training will **pour** (tell) under parent pressure to “explain until she gets it,” reproducing the bias Roscoe/Chi document — and colliding with Chi et al. (2001) prompt≈didactic lesson (SAFE-HITL). Training must suppress telling-as-default, not celebrate “natural explainers.”

Kill: Hire for “great at explaining.” Kill tutor-learning / “tutors grow too” as the student-facing value prop (nice externality; not the NS).

LXXI.6 Near-peer is the default grain — expert is escalate, not SKU

HYPOTHESIS (grain map):

Grain	When it earns keep	When it fails
Same-age peer	Optional structured practice pairs; CoP peripheral work with old-timer (SAFE-CoP)	Paid product; unmoderated forums; peer ranks (SAFE-COMPARE)
Near-peer (college→HS)	Default paid human; Map-briefed; prompt-first; FEI QA	Vibes-only; knowledge-telling; trait expectancy briefs
Expert (teacher/coach)	Coach-of-tutors; hard escalate packets; rare clinical/pedagogy consult	Every-session expert SKU; 2σ marketing; lecture-first “master teacher”

FOUNDER BELIEF: Near-peer *rapport* + *recent struggle memory* can support belonging/recognition (Part LXI) without expert costume — *if* Map prevents “I remember how hard this was” from becoming soft-wrong avoidance (politeness tax, Part LXVIII).

SPECULATION: Parents may *say* they want experts in surveys and *choose* near-peer+Map+transfer report in CBC (test under SAFE-WTP; do not treat Likert “prefer certified teacher” as WTP).

Wound: Classroom peer-tutoring metas ≠ remote 1:1 ACT prep. Borrow structure moderators (Leung), not K–8 effect sizes as ACT guarantees.

LXXI.7 Hire bar (product-facing)

HYPOTHESIS — minimum viable tutor screen (ordered):

- 1. **Content floor:** Can work target ACT-band items at recommended level without Solver crutches (spot-check).
- 2. **Map fidelity drill:** Given a fake brief, names concept/bridge/format and next move — no trait labels (SAFE-EXPECTANCY).
- 3. **Prompt>pour drill:** On a soft-wrong, asks before explains; refuses “got it?” as close (SAFE-HITL).
- 4. **Repair drill:** On planted AI wrongness, owns + Map-checks + destaked reattempt path (SAFE-REPAIR).
- 5. **School brand / GPA:** Tie-break only — never primary.

Ops: Expert grain = internal coach + sampled session QA (Kraft coaching fidelity wound already in LXVIII), not a parent-facing “Expert Tutor” upsell that implies every minute is master-teacher.

LXXI.8 Positioning

Competitor pattern	What they optimize	MindCraft TUTORGRAIN counter
Marketplace “elite tutors”	Credentials / hourly	Near-peer + Map + fidelity QA; publish grain honesty
Peer Discord / homework help	Availability / vibe	Structured session ≠ chat; CoP ban without repertoire
AI-only tutors	Fluency / cost	Human near-peer for escalate + recognition; solo transfer NS
School peer programs	Coverage / cost	Cross-age trained operators, not same-age homework swap

Marketing that survives: “College tutors trained on your living Map — not random homework help.” / “We hire for how they coach struggle, not the logo on their hoodie.” / “Expert coaches train our tutors; your session is aimed near-peer time, not a TED talk.”

Marketing that dies: “PhD/Ivy tutors”; “2-sigma expert instruction”; “peers teaching peers” as unpaid Discord; “any human is better than AI” without grain+playbook.

LXXI.9 SAFE-TUTORGRAIN stack (surviving doctrine)

1. **Default grain = trained near-peer** (college→HS), not same-age peer product and not expert-every-call.
2. **Structure > pairing** — Leung/Topping: role design and training beat “put humans together.”
3. **No 2σ / expert-magic claims** — VanLehn human $\approx$ 0.79; Bloom folklore banned.
4. **Expertise $\neq$ student model** — Chi/Siler/J Jeong monitoring wound; Map carries diagnosis.
5. **Suppress knowledge-telling** — Roscoe/Chi; prompt>pour hire+QA.
6. **Hire bar = fidelity drills**, not school brand.
7. **Expert = escalate + coach-of-tutors**, not default SKU.
8. **Tutor learning $\neq$ student NS** — externality only.
9. **Ties:** SAFE-HITL, SAFE-ONTOLOGY, SAFE-EXPECTANCY, SAFE-REPAIR, SAFE-CoP, SAFE-WTP.

LXXI.10 Experiments (GRAIN family)

ID	Contrast	Primary endpoints	Kill / survive rule
GRAIN-1	Trained near-peer vs untrained near-peer (same Map)	Soft-wrong rate; talk ratio; solo_transfer_pass	Untrained $\approx$ trained $\rightarrow$ training theater
GRAIN-2	Near-peer+Map vs expert-credential tutor without Map brief	Map agreement; FEI events; transfer	Expert-no-Map $\geq$ near-peer+Map $\rightarrow$ revisit grain bet
GRAIN-3	Near-peer+HITL vs same-age structured peer pair	Transfer; challenge_accept; safety codes	Peer-pair $\geq$ near-peer on transfer $\rightarrow$ paid grain unjustified
GRAIN-4	Hire-by-fidelity-drills vs hire-by-school-brand	30d rubric fidelity; escalate appropriateness	Brand hires win fidelity $\rightarrow$ drills wrong
GRAIN-5	Parent CBC: near-peer+Map report vs “certified expert” attribute	Choice share; WTP	Expert attribute dominates at equal price $\rightarrow$ message/package work (not auto-hire experts)
GRAIN-QUAL	10 tutor + 10 parent interviews on “who should tutor Maya”	Codebook: credential vs process trust	Feeds copy + GRAIN-5

Ties: HITL-, ONTO-, EXP-, REPAIR-, WTP-, CoP-.

LXXI.11 Commercial implications

Copy: Lead with *trained near-peer on your Map*. Never lead with Ivy. Never imply every tutor is a master teacher. Never call Discord “tutoring.”

Product: Matching defaults to near-peer; expert badge only on coach-of-tutors / escalate path. Instrument GRAIN-1/2 before raising expert SKUs. Tutor application = drills in LXXI.7.

Growth: B2B story = consistent cross-age operators under one playbook (cheaper than expert marketplace, more accountable than peer programs). Consumer story = “someone who recently walked this path — with a diagnosis, not vibes.”

Vision: Identity transformation needs honest witnesses (LXI). Near-peers can be witnesses without colonizing as gurus — if grain and playbook stay honest for thirty years of hire scale (ties id 75 workforce drift).

LXXI.12 Confidence table

Claim	Label	Confidence
School tutoring programs help tutees (and often tutors) on achievement/attitudes (Cohen et al. 1982)	FACT	High
Peer tutoring positive on achievement; structure/role moderators matter (Leung 2015)	FACT	High
HE peer tutoring has evidenced formats; cross-year among stronger lines (Topping 1996)	FACT	Medium–High (review, heterogeneous)
Human tutoring $\sim d=0.79$; ITS $\sim d=0.76$; not Bloom 2.0 (VanLehn 2011)	FACT	High
Novice tutors monitor understanding poorly (Chi, Siler & Jeong 2004)	FACT	High
Peer tutors show knowledge-telling bias; knowledge-building rarer (Roscoe & Chi 2007/2008)	FACT	High
Trained Map-briefed near-peers beat expert-no-Map and peer-no-structure on FEI/transfer	HYPOTHESIS	Medium
SAFE-TUTORGRAIN is right commercial mapping	FOUNDER BELIEF	Medium

LXXI.13 What this chapter kills

1. **Kill:** Ivy / PhD / “expert tutor” résumé as quality proof or default SKU.

2. **Kill:** Bloom 2-sigma human-tutor marketing (reinforce standing ban via VanLehn).
3. **Kill:** Unstructured peer Discord / classmate help $\equiv$ MindCraft tutoring.
4. **Kill:** “Great explainers” hire bar (knowledge-telling bias).
5. **Kill:** Expertise $\equiv$ accurate diagnosis (Chi monitoring wound; Map required).
6. **Kill:** Tutor-learning / “tutors grow too” as student North Star.
7. **Wound:** Importing K–8 peer-meta effect sizes as ACT score guarantees.
8. **Wound:** Reading VanLehn ITS $\approx$ human as license for AI-only brand without FEI.
9. **Survive:** Cohen/Leung/Topping structure-eligible peer/near-peer; VanLehn anti-2 σ ; Chi monitoring + Roscoe/Chi telling bias; GRAIN-1...5; trained near-peer default + expert-as-escalate.

Doctrine until data: MindCraft hires **trainable near-peer operators**, not peer chaos and not expert cosplay. Expertise grain is a product decision with falsifiers — résumé theater is not a research conclusion.

Part LXXII — Generated Question Validity & Key Risk

Chapter status: Living evidence + bank/ops brief — Researcher tick 2026-08-03

Primary question: When may MindCraft *generate* practice/diagnostic items (LLM or template), and what validity / answer-key failure modes make generation a product hazard rather than a coverage moat?

Owners: Engine (generation / verify) · Product (Practice / diagnostic bank) · Trust & Safety · Positioning vs “AI made infinite questions” competitors

Commercial job: Ship a **SAFE-GENQ** doctrine: generation is a *coverage accelerator under a verify loop*, never a substitute for keyed truth; bad keys poison the mastery graph and identity evidence; sell *bank trust*, not item count.

LXXII.1 Why this chapter exists

Parts XXXIII / LX / LXVII already kill fluent wrongness as tutoring, apology-as-repair, and chat-as-diagnosis. The remaining bank-specific hazard is quieter: **MindCraft’s own generation path can invent stems that look ACT-ready while shipping wrong keys**. Internal verify already dropped ~30% of a batch (104 kept / 45 dropped) — too high to scale into the live bank. Diagnostic hide-correctness (C4) and /record-outcomes make a wrong key *especially* dangerous: the product records “failure” against a false ground truth and rewrites Maya’s Map.

FOUNDER BELIEF under audit: Ontology-joined generation + deterministic key verify can fill uncovered concepts/formats faster than human authoring — *if and only if* shipping gates treat key validity as a hard fail, not a soft QA note.

Claims we refuse as doctrine:

1. LLM-generated items $\equiv$ ready-to-ship practice without independent key verify.
2. Item-count / coverage % as North Star or marketing hero.
3. Unverified keys as diagnostic or mastery ground truth.
4. “AI wrote a thousand questions” as trust or ACT readiness.

5. Isomorphic clones / surface variants $\equiv$ far transfer or exam readiness.
6. Fluency of stem / explanation $\equiv$ arithmetic or keyed correctness.
7. Verify-off generation to “move fast” on identity or FEI claims.

LXXII.2 Constructs (bank language)

Construct	Research meaning	MindCraft analogue	Failure mode if misused
Automatic item generation (AIG)	Model/template-driven item production with cognitive + psychometric intent	ml/generation/essence $\rightarrow$ LLM $\rightarrow$ format-tagged Question	Free-form LLM dump without item model
Item model / template	Spec of which features may vary	Layer-3 joins + essence prompts; C5 Question schema	Prompt-only “models” with no constrained slots
Key / scoring inference	Map observed response $\rightarrow$ scored outcome	correctAnswer / options; bank-key verify (AIT-6)	Wrong key $\rightarrow$ poisoned events
Isomorph / clone	Surface variants intended as psychometrically similar	Format/surface rewrites of a parent	Clone flood $\equiv$ learning
Verification / review loop	Content+logic check before operational use	--verify pass; drop list; human spot-check	Soft “looks fine” gate
Interpretation/use argument (IUA)	Kane: validate <i>uses</i> of scores, not “the test”	Gap-scan / practice outcomes $\rightarrow$ Map / FEI claims	Ship items $\rightarrow$ claim identity without key evidence
SAFE-GENQ	Generation under verify doctrine	See LXXII.9	Coverage theater

Operational definition (HYPOTHESIS): A generated item is *shippable* only when (a) it matches C5 schema and ontology conceptId/format, (b) the keyed answer survives an independent check (symbolic/arithmetic and/or SME), (c) distractors are not merely fluent nonsense, (d) it is barred from hide-correctness diagnostic and mastery writes until (b) passes, and (e) batch drop rate is below an explicit gate (target $\sim 10\%$ before scale; current $\sim 30\%$ = **do not scale**).

LXXII.3 What classical AIG actually earned

FACT (item models as AIG spine): Gierl & Lai (2012, *International Journal of Testing*, 12(3), 273–298, doi:10.1080/15305058.2011.635830) — AIG is a two-step practice: specialists build *item models* (templates specifying manipulable features), then algorithms instantiate many items. Benefits include volume and potential parameter prediction — not “LLM fluency.”

FACT (review process for generated items): Gierl & Lai (2016, *Educational Measurement: Issues and Practice*, 35(4), 6–20, doi:10.1111/emip.12129) — generated items still need a structured review focused on the *content and logic in the generation procedure*, not vibes after the fact. Illustration domains include mathematics (CCSS-linked) and surgical education.

FACT (quality vs traditional items — distractor wound): Gierl, Lai & Turner (2012, *Medical Education*, 46(8), 757–765, doi:10.1111/j.1365-2923.2012.04289.x) and follow-on quality comparisons (e.g. Gierl et al., *Medical Education*, 2013/quality indicators line) — AIG items can approach traditional quality on many indicators; **distractors** are a recurring weak spot (implausible options more often than hand-authored items). Experts cannot reliably “eyeball which is AIG.”

FACT (feasibility / Kane framing in medical AIG): Reviews of AIG in medical assessment (e.g. Falcão et al., 2022, *Advances in Health Sciences Education*, doi:10.1007/s10459-022-10092-z) map validity work onto Kane-style inferences (scoring → generalization → extrapolation → implication). Scoring/construction clarity and psychometric spread can support use; **implication** (high-stakes decisions) is not free.

Commercial implication: MindCraft may borrow *strong-theory AIG discipline* (constrained models + review of the generator, not only the stem). It may **not** borrow “we generated thousands, therefore bank is valid.”

LXXII.4 Validity is about use — Kane binds the bank

FACT (argument-based validity): Kane (2013, *Journal of Educational Measurement*, 50(1), 1–73, doi:10.1111/jedm.12000) — what is validated are *interpretations and uses* of scores, not the instrument as a fetish object. More ambitious claims need more evidence. Negative consequences can kill a *use* even if a thinner interpretation survives.

Applied to MindCraft (HYPOTHESIS):

Use	Ambition	Key-risk if wrong
Low-stakes practice with reveal + repair	Low–medium	Trust ding; recoverable via SAFE-REPAIR
Hide-correctness diagnostic → /seed-assessment / outcomes	High	Silent Map corruption; false weakness
Parent “diagnosis report” / FEI evidence	High	WTP and trust burn (SAFE-WTP / SAFE-REPAIR)
ACT-readiness or identity claims	L3+ (banned from L1 item A/Bs)	Claim-ladder violation (XXXIV)

Kill: “Items exist in the bank” ≡ validated for diagnostic or identity use.

Survive: Separate *coverage use* (practice under reveal) from *evidence use* (diagnostic / mastery writes) until key gates pass.

LXXII.5 LLM generation is not classical AIG (key failure rates)

Classical AIG manipulates constrained slots inside an item model. LLM generation samples fluent text that *looks* like an item. That difference matters for keys.

FACT (ungrounded GPT math error rates): Bastani et al. (2025, *PNAS*, doi:10.1073/pnas.2422633122) — GPT Base correct only ~51% when repeatedly asked for answers on practice problems; logical errors ~42%, arithmetic ~8%; problem-specific heterogeneity. Students often fail to detect or refuse to check. Unguarded use ↑ concurrent practice then ↓ solo exam.

FACT (classroom ChatGPT wrongness): Primary-school robot/ChatGPT study (IEEE AIRC 2024, doi:10.1109/airc61399.2024.10672220) — LLM inaccurate on ~13/29 curriculum maths prompts (~55% accuracy); students' ability to spot errors tracked *their own* correctness — prior knowledge gates detection.

FACT / REVIEW (LLM math learning concerns): Findings-EMNLP analyses (e.g. “Three Questions...” 2023, doi:10.18653/v1/2023.findings-emnlp.201) — LLMs can mis-grade human work, invent wrong intermediates, and need tool/augmentation discipline; fluent rationale ≠ valid math.

FACT (internal engineering audit): MindCraft generation --verify pass retained 104 / dropped 45 (~30% bad-key rate). Lab standing order already: drop rate must fall before --tested --formats all scale and before syncGeneratedQuestions into the live bank.

Kill: Prompt-scale generation without verify as “coverage solution.”

Wound: Treating Bastani's *tutor* wrongness as identical to *item-authoring* wrongness — related mechanism (fluent false math), different product surface (bank key vs live chat).

Survive: AIT-6 bank-key verify as load-bearing; generation inherits PWC — LLM is bookend, deterministic check is spine.

LXXII.6 Clones, coverage, and the fake-mastery path

FACT (isomorph intent vs clone pejorative): AIG literature distinguishes isomorphs (psychometrically similar variants) from free variants (Gierl & Lai ITEMS module / Gierl & Haladyna AIG line; Bejar item-modeling tradition). Specialists often disparage shallow 1-layer “clones” as detectable and pedagogically thin.

HYPOTHESIS (MindCraft): Format-tagged generation that only swaps surface nouns while freezing the same arithmetic skeleton inflates questionCount without buying SAFE-TRANSFER / interleaving discrimination. Clone floods can also teach answer patterns rather than strategy selection.

Ties: SAFE-TRANSFER (story≠hop; blocked≠ready), SAFE-MISCON (distractors must be diagnostic), SAFE-ONTOLOGY (coverage fills holes; coverage ≠ product), SAFE-CALIB (hide-correctness only on keyed truth).

Kill: “1,500 questions” / “AI-generated bank” as the marketing hero metric.

Survive: Publish *verify rate*, *key-fail taxonomy*, and *concepts still unverified* more loudly than raw count when talking to serious buyers.

LXXII.7 Failure taxonomy (ops)

Failure class	Example	Student harm	Graph harm	Gate
Arithmetic key fail	Stem OK; keyed option wrong	Learns false fact; trust burn	False miss/hit events	Hard fail
Logical key fail	Wrong method encoded as correct	Misconception installed	Bridge/concept mis-update	Hard fail
Ambiguous stem	Multiple defensibly correct	Frustration / “gotcha” affect	Noisy outcomes	Soft fail / rewrite
Implausible distractors	Options no Maya would choose	Easy gaming; shallow practice	Overconfident mastery	Soft fail
Format lie	Tagged diagram but text-only / unreadable	Format-axis diagnosis false	Format-gap severity wrong	Hard fail for format claims
Ontology misjoin	Wrong conceptId / ingredient	Wrong path / cards	Poisoned recommendations	Hard fail
Clone-only batch	N surface twins	Illusion of variety	Inflated exposure counts	Soft fail / cap

FOUNDER BELIEF: Arithmetic/logical key fails are *product-stopping* for diagnostic and mastery writes; ambiguous/implausible are *batch-quality* issues that still block “science-backed bank” copy.

LXXII.8 Positioning

Competitor pattern	What they optimize	MindCraft GENQ counter
“Infinite AI questions”	Volume / novelty	Verify gate; publish drop reasons
Unguarded ChatGPT homework	Fluency / speed	Bank keys + Map; no live key from chat
Static textbook banks	Stability	Generate only into holes; human/ACT seeds remain gold
Clone drill apps	Throughput	Cap isomorphs; measure delayed mix not count

Marketing that survives: “Every practice key is checked before it can rewrite your Map.” / “We generate to fill holes — we ship only what survives verify.” / “Coverage without corrupted diagnosis.”

Marketing that dies: “AI-generated question bank” as hero; “unlimited unique ACT questions”; “our AI never gets the answer wrong”; using generation volume to imply exam readiness or identity change.

LXXII.9 SAFE-GENQ stack (surviving doctrine)

1. **Verify-before-ship** — no live-bank sync without independent key check.
 2. **Key fails are hard fails** — arithmetic/logical wrongness blocks diagnostic + /record-outcomes paths.
 3. **IUA honesty (Kane)** — bank existence $\neq$ validated diagnostic/identity use.
 4. **LLM $\neq$ classical AIG** — constrain with essence/ontology joins; never prompt-and-pray.
 5. **Drop-rate gate** — do not scale while verify drop ■ ~10%; current ~30% = pause scale.
 6. **Coverage $\neq$ North Star** — fill ontology holes; measure transfer / FEI, not questionCount.
 7. **Clone discipline** — isomorph caps; format tags must be true.
 8. **Repair continuity** — if a shipped key is later found wrong: SAFE-REPAIR + event quarantine playbook.
 9. **Ties:** SAFE-REPAIR, SAFE-ONTOLOGY, SAFE-CALIB (C4), SAFE-TRANSFER, SAFE-MISCON, PWC/AIT-6, SAFE-LABMETA (no cite-wash of “AI bank”).
-

LXXII.10 Experiments (GENQ family)

ID	Contrast	Primary endpoints	Kill / survive rule
GENQ-1	Verify-on vs verify-off generated items in <i>revealed</i> practice	Key-error rate spotted by students; retry_120s; trust items	Verify-off $\approx$ verify-on on trust $\rightarrow$ verify theater
GENQ-2	Verified generated vs human/ACT-seed items (same concept/format)	Delayed mix / solo_transfer_pass	Generated ■ seed $\rightarrow$ generation only for holes, not replacement
GENQ-3	Diagnostic eligibility: verified-only vs mixed bank under C4 hide-correctness	Post-scan Map agreement with tutor gold; false-weakness rate	Mixed bank $\uparrow$ false weakness $\rightarrow$ keep diagnostic sealed
GENQ-4	Prompt-hardened arithmetic vs baseline generator	Verify drop rate; residual fail taxonomy	Drop still $\geq 20\%$ after harden $\rightarrow$ architecture change (tool-calc / template-first)
GENQ-5	Parent CBC: “verified bank” attribute vs “AI-generated unlimited Qs”	Choice share; WTP	Unlimited-AI attribute wins $\rightarrow$ message work, not ship unverified
GENQ-QUAL	10 tutor reviews of dropped vs kept items	Codebook: fail classes; SME disagreement rate	Feeds GENQ-4 prompt/tool design

Ties: AIT-6, REPAIR-, ONTO-, CAL-, TR-, WTP-, LABMETA-.

LXXII.11 Commercial implications

Copy: Lead with *checked keys* and *honest coverage*. Never lead with generation volume. Never imply generated $\equiv$ ACT-official. Never claim diagnostic science on unverified items.

Product: Keep generated JSON inert until drop rate clears gate; seal C4 diagnostic + mastery writes to verified keys only; instrument key-fail reports from tutors; quarantine path when a live key is revoked.

Growth: District/parent trust packet (ties id 73) should show verify rates and refusal to ship bad keys — same posture as SAFE-PRIVACY honesty. Competitors selling infinite AI drills are the foil: *we refuse to poison the graph*.

Vision: Thirty-year identity product needs a bank Maya can *believe*. A single wrong key in a hide-correctness diagnostic teaches the company the wrong student. Generation remains a tool for holes under ontology joins — never the adjudicator of truth.

LXXII.12 Confidence table

Claim	Label	Confidence
AIG rests on item models + algorithmic instantiation (Gierl & Lai 2012)	FACT	High
Generated items require procedure-focused review (Gierl & Lai 2016 EMIP)	FACT	High
AIG distractors often weaker than traditional (Gierl/Lai/Turner medical AIG line)	FACT	Medium–High
Validity concerns interpretations/uses; ambitious uses need more evidence (Kane 2013)	FACT	High
Unguarded GPT Base ~51% correct on Bastani practice answer queries; harms solo exam	FACT	High
ChatGPT/robot maths answers ~55% accurate in cited primary classroom study	FACT	Medium (single study)
MindCraft verify drop ~30% on audited batch	FACT (internal audit)	High for that batch; medium as ongoing rate
Prompt hardening + tool-calc can push drop $\backslash < 10\%$ without abandoning LLM stems	HYPOTHESIS	Medium
SAFE-GENQ is the right commercial mapping	FOUNDER BELIEF	Medium–High

LXXII.13 What this chapter kills

1. **Kill:** LLM-generated items $\equiv$ ready-to-ship bank without independent key verify.
2. **Kill:** Item-count / “AI wrote N questions” as hero metric or ACT/identity proof.
3. **Kill:** Unverified keys as diagnostic (C4) or /record-outcomes ground truth.
4. **Kill:** Fluency of stem/explanation $\equiv$ keyed correctness.
5. **Kill:** Clone/isomorph floods $\equiv$ transfer or exam readiness.
6. **Kill:** Scaling generation while verify drop remains $\sim 30\%$.
7. **Wound:** Equating classical constrained AIG validity evidence with unconstrained LLM item dumps.
8. **Wound:** Using Bastani chat-tutor error rates as a precise forecast of *authoring* error without batch audits.
9. **Survive:** Gierl/Lai item-model discipline; Kane IUA for bank *uses*; Bastani + classroom LLM wrongness as hazard priors; internal verify audit; GENQ-1...5; verify-before-ship + diagnostic seal.

Doctrine until data: MindCraft generates to **fill ontology holes**, ships only **verified keys**, and treats key failure as a **graph integrity** incident — not a content backlog inconvenience.

Part LXXIII — District Procurement & Privacy Review as GTM

Chapter status: Living evidence + GTM/ops brief — Researcher tick 2026-08-03

Primary question: What does a U.S. K–12 *district* actually require before MindCraft can sell into schools — and how should that privacy/procurement gate reshape packaging, copy, and roadmap?

Owners: GTM / partnerships · Legal/trust · Product (data inventory) · Parent B2C (contrast) · Engine (affect / PII surfaces)

Commercial job: Ship a **SAFE-PROCURE** doctrine: district revenue is gated by a *trust packet* (DPA + data map + deletion/breach + no-biometric / no-ad use), not by Bloom theater or “FERPA compliant” badges; privacy review *is* the sales process.

LXXIII.1 Why this chapter exists

Parts XXVII / LIX sell parents. Part LXIX constrains affect telemetry. Part LX kills “never hallucinates.” Part LXXII seals the bank. None of that closes a **district** deal. District buyers run a different machine: legal review, data-privacy agreements, security questionnaires, sometimes state-specific overlays — often *before* a pilot classroom ever opens.

FOUNDER BELIEF under audit: MindCraft’s identity / FEI story only reaches school budgets if procurement friction is treated as a *product surface* (shippable packet + NDPA readiness), not a lawyer afterthought once a champion teacher clicks “I agree.”

Claims we refuse as doctrine:

1. “FERPA compliant” / school-official badge $\equiv$ district-ready or ethical clearance.

2. Teacher click-wrap ToS as a substitute for LEA-controlled contracts.
3. Biometric / face-voice “empathy” as a school differentiator.
4. Free teacher tools bypass the same privacy scrutiny as paid tools.
5. Marketing-first RFP chase before a data inventory + DPA pack exists.
6. Parent B2C WTP copy pasted into district RFPs as the trust story.
7. Anxiety Score™ / sell-affect / ad targeting of pupil records (SAFE-PRIVACY collision).

LXXIII.2 Constructs (procurement language)

Construct	Research / policy meaning	MindCraft analogue	Failure mode if misused
LEA / district buyer	Local educational agency that owns procurement + policy	School/district SKU path	Selling only to a teacher’s card
DPA / NDPA	Data privacy agreement; SDPC National Data Privacy Agreement template	Signed DPA + Exhibit A data map	One-off legal chaos per district
School official (FERPA)	Exception allowing disclosure under direct control for legitimate educational interest	Contractual school-official posture	Exception ≡ ethics / marketing stamp
Click-wrap ToS	Teacher-accepted provider terms	Banned as primary school path	Binds LEA to hostile terms
Trust packet	Pre-assembled answers buyers ask	DPA + inventory + retention/deletion + breach + subprocessors + ERT ban	PDF theater without product truth
State overlay	SOPPA, CA EDC 49073.1, etc.	Exhibit G / state terms readiness	Assume federal floor is enough
SAFE-PROCURE	Doctrine: privacy review = GTM gate	See LXXIII.9	Compliance cosplay

Operational definition (HYPOTHESIS): MindCraft is *district-sellable* when (a) a standard NDPA-compatible DPA is ready, (b) Exhibit A lists every student-data category actually collected (including check-in affect and mastery events), (c) purpose limitation + no targeted advertising + deletion-on-exit are product-true, (d) biometric ERT is contractually and productively banned, and (e) free classroom pilots use the same approval path as paid — not a ToS side door.

LXXIII.3 Federal floor: FERPA / PPRA / COPPA bind the *school*, not your badge

FACT (ED PTAC — online services guidance): U.S. Department of Education Privacy Technical Assistance Center (PTAC), *Protecting Student Privacy While Using Online Educational Services: Requirements and Best Practices* (Feb 2014) — when schools use third-party online educational services as part of school activity, FERPA (and often PPRA/COPPA) still apply; best practice is written agreements covering collection, use, retention, disclosure, destruction, security, breach, and access — not informal app installs.

FACT (ED PTAC — Model Terms of Service): PTAC, *Protecting Student Privacy While Using Online Educational Services: Model Terms of Service* (Mar 2016; ERIC ED585306) — click-wrap TOS that a teacher accepts can govern what a provider may collect, use, and share; those terms may conflict with FERPA/PPRA/best practice. PTAC supplies a checklist of “GOOD!” provisions (purpose limitation, security, modification notice/consent, breach response, etc.) for evaluating provider agreements.

FACT (CoSN toolkit — partnering / click-wrap risk): Consortium for School Networking (CoSN), *Protecting Privacy in Connected Learning* toolkit Part 2 (*Partnering with Service Providers*, updated ~2023) — flowchart for FERPA/PPRA/COPPA decisions; warns that teacher click-wrap can bind the school system to terms that misalign with district policy and law; ED recommends free online educational services go through the *same or similar* approval process as paid services; FTC best practice: COPPA appropriateness decisions at district/admin level, not delegated to individual teachers.

Commercial implication: District GTM is blocked or delayed by *agreement shape*, not by whether MindCraft’s landing page says “privacy-first.” A champion teacher who click-wraps MindCraft can create legal risk for the LEA — and create a *bad* sales motion for us (unsigned, non-compliant, later ripped out).

Kill: “Teachers love us / free classroom → privacy doesn’t apply.”

Survive: Same approval rigor for free pilots as for paid SKUs.

LXXIII.4 Standardization: SDPC / NDPA as the industry contracting rail

FACT (SDPC National DPA): Student Data Privacy Consortium (SDPC / Access 4 Learning Community) publishes the **National Data Privacy Agreement (NDPA)** (v2 announced Apr 2024; subsequent point releases) to set common expectations between LEAs and providers and reduce one-off contracts across ~13,000 U.S. districts. NDPA explicitly frames Provider as **School Official** with legitimate educational interest under FERPA when LEA policy permits, under LEA direct control; DPA terms take precedence over conflicting ToS/privacy policies for Student Data treatment; cites FERPA (20 U.S.C. § 1232g / 34 C.F.R. Part 99), PPRA (20 U.S.C. § 1232h), COPPA (15 U.S.C. §§ 6501–6506).

FACT (registry scale — SDPC claims): SDPC states that since 2016, hundreds of thousands of standard DPAs have been executed/subscribed (public materials cite 275,000+ and a large Resource Registry of signed agreements). Treat the *existence of a shared rail* as FACT for GTM planning; treat exact dollar-savings marketing math as vendor/community estimate (not MindCraft-validated).

HYPOTHESIS (MindCraft): Refusing NDPA-compatible contracting while chasing district RFPs maximizes legal cost and cycle time. The moat is not “we invented a unique privacy PDF” — it is product truth that maps cleanly into Exhibit A (data categories) and Exhibit B (schedule of data) without lying about affect check-ins, LLM subprocessors, or tutoring session artifacts.

Kill: Custom 40-page privacy essay as the only offer.

Survive: NDPA-ready pack + honest Exhibit A as table stakes; state exhibits as overlays.

LXXIII.5 State overlays are not optional footnotes

FACT (California contract mandates): California Education Code § 49073.1 (AB 1584) — LEA contracts for digital storage of pupil records and/or digital educational software that accesses/stores/uses pupil records must include specified clauses, including: LEA ownership/control of pupil records; pupil-generated content retention options; prohibition on use beyond contract purposes; parent/eligible-pupil review procedures; security/confidentiality actions (training); unauthorized-disclosure notification procedures; certification of non-retention after contract end (with enforcement description); joint FERPA compliance description; prohibition on using PII in pupil records for targeted advertising. Contracts missing required terms are vulnerable to cure/void dynamics under the statute’s structure.

FACT (Illinois SOPPA): Illinois Student Online Personal Protection Act (105 ILCS 85/) — regulates covered information collected by operators for K–12 school purposes; requires written agreements covering data practices between schools and operators; defines covered information broadly (educational record fields, biometric information, evaluations, identifiers, voice recordings, geolocation, etc.); operator prohibitions include targeted advertising based on acquired information, amassing profiles except for K–12 school purposes, and selling/renting student information (with acquisition exceptions). Illinois districts commonly execute DPAs via SDPC rails under SOPPA.

SPECULATION (biometric tightening): Proposed Illinois school-code bills (e.g. SB 1239 line) seek to ban district acquisition of biometric systems for students. Direction of travel alone makes face/voice emotion AI a district GTM poison pill — consistent with SAFE-PRIVACY science kills.

Commercial implication: A “national” privacy page that ignores CA 49073.1 ad bans or SOPPA covered-info breadth fails the first serious CA/IL review. Build the packet for the *hard* states; lighter states ride along.

LXXIII.6 FERPA badge ≠ ethics, ≠ sales readiness (Zeide wound)

FACT (FERPA’s institutional delegation): Zeide (2016, *Drexel Law Review*, 8, 339–480; SSRN:2821837) — FERPA’s access-control model delegates most privacy decisions to institutions via broad exceptions (notably school-official); transparency/oversight/accountability are thin; notice-and-consent fits poorly under compulsory schooling.

FACT (purpose-limitation theater): Zeide (2017, *University of Miami Law Review*, 71(2), 494–533) — “education purpose” limits are often procedural (authorized by schools), not substantive student-interest protections.

FACT (contextual integrity): Nissenbaum (2010, *Privacy in Context*, Stanford University Press; also 2004, *Washington Law Review*, 79, 119) — privacy = appropriate information flows for context norms, not merely “stayed inside FERPA.”

Ties to LXIX: Even a signed DPA does not make Anxiety Score™, emotion-trait tutor cards, or ad cohorts appropriate. Procurement clears *permission*; doctrine still constrains *use*.

Kill: Homepage “FERPA compliant” as the district hero claim.

Survive: “We sign NDPA-class DPAs; LEA owns pupil records; no targeted ads; no biometrics; deletion on exit; here’s the data map.”

LXXIII.7 What buyers actually ask (trust packet contents)

HYPOTHESIS (MindCraft trust packet = GTM SKU): Ship one downloadable + SE-backed artifact covering:

- 1. **DPA / NDPA stance** — standard agreement + state exhibits.
- 2. **Data inventory (Exhibit A)** — account PII; practice/diagnostic events; Map state; session artifacts; optional affect check-in; subprocessors; explicit non-collection (face/voice biometrics, continuous mood).
- 3. **Purpose limitation** — pedagogy / mastery / FEI ops — not ads or enrichment.
- 4. **Retention / deletion** — TTLs (affect); deletion-on-exit certification.
- 5. **Security + breach** — access control, encryption-in-transit, ML auth, notify timeline.
- 6. **AI honesty** — SAFE-REPAIR + SAFE-GENQ; no “AI never wrong.”
- 7. **Rights path** — LEA/parent review via district intermediary (PTAC preference).
- 8. **Pilot ≠ click-wrap** — written approval for free trials (CoSN/ED).

FOUNDER BELIEF: The packet is more load-bearing for district ARR than another hero video. Parent WTP (SAFE-WTP) is a parallel motion — do not conflate CBC tests with LEA legal review.

LXXIII.8 Positioning vs competitors

Competitor pattern	MindCraft PROCURE counter
Free teacher viral + click-wrap	Same approval as paid; DPA-first pilots
“FERPA compliant” badge	NDPA + Exhibit A + no-ad / no-bio clauses
Emotion-AI engagement	Contractual + product ERT ban (SAFE-PRIVACY)
Infinite AI tutor / bank	Repair + verified keys in trust packet

Marketing that survives: “Built for district DPAs.” / “LEA owns the records.” / “No biometrics. No targeted ads on pupil data.” / “Affect check-ins optional, short-lived, never sold.”

Marketing that dies: “FERPA compliant AI tutor”; “teachers just click to start”; “cameras see struggle”; “unlimited free classroom without IT.”

LXXIII.9 SAFE-PROCURE stack (surviving doctrine)

- 1. **Privacy review = GTM gate** — no district revenue theater without a trust packet.
- 2. **NDPA-ready contracting** — prefer standard DPA rails over one-off essays.
- 3. **Click-wrap is not the school path** — written LEA-controlled agreements; free ≠ unvetted.

4. **Exhibit A honesty** — inventory includes mastery events, session artifacts, optional affect; no hidden LLM training on covered student content without contractual allowance.
5. **No targeted advertising on pupil records** — product + contract (CA 49073.1 posture nationally).
6. **Biometric / ERT ban** — procurement poison + science kill (LXIX).
7. **FERPA badge ≠ ethics** — Zeide/Nissenbaum; compliance floor under SAFE-PRIVACY.
8. **Separate motions** — parent B2C WTP ≠ district legal review; both needed.
9. **Ties:** SAFE-PRIVACY, SAFE-REPAIR, SAFE-GENQ, SAFE-HITL, SAFE-WTP, SAFE-LABMETA (no “science-backed” without falsifiers in RFP claims).

LXXIII.10 Experiments (PROCURE family)

ID	Contrast	Primary endpoints	Kill / survive rule
PROCURE-1	Trust-packet-first outreach vs feature-demo-first to district IT/privacy	Time-to-security-questionnaire complete; DPA signature rate	Demo-first wins cycle time <i>and</i> DPA rate → packet overbuilt
PROCURE-2	NDPA-standard offer vs custom long-form DPA only	Legal cycle days; counsel revision count	Custom always faster → revisit NDPA hypothesis
PROCURE-3	Free pilot requiring LEA approval vs teacher click-wrap	Pilot survival at 90d; forced-removal rate; champion NPS	Click-wrap ↑ short adoption but ↑ removals → keep approval gate
PROCURE-4	Parent B2C message vs district trust-packet message on same landing	Qualified district lead rate vs parent trial rate	One message wins both → still keep packets separate operationally
PROCURE-5	Explicit “no biometric / no pupil-data ads” attribute in district CBC-style choice	Choice share among admins	Attribute flat → still keep as risk insurance, not hero
PROCURE-QUAL	10 interviews: district privacy officers / tech directors	Codebook: blockers, Exhibit A surprises, AI red lines	Feeds packet v1 contents

Ties: PRIV-, WTP-, REPAIR-, GENQ-, HITL-.*.

LXXIII.11 Commercial implications

Copy: Lead district pages with *contractable trust* (DPA, ownership, no ads, no biometrics, deletion). Never lead with Bloom 2-sigma, “AI that never errs,” or emotion cameras. Keep parent score/WTP language off the procurement packet.

Product: Maintain a living data inventory tied to real collections; short-TTL affect; refuse ERT features that would fail SOPPA/biometric reviews; document subprocessors; make deletion/export paths operable, not aspirational.

Growth: Treat trust-packet completion as a pipeline stage (like SOC questionnaire). Instrument days-in-legal. Prefer SDPC/NDPA participation over inventing a boutique privacy brand.

Vision: Thirty-year identity company that cannot clear a district privacy desk does not get the students who most need structured help. Procurement hygiene is not anti-vision — it is how the vision enters schools without becoming surveillance edtech.

LXXIII.12 Confidence table

Claim	Label	Confidence
PTAC 2014/2016: written agreements + Model ToS checklist; click-wrap risk	FACT	High
CoSN toolkit: free apps need similar approval; teacher click-wrap risk; COPPA admin-level decisions	FACT	High
SDPC NDPA exists as national DPA rail; school-official framing; DPA > conflicting ToS for Student Data	FACT	High
SDPC scale/savings marketing numbers	FACT (org claim) / treat cautiously	Medium
CA EDC § 49073.1 mandatory contract clauses incl. no targeted ads	FACT	High
IL SOPPA regulates operators + written agreements; broad covered information	FACT	High
Zeide 2016/2017 FERPA limits; Nissenbaum contextual integrity	FACT	High
Pending IL biometric school bans will pass everywhere soon	SPECULATION	Low–Medium
Trust-packet-first shortens district cycle vs demo-first	HYPOTHESIS	Medium
SAFE-PROCURE is the right commercial mapping	FOUNDER BELIEF	Medium–High

LXXIII.13 What this chapter kills

1. **Kill:** “FERPA compliant” badge $\equiv$ district-ready or ethical clearance.
2. **Kill:** Teacher click-wrap / free viral classroom as the primary school GTM path.
3. **Kill:** Biometric / face-voice empathy features as school differentiators.
4. **Kill:** Marketing-first RFP chase without DPA + Exhibit A inventory.
5. **Kill:** Parent B2C WTP copy as the district trust story.
6. **Kill:** Targeted advertising / sell-affect on pupil records.
7. **Wound:** Treating NDPA signature alone as proof of contextual integrity (still need SAFE-PRIVACY use limits).
8. **Wound:** Assuming federal floor covers CA/IL overlays without Exhibit G work.
9. **Survive:** PTAC Model ToS discipline; CoSN free=paid approval; SDPC/NDPA rail; CA 49073.1 / SOPPA overlays; Zeide/Nissenbaum wound on badge ethics; PROCURE-1...5; trust packet as GTM gate.

Doctrine until data: MindCraft sells into districts by **shipping contractable truth** — NDPA-ready DPA, honest data map, no biometrics, no pupil-data ads, deletion on exit — and treats privacy review as the **front door**, not the cleanup crew.

Part LXXIV — Spaced Retrieval Schedules in Product UX

Chapter status: Living evidence + product/UX brief — Researcher tick 2026-08-04

Primary question: How should MindCraft schedule *returns* to concepts (lags, expanding vs equal spacing, exam-calendar anchors) so competence evidence persists — without shipping streak theater or “optimal brain algorithm” marketing?

Owners: Product · Engine (path / practice) · Growth/copy · Red Team

Commercial job: Ship a **SAFE-SCHED** doctrine: schedule retrieval to the *retention horizon* (homework days vs ACT weeks), prefer delayed first retrieval + equal-ish lags over expanding-SRS cosplay, and sell durable returns — never “our AI found your perfect interval.”

Builds on: Part XXIX (MoC / spacing substrate), XXXIX (interleaving), XLI (desirable difficulties $\times$ anxiety), XLVII (timing/sleep), LXIII (exam dual rail).

LXXIV.1 Why a second spacing chapter

Part XXIX established spacing and retrieval as *identity substrate* (Memory of Competence). This tick answers the *product* question XXIX left open: **what schedule ships in UX** when Maya has an ACT date, a homework week, and a mastery graph that currently clears concepts more than it *returns* them?

FOUNDER BELIEF under audit: MindCraft can out-position Khan/Duo/Brilliant/ChatGPT on durable competence if the practice surface schedules scientifically honest returns — not if we only mass missions until a local mastery mark.

Claims we refuse as doctrine:

1. Expanding-SRS / “optimal interval AI” as hero marketing.
2. Streaks / daily-open as the spacing engine (SAFE-HABIT collision).
3. Cram week $\equiv$ ACT-ready (SAFE-EXAM / blocked-accuracy kill).
4. Overlearning tonight (40 more of the same) as retention strategy.
5. Restudy / tip re-read as “review mode” without generation.
6. Shuffle / random mix $\equiv$ spaced retrieval science.
7. Guaranteed ACT points from any spacing schedule.

LXXIV.2 Constructs (schedule language)

Construct	Research meaning	MindCraft analogue	Failure mode if misused
ISI / lag	Gap between study/practice episodes	Days between concept returns	One-night massing sold as mastery
Retention interval (RI)	Delay to criterial test	Homework due / ACT date / 28d MoC check	Optimize for tonight’s score only
Massed vs spaced	Same total practice, different distribution	Mission tonight vs split across sessions	“Finish the pad” as North Star
Expanding schedule	Growing ISIs (e.g. 1-5-9)	Naive SRS / Duo-like ease	Short-term glow; long-term weaker than equal
Equal / fixed lag	Similar ISIs (e.g. 5-5-5)	Calendar returns matched to RI	Feels “less smart” than expanding UI
Retrieval practice	Generate from memory, not restudy	Soft-wrong reconstruct; spaced quiz	Hint dump / worked-example re-watch
SAFE-SCHED	Doctrine: horizon-matched returns	See LXXIV.9	Algorithm cosplay

Operational definition (HYPOTHESIS): After transfer_pass (or a solid concept exposure), MindCraft schedules at least one **retrieval return** at an ISI scaled to the student’s next use horizon ($\approx$ homework days; $\approx$ 10–20% of RI for multi-month ACT prep per Cepeda ridgeline intuition), with the *first* return delayed enough to be effortful — not an immediate re-quiz while the worked example is still in working memory.

LXXIV.3 Spacing is real; optimal ISI scales with how long you need it

FACT (meta-synthesis): Cepeda, Pashler, Vul, Wixted & Rohrer (2006, *Psychological Bulletin*, 132(3), 354–380, doi:10.1037/0033-2909.132.3.354) — 839 assessments / 317 experiments / 184 articles on distributed practice in verbal recall; spacing and lag effects are robust; the ISI that maximizes final-test retention **increases as the retention interval increases**.

FACT (temporal ridgeline): Cepeda, Vul, Rohrer, Wixted & Pashler (2008, *Psychological Science*, 19(11), 1095–1102, doi:10.1111/j.1467-9280.2008.02209.x) — >1,350 learners; gaps up to ~3.5 months; final tests up to ~1 year; at a given RI, performance rises then falls as ISI grows; optimal gap *increases* with RI, but as a *proportion* of RI the optimal gap *shrinks* (roughly ~20–40% of a 1-week RI → ~5–10% of a 1-year RI). Educational practices that ignore this interaction are inefficient.

FACT (math massing vs distributing): Rohrer & Taylor (2006, *Applied Cognitive Psychology*, 20(9), 1209–1224, doi:10.1002/acp.1266) — college students; distributing 10 math problems across two sessions (1-week gap) vs massing 10 in one session: spacing benefit **nil at 1-week test, large at 4-week test**; extra same-session problems (overlearning) did **not** help 1- or 4-week retention. Textbook massing + overlearning is the wrong default format.

Commercial implication: Parent/student copy that celebrates “cleared the topic tonight” without a scheduled return is selling a 1-week illusion. Product must show *return appointments* (Map / Practice path), not only first-clear fireworks.

Kill: Overlearn-tonight / finish-the-pad as retention strategy.

Survive: Same problem budget, split across sessions matched to RI.

LXXIV.4 Expanding retrieval is not the long-horizon winner

FACT (tradition): Landauer & Bjork (1978) proposed expanding retrieval as a superior long-term schedule — still the folk intuition behind many SRS UIs.

FACT (lab reverse at delay): Karpicke & Roediger (2007, *Journal of Experimental Psychology: Learning, Memory, and Cognition*, 33(4), 704–719, doi:10.1037/0278-7393.33.4.704) — vocabulary pairs; expanding (e.g. 1-5-9) beat equal spacing at a **10-min** final test (replicating Landauer & Bjork short delays); **equal spacing (e.g. 5-5-5) beat expanding at 2 days**, with or without feedback. Experiment 3: delaying the *first* retrieval improved long-term retention; expanding the later intervals added little once first-test delay was controlled.

FACT (review nuance): Roediger & Karpicke line (incl. 2010 *Memory & Cognition* text materials; 2011 spaced-retrieval review) — expanding advantages are mixed; longer delays often favor equal spacing or null; confounds of first-test position are load-bearing.

HYPOTHESIS (MindCraft): Shipping an “expanding SRS brain” as brand differentiates on *algorithm mystique*, not on durable math competence. For ACT horizons (weeks–months), prefer **horizon-matched equal-ish lags** + delayed first return after acquisition; use short expanding only for *stabilizing* a fragile new skill in-session or next-day — never as the hero story.

Kill: “Our AI found your perfect expanding interval.”

Survive: Calendar honesty: “Come back in N days because that gap matches your exam/use horizon.”

LXXIV.5 Retrieval vs restudy — and the math caveat

FACT (testing effect — prose / education): Roediger & Karpicke (2006a, *Psychological Science*, 17(3), 249–255, doi:10.1111/j.1467-9280.2006.01693.x) — repeated testing beats repeated studying on delayed retention; restudy can

inflate immediate confidence without durable gain. Roediger & Karpicke (2006b, *Perspectives on Psychological Science*, 1(3), 181–210, doi:10.1111/j.1745-6916.2006.00012.x) — selective review of testing as learning.

FACT (classroom math spacing × amount): Lyle, Bego, Hopkins, Hieb & Ralston (2020, *Educational Psychology Review*, 32, 277–295, doi:10.1007/s10648-019-09489-x) — engineering precalculus; within-subjects amount × spacing of quiz retrieval; **within-semester** retention benefited from more practice *and* more spacing; **across-semester** (~4 weeks into calculus) benefited **exclusively from increased spacing**. Spacing beats piling more same-window quizzes for transfer-to-next-course durability.

FACT / contradicting nuance (math meta 2025): Murray, Horner & Göbel (2025, *Educational Psychology Review*, 37, article 75, doi:10.1007/s10648-025-10035-1) — spaced vs massed in math: overall $g \approx 0.28$ (27 studies / 53 ES; robust small–medium); isolated-material subset larger ($g \approx 0.43$) than course-embedded ($g \approx 0.24$). Testing vs restudy in math: $g \approx 0.18$ across 7 studies / 32 ES, but **95% CI crossed zero** — math retrieval-vs-restudy evidence is thinner/less conclusive than spacing-vs-massing. Effects may be **smaller than verbal-domain folklore**.

Red Team wound: Do not advertise “science-backed testing effect in math” as if g were large and settled. Spacing-vs-massing is the stronger commercial claim; retrieval-vs-passive-review remains **HYPOTHESIS-strength for math procedures** pending MindCraft FEI data — still prefer generation over tip-restudy as *design default* (aligned with XXIX / XL), but measure it.

Kill: “Retrieval practice always $+0.8\sigma$ in math — trust us.”

Survive: Spaced returns beat cram nights; instrument generation vs restudy inside SCHED experiments.

LXXIV.6 ACT calendar vs cram week (product dual rail)

FACT (from Cepeda ridgeline + Rohrer 4-week): When RI is exam-scale (weeks–months), massed “ACT bootcamp week” optimizes the wrong end of the lag×RI surface. Rohrer’s 1-week null / 4-week large gap is the cautionary tale for last-week cramming sold as readiness.

Ties to LXIII (SAFE-EXAM): Destake-learn / dose-prove. Spacing belongs on the *learn* rail; timed mixed sets belong on the *prove* rail near the exam — not day-one flood.

HYPOTHESIS — Exam Horizon Scheduler:

1. Onboarding / gap-scan captures (or estimates) exam date → RI.
2. Engine sets default ISI $\approx f(\text{RI})$ using Cepeda-style proportions (conservative band; A/B later) for concepts on the ACT path.
3. After transfer_pass, enqueue a **Return** card on Practice/Map (not a streak nag).
4. Fail return → classify forgot vs never-had (XXIX.6); never punish with Calm Score™ / grit meter.
5. Final 7–14 days: increase *mixed* / *timed prove* density; do **not** abandon spaced review of weak bridges.

SPECULATION: Parents will pay for “shows up again before the test” if the dashboard shows *scheduled returns completed* and *28d retention probes* — more than for “day streak.” Validate via WTP CBC (LIX), not founder vibes.

LXXIV.7 Anxiety and desirable difficulty (dose the lag)

FACT / tradition: Spacing and delayed retrieval are desirable difficulties (Bjork line) — they raise short-term failure rates.

Ties to XLI (SAFE-DD) / LVIII (SAFE-EXPOSE): Equip → destake → dose. Do not drop a cold 3-week lag on a high-anxiety novice’s first success.

HYPOTHESIS — Sequenced Scheduling:

1. Acquisition block: short lags OK (same session / next day) until `retry_120s` healthy.
2. Then stretch ISI toward RI-matched equal spacing.
3. Soften targets under high affective stress (LXIX modifiers) — shorter lag, not zero challenge forever.
4. Never frame lag failure as character (“you broke your streak”).

LXXIV.8 Competitive positioning (what not to out-copy)

Competitor pattern	Mechanism	MindCraft response
Duo-like expanding SRS + streak	Habit + short-lag expanding	Equal-ish horizon lags; streak ≠ learning (XXI / XLIII)
Khan mastery clear	Local accuracy	Return appointments + MoC 7/28d
Brilliant delight / puzzle binge	Engagement massing	Delight <i>inside</i> spaced sessions, not instead of them
ChatGPT cram chat	Fluency tonight	Retrieval generation + scheduled re-ask of decisive step
ACT cram packs	Massed week	Horizon scheduler + prove rail late

Competitive wedge (FOUNDER BELIEF): Sell “your Map remembers to bring it back” — inspectable returns — vs “our model optimized your brain.”

LXXIV.9 Doctrine — SAFE-SCHED (provisional)

1. **Horizon-match ISI to RI** (homework vs ACT); do not use one universal “daily” as science.
2. **Delay first retrieval** after acquisition; prefer equal-ish multi-day lags for long RI over expanding-SRS hero UI.
3. **Generation > restudy** in Review/Return modes (measure; don’t overclaim math testing-effect size).
4. **Overlearning tonight ≠ retention**; redistribute problem budget across sessions.
5. **Returns are product surface** (Practice/Map cards), not push-nag streaks.
6. **Exam dual rail**: spaced learn early; timed mixed prove late — never day-one flood.
7. **Anxiety-aware lag stretch** (SAFE-DD); lag failure → diagnose forgot/never-had.
8. **Copy ban**: no “perfect interval AI,” no streak-as-spacing, no guaranteed ACT points from schedule.

Confidence: High that spacing-vs-massing is real (incl. math). Medium that equal > expanding at ACT lags (lab verbal evidence strong; MindCraft math FEI pending). Medium-low that math testing-vs-restudy effect size is large (Murray et al. 2025 wound). Medium that horizon-scheduler UX beats streak UX on retention *and* WTP.

LXXIV.10 Experiments

ID	Question	Design	Primary
SCHED-1	Horizon-matched equal lag vs massed clear-after-mastery	A/B after transfer_pass	retention @ 7d / 28d; solo_transfer_pass
SCHED-2	Expanding UI schedule vs equal-lag returns (same total retrievals)	A/B	14d / 28d retention; first-return success
SCHED-3	Delayed first return (≈1–3d) vs immediate same-session re-quiz	A/B	7d retention; retry_120s
SCHED-4	Generation Return (decisive step) vs restudy tip/example	A/B	retention; transfer_pass
SCHED-5	ACT-date scheduler on vs off (path concepts)	A/B / cohort	mixed delayed accuracy; exam prove set
SCHED-QUAL	10 Maya + 10 parent interviews: return-card language vs streak language	Qual	message resonance; WTP attributes

Falsifier: Challenge-seeking / streak / session minutes rise while 28-day MoC collapses → schedule theater; demote SAFE-SCHED claims and kill streak-as-spacing copy.

Pre-register: SCHED-* before shipping “SRS / smart review” marketing (SAFE-LABMETA).

LXXIV.11 So what for MindCraft commercially

- **Copy:** “We bring topics back before they fade” / “returns matched to your test date” — not “AI-optimized spaced repetition.”
- **Product:** Return queue on Practice + Map; exam-date field drives ISI; Review mode must force generation.
- **Metric:** return_due, return_completed, retention@7/28 — co-primary with FEI; streak demoted.
- **Kill list:** expanding-SRS brand, cram≡ready, overlearn-tonight, restudy-as-review, shuffle≡spacing.
- **Growth:** Parent dashboard shows scheduled returns + retention probes (feeds LIX CBC).
- **Vision:** Identity needs autobiographical *durable* evidence; scheduling is how the product keeps the evidence from evaporating between sessions.

References (verified)

- Cepeda, N. J., Pashler, H., Vul, E., Wixted, J. T., & Rohrer, D. (2006). Distributed practice in verbal recall tasks: A review and quantitative synthesis. *Psychological Bulletin*, 132(3), 354–380. <https://doi.org/10.1037/0033-2909.132.3.354>
- Cepeda, N. J., Vul, E., Rohrer, D., Wixted, J. T., & Pashler, H. (2008). Spacing effects in learning: A temporal ridgeline of optimal retention. *Psychological Science*, 19(11), 1095–1102. <https://doi.org/10.1111/j.1467-9280.2008.02209.x>
- Karpicke, J. D., & Roediger, H. L. (2007). Expanding retrieval practice promotes short-term retention, but equally spaced retrieval enhances long-term retention. *Journal of Experimental Psychology: Learning, Memory, and Cognition*, 33(4), 704–719. <https://doi.org/10.1037/0278-7393.33.4.704>
- Landauer, T. K., & Bjork, R. A. (1978). Optimum rehearsal patterns and name learning. In M. M. Gruneberg, P. E. Morris, & R. N. Sykes (Eds.), *Practical aspects of memory* (pp. 625–632). Academic Press.
- Lyle, K. B., Bego, C. R., Hopkins, R. F., Hieb, J. L., & Ralston, P. A. S. (2020). How the amount and spacing of retrieval practice affect the short- and long-term retention of mathematics knowledge. *Educational Psychology Review*, 32, 277–295. <https://doi.org/10.1007/s10648-019-09489-x>
- Murray, E., Horner, A. J., & Göbel, S. M. (2025). A meta-analytic review of the effectiveness of spacing and retrieval practice for mathematics learning. *Educational Psychology Review*, 37, 75. <https://doi.org/10.1007/s10648-025-10035-1>
- Roediger, H. L., & Karpicke, J. D. (2006a). Test-enhanced learning: Taking memory tests improves long-term retention. *Psychological Science*, 17(3), 249–255. <https://doi.org/10.1111/j.1467-9280.2006.01693.x>
- Roediger, H. L., & Karpicke, J. D. (2006b). The power of testing memory: Basic research and implications for educational practice. *Perspectives on Psychological Science*, 1(3), 181–210. <https://doi.org/10.1111/j.1745-6916.2006.00012.x>
- Rohrer, D., & Taylor, K. (2006). The effects of overlearning and distributed practise on the retention of mathematics knowledge. *Applied Cognitive Psychology*, 20(9), 1209–1224. <https://doi.org/10.1002/acp.1266>

Part LXXV — Tutor Workforce Pipeline & Quality Drift

Chapter status: Living evidence + ops/GTM brief — Researcher tick 2026-08-04

Primary question: How does MindCraft keep *FEI fidelity* across hire → onboard → tenure → churn when college near-peers burn out, relapse into knowledge-telling, or scale faster than coaching capacity?

Owners: Tutor ops · People/hiring · Product (tutor console + QA) · Finance (unit economics) · Positioning vs marketplace “we hired N tutors” vanity

Commercial job: Ship a **SAFE-WORKFORCE** doctrine: the product is not a warm bench of humans — it is a *pipeline that preserves Map-briefed prompt>pour fidelity over tenure*. Headcount, tenure months, and “wellness” theater are not North Stars; fidelity@30/90d + solo transfer under coached tutors are.

LXXV.1 Why this chapter exists

Parts LXVIII (SAFE-HITL) and LXXI (SAFE-TUTORGRAIN) answer *how sessions run* and *who to hire*. Neither answers the scale failure mode every tutoring company hits: **quality drifts after the photo of the hire cohort**. New tutors pass drills, then under parent pressure and fatigue they pour explanations, skip Map briefs, soft-pedal soft-wrongs, and quit — forcing rehire that resets the Kraft scale wound. Competitors sell marketplace volume (“1,000 expert tutors”) or peer Discord availability. MindCraft’s bet dies if pipeline ops treat onboarding as a one-time PDF.

FOUNDER BELIEF under audit: A smaller coached near-peer bench with measurable fidelity over tenure beats a larger uncoached bench on `solo_transfer_pass` and parent-auditable quality — at lower long-run CAC of *quality sessions*, even if raw tutor count looks worse in a pitch deck.

Claims we refuse as doctrine:

1. Tutor headcount / hours-booked as quality or growth NS.
2. Tenure months $\equiv$ pedagogical improvement (seniority theater).
3. One-off onboarding $\equiv$ coaching (Kraft scale wound already in LXVIII).
4. Wellness posters / yoga stipends as substitute for workload + QA design.
5. “We retain tutors” without fidelity instrumentation.
6. Burnout as moral failure of individual tutors (structure > slogans; ties SAFE-STRUCTURE).
7. Marketplace “elite tutor” SKU as pipeline strategy.
8. Bloom 2σ / Ivy hire as attrition insurance.

LXXV.2 Constructs (ops language)

Construct	Research meaning	MindCraft analogue	Failure mode if misused
Pipeline	Hire $\rightarrow$ train $\rightarrow$ deploy $\rightarrow$ coach $\rightarrow$ retain/exit	Funnel with fidelity gates at each stage	Resume pile $\equiv$ pipeline
Quality drift	Practice fidelity decays vs playbook after deploy	Rubric score $\downarrow$ @30/90d without coaching	Blame “bad hires” only
Burnout	Exhaustion, depersonalization, reduced accomplishment	Tutor quit intent + session softness	Wellness theater without load redesign
Dosage / tutor type	Who delivers how often (Nickow framework)	Near-peer session cadence + Map support	Import K–12 school-day ES as ACT guarantee
Coaching at scale	Sustained observation-linked PD	Clip/trace QA cycles (HITL-3)	Efficacy-trial cosplay at marketplace scale
SAFE-WORKFORCE	Doctrine for durable human quality	See LXXV.9	Headcount vanity

Operational definition (HYPOTHESIS): A MindCraft tutor workforce is *pipeline-honest* when (a) hire screens use fidelity drills (SAFE-TUTORGRAIN), (b) deploy requires Map-brief open before first explain, (c) coaching cycles continue past onboarding with rubric samples, (d) burnout/load signals trigger *caseload* or *schedule redesign* not slogans, (e) exit interviews + fidelity trends feed hire-bar revision, and (f) marketing never equates bench size with learning quality.

LXXV.3 Tutoring works — and *who + how often* moderates (not “any human”)

FACT (field-experiment meta): Nickow, Oreopoulos & Quan (2024, *American Educational Research Journal*, 61(1), 74–107, doi:10.3102/00028312231208687) — systematic review/meta of tutoring RCTs; overall pooled effect $\approx$ **0.288 SD**. Effects tend to be largest for programs using **teachers or paraprofessionals**, earlier grades, ≥ 3 **days/week**, and **during school**. Nonprofessional/parent lines are weaker on average in the program typology.

Wound (mandatory): This is PreK–12 *program* evidence, often school-embedded high-dosage models — not remote college→HS ACT prep marketplaces. Borrow the *moderator logic* (tutor accountability, frequency, structure), not the pooled ES as MindCraft ACT insurance. Near-peer college tutors are closer to trained non-teachers than to certified teachers — which is why SAFE-HITL playbooks + Map must carry what credential would otherwise proxy.

Cross-link: VanLehn human $\approx$ 0.79 / ITS $\approx$ 0.76 (LXXI); Graesser naturalistic under-strategy baseline (LXVIII). Pipeline must manufacture structure the meta says matters — it will not appear from “hire nice college kids.”

Commercial implication: Sell *structured, frequent-enough, Map-accountable sessions*, not “access to humans.” Kill pitch slides that lead with tutor count.

LXXV.4 Coaching lifts practice — and collapses when scaled without fidelity

FACT (coaching meta): Kraft, Blazar & Hogan (2018, *Review of Educational Research*, 88(4), 547–588, doi:10.3102/0034654318759268) — 60 causal-design studies; pooled coaching $\approx$ **0.49 SD** on instructional practice and $\approx$ **0.18 SD** on achievement. **Scale wound:** larger effectiveness trials show only a **fraction** of small efficacy-trial effects — taking coaching to scale while keeping quality is the hard problem.

HYPOTHESIS (MindCraft mapping): Quality drift is what the Kraft scale wound *looks like from inside the company*: early cohort gets founder/coach attention (efficacy); hire wave two gets a Loom + checklist (effectiveness collapse). HITL-3 (coached vs onboarding-only) is the product falsifier; WORK-* experiments extend it across tenure.

Kill: “We trained them once.” Kill headcount OKRs that outrun coach capacity (coaches-per-tutor ratio as a hard constraint).

LXXV.5 Burnout predicts quit — and harms the interaction students feel

FACT (attrition meta): Madigan & Kim (2021, *Teaching and Teacher Education*, 105, 103425, doi:10.1016/j.tate.2021.103425) — teacher burnout dimensions relate positively to intentions to quit (exhaustion $r^+ \approx$ **.41**, depersonalization $\approx$ **.32**, reduced accomplishment $\approx$ **.21**); job satisfaction relates negatively ($r^+ \approx$ **-.40**); burnout +

satisfaction together $\approx$ 27% of quit-intent variance, with burnout carrying more relative weight than satisfaction confers protection; relationships appear to have strengthened over time in their moderation checks.

FACT (exhaustion $\rightarrow$ instructional quality $\rightarrow$ students): Klusmann, Aldrup, Roloff, Lüdtke & Hamre (2022, *Journal of Educational Psychology*, 114(6), 1442–1460, doi:10.1037/edu0000703) — large multilevel sample; teachers’ emotional exhaustion associated with lower student self-concept, interest, and (in German) achievement; exhaustion linked to less emotional support and classroom organization; instructional-quality paths partially mediate exhaustion $\rightarrow$ student outcomes.

HYPOTHESIS: Fatigued MindCraft tutors will (a) skip Map briefs, (b) relapse into knowledge-telling (Roscoe/Chi), (c) overuse politeness soft-pedaling (Person et al. 1995), and (d) quit — each is a *fidelity* failure, not only an HR failure. Wellness Slack channels without caseload caps do not fix Klusmann’s mediation path.

Wound: Teacher classroom metas $\neq$ gig college tutors. Use burnout $\rightarrow$ quit and burnout $\rightarrow$ interaction-quality as *mechanisms to instrument*, not K–12 turnover rates as MindCraft forecasts.

Kill: “Tutors just need grit.” Kill Burnout Score™ as student-facing or parent SKU. Kill treating retention % as proof of quality.

LXXV.6 Drift mechanisms inside MindCraft’s model

HYPOTHESIS — five drift channels (ordered by ops leverage):

- 1. **Brief skip** — under time pressure, tutor opens chat not Map (map_brief_open ↓).
- 2. **Pour relapse** — parent “just explain it” + fatigue $\rightarrow$ knowledge-telling (SAFE-TUTORGRAIN).
- 3. **False-mastery probes** — “got it?” returns as session closer (Person et al. 1994).
- 4. **Soft-wrong avoidance** — politeness tax rises when exhausted (rapport over diagnosis).
- 5. **QA starvation** — hire rate > coach sample rate $\rightarrow$ Kraft effectiveness collapse.

SPECULATION: College tutors’ academic calendars create *batch churn* (semesters) — pipeline must assume partial cohort replacement yearly; design for re-certify drills, not permanent tenure mythology.

FOUNDER BELIEF: Instrumenting drift channels early is cheaper than brand repair after parents notice “tutor #4 this month who just lectures.”

LXXV.7 Pipeline stages (product-facing)

HYPOTHESIS — minimum viable workforce pipeline:

Stage	Gate	Fail action
Source	Near-peer grain; content floor	Reject expert-SKU theater as default
Screen	Fidelity drills (LXXI.7)	School brand never primary

Onboard	Shadow + scored Map-brief session	No deploy on deck-only
Deploy	map_brief_open required in console	Soft-block explain-first
Coach	Biweekly sample @30/90d rubric	Caseload pause if fidelity↓
Load	Cap concurrent FEI-heavy students	Redesign schedule before pep talk
Exit	Fidelity trend + reason codes	Feed hire/coach revision

Ties: SAFE-HITL playbook; SAFE-EXPECTANCY (no trait briefs even when tired); SAFE-REPAIR escalate packets must not become infinite unpaid emotional labor.

LXXV.8 Positioning

Competitor pattern	What they optimize	MindCraft WORKFORCE counter
Marketplace volume	Tutor count / credentials	Fidelity@tenure + Map audit trail
Peer Discord help	Availability	Structured pipeline ≠ chat bench
AI-only	Cost / 24-7	Humans kept only if WORK-* fidelity holds
“Ivy tutors” brand	Status	Hire-by-drill; coach-of-tutors internal

Marketing that survives: “Every tutor is coached against the same Map your practice uses.” / “We measure whether tutors still make students think at day 90 — not how many we hired.” / “Smaller bench, higher fidelity.”

Marketing that dies: “1,000 tutors ready now”; tenure-as-quality; wellness-as-pedagogy; retention-% hero metrics; expert marketplace cosplay.

LXXV.9 SAFE-WORKFORCE stack (surviving doctrine)

1. **Fidelity over tenure > headcount** — instrument rubric@30/90d before scaling bench.
2. **Coaching is continuous** — onboarding ≠ PD (Kraft scale wound).
3. **Coach capacity gates hiring** — no hire wave without sample-QA bandwidth.
4. **Burnout** → **redesign load**, not grit posters (Madigan/Klusmann mechanisms).
5. **Drift channels named** — brief skip, pour relapse, got-it?, soft-wrong avoid, QA starve.
6. **Grain stays near-peer default** — expert = escalate/coach (SAFE-TUTORGRAIN).
7. **Retention without fidelity is vanity** — keep rate subordinate to FEI QA.

8. **No Burnout Score™ / wellness theater as product.**

9. **Ties:** SAFE-HITL, TUTORGRAIN, EXPECTANCY, REPAIR, STRUCTURE, WTP (auditable quality attribute).

LXXV.10 Experiments (WORK family)

ID	Contrast	Primary endpoints	Kill / survive rule
WORK-1	Continuous coaching vs onboarding-only (extend HITL-3 to 90d)	Rubric fidelity@30/90d; solo_transfer_pass	No fidelity lift → coaching theater
WORK-2	Hire-rate capped by coach capacity vs uncapped surge hire	Fidelity variance; escalate appropriateness	Surge ≥ capped on fidelity → capacity gate wrong
WORK-3	Caseload cap + schedule redesign vs wellness-module add-on	Quit intent proxy; brief-skip rate; pour rate	Wellness ≥ cap on drift → load redesign unjustified
WORK-4	Console hard-require Map-brief vs optional brief	map_brief_open; soft-wrong rate; talk ratio	Optional ≈ required → UX nag enough
WORK-5	Parent CBC: “coached fidelity report” vs “more tutors / Ivy” attributes	Choice share; WTP	Volume/Ivy dominates → packaging work (not auto-abandon fidelity)
WORK-QUAL	10 tutor + 10 parent interviews on drift/churn trust	Codebook: what “quality held” means	Feeds copy + WORK-5

Ties: HITL-, GRAIN-, EXP-, WTP-, REPAIR-.*.

LXXV.11 Commercial implications

Copy: Lead with *coached operators whose quality is checked over time*. Never lead with bench size. Never imply seniority months = better FEI.

Product: Tutor console enforces Map-brief; ops dashboard shows fidelity@tenure and coach utilization; pause hiring when QA bandwidth saturated. Instrument WORK-1...4 before marketplace expansion narrative.

Growth: B2B story = consistent quality under turnover (schools hate random substitutes). Consumer story = “your tutor is still working from your Map at week 12 — we check.” Unit economics: count *fidelity-weighted session capacity*, not raw tutor FTEs.

Vision: Thirty-year identity witnesses (LXI) require a workforce that does not quietly become lecturebots. Pipeline honesty is how human recognition stays falsifiable as the company scales.

LXXV.12 Confidence table

Claim	Label	Confidence
Tutoring RCTs pool ~0.288 SD; stronger for teacher/para, $\geq 3\times$ /week, during school (Nickow et al. 2024)	FACT	High
Coaching ~0.49 SD instruction / ~0.18 SD achievement; scale shrinks effects (Kraft et al. 2018)	FACT	High
Burnout dimensions predict quit intent; burnout outweighs satisfaction as risk (Madigan & Kim 2021)	FACT	High
Exhaustion linked to worse instructional support/organization and student outcomes via quality paths (Klusmann et al. 2022)	FACT	High (classroom context)
MindCraft drift channels (brief skip / pour / got-it / soft-wrong / QA starve) drive FEI loss	HYPOTHESIS	Medium
Coach-capacity-gated hiring + continuous QA beats surge headcount on transfer	HYPOTHESIS	Medium
SAFE-WORKFORCE is right commercial mapping	FOUNDER BELIEF	Medium
Semester batch churn dominates college-tutor attrition	SPECULATION	Low–Medium

LXXV.13 What this chapter kills

1. **Kill:** Tutor headcount / hours-booked / retention-% as North Stars.
2. **Kill:** Tenure months $\equiv$ quality improvement.
3. **Kill:** One-off onboarding $\equiv$ coaching (reinforce Kraft wound).
4. **Kill:** Wellness theater / grit posters as burnout strategy.
5. **Kill:** Burnout Score™ or emotion-wellness SKU for tutors-as-product.
6. **Kill:** Marketplace volume / Ivy bench as pipeline strategy.

7. **Wound:** Importing Nickow pooled ES or K–12 dosage as ACT score guarantees for remote near-peer.
8. **Wound:** Treating Madigan/Klusmann teacher samples as exact quit-rate forecasts for gig tutors.
9. **Survive:** Nickow moderator logic (structure/accountability/frequency); Kraft coaching-with-scale-humility; Madigan burnout→quit; Klusmann exhaustion→interaction quality; WORK-1...5; fidelity-over-tenure pipeline.

Doctrine until data: MindCraft’s human layer is a **coached fidelity pipeline**, not a headcount metric. If quality drifts after hire, the Constitution’s HITL/TUTORGRAIN rows were theater — and parents will notice before the pitch deck does.

Part LXXVI — Forgetting Curves as Product Honesty

Chapter status: Living evidence + product/GTM brief — Researcher tick 2026-08-04

Primary question: How should MindCraft show *decay risk and delayed competence* so parents and students trust retention claims — without shipping mastery fireworks, fake Ebbinghaus branding, or “we cured forgetting” theater?

Owners: Product (Map / Practice / parent dashboard) · Growth/copy · Engine (mastery decay honesty) · Red Team

Commercial job: Ship a **SAFE-FORGET** doctrine: treat forgetting as expected physics of memory, surface retention probes and scheduled returns as the product truth, and refuse copy that equates tonight’s clear with durable identity evidence.

Builds on: Part XXIX (Memory of Competence), L (SAFE-CALIB), LXII (SAFE-LONGID), LXIII (SAFE-EXAM), LXXIV (SAFE-SCHED).

LXXVI.1 Why this chapter exists

SAFE-SCHED answered *how to space returns*. This tick answers the *honesty* problem: edtech UIs celebrate local mastery (“topic cleared,” “100% tonight,” green graph) while human memory predictably fades — and learners’ own forecasts systematically ignore that fade. Competitors sell fireworks. MindCraft’s identity thesis dies if the dashboard lies about time.

FOUNDER BELIEF under audit: Parents will pay more for a product that *admits forgetting and schedules against it* (inspectable 7d/28d probes + returns) than for a product that looks permanently mastered after one good session — if the honesty is framed as competence protection, not as “you’re failing.”

Claims we refuse as doctrine:

1. Mastery fireworks / “cleared forever” after a blocked session.
2. Personal Ebbinghaus curve / “forgetting score” as hero metric or shame meter.
3. Streaks as proof memory held (SAFE-HABIT / SAFE-SCHED collision).
4. Cram-week green boards ≡ ACT-ready (SAFE-EXAM).
5. Student self-rating of “I still know it” without delayed retrieval probe.
6. Guaranteed retention % or ACT points from any decay UI.
7. Fake neuroscience (“we reverse the forgetting curve”).

LXXVI.2 Constructs (honesty language)

Construct	Research meaning	MindCraft analogue	Failure mode if misused
Forgetting curve	Retention declines with delay (savings / recall)	Expected fade between returns	Brand cosplay; shame curve
Retention interval (RI)	Delay to criterial test	Homework / ACT / MoC probe	Optimize only for tonight
Judgment of learning (JOL)	Predicted future recall	Confidence / “got it?” / parent belief	Fluency-as-mastery
Foresight bias	Study-time cues inflate JOLs	Clear-screen glow after hints	Sell glow as durable learning
Learning vs performance	Current fluency $\neq$ later access	Blocked accuracy vs delayed mix	Leaderboard of tonight
SAFE-FORGET	Doctrine: show decay risk honestly	See LXXVI.9	Mastery-fireworks NS

Operational definition (HYPOTHESIS): MindCraft is *forgetting-honest* when (a) concept “cleared” states expire or downgrade without a delayed retrieval return, (b) parent/student surfaces show last-proof age + next return due, (c) JOLs/confidence are never the sole retention claim, (d) copy frames returns as protecting competence (not character), and (e) marketing never implies permanent mastery from one session.

LXXVI.3 Forgetting is real — and not a marketing metaphor

FACT (classic replication): Murre & Dros (2015, *PLoS ONE*, 10(7), e0120644, doi:10.1371/journal.pone.0120644) — successful replication of Ebbinghaus’s savings forgetting curve (1880/1885); one subject, ~70 hours of list learning/relearning across delays from 20 min to 31 days; results similar to the original; curve is robust and not perfectly smooth (likely upward jump near the 24-hour point).

FACT (form of forgetting): Wixted (2004, *Annual Review of Psychology*, 55, 235–269, doi:10.1146/annurev.psych.55.090902.141555) — review of psychology/neuroscience of forgetting; candidate functions (power, logarithmic) share decelerating proportional decay; exponential is a poor description of typical empirical forgetting functions. **Wound:** Do not ship a branded “your personal power-law curve” — the science supports *fade with time*, not a student-facing physics widget.

Cross-link (spacing as intervention): Cepeda et al. (2006) and Rohrer & Taylor (2006) — already SAFE-SCHED: distributing practice beats massing for longer RIs; overlearning tonight does not buy 4-week math retention. Forgetting honesty without scheduled returns is merely a sad dashboard.

Commercial implication: The Map must age evidence. A green node with no delayed probe is a *performance snapshot*, not a competence claim.

Kill: “Once cleared, always mastered.”

Survive: Time-stamped last proof + expected fade until return.

LXXVI.4 Learners (and products) ignore the fade they “know about”

FACT (JOLs ignore retention interval): Koriat, Bjork, Sheffer & Bar (2004, *Journal of Experimental Psychology: General*, 133(4), 643–656, doi:10.1037/0096-3445.133.4.643) — when JOLs are experience-based (fluency at study), item-by-item and aggregate predictions are largely **insensitive to anticipated retention interval**, while actual recall shows steep forgetting; people can access theory-based forgetting knowledge mainly when RI is made salient (within-subject RI manipulation or framing in terms of *forgetting* rather than remembering). Classic dissociation: predicted week-later recall can stay high while actual recall collapses.

FACT (foresight / illusion of competence): Koriat & Bjork (2005, *Journal of Experimental Psychology: Learning, Memory, and Cognition*, 31(2), 187–194, doi:10.1037/0278-7393.31.2.187) — study-time JOLs suffer an inherent study–test mismatch (answer present at study, absent at test); a posteriori associations inflate competence judgments relative to later cued recall.

FACT (remediation path): Koriat & Bjork (2006, *Memory & Cognition*, 34(5), 959–972, doi:10.3758/BF03193244) — foresight bias can be reduced by study–test experience (especially *test* experience) and by delayed JOLs; overconfidence has behavioral costs for study-time allocation.

HYPOTHESIS (MindCraft mapping): Session-end “How confident?” prompts and parent dashboards that show tonight’s percent correct reproduce Koriat’s experience-based JOL trap unless the UI forces theory-based forgetting salience (last-proof age, days-to-exam RI, scheduled return) and delayed retrieval probes.

Kill: Confidence slider / thumbs-up as retention KPI.

Survive: Delayed generation Returns as the honesty engine (ties SAFE-CALIB + SAFE-SCHED).

LXXVI.5 Performance glow is the competitor’s product — and our trap

FACT / tradition (desirable difficulties): Bjork & Bjork (2011) — “Making things hard on yourself, but in a good way,” in *Psychology and the real world*; conditions that raise current performance (massing, blocking, rereading fluency) often undermine long-term retention/transfer; fluency and familiarity create illusions of comprehension; desirable difficulties (spacing, interleaving, testing) feel worse early and work better later.

FACT (math massing wound — reuse): Rohrer & Taylor (2006, *Applied Cognitive Psychology*, 20(9), 1209–1224) — spacing benefit nil at 1-week test, large at 4-week; same-session overlearning did not help delayed retention.

HYPOTHESIS: Khan-style clear / Duo-style streak / Brilliant binge / ChatGPT cram-chat all sell *performance glow*. If MindCraft’s Map fireworks the same glow without aging nodes and MoC probes, we are not differentiated — we are cosplaying competitors with a prettier ontology.

FOUNDER BELIEF: The commercial wedge is *inspectable honesty*: “Last proved 18 days ago — return due before your ACT” beats “Mastered!” on trust once parents understand the stakes (validate with WTP CBC, LIX).

LXXVI.6 Product surface — forgetting-honest UX

HYPOTHESIS — minimum honest surfaces:

Surface	Must show	Must not show as NS
Concept node	Last proof timestamp; proof type (blocked vs delayed mix / transfer)	Permanent green from one session
Practice Return	Due date matched to RI; generation task	Streak nag / shame fade meter
Parent dashboard	Returns completed; retention@7/28 probe rates	“Your child never forgets”
Post-mission	“Held tonight — still need a return”	Confetti ≡ durable mastery
Exam rail	Dual rail: learn spaced / prove timed (LXIII)	Cram board as readiness

SPECULATION: Framing returns as “protecting what you earned” outperforms “you forgot” on anxiety-sensitive Mayas — test in FORGET-QUAL; never soften standards into permanent hint crutches (SAFE-EXPOSE / SAFE-DD).

Engine note (HYPOTHESIS): Temporal decay toward prior in the ML mastery model is aligned in spirit with SAFE-FORGET — but product UI must not expose raw Beta posteriors as a “Forgetting Score™.” Show behavioral probe outcomes, not spooky brain meters.

LXXVI.7 Competitive positioning

Competitor pattern	What they sell	MindCraft FORGET counter
Mastery clear fireworks	Tonight accuracy	Aged evidence + delayed mix
Streak / XP	Habit continuity	Returns ≠ streaks
Expanding SRS mystique	Algorithm brain	Horizon-matched honesty (SAFE-SCHED)
Cram packs	Short-RI performance	RI-matched returns + prove rail
Chat fluency	Feeling of knowing	Generation probe after delay

Marketing that survives: “We bring it back before it fades.” / “Cleared tonight ≠ ready on test day — here’s your return.” / “Parents see last proof age, not just progress bars.”

Marketing that dies: Personal Ebbinghaus brand; cured-forgetting claims; mastery confetti as North Star; retention guarantees; neuroscience theater.

LXXVI.8 Parent WTP and identity clocks

Ties: SAFE-WTP (CBC attributes), SAFE-LONGID (8/26/52w fade checks), SAFE-IDMEASURE (Likert ≠ durable identity).

HYPOTHESIS: A parent CBC attribute “shows when skills last proved + schedules returns” should beat “shows topics mastered this week” and “daily streak” for high-stakes ACT segments — without implying score guarantees.

FOUNDER BELIEF: Identity transformation requires *autobiographical competence that survives delay*. Forgetting honesty is how the product refuses to mint false identity evidence.

Wound: Do not claim parents “only care about honesty” — some buy cram. Offer dual rail; measure choice share; do not moralize buyers who pick cram attributes (still refuse to *market* cram as durable identity).

LXXVI.9 SAFE-FORGET stack (surviving doctrine)

1. **Fade is expected** — design for decay; do not brand-shame it.
 2. **Age the evidence** — last-proof timestamps; clear states expire without delayed return.
 3. **JOLs ≠ retention** — never sell confidence/fluency as held knowledge (Koriat line).
 4. **Delayed probes are the truth surface** — MoC 7d/28d + generation Returns.
 5. **Frame returns as competence protection**, not character failure.
 6. **No Forgetting Score™ / personal Ebbinghaus hero UI.**
 7. **Ties:** SAFE-SCHED (how to schedule), SAFE-CALIB (how to elicit confidence), SAFE-EXAM (dual rail), SAFE-LONGID (fade checks), SAFE-WTP (package honesty).
 8. **Copy ban:** cured forgetting; cleared forever; streak-as-memory; guaranteed retention %.
-

LXXVI.10 Experiments (FORGET family)

ID	Contrast	Primary endpoints	Kill / survive rule
FORGET-1	Aged-node UI (last proof + due return) vs permanent-clear fireworks	retention@7/28; return_completed; false-mastery quit	Fireworks ≥ aged on retention → honesty UX not load-bearing
FORGET-2	Post-mission “held tonight / return needed” vs confetti-only	7d return success; student JOL calibration error	Confetti ≥ honest on calibration → copy work
FORGET-3	Parent dashboard: proof-age + returns vs topics-mastered-this-week	CBC choice share; WTP (with LIX)	Mastered-week dominates → packaging, not auto-abandon honesty

FORGET-4	Theory-framed forgetting salience (days-to-exam RI shown) vs fluency-only confidence prompt	JOL–actual gap at delayed probe	Salience null → need stronger probe design
FORGET-5	Clear-state expiry without return vs sticky green	delayed mix accuracy; solo_transfer_pass	Sticky ≥ expiry → expiry may be too harsh / needs soft downgrade
FORGET-QUAL	10 Maya + 10 parent interviews: “protect what you earned” vs “you forgot” language	Codebook: threat vs agency	Feeds copy; anxiety-aware framing

Ties: SCHED-, CAL-, WTP-, LONG-, EXAM-*.

Falsifier: Honest decay UI raises anxiety / dropout while 28d MoC stays flat → wound framing; do not kill the decay science, redesign threat language (SAFE-DD).

LXXVI.11 Commercial implications

Copy: Lead with *time-honest competence* — returns before fade, proof age visible. Never lead with permanent mastery fireworks or neuroscience curves.

Product: Map ages nodes; Practice Return queue is first-class; parent view shows returns + retention probes; post-mission copy separates tonight’s performance from delayed readiness.

Metric: last_proof_age_days, return_due, return_completed, retention@7/28 co-primary with FEI; demote session accuracy and streak as learning claims.

Growth: Differentiate vs Khan/Duo/Brilliant/ChatGPT on *inspectable durability*, not content breadth or habit theater. District story: evidence that does not silently expire into false readiness.

Vision: Thirty-year identity witnesses need memories that survive delay. Forgetting honesty is how Minecraft refuses to sell a green lie.

LXXVI.12 Confidence table

Claim	Label	Confidence
Ebbinghaus-style savings forgetting replicates (Murre & Dros 2015)	FACT	High

Empirical forgetting often decelerating (power/log-like); not a student UI mandate (Wixted 2004)	FACT	High
Experience-based JOLs largely ignore anticipated RI; theory framing can restore sensitivity (Koriat et al. 2004)	FACT	High
Study-time foresight bias inflates competence; test experience / delayed JOLs help (Koriat & Bjork 2005/2006)	FACT	High
Fluency/massing create performance–learning traps (Bjork & Bjork 2011; Rohrer & Taylor 2006)	FACT	High
Aged-node + return UX beats fireworks on 28d retention and parent WTP	HYPOTHESIS	Medium
“Protect what you earned” framing beats shame-fade for anxious Mayas	HYPOTHESIS	Medium
SAFE-FORGET is right commercial mapping	FOUNDER BELIEF	Medium
Raw engine decay posteriors should stay off student-facing meters	SPECULATION	Medium

LXXVI.13 What this chapter kills

1. **Kill:** Mastery fireworks / cleared-forever as durable competence claims.
2. **Kill:** Personal Ebbinghaus / Forgetting Score™ hero metrics.
3. **Kill:** Streak / tonight-accuracy as memory-held proof.
4. **Kill:** Confidence / “got it?” as retention KPI.
5. **Kill:** Cured-forgetting / reverse-the-curve neuroscience ads.
6. **Kill:** Guaranteed retention % or ACT points from decay UI.
7. **Wound:** Importing nonsense-syllable savings percentages as ACT readiness numbers.
8. **Wound:** Shame-framed fade meters that spike anxiety without lifting MoC.
9. **Survive:** Murre/Dros robustness of forgetting; Koriat JOL–RI dissociation; Bjork learning≠performance; Rohrer delayed math spacing; FORGET-1...5; aged evidence + scheduled returns as honesty stack (with SAFE-SCHED).

Doctrine until data: MindCraft sells **time-honest competence evidence**, not permanent green. If the Map cannot age a claim, the claim is marketing — and the Constitution’s identity thesis is theater.

Part LXXVII — Adaptive Spacing Algorithms vs Fixed Calendars

Chapter status: Living evidence + product/engine brief — Researcher tick 2026-08-04

Primary question: When (if ever) should MindCraft replace horizon-matched *fixed* return calendars with *adaptive* lag models (SM-2 / FSRS-style / half-life regression), and how do we do that without shipping black-box “perfect brain interval” brand cosplay?

Owners: Engine (return scheduler) · Product (Practice/Map Returns) · Growth/copy · Red Team

Commercial job: Ship a **SAFE-ADAPT** doctrine: keep SAFE-SCHED’s inspectable horizon calendar as the *default*; allow adaptive *modulation inside declared bands*; refuse SM-2/FSRS/Duo-HLR as hero marketing or as a substitute for transfer probes.

Builds on: Part XXIX (MoC), LXXIV (SAFE-SCHED), LXXVI (SAFE-FORGET), LXIII (SAFE-EXAM), LXVII (SAFE-ONTOLOGY).

LXXVII.1 Why a third spacing chapter

SAFE-SCHED answered *what schedule shape* (ISI×RI, delayed first return, equal-ish lags). SAFE-FORGET answered *how to show decay honestly*. Competitors and eng blogs answer with a third move: **personalize the lag with an opaque model** — SM-2 ease factors, FSRS stability/difficulty, Duolingo half-life regression — and sell that as the science.

FOUNDER BELIEF under audit: Parents will pay for *inspectable returns matched to an exam/homework horizon* more than for “our AI found your perfect interval”; adaptive models may still *win retention efficiency* if they stay subordinate, auditable, and measured on delayed math transfer — not flashcard MAE.

Claims we refuse as doctrine:

- 1. SM-2 / Anki / FSRS as the MindCraft brand story.
 - 2. Black-box “optimal interval AI” with no parent/student-visible due date rationale.
 - 3. Flashcard recall-prediction gains ≡ ACT / solo_transfer readiness.
 - 4. Engagement / DAU lift from adaptive SRS ≡ learning or identity (Settles wound).
 - 5. Expanding adaptive schedules as default for multi-month RI (SAFE-SCHED kill retained).
 - 6. Streaks / daily open as the adaptive engine.
 - 7. Guaranteed retention % or ACT points from any scheduler.
-

LXXVII.2 Constructs (algorithm language)

Construct	Research / eng meaning	MindCraft analogue	Failure mode if misused
-----------	------------------------	--------------------	-------------------------

Fixed / calendar lag	Same ISI rule for a cohort (or RI band)	Exam-horizon Return queue	One-size misses item difficulty
Adaptive scheduler	Next ISI from learner×item history	Optional modulator on due date	Opaque brain-AI cosplay
SM-2	Heuristic ease-factor expanding schedule (Wo■niak / SuperMemo 2)	Anti-pattern as hero	Ease hell; vocab-shaped UX
FSRS / DSR	Difficulty–Stability–Retrievability models; predict recall prob	Possible <i>engine</i> candidate later	Flashcard MAE sold as math science
Half-life regression (HLR)	Trainable decay half-life (Settles & Meeder)	Lexeme-memory, not FEI	Engagement-as-proof
Personalized review	Select <i>what</i> to review given history (Lindsey et al.)	Priority among due Returns	Random “smart” without lag honesty
SAFE-ADAPT	Doctrine: calendar first; adapt inside bands	See LXXVII.9	Black-box SRS brand

Operational definition (HYPOTHESIS): MindCraft may adapt a concept’s next return ISI only inside a **declared band** derived from RI (e.g. $\pm 30\%$ of the SAFE-SCHED horizon lag), using inspectable inputs (last proof age, fail/pass on generation Returns, bridge-gap severity) — never a silent SM-2/FSRS clone that rewrites due dates without showing *why* and *when*.

LXXVII.3 Adaptive scheduling can beat fixed spacing — in the domains studied

FACT (lab optimization): Pavlik & Anderson (2008, *Journal of Experimental Psychology: Applied*, 14(2), 101–117, doi:10.1037/1076-898X.14.2.101) — ACT-R–based scheduler that chooses practice timing to maximize long-term gain per unit practice time outperformed comparison schedules on recall accuracy and latency for paired-associate / fact learning; optimal intervals *lengthen* as items stabilize.

FACT (classroom personalized review): Lindsey, Shroyer, Pashler & Mozer (2014, *Psychological Science*, 25(3), 639–647, doi:10.1177/0956797613504302) — middle-school foreign-language course; semester-long retrieval software; on a cumulative exam ~1 month after semester end, **personalized** review beat **massed** (~+16.5% retention) and beat **one-size-fits-all spaced** (~+10.0%). Personalization here selected *which* material to review given inferred memory strength — not merely expanding flashcard ease cosplay.

FACT (trainable production SRS): Settles & Meeder (2016, *ACL*, 1848–1858, doi:10.18653/v1/P16-1174) — half-life regression on Duolingo data reduced recall-prediction error >45% vs several baselines; an operational study reported ~12% **higher daily student engagement**. HLR generalizes Leitner/Pimsleur-style ladders with fitted features.

Commercial implication: “Adapt > fixed” is a *real* claim in vocab / flashcard / lexeme settings. MindCraft cannot Red-Team it out of existence. The fight is **domain transfer + brand honesty**, not denial of personalization.

Kill: “All adaptive SRS is pseudoscience.”

Survive: Personalized review can beat generic spacing for item memory — measure whether that transfers to MindCraft’s FEI endpoints.

LXXVII.4 SM-2 and FSRS are engineering lineages — not math-identity science

FACT (SM-2 origin): Woźniak’s SuperMemo SM-2 algorithm (documented 1987; SuperMemo 2 era) sets early intervals (classically 1 day then 6) then multiplies by an ease factor updated from 0–5 quality ratings — a hand-tuned expanding heuristic that powered decades of SRS apps (incl. Anki’s legacy default). It is **not** a peer-reviewed math-education RCT.

FACT (modern flashcard ML): Ye, Su & Cao (2022, *KDD* ’22, 4381–4390, doi:10.1145/3534678.3539081) — MaiMemo-scale memory model + stochastic shortest-path scheduler; reported ~12.6% improvement over SOTA methods on review-cost / memorization efficiency for language learning; deployed in production. Open FSRS builds on related DSR/DHP memory-state ideas and is widely used to **predict flashcard recall probability** and reduce review load vs SM-2 on Anki-style logs (engineering benchmarks; recall calibration ≠ transfer).

HYPOTHESIS (domain mismatch): Math FEI depends on *ingredient/bridge transfer under variation*, not binary “card recalled?” grades. Mapping a concept node to an SM-2 “card” collapses diagnosis (SAFE-ONTOLOGY), soft-wrong science (SAFE-MISCON), and solo transfer into a flashcard grade — the wrong grain.

Wound: Cite FSRS/Anki MAE wins as *engine inspiration* for predicting return urgency; do **not** cite them as evidence MindCraft will raise ACT or solo_transfer_pass.

Kill: “We use FSRS / SM-2 — therefore science-backed ACT prep.”

Survive: Optional recall-probability models as *priority signals* among already-due Returns.

LXXVII.5 Engagement and prediction accuracy are not identity outcomes

FACT / Red Team wound (Settles): The production win highlighted for HLR includes **daily engagement +12%**. Engagement is not FEI. Duo-shaped products can improve revisit rates while still selling streak theater and short-horizon expanding schedules that SAFE-SCHED already rejects for long RI.

Cross-link (SAFE-FORGET): Koriat & Bjork foresight bias — learners (and UIs) confuse fluent recent success with future recall. An adaptive model that *schedules* well can still be wrapped in mastery fireworks that *communicate* poorly.

Cross-link (SAFE-SCHED / Karpicke & Roediger 2007): Expanding schedules that look “smart” can lose to equal spacing at longer delays. Adaptive ≠ expanding; if the learner model quietly expands first intervals for dopamine, it collides with equal-ish long-RI doctrine.

Commercial implication: Never advertise “more daily opens from our scheduler” as learning proof. Instrument retention_probe_7d/28d, solo_transfer_pass, and return_completed — co-primary with any ADAPT A/B.

LXXVII.6 Fixed calendars remain the right *product* default

FACT (reuse — ISI×RI): Cepeda et al. (2006; 2008) — optimal lag scales with retention interval; educational systems that ignore RI are inefficient. A **visible calendar** tied to homework/ACT date is the UX that makes ISI×RI *legible* to parents (SAFE-WTP / SAFE-FORGET proof-age).

HYPOTHESIS — Calendar-first stack:

1. **Layer A (default):** Horizon-matched equal-ish lags from exam/homework RI (SAFE-SCHED Exam Horizon Scheduler).
2. **Layer B (optional adapt):** Within $\pm$ band, pull forward Returns with weak last proof / bridge-gap / failed generation; push back only after strong delayed probe — never past RI honesty.
3. **Layer C (selection):** When many Returns are due, personalized *priority* (Lindsey-style) picks which concept to surface first — still showing due dates.
4. **Hard ban:** Silent SM-2 ease hell; student-facing “easiness 2.1”; FSRS Score™; streak as scheduler.

SPECULATION: Inspectable Layer B will beat black-box Layer-B-as-brand on parent WTP even if flashcard MAE is slightly worse — test via CBC (LIX), not founder vibes.

LXXVII.7 Anxiety, load, and adaptive aggression

Ties to SAFE-DD / SAFE-EXPOSE / SAFE-PRIVACY: Adaptive systems that schedule reviews at low predicted retrievability maximize desirable difficulty — and can spike threat for high-anxiety Maya.

HYPOTHESIS: Cap adaptive “pull forward” under high self-authored stress modifiers (shorten challenge density / soften stakes, not erase Returns). Never use affect to *delete* horizon honesty.

Kill: Anxiety Score™ that auto-masses practice “for comfort.”

Survive: Affect → lag *aggression* caps inside SAFE-SCHED bands.

LXXVII.8 Competitive positioning

Competitor pattern	Mechanism	MindCraft response
Anki / FSRS mystique	Recall-prob optimization	Engine-only; no SRS brand
Duo HLR + streak	Engagement + lexeme half-life	Returns + MoC probes; streak demoted
Khan clear-and-forget	Local mastery	Aged evidence + calendar Returns
“AI perfect interval” edtech ads	Authority cosplay	Inspectable due + why chip
ACT cram packs	Massed week	Horizon calendar + late prove rail

Competitive wedge (FOUNDER BELIEF): Sell “your Map shows the next return and why” — not “smarter than Anki.”

LXXVII.9 Doctrine — SAFE-ADAPT (provisional)

1. **Calendar-first:** SAFE-SCHED horizon lags are the default product surface.
2. **Adapt inside bands:** Personalization may modulate ISI only within declared RI-derived bounds; show due date + reason.
3. **Selection ≠ mystique:** Prefer Lindsey-style *what to review among due* over silent SM-2 *rewrite of science*.
4. **Flashcard lineage ≠ FEI proof:** SM-2 / FSRS / HLR evidence is vocab/item-memory; do not claim ACT/transfer from it.
5. **Engagement ≠ learning:** Ban DAU/streak lifts as scheduler success.
6. **No black-box brand:** No “perfect interval AI,” FSRS Score™, or easiness theater.
7. **Anxiety caps:** Stress modifiers limit adaptive aggression; never erase returns.
8. **Measure FEI:** ADAPT-* wins only if delayed probes / solo transfer hold or rise vs fixed calendar.

Confidence: High that personalized review can beat massed and generic spaced for *item* memory (Lindsey; Pavlik). Medium that MindCraft Layer-B banding helps retention without hurting trust. Low–medium that off-the-shelf FSRS on concept nodes helps `solo_transfer_pass` (domain mismatch). High that SRS-as-brand collides with SAFE-SCHED / SAFE-FORGET / competitive wedge.

LXXVII.10 Experiments

ID	Question	Design	Primary
ADAPT-1	Fixed horizon calendar vs calendar + within-band lag modulation	A/B after <code>transfer_pass</code>	<code>retention@7/28</code> ; <code>solo_transfer_pass</code>
ADAPT-2	Due-set priority (personalized selection) vs FIFO due queue (same lags)	A/B	28d retention; <code>return_completed</code>
ADAPT-3	Inspectable “due + why” UI vs opaque “smart review” queue	A/B + parent CBC	trust / WTP; retention (non-inferiority)
ADAPT-4	FSRS-like urgency score as <i>priority only</i> vs full ISI rewrite	A/B	28d retention; complaint/confusion rate
ADAPT-5	Adaptive pull-forward under high stress: capped vs uncapped	A/B	anxiety state; <code>retry_120s</code> ; retention

ADAPT-QUAL	10 Maya + 10 parent: “smart algorithm” vs “returns on your calendar” copy	Qual	message resonance; SRS-brand allergy
------------	--	------	---

Falsifier: Adaptive arm raises opens/streaks while retention_probe_28d or solo_transfer_pass falls → demote adaptive claims; keep calendar-first.

Pre-register: ADAPT-* before any “AI spaced repetition” marketing (SAFE-LABMETA).

LXXVII.11 So what for MindCraft commercially

- **Copy:** “Returns on your calendar — adjusted when evidence says so” — never “FSRS-powered brain optimization.”
- **Product:** Ship Layer A Returns first; Layer B/C behind flags; every due chip shows horizon basis.
- **Engine:** Decay + proof-age already exist — wire them to *visible* due dates before importing Anki math.
- **Metric:** Co-primary retention/transfer; demote scheduler engagement vanity.
- **Kill list:** SM-2/FSRS hero, black-box perfect-interval AI, engagement≡learning, flashcard MAE≡ACT.
- **Growth:** Parent dashboard = next returns + why — feeds WTP without SRS cosplay.
- **Vision:** Identity needs durable competence evidence; algorithms may *protect* evidence efficiently, but the product promise is honesty and transfer — not scheduler mystique.

References (verified)

- Cepeda, N. J., Pashler, H., Vul, E., Wixted, J. T., & Rohrer, D. (2006). Distributed practice in verbal recall tasks: A review and quantitative synthesis. *Psychological Bulletin*, 132(3), 354–380. <https://doi.org/10.1037/0033-2909.132.3.354>
- Cepeda, N. J., Vul, E., Rohrer, D., Wixted, J. T., & Pashler, H. (2008). Spacing effects in learning: A temporal ridge of optimal retention. *Psychological Science*, 19(11), 1095–1102. <https://doi.org/10.1111/j.1467-9280.2008.02209.x>
- Karpicke, J. D., & Roediger, H. L. (2007). Expanding retrieval practice promotes short-term retention, but equally spaced retrieval enhances long-term retention. *Journal of Experimental Psychology: Learning, Memory, and Cognition*, 33(4), 704–719. <https://doi.org/10.1037/0278-7393.33.4.704>
- Lindsey, R. V., Shroyer, J. D., Pashler, H., & Mozer, M. C. (2014). Improving students’ long-term knowledge retention through personalized review. *Psychological Science*, 25(3), 639–647. <https://doi.org/10.1177/0956797613504302>
- Pavlik, P. I., & Anderson, J. R. (2008). Using a model to compute the optimal schedule of practice. *Journal of Experimental Psychology: Applied*, 14(2), 101–117. <https://doi.org/10.1037/1076-898X.14.2.101>
- Settles, B., & Meeder, B. (2016). A trainable spaced repetition model for language learning. In *Proceedings of the 54th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)* (pp. 1848–1858). ACL. <https://doi.org/10.18653/v1/P16-1174>
- Woźniak, P. A. (1987/archived). *SuperMemo 2: Algorithm* [technical documentation]. <https://www.super-memory.org/archive/english/ol/sm2.htm>

- Ye, J., Su, J., & Cao, Y. (2022). A stochastic shortest path algorithm for optimizing spaced repetition scheduling. In *Proceedings of the 28th ACM SIGKDD Conference on Knowledge Discovery and Data Mining* (pp. 4381–4390). ACM. <https://doi.org/10.1145/3534678.3539081>
-

Part LXXVIII — Cram Products vs Durable-Identity Positioning

Chapter status: Living evidence + GTM/product brief — Researcher tick 2026-08-04

Primary question: How should MindCraft compete with massed “cram pack / bootcamp week / guaranteed points” test-prep without either (a) becoming another cram SKU or (b) moralizing parents who buy short-horizon score relief?

Owners: Growth/copy · Product (Practice / Returns / exam dual rail) · Parent WTP · Red Team

Commercial job: Ship a **SAFE-DURABLE** doctrine: dual-rail GTM (durable learn rail + late prove rail); refuse cram-ready and package ACT guarantees; price and message *returns of competence* while still selling honest timed practice when the calendar demands it.

Builds on: Parts XXVII (markets), XXIX (MoC), LIX (SAFE-WTP), LXIII (SAFE-EXAM), LXXIV (SAFE-SCHED), LXXVI (SAFE-FORGET), LXXVII (SAFE-ADAPT).

LXXVIII.1 Why this chapter exists

The ACT/SAT prep market already sells what parents can feel in a weekend: massed drills, strategy tips, score-lift ads, and package tiers with point targets. MindCraft’s thesis is identity transformation via durable competence evidence. Those two value props collide in the same checkout moment.

FOUNDER BELIEF under audit: Parents will pay a durable-identity premium *if* the product makes delayed competence legible (returns, aged proof, solo transfer) *and* still offers an honest late-stage “prove under stakes” rail — not if we only preach spacing while competitors own the anxiety week.

Claims we refuse as doctrine:

1. Cram-week / bootcamp massing $\equiv$ ACT-ready or “math person.”
 2. Guaranteed ACT/SAT point packages as MindCraft SKUs.
 3. Overlearn-tonight / marathon session as retention science.
 4. Score-only parent messaging as the default brand (SAFE-WTP wound retained).
 5. Shame-parenting (“if you bought Kaplan you failed your kid”).
 6. Pure anti-exam idealism that deletes timed practice (SAFE-EXAM dual rail retained).
 7. Streak / DAU / tonight-accuracy as proof of durable positioning.
-

LXXVIII.2 Constructs (market language)

Construct	Research / market meaning	MindCraft analogue	Failure mode if misused
Massed / cram	Practice compressed into short RI before test	Anti-pattern as <i>learn</i> rail	Fluency illusion; fast fade
Distributed / spaced	Same effort across ISIs matched to RI	Returns / SAFE-SCHED	Sold as moral purity, not product
Overlearning	Extra same-session reps after criterion	Blocked drill after clear	Wastes time vs spacing (Rohrer)
Cram product	Commercial SKU selling massed score relief	Competitor default	Owens anxiety week
Durable-identity offer	Pay for lasting competence + self-story	FEI + Returns + Map	Soft brand with no exam rail
Dual rail	Destake-to-learn / dose-to-prove	SAFE-EXAM	Learn rail only → lost GTM
SAFE-DURABLE	Positioning doctrine (this chapter)	See LXXVIII.9	Cram cosplay or exam denial

Operational definition (HYPOTHESIS): A MindCraft commercial offer is *durable-positioned* iff (1) default Practice schedule is horizon-matched spaced Returns, (2) parent-facing proof includes aged evidence / delayed probes, (3) any exam-week intensity is labeled *prove rail* not *mastery*, and (4) no package promises a point delta.

LXXVIII.3 Spacing beats massing for durable retention — including math

FACT (meta, verbal recall): Cepeda, Pashler, Vul, Wixted & Rohrer (2006, *Psychological Bulletin*, 132(3), 354–380, doi:10.1037/0033-2909.132.3.354) — quantitative synthesis of distributed practice; spacing > massing across retention intervals; lag effects depend on RI (ISI×RI).

FACT (optimal lag): Cepeda, Vul, Rohrer, Wixted & Pashler (2008, *Psychological Science*, 19(11), 1095–1102, doi:10.1111/j.1467-9280.2008.02209.x) — temporal ridgeline: optimal gap scales with desired retention interval (~10–20% of RI as a rough educational heuristic in the spacing literature).

FACT (math practice): Rohrer & Taylor (2006, *Applied Cognitive Psychology*, 20(9), 1209–1224, doi:10.1002/acp.1266) — college students learning a math procedure; distributed practise across sessions beat massing at a **4-week** test (large benefit), with little/no spacing advantage at a 1-week test in Exp. 1; Exp. 2: tripling same-session problems (**overlearning**) did **not** raise 1- or 4-week scores.

FACT (math meta): Murray, Horner & Göbel (2025, *Educational Psychology Review*, 37, article 75, doi:10.1007/s10648-025-10035-1) — spaced vs massed practice in mathematics: robust small-to-medium overall effect (**g** = **0.28**; 27 studies); larger in isolated-learning designs (**g** = 0.43) than course-embedded (**g** = 0.24). Testing-vs-restudy effect in math was small and not robust (**g** = 0.18; CI crossed zero).

FACT (utility rating): Dunlosky, Rawson, Marsh, Nathan & Willingham (2013, *Psychological Science in the Public Interest*, 14(1), 4–58, doi:10.1177/1529100612453266) — distributed practice rated **high utility**; many popular student strategies (rereading, highlighting) rated low.

Commercial implication: Durable retention is the scientifically stronger long-horizon claim. Effect sizes in *math* are real but often modest — do **not** advertise “2× retention” or Bloom-scale miracles from spacing alone.

Kill: Cram ≡ science-backed durable ACT readiness.

Survive: Spacing/distribution as the default *learn* schedule; measure delayed math outcomes, not vibes.

LXXVIII.4 Learners (and buyers) prefer the wrong schedule

FACT (metacognitive illusion): Kornell (2009, *Applied Cognitive Psychology*, 23(9), 1297–1317, doi:10.1002/acp.1537) — GRE-type flashcard study; spacing beat massing and beat last-day cramming; spacing helped ~90% of participants, yet after the first session ~72% **believed massing had been more effective**.

Cross-link (SAFE-FORGET): Koriat & Bjork foresight bias — fluent recent performance feels like lasting knowledge. Cram products sell that feeling on purpose.

HYPOTHESIS (parent market): Parent WTP for cram packs is partly purchase of *felt readiness* and social proof (“we did a program”), not Bayesian belief about 4-week retention. CBC must separate “weekend intensive” vs “returns on a calendar” vs “timed prove week” attributes (SAFE-WTP).

Commercial implication: Do not win by shaming the illusion. Win by making delayed competence *visible* (proof age, return chips, retention probes) so the durable offer has a felt artifact competitors lack.

Kill: “Parents only buy scores, so we must become a cram brand.”

Wound: Score anxiety is real demand — dual-rail, don’t deny.

Survive: Legible MoC > moral lecture.

LXXVIII.5 Test-prep market reality (without becoming it)

FACT (industry shape): College-admissions test prep remains a large, multi-billion-dollar category with formal prep (courses, online, 1:1 tutoring) and informal prep (books, practice tests). Brookings’ synthesis (Reber, Roller & Sanderson, 2025, “Exam ready”) notes paid prep is uneven by income; individual tutoring/coaching is often more associated with gains than lighter formats; effects of commercial prep are generally **modest** relative to aggressive advertising; rigorous evidence for high-priced “guaranteed score” packages is weak/absent.

FACT (ACT retest prep activities): Moore, Sanchez & San Pedro (ACT Research Report R1710, 2018) — among students who retested, hours with a **private tutor/consultant** related to composite retest gains in propensity-adjusted models; many other listed prep activities were not statistically related. Absolute adjusted differences were **small** (sub-point on the ACT composite scale in the reported contrasts) — not “+4 guaranteed.”

SPECULATION / market observation: Cram packs and bootcamps will keep winning anxiety-week SEO and school-hallway word-of-mouth. MindCraft should not spend brand capital trying to out-Kaplan Kaplan on massed volume.

Commercial implication: Position against *mechanism* (massed fluency theater), not against *parents*. Offer: diagnosis + spaced returns + Map-briefed humans + late prove rail. Package what can be instrumented (returns completed, retention probes, solo transfer) — never point guarantees.

LXXVIII.6 Dual-rail GTM (durable + prove)

Reuse (SAFE-EXAM / SAFE-SCHED): Destake-to-learn early; dose-to-prove late; readiness = delayed mixed under stakes — not blocked tonight accuracy.

HYPOTHESIS — SAFE-DURABLE product stack:

- 1. **Learn rail (default SKU story):** Horizon-matched Returns; aged evidence; bridge-gap honesty; FEI session spine.
- 2. **Prove rail (calendar-triggered):** Timed mixed sets, stakes ladder, exam affect handling — labeled as pressure practice, not “you’re mastered.”
- 3. **Cram-intercept UX:** If user requests “only this week,” show tradeoff: massing can lift short-RI performance while risking 28d fade — then offer denser Returns inside the remaining horizon rather than silent overlearn-tonight.
- 4. **Parent surface:** Next returns + last proof age + what “ready” means (delayed mix) — CBC-tested vs score-guarantee competitor copy.

Kill: Exam-week massing sold as the MindCraft method.

Survive: Honest prove rail subordinate to durable learn rail.

LXXVIII.7 Competitive positioning

Competitor pattern	Mechanism sold	MindCraft response
Kaplan / Princeton Review packs	Massed content + strategies + score ads	Dual rail; no point guarantees
Local bootcamp week	Anxiety-week intensity	Cram-intercept + denser Returns
Duo-style daily streak prep	Habit / engagement	Demote streak; Returns NS
Khan clear-and-move-on	Local mastery fireworks	Aged evidence + generation Returns
ChatGPT cram chat	Fluent coverage theater	PWC + solo_transfer; no monologue⇒ready
“AI perfect interval” apps	Scheduler mystique	SAFE-SCHED / SAFE-ADAPT calendar-first

Competitive wedge (FOUNDER BELIEF): Sell “still sharp when it counts — because we returned” — not “more problems this weekend.”

LXXVIII.8 Equity and structure (do not moralize scarcity)

Cross-link (SAFE-STRUCTURE / SAFE-WTP): Higher-income families buy more formal prep (Brookings; ACT ecological briefs). Mocking cram purchases shames constrained families who may only afford a short intensive.

HYPOTHESIS: Durable positioning must include *access-honest* paths (shorter calendars with honest RI matching; free/LEA rails under SAFE-PROCURE) — not “only year-long subscriptions are ethical.”

Kill: “Real parents buy identity programs; cram buyers are unserious.”

Survive: Structure and calendar honesty over character judgment.

LXXVIII.9 Doctrine — SAFE-DURABLE (provisional)

- 1. **Default = durable learn rail:** Horizon-matched spaced Returns; generation probes; Map diagnosis.
- 2. **Exam intensity = prove rail:** Timed/stakes practice labeled honestly; never rebranded as mastery fireworks.
- 3. **No point-guarantee packages:** Ban ACT/SAT +N SKUs and “bootcamp ready in 7 days” as brand claims.
- 4. **Cram-intercept, don’t cram-cosplay:** If massing is requested, show RI tradeoff and densify Returns — do not ship overlearn-tonight as science.
- 5. **Legible MoC for WTP:** Parent value prop = proof age + returns + delayed mix — CBC before score-only default.
- 6. **Modest math effect honesty:** Cite spacing as real (Murray g~0.28 class) — never 2-sigma or cured-forgetting.
- 7. **No shame GTM:** Compete with products, not with parents’ anxiety purchases.
- 8. **Measure FEI + retention, not bootcamp NPS:** CRAM-* wins only if delayed probes / solo transfer hold vs massed-pack messaging arms.

Confidence: High that spacing/distribution beats massing for long-RI retention (Cepeda; Rohrer & Taylor; Murray et al.). High that learners mis-prefer massing (Kornell). Medium that parent CBC will pay a premium for legible returns vs cram when price is matched. Low–medium that dual-rail alone beats entrenched score ads without distribution/partner channels. High that point-guarantee packaging is a kill under SAFE-WTP / SAFE-EXAM / claim ladder.

LXXVIII.10 Experiments

ID	Question	Design	Primary
CRAM-1	Durable-returns parent CBC vs cram-pack / +points CBC attributes	Parent CBC (SAFE-WTP style)	WTP; choose-durable rate

CRAM-2	Cram-intercept UX (tradeoff + denser Returns) vs silent massed week mode	A/B near exam date	retention_probe_28d; solo_transfer_pass; session load
CRAM-3	Dual-rail labeled prove week vs unlabeled “final mastery blitz”	A/B	calibration; anxiety state; delayed mix
CRAM-4	Score-lift ad creative vs returns/proof-age creative (same offer)	Message test	CTR/qualified lead; complaint; misexpectation
CRAM-5	Massed equal-time week vs horizon Returns (same total minutes)	A/B after transfer_pass	retention@7/28; blocked vs mixed
CRAM-QUAL	10 parents: “weekend intensive” vs “calendar returns + prove week” stories	Qual	resonance; shame allergy; price anchors

Falsifier: Durable messaging wins surveys but massed arm wins retention@28d and solo_transfer → revisit schedule fidelity before brand claims.

Falsifier: Cram-intercept raises short-term satisfaction while 28d retention collapses → demote intercept; strengthen defaults.

Pre-register: CRAM-* before any anti-Kaplan attack ads or “we don’t cram” hero (SAFE-LABMETA).

LXXVIII.11 So what for MindCraft commercially

- **Copy:** “Built for what you still know later — plus timed practice when the date is close.” Never “+4 ACT points” or “crush it in a weekend.”
 - **Product:** Keep SAFE-SCHED Returns as default; ship cram-intercept + prove-rail labeling; surface proof age on parent dashboard.
 - **Pricing:** Packages = horizon length + tutor/Map depth — not score deltas.
 - **Metric:** Co-primary retention_probe_28d, solo_transfer_pass, return_completed; demote bootcamp NPS / tonight accuracy.
 - **Kill list:** Cram⇒ready; point guarantees; overlearn-tonight; shame-parent GTM; streak-as-durable.
 - **Growth:** Compete on legible durability in a market that sells fluency theater.
 - **Vision:** Identity needs competence that survives the test date — and a company that can say so without lying about anxiety week.
-

References (verified)

- Cepeda, N. J., Pashler, H., Vul, E., Wixted, J. T., & Rohrer, D. (2006). Distributed practice in verbal recall tasks: A review and quantitative synthesis. *Psychological Bulletin*, 132(3), 354–380. <https://doi.org/10.1037/0033-2909.132.3.354>
- Cepeda, N. J., Vul, E., Rohrer, D., Wixted, J. T., & Pashler, H. (2008). Spacing effects in learning: A temporal ridgeline of optimal retention. *Psychological Science*, 19(11), 1095–1102. <https://doi.org/10.1111/j.1467-9280.2008.02209.x>
- Dunlosky, J., Rawson, K. A., Marsh, E. J., Nathan, M. J., & Willingham, D. T. (2013). Improving students' learning with effective learning techniques: Promising directions from cognitive and educational psychology. *Psychological Science in the Public Interest*, 14(1), 4–58. <https://doi.org/10.1177/1529100612453266>
- Kornell, N. (2009). Optimising learning using flashcards: Spacing is more effective than cramming. *Applied Cognitive Psychology*, 23(9), 1297–1317. <https://doi.org/10.1002/acp.1537>
- Moore, R., Sanchez, E., & San Pedro, M. O. (2018). *Investigating test prep impact on score gains using quasi-experimental propensity score matching* (ACT Research Report R1710). ACT.
- Murray, E., Horner, A. J., & Göbel, S. M. (2025). A meta-analytic review of the effectiveness of spacing and retrieval practice for mathematics learning. *Educational Psychology Review*, 37, 75. <https://doi.org/10.1007/s10648-025-10035-1>
- Rohrer, D., & Taylor, K. (2006). The effects of overlearning and distributed practise on the retention of mathematics knowledge. *Applied Cognitive Psychology*, 20(9), 1209–1224. <https://doi.org/10.1002/acp.1266>
- Reber, S., Roller, Z., & Sanderson, Q. (2025, July 15). Exam ready: Who uses college admissions test prep and does it work? *Brookings* (Center for Economic Security and Opportunity). <https://www.brookings.edu/articles/exam-ready-who-uses-college-admissions-test-prep-and-does-it-work/>

Part LXXIX — Feedback Timing: Immediate vs Delayed

Chapter status: Living evidence + product UX brief — Researcher tick 2026-08-04

Primary question: When should MindCraft reveal correctness, coach on a soft-wrong, or withhold judgment — and what timing claims are scientifically indefensible in marketing?

Owners: Product (Practice / Solver / Diagnostic) · Coach UX · Tutor ops · Brand & copy · Red Team

Commercial job: Ship a **SAFE-FBTIME** doctrine: mode-conditional timing (diagnostic hide / learn micro-delay + soft-wrong / prove-rail KR); sell *task-focused informative feedback*, never “instant AI feedback” or “science says always delay” as the brand.

Builds on: Parts XXVI (CLT/tutoring), XL (SAFE-SE), XLI (SAFE-DD), XLIX (SAFE-MISCON), L (SAFE-CALIB), LI (SAFE-DP), LVIII (SAFE-EXPOSE), LXIII (SAFE-EXAM), LXXVIII (SAFE-DURABLE).

LXXIX.1 Why this chapter exists

Edtech defaults to a green check the instant a student taps an answer. Tutors default to explaining before the student finishes thinking. Competitors market “instant personalized feedback” as the product. MindCraft already needs the opposite in places (hide-correctness diagnostic; attempt-first SE; soft-wrong springboards) — but without a timing law, those become inconsistent vibes.

FOUNDER BELIEF under audit: A short, intentional delay before correctness reveal (or coach pour) raises constructive retry and transfer *for learn-rail practice*, especially after soft-wrongs — while instant KR can be right for novices acquiring a procedure and for labeled prove-rail calibration. Timing is a *mode switch*, not a universal moral.

Claims we refuse as doctrine:

1. “Instant feedback” as the MindCraft hero claim.
2. “Always delay feedback” as desirable-difficulty cosplay.
3. Instant right/wrong $\equiv$ mastery or identity progress.
4. Feedback $\equiv$ learning (ignoring Kluger’s one-third harm rate).
5. Soft-wrong delay as Calm Score™ / therapy theater.
6. Explanation-first pour timed as “support.”
7. Streak / speed-to-green-check as North Star.

LXXIX.2 Constructs

Construct	Research meaning	MindCraft analogue	Failure mode
Immediate feedback	KR/KCR/EF right after a response	Instant check / Solver reveal	Dependency; shallow processing
Delayed feedback	After items, minutes, or sessions	End-of-set review; proof age	Frustration; abandoned retry
Micro-delay	Seconds of forced attempt / SE before reveal	Soft-wrong coach gate	Feels like lag / broken UX
Hide-correctness	No right/wrong reveal in-session	Diagnostic C4	Corrupts Map if keys wrong
Self-focused FI	Praise, ego, comparative rank	“You’re a math person!” toast	Attention off task (FIT)
Task/process FI	How to revise the strategy	Soft-wrong $\rightarrow$ ingredient card	Monologue without generation
SAFE-FBTIME	Mode-conditional timing law	This chapter	Timing silver-bullet ads

Operational definition (HYPOTHESIS): Timing policy is *SAFE-FBTIME-honest* iff (1) diagnostic never reveals correctness in-session, (2) learn-rail soft-wrongs require attempt + brief SE/micro-delay before coach, (3) prove-rail may use immediate KR under explicit stakes label, (4) feedback content is task/process-focused, and (5) no marketing claims a universal optimal delay.

LXXIX.3 Feedback is powerful — and often harmful

FACT (meta): Kluger & DeNisi (1996, *Psychological Bulletin*, 119(2), 254–284, doi:10.1037/0033-2909.119.2.254) — 607 effect sizes; average FI effect $d \approx 0.41$, but **over one-third of interventions decreased performance**. Feedback Intervention Theory (FIT): effectiveness falls as attention moves from task learning toward self / meta-task processes.

FACT (education review): Hattie & Timperley (2007, *Review of Educational Research*, 77(1), 81–112, doi:10.3102/003465430298487) — feedback is among the strongest influences on achievement *when* it answers where am I going / how am I going / where to next; timing is a thorny moderator, not a free lunch. Self-level feedback is typically least useful for learning.

FACT (formative guidelines): Shute (2008, *Review of Educational Research*, 78(1), 153–189, doi:10.3102/0034654307313795) — formative feedback should be nonevaluative, supportive, timely, and specific; timing interacts with learner and task (immediate often better for difficult new material / low prior knowledge; delayed can support transfer and avoid over-reliance).

Commercial implication: Do not sell “more feedback faster.” Sell *task-focused next move*. Instant ego toasts and fluency theater violate FIT.

Kill: Feedback volume or speed as product quality.

Survive: Task/process feedback with mode-aware timing.

LXXIX.4 Timing is not a silver bullet (classic split + modern digital)

FACT (classic meta): Kulik & Kulik (1988, *Review of Educational Research*, 58(1), 79–97, doi:10.3102/00346543058001079) — 53 studies; **applied classroom** quizzes usually favored **immediate** feedback; **lab** / **test-content acquisition** studies often favored **delayed**. Timing effects are context-bound.

FACT (test learning): Butler, Karpicke & Roediger (2007, *Journal of Experimental Psychology: Applied*, 13(4), 273–281, doi:10.1037/1076-898X.13.4.273) — after MC practice on prose passages, **delayed** feedback beat immediate on later cued recall; authors attribute benefit partly to **spaced re-presentation** of information.

FACT (digital feedback meta): Brummer, de Boer, Mouw & Strijbos (2024, *Learning Environments Research*, 27(3), 453–476, doi:10.1007/s10984-024-09501-4) — 116 interventions; overall $g \approx 0.41$ vs controls; both immediate and delayed timing showed strong positive effects vs no-feedback baselines; authors argue **clarity/consistency of the timing regime** matters more than a universal clock; meta-regression favored **feedback focus** (process) over timing as the practical lever.

FACT (emerging CAL timing meta — preprint): Kandemir, Esposito, Gurgand & Ramus (preprint, HAL:hal-05546645; 51 studies / 160 ES, 1988–2024) — average immediate-vs-delayed contrast $g = 0.03$, 95% CI [−0.08, 0.13], *ns*; substantial heterogeneity (level, domain, response-time constraints). Cite as preprint, not journal-of-record; use as *warning against timing absolutism*, not as a kill of all delay effects.

Commercial implication: “Our AI gives instant feedback” and “we delay because science” are both overclaims. Position on **mode + focus + generation**, not milliseconds.

Kill: Universal immediate-or-delayed brand claim.

Survive: Consistent, labeled policies per rail (diagnostic / learn / prove).

LXXIX.5 Math-specific: prior knowledge flips the sign

FACT: Fyfe & Rittle-Johnson (2016a, *Journal of Educational Psychology*, 108(1), 82–97, doi:10.1037/edu0000053) — elementary math; when strategy knowledge was induced, **immediate verification feedback hurt** relative to no feedback; without induced knowledge, feedback helped.

FACT: Fyfe & Rittle-Johnson (2016b, *Journal of Experimental Child Psychology*, 147, 140–151, doi:10.1016/j.jecp.2016.03.009) — after strategy instruction, computer feedback (immediate or summative) beat no feedback for low prior knowledge on next-day posttest; **immediate** was particularly effective for mastery across knowledge levels in that design (correct-answer feedback without explicit right/wrong judgment).

FACT (CBLE synthesis): van der Kleij, Feskens & Eggen (2015, *Review of Educational Research*, 85(4), 475–511, doi:10.3102/0034654314564881) — elaborated feedback (EF) produced larger effects than KR/KCR especially for higher-order outcomes; **delayed timing** associated with smaller effects in their models — again: content richness and outcome grain matter more than a slogan about clocks.

HYPOTHESIS (MindCraft): Gap-scan “hard” / low event-count concepts → lean immediate task-focused KR after attempt; higher-knowledge / soft-wrong discrimination items → micro-delay + SE prompt before reveal; never pour explanation-first (Experiment D / SAFE-SE).

Kill: One timing default for all Map nodes.

Survive: Confidence/prior-knowledge tiered timing (ties SAFE-CALIB).

LXXIX.6 Soft-wrong delay vs green-check addiction

Reuse (SAFE-MISCON / SAFE-SE / SAFE-CALIB): Soft-wrongs are diagnostic springboards; student-generated why beats AI monologue; hide-correctness protects the graph when keys are sacred.

HYPOTHESIS — product stack:

1. **Diagnostic rail:** No correctness reveal; record outcome silently; never sell “instant grade.”
2. **Learn rail:** Submit → 3–8s SE / “what went wrong?” gate → soft-wrong coach card (task/process) → optional Map link.
3. **Prove rail (SAFE-EXAM / SAFE-DURABLE):** Immediate KR allowed; label as pressure practice; no mastery fireworks.
4. **Tutor HITL:** Pause before pour; Map brief; prompt>explain (SAFE-HITL).
5. **Parent surface:** Show *quality of next move* (retry, transfer), not “feedback latency SLA.”

Kill: Green-check dopamine as engagement KPI.

Wound: Some students hate micro-delay — instrument abandon rate; adapt length by prior knowledge, don’t delete the gate.

Survive: Attempt-first + task feedback as brandable difference vs ChatGPT instant essay.

LXXIX.7 Competitive positioning

Pattern	Timing sold	MindCraft response
Duo / game apps	Instant KR + XP	Demote green-check; FEI metrics
Khan hints	Fast reveal / explain	Soft-wrong SE before hint binge
Brilliant	Puzzle delight, often immediate	Keep delight; add transfer delay probes
ChatGPT tutors	Instant fluent critique	PWC + micro-delay generation; no monologue=feedback
Human tutors (untrained)	Explain-first “help”	HITL pause + Map
Cram packs	Instant score theater	SAFE-DURABLE dual rail

Competitive wedge (FOUNDER BELIEF): “We wait just long enough for *your* reason — then we diagnose” beats “instant AI feedback.”

LXXIX.8 Doctrine — SAFE-FBTIME (provisional)

1. **Mode-conditional clocks:** Diagnostic hide; learn micro-delay + SE; prove immediate KR labeled.
2. **Focus > timing:** Task/process feedback; ban self-level ego toasts as learning feedback.
3. **Content > clock:** Prefer informative next-step / soft-wrong contrast over bare KR when the goal is transfer.
4. **Prior-knowledge switch:** Novice acquisition may need faster KR; higher knowledge needs delay/generation.
5. **No timing silver-bullet ads:** Ban “instant feedback AI” and “always delayed = science.”
6. **Kluger honesty:** Instrument harm (abandon, worse delayed mix) — feedback can hurt.
7. **Keys first:** Never delay-or-reveal theater on unverified items (SAFE-GENQ).
8. **Measure FEI not latency NPS:** `retry_120s`, `solo_transfer_pass`, `soft-wrong→drill`; demote `ms-to-green-check`.

Confidence: High that FIs are heterogeneous and can harm (Kluger & DeNisi). High that timing effects are context-dependent (Kulik & Kulik; Brummer et al.; Kandemir preprint null average). High that delayed feedback can aid retention via spacing in test settings (Butler et al.). Medium that MindCraft micro-delay+SE lifts transfer vs instant KR on soft-wrongs (needs FB-*). Medium-high that prior knowledge moderates sign (Fyfe line). High that instant-feedback brand claims are a kill.

LXXIX.9 Experiments

ID	Question	Design	Primary
----	----------	--------	---------

FB-1	Learn-rail micro-delay+SE vs instant KR after soft-wrong	A/B within concept	retry_120s; solo_transfer_pass; abandon
FB-2	Prior-knowledge tiered timing vs one global delay	Stratify by gap-scan / eventCount	Delayed mix; mastery Δ ; frustration
FB-3	Diagnostic hide vs immediate KR diagnostic	A/B (C4 integrity)	Post-scan Map agreement; false weakness
FB-4	Prove-rail immediate KR vs delayed end-of-set KR	A/B near exam	Calibration; anxiety; delayed mix under stakes
FB-5	Process soft-wrong card vs self-praise toast (same timing)	A/B	Transfer; talk/attention proxies
FB-QUAL	10 Maya: “wait for my reason” vs “instant check” preference	Qual	Allergy to lag; trust; identity language

Falsifier: Micro-delay raises survey “thoughtfulness” but lowers transfer and raises abandon → shorten gate or restrict to high-knowledge tier.

Falsifier: Instant KR wins delayed mix for all segments → demote delay doctrine; keep only diagnostic hide + SE prompts without clock.

Pre-register: FB-* before any “instant AI feedback” or “we scientifically delay” campaign (SAFE-LABMETA).

LXXIX.10 So what for MindCraft commercially

- **Copy:** “Feedback that waits for your thinking — then diagnoses the bridge.” Never “instant personalized feedback” as hero.
- **Product:** Ship mode switches in Practice/Diagnostic/Prove; soft-wrong SE gate; tutor pause cue.
- **Positioning:** Against ChatGPT/Duo instant-KR theater; for attempt-first competence evidence.
- **Metric:** Co-primary retry_120s, solo_transfer_pass, soft-wrong→card; demote latency NPS / green-check rate.
- **Kill list:** Instant-feedback brand; always-delay cosplay; feedback≡learning; ego toasts as FI; green-check NS.
- **Growth:** Differentiate on *quality of struggle recovery*, not feedback speed SLA.
- **Vision:** Identity forms when students author the next move — clocks serve that authorship.

References (verified)

- Brummer, L., de Boer, H., Mouw, J. M., & Strijbos, J. W. (2024). A meta-analysis of the effects of context, content, and task factors of digitally delivered instructional feedback on learning performance. *Learning Environments Research*, 27(3), 453–476. <https://doi.org/10.1007/s10984-024-09501-4>
- Butler, A. C., Karpicke, J. D., & Roediger, H. L., III. (2007). The effect of type and timing of feedback on learning from multiple-choice tests. *Journal of Experimental Psychology: Applied*, 13(4), 273–281. <https://doi.org/10.1037/1076-898X.13.4.273>
- Fyfe, E. R., & Rittle-Johnson, B. (2016a). Feedback both helps and hinders learning: The causal role of prior knowledge. *Journal of Educational Psychology*, 108(1), 82–97. <https://doi.org/10.1037/edu0000053>
- Fyfe, E. R., & Rittle-Johnson, B. (2016b). The benefits of computer-generated feedback for mathematics problem solving. *Journal of Experimental Child Psychology*, 147, 140–151. <https://doi.org/10.1016/j.jecp.2016.03.009>
- Hattie, J., & Timperley, H. (2007). The power of feedback. *Review of Educational Research*, 77(1), 81–112. <https://doi.org/10.3102/003465430298487>
- Kandemir, E. N., Esposito, E., Gurgand, L., & Ramus, F. (preprint). A meta-analysis of the impact of feedback timing on learning outcomes in computer-assisted learning. HAL:hal-05546645. <https://hal.science/hal-05546645>
- Kluger, A. N., & DeNisi, A. (1996). The effects of feedback interventions on performance: A historical review, a meta-analysis, and a preliminary feedback intervention theory. *Psychological Bulletin*, 119(2), 254–284. <https://doi.org/10.1037/0033-2909.119.2.254>
- Kulik, J. A., & Kulik, C.-L. C. (1988). Timing of feedback and verbal learning. *Review of Educational Research*, 58(1), 79–97. <https://doi.org/10.3102/00346543058001079>
- Shute, V. J. (2008). Focus on formative feedback. *Review of Educational Research*, 78(1), 153–189. <https://doi.org/10.3102/0034654307313795>
- van der Kleij, F. M., Feskens, R. C. W., & Eggen, T. J. H. M. (2015). Effects of feedback in a computer-based learning environment on students’ learning outcomes: A meta-analysis. *Review of Educational Research*, 85(4), 475–511. <https://doi.org/10.3102/0034654314564881>

Part LXXX — Bridge-Gap Product Narrative

Chapter status: Living evidence + Map/copy brief — Researcher tick 2026-08-04

Primary question: When should MindCraft narrate failure as a *connection* problem (knows A, knows B, fails A→B) rather than “weak concept,” and what bridge claims survive competitive honesty?

Owners: Product (Map / Practice / Solver) · Engine (bridges / pathfinder) · Brand & copy · Tutor ops · Red Team

Commercial job: Ship a **SAFE-BRIDGE** doctrine: name and treat directed bridges as first-class diagnosis objects; sell *connection failures* as the honest weak-spot story; never sell bridge-count, chat re-explain, or “we connect everything” as the brand.

Builds on: Parts XXXV (competitive audits), XXXIX (interleaving), XL (SAFE-SE), XLVIII (SAFE-TRANSFER), XLIX (SAFE-MISCON), LXVII (SAFE-ONTOLOGY), LXVIII (SAFE-HITL), LXXIX (SAFE-FBTIME).

LXXX.1 Why this chapter exists

Edtech and tutors default to a *node* story: “She’s weak at quadratic equations.” Parents buy node stories. ChatGPT re-explains the node. Progress dashboards light up nodes. MindCraft’s engine already encodes something sharper — **bridges** (directed enabling relations) and bridge-gap severity when two sides look mastered but the join fails — yet Map/copy still risk collapsing that into “another weak topic.”

FOUNDER BELIEF under audit: Many high-friction ACT and classroom failures are *relational*: the student can execute each side in isolation and still cannot map, compare, or carry structure across the join. Naming that failure is commercially load-bearing — it is the honest difference from content libraries and fluent chat — *if* we instrument transfer across the bridge and refuse costume claims.

Claims we refuse as doctrine:

- 1. Bridge-count / ontology edge-count $\equiv$ product quality or moat.
- 2. “We fix connections” / “no more gaps between topics” as absolute marketing.
- 3. Re-explaining concept A (or B) $\equiv$ bridging $A \rightarrow B$.
- 4. Prerequisite absence $\equiv$ bridge gap (different diagnosis; different CTA).
- 5. Story-skin change $\equiv$ bridge practice.
- 6. Bridge-gap severity as Anxiety Score™ / shame Map theater.
- 7. Chat memory of “we covered both” $\equiv$ knowledge integration.

LXXX.2 Constructs

Construct	Research meaning	MindCraft analogue	Failure mode
Fragmented knowledge	Ideas present but weakly linked	High node mastery, low join performance	Celebrate node greens
Bridge (product)	Directed enabling relation $A \rightarrow B$	Ontology bridges; 1.5× priority in runtime	Treat as another node
Bridge gap	Can do A and B separately; fails tasks needing the join	isBridgeGap / severity on Map & /recommend	Mislabel as “weak A”
High-road / bridging instruction	Mindful abstraction + search for connections	Compare/SE prompts; bridge cards	Abstraction monologue without instances
Hugging	Practice that approximates target performance	ACT-like items on the join	Only hug one side
Relational comparison	Aligned mapping $source \leftrightarrow target$	Side-by-side contrast cards; tutor gestures-as-UX	Analogy without cognitive supports
SAFE-BRIDGE	Connection-first narrative + measurement	This chapter	Bridge cosplay ads

Operational definition (HYPOTHESIS): A failure is a *bridge gap* iff (1) recent evidence shows non-trivial competence on both endpoint concepts/ingredients, (2) items that require the directed join fail at elevated rate relative

to endpoint-matched controls, and (3) remediation targets the *relation* (compare, map, carry invariant) rather than re-teaching an endpoint in isolation. Marketing may claim bridge diagnosis only when (1)–(3) are instrumented.

LXXX.3 Knowing both sides ≠ using the join

FACT (spontaneous transfer failure): Gick & Holyoak (1980, *Cognitive Psychology*, 12(3), 306–355, doi:10.1016/0010-0285(80)90013-4) — learners who possessed an analogous “fortress” story often failed to apply the convergence principle to Duncker’s radiation problem until a **hint** made the connection available. Structural isomorphism does not guarantee noticed transfer.

FACT (transfer mechanisms): Salomon & Perkins (1989, *Educational Psychologist*, 24(2), 113–142, doi:10.1207/s15326985ep2402_1) — **low-road** transfer (automatic triggering after varied practice) vs **high-road** transfer (mindful abstraction and deliberate search for connections). Findings of transfer/nontransfer track whether those conditions were met — not whether “the topics were covered.”

FACT (instructional tactics): Perkins & Salomon (1988, *Educational Leadership*, 46(1), 22–32; summarized in Perkins & Salomon, 1992, “Transfer of Learning,” *International Encyclopedia of Education*) — **hugging** (approximate the target) and **bridging** (prompt abstraction, metacognition, planned connections). Education gets transfer when designed for it.

FACT (taxonomy honesty): Barnett & Ceci (2002, *Psychological Bulletin*, 128(4), 612–637, doi:10.1037/0033-2909.128.4.612) — “far transfer” is not one effect size; content×context dimensions must be named. Bridge-gap claims should specify *which* dimension moved (representation, problem type, temporal delay, solo vs scaffolded).

Commercial implication: “She’s done both units” is not readiness. Sell *join evidence*. Chat that re-explains each unit is the wrong remediator for a join failure.

Kill: Coverage ≡ connection.

Survive: Hinted / designed bridging as the product job (ties SAFE-TRANSFER).

LXXX.4 Knowledge integration — links are the learning object

FACT (science-education synthesis): Eylon & Linn (1988, *Review of Educational Research*, 58(3), 251–301, doi:10.3102/00346543058003251) — perspectives on concept learning and problem solving converge on mechanisms of change in knowledge organization; fleeting coverage of many topics underperforms deeper integration work.

FACT (KI framework): Linn & Eylon (2006, in Alexander & Winne, Eds., *Handbook of Educational Psychology*, 2nd ed., pp. 511–544) — knowledge integration as eliciting ideas, adding normative ideas, **distinguishing**, and reflecting toward a more coherent network. Fragmented “pieces” can coexist without productive links.

HYPOTHESIS (MindCraft): Soft-wrongs that fire when both endpoints look “green” are high-value KI events: elicit the student’s mapping, add the bridge card, force distinguish (“what stays the same across these two?”), then reflect into Map as bridge evidence — not as a third unrelated node.

Kill: More topics completed as integration proof.

Survive: Distinguish + reflect on the *relation* as the object of learning (Marton-style discernment in XLVIII applied to bridges).

LXXX.5 Analogies without supports are theater

FACT (classroom video): Richland, Zur & Holyoak (2007, *Science*, 316(5828), 1128–1129, doi:10.1126/science.1142103) — TIMSS 1999 video sample: U.S., Hong Kong, and Japanese eighth-grade lessons all used relational comparisons frequently, but **U.S. teachers provided fewer cognitive supports** (familiar source, visual source kept visible, spatial cues, gestures, imagery) that help students notice the mapping. Analogy *mention* ≠ analogy *learning*.

FACT (instructional agenda): Richland, Stigler & Holyoak (2012, *Educational Psychologist*, 47(3), 189–203, doi:10.1080/00461520.2012.667065) — many post-K–12 students still treat math as memorized (often incorrect) procedures rather than a relational system; refining instruction so students draw connections through **relational comparisons** is a research-backed path to flexible knowledge; U.S. classrooms under-use the supports that make comparisons work.

Product translation:

1. Bridge cards must **co-present** both sides (source visible while targeting the join) — not sequential monologues.
2. Prompt alignment (“point to what maps to what”) before pouring explanation.
3. Tutor HITL briefs: gesture/spatial UX analogues (highlight shared structure), not “explain quadratic again.”
4. Interleave near-miss types that share surface features but differ on the bridge (XXXIX) — discrimination of the join.

Kill: “We use analogies” as pedagogy claim without co-presence + student mapping.

Wound: Some students find dual displays overwhelming (CLT) — dose with SAFE-DD / prior-knowledge tiers.

Survive: Supported relational comparison as Map→Practice CTA copy.

LXXX.6 Engine reality vs UI honesty

Reuse (SAFE-ONTOLOGY): Diagnosis-before-dialogue; bridges prioritized over isolated nodes in ingredient runtime; chat optimizes next utterance, not falsifiable join state.

HYPOTHESIS — product stack:

1. **Map:** Render bridge gaps as *edges* (or edge callouts), not only as red nodes. Copy: “You can do X and Y — the hard part is connecting them.”
2. /recommend / **PawHub Practice:** Prefer playable bridge-gap missions when severity dominates plain 1–mastery (ties C1 worstWeakness).
3. **Practice:** Object of learning = the bridge; VT: hold join invariant, vary one endpoint surface at a time; measure solo_transfer_pass on join items after delay.

4. **Solver / coach:** Soft-wrong → “which connection broke?” SE before endpoint re-teach (SAFE-SE / SAFE-FBTIME).
5. **Tutor ops:** Map brief must list top bridge gaps; prompt>pour on the relation (SAFE-HITL).
6. **Parent surface:** “Connection practice” / proof on joins — not “another weak chapter” shame.

SPECULATION: Competitors can screenshot a graph; they cannot ship longitudinal bridge evidence + solo join transfer without the spine — *if* MindCraft instruments and shows those metrics honestly (same data-moat logic as LXVII).

Kill: Node-only weakness narrative when join evidence exists.

Survive: Connection-first CTA when endpoints are non-naive.

LXXX.7 Competitive positioning

Pattern	Failure story sold	MindCraft response
Khan topic paths	Weak skill node	Name bridge when both skills green
Duo skill tree	Isolated skill XP	Ban XP-as-connection; join probes
Brilliant puzzles	Insight inside one puzzle	Keep delight; add explicit compare across puzzles
ChatGPT tutors	Re-explain whichever topic was named	Refuse re-teach-as-bridge; force mapping
Tutor shops	“Review chapter 5 again”	Brief on A→B join; TR-2 style compare
Cram packs	Massed node blitz	SAFE-DURABLE: join returns beat massed re-cover

Competitive wedge (FOUNDER BELIEF): “We find the connection you think you already have” beats “we explain every topic again.”

LXXX.8 Doctrine — SAFE-BRIDGE (provisional)

1. **Connection > isolated node** when endpoint evidence is non-trivial and join items fail.
2. **Name the edge in copy** — Map, Practice CTA, tutor brief, parent note.
3. **Remediate the relation** — co-present compare, SE mapping, bridge cards; ban endpoint re-teach as default.
4. **Supports required** — visible source+target, alignment prompts; analogy mention ≠ pedagogy.
5. **Measure join transfer** — delayed / format-hop / solo items on the bridge; demote node greens as readiness.
6. **Distinguish diagnoses** — missing prerequisite ≠ bridge gap ≠ format gap (Part LXXXII queued).
7. **No bridge-count moat ads** — edges are substrate; FEI + solo_transfer_pass on joins are the claim.

8. **Severity with dignity** — no shame heatmaps; competence-protection framing (SAFE-FORGET honesty applies to aged bridge evidence too).

Confidence: High that spontaneous far/join transfer is rare without cues (Gick & Holyoak; Salomon & Perkins). High that hugging/bridging are designable tactics (Perkins & Salomon). High that U.S.-style analogy use often lacks cognitive supports (Richland et al. 2007). High that knowledge integration is about links, not coverage (Eylon & Linn; Linn & Eylon). Medium that MindCraft’s bridge-gap detector matches the operational definition above (needs BRIDGE-* calibration against human raters). Medium that connection-first copy lifts WTP vs node-weakness copy (needs parent CBC). High that re-explain=bridge and bridge-count=moat are kills.

LXXX.9 Experiments

ID	Question	Design	Primary
BRIDGE-1	Bridge-card compare/SE vs endpoint re-teach after detected gap	A/B within bridge id	Join-item accuracy; solo_transfer_pass; time-to-remediate
BRIDGE-2	Map edge callout CTA vs node-weakness CTA	A/B	Mission start; completion; post join probe
BRIDGE-3	Detector vs human tutor labels on same sessions	Agreement study	Precision/recall of isBridgeGap; false “weak A”
BRIDGE-4	Co-present dual display vs sequential explain	A/B	Mapping SE quality; transfer; CLT abandon
BRIDGE-5	Parent CBC: “connection gap” vs “weak topic” message	CBC	WTP; trust; stigma
BRIDGE-QUAL	10 Maya: “I knew both but...” language	Qual	Identity language; allergy to Map edges

Falsifier: Endpoint re-teach matches or beats bridge remediation on delayed join items → demote connection-first default; keep bridges as secondary tags only.

Falsifier: Detector disagreement with tutors > threshold → freeze severity-driven PawHub routing until calibrated.

Falsifier: Connection copy raises shame/abandon without transfer lift → rewrite dignity framing; reduce edge heat.

Pre-register: BRIDGE-* before any “we fix the connections textbooks miss” campaign (SAFE-LABMETA).

LXXX.10 So what for MindCraft commercially

- **Copy:** “The gap isn’t the chapter — it’s the connection.” Never “master every topic” as the hero.

- **Product:** Edge-aware Map; bridge missions in Practice; Solver SE on joins; tutor briefs list bridges first when severity wins.
 - **Positioning:** Against chat re-explain and skill-tree XP; for inspectable join diagnosis (with SAFE-ONTOLOGY).
 - **Metric:** Co-primary join solo_transfer_pass, bridge remediation success, detector↔tutor agreement; demote node greens / bridge-count.
 - **Kill list:** Bridge-count moat; re-explain⇒bridge; coverage⇒connection; absolute “we connect everything”; shame-edge theater.
 - **Growth:** Parent WTP for *named connection practice* after both sides look fine — a story Khan/Duo/ChatGPT do not own.
 - **Vision:** Identity as a mathematical thinker is relational competence — seeing structure *between* ideas, not collecting lit nodes.
-

References (verified)

- Barnett, S. M., & Ceci, S. J. (2002). When and where do we apply what we learn? A taxonomy for far transfer. *Psychological Bulletin*, 128(4), 612–637. <https://doi.org/10.1037/0033-2909.128.4.612>
 - Eylon, B., & Linn, M. C. (1988). Learning and instruction: An examination of four research perspectives in science education. *Review of Educational Research*, 58(3), 251–301. <https://doi.org/10.3102/00346543058003251>
 - Gick, M. L., & Holyoak, K. J. (1980). Analogical problem solving. *Cognitive Psychology*, 12(3), 306–355. [https://doi.org/10.1016/0010-0285\(80\)90013-4](https://doi.org/10.1016/0010-0285(80)90013-4)
 - Linn, M. C., & Eylon, B.-S. (2006). Science education: Integrating views of learning and instruction. In P. A. Alexander & P. H. Winne (Eds.), *Handbook of educational psychology* (2nd ed., pp. 511–544). Lawrence Erlbaum.
 - Perkins, D. N., & Salomon, G. (1988). Teaching for transfer. *Educational Leadership*, 46(1), 22–32.
 - Perkins, D. N., & Salomon, G. (1992). Transfer of learning. In *International encyclopedia of education* (2nd ed.). Pergamon.
 - Richland, L. E., Stigler, J. W., & Holyoak, K. J. (2012). Teaching the conceptual structure of mathematics. *Educational Psychologist*, 47(3), 189–203. <https://doi.org/10.1080/00461520.2012.667065>
 - Richland, L. E., Zur, O., & Holyoak, K. J. (2007). Cognitive supports for analogies in the mathematics classroom. *Science*, 316(5828), 1128–1129. <https://doi.org/10.1126/science.1142103>
 - Salomon, G., & Perkins, D. N. (1989). Rocky roads to transfer: Rethinking mechanisms of a neglected phenomenon. *Educational Psychologist*, 24(2), 113–142. https://doi.org/10.1207/s15326985ep2402_1
-

Part LXXXI — Cold-Start Without Fake Mastery

Chapter status: Living evidence + onboarding/copy brief — Researcher tick 2026-08-05

Primary question: How should MindCraft initialize a student graph on day one so recommendations feel personal *without* manufacturing mastery fireworks, inflated greens, or “we already know you” theater?

Owners: Product (gap-scan / PawHub / Map) · Engine (seed-assessment / priors / event sources) · Brand & copy · Parent trust · Red Team

Commercial job: Ship a **SAFE-COLD** doctrine: humble priors + labeled seed evidence + hide-correctness probes; sell *honest mapping*, never day-one mastery or confidence⇒mastery.

Builds on: Parts L (SAFE-CALIB), LXVII (SAFE-ONTOLOGY), LXXII (SAFE-GENQ), LXXVI (SAFE-FORGET), LXXIX (SAFE-FBTIME), LXXX (SAFE-BRIDGE). Product seam: /seed-assessment, C4 hide-correctness diagnostic, population_priors, eventCount === 0 → Learn-next.

LXXXI.1 Why this chapter exists

Every adaptive product faces the same commercial temptation at cold start: empty state feels broken, so the UI invents competence — green trees, “You’re already strong at...,” personalized paths from zero observations. Parents buy the fireworks. Students learn the app rewards confident self-report and hint-skipping. The graph then *lies* for weeks.

MindCraft’s gap-scan already does something sharper: confidence ratings → synthetic assessment events (source=onboarding_assessment) that *replace* prior seed, plus optional hide-correctness probes. That is the right *family* of move — **if** we refuse to market the seed as mastery, refuse easy-default skips, and keep Map language in “uncertain / untouched / early signal” until delayed proof arrives (SAFE-FORGET).

FOUNDER BELIEF under audit: The first session’s job is *diagnosis hygiene*, not delight-as-competence. Personalization that survives is “we asked hard questions and stayed honest about uncertainty,” not “your Map lit up on minute five.”

Claims we refuse as doctrine:

1. Day-one green Map / mastery fireworks ≡ personalization.
 2. Self-rated “easy” ≡ mastery or ACT readiness.
 3. Population prior alone ≡ student-specific diagnosis worth selling.
 4. Skip-scan / “I’ll explore” defaults that write easy/high mastery.
 5. Onboarding chat fluency or AI pep ≡ calibrated prior.
 6. Cold-start “Identity Score™” or “math person already” splash.
 7. Unverified generated keys as first diagnostic ground truth (SAFE-GENQ).
-

LXXXI.2 Constructs

Construct	Research meaning	MindCraft analogue	Failure mode
Cold start	No / little learner history for KT or recs	New users/{uid}; empty graph	Fake greens to fill void
Prior (P(L0))	Initial P(know skill) before observations	Ontology population_priors; Beta seeds	Sell prior as “your level”
Seed evidence	Labeled synthetic events from onboarding	/seed-assessment confidence map	Treat seed as FEI mastery
Probe evidence	Observed correctness without reveal theater	C4 hide-correctness items	Reveal keys → answer-hunting

Calibration	Confidence matches accuracy long-run	SAFE-CALIB; hard/kinda/easy	Confidence $\equiv$ truth
Fake mastery	UI/state claims competence without durable proof	Fireworks; green untouched nodes	Trust debt; wrong path
SAFE-COLD	Humble init + labeled seed + probes	This chapter	Personalization cosplay

Operational definition (HYPOTHESIS): A cold-start is *honest* iff (1) every non-prior belief is tagged by source (onboarding_assessment / practice / session), (2) Map/copy distinguish untouched, seed-only, and delayed-proof states, (3) no mastery confetti or “you’ve mastered X” until criteria beyond seed confidence, and (4) skip paths do not write high mastery. Marketing may claim “personalized from day one” only as *personalized uncertainty + next probe*, not personalized mastery.

LXXXI.3 Knowledge tracing needs a prior — priors are not biographies

FACT (BKT spine): Corbett & Anderson (1995, *User Modeling and User-Adapted Interaction*, 4, 253–278, doi:10.1007/BF01099821) — knowledge tracing maintains P(learned) per production rule from observed attempts (prior, learn, guess, slip). Mastery-based sequencing depends on those estimates; the model is only as honest as its updates and its **initial prior**.

FACT (individualizing P(L0)): Pardos & Heffernan (2010, in *UMAP 2010*, LNCS 6075, pp. 255–266, doi:10.1007/978-3-642-13470-8_24) — standard deployed BKT often used skill-specific but not student-specific initial knowledge; individualizing the prior (especially pooling information across skills for each student) improved prediction on ASSISTments problem sets (lower error on 33/42 sets in their evaluation). Cold start is a *parameter* problem, not a license to invent greens.

FACT (prior sensitivity): Students enter with heterogeneous P(L0); a single global “everyone starts at 0.1” or “everyone starts mastered” both mis-sequence practice (BKT robustness / prior-variation literature; Pardos & Heffernan individualization agenda).

HYPOTHESIS (MindCraft): Ontology population_priors are the right *default* P(L0) layer — inspectable, shared, not personalized theater. Student-specific movement must come from labeled seed + probes + practice, with decay toward prior (align copy with engine honesty).

Kill: Population prior sold as “we know Maya.”

Survive: Prior as humble shared baseline; personalization = updates with provenance.

LXXXI.4 Self-report is necessary and systematically biased

Gap-scan asks hard / kinda / easy. That is commercially necessary (coverage across ~29 ACT concepts without a three-hour exam) and epistemically dangerous.

FACT (overconfidence): Fischhoff, Slovic & Lichtenstein (1977, *Journal of Experimental Psychology: Human Perception and Performance*, 3(4), 552–564) — people are often wrong when “certain”; calibration curves show hit rates below stated confidence.

FACT (reasons bias): Koriat, Lichtenstein & Fischhoff (1980, *Journal of Experimental Psychology: Human Learning and Memory*, 6(2), 107–118) — confidence is inflated by evidence for the chosen answer; prompting contradicting reasons improves appropriateness of confidence.

FACT (unskilled–unaware pattern): Kruger & Dunning (1999, *Journal of Personality and Social Psychology*, 77(6), 1121–1134, doi:10.1037/0022-3514.77.6.1121) — bottom-quartile performers grossly overestimated ability. (Directional product risk — not a student-facing label.)

FACT (math metacognition): Erickson & Heit (2015, *Frontiers in Psychology*, 6, 742, doi:10.3389/fpsyg.2015.00742) — students overpredicted math test scores alongside math anxiety; not simply “low confidence everywhere.”

FACT (item-level blindness): Lindsey & Nagel (2015, *Physical Review Special Topics — Physics Education Research*, 11, 020103, doi:10.1103/PhysRevSTPER.11.020103) — at item level, students of all abilities struggle to foresee which questions they will miss — self-assessment is a blunt prior, not a key.

Commercial implication: Map “easy” → L3 gating is coherent; Map “easy” → green mastery badge is a lie. Pair confidence seed with hide-correctness probes (SAFE-CALIB / C4); never reveal keys during diagnostic (SAFE-GENQ).

Kill: Confidence≡mastery; “raise onboarding confidence” as success metric.

Survive: Confidence as *routing prior* + miss-tier signal; probes as corrective evidence.

LXXXI.5 Gaming the onboarding is fake mastery’s twin

FACT (gaming + harm): Baker, Corbett, Koedinger & Wagner (2004, *CHI ’04*, pp. 383–390) — “gaming the system” = succeeding by exploiting system properties rather than learning; gaming frequency correlated negatively with learning, distinct from other off-task behavior.

HYPOTHESIS (product): Cold-start analogues include tapping “easy” on everything, skipping gap-scan, or hint-binging first Practice to farm greens. Designs that reward scan-completion over scan-honesty invite this.

Product translation:

1. No streak/XP for “finished diagnostic.”
2. Skip-scan must not write optimistic mastery; prefer untouched + population prior.
3. Retake gap-scan *replaces* onboarding seed — do not stack optimistic layers.
4. PawHub Practice/Learn-next read eventCount honesty — do not celebrate seed as conquered islands.

Kill: Diagnostic completion % as North Star.

Wound: Anxious under-raters (“hard” everywhere) — SAFE-CALIB underconfidence unlock still applies.

Survive: Friction that makes honest mapping cheaper than fake competence.

LXXXI.6 Adaptive testing envy — shorter ≠ fireworks

FACT (CAT premise): Adaptive testing selects informative items given current estimates and can shorten tests when banks/models are calibrated (CAT survey/primer tradition). Early steps still start from a prior; stopping rules are about *measurement error*, not celebration.

HYPOTHESIS: Borrow CAT’s *information* ethic (ask what reduces uncertainty) without placement trophies or “you placed into Gold” skins. Gap-scan + sparse probes ≈ poor-man’s adaptive placement; sell “where to start,” not a belt.

Kill: Placement-league / belt / IQ-cosplay cold start.

Survive: Uncertainty-reducing item selection as the onboarding craft.

LXXXI.7 Engine + UI stack (SAFE-COLD)

Reuse: SAFE-ONTOLOGY (inspectable state); SAFE-CALIB + C4 hide-correctness; SAFE-FORGET (age evidence; no fireworks≡durable); SAFE-FBTIME (diagnostic hide); SAFE-GENQ (verified keys only for probes that write mastery).

HYPOTHESIS — product stack:

1. **Prior layer:** Population / ontology priors visible as *shared* baselines — never “your mastery.”
2. **Seed layer:** /seed-assessment writes labeled synthetic events; UI copy: “your confidence map” / “starting signals,” not “mastered.”
3. **Probe layer:** Short hide-correctness set across concepts/formats; record outcomes without reveal; optional contradicting-reason micro-prompt on high-confidence picks (Koriat et al. mechanism, light-touch).
4. **Map:** Untouched ≠ red failure; seed-only ≠ green mastery; reserve strong chrome for delayed/generation proof.
5. **PawHub:** Learn-next = first eventCount === 0 playable on exam path; Practice = worst playable weakness — both read honesty, not confetti.
6. **Parent:** “We started with an honest map” CBC arm vs “personalized genius path from minute one.”

SPECULATION: Competitors can ship prettier empty-state trees; MindCraft’s wedge is *provenance* — every early claim has a source tag parents and tutors can understand.

LXXXI.8 Competitive positioning

Pattern	Cold-start move	MindCraft response
Duo skill tree	Early greens / crowns	Untouched honesty; ban XP-as-mastery
Khan mastery	Energy + % bars from thin attempts	Seed ≠ mastery bar; delayed proof
Brilliant	Delight puzzle → “you’re clever”	Delight ok; no cleverness≡graph

ChatGPT tutors	Instant “I know your level” from chat	Chat ≠ prior; refuse fluency prior
Bootcamp quizzes	Placement score theater	Uncertainty map; no belt
Adaptive apps	“AI personalized path” from 3 clicks	Label seed; show what was never probed

Competitive wedge (FOUNDER BELIEF): “We refuse to fake knowing you” is sellable trust — especially to parents burned by apps that looked personal and taught nothing.

LXXXI.9 Doctrine — SAFE-COLD (provisional)

1. **Humble prior first** — population/ontology priors initialize; they are not biography.
2. **Labeled seed only** — confidence → synthetic events with source tag; replace-on-retake, don’t stack lies.
3. **Probes without reveal** — C4 hide-correctness; verified keys only; no answer-hunting diagnostic.
4. **Confidence ≠ mastery** — hard/kinda/easy routes level + struggle prior; never crowns.
5. **Map chrome matches evidence age** — untouched / seed-only / proved (SAFE-FORGET).
6. **No skip-to-green** — skip leaves untouched + prior; completion is not a learning KPI.
7. **Anti-gaming** — no streak/XP for scan finish; watch uniform-easy patterns as QA.
8. **Copy:** “Honest starting map” / “where to begin” — never “you’ve already mastered” / Identity Score™.

Confidence: High that BKT-style systems need explicit priors and that individualizing P(L0) matters (Corbett & Anderson; Pardos & Heffernan). High that confidence is miscalibrated and especially risky for low performers (Fischhoff et al.; Kruger & Dunning; Erickson & Heit). High that gaming harms learning when systems can be exploited (Baker et al.). Medium that MindCraft’s current seed→event map is well-calibrated to later practice (needs COLD-*). Medium that honest-cold-start copy wins parent WTP vs fireworks (needs CBC). High that day-one mastery fireworks and confidence=mastery are kills.

LXXXI.10 Experiments

ID	Question	Design	Primary
COLD-1	Seed-only vs seed+hide-correctness probes	A/B new users	7d /recommend stability; probe–seed disagreement rate; abandon
COLD-2	Map chrome: seed-as-mastery greens vs honest untouched/seed labels	A/B	Trust items; wrong-path starts; parent CBC
COLD-3	Uniform-easy / skip patterns → QA flag vs ignore	Observational + intervene	Later miss rate; gaming-like behavior

COLD-4	Contradicting-reason micro-prompt on “easy” ratings	A/B	Calibration vs probes; time cost
COLD-5	Parent CBC: “honest map” vs “personalized from minute one”	CBC	WTP; trust; stigma
COLD-QUAL	10 Maya: what felt fake vs useful in first session	Qual	Language for copy; fireworks allergy

Falsifier: Seed-as-mastery greens raise 7d retention *and* delayed transfer vs honest chrome → revisit chrome (unlikely; if retention-only, treat as engagement trap).

Falsifier: Probes add abandon without improving recommend quality → shorten probe set; keep seed.

Falsifier: Honest-map copy loses WTP with no trust gain → reframe dignity without fireworks; do not restore fake greens.

Pre-register: COLD-* before any “personalized genius path from day one” campaign (SAFE-LABMETA).

LXXXI.11 So what for MindCraft commercially

- **Copy:** “An honest starting map — not a fake report card.” Never day-one “mastered” / belts / Identity Score™.
 - **Product:** Prior → labeled seed → hide-correctness probes → provenance on Map; PawHub reads eventCount and seed humility.
 - **Positioning:** Against Duo greens, chat “I know you,” and placement theater; for inspectable cold-start (with SAFE-ONTOLOGY).
 - **Metric:** Probe–seed disagreement, recommend stability 7d, skip/uniform-easy rate, parent trust; demote diagnostic completion % and early green counts.
 - **Kill list:** Fake mastery fireworks; confidence≡mastery; skip-to-green; population-prior-as-biography; unverified diagnostic keys.
 - **Growth:** Trust compounds when week-two practice matches week-one honesty — fireworks create churn when the lie shows.
 - **Vision:** Identity transformation starts with *accurate self-location*, not premature crowning.
-

References (verified)

- Baker, R. S., Corbett, A. T., Koedinger, K. R., & Wagner, A. Z. (2004). Off-task behavior in the Cognitive Tutor classroom: When students game the system. In *Proceedings of CHI 2004* (pp. 383–390). ACM.
- Corbett, A. T., & Anderson, J. R. (1995). Knowledge tracing: Modeling the acquisition of procedural knowledge. *User Modeling and User-Adapted Interaction*, 4, 253–278. <https://doi.org/10.1007/BF01099821>
- Erickson, S., & Heit, E. (2015). Metacognition and confidence: Comparing math to other academic subjects. *Frontiers in Psychology*, 6, 742. <https://doi.org/10.3389/fpsyg.2015.00742>

- Fischhoff, B., Slovic, P., & Lichtenstein, S. (1977). Knowing with certainty: The appropriateness of extreme confidence. *Journal of Experimental Psychology: Human Perception and Performance*, 3(4), 552–564.
- Koriat, A., Lichtenstein, S., & Fischhoff, B. (1980). Reasons for confidence. *Journal of Experimental Psychology: Human Learning and Memory*, 6(2), 107–118.
- Kruger, J., & Dunning, D. (1999). Unskilled and unaware of it: How difficulties in recognizing one’s own incompetence lead to inflated self-assessments. *Journal of Personality and Social Psychology*, 77(6), 1121–1134. <https://doi.org/10.1037/0022-3514.77.6.1121>
- Lindsey, B. A., & Nagel, M. L. (2015). Do students know what they know? Exploring the accuracy of students’ self-assessments. *Physical Review Special Topics — Physics Education Research*, 11, 020103. <https://doi.org/10.1103/PhysRevSTPER.11.020103>
- Pardos, Z. A., & Heffernan, N. T. (2010). Modeling individualization in a Bayesian networks implementation of knowledge tracing. In P. De Bra, A. Kobsa, & D. Chin (Eds.), *UMAP 2010* (LNCS 6075, pp. 255–266). Springer. https://doi.org/10.1007/978-3-642-13470-8_24

Part LXXXII — Format-Axis Product Narrative

Chapter status: Living evidence + Map/Practice/copy brief — Researcher tick 2026-08-05

Primary question: When should MindCraft narrate weakness as a *format* failure (can do the idea symbolically, fails the diagram — or vice versa) rather than “weak concept,” and what FormatId claims survive competitive honesty?

Owners: Product (Practice / Map / gap-scan) · Engine (format gaps / getQuestions format param) · Bank & generation · Brand & copy · Red Team

Commercial job: Ship a **SAFE-FORMAT** doctrine: treat input representation as a first-class diagnosis axis; sell *format-flexible competence* with conversion evidence; never sell format-count, decorative diagrams, or “we cover every format” as the brand.

Builds on: Parts XXVI (CLT), XXXIX (interleaving), XL (SAFE-SE), XLVIII (SAFE-TRANSFER / VT), LXIV (SAFE-ELL language vs math load), LXVII (SAFE-ONTOLOGY), LXXII (SAFE-GENQ), LXXX (SAFE-BRIDGE), LXXXI (SAFE-COLD). Product seams: FormatId / C2–C3, format-tagged bank, format-aware CTA, Layer-4 representation_profile.

LXXXII.1 Why this chapter exists

Edtech and tutors default to a *topic* story: “She’s weak at functions.” Parents buy topics. Progress dashboards light up topics. MindCraft’s bank and recommend stack already encode something sharper — **formats** (word_problem, diagram, number_line, symbolic_expression, coordinate_graph, ...) and format-preferring practice — yet Map/copy still risk collapsing “missed the graph” into “weak functions.”

FOUNDER BELIEF under audit: A large share of ACT/classroom fails are *representational*: the student can execute the idea in one vessel and still cannot read, translate, or produce it in another. Naming that failure is commercially load-bearing — it is an honest difference from symbolic drill packs and fluent chat that only ever sees text — *if* we instrument conversion and refuse costume claims.

Claims we refuse as doctrine:

1. Format-count / FormatId vocabulary size $\equiv$ product quality or moat.
 2. “We teach every format” / “visual learning for all” as absolute marketing.
 3. Pretty diagram / story art $\equiv$ FormatId coverage or learning.
 4. Concept mastery on one format $\equiv$ format-flexible readiness.
 5. Format shuffle / random vessel $\equiv$ science without conversion practice.
 6. Format Personality / learning-style quiz as North Star.
 7. Unverified generated figures as diagnostic ground truth (SAFE-GENQ).
-

LXXXII.2 Constructs

Construct	Research meaning	MindCraft analogue	Failure mode
External representation	Visible encoding of math ideas	Item stem vessel; Map format tags	Treat as decoration
Format / vessel	Product vocab for input representation	FormatId; Layer-4 representation_profile	Collapse into topic only
Treatment	Transform within one register	Symbolic algebra within symbols	Call treatment “transfer”
Conversion	Translate across registers	Diagram $\leftrightarrow$ equation practice	Re-teach topic instead
Representational competence	Knowing how a vessel encodes content	Format fluency + conversion prompts	Assume vessels are transparent
Format gap	Strong in format A, weak in B for same concept	Format severity in /recommend / worstWeakness	Mislabel as “weak concept”
SAFE-FORMAT	Format-first narrative + conversion measurement	This chapter	Format cosplay ads

Operational definition (HYPOTHESIS): A failure is a *format gap* iff (1) recent evidence shows non-trivial competence on the same concept in at least one other FormatId, (2) items in the weak FormatId fail at elevated rate relative to format-matched controls on that concept, and (3) remediation targets *reading/producing/converting* that vessel — not only re-explaining the concept in the student’s strong format. Marketing may claim format diagnosis only when (1)–(3) are instrumented.

LXXXII.3 Registers differ — conversion is the hard cognitive cut

FACT (semiotic registers): Duval (2006, *Educational Studies in Mathematics*, 61(1–2), 103–131) — mathematical activity requires coordinating semiotic representations; **treatment** (within one register) and **conversion** (across registers) are distinct cognitive processes; conversion is often the deciding factor for comprehension failures even when

within-register work looks fine.

FACT (computational efficiency $\neq$ equivalence): Larkin & Simon (1987, *Cognitive Science*, 11(1), 65–100, doi:10.1111/j.1551-6708.1987.tb00863.x) — diagrammatic and sentential representations can be informationally equivalent yet differ sharply in search cost, pattern recognition, and which inferences are “free”; a diagram is sometimes worth ten thousand words *for a given operator set* — not as a marketing slogan that “visuals always help.”

Commercial implication: “She knows linear equations” is incomplete. Sell *which vessels she can use* and whether she can convert. Chat that only ever explains in symbols cannot remediate a diagram-read failure.

Kill: Single-format accuracy $\equiv$ concept mastery.

Survive: Conversion as a named learning object (ties SAFE-TRANSFER / SAFE-BRIDGE when the join is also a register join).

LXXXII.4 Multiple representations help only under design + competence conditions

FACT (functional taxonomy): Ainsworth (1999, *Computers & Education*, 33(2–3), 131–152, doi:10.1016/S0360-1315(99)00029-9) — MERs serve distinct functions: **complement**, **constrain**, and **construct**; conflicting evaluation results track unrecognized function mix and unsupported translation across representations.

FACT (DeFT): Ainsworth (2006, *Learning and Instruction*, 16(3), 183–198, doi:10.1016/j.learninstruc.2006.03.001) — effectiveness depends on design parameters, pedagogical functions, and the cognitive **tasks** of understanding each representation *and relating them*; learners often treat representations in isolation and fail to integrate.

FACT (representation dilemma): Rau (2017, *Educational Psychology Review*, 29(4), 717–761, doi:10.1007/s10648-016-9365-3) — students often must learn content they do not yet understand from visuals they do not yet understand; benefit from multiple visuals requires **representational competencies**; more visuals $\neq$ automatically better — conditions matter (T+SV vs T+MV hypotheses).

FACT (support moderator): Rexig, Kuhn, Becker & Malone (2024, *Educational Psychology Review*, 36, article 124, doi:10.1007/s10648-024-09958-y) — meta-analysis (132 effect sizes): learning with more than two external representations can show small performance and cognitive-load advantages vs two, but benefits **primarily depend on appropriate support** — stacking formats without scaffolds is not a product strategy.

HYPOTHESIS (MindCraft): Format-axis Practice should (a) diagnose per FormatId, (b) remediate with co-present conversion supports (source vessel visible while targeting the weak one), and (c) measure delayed format-hop transfer — not dump three diagrams beside a stem for “engagement.”

Kill: “More formats = better learning” without support + conversion tasks.

Wound: Dual/multi displays can overload (CLT / XXVI) — dose by prior knowledge and anxiety (SAFE-DD).

Survive: Supported conversion as Map→Practice CTA.

LXXXII.5 Coherence: decoration is not FormatId coverage

FACT (extraneous material hurts): Mayer, Heiser & Lonn (2001, *Journal of Educational Psychology*, 93(1), 187–198, doi:10.1037/0022-0663.93.1.187) — adding interesting but extraneous material can reduce understanding (coherence / cognitive constraints on multimedia).

HYPOTHESIS: Story-world chrome that does not encode the mathematical structure is *not* a diagram FormatId item. Bank tags must distinguish decorative illustration from structurally necessary diagrams/graphs/number lines. Generation that invents pretty figures without keyed structure fails SAFE-GENQ and SAFE-FORMAT together.

Kill: Story art $\equiv$ visual-learning claim.

Survive: Structural FormatId tagging + verified figures only when they carry the math.

LXXXII.6 Engine + UI stack (SAFE-FORMAT)

Reuse: SAFE-ONTOLOGY (inspectable state); SAFE-TRANSFER / VT (vary representation, hold invariant); SAFE-BRIDGE (when join = register join); SAFE-ELL (language load $\neq$ math load; format triangulation); SAFE-GENQ (verified keys/figures); SAFE-COLD (format probes in hide-correctness without fireworks); C1 severity comparable across concept/bridge/format gaps.

HYPOTHESIS — product stack:

- 1. **Bank:** Every practice/diagnostic item carries FormatId; prefer format-matched draws with concept-pool fallback (C3).
- 2. **State:** Maintain `representation_profile` (or equivalent per-format evidence), not only concept mastery.
- 3. **Map / /recommend:** Surface format gaps when severity dominates plain 1–mastery; copy: “You can do this symbolically — the hard part is the graph” (or reverse).
- 4. **PawHub / Practice CTA:** Format-aware mission labels; object of learning = vessel + conversion, not only topic name.
- 5. **Cards / tutor HITL:** Co-present source and target formats; SE prompt on what maps to what (SAFE-SE); never explain-only in the strong format.
- 6. **Parent:** “Format flexibility” CBC arm vs “more visuals” / learning-style quiz.

SPECULATION: Competitors can ship prettier diagram packs; MindCraft’s wedge is *falsifiable format state* — which vessels fail, with conversion missions, not a VARK costume.

LXXXII.7 Competitive positioning

Pattern	Format move	MindCraft response
Khan / drill apps	Mostly symbolic or text stems	Format profile + conversion missions
Duo	Skill tree greens ignore vessel	Ban format-blind mastery crowns
Brilliant	Delight visuals as brand	Delight ok; tag structure vs decoration

ChatGPT tutors	Text-only explanations	Text fluency ≠ diagram competence
Learning-style apps	“You’re a visual learner”	Kill style quiz; sell conversion evidence
ACT cram PDFs	Mixed vessels unlabeled	Label FormatId; dual-rail prove vs learn

Competitive wedge (FOUNDER BELIEF): “We find when the idea fails in the vessel” is sellable diagnosis — especially vs chat that never sees the graph the student cannot read.

LXXXII.8 Doctrine — SAFE-FORMAT (provisional)

- 1. Format is first-class evidence** — tag items; update per-FormatId competence separately from concept greens.
- 2. Conversion is the object** — remediate with co-present vessels + mapping prompts; measure format-hop transfer.
- 3. Concept ≠ format-flexible** — single-vessel success never markets as full mastery.
- 4. Support before stack** — more FormatIds without scaffolds is load, not pedagogy (Ainsworth / Rau / Rexigel et al.).
- 5. Decoration ≠ diagram** — structural encodings only count as FormatId coverage.
- 6. No Format Personality™** — ban learning-style quiz as North Star.
- 7. Severity honesty** — format gaps compete in worstWeakness with concept/bridge gaps (C1).
- 8. Copy:** “Where the idea breaks in the graph / word problem / number line” — never “we teach every format” / format-count moat.

Confidence: High that treatment≠conversion and conversion drives many math comprehension failures (Duval). High that MERs are conditional and translation is hard (Ainsworth 1999/2006; Rau 2017). High that unsupported multi-rep stacking is not automatic gain (Rexigel et al. 2024). High that diagram vs sentential efficiency differs by operators (Larkin & Simon). Medium that MindCraft’s current FormatId bank density is enough for reliable per-student format profiles (needs FORMAT-* + coverage audit). Medium that format-gap copy wins parent WTP vs “more visuals” (needs CBC). High that format-count moat, style quizzes, and decoration≡coverage are kills.

LXXXII.9 Experiments

ID	Question	Design	Primary
FORMAT-1	Format-gap CTA vs concept-only CTA when severity says format	A/B	Format-hop solo_transfer_pass; abandon; hint binge
FORMAT-2	Co-present conversion card vs strong-format re-explain	A/B within concept	Delayed weak-format accuracy

FORMAT-3	Structural diagram items vs decorative-art distractors	Bank A/B or audit	Error patterns; false “diagram” tags
FORMAT-4	Multi-format stack with supports vs unsupported triple display	A/B	Load proxies; transfer; time
FORMAT-5	Parent CBC: “format flexibility” vs “visual learning / every format”	CBC	WTP; trust; stigma
FORMAT-QUAL	10 Maya: when did a graph/word problem feel like a different subject?	Qual	Copy language; vessel shame

Falsifier: Format-gap CTAs raise engagement but not delayed weak-format transfer → demote copy; keep tagging for research only.

Falsifier: Co-present conversion increases load/abandon without transfer gain → dose/scaffold; do not restore explain-in-strong-format as default.

Falsifier: “Every format” marketing wins WTP with no trust cost → still ban as doctrine (overclaim); test quieter “graph vs symbols” specificity.

Pre-register: FORMAT-* before any “visual learning / every format / Format Score™” campaign (SAFE-LABMETA).

LXXXII.10 So what for MindCraft commercially

- **Copy:** “We find where the idea breaks in the format — graph, symbols, words — and practice the conversion.” Never format-count, VARK, or “we teach every format.”
- **Product:** FormatId on bank + representation_profile + format-aware / recommend CTA; co-present conversion cards; Map language for format gaps alongside bridges.
- **Positioning:** Against symbolic-only drills, chat text fluency, and learning-style quizzes; for inspectable format diagnosis (with SAFE-ONTOLOGY).
- **Metric:** Format-hop transfer, format-gap mission starts, false decorative tags, parent CBC; demote format-count and style-quiz completion.
- **Kill list:** Format-count moat; decoration≡coverage; concept≡format-flexible; shuffle≡science; Format Personality™.
- **Growth:** ACT parents already feel “she freezes on word problems / graphs” — name it, instrument it, sell the remediator.
- **Vision:** Identity as a math thinker includes *moving ideas across vessels*, not only green topic nodes.

References (verified)

- Ainsworth, S. (1999). The functions of multiple representations. *Computers & Education*, 33(2–3), 131–152. [https://doi.org/10.1016/S0360-1315\(99\)00029-9](https://doi.org/10.1016/S0360-1315(99)00029-9)
 - Ainsworth, S. (2006). DeFT: A conceptual framework for considering learning with multiple representations. *Learning and Instruction*, 16(3), 183–198. <https://doi.org/10.1016/j.learninstruc.2006.03.001>
 - Duval, R. (2006). A cognitive analysis of problems of comprehension in a learning of mathematics. *Educational Studies in Mathematics*, 61(1–2), 103–131.
 - Larkin, J. H., & Simon, H. A. (1987). Why a diagram is (sometimes) worth ten thousand words. *Cognitive Science*, 11(1), 65–100. <https://doi.org/10.1111/j.1551-6708.1987.tb00863.x>
 - Mayer, R. E., Heiser, J., & Lonn, S. (2001). Cognitive constraints on multimedia learning: When presenting more material results in less understanding. *Journal of Educational Psychology*, 93(1), 187–198. <https://doi.org/10.1037/0022-0663.93.1.187>
 - Rau, M. A. (2017). Conditions for the effectiveness of multiple visual representations in enhancing STEM learning. *Educational Psychology Review*, 29(4), 717–761. <https://doi.org/10.1007/s10648-016-9365-3>
 - Rexigel, E., Kuhn, J., Becker, S., & Malone, S. (2024). The more the better? A systematic review and meta-analysis of the benefits of more than two external representations in STEM education. *Educational Psychology Review*, 36, 124. <https://doi.org/10.1007/s10648-024-09958-y>
-

Part LXXXIII — Parent Dashboard Honesty UX

Chapter status: Living evidence + parent-surface brief — Researcher tick 2026-08-05

Primary question: How should MindCraft’s parent-facing surface show *proof age, returns, and competence evidence* so families can act — without shame heatmaps, tonight-% fireworks, or surveillance cosplay?

Owners: Product (ParentDashboard / MoC surfaces) · Growth/copy · Brand · Privacy · Red Team

Commercial job: Ship a **SAFE-PDASH** doctrine: treat the parent dashboard as a *belief-correction + action* instrument (accurate effort/competence beliefs → autonomy-supportive next step), never as a gradebook clone, streak theater, or shame engine.

Builds on: Parts XXVII (parent anxiety), LIX (SAFE-WTP), LXIII (SAFE-EXAM), LXVI (SAFE-STRUCTURE), LXIX (SAFE-PRIVACY), LXXVI (SAFE-FORGET), LXXVIII (SAFE-DURABLE), LXXXI (SAFE-COLD), LXXXII (SAFE-FORMAT). Product seam: parent proof-age / returns / MoC without selling affect or raw wrong-answer feeds.

LXXXIII.1 Why this chapter exists

SAFE-FORGET defined *what* honesty means (age evidence; JOLs≠retention). SAFE-WTP defined *how to price* legible MoC. This tick answers the *parent UX* problem: edtech parent portals default to (a) real-time greens and missing-assignment panic, (b) gamification screenshots for adults, or (c) AI-written “supportive” digests that flatten diagnosis. MindCraft’s identity thesis dies if parents only see tonight’s glow — and dies differently if the dashboard arms controlling homework battles.

FOUNDER BELIEF under audit: Parents will pay and stay for a dashboard that *corrects optimistic misbeliefs with dated competence evidence* and offers one autonomy-supportive action (“return due before ACT — here’s the mission”), outperforming both “100% mastered!” theater and red-alert shame — if copy never implies character failure or guaranteed points.

Claims we refuse as doctrine:

1. Tonight-% / streak / XP / minutes as parent North Star.
2. Mastery fireworks / “cleared forever” on the parent surface.
3. Shame heatmaps, public ranks, or “your child is behind classmates” as default chrome (SAFE-COMPARE collision).
4. Real-time wrong-answer surveillance / live session stalking as “involvement.”
5. Anxiety Score™ / mood telemetry on parent portal (SAFE-PRIVACY).
6. Gradebook clone / missing-assignment spam as MindCraft differentiation.
7. AI weekly letter ≡ diagnosis (fluency cosplay).
8. Guaranteed ACT points or “never forgets” from any parent widget.
9. Controlling homework nag prompts as the product’s primary parent loop.

LXXXIII.2 Constructs

Construct	Research meaning	MindCraft analogue	Failure mode
Information friction	Parents hold biased beliefs about effort/progress	Optimistic “she’s fine” vs dated Map	Leave bias untouched OR flood noise
Belief correction	Accurate cumulative facts update parent priors	Last-proof age; returns due; probe rates	Relative-rank shame as “correction”
Actionable message	Specific, improvable next step	One mission CTA / return due	Vague “encourage more math”
Autonomy support	Parent involvement that supports student agency	Share return + invite student ownership	Controlling nag / do-it-for-them
MoC surface	Memory of Competence made legible	Proof age + delayed probe pass	Tonight accuracy as MoC
SAFE-PDASH	Parent UX doctrine for honest MoC	This chapter	Surveillance / fireworks portal

Operational definition (HYPOTHESIS): A parent dashboard is *honesty-complete* when (a) every competence claim shows last-proof age and proof type (blocked vs delayed/transfer), (b) at least one next action is a scheduled return or format/bridge mission — not “open the app for XP,” (c) comparative classmate ranks are off by default, (d) affect/biometric data never appear, (e) copy frames returns as protecting earned competence (SAFE-FORGET), and (f) the surface never claims permanent mastery or guaranteed scores.

LXXXIII.3 Parents systematically misbelieve — correction is load-bearing

FACT (effort/grade bias): Bergman (2021, *Journal of Political Economy*, 129(1), 286–322, doi:10.1086/711410) — field experiment providing biweekly detailed information on missed assignments and grades; parents hold **upwardly biased** beliefs about child effort; information attenuates bias and raises achievement via more accurate beliefs and lower monitoring costs (persuasion-game framing).

FACT (high-frequency automated alerts): Bergman & Chan (2021, *Journal of Human Resources*, 56(1), 125–158, doi:10.3368/jhr.56.1.1118-9837R1) — weekly automated texts on missed assignments, grades, and absences reduced course failures (~27%) and raised attendance (~12%) at very low variable cost; larger effects for below-median GPA and high-school students; **no** significant state-test lift — honesty about process metrics ≠ automatic standardized-score miracle.

FACT (cumulative absences misbelief): Rogers & Feller (2018, *Nature Human Behaviour*, 2, 335–342, doi:10.1038/s41562-018-0328-1) — personalized reports of *total accumulated* absences reduced chronic absenteeism (~10%+ in best arms) partly by correcting parents' underestimates; relative classmate-comparison framing did **not** add impact in their design — social-comparison “you’re worse than peers” is not the active ingredient.

Commercial implication: Sell *dated, specific competence facts* (proof age, return due, which gap) — not relative shame leagues. Do not market parent alerts as ACT-point engines (SAFE-EXAM / SAFE-DURABLE).

Kill: Classmate rank as default parent “truth.”

Survive: Cumulative, personalized, actionable honesty (last proof; next return).

LXXXIII.4 Improvement-framed, light-touch messages beat praise theater

FACT (teacher→parent weekly notes): Kraft & Rogers (2015, *Economics of Education Review*, 47, 49–63, doi:10.1016/j.econedurev.2015.04.001) — weekly individualized teacher messages raised credit-earning probability (~6.5 pp; ~41% reduction in failing to earn credit in their credit-recovery setting); effects driven largely by preventing dropout; **improvement-focused** messages showed the largest point estimates vs positive-continue messages (study underpowered to fully separate arms).

HYPOTHESIS (MindCraft mapping): Parent push copy should prefer one improvable object (“Return due in 3 days: linear equations — delayed mix”) over ego praise (“Amazing streak!”) or vague warmth. Improvement ≠ shame: name the skill/vessel/bridge, not the child’s character.

Kill: Streak/XP push as parent engagement.

Wound: Pure positivity digests that leave beliefs untouched.

Survive: Single improvement CTA tied to Map evidence.

LXXXIII.5 Involvement type matters — control is not support

FACT (homework involvement meta): Xu, Guo, Feng, Ma, Zhang, Núñez & Fan (2024, *Psicothema*, 36(1), 1–14, doi:10.7334/psicothema2023.92) — three-level meta-analysis (28 studies, 252 effects, N≈378k): overall parental homework involvement weakly **negatively** related to achievement ($r \approx -0.064$); **autonomy support** positively related; content support, parental control, frequency, and mixed dimensions largely unrelated or not helpful.

HYPOTHESIS: A dashboard that invites “sit beside and correct every item” or nag-frequency maximization recreates controlling involvement. Prefer surfaces that help parents *structure time and witness returns* while the student owns the attempt (ties SAFE-STRUCTURE bandwidth hygiene; SAFE-REWARD anti-controlling rewards).

Kill: Parent-as-second-tutor homework police UX.

Survive: Autonomy-supportive prompts (“Ask what felt hard on the return — don’t solve it for them”).

LXXXIII.6 Tech-mediated engagement is not magic — design the channel

FACT (systematic review caution): See, Gorard, Lu, Dong & Siddiqui (2021, *Educational Research and Evaluation*, doi:10.1080/13803611.2021.1924791) — promising evidence concentrates on **phone/text/email** school–parent communication that is two-way, personalized, and positive/constructive; evidence that generic online portals/devices alone improve outcomes is weak or inconclusive; research quality often low.

Wound for MindCraft: Shipping a pretty ParentDashboard SPA is not, by itself, an achievement intervention. The active ingredients are belief-correcting facts + actionable improvement framing + low friction — not dashboard feature count.

SPECULATION: MindCraft’s wedge vs school gradebooks is MoC (proof age / delayed probes / format-bridge gaps), not faster missing-homework alerts. Competing on gradebook UX is a trap.

LXXXIII.7 Product surface — SAFE-PDASH stack

Reuse: SAFE-FORGET (age nodes; no fireworks); SAFE-DURABLE (legible MoC for WTP); SAFE-WTP (CBC arms); SAFE-COLD (no day-one greens); SAFE-FORMAT / SAFE-BRIDGE (name the real gap); SAFE-PRIVACY (no affect on parent view); SAFE-EXAM (dual rail labels); SAFE-COMPARE (ranks off).

HYPOTHESIS — minimum honest parent chrome:

Element	Must show	Must not
Proof strip	Last proof age; proof type; concept/format/bridge label	Permanent green / Identity Score™
Return agenda	Next 1–3 returns with due+why	Streak calendar as learning
One CTA	Open student’s due mission / share talk prompt	“Buy more AI minutes” as primary
Weekly digest	Δ proof age; returns completed; one improvement sentence	AI essay of empty praise
Exam rail	Prove vs learn labeled intensity	Cram board as readiness

Privacy	Progress + MoC only	Mood, voice, face, raw chat dumps
---------	---------------------	-----------------------------------

FOUNDER BELIEF: Parent WTP rises when the dashboard makes *decay risk and the return plan* legible without humiliating Maya — validate with PDASH-5 CBC vs fireworks portal and vs shame-rank portal.

LXXXIII.8 Competitive positioning

Pattern	Parent move	MindCraft response
School SIS portals	Grades / missing work	MoC + returns; don’t clone SIS
Duo / gamified apps	Streak screenshots to parents	Ban streak-as-learning for adults
AI tutor dashboards	Tonight mastery + chat summary	Aged evidence; verify claims
Test-prep portals	Score predictors / point guarantees	Dual-rail; no guarantees
Surveillance SKUs	Live wrong-answer feed	Aggregate MoC; student dignity

Competitive wedge (FOUNDER BELIEF): “See what still holds — and what’s due back” is sellable trust language against both fireworks and panic portals.

LXXXIII.9 Doctrine — SAFE-PDASH (provisional)

- Beliefs before chrome** — correct optimistic misbeliefs with dated competence facts, not relative shame.
- Proof age is the hero metric on the parent surface** — tonight-% is subordinate.
- One improvement CTA** — return / format / bridge mission, Kraft-style specificity.
- Autonomy support > control** — never optimize nag frequency or parent-as-solver.
- No fireworks / no Ebbinghaus Score™ / no Anxiety Score™** on parent view.
- Ranks off by default** — Rogers–Feller: totals beat relative framing for absences; don’t import league toxicity.
- Not a gradebook** — differentiate on MoC and diagnosis, not SIS feature parity.
- Copy:** “Last proved 18 days ago — return due before the exam.” Never “mastered forever” / “behind her class” / guaranteed points.

Confidence: High that parents hold upwardly biased effort/progress beliefs and that information can move outcomes (Bergman 2021; Bergman & Chan 2021; Rogers & Feller 2018). High that improvement-framed light-touch messages can matter (Kraft & Rogers 2015). High that autonomy-supportive involvement beats controlling homework help on average (Xu et al. 2024). Medium that MindCraft’s MoC framing wins parent WTP vs fireworks and vs shame (needs PDASH-* / ties WTP). High that streak/XP/surveillance/affect on parent portal are kills. Medium that portal-alone engagement without push facts underperforms (See et al. review caution).

LXXXIII.10 Experiments

ID	Question	Design	Primary
PDASH-1	Proof-age strip + return CTA vs tonight-% fireworks parent home	A/B	Parent open→student return start; 7d probe; cancel
PDASH-2	Improvement digest vs positivity-only digest	A/B	Return completion; parent–child talk prompt use
PDASH-3	Absolute proof-age facts vs classmate rank chrome	A/B	Trust; stigma; Maya anxiety proxy; retention
PDASH-4	Autonomy-support talk card vs “check their work tonight” nag	A/B	Controlling-involvement self-report; abandon
PDASH-5	Parent CBC: honest MoC portal vs fireworks vs shame-rank	CBC	WTP; trust; stigma (SAFE-WTP)
PDASH-QUAL	10 parents: what would make you trust a math app report?	Qual	Copy; shame triggers; desired actions

Falsifier: Proof-age honesty raises trust surveys but lowers renewal vs fireworks → keep honesty in product doctrine; change *framing/dose*, do not restore permanent-mastery lies.

Falsifier: Improvement digests increase conflict/controlling homework without return gains → soften tone; keep student-owned CTA; do not add live wrong-answer feeds.

Falsifier: Rank chrome wins short-run opens with trust cost → ban as default (SAFE-COMPARE).

Pre-register: PDASH-* before any “parent portal / live progress / Family Leaderboard” campaign (SAFE-LABMETA).

LXXXIII.11 So what for MindCraft commercially

- **Copy:** “See what still holds — and what’s due back.” Never streak-for-parents, forever-mastered, classmate shame, or point guarantees.
- **Product:** Parent home = proof strip + return agenda + one improvement CTA; weekly digest optional; ranks/affect off.
- **Positioning:** Against SIS clones, Duo screenshot culture, AI mastery letters, and panic surveillance SKUs; for inspectable MoC.
- **Metric:** Parent open→return start, returns completed after digest, proof-age CBC WTP; demote parent DAU on fireworks and nag frequency.
- **Kill list:** Tonight-% NS; shame ranks; live wrong-answer stalk; Anxiety Score™; gradebook cosplay; controlling nag loop.

- **Growth:** ACT parents already fear silent fade — name proof age, sell the return plan (SAFE-FORGET × SAFE-WTP).
 - **Vision:** Family trust is part of identity transformation — adults witness competence honestly without colonizing the struggle.
-

References (verified)

- Bergman, P. (2021). Parent-child information frictions and human capital investment: Evidence from a field experiment. *Journal of Political Economy*, 129(1), 286–322. <https://doi.org/10.1086/711410>
 - Bergman, P., & Chan, E. W. (2021). Leveraging parents through low-cost technology: The impact of high-frequency information on student achievement. *Journal of Human Resources*, 56(1), 125–158. <https://doi.org/10.3368/jhr.56.1.1118-9837R1>
 - Kraft, M. A., & Rogers, T. (2015). The underutilized potential of teacher-to-parent communication: Evidence from a field experiment. *Economics of Education Review*, 47, 49–63. <https://doi.org/10.1016/j.econedurev.2015.04.001>
 - Rogers, T., & Feller, A. (2018). Reducing student absences at scale by targeting parents’ misbeliefs. *Nature Human Behaviour*, 2, 335–342. <https://doi.org/10.1038/s41562-018-0328-1>
 - See, B. H., Gorard, S., Lu, B., Dong, L., & Siddiqui, N. (2021). A systematic review of the impact of technology-mediated parental engagement on student outcomes. *Educational Research and Evaluation*. <https://doi.org/10.1080/13803611.2021.1924791>
 - Xu, J., Guo, S., Feng, Y., Ma, Y., Zhang, Y., Núñez, J. C., & Fan, H. (2024). Parental homework involvement and students’ achievement: A three-level meta-analysis. *Psicothema*, 36(1), 1–14. <https://doi.org/10.7334/psicothema2023.92>
-

Part LXXXIV — Tutor Talk-Ratio Instrumentation

Chapter status: Living evidence + HITL telemetry brief — Researcher tick 2026-08-05

Primary question: How should MindCraft instrument *who talks, who constructs, and who waits* in human (and AI-assisted) tutoring so prompt>pour becomes measurable ops fidelity — without turning silence into a vanity North Star?

Owners: Tutor ops / HITL QA · Product (session telemetry) · Workforce coaching · Brand/copy · Red Team

Commercial job: Ship a **SAFE-TALK** doctrine: treat talk-ratio as a *fidelity signal for student construction* (scaffolding density, wait time, student why-before-pour), never as a Talk Ratio Score™, equal-time dogma, or “our tutors are quiet” marketing costume.

Builds on: Parts XL (SAFE-SE), LI (SAFE-DP), LIV (SAFE-AAR), LXVIII (SAFE-HITL), LXXI (SAFE-TUTORGRAIN), LXXV (SAFE-WORKFORCE), LXXIX (SAFE-FBTIME). Product seam: HITL QA telemetry that makes “prompt>pour” auditable without surveilling affect.

LXXXIV.1 Why this chapter exists

SAFE-HITL already requires talk ratio favoring student construction over tutor monologue — but the ops stack still lacks a *shippable instrumentation contract*. Without one, coaches grade vibes (“warm energy”), tutors lecture to fill silence anxiety, and marketing invents “Socratic AI” claims nobody can audit. Workforce fidelity (SAFE-WORKFORCE) also needs observable session features beyond tenure and star ratings.

FOUNDER BELIEF under audit: Parents and districts will trust MindCraft’s human loop more when we can show *prompt density + student constructive talk + wait discipline* correlating with solo transfer — not when we advertise “50/50 talk balance” or tutor silence as virtue.

Claims we refuse as doctrine:

1. Raw tutor-talk-% or student-talk-% as North Star / Talk Ratio Score™.
2. Equal 50/50 airtime as pedagogical law (context and ICAP mode matter).
3. Silence theater (long pauses with no elicitation) as “Socratic.”
4. AI monologue length dressed as “dialogue” via fake student-token counts.
5. “Got it?” check-understanding as student talk (Graesser frame step 5 ≠ construction).
6. Warmth / energy / Ivy brand as substitutes for talk instrumentation (SAFE-HITL / GRAIN kills).
7. Affect/voice emotion AI layered on talk metrics (SAFE-PRIVACY).
8. Bloom 2-sigma marketing from “more student talk.”

LXXXIV.2 Constructs

Construct	Research meaning	MindCraft analogue	Failure mode
IRF / IRE	Teacher Initiation–Response–Feedback/Evaluation triad	Default pour loop if F = lecture	Closed IRF forever
5-step tutoring frame	Tutor Q → answer → short FB → collaborative improve → check	Session skeleton	Skip improve; pour explain
Wait time 1 / 2	Pause after ask / after student turn (~3s+)	Coachable micro-behavior	Sub-1s snatch-back
ICAP mode	Passive < Active < Constructive < Interactive	Student why / soft-wrong / join	Words ≠ constructive
Prompt>pour	Scaffold/elicitor before explanation dump	HITL playbook law	Explain-first brand
SAFE-TALK	Instrumentation doctrine for construction fidelity	This chapter	Ratio vanity / silence cosplay

Operational definition (HYPOTHESIS): A session is *talk-instrumented* when MindCraft can report, per session sample: (a) tutor vs student floor time and turn counts, (b) **prompt density** (eliciting moves / tutor turns), (c) **student constructive share** (why/self-explain/soft-wrong articulations — not yes/no or “got it”), (d) median wait time 1 and 2,

(e) Map-brief opened before first explain (map_brief_open), and (f) at least one solo attempt evidence after scaffolding — *without* claiming any single ratio proves mastery.

LXXXIV.3 Classroom IRF shows why “teacher talks most” is the default failure

FACT (IRF as classroom grammar): Sinclair & Coulthard (1975) described the Initiation–Response–Feedback exchange as a fundamental unit of classroom discourse — teacher-led, often evaluative, structurally favoring adult control of the floor.

FACT (talk imbalance persists in IRF-heavy rooms): Contemporary discourse analyses continue to find teacher-dominated floor time under IRF (often ~60%+ teacher share in coded secondary lessons) — illustrative of the pattern, not a universal constant.

Commercial implication: If MindCraft tutors recreate classroom IRF with longer lectures, the “human advantage” collapses into expensive school-talk. Instrument *pour-IRF* (short student token → explanation dump) vs *scaffold-IRF* (extended student construction → brief feedback → next prompt).

Kill: Marketing “conversation-based tutoring” without measuring construction.

Survive: IRF as descriptive grammar; redesign the Feedback move toward elicit-and-verify.

LXXXIV.4 One-to-one tutoring is collaborative — but tutors still drive the frame

FACT (naturalistic tutoring dialogue): Graesser, Person & Magliano (1995, *Applied Cognitive Psychology*, 9(6), 495–522, doi:10.1002/acp.2350090604) — normal (unskilled) tutors show collaborative multi-turn problem solving; a pervasive **5-step frame**: tutor asks → student answers → short feedback → collaborative improvement of the answer → tutor checks understanding. Median turns per tutor-initiated question were far above one (multi-turn collaborative answers). Deep diagnosis of idiosyncratic knowledge deficits was often underdeveloped.

FACT (questioning density vs classroom): Graesser & Person (1994, *American Educational Research Journal*, 31(1), 104–137, doi:10.3102/00028312031001104) — students ask far more questions in tutoring than in classrooms (order-of-magnitude higher rates in their corpora), yet tutors still control agenda via curriculum scripts and dialogue frames.

HYPOTHESIS (MindCraft mapping): Success is not “maximize student words.” It is *keep the collaborative improve step alive* (step 4) and refuse to skip from short feedback into monologue. Check-understanding (“got it?”) is the weakest step for identity/FEI — prefer answer evidence / soft-wrong / Map probe (SAFE-HITL).

Kill: “Got it?” as mastery talk.

Wound: Counting all student tokens equally (procedural yes ≈ constructive why).

Survive: Multi-turn collaborative improvement as the auditable unit.

LXXXIV.5 Construction beats tutor explanation volume — Chi’s interaction test

FACT (prompt-constrained tutoring): Chi, Siler, Jeong, Yamauchi & Hausmann (2001, *Cognitive Science*, 25(4), 471–533, doi:10.1207/s15516709cog2504_1) — analyses of human tutoring support tutor-, student-, and interaction-centered contributions; critically, when tutors were **suppressed from giving explanations and feedback** and restricted to prompting, students learned **as effectively** as in freer tutoring, attributed to deeper/more scaffolding episodes and greater student control (with reading-ability limits on what self-study recovered).

FACT (ICAP hierarchy): Chi & Wylie (2014, *Educational Psychologist*, 49(4), 219–243, doi:10.1080/00461520.2014.965823) — engagement modes Passive < Active < Constructive < Interactive predict increasing learning; overt behaviors must map to knowledge-change processes — “active” clicking/speaking is not automatically constructive.

Commercial implication: SAFE-HITL’s prompt>pour is not a vibe — it is Chi-compatible. **Instrument prompts that elicit construction**, not tutor silence. An AI that “talks less” while leaving the student passive fails ICAP; a near-peer who prompts why-before-reveal can win without Ivy monologue (SAFE-TUTORGRAIN).

Kill: Tutor-talk-% minimization as the product KPI.

Survive: Prompt density × constructive student share as fidelity features feeding tutor_fidelity_*.

LXXXIV.6 Wait time is cheap, coachable, and repeatedly evidenced

FACT (baseline snatch-back): Rowe (1974, *Journal of Research in Science Teaching*, 11(2), 81–94, doi:10.1002/tea.3660110202; see also Rowe 1986, *Journal of Teacher Education*, 37(1), 43–50, doi:10.1177/002248718603700110) — teachers typically wait <1 **second** after questions and after student responses; extending wait time 1 and wait time 2 toward ~3+ seconds associates with longer, more complex, more speculative student responses and broader participation.

FACT (review): Tobin (1987, *Review of Educational Research*, 57(1), 69–95, doi:10.3102/00346543057001069) — synthesizing classroom wait-time studies, longer pauses support higher-quality discourse when implemented with accountable responding norms (not isolated stopwatch theater).

HYPOTHESIS: Median WT1/WT2 on sampled sessions is a **coachable leading indicator** of pour-vs-prompt culture — especially for anxious Mayas who need think-space before soft-wrong (ties SAFE-DD / SAFE-EXPOSE / SAFE-FBTIME).

Kill: Stopwatch cosplay without elicitation quality.

Survive: Banded wait targets (e.g., WT1/WT2 medians ≥3s on open prompts) inside FEI QA rubrics.

LXXXIV.7 Dialogic teaching ≠ talk-count theater

FACT (large dialogic-teaching trial): Alexander (2018, *Research Papers in Education*; see also Alexander’s dialogic-teaching programme documentation) — developing dialogic teaching is a multi-principle pedagogy (collective, reciprocal, supportive, cumulative, purposeful), not a single talk-ratio knob. Implementation requires

repertoire and epistemology, not merely “more student talk.”

Wound for MindCraft: Importing “dialogic” as brand language without Map-grounded cumulation (SAFE-ONTOLOGY / SAFE-BRIDGE) recreates supportive chat that never remediates joins.

SPECULATION: Online one-to-one math corpora will keep finding tutor-dominant initiation with sparse conceptual talk unless products force scaffold ladders — Map brief + soft-wrong gates should move sequences more than a talk-% dashboard alone.

LXXXIV.8 Product surface — SAFE-TALK instrumentation stack

Reuse: SAFE-HITL (prompt>pour; Map brief; FEI QA); SAFE-WORKFORCE (fidelity@tenure); SAFE-SE (student-generated why); SAFE-FBTIME (micro-delay before reveal); SAFE-AAR (improve-how close); SAFE-PRIVACY (no emotion AI on voice); SAFE-GRAIN (hire by fidelity not Ivy pour).

HYPOTHESIS — minimum honest session telemetry (sampled + consented):

Signal	Definition	Use	Must not
tutor_floor_share	Tutor speaking time / (tutor+student)	Drift monitor	Sole NS
student_constructive_share	Why/SE/soft-wrong turns / student turns	Primary fidelity	Count “yeah/got it”
prompt_density	Eliciting prompts / tutor turns	Coach target	Reward vague “what do you think?” spam
wait_ms_p50	Median WT1 and WT2	Coach micro-skill	Shame tutors live mid-session
pour_streak_max	Longest consecutive tutor explain turns	Detect monologue	Ban all explanation
map_brief_open	Brief before first explain	HITL hard gate	Trait-label briefs
annotate_before_reveal	Why before key/coach	SE×FBTIME	AI monologue first

Ops rule (FOUNDER BELIEF): Coaches review **clips where pour_streak_max is high and constructive share is low**, not leaderboards of who talked least. Parent/district packets may show *aggregate fidelity bands* (“sessions coached for prompt-first practice”) — never per-child Talk Ratio Score™.

AI sessions: Apply the same schema to Solver/coach turns — fluent AI floor share with low student constructive share is a **sycophancy/pour risk** (Part XXXIII), not “personalized dialogue.”

LXXXIV.9 Competitive positioning

Pattern	Talk move	MindCraft response
Classroom / Zoom school help	IRF pour	Instrument scaffold improve-step
ChatGPT tutor	Fluent monologue	Ban fluency≡dialogue; require student why gates
Marketplace “Ivy explainers”	Expert pour as quality	Hire/coach on prompt density (GRAIN×HITL)
Duo / app streaks	Engagement without discourse	Irrelevant; don’t import streak as talk proxy
“Socratic AI” ads	Brand claim	Show prompt density + constructive share or stay quiet

Competitive wedge (FOUNDER BELIEF): “We coach tutors (and AI) to make students construct — and we can show the tape metrics” beats explain-first marketplaces and unmeasured Socratic cosplay.

LXXXIV.10 Doctrine — SAFE-TALK (provisional)

1. **Construction > airtime** — student constructive share and prompt density outrank raw talk-%.
2. **Prompt>pour is measurable** — elicit before explain; cap pour streaks in QA.
3. **Wait time is a micro-skill** — band WT1/WT2; coach snatch-back.
4. **“Got it?” ≠ evidence** — prefer soft-wrong, annotated why, Map probe.
5. **No Talk Ratio Score™ / silence NS** — fidelity feature for coaching, not vanity KPI.
6. **No emotion AI on talk** — acoustic floor time ok; affect inference banned (SAFE-PRIVACY).
7. **Same schema for human and AI** — fluent pour fails either channel.
8. **Copy:** “Prompt first — then prove.” Never “our tutors talk half the time” / Bloom-from-talk / “Socratic guaranteed.”

Confidence: High that IRF defaults to teacher control (Sinclair–Coulthard) and that naturalistic tutoring uses multi-turn collaborative frames under tutor agenda control (Graesser et al. 1995; Graesser & Person 1994). High that student construction can carry learning when tutor explanation is constrained (Chi et al. 2001) and that ICAP ranks engagement modes (Chi & Wylie 2014). High that wait-time extension changes discourse quality (Rowe; Tobin). Medium that MindCraft’s telemetry bundle predicts solo transfer (needs TALK-*). High that talk-% NS / silence theater / got-it checks are kills. Medium that dialogic-brand copy without Map cumulation underperforms (Alexander).

LXXXIV.11 Experiments

ID	Question	Design	Primary
TALK-1	Prompt-first playbook + wait coaching vs explain-first business-as-usual	Tutor RCT / stepped-wedge	student_constructive_share; solo_transfer_pass; pour_streak_max

TALK-2	Live talk dashboard for coaches vs weekly clip QA only	Ops A/B	tutor_fidelity_30d; coach time; tutor reactance
TALK-3	Soft-wrong + annotate-before-reveal gate vs free talk	Session A/B	SE quality; transfer; time-on-task
TALK-4	AI coach with prompt-density cap vs unconstrained monologue	Product A/B	Constructive share; hallucination repair need; abandon
TALK-5	Parent CBC: “coached prompt-first tutors” vs “Ivy explainers” vs “Socratic AI”	CBC	WTP; trust (ties SAFE-WTP / GRAIN)
TALK-QUAL	10 tutors: what makes silence feel unsafe?	Qual	Snatch-back drivers; playbook friction

Falsifier: Raising constructive share / wait time without transfer lift → keep as process hygiene; do not market talk metrics as learning proof; retarget prompts to Map gaps.

Falsifier: Talk dashboards increase tutor reactance / fake short student turns → remove live shame UI; keep sampled clip QA.

Falsifier: AI prompt-caps raise abandon without transfer gain → soften cap; keep why-before-reveal on soft-wrong paths only.

Pre-register: TALK-* before any “Socratic / talk-balanced tutoring” campaign (SAFE-LABMETA).

LXXXIV.12 So what for MindCraft commercially

- **Copy:** “Prompt first — then prove.” Never Talk Ratio Score™, 50/50 dogma, silence theater, or Bloom-from-talk.
- **Product:** Sampled session telemetry — constructive share, prompt density, wait bands, pour-streak — wired into HITL QA and workforce fidelity.
- **Positioning:** Against Ivy pour marketplaces and unmeasured Socratic AI; for auditable construction.
- **Metric:** student_constructive_share, prompt_density, wait_ms_p50, pour_streak_max feeding tutor_fidelity_*; demote raw talk-% and energy ratings.
- **Kill list:** Talk-% NS; got-it≡mastery; emotion-AI talk scoring; AI monologue-as-dialogue ads.
- **Growth:** District/parent trust packets can cite coaching-for-construction without biometric theater (SAFE-PROCURE × SAFE-HITL).
- **Vision:** Identity transformation requires the student to *speak the math into being* — instrumentation exists to protect that, not to gamify silence.

References (verified)

- Alexander, R. J. (2018). Developing dialogic teaching: Genesis, process, trial. *Research Papers in Education*. <https://doi.org/10.1080/02671522.2018.1481140>
- Chi, M. T. H., Siler, S. A., Jeong, H., Yamauchi, T., & Hausmann, R. G. (2001). Learning from human tutoring. *Cognitive Science*, 25(4), 471–533. [https://doi.org/10.1016/S0364-0213\(01\)00044-1](https://doi.org/10.1016/S0364-0213(01)00044-1)
- Chi, M. T. H., & Wylie, R. (2014). The ICAP framework: Linking cognitive engagement to active learning outcomes. *Educational Psychologist*, 49(4), 219–243. <https://doi.org/10.1080/00461520.2014.965823>
- Graesser, A. C., & Person, N. K. (1994). Question asking during tutoring. *American Educational Research Journal*, 31(1), 104–137. <https://doi.org/10.3102/00028312031001104>
- Graesser, A. C., Person, N. K., & Magliano, J. P. (1995). Collaborative dialogue patterns in naturalistic one-to-one tutoring. *Applied Cognitive Psychology*, 9(6), 495–522. <https://doi.org/10.1002/acp.2350090604>
- Rowe, M. B. (1974). Wait-time and rewards as instructional variables, their influence on language, logic, and fate control: Part one—Wait-time. *Journal of Research in Science Teaching*, 11(2), 81–94. <https://doi.org/10.1002/tea.3660110202>
- Rowe, M. B. (1986). Wait time: Slowing down may be a way of speeding up! *Journal of Teacher Education*, 37(1), 43–50. <https://doi.org/10.1177/002248718603700110>
- Sinclair, J. McH., & Coulthard, R. M. (1975). *Towards an analysis of discourse: The English used by teachers and pupils*. Oxford University Press.
- Tobin, K. (1987). The role of wait time in higher cognitive level learning. *Review of Educational Research*, 57(1), 69–95. <https://doi.org/10.3102/00346543057001069>

Part LXXXV — Story-World Cognitive Load vs Identity Payoff

Chapter status: Living evidence + narrative-wrap design brief — Researcher tick 2026-08-05

Primary question: When does a story world *pay for its working-memory tax* with identity/belonging/transfer — and when is it seductive load that hurts learning while looking branded?

Owners: Story/content ops · Product (Solver / chapter / practice stems) · Learning science · Brand/copy · Equity (Part XXXVI) · Red Team

Commercial job: Ship a **SAFE-STORYLOAD** doctrine: narrative wrap is allowed only when it encodes the math structure, keeps language load banded, and graduates via learning+belonging experiments — never as franchise lore, immersion theater, or “story = engagement = learning.”

Builds on: Parts XXVI (CLT), XXVIII (invention-story / HIST), XXXVI (equity of worlds), XLVIII (SAFE-TRANSFER; story≠far-transfer), LXIV (SAFE-ELL language vs math load), LXXXII (SAFE-FORMAT). Product seam: story-module / stem reskin must freeze the math and wrap it — WORLD_VISION — under load budgets.

LXXXV.1 Why this chapter exists

WORLD_VISION and brand thesis bet that math-inside-a-world beats naked drill for Maya’s identity. Parts XXVIII and XXXVI already constrain *whose* story and whether invention lore is seductive trivia. They do not settle the **cognitive-load bargain**: every extra character beat, lore paragraph, and decorative cutscene consumes working

memory that might have built the schema.

FOUNDER BELIEF under audit: Story-first is a defensible differentiator *if and only if* the wrap reduces threat / raises situational interest / encodes structure — without exceeding WM for novices and anxious Mayas.

Claims we refuse as doctrine:

1. Immersion / lore depth as North Star or Story Engagement Score™.
2. “Story always helps learning” / narrative $\equiv$ transfer (SAFE-TRANSFER kill).
3. Seductive franchise trivia as “world building.”
4. Decorative art or soundtrack $\equiv$ FormatId / pedagogy (SAFE-FORMAT).
5. Long pre-attempt lore dumps before first micro-win.
6. ELL/language-heavy wraps that conflate reading load with math ability (SAFE-ELL).
7. Token skin-tone diversity as load justification (XXXVI).
8. Story feature shipping without HIST / EQ / STORYLOAD experiment plan.

LXXXV.2 Constructs

Construct	Research meaning	MindCraft analogue	Failure mode
Intrinsic load	Element interactivity of the math itself	Ontology ingredient / problem structure	Cannot “story away” hard joins
Extraneous load	Presentation that does not aid schema	Lore dump, split lore/diagram, BGM	Brand chrome tax
Seductive details	Interesting but irrelevant add-ons	Side quests, cute facts, unrelated plot	Interest without learning
Coherence principle	Exclude extraneous words/pictures/sounds	Tight wrap; silence optional	Multimedia clutter
Situational → individual interest	Triggered then maintained interest	World re-entry + role agency	Trigger without maintain
Structure-encoding wrap	Narrative makes the invariant memorable	Story-first question design	Scene then unrelated textbook ask
SAFE-STORYLOAD	Load-budgeted identity wrap	This chapter	Immersion NS / lore hero

Operational definition (HYPOTHESIS): A story module is *load-honest* when MindCraft can show: (a) the mathematical invariant the wrap encodes (structure test), (b) pre-attempt narrative word/screen budget within a band by level and language status, (c) time-to-first-attempt under a cap, (d) learning outcomes on retention/transfer not worse than matched plain stem (non-inferiority), and (e) optional belonging/interest lift — *without* claiming immersion proves identity.

LXXXV.3 CLT is the physics; story is optional mass

FACT (CLT core): Sweller (1988, *Cognitive Science*, 12(2), 257–285, doi:10.1207/s15516709cog1202_4) — problem-solving search can impose heavy load that impedes schema acquisition; instructional design must respect limited working memory.

FACT (intrinsic vs extraneous tradition): Chandler & Sweller and later CLT reviews distinguish load from the material’s element interactivity vs load from how material is organized/presented (see also Sweller, van Merriënboer & Paas, 1998, *Educational Psychology Review*, 10, 251–296). Extraneous load is the designer’s fault; intrinsic load is managed by sequencing and prior knowledge — not by prettier lore.

Commercial implication: MindCraft cannot sell “epic world” as pedagogy. The world is a *presentation choice* that must clear an extraneous-load budget. If the wrap does not redirect attention to schema-relevant relations, it is cost.

Kill: Story immersion as learning science.

Survive: CLT as hard constraint on narrative chrome.

LXXXV.4 Seductive details and coherence — the default kill for franchise lore

FACT (seductive details experiments): Harp & Mayer (1998, *Journal of Educational Psychology*, 90(3), 414–434, doi:10.1037/0022-0663.90.3.414) — undergraduates reading science text *with* interesting but irrelevant details recalled fewer main ideas and generated fewer transfer solutions than those without; diversion toward inappropriate organizing schemas is a leading account (not mere “fun distraction”).

FACT (meta-analyses): Rey (2012, *Educational Research Review*, 7(3), 216–237, doi:10.1016/j.edurev.2012.05.003) — across ~39 experimental effects, seductive details show small-to-medium negative effects on retention and medium negative effects on transfer; moderators exist (time pressure, detail type, domain) but mean direction is caution. Sundararajan & Adesope (2020, *Educational Psychology Review*, 32, 707–734, doi:10.1007/s10648-020-09522-4) — later meta likewise supports a negative seductive-details pattern with coherence-relevant moderators.

FACT (coherence principle): Mayer’s multimedia coherence principle — people learn more deeply when extraneous words, pictures, and sounds are excluded (handbook statements; e.g., Moreno & Mayer, 2000, *Journal of Educational Psychology*, 92(1), 117–125, doi:10.1037/0022-0663.92.1.117 on irrelevant sounds). Cute BGM, floating stickers, and lore sidebars are not free.

HYPOTHESIS (MindCraft mapping): Pre-question character backstory that does not encode the invariant is a textbook seductive detail. Structure-encoding narrative (the plot *is* the proportion / the journey *is* the function composition) is not “irrelevant” — it is germane packaging — **if** it stays short and co-located with the ask (no split-attention lore panel vs equation).

Kill: Lore-first chapter opens; museum trivia; soundtrack-as-engagement.

Wound: “Interestingness” as A/B primary without retention/transfer.

Survive: Coherence-first wrap; seductive-detail gate in content ops (aligns XXXVI structure test).

LXXXV.5 Interest development — the identity-side payoff (conditional)

FACT (four-phase interest): Hidi & Renninger (2006, *Educational Psychologist*, 41(2), 111–127, doi:10.1207/s15326985ep4102_4) — interest develops from triggered situational → maintained situational → emerging individual → well-developed individual interest; supports that mismatch phase can impede.

FACT (interest as context, not magic): Renninger, Ewen & Lasher (2002, *Learning and Instruction*, 12(4), 467–490, doi:10.1016/S0959-4752(01)00012-3) — inserting well-developed individual-interest contexts into text/word problems can scaffold meaning-focus, but can also *mask* difficulties or feel harder; teacher/product support still matters. Students often respond more to difficulty than to context alone.

FACT (personalization mixed): Walkington (2013, *Journal of Educational Psychology*, 105(4), 932–945, doi:10.1037/a0031882) — interest-personalized algebra contexts can improve performance/learning relative to generic problems in adaptive tutors (depth of connection matters). Van de Weijer-Bergsma & Van der Ven (2021, *Learning and Individual Differences*, 87, 101982, doi:10.1016/j.lindif.2021.101982) — two primary-school personalization studies found **no** effect of personalization on performance, enjoyment, or cognitive load (though enjoyment and load predicted performance). Personalization is not an automatic win.

HYPOTHESIS: MindCraft’s identity payoff is less “pick Maya’s favorite fandom” and more *maintained situational interest via agency in a coherent world* (role, re-entry, recognition of competence) tied to FEI attempts — not surface name-swap personalization.

Kill: Surface personalization / fandom skin as product hero.

Survive: Interest as a *possible* mediator of persistence (retry_120s) when wrap encodes structure and stays in load band; always co-primary with transfer.

LXXXV.6 Language load, anxiety, and format — compound taxes

FACT / doctrine reuse: SAFE-ELL — separate language load from math load; heavy prose can false-flag ability. SAFE-DD / Ashcraft line — anxiety consumes WM; extraneous narrative tax hits threatened Mayas hardest. SAFE-FORMAT — decorative story art ≠ diagram FormatId.

HYPOTHESIS: For high-anxiety / ELL / low-prior students, pre-attempt narrative budgets must be *shorter*, not richer. Offer plain-stem toggle or “skip to math” without shame; never gate belonging on lore completion.

Commercial implication: Parent copy cannot say “immersive story worlds” without “math stays clear — story carries the idea, not the obstacle.” District reviewers will ask about reading demand (SAFE-PROCURE packet honesty).

LXXXV.7 Product surface — SAFE-STORYLOAD design contract

Reuse: XXVI attempt-grain; XXVIII HIST structure; XXXVI equity rubric; XLVIII name-the-hop (story world ≠ transfer world); LXIV linguistic modify without math dilute.

HYPOTHESIS — merge gates for story modules / themed stems:

Gate	Pass criterion	Fail →
------	----------------	--------

Structure encode	Writer names the invariant the narrative makes memorable	Cut or rewrite
Ask binding	Numbers/relations in scene match the frozen item (no scene→unrelated textbook ask)	Reskin fail
Pre-attempt budget	Words/screens/time-to-first-attempt within level×language band	Compress; offer skip
Coherence	No BGM/stickers/side lore required to solve	Delete extraneous
Non-inferiority learning	Retention/transfer $\geq$ plain stem within pre-reg margin (STORYLOAD-1)	Demote to optional flavor
Belonging optional lift	EQ/HIST metrics only after learning gate	Do not ship on vibes
Equity scan	XXXVI stereotype / plurality checks	Rewrite

UX rules (FOUNDER BELIEF):

1. First micro-attempt comes fast — lore after or interleaved as completion prompts, not walls of text.
2. “Skip story” is always available; never shame.
3. World reveal scales with mastery (WORLD_VISION Horizons) — novices see *less* lore, not more.
4. Exam dual-rail (SAFE-EXAM): prove mode uses exam-like stems; story wrap is learn-rail default, not timed ACT costume.

LXXXV.8 Competitive positioning

Pattern	Story move	MindCraft response
Naked drill apps	Zero wrap	Wrap only if structure+load gates pass
Edutainment lore	Franchise immersion	Coherence + seductive-detail kill
Generic word-problem skins	Surface personalization	Depth/agency > name-swap; null results possible
Khan culture-thin library	Minimal narrative	Narrative risk accepted <i>with</i> instrumentation
Brilliant delight puzzles	Aesthetic engagement	Delight subordinate to transfer; no immersion NS

Competitive wedge (FOUNDER BELIEF): “Stories that carry the math — without drowning working memory” beats both sterile banks and lore-bloated games.

LXXXV.9 Doctrine — SAFE-STORYLOAD (provisional)

1. **CLT budgets narrative** — extraneous lore is a tax; intrinsic math difficulty is not cured by plot.
2. **Structure-encoding or cut** — wrap must make the invariant memorable; else seductive detail.
3. **Coherence > immersion** — exclude irrelevant words/pictures/sounds; no Story Engagement Score™.
4. **Interest is conditional** — situational interest may aid persistence; personalization is mixed evidence; never primary KPI.
5. **Novices get less lore** — mastery unlocks world depth (Horizons), not day-one encyclopedias.
6. **Language ≠ math** — band prose; skip-story; ELL hygiene (SAFE-ELL).
7. **Story ≠ transfer** — still name hops; measure `transfer_pass` / `solo_transfer_pass` (SAFE-TRANSFER).
8. **Copy:** “The story carries the idea — the math stays clear.” Never immersion hero / lore-as-learning / guaranteed identity-from-world.

Confidence: High that seductive details and coherence violations can impair retention/transfer (Harp & Mayer; Rey; Sundararajan & Adesope; Moreno & Mayer). High that WM limits constrain instructional search (Sweller). Medium that MindCraft structure-encoding wraps will clear non-inferiority (needs STORYLOAD-). *Medium–low that wrap lifts belonging* and* learning jointly (HIST-EQ / EQ still open). High that immersion NS / lore dumps / surface fandom skins are kills. Medium that skip-story + mastery-gated reveal protects anxious/ELL segments.

LXXXV.10 Experiments

ID	Question	Design	Primary
STORYLOAD-1	Structure-encoding wrap vs plain stem (same items)	A/B	Retention; <code>transfer_pass</code> ; time-to-first-attempt; self-load rating
STORYLOAD-2	Pre-attempt lore long vs short budget	A/B	Learning; abandon; <code>retry_120s</code>
STORYLOAD-3	Skip-story available vs forced wrap	A/B	Use-rate; learning by anxiety/ELL segment
STORYLOAD-4	Mastery-gated world reveal vs day-one full lore	A/B	Learning; belonging; session length
STORYLOAD-5	Parent CBC: “clear story that carries the math” vs “immersive world” vs “no story / more practice”	CBC	WTP; trust (ties SAFE-WTP)
STORYLOAD-QUAL	10 Mayas: when does story feel like help vs homework reading?	Qual	Friction themes; skip motives

Falsifier: Wrap loses to plain stem on transfer beyond margin → demote story to optional cosmetic; keep brand elsewhere.

Falsifier: Short budget kills belonging with no learning gain → redesign agency/role, not add lore.

Falsifier: Skip-story users dominate and convert better → default plain; story opt-in.

Falsifier: Immersion CBC wins WTP but STORYLOAD-1 fails learning → do not ship immersion ads; dual-rail honest packaging only.

Pre-register: STORYLOAD-* / continue HIST-EQ before any “immersive story worlds improve learning” campaign (SAFE-LABMETA).

LXXXV.11 So what for MindCraft commercially

- **Copy:** “The story carries the idea — the math stays clear.” Never immersion NS, lore-as-pedagogy, or identity-from-world guarantees.
 - **Product:** Structure-encode + ask-bind + pre-attempt budget + skip-story + mastery-gated reveal as merge gates for story modules/stems.
 - **Positioning:** Against sterile banks *and* edutainment lore bloat; for load-honest narrative wrap.
 - **Metric:** Non-inferiority on transfer vs plain; time-to-first-attempt; optional belonging; demote Story Engagement / lore completion.
 - **Kill list:** Franchise trivia dumps; forced lore walls; BGM/sticker coherence violations; surface fandom personalization hero; story=transfer ads.
 - **Growth:** Parent/district packets show learning gates before world screenshots; equity rubric visible as craft not costume (XXXVI).
 - **Vision:** Identity transformation uses story as *meaning technology under CLT* — worlds reveal as competence grows, they do not replace schema construction.
-

References (verified)

- Chandler, P., & Sweller, J. (1991). Cognitive load theory and the format of instruction. *Cognition and Instruction*, 8(4), 293–332. https://doi.org/10.1207/s1532690xc0804_2
- Harp, S. F., & Mayer, R. E. (1998). How seductive details do their damage: A theory of cognitive interest in science learning. *Journal of Educational Psychology*, 90(3), 414–434. <https://doi.org/10.1037/0022-0663.90.3.414>
- Hidi, S., & Renninger, K. A. (2006). The four-phase model of interest development. *Educational Psychologist*, 41(2), 111–127. https://doi.org/10.1207/s15326985ep4102_4
- Moreno, R., & Mayer, R. E. (2000). A coherence effect in multimedia learning: The case for minimizing irrelevant sounds in the design of multimedia instructional messages. *Journal of Educational Psychology*, 92(1), 117–125. <https://doi.org/10.1037/0022-0663.92.1.117>
- Renninger, K. A., Ewen, L., & Lasher, A. K. (2002). Individual interest as context in expository text and mathematical word problems. *Learning and Instruction*, 12(4), 467–490. [https://doi.org/10.1016/S0959-4752\(01\)00012-3](https://doi.org/10.1016/S0959-4752(01)00012-3)
- Rey, G. D. (2012). A review of research and a meta-analysis of the seductive detail effect. *Educational Research Review*, 7(3), 216–237. <https://doi.org/10.1016/j.edurev.2012.05.003>
- Sundararajan, N., & Adesope, O. (2020). Keep it coherent: A meta-analysis of the seductive details effect. *Educational Psychology Review*, 32, 707–734. <https://doi.org/10.1007/s10648-020-09522-4>

- Sweller, J. (1988). Cognitive load during problem solving: Effects on learning. *Cognitive Science*, 12(2), 257–285. https://doi.org/10.1207/s15516709cog1202_4
 - Sweller, J., van Merriënboer, J. J. G., & Paas, F. G. W. C. (1998). Cognitive architecture and instructional design. *Educational Psychology Review*, 10(3), 251–296. <https://doi.org/10.1023/A:1022193728205>
 - Van de Weijer-Bergsma, E., & Van der Ven, S. H. G. (2021). Why and for whom does personalizing math problems enhance performance? Testing the mediation of enjoyment and cognitive load at different ability levels. *Learning and Individual Differences*, 87, 101982. <https://doi.org/10.1016/j.lindif.2021.101982>
 - Walkington, C. (2013). Using adaptive learning technologies to personalize instruction to student interests: The impact of relevant contexts on performance and learning outcomes. *Journal of Educational Psychology*, 105(4), 932–945. <https://doi.org/10.1037/a0031882>
-

Part LXXXVI — Competence Evidence as Social Proof

Chapter status: Living evidence + growth/copy brief — Researcher tick 2026-08-05

Primary question: What *witnessable competence artifacts* (solo-transfer / delayed proof) should MindCraft use as social proof for parents, tutors, and Maya — instead of vanity counts, vague praise, or mastery fireworks?

Owners: Growth/copy · Product (shareable MoC / portfolio surfaces) · Parent UX · Brand · Red Team

Commercial job: Ship a **SAFE-PROOF** doctrine: social proof must be *dated, specific, solo-capable competence evidence* others can inspect — never crowd size, star ratings, Identity Score™, or “she’s a math person now” testimonials without a named hop.

Builds on: Parts XXV (self-efficacy), XXVII/LIX (parent WTP), XLII (SAFE-COMPARE), XLVIII (SAFE-TRANSFER), LXI (SAFE-IDMEASURE), LXXVI (SAFE-FORGET), LXXVIII (SAFE-DURABLE), LXXXIII (SAFE-PDASH). Product seam: shareable solo_transfer_pass / proof-age artifacts that correct buyer uncertainty without shame leagues or guaranteed-points case studies.

LXXXVI.1 Why this chapter exists

SAFE-PDASH told parents *what to show* on a dashboard (proof age, returns, one autonomy CTA). SAFE-WTP told us to price legible MoC. Growth still defaults to the edtech social-proof stack: “Trusted by 50k families,” five-star walls, Ivy-tutor logos, AI-written “Maya feels confident!” quotes. That stack sells *affiliation and affect*, not *competence*. MindCraft’s identity thesis needs a different proof object: a third party (parent, tutor, skeptic buyer) can *witness* that the student solved a delayed / format-hop / join item **without** Solver open — and the claim names the concept/bridge/format and proof age.

FOUNDER BELIEF under audit: Parents and Maya buy (and stay) when social proof looks like *inspectable competence*, not herd size — if artifacts stay private-by-default, autonomy-supportive, and never imply guaranteed ACT points.

Claims we refuse as doctrine:

1. User-count / DAU / “#1 app” / star-average as primary learning social proof.
2. Vague testimonials (“amazing!” / “so confident!”) without a named competence object.

3. Mastery fireworks / tonight-% screenshots as shareable proof (SAFE-FORGET / SAFE-COLD).
4. Classmate ranks / Family Leaderboard as social proof (SAFE-COMPARE).
5. Identity Score™ / two-item HSLS pre-post as marketing proof (SAFE-IDMEASURE).
6. Guaranteed ACT-point case studies or “+N points in 8 weeks” packages (SAFE-EXAM / SAFE-DURABLE).
7. Empty AI praise / “math person” recognition without repertoire evidence (SAFE-IDMEASURE kill).
8. Forced public portfolios or viral share as North Star.

LXXXVI.2 Constructs

Construct	Research meaning	MindCraft analogue	Failure mode
Social proof (persuasion)	Under uncertainty, people look to others’ actions/opinions	Landing / parent CBC / tutor sales	Herd vanity without competence
Descriptive / provincial norms	“People like you / in this situation do X”	Similar-Maya case cards	Boomerang if framed as rank
Enactive mastery	Own successful performance → strongest efficacy source	solo_transfer_pass events	Scaffolded greens sold as mastery
Vicarious experience	Similar others succeed → “I can too”	Peer/near-peer artifacts; tutor film clips	Dissimilar Ivy hero ≠ Maya model
Witnessable artifact	Third party can inspect the competence claim	Shareable proof card: item type, hop, age, solo flag	Screenshot of streak / XP
SAFE-PROOF	Competence-first social proof doctrine	This chapter	Vanity / affect / points theater

Operational definition (HYPOTHESIS): A MindCraft social-proof unit is *SAFE-PROOF complete* when it (a) names a competence object (concept / bridge / FormatId), (b) states proof type (blocked vs delayed / transfer / solo), (c) shows proof age or date, (d) comes from a *similar* role (parent-of-anxious-Maya, near-peer tutor, not celebrity), (e) never claims guaranteed scores or permanent mastery, and (f) is share-opt-in with student dignity defaults (SAFE-PRIVACY / SAFE-PDASH).

LXXXVI.3 Social proof is real — and dangerous when it is only a crowd

FACT (principle + field norms): Cialdini’s social-proof principle — under uncertainty, people use others’ behavior as evidence of correct action — is the commercial spine of testimonials and “most guests...” signs (Cialdini, 2007/2009, *Influence*; field validation in Goldstein, Cialdini & Griskevicius, 2008, *Journal of Consumer Research*, 35(3), 472–482, doi:10.1086/586910). Descriptive-norm appeals beat generic exhortation; **provincial** norms (closest situational match) outperform broad “most people” claims.

FACT (boomerang risk): Schultz, Nolan, Cialdini, Goldstein & Griskevicius (2007, *Psychological Science*, 18(5), 429–434, doi:10.1111/j.1467-9280.2007.01917.x) — descriptive norms alone can *raise* undesired behavior among already-good performers (boomerang); adding an injunctive cue (approval of the desired behavior) reconstructs the message.

Commercial implication: “Most students on MindCraft practice daily” is weak and can shame or normalize the wrong thing. Prefer *similar-situation competence norms* (“Parents of ACT-anxious juniors who saw a delayed solo pass on linear systems...”) plus injunctive frame that values the *return / solo attempt*, not volume. Never use “you’re behind the average” as social proof (SAFE-COMPARE / SAFE-PDASH).

Kill: Crowd-size hero; classmate-average shame as proof.

Survive: Provincial, role-matched, competence-named norms; injunctive respect for honest returns.

LXXXVI.4 Mastery beats persuasion — artifacts must be enactive, not verbal

FACT (efficacy sources): Bandura (1977, *Psychological Review*, 84(2), 191–215, doi:10.1037/0033-295X.84.2.191) — efficacy expectations arise from performance accomplishments, vicarious experience, verbal persuasion, and physiological states; **performance accomplishments** are the most dependable.

FACT (school review): Usher & Pajares (2008, *Review of Educational Research*, 78(4), 751–796, doi:10.3102/0034654308321456) — across academic contexts, mastery experience is typically the most influential self-efficacy source; source strength varies by gender, ethnicity, ability, and domain; measurement of “sources” is often messy — product must not invent a Source Score™.

HYPOTHESIS (MindCraft mapping): Parent/tutor *verbal* persuasion (“you’re so smart”) and AI praise are the weakest durable channel. The growth surface should show **enactive** evidence: a dated solo delayed pass. Vicarious channel = similar Maya’s artifact or a near-peer tutor film of *student construction* (SAFE-TALK / SAFE-FILM), not an Ivy explainer reel.

Kill: Praise-wall social proof; “AI says she’s a math person.”

Survive: Shareable mastery events tied to FEI metrics (`solo_transfer_pass`, `format_hop_solo_transfer_pass`, `join_solo_transfer_pass`).

LXXXVI.5 Portfolios are the right metaphor — if student-owned and reflective

FACT (portfolio definition): Paulson, Paulson & Meyer (1991, *Educational Leadership*, 48(5), 60–63) — a portfolio is a *purposeful* collection exhibiting efforts, progress, and achievements; it must include student participation in selection, criteria, and **self-reflection** — “done *by* the student, not *to* the student.”

HYPOTHESIS: MindCraft’s shareable proof card / Field Journal export is a *thin portfolio unit*: one (or few) items with reflection (“what I can do alone now”) beats a dump of every green check. Parents witness competence; Maya owns selection — autonomy support, not surveillance scrapbook (SAFE-PDASH).

Wound: Math portfolio literature’s outcome claims are thinner than rhetoric (see critiques of large-scale portfolio programs). Do not market “portfolios raise ACT scores.” Market *inspectability* and belief correction.

Kill: Auto-publishing every session to parents/social.

Survive: Student-selected competence exhibits with age + solo flag.

LXXXVI.6 Honesty beats one-sided hype — two-sided proof builds credibility

FACT (two-sided ads): Eisend (2006, *International Journal of Research in Marketing*, 23(2), 187–198, doi:10.1016/j.ijresmar.2005.11.001) — meta-analysis: two-sided advertising (acknowledging a limitation) can enhance credibility and persuasion under conditions (e.g., when the negative is refutable or less important than gains) versus purely one-sided puffery.

HYPOTHESIS: Parent-facing case cards should admit dual-rail truth: “Blocked accuracy rose in week 2; delayed solo transfer on systems took until week 6 — returns still due.” That is SAFE-FORGET honesty as *marketing craft*, not a confession booth. Guaranteed-points one-sided case studies are both banned and less credible to skeptical buyers.

Kill: “+4 ACT points guaranteed” social proof.

Survive: Dated competence + remaining gap + next return CTA.

LXXXVI.7 Product surface — SAFE-PROOF artifact contract

Reuse: SAFE-TRANSFER (name the hop); SAFE-FORGET (age evidence); SAFE-PDASH (autonomy; no shame ranks); SAFE-WTP (CBC arms); SAFE-IDMEASURE (no Identity Score™ splash); SAFE-DURABLE (no cram-as-ready); SAFE-COMPARE (ranks off); SAFE-PRIVACY (opt-in share; no affect).

HYPOTHESIS — minimum shareable proof card fields:

Field	Required	Banned substitute
Competence object	Concept / bridge / FormatId name	“Math improved!”
Proof type	Delayed / transfer / solo flag	Tonight blocked %
Proof age	Last-proof date or days since	“Mastered forever”
Similarity frame	Role match (anxious ACT junior / ELL / etc.)	Celebrity / Ivy-only
Student voice	Optional short reflection	AI-written confidence blurb
Share control	Opt-in; revoke; no public default	Forced feed / viral NS
Claim hygiene	No score guarantee; no permanent mastery	+points package language

Surfaces: Parent digests (one card); tutor briefing (film + artifact); marketing landing (anonymized provincial cases); in-app MoC share sheet — not App Store star walls as the learning claim.

LXXXVI.8 Competitive positioning

Competitor pattern	Social-proof move	MindCraft response
Duo / streak apps	League ranks, streak screenshots	Competence cards; ranks off
ChatGPT tutor ads	Fluency demos, “feels like a tutor”	Solo-pass bar; PWC
Cram/bootcamp	+points testimonials	Dual-rail dated MoC; no guarantees
Khan	Library scale / “free for all”	Diagnosis + witnessable hops, not item-count
Generic edtech	Star walls / DAU	Provincial competence norms

Competitive wedge (FOUNDER BELIEF): “Proof you can see — a solo pass on the real gap, with an age stamp” beats both herd vanity and guaranteed-points theater.

LXXXVI.9 Doctrine — SAFE-PROOF (provisional)

1. **Competence > crowd** — social proof names a hop and solo/delayed evidence; demote user-count / stars as learning claims.
2. **Enactive > verbal** — prefer mastery artifacts over praise walls and AI “math person” lines (Bandura; Usher & Pajares).
3. **Provincial similarity** — match role/situation; avoid broad “everyone” and classmate shame (Goldstein et al.; Schultz et al.).
4. **Age the claim** — proof age required; no permanent-mastery fireworks (SAFE-FORGET).
5. **Student-owned exhibit** — thin portfolio: select + reflect; not done *to* the student (Paulson et al.).
6. **Two-sided honesty** — admit returns due / remaining gaps; ban +points guarantees (Eisend; SAFE-DURABLE).
7. **Opt-in dignity** — share is voluntary; no viral NS; no affect telemetry on cards (SAFE-PRIVACY).
8. **Copy:** “See what she can do alone — dated, on the real gap.” Never star-hero / Identity Score™ / guaranteed ACT case study.

Confidence: High that social proof / descriptive norms move behavior under uncertainty (Cialdini tradition; Goldstein et al. 2008) and that descriptive-only norms can boomerang (Schultz et al. 2007). High that mastery experiences dominate verbal persuasion for efficacy (Bandura 1977; Usher & Pajares 2008). Medium that MindCraft proof cards beat star walls on parent WTP (needs PROOF-* / ties WTP/PDASH). High that vanity counts, Identity Score™, +points testimonials, and forced public portfolios are kills. Medium that student-selected thin portfolios raise retention without becoming homework theater.

LXXXVI.10 Experiments

ID	Question	Design	Primary
PROOF-1	Competence-card landing vs star-wall / user-count landing	A/B	Signup; trust; time-on-proof
PROOF-2	Provincial similar-Maya case vs generic testimonial	A/B	Conversion; perceived relevance
PROOF-3	Two-sided dated MoC case vs one-sided +points claim	A/B / CBC	WTP; credibility; regret
PROOF-4	Opt-in share of solo-transfer card vs auto parent push of greens	A/B	Share rate; student stigma; parent action quality
PROOF-5	Parent CBC: “dated solo proof” vs “streak screenshots” vs “ACT points guarantee”	CBC	WTP; trust (SAFE-WTP)
PROOF-QUAL	10 parents + 10 Mayas: what feels like real proof vs hype?	Qual	Artifact fields; shame triggers

Falsifier: Star-wall wins conversion and 90-day retention equally → still ban stars as *learning* claim; segregate storefront chrome from MoC doctrine.

Falsifier: Proof cards raise signup but increase controlling parent nag → redesign share copy toward autonomy (SAFE-PDASH).

Falsifier: Two-sided honesty loses CBC badly with no credibility gain → keep ban on guarantees; soften wording, do not restore +points ads.

Falsifier: Students refuse all sharing → keep private MoC; marketing uses anonymized aggregates only.

Pre-register: PROOF-* before any “trusted by N families / +points stories” campaign (SAFE-LABMETA).

LXXXVI.11 So what for MindCraft commercially

- **Copy:** “See what she can do alone — dated, on the real gap.” Never crowd-size learning claims, vague confidence quotes, or guaranteed ACT case studies.
- **Product:** Shareable proof-card schema (object, hop type, age, solo flag, opt-in); thin portfolio select+reflect; parent digest = one card + one return CTA.
- **Positioning:** Against Duo screenshot culture, bootcamp +points testimonials, and ChatGPT fluency demos; for witnessable solo competence.

- **Metric:** Proof-card view→action; share opt-in rate; parent CBC preference; co-primary solo_transfer_pass — demote star NPS / DAU-as-proof.
 - **Kill list:** User-count hero; star walls as pedagogy; Identity Score™ splash; Family Leaderboard proof; AI praise testimonials; forced public portfolios.
 - **Growth:** Landing and sales decks lead with provincial competence artifacts; district packets can show anonymized MoC samples without biometric theater (SAFE-PROCURE).
 - **Vision:** Identity transformation becomes *socially legible* when competence can be witnessed — recognition of practice, not applause for streaks.
-

References (verified)

- Bandura, A. (1977). Self-efficacy: Toward a unifying theory of behavioral change. *Psychological Review*, 84(2), 191–215. <https://doi.org/10.1037/0033-295X.84.2.191>
 - Cialdini, R. B. (2009). *Influence: Science and practice* (5th ed.). Pearson. (Social-proof principle; earlier editions from 1984 onward.)
 - Eisend, M. (2006). Two-sided advertising: A meta-analysis. *International Journal of Research in Marketing*, 23(2), 187–198. <https://doi.org/10.1016/j.ijresmar.2005.11.001>
 - Goldstein, N. J., Cialdini, R. B., & Griskevicius, V. (2008). A room with a viewpoint: Using social norms to motivate environmental conservation in hotels. *Journal of Consumer Research*, 35(3), 472–482. <https://doi.org/10.1086/586910>
 - Paulson, F. L., Paulson, P. R., & Meyer, C. A. (1991). What makes a portfolio a portfolio? *Educational Leadership*, 48(5), 60–63.
 - Schultz, P. W., Nolan, J. M., Cialdini, R. B., Goldstein, N. J., & Griskevicius, V. (2007). The constructive, destructive, and reconstructive power of social norms. *Psychological Science*, 18(5), 429–434. <https://doi.org/10.1111/j.1467-9280.2007.01917.x>
 - Usher, E. L., & Pajares, F. (2008). Sources of self-efficacy in school: Critical review of the literature and future directions. *Review of Educational Research*, 78(4), 751–796. <https://doi.org/10.3102/0034654308321456>
-

Part LXXXVII — Bank Coverage Honesty in Marketing

Chapter status: Living evidence + GTM/copy brief — Researcher tick 2026-08-05

Primary question: How should MindCraft *advertise* question-bank coverage (ACT track, concept×level×format) without item-count heroics, “complete ACT bank” puffery, or GENQ volume theater — while still selling diagnosis trust?

Owners: Growth/copy · Product (Practice / Admin coverage) · Engine (bank audit / GENQ) · District GTM · Brand · Red Team

Commercial job: Ship a **SAFE-COVER** doctrine: coverage claims are *blueprint-honest matrices* + *use-tier labels*, never raw questionCount, never implied ACT-official completeness, never unverified generated volume as readiness.

Builds on: Parts XXXIV (claim ladder), XXXV (competitive audits), LIX (SAFE-WTP), LX (SAFE-REPAIR), LXVII (SAFE-ONTOLOGY), LXXII (SAFE-GENQ), LXXIII (SAFE-PROCURE), LXXXII (SAFE-FORMAT), LXXXVI (SAFE-PROOF). Product seams: actOntologyCoverage.json, Admin Testing coverage table, questionBank merge,

LXXXVII.1 Why this chapter exists

SAFE-GENQ already kills “AI wrote N questions” and unverified keys as diagnostic ground truth. Growth still faces a quieter hazard: **coverage theater** — leading landing pages, parent decks, and district packets with “1,500+ ACT questions,” “covers every ACT topic,” or “full ontology bank” while the live audit shows *partial* level coverage, ontology-only ACT-tested slots with zero static items, GCSE-tagged Eedi volume that must not be sold as ACT inventory, and generated JSON that is **inert until verify clears**. Item count is a *warehouse metric*; identity products sell *usable, keyed, blueprint-aligned practice for a named use*.

FOUNDER BELIEF under audit: Honest coverage tables (what exists, at which level/format, for which *use*) convert serious parents and LEAs better than inflated totals — and protect Map integrity when Maya’s gap-scan only draws from sealed keys.

Claims we refuse as doctrine:

1. Item-count / “N questions” as marketing hero or North Star.
 2. “Complete ACT coverage” / “every concept” without concept×level (and format) honesty.
 3. Treating GCSE/Eedi/other-examTag inventory as ACT-prep proof.
 4. Counting unverified / unsynced generated items in marketed coverage.
 5. Coverage % ≡ exam readiness, identity change, or FEI proof.
 6. Ontology node count / Layer-1 breadth ≡ playable practice bank.
 7. “AI fills any hole instantly” as coverage substantiation (SAFE-GENQ).
 8. Stale coverage audits sold as current product truth.
-

LXXXVII.2 Constructs

Construct	Research meaning	MindCraft analogue	Failure mode
Domain representation	Content samples the intended domain	Ontology×bank join; ACT-tested set	Node list without items
Alignment	Agreement of assessment with expectations	Concept/level/format vs ACT blueprint intent	Count without DOK/range
Use / IUA	Kane: validate <i>interpretations and uses</i>	Practice reveal vs C4 diagnostic vs parent report	Bank exists ≡ ACT-ready
Coverage table	Public matrix of what is shippable where	Admin actOntologyCoverage + marketing cut	Totals without gaps

Use-tier label	Which claims a slot may support	Practice / diagnostic-sealed / not playable	One number for all uses
SAFE-COVER	Blueprint-honest coverage GTM	This chapter	Item-count hero

Operational definition (HYPOTHESIS): A MindCraft coverage claim is *SAFE-COVER complete* when it (a) names the exam/track and ontology join, (b) reports **playable** counts by concept×level (and format when claimed), (c) lists **gaps** as first-class (partial / ontology-only / verify-blocked), (d) separates practice-eligible from diagnostic-sealed inventory, (e) excludes unverified GENQ and wrong-examTag inflation, and (f) never implies official ACT endorsement or guaranteed score lift from bank size.

LXXXVII.3 Validity is about uses — coverage is not a fetish object

FACT (argument-based validity): Kane (2013, *Journal of Educational Measurement*, 50(1), 1–73, doi:10.1111/jedm.12000) — what is validated are *interpretations and uses* of scores, not “the test” as a thing. More ambitious claims need more evidence; negative consequences can kill a use even if a thinner interpretation survives.

FACT (Standards framing): AERA, APA, & NCME (2014). *Standards for Educational and Psychological Testing*. Washington, DC: AERA. — Validity is the degree to which evidence and theory support interpretations for proposed uses; statements should not claim unqualified “validity of the test.” Content-related evidence must speak to the intended construct and use.

Applied (HYPOTHESIS):

Marketed claim	Ambition	Substantiation bar
“Practice items exist for topic X at L2”	Low	Playable keyed items in bank for that slot
“Gap-scan covers ACT-tested concepts”	High	Diagnostic-sealed keys + honest missing slots
“Full ACT Math coverage”	Extreme	Blueprint alignment + range/balance evidence — do not claim while partial/ontology-only gaps remain
“Bank proves identity / +ACT points”	Banned	Claim-ladder / SAFE-EXAM / SAFE-DURABLE

Kill: “Questions in the bank” ≡ validated ACT-readiness or identity use.

Survive: Use-tier coverage language tied to Kane IUA (ties SAFE-GENQ).

LXXXVII.4 Content evidence is representation — not warehouse volume

FACT (content-validity elements): Sireci (1998b, *Social Indicators Research*, 45(1–3), 83–117, doi:10.1023/A:1006985528729) — content validity concerns domain definition, **domain representation**, domain relevance, and appropriateness of development procedures — not how many items sit in a folder.

FACT (alignment as content evidence): Martone & Sireci (2009, *Review of Educational Research*, 79(4), 1332–1361, doi:10.3102/0034654309341375) — alignment studies evaluate congruence among curriculum, instruction, and assessment; content approaches (including SME congruence ratings) are a primary method family for gathering content-related evidence.

FACT (Webb criteria): Webb (2007, *Applied Measurement in Education*, 20(1), 7–25, doi:10.1080/08957340709336728) — alignment judgments use categorical concurrence, depth-of-knowledge consistency, range-of-knowledge correspondence, and balance of representation; “good enough” thresholds are pragmatic and partly subjective — volume alone never settles alignment.

Commercial implication: Marketing may show a **matrix** (concepts × levels × formats) and name missing cells. It may **not** treat “1,943 static questions” (current audit-scale warehouse) as proof the ACT blueprint is represented in depth, range, and balance. Sparse L3 on an ACT-frequent topic is a *coverage honesty* issue, not a footnote.

Kill: Item-count hero; ontology breadth ≡ bank representation.

Survive: Publish representation gaps; regenerate actOntologyCoverage.json before any coverage campaign.

LXXXVII.5 Alignment is hard — raters disagree; do not overclaim

FACT (rater variability): Herman, Webb & Zuniga (2007, *Applied Measurement in Education*, 20(1), 101–126, doi:10.1080/08957340709336732) — alignment pictures can change with which raters are used; agreement on topic/DOK is nontrivial. Large-scale programs still treat third-party alignment evidence as ethically and regulatorily important (WebbAlign / peer-review practice summarizing Webb 1997→; AERA/APA/NCME 2014 content evidence).

HYPOTHESIS: MindCraft’s internal conceptId join is a *lightweight congruence audit*, not a full Webb study. Selling “aligned to ACT” as if a DOK/range panel ran is cite-wash (SAFE-LABMETA). Prefer: “Mapped to our ACT-tested ontology concepts; here is the playable matrix and holes.”

Wound: Equating bank-alias joins with formal standards–assessment alignment studies.

Survive: Honest internal coverage tables as *minimum* GTM; commission real alignment only if making high-ambition district claims.

LXXXVII.6 Advertising law: prior substantiation for objective coverage claims

FACT (FTC prior substantiation): Federal Trade Commission (1984). *FTC Policy Statement Regarding Advertising Substantiation* (appended to *Thompson Medical Co.*, 104 F.T.C. 648, 839) — advertisers must have a **reasonable basis** for objective claims *before* dissemination; express/implied claims that suggest a level of proof require at least that level of support; failure to possess prior substantiation can violate Section 5 of the FTC Act.

FACT (consumer guidance): FTC Advertising FAQs (business guidance) — ads must be truthful and non-deceptive; advertisers need evidence for express *and* implied claims reasonable consumers take away; “tests prove”-class claims need the advertised level of proof.

Applied (HYPOTHESIS): “Complete ACT question bank,” “covers every tested topic,” and “1,500+ ACT questions” (when inventory mixes GCSE tags, partial levels, or unsynced GENQ) are **objective** coverage claims. Implied takeaway for parents often includes “enough for ACT prep / diagnosis.” Substantiation = dated coverage matrix + examTag discipline + verify seal — not a blog total.

Kill: Running coverage ads from stale audits or warehouse sums.

Survive: Pre-claim packet: regenerate audit → marketing cut of full/partial/ontology-only → use-tier footnotes.

LXXXVII.7 MindCraft bank reality (internal FACT for GTM)

FACT (product audit snapshot — regenerate before campaigns): `app/src/data/actOntologyCoverage.json` (audit summary fields) has reported on the order of ~29 ACT-tested concepts, mixed **full** / **partial** / **ontology_only** statuses, and nonzero **gapsNeedingContent** (e.g. partial level holes; ontology-only slots such as `strategy/representation_translation/basic_equations` with zero static bank match). Static warehouse totals can look large while L1/L2/L3 playable cells remain empty. CLAUDE.md / bank merge notes: multi-source bank (~static + ACT master + Eedi GCSE-tagged + generated stubs); several advanced concepts still zero-coverage; generation verify historically ~30% drop — **do not count unsynced GENQ**.

HYPOTHESIS: Parent- and LEA-facing honesty that *names* partial cells increases trust relative to “1,500+ questions” hero — testable via COVER- / *ties WTP*-.

Kill: Marketing the warehouse total without the gap list.

Survive: Admin Testing coverage table as the source of truth for copy; “we show the holes” as brand.

LXXXVII.8 Product surface — SAFE-COVER claim contract

Claim surface	Required fields	Banned substitute
Landing / ads	Track name; playable concept count; link or footnote to gaps	Bare “N questions”
Parent deck	Full / partial / missing summary; diagnostic seal note	“Complete ACT bank”
District packet	Dated audit; examTag policy; GENQ verify rate (SAFE-GENQ/PROCURE)	FERPA badge as coverage
In-app Learn/Practice	Playable gate (<code>hasPlayableQuestions</code>); no fake full path	Path nodes without items

Competitive foil	Diagnosis + sealed keys + honest matrix	Out-Khan on library size
------------------	---	--------------------------

Reuse: SAFE-GENQ (verify-before-count); SAFE-FORMAT (format claims need true FormatId stock); SAFE-ONTOLOGY (graph ≠ bank); SAFE-COLD (gap-scan only sealed keys); SAFE-PROOF (competence artifacts > library vanity); SAFE-WTP (CBC arms: honest matrix vs item-count hero).

LXXXVII.9 Competitive positioning

Competitor pattern	Coverage move	MindCraft response
Khan-scale library	Breadth / free forever	Diagnosis + playable honesty; do not out-count
Infinite AI drills	Volume / novelty	Verify gate; refuse poisoned Map (SAFE-GENQ)
Cram packs	“X tests worth of items”	Dual-rail; no completeness theater (SAFE-DURABLE)
ChatGPT homework	Unlimited stems	Keyed bank + seal; fluency ≠ coverage
Static publishers	Stable catalogs	Show ontology holes we still owe

Competitive wedge (FOUNDER BELIEF): “We publish the coverage matrix — including empty cells — and only sealed keys rewrite the Map” beats both library vanity and infinite-AI theater.

LXXXVII.10 Doctrine — SAFE-COVER (provisional)

1. **Matrix > total** — market concept×level (×format when claimed); demote item-count hero.
2. **Gaps are features** — partial and ontology-only cells appear in GTM, not only in eng docs.
3. **Use-tier labels** — practice vs diagnostic-sealed vs not playable (Kane IUA).
4. **ExamTag honesty** — do not sell GCSE/other inventory as ACT coverage.
5. **GENQ discipline** — unverified/unsynced generated items never enter marketed counts (SAFE-GENQ).
6. **Prior substantiation** — regenerate audit before coverage ads (FTC reasonable-basis posture).
7. **No ACT-official / complete-bank puffery** while holes remain.
8. **Copy:** “Here’s what you can practice now — and what we’re still filling.” Never “1,500+ ACT questions” as the learning claim.

Confidence: High that Kane/Standards bind *uses* not warehouses (Kane 2013; AERA/APA/NCME 2014). High that content evidence cares about representation/alignment, not raw N (Sireci 1998b; Martone & Sireci 2009; Webb 2007). High that FTC expects prior substantiation for objective ads (FTC 1984). Medium that honest matrices beat item-count heroes on parent/LEA conversion (needs COVER-*). High that counting unverified GENQ or wrong-examTag stock is a kill. Medium that lightweight ontology joins suffice for early GTM without full Webb panels.

LXXXVII.11 Experiments

ID	Question	Design	Primary
COVER-1	Honest matrix landing vs item-count hero landing	A/B	Signup; trust; bounce on gaps
COVER-2	Parent CBC: “shows holes + sealed keys” vs “1,500+ ACT questions” vs “AI unlimited”	CBC	WTP; credibility (SAFE-WTP)
COVER-3	District packet with dated coverage audit vs brochure totals	Sales A/B / qual	Advance-to-DPA; objection themes
COVER-4	In-app gap callout (“L3 not ready”) vs silent fail / dynamic-gen promise	A/B	Session start; trust; support tickets
COVER-5	Regenerate-audit cadence before campaigns vs stale JSON	Ops	Claim–product mismatch incidents
COVER-QUAL	10 parents + 5 tutors: what “covers ACT” means to them	Qual	Implied-claim codebook for FTC-style review

Falsifier: Item-count hero wins conversion *and* 90-day retention equally → still ban count as *learning/ACT completeness* claim; segregate storefront vanity from coverage doctrine.

Falsifier: Publishing gaps tanks signup with no trust gain → redesign framing (roadmap + sealed keys), do not restore “complete bank” lies.

Falsifier: Parents ignore matrices → keep doctrine for LEA/trust; simplify consumer cut to three numbers (full / partial / missing) + one CTA.

Pre-register: COVER-* before any “complete ACT / N questions / AI fills all holes” campaign (SAFE-LABMETA).

LXXXVII.12 So what for MindCraft commercially

- **Copy:** “Here’s what you can practice now — and what we’re still filling.” Lead with matrix + sealed keys. Never item-count / complete-ACT / infinite-AI coverage heroes.
- **Product:** Keep Admin coverage table + `actOntologyCoverage.json` as GTM source of truth; playable gates; use-tier labels on Learn/Practice; regenerate audit in release checklist before coverage claims.

- **Positioning:** Against Khan library vanity and AI-infinite drills; for blueprint-honest, verify-sealed practice that cannot silently poison diagnosis.
 - **Metric:** Coverage-claim incidents (mismatched ads); matrix CTR→session; parent CBC preference; co-primary playable-slot fill rate — demote questionCount / generation volume.
 - **Kill list:** Item-count hero; complete-ACT puffery; GCSE-as-ACT; counting unsynced GENQ; ontology-nodes=bank; stale-audit campaigns.
 - **Growth:** District trust packets include dated coverage + verify rates (SAFE-PROCURE × SAFE-GENQ); parent decks show holes as integrity, not weakness.
 - **Vision:** A thirty-year identity company earns belief by refusing warehouse theater — Maya’s Map only moves on keys that exist and are true.
-

References (verified)

- American Educational Research Association, American Psychological Association, & National Council on Measurement in Education. (2014). *Standards for educational and psychological testing*. AERA.
 - Federal Trade Commission. (1984). FTC policy statement regarding advertising substantiation. Appended to *Thompson Medical Co.*, 104 F.T.C. 648, 839.
<https://www.ftc.gov/legal-library/browse/ftc-policy-statement-regarding-advertising-substantiation>
 - Herman, J. L., Webb, N. M., & Zuniga, S. A. (2007). Measurement issues in the alignment of standards and assessments: A case study. *Applied Measurement in Education*, 20(1), 101–126.
<https://doi.org/10.1080/08957340709336732>
 - Kane, M. T. (2013). Validating the interpretations and uses of test scores. *Journal of Educational Measurement*, 50(1), 1–73. <https://doi.org/10.1111/jedm.12000>
 - Martone, A., & Sireci, S. G. (2009). Evaluating alignment between curriculum, assessment, and instruction. *Review of Educational Research*, 79(4), 1332–1361. <https://doi.org/10.3102/0034654309341375>
 - Sireci, S. G. (1998b). The construct of content validity. *Social Indicators Research*, 45(1–3), 83–117.
<https://doi.org/10.1023/A:1006985528729>
 - Webb, N. L. (2007). Issues related to judging the alignment of curriculum standards and assessments. *Applied Measurement in Education*, 20(1), 7–25. <https://doi.org/10.1080/08957340709336728>
-

Part LXXXVIII — Worked-Example Fading in Solver UX

Chapter status: Living evidence + Solver/Practice UX brief — Researcher tick 2026-08-06

Primary question: How should MindCraft fade guidance from full worked examples → completion steps → solo solve in Solver (and Practice help rails) without permanent solution-dump, expertise-reversal tax, or “instant full answer” as the product brand?

Owners: Product (Solver / Practice) · Engine (ingredient cards / coach) · HITL tutors · Brand · Red Team

Commercial job: Ship a **SAFE-FADE** doctrine: guidance level is a *diagnosed state*, not a permanent chrome; attempt grain before reveal; completion problems as the default bridge; solo transfer as the proof that fading worked.

Builds on: Parts XXVI (CLT / fading), XL (SAFE-SE), LI (SAFE-DP), LXXIX (SAFE-FBTIME), LXXXIV (SAFE-TALK), LXXXV (SAFE-STORYLOAD), LXXXVI (SAFE-PROOF). Product seams: Solver help, ingredient

cards, soft-wrong / SE gates, solo_transfer_pass, hint binge guards.

LXXXVIII.1 Why this chapter exists

Homework Help / Solver is the surface where competitors and ChatGPT win on *fluency*: paste problem → full worked solution → green check feeling. MindCraft already bans unguarded Solver hero (PWC) and AI-monologue≡SE. The remaining product hazard is quieter: **permanent full worked examples** (or always-on step dumps) that help absolute novices once, then become redundant load — and train help-seeking that never fades.

FOUNDER BELIEF under audit: A Solver that *fades* (full example → missing last steps → missing more → full problem) converts “I needed help” into “I can finish the join myself” — the identity move — better than unlimited perfect solutions. That claim is commercial only if we instrument attempt grain and solo transfer, not solution views.

Claims we refuse as doctrine:

1. Always-full-worked / solution-first Solver as learning North Star.
 2. Never-fade step scaffolding (“we always show every step”).
 3. Completion-of-viewing ≡ mastery (watching steps = doing math).
 4. Instant full answer as brand differentiator vs ChatGPT.
 5. Expertise-blind help (same dump for L1 and near-fluent).
 6. Hint binge / reveal streak without fade-up after success.
 7. Fade Score™ / guidance-% as vanity NS.
 8. Forward-only fade as dogma without local FADE-* check (backward often wins in lab).
-

LXXXVIII.2 Constructs

Construct	Research meaning	MindCraft analogue	Failure mode
Worked example	Full problem + solution steps for study	Full coach card / Solver reveal	Permanent dump
Completion problem	Partial solution; learner finishes	Missing-step Solver; last-step blank	Fake blanks with answer peek
Fading	Successive removal of worked steps	Guidance ladder per concept/format	One-shot jump to solo
Expertise reversal	Novice-helpful guidance hurts experts	Same full example for advanced	Redundant load / boredom
Attempt grain	Learner generates before full reveal	SE / soft-wrong / blank-step gate	Watch-only “learning”
SAFE-FADE	Diagnosed fade ladder + solo proof	This chapter	Solution theater

Operational definition (HYPOTHESIS): A Solver/help path is *SAFE-FADE complete* when it (a) defaults to **attempt or completion** before full worked reveal on learn rail, (b) fades **backward** (omit final steps first) unless

FADE- *prefers otherwise for that schema*, (c) *raises guidance only on diagnosed struggle (not permanent chrome)*, (d) *fades up** after consecutive solo successes, (e) *never counts solution-view as mastery*, and (f) *co-primary proves with solo_transfer_pass / delayed mixed accuracy — not hint satisfaction NPS*.

LXXXVIII.3 Novices need worked examples — search is the tax

FACT (worked-example effect): Sweller & Cooper (1985, *Cognition and Instruction*, 2(1), 59–89, doi:10.1207/s1532690xc0201_3) — for algebra novices, studying worked examples as a substitute for conventional problem solving reduced time and errors on subsequent similar problems; conventional search imposed heavy load that retarded schema acquisition.

FACT (CLT framing): Sweller (1988, *Cognitive Science*, 12(2), 257–285, doi:10.1207/s15516709cog1202_4) — problem-solving search can consume working memory needed for schema construction; reducing extraneous search supports learning of complex material.

Applied (HYPOTHESIS): Day-one Solver that *only* withholds all guidance (pure struggle theater) violates CLT for true novices. SAFE-FADE is not “never help” — it is **help that starts complete enough to build schema, then steps down**.

Kill: Discovery-only / no-example Solver for first acquisition.

Survive: Full worked study *as a stage*, not as the product forever.

LXXXVIII.4 Expertise reverse — permanent examples become clutter

FACT (expertise reversal): Kalyuga, Ayres, Chandler & Sweller (2003, *Educational Psychologist*, 38(1), 23–31, doi:10.1207/s15326985ep3801_4) — instructional methods highly effective for inexperienced learners can lose effectiveness or become detrimental as expertise grows, because guidance becomes redundant and loads working memory with unnecessary cross-checking.

FACT (when solving beats examples): Kalyuga, Chandler, Tuovinen & Sweller (2001, *Journal of Educational Psychology*, 93(3), 579–588, doi:10.1037/0022-0663.93.3.579) — as learners gain expertise in a domain, problem solving can become superior to studying worked examples (classic reverse pattern under CLT).

Commercial implication: Marketing “we show every step, always” is a **novice costume** that taxes Maya once she has partial schemas — and looks identical to ChatGPT solution dump. Map/level/format evidence must *lower* default guidance after demonstrated competence.

Kill: Expertise-blind always-full-worked brand.

Survive: Guidance that tracks diagnosed grain (ties SAFE-COLD / SAFE-FORMAT / SAFE-ONTOLOGY).

LXXXVIII.5 Completion problems bridge example → solve

FACT (completion as bridge): van Merriënboer and colleagues introduced *completion problems* — given state, goal, and partial solution the learner must finish — as intermediates between fully worked examples and conventional problems; completion strategies emphasize completing increasingly larger parts of incomplete solutions (see van Merriënboer, 1990; van Merriënboer & de Croock, 1992, in programming; reviewed in Paas & van Merriënboer instructional-control work).

FACT (mental effort / transfer): Paas (1992, *Journal of Educational Psychology*, 84(4), 429–434, doi:10.1037/0022-0663.84.4.429) — in statistics problem solving, training with completion problems or worked examples yielded better transfer with lower test-phase mental effort than conventional problem training (completion ≈ worked examples on effort during training; both beat conventional on transfer efficiency).

Applied (HYPOTHESIS): Solver’s default “help” should often be a **blank last step / blank middle join**, not a monologue. Ingredient cards already suggest multi-representation scaffolds — fade means those scaffolds lose steps over successes, not that more cards appear forever.

Kill: Help ≡ paste full solution.

Survive: Help ≡ completion problem with named missing grain (procedure, bridge, format).

LXXXVIII.6 Smooth fading beats static example–problem pairs

FACT (fading procedure): Renkl, Atkinson, Maier & Staley (2002, *Journal of Experimental Education*, 70(4), 293–315, doi:10.1080/00220970209599510) — successive integration of problem-solving elements into example study (complete example → increasingly incomplete → to-be-solved) fostered near-transfer learning vs traditional example–problem pairs; errors during learning mediated the effect; **backward fading** (omit last solution steps first) outperformed forward fading.

FACT (CLT transition argument): Renkl & Atkinson (2003, *Educational Psychologist*, 38(1), 15–22, doi:10.1207/S15326985EP3801_3) — as intrinsic load falls with schema growth, problem-solving demands can increase without overload; fading structures that transition; later stages change which activities are germane vs extraneous.

Applied (FOUNDER BELIEF → testable): MindCraft Solver ladder:

Stage	UX	Gate to next
E0 Full example	Study + SE prompt on a marked step	SE / principle tag
E1 Backward fade	Last step blank	Correct completion
E2 Deeper fade	Last two / critical join blank	Correct + no peek
E3 Full problem	Solo attempt	solo_transfer_pass delayed

Raise guidance only after miss class / pour streak — never as default forever.

Kill: Static example then unrelated dump; jump from E0 to E3 without completions.

Survive: Backward-first fade ladder with local FADE-* override rights.

LXXXVIII.7 Attempt grain — watching is not doing

FACT (passive example risk): Worked-example literature repeatedly notes that *passive* viewing weakens the effect; prompting self-explanation and active study of solutions strengthens learning (ties Chi / Renkl SE line — Part XL). MIT TLL synthesis of Sweller & Cooper (1985) and later work: insufficient examples and passive demos neutralize the worked-example effect.

HYPOTHESIS: Solver telemetry that optimizes `solution_reveal_rate` or time-on-help will invent help abuse (next chapter 89). Co-primary must be **student-generated steps before reveal** + delayed solo.

Reuse: SAFE-SE (why before wrap); SAFE-FBTIME (micro-delay on learn rail); SAFE-TALK (construction > pour); SAFE-PROOF (solo artifact).

Kill: View-complete $\equiv$ session success.

Survive: `fade_step_attempt` events before unlock of next worked line.

LXXXVIII.8 Product surface — SAFE-FADE claim contract

Surface	Required behavior	Banned substitute
Solver first open (learn)	Attempt or completion before full dump	Instant ChatGPT-style full solution
After consecutive solos	Fade <i>up</i> (less guidance)	Keep dumping “to be helpful”
After struggle	Fade <i>down</i> one stage; name the missing grain	Permanent max scaffolding
Ingredient / coach cards	Steps can blank; SE on blank	Always-complete card deck
Marketing	“We fade help as you prove the join”	“Unlimited full solutions”
Prove / exam rail	Minimal worked reveal; KR timing (SAFE-FBTIME)	Learn-rail full examples under stakes

Competitive foil: Khan/Brilliant often lean example+practice; ChatGPT leans full solution fluency. MindCraft differentiates on **diagnosed fade + solo transfer**, not on prettier monologues.

LXXXVIII.9 Competitive positioning

Competitor pattern	Guidance move	MindCraft response
ChatGPT homework	Full fluent solution	Completion + fade + Map-check (SAFE-REPAIR)
Instant-feedback AI	Always show steps	Mode-conditional; attempt grain

Cram packs	Worked keys as content	Dual-rail; fade on learn, prove alone
“Socratic” bots	Endless questions <i>or</i> dump	Construction metrics; not talk-% (SAFE-TALK)
Static textbooks	Fixed example then exercises	Adaptive fade from evidence state

Competitive wedge (FOUNDER BELIEF): “Help that disappears when you can carry the step” beats both abandonment struggle and infinite solution theater.

LXXXVIII.10 Doctrine — SAFE-FADE (provisional)

1. **Stage, don’t costume** — full worked examples are an *acquisition stage*, not the brand forever.
2. **Completion is the default bridge** — blank critical / final steps before full reveal.
3. **Backward fade first** — omit last steps first unless FADE-* shows forward better for that schema.
4. **Expertise-aware** — raise guidance on struggle; lower after consecutive solos (expertise reversal).
5. **Attempt grain** — student-generated step / SE before unlocking more worked lines.
6. **Proof = solo transfer** — not solution views, Fade Score™, or help NPS.
7. **Copy:** “We fade the steps as you prove them.” Never “unlimited perfect solutions” or “we always show every step.”
8. **Tie rails** — learn rail may use E0–E2; prove/exam rail stays near E3 (SAFE-EXAM / SAFE-DURABLE).

Confidence: High that worked examples help novices vs conventional search (Sweller & Cooper 1985; Sweller 1988). High that expertise reverses example benefits (Kalyuga et al. 2003; Kalyuga et al. 2001). High that fading/completion structures transition better than static pairs (Renkl et al. 2002; Renkl & Atkinson 2003; Paas 1992). Medium that backward fade is the right default for MindCraft algebra/bridge schemas (needs FADE-*). Medium that parent WTP prefers fade story over unlimited solutions (needs WTP/FADE CBC). High that permanent full-dump Solver is commercially indistinguishable from ChatGPT and anti-identity.

LXXXVIII.11 Experiments

ID	Question	Design	Primary
FADE-1	Backward completion ladder vs always-full-worked Solver	A/B within concept	Near transfer; solo_transfer_pass; hint binge
FADE-2	Backward vs forward fade order	A/B	Transfer; errors during learning
FADE-3	Expertise-aware fade-up after 2 solos vs fixed E1 forever	A/B	Time-on-task; transfer; boredom/skip

FADE-4	SE-on-blank-step vs passive completion	A/B	Transfer; SE quality
FADE-5	Parent CBC: “fades help as you prove” vs “unlimited full solutions” vs “never shows answers”	CBC	WTP; trust (SAFE-WTP)
FADE-QUAL	10 Maya: when does a shown step feel like rescue vs theft of the win?	Qual	Fade UX codebook

Falsifier: Always-full-worked wins transfer *and* 26w identity equally → still ban as *brand* vs ChatGPT; segregate acquisition-stage examples from marketing hero.

Falsifier: Aggressive fade tanks novice retention with no solo gain → hold longer at E0/E1; do not restore permanent dump.

Falsifier: Forward fade beats backward on MindCraft bank → update ladder default; keep completion doctrine.

Pre-register: FADE-* before any “AI shows every step / unlimited solutions / never-fade scaffold” campaign (SAFE-LABMETA).

LXXXVIII.12 So what for MindCraft commercially

- **Copy:** “We fade the steps as you prove them.” Lead with completion → solo, not unlimited solutions.
 - **Product:** Solver/help ladder E0→E3; blank-step attempt events; fade-up after solos; prove rail near E3; ingredient cards lose steps over success.
 - **Positioning:** Against ChatGPT solution dump and struggle-theater apps; for diagnosed guidance that earns identity via solo carry.
 - **Metric:** fade_stage, fade_step_attempt, hint_binge, co-primary solo_transfer_pass — demote solution_reveal_rate as success.
 - **Kill list:** Always-full-worked hero; never-fade scaffolding; view≡mastery; expertise-blind dump; Fade Score™ NS; unlimited-solutions ads.
 - **Growth:** Parent decks sell *earned independence of steps*; tutors QA to completion prompts not monologues (SAFE-HITL × SAFE-TALK).
 - **Vision:** Thirty-year identity company teaches Maya to need the worked line less — help that graduates her, not a solution vending machine.
-

References (verified)

- Kalyuga, S., Ayres, P., Chandler, P., & Sweller, J. (2003). The expertise reversal effect. *Educational Psychologist*, 38(1), 23–31. https://doi.org/10.1207/s15326985ep3801_4

- Kalyuga, S., Chandler, P., Tuovinen, J., & Sweller, J. (2001). When problem solving is superior to studying worked examples. *Journal of Educational Psychology*, 93(3), 579–588. <https://doi.org/10.1037/0022-0663.93.3.579>
- Paas, F. G. W. C. (1992). Training strategies for attaining transfer of problem-solving skill in statistics: A cognitive-load approach. *Journal of Educational Psychology*, 84(4), 429–434. <https://doi.org/10.1037/0022-0663.84.4.429>
- Renkl, A., & Atkinson, R. K. (2003). Structuring the transition from example study to problem solving in cognitive skill acquisition: A cognitive load perspective. *Educational Psychologist*, 38(1), 15–22. https://doi.org/10.1207/S15326985EP3801_3
- Renkl, A., Atkinson, R. K., Maier, U. H., & Staley, R. (2002). From example study to problem solving: Smooth transitions help learning. *Journal of Experimental Education*, 70(4), 293–315. <https://doi.org/10.1080/00220970209599510>
- Sweller, J. (1988). Cognitive load during problem solving: Effects on learning. *Cognitive Science*, 12(2), 257–285. https://doi.org/10.1207/s15516709cog1202_4
- Sweller, J., & Cooper, G. A. (1985). The use of worked examples as a substitute for problem solving in learning algebra. *Cognition and Instruction*, 2(1), 59–89. https://doi.org/10.1207/s1532690xci0201_3

Part LXXXIX — Help-Seeking vs Help Abuse

Chapter status: Living evidence + Solver/Practice/HITL brief — Researcher tick 2026-08-06

Primary question: When should MindCraft treat help use as adaptive self-regulation versus executive help abuse (or help avoidance), and what product/copy rules follow without turning “unlimited hints” or “never help” into the brand?

Owners: Product (Solver / Practice / coach) · Engine (hint ladder / telemetry) · HITL tutors · Brand · Red Team

Commercial job: Ship a **SAFE-HELP** doctrine: instrumental help that graduates into solo transfer; detect and discourage executive binge / bottom-out racing; invite help when avoidance is the failure; prove with deliberate help metrics + solo_transfer_pass, not hint satisfaction NPS.

Builds on: Parts XXVI (fading), XL (SAFE-SE), XLV (SAFE-RESILIENCE), LI (SAFE-DP), LVIII (SAFE-EXPOSE), LXXIX (SAFE-FBTIME), LXXXIV (SAFE-TALK), LXXXVIII (SAFE-FADE). Product seams: hint_binge, ai_reveal_rate, help_recruit_then_solo, fade ladder E0–E3, Solver reveal gates. Next sibling: id 91 hint economy (cost of peek).

LXXXIX.1 Why this chapter exists

SAFE-FADE answered *what* guidance looks like as expertise grows. This chapter answers *when students ask for it* — and when that ask is learning versus gaming.

FOUNDER BELIEF under audit: A product that makes help *always cheap and bottom-out-complete* trains Maya to recruit the answer, not the join. A product that shames all help trains avoidance and freeze. MindCraft’s identity claim requires **instrumental** help-seeking (help to understand the next step) and refuses **executive** help-seeking (help to finish without thinking) as success theater.

Claims we refuse as doctrine:

1. Unlimited hints / always-open bottom-out as learning North Star.
2. Hint volume or help NPS as proof of pedagogy quality.
3. Help avoidance as grit (struggle theater that bans support).
4. Metacognitive nag alone as guaranteed domain gains (Help Tutor caution).
5. Gaming detectors as shame dashboards for parents.
6. “AI always knows when you need help” black-box contingent brand.
7. Help Score™ / Talk-to-hint ratio vanity NS.
8. Permanent tutor dependency sold as “supportive.”

LXXXIX.2 Constructs

Construct	Research meaning	MindCraft analogue	Failure mode
Instrumental help	Help aimed at learning how to proceed	Principle hint / completion blank / SE gate	Rare without design
Executive help	Help aimed at getting the answer done	Bottom-out dump / ChatGPT paste	Common default
Help abuse	Maladaptive overuse (e.g. hint racing)	Rapid reveal chain; hint_binge	Fake fluency
Help avoidance	Under-use when help would help	Stuck loops; refuse coach	Freeze / shame
Try-step abuse	Hasty guessing without sense-making	Spam submit / click farm	Gaming
Contingent help	Scaffold level tracks need	Fade down/up from evidence	Fixed forever chrome
SAFE-HELP	Instrumental + contingent + solo proof	This chapter	Hint vending

Operational definition (HYPOTHESIS): A help path is *SAFE-HELP complete* when it (a) distinguishes instrumental vs executive use in telemetry, (b) makes bottom-out / full reveal costly in *effort or stage* (not humiliation), (c) invites help after repeated miss without pouring the answer first, (d) co-primary proves with `help_recruit_then_solo` and `solo_transfer_pass`, and (e) never markets unlimited answer-complete help as identity transformation.

LXXXIX.3 Instrumental vs executive — the commercial fork

FACT (classroom help typology): Classic help-seeking research distinguishes help that supports learning how to solve from help that merely completes the task — often labeled instrumental vs executive (Nelson-Le Gall, 1981; Karabenick & Knapp, 1991; summarized in Roll, Aleven, McLaren & Koedinger, 2011, *Learning and Instruction*, 21(2), 267–280, doi:10.1016/j.learninstruc.2010.07.004).

Applied (HYPOTHESIS): ChatGPT-style Solver and “unlimited hints” ads optimize for *executive* completion feeling. Parent WTP for “my kid got unstuck” must be re-aimed at *instrumental* unstuck — named next grain + student-generated step — or MindCraft collapses into homework completion SaaS.

Kill: Help $\equiv$ finish the worksheet.

Survive: Help $\equiv$ recruit a scaffold that still requires Maya’s next move (ties SAFE-FADE / SAFE-SE).

LXXXIX.4 Help abuse is frequent — and often anti-learning

FACT (early Cognitive Tutor logs): Aleven & Koedinger (2000/2001; reviewed in Aleven, McLaren, Roll & Koedinger, 2016, *IJAIED*, 26(1), 205–223, doi:10.1007/s40593-015-0089-1) — students often raced hint levels to the bottom-out answer (many pre-bottom levels viewed <1s) and under-asked after errors; on-demand help frequency correlated negatively with learning gains (selection + misuse both plausible).

FACT (model taxonomy): Aleven, McLaren, Roll & Koedinger (2006, *IJAIED*; also Aleven et al., 2004, ITS) — executable model with Help Abuse, Help Avoidance, and Try-Step Abuse; early offline eval classified a large share of actions as unproductive (Aleven et al., 2004 ~72% under a strict model — model-dependent, not a marketing statistic).

FACT (gaming): Baker and colleagues treat rapid help abuse / exploiting tutor regularities as primary “gaming the system,” associated with reduced learning (Baker, Corbett & Koedinger line; Roll et al., 2011; Aleven et al., 2016).

Applied (FOUNDER BELIEF $\rightarrow$ testable): MindCraft’s existing `hint_binge` (≥ 3 hints without independent solve) and `ai_reveal_rate` are the right *family* of detectors — but must not become parent shame ranks (SAFE-PDASH). Use them for coach fade, tutor QA, and product gates.

Kill: Hint clicks as engagement success.

Survive: Deliberate hint reading time + fewer levels when help is taken (Help Tutor behavioral signature).

LXXXIX.5 Help avoidance is the other failure — not grit

FACT: The same Aleven/Roll line documents frequent help avoidance — students who need help most often seek it least or seek it poorly (Aleven et al., 2003 overview; Wood & Wood, 1999; Nelson-Le Gall et al., 1990 on poorer self-assessment among lower prior-knowledge students — cited throughout the Help Tutor program).

FACT (contingent tutoring): Wood & Wood (1999, *Computers & Education*, 33(2–3), 153–169) — contingent tutoring adjusts help level with success/failure; process measures of help-seeking relate to prior knowledge and outcomes; help seeking is part of constructing a workable zone of support, not a character flaw.

Applied (HYPOTHESIS): After repeated soft-wrong / pour risk, MindCraft should *offer* instrumental help (completion blank, principle card) — not celebrate silent freeze as resilience. SAFE-RESILIENCE already recruits support; SAFE-HELP specifies the *kind*.

Kill: “Real math kids never ask for hints.”

Survive: Help invite after struggle; still require attempt grain before bottom-out.

LXXXIX.6 Metacognitive feedback improves help skill — not always domain gain

FACT (Help Tutor studies): Roll et al. (2006 ITS; 2011, *Learning and Instruction*, 21(2), 267–280, doi:10.1016/j.learninstruc.2010.07.004) — metacognitive feedback made help use more deliberate (more time per level; fewer levels), sometimes persisting after feedback ended — **but domain learning gains did not reliably improve** (Aleven et al., 2016: “help helps, but only so much”).

FACT (theoretical update): Aleven et al. (2016) — principle-based hints help when students sense-make; bottom-out can act like a worked example *if* self-explained (Shih, Koedinger & Scheines, 2008); otherwise executive racing yields thin learning.

Commercial implication: Do not sell “we tutor metacognition so scores rise.” Sell **instrumental help + SE/fade gates + solo proof**. Metacognitive nags are support tools, not the FEI product.

Kill: Help-Tutor cosplay as ACT-point engine.

Survive: Deliberate-help UX + SAFE-SE on hints + FADE ladder; HELP-* before metacognition brand campaigns.

LXXXIX.7 Contingent scaffolding — system share of the load

FACT: Contingent / adaptive hint level (Wood & Wood, 1999; Luckin & du Boulay, 1999 — cited in Roll et al., 2011) reduces opportunities for some maladaptive patterns by matching help grain to estimated need — a design cousin of Cognitive Tutor delays between rapid hint requests.

Applied (HYPOTHESIS): MindCraft should prefer **contingent stage** (SAFE-FADE E0–E3 driven by struggle/solo evidence) over unlimited student-chosen bottom-out. Student agency remains: they can request the next instrumental grain; the system refuses instant executive completion on learn rail without attempt.

SPECULATION: Fully automatic “AI chooses every hint” without inspectable why will read as surveillance pedagogy to parents and as black-box to tutors — prefer named stage + reason chip (ties SAFE-ADAPT inspectability pattern).

Kill: Unlimited student-controlled answer dump as personalization.

Survive: Contingent ladder + requestable next grain + visible stage.

LXXXIX.8 Product surface — SAFE-HELP claim contract

Surface	Required behavior	Banned substitute
Solver learn rail	Instrumental first; bottom-out behind attempt/SE	Instant full answer
Practice hints	Levelled hints; binge gate; fade-up after solos	Hint spam = progress

After 2–3 misses	Offer help / completion; don't shame ask	Silent freeze as virtue
Tutor HITL	Prompt recruit of student why before pour	Tutor as answer key
Parent view	Aggregate deliberate-help / solo — not hint shame	Live hint-stalk dashboard
Marketing	“Help that teaches the next step”	“Unlimited hints / never stuck”
Prove / exam rail	Minimal executive help; KR timing (SAFE-FBTIME)	Learn-rail binge under stakes

Competitive foil: ChatGPT wins executive completion. Struggle-only apps win avoidance theater. Duo-like products win engagement loops. MindCraft differentiates on **instrumental, contingent help that still demands Maya's construction** — then disappears into solo transfer.

LXXXIX.9 Doctrine — SAFE-HELP (provisional)

1. **Instrumental > executive** — help must leave a student-generated step; bottom-out is a last grain, not the brand.
2. **Abuse and avoidance are both failures** — binge racing and freeze both fail FEI; design for both detectors.
3. **Contingent stage, not unlimited menu** — help level tracks evidence (SAFE-FADE); student requests next grain, not full dump by default.
4. **Sense-making gate** — hints without SE/attempt are incomplete (SAFE-SE × Aleven 2016).
5. **Metacognition is not score magic** — deliberate-help skill can improve without domain gains; co-primary stays transfer/solo.
6. **Dignity telemetry** — hint_binge / reveal rate for product + tutor QA; never parent shame ranks.
7. **Copy:** “Ask for the next step — not the finished answer.” Never “unlimited hints so you're never stuck.”
8. **Tie rails** — learn rail allows instrumental ladders; prove rail stays near solo (SAFE-EXAM / SAFE-DURABLE).

Confidence: High that executive hint racing/gaming is common and often anti-learning (Aleven/Koedinger; Baker; Aleven et al., 2016). High that instrumental vs executive is the commercial fork (Nelson-Le Gall / Karabenick via Roll et al., 2011). High that contingent tutoring beats free bottom-out as design prior (Wood & Wood, 1999). High that Help Tutor-style feedback can change help behavior without guaranteed domain gains (Roll et al., 2011; Aleven et al., 2016) — medium that SE+fade closes that gap (HELP-*). Medium parent CBC preference for instrumental copy (HELP-5). High that unlimited-hints brand ≈ ChatGPT completion.

LXXXIX.10 Experiments

ID	Question	Design	Primary
HELP-1	Contingent instrumental ladder vs unlimited bottom-out Solver	A/B within concept	solo_transfer_pass; hint_binge; near transfer

HELP-2	SE-required before next hint level vs free drill-down	A/B	Deliberate hint time; transfer; binge
HELP-3	Help-invite after 2 misses vs no invite (avoidance arm)	A/B	Help recruit; eventual solo; anxiety item
HELP-4	Metacognitive binge message vs silent gate vs neither	A/B/C	Help behavior; domain transfer (expect possible null on domain)
HELP-5	Parent CBC: “next-step help” vs “unlimited hints” vs “we never show answers”	CBC	WTP; trust (SAFE-WTP)
HELP-QUAL	10 Maya: when does a hint feel like a tool vs a theft of the win?	Qual	Instrumental/executive codebook

Falsifier: Unlimited bottom-out wins delayed solo transfer *and* 26w identity equally → still ban as *brand* vs ChatGPT; segregate rare emergency reveal from marketing.

Falsifier: Aggressive binge gates raise avoidance and drop retention with no solo gain → soften gate; keep instrumental-first doctrine.

Falsifier: Metacognitive messages alone raise ACT-relevant transfer → allow limited HELP-Tutor-style copy; still co-primary solo, not Help Score™.

Pre-register: HELP-* before any “unlimited hints / never stuck / AI knows exactly when to help” campaign (SAFE-LABMETA). Sibling id 91 (hint economy) owns peek-cost pricing experiments — do not duplicate as HELP-\$.

LXXXIX.11 So what for MindCraft commercially

- **Copy:** “Ask for the next step — not the finished answer.” Lead with instrumental help + solo proof.
- **Product:** Contingent hint/fade ladder; SE/attempt before bottom-out; help-invite on avoidance; binge gates without shame UX.
- **Positioning:** Against ChatGPT executive completion and struggle-theater grit; for self-regulated help that graduates into independence.
- **Metric:** hint_binge, ai_reveal_rate, help_recruit_then_solo, deliberate hint time, co-primary solo_transfer_pass — demote hint count / help NPS as success.
- **Kill list:** Unlimited-hints hero; help NPS NS; avoidance-as-grit; Help Score™; parent hint-stalk; metacognition≡ACT points.
- **Growth:** Parent decks sell *earned independence of help*; tutors QA to instrumental prompts (SAFE-HITL × SAFE-TALK).
- **Vision:** Thirty-year identity company teaches Maya *when and how* to recruit support — then to need less of it — not a hint vending machine or a silence cult.

References (verified)

- Aleven, V., & Koedinger, K. R. (2000/2001). Investigations into help seeking and learning with a Cognitive Tutor (and related early EDM reports). PACT / CMU technical line; synthesized in Aleven et al. (2016).
 - Aleven, V., McLaren, B., Roll, I., & Koedinger, K. (2004). Toward tutoring help seeking: Applying cognitive modeling to meta-cognitive skills. In *ITS 2004* (pp. 227–239). Springer.
 - Aleven, V., McLaren, B., Roll, I., & Koedinger, K. (2006). Toward meta-cognitive tutoring: A model of help seeking with a Cognitive Tutor. *International Journal of Artificial Intelligence in Education*.
 - Aleven, V., McLaren, B. M., Roll, I., & Koedinger, K. R. (2016). Help helps, but only so much: Research on help seeking with intelligent tutoring systems. *International Journal of Artificial Intelligence in Education*, 26(1), 205–223. <https://doi.org/10.1007/s40593-015-0089-1>
 - Baker, R. S., Corbett, A. T., Koedinger, K. R., & Roll, I. (2005/2008). Detecting / adapting to when students game an intelligent tutoring system (gaming-the-system line; help abuse as gaming category).
 - Karabenick, S. A., & Knapp, J. R. (1991). Relationship of academic help seeking to the use of learning strategies and other instrumental achievement behavior. *Journal of Educational Psychology* (instrumental vs executive framing in help-seeking literature).
 - Nelson-Le Gall, S. (1981). Help-seeking: An understudied problem-solving skill in children. *Developmental Review*.
 - Roll, I., Aleven, V., McLaren, B. M., Ryu, E., Baker, R. S. J. d., & Koedinger, K. R. (2006). The Help Tutor: Does metacognitive feedback improve students' help-seeking actions? In *ITS 2006*.
 - Roll, I., Aleven, V., McLaren, B. M., & Koedinger, K. R. (2011). Improving students' help-seeking skills using metacognitive feedback in an intelligent tutoring system. *Learning and Instruction*, 21(2), 267–280. <https://doi.org/10.1016/j.learninstruc.2010.07.004>
 - Shih, B., Koedinger, K. R., & Scheines, R. (2008). A response time model for bottom-out hints as worked examples (EDM / Cognitive Tutor analyses; cited in Aleven et al., 2016).
 - Wood, H., & Wood, D. (1999). Help seeking, learning and contingent tutoring. *Computers & Education*, 33(2–3), 153–169.
-

Part XC — Product Analytics for FEI North Stars

Chapter status: Living evidence + instrumentation brief — Researcher tick 2026-08-06

Primary question: How should MindCraft instrument, govern, and commercially use FEI North Star events (retry_120s, challenge_accept, transfer_pass, solo_transfer_pass) without recreating vanity-metric theater or Goodhart-corrupted dashboards?

Owners: Product analytics · Engineering · Brand · HITL ops · Red Team · Founders

Commercial job: Ship a **SAFE-INSTRUMENT** doctrine: one FEI-aligned scoreboard that predicts durable competence and identity movement; demote DAU/streak/XP/accuracy as North Stars; validate leading events before marketing them; protect against metric gaming by students *and* by the company.

Builds on: Parts I.3 / VIII / IX Experiment C / XXI (Metrics Dictionary), XXXIV (claim ladder), XLIII (SAFE-HABIT), LXX (SAFE-LABMETA), LXXXVI (SAFE-PROOF). Immediate product ask in `NEXT_LAB.md`: align analytics with Part XXI. Sibling queue: 91 hint economy (peek cost once help metrics exist).

XC.1 Why this chapter exists

The Constitution already *named* the right events. Production still mostly optimizes what is easy to count: opens, minutes, streaks, tonight-% correct, tokens, “felt helpful.” That gap is not a documentation problem — it is a **commercial operating system** problem.

FOUNDER BELIEF under audit: If MindCraft cannot see Fear → Evidence → Identity in the telemetry, the company will sell the competitors’ story by accident: engagement theater (Duo-like), answer fluency (ChatGPT-like), or blocked accuracy (cram-pack-like). Identity transformation without instrumented FEI is a speech, not a product.

Claims we refuse as doctrine:

1. DAU / streak / XP / time-on-app as learning North Stars.
 2. Blocked-session accuracy or Map green-% as readiness.
 3. Coach thumbs-up / help NPS / “felt helpful” without `solo_transfer_pass`.
 4. Shipping Experiment A–D without pre-registered primary events.
 5. A single composite FEI Score™ sold to parents as proof.
 6. Optimizing `retry_120s` alone until students click-farm retries.
 7. Marketing “science-backed metrics” without Kane-style interpretation/use arguments.
 8. Analytics as surveillance shame for parents or tutors (SAFE-PDASH / SAFE-HITL dignity).
-

XC.2 Constructs

Construct	Research meaning	MindCraft analogue	Failure mode
Vanity metric	Looks good; does not change decisions (Croll & Yoskovitz)	Raw DAU, downloads, streak length	Fake growth
OMTM / NSM	Focus metric for stage vs enduring success metric	Stage OMTM = <code>retry_120s</code> ; NSM stack = challenge + transfer	Metric soup
Goodhart / Campbell	Target corrupts measure and process	Optimize accuracy → gaming	Corrupted FEI
Gaming the system	Exploit software regularities without learning (Baker)	Hint race, spam-retry, soft-wrong farm	Fake persistence
IUA (Kane)	Claims about <i>uses</i> of scores need evidence proportional to ambition	Parent deck claiming “identity” from 7-day retry	Overclaim
SAFE-INSTRUMENT	FEI events as governed product law	This chapter	Dashboard cosplay

Operational definition (HYPOTHESIS): Analytics are *SAFE-INSTRUMENT complete* when (a) the four FEI core events fire reliably in production, (b) Experiment C–style lagged prediction ranks them above streak/DAU for 8-week challenge-seeking, (c) anti-gaming companions (`hint_binge`, `ai_reveal_rate`, motive tags) are co-primary gates, (d) marketing uses only IUA-backed interpretations (L0–L2 until L3 evidence), and (e) no composite FEI Score™ is sold as a product hero.

XC.3 The four events — what they measure

FACT (internal doctrine spine): Part XXI already defines:

Event	Definition (abridged)	FEI role
<code>retry_120s</code>	New attempt within 120s of soft-wrong	Persistence under safety
<code>challenge_accept</code>	Chose harder level when offered (+ motive code)	Challenge-seeking
<code>transfer_pass</code>	Correct on varied item after mastery mark	Anti-fake-mastery
<code>solo_transfer_pass</code>	Transfer with AI/Solver closed or denied	Anti-crutch

HYPOTHESIS (working North Star composite, not a Score™): Challenge-seeking under safety *with* transfer — mastery-motive `challenge_accept`, return after miss (`retry_120s`), delayed mixed / solo transfer — predicts identity movement better than streaks (I.3; Experiment C prior).

Applied: Ship these four first. Secondary Part XXI events (`coach_shown`, `write_exit_to_retry`, `hint_binge`, ...) are *guards and mediators*, not substitute North Stars.

Kill: “We measure learning” while only shipping session accuracy + DAU.

Survive: Named FEI events with motive tags and transfer co-primaries.

XC.4 Vanity metrics lose — on purpose

FACT (product analytics practice): Lean Analytics distinguishes vanity metrics (impressive, non-actionable totals) from actionable metrics that change behavior; it further advocates focusing on a stage-appropriate One Metric That Matters rather than “data puking” dashboards (Croll & Yoskovitz, *Lean Analytics*, O’Reilly; OMTM essays at leananalyticsbook.com / O’Reilly).

Applied (FOUNDER BELIEF → testable): MindCraft’s *enduring* North Star is not a single SaaS OMTM forever — identity is lagging — but early-stage focus should pick **one FEI leading indicator at a time** (today: `retry_120s` under coach experiments) while keeping the **NSM stack** (`challenge_accept` + `transfer_pass` / `solo_transfer_pass`) as the company scoreboard that OMTMs must ladder into.

Commercial implication: Investor and parent decks may show growth oxygen (activation, retention) in an appendix. The *hero* learning claim must name FEI events or dated solo-transfer artifacts (SAFE-PROOF) — never streak length.

Kill: Streak / DAU / XP velocity as learning proof.

Survive: FEI NSM stack + stage OMTM discipline.

XC.5 When measures become targets — Campbell & Goodhart

FACT: Campbell’s law — the more a quantitative social indicator is used for decision-making, the more subject it is to corruption pressures and the more apt it is to distort the processes it monitors (Campbell, 1979, *Evaluation and Program Planning*, 2(1), 67–90, doi:10.1016/0149-7189(79)90048-X).

FACT: Goodhart’s law (and Strathern’s popular paraphrase) — when a measure becomes a target, it ceases to be a good measure — is the economics twin of the same corruption dynamic (Goodhart, 1975 monetary-policy statement; Strathern, 1997 paraphrase widely cited in accountability debates).

Applied (HYPOTHESIS): If tutors are paid on `retry_120s`, students will be coached to tap retry without sense-making. If marketing celebrates rising `challenge_accept` without motive codes, appearance-climate CTA wins (SAFE / TARGET / Part XXXVIII). If `transfer_pass` is gamed with isomorphic clones, fake mastery returns.

Design rule: Never attach high-stakes pay, shame ranks, or public heroboards to a *single* FEI event. Use **co-primary gates**: raise retry *without* raising `hint_binge` / `ai_reveal_rate`; raise `challenge_accept` *with* mastery-motive tag; raise transfer *with* declared hop + delayed mix (XXI.4).

Kill: FEI Score™ / single-KPI tutor commissions / parent shame ranks on retry.

Survive: Multi-metric decision rule + qualitative spot-checks (SAFE-LABMETA).

XC.6 Student gaming is also company risk

FACT: In intelligent tutoring systems, “gaming the system” — succeeding by exploiting software regularities (hint abuse, systematic guessing) rather than learning — is common and associated with poorer learning (Baker, Corbett, Koedinger & Wagner, 2004 line; Baker et al., 2009, AIED — tutor features explained a large share of gaming variance; Cocea, HersHKovitz & Baker — harmful gaming wastes learning opportunities).

FACT: Help-abuse / executive racing is already Constitution law under SAFE-HELP (Aleven/Koedinger; Roll et al., 2011; Aleven et al., 2016) — analytics must treat `hint_binge` and `help_executive_race` as *first-class companions* to FEI wins, not afterthoughts.

Applied: Instrument anti-gaming detectors *in the same ship* as FEI events. A week where `retry_120s` rises and `hint_binge` rises is **not** a FEI win — it is Campbell corruption in progress.

Kill: Celebrate raw retry/challenge lifts without gaming companions.

Survive: XXI.4 decision rule as shipping gate.

XC.7 Validity of uses — Kane before decks

FACT: Kane’s argument-based approach holds that what is validated are proposed *interpretations and uses* of scores, not the numbers themselves; more ambitious claims require more evidence; evaluating score *uses* requires evaluating consequences (Kane, 2013, *Journal of Educational Measurement*, 50(1), 1–73, doi:10.1111/jedm.12000; see also Kane, 1992, *Psychological Bulletin*).

Applied (HYPOTHESIS): Emitting `retry_120s` is an L0/L1 engineering fact. Saying “FEI is working” from a 2-week A/B is L1. Saying “we build math identity” from those events without 26w triangulation is an L3 claim without warrant (SAFE-LONGID / claim ladder XXXIV). Parent-facing “proof” should prefer dated solo-transfer artifacts (SAFE-PROOF) over composite meters.

Commercial implication: Marketing and sales get an **IUA card** per claim: event → interpretation → allowed use → required evidence → banned overclaim. No card → no ad.

Kill: “Science-backed identity growth” from unvalidated dashboards.

Survive: Claim-laddered FEI reporting + Experiment C lagged prediction.

XC.8 Product surface — instrumentation contract

Surface	Required event / behavior	Banned substitute
Soft-wrong path	<code>retry_120s</code> , <code>coach_shown</code> , <code>binge guards</code>	Thumbs-up only
Level offer	<code>challenge_accept</code> + motive enum (mastery/appearance/norm)	Harder-click without why
Mastery mark	Delayed <code>transfer_pass</code> probe scheduled	Instant blocked green
Solver / AI	<code>solo_transfer_pass</code> / <code>ai_reveal_rate</code>	“Felt helpful” chat NPS
Experiment A	Pre-reg primary = <code>retry_120s</code>	Engagement-only success
Parent deck	Proof-age + named hop artifact	Streak / tonight-% hero
Tutor QA	FEI sample + talk/help guards	Hours booked NS
Board OMTM	One FEI leading metric per stage	Metric soup / DAU hero

Competitive foil: Duo wins streak oxygen. ChatGPT wins executive completion counts. Cram packs win blocked accuracy. Minecraft differentiates by publishing (internally first, externally carefully) **persistence under safety + challenge with motive + delayed solo transfer** — the FEI scoreboard competitors cannot honestly claim without changing their products.

XC.9 Doctrine — SAFE-INSTRUMENT (provisional)

1. **Ship the four** — `retry_120s`, `challenge_accept` (+motive), `transfer_pass`, `solo_transfer_pass` before any new vanity KPI.
2. **NSM stack, stage OMTM** — enduring scoreboard is FEI+transfer; current sprint focuses one leading event.
3. **Co-primary gates** — XXI.4: no FEI “win” if binge/reveal rises or transfer falls.
4. **Anti-Goodhart** — no single-event pay, shame, or public heroboard; rotate checks; keep qualitative audits.
5. **Anti-gaming companions** — `hint_binge`, `help_executive_race`, `ai_reveal_rate` ship with FEI.
6. **Kane IUA for uses** — decks/ads need interpretation/use cards; ambition $\leq$ evidence (L0–L4).
7. **Dignity** — analytics for product learning and tutor coaching, not parent/student humiliation.
8. **Copy**: “We measure return after a miss, choosing harder for mastery reasons, and solving later without the model.” Never “we’re #1 for engagement.”

Confidence: High that vanity engagement metrics are weak durable-learning proxies (Lean Analytics practice + Constitution HID / SAFE-HABIT). High that Campbell/Goodhart corruption applies once FEI events become high-stakes targets. High that ITS gaming literature predicts students will exploit naive persistence metrics (Baker line). High that Kane IUA is the right hygiene for marketing uses. Medium that Experiment C will rank FEI events above streaks for 8-week challenge-seeking (pre-registered prior, not yet run). Medium on exact 120s window / motive enum reliability until instrumentation QA.

XC.10 Experiments

ID	Question	Design	Primary
INSTR-1	Do <code>retry_120s</code> / mastery-motive <code>challenge_accept</code> / <code>transfer_pass</code> beat streak & DAU predicting 8w challenge-seeking?	Lagged observational (Experiment C densify)	Ranked predictive validity; calibration plots
INSTR-2	Does publishing <code>retry_120s</code> as tutor KPI without co-gates raise binge/spam-retry?	Ops A/B (gate vs single KPI)	<code>hint_binge</code> ; true retry quality sample
INSTR-3	Soft-wrong → coach instrumentation completeness vs dark funnel	Audit + fix ship	Event fire rate; missingness by surface
INSTR-4	Motive-coded <code>challenge_accept</code> vs uncoded harder-click	A/B CTA	Mastery-motive share; later transfer
INSTR-5	Parent CBC: FEI report card vs streak/accuracy hero deck	CBC (SAFE-WTP)	WTP; trust; comprehension

INSTR-QUAL	10 Maya + 5 tutors: which dashboard numbers feel honest vs gamable?	Qual	Corruption codebook
------------	--	------	---------------------

Falsifier: Streak/DAU predict 8w challenge-seeking and 26w identity as well as FEI events → still ban as *marketing* North Stars (habit≠identity); allow as oxygen metrics only.

Falsifier: Co-primary gates stall shipping for >1 quarter with no better learning outcomes → simplify to two-event gate (retry_120s + solo_transfer_pass) temporarily; keep anti-gaming.

Falsifier: Motive tags unreliable / demand-laden → drop parent-facing motive; keep internal coding research-only.

Pre-register: INSTR-* / Experiment C before any “our North Star proves identity” campaign (SAFE-LABMETA). Do not invent FEI Score™ packaging experiments that recreate Belief Score™ failures (SAFE-CALIB).

XC.11 So what for MindCraft commercially

- **Copy:** Lead with “return after a miss,” “choose harder to learn,” “prove it later without the model.” Ban engagement/streak/accuracy hero claims as learning proof.
- **Product:** Instrument Part XXI FEI core + gaming companions this quarter; wire Experiment A primary to retry_120s.
- **Positioning:** Against Duo streak oxygen, GPT completion counts, and cram accuracy — for a FEI scoreboard tied to solo transfer.
- **Metric law:** NSM = challenge-seeking under safety with transfer; OMTM rotates among FEI leadings; vanity demoted.
- **Kill list:** DAU/streak/XP NS; blocked accuracy NS; FEI Score™; single-KPI tutor pay; thumbs-up=learning; science-backed identity ads without IUA.
- **Growth:** Parent/LEA decks sell dated proof + FEI process honesty (SAFE-PROOF × SAFE-PDASH); sales enablement uses IUA cards.
- **Vision:** A thirty-year identity company measures the loop that builds identity — fear metabolized into evidence — not the vanity that only looks like school.

References (verified)

- Baker, R. S., Corbett, A. T., Koedinger, K. R., & Wagner, A. Z. (2004). Off-task behavior in the Cognitive Tutor classroom: When students “game the system.” In *Proceedings of ITS / related Cognitive Tutor classroom studies* (gaming-the-system foundational line).
- Baker, R. S. J. d., de Carvalho, A. M. J. B., Raspat, J., Aleven, V., Corbett, A. T., & Koedinger, K. R. (2009). Educational software features that encourage and discourage “gaming the system.” In *AIED 2009* (pp. 475–482). IOS Press. doi:10.3233/978-1-60750-028-5-475.
- Campbell, D. T. (1979). Assessing the impact of planned social change. *Evaluation and Program Planning*, 2(1), 67–90. [https://doi.org/10.1016/0149-7189\(79\)90048-X](https://doi.org/10.1016/0149-7189(79)90048-X)
- Cocea, M., HersHKovitz, A., & Baker, R. S. J. d. (2009). The impact of off-task and gaming behaviors on learning: Immediate or aggregate? (EDM / ITS analyses of harmful gaming).

- Croll, A., & Yoskovitz, B. (2013). *Lean Analytics: Use Data to Build a Better Startup Faster*. O'Reilly Media. (Vanity vs actionable metrics; One Metric That Matters.)
 - Goodhart, C. A. E. (1975). Problems of monetary management: The UK experience. (Origin of Goodhart's law; later paraphrased in accountability literature.)
 - Kane, M. T. (1992). An argument-based approach to validity. *Psychological Bulletin*, 112(3), 527–535.
 - Kane, M. T. (2013). Validating the interpretations and uses of test scores. *Journal of Educational Measurement*, 50(1), 1–73. <https://doi.org/10.1111/jedm.12000>
 - Strathern, M. (1997). “Improving ratings”: Audit in the British university system. *European Review* (popular “when a measure becomes a target...” paraphrase of Goodhart in accountability contexts).
 - Internal Constitution: Parts I.3, VIII.3, IX Experiment C, XXI Metrics Dictionary; SAFE-HABIT / SAFE-HELP / SAFE-PROOF / SAFE-LABMETA / claim ladder XXXIV.
-

Part XCI — Hint Economy & Contingent Scaffolding

Chapter status: Living evidence + Solver/Practice hint-ladder brief — Researcher tick 2026-08-06

Primary question: How should MindCraft price the *cost of a peek* (effort, stage, and timing — not humiliation) and make scaffolding *contingent* (up after miss, down after success) so soft principle hints beat free bottom-out racing — without renaming ChatGPT dump as “adaptive help”?

Owners: Product (Solver / Practice hints) · Engine (fade×hint state) · HITL tutors · Brand · Red Team

Commercial job: Ship a **SAFE-HINT** doctrine: soft before hard; contingency before menu; peek costs construction; fade×hint as one state machine; prove with solo transfer, not hint-click engagement.

Builds on: Parts XXVI (fading), XL (SAFE-SE), LXXXVIII (SAFE-FADE), LXXXIX (SAFE-HELP), XC (SAFE-INSTRUMENT). Product seams: hint levels, hint_binge, help_deliberate_ms, help_executive_race, fade_stage, SE gates. Sibling: id 92 explanation length (token tax on monologue hints).

XCI.1 Why this chapter exists

SAFE-HELP answered *when* help is instrumental vs abusive. SAFE-FADE answered *what stage* of worked guidance belongs on the learn rail. This chapter answers the *microeconomics of the hint button*: peek cost, soft vs hard levels, and contingency (Wood) + fading (van de Pol) so Maya constructs instead of shopping for answers.

FOUNDER BELIEF under audit: Free, instant, hard peeks train executive completion. Making the next *principle* cheap and the *answer* expensive — in effort, not shame — converts Solver from a dump into an identity machine, if parents hear “help that still asks you to think” and students still reach solo transfer.

Claims we refuse as doctrine:

1. Costless unlimited peeks / free bottom-out as learning North Star.
2. Soft-vs-hard confusion: treating static “hard scaffolds” (Saye & Brush) as the same as hard (=answer) hints.
3. Hint clicks, Hint Score™, or XP-for-hints as engagement success.
4. Non-contingent fixed chrome (same hard dump forever, or never any hint).

5. Fade ladder and hint ladder as two disconnected products.
6. Humiliation economy (shame timers, public peek counts) as “cost.”
7. “AI always picks the perfect hint” black-box brand without inspectable stage.
8. Points / streaks unlocked by peeking.

XCI.2 Constructs (keep the two “soft/hard” axes separate)

Construct	Research meaning	MindCraft analogue	Failure mode
Scaffolding (classic)	Temporary support that enables success beyond unassisted ability	Coach / hint / tutor prompt	Permanent crutch
Contingency	Increase control after failure; decrease after success	Fade up/down + hint level	Fixed menu
Fading / transfer of responsibility	Support shrinks as competence grows	SAFE-FADE E0–E3	Never fade / never help
Soft hint (ITS)	Principle / goal / critical-feature prompt; not the answer	Principle card; “which join?”	Vague pep talk
Hard hint (ITS)	Near-answer / bottom-out	Full step reveal / expression dump	Executive race
Hard scaffold (Saye & Brush)	Static, preplanned support in the environment	Ingredient card templates; completion blanks	Mistaken for “hard hint”
Soft scaffold (Saye & Brush)	Dynamic human/peer contingent aid	Tutor HITL; adaptive coach	Mistaken for “soft hint”
Peek cost	Friction that preserves sense-making before answer	SE gate; dwell; stage lock	Shame timer
SAFE-HINT	Contingent soft→hard ladder with priced peeks	This chapter	Free dump

Operational definition (HYPOTHESIS): A hint path is *SAFE-HINT complete* when it (a) opens at a soft/principle grain unless diagnosis already requires harder support, (b) advances hard grain only after miss or deliberate request *with* attempt/SE, (c) decreases control after consecutive solo successes (contingency × fade), (d) makes bottom-out costly in *construction effort* (not humiliation), (e) never awards XP/streak for peeks, and (f) co-primary proves with `solo_transfer_pass` and lower `help_executive_race`.

XCI.3 Scaffolding is contingent by definition — not a hint vending machine

FACT (origin metaphor): Wood, Bruner & Ross (1976, *Journal of Child Psychology and Psychiatry*, 17(2), 89–100) — tutoring as scaffolding: recruitment, reduction of degrees of freedom, direction maintenance, marking critical features, frustration control, and demonstration when the learner can recognize the solution.

FACT (three characteristics): van de Pol, Volman & Beishuizen (2010, *Educational Psychology Review*, 22, 271–296) — scaffolding reviews converge on **contingency**, **fading**, and **transfer of responsibility**. Contingency (Wood, Wood & Middleton, 1978 line, carried forward): increase degree of control after failure; decrease after success.

FACT (computer contingent tutoring): Wood & Wood (1999, *Computers & Education*, 33(2–3), 153–169) — QUADRATIC algebra tutor implementing contingent support; help-seeking process measures predict learning beyond pretest; prior knowledge moderates how help use relates to outcomes.

Applied (HYPOTHESIS): MindCraft’s hint button must encode *contingency*, not a flat shelf of equally cheap peeks. After two soft-wrongs, offer the next soft grain (mark critical feature / reduce degrees of freedom). After a clean solo, lower default control on the next isomorphic item. That is the same state machine as SAFE-FADE — not a second product.

Kill: Unlimited student-chosen answer menu labeled “scaffolding.”

Survive: Contingent control that moves with evidence.

XCI.4 Soft before hard — Cognitive Tutor ladder evidence

FACT (hint sequences): Cognitive Tutors typically expose multi-level on-demand hints from general/goal reminders → increasingly specific → **bottom-out** near-answer (Aleven & Koedinger, 2000/2001; reviewed in Aleven, McLaren, Roll & Koedinger, 2016, *IJAIED*, 26(1), 205–223, doi:10.1007/s40593-015-0089-1).

FACT (abuse pattern): Students often race to bottom-out; large shares of pre-bottom levels viewed under ~1s; when help is requested, students frequently advance through all levels to the answer (Aleven & Koedinger, 2000/2001; Aleven et al., 2016). Help abuse correlates negatively with learning in multiple analyses; Baker’s gaming line treats rapid help abuse as a primary gaming category.

FACT (caveat): Shih, Koedinger & Scheines (2008) — time spent with bottom-out can correlate positively with learning when treated as a worked-example opportunity (self-explanation), not as copy-paste (cited in Aleven et al., 2016). Soft-before-hard still holds; hard-without-sense-making fails.

Naming hygiene (FACT about literature collision): Saye & Brush (2002) call *hard scaffolds* the static, preplanned supports and *soft scaffolds* the dynamic teacher/peer supports. That axis is orthogonal to ITS soft/hard *hint specificity*. Product copy must not say “hard scaffolding” when it means “answer dump,” or parents will hear cruelty; tutors must not hear “soft scaffolding” and think “vague AI pep talk.”

Applied (FOUNDER BELIEF → testable): Default open = soft hint (critical feature / principle / “name the bridge”). Hard hint = gated. Bottom-out = last grain + SE/attempt tax (SAFE-HELP × SAFE-SE).

Kill: Soft and hard as brand poetry without ladder semantics.

Survive: Named levels: Goal → Principle → Join cue → Step sketch → Bottom-out.

XCI.5 The cost of a peek — effort economy, not shame economy

HYPOTHESIS (hint economy): If hard peeks are free in time, clicks, and social cost, rational students (and anxious ones) buy the answer. Learning products that “remove friction from help” optimize completion SaaS. Identity products must put friction on *executive* peeks while keeping *instrumental* peeks low-cost.

Legitimate costs (product priors):

1. **Construction cost** — write a why / fill a blank / name the FormatId or bridge before next hard level (SAFE-SE / SAFE-FADE).
2. **Stage cost** — hard levels unavailable until soft grain shown *or* diagnosed struggle (contingency).
3. **Dwell cost** — minimum readable time before next level advances (anti-race; Help Tutor behavioral signature).
4. **Budget cost (optional)** — limited hard peeks per session that refresh with solo successes (not daily streak theater).

Illegitimate costs (kills): Public peek shaming; parent live stalk of hint count; countdown humiliation; XP penalties that feel like punishment for asking; Anxiety Score™ gated on peeks.

FACT (support): Roll et al. (2011, *Learning and Instruction*, 21(2), 267–280, doi:10.1016/j.learninstruc.2010.07.004) — metacognitive feedback increased deliberate help (more time, fewer levels) without reliable domain-gain magic (Aleven et al., 2016). Cost that changes *behavior* is not automatically an ACT-point engine — co-primary stays transfer/solo (SAFE-INSTRUMENT).

Kill: Free hard peek as brand kindness.

Survive: Cheap soft peek + expensive hard peek in *effort*.

XCI.6 Fade × hint — one state machine

HYPOTHESIS: Running SAFE-FADE (E0 full example → E3 solo) and a separate always-on hard-hint menu recreates ChatGPT inside the “faded” product. The hint ladder *is* the local control knob inside a fade stage:

Fade stage	Default hint open	Hard peek policy
E0 (full example study)	Soft marks on the example	Bottom-out redundant — ban dump-over-example
E1 (completion)	Soft → principle on blank	Hard only after miss + SE on blank
E2 (sparse scaffold)	Soft first	One hard grain / attempt budget
E3 (near solo)	Soft invite after 2 misses	Hard rare; prove rail near zero

SPECULATION: Inspectable chip (“Stage E1 · next soft: name the join”) beats black-box “AI chose your hint” for parent trust and tutor QA (SAFE-ADAPT inspectability pattern).

Kill: Fade theater with free hard menu underneath.

Survive: Coupled fade×hint state; tutor briefs read the same chip.

XCI.7 Product surface — SAFE-HINT claim contract

Surface	Required behavior	Banned substitute
Solver learn rail	Soft-first ladder; hard behind construction cost	Instant bottom-out
Practice hints	Contingent up/down; dwell anti-race	Hint spam = XP
Prove / exam rail	Soft only or none; no hard peeks as default	Learn-rail dump under stakes
Tutor HITL	Soft prompt before pour; match stage chip	Tutor as hard-hint vending
Parent view	Aggregate deliberate-help / solo — not peek ranks	Live “hints used” shame
Marketing	“Next-step hints that still ask you to think”	“Unlimited free hints / never stuck”
Analytics	help_deliberate_ms, help_executive_race, peek-budget use	Hint Score™ / peeks-as-DAU

Competitive foil: ChatGPT = costless hard peek. Struggle-only apps = infinite soft silence. Duo-like loops = XP for peeks. MindCraft = priced peeks under contingency that graduate into solo transfer.

XCI.8 Doctrine — SAFE-HINT (provisional)

1. **Soft before hard** — principle / critical-feature before bottom-out; name levels in UX.
2. **Contingency is law** — control up after failure, down after success (Wood; van de Pol).
3. **Peek cost = construction** — SE/blank/dwell/stage; never humiliation or parent stalk.
4. **Fade × hint = one machine** — no free hard menu under a “faded” costume.
5. **Two soft/hard vocabularies stay separate** — ITS hint specificity ≠ Saye & Brush hard/soft scaffolds.
6. **No reward for peeking** — no XP, streak, or Hint Score™ for reveals.
7. **Inspectable stage** — chip states why this grain; ban perfect-AI-hint mystique.
8. **Proof** — solo_transfer_pass + lower executive race; deliberate soft use OK.

Confidence: High that contingency/fading/transfer-of-responsibility define scaffolding (van de Pol et al., 2010; Wood et al., 1976). High that free bottom-out racing is common and often anti-learning (Aleven/Koedinger; Aleven et al., 2016; Baker). High that soft-before-hard is the right default. Medium that specific peek-budget numbers generalize — run HINT-*. High that shame economies violate SAFE-PDASH. Medium parent CBC for “next-step cost” copy (HINT-5). High that fade×hint decoupling recreates ChatGPT.

XCI.9 Experiments

ID	Question	Design	Primary
HINT-1	Soft-first contingent ladder vs free hard menu	A/B within concept	solo_transfer_pass; help_executive_race; near transfer
HINT-2	SE/blank cost before hard vs dwell-only vs neither	A/B/C	Deliberate ms; binge; transfer
HINT-3	Coupled fade×hint vs faded stage + free hard submenu	A/B	Solo transfer; hard-peek rate
HINT-4	Peek budget (N hard/session, refill on solo) vs unlimited hard	A/B	Executive race; avoidance; retention
HINT-5	Parent CBC: “next-step hints” vs “unlimited free hints” vs “we never hint”	CBC	WTP; trust (SAFE-WTP)
HINT-QUAL	10 Maya: when does a peek feel earned vs stolen?	Qual	Soft/hard phenomenology codebook

Falsifier: Free hard menu wins delayed solo transfer *and* 26w identity equally → still ban as *brand* vs ChatGPT; allow rare emergency reveal off-learn-rail.

Falsifier: Peek budgets spike avoidance and drop retention with no solo gain → drop budgets; keep soft-first + SE cost.

Falsifier: Soft-first harms true novices vs E0 full example → keep E0 stage; do not open with hard silence.

Pre-register: HINT-* before any “unlimited free hints / AI perfect peeks / Hint Score™” campaign (SAFE-LABMETA). Do not duplicate HELP-1 pricing arms — HINT owns peek-cost; HELP owns instrumental vs executive typology.

XCI.10 So what for MindCraft commercially

- **Copy:** “Next-step hints that still ask you to think.” Lead with soft-first + solo proof. Never “unlimited free hints so you’re never stuck.”
- **Product:** Contingent soft→hard ladder coupled to fade stage; construction cost on hard peeks; inspectable stage chip.
- **Positioning:** Against ChatGPT costless dump and struggle-theater silence; for an *economy of peeks* that transfers responsibility.
- **Metric:** help_deliberate_ms, help_executive_race, hard-peek rate by stage, co-primary solo_transfer_pass — demote raw hint count / Hint Score™.

- **Kill list:** Free hard peeks as kindness brand; Hint Score™ / XP-for-hints; shame timers; black-box perfect-hint AI; fade costume with free dump underneath.
 - **Growth:** Parent decks sell earned independence of peeks; tutors QA soft-before-pour (SAFE-HITL × SAFE-TALK).
 - **Vision:** Teach Maya to recruit the smallest scaffold that works — then to need less of it.
-

References (verified)

- Aleven, V., & Koedinger, K. R. (2000). Limitations of student control: Do students know when they need help? In *ITS 2000* / related PACT reports; see also Aleven & Koedinger (2001) help-seeking investigations.
 - Aleven, V., McLaren, B. M., Roll, I., & Koedinger, K. R. (2016). Help helps, but only so much: Research on help seeking with intelligent tutoring systems. *International Journal of Artificial Intelligence in Education*, 26(1), 205–223. <https://doi.org/10.1007/s40593-015-0089-1>
 - Baker, R. S., Corbett, A. T., Koedinger, K. R., & colleagues. Gaming-the-system line (help abuse as a primary gaming category; multiple EDM/AIED papers 2004–2013).
 - Roll, I., Aleven, V., McLaren, B. M., & Koedinger, K. R. (2011). Improving students’ help-seeking skills using metacognitive feedback in an intelligent tutoring system. *Learning and Instruction*, 21(2), 267–280. <https://doi.org/10.1016/j.learninstruc.2010.07.004>
 - Saye, J. W., & Brush, T. (2002). Scaffolding critical reasoning about history and social issues in multimedia-supported learning environments. *Educational Technology Research and Development*, 50(3), 77–96. (Hard vs soft *scaffolds* axis — distinct from ITS soft/hard hint specificity.)
 - Shih, B., Koedinger, K. R., & Scheines, R. (2008). A response time model for bottom-out hints as worked examples (EDM analyses; cited in Aleven et al., 2016).
 - van de Pol, J., Volman, M., & Beishuizen, J. (2010). Scaffolding in teacher–student interaction: A decade of research. *Educational Psychology Review*, 22, 271–296. <https://doi.org/10.1007/s10648-010-9127-6>
 - Wood, D., Bruner, J. S., & Ross, G. (1976). The role of tutoring in problem solving. *Journal of Child Psychology and Psychiatry*, 17(2), 89–100.
 - Wood, H., & Wood, D. (1999). Help seeking, learning and contingent tutoring. *Computers & Education*, 33(2–3), 153–169.
-

Part XCII — Explanation Length vs Germane Load

Chapter status: Living evidence + coach/Solver length brief — Researcher tick 2026-08-06

Primary question: When does a coach or Solver *explanation* create germane processing (schema-building attention to the join) versus a **token tax** — extraneous words that crowd working memory, suppress self-explanation, and cosplay as “thorough teaching”?

Owners: Product (coach copy / Solver wrap) · Engine (prompt budgets) · HITL tutors · Brand · Red Team

Commercial job: Ship a **SAFE-EXPLAIN** doctrine: short principle-first coach beats monologue; length follows expertise and element interactivity; student-generated why beats provided essay; prove with transfer/solo, not word count or “AI explained fully” satisfaction.

Builds on: Parts XXVI (CLT / fading), XL (SAFE-SE), LXXXV (SAFE-STORYLOAD), LXXXVIII (SAFE-FADE), LXXXIX (SAFE-HELP), XCI (SAFE-HINT). Sibling queue: id 93 retrieval failure modes. Product seams: coach card max length, soft-hint grain, SE-before-wrap, monologue-cap telemetry.

XCII.1 Why this chapter exists

SAFE-HINT priced the *peek*. SAFE-FADE staged the *worked example*. SAFE-SE required the *student why*. This chapter prices the *words after the why* (and the words that replace it): how long may a coach speak before thoroughness becomes overload?

Default AI tutor UX optimizes for fluent completeness — multi-paragraph solutions that feel caring. Default parent anxiety asks for “more explanation.” Default LLM systems reward tokens. MindCraft’s identity claim needs the opposite default: **the shortest explanation that marks the critical feature and leaves room for construction**.

FOUNDER BELIEF under audit: Verbose AI wrap trains passive receipt. Principle-length coach + SE + fade stage trains transferable competence — if parents hear “clear, not endless,” and students still reach solo transfer.

Claims we refuse as doctrine:

- 1. Longer explanation $\equiv$ better teaching / more learning.
 - 2. Token count, “full solution essay,” or Explanation Score™ as quality North Stars.
 - 3. Always-on monologue wrap that displaces self-explanation (Kill #8 / SAFE-SE).
 - 4. Same explanation length for novice and near-expert (expertise reversal ignore).
 - 5. Seductive lore / pep talk / franchise trivia inside the coach card (SAFE-STORYLOAD).
 - 6. “AI explained everything” as marketing proof of pedagogy.
 - 7. Unlimited free hard peeks *plus* essay dumps (SAFE-HINT $\times$ this chapter).
 - 8. Streaks/XP for reading long explanations.
-

XCII.2 Constructs (keep load types honest)

Construct	Research meaning	MindCraft analogue	Failure mode
Intrinsic load	Element interactivity of the math itself	Concept + bridge + FormatId difficulty	Soften the math instead of the wrap
Extraneous load	Processing that does not build the schema	Redundant essay, seductive asides, duplicate modalities	“Helpful” fluff
Germane processing	WM devoted to intrinsic structure (modern CLT: redistribute, not a third additive pile)	Naming the join; SE; marking critical features	Calling any effort “germane”

Token tax	Product term: words that consume WM without advancing schema	Multi-paragraph Solver dump	Fluency-as-care
Instructional explanation	Provided why/principle/operator text	Coach card; soft hint text	Essay replacing attempt
Self-explanation (SE)	Learner-generated why	SAFE-SE prompts	AI monologue labeled SE
SAFE-EXPLAIN	Length-contingent, principle-first, SE-preserving coach	This chapter	Monologue costume

Operational definition (HYPOTHESIS): An explanation path is *SAFE-EXPLAIN complete* when it (a) defaults to principle / critical-feature length ($\approx 1\text{--}3$ sentences or one marked step), (b) expands only after miss, demand, or diagnosed high element interactivity, (c) never replaces a required SE with a provided essay, (d) shortens as fade stage / expertise rises, (e) strips seductive and redundant material, and (f) co-primaries prove with `solo_transfer_pass` / near transfer — not dwell-on-essay or thumbs-up on “thoroughness.”

XCII.3 Germane load is not “more words”

FACT (modern CLT): Sweller, van Merriënboer & Paas (2019, *Educational Psychology Review*, 31, 261–292, doi:10.1007/s10648-019-09465-5) reconceptualize germane load: it is working-memory resources devoted to *dealing with intrinsic load*, not a third additive load that rises when you add “learning activities.” Reducing extraneous load frees capacity for intrinsic processing; germane is redistributive.

FACT (element interactivity): Sweller (2010, *Educational Psychology Review*, 22, 123–138, doi:10.1007/s10648-010-9128-5) — intrinsic load tracks element interactivity relative to expertise; instructional design should manage extraneous load and optimize germane *processing of* that intrinsic structure.

Applied (HYPOTHESIS): A 600-token coach monologue that re-derives every algebraic step the student already executed is mostly **extraneous** for a mid-fade learner — even if every sentence is “about math.” Germane would be: mark the failed join, name the FormatId shift, invite one SE. Length is not a virtue metric.

Kill: “We maximized germane load by adding more explanation.”

Survive: Cut extraneous wrap so WM can work the intrinsic join.

XCII.4 Provided explanations: small effects; SE often wins

FACT (meta-analysis): Wittwer & Renkl (2010, *Educational Psychology Review*, 22, 393–409, doi:10.1007/s10648-010-9136-5) — 21 studies / 28 comparisons of worked examples *with* vs *without* instructional explanations. Overall weighted mean effect **small** ($d = 0.16$). Benefits stronger for **conceptual knowledge** ($d \approx 0.36$) than for transfer/problem-solving measures (often ns). When controls were **prompted to self-explain**, instructional

explanations showed **no advantage** (subset $d \approx -0.01$).

FACT (withholding can beat providing): Richey & Nokes-Malach (2013, *Learning and Instruction*, 25, 22–34, doi:10.1016/j.learninstruc.2012.11.006) — across three experiments, students who had stepwise instructional explanations *withheld* showed **greater conceptual learning** than students who received them; achievement goals predicted learning more under withholding (constructive effort was volitional).

FACT (framework): Wittwer & Renkl (2008, *Educational Psychologist*, 43(1), 49–64) — instructional explanations often fail when mistimed, misleveled, or unused; design must fit generation *and* use conditions — not assume “explain more.”

Applied (FOUNDER BELIEF → testable): MindCraft’s coach should prefer **SE prompts + short principle marks** over default essay wraps. Long provided explanations are a *last* grain (like hard hints), not the brand voice.

Kill: AI monologue $\equiv$ instructional science; “we explain thoroughly” as FEI proof.

Survive: Minimal instructional text that enables construction; expand only when SE stalls.

XCII.5 Coherence, redundancy, and the token tax

FACT (coherence): Mayer’s coherence principle — people learn better when extraneous material is excluded (supported across many multimedia tests; median ES often large in Mayer handbook summaries). Seductive details (interesting but irrelevant words/images/sounds) hurt retention/transfer (Harp & Mayer line; Moreno & Mayer, 2000, *Journal of Educational Psychology*, 92, 117–125 on minimizing irrelevant sounds).

FACT (redundancy): Chandler & Sweller (1991) / Sweller line — presenting the same information in unnecessary duplicate forms increases extraneous load. Mayer restricted redundancy (narration + on-screen text with animation) similarly taxes capacity (Mayer, Heiser & Lonn / Moreno & Mayer series; median ES ~ 0.7 in “nine ways” summaries).

Applied (HYPOTHESIS): Solver surfaces that show (1) full worked steps, (2) a paragraph restating those steps, and (3) a chatty pep layer create **content redundancy + seductive fluff**. Prefer one representation at a time: sparse step marks *or* short principle — not both plus lore.

Product prior: Soft hint = principle sentence. Hard peek = one revealed step. Bottom-out = minimal completed blank — not a TED talk. Story wrap stays outside the coach card budget (SAFE-STORYLOAD).

Kill: Duplicate text+voice essays; “engaging” asides inside the critical coach moment.

Survive: One clear signal; cut the rest.

XCII.6 Expertise reversal: length must fade

FACT: Kalyuga, Ayres, Chandler & Sweller (2003, *Educational Psychologist*, 38(1), 23–31) — instructional supports that help novices can **hurt** relative experts (expertise reversal). Detailed guidance becomes redundant; experts waste WM reconciling external text with already-formed schemas.

FACT (design implication): Kalyuga (2007, *Educational Psychology Review*, 19, 509–539) — tailor and fade guidance as expertise rises; fixed high-guidance designs are not “always safer.”

Applied (HYPOTHESIS): Couple explanation *length* to SAFE-FADE stage and SAFE-HINT contingency:

Stage / expertise	Default coach length	Expand when
E0 novice example study	Short principle + marked critical feature on the example	Student SE fails / high anxiety + miss
E1 completion	One cue for the blank	Second miss
E2 sparse	Principle only	Demand + dwell
E3 near solo / prove rail	Invite SE; near-zero provided essay	Rare emergency reveal

Kill: Same multi-paragraph wrap for every student forever.

Survive: Length fades with responsibility transfer.

XCII.7 Assistance dilemma — not “more help”

FACT: Koedinger & Aleven (2007, *Educational Psychology Review*, 19, 239–264) — the **assistance dilemma**: how much information to provide vs withhold so that learners still do the constructive work that builds robust knowledge.

Bridge to product: SAFE-HELP / SAFE-HINT already reject unlimited executive dumps. SAFE-EXPLAIN rejects *verbose instrumental* dumps — explanations that look helpful, consume tokens, and still steal the generative work Richey & Nokes-Malach show matters for conceptual gain.

SPECULATION: LLM tutors systematically bias toward over-assistance because fluency and completeness are RLHF-shaped. Without an explicit length budget and SE gate, MindCraft will drift ChatGPT-ward even with “principle” system prompts.

XCII.8 Product surface — SAFE-EXPLAIN claim contract

Surface	Required behavior	Banned substitute
Soft coach / soft hint	$\leq$ principle grain; mark join/FormatId	Multi-paragraph lecture
Soft-wrong path	SE before wrap; wrap $\leq$ short correction	Essay before student why
Solver learn rail	Length tied to fade stage	Always-full verbal restatement of steps
Prove / exam rail	Minimal KR; no lecture	Post-item monologue
Tutor HITL	Prompt $>$ pour; pour stays short (SAFE-TALK)	Tutor as essay machine

Parent view	“Clear next step” copy — not word-count care	“AI explained for 10 minutes” proof
Marketing	“Short coaches that leave room to think”	“Unlimited thorough AI explanations”
Analytics	coach_tokens, se_before_wrap, wrap-expand rate	Explanation Score™ / tokens-as-quality

Competitive foil: ChatGPT = costless long dump. Khan video = long lecture outside attempt loop. Duo = short but often shallow XP copy. MindCraft = **length-contingent principle coach** under FEI.

XCII.9 Doctrine — SAFE-EXPLAIN (provisional)

1. **Short before long** — default principle / critical-feature length; expand only on evidence of need.
2. **Germane ≠ more text** — free WM for intrinsic joins; do not advertise “added germane load.”
3. **SE before wrap** — provided essay never replaces student why (SAFE-SE).
4. **Length fades with expertise** — couple to fade×hint state (SAFE-FADE / SAFE-HINT).
5. **Coherence > charm** — strip seductive and redundant coach material.
6. **No token North Star** — ban Explanation Score™ / tokens-read / “thoroughness” thumbs as learning KPIs.
7. **Assistance dilemma honesty** — withhold when construction is the learning act.
8. **Proof** — solo_transfer_pass / conceptual items / lower wrap-expand binge — not essay dwell.

Confidence: High that modern CLT treats germane as processing of intrinsic structure, not additive word count (Sweller et al., 2019; Sweller, 2010). High that instructional explanations’ average gain is small and often loses to SE prompts (Wittwer & Renkl, 2010). High that withholding stepwise explanations can improve conceptual outcomes (Richey & Nokes-Malach, 2013). High that coherence/redundancy and expertise reversal caution against always-longer guidance (Mayer line; Kalyuga et al., 2003). Medium on exact token caps / sentence budgets — run EXPLAIN-*. High that LLM defaults drift verbose without hard budgets.

XCII.10 Experiments

ID	Question	Design	Primary
EXPLAIN-1	Principle-short coach vs multi-paragraph monologue (same facts)	A/B within concept	solo_transfer_pass; conceptual item; wrap dwell
EXPLAIN-2	SE-before-wrap vs wrap-first essay	A/B	SE completeness; transfer; binge-expand
EXPLAIN-3	Length coupled to fade stage vs fixed long wrap	A/B	Solo transfer; expertise×length interaction

EXPLAIN-4	Soft-wrong: short mark vs long re-teach	A/B	Retry_120s; near transfer; anxiety self-report (secondary)
EXPLAIN-5	Parent CBC: “clear short coach” vs “thorough AI essays” vs “we never explain”	CBC	WTP; trust (SAFE-WTP)
EXPLAIN-QUAL	10 Maya: when does an explanation feel clarifying vs drowning?	Qual	Length phenomenology codebook

Falsifier: Fixed long monologue wins delayed solo transfer *and* 26w identity equally → still ban as *brand* vs ChatGPT; allow expandable “Show more” off default.

Falsifier: Ultra-short coach harms true novices on high-interactivity first exposure → keep E0 slightly richer marks; do not open with silence theater.

Falsifier: Parents refuse short-coach CBC entirely → redesign copy (“clear next step”) before surrendering to essay brand.

Pre-register: EXPLAIN- before any “unlimited thorough AI explanations / Explanation Score™ / longest tutor wins” campaign (SAFE-LABMETA). Do not duplicate HINT hard-peek arms — EXPLAIN owns provided-text length*; HINT owns soft→hard specificity.

XCII.11 So what for MindCraft commercially

- **Copy:** “Short coaches that leave room to think.” Never “the most thorough AI tutor.”
- **Product:** Default principle grain; SE-before-wrap; length tied to fade×hint; hard token/sentence budgets on coach prompts.
- **Positioning:** Against ChatGPT essay dumps and lecture-as-care; for *assistance-dilemma-honest* brevity that transfers.
- **Metric:** coach_tokens, se_before_wrap, wrap-expand rate — demote tokens / thoroughness NPS / Explanation Score™.
- **Kill list:** Longer≡better; monologue≡SE; token North Star; seductive coach fluff; fixed long wrap across expertise.
- **Growth:** Parent decks sell clarity and independence, not word count; tutors QA pour length (SAFE-TALK).
- **Vision:** Teach Maya to need *fewer* words from the coach — and to supply her own.

References (verified)

- Chandler, P., & Sweller, J. (1991). Cognitive load theory and the format of instruction. *Cognition and Instruction*, 8(4), 293–332. (Redundancy / split-attention lineage.)

- Kalyuga, S. (2007). Expertise reversal effect and its implications for learner-tailored instruction. *Educational Psychology Review*, 19, 509–539.
 - Kalyuga, S., Ayres, P., Chandler, P., & Sweller, J. (2003). The expertise reversal effect. *Educational Psychologist*, 38(1), 23–31.
 - Koedinger, K. R., & Aleven, V. (2007). Exploring the assistance dilemma in experiments with Cognitive Tutors. *Educational Psychology Review*, 19, 239–264.
 - Moreno, R., & Mayer, R. E. (2000). A coherence effect in multimedia learning: The case for minimizing irrelevant sounds in the design of multimedia instructional messages. *Journal of Educational Psychology*, 92(1), 117–125.
 - Richey, J. E., & Nokes-Malach, T. J. (2013). How much is too much? Learning and motivation effects of adding instructional explanations to worked examples. *Learning and Instruction*, 25, 22–34. <https://doi.org/10.1016/j.learninstruc.2012.11.006>
 - Sweller, J. (2010). Element interactivity and intrinsic, extraneous, and germane cognitive load. *Educational Psychology Review*, 22(2), 123–138. <https://doi.org/10.1007/s10648-010-9128-5>
 - Sweller, J., van Merriënboer, J. J. G., & Paas, F. (2019). Cognitive architecture and instructional design: 20 years later. *Educational Psychology Review*, 31, 261–292. <https://doi.org/10.1007/s10648-019-09465-5>
 - Wittwer, J., & Renkl, A. (2008). Why instructional explanations often do not work: A framework for understanding the effectiveness of instructional explanations. *Educational Psychologist*, 43(1), 49–64.
 - Wittwer, J., & Renkl, A. (2010). How effective are instructional explanations in example-based learning? A meta-analytic review. *Educational Psychology Review*, 22, 393–409. <https://doi.org/10.1007/s10648-010-9136-5>
-

Part XCIII — Retrieval Failure Modes in Practice UX

Chapter status: Living evidence + practice-session brief — Researcher tick 2026-08-06

Primary question: When a learner fails to retrieve mid-practice, which failure mode is occurring — tip-of-tongue / partial access, productive struggle, freeze/blanking, or avoidance — and what product response preserves construction without cosplaying “struggle theater” or dumping answers?

Owners: Product (Practice / Solver timers & hints) · Engine (spacing / soft-wrong) · HITL tutors · Brand · Red Team

Commercial job: Ship a **SAFE-RETRIEVE** doctrine: name the failure mode before treating it; wait + partial cue for TOT-class stalls; contingent soft scaffold for blank/freeze; unsuccessful attempt + feedback can potentiate learning; prove with retry/transfer — never Struggle Score™, blank-time NS, or “always stew / never stuck.”

Builds on: Parts XXIX (spacing/retrieval), XLI (desirable difficulties × anxiety), XLIX (productive misconceptions), LXXXVIII–XCII (fade / help / instrument / hint / explain). Sibling queue: id 94 productive failure sequencing. Product seams: dwell timers, tip-cue UI, freeze detection, soft-wrong path, hint contingency.

XCIII.1 Why this chapter exists

Practice UX collapses many stalls into one button: “Need a hint?” Parents hear “struggle” and either panic or romanticize it. Competitors either (a) intervene instantly (ChatGPT dump), (b) gamify stuckness (streaks, pep), or (c) preach productive struggle without instruments.

MindCraft’s FEI loop needs a finer grain: **retrieval failure is not one event**. A tip-of-tongue state that motivates continued search is not the same as working-memory blanking under anxiety, which is not the same as Kapur-style generation before instruction, which is not the same as help-avoidance.

FOUNDER BELIEF under audit: Distinguishing failure modes lets MindCraft keep desirable difficulty without flooding Maya after a freeze — and without selling “we never interrupt struggle” as brand.

Claims we refuse as doctrine:

1. All struggle $\equiv$ productive struggle / desirable difficulty.
2. Blank-time, dwell-seconds, or Struggle Score™ as learning North Stars.
3. Always-wait / never-intervene (silence theater $\equiv$ pedagogy).
4. Instant hard peek / full solution on first stall (executive dump).
5. Unsuccessful retrieval $\equiv$ wasted / harmful (so skip attempt and show answer).
6. “Never stuck” marketing that legitimizes answer delivery.
7. Tip-of-tongue pep copy without a retrieval path (affirmation without cue).
8. Streaks/XP for enduring blank screens.

XCIII.2 Constructs (keep modes distinct)

Construct	Research meaning	MindCraft analogue	Failure mode
Retrieval attempt	Effort to access stored knowledge	First try; soft-wrong retry	Skip attempt → restudy tip
TOT (tip-of-the-tongue)	Feeling that an unretrieved item is known / recoverable; metacognitive state dissociable from access	Partial recall; “I almost have it”; letter/structure cue hunger	Treat as total ignorance
Omission / blanking	No candidate produced; freeze	Empty response; rapid panic peek	Force endless wait
Commission error	Wrong candidate produced	Soft-wrong path	Shame / instant dump
Productive struggle	Effort to make sense of math that is not immediately apparent (Hiebert & Grouws)	Sense-making dwell with progress markers	Romanticize despair
Productive failure (PF)	Generation/exploration before consolidation (Kapur) — design, not any fail	Optional struggle-first sequences (id 94)	Label every miss “PF”
Anxiety–WM tax	Anxiety consumes WM needed for math (Ashcraft line)	Freeze under timed/exam chrome	“Try harder” as fix

SAFE-RETRIEVE	Mode-contingent response to retrieval failure	This chapter	One-size struggle UX
---------------	---	--------------	----------------------

Operational definition (HYPOTHESIS): A stall is *SAFE-RETRIEVE handled* when the product (a) attempts mode classification (TOT/partial vs blank/freeze vs productive progress vs avoidance), (b) defaults to wait + micro-cue for TOT-class, (c) offers contingent soft instrumental scaffold for blank/freeze after a bounded wait (SAFE-HINT), (d) treats commission errors via soft-wrong + SE (not dump), (e) never scores blank-seconds as virtue, and (f) co-primaries prove with `retry_120s / solo_transfer_pass` — not Struggle Score™.

XCIII.3 TOT is metacognitive — not “they don’t know it”

FACT: Brown & McNeill (1966) established laboratory elicitation of tip-of-the-tongue states as distinct from simple forgetting — partial phonological/orthographic information often available when the full target is not.

FACT (review): Schwartz & Metcalfe (2011, *Memory & Cognition*, 39, 737–749, doi:10.3758/s13421-010-0066-8) synthesize direct-access and heuristic–metacognitive accounts: TOTs accompany some failed or slow retrievals; the *feeling* is informed by cues (partial info, related info, cue familiarity) and can change retrieval *behavior* (persist in search). TOT experience and retrieval success are dissociable.

Applied (HYPOTHESIS): In Practice, a student who can name “quadratic formula starts with negative b...” or sketch a triangle but cannot finish is often in a **TOT-class stall**, not a cold start. Correct product move: **bounded wait + partial cue** (first letter of step label, FormatId reminder, “name the bridge”) — not full worked example, not “you don’t know this — here’s the lecture.”

Kill: Instant dump on first hesitation; “blank = never learned.”

Survive: Treat recoverable inaccessibility as a search problem; cue, don’t replace.

XCIII.4 Unsuccessful retrieval can potentiate later learning

FACT: Kornell, Hays & Bjork (2009, *Journal of Experimental Psychology: Learning, Memory, and Cognition*, 35(4), 989–998) — unsuccessful retrieval attempts followed by study of the correct answer enhanced subsequent learning relative to study-only controls across multiple experiments (failed tests were not mere harm).

FACT: Richland, Kornell & Kao (2009, *Journal of Experimental Psychology: Applied*, 15(3), 243–257) — pretesting (often unsuccessful) on prose material improved later retention of those facts vs non-pretested content (pretesting / failed-test potentiation).

FACT (synthesis): Vaughn, Hausman & Kornell / Kornell & Vaughn line (e.g. Kornell & Vaughn, 2016, *Psychology of Learning and Motivation*, Vol. 65) — retrieval practice benefits appear even when the attempt fails, provided feedback/study follows; a two-stage view (attempt → process answer) undercuts “only successful retrieval counts.”

Applied (FOUNDER BELIEF → testable): MindCraft should **require an attempt** before heavy help (already SAFE-HELP / SAFE-HINT). Marketing must not imply “wrong first tries damage the brain.” Parent copy: “Trying before the coach — even when stuck — sets up the learning that follows.”

Contradicting pressure (FACT/HYPOTHESIS): Anxiety–WM literature (Ashcraft & Kirk, 2001; Ashcraft, 2002) — under high anxiety, forced struggle without support can collapse performance. Mode matters: potentiation assumes a *search* is happening, not a panic freeze. Hence SAFE-RETRIEVE’s freeze branch (soft scaffold after bound), not endless blank.

Kill: Skip-attempt restudy as default; “errors always scar memory” fear marketing.

Survive: Attempt → feedback; shorten wait when freeze signals dominate.

XCIII.5 Productive struggle ≠ every stuck screen

FACT: Hiebert & Grouws (2007) define productive struggle as intellectual effort to make sense of mathematics that is not immediately apparent — sense-making, not despair.

FACT: Warshauer (2015, *Journal of Mathematics Teacher Education*, 18, 375–400, doi:10.1007/s10857-014-9286-3) — 186 classroom struggle episodes; taxonomy of struggle types (get started, carry out process, explain, misconception/error) and teacher responses (telling → affordance). Productivity judged by maintaining cognitive demand, addressing the struggle, and building on student thinking — not by duration alone.

FACT (PF design): Kapur (2014, *Cognitive Science*, 38, 1008–1022) — productive *failure* is sequenced generation then consolidation, not any unsuccessful miss. Full PF sequencing is queue id 94; this chapter bans mislabeling every miss as “productive failure.”

Applied (HYPOTHESIS): Detect **progress markers** during dwell (partial work, SE draft, diagram marks). Progress + TOT → wait/cue. No progress + peek urgency / affect → soft scaffold. Freeze → earlier help (optional SAFE-EXPOSE), not Struggle Score™.

Kill: “We maximize productive struggle” measured as blank minutes.

Survive: Sense-making effort with demand maintained; telling as last grain (Warshauer continuum ≈ SAFE-HINT soft→hard).

XCIII.6 Anxiety blanking — different medicine than TOT

FACT: Ashcraft (2002, *Current Directions in Psychological Science*, 11(5), 181–185) — math anxiety as tension/fear that interferes with performance; associated educational and cognitive consequences.

FACT: Ashcraft & Kirk (2001, *Journal of Experimental Psychology: General*, 130(2), 224–237) — math anxiety relates to reduced working-memory capacity available for math; dual-task-like tax.

Bridge to Part XLI: Desirable difficulties that raise challenge without WM support can become *undesirable* under anxiety. Timed chrome, shame clocks, and public blank timers are anti-patterns (also SAFE-EXAM / SAFE-PDASH).

Applied (HYPOTHESIS):

Signal (product proxies)	Likely mode	Default response
--------------------------	-------------	------------------

Partial work / self-report “almost” / slow edits	TOT / partial	Wait band + micro-cue; no full dump
Empty + rapid rage-click hint / affect spike	Freeze / blank	Soft instrumental after short bound; lower timer salience
Wrong worked path	Commission	Soft-wrong + SE before wrap (SAFE-EXPLAIN)
Idle + no interaction + skip	Avoidance	Help-invite (SAFE-HELP); do not score as grit
Generation phase before teach (designed)	PF (id 94)	Separate experiment arm — not default Practice miss

Kill: One dwell timer policy for all stalls; shame countdown.

Survive: Mode-contingent clocks and scaffolds.

XCIII.7 Product surface — SAFE-RETRIEVE claim contract

Surface	Required behavior	Banned substitute
Practice first stall	Classify; TOT → cue; freeze → soft soon	Instant hard peek
Soft-wrong	Attempt counted as retrieval work	Treat miss as moral failure
Hint ladder	Soft before hard; SE cost (SAFE-HINT)	Free bottom-out menu
Timers	Bound waits; no public blank shame	Struggle Score™ / blank-seconds XP
Coach copy	“Keep searching” only with a cue path	Empty pep / “embrace struggle” poster
Exam / prove rail	KR timing per SAFE-FBTIME / SAFE-EXAM	Flood after first blank
Tutor HITL	Prompt; name mode in brief	“Let them suffer” cosplay
Parent view	“Attempt before coach” honesty	“Never stuck” or “max struggle minutes”
Marketing	Recoverable difficulty + mode-smart help	Struggle theater / answer-dump hero
Analytics	stall_mode, cue_shown, attempt_before_help, retry_120s	Struggle Score™ / dwell-as-virtue

Competitive foil: ChatGPT = zero attempt dump. Duo = streak through shallow items. “Grit” apps = blank-time theater. MindCraft = **mode-contingent retrieval support** under FEI.

XCIII.8 Doctrine — SAFE-RETRIEVE (provisional)

1. **Name the mode** — TOT/partial ≠ blank/freeze ≠ PF design ≠ avoidance.
2. **Attempt before heavy help** — unsuccessful retrieval + feedback can potentiate (Kornell/Bjork line).
3. **TOT** → **wait + micro-cue** — preserve search; do not replace with lecture.
4. **Freeze** → **bounded soft scaffold** — anxiety–WM honesty; no endless stew.
5. **Progress marks productive struggle** — not clock alone (Warshauer / Hiebert).
6. **No Struggle Score™** — ban blank-seconds / dwell-XP / “maximize struggle” NS.
7. **No never-stuck dump brand** — ban executive rescue as identity.
8. **Proof** — retry_120s, near/solo transfer, lower freeze-binge peeks — not pep NPS.

Confidence: High that TOT is a studied metacognitive state distinct from mere absence of knowledge (Brown & McNeill; Schwartz & Metcalfe, 2011). High that failed retrieval + feedback can enhance later learning (Kornell, Hays & Bjork, 2009; Richland et al., 2009). High that productive struggle is sense-making effort, not duration (Hiebert & Grouws; Warshauer, 2015). High that anxiety taxes WM (Ashcraft line). Medium on automatic stall-mode classifiers in-product — run RETRIEVE-* + qual. High that one-size “always wait” or “always hint” will fail FEI.

XCIII.9 Experiments

ID	Question	Design	Primary
RETRIEVE-1	TOT micro-cue vs instant soft hint vs wait-only on partial-work stalls	A/B/C	retry_120s; solve without hard peek; near transfer
RETRIEVE-2	Freeze: short-bound soft scaffold vs long forced wait vs instant hard peek	A/B/C	Solo transfer; anxiety self-report; peek binge
RETRIEVE-3	Require attempt before any help vs optional skip-to-coach	A/B	Delayed retention; attempt_before_help rate
RETRIEVE-4	Progress-aware dwell (partial marks extend wait) vs fixed dwell timer	A/B	Transfer; freeze false-positives
RETRIEVE-5	Parent CBC: “attempt then coach” vs “never stuck” vs “max struggle time”	CBC	WTP; trust (SAFE-WTP)

RETRIEVE-QUAL	10 Maya: phenomenology of almost-have-it vs blank panic vs fruitful stuck	Qual	Mode codebook
---------------	--	------	---------------

Falsifier: Instant hard peek wins delayed solo transfer *and* 26w identity equally on freeze *and* TOT stalls → still ban as brand vs ChatGPT; allow emergency reveal off default.

Falsifier: Micro-cues never beat wait-only for TOT-class → keep wait bands; demote cue chrome.

Falsifier: Parents reject attempt-first CBC entirely → rewrite (“short try, then a nudge”) before surrendering to never-stuck.

Pre-register: RETRIEVE- *before any* “maximize productive struggle minutes / Struggle Score™ / never interrupt / never stuck” campaign (SAFE-LABMETA). Do not steal Kapur PF sequencing arms — those belong to id 94; this family owns in-session stall response*.

XCIII.10 So what for MindCraft commercially

- **Copy:** “When you’re stuck, we read the stuck — almost-there gets a nudge; freeze gets a next step.” Never “unlimited struggle” or “never stuck.”
- **Product:** Stall-mode proxies; TOT cues; freeze-bound soft hints; attempt-before-help gate; no blank-time XP.
- **Positioning:** Against answer-dump tutors *and* grit-theater apps; for recoverable, mode-smart difficulty.
- **Metric:** stall_mode, cue_shown, attempt_before_help, retry_120s — demote Struggle Score™ / dwell-as-virtue.
- **Kill list:** All-struggle≡productive; blank-time NS; always-wait; instant dump; never-stuck hero; errors-always-scar fear.
- **Growth:** Parent decks sell attempt→coach honesty; tutors brief on mode naming (SAFE-HITL / SAFE-TALK).
- **Vision:** Maya learns to stay with *almost* — and to accept a soft scaffold when the mind blanks — without becoming either helpless or performatively stuck.

References (verified)

- Ashcraft, M. H. (2002). Math anxiety: Personal, educational, and cognitive consequences. *Current Directions in Psychological Science*, 11(5), 181–185. <https://doi.org/10.1111/1467-8721.00196>
- Ashcraft, M. H., & Kirk, E. P. (2001). The relationships among working memory, math anxiety, and performance. *Journal of Experimental Psychology: General*, 130(2), 224–237.
- Brown, R., & McNeill, D. (1966). The “tip of the tongue” phenomenon. *Journal of Verbal Learning and Verbal Behavior*, 5(4), 325–337.
- Hiebert, J., & Grouws, D. A. (2007). The effects of classroom mathematics teaching on students’ learning. In F. K. Lester Jr. (Ed.), *Second handbook of research on mathematics teaching and learning* (pp. 371–404). Information Age.
- Kapur, M. (2014). Productive failure in learning math. *Cognitive Science*, 38(5), 1008–1022. <https://doi.org/10.1111/cogs.12107>

- Kornell, N., Hays, M. J., & Bjork, R. A. (2009). Unsuccessful retrieval attempts enhance subsequent learning. *Journal of Experimental Psychology: Learning, Memory, and Cognition*, 35(4), 989–998.
 - Kornell, N., & Vaughn, K. E. (2016). How retrieval attempts affect learning: A review and synthesis. In B. H. Ross (Ed.), *Psychology of Learning and Motivation* (Vol. 65, pp. 183–215). Academic Press.
 - Richland, L. E., Kornell, N., & Kao, L. S. (2009). The pretesting effect: Do unsuccessful retrieval attempts enhance learning? *Journal of Experimental Psychology: Applied*, 15(3), 243–257.
 - Schwartz, B. L., & Metcalfe, J. (2011). Tip-of-the-tongue (TOT) states: Retrieval, behavior, and experience. *Memory & Cognition*, 39(5), 737–749. <https://doi.org/10.3758/s13421-010-0066-8>
 - Warshauer, H. K. (2015). Productive struggle in middle school mathematics classrooms. *Journal of Mathematics Teacher Education*, 18(4), 375–400. <https://doi.org/10.1007/s10857-014-9286-3>
-

Part XCIV — Productive Failure vs Guided Success Sequencing

Chapter status: Living evidence + sequencing brief — Researcher tick 2026-08-07

Primary question: When should MindCraft put **generation / struggle-first** before instruction (productive failure / PS-I), and when should it put **guided success / instruction-first** (I-PS, fade, worked examples) — without mislabeling every miss as “PF” or selling fail-first as brand?

Owners: Product (Learn / Practice sequence flags) · Engine (mission templates) · HITL tutors · Brand · Red Team

Commercial job: Ship a **SAFE-PF** doctrine: PF is a *designed two-phase sequence* (generate → consolidate on student RSMs), not any failed attempt; fidelity and priors gate when struggle-first beats teach-first; prove with conceptual + transfer — never PF Score™, fail-first-always, or discovery-without-consolidation.

Builds on: Parts XXVI (CLT / fading), XLI (DD × anxiety), XLIX (productive misconceptions), LXXXVIII–XCIII (fade / help / hint / explain / retrieve). Sibling: XCIII owns *in-session stall response*; this chapter owns *mission-level generation→instruction order*.

XCIV.1 Why this chapter exists

Brand temptation: “We let students fail productively.” Competitor temptation: “We never let them struggle — answers now.” Both are costumes.

Kapur’s productive failure (PF) is a **learning design**: students generate representations and solution methods (RSMs) on a novel complex problem *before* teacher-led consolidation that compares student RSMs to canonical ones. That is not (a) blank-screen theater, (b) ChatGPT dump after three seconds, or (c) relabeling every Practice miss as science.

MindCraft already has guided rails (SAFE-FADE, SAFE-HINT, SAFE-EXPLAIN). The open commercial question is **when Learn-next / new-concept missions should invert the order** — generate first — and when that would waste WM, spike anxiety, or become discovery cosplay.

FOUNDER BELIEF under audit: A high-fidelity PF arm on *selected* novel concepts will beat tell-and-practice on conceptual understanding and transfer without becoming the default for every Practice item or exam drill.

Claims we refuse as doctrine:

1. Every unsuccessful attempt $\equiv$ productive failure.
2. Fail-first / discovery-always as brand or North Star.
3. Generation without Phase-2 consolidation that builds on student work.
4. PF Score™ / generation-minutes / “maximize failure” KPIs.
5. Unguided discovery as identity (Kirschner/Sweller/Clark critique ignored).
6. Struggle-first for high-anxiety freeze states without SAFE-RETRIEVE exit.
7. PF as guaranteed ACT points or “2-sigma from failing first.”
8. Relabeling soft-wrong Practice as Kapur PF.

XCIV.2 Constructs (keep sequences distinct)

Construct	Research meaning	MindCraft analogue	Not the same as
I-PS / DI	Instruction (or strong guidance) then problem solving	Coach/fade $\rightarrow$ practice; worked-example first	“Never let them try”
PS-I	Problem solving then instruction (umbrella)	Generate $\rightarrow$ teach	Any miss
Productive failure (PF)	High-fidelity PS-I: generation+exploration then consolidation assembling student RSMs $\rightarrow$ canonical	Opt-in Learn mission: invent methods $\rightarrow$ Map/coach C&C	Blank stall (XCIII)
Invention / PFL	Invent methods to prepare for future learning from resources (Schwartz & Martin)	Invent-a-measure / invent-a-rule cards before lecture	Unscaffolded forever
Guided success	Strong guidance so first successes are correct under support	SAFE-FADE completion $\rightarrow$ solo	Mastery fireworks
SAFE-PF	When to run PF vs guided success; fidelity gates	This chapter	Struggle Score™

Operational definition (HYPOTHESIS): A mission is *SAFE-PF compliant* when (a) the target concept is novel enough that students lack the canonical method, (b) Phase 1 requires multiple RSM attempts with prior-knowledge hooks (not empty invent-from-nothing), (c) Phase 2 **must** compare/contrast student-generated work with canonical forms (tutor or coach C&C — not a generic lecture that ignores what they produced), (d) anxiety/freeze exits to SAFE-RETRIEVE soft scaffold rather than endless generation, (e) outcomes judged on conceptual + transfer items (and later solo transfer), not on Phase-1 correctness or generation time, and (f) default Practice on already-taught skills stays guided/fade — not fail-first.

XCIV.3 PF is a design — not romanticized failure

FACT: Kapur & Bielaczyc (2012, *Journal of the Learning Sciences*, 21(1), 45–83, doi:10.1080/10508406.2011.591717) — design principles for PF across Singapore secondary math classrooms: collaborative complex problem solving *without* instructional scaffolds until teacher-led consolidation; PF students often “failed” to produce canonical solutions in Phase 1 yet outperformed direct-instruction peers on posttests, including representational flexibility.

FACT: Kapur (2014, *Cognitive Science*, 38(5), 1008–1022, doi:10.1111/cogs.12107) — controlled contrast of PF (problem solving then instruction) vs DI (instruction then problem solving) for learning standard deviation; PF outperformed DI; number of student-generated solutions predicted learning. PF is explicitly defined as reversing DI’s teach-then-practice order while still providing instruction afterward (not pure discovery).

FACT (mechanisms summary): Kapur’s program states four interdependent mechanisms — activation and differentiation of prior knowledge, attention to critical features, explanation/elaboration, and organization/assembly into the target concept — requiring both generation *and* consolidation that uses student RSMs (see Kapur design summaries; Kapur & Bielaczyc, 2012).

Kill: “Fail first” marketing with no consolidation that cites student work.

Survive: Two-phase design; generation diversity as process signal, not as KPI vanity.

XCIV.4 Meta-analytic grain — when PS-I wins, and when it does not

FACT: Sinha & Kapur (2021, *Review of Educational Research*, 91(5), 761–798, doi:10.3102/00346543211019105) — meta-analysis of 53 studies / 166 comparisons of PS-I vs I-PS: overall moderate advantage for PS-I (Hedge’s $g \approx 0.36$, 95% CI ~ 0.20 – 0.51). Effects stronger when implementations match PF fidelity principles (g in roughly 0.37–0.58 range in high-fidelity subgroups). Contrasting trends: younger learners (roughly grades 2–5) and domain-general skill targets more often favored I-PS. Publication-bias-adjusted estimates in the paper are larger still — treat magnitude as uncertain, direction as the commercial takeaway.

FACT: Loibl, Roll & Rummel (2017, *Educational Psychology Review*, 29(4), 693–715) — theory of *when and how* problem solving followed by instruction supports learning: benefits hinge on preparation (activation, awareness of knowledge gaps) and on instruction that makes critical features discernable — often by building on learner solutions / contrasting cases.

FACT: Loibl & Rummel (2014, *Learning and Instruction*, 34, 74–85) — “knowing what you don’t know” can make failure productive: gap awareness during/after generation is part of the mechanism, not mere time-on-wrong-answers.

Applied (HYPOTHESIS): MindCraft should **not** ship universal fail-first. Prefer PF arms when: (1) concept is new, (2) student has enough prior resources to generate *suboptimal* RSMs (SAFE-COLD / Map priors), (3) consolidation UI/tutor can show *their* attempts vs canonical, (4) age/domain fit secondary math-style conceptual targets — not every elementary procedure drill, (5) FEI co-primaries include conceptual/transfer, not Phase-1 accuracy.

Kill: One sequence for all missions; “science says always invent first.”

Survive: Fidelity-gated PS-I; I-PS default when priors thin, anxiety high, or skill is procedural fluency under exam chrome (SAFE-EXAM).

XCIV.5 Invention prepares for future learning — assessments that include learning

FACT: Schwartz & Martin (2004, *Cognition and Instruction*, 22(2), 129–184, doi:10.1207/s1532690xci2202_1) — ninth-grade statistics invention activities looked inefficient on standard assessments (students often failed to invent canonical methods) but, when coupled with subsequent learning resources, yielded stronger preparation-for-future-learning outcomes than tell-and-practice — especially on transfer tests that included a learning resource.

Applied (FOUNDER BELIEF → testable): Parent and tutor copy must allow **ugly early work**. Measuring Learn missions only by Phase-1 correctness will kill PF before Phase 2. Product analytics: separate pf_phase = generate | consolidate; never score generation accuracy as mastery fireworks (SAFE-COLD / SAFE-INSTRUMENT).

Kill: Early-wrong shame on invent missions; gradebook cosplay of generation.

Survive: “Invent then compare” honesty; MoC after consolidation + transfer.

XCIV.6 The guided-instruction critique is a boundary, not a ban on PF

FACT: Kirschner, Sweller & Clark (2006, *Educational Psychologist*, 41(2), 75–86) — argue minimally guided instruction is generally less effective/efficient than strong guidance for novices, citing cognitive architecture and WM limits; guidance advantage recedes as prior knowledge rises. This is the load-bearing critique against pure discovery.

How PF answers without hand-waving (HYPOTHESIS + Kapur framing): Classic PF is *not* “minimal guidance forever.” It is delayed strong guidance: Phase 1 generation uses prior knowledge on complex-but-accessible problems; Phase 2 is teacher-led assembly. Pure discovery without consolidation fails the Kirschner test *and* fails Kapur’s own design. MindCraft must not ship Phase-1-only “explore the world” lore walls (also SAFE-STORYLOAD).

Bridge to XLI / XCIII: Under high math anxiety, forced generation without an exit can tax WM (Ashcraft line) and become freeze — use SAFE-RETRIEVE soft bound, then resume consolidation; do not romanticize panic as PF fidelity.

Kill: Discovery-without-teach brand; “minimal guidance = identity.”

Survive: Hybrid PS-I with mandatory consolidation; guided success when priors or affect demand it.

XCIV.7 Product surface — SAFE-PF claim contract

Surface	Required behavior	Banned substitute
Learn-next (novel concept)	Optional PF mission flag: generate RSMs → C&C consolidate	Always invent-first for every topic

Practice (taught skill)	Guided/fade/retrieve rails (SAFE-FADE / HINT / RETRIEVE)	Relabel misses as PF
Coach / Solver	Phase 2 builds on <i>student</i> attempts (show their RSMs)	Generic lecture ignoring generation
Tutor HITL	Brief: elicit 2+ methods, then compare; no dump in Phase 1	“Let them suffer” or explain-first always
Map / ontology	Target novel join; priors must support generation	Cold invent with zero hooks
Exam rail	Prefer guided success / retrieval under time	Fail-first timed flood
Parent view	“Ugly invent → clear compare” timeline	Generation % / PF Score™
Marketing	Designed struggle-then-teach on new ideas	“We make you fail” / guaranteed points
Analytics	pf_phase, RSM count, consolidate_on_student_work, transfer	Generation-minutes NS

Competitive foil: ChatGPT = instruction/dump without generation. Pure discovery apps = generation without assembly. Duo = shallow success loops. MindCraft = **sequence honesty** — invent when priors allow; guide when they don’t; always consolidate before claiming learning.

XCIV.8 Doctrine — SAFE-PF (provisional)

1. **PF ≠ every miss** — reserve the label for designed generation→consolidation sequences (Kapur).
2. **Two phases or it is not PF** — Phase 2 must build on student RSMs (C&C), not ignore them.
3. **Fidelity over slogans** — generation diversity + gap awareness + feature discernment (Sinha & Kapur; Loibl et al.).
4. **Gate the arm** — novel concept + adequate priors + secondary-math-style conceptual targets; default I-PS/guided when thin priors, young procedural drills, or exam chrome.
5. **Anxiety exit** — freeze → SAFE-RETRIEVE soft scaffold; do not endless-stew.
6. **No PF Score™** — ban generation-minutes / fail-count / “maximize productive failure” NS.
7. **Proof** — conceptual understanding + near/far transfer + later solo_transfer_pass — not Phase-1 correctness.
8. **Positioning** — against answer-dump *and* discovery-cosplay; for inspectable sequence choice under FEI.

Confidence: High that Kapur PF is a documented two-phase design with classroom and lab evidence (Kapur & Bielaczyc, 2012; Kapur, 2014). High that meta-analytic PS-I advantage exists with fidelity and population moderators (Sinha & Kapur, 2021). High that invention can prepare future learning despite ugly early assessments (Schwartz & Martin, 2004). High that pure minimal guidance is a bad brand for novices (Kirschner et al., 2006). Medium on MindCraft’s ability to automate high-fidelity C&C in AI coach without tutor HITL — run PF-* with human-led arms first. Medium on anxiety × PF interaction in-product — co-design with SAFE-RETRIEVE / SAFE-EXPOSE.

XCIV.9 Experiments

ID	Question	Design	Primary
PF-1	High-fidelity PF (generate→C&C on student RSMs) vs DI (teach→practice) on one novel ACT-adjacent concept	A/B	Conceptual post; near transfer; delayed retention
PF-2	Consolidation that cites student RSMs vs generic lecture after same generation	A/B	Conceptual gain; consolidate_on_student_work
PF-3	PF vs guided-success when Map priors low (thin prior knowledge)	A/B	Transfer; freeze/peek binge; dropout
PF-4	PF with SAFE-RETRIEVE freeze exit vs PF with forced long generation	A/B	Solo transfer; anxiety self-report
PF-5	Parent CBC: “invent then compare” vs “never fail” vs “fail-first grit”	CBC	WTP; trust (SAFE-WTP)
PF-QUAL	10 Maya + 5 tutors: phenomenology of invent missions vs explain-first	Qual	Sequence codebook; shame markers

Falsifier: DI beats high-fidelity PF on conceptual *and* transfer for our secondary cohort on pre-registered concepts → keep PF as lab-only; do not market struggle-first.

Falsifier: Generic post-generation lecture equals student-RSM C&C → demote fidelity chrome; still ban Phase-1-only.

Falsifier: Parents reject invent-then-compare CBC → rewrite (“short try at inventing, then we show how experts organize it”) before surrendering to never-fail brand.

Pre-register: PF- *before any “fail productively / discovery identity / PF Score™ / always invent first” campaign (SAFE-LABMETA).* Do not steal RETRIEVE- stall arms — those own mid-item stuckness, not mission order.

XCIV.10 So what for MindCraft commercially

- **Copy:** “On new ideas, try inventing a method — then we compare yours to the standard one.” Never “we make you fail” or “never struggle.”

- **Product:** mission_sequence = pf | guided; Phase flags; student-RSM capture for C&C; freeze exit into SAFE-RETRIEVE.
 - **Positioning:** Against ChatGPT dumps *and* unguided discovery cosplay; for sequenced construction with inspectable consolidation.
 - **Metric:** conceptual/transfer lifts under PF-1/2; rsm_count; consolidate_on_student_work — demote generation-minutes / PF Score™.
 - **Kill list:** Every-miss=PF; fail-first-always; generation-without-consolidation; PF Score™; discovery-as-identity; anxiety-forced stew; ACT-points-from-PF ads.
 - **Growth:** Tutor playbooks for elicit→compare (feeds id 95 cognitive apprenticeship); parent decks sell invent-then-compare honesty.
 - **Vision:** Maya learns that early wrong methods can be *raw material for understanding* — not proof she is not a math person — when the product finishes the job with consolidation.
-

References (verified)

- Kapur, M. (2014). Productive failure in learning math. *Cognitive Science*, 38(5), 1008–1022. <https://doi.org/10.1111/cogs.12107>
 - Kapur, M., & Bielaczyc, K. (2012). Designing for productive failure. *Journal of the Learning Sciences*, 21(1), 45–83. <https://doi.org/10.1080/10508406.2011.591717>
 - Kirschner, P. A., Sweller, J., & Clark, R. E. (2006). Why minimal guidance during instruction does not work: An analysis of the failure of constructivist, discovery, problem-based, experiential, and inquiry-based teaching. *Educational Psychologist*, 41(2), 75–86. https://doi.org/10.1207/s15326985ep4102_1
 - Loibl, K., & Rummel, N. (2014). Knowing what you don’t know makes failure productive. *Learning and Instruction*, 34, 74–85.
 - Loibl, K., Roll, I., & Rummel, N. (2017). Towards a theory of when and how problem solving followed by instruction supports learning. *Educational Psychology Review*, 29(4), 693–715.
 - Schwartz, D. L., & Martin, T. (2004). Inventing to prepare for future learning: The hidden efficiency of encouraging original student production in statistics instruction. *Cognition and Instruction*, 22(2), 129–184. https://doi.org/10.1207/s1532690xc2202_1
 - Sinha, T., & Kapur, M. (2021). When problem solving followed by instruction works: Evidence for productive failure. *Review of Educational Research*, 91(5), 761–798. <https://doi.org/10.3102/00346543211019105>
-

Part XCV — Cognitive Apprenticeship in Tutor Playbooks

Chapter status: Living evidence + HITL ops brief — Researcher tick 2026-08-07

Primary question: How should MindCraft encode **cognitive apprenticeship** (modeling → coaching → scaffolding → fading + articulation/reflection/exploration) into tutor playbooks and coach UX — without guild cosplay, demo-only “teaching,” or never-fade permanent coaching?

Owners: HITL ops · Tutor training · Coach / Solver product · Brand · Red Team

Commercial job: Ship a **SAFE-APPRENTICE** doctrine: make expert *thinking* visible, then transfer responsibility on a fade schedule; prove with solo transfer and articulation quality — never Master Score™, modeling-as-lecture, or

“we’re an apprenticeship” costume ads.

Builds on: Parts XXVI (CLT / tutoring calibration), LI (deliberate practice spine), LII (CoP careful transfer), LIV (AAR), LXVIII (HITL), LXXI (tutor grain), LXXXIV (talk-ratio), LXXXVIII–XCI (fade / help / hint), XCII–XCIV (explain / retrieve / PF). Sibling: SAFE-FADE owns *Solver example stage*; this chapter owns the **full method stack + sociology of the session** for human tutors and the AI coach that must not impersonate a guild.

XCV.1 Why this chapter exists

Brand temptation: “Learn like an apprentice under a master.” Competitor temptation: “Watch the AI solve it, then mimic.” Both fail the Collins test.

Cognitive apprenticeship was proposed because school often *hides* thinking: students see answers, not the control decisions, heuristics, and belief checks experts use. Collins, Brown, and Newman’s fix is not medieval cosplay — it is a design for making cognition visible, supporting performance, then withdrawing support so the learner owns the craft.

MindCraft already fades worked examples (SAFE-FADE), prices peeks (SAFE-HINT), and measures construction (SAFE-TALK). The open commercial question is whether tutor **playbooks** and coach **phase labels** implement the *whole* method stack — or whether we slap “apprenticeship” on explain-first sessions and call it science.

FOUNDER BELIEF under audit: Sessions briefed as model→coach→fade with forced student articulation will beat explain-pour + “got it?” on solo_transfer_pass and transfer without raising talk-% vanity.

Claims we refuse as doctrine:

1. Guild / master / “journeyman math” costume as pedagogy proof.
 2. Modeling ≡ lecture or AI monologue without student attempt.
 3. Scaffolding without fading (permanent coach dependency).
 4. Talk-ratio alone ≡ cognitive apprenticeship (SAFE-TALK already killed airtime dogma).
 5. Apprenticeship Score™ / “hours under master” North Stars.
 6. Discord / community ≡ apprenticeship (SAFE-CoP boundary).
 7. Reciprocal-teaching costume without turn-taking dialogue fidelity.
 8. “Cognitive apprenticeship” marketing without inspectable method phases in the session log.
-

XCV.2 Constructs (keep methods distinct)

Construct	Research meaning	MindCraft analogue	Not the same as
Modeling	Expert externalizes thinking while performing	Think-aloud on one item; coach “watch my control moves” card	Dumping a full solution without process talk
Coaching	Observe learner; hints, feedback, reminders, new tasks in situ	Tutor prompts mid-attempt; soft→hard peeks	Pre-scripted lecture

Scaffolding	Temporary support controlling hard elements so learner can succeed on rest	Completion problems; Map-briefed hints; Wood functions	Permanent answer key
Fading	Gradual removal of support toward independence	SAFE-FADE E0→E3; less tutor airtime	Abrupt “you’re on your own” shame
Articulation	Learner states knowledge/reasoning	SE prompts; postmortem why (SAFE-ANNOT cousin)	“Got it?” yes/no
Reflection	Compare own process to expert/peer	AAR light; film-clip of <i>their</i> attempt	Generic tip flood
Exploration	Learner sets goals / invents problems	Challenge accept; optional PF generate phase	Unguided forever
SAFE-APPRENTICE	When/how to run the stack in HITL + coach	This chapter	Guild branding

Operational definition (HYPOTHESIS): A session is *SAFE-APPRENTICE* compliant when (a) at least one short modeling segment makes *control/heuristic* moves audible (not only algebra steps), (b) the majority of problem time is student performance under coaching, (c) scaffolds are logged and *fade within the session or across the mission*, (d) the student articulates a principle or misconception before bottom-out reveal, (e) reflection compares *their* attempt to a canonical path (not a tip dump), (f) outcomes judged on solo transfer + articulation quality, not modeling minutes or master praise, and (g) AI coach phases use the same labels — monologue without attempt fails the gate.

XCV.3 The framework — methods that make thinking visible

FACT: Collins, Brown, & Newman (1989, in L. B. Resnick, Ed., *Knowing, learning, and instruction: Essays in honor of Robert Glaser*, pp. 453–494, Erlbaum) — propose *cognitive apprenticeship* to teach the crafts of reading, writing, and mathematics by adapting traditional apprenticeship (observation, coaching, practice / modeling, coaching, fading) so that otherwise invisible cognitive and metacognitive processes become observable. Core methods: modeling, coaching, scaffolding (with fading), articulation, reflection, exploration. Content types include domain knowledge, heuristic strategies, control strategies, and learning strategies. Sequencing principles include global before local skills, increasing complexity and diversity.

FACT: Collins, Brown, & Holum (1991, *American Educator*, Winter) — practitioner restatement: traditional apprenticeship’s four aspects are modeling, scaffolding, fading, and coaching; cognitive apprenticeship extends these to school subjects and adds articulation, reflection, and exploration so students generalize and take autonomy. Emphasizes that in schooling, thinking is often invisible to both student and teacher — the design problem is visibility + transfer of responsibility.

FACT: Brown, Collins, & Duguid (1989, *Educational Researcher*, 18(1), 32–42, doi:10.3102/0013189X018001032) — argue knowledge is situated in activity, context, and culture; propose cognitive apprenticeship as an alternative to didactic methods that treat concepts as context-free. Cite Schoenfeld’s math problem-solving teaching and other

classroom designs as exemplars of entering a *culture of practice* (tool use under authentic activity) rather than memorizing inert procedures.

Applied (HYPOTHESIS): MindCraft tutor briefs must name **phase** (model | coach | scaffold | fade | articulate | reflect | explore) in the session checklist — not “be Socratic” as a vibe. Product coach UI should show the same phase chip so AI and human share a fidelity language (SAFE-HITL / SAFE-TALK).

Kill: “Apprenticeship” as brand adjective with no phase log.

Survive: Inspectable method stack; thinking made audible; responsibility transferred.

XCV.4 Math exemplar — Schoenfeld’s problem-solving teaching

FACT: Schoenfeld (1985, *Mathematical Problem Solving*, Academic Press) — framework for expert/novice math problem solving across resources, heuristics, control, and belief systems; documents teaching methods that make control and heuristic selection explicit rather than treating “problem solving” as extra practice items.

FACT (as cited in the CA literature): Collins et al. (1989/1991) and Brown et al. (1989) treat Schoenfeld’s college problem-solving teaching as a cognitive-apprenticeship exemplar: modeling heuristic selection and control on live problems, coaching students as they attempt, scaffolding/fading support, and pushing reflection on process — aimed at entering mathematical *practice*, not only accumulating procedures.

Applied (FOUNDER BELIEF → testable): Tutor playbooks for ACT-adjacent sessions should script **one** short model of *control talk* (“I’m stuck — I’ll try a simpler case / check units / draw the join”) before coaching the student’s attempt. Modeling only algebra steps without control talk fails Schoenfeld’s point and becomes SAFE-EXPLAIN monologue risk.

Kill: Heuristic tip card flood without student control practice.

Survive: Audible control moves + student ownership of the next attempt.

XCV.5 Scaffolding is a function set — not a permanent crutch

FACT: Wood, Bruner, & Ross (1976, *Journal of Child Psychology and Psychiatry*, 17(2), 89–100, doi:10.1111/j.1469-7610.1976.tb00381.x) — characterize tutoring as a *scaffolding* process: the tutor controls elements beyond the learner’s capacity so the learner can complete what is within reach. Functions include recruitment, reduction of degrees of freedom, direction maintenance, marking critical features, frustration control, and demonstration.

FACT: Palincsar & Brown (1984, *Cognition and Instruction*, 1(2), 117–175, doi:10.1207/s1532690xci0102_1) — reciprocal teaching: tutor and students take turns leading dialogue (summarize, question, clarify, predict); adult models then fades leadership; improves comprehension and transfer relative to typical classroom practice in reported studies. Collins et al. treat this as a canonical CA method stack for reading — turn-taking + fade of expert lead.

FACT (review grain): Dennen (2004, “Cognitive apprenticeship in educational practice,” in *Handbook of Research on Educational Communications and Technology*) — synthesizes research on scaffolding, modeling, mentoring, and

coaching as CA strategies; notes modeling/coaching/fading as predominant methods, with scaffolding as part of coaching; fading is gradual (hints become less frequent/detailed), not abrupt abandonment. Warns that implementations vary — label ≠ fidelity.

Bridge to SAFE-FADE / SAFE-HINT / SAFE-HELP: Worked-example fading and contingent peeks are *product instantiations* of scaffolding+fading. Cognitive apprenticeship adds the **sociology**: who leads the dialogue, when the expert models, when the learner must articulate. Reciprocal-style turn-taking on a hard item (student leads a “what do we know / what’s the join?” beat) is a HITL technique, not a Discord forum.

Kill: Permanent scaffold as kindness brand; reciprocal-teaching costume without turn-taking.

Survive: Wood-function scaffolds with scheduled fade; student-led beats.

XCV.6 Boundaries — situated culture ≠ lore franchise; CA ≠ discovery

FACT / boundary (Brown et al., 1989): Situated cognition argues for authentic activity and culture of practice — not that any immersive story world equals apprenticeship. SAFE-STORYLOAD already budgets narrative load; CA does not license lore walls as “culture.”

Boundary with SAFE-PF: Exploration in Collins’s sense can include learner-set goals; Kapur PF is a *specific* generate→consolidate design. Do not rename every explore beat as PF, and do not run invent-first when the playbook called for modeling first on thin priors.

Boundary with Kirschner/Sweller/Clark (2006) critique (already in XCIV): Cognitive apprenticeship is *not* minimal guidance forever. Modeling and scaffolding are strong guidance; fading withdraws them as competence grows. Branding “apprenticeship” while shipping unguided stew fails both CA and the guidance literature.

Kill: Lore≡culture-of-practice; CA as discovery cosplay; never-guide brand.

Survive: Guided visibility → fade; authenticity = real problem work + audible thinking, not franchise trivia.

XCV.7 Product surface — SAFE-APPRENTICE claim contract

Surface	Required behavior	Banned substitute
Tutor playbook	Phase checklist: model (short) → coach → fade; articulation required	“Be warm / Socratic” vibes only
Live session QA	Log ca_phase; sample for modeling-without-attempt / never-fade	Talk-% only QA
Coach / Solver	Phase chip aligned with fade×hint; SE before wrap	AI monologue labeled “modeling”
Map brief	Heuristic/control focus for the target join	Tip flood / lore dump

AAR / Notes	Reflection on <i>student</i> process vs canonical	Tip-of-the-day dump
Parent copy	“We show how experts think, then you take the wheel”	Master/guild / hours-under-tutor
Marketing	Method stack + solo proof	“Cognitive apprenticeship™” without phases
Analytics	Phase dwell, fade events, articulation rate, solo transfer	Apprenticeship Score™ / modeling-minutes NS

Competitive foil: ChatGPT = fluent modeling without coaching/fade/articulation. Khan = content without audible control craft. Duo = practice without expert thinking made visible. MindCraft = **visible craft + transferred responsibility** under FEI.

XCV.8 Doctrine — SAFE-APPRENTICE (provisional)

1. **Visibility first** — model *control/heuristic* moves, not only steps (Collins; Schoenfeld).
2. **Performance under coaching** — most minutes = student attempt with in-situ support (coaching ≠ lecture).
3. **Scaffold then fade** — Wood functions + SAFE-FADE/HINT; no permanent crutch.
4. **Articulation required** — student-generated why before bottom-out (SAFE-EXPLAIN / SE).
5. **Reflection on their work** — compare learner process to canonical (AAR/ANNOT light).
6. **Exploration gated** — challenge/PF only when priors/affect allow (SAFE-PF / RETRIEVE).
7. **No costume metrics** — ban Apprenticeship Score™, guild cosplay, modeling-minutes NS.
8. **Shared language** — human tutor and AI coach use the same phase labels; fidelity > brand word.

Confidence: High that Collins et al. (1989/1991) define a method stack (model/coach/scaffold/fade + articulate/reflect/explore) aimed at making cognition visible. High that Brown et al. (1989) situate CA as an alternative to inert didactic knowledge. High that Schoenfeld (1985) supplies a math-relevant control/heuristic/belief frame used as a CA exemplar in that literature. High that Wood et al. (1976) and Palincsar & Brown (1984) specify scaffolding and reciprocal fade-of-leadership mechanisms. Medium that MindCraft HITL can hit phase fidelity at scale without over-scripting warmth away — run APPRENTICE-* with QA sampling. Medium that AI coach “modeling” can externalize control without becoming monologue — co-test with SAFE-EXPLAIN length caps.

XCV.9 Experiments

ID	Question	Design	Primary
APPRENTICE-1	Playbook with explicit model→coach→fade+ articulate vs explain-pour control	A/B sessions	solo_transfer_pass; articulation rubric

APPRENTICE-2	Short control-talk model vs steps-only model before same coaching	A/B	Near transfer; student control moves coded
APPRENTICE-3	Within-session fade schedule vs constant high scaffold	A/B	Solo transfer; peek binge; dependency
APPRENTICE-4	Reciprocal student-led beat (1 item) vs tutor-led whole session	A/B	Construction share; transfer
APPRENTICE-5	Parent CBC: “show thinking then you take over” vs “master tutor explains” vs “AI always models”	CBC	WTP; trust (SAFE-WTP)
APPRENTICE-QUAL	10 tutors + 10 Maya: phenomenology of phase checklist vs vibe coaching	Qual	Fidelity codebook; costume markers

Falsifier: Explain-pour equals or beats model→coach→fade on solo transfer *and* articulation for our secondary cohort → keep CA language internal; do not market apprenticeship.

Falsifier: Steps-only model equals control-talk model → demote control-talk chrome; still ban monologue-as-modeling.

Falsifier: Parents prefer master-explains CBC and reject take-the-wheel → rewrite copy before surrendering to explain-first brand.

Pre-register: APPRENTICE- before any “cognitive apprenticeship / guild / master tutor / Apprenticeship Score™” campaign (SAFE-LABMETA). Do not steal FADE-/HINT-/TALK-/PF-* arms — those own example stage, peek pricing, construction telemetry, and invent-first sequencing; this family owns **playbook phase fidelity** across the stack.

XCV.10 So what for MindCraft commercially

- **Copy:** “We make the thinking visible — then you take the wheel.” Never “study under a master” guild cosplay or “watch the AI solve everything.”
- **Product:** ca_phase on session/coach events; playbook checklist; articulation gate before hard reveal; fade schedule linked to SAFE-FADE/HINT.
- **Positioning:** Against ChatGPT monologue-as-tutor *and* content libraries without craft visibility; for inspectable apprenticeship *methods*, not costumes.
- **Metric:** solo transfer + articulation quality + fade events — demote modeling-minutes / Apprenticeship Score™ / talk-% alone.
- **Kill list:** Guild cosplay; modeling≡lecture; never-fade; talk-%≡CA; Apprenticeship Score™; Discord≡apprenticeship; reciprocal costume without turns; CA ads without phase logs.

- **Growth:** Tutor ops hire/train on phase fidelity (feeds SAFE-HITL / WORKFORCE); parent decks sell take-the-wheel honesty.
 - **Vision:** Maya hears how a competent solver *manages* stuckness — then does it herself — so identity shifts from “I need someone to show me” to “I can run the craft.”
-

References (verified)

- Brown, J. S., Collins, A., & Duguid, P. (1989). Situated cognition and the culture of learning. *Educational Researcher*, 18(1), 32–42. <https://doi.org/10.3102/0013189X018001032>
 - Collins, A., Brown, J. S., & Holum, A. (1991). Cognitive apprenticeship: Making thinking visible. *American Educator*, 15(3), 6–11, 38–46.
 - Collins, A., Brown, J. S., & Newman, S. E. (1989). Cognitive apprenticeship: Teaching the crafts of reading, writing, and mathematics. In L. B. Resnick (Ed.), *Knowing, learning, and instruction: Essays in honor of Robert Glaser* (pp. 453–494). Lawrence Erlbaum Associates.
 - Dennen, V. P. (2004). Cognitive apprenticeship in educational practice: Research on scaffolding, modeling, mentoring, and coaching as instructional strategies. In D. H. Jonassen (Ed.), *Handbook of research on educational communications and technology* (2nd ed., pp. 813–828). Lawrence Erlbaum Associates.
 - Palincsar, A. S., & Brown, A. L. (1984). Reciprocal teaching of comprehension-fostering and comprehension-monitoring activities. *Cognition and Instruction*, 1(2), 117–175. https://doi.org/10.1207/s1532690xci0102_1
 - Schoenfeld, A. H. (1985). *Mathematical problem solving*. Academic Press.
 - Wood, D., Bruner, J. S., & Ross, G. (1976). The role of tutoring in problem solving. *Journal of Child Psychology and Psychiatry*, 17(2), 89–100. <https://doi.org/10.1111/j.1469-7610.1976.tb00381.x>
-

Part XCVI — Attention Residue & Device Distraction in Practice

Chapter status: Living evidence + practice UX brief — Researcher tick 2026-08-07

Primary question: How should MindCraft treat **attention residue** and **device/media distraction** during practice and tutoring so FEI attempts stay cognitively available — without Focus Score™ theater, phone-shame branding, or pretending “mobile-first” erases switching costs?

Owners: Practice UX · Coach / session design · Parent honesty · Brand · Red Team

Commercial job: Ship a **SAFE-ATTN** doctrine: protect working memory for the attempt; make task-switch tax inspectable; offer opt-in focus modes and tech breaks — never monetize multitasking heroics or sell distraction-free as guaranteed ACT points.

Builds on: Parts XXVI (CLT), XLI (DD × anxiety), XLVII (sleep/stress), LXIII (exam pressure), LXXXV (story-load), XCIII (retrieval failure modes), XCV (apprenticeship phases need available attention). Sibling surfaces: habit cues (XLIII) without streak-as-focus; privacy (LXIX) if we ever log “distraction” signals.

XCVI.1 Why this chapter exists

Brand temptation: “Our app is so engaging students never look away.” Competitor temptation: “Practice anywhere — phone in hand, tabs open, notifications on.” Both ignore the load-bearing constraint: mathematical thinking is a limited-capacity resource, and unfinished or competing cognitions steal it.

Sophie Leroy’s *attention residue* names the leftover cognitions about Task A that persist after switching to Task B. Smartphones add a second tax: even *unused* phones can drain available capacity (Ward et al.). Media multitasking while studying is common and reliably associated with worse academic outcomes in reviews and meta-analyses — not because teens are morally weak, but because switching and divided attention interfere with encoding and control.

MindCraft already fights fake struggle theater (SAFE-RETRIEVE), lore load (SAFE-STORYLOAD), and monologue token tax (SAFE-EXPLAIN). The open question: do we **design for residue and device presence**, or ship notification “engagement” and call interrupted attempts grit?

FOUNDER BELIEF under audit: Fewer mid-attempt switches and lower phone salience will raise `solo_transfer_pass` and cut freeze/binge-peek vs notification-on defaults — without a Focus Score™.

Claims we refuse as doctrine:

1. Focus Score™ / Deep Work Score™ as North Star.
2. Phone-shame or “kids these days” moralizing as pedagogy.
3. “Mobile-first multitasking is fine for math” without residue accounting.
4. Always-on push notifications as engagement science.
5. Streak / DAU protection via interrupt-driven re-engagement during an attempt.
6. Presence of phone ≡ proof of modern learning.
7. Focus-mode ads that promise guaranteed ACT points or “cured distraction.”
8. Covert keystroke/gaze distraction surveillance without SAFE-PRIVACY gates.

XCVI.2 Constructs (keep mechanisms distinct)

Construct	Research meaning	MindCraft analogue	Not the same as
Attention residue	Cognitions about prior Task A persist into Task B	Mid-mission chat ping; unfinished homework tab; tutor Slack	Mere time-on-app
Task-switch cost	Performance drop when changing task sets	Alt-tab mid-item; notification → return	Deliberate interleaving of <i>problem types</i> (Part XXXIX)
Media multitasking	Concurrent media streams / off-task tech while studying	Phone + practice; social tab + Solver	Dual coding <i>within</i> one math item (Part XCVII queue)
Phone presence / brain drain	Own phone nearby reduces available WMC/Gf even unused	Phone on desk during timed set	Using phone <i>as</i> the practice device intentionally

SAFE-ATTN	When/how to protect attempt windows	This chapter	Productivity-guru cosplay
-----------	-------------------------------------	--------------	---------------------------

Operational definition (HYPOTHESIS): A mission is *SAFE-ATTN compliant* when (a) the product does not fire engagement pushes during an active attempt, (b) optional focus mode suppresses non-essential chrome and suggests phone-away for desk sessions, (c) tech-break / resume affordances exist so unfinished social cognitions can close cleanly (Leroy completion logic), (d) analytics separate attempt_interrupted from productive struggle modes (SAFE-RETRIEVE), (e) outcomes judged on transfer and construction — not Focus Score™ or minutes-without-switch vanity, and (f) any distraction telemetry is opt-in, non-biometric, and privacy-bound.

XCVI.3 Attention residue — unfinished cognitions steal the next task

FACT: Leroy (2009, *Organizational Behavior and Human Decision Processes*, 109(2), 168–181, doi:10.1016/j.obhdp.2009.04.002) — defines *attention residue* as the persistence of cognitive activity about Task A after one has stopped working on Task A and is performing Task B. Two experiments show people struggle to fully transition attention away from unfinished prior work; subsequent Task B performance suffers. Finishing Task A before switching helps, but is not always enough: time pressure while finishing the prior task aids disengagement and contributes to higher subsequent performance. The mechanism is not “lazy switching” folklore — it is residual cognitions occupying capacity needed for the next task.

Applied (HYPOTHESIS): Mid-attempt notifications, coach chat bubbles, and “come back!” streaks create *unfinished social/app tasks* that leave residue on the math item. Product should prefer **close-the-loop** patterns: finish or park the interrupt (snooze with explicit resume) before demanding germane load on the problem.

Kill: Interrupt-driven engagement during active attempts as “habit science.”

Survive: Protect attempt windows; allow clean completion/park of interrupts.

XCVI.4 Device distraction & media multitasking — field pattern, not vibes

FACT: Rosen, Carrier, & Cheever (2013, *Computers in Human Behavior*, 29(3), 948–958, doi:10.1016/j.chb.2012.12.001) — observed 263 middle-school, high-school, and university students studying 15 minutes at home. Participants averaged **less than six minutes** on task before switching, most often to technological distractions (social media, texting) or task-switching preference. Students who accessed Facebook had lower GPAs than those who avoided it; higher use of study strategies predicted more on-task behavior. Authors discuss short “technology breaks” and metacognitive strategies about when interruptions hurt — not total abstinence theater.

FACT: May & Elder (2018, *International Journal of Educational Technology in Higher Education*, 15, 13, doi:10.1186/s41239-018-0096-z) — literature review: media multitasking interferes with attention and working memory and is associated with worse GPA, test performance, recall, reading comprehension, note-taking, self-regulation, and efficiency, in class and while studying. Students often **misjudge** how much multitasking will cost them. Self-regulation support is a promising intervention direction.

FACT: Kates, Wu, & Coryn (2018, *Computers & Education*, 127, 107–112, doi:10.1016/j.compedu.2018.08.012) — meta-analysis of mobile phone use (non-educational-improvement use) and academic outcomes over 2008–2017:

overall average effect $r = -0.162$ (95% CI -0.196 to -0.128) — a **small negative** association, not a cartoon catastrophe and not zero. Effect sizes vary by design and construct; do not market “phones destroy learning” as a 2-sigma claim.

Applied (HYPOTHESIS): MindCraft should treat off-task media as a **moderated tax** on FEI attempts: design for fewer switches and better metacognition, instrument interruption, and avoid both denial (“multitask freely”) and moral panic (“ban phones or fail”).

Kill: Phone-apocalypse marketing; “multitasking digital natives” as free pass.

Survive: Small-to-medium evidenced tax; self-regulation affordances; tech breaks.

XCVI.5 Mere presence — the phone can tax you without a tap

FACT: Ward, Duke, Gneezy, & Bos (2017, *Journal of the Association for Consumer Research*, 2(2), 140–154, doi:10.1086/691462) — “brain drain” hypothesis: mere presence of one’s own smartphone can occupy limited-capacity cognitive resources. Across experiments, available working memory capacity and fluid intelligence measures were worse when the phone was more salient (e.g., desk vs other room), even when participants were not using the phone and did not report thinking about it. Costs were highest for those highest in smartphone dependence.

Applied (FOUNDER BELIEF → testable): For desk/laptop practice and tutor sessions, default copy can invite “phone in another room / face-down out of reach” as a *capacity* tip — framed as working-memory hygiene, not character judgment. On phone-as-device sessions, the tradeoff flips: the practice surface *is* the phone — then minimize *competing* apps/notifications rather than pretending presence-drain vanishes.

Boundary: Ward et al. are lab cognitive-capacity tasks, not ACT RCTs. Do not claim “phone in bag = +N ACT points.”

Kill: Guaranteed-score focus kits; shaming dependent users in parent dashboards.

Survive: Optional salience reduction; honesty that unused phones can still cost capacity.

XCVI.6 Cognitive control — heavy multitaskers are not automatically better switchers

FACT: Ophir, Nass, & Wagner (2009, *Proceedings of the National Academy of Sciences*, 106(37), 15583–15587, doi:10.1073/pnas.0903620106) — heavy media multitaskers (vs light) showed greater susceptibility to interference from irrelevant environmental stimuli and memory representations, and *worse* performance on a task-switching test — consistent with a breadth-biased control style that filters poorly. Association, not proven causation from multitasking volume alone; still kills the folk claim “kids who multitask all day are better at switching for homework.”

Bridge to SAFE-RETRIEVE / SAFE-EXPLAIN: Blanking and monologue wraps already burn WM. Residue + media interference compound those failure modes. Do not label notification-induced stalls as productive struggle.

Kill: Multitasking-as-talent brand.

Survive: Filter-friendly session design; one primary cognitive thread during attempts.

XCVI.7 Product surface — SAFE-ATTN claim contract

Surface	Required behavior	Banned substitute
Practice attempt	No marketing/streak pushes mid-item; visible “attempt in progress” lock	DAU pings during problem
Focus mode (opt-in)	Suppress non-essential chrome; suggest phone-away for desk; timed tech break	Focus Score™; forced logout shame
Interrupt handling	Snooze/park with resume; log attempt_interrupted	Silent tab-switch ignored in analytics
Tutor sessions	Brief “devices down for attempt windows”; tech break between items	Phone-confiscation cosplay; emotion-AI attention cops
Parent copy	“Deep attempts need quiet WM — here’s an optional focus tip”	Shame ranks for “phone addiction”; Anxiety Score™
Marketing	Capacity honesty; self-regulation aids	“Never distracted again” / guaranteed points
Analytics	Interrupt rate, resume latency, solo transfer	Focus Score™ / minutes-on-task NS alone
Privacy	No covert gaze/mic distraction scoring	Empathy-camera attention (SAFE-PRIVACY ban)

Competitive foil: Duo = streak/notification engagement that may interrupt encoding. ChatGPT = infinite tab switching without attempt protection. Khan = video + phone parallel common. MindCraft = **attempt windows protected for construction**, with optional focus hygiene — still FEI, not productivity-guru SKU.

XCVI.8 Doctrine — SAFE-ATTN (provisional)

1. **Residue is real** — unfinished cognitions tax Task B (Leroy); close or park interrupts before demanding the join.
2. **Media multitasking costs learning on average** — reviews/meta show negative academic associations (May & Elder; Kates et al.); size is often small-to-moderate, not apocalyptic.
3. **Presence can drain** — unused phones may reduce available capacity (Ward et al.); offer desk hygiene tips without shame.
4. **Heavy multitasking ≠ better switching** — Ophir et al. association cuts against “digital native switcher” ads.
5. **Protect FEI attempt windows** — no engagement pushes mid-item; separate interrupt telemetry from productive struggle.
6. **Tech breaks > total war** — Rosen et al. style short breaks + metacognition beat abstinence theater.
7. **No Focus Score™ NS** — prove with solo transfer / construction, not vanity focus minutes.
8. **Privacy bound** — no biometric attention policing; opt-in only for any distraction signals.

Confidence: High that Leroy (2009) establishes attention residue as a measured Task A→B performance mechanism in experiments. High that Ward et al. (2017) show mere phone presence can reduce available WMC/Gf in lab tasks. High that May & Elder (2018) and Kates et al. (2018) document negative media-multitasking / phone-use associations with academic outcomes at review/meta grain. High that Rosen et al. (2013) observe frequent short on-task bouts and tech-driven switches in naturalistic studying. Medium that Ophir et al. (2009) generalize to MindCraft HS cohorts — association study, replication debates exist in the broader MMI literature; use as anti-folk-claim, not as trait diagnosis. Medium that product focus-mode + push suppression will move FEI metrics — run ATTN-* before Focus brand campaigns.

XCVI.9 Experiments

ID	Question	Design	Primary
ATTN-1	Mid-attempt push suppression vs default engagement pushes	A/B	solo_transfer_pass; interrupt rate; peek binge
ATTN-2	Opt-in focus mode (chrome quiet + phone-away tip) vs control	A/B	Transfer; resume latency; self-report residue
ATTN-3	Scheduled tech-break between items vs no break / vs mid-item free switch	A/B	On-task bout length; transfer
ATTN-4	Desk session: phone other-room tip vs phone-on-desk (observational + randomized tip)	A/B / hybrid	Capacity proxies; FEI events
ATTN-5	Parent CBC: “protect attempt windows” vs “streaks keep them coming back” vs “Focus Score™”	CBC	WTP; trust (SAFE-WTP)
ATTN-QUAL	10 Maya: phenomenology of residue after chat/social interrupt during a hard item	Qual	Interrupt codebook; shame markers

Falsifier: Push-on during attempts equals or beats push-off on solo transfer *and* retention for secondary cohort → keep ATTN internal; do not market focus windows.

Falsifier: Focus mode harms anxious students (avoidance/shame) more than it helps transfer → demote forced focus; keep optional + anxiety-gated (XLI).

Falsifier: Parents prefer streak-interrupt CBC and reject attempt-protection copy → rewrite GTM before surrendering to notification-growth brand.

Pre-register: ATTN- before any “Focus Score™ / deep work / phone-free guaranteed gains / distraction-cured” campaign (SAFE-LABMETA). Do not steal RETRIEVE-/HABIT-/PRIV- arms — those own struggle-mode taxonomy, cue design, and affect telemetry ethics; this family owns **attempt-window protection and device/residue hygiene**.

XCVI.10 So what for MindCraft commercially

- **Copy:** “Hard math needs a clear head — we protect your attempt window.” Never “multitask like a pro” or “we’ll shame the phone out of you.”
 - **Product:** Mid-attempt push lock; opt-in focus mode; tech-break + resume; attempt_interrupted event distinct from freeze/TOT modes.
 - **Positioning:** Against notification-growth edtech and productivity-guru Focus Score™ SKUs; for capacity-honest practice that still ships FEI proof.
 - **Metric:** solo transfer + interrupt/resume quality — demote Focus Score™ / raw minutes-on-task / streak pings during items.
 - **Kill list:** Focus Score™ NS; phone-shame; multitasking-as-talent; mid-attempt DAU pushes; presence≡modern; guaranteed ACT from focus kits; covert attention surveillance.
 - **Growth:** Parent decks sell attempt hygiene as care, not control; tutor ops brief device norms between items.
 - **Vision:** Maya can finish a join without half her working memory still arguing with a group chat — so evidence of competence can land as identity, not as another interrupted almost.
-

References (verified)

- Kates, A. W., Wu, H., & Coryn, C. L. S. (2018). The effects of mobile phone use on academic performance: A meta-analysis. *Computers & Education*, 127, 107–112. <https://doi.org/10.1016/j.compedu.2018.08.012>
 - Leroy, S. (2009). Why is it so hard to do my work? The challenge of attention residue when switching between work tasks. *Organizational Behavior and Human Decision Processes*, 109(2), 168–181. <https://doi.org/10.1016/j.obhdp.2009.04.002>
 - May, K. E., & Elder, A. D. (2018). Efficient, helpful, or distracting? A literature review of media multitasking in relation to academic performance. *International Journal of Educational Technology in Higher Education*, 15, 13. <https://doi.org/10.1186/s41239-018-0096-z>
 - Ophir, E., Nass, C., & Wagner, A. D. (2009). Cognitive control in media multitaskers. *Proceedings of the National Academy of Sciences*, 106(37), 15583–15587. <https://doi.org/10.1073/pnas.0903620106>
 - Rosen, L. D., Carrier, L. M., & Cheever, N. A. (2013). Facebook and texting made me do it: Media-induced task-switching while studying. *Computers in Human Behavior*, 29(3), 948–958. <https://doi.org/10.1016/j.chb.2012.12.001>
 - Ward, A. F., Duke, K., Gneezy, A., & Bos, M. W. (2017). Brain drain: The mere presence of one’s own smartphone reduces available cognitive capacity. *Journal of the Association for Consumer Research*, 2(2), 140–154. <https://doi.org/10.1086/691462>
-

Part XCVII — Dual Coding & Diagram FormatId Load

Chapter status: Living evidence + FormatId / figure UX brief — Researcher tick 2026-08-07

Primary question: When do words+pictures (dual coding / multimedia principles) help MindCraft FormatId surfaces (diagram, coordinate_graph, number_line, ...), and when do figures add **extraneous load**, split attention, or learning-styles costume — so we ship conversion competence without “visual learning” theater?

Owners: Practice / Solver figure UX · Bank & generation (GENQ figures) · Map format CTAs · Brand · Red Team

Commercial job: Ship a **SAFE-DUAL** doctrine: structural dual-channel design under CLT budgets; diagrams earn their keep via integration and conversion evidence; never sell Dual Coding Score™, always-more-visuals, animation-as-default, or visual-learner meshing as science.

Builds on: Parts XXVI (CLT), LXXXII (SAFE-FORMAT), LXXXV (SAFE-STORYLOAD), XCII (SAFE-EXPLAIN), LXXII (SAFE-GENQ). Sibling: XCVI (device residue ≠ pictorial channel). Product seams: FormatId, format-tagged bank, figure/render specs, coach wrap next to stems.

XCVII.1 Why this chapter exists

Brand temptation: “We use dual coding — words and pictures — so students learn twice as deep.” Competitor temptation: decorate every stem with stock art, autoplay animations, or a “visual learner” quiz. Both confuse **Paivio/Mayer dual-channel design** with **pretty multimedia** and **learning-styles meshing** (already banned under SAFE-FORMAT).

MindCraft’s format axis already treats vessels as diagnosis objects. Open load question: a diagram FormatId item can be the *target competence* (read/produce the vessel) or a *coach scaffold* (words+figure explaining a symbolic item). Those jobs have different design rules. Stacking narration + on-screen text + busy figure can *hurt*. Separating legend far from the figure forces split attention. Decorative story chrome next to a graph re-imports SAFE-STORYLOAD failures.

FOUNDER BELIEF under audit: Contiguous, coherent, structurally labeled words+diagrams (plus conversion practice across FormatIds) will raise `format_hop_solo_transfer_pass` vs decoration-heavy or styles-matched “visual track” defaults — without a Dual Coding Score™.

Claims we refuse as doctrine:

1. Dual Coding Score™ / Visual Engagement Score™ as North Star.
 2. Always-add-a-picture ≡ dual coding science.
 3. Decorative art / franchise chrome ≡ FormatId coverage or learning.
 4. Learning-styles meshing (“visual learners get diagrams”) as product law.
 5. Animation/video always better than static annotated diagrams.
 6. Stacked animation + narration + on-screen text as “thorough dual coding.”
 7. Split text beside distant figure as “more information.”
 8. Unverified generated figures as diagnostic ground truth (SAFE-GENQ).
 9. Guaranteed ACT points from “dual coding method” ads.
-

XCVII.2 Constructs (keep mechanisms distinct)

Construct	Research meaning	MindCraft analogue	Not the same as
Dual coding (Paivio)	Separate verbal + nonverbal systems; referential links strengthen memory	Stem + structural figure; coach labels tied to parts	Learning-styles preference
CTML dual-channel	Limited pictorial + verbal channels; select—organize—integrate	Contiguous labels; modality-aware coach	Unlimited media stack
Multimedia principle	Words+pictures can beat words alone <i>when designed</i>	Dual-channel coach cards	Always-more-visuals
Split attention	Mentally integrating separated sources burns WM	Legend far from graph; footnote under fold	Deliberate conversion practice
Redundancy / coherence	Duplicate channels or seductive extras can hurt	Narration + full on-screen dump + figure	Thoroughness theater
FormatId vessel	Input representation as diagnosis object	diagram / coordinate_graph / ...	Decoration
Conversion	Translate across registers (Duval; SAFE-FORMAT)	Diagram↔equation hops	Treatment within one vessel
SAFE-DUAL	When/how figures load under dual-channel + CLT	This chapter	Dual-coding costume ads

Operational definition (HYPOTHESIS): A figure treatment is *SAFE-DUAL compliant* when (a) the pictorial element is **structurally necessary** for the learning object (vessel reading, conversion, or constrained MER function — Ainsworth), not decorative, (b) corresponding words and picture parts are **spatially/temporally contiguous**, (c) extraneous words/sounds/chrome are cut (coherence), (d) redundant on-screen text is not stacked on narration+animation without cause, (e) outcomes are judged on retention/transfer and format_hop_solo_transfer_pass — not Dual Coding Score™ or “visual minutes,” and (f) any generated figure used diagnostically is verify-gated (SAFE-GENQ).

XCVII.3 Dual coding is a memory architecture — not a marketing sticker

FACT: Paivio (1986, *Mental Representations: A Dual Coding Approach*, Oxford University Press) — performance on memory and related tasks is mediated by interconnected **verbal** and **nonverbal (imagery)** systems; concrete, imageable materials and referential links between systems support richer encoding than verbal-only routes alone. Dual coding is a theory of representation, not a license to wallpaper every worksheet.

FACT: Clark & Paivio (1991, *Educational Psychology Review*, 3(3), 149–210, doi:10.1007/BF01320076) — review dual coding implications for education: concreteness, imagery, and dual verbal–nonverbal routes matter for learning design; still not equivalent to diagnosing “visual learners” and matching preference.

Applied (HYPOTHESIS): MindCraft should use dual coding as a **design constraint** (referential links between words and structural pictures) when the math object is spatial/graphic — and refuse dual-coding *brand* when the “picture” is unrelated atmosphere.

Kill: “Science-backed dual coding” ads that mean stock illustration.

Survive: Referential word↔structure links on FormatId surfaces that require them.

XCVII.4 Multimedia learning: limited channels, active integration

FACT: Mayer & Moreno (2003, *Educational Psychologist*, 38(1), 43–52, doi:10.1207/S15326985EP3801_6) — cognitive theory of multimedia learning: dual-channel, limited-capacity, and active-processing assumptions; overload scenarios and theory-based load-reduction methods. Cognitive load is central — more media ≠ more learning.

FACT: Moreno & Mayer (1999, *Journal of Educational Psychology*, 91(2), 358–368, doi:10.1037/0022-0663.91.2.358) — modality and contiguity: better learning when corresponding words and pictures are coordinated. Do not treat “add text everywhere” as dual coding.

FACT: Mayer, Heiser, & Lonn (2001, *Journal of Educational Psychology*, 93(1), 187–198, doi:10.1037/0022-0663.93.1.187) — extraneous/seductive material can impair learning (**coherence**). Seductive details are not enrichment.

Bridge to SAFE-STORYLOAD / SAFE-EXPLAIN: Lore beside a graph and long monologue wraps compete for the same limited channels. Coach text next to a figure must be short, principle-linked, and contiguous.

Kill: Seductive detail as engagement; monologue+busy figure as thoroughness.

Survive: Weeded, contiguous multimedia under capacity limits.

XCVII.5 Split attention — separated sources tax integration

FACT: Chandler & Sweller (1991, *Cognition and Instruction*, 8(4), 293–332, doi:10.1207/s1532690xci0804_2) — when learners must mentally integrate mutually referring sources (diagram + separate text), **split-source** formats can impose heavy extraneous load; **physically integrating** related information often improves learning when integration is necessary. Nonessential add-ons can hurt even when “integrated.”

Applied (HYPOTHESIS): Keep axis labels, key lengths, and callouts **on or adjacent to** the figure; avoid legend-in-sidebar / explanation-three-scrolls-away for novices. Coach cards that reference “the red segment” must point or highlight.

Kill: Split-pane “more information density” as pedagogy.

Survive: Integrated labeling; necessity test before adding commentary layers.

XCVII.6 Diagrams are active constructions — not free downloads

FACT: Hegarty (1992, *Journal of Experimental Psychology: Learning, Memory, and Cognition*, 18(5), 1084–1102, doi:10.1037/0278-7393.18.5.1084) — *mental animation*: inferring kinematics from static mechanical diagrams is an active, capacity-limited process (piecemeal causal sequence); harder against causal order. Diagrams do not pour understanding in; learners build models.

FACT: Hegarty, Kriz, & Cate (2003, *Cognition and Instruction*, 21(4), 209–249, doi:10.1207/s1532690xci2104_1) — mental vs external animation in understanding mechanical systems; spatial ability and knowledge shape what is learned from displays; comprehension is jointly determined by display and learner. Animations are not automatically superior to static displays that support mental animation.

FACT: Mayer, Hegarty, Mayer, & Campbell (2005, *Journal of Experimental Psychology: Applied*, 11(4), 256–265, doi:10.1037/1076-898X.11.4.256) — across four topics, paper-based **static diagrams + text** often matched or beat computer **narrated animations** on retention/transfer comparisons — supporting a *static media* hypothesis: static annotated illustrations can reduce extraneous processing and foster germane processing relative to narrated animation under those materials.

Applied (FOUNDER BELIEF → testable): Default FormatId scaffolds for geometry/graph items should prefer **static, annotated, contiguous** figures; treat animation as an opt-in for specific dynamic systems, not brand “modern = moving.” Coordinate-graph FormatId competence is reading/producing the static vessel ACT uses — not watching a fly-through.

Kill: Animation-first Solver / “immersive diagram video” as default dual coding.

Survive: Static annotated figures; animation only when motion is the learning object and load is managed.

XCVII.7 Learning styles meshing is not dual coding

FACT: Pashler, McDaniel, Rohrer, & Bjork (2008, *Psychological Science in the Public Interest*, 9(3), 105–119, doi:10.1111/j.1539-6053.2009.01038.x) — review finds inadequate evidence for the **meshing hypothesis** (match instruction to preferred style, e.g., visual vs verbal learners). Preferring pictures is not the same as needing pictures for a content-appropriate dual-channel design.

Bridge to SAFE-FORMAT: Format Personality™ / visual-learner tracks remain killed. Dual coding says: for *this math object*, use words and structural pictures under capacity rules — not “Maya self-reports visual, so only diagrams forever.”

Kill: Style quiz → vessel lock.

Survive: Task-demand FormatId diagnosis + conversion practice.

XCVII.8 Product surface — SAFE-DUAL claim contract

Surface	Required behavior	Banned substitute
Diagram / graph stems	Structural figure; labels contiguous; no decorative chrome in critical area	Stock art “for dual coding”
Coach wrap beside figure	Short principle/misconception; points to parts; fades with expertise (SAFE-EXPLAIN)	Monologue + figure dump
Conversion practice	Explicit diagram↔symbolic hops with supports (SAFE-FORMAT)	Re-teach topic in strong vessel only
Animation	Opt-in; motion must be the object; segment + weed	Autoplay narrated video default
Story wrap	Structure-encoding only; lore off the critical path (SAFE-STORYLOAD)	Franchise sticker on axes
Generation	Verify figure+key before diagnostic use (SAFE-GENQ)	Unchecked LLM diagrams as ground truth
Marketing	“We design words+figures so the structure is readable”	Dual Coding Score™ / visual-learner SKU / guaranteed points
Analytics	format_hop_solo_transfer_passe; figure-load A/Bs	Visual Engagement Score™ / minutes-with-picture NS

Competitive foil: Brilliant = delight visuals that may under-specify conversion. Khan = video+worksheet split attention common. ChatGPT = text-only fluency without vessel diagnosis. Duo = decoration without dual-channel discipline. MindCraft = **FormatId honesty + contiguous structural dual-channel design** — still FEI, not multimedia theater.

XCVII.9 Doctrine — SAFE-DUAL (provisional)

1. **Dual coding ≠ decoration** — Paivio/Clark referential links require structural nonverbal content, not wallpaper.
2. **Channels are limited** — Mayer/Moreno overload rules; weed seductive detail (coherence).
3. **Integrate or pay the tax** — Chandler/Sweller split attention; keep words and picture parts contiguous when mutual reference is required.
4. **Diagrams demand construction** — Hegarty mental animation; do not treat viewing as mastery.
5. **Static annotated often wins** — Mayer et al. (2005) static media results; animation is conditional, not default modernity.
6. **No styles meshing** — Pashler et al.; task demand + FormatId gaps, not preference quizzes.
7. **Conversion remains the object** — SAFE-FORMAT: dual-channel scaffolds serve hops, not single-vessel comfort forever.
8. **No Dual Coding Score™** — prove with transfer and format hops; GENQ figures verify-gated.

Confidence: High that Paivio (1986) and Clark & Paivio (1991) establish dual verbal–nonverbal systems as a serious framework. High that Mayer & Moreno (2003) and Moreno & Mayer (1999) ground limited-capacity dual-channel principles. High that Mayer, Heiser, & Lonn (2001) show extraneous material can hurt. High that Chandler & Sweller (1991) demonstrate split-attention effects when sources must combine. High that Hegarty (1992) and Hegarty et al. (2003) show diagram comprehension is active and capacity-/knowledge-sensitive. High that Mayer et al. (2005) support static annotated media vs narrated animation on several science topics — **medium** generalization to ACT math FormatIds (run DUAL-*). High that Pashler et al. (2008) undermine learning-styles meshing as product law. Medium that contiguous-figure defaults will move `format_hop_solo_transfer_pass` — instrument before dual-coding brand campaigns.

XCVII.10 Experiments

ID	Question	Design	Primary
DUAL-1	Contiguous on-figure labels vs split legend/sidebar	A/B	Transfer; time-to-first-correct; peek binge
DUAL-2	Structural figure + short coach vs decorative figure + same coach vs text-only	A/B	Retention; <code>format_hop_solo_transfer_pass</code>
DUAL-3	Static annotated diagram vs narrated animation for same geometry/graph lesson	A/B	Transfer; extraneous-load self-report; interrupt rate
DUAL-4	Coherence weed (no seductive story chrome on figure) vs chrome-on	A/B	Transfer; story-trivia lure errors (SAFE-STORYLOAD)
DUAL-5	Parent CBC: “contiguous dual-channel figures” vs “visual learner dual coding” vs Dual Coding Score™	CBC	WTP; trust (SAFE-WTP)
DUAL-QUAL	10 Maya: where figures help vs overwhelm on ACT-like diagram items	Qual	Load codebook; conversion language

Falsifier: Decorative ≈ structural on format-hop transfer → demote figure strictness; keep FORMAT diagnosis.

Falsifier: Narrated animation reliably beats static annotated on ACT vessels without peek/avoidance → allow animation for those families only.

Falsifier: Parents prefer visual-learner CBC over contiguity copy → rewrite GTM; do not ship styles meshing.

Pre-register: DUAL- before dual-coding / visual-learner / Dual Coding Score™ / animation-first / multimedia ACT-guarantee campaigns (SAFE-LABMETA). Do not steal FORMAT-/STORYLOAD-/EXPLAIN-/GENQ-/ATTN-arms.

XCVII.11 So what for MindCraft commercially

- **Copy:** “Figures that carry structure — labels where you look.” Never “dual coding for visual learners” or Dual Coding Score™.
 - **Product:** Contiguous structural diagrams; weeded chrome; static-first; conversion hops; verify-gated GENQ figures.
 - **Positioning:** Against decoration-edtech and styles quizzes; for FormatId-honest dual-channel design under CLT.
 - **Metric:** format_hop_solo_transfer_pass + figure A/Bs — demote visual-minutes / Dual Coding Score™.
 - **Kill list:** Dual Coding Score™; always-add-picture; decoration≡learning; styles meshing; animation-default; redundancy stack; split legends; unverified figure diagnostics; ACT guarantees from multimedia ads.
 - **Growth:** Parent decks show contiguous vs split figure hygiene, not “we’re a visual platform.”
 - **Vision:** Maya can read the graph vessel and convert it to symbols without WM spent on chrome, distant legends, or a false “visual learner” identity that never leaves her comfort vessel.
-

References (verified)

- Chandler, P., & Sweller, J. (1991). Cognitive load theory and the format of instruction. *Cognition and Instruction*, 8(4), 293–332. https://doi.org/10.1207/s1532690xc0804_2
- Clark, J. M., & Paivio, A. (1991). Dual coding theory and education. *Educational Psychology Review*, 3(3), 149–210. <https://doi.org/10.1007/BF01320076>
- Hegarty, M. (1992). Mental animation: Inferring motion from static displays of mechanical systems. *Journal of Experimental Psychology: Learning, Memory, and Cognition*, 18(5), 1084–1102. <https://doi.org/10.1037/0278-7393.18.5.1084>
- Hegarty, M., Kriz, S., & Cate, C. (2003). The roles of mental animations and external animations in understanding mechanical systems. *Cognition and Instruction*, 21(4), 209–249. https://doi.org/10.1207/s1532690xc02104_1
- Mayer, R. E., Hegarty, M., Mayer, S., & Campbell, J. (2005). When static media promote active learning: Annotated illustrations versus narrated animations in multimedia instruction. *Journal of Experimental Psychology: Applied*, 11(4), 256–265. <https://doi.org/10.1037/1076-898X.11.4.256>
- Mayer, R. E., Heiser, J., & Lonn, S. (2001). Cognitive constraints on multimedia learning: When presenting more material results in less understanding. *Journal of Educational Psychology*, 93(1), 187–198. <https://doi.org/10.1037/0022-0663.93.1.187>
- Mayer, R. E., & Moreno, R. (2003). Nine ways to reduce cognitive load in multimedia learning. *Educational Psychologist*, 38(1), 43–52. https://doi.org/10.1207/S15326985EP3801_6
- Moreno, R., & Mayer, R. E. (1999). Cognitive principles of multimedia learning: The role of modality and contiguity. *Journal of Educational Psychology*, 91(2), 358–368. <https://doi.org/10.1037/0022-0663.91.2.358>
- Paivio, A. (1986). *Mental representations: A dual coding approach*. Oxford University Press.

- Pashler, H., McDaniel, M., Rohrer, D., & Bjork, R. (2008). Learning styles: Concepts and evidence. *Psychological Science in the Public Interest*, 9(3), 105–119. <https://doi.org/10.1111/j.1539-6053.2009.01038.x>
-

Part XCVIII — Metacognitive Monitoring in Gap-Scan

Chapter status: Living evidence + Diagnostic / gap-scan UX brief — Researcher tick 2026-08-07 (UTC hour 12 ≡ Red Team slot, but ch98 never written → prefer Researcher per rotation; researcher count since synthesizer v1.12 = 6 → Researcher)

Primary question: How should MindCraft treat *metacognitive monitoring* inside gap-scan and related diagnostics — as a Monitoring Score™ to gamify, a confidence vibe to inflate, or a **control signal** that classifies knowledge vs confident misconception vs uninformed miss and routes FEI without corrupting the graph?

Owners: Product (Diagnostic C4 / gap-scan / soft-wrong) · Engine (evidence_update_policy / confidence tiers) · Brand · Red Team

Commercial job: Ship a **SAFE-MONITOR** doctrine that densifies SAFE-CALIB (Part L) and SAFE-COLD (Part LXXXI): elicit the right judgment type, separate absolute vs relative accuracy, use certainty to sort misconception from lack-of-knowledge, refuse Monitoring Score™ / “raise metacognition” theater, and prove value with confidence_miss_tier + weakness stability — not Belief Score™.

Builds on: Parts L (SAFE-CALIB), XLIX (SAFE-MISCON), LXXXI (SAFE-COLD), XXXIII (AI overtrust as *system* miscalibration), XC (FEI instrumentation). Product seams: gap-scan confidence ratings, hide-correctness probes, /seed-assessment, /record-outcomes, Layer-4 event schemas.

XCVIII.1 Why this chapter exists

Part L already bans “raise confidence” as North Star and wires item-level confidence to feedback tiers. Gap-scan still tempts a second costume: **metacognition theater** — a Monitoring Score™, a “know thyself” splash, or treating a single easy/kinda/hard slider as proof of self-regulated learning.

MindCraft’s commercial wedge is not that students *feel* more metacognitive. It is that the product **uses monitoring judgments as diagnostic fuel**: distinguishing confident wrong (misconception / hypercorrection path) from low-confidence wrong (uninformed / scaffold path), and refusing to mint mastery fireworks from either. Competitors can ship a confidence emoji; few ship monitoring → classification → graph policy → FEI.

FOUNDER BELIEF under audit: Honest gap-scan requires *elicited* monitoring under destaked conditions (C4), not louder certainty — and monitoring without control (no change to routing/updates) is empty UX.

Claims we refuse as doctrine:

1. Monitoring Score™ / Metacognition % as North Star or parent KPI.
2. “We teach metacognition” ads without judgment→control wiring.
3. Global easy/kinda/hard alone ≡ calibrated monitoring science.
4. Immediate post-item confidence ≡ delayed judgment-of-learning accuracy.
5. Confidence slider without miss-class routing (Foster CA-alone null — Part L).

6. Raise-metacognition / raise-confidence as identity transformation.
7. Dunning–Kruger shaming as brand voice (Part L wound).
8. Guaranteed ACT points from “metacognitive gap scan” packaging.

XCVIII.2 Constructs (keep mechanisms distinct)

Construct	Research meaning	MindCraft analogue	Not the same as
Monitoring	Judging own knowledge/performance (Nelson & Narens)	Item confidence; concept ease ratings	Identity survey
Control	Acting on those judgments (study, skip, restudy)	Soft-wrong route; challenge unlock; graph weight	Decorative slider
EOL / JOL / FOK / RCJ	Ease-of-learning; judgment of learning; feeling-of-knowing; retrospective confidence	Gap-scan ease $\approx$ EOL-ish; post-answer surety $\approx$ RCJ; restudy prompts $\approx$ JOL	One fused “metacog” number
Absolute accuracy	Match of mean confidence to mean accuracy	Bias / overconfidence index	Discrimination
Relative accuracy	Rank-order: higher conf $\rightarrow$ more likely correct	Resolution / Gamma-style discrimination	Calibration % vanity
CRI / three-tier logic	Certainty separates misconception vs don’t-know	High-conf wrong vs low-conf wrong classes	Always-reveal keys
SAFE-MONITOR	Gap-scan monitoring doctrine under SAFE-CALIB	This chapter	Metacognition costume ads

Operational definition (HYPOTHESIS): Gap-scan monitoring is *SAFE-MONITOR compliant* when (a) judgments are **item- or concept-grain**, not only global vibe, (b) judgment **type** is labeled (prospective ease vs retrospective confidence), (c) outcomes + confidence jointly classify miss types for routing and graph updates, (d) C4 hide-correctness prevents answer-hunting that collapses monitoring into key-chasing, (e) marketing never sells Monitoring Score™, and (f) success is measured by weakness stability, miss-class hit rate, and later FEI — not by rising self-rated “metacognition.”

XCVIII.3 Monitoring and control are a system — not a sticker

FACT: Nelson & Narens (1990, in Bower, *The Psychology of Learning and Motivation*, Vol. 26, Academic Press) — metamemory framework separates **monitoring** (EOL, JOL, FOK, retrospective confidence) from **control** (allocation of

study, termination, strategy). Judgment types are not highly interchangeable; they can index different aspects of memory.

Product translation: Gap-scan's concept-level easy/kinda/hard is closer to a coarse **EOL / expected competence** signal for seeding (/seed-assessment). Practice soft-wrong "how sure?" is closer to **retrospective confidence**. Do not market them as the same "metacognition feature," and do not assume one slider fixes both.

Kill: One Metacognition Score™ merging EOL + RCJ + identity Likert.

Survive: Named judgment types → named control actions.

XCVIII.4 Delayed judgments beat fluency illusions (when the object is future recall)

FACT: Nelson & Dunlosky (1991, *Psychological Science*, 2(4), 267–270, doi:10.1111/j.1467-9280.1991.tb00147.x) — judgments of learning made after a short delay (cue-only) predict later recall far better than immediate JOLs (**delayed-JOL effect**). Immediate JOLs often track short-term accessibility, not durable learning.

FACT: Rhodes & Tauber (2011, *Psychological Bulletin*, 137(5), 791–813, doi:10.1037/a0021705) — meta-analysis: delaying JOLs yields large gains in **relative** accuracy (gamma; $g \approx 0.93$) and a small memorial benefit ($g \approx 0.08$). Delay is not magic for every product metric; it is a robust monitoring-accuracy lever for recall prediction.

Applied (HYPOTHESIS): MindCraft should not treat *immediate* post-answer confidence as a durable-mastery oracle. For restudy / spacing decisions (SAFE-SCHED / SAFE-FORGET), prefer delayed or cue-only "will you still get this?" prompts over fluency-right-after-feedback. For **miss classification** in gap-scan, retrospective confidence at answer time remains useful (CRI logic) even when it is a weak JOL for long-term retention.

Kill: Instant "I'm sure" after green check $\equiv$ durable learning.

Survive: Separate diagnostic RCJ (classify miss) from delayed JOL (schedule return).

XCVIII.5 Math students mis-monitor — often toward overconfidence

FACT: Erickson & Heit (2015, *Frontiers in Psychology*, 6, 742, doi:10.3389/fpsyg.2015.00742) — high-school and undergrad samples showed **overconfidence** predicting math performance, with greater overconfidence in math than in comparison subjects in their designs; anxiety and overconfidence can co-exist and both impair adaptive monitoring/control.

FACT: Lingel, Lenhart & Schneider (2019, *ZDM*, ERIC EJ1222655) — seventh-graders showed **pervasive overconfidence**; monitoring precision differed by scale format and prospective vs retrospective timing; absolute, relative, and diagnostic accuracy indices are not interchangeable.

FACT: Foster (2021/2022, *IJSME*, doi:10.1007/s10763-021-10207-9) — confidence assessment alone is **non-inferior** on attainment and **not** a proven attainment booster (Part L load-bearing null).

Commercial implication: Gap-scan that only asks "how confident are you in algebra?" without item probes will inherit overconfidence bias and mint false greens (anti-SAFE-COLD). Item/concept probes + hide-correctness + confidence×outcome classification are the wedge; "metacognitive onboarding" copy without that spine is costume.

Kill: Confidence-only placement $\equiv$ diagnosis.

Survive: Monitoring judgments as *inputs* to classification under C4.

XCVIII.6 Certainty of response — misconception vs don't-know

FACT: Hasan, Bagayoko & Kelley (1999, *Physics Education*, 34(5), 294–299, doi:10.1088/0031-9120/34/5/304) — **Certainty of Response Index (CRI)** with MC items distinguishes misconceptions (wrong + high certainty) from lack of knowledge (wrong + low certainty) and from knowledgeable correct responses.

HYPOTHESIS for MindCraft: Gap-scan and FEI probes should treat high-confidence wrongs as **SAFE-MISCON** / **hypercorrection** candidates (Butterfield & Metcalfe lineage — Part L), and low-confidence wrongs as **uninformed** / **scaffold** candidates — never identical mastery deltas (Layer-4 evidence_update_policy).

Wound: CRI was developed in physics MC contexts; HS ACT items and Likert easy/kinda/hard are coarser. Survive the *logic*, not worksheet cosplay.

Kill: Equal graph update for every wrong.

Survive: Confidence $\times$ correctness classes as first-class events.

XCVIII.7 Practice-test feedback does not automatically fix miscalibration

FACT: Persistent miscalibration literature (e.g. introductory biology practice-test feedback studies — Niedjelski & colleagues line summarized in open reports such as PMC8442020 / *CBE—Life Sciences Education* calibration work) shows practice testing can improve average calibration for many students while **low performers often remain overconfident** and high performers can tip underconfident — feedback alone is incomplete.

Product translation: Do not promise “gap-scan calibrates you.” Promise: gap-scan **surfaces** mismatch under destaked conditions, then practice springboards + FEI do the hard work. Parent copy: “we separate sure-and-wrong from unsure,” not “we fix metacognition in one scan.”

Kill: One diagnostic $\Rightarrow$ calibrated learner ads.

Survive: Scan as evidence intake; calibration as longitudinal process (SAFE-CALIB + SAFE-LONGID).

XCVIII.8 Gap-scan product spine (SAFE-MONITOR)

1. **Label the judgment** — Concept ease (seed) $\neq$ item retrospective confidence (probe) $\neq$ delayed JOL (return schedule).
2. **Destake** — C4 hide-correctness during scan; no public shame ranks (SAFE-COMPARE).
3. **Classify** — correct/incorrect $\times$ low/med/high $\rightarrow$ guess / slip / confident misconception / underconfident correct.
4. **Control** — route: hypercorrection / contradicting-reason for high-conf miss; hug-then-dose for low-conf miss; challenge unlock for underconf correct (SAFE-CALIB stack).

- 5. **Weight evidence** — high-conf wrong $\neq$ low-conf wrong in mastery updates (engine).
- 6. **Refuse fireworks** — no day-one greens from ease ratings alone (SAFE-COLD).
- 7. **Measure** — confidence_miss_tier, bias Δ , worstweakness stability, later solo_transfer_pass — never Monitoring Score™.

Marketing language that survives: “Know what you know,” “sure-and-stuck vs guessing,” “honest scan,” “confidence that matches skill.”

Marketing language that dies: “Metacognition Score™,” “we raise metacognition,” “calibrated in one quiz,” “guaranteed ACT from monitoring method.”

XCVIII.9 Competitive wedge (brief)

Khan (explain-first, weak miss classification), Duo (binary correctness + streaks), Brilliant (scaffold dopamine), ChatGPT tutors (fluent **system** certainty $\rightarrow$ student overtrust — Part XXXIII) all under-instrument *student* monitoring $\rightarrow$ control. MindCraft’s L1-safe line: gap-scan is not a vibe check — it is **monitoring that changes what happens next**.

XCVIII.10 Claim ladder

Claim	Max ladder without new data
Monitoring $\neq$ control; judgment types dissociate (Nelson & Narens)	L1
Delayed JOLs improve relative accuracy vs immediate (Nelson & Dunlosky; Rhodes & Tauber)	L1
Math learners often overconfident; indices are plural (Erickson & Heit; Lingel et al.)	L1
CRI/confidence $\times$ accuracy separates misconception vs don’t-know (Hasan et al.)	L1
CA alone does not raise attainment (Foster 2021) — monitoring must wire to control	L1
C4 + confidence $\times$ outcome improves weakness ranking vs confidence-only or reveal-as-you-go	L1 product bet (MONITOR-* / CAL-4)
Monitoring Score™ / one-scan calibration transforms identity	Banned

XCVIII.11 Experiments

ID	Question	Design	Primary
----	----------	--------	---------

MONITOR-1	Item RCJ + outcome classification vs ease-only seed	2-arm	worst weakness stability; next-session accuracy
MONITOR-2	High-conf-miss hypercorrection path vs uniform feedback	2-arm	delayed isomorphic correction; retry_120s
MONITOR-3	Delayed cue-only JOL for return scheduling vs immediate post-feedback “sure?”	2-arm	spacing adherence; delayed probe accuracy (SAFE-SCHED)
MONITOR-4	Hide-correctness + confidence vs reveal-as-you-go (shared CAL-4 / MIS-4)	2-arm	completion; affect; weakness stability
MONITOR-5	Parent CBC: “sure-and-wrong diagnosis” vs “Metacognition Score™” vs “raises confidence”	CBC	WTP; trust (SAFE-WTP)
MONITOR-QUAL	10 Maya: meaning of easy/kinda/hard vs item “how sure?”; shame after high-conf miss	Qual	codebook: threat vs curiosity

Falsifier: Ease-only seed matches item-classified routing on weakness stability → demote probe cost.

Falsifier: Uniform feedback ≥ tiered after high-conf miss → demote MONITOR/CAL tiering.

Falsifier: Parents prefer Metacognition Score™ CBC → rewrite GTM; do not ship score NS.

Pre-register: MONITOR- before metacognition / Monitoring Score™ / one-scan-calibration / guaranteed-ACT-from-monitoring campaigns (SAFE-LABMETA). Do not steal CAL-/COLD-/MIS-/SCHED-* arms — densify them.

XCVIII.12 So what for MindCraft commercially

- **Copy:** “Honest scan — sure-and-stuck vs guessing.” Never Monitoring Score™ or “we raise metacognition.”
- **Product:** Named judgments; C4; confidence×outcome classes; tiered control; no day-one fireworks from ease alone.
- **Positioning:** Against confidence-theater edtech and fluent AI certainty; for monitoring that changes routing and graph weight.
- **Metric:** confidence_miss_tier + weakness stability + FEI — demote Metacognition %.
- **Kill list:** Monitoring Score™; metacognition ads without control wiring; ease-only ≡ calibration science; immediate surety ≡ delayed JOL; equal wrong updates; one-scan calibration; ACT guarantees from monitoring

packaging.

- **Growth:** Parent decks show a 2×2 (correct/wrong × sure/unsure) with routing consequences, not a metacog leaderboard.
 - **Vision:** Maya leaves gap-scan knowing the product heard *how sure she was*, not that she became a “metacognitive person” by finishing a quiz.
-

XCVIII.13 Doctrine — SAFE-MONITOR (provisional)

1. **Monitoring without control is costume** — Nelson & Narens: judgments must change study/routing/updates.
2. **Name the judgment** — EOL-ish ease ≠ RCJ ≠ delayed JOL.
3. **Math overconfidence is expected** — design for bias, do not shame (Erickson & Heit; Lingel).
4. **Certainty sorts miss types** — Hasan CRI logic → SAFE-MISCON vs uninformed paths.
5. **Delay when predicting future recall** — Nelson & Dunlosky; Rhodes & Tauber — not for every UI click.
6. **CA/slider alone ≠ attainment** — Foster null; wire to SAFE-CALIB.
7. **No Monitoring Score™** — prove with classification quality + FEI.
8. **Densifies SAFE-COLD** — humble seed + probes; no fireworks from self-ratings alone.

Confidence: High that Nelson & Narens (1990) establish monitoring/control and dissociable judgment types. High that Nelson & Dunlosky (1991) and Rhodes & Tauber (2011) establish delayed-JOL relative-accuracy benefits. High that Erickson & Heit (2015) and Lingel et al. (2019) document math monitoring/overconfidence patterns and measure plurality. High that Hasan et al. (1999) justify certainty×accuracy misconception sorting. High that Foster (2021) blocks CA-alone attainment claims (Part L). Medium that MindCraft’s specific gap-scan UX will move weakness stability — instrument MONITOR-* before metacognition brand campaigns.

References (verified)

- Erickson, S., & Heit, E. (2015). Metacognition and confidence: Comparing math to other academic subjects. *Frontiers in Psychology*, 6, 742. <https://doi.org/10.3389/fpsyg.2015.00742>
 - Foster, C. (2022). Confidence assessment and achievement in mathematics: A quasi-experimental study. *International Journal of Science and Mathematics Education*, 20, 1411–1429. <https://doi.org/10.1007/s10763-021-10207-9>
 - Hasan, S., Bagayoko, D., & Kelley, E. L. (1999). Misconceptions and the Certainty of Response Index (CRI). *Physics Education*, 34(5), 294–299. <https://doi.org/10.1088/0031-9120/34/5/304>
 - Lingel, K., Lenhart, J., & Schneider, W. (2019). Metacognition in mathematics: Do different metacognitive monitoring measures make a difference? *ZDM*, 51, 587–600. ERIC EJ1222655.
 - Nelson, T. O., & Dunlosky, J. (1991). When people’s judgments of learning (JOLs) are extremely accurate at predicting subsequent recall: The “delayed-JOL effect.” *Psychological Science*, 2(4), 267–270. <https://doi.org/10.1111/j.1467-9280.1991.tb00147.x>
 - Nelson, T. O., & Narens, L. (1990). Metamemory: A theoretical framework and new findings. In G. H. Bower (Ed.), *The psychology of learning and motivation* (Vol. 26, pp. 125–173). Academic Press.
 - Rhodes, M. G., & Tauber, S. K. (2011). The influence of delaying judgments of learning on metacognitive accuracy: A meta-analytic review. *Psychological Bulletin*, 137(5), 791–813. <https://doi.org/10.1037/a0021705>
-

Part XCIX — Opportunity Cost of Tutor Minutes vs AI Scaffold

Chapter status: Living evidence + HITL economics / session-design brief — Researcher tick 2026-08-07 (UTC hour 18 ≡ Red Team slot, but ch99 never written → prefer Researcher per rotation; researcher count since synthesizer v1.12 = 7 → Researcher)

Primary question: When should MindCraft spend scarce *human tutor minutes* versus AI coach / Solver / Map scaffolds — and how do we sell that allocation honestly without “AI replaces tutors,” “unlimited human,” or hours-booked vanity?

Owners: Tutor ops · Product (session brief / coach routing) · Pricing / GTM · Brand · Red Team

Commercial job: Ship a **SAFE-ROI** doctrine that densifies SAFE-HITL / SAFE-TUTORGRAIN / SAFE-WORKFORCE / SAFE-HELP / SAFE-REPAIR: treat tutor time as a high-dosage *scarce input* with measurable opportunity cost; default AI for rehearsal and contingent soft scaffolds; reserve humans for diagnosis handoff, affective destake, bridge articulation, and repair escalation; prove value with FEI per tutor-minute — never Session ROI Score™ theater or “AI = free tutoring.”

Builds on: Parts LXVIII (SAFE-HITL), LXXI (SAFE-TUTORGRAIN), LXXV (SAFE-WORKFORCE), XXXIII (AI trust / Bastani), LX (SAFE-REPAIR), LXXXIX (SAFE-HELP), XCI (SAFE-HINT), XC (SAFE-INSTRUMENT), LXVII (SAFE-ONTOLOGY). Product seams: Map brief before pour, coach cards, Solver fade, human escalate packets, session length / booking UX.

XCIX.1 Why this chapter exists

SAFE-HITL already says humans are operators on the Map, not warmth stickers. SAFE-TUTORGRAIN already says near-peer fidelity beats Ivy pour. Neither answers the commercial knife-fight every parent and district will force:

> “If the AI is so good, why am I paying for tutors?”

> “If tutors matter so much, why not unlimited human chat?”

Wrong answers kill the company twice: (a) **AI-replaces-tutors** cosplay that Bastani-style crutches and hallucination repair already falsify as learning strategy; (b) **human-minutes-as-unlimited entitlement** that bankrupts unit economics and drives pour-relapse under parent pressure (SAFE-WORKFORCE).

MindCraft’s wedge is not cheaper chat or more booked hours. It is **honest allocation**: AI for high-frequency, low-stakes, Map-aware practice loops; humans for scarce moments where contingent prompting + diagnosis handoff + accountability change FEI — then fade back to solo.

FOUNDER BELIEF under audit: Tutor minutes are the dearest inventory in the product. Spending them on re-explaining what the Map already named, or on answer-dumping that Solver could gate, is the real product failure — not “not enough AI.”

Claims we refuse as doctrine:

1. AI replaces tutors / “tutoring is free now.”
2. Unlimited human minutes as brand promise or default SKU.

3. Hours booked / talk time / session count as North Star.
4. Session ROI Score™ as gamified ops KPI.
5. VanLehn ITS≈human as license for unguarded ChatGPT “tutors.”
6. Bastani GPT Tutor safeguards ≡ FEI without solo transfer.
7. High-dosage school RCT ES as guaranteed ACT points from thin marketplace sessions.
8. Warmth, Ivy résumé, or “always-on AI” as substitutes for allocation honesty.

XCIX.2 Constructs (keep mechanisms distinct)

Construct	Research meaning	MindCraft analogue	Not the same as
High-dosage tutoring	Frequent, structured 1:1/small-group instruction	Near-peer sessions + Map brief	Occasional homework chat
Opportunity cost	Next-best use of scarce tutor minutes	FEI events per tutor-minute; deferred students	Raw cost/hour only
AI scaffold	Step/hint/coach without human	Solver fade, soft hints, Map cards	Unguarded answer dump
Guardrailed generative tutor	Pedagogy wrap that blocks crutches	SAFE-HINT / SAFE-FADE / SAFE-HELP stack	GPT Base / ChatGPT paste
Complementarity	Human + system beat either alone when roles differ	HITL on Map; AI for rehearsal	Human reading AI script
SAFE-ROI	Allocation doctrine for minutes vs AI	This chapter	Efficiency theater score

Operational definition (HYPOTHESIS): A MindCraft session stack is *ROI-honest* when (a) pre-session Map brief names the 1–3 human-worthy moves (bridge gap, soft-wrong prompt, exposure destake, repair escalate), (b) AI owns rehearsal / soft peeks / fade steps outside that list, (c) tutor talk favors construction prompts over re-lecture (SAFE-TALK), (d) marketing never claims AI≡tutor or unlimited humans, and (e) success is *solo_transfer_pass* and related FEI **per paid tutor-minute**, not hours booked or tokens generated.

XCIX.3 Tutoring works — and dosage / labor model set the cost curve

FACT: Nickow, Oreopoulos & Quan (2020, NBER WP 27476) — experimental PreK–12 tutoring meta-analysis: pooled ~**0.37 SD**; teacher/paraprofessional programs tend to beat nonprofessional/parent models; in-school tends to beat after-school.

FACT: Guryan et al. (2023, *AER*, 113(3), 738–765, doi:10.1257/aer.20210434) — two Chicago RCTs of Saga-style **high-dosage** secondary math tutoring with paraprofessionals: about **0.18–0.40 SD** math gains, course-grade benefits, later-year persistence; cost held via paras (~\$3.2–4.8k/student-year in related Saga/NBER WP 28531 writeups), not

daily certified teachers.

FACT: Kraft & Falken (2021, *AERA Open*, doi:10.1177/23328584211042858) — individualized tutoring is powerful but **supply- and budget-constrained**; targeted national K–8 Title I scale scenarios are **billions/year** — software does not make human minutes free.

Wound: School-embedded high-frequency programs $\neq$ remote ACT marketplace 1–2 $\times$ /week. Borrow dosage + structure + labor-arbitrage logic — not pooled ES as score insurance.

Kill: “Our tutors are 0.37 σ .” Kill “AI makes high-dosage free.”

Survive: Humans work when dosage/structure are real — so do not waste the minutes you can afford.

XCIX.4 Cheap / thin human time is not automatic FEI

FACT: Kraft, List, Livingston & Sadoff (2022, *AEA Papers & Proceedings*, 112, 614–618, doi:10.1257/pandp.20221038) — randomized pilot of **online college-volunteer** tutoring for middle schoolers: ITT estimates about **+0.07 SD math** and **+0.04 SD reading**, positive but **not statistically significant**; authors emphasize much **lower cost** than high-dosage in-person models and a **dosage** interpretation (more hours $\rightarrow$ larger effects) under COVID constraints.

Product translation: Near-peer online supply can be cost-efficient and still under-deliver if dosage, structure, and Map fidelity are thin. SAFE-TUTORGRAIN + SAFE-HITL are not optional garnish on “we have college tutors.” Opportunity cost cuts both ways: expensive pour sessions *and* cheap unstructured chats can both waste the learning budget.

Kill: Marketplace volume / volunteer availability $\equiv$ high-impact tutoring.

Survive: Fidelity $\times$ dosage $\times$ Map brief as the human product.

XCIX.5 ITS parity is not ChatGPT license

FACT: VanLehn (2011, *Educational Psychologist*, 46(4), 197–221, doi:10.1080/00461520.2011.611369) — review finds adult human tutoring effect vs no tutoring around **$d \approx 0.79$** and step-based intelligent tutoring systems around **$d \approx 0.76$** — roughly **comparable** in that literature, undermining Bloom-style 2σ folklore and crude “humans always crush software” claims. Interaction appears to plateau around step-based granularity.

Wound: VanLehn’s ITS are **constrained pedagogical systems**, not unconstrained LLM chat. Treating ChatGPT paste as “ITS-proven” is cite-wash.

HYPOTHESIS: MindCraft AI value is closest to VanLehn when it is **step/hint/Map-constrained** (SAFE-HINT / SAFE-FADE / SAFE-ONTOLOGY). Open generative dump is a different object — often a crutch.

Kill: “Science says AI = human tutor” as ChatGPT marketing.

Survive: Constrained scaffolds can carry large shares of practice load; humans still own accountability, diagnosis negotiation, and repair.

XCIX.6 Unguarded generative AI can buy practice scores and sell learning debt

FACT: Bastani, Bastani, Sungu, Ge, Kabakc■ & Mariman (2025, *PNAS*, doi:10.1073/pnas.2422633122; earlier SSRN “Generative AI Can Harm Learning”) — ~1,000 HS math students; GPT Base and GPT Tutor both **raise practice performance** sharply; when AI is removed, **GPT Base students do ~17% worse** than never-AI controls; GPT Tutor safeguards largely **mitigate the harm** but do not automatically produce large positive exam gains vs control. Students often use Base as a **solution crutch**; many do not perceive the learning loss.

Commercial implication: Selling “AI tutor that gets homework done” can raise short-run satisfaction while destroying solo transfer — the opposite of FEI / SAFE-PROOF. Human minutes spent *checking copied AI work* are negative ROI. Human minutes spent *enforcing attempt→soft scaffold→solo* are high ROI.

Kill: Practice-with-AI accuracy $\equiv$ learning. Kill unlimited-AI-homework hero.

Survive: Guardrails + solo gates; humans as anti-crutch operators when AI fails or when dependency forms (SAFE-HELP / SAFE-REPAIR).

XCIX.7 When human minutes beat AI (allocation spine)

HYPOTHESIS — spend human minutes when ≥ 1 holds: Map diagnosis handoff (bridge/format/high-conf miss); affective destake / graded exposure (not therapy cosplay); construction prompts on soft-wrong (SAFE-TALK); repair escalate after AI fail (SAFE-REPAIR); witness solo transfer (SAFE-PROOF); fade / “your turn” gate (SAFE-FADE).

HYPOTHESIS — default AI / async Practice when: Spaced rehearsal on named weaknesses; soft contingent peeks after attempt; faded completion examples; FormatId conversion drills; parent MoC without live stalk.

Kill: Universal “human for every emotion” or “AI for every stuck.” **Survive:** Pre-session **minute budget** tied to Map tasks.

XCIX.8 Unit economics without cruelty

FOUNDER BELIEF: Parents buy a bundle — trusted diagnosis + practice that doesn’t fake mastery + humans for hard joins. Pure-AI underprices scarce input; pure-human overprices pour time.

SPECULATION: SKU ladder = Map/Practice/guardrailed coach → scheduled near-peer FEI sessions → expert escalate packets — never unlimited everything. SAFE-WTP CBC must test “tutors for joins; AI for reps” vs unlimited-AI vs unlimited-human.

Kill: Hours-booked discounting as the only lever. **Survive:** Outcome-linked session design + honest scarcity.

XCIX.9 SAFE-ROI product spine

1. **Minute budget on the brief** — Top 1–3 human moves named before start.
2. **AI owns the rest** — rehearsal / soft peeks / fade outside the budget.
3. **Anti-crutch gate** — no human or AI bottom-out before attempt (SAFE-HELP × Bastani).
4. **Repair path** — escalate packet, not infinite free human (SAFE-REPAIR).
5. **Instrument** — FEI per tutor-minute; pour-relapse; escalate precision; solo after AI.
6. **Refuse vanity** — no hours NS; no Session ROI Score™ that punishes destake time.
7. **Copy** — “Humans for the hard joins. AI for the reps.” Never tutoring-is-free / AI-replaces-tutors.

Survives: “Tutor time on what the Map says matters,” “reps without crutches,” “solo after the scaffold.”

Dies: “Unlimited AI tutor,” “unlimited human help,” “AI = private tutor,” “ 0.37σ guaranteed,” “hours = progress.”

XCIX.10 Competitive wedge (brief)

Khan/Brilliant: content + solo scaffolds (weak human dosage). Duo: streaks (weak FEI). ChatGPT tutors: fluent availability (Bastani crutch; SAFE-REPAIR debt). Marketplace experts: résumé + booked hours. MindCraft L1 line: **allocate scarce humans with Map honesty; constrained AI for volume without learning debt.**

XCIX.11 Claim ladder

Claim	Max ladder without new data
Experimental tutoring often large (~0.37 SD pooled; Nickow et al.)	L1
High-dosage para tutor models can move secondary math (Guryan et al. 2023)	L1
Scale is budget/supply constrained (Kraft & Falken 2021)	L1
Thin online volunteer dosage can be underpowered (Kraft et al. 2022)	L1
Step-based ITS $\approx$ human in VanLehn (2011) review — not LLM chat	L1
Unguarded GPT can harm post-AI exam performance (Bastani et al. 2025)	L1
Map-briefed minute budgets raise FEI per tutor-minute vs unstructured pour	L1 product bet (ROI-*)
Session ROI Score™ / AI-replaces-tutors transforms learning	Banned

XCIX.12 Experiments

ID	Question	Design	Primary
ROI-1	Map-briefed minute budget vs unstructured “help with homework” session	2-arm tutors	FEI per tutor-minute; pour-relapse
ROI-2	AI-first rehearsal + human join vs human-first explain-all vs AI-only	3-arm	solo_transfer_pass; help-abuse rate
ROI-3	Bastani-style hard-dump blocked vs allowed in Solver during practice week	2-arm	delayed no-AI probe (shared HELP/HINT)
ROI-4	Escalate-to-human packet precision vs open “book more hours” after AI fail	2-arm	repair success; cost per resolved gap
ROI-5	Parent CBC: “humans for joins / AI for reps” vs “unlimited AI tutor” vs “unlimited human chat”	CBC	WTP; trust (SAFE-WTP)
ROI-QUAL	10 Maya + 10 parents: when they want a human vs AI; shame if human “only” 20 focused minutes	Qual	codebook: entitlement vs relief

Falsifier: Unstructured human sessions $\geq$ Map-budgeted on FEI/minute $\rightarrow$ demote brief ritual cost.

Falsifier: AI-only $\geq$ AI+human-join on solo transfer for diagnosed bridge gaps $\rightarrow$ narrow human SKU (do not overclaim).

Falsifier: Parents prefer unlimited-AI CBC and show equal retention $\rightarrow$ rewrite GTM; still do not ship crutch UX.

Pre-register: ROI-* before “AI replaces tutors,” “unlimited human,” hours-as-NS, or Session ROI Score™ campaigns (SAFE-LABMETA). Do not steal HITL/HELP/HINT/REPAIR/TALK arms — densify them.

XCIX.13 So what for MindCraft commercially

- **Copy:** “Humans for the hard joins. AI for the reps.” Never tutoring-is-free or AI=private-tutor.
- **Product:** Map minute budget; AI for rehearsal; human for diagnosis/destake/construction/repair/witness.
- **Positioning:** Against ChatGPT crutches and hours vanity; for scarce-human + constrained-AI complementarity.
- **Metric:** FEI per tutor-minute; escalate precision; solo-after-AI — demote hours / Session ROI Score™.

- **Kill list:** AI replaces tutors; unlimited human; hours NS; VanLehn-as-ChatGPT; Bastani practice=learning; 0.37σ ACT guarantee.
- **Growth:** Parent decks show 3 human moves + AI reps plan — not a timesheet.
- **Vision:** Maya experiences aimed human minutes, then owns AI reps without crutches.

XCIX.14 Doctrine — SAFE-ROI (provisional)

1. **Tutor minutes are scarce inventory** — Nickow/Guryan/Kraft: tutoring works and costs real labor.
2. **Thin human time ≠ automatic impact** — Kraft et al. 2022 dosage wound.
3. **Constrained AI ≠ ChatGPT dump** — VanLehn ITS parity does not license crutches.
4. **Unguarded generative help can create learning debt** — Bastani et al. 2025.
5. **Allocate by Map task** — humans for joins/destake/repair/witness; AI for reps/fade/soft peeks.
6. **Measure FEI per tutor-minute** — not hours booked.
7. **No Session ROI Score™ / AI-replaces-tutors / unlimited-human brand.**
8. **Densifies SAFE-HITL × SAFE-HELP × SAFE-REPAIR** — complementarity with honesty.

Confidence: High that Nickow et al. (2020) and Guryan et al. (2023) establish large experimental tutoring impacts under structured dosage. High that Kraft & Falken (2021) establish scale/cost constraints. High that Kraft et al. (2022) show thin online volunteer models can be underpowered. High that VanLehn (2011) places step-based ITS near human averages without rescuing LLM dumps. High that Bastani et al. (2025) show unguarded GPT can harm post-removal performance. Medium that MindCraft’s specific minute-budget UX raises FEI/minute — instrument ROI-* before AI-replaces-tutors or unlimited-human campaigns.

References (verified)

- Bastani, H., Bastani, O., Sungu, A., Ge, H., Kabakc■, Ö., & Mariman, R. (2025). Generative AI without guardrails can harm learning: Evidence from high school mathematics. *Proceedings of the National Academy of Sciences*, 122. <https://doi.org/10.1073/pnas.2422633122>
 - Guryan, J., Ludwig, J., Bhatt, M. P., Cook, P. J., Davis, J. M. V., Dodge, K., Farkas, G., Fryer, R. G., Jr., Mayer, S., Pollack, H., Steinberg, L., & Stoddard, G. (2023). Not too late: Improving academic outcomes among adolescents. *American Economic Review*, 113(3), 738–765. <https://doi.org/10.1257/aer.20210434>
 - Kraft, M. A., & Falken, G. T. (2021). A blueprint for scaling tutoring and mentoring across public schools. *AERA Open*, 7, 1–21. <https://doi.org/10.1177/23328584211042858>
 - Kraft, M. A., List, J. A., Livingston, J. A., & Sadoff, S. (2022). Online tutoring by college volunteers: Experimental evidence from a pilot program. *AEA Papers & Proceedings*, 112, 614–618. <https://doi.org/10.1257/pandp.20221038>
 - Nickow, A., Oreopoulos, P., & Quan, V. (2020). The impressive effects of tutoring on PreK–12 learning: A systematic review and meta-analysis of the experimental evidence (NBER Working Paper 27476). <https://doi.org/10.3386/w27476>
 - VanLehn, K. (2011). The relative effectiveness of human tutoring, intelligent tutoring systems, and other tutoring systems. *Educational Psychologist*, 46(4), 197–221. <https://doi.org/10.1080/00461520.2011.611369>
-

Part C — Worked-Example vs Problem-Solving Timing Across Expertise

Chapter status: Living evidence + Practice/Solver schedule brief — Researcher tick 2026-08-08 (UTC hour 0 $\equiv$ Red Team slot, but ch100 never written $\rightarrow$ prefer Researcher per rotation; researcher count since synthesizer v1.13 = 0 $\rightarrow$ Researcher)

Primary question: When should MindCraft schedule full worked examples vs completion vs conventional problem solving — especially when a fixed E0 $\rightarrow$ E3 fade ladder collides with *prior knowledge that is already partial, format-specific, or misdiagnosed?*

Owners: Product (Practice / Solver fade machine) · Engine (Map grain / level / format) · HITL tutors · Brand · Red Team

Commercial job: Ship **SAFE-EXPTIME** densifying SAFE-FADE / SAFE-COLD / SAFE-FORMAT / SAFE-PF / SAFE-HINT: timing is diagnosed, not calendar cosplay; fixed ladders yield to rapid expertise checks; adaptive > fixed when grain is real — never Timing Score™, universal E0 $\rightarrow$ E3 destiny, or “struggle-first always builds grit.”

Builds on: Parts XXVI (CLT), LXXXVIII (SAFE-FADE), LXXXI (SAFE-COLD), LXXXII (SAFE-FORMAT), XCIV (SAFE-PF), XCI (SAFE-HINT), XCV (SAFE-APPRENTICE), XCIII (SAFE-RETRIEVE). Product seams: fade_stage E0–E3, gap-scan / Map eventCount, FormatId routing, completion blanks, solo transfer proof.

C.1 Why this chapter exists

SAFE-FADE already killed always-full-worked / never-fade / view $\equiv$ mastery. What remains is the *schedule collision*: gap-scan “hard” $\rightarrow$ E0 full worked while Maya already half-knows the procedure but blanks on diagram FormatId $\rightarrow$ expertise-reversal tax + peek binge; or Map “kinda” $\rightarrow$ early E3 $\rightarrow$ freeze / help-abuse.

Competitors sell always-every-step (ChatGPT) or always-struggle-first grit theater. MindCraft’s wedge is **when** — examples when schemas are thin; solving when they exist; completion as bridge; re-time when concept $\neq$ format $\neq$ bridge expertise.

FOUNDER BELIEF under audit: A fixed E0 $\rightarrow$ E3 ladder without prior-knowledge checks is expertise-blind pedagogy with a curriculum skin. Honesty is “help matches *this* grain,” not “we have a science fade sequence.”

Claims we refuse as doctrine:

1. One universal E0 $\rightarrow$ E3 schedule for all concepts/formats/students.
 2. Always-example-first or always-problem-first as brand method.
 3. Fixed fade $\equiv$ adaptive fade (calendar costume without checks).
 4. Gap-scan confidence alone $\equiv$ fade stage (confidence $\neq$ schema; SAFE-CALIB / SAFE-MONITOR).
 5. Example Timing Score™ / %worked as North Star.
 6. Expertise reversal as license for zero guidance day one.
 7. “We adapt perfectly” black-box AI timing hero.
 8. Problem-solving first $\equiv$ PF without consolidation gates (SAFE-PF).
-

C.2 Constructs (keep mechanisms distinct)

Construct	Research meaning	MindCraft analogue	Not the same as
Worked-example effect	Studying full solutions beats unguided search for novices	E0 full coach / Solver study	Permanent dump
Expertise reversal	Novice-helpful guidance hurts / wastes experts	Same E0 after demonstrated grain	“They got bored” only
Completion / fading	Successive blanks until full solve	E1–E2 missing steps	Fake blanks + instant peek
Fixed fade	Same omit schedule for all	Default E0→E3 by level	Diagnosed timing
Adaptive fade	Fade point tracks measured understanding	Map×FormatId×recent solo	Black-box “AI knows”
SAFE-EXPTIME	Diagnose prior grain → choose example/completion/solve; recheck	This chapter	Timing Score™

Operational definition (HYPOTHESIS): A path is *EXPTIME-honest* when it (a) sets initial fade from **grain-local evidence** (concept×format×bridge), (b) **rechecks** via completion/SE/solo to move E up/down, (c) prefers **adaptive over fixed** when instrumented, (d) never equates gap-scan Likert with schema readiness, and (e) proves with solo_transfer_pass / delayed mixed accuracy — not %worked or explanation NPS.

C.3 Novices: worked examples beat unguided problem solving

FACT: Sweller & Cooper (1985, *Cognition and Instruction*, 2(1), 59–89, doi:10.1207/s1532690xc0201_3) — algebra novices who studied worked examples as a substitute for conventional problem solving showed reduced time and errors on subsequent similar problems; conventional search imposed heavy load that retarded schema acquisition.

FACT: Sweller (1988, *Cognitive Science*, 12(2), 257–285, doi:10.1207/s15516709cog1202_4) — problem-solving search can consume working memory needed for schema construction.

Applied (HYPOTHESIS): Day-one E3-only (“just try”) for true novices on a new grain violates CLT. SAFE-EXPTIME inherits SAFE-FADE’s stage-not-struggle-theater rule: start complete enough to build schema.

Kill: Problem-solving-first as universal brand for acquisition.

Survive: Full worked study as *early-stage* tool on thin grain.

C.4 Experts: problem solving can beat more examples

FACT: Kalyuga, Chandler, Tuovinen & Sweller (2001, *Journal of Educational Psychology*, 93(3), 579–588, doi:10.1037/0022-0663.93.3.579) — as learners gain domain expertise, problem solving can become superior to studying worked examples (classic reverse pattern under CLT).

FACT: Kalyuga, Ayres, Chandler & Sweller (2003, *Educational Psychologist*, 38(1), 23–31, doi:10.1207/s15326985ep3801_4) — methods highly effective for inexperienced learners can lose effectiveness or become detrimental as expertise grows because guidance becomes redundant processing.

FACT (review): Kalyuga (2007, *Educational Psychology Review*, 19(4), 509–539, doi:10.1007/s10648-007-9054-3) — expertise reversal is a robust ATI-style pattern within CLT; learner-tailored systems must match support to knowledge state, not average syllabus week.

Commercial implication: Marketing “we walk every step with you” past partial mastery is ChatGPT costume with a curriculum skin. Map must *lower* default guidance after demonstrated competence on that grain.

Kill: Expertise-blind always-example brand.

Survive: Problem solving / sparse completion when grain evidence shows schemas exist.

C.5 Transition tool: fading / completion, not cliff jumps

FACT: Renkl & Atkinson (2003, *Educational Psychologist*, 38(1), 15–22, doi:10.1207/S15326985EP3801_3) — theoretical CLT account: intrinsic load falls across skill acquisition so problem-solving elements can be *gradually* integrated into example study; propose fading from complete examples → incomplete → independent solving.

FACT: Renkl, Atkinson, Maier & Staley (2002, *Journal of Experimental Education*, 70(4), 293–315, doi:10.1080/00220970209599510) — field + lab: successive fading of worked steps outperformed traditional example–problem pairs on near transfer; mediation via fewer errors during learning under fading.

FACT: Renkl, Atkinson & Große (2004, *Instructional Science*, 32(1–2), 59–82, doi:10.1023/B:TRUC.0000021815.74806.f6) — fading associated with fewer *unproductive* learning events; position of faded steps mattered less than that fading triggered principle-focused processing on faded content.

Applied (HYPOTHESIS): MindCraft E1–E2 are not “half a ChatGPT dump.” They are **completion problems** with named missing grain (procedure step, bridge join, format conversion). Cliff jumps E0→E3 without completion are the UX that creates freeze or peek binge.

Kill: Binary help-on / help-off without completion bridge.

Survive: Backward-leaning completion ladder as default bridge (SAFE-FADE), re-timed per C.6–C.7.

C.6 Fixed fade is not enough — adaptive beats fixed in tutored settings

FACT: Salden, Aleven, Schwonke & Renkl (2010, *Instructional Science*, 38(3), 289–307, doi:10.1007/s11251-009-9107-8) — lab + classroom Geometry Cognitive Tutor studies: example-enhanced tutors with **adaptive** fading of worked steps outperformed **fixed** fading, which outperformed standard tutored problem solving alone on learning/retention outcomes discussed in the paper.

Wound: Cognitive Tutor grain $\neq$ MindCraft Map grain — transfer the *principle* (fade to understanding), not CMU costume ads. Shipping a fixed E0→E3 animation labeled “adaptive” is cite-wash.

Kill: Fixed ladder $\equiv$ adaptive science.

Survive: Default ladder + instruments that move stage from evidence.

C.7 Rapid expertise checks — how to time without a Timing Score™

FACT: Kalyuga & Sweller (2004, *Journal of Educational Psychology*, 96(3), 558–568, doi:10.1037/0022-0663.96.3.558) — schema-based rapid knowledge measures (algebra/geometry) correlate highly with traditional tests at much shorter time; used to select instructional design matched to expertise (expertise-reversal mitigation).

FACT: Kalyuga & Sweller (2005, *Educational Technology Research and Development*, 53(3), 83–93, doi:10.1007/BF02504800) — learner-adapted instruction using rapid tests + cognitive-load indices yielded higher knowledge and cognitive-efficiency gains than a yoked non-adapted control in an elementary algebra tutor.

Applied (HYPOTHESIS): MindCraft already has cheap proxies — completion success, soft-wrong class, FormatId miss pattern, delayed solo, SE before wrap — that can play the *role* of rapid assessment without inventing Example Timing Score™. Gap-scan confidence is a *seed prior*, not a fade oracle (SAFE-COLD / SAFE-MONITOR).

Kill: One Likert → permanent E stage. Kill black-box “perfect interval of fading.”

Survive: Inspectable rules: e.g., two consecutive completion successes on grain → fade up; freeze/TOT with no progress → soft bound then temporary fade down (SAFE-RETRIEVE), not instant hard dump.

C.8 Collision cases MindCraft will actually hit

HYPOTHESIS — prior knowledge $\neq$ zero on concept, zero on format: Student solves symbolic `linear_equations` but fails `coordinate_graph` isomorphic items. Correct timing: E0–E1 on *format conversion* grain, not re-teaching full algebra example dump (SAFE-FORMAT $\times$ SAFE-DUAL).

HYPOTHESIS — endpoints OK, bridge weak: Concept nodes look “easy”; join fails. Timing: co-present supports + completion on the *relation*, not more isolated worked examples (SAFE-BRIDGE).

HYPOTHESIS — overconfident cold start: High confidence, low probe accuracy. Timing: hide-correctness densifies evidence; do not gift E3 because they said “easy” (SAFE-CALIB / SAFE-MONITOR / SAFE-COLD).

HYPOTHESIS — exam dual rail: Late ACT prove week may need more problem solving / mixed retrieval than worked study — without relabeling cram as mastery (SAFE-EXAM / SAFE-DURABLE / SAFE-SCHED).

C.9 Product rules (MindCraft)

1. **Initial stage from grain-local evidence** — concept×format×bridge, not global belt.
2. **Default ladder exists** — E0 study → E1–E2 completion → E3 solo — but is a *prior*, not destiny.
3. **Adaptive overrides fixed** — success fades up; diagnosed freeze/miss-class can fade down one step, not flood.
4. **Completion is the commercial middle** — blank named join beats monologue.
5. **Attempt grain before reveal** — SE / soft-wrong before unlocking fuller worked (SAFE-FADE × SAFE-HINT).
6. **Proof = solo transfer / delayed mixed** — never %time-on-examples or explanation NPS.
7. **Tutor ops:** HITL brief names current fade stage + why; humans do not re-pour E0 after Map says E2 (SAFE-APPRENTICE / SAFE-ROI).
8. **Copy:** “Help that steps down as you step up” — not “always every step” or “always struggle first.”

C.10 Confidence table

Claim	Label	Confidence
Worked examples beat unguided PS for novices (Sweller & Cooper 1985; Sweller 1988)	FACT	High
Problem solving can beat examples as expertise rises (Kalyuga et al. 2001; 2003)	FACT	High
Fading/completion bridges example→solve better than cliff or static pairs (Renkl & Atkinson line)	FACT	High
Adaptive fade > fixed fade > PS-only in Cognitive Tutor studies (Salden et al. 2010)	FACT	High (wound: tutor context)
Rapid dynamic assessment can improve efficiency of adapted instruction (Kalyuga & Sweller 2004/2005)	FACT	High
MindCraft Map×FormatId proxies can substitute for lab rapid tests	HYPOTHESIS	Medium — EXPTIME-*
Fixed E0→E3 without recheck harms partial-knowledge Maya	HYPOTHESIS	Medium–High product bet
Example Timing Score™ / always-struggle-first improves FEI	Banned	—

C.11 Experiments

ID	Question	Design	Primary
EXPTIME-1	Fixed E0→E3 vs adaptive fade (completion success + miss-class rules) on same concept	2-arm	solo_transfer_pass; time-to-solo; peek binge
EXPTIME-2	Concept-matched vs concept×format grain for initial stage	2-arm	format conversion accuracy; redundant-example rate
EXPTIME-3	Gap-scan confidence→E stage vs probe-evidence→E stage	2-arm	calibration of stage; early quit; FEI
EXPTIME-4	Adaptive fade vs always-E0 worked vs always-E3 solo on mixed prior-knowledge cohort	3-arm	delayed mixed accuracy; mental effort proxy
EXPTIME-5	Parent CBC: “steps down as you step up” vs “always every step” vs “struggle-first grit”	CBC	WTP; trust (SAFE-WTP)
EXPTIME-QUAL	10 Maya: when full examples felt helpful vs insulting; when solo felt fair vs flooding	Qual	codebook: reversal pain / freeze

Falsifier: Fixed ladder $\geq$ adaptive on solo transfer across prior-knowledge strata → keep fixed as ops default; demote adaptive cost.

Falsifier: Always-E0 $\geq$ adaptive for partial-knowledge students on delayed tests → wound expertise-reversal product claim for our bank.

Falsifier: Parents prefer always-every-step CBC with equal retention → rewrite copy; still do not ship permanent dump.

Pre-register: EXPTIME-* before “perfect adaptive fade AI,” struggle-first, or Timing Score™ campaigns. Densify FADE/HINT/COLD/FORMAT — do not fork a second fade machine.

C.12 So what for MindCraft commercially

- **Copy:** “Help that steps down as you step up — on *this* skill, not a fake level.” Never always-every-step or grit-first acquisition.
- **Product:** Grain-local initial E stage; completion bridge; adaptive up/down from evidence; inspectable why.
- **Positioning:** Against ChatGPT permanent worked solutions *and* against struggle-theater apps that skip schema building.

- **Metric:** Stage appropriateness (miss-class × stage), time-to-honest-solo, solo_transfer_pass — demote %worked and Timing Score™.
 - **Kill list:** Universal E0→E3 destiny; fixed≡adaptive; confidence≡stage; always-example / always-PS brand; black-box perfect fade AI.
 - **Growth / vision:** Parent demo shows the same problem at E0 vs E2 for different Map states — Maya feels guided when thin and trusted when ready; identity is the handoff, not dump or flood.
-

C.13 Doctrine — SAFE-EXPTIME (provisional)

1. **Timing tracks expertise grain** — concept×format×bridge, not global belt.
2. **Examples early, solving late — on that grain** — Sweller/Kalyuga line.
3. **Completion fading is the bridge** — Renkl/Atkinson; not cliff jumps.
4. **Adaptive > fixed when evidence exists** — Salden et al. 2010 principle.
5. **Rapid rechecks beat Likert destiny** — Kalyuga & Sweller rapid assessment role.
6. **Densifies SAFE-FADE** — schedule collision chapter, not a second chrome.
7. **Proof = solo / delayed mixed** — not explanation time share.
8. **No Example Timing Score™ / always-struggle-first / black-box perfect-fade hero.**

Confidence: High on worked-example effect, expertise reversal, fading, and adaptive>fixed in published tutor studies. Medium that Map×FormatId rules will match Salden-class gains — instrument EXPTIME-* before adaptive-fade marketing.

References (verified)

- Kalyuga, S. (2007). Expertise reversal effect and its implications for learner-tailored instruction. *Educational Psychology Review*, 19(4), 509–539. <https://doi.org/10.1007/s10648-007-9054-3>
- Kalyuga, S., Ayres, P., Chandler, P., & Sweller, J. (2003). The expertise reversal effect. *Educational Psychologist*, 38(1), 23–31. https://doi.org/10.1207/s15326985ep3801_4
- Kalyuga, S., Chandler, P., Tuovinen, J., & Sweller, J. (2001). When problem solving is superior to studying worked examples. *Journal of Educational Psychology*, 93(3), 579–588. <https://doi.org/10.1037/0022-0663.93.3.579>
- Kalyuga, S., & Sweller, J. (2004). Measuring knowledge to optimize cognitive load factors during instruction. *Journal of Educational Psychology*, 96(3), 558–568. <https://doi.org/10.1037/0022-0663.96.3.558>
- Kalyuga, S., & Sweller, J. (2005). Rapid dynamic assessment of expertise to improve the efficiency of adaptive e-learning. *Educational Technology Research and Development*, 53(3), 83–93. <https://doi.org/10.1007/BF02504800>
- Renkl, A., & Atkinson, R. K. (2003). Structuring the transition from example study to problem solving in cognitive skill acquisition: A cognitive load perspective. *Educational Psychologist*, 38(1), 15–22. https://doi.org/10.1207/S15326985EP3801_3
- Renkl, A., Atkinson, R. K., & Große, C. S. (2004). How fading worked solution steps works — a cognitive load perspective. *Instructional Science*, 32(1–2), 59–82. <https://doi.org/10.1023/B:TRUC.0000021815.74806.f6>
- Renkl, A., Atkinson, R. K., Maier, U. H., & Staley, R. (2002). From example study to problem solving: Smooth transitions help learning. *Journal of Experimental Education*, 70(4), 293–315. <https://doi.org/10.1080/00220970209599510>

- Salden, R. J. C. M., Aleven, V., Schwonke, R., & Renkl, A. (2010). The expertise reversal effect and worked examples in tutored problem solving. *Instructional Science*, 38(3), 289–307. <https://doi.org/10.1007/s11251-009-9107-8>
 - Sweller, J. (1988). Cognitive load during problem solving: Effects on learning. *Cognitive Science*, 12(2), 257–285. https://doi.org/10.1207/s15516709cog1202_4
 - Sweller, J., & Cooper, G. A. (1985). The use of worked examples as a substitute for problem solving in learning algebra. *Cognition and Instruction*, 2(1), 59–89. https://doi.org/10.1207/s1532690xci0201_3
-

Part CI — Self-Regulated Learning Cycles (Zimmerman) in Practice UX

Chapter status: Living evidence + Practice session UX brief — Researcher tick 2026-08-08 (UTC hour 3; hour%6≠0; researcher count since synthesizer v1.13 = 1 → Researcher)

Primary question: How should MindCraft embed Zimmerman’s cyclical SRL phases — forethought, performance, self-reflection — into Practice without minting an SRL Score™, reflection-journal theater, or streak-as-self-regulation costume?

Owners: Product (Practice / Notes / Map handoffs) · HITL tutors · Brand · Red Team

Commercial job: Ship **SAFE-SRL** densifying SAFE-MONITOR / SAFE-CALIB / SAFE-ATTN / SAFE-HELP / SAFE-APPRENTICE / SAFE-AAR: phase-timed micro-prompts that change control (goal grain, strategy choice, attribution→next plan), prove value with FEI + strategy-adaptive retries — never “we teach metacognition” splash without wiring.

Builds on: Parts XXV (self-efficacy), XXX (attribution), L (SAFE-CALIB), LIV (SAFE-AAR), XCV (SAFE-APPRENTICE), XCVI (SAFE-ATTN), XCVIII (SAFE-MONITOR), XC (FEI instrumentation). Product seams: mission launch, attempt windows, soft-wrong, AAR/annotation, Map return CTAs.

CI.1 Why this chapter exists

Competitors sell either autopilot (streak → next item) or post-hoc journaling (“reflect on your learning”). Both skip the load-bearing claim in social-cognitive SRL: regulation is a **cycle** that must fire *around* a real task — goals and plans *before* effort, monitoring and control *during*, judgment and adaptive inference *after* — and the after must rewrite the before.

MindCraft already has pieces (gap-scan ease, soft-wrong surety, film/AAR cards, Map). What is missing is doctrine that forbids fusing them into one vanity SRL meter while requiring **phase→control** wiring so Maya’s next attempt actually changes.

FOUNDER BELIEF under audit: Identity transformation needs students who can run a thin cycle without a tutor present — but the product earns that claim by scaffolding *events*, not by advertising Self-Regulation Scores.

Claims we refuse as doctrine:

1. SRL Score™ / Metacognition % / Reflection streak as North Star.

2. “We teach self-regulated learning” ads without phase-timed control actions.
3. Global SRL surveys $\equiv$ event regulation on *this* problem.
4. Forethought theater (goal picker) that never constrains strategy or difficulty.
5. Endless end-of-session journals as proof of reflection.
6. Streak / DAU $\equiv$ self-regulation.
7. Ability attributions (“you’re not a math person”) as brand voice; defensive withdrawal as grit.
8. Guaranteed ACT points from “SRL-powered Practice.”

CI.2 Constructs (keep mechanisms distinct)

Construct	Research meaning	MindCraft analogue	Not the same as
Forethought	Task analysis, goals, strategic planning, motivational beliefs before effort	Mission brief: grain goal + plan chip + efficacy check	Vague “get better at math”
Performance	Self-control + self-observation during the task	Attempt window; soft hints; monitoring prompts mid-stall	Mid-attempt DAU push
Self-reflection	Self-judgment (eval, attribution) + self-reaction (adaptive vs defensive)	Soft-wrong class + one adaptive inference $\rightarrow$ next plan	Shame reel / Calm Score™
Cyclical link	Reflection alters subsequent forethought	Miss $\rightarrow$ strategy attribution $\rightarrow$ new plan chip on retry	XP toast that resets nothing
Microanalysis	Brief, task-timed questions at phase boundaries	1–2 prompts pre/during/post attempt	20-item SRL inventory
SAFE-SRL	Practice doctrine for phase $\rightarrow$ control cycles under FEI	This chapter	SRL costume ads

Operational definition (HYPOTHESIS): A Practice path is *SAFE-SRL compliant* when (a) each FEI-relevant attempt has at least thin forethought (specific goal + strategy intent), (b) performance-phase support protects construction (SAFE-ATTN / SAFE-HINT) rather than interrupting for surveys, (c) reflection elicits **controllable** attributions and one adaptive next step, (d) that next step is *written into* the following forethought (visible plan change), and (e) success is measured by strategy-adaptive `retry_120s`, `challenge_accept` motives, and `solo_transfer_pass` — not by reflection completion rates.

CI.3 The cyclical model is company vocabulary — not a sticker

FACT: Zimmerman (2000, in Boekaerts, Pintrich & Zeidner, *Handbook of Self-Regulation*, Academic Press, pp. 13–39, doi:10.1016/B978-012109890-2/50031-7) — social-cognitive account of attaining self-regulation; self-regulatory processes organized as interdependent **forethought** → **performance** → **self-reflection** cycles, with motivational beliefs (e.g., self-efficacy) energizing the loop.

FACT: Zimmerman (2002, *Theory Into Practice*, 41(2), 64–70, doi:10.1207/s15430421tip4102_2) — practitioner overview of the same three-phase structure: processes *before*, *during*, and *after* learning efforts; reflections from prior efforts feed subsequent forethought (e.g., dissatisfaction → lowered efficacy → diminished later effort unless attributions and plans are repaired).

FACT: Zimmerman & Moylan (2009, in Hacker, Dunlosky & Graesser, *Handbook of Metacognition in Education*, Routledge, pp. 299–315) — refined cyclical model emphasizing the intersection of metacognition and motivation; performance-phase volitional/self-control processes sit alongside monitoring.

FACT (review): Panadero (2017, *Frontiers in Psychology*, 8, 422, doi:10.3389/fpsyg.2017.00422) — places Zimmerman’s cyclical phases among major SRL models; phases are descriptive process architecture, not a license for a single composite “SRL ability” score.

Kill: One SRL Score™ collapsing phases.

Survive: Named phase moments with named controls.

CI.4 Forethought: specific goals and strategies beat vibe goals

FACT: Cleary & Zimmerman (2001, *Journal of Applied Sport Psychology*, 13(2), 185–206, doi:10.1080/104132001753149883) — adolescent basketball free-throw microanalysis: experts set more **specific** goals, selected more **technique-oriented** strategies, and reported higher self-efficacy than non-experts and novices; forethought processes intercorrelated; self-reflection attributions predicted later strategy selection (cyclical evidence).

Applied (HYPOTHESIS): Mission launch should ask for a *grain-local* goal (“finish this FormatId conversion without hard peek”) and a strategy chip (diagram first / equation first / bridge checklist) — not “crush algebra.” Efficacy prompts stay task-specific (SAFE-MONITOR / Bandura grain), never identity Likert.

Commercial implication: Copy “plan the attempt, not the streak.” Competitors can autoplay the next item; MindCraft sells a 10-second forethought that changes what counts as success on *this* card.

Kill: Goal theater with no strategy constraint.

Survive: Specific process goals + strategy intent as first-class events.

CI.5 Performance: protect the attempt; monitor without survey spam

FACT: Zimmerman’s performance phase includes self-control (attention, strategies) and self-observation (Zimmerman, 2002; Zimmerman & Moylan, 2009). Heavy instrumentation mid-task can become the task.

HYPOTHESIS (product): Mid-attempt UX follows SAFE-ATTN + SAFE-RETRIEVE + SAFE-HINT: protect the window; mode-contingent stalls; soft before hard. Metacognitive monitoring prompts (if any) are single, stall-triggered,

and routed to control — not a running Metacognition Score.

Wound: Cleary/Zimmerman athletic microanalysis pauses are interviewer-administered; app prompts must be thinner. Prefer silent telemetry + rare prompts over constant “how are you regulating?” popups.

Kill: Mid-attempt SRL questionnaire / Focus Score™.

Survive: Construction-first performance phase with contingent help.

CI.6 Self-reflection: controllable attributions → adaptive inferences

FACT: Zimmerman (2002) — attributing a poor math score to controllable process causes (wrong strategy) sustains motivation because a different strategy may work; ability attributions and defensive reactions (withdraw, skip tests) undermine further learning cycles.

FACT: Cleary, Zimmerman & Keating (2006, *Research Quarterly for Exercise and Sport*, 77(2), 251–262, doi:10.1080/02701367.2006.10599358) — additive training across phases: groups trained in more phases showed better free-throw accuracy and more frequent self-correction after misses; three-phase trainees showed the most adaptive profile (strategic attributions, adaptive inferences, process criteria in self-evaluation).

Applied (HYPOTHESIS): Soft-wrong / post-mission reflection should offer **process/strategy** attribution chips (“picked wrong representation,” “skipped bridge check”) and one adaptive inference that becomes the next forethought plan — densifying SAFE-AAR / SAFE-ANNOT without military or grandmaster cosplay.

Kill: Ability shame; reflection without next-plan writeback; defensive skip framed as grit.

Survive: Strategy attribution → visible plan change on retry.

CI.7 Microanalysis > inventories for product truth

FACT: Cleary, Callan & Zimmerman (2012, *Education Research International*, 2012, Article 428639, doi:10.1155/2012/428639) — SRL **microanalysis:** brief, task-specific questions administered at temporal phase boundaries (before/during/after a defined task); targets subprocesses (goal-setting, strategic planning, monitoring, self-evaluation, attributions); designed as contextualized event assessment rather than decontextualized trait surveys.

Commercial implication: Instrument Practice with microanalytic *events* (forethought_goal_type, strategy_intent, attribution_class, adaptive_inference_applied) under SAFE-INSTRUMENT. Do not ship a 40-item “Are you a self-regulated learner?” onboarding that parents confuse with proof.

Kill: Survey SRL ≡ FEI readiness.

Survive: Phase-boundary micro-prompts as telemetry + light UX.

CI.8 Interventions work when strategy + motivation + reflection combine — with caveats

FACT: Dignath & Büttner (2008, *Metacognition and Learning*, 3(3), 231–264, doi:10.1007/s11409-008-9029-x) — meta-analysis of SRL interventions at primary and secondary levels: training can improve academic performance, strategy use, and motivation; effects vary by components (metacognitive reflection often important) and context (subject and grade moderators).

FACT: Dignath, Büttner & Langfeldt (2008, *Educational Research Review*, 3(2), 101–129, doi:10.1016/j.edurev.2008.02.003) — primary-school focused meta-analysis: SRL programmes can be effective even early; combinations of metacognitive and motivational elements matter; mathematics effects notable in primary samples in their synthesis — not a blank check for HS ACT apps.

Wound (HYPOTHESIS): Lab and school interventions $\neq$ 12-minute app sessions. MindCraft should claim **cycle scaffolding**, not “SRL training effect sizes transfer to our DAU.” Secondary-school moderators in Dignath & Büttner warn against copying primary-math packages wholesale into ACT prep.

Kill: “Research shows SRL raises grades — therefore our reflection journal works.”

Survive: Multi-phase micro-scaffolds + FEI outcomes as the honesty bar.

CI.9 Where competitors fail the cycle

Competitor pattern	Broken phase	MindCraft counter
Duo-style streak autopilot	Weak forethought; reflection = XP	Goal+strategy chips; adaptive inference writeback
ChatGPT dump	No student performance phase	Attempt-first; SAFE-FADE / SAFE-HINT
Brilliant delight streak	Reflection as dopamine	Process attribution; solo proof
School portal journals	Reflection detached from next task	One inference → next mission plan
“Metacognition module” quiz	Survey $\neq$ event	Microanalysis at phase boundaries

FOUNDER BELIEF: The wedge is not more content about SRL. It is a Practice loop that is an SRL cycle thin enough for Maya and rich enough for tutors to coach (SAFE-APPRENTICE articulate step).

CI.10 Decision table (ship / wound / kill)

Claim	Verdict	Notes
Three-phase cycle as UX architecture	Ship (provisional)	Zimmerman line; wire controls
Microanalytic events > trait SRL surveys	Ship	Cleary et al. 2012

Specific process goals + strategy intent	Ship	Cleary & Zimmerman 2001 analogue
Controllable attributions → plan writeback	Ship	Zimmerman 2002; Cleary et al. 2006
Multi-phase training additive	Wound→test	Motor-skill RCT; math transfer open
Meta-analytic SRL → ACT points	Kill as ad	Context/moderator wound
SRL Score™ / reflection streak NS	Kill	Standing ban family
Streak ≡ self-regulation	Kill	HID / SAFE-INSTRUMENT

CI.11 Experiments

ID	Question	Design	Primary
SRL-1	Forethought chips (goal+strategy) vs autopilot next-item on same bank	2-arm	solo_transfer_pass; peek binge; time-to-honest-attempt
SRL-2	Post-miss strategy attribution + plan writeback vs XP-only reflection	2-arm	adaptive retry_120s; strategy-change rate; early quit
SRL-3	Microanalytic 3-prompt cycle vs 8-item end journal	2-arm	FEI; completion without fatigue; tutor QA fidelity
SRL-4	Ability-attribution UI vs process-attribution UI after soft-wrong	2-arm	challenge_accept motives; defensive skip
SRL-5	Parent CBC: “plans the attempt” vs “builds metacognition habits” vs streak hero	CBC	WTP; trust (SAFE-WTP)
SRL-QUAL	10 Maya: when planning felt useful vs naggy; when reflection changed next try	Qual	codebook: theater vs control

Falsifier: Autopilot ≥ chip forethought on solo transfer and peek binge → demote forethought cost; keep optional.

Falsifier: Journal arm ≥ microanalytic cycle on FEI with less friction → wound microanalysis product claim; keep event telemetry only.

Falsifier: Parents prefer metacognition-habit CBC with equal retention → rewrite copy; still do not ship SRL Score™.

Pre-register: SRL-* before “AI metacognition coach,” reflection streaks, or SRL Score™ campaigns. Densify MONITOR/AAR/HELP/ATTN — do not fork a second dashboard religion.

CI.12 So what for MindCraft commercially

- **Copy:** “Plan the attempt. Adjust the strategy. Prove it alone.” Never “unlock your inner self-regulated learner” or reflection-streak hero.
 - **Product:** Phase-boundary micro-prompts with writeback; strategy attribution chips; protect mid-attempt.
 - **Positioning:** Against streak autopilot and against journal theater — MindCraft is the FEI gym with a visible cycle.
 - **Metric:** Strategy-adaptive retries, motive-coded challenge accepts, solo transfer — demote reflection completion % and SRL Score™.
 - **Kill list:** SRL Score™; survey≡event; streak≡SRL; forethought without control; ability-shame reflection; ACT points from metacognition modules.
 - **Growth / vision:** Parent/tutor demo shows Maya miss → process attribution → new strategy chip → cleaner retry — identity is owning the cycle, not collecting journal badges.
-

CI.13 Doctrine — SAFE-SRL (provisional)

1. **SRL is a cycle around a real attempt** — forethought → performance → reflection → rewritten forethought.
2. **Phase→control or it is costume** — every prompt must change goal, strategy, help, or next plan.
3. **Specific process goals beat vibe goals** — grain-local, technique-oriented.
4. **Performance protects construction** — no mid-attempt survey religion (SAFE-ATTN).
5. **Controllable attributions → adaptive inferences** — ability shame and defensive skip are fails.
6. **Microanalysis > inventories** for product truth (Cleary lineage).
7. **Proof = FEI + strategy change** — not reflection NPS or SRL Score™.
8. **No SRL Score™ / streak-as-SRL / metacognition-module ACT guarantees.**

Confidence: High on cyclical architecture, microanalysis vs surveys, and attribution/adaptive-inference logic. Medium that thin app chips will match multi-phase training gains — instrument SRL-* before metacognition marketing.

References (verified)

- Cleary, T. J., Callan, G. L., & Zimmerman, B. J. (2012). Assessing self-regulation as a cyclical, context-specific phenomenon: Overview and analysis of SRL microanalytic protocols. *Education Research International*, 2012, Article 428639. <https://doi.org/10.1155/2012/428639>
- Cleary, T. J., & Zimmerman, B. J. (2001). Self-regulation differences during athletic practice by experts, non-experts, and novices. *Journal of Applied Sport Psychology*, 13(2), 185–206. <https://doi.org/10.1080/104132001753149883>

- Cleary, T. J., Zimmerman, B. J., & Keating, T. (2006). Training physical education students to self-regulate during basketball free throw practice. *Research Quarterly for Exercise and Sport*, 77(2), 251–262. <https://doi.org/10.1080/02701367.2006.10599358>
- Dignath, C., & Büttner, G. (2008). Components of fostering self-regulated learning among students. A meta-analysis on intervention studies at primary and secondary school level. *Metacognition and Learning*, 3(3), 231–264. <https://doi.org/10.1007/s11409-008-9029-x>
- Dignath, C., Büttner, G., & Langfeldt, H. P. (2008). How can primary school students learn self-regulated learning strategies most effectively? A meta-analysis on self-regulation training programmes. *Educational Research Review*, 3(2), 101–129. <https://doi.org/10.1016/j.edurev.2008.02.003>
- Panadero, E. (2017). A review of self-regulated learning: Six models and four directions for research. *Frontiers in Psychology*, 8, 422. <https://doi.org/10.3389/fpsyg.2017.00422>
- Zimmerman, B. J. (2000). Attaining self-regulation: A social cognitive perspective. In M. Boekaerts, P. R. Pintrich, & M. Zeidner (Eds.), *Handbook of self-regulation* (pp. 13–39). Academic Press. <https://doi.org/10.1016/B978-012109890-2/50031-7>
- Zimmerman, B. J. (2002). Becoming a self-regulated learner: An overview. *Theory Into Practice*, 41(2), 64–70. https://doi.org/10.1207/s15430421tip4102_2
- Zimmerman, B. J., & Moylan, A. R. (2009). Self-regulation: Where metacognition and motivation intersect. In D. J. Hacker, J. Dunlosky, & A. C. Graesser (Eds.), *Handbook of metacognition in education* (pp. 299–315). Routledge.

Part CII — Generative Learning Activities (Fiorella/Mayer) in Practice

Chapter status: Living evidence + Practice UX brief — Researcher tick 2026-08-08 (UTC hour 6 ≡ Red Team slot, but ch102 never written → prefer Researcher per rotation; researcher count since synthesizer v1.13 = 2 → Researcher)

Primary question: How should MindCraft ship Fiorella/Mayer generative learning activities — summarize, map, draw, imagine, self-test, self-explain, teach, enact — inside Practice without minting a Generative Learning Score™, journal theater, or “learning-by-doing” costume that never forces select–organize–integrate?

Owners: Product (Practice / Notes / Map / Solver handoffs) · Brand · HITL tutors · Red Team

Commercial job: Ship **SAFE-GENERATE** densifying SAFE-EXPLAIN / SAFE-DUAL / SAFE-SRL / SAFE-SE / SAFE-RETRIEVE / SAFE-ATTN: thin, timed generative asks that change what the student *constructs*, prove value with solo/delayed transfer — never activity-checklist ads or AI monologue sold as “generation.”

Builds on: Parts XXVI (CLT), XL (self-explanation), XCII (SAFE-EXPLAIN), XCVII (SAFE-DUAL), XCVIII (SAFE-MONITOR), CI (SAFE-SRL), XCIII (SAFE-RETRIEVE). Product seams: soft-wrong, coach wrap, FormatId conversion, teach-back after grain, Map return.

CII.1 Why this chapter exists

Competitors sell either (a) passive consumption (watch solution → green check) or (b) activity theater (summarize this lesson / draw anything / teach a friend for XP). Generative learning research says the load-bearing move is

sense-making: the learner selects relevant structure, organizes it, and integrates it with prior knowledge so transfer becomes possible. MindCraft’s brand risk is claiming “active learning” while shipping either AI essays the student reads or checkbox activities that never reorganize math structure.

FOUNDER BELIEF under audit: Identity as a mathematical thinker requires *generating* explanations, sketches, and teach-backs — not collecting them from the model. The product earns that claim by wiring generative asks to FEI attempts, not by advertising eight strategies as a curriculum module.

Claims we refuse as doctrine:

1. Generative Learning Score™ / Activity Completion % as North Star.
2. “We use generative learning science” ads without SOI-forcing prompts.
3. AI wrap / monologue ≡ student generation (already SAFE-EXPLAIN / kill #8).
4. Verbatim copy-summary / highlight theater as summarizing.
5. Unguided free-draw as dual coding / format mastery.
6. Teach-back as social XP or influencer cosplay without content quality.
7. Eight-strategy checklist before every item (germane-load bomb).
8. Guaranteed ACT points from “generative Practice.”

CII.2 Constructs (keep mechanisms distinct)

Construct	Research meaning	MindCraft analogue	Not the same as
Generative learning	Reorganize + integrate with prior knowledge for transfer (Wittrock; Fiorella & Mayer)	Student-produced structure after/during attempt	Passive solution view
SOI	Select → Organize → Integrate (Mayer)	Prompt that forces pick/structure/link	“Reflect in 3 sentences” vibe
GLA / eight strategies	Summarize, map, draw, imagine, self-test, self-explain, teach, enact	Sparse toolkit chips, grain-timed	Full eight before every card
Constructive (ICAP)	Generate ideas beyond given material	Soft-wrong SE; teach-back; structural draw	Clicking / underlining / rewatching
SAFE-GENERATE	Practice doctrine for SOI-forcing micro-GLAs under FEI	This chapter	Activity costume ads

Operational definition (HYPOTHESIS): A Practice path is *SAFE-GENERATE compliant* when (a) generative asks require student production that cannot be satisfied by pasting the coach text, (b) timing protects the attempt window (SAFE-ATTN) — generation is pre-brief, stall-contingent, or post-attempt, never mid-solve spam, (c) strategy choice is grain-local (e.g., draw for diagram FormatId; teach-back after bridge grain), (d) quality gates exist (structure check, key idea coverage) without a vanity Generative Score, and (e) success is measured by solo_transfer_pass / delayed

mixed / format-hop — not GLA completion rates.

CII.3 Generative learning is company vocabulary — not a sticker pack

FACT: Wittrock (1974, *Educational Psychologist*, 11(2), 87–95, doi:10.1080/00461527409529129) — generative model: learners actively construct meanings consistent with prior knowledge; learning with understanding is generation and transfer of meaning from background and experience, not passive reception.

FACT: Wittrock (1989, *Educational Psychologist*, 24(4), 345–376, doi:10.1207/s15326985ep2404_2) — generative processes of comprehension: generation, motivation, attention, memory; mind as active constructor, not passive consumer.

FACT: Mayer (1996, *Educational Psychology Review*, 8(4), 357–371, doi:10.1007/BF01463939) — SOI model: select relevant information, organize into coherent structure, integrate with prior knowledge for sense-making from expository text.

FACT: Fiorella & Mayer (2016, *Educational Psychology Review*, 28(4), 717–741, doi:10.1007/s10648-015-9348-9; online 2015) — eight generative strategies with median effect sizes $\geq \sim 0.40$ across reviewed literatures (summarizing $d \approx 0.50$; mapping ~ 0.62 – 1.07 by organizer type; drawing ~ 0.40 ; imagining ~ 0.65 ; self-testing ~ 0.57 ; self-explaining ~ 0.61 ; teaching ~ 0.77 ; enacting ~ 0.51); grounded in Wittrock + Mayer SOI; boundary conditions emphasized per strategy.

FACT (book): Fiorella & Mayer (2015, *Learning as a Generative Activity*, Cambridge University Press, doi:10.1017/CBO9781107707085) — same eight strategies with practical recommendations and boundary conditions for implementation.

FACT: Chi & Wylie (2014, *Educational Psychologist*, 49(4), 219–243, doi:10.1080/00461520.2014.965823) — ICAP: Passive < Active < Constructive < Interactive for learning outcomes; generative processing aligns with constructive/interactive modes, not passive receipt or mere manipulative “active” behaviors (underlining, verbatim notes).

Kill: Sticker pack of eight activities without SOI.

Survive: Named generative moves that force construction.

CII.4 Self-explain and self-test already live — do not rebrand them as a new religion

FACT: Self-explaining and self-testing are two of the eight GLAs with strong support (Fiorella & Mayer, 2016). Self-explanation science is already Part XL / SAFE-EXPLAIN densifier; retrieval/self-testing is SAFE-RETRIEVE / SAFE-SCHED.

HYPOTHESIS (product): SAFE-GENERATE does **not** fork a parallel SE or retrieval stack. It adds the *missing* GLAs (summarize-in-own-words after grain, structural map/draw, teach-back) and forbids calling coach monologue “self-explanation.”

Commercial implication: Marketing may say “students generate, not consume” — never “powered by eight generative strategies™.” Competitors can list Dunlosky techniques; MindCraft sells *when* generation happens relative to FEI attempts.

Kill: Eight-box curriculum module.

Survive: Sparse GLA chips densifying existing SE/retrieval doctrine.

CII.5 Summarize and teach: production quality beats presence

FACT: Effective summarizing requires restating main ideas in one’s own words — not verbatim copying (Fiorella & Mayer, 2016). Verbatim notes/highlighting map to ICAP *active*, not *constructive* (Chi & Wylie, 2014).

FACT: Fiorella & Mayer (2013, *Contemporary Educational Psychology*, 38(4), 281–288, doi:10.1016/j.cedpsych.2013.06.001) — teaching expectancy improved immediate comprehension vs study-for-test ($d=0.59$ in reported contrast); preparing-to-teach alone did not reliably sustain 1-week delay benefits; students who both prepared and actually taught showed delayed benefits — explanation act matters for durability.

FACT: Fiorella & Mayer (2014, *Contemporary Educational Psychology*, 39(2), 75–85, doi:10.1016/j.cedpsych.2014.01.001) — expectations and explanations in learning-by-teaching: long-term benefits tied to generating explanations, not expectancy alone; quality of explanation (knowledge-building vs knowledge-telling) is critical.

Applied (HYPOTHESIS): After a bridge/grain win (or soft-wrong repair), offer a **teach-fictive-peer** card: 2–4 sentences or voice note explaining the *structure* (what to select, how pieces connect) without the solution stem visible. Reject paste-from-coach; optional rubric: names the bridge / FormatId conversion / key constraint. Summarize prompts are grain-local (“one sentence: what changed when you moved formats?”), not end-of-session essay.

Commercial implication: Parent/tutor demo shows Maya teaching the bridge back — competence proof (SAFE-PROOF) without influencer “teach TikTok” cosplay.

Kill: Teach-for-XP / expectancy-only theater.

Survive: Actual explanation production with delayed-transfer proof.

CII.6 Draw / map / imagine: guidance and FormatId, or cognitive tax

FACT: Fiorella & Zhang (2018, *Educational Psychology Review*, 30(3), 1115–1137, doi:10.1007/s10648-018-9444-8) — learner-generated drawing helps vs reading-only or text-focused strategies; effects vs instructor-provided illustrations are mixed and often require high guidance (partially provided figures, compare-to-model); drawing is cognitively demanding; accuracy and monitoring mediate benefits.

FACT (age/support boundary): Brod (2021, *Educational Psychology Review*, 33, 1295–1318, doi:10.1007/s10648-020-09571-9) — reviews which generative strategies have evidence at which school ages; support and prerequisites matter; drawing benefits vs provided illustrations often need substantial instructional support.

HYPOTHESIS (product): Drawing/mapping is **FormatId-contingent** and scaffolded: blank canvas only when grain is spatial/diagrammatic and prior model exists; prefer partially provided skeletons (axes, triangle template, number-line ticks) + “complete the structure” — aligns SAFE-DUAL (structural figures) and SAFE-EXPTIME (expertise-sensitive generation). Imagining (mental sketch) is a low-friction alternative when motor/UI cost is high — still require a check (label three parts aloud).

Kill: Always-draw / decoration doodle $\equiv$ dual coding.

Survive: Guided structural generation on diagram/graph grains.

CII.7 Boundary conditions: germane load, anxiety, and timing

FACT: Generative strategies can fail under high element interactivity, poor training, or when compared to strong provided models (Fiorella & Mayer, 2016; Fiorella & Zhang, 2018). CLT / expertise reversal (Parts XXVI, C / SAFE-EXPTIME) imply novices may need worked structure before free generation.

HYPOTHESIS: SAFE-GENERATE inherits SAFE-DD $\times$ anxiety and SAFE-EXPOSE: destake generation (no public shame reel); dose GLAs (one micro-ask, not eight); never flood anxious Maya with teach-back on day-one ACT items. Mid-attempt generation spam is banned (SAFE-ATTN). Prefer post-attempt SOI prompts or stall-contingent SE already in SAFE-EXPLAIN.

Contradiction to watch: Activity completion can *feel* productive while transferring poorly if generation is shallow (knowledge-telling). Red Team must kill any GLA completion NS.

Kill: Generation-everywhere / struggle-as-generation.

Survive: Timed, scaffolded, grain-local GLAs.

CII.8 Competitive contrast (mechanism lens)

Competitor pattern	Mechanism	MindCraft counter
Solution video $\rightarrow$ green	Passive / active ICAP	Attempt-first + post generative ask
Unlimited AI essay	Coach generates; student consumes	SE/teach-back <i>by student</i> before wrap (SAFE-EXPLAIN)
“Active learning” checklist	Activity theater	SOI-forcing chips + FEI proof
Free-draw sticker	Unguided drawing tax	Guided FormatId skeletons (SAFE-DUAL)
Teach-for-streak	Social XP	Teach-fictive-peer + delayed transfer

Positioning line (FOUNDER BELIEF $\rightarrow$ copy test): “You generate the structure. The coach doesn’t do it for you.”

CII.9 Confidence matrix

Claim	Label	Confidence	Notes
Generative $\neq$ passive consumption; SOI is the spine	FACT (theory) + HYPOTHESIS (product map)	High	Wittrock; Mayer SOI; Fiorella & Mayer
Eight strategies median $d \geq \sim 0.4$ in reviewed literatures	FACT (review summary)	High on direction; medium on portability to ACT Practice UI	Do not quote as guaranteed product lift
Teach act > expectancy alone for delay	FACT	High	Fiorella & Mayer 2013/2014
Unguided draw often loses to provided illustrations	FACT (boundary)	High	Fiorella & Zhang 2018
Thin app GLAs $\rightarrow$ solo transfer lift	HYPOTHESIS	Medium	Needs GENACT-*
Generative Learning Score™ helps learning	SPECULATION (reject)	—	Kill

CII.10 Experiments

ID	Question	Design	Primary
GENACT-1	Post-grain teach-fictive-peer vs solution-view-only	2-arm	delayed solo_transfer_pass; explanation quality rubric
GENACT-2	Guided structural draw (skeleton) vs free-draw vs provided figure only on diagram FormatId	3-arm	format-hop accuracy; time-on-tax; draw accuracy
GENACT-3	One-sentence SOI summarize after soft-wrong vs XP-only continue	2-arm	misconception recurrence; adaptive retry_120s
GENACT-4	Eight-strategy menu every item vs sparse grain-triggered GLA	2-arm	FEI; quit; peek binge (load)
GENACT-5	Parent CBC: “generates the explanation” vs “active learning strategies” vs streak hero	CBC	WTP; trust (SAFE-WTP)

GENACT-QUAL	10 Maya: which generative ask felt useful vs busywork; teach-back shame	Qual	codebook: SOI vs costume
-------------	---	------	--------------------------

Falsifier: Solution-view $\geq$ teach-back on delayed solo transfer $\rightarrow$ demote teach-back cost; keep optional Notes export only.

Falsifier: Free-draw $\geq$ guided skeleton on format-hop with less time $\rightarrow$ wound guidance claim; still ban decoration doodle ads.

Falsifier: Eight-menu $\geq$ sparse triggers on FEI with equal quit $\rightarrow$ still kill menu-as-brand; instrument load.

Pre-register: GENACT-* before “generative learning OS,” activity Score™, or eight-strategy splash campaigns. Denisify EXPLAIN/DUAL/SRL/SE — do not fork a second pedagogy religion.

CII.11 So what for MindCraft commercially

- **Copy:** “Generate the structure. Then prove it alone.” Never “complete all eight generative strategies” or Generative Learning Score™.
- **Product:** Grain-triggered teach-back, one-sentence SOI summary, guided FormatId skeletons; protect attempt windows.
- **Positioning:** Against passive AI tutors and against busywork activity apps — MindCraft is the FEI gym where *Maya* constructs meaning.
- **Metric:** Solo/delayed transfer, format-hop, explanation quality samples — demote GLA completion %.
- **Kill list:** Generative Learning Score™; AI monologue $\equiv$ generation; verbatim summary theater; unguided doodle $\equiv$ dual coding; teach-for-XP; eight-menu load bomb; ACT guarantees from GLA branding.
- **Growth / vision:** Parent proof artifact = student’s teach-back or structural sketch on a bridge grain — identity evidence without star-walls.

CII.12 Doctrine — SAFE-GENERATE (provisional)

1. **Generation = SOI** — select, organize, integrate; else costume.
2. **Student produces; coach does not substitute** — densifies SAFE-EXPLAIN.
3. **Sparse, grain-triggered GLAs** — not eight-before-every-item.
4. **Teach-back requires the explanation act** — expectancy/XP alone insufficient.
5. **Draw/map need guidance on FormatId grains** — densifies SAFE-DUAL; no decoration.
6. **Protect performance windows** — no mid-solve generative spam (SAFE-ATTN).
7. **Proof = transfer / format-hop / quality sample** — not activity completion.
8. **No Generative Learning Score™ / AI $\equiv$ generation / eight-strategy ACT ads.**

Confidence: High on SOI/ICAP constructive requirement, teach-back durability evidence, and drawing boundary conditions. Medium that thin mobile GLAs match lab effect sizes — instrument GENACT-* before generative-learning marketing.

References (verified)

- Brod, G. (2021). Generative learning: Which strategies for what age? *Educational Psychology Review*, 33, 1295–1318. <https://doi.org/10.1007/s10648-020-09571-9>
- Chi, M. T. H., & Wylie, R. (2014). The ICAP framework: Linking cognitive engagement to active learning outcomes. *Educational Psychologist*, 49(4), 219–243. <https://doi.org/10.1080/00461520.2014.965823>
- Fiorella, L., & Mayer, R. E. (2013). The relative benefits of learning by teaching and teaching expectancy. *Contemporary Educational Psychology*, 38(4), 281–288. <https://doi.org/10.1016/j.cedpsych.2013.06.001>
- Fiorella, L., & Mayer, R. E. (2014). Role of expectations and explanations in learning by teaching. *Contemporary Educational Psychology*, 39(2), 75–85. <https://doi.org/10.1016/j.cedpsych.2014.01.001>
- Fiorella, L., & Mayer, R. E. (2015). *Learning as a generative activity: Eight learning strategies that promote understanding*. Cambridge University Press. <https://doi.org/10.1017/CBO9781107707085>
- Fiorella, L., & Mayer, R. E. (2016). Eight ways to promote generative learning. *Educational Psychology Review*, 28(4), 717–741. <https://doi.org/10.1007/s10648-015-9348-9>
- Fiorella, L., & Zhang, Q. (2018). Drawing boundary conditions for learning by drawing. *Educational Psychology Review*, 30(3), 1115–1137. <https://doi.org/10.1007/s10648-018-9444-8>
- Mayer, R. E. (1996). Learning strategies for making sense out of expository text: The SOI model for guiding three cognitive processes in knowledge construction. *Educational Psychology Review*, 8(4), 357–371. <https://doi.org/10.1007/BF01463939>
- Wittrock, M. C. (1974). Learning as a generative process. *Educational Psychologist*, 11(2), 87–95. <https://doi.org/10.1080/00461527409529129>
- Wittrock, M. C. (1989). Generative processes of comprehension. *Educational Psychologist*, 24(4), 345–376. https://doi.org/10.1207/s15326985ep2404_2

Part CIII — Error Climate & Psychological Safety in Tutoring

Chapter status: Living evidence + tutor/ops brief — Researcher tick 2026-08-08 (UTC hour 9; hour%6≠0; researcher count since synthesizer v1.13 = 3 → Researcher)

Primary question: How should MindCraft destake mistakes in tutoring and Practice without lowering standards, minting a Safety Score™, or confusing warm vibes with an *error climate* that actually changes adaptive reactions to errors?

Owners: HITL tutors · Product (Practice soft-wrong / Map) · Brand · Parent messaging · Red Team

Commercial job: Ship **SAFE-ERRCLIMATE** densifying SAFE-MISCON / SAFE-HELP / SAFE-TALK / SAFE-AAR / SAFE-PF / SAFE-EXPOSE: observable tutor and product responses to errors that raise adaptive retry and diagnosis — never “psychologically safe” posters, shame reels, or Comfort Score™ as North Star.

Builds on: Parts XXIV (affect/anxiety), XXXVIII (goal climate), XLIX (SAFE-MISCON), LIV (SAFE-AAR), LVIII (SAFE-EXPOSE), LXXIX (SAFE-FBTIME), LXXXIV (SAFE-TALK), LXXXIX (SAFE-HELP), XCIII–XCIV (retrieve / PF). Product seams: soft-wrong, hide-correctness diagnostic, tutor playbooks, parent dash (no live wrong stalk).

CIII.1 Why this chapter exists

Brand temptation: “We’re a psychologically safe space for math.” Competitor temptation: public leaderboards, red-X theater, or Bermuda-triangle redirection that abandons the student who erred. Organizational literature (Edmondson) made “psychological safety” a buzzword; classroom literature (Steuer/Dresel/Tulis/Soncini) names a sharper construct for MindCraft: **error climate** — how errors are perceived, evaluated, and used in the learning setting.

MindCraft already destakes some stakes (soft-wrong, C4 hide-correctness). The open commercial question is whether tutors and product *behaviors after an error* systematically look like a positive error climate — or like empty kindness / shame / redirect-away — and whether parents/tutors can witness that without a Safety Score™.

FOUNDER BELIEF under audit: Identity as a mathematical thinker requires interpersonal risk (admit confusion, try a wrong path, ask for help). That risk is only rational when the environment treats errors as diagnostic fuel, not as verdicts on worth — *and* still demands standards via FEI proof.

Claims we refuse as doctrine:

1. Psychological Safety Score™ / Comfort % / “vibes” as North Star.
2. “Psychologically safe” marketing without observable error-handling behaviors.
3. Destaking $\equiv$ lowering standards / never grading / never timed practice.
4. Shame reels, public wrong-answer stalk, family leaderboard as error theater.
5. Bermuda-triangle redirect (leave the erring student; crown the next hand).
6. Empty “mistakes are beautiful” posters without corrective feedback / Map route.
7. Therapy/CBT cosplay or Calm Score™ as error climate (already SAFE-EXPOSE).
8. Guaranteed ACT points from “safe tutoring.”

CIII.2 Constructs (keep mechanisms distinct)

Construct	Research meaning	MindCraft analogue	Not the same as
Team psychological safety	Shared belief the setting is safe for interpersonal risk-taking (Edmondson)	Tutor–student microclimate: ask, admit miss, try again	Warmth / niceness alone
Classroom error climate	Perception, evaluation, utilization of errors as learning elements (Steuer et al.)	Soft-wrong + stay-with-student repair + Map classify	Poster slogans
Adaptive affective-motivational reaction	Maintain motivation/functional affect after error	Soft-wrong copy; destaked framing	Empty praise
Adaptive action reaction	Effort, strategy change, work-it-out after error	retry_120s; Map route; AAR grain	Ignore-and-move-on

Positive error handling	Constructive teacher responses that invite learning from the error (Soncini et al.)	Tutor script: stay, diagnose, scaffold, re-attempt	Redirect / ignore / “No!” only
SAFE-ERRCLIMATE	Product+HITL doctrine for error climate under FEI	This chapter	Safety Score™

Operational definition (HYPOTHESIS): A session or Practice path is *SAFE-ERRCLIMATE compliant* when (a) errors trigger stay-with-student diagnosis (not Bermuda-triangle redirect), (b) evaluation stakes are separated from learning stakes where designed (soft-wrong / C4), (c) tutor/coach responses avoid ridicule and avoid empty “you’re fine” without Map/feedback, (d) adaptive *action* is instrumented (retry_120s, challenge_accept, misconception recurrence) — not comfort surveys alone, (e) high-stakes dose still exists on the exam rail (SAFE-EXAM), and (f) parents see proof age / repair artifacts — not live wrong stalk (SAFE-PDASH).

CIII.3 Psychological safety enables learning behavior — it is not the product

FACT: Edmondson (1999, *Administrative Science Quarterly*, 44(2), 350–383, doi:10.2307/2666999) — team psychological safety = shared belief that the team is safe for interpersonal risk-taking; in a multimethod field study of 51 manufacturing teams, psychological safety associated with learning behavior; learning behavior mediated the link to team performance; team efficacy did not predict learning when controlling for psychological safety.

FACT (boundary): Edmondson’s construct is group-level and organizational. Direct port to 1:1 tutoring or solo Practice is a **HYPOTHESIS**, not a license to sell corporate safety surveys as math pedagogy.

Applied (HYPOTHESIS): For MindCraft, the *portable* piece is the micro-calculus students run before risk: Will admitting “I don’t get the bridge” cost status? If yes, they hide, peek, or freeze (SAFE-HELP / SAFE-RETRIEVE). Tutor briefing and product copy should make interpersonal risk cheap *during learning*, then demand competence evidence *after*.

Kill: Safety Score™ / engagement-as-safety.

Survive: Risk-taking behaviors (ask, retry, expose misconception) as process signals under FEI.

CIII.4 Error climate is the classroom-specific spine — eight facets, one commercial job

FACT: Steuer, Rosentritt-Brunn, & Dresel (2013, *Contemporary Educational Psychology*, 38(3), 196–210, doi:10.1016/j.cedpsych.2013.03.002) — N=1,116 students in 56 German secondary math classrooms; validated multidimensional *perceived error climate* (eight subdimensions → superordinate factor); predicted adaptivity of individual error reactions above classroom goal structures and personal motivation; effort effects partially mediated through those reactions.

FACT (product checklist): Facets include teacher error tolerance; irrelevance of errors for assessment in learning contexts; teacher support after errors; absence of negative teacher/classmate reactions; analysis of errors; functionality

of errors for learning — jointly: errors as integral vs threat.

FACT: Dresel, Daumiller, Spear, Janke, Dickhäuser, & Steuer (2024, *British Journal of Educational Psychology*, doi:10.1111/bjep.12697) — 2-year longitudinal N=1,641 in 69 math classrooms (grades 5–6, five waves); adaptive affective-motivational and action reactions to errors *declined* on average; decline varied by classroom; more positive perceived error climate buffered the decline. Teacher support for error adaptivity is load-bearing.

Commercial implication: Error climate is a **designed** property of tutor ops + Practice UX — not a natural developmental gift. Longitudinal decline without climate support is a GTM argument for HITL briefing quality (SAFE-HITL), not Comfort Score™ ads.

Kill: Poster-only “mistakes welcome.”

Survive: Facet-mapped behaviors in playbooks and soft-wrong.

CIII.5 Teacher/tutor responses are causal levers — not just correlates

FACT: Tulis (2013, *Teaching and Teacher Education*, 33, 56–68, doi:10.1016/j.tate.2013.02.003) — observational/video studies of teachers’ error-management routines; responses shape beliefs about learning from errors and domain emotions; redirecting away from the erring student (“Bermuda triangle of error correction,” Oser & Spychiger) wastes the learning opportunity.

FACT: Soncini, Matteucci, & Butera (2021, *European Journal of Psychology of Education*, 36, 719–738, doi:10.1007/s10212-020-00494-1) — experimental pre-post (N=108 fifth-graders): positive vs neutral error-handling scripts; positive handling (stay-and-support) increased perceived trustful error climate vs ignore / redirect / bare “incorrect.” First experimental demo that manipulating teacher error strategies shifts climate perception.

FACT: Soncini, Visintin, Matteucci, Tomasetto, & Butera (2022, *Learning and Instruction*, 80, 101627, doi:10.1016/j.learninstruc.2022.101627) — N=563 Italian middle-schoolers, 32 math classes; multilevel path: positive error climate → adaptive reactions → math grades (direction commercial; magnitude uncertain).

Applied (HYPOTHESIS): Tutor QA should score *error episodes*, not warmth ratings: stay / diagnose / scaffold / re-attempt — or redirect / dump / shame? AI coach: soft-wrong + short principle (SAFE-EXPLAIN) + retry gate beats monologue + green-check theater.

Kill: Warmth NPS ≡ error climate.

Survive: Stay-diagnose-scaffold-reattempt as the HITL unit.

CIII.6 Learning from errors needs corrective feedback and destaked practice — not error romance

FACT: Metcalfe (2017, *Annual Review of Psychology*, 68, 465–489, doi:10.1146/annurev-psych-010416-044022) — review: error avoidance is common in classrooms but often counterproductive for typical learners; errorful learning *with corrective feedback* benefits learning; high-confidence errors are especially correctable (hypercorrection); teachers also gain diagnostic information from errors; if optimal high-stakes performance is the goal, encourage

commit-and-correct in *low-stakes* learning situations rather than avoid errors at all costs.

HYPOTHESIS (product): SAFE-ERRCLIMATE pairs with SAFE-EXAM dual rail: destake-to-learn (Practice soft-wrong, C4 diagnostic) / dose-to-prove (timed mixed sets). Error climate without corrective Map feedback is kindness costume; feedback without destaking for anxious Maya is flooding (SAFE-EXPOSE / SAFE-DD).

Contradiction to watch: Celebrating errors without standards can become low-rigor brand. Red Team kills “mistakes are the curriculum” when FEI proof disappears.

Kill: Error romance / never-correct / forever-low-stakes.

Survive: Low-stakes commit→correct; high-stakes prove later.

CIII.7 Boundary conditions: standards, public exposure, AI sycophancy

HYPOTHESIS: Positive error climate ≠ absence of evaluation forever. Steuer’s “irrelevance of errors for assessment” applies to *learning episodes*, not to deleting ACT reality. Public classmate ridicule facets do not license private shame dashboards for parents (SAFE-PDASH).

HYPOTHESIS (AI risk): Sycophantic “great try!” without diagnosing the misconception is *negative* error climate dressed as warmth (SAFE-REPAIR / Part XXXIII). Fluent wrongness plus praise destroys both safety *and* learning.

HYPOTHESIS (anxiety): Destaking helps approach; it does not replace graded exposure ladders when avoidance is entrenched (SAFE-EXPOSE). Dose GLAs/PF carefully (SAFE-GENERATE / SAFE-PF).

Kill: Comfort-only AI; public error humiliation; standards-free “safe space.”

Survive: Destaked diagnosis + standards on FEI.

CIII.8 Competitive contrast (mechanism lens)

Competitor pattern	Mechanism	MindCraft counter
Red-X leaderboard / shame reel	Negative classmate/teacher reaction facets	Private soft-wrong; no live wrong stalk
Unlimited “you’re amazing” AI	Warmth without analysis of errors	Soft-wrong + Map classify + short principle
Bermuda-triangle classroom	Leave erring student	Stay-with-student tutor QA
“Safe space” poster / Comfort Score™	Slogan without handling scripts	Observable error-handling playbook
Day-one timed ACT flood	High-stakes without destake-to-learn	SAFE-EXAM dual rail

Positioning line (FOUNDER BELIEF → copy test): “Mistakes get diagnosed here — not mocked, not ignored, not praised into fog.”

CIII.9 Confidence matrix

Claim	Label	Confidence	Notes
Psych safety ↔ learning behavior in teams	FACT	High (org context)	Edmondson 1999; port to tutoring = HYPOTHESIS
Multidimensional math error climate predicts adaptive reactions	FACT	High	Steuer et al. 2013
Error climate buffers decline in error adaptivity over years	FACT	High (school longit.)	Dresel et al. 2024
Teacher handling strategies can shift perceived climate (causal)	FACT	High on climate perception	Soncini et al. 2021; short lesson
Climate → grades via adaptive reactions	FACT (correlational path)	Medium–high direction	Soncini et al. 2022; not ACT RCT
Errorful+feedback > errorless (typical learners)	FACT (review)	High on direction	Metcalfe 2017
Tutor stay-diagnose scripts raise retry_120s / FEI	HYPOTHESIS	Medium	Needs ERRCLIM-*
Safety Score™ improves learning	SPECULATION (reject)	—	Kill

CIII.10 Experiments

ID	Question	Design	Primary
ERRCLIM-1	Stay-diagnose-reattempt tutor script vs redirect/ignore script on soft-wrong episodes	2-arm HITL	retry_120s; misconception recurrence; student risk-taking (ask/admit)
ERRCLIM-2	Soft-wrong + Map classify vs soft-wrong + praise-only	2-arm Practice	delayed mix; adaptive action items; quit

ERRCLIM-3	Parent CBC: “mistakes get diagnosed” vs “psychologically safe space” vs streak hero	CBC	WTP; trust (SAFE-WTP)
ERRCLIM-4	Tutor QA rubric on error episodes vs warmth NPS only	Ops A/B	FEI per session; playbook fidelity; burnout proxy
ERRCLIM-5	Public vs private error visibility in small-group / parent views	2-arm	help-seeking; shame; challenge_accept
ERRCLIM-QUAL	10 Maya + 10 tutors: what made a miss feel safe vs humiliating vs useless	Qual	codebook: stay / redirect / praise-fog / diagnosis

Falsifier: Praise-only $\geq$ stay-diagnose on delayed mix with equal quit $\rightarrow$ wound diagnosis cost; still ban Bermuda-triangle and Safety Score™.

Falsifier: Warmth NPS predicts FEI as well as error-episode QA $\rightarrow$ still prefer behavioral QA for ops; do not make NPS the learning NS.

Falsifier: “Safe space” CBC $\geq$ “diagnosed mistakes” on WTP with equal trust $\rightarrow$ still ban Safety Score™; test copy hybrids.

Pre-register: ERRCLIM-* before psychologically-safe splash campaigns, Comfort Score™, or “we never make kids wrong” brand. Densify MISCON/HELP/TALK/AAR — do not fork a therapy religion.

CIII.11 So what for MindCraft commercially

- **Copy:** “Mistakes get diagnosed — then you prove the fix alone.” Never Psychological Safety Score™ or “never wrong here.”
- **Product:** Soft-wrong + stay-with-repair + Map misconception route; hide-correctness where designed; no parent live-wrong stalk.
- **Positioning:** Against shame apps and against comfort-only AI — MindCraft is the FEI gym where errors are fuel under standards.
- **Metric:** retry_120s, challenge_accept, misconception recurrence, tutor error-episode fidelity — demote comfort surveys.
- **Kill list:** Safety Score™; poster-only safe space; Bermuda-triangle redirect; shame reels; praise-fog AI; standards-free destake; ACT guarantees from climate branding.
- **Growth / vision:** Parent/tutor demo = one repaired bridge grain (before/after Map) — identity evidence from risk taken and corrected.

CIII.12 Doctrine — SAFE-ERRCLIMATE (provisional)

1. **Error climate > safety slogan** — facets and behaviors, not vibes.
2. **Stay with the student who erred** — no Bermuda-triangle redirect.
3. **Destake learning episodes; keep FEI standards** — dual rail with SAFE-EXAM.
4. **Adaptive action beats comfort survey** — instrument retry / repair / transfer.
5. **Corrective diagnosis required** — densifies SAFE-MISCON / SAFE-EXPLAIN; no praise-fog.
6. **Tutor QA on error episodes** — densifies SAFE-TALK / SAFE-HITL / SAFE-AAR.
7. **Private repair; no shame reel / live wrong stalk** — densifies SAFE-PDASH.
8. **No Psychological Safety Score™ / Comfort % / safe-space ACT ads.**

Confidence: High on error-climate construct validity, longitudinal buffering, and experimental handling→climate perception. Medium that 1:1 MindCraft scripts match classroom effect sizes — instrument ERRCLIM-* before psychologically-safe marketing.

References (verified)

- Dresel, M., Daumiller, M., Spear, J., Janke, S., Dickhäuser, O., & Steuer, G. (2024). Learning from errors in mathematics classrooms: Development over 2 years in dependence of perceived error climate. *British Journal of Educational Psychology*. <https://doi.org/10.1111/bjep.12697>
 - Edmondson, A. (1999). Psychological safety and learning behavior in work teams. *Administrative Science Quarterly*, 44(2), 350–383. <https://doi.org/10.2307/2666999>
 - Metcalfe, J. (2017). Learning from errors. *Annual Review of Psychology*, 68, 465–489. <https://doi.org/10.1146/annurev-psych-010416-044022>
 - Soncini, A., Matteucci, M. C., & Butera, F. (2021). Error handling in the classroom: An experimental study of teachers' strategies to foster positive error climate. *European Journal of Psychology of Education*, 36, 719–738. <https://doi.org/10.1007/s10212-020-00494-1>
 - Soncini, A., Visintin, E. P., Matteucci, M. C., Tomasetto, C., & Butera, F. (2022). Positive error climate promotes learning outcomes through students' adaptive reactions towards errors. *Learning and Instruction*, 80, 101627. <https://doi.org/10.1016/j.learninstruc.2022.101627>
 - Steuer, G., Rosentritt-Brunn, G., & Dresel, M. (2013). Dealing with errors in mathematics classrooms: Structure and relevance of perceived error climate. *Contemporary Educational Psychology*, 38(3), 196–210. <https://doi.org/10.1016/j.cedpsych.2013.03.002>
 - Tulis, M. (2013). Error management behavior in classrooms: Teachers' responses to student mistakes. *Teaching and Teacher Education*, 33, 56–68. <https://doi.org/10.1016/j.tate.2013.02.003>
-

Part CIV — Homework Help vs Practice Identity Conflict

Chapter status: Living evidence + Solver/Practice/parent brief — Researcher tick 2026-08-08 (UTC hour 12 ≡ Red Team slot, but ch104 never written → prefer Researcher per rotation; researcher count since synthesizer v1.13 = 4 → Researcher)

Primary question: How should MindCraft sell and ship **Solver** (homework help) without training executive completion identity — while **Practice** remains the FEI gym — without pretending parents do not need deadline help?

Owners: Product (Solver / Practice / fade) · Brand · Parent messaging · HITL · Red Team

Commercial job: Ship **SAFE-HWHELP** densifying SAFE-HELP / SAFE-FADE / SAFE-ROI / SAFE-DURABLE / SAFE-PROOF: two labeled rails (help-to-submit vs practice-to-prove); instrumental scaffolds only; solo transfer as the identity proof — never Homework Completion Score™, Chegg-clone ads, or “we finish your worksheet” as the brand.

Builds on: Parts XXXIII (AI crutch / PWC), XXXV (competitive audits), LIX (parent WTP), LXXVIII (SAFE-DURABLE dual-rail), LXXXVIII–LXXXIX (SAFE-FADE / SAFE-HELP), XCI (SAFE-HINT), XCIX (SAFE-ROI), CI–CIII (SRL / generate / error climate). Product seams: Solver, Practice sessions, fade ladder, solo_transfer_pass, parent dash (no live homework stalk).

CIV.1 Why this chapter exists

MindCraft already has two surfaces that collide in parental speech: “Help my kid finish tonight’s homework” and “Make my kid a mathematical thinker.” Chegg / Course Hero / ChatGPT optimized the first. Identity products that refuse the first lose GTM. Identity products that *become* the first lose the thesis.

FOUNDER BELIEF under audit: Homework help is a legitimate acquisition and retention job — but it is not the North Star. If Solver completion % rises while solo_transfer_pass falls, MindCraft has become a nicer Chegg.

Claims we refuse as doctrine:

1. Homework Completion Score™ / tonight-% submitted as learning NS.
 2. “We finish your worksheet” / unlimited answer dump as brand.
 3. Solver-open accuracy ≡ exam readiness / identity.
 4. Ban all homework help (purity theater that loses parents).
 5. Shame parents who need deadline relief.
 6. Practice and Solver as one undifferentiated “AI tutor” SKU.
 7. Bastani GPT Tutor cosplay as MindCraft Solver license (already XXXIII RT).
 8. Guaranteed grades / ACT points from Solver sessions.
-

CIV.2 Constructs (keep rails distinct)

Construct	Research meaning	MindCraft analogue	Failure mode
Homework as practice	Assigned work that builds skill when students actually work it	Practice rail / FEI gym	Outsourced answers
Executive homework help	Help that completes the product for the student	Bottom-out Solver dump / Chegg paste	Fake tonight-%
Instrumental homework help	Help that restores process so student can finish	Fade stage / principle hint / completion blank	Rare without design

Assisted performance	Score <i>with</i> tool open	Solver session accuracy	Mistaken for mastery
Unassisted transfer	Score / proof with tool closed	solo_transfer_pass	Skipped in GTM
SAFE-HWHELP	Dual-rail doctrine: help labeled; identity on Practice+solo	This chapter	One-SKU Chegg

Operational definition (HYPOTHESIS): A Solver path is *SAFE-HWHELP* compliant when (a) UI/copy labels help vs practice rails, (b) default help is instrumental (SAFE-HELP / SAFE-FADE), (c) full reveal is costly and logged, (d) every help episode offers / requires a follow-on Practice or solo grain when feasible, (e) parent messaging separates “unstuck tonight” from “proved alone,” and (f) product never markets assisted completion as identity transformation.

CIV.3 Homework can help achievement — when students do the work

FACT: Cooper, Robinson, & Patall (2006, *Review of Educational Research*, 76(1), 1–62, doi:10.3102/00346543076001001) — synthesis of U.S. homework research 1987–2003; despite design flaws across study types, evidence generally supports a positive homework→achievement influence, stronger in grades 7–12 than K–6, and stronger when students (not parents) report homework time.

Applied (HYPOTHESIS): The commercial lesson is *not* “more homework widgets.” It is that the achievement link assumes **student cognitive work**. Answer marketplaces sever that assumption while preserving the grade signal.

Kill: Tonight-submitted-% as MindCraft learning NS.

Survive: Homework as a *context* parents care about — routed into instrumental help + Practice proof.

CIV.4 When solutions are outsourced, HW grades decouple from exams

FACT (course evidence): DeGoede (2020, ASEE Annual Conference, doi:10.18260/1-2--33979) — engineering dynamics sections: pre–web-solutions era showed strong homework–exam correlations (reported $R^2 \approx 0.74$ – 0.98); after extensive self-reported Chegg-style solution use, homework–exam correlation collapsed ($R^2 \approx 0.08$) while homework scores stayed high / rose and exam performance weakened on average — correctness-graded HW without process control stopped indexing learning.

FACT (integrity audit): Pickering & Schuller (2024, *Journal of Academic Ethics*, doi:10.1007/s10805-024-09551-6) — audit in an Australian engineering school: Chegg content implicated across a majority of audited units; ~50% of questions answered within ~1.5 hours — fast enough for assignment and online-exam misuse; authors distinguish “homework-helper” branding from cheat-vulnerable business models.

FACT (quasi-experimental conference evidence — label carefully): Roshak (AEA 2023 poster / working paper, “Technology Undermining Education: The Case of Chegg”) — finance course; ~25% of students (incl. ~15% of high scorers) copied obviously wrong Chegg answers; DMCA takedown variation reducing solution availability associated with better subsequent exam performance especially among high-scoring students. **Not** a peer-reviewed journal RCT

— treat magnitude as provisional; mechanism (temptation to shirk process work) aligns with DeGoede / Pickering.

Applied (FOUNDER BELIEF → testable): Parent delight at “homework done” is orthogonal to durable identity. If MindCraft Solver raises assisted completion without solo proof, it recreates the Chegg grade–learning split inside our product.

Kill: Solver completion NPS / homework-finished toast as FEL.

Survive: Instrumental unstuck + required prove-alone follow-through.

CIV.5 Generative help recreates the crutch at AI speed

FACT: Bastani, Bastani, Sungu, Ge, Kabakc■, & Mariman (2025, *PNAS*, doi:10.1073/pnas.2422633122) — ~1,000 high-school math students; GPT Base raised practice grades (~+48%) but harmed later unassisted exams (~-17% vs never-AI); GPT Tutor (hint safeguards + teacher keys in prompt) largely neutralized harm but did **not** beat no-AI on learning; students used unfettered GPT as a crutch; perception ≠ exam learning.

Already company law (XXXIII RT): Bastani Tutor ■ MindCraft PWC; guards ≠ gain. This chapter densifies the *product-surface* implication: **Solver that behaves like GPT Base is commercially toxic to the identity claim**, even if parents love tonight’s answers.

HYPOTHESIS: Deadline-mode Solver without fade/solo gates will show Bastani Base signature: high assisted score, low solo_transfer_pass.

Kill: Unlimited ChatGPT-style homework hero.

Survive: Guarded instrumental Solver + Practice gym (SAFE-FADE / SAFE-HELP / PWC).

CIV.6 Instrumental vs executive help is the identity fork inside Solver

FACT: Nelson-Le Gall (1981, *Developmental Review*, 1(3), 224–246, doi:10.1016/0273-2297(81)90019-8) — distinguishes **executive** help-seeking (get the goal done for me) from **instrumental** help-seeking (get the minimal process help so I can finish myself); continued executive reliance risks dependency; instrumental help is a path from other-regulation to self-regulation.

FACT (ITS line, already SAFE-HELP): Aleven / Roll / Baker — on-demand help abuse (hint racing to bottom-out) often correlates with worse learning; help avoidance is the twin failure (summarized in Roll, Aleven, McLaren & Koedinger, 2011, *Learning and Instruction*, 21(2), 267–280, doi:10.1016/j.learninstruc.2010.07.004; Aleven et al., 2016 IJAIED help-seeking review).

Applied (HYPOTHESIS): Homework-help mode is where executive pressure peaks (due tonight). Product must make instrumental the default path and executive the expensive exception — not ban help, not shame need, not celebrate dump.

Kill: “Never stuck” answer-complete Solver ads.

Survive: Principle → completion → (costly) reveal; then Practice grain.

CIV.7 Parent messaging: two jobs, two sentences

FOUNDER BELIEF: Parents buy relief *and* identity. Conflating them produces either Chegg competitors (relief-only) or purity products parents bounce from (identity-only).

HYPOTHESIS (copy): Winning CBC pair (test in HWHELP-3 / densifies SAFE-WTP):

1. Help rail: “Get unstuck without pasting the answer.”
2. Practice rail: “Then prove the same join alone.”

Kill: Single splash “AI homework tutor that does it for them.”

Survive: Dual-rail honesty (mirrors SAFE-DURABLE exam/durable dual-rail).

Contradiction to watch: If instrumental Solver is too slow under deadline, students flee to ChatGPT. Red Team: measure time-to-unstuck *and* solo transfer — not purity of friction.

CIV.8 Competitive contrast (mechanism lens)

Competitor pattern	Mechanism	MindCraft counter
Chegg / Course Hero dump	Executive completion; HW—exam decouple	Instrumental Solver + solo gate
ChatGPT Base homework	Crutch; assisted≠learning (Bastani)	Guarded fade; no Base hero
“AI tutor” one SKU	Collapses rails	Labeled Help vs Practice
Shame-the-cheater moral brand	Loses parent WTP for relief	Dignity + standards; no parent stalk
Unlimited free solutions hero	Executive default	Costly reveal + FEI follow-on

Positioning line (FOUNDER BELIEF → copy test): “We help you get unstuck — then you prove you can do it without us.”

CIV.9 Confidence matrix

Claim	Label	Confidence	Notes
Homework can relate positively to achievement (esp. secondary)	FACT	High on direction; methods mixed	Cooper et al. 2006

Web solution outsourcing can break HW–exam coupling	FACT	High (course / audit grain)	DeGoede 2020; Pickering & Schuller 2024
Chegg temptation can harm later exams (high scorers)	FACT (working paper)	Medium until journal replication	Roshak AEA 2023
Unguarded generative practice can harm solo exams	FACT	High	Bastani et al. 2025 PNAS
Executive vs instrumental help-seeking is load-bearing	FACT (construct)	High	Nelson-Le Gall 1981; ITS help literature
Dual-rail Solver/Practice raises solo_transfer_pass vs dump Solver	HYPOTHESIS	Medium	Needs HWHELP-*
Homework Completion Score™ improves identity	SPECULATION (reject)	—	Kill

CIV.10 Experiments

ID	Question	Design	Primary
HWHELP-1	Instrumental Solver (fade+costly reveal) vs dump Solver on same homework items	2-arm	assisted accuracy; solo_transfer_pass 24–72h; time-to-unstuck
HWHELP-2	Help-then-Practice forced follow-on vs help-only	2-arm	help_recruit_then_solo; delayed mix; quit
HWHELP-3	Parent CBC: dual-rail copy vs “finishes homework” vs “no homework help / purity”	CBC	WTP; trust; brand fit (SAFE-WTP)
HWHELP-4	Deadline-mode soft reveal (one grain) vs hard ban on full dump	2-arm	ChatGPT escape rate; solo pass; parent NPS
HWHELP-5	Tutor HITL: homework triage → Map grain vs full re-teach	Ops A/B	FEI per tutor-minute (SAFE-ROI); dependency

HWHELP-QUAL	10 Maya + 10 parents: what “help” meant last Sunday night vs what “got better at math” meant	Qual	codebook: executive / instrumental / shame / relief
-------------	--	------	---

Falsifier: Dump Solver $\geq$ instrumental on delayed solo with equal quit $\rightarrow$ wound fade cost; still ban Completion Score™ and Chegg-clone ads.

Falsifier: Purity (no Solver) $\geq$ dual-rail on parent WTP with equal solo $\rightarrow$ still keep instrumental help for acquisition; do not shame deadline need.

Falsifier: Dual-rail copy loses CBC to “finishes homework” by large margin $\rightarrow$ iterate copy/time-to-unstuck; do not adopt dump brand.

Pre-register: HWHELP-* before “AI homework helper” splash, Completion Score™, or Solver-open accuracy as readiness claims. Densify HELP/FADE/ROI/DURABLE — do not fork a Chegg religion.

CIV.11 So what for MindCraft commercially

- **Copy:** “Unstuck without the paste — then prove it alone.” Never Homework Completion Score™ or “we do the worksheet.”
- **Product:** Labeled Help vs Practice; Solver defaults to instrumental fade; costly full reveal; optional/forced Practice follow-on grain; instrument assisted vs solo separately.
- **Positioning:** Against Chegg dumps and against purity products that abandon Sunday-night parents — MindCraft is the dual-rail FEI gym.
- **Metric:** solo_transfer_pass, help_recruit_then_solo, time-to-unstuck, ChatGPT-escape proxy — demote tonight-% / Solver NPS as learning NS.
- **Kill list:** Completion Score™; dump hero; assisted $\Rightarrow$ ready; purity ban on help; parent homework stalk; Bastani-cosplay Solver ads; grade/ACT guarantees from Solver.
- **Growth / vision:** Acquisition may start on Help; retention and brand proof live on Practice identity evidence.

CIV.12 Doctrine — SAFE-HWHELP (provisional)

1. **Two rails, two jobs** — Help = instrumental unstuck; Practice = FEI identity gym.
2. **Executive completion is not learning** — densifies SAFE-HELP / Nelson-Le Gall.
3. **Assisted $\neq$ unassisted** — always separate metrics (Bastani signature watch).
4. **Default instrumental; costly dump** — densifies SAFE-FADE / SAFE-HINT.
5. **Help episodes should recruit solo proof** — densifies SAFE-PROOF / SAFE-ROI.
6. **Parent honesty without shame** — dual-rail copy; no live stalk (SAFE-PDASH).
7. **No Chegg-clone brand / Bastani Base hero / Completion Score™.**
8. **No purity theater that bans all deadline help.**

Confidence: High that outsourcing solutions and unguarded generative help can inflate assisted scores while harming or failing to improve solo learning. Medium that MindCraft’s dual-rail UX will beat dump competitors on parent WTP and solo transfer simultaneously — instrument HWHELP-* before homework-helper splash campaigns.

References (verified)

- Aleven, V., Roll, I., McLaren, B. M., & Koedinger, K. R. (2016). Help helps, but only so much: Research on help seeking with intelligent tutoring systems. *International Journal of Artificial Intelligence in Education*, 26(1), 205–223. <https://doi.org/10.1007/s40593-015-0089-1>
 - Bastani, H., Bastani, O., Sungu, A., Ge, H., Kabakc■, Ö., & Mariman, R. (2025). Generative AI without guardrails can harm learning: Evidence from high school mathematics. *Proceedings of the National Academy of Sciences*, 122(26), e2422633122. <https://doi.org/10.1073/pnas.2422633122>
 - Cooper, H., Robinson, J. C., & Patall, E. A. (2006). Does homework improve academic achievement? A synthesis of research, 1987–2003. *Review of Educational Research*, 76(1), 1–62. <https://doi.org/10.3102/00346543076001001>
 - DeGoede, K. (2020). A Chegg® era model for HW. *2020 ASEE Virtual Annual Conference*. <https://doi.org/10.18260/1-2--33979>
 - Nelson-Le Gall, S. (1981). Help-seeking: An understudied skill in children. *Developmental Review*, 1(3), 224–246. [https://doi.org/10.1016/0273-2297\(81\)90019-8](https://doi.org/10.1016/0273-2297(81)90019-8)
 - Pickering, E., & Schuller, C. (2024). Chegg’s growth, response rate, and prevalence as a cheating tool: Insights from an audit within an Australian engineering school. *Journal of Academic Ethics*. <https://doi.org/10.1007/s10805-024-09551-6>
 - Roll, I., Aleven, V., McLaren, B. M., & Koedinger, K. R. (2011). Improving students’ help-seeking skills using metacognitive feedback in an intelligent tutoring system. *Learning and Instruction*, 21(2), 267–280. <https://doi.org/10.1016/j.learninstruc.2010.07.004>
 - Roshak, K. (2023). Technology undermining education: The case of Chegg. AEA Annual Meeting poster / working paper. <https://www.aeaweb.org/conference/2023/program/paper/ZbGk9d39>
-

Part CV — Gesture & Embodiment in Math Practice

Chapter status: Living evidence + Practice / FormatId / tutor brief — Researcher tick 2026-08-08 (UTC hour 15; hour%6≠0; researcher count since synthesizer v1.13 = 5 → Researcher)

Primary question: When should MindCraft prompt **gesture, enactment, or finger-trace** as part of Practice / coach / HITL — and when is “embodied learning” just kinesthetic-style theater or camera surveillance cosplay?

Owners: Product (Practice / FormatId / GLA draw) · Tutor playbooks · Brand · Privacy · Red Team

Commercial job: Ship **SAFE-GESTURE** densifying SAFE-DUAL / SAFE-GENERATE / SAFE-FORMAT / SAFE-APPRENTICE: strategy-specific, optional, transfer-tested motion scaffolds — never Embodiment Score™, learning-styles meshing, mandatory dance-math, or motion-capture empathy SKUs.

Builds on: Parts XXVI (CLT), XLVIII (transfer/variation), LXXXII (SAFE-FORMAT), LXXXV (story load), XCVII (SAFE-DUAL), CII (SAFE-GENERATE draw/map), XCV (SAFE-APPRENTICE modeling). Product seams: FormatId diagram/coordinate_graph grains, generative draw prompts, tutor modeling→fade, no biometric affect

cameras (SAFE-PRIVACY).

CV.1 Why this chapter exists

Gesture research is one of the few embodiment literatures with **causal math experiments** (not just philosophy). Founders and designers love the slogan “get kids moving.” Competitors ship AR overlays, kinesthetic playlists, and “whole-body math.” MindCraft’s risk is either (a) ignoring a real transfer lever on FormatId / equivalence / geometry grains, or (b) branding motion as identity while measuring engagement theater.

FOUNDER BELIEF under audit: Hands can carry strategy information speech misses — but only when the motion *represents the math idea*, not when the product forces any movement for dopamine.

Claims we refuse as doctrine:

1. Embodiment Score™ / Gesture Score™ / Motion Minutes as learning NS.
 2. Kinesthetic / “hands-on learner” meshing quizzes as placement.
 3. Always-add-gesture / mandatory enactment on every item.
 4. Manipulative action $\equiv$ gesture $\equiv$ abstraction (false; transfer differs).
 5. Camera / pose / motion-capture “embodiment AI” as empathy or mastery.
 6. AR/VR body immersion $\equiv$ FEI identity.
 7. Forced produce-gesture for spontaneous non-gesturers with no alternative scaffold.
 8. Guaranteed ACT points / “move to learn” splash as science-backed readiness.
-

CV.2 Constructs (keep motion kinds distinct)

Construct	Research meaning	MindCraft analogue	Failure mode
Gesture	Representational hand movement that can encode strategy without manipulating objects	Coach prompt: “show the groups with your hands”; tutor model	Flailing for engagement
Action on objects	Physically move tokens / tiles	Rare in SPA; optional desk tokens	Concrete lock-in; weak transfer
Finger-trace / draw	Trace figure or sketch structure	FormatId diagram grain; SAFE-GENERATE draw	Decoration without structure
Dynamic depictive gesture	Enact transformational ops (rotate, dilate, regroup)	Geometry / functions FormatId coach	Costume without language

SAFE-GESTURE	Strategy-tied, optional, transfer-tested motion scaffolds	This chapter	Motion theater SKU
---------------------	---	--------------	--------------------

Operational definition (HYPOTHESIS): A Practice/coach path is *SAFE-GESTURE compliant* when (a) any prompted motion encodes a named strategy or FormatId structure (not “wiggle”), (b) produce vs observe is chosen with learner readiness in mind, (c) transfer/solo_transfer_pass is the success metric (not motion compliance), (d) non-producers get an alternate scaffold (manipulative analogy, draw, or speech SE), (e) no camera pose scoring, and (f) marketing never claims kinesthetic-style personalization.

CV.3 Gesture can create ideas — not just express them

FACT: Goldin-Meadow, Cook, & Mitchell (2009, *Psychological Science*, 20(3), 267–272, doi:10.1111/j.1467-9280.2009.02297.x) — children required to produce **correct** gestures during a math lesson learned more than those producing partially correct gestures, who learned more than no-gesture controls; the effect was mediated by children incorporating gesture-only strategy information into later speech. Mechanism claim: learners can extract meaning from their own hand movements.

Applied (HYPOTHESIS): MindCraft coach prompts that specify a *correct strategy gesture* (e.g., grouping for equivalence) can seed speechable strategy — fluff “use your hands!” without strategy content will not.

Kill: Generic “embodied warm-up” before every mission.

Survive: Strategy-specified gesture prompts on selected grains (esp. when speech is stuck).

CV.4 Gesture can make learning last

FACT: Cook, Mitchell, & Goldin-Meadow (2008, *Cognition*, 106(2), 1047–1058, doi:10.1016/j.cognition.2007.04.010) — children who gestured while learning a math concept showed better **retention** than speech-only learners even when immediate post-lesson gains looked similar; gesture offered an embodied alternate representation that supported consolidation.

Applied (FOUNDER BELIEF → testable): For high-value concepts on the ACT path, a short produce-gesture or finger-trace step at encoding may raise delayed transfer_pass more than adding another AI sentence (densifies SAFE-EXPLAIN / SAFE-SCHED).

Kill: Motion Minutes / streak-of-gestures as retention proof.

Survive: Sparse gesture at encoding + scheduled return check (SAFE-FORGET honesty).

CV.5 Gesture beats action for transfer — action is not “more embodied = better”

FACT: Novack, Congdon, Hemani-Lopez, & Goldin-Meadow (2014, *Psychological Science*, 25(4), 903–910, doi:10.1177/0956797613518351) — third-graders taught a mathematical-equivalence strategy via (a) action on objects,

(b) concrete gesture miming that action, or (c) abstract gesture: all helped on trained problems, but **only gesture conditions** supported generalization to novel problems better than would be expected from action alone; gesture promoted transfer better than direct action on objects.

Commercial implication (HYPOTHESIS): Selling “hands-on manipulatives forever” or AR drag-tokens as the embodiment hero can **hurt** the identity thesis if students stay glued to concrete moves. Abstract/representational gesture (or structural draw) is often the better product bet for far transfer.

Kill: Manipulative theater / “move the blocks = mastery” ads.

Survive: Representational gesture or FormatId sketch that abstracts the same strategy.

CV.6 Individual differences: forced gesture can harm non-gesturers

FACT: Congdon, Wakefield, Novack, Hemani-Lopez, & Goldin-Meadow (2024, *Cognitive Science*, 48(7), e13479, doi:10.1111/cogs.13479) — whether children spontaneously gestured before a math lesson predicted who benefited from seeing vs doing gesture during instruction; for spontaneous gesturers, both seeing and doing gesture beat a comparison manipulative action on generalization/retention; for **non-spontaneous gesturers, doing gesture was less effective** than the comparison action (deficit specific to gesture — they still benefited from the manipulative action).

Applied (HYPOTHESIS): One-size mandatory “do this hand move” UX is product-hostile for a slice of Maya. Offer: observe tutor/coach gesture, produce gesture, **or** alternate structural action/draw — never force produce-only.

Kill: Forced gesture compliance gates / Gesture Score™.

Survive: Optional strategy-motion with alternate FormatId scaffold (densifies SAFE-FORMAT).

CV.7 Embodiment is real in teaching — and easy to costume

FACT (review / theoretical synthesis): Alibali & Nathan (2012, *Journal of the Learning Sciences*, 21(2), 247–286, doi:10.1080/10508406.2011.611446) — learners’ and teachers’ gestures evidence that mathematical knowledge is grounded in perception/action and environment; pointing, representational, and metaphoric gestures are instructional signals, not decoration.

FACT (geometry / proof line): Nathan & Walkington (2017, *Cognitive Research: Principles and Implications*, 2, 9, doi:10.1186/s41235-016-0040-5) — grounded and embodied mathematical cognition (GEMC): action can foster insight, but valid proof needs action coordinated with language that propositionalizes general properties. Nathan, Schenck, Vinsonhaler, Michaelis, Swart, & Walkington (2021, *Journal of Educational Psychology*, 113(5), 929–948, doi:10.1037/edu0000638) — dynamic depictive gestures (enact transformational ops) associate with geometry intuition, insight, and proof beyond speech alone, controlling for expertise and spatial ability.

Applied (HYPOTHESIS): HITL tutors should **model** strategy gestures (SAFE-APPRENTICE) then fade; coach copy for diagram FormatId can invite “show the transformation,” then require a one-sentence general claim (gesture + language), not mute waving.

Kill: AR immersion / body-tracking as proof of learning.

Survive: Tutor modeling + student optional enact + verbalize general rule.

CV.8 Link to FormatId, dual coding, and generative draw

HYPOTHESIS (product seam): Gesture and finger-trace are most load-honest on **diagram** / **coordinate_graph** / **number_line** FormatIds and on bridge grains where spatial structure is the object — not on every symbolic multiple-choice. Densifies SAFE-DUAL (structural figures) and SAFE-GENERATE (guided draw): a student-produced sketch or trace is often the SPA-native embodiment, without cameras.

Contradiction to watch: Extra motion + story wrap + figure can explode extraneous load (SAFE-STORYLOAD / CLT). Red Team: if gesture prompts raise quit or lower accuracy without delayed transfer gains, cut them.

Kill: “We teach every format with full-body VR.”

Survive: Sparse, FormatId-triggered motion/draw scaffolds with transfer checks.

CV.9 Competitive contrast (mechanism lens)

Competitor pattern	Mechanism	MindCraft counter
Kinesthetic learning-style quiz	Meshing myth	No styles placement; grain-local scaffolds
Always-move / dance-math engagement	Motion as DAU	Strategy-tied optional gesture only
Manipulative-only “hands-on” brand	Concrete lock-in	Prefer representational gesture/draw for transfer
Pose-camera “embodied AI”	Surveillance cosplay	Ban cameras; SAFE-PRIVACY
AR/VR immersion hero	Presence $\neq$ FEI	Diagram FormatId + solo proof
Tutor who only talks	Missed modeling channel	Apprentice model gesture → fade

Positioning line (FOUNDER BELIEF → copy test): “Show the idea with your hands — then prove it without the props.”

CV.10 Confidence matrix

Claim	Label	Confidence	Notes
Correct strategy gesture can improve learning via gesture→speech uptake	FACT	High (lab grain)	Goldin-Meadow et al. 2009

Gesture during encoding can improve retention vs speech-only	FACT	High (lab grain)	Cook et al. 2008
Gesture can outperform object action on transfer/generalization	FACT	High (lab grain)	Novack et al. 2014
Forced produce-gesture can hurt spontaneous non-gesturers vs alternate action	FACT	Medium–high	Congdon et al. 2024
Teacher/learner gesture indexes embodied math thinking	FACT (synthesis)	High as description	Alibali & Nathan 2012
Dynamic gesture tracks geometry insight/proof beyond speech	FACT (assoc./design)	Medium–high	Nathan/Walkington line
MindCraft FormatId gesture/draw prompts raise solo_transfer_pass	HYPOTHESIS	Medium	Needs GESTURE-*
Embodiment Score™ / styles meshing improves identity	SPECULATION (reject)	—	Kill

CV.11 Experiments

ID	Question	Design	Primary
GESTURE-1	Strategy-correct produce-gesture vs speech-only SE on equivalence / regroup grains	2-arm	immediate accuracy; delayed transfer_pass; gesture→speech uptake code
GESTURE-2	Abstract gesture vs object-action scaffold vs draw-only on same strategy	3-arm	near + far transfer; quit
GESTURE-3	Observe-coach-gesture vs produce-gesture vs choice (observe/produce/draw) stratified by pre spontaneous gesture	Factorial	generalization/retention ; harm check on non-gesturers

GESTURE-4	FormatId diagram: finger-trace/draw prompt vs redundant animated overlay	2-arm	format-hop solo_transfer_pass; load/quit (SAFE-DUAL)
GESTURE-5	Tutor HITL: model strategy gesture + fade vs speech-only model	Ops A/B	FEI per minute; student produce rate; dependency
GESTURE-QUAL	10 Maya: when did hands help vs feel silly/shameful; camera comfort	Qual	codebook: help / shame / preference observe vs do

Falsifier: Speech-only $\geq$ strategy-gesture on delayed transfer with equal quit $\rightarrow$ demote gesture prompts; keep ban on Embodiment Score™.

Falsifier: Forced produce $\geq$ choice arm for non-gesturers $\rightarrow$ still offer alternates; Congdon harm check stands as Red Team watch.

Falsifier: Draw-only $\geq$ gesture on FormatId hops $\rightarrow$ prefer SPA-native draw; treat gesture as HITL-first.

Pre-register: GESTURE-* before “embodied learning” splash, kinesthetic quiz, or camera motion SKUs. Densify DUAL/GENERATE/FORMAT/APPRENTICE — do not fork a movement religion.

CV.12 So what for MindCraft commercially

- **Copy:** “Show the structure — then solve without the props.” Never Embodiment Score™, kinesthetic-style placement, or “move to learn = ACT ready.”
- **Product:** Optional strategy-gesture / finger-trace / draw on selected FormatId and bridge grains; observe/produce/draw choice; no pose cameras; tutor model $\rightarrow$ fade.
- **Positioning:** Against learning-styles and AR immersion theater; for representational motion that earns transfer.
- **Metric:** Delayed transfer_pass / solo_transfer_pass, format-hop success, quit under motion prompts — demote Motion Minutes / gesture compliance %.
- **Kill list:** Embodiment Score™; styles meshing; always-gesture; action $\equiv$ gesture $\equiv$ abstraction; motion-capture SKU; AR body hero; forced produce-only; ACT guarantees from gesture kits.
- **Growth / vision:** Embodiment is a **scaffold channel** for identity evidence (solo proof), not a brand costume.

CV.13 Doctrine — SAFE-GESTURE (provisional)

1. **Strategy-specific motion only** — encode the math idea; no flailing theater.
2. **Gesture $\neq$ action on objects** — prefer representational gesture/draw when transfer is the goal (Novack).
3. **Retention is a reason, not a streak** — sparse encode-time gesture; check delayed return (Cook).
4. **Respect non-gesturers** — offer observe / produce / draw|action alternates (Congdon).
5. **Language closes the loop** — gesture + propositionalize (GEMC); densifies SAFE-GENERATE / SE.
6. **FormatId-triggered, load-honest** — densifies SAFE-DUAL / SAFE-FORMAT / SAFE-STORYLOAD.

7. **HITL models then fades** — densifies SAFE-APPRENTICE; no permanent motion coach.

8. **No Embodiment Score™ / styles meshing / camera pose SKUs / ACT-from-motion ads.**

Confidence: High that strategy-linked gesture can improve learning, retention, and transfer in controlled math tasks, and that action≠gesture for generalization. Medium that SPA prompts will replicate lab gains without load/quit costs — instrument GESTURE-* before embodied-learning marketing. High confidence to kill styles meshing and surveillance embodiment SKUs regardless.

References (verified)

- Alibali, M. W., & Nathan, M. J. (2012). Embodiment in mathematics teaching and learning: Evidence from learners' and teachers' gestures. *Journal of the Learning Sciences*, 21(2), 247–286. <https://doi.org/10.1080/10508406.2011.611446>
 - Congdon, E. L., Wakefield, E. M., Novack, M. A., Hemani-Lopez, N., & Goldin-Meadow, S. (2024). Learners' spontaneous gesture before a math lesson predicts the efficacy of seeing versus doing gesture during the lesson. *Cognitive Science*, 48(7), e13479. <https://doi.org/10.1111/cogs.13479>
 - Cook, S. W., Mitchell, Z., & Goldin-Meadow, S. (2008). Gesturing makes learning last. *Cognition*, 106(2), 1047–1058. <https://doi.org/10.1016/j.cognition.2007.04.010>
 - Goldin-Meadow, S., Cook, S. W., & Mitchell, Z. A. (2009). Gesturing gives children new ideas about math. *Psychological Science*, 20(3), 267–272. <https://doi.org/10.1111/j.1467-9280.2009.02297.x>
 - Nathan, M. J., & Walkington, C. (2017). Grounded and embodied mathematical cognition: Promoting mathematical insight and proof using action and language. *Cognitive Research: Principles and Implications*, 2, 9. <https://doi.org/10.1186/s41235-016-0040-5>
 - Nathan, M. J., Schenck, K. E., Vinsonhaler, R., Michaelis, J. E., Swart, M. I., & Walkington, C. (2021). Embodied geometric reasoning: Dynamic gestures during intuition, insight, and proof. *Journal of Educational Psychology*, 113(5), 929–948. <https://doi.org/10.1037/edu0000638>
 - Novack, M. A., Congdon, E. L., Hemani-Lopez, N., & Goldin-Meadow, S. (2014). From action to abstraction: Using the hands to learn math. *Psychological Science*, 25(4), 903–910. <https://doi.org/10.1177/0956797613518351>
-

Part CVI — Peer Explanation Quality (Knowledge-Building vs Telling)

Chapter status: Living evidence + Practice teach-back / HITL QA / generative explain — Researcher tick 2026-08-08 (UTC hour 18 ≡ Red Team slot, but ch106 never written → prefer Researcher per rotation; researcher count since synthesizer v1.13 = 6 → Researcher)

Primary question: When MindCraft asks a student (or near-peer tutor) to **explain**, what makes that explanation *learning-producing* vs a fluent restatement that feels like teaching but builds little?

Owners: Product (Practice teach-back / Notes / Solver) · Tutor playbooks / HITL QA · Brand · Red Team

Commercial job: Ship **SAFE-PEERX** densifying SAFE-GENERATE / SAFE-EXPLAIN / SAFE-HELP / SAFE-TUTORGRAIN / SAFE-APPRENTICE: knowledge-building explanation acts with rubrics and transfer checks — never Peer Explanation Score™, teach-for-XP, Discord-chat≡tutoring, or AI monologue≡peer teaching.

Builds on: Parts XL (SE prompts), LXXI (tutor grain), LXXXIV (talk ratio), LXXXVIII–XCI (fade/help/hint), XCII (SAFE-EXPLAIN), CII (SAFE-GENERATE teach-back), XCV (SAFE-APPRENTICE articulation). Product seams: generative teach-back prompts, tutor QA codes, soft→hard help, solo gate after explain.

CVI.1 Why this chapter exists

“Explain it to a friend,” “teach-back,” and peer tutoring are beloved growth slogans. Competitors ship Discord study rooms, AI “tutor mode,” and XP for explanations. The peer-tutoring literature is clear and awkward for marketers: **tutors often default to knowledge-telling** (summarize / recite / dump steps), and that bias caps the famous “tutor learning” benefit. MindCraft’s risk is shipping teach-back theater that measures words spoken, not knowledge built — then selling it as identity transformation.

FOUNDER BELIEF under audit: Explanation can revise identity *when* it forces monitoring, integration, and repair — not when it rehearses a script the coach already said.

Claims we refuse as doctrine:

1. Peer Explanation Score™ / Teach Minutes / word-count as learning NS.
 2. Teach-for-XP / explanation streaks as growth.
 3. Knowledge-telling monologue ≡ tutoring quality (student or AI).
 4. Discord / forum peer chat ≡ knowledge-building community.
 5. AI dump-as-teach-back or “explain like I’m 5” fluency ≡ peer teaching.
 6. Unlimited explain-to-unlock without solo transfer gate.
 7. Peer tutoring branding as Bloom-2σ or ACT-point guarantee.
 8. Length / thoroughness of peer talk as quality (conflicts SAFE-EXPLAIN).
-

CVI.2 Constructs (keep telling vs building distinct)

Construct	Research meaning	MindCraft analogue	Failure mode
Knowledge-telling	Summarize / paraphrase source with little monitoring or elaboration	Restate steps; AI wrap verbatim; “say the formula”	Fluency illusion; tutor learning underwhelms
Knowledge-building (reflective)	Self-monitor, integrate prior+new, elaborate, repair gaps while explaining / answering	Teach-back with why/how + misconception contrast; inferential Q	Over-prompt load; empty metacog slogans
Constructive use of help	Receiver applies elaborated help to solve, not just hears it	After hint/explain: attempt on <i>this</i> item before next	Passive listening; help abuse
Explanation quality (tutor→tutee)	Conceptual + procedural adequacy; strategy variety	HITL QA rubric; near-peer brief	Warm dump; “got it?” theater

SAFE-PEERX	Rubric-gated explain acts that force building, then solo proof	This chapter	Peer-chat / XP explain SKU
------------	--	--------------	----------------------------

Operational definition (HYPOTHESIS): A Practice/HITL path is *SAFE-PEERX compliant* when (a) explain prompts require at least one building move (why this step, what would go wrong, how this joins a prior concept), (b) telling-only responses are incomplete for unlock, (c) after explain, a solo or format-hop attempt is required (solo_transfer_pass / transfer check), (d) HITL QA codes building vs telling (not talk-% alone), (e) AI never substitutes its monologue for the student’s explanation act, and (f) marketing never equates peer rooms or teach-XP with FEI identity.

CVI.3 Tutor learning is real — and usually underwhelming without building

FACT (review): Roscoe & Chi (2007, *Review of Educational Research*, 77(4), 534–574, doi:10.3102/0034654307309920) — peer tutors can benefit academically from tutoring, but gains are often modest; explaining and questioning are hypothesized to help via **reflective knowledge-building** (self-monitoring, integrating new and prior knowledge, elaborating/constructing knowledge); across studies, peer tutors show a pervasive **knowledge-telling bias**, focusing on delivering knowledge rather than developing it, so the true potential for tutor learning is rarely achieved.

Applied (HYPOTHESIS): Shipping “peer teach” or “explain to unlock” without a building prompt is likely to recreate the bias Roscoe & Chi document — MindCraft would pay UX cost for theater.

Kill: “Peer tutoring = automatic deeper learning” splash.

Survive: Explicit knowledge-building scaffolds inside explain acts.

CVI.4 What elicits building: inferential questions beat summary prompts

FACT (empirical): Roscoe & Chi (2008, *Instructional Science*, 36(4), 321–350, doi:10.1007/s11251-007-9034-5) — untrained peer tutors learned most when instructional activity included reflective knowledge-building (monitor understanding, generate inferences to repair misunderstandings, elaborate on source material); tutors still showed a knowledge-telling bias (summaries with little elaboration); building was more often elicited by **tutee interaction**, especially when tutee questions contained an inference or required an inferential answer — those questions pulled elaborative, metacognitive tutor responses.

Applied (HYPOTHESIS): Teach-back UX should simulate a good tutee: ask for an inference (“why not divide both sides here?” / “what changes if the diagram flips?”) rather than “summarize the solution.” Densifies SAFE-GENERATE teach-back and SAFE-EXPLAIN (short principle + student construction).

Kill: “Explain the steps” as the only teach-back template.

Survive: Inferential prompts that force repair and elaboration.

CVI.5 Receiving help is not learning — constructive activity is

FACT: Webb, Troper, & Fall (1995, *Journal of Educational Psychology*, 87(3), 406–423, doi:10.1037/0022-0663.87.3.406) — in collaborative small groups, **constructive activity** (actively using received explanations to solve / explain) predicts learning beyond merely being told; help that is not applied to the problem at hand underperforms.

FACT (helping continuum / math groups): Webb (2002 / TIP line; see also Webb 1991 JRME review, *Journal for Research in Mathematics Education*, 22(5), 366–389, doi:10.5951/jresmetheduc.22.5.0366) — level of elaboration of help given and received, and responsiveness to need, matter; answer-only help is weak; labeled conceptual/numerical explanations plus opportunity to apply them matter for posttest success.

Applied (HYPOTHESIS): After any peer/AI/tutor explanation in MindCraft, the next beat is **student construction on this item** (or a near transfer), not a green check for having heard the speech. Densifies SAFE-HELP / SAFE-HINT / SAFE-HWHELP (instrumental help → solo proof).

Kill: Listen-complete / watch-explain unlock without attempt.

Survive: Explain → immediate apply → then delayed mix.

CVI.6 Tutor ability grain: explanation quality is not evenly distributed

FACT: Fuchs, Fuchs, Karns, Hamlett, Dutka, & Katzaroff (1996, *American Educational Research Journal*, 33(3), 631–664, doi:10.3102/00028312033003631) — after training and sustained peer-tutoring practice, **high-achieving** peer tutors produced higher-rated conceptual, procedural, and overall explanations, used more explanatory strategies, scored higher on conceptual orientation, and produced better tutee performance than medium-achieving tutors tutoring the same classmate with a math learning disability.

Applied (HYPOTHESIS): Near-peer / community explain features need **content readiness + rubric**, not vibes. Densifies SAFE-TUTORGRAIN (hire/ops bar) and SAFE-HITL (QA on explanation quality). Do not assume any peer room raises FEI.

Kill: Open Discord “everyone teach everyone” as pedagogy.

Survive: Matched explain tasks + QA codes; expert tutors for fragile joins.

CVI.7 Actually explaining beats mere expectancy for durable learning

FACT: Fiorella & Mayer (2013, *Contemporary Educational Psychology*, 38(4), 281–288, doi:10.1016/j.cedpsych.2013.06.001) — preparing to teach can help immediate comprehension vs studying for a test, but **actually teaching** (generating the explanation) is critical for longer-term retention; expectancy alone is weaker after delay.

Applied (FOUNDER BELIEF → testable): MindCraft teach-back must require the **produced explanation act** (typed/spoken student artifact), not a “prepare to teach later” checkbox or an AI that speaks for them. Densifies SAFE-GENERATE (student produces; coach does not substitute).

Kill: Expectancy theater / “you’re the tutor today” badge without explain artifact.

Survive: Captured building explanation + delayed transfer_pass.

CVI.8 Link to SE, AI monologue, and talk-ratio

HYPOTHESIS (product seam): Knowledge-building peer/self explain is the same family as Chi-style self-explanation (Part XL) and Fiorella/Mayer generative teach (Part CII) — the enemy is the same as SAFE-EXPLAIN’s kill list: long provided monologues that steal the generative move. Talk-ratio instrumentation (SAFE-TALK) without a **building vs telling code** will reward airtime, not cognition.

Contradiction to watch: Over-structured explain rubrics can balloon germane+extraneous load or create compliance scripts (students paste template sentences). Red Team: if building prompts raise quit or lower delayed transfer vs short SE, simplify the rubric; keep the ban on telling-as-quality.

Kill: Talk-% / Explanation Score™ / AI essay teach-back.

Survive: Sparse building moves + solo proof + HITL sample codes.

CVI.9 Competitive contrast (mechanism lens)

Competitor pattern	Mechanism	MindCraft counter
Discord study room / peer dump	Social presence ≠ building	Rubriced teach-back + solo gate
Teach-for-XP / explain streaks	Engagement theater	No XP for telling; FEI metrics
AI “tutor mode” monologue	Knowledge-telling at scale	Student produces; AI asks inferential Qs
Unlimited homework explain dump	Executive help without construction	SAFE-HWHELP dual-rail + apply beat
“We have peer tutors” brand	Cosplay without grain/QA	Fuchs grain + Roscoe codes in ops

Positioning line (FOUNDER BELIEF → copy test): “Don’t recite the steps — build the why, then prove it alone.”

CVI.10 Confidence matrix

Claim	Label	Confidence	Notes
Peer tutors often show knowledge-telling bias; building is rarer than slogans imply	FACT	High	Roscoe & Chi 2007

Reflective knowledge-building in tutor activity associates with better tutor learning	FACT	High	Roscoe & Chi 2007/2008
Inferential tutee questions elicit more elaborative/metacognitive tutor responses	FACT	High (lab)	Roscoe & Chi 2008
Constructive use of received help predicts learning beyond hearing help	FACT	High	Webb et al. 1995; Webb reviews
Higher-achieving peer tutors produce higher-quality math explanations and better tutee outcomes	FACT	High (trained classroom dyads)	Fuchs et al. 1996
Actually teaching beats teaching-expectancy alone for delayed retention	FACT	High (lab grain)	Fiorella & Mayer 2013
MindCraft building-gated teach-back raises solo_transfer_pass vs telling-only	HYPOTHESIS	Medium	Needs PEERX-*
Peer Explanation Score™ / Discord≡building improves identity	SPECULATION (reject)	—	Kill

CVI.11 Experiments

ID	Question	Design	Primary
PEERX-1	Building teach-back (why/how + misconception) vs telling-only (“restate steps”) unlock	2-arm	delayed transfer_pass; building-code rate; quit
PEERX-2	Inferential coach questions vs summary prompts during teach-back	2-arm	elaboration/metacog codes; tutor-or-student learning proxy; solo_transfer_pass

PEERX-3	Explain→immediate apply on same item vs explain→next item without apply	2-arm	post-item accuracy; hint binge; HELP co-gate
PEERX-4	HITL QA with building-vs-telling rubric vs talk-% / warmth-only QA	Ops A/B	FEI per tutor-minute; tutee transfer; dependency
PEERX-5	Near-peer matched explain (content-ready) vs open peer room without rubric	2-arm	explanation quality ratings; tutee outcomes; harm/quit
PEERX-QUAL	10 Maya + 5 tutors: when did “explain it” feel fake vs clarifying; template compliance	Qual	codebook: telling / building / shame / load

Falsifier: Telling-only $\geq$ building prompts on delayed transfer with equal quit → demote building rubrics; keep ban on Peer Explanation Score™ and Discord $\equiv$ pedagogy.

Falsifier: Inferential prompts raise quit without transfer gain → shorten to one building move (SAFE-EXPLAIN length discipline).

Falsifier: Open peer room $\geq$ rubriced teach-back on FEI → still require solo gate; Fuchs/Roscoe still constrain marketing claims.

Pre-register: PEERX-* before peer-tutoring splash, teach-for-XP, or “AI peer teacher” campaigns. Densify GENERATE/EXPLAIN/HELP/TUTORGRAIN/APPRENTICE — do not fork a peer-religion SKU.

CVI.12 So what for MindCraft commercially

- **Copy:** “Build the why — then prove it alone.” Never Peer Explanation Score™, teach-for-XP, or “our Discord is the tutoring.”
- **Product:** Teach-back requires ≥ 1 knowledge-building move; inferential coach questions; explain→apply beat; AI asks, student speaks; HITL codes building vs telling.
- **Positioning:** Against peer-chat theater and AI monologue-as-teaching; for explanation quality that earns transfer.
- **Metric:** Building-code rate, post-explain apply success, delayed transfer_pass / solo_transfer_pass — demote Teach Minutes / word-count / talk-%.
- **Kill list:** Peer Explanation Score™; teach-for-XP; telling $\equiv$ quality; Discord $\equiv$ building; AI dump teach-back; unlock-without-solo; peer-tutoring ACT guarantees.
- **Growth / vision:** Explanation is a **construction channel** for identity evidence, not a social or fluency costume.

CVI.13 Doctrine — SAFE-PEERX (provisional)

1. **Building over telling** — explain acts must force monitor/integrate/elaborate/repair, not recite (Roscoe & Chi).
2. **Inferential prompts** — ask why/how/what-if; do not default to “summarize the steps.”
3. **Apply after explain** — constructive use of help is the learning event (Webb).
4. **Capture the act** — student produces the explanation; expectancy badges and AI substitutes do not count (Fiorella & Mayer).
5. **Grain + QA** — near-peer explain needs readiness and building-vs-telling codes (Fuchs; densifies SAFE-TUTORGRAIN/HITL).
6. **Length stays load-honest** — densifies SAFE-EXPLAIN; no word-count quality.
7. **Solo proof closes the loop** — densifies SAFE-GENERATE / SAFE-HELP / SAFE-PROOF.
8. **No Peer Explanation Score™ / teach-for-XP / Discord≡pedagogy / AI-monologue≡peer teaching / ACT-from-peer-tutor ads.**

Confidence: High that knowledge-telling bias is the default and that building moves + constructive application are the evidenced levers. Medium that SPA teach-back rubrics will raise delayed transfer without quit — instrument PEERX-* before peer-tutoring marketing. High confidence to kill Peer Explanation Score™, teach-for-XP, and Discord≡building regardless.

References (verified)

- Fiorella, L., & Mayer, R. E. (2013). The relative benefits of learning by teaching and teaching expectancy. *Contemporary Educational Psychology*, 38(4), 281–288. <https://doi.org/10.1016/j.cedpsych.2013.06.001>
 - Fuchs, L. S., Fuchs, D., Karns, K., Hamlett, C. L., Dutka, S., & Katzaroff, M. (1996). The relation between student ability and the quality and effectiveness of explanations. *American Educational Research Journal*, 33(3), 631–664. <https://doi.org/10.3102/00028312033003631>
 - Roscoe, R. D., & Chi, M. T. H. (2007). Understanding tutor learning: Knowledge-building and knowledge-telling in peer tutors’ explanations and questions. *Review of Educational Research*, 77(4), 534–574. <https://doi.org/10.3102/0034654307309920>
 - Roscoe, R. D., & Chi, M. T. H. (2008). Tutor learning: The role of explaining and responding to questions. *Instructional Science*, 36(4), 321–350. <https://doi.org/10.1007/s11251-007-9034-5>
 - Webb, N. M. (1991). Task-related verbal interaction and mathematics learning in small groups. *Journal for Research in Mathematics Education*, 22(5), 366–389. <https://doi.org/10.5951/jresmetheduc.22.5.0366>
 - Webb, N. M., Troper, J. D., & Fall, R. (1995). Constructive activity and learning in collaborative small groups. *Journal of Educational Psychology*, 87(3), 406–423. <https://doi.org/10.1037/0022-0663.87.3.406>
-

Part CVII — Curiosity & Prediction Before Reveal

Chapter status: Living evidence + Practice soft-wrong / predict→feedback — Researcher tick 2026-08-09 (UTC hour 0 ≡ Red Team slot, but ch107 never written → prefer Researcher per rotation; researcher count since synthesizer v1.13 = 7 → Researcher)

Primary question: When should MindCraft force a **committed prediction** (or soft-wrong commit) *before* revealing correctness / coach content — and when is “curiosity” just cliffhanger entertainment that never updates knowledge?

Owners: Product (Practice / Diagnostic / Solver feedback UX) · Tutor playbooks · Brand · Red Team

Commercial job: Ship **SAFE-PREDICT** densifying SAFE-FBTIME / SAFE-MONITOR / SAFE-RETRIEVE / SAFE-MISCON / SAFE-PF: prediction commits that light curiosity and surprise *for expectancy-violating joins*, then encode the repair — never Curiosity Score™, cliffhanger-without-commit, instant-reveal entertainment, or curiosity-as-North-Star ads.

Builds on: Parts LXXIX (SAFE-FBTIME soft-wrong delay), XLIX (misconceptions as productive), XCIII (SAFE-RETRIEVE stall modes), XCVIII (SAFE-MONITOR confidence×outcome), XCIV (SAFE-PF generate→consolidate), LXXXV (story load). Product seams: soft-wrong predict→feedback, gap-scan hide-correctness, coach reveal gates, story beats that ask “what happens next?” only when math is the stake.

CVII.1 Why this chapter exists

Edtech loves cliffhangers: tease the answer, drip lore, celebrate “curiosity.” Competitors ship instant green checks, spoilers-as-help, and Curiosity Score™ dashboards. Classroom and lab literatures converge on a sharper claim: **curiosity and surprise help when the learner has a stake in an outcome** — typically by *generating an explicit prediction* before feedback. Passive watch / instant dump / post-hoc “were you surprised?” theater underdeliver.

FOUNDER BELIEF under audit: Soft-wrong and predict-before-reveal are identity technology when they create a felt knowledge gap the student then *repairs* — not when they maximize dopamine for the next swipe.

Claims we refuse as doctrine:

1. Curiosity Score™ / Engagement Score™ / self-rated curiosity as learning NS.
 2. Cliffhanger / lore tease without a committed prediction on the math outcome.
 3. Instant reveal / instant coach dump as “curiosity-friendly UX.”
 4. Post-hoc surprise rating $\equiv$ prediction-induced surprise.
 5. Prediction theater (guess for XP) without encoding the correct join.
 6. Trivia-curiosity cosplay as ACT readiness.
 7. Always-predict dogma on every item (load + freeze risk).
 8. Curiosity kits / “we spark wonder” ACT-point guarantees.
-

CVII.2 Constructs (keep deprivation, commit, and surprise distinct)

Construct	Research meaning	MindCraft analogue	Failure mode
Information-gap curiosity	Cognitively induced deprivation from a perceived knowledge gap	Soft-wrong commit; “what will the Map say?” before reveal	Trivia tease; lore FOMO
Epistemic curiosity (state)	Desire to know an answer; reward anticipation for information	Anticipation after prediction, before KR	Trait Curiosity Score™

Explicit prediction	Learner commits to an outcome before feedback	Choice / numeric / qualitative commit on this item	Silent watch; postdiction
Expectancy violation / surprise	Outcome $\neq$ prediction; arousal that can aid updating	Soft-wrong $\rightarrow$ contrast card	Shame flood; ignore update
SAFE-PREDICT	Commit $\rightarrow$ curiosity/surprise $\rightarrow$ encode repair $\rightarrow$ transfer	This chapter	Cliffhanger SKU

Operational definition (HYPOTHESIS): A Practice/Diagnostic path is *SAFE-PREDICT compliant* when (a) on learn-rail soft-wrongs and selected bridge/format probes, the student commits a prediction or answer *before* correctness/coach reveal, (b) cliffhangers without commit are banned, (c) expectancy-violating feedback is paired with a short repair encode (SE / soft-wrong contrast), not only a green/red flash, (d) curiosity is never a North Star metric, (e) predict is mode-gated (not every item; respect freeze/anxiety), and (f) marketing never sells wonder kits as exam points.

CVII.3 Curiosity is an information-gap drive — and it is transient

FACT (review): Loewenstein (1994, *Psychological Bulletin*, 116(1), 75–98, doi:10.1037/0033-2909.116.1.75) — curiosity as cognitively induced deprivation from perceiving a gap in knowledge; intense yet often transient and easily satisfied; situational determinants matter more than trait slogans.

Applied (HYPOTHESIS): Product should *open a gap the student cares about* (commit to “what will this expression simplify to?” / “which FormatId is needed?”) rather than decorate with wonder copy. Densifies Insight 6 (history frames) only when the gap is mathematical, not franchise trivia (SAFE-STORYLOAD).

Kill: Trait Curiosity Score™ / perpetual-wonder brand.

Survive: Situational gaps created by prediction commits.

CVII.4 State curiosity predicts memory for surprising answers

FACT (empirical + fMRI): Kang, Hsu, Krajbich, Loewenstein, McClure, Wang, & Camerer (2009, *Psychological Science*, 20(8), 963–973, doi:10.1111/j.1467-9280.2009.02402.x) — curiosity while reading trivia questions correlated with caudate activity (reward anticipation); higher curiosity predicted willingness to spend scarce resources for answers; curiosity enhanced memory-related activity after incorrect guesses; higher initial curiosity predicted better recall of surprising answers 1–2 weeks later.

Applied (HYPOTHESIS): Soft-wrong + delayed reveal is not cruelty theater — it can be the curiosity window Kang et al. document, *if* the student has a prediction stake. Densifies SAFE-FBTIME (micro-delay after commit) and SAFE-MISCON (productive wrong).

Kill: Instant dump that collapses the gap before encoding.

Survive: Commit → brief anticipation → informative repair.

CVII.5 Generating a prediction stimulates curiosity (vs generating an example)

FACT: Brod & Breitwieser (2019, *npj Science of Learning*, 4, Article 17, doi:10.1038/s41539-019-0056-y) — in a numerical-facts task, generating a **prediction** before the answer raised the share of high-curiosity ratings vs generating an **example**; high curiosity associated with better memory for the correct answer; pupillary data showed larger dilation during anticipation (and further enhancement when predicting).

Nuance (FACT from Brod's later review): Brod (2021, *Psychonomic Bulletin & Review*, 28, 1839–1848, doi:10.3758/s13423-021-01904-1) — predicting is more than active retrieval; benefits also track how the ensuing feedback is processed (curiosity/surprise pathways). Overall retention gaps between predict vs example can be small even when curiosity rises — so product must optimize **encoding of the correct outcome**, not curiosity ratings alone.

Applied (HYPOTHESIS): Prefer “predict the result / choose the trap answer” over “give an example of the concept” as the pre-reveal generative move when the goal is curiosity→memory for the join. Still require repair encode (SAFE-EXPLAIN length discipline).

Kill: Curiosity ratings as success; example-prompt as curiosity substitute.

Survive: Prediction commit + measured delayed recall of the repaired join.

CVII.6 Surprise needs a prior prediction — postdiction is not the same

FACT: Brod, Hasselhorn, & Bunge (2018, *Learning and Instruction*, 55, 22–31, doi:10.1016/j.learninstruc.2018.01.013) — generating predictions (vs post-hoc evaluations) facilitated geography learning; pupillometry showed expectancy-violating outcomes produced a surprise response **only when a prediction was made beforehand**, and surprise strength related positively to learning; soccer-match prediction similarly improved memory specifically for expectancy-violating events.

FACT (follow-on): Theobald & Brod / Brod line (2022, *Psychonomic Bulletin & Review*, doi:10.3758/s13423-022-02124-x) — explicitly predicting outcomes enhances learning of expectancy-violating information relative to post-hoc evaluation; U-shaped expectedness–memory relation under prediction; pupillary surprise linked to successful learning of highly unexpected outcomes.

Applied (HYPOTHESIS): Soft-wrong UX must capture the commit *before* reveal. Asking “were you surprised?” after dumping the key is postdiction cosplay. Densifies SAFE-FBTIME and SAFE-MONITOR (confidence×outcome needs a prior judgment).

Kill: Post-hoc surprise theater / reveal-first “engagement.”

Survive: Explicit prediction → expectancy-violating feedback → encode.

CVII.7 Classroom grain: predict before demo beats passive watch

FACT: Crouch, Fagen, Callan, & Mazur (2004, *American Journal of Physics*, 72(6), 835–838, doi:10.1119/1.1707018) — students who passively observe classroom demonstrations understand underlying concepts no better than students who see no demonstration; students who **predict the outcome before seeing it** show significantly greater understanding.

Applied (HYPOTHESIS): Story-world “show the magic” beats and Solver worked-example reveals that play without a student prediction inherit the passive-demo failure mode. Densifies SAFE-STORYLOAD / SAFE-FADE / SAFE-EXPTIME: prediction or attempt before full reveal.

Kill: Spectacle-first pedagogy; watch-the-demo $\equiv$ learning.

Survive: Predict (or attempt) before reveal on demos, figures, and coach dumps.

CVII.8 Mode gates: when NOT to force predict

HYPOTHESIS (product seam): Always-predict is load and anxiety risk. Gate prediction commits to: (1) learn-rail soft-wrongs and bridge/format discrimination items, (2) story/demo beats where the math outcome is the stake, (3) optional “predict the trap” on misconception-rich stems. Skip or soften when SAFE-RETRIEVE freeze is active, exam prove-rail needs immediate KR under labeled stakes (SAFE-FBTIME), or novice acquisition needs short worked-example before struggle (SAFE-EXPTIME / SAFE-PF guided default).

Contradiction to watch: Prediction commits can raise quit if shame-framed (“you were wrong”) or if every item becomes a cliffhanger tax. Red Team: if predict-before-reveal raises abandon without delayed transfer gain, narrow to soft-wrong subset; keep ban on Curiosity Score™ and cliffhanger-without-commit.

Kill: Always-predict dogma / shame-reveal.

Survive: Mode-conditional predict + destaked soft-wrong climate (SAFE-ERRCLIMATE).

CVII.9 Competitive contrast (mechanism lens)

Competitor pattern	Mechanism	MindCraft counter
Instant green check / answer dump	Collapses information gap	Commit $\rightarrow$ micro-anticipation $\rightarrow$ repair
Lore cliffhanger without math stake	Entertainment gap $\neq$ knowledge gap	Predict the <i>math</i> outcome (SAFE-STORYLOAD)
Curiosity Score™ / wonder surveys	Trait/state vanity	Delayed transfer_pass after repair
Passive demo / animation reveal	Crouch/Mazur passive fail	Predict before figure/demo play
Post-hoc “surprise!” toast	Postdiction $\neq$ prediction	Capture commit pre-reveal
Guess-for-XP	Engagement theater	No XP for guesses; FEI metrics

Positioning line (FOUNDER BELIEF $\rightarrow$ copy test): “Commit before the reveal — then keep the repair.”

CVII.10 Confidence matrix

Claim	Label	Confidence	Notes
Curiosity arises from perceived information gaps; often intense but transient	FACT	High	Loewenstein 1994
State curiosity associates with better memory for surprising answers	FACT	High (trivia grain)	Kang et al. 2009
Generating predictions raises curiosity vs generating examples	FACT	High (lab)	Brod & Breitwieser 2019
Surprise/pupillary response to expectancy violation requires prior prediction (vs postdiction)	FACT	High	Brod et al. 2018; 2022 line
Passive demo observation $\approx$ no demo; predict-before-demo improves understanding	FACT	High (physics classroom)	Crouch et al. 2004
Curiosity ratings alone are insufficient product NS (retention gaps can be small)	FACT / HYPOTHESIS	High / Medium	Brod 2021 review nuance
MindCraft predict-befo re-soft-wrong-reveal raises delayed transfer vs instant dump	HYPOTHESIS	Medium	Needs PREDICT-*
Curiosity Score™ / cliffhanger SKU improves identity	SPECULATION (reject)	—	Kill

CVII.11 Experiments

ID	Question	Design	Primary
----	----------	--------	---------

PREDICT-1	Soft-wrong: predict/commit → micro-delay → repair card vs instant KR dump	2-arm within concept	delayed transfer_pass; retry_120s; quit
PREDICT-2	Predict-before-reveal vs post-hoc “were you surprised?” after same reveal	2-arm	encoding of correct join; pupillary/attention proxy optional
PREDICT-3	Prediction commit vs “give an example” generative prompt before feedback	2-arm	curiosity proxy; delayed recall of answer; load/quit
PREDICT-4	Story/demo beat: predict math outcome before play vs passive watch	2-arm	concept check; STORYLOAD co-gate; quit
PREDICT-5	Always-predict vs soft-wrong-only predict gate	2-arm	transfer; abandon; freeze rate (RETRIEVE)
PREDICT-QUAL	10 Maya: when did cliffhangers feel manipulative vs helpful; shame after wrong predict	Qual	codebook: gap / shame / load / repair

Falsifier: Instant dump $\geq$ predict-before-reveal on delayed transfer with equal quit → demote broad predict UX; keep ban on Curiosity Score™ and cliffhanger-without-commit.

Falsifier: Predict raises quit without transfer → soft-wrong-only gate + destake copy (SAFE-ERRCLIMATE).

Falsifier: Example prompt $\geq$ prediction on curiosity *and* transfer → still prefer commit when expectancy violation is the learning event (Brod 2018); do not ship Curiosity Score™.

Pre-register: PREDICT-* before “spark curiosity” campaigns, cliffhanger SKUs, or wonder-kit ACT ads. Densify FBTIME/MONITOR/RETRIEVE/MISCON/PF — do not fork a curiosity religion.

CVII.12 So what for MindCraft commercially

- **Copy:** “Commit before the reveal — then keep the repair.” Never Curiosity Score™, wonder-kit ACT points, or lore cliffhangers without a math stake.
- **Product:** Soft-wrong and selected bridge/format items: answer/predict commit → brief anticipation → soft-wrong contrast / SE repair; demos and figures: predict before play; mode-gate always-predict.
- **Positioning:** Against instant-dump and entertainment cliffhangers; for prediction-induced curiosity and surprise that updates knowledge.

- **Metric:** Post-repair delayed transfer_pass / solo_transfer_pass, retry_120s, soft-wrong→encode success — demote curiosity Likert, cliffhanger CTR, ms-to-green.
 - **Kill list:** Curiosity Score™; Engagement-from-curiosity NS; cliffhanger-without-commit; postdiction theater; instant-reveal entertainment; guess-for-XP; always-predict dogma; curiosity ACT guarantees.
 - **Growth / vision:** Curiosity is a **temporary information-gap state** that MindCraft *earns* with commits and spends on repairs — not a brand emotion to maximize.
-

CVII.13 Doctrine — SAFE-PREDICT (provisional)

1. **Gap before reveal** — open a knowledge gap the student cares about; do not collapse it with instant dump (Loewenstein; Kang).
2. **Commit is the wick** — explicit prediction/answer before feedback; postdiction does not create the same surprise (Brod et al.).
3. **Encode the repair** — curiosity without encoding the correct join is incomplete (Brod 2021 nuance; SAFE-EXPLAIN / soft-wrong contrast).
4. **Demo/figure predict** — no passive spectacle pedagogy (Crouch et al.).
5. **Mode gates** — soft-wrong / bridge / format / story-math stakes; soften under freeze and labeled prove-rail (SAFE-RETRIEVE / SAFE-FBTIME / SAFE-EXPTIME).
6. **Destake the miss** — wrong predictions feed climate of diagnosis, not shame (SAFE-ERRCLIMATE).
7. **Densify, don't fork** — FBTIME delay, MONITOR prior judgment, MISCON productive wrong, PF generate→consolidate.
8. **No Curiosity Score™ / cliffhanger-without-commit / postdiction theater / guess-for-XP / curiosity ACT ads.**

Confidence: High that information-gap curiosity and prediction-before-feedback are real levers (Loewenstein; Kang; Brod line; Crouch/Mazur). Medium that SPA soft-wrong predict gates will raise delayed transfer without quit — instrument PREDICT-* before curiosity marketing. High confidence to kill Curiosity Score™, cliffhanger-without-commit, and postdiction theater regardless.

References (verified)

- Brod, G. (2021). Predicting as a learning strategy. *Psychonomic Bulletin & Review*, 28, 1839–1848. <https://doi.org/10.3758/s13423-021-01904-1>
- Brod, G., & Breitwieser, J. (2019). Lighting the wick in the candle of learning: Generating a prediction stimulates curiosity. *npj Science of Learning*, 4, Article 17. <https://doi.org/10.1038/s41539-019-0056-y>
- Brod, G., Hasselhorn, M., & Bunge, S. A. (2018). When generating a prediction boosts learning: The element of surprise. *Learning and Instruction*, 55, 22–31. <https://doi.org/10.1016/j.learninstruc.2018.01.013>
- Crouch, C. H., Fagen, A. P., Callan, J. P., & Mazur, E. (2004). Classroom demonstrations: Learning tools or entertainment? *American Journal of Physics*, 72(6), 835–838. <https://doi.org/10.1119/1.1707018>
- Kang, M. J., Hsu, M., Krajbich, I. M., Loewenstein, G., McClure, S. M., Wang, J. T., & Camerer, C. F. (2009). The wick in the candle of learning: Epistemic curiosity activates reward circuitry and enhances memory. *Psychological Science*, 20(8), 963–973. <https://doi.org/10.1111/j.1467-9280.2009.02402.x>

- Loewenstein, G. (1994). The psychology of curiosity: A review and reinterpretation. *Psychological Bulletin*, 116(1), 75–98. <https://doi.org/10.1037/0033-2909.116.1.75>
 - Theobald, M., & Brod, G. (2021/2022). Explicitly predicting outcomes enhances learning of expectancy-violating information. *Psychonomic Bulletin & Review*. <https://doi.org/10.3758/s13423-022-02124-x>
-

Part CVIII — Analogical Reasoning & Bridge Teaching

Chapter status: Living evidence + Map/Practice bridge pedagogy — Researcher tick 2026-08-09 (UTC hour 6 ≡ Red Team slot, but ch108 never written → prefer Researcher per rotation; researcher count since synthesizer v1.14 = 0 → Researcher)

Primary question: When should MindCraft teach *joins* (ontology bridges, FormatId hops, near-miss discrimination) as **structure-mapped analogies** — co-present source↔target, student alignment, relational encoding — versus “we use analogies” costume, story-lore surface matches, or AI “this is like...” dumps?

Owners: Product (Map / Practice / Solver contrast cards) · Engine (bridges / pathfinder) · Tutor playbooks · Brand · Red Team

Commercial job: Ship **SAFE-ANALOG** densifying SAFE-BRIDGE / SAFE-TRANSFER / SAFE-DUAL / SAFE-GESTURE / SAFE-PREDICT: relational comparison with cognitive supports as the bridge-remediation method — never Analogy Score™, analogy-mention≡pedagogy, surface-story cosplay, or mapping ACT guarantees.

Builds on: Parts LXXX (SAFE-BRIDGE), XLVIII (SAFE-TRANSFER), XXXIX (interleaving), XL (SAFE-SE), XCVII (SAFE-DUAL), CV (SAFE-GESTURE), CVII (SAFE-PREDICT), LXVII (SAFE-ONTOLOGY). Product seams: bridge cards, Map edge CTAs, side-by-side contrast, tutor mapping prompts, story wraps that carry structure not costume.

CVIII.1 Why this chapter exists

SAFE-BRIDGE already kills bridge-count ads and re-explain≡bridge. The open product question is *how* to teach the join once diagnosed. Edtech defaults: (1) mention an analogy in coach copy, (2) skin both sides with the same story franchise, (3) dump a fluent “this is just like...” monologue, (4) celebrate “higher-order thinking” without a mechanism. Cognitive science is sharper: **analogy is structural alignment** — mapping systems of relations — and classrooms routinely under-support the mapping even when they “use analogies.”

FOUNDER BELIEF under audit: MindCraft’s commercial wedge (“the gap is the connection”) only ships if Practice forces *noticed structure* across the join — not if Map draws pretty edges and chat narrates likeness.

Claims we refuse as doctrine:

1. Analogy Score™ / Mapping Minutes / Higher-Order Thinking Score™ as learning NS.
2. Analogy *mention* or story-skin match ≡ analogy *learning*.
3. Surface similarity (same characters, same numbers vibe) ≡ relational alignment.
4. AI monologue “this is like X” without student-produced mapping.
5. Always-analogy dogma on every item (load + freeze risk).

6. Cross-domain metaphor cosplay as ACT readiness.
7. Discord/forum “peer analogy share” $\equiv$ supported comparison.
8. Analogy kits / “we teach with analogies” ACT-point guarantees.

CVIII.2 Constructs (keep relation, surface, and support distinct)

Construct	Research meaning	MindCraft analogue	Failure mode
Structure-mapping	Align base $\leftrightarrow$ target by shared relational structure; systematicity favors interconnected systems	Bridge card: map roles across $A \rightarrow B$	Attribute match; story costume
Analogical encoding	Learning by comparing two cases that share structure	Side-by-side join items; explicit compare prompt	Study cases separately
Cognitive supports	Familiar source; source kept visible; spatial cues; comparative gestures; imagery	Co-present panels; highlight correspondence; gesture/draw (SAFE-GESTURE)	Sequential monologue
Relational comparison	Noticing shared structure despite surface difference	Near-miss interleave that shares surface, differs on join	Surface lure trap
SAFE-ANALOG	Supported structure-mapping as bridge pedagogy	This chapter	Analogy costume SKU

Operational definition (HYPOTHESIS): A bridge remediation is *SAFE-ANALOG compliant* when (a) source and target are co-present during comparison, (b) the student produces an alignment (what maps to what / what stays invariant) before or during coach wrap, (c) supports reduce WM load (visibility, spatial cues, optional gesture), (d) surface-lure near-misses test discrimination, (e) delayed join solo_transfer_pass is the claim metric, and (f) marketing never sells Analogy Score™ or analogy-mention as pedagogy.

CVIII.3 Structure-mapping: relations, not attributes

FACT (theory): Gentner (1983, *Cognitive Science*, 7(2), 155–170, doi:10.1207/s15516709cog0702_3) — structure-mapping: analogy interprets a base–target pair by mapping **relations among objects** (and preferring systematic higher-order relational systems), not by matching object attributes; syntax of the representation, not domain content alone, constrains the mapping.

FACT (review extension): Gentner & Markman (1997, *American Psychologist*, 52(1), 45–56, doi:10.1037/0003-066X.52.1.45) — analogy and similarity share structural alignment; comparison highlights common relational structure and can generate candidate inferences from the aligned system.

Applied (HYPOTHESIS): Bridge UX must ask for *role correspondence* (“what plays the role of the slope here?” / “which quantity is the conserved quantity across both?”) — not “these look similar” or “same story world.” Densifies SAFE-BRIDGE remediation of the *relation*.

Kill: Attribute/story match as bridge proof.

Survive: Student-named relational alignment as the learning act.

CVIII.4 Spontaneous transfer fails without noticing the structure

FACT (classic): Gick & Holyoak (1980, *Cognitive Psychology*, 12(3), 306–355, doi:10.1016/0010-0285(80)90013-4) — learners who possessed an analogous “fortress” story often failed to apply the convergence principle to Duncker’s radiation problem until a **hint** made the connection available. Structural isomorphism does not guarantee noticed transfer.

Applied (HYPOTHESIS): Map edge callouts and Practice bridge missions are the product “hint” — they must cue *search for the join*, not re-teach an endpoint. Densifies SAFE-BRIDGE / SAFE-TRANSFER; pairs with SAFE-PREDICT when the student predicts which mapping holds before reveal.

Kill: Coverage of both sides $\equiv$ transfer readiness.

Survive: Designed noticing + mapping prompts on the join.

CVIII.5 Analogical encoding: compare two cases, don’t study them apart

FACT (experiments): Gentner, Loewenstein, & Thompson (2003, *Journal of Educational Psychology*, 95(2), 393–408, doi:10.1037/0022-0663.95.2.393) — analogical encoding (drawing a comparison across examples) improved schema abstraction and transfer vs no case study; **comparing two cases beat studying the two cases separately**; increasing comparison support raised transfer in a dynamic negotiation exercise.

FACT (earlier line): Loewenstein, Thompson, & Gentner (1999, *Psychonomic Bulletin & Review*, 6(4), 586–597, doi:10.3758/BF03212967) — comparing two cases during study made transfer of a negotiation principle far more likely than presenting the same cases separately.

Applied (HYPOTHESIS): Default bridge Practice should be **side-by-side compare** (two join instances or endpoint→join pair) with an explicit “what is the same structure?” SE — not sequential node drills that happen to mention both concepts. Densifies SAFE-SE and SAFE-GENERATE (student produces the schema language).

Kill: Separate restudy of A then B as bridge remediation.

Survive: Forced compare across co-present cases as the encoding move.

CVIII.6 Classroom math: analogy mention $\neq$ supported analogy

FACT (TIMSS video): Richland, Zur, & Holyoak (2007, *Science*, 316(5828), 1128–1129, doi:10.1126/science.1142103) — U.S., Hong Kong, and Japanese eighth-grade lessons all used relational comparisons frequently, but **U.S. teachers provided fewer cognitive supports** (familiar source; visual source kept visible; spatial cues; gestures; imagery) that help students notice the mapping.

FACT (agenda): Richland, Stigler, & Holyoak (2012, *Educational Psychologist*, 47(3), 189–203, doi:10.1080/00461520.2012.667065) — many students treat math as memorized procedures rather than a relational system; refining instruction so students draw connections through **relational comparisons** is a research-backed path; U.S. classrooms under-use the supports that make comparisons work.

FACT (cuing RCT grain): Richland & McDonough (2010, *Contemporary Educational Psychology*, 35(1), 28–43, doi:10.1016/j.cedpsych.2009.09.001) — in two math contexts, providing comparative gestures and **visibly aligned** source/target during an instructional analogy (vs same analogs without such cues) improved later discrimination among potential analogs and reduced susceptibility to misleading surface features — immediately and after a one-week delay. Mere inclusion of a comparison without high support was not enough for expert-like discrimination.

Applied (HYPOTHESIS): Bridge cards and Solver soft-wrongs must: (1) co-present both representations, (2) highlight correspondence (spatial alignment / arrows / shared labels), (3) invite comparative gesture or “point what maps” (SAFE-GESTURE optional), (4) then encode with short SE — not dump a second worked example. Densifies SAFE-DUAL (contiguous structural figures) and SAFE-BRIDGE BRIDGE-4.

Kill: “We use analogies” brand without co-presence + cues.

Survive: High-support compare as the default bridge remediator.

CVIII.7 Higher-order thinking is relational reasoning — with supports

FACT (review): Richland & Simms (2015, *WIREs Cognitive Science*, 6, 177–192, doi:10.1002/wcs.1336) — analogical/relational reasoning (align, compare, map structured relationships) is a productive cognitive underpinning of “higher-order thinking”; in math/science/history, expert-like knowledge treats concepts as systems of relationships; educational analogies require **explicit support** so learners notice relational relevance, have processing resources to hold relations, and recognize both similarities and differences (where the analogy breaks).

Applied (HYPOTHESIS): Parent/school copy may say “relational thinking” only when the product instruments mapping acts and join transfer — never Higher-Order Thinking Score™. Flag where the analogy breaks (SAFE-MISCON / soft-wrong contrast) so students do not over-map attributes.

Kill: HOT Score™ / Bloom-verbs theater as analogy pedagogy.

Survive: Supported relational mapping + break-points named.

CVIII.8 Load, expertise, and when NOT to force analogy

HYPOTHESIS (product seam): Always-analogy is CLT and anxiety risk. Gate high-support compare to: (1) detected bridge gaps (SAFE-BRIDGE operational definition), (2) FormatId hops that share concept but not vessel (SAFE-FORMAT), (3) near-miss interleaves that share surface and differ on join (XXXIX). Soften when: novice

acquisition needs a short worked example before dual display (SAFE-EXPTIME), SAFE-RETRIEVE freeze is active, or dual panels raise abandon (dose supports; sequential with source still visible beats disappearing source).

Contradiction to watch: Dual displays can overwhelm; surface-rich story worlds can lure attribute mapping (SAFE-STORYLOAD). Red Team: if high-support compare raises quit without delayed join transfer, narrow to bridge-severity missions only; keep ban on Analogy Score™ and mention≡pedagogy.

Kill: Always-analogy dogma / lore-as-structure.

Survive: Severity- and mode-gated compare with load-honest supports.

CVIII.9 Competitive contrast (mechanism lens)

Competitor pattern	Mechanism	MindCraft counter
ChatGPT “this is like...” dump	Fluent attribute/story likeness	Student mapping before wrap
Khan sequential topic videos	Coverage without join encode	Co-present bridge compare
Duo skill-tree XP	Isolated skill greens	Join probes; ban XP-as-connection
Brilliant single-puzzle insight	Delight inside one frame	Keep delight; add explicit cross-item compare
Story franchise same characters	Surface match cosplay	Structure-first; story only as carrier (SAFE-STORYLOAD)
Tutor “review chapter 5 again”	Endpoint re-teach	Mapping brief on A→B

Positioning line (FOUNDER BELIEF → copy test): “We don’t just mention the connection — we make you map it.”

CVIII.10 Confidence matrix

Claim	Label	Confidence	Notes
Analogy = structure-mapping of relations (systematicity), not attribute match	FACT	High	Gentner 1983; Gentner & Markman 1997
Spontaneous analogical transfer often fails without noticing/hint	FACT	High	Gick & Holyoak 1980
Comparing two cases beats studying them separately for schema/transfer	FACT	High (negotiation grain; mechanism general)	Gentner et al. 2003; Loewenstein et al. 1999

U.S. math lessons under-provide cognitive supports for analogies	FACT	High (TIMSS video sample)	Richland et al. 2007
High cuing (gesture + visible alignment) improves later analog discrimination in math	FACT	High (lab)	Richland & McDonough 2010
Higher-order thinking productively framed as relational reasoning needing supports	FACT (review)	High	Richland & Simms 2015
MindCraft co-present bridge compare raises delayed join solo_transfer_pass vs sequential re-teach	HYPOTHESIS	Medium	Needs ANALOG- / <i>densifies</i> BRIDGE-
Analogy Score™ / mention=pedagogy improves identity	SPECULATION (reject)	—	Kill

CVIII.11 Experiments

ID	Question	Design	Primary
ANALOG-1	Co-present structure-map bridge card (align + SE) vs sequential endpoint re-teach	2-arm within bridge id	delayed join solo_transfer_pass; time-to-remediate; quit
ANALOG-2	High-support cues (visible source + spatial alignment + map prompt) vs compare-without-cues	2-arm	surface-lure errors; week-delay discrimination (Richland/McDonough grain)
ANALOG-3	Student-produced mapping before coach wrap vs AI “this is like...” monologue first	2-arm	mapping quality; transfer; EXPLAIN length co-gate
ANALOG-4	Story-skin match alone vs structure-map with story carrier	2-arm	attribute overmap rate; join transfer; STORYLOAD quit

ANALOG-5	Always-compare vs bridge-severity-gated compare	2-arm	transfer; abandon; freeze (RETRIEVE)
ANALOG-QUAL	10 Maya: when did dual panels help vs overwhelm; “I knew both but...” language	Qual	codebook: map / load / surface lure / dignity

Falsifier: Sequential re-teach $\geq$ high-support compare on delayed join transfer with equal quit $\rightarrow$ demote dual-panel default; keep ban on Analogy Score™ and mention $\equiv$ pedagogy.

Falsifier: High support raises quit without transfer $\rightarrow$ gate to severity-only; reduce panel density; keep source visible in sequential fallback.

Falsifier: AI likeness monologue $\geq$ student mapping on transfer $\rightarrow$ still require mapping act for FEI proof; do not ship Analogy Score™.

Pre-register: ANALOG- before “we teach with analogies” / higher-order-thinking campaigns. Densify BRIDGE/TRANSFER/DUAL/GESTURE/PREDICT — do not fork an analogy religion. Experiment family **ANALOG-** is distinct from company-law **BRIDGE-** (diagnosis/narrative); ANALOG owns pedagogy of the join*.

CVIII.12 So what for MindCraft commercially

- **Copy:** “We don’t just mention the connection — we make you map it.” Never Analogy Score™, HOT Score™, or “powered by analogies” ACT points.
- **Product:** Bridge missions = co-present compare + alignment prompt + short SE; Map CTA opens dual-panel join practice; Solver soft-wrong asks which correspondence broke; tutors brief on mapping, not re-explain.
- **Positioning:** Against chat likeness dumps and skill-tree coverage; for inspectable relational encoding on diagnosed bridges.
- **Metric:** Delayed join solo_transfer_pass / transfer_pass, surface-lure error rate, mapping-act completion — demote analogy Likert, Mapping Minutes, edge-count.
- **Kill list:** Analogy Score™; Mapping Minutes NS; mention $\equiv$ pedagogy; surface-story $\equiv$ structure; AI likeness dump; always-analogy dogma; Discord $\equiv$ supported compare; analogy ACT guarantees.
- **Growth / vision:** Identity as a mathematical thinker is **relational competence** — aligning structure across situations — instrumented as joins students can map, not analogies the brand can name.

CVIII.13 Doctrine — SAFE-ANALOG (provisional)

1. **Relations over attributes** — map roles and systematic structure; ban story/attribute match as proof (Gentner; Gentner & Markman).
2. **Notice the join** — design the hint; coverage $\neq$ transfer (Gick & Holyoak; SAFE-BRIDGE).
3. **Encode by compare** — co-present cases; comparing beats separate study (Gentner et al. 2003).
4. **Supports required** — source visible; spatial alignment; optional comparative gesture; student alignment act (Richland et al. 2007/2010).

5. **Name the break** — where the analogy fails; prevent overmap (Richland & Simms; SAFE-MISCON).
6. **Mode gates** — bridge/format/near-miss severity; soften under freeze and novice EXPTIME load.
7. **Densify, don't fork** — BRIDGE diagnosis, TRANSFER measurement, DUAL contiguity, GESTURE optional, PREDICT commit-before-reveal on mappings.
8. **No Analogy Score™ / mention≡pedagogy / surface-story cosplay / AI likeness dump / always-analogy dogma / analogy ACT ads.**

Confidence: High that structure-mapping and supported comparison are real levers (Gentner line; Richland line; Gick & Holyoak). Medium that SPA dual-panel bridge cards will raise delayed join transfer without quit — instrument ANALOG-* before analogy marketing. High confidence to kill Analogy Score™, mention≡pedagogy, and surface-story≡structure regardless.

References (verified)

- Gentner, D. (1983). Structure-mapping: A theoretical framework for analogy. *Cognitive Science*, 7(2), 155–170. https://doi.org/10.1207/s15516709cog0702_3
 - Gentner, D., Loewenstein, J., & Thompson, L. (2003). Learning and transfer: A general role for analogical encoding. *Journal of Educational Psychology*, 95(2), 393–408. <https://doi.org/10.1037/0022-0663.95.2.393>
 - Gentner, D., & Markman, A. B. (1997). Structure mapping in analogy and similarity. *American Psychologist*, 52(1), 45–56. <https://doi.org/10.1037/0003-066X.52.1.45>
 - Gick, M. L., & Holyoak, K. J. (1980). Analogical problem solving. *Cognitive Psychology*, 12(3), 306–355. [https://doi.org/10.1016/0010-0285\(80\)90013-4](https://doi.org/10.1016/0010-0285(80)90013-4)
 - Loewenstein, J., Thompson, L., & Gentner, D. (1999). Analogical encoding facilitates knowledge transfer in negotiation. *Psychonomic Bulletin & Review*, 6(4), 586–597. <https://doi.org/10.3758/BF03212967>
 - Richland, L. E., & McDonough, I. M. (2010). Learning by analogy: Discriminating between potential analogs. *Contemporary Educational Psychology*, 35(1), 28–43. <https://doi.org/10.1016/j.cedpsych.2009.09.001>
 - Richland, L. E., & Simms, N. (2015). Analogy, higher order thinking, and education. *WIREs Cognitive Science*, 6, 177–192. <https://doi.org/10.1002/wcs.1336>
 - Richland, L. E., Stigler, J. W., & Holyoak, K. J. (2012). Teaching the conceptual structure of mathematics. *Educational Psychologist*, 47(3), 189–203. <https://doi.org/10.1080/00461520.2012.667065>
 - Richland, L. E., Zur, O., & Holyoak, K. J. (2007). Cognitive supports for analogies in the mathematics classroom. *Science*, 316(5828), 1128–1129. <https://doi.org/10.1126/science.1142103>
-

Part CIX — Writing-to-Learn in Math Practice

Chapter status: Living evidence + Notes/Practice pedagogy — Researcher tick 2026-08-09 (UTC hour 9; hour%6≠0; researcher count since synthesizer v1.14 = 1 → Researcher)

Primary question: When should MindCraft treat **writing** (short explanation drafts, learning-protocol prompts, Notes captures) as a learning mechanism versus diary theater, essay cosplay, AI-polished Notes, or a Writing Score™ that never touches FEI?

Owners: Product (Notes / Practice / Solver soft-wrong / Field Journal) · Brand · HITL tutors · Red Team

Commercial job: Ship **SAFE-WRITE** densifying SAFE-GENERATE / SAFE-SE / SAFE-EXPLAIN / SAFE-SRL / SAFE-PEERX / SAFE-PROOF: short, prompted, cognitive+metacognitive writing tied to attempts and transfer — never journal minutes, freewrite spirituality, long-essay depth theater, or AI rewrite sold as student writing-to-learn.

Builds on: Parts XL (SAFE-SE), XCII (SAFE-EXPLAIN), CII (SAFE-GENERATE), CI (SAFE-SRL), CVI (SAFE-PEERX), LXXXVI (SAFE-PROOF), LVII (SAFE-ANNOT), CIV (SAFE-HWHELP). Product seams: Notes rail, Field Journal, soft-wrong explain box, teach-back text, post-mission protocol chips.

CIX.1 Why this chapter exists

Edtech defaults for “writing in math”: (1) open diary / reflection streaks, (2) long essays as “deeper learning,” (3) AI that rewrites the student’s Notes into fluent paragraphs, (4) word-count or Journal Minutes as engagement KPIs. Writing-across-the-curriculum lore often treats writing as *inherently* educative. The empirical record is sharper and more conditional: writing can help a little, mostly when prompts force cognitive and metacognitive strategies, treatments last long enough to matter, and assignments stay short enough not to bury content learning under composition load — and even then, early reviews found the “writing always causes learning” claim oversold.

MindCraft already ships student-produced language via SAFE-SE / SAFE-GENERATE / SAFE-PEERX. The open product question is how the **Notes** rail and Field Journal participate without becoming a scrapbook of feelings or a Chegg-adjacent essay mill.

FOUNDER BELIEF under audit: Identity as a mathematical thinker includes *authoring* concise mathematical explanations — but only if the writing reorganizes structure tied to FEI attempts, not if the brand sells journaling vibes.

Claims we refuse as doctrine: Writing Score™ / Journal Minutes / word-count NS; freewrite diary ≡ math learning; longer essays ≡ deeper learning; AI rewrite of Notes ≡ WTL; Notes scrapbook without attempt linkage; always-journal dogma; reflection streak ≡ SRL; “we journal” ACT-point guarantees.

CIX.2 Constructs (keep writing-as-learning distinct from composition theater)

Construct	Research meaning	MindCraft analogue	Failure mode
Writing-to-learn (WTL)	Writing used to promote content learning, not primarily to teach composition	Short explain / protocol chips after grain	Essay unit cosplay
Learning protocol / journal	Follow-up writing that elicits cognitive + metacognitive strategies	Post-mission Notes prompts; Field Journal cards	Mood diary
Metacognitive prompt	Cue monitoring, planning, gap-naming	“What still confuses you about the join?”	“How do you feel?”

Cognitive prompt	Cue organize, elaborate, example, contrast	“State the invariant in your own words”	Verbatim copy of coach
SAFE-WRITE	Prompted short WTL under FEI	This chapter	Journal costume SKU

Operational definition (HYPOTHESIS): A Notes/Practice path is *SAFE-WRITE compliant* when (a) writing asks are short and prompted (cognitive and/or metacognitive), (b) the student produces language that cannot be satisfied by pasting the coach wrap, (c) the artifact links to a concept/bridge/FormatId attempt, (d) success is measured by delayed solo_transfer_pass / explanation quality gates — not Journal Minutes, and (e) AI may *critique or cue* but never substitute authorship for the learning claim.

CIX.3 Writing can be a mode of learning — not a magic mode

FACT (foundational claim): Emig (1977, *College Composition and Communication*, 28(2), 122–128, doi:10.2307/356095) — writing uniquely supports certain learning strategies as process and product (originating and graphically recording a verbal construct), contrasting writing with listening, reading, and talking; a theoretical case for writing as a heuristic mode of learning, not an RCT.

FACT (contradiction / caution): Ackerman (1993, *Written Communication*, 10(3), 334–370, doi:10.1177/0741088393010003002) — reading of ~35 empirical studies did **not** deliver the long-sought validation that writing inherently promotes learning; typical “writing as unique learning tool” accounts ignore classroom/cultural/media confounds; writing aids learning mainly when situationally supported and valued within disciplinary instruction — not as a freestanding mantra.

Applied (HYPOTHESIS): Brand may say students *write to understand* only when product writing is discipline-tied (math structure, joins, misconceptions) and instrumented — never “we believe in writing” lifestyle ads.

Kill: Writing=learning slogan without task design.

Survive: Emig’s uniqueness claim as motivation; Ackerman’s situational constraint as product law.

CIX.4 Meta-analytic grain: small gains, short prompts, metacognition, duration

FACT (meta-analysis): Bangert-Drowns, Hurley, & Wilkinson (2004, *Review of Educational Research*, 74(1), 29–58, doi:10.3102/00346543074001029) — 48 school-based WTL programs: writing can have a **small, positive** impact on conventional achievement measures; **enhanced** by metacognitive prompts and longer treatment duration; **reduced** by implementation in Grades 6–8 and by **longer writing assignments**.

Applied (HYPOTHESIS): MindCraft default = micro-writes (1–3 sentences / structured chips), repeated across weeks of Practice, with metacognitive gap-naming — not weekly five-paragraph “math essays.” Middle-school Maya cohorts need extra load honesty (pair with SAFE-STORYLOAD / SAFE-EXPLAIN length gates).

Kill: Longer=deeper essay theater; one-shot journal day as WTL science.

Survive: Short + metacognitive + sustained as the WTL dose recipe.

CIX.5 Mechanisms matter more than “did they write?”

FACT (review): Klein (1999, *Educational Psychology Review*, 11(3), 203–270, doi:10.1023/A:1006205501054) — reframes WTL from “does writing cause learning?” to **which cognitive processes** differentiate helpful vs inert writing; clusters theories into spontaneous generation at utterance, forward search (review externalized text → inferences), genre-based organization, and backward search (goals/subgoals / knowledge-transforming); evidence uneven; genre hypotheses best tested at the time; mechanisms are complementary, not mutually exclusive.

Applied (HYPOTHESIS): Notes UX should force forward search (reread own soft-wrong draft → revise one claim) and genre micro-moves (contrast two FormatIds; name a misconception; map a bridge) — not free utterance alone. Densifies SAFE-GENERATE SOI and SAFE-SRL reflection→control.

Kill: Any writing box ≡ WTL.

Survive: Prompted processes that transform knowledge representation.

CIX.6 Learning protocols: cognitive + metacognitive prompts beat naive journals

FACT (experiment): Berthold, Nückles, & Renkl (2007, *Learning and Instruction*, 17(5), 566–579, doi:10.1016/j.learninstruc.2007.09.003) — cognitive prompts, or combined cognitive+metacognitive prompts, improved learning vs no prompts; pure metacognitive prompts alone did not improve outcomes; effects mediated by cognitive learning strategies.

FACT (experiment): Nückles, Hübner, & Renkl (2009, *Learning and Instruction*, 19(3), 259–271, doi:10.1016/j.learninstruc.2008.05.002) — learning protocols as follow-up coursework to foster self-regulated learning; prompts activate cognitive and metacognitive strategies students often fail to apply spontaneously when journaling “naively.”

FACT (math journals): Glogger, Schwonke, Holzäpfel, Nückles, & Renkl (2012, *Journal of Educational Psychology*, 104(2), 452–468, doi:10.1037/a0026683) — ninth-grade mathematics learning journals over six weeks (N=236): **quality and quantity of cognitive strategies** predicted learning outcomes controlling for prior knowledge; learners combining cognitive **plus** metacognitive strategies were particularly successful; using mainly one type (or hardly any) performed poorly. Biology replication (N=144) essentially agreed.

Applied (HYPOTHESIS): Field Journal / Notes chips default to **paired** prompts: (cognitive) “Restate the rule/invariant / give a near-miss contrast” + (metacognitive) “What would you try differently on a similar join?” Ban pure vibe-check metacognition without cognitive content. Instrument strategy quality proxies (key idea coverage, misconception named) not Journal Minutes.

Kill: Mood/reflection-only Notes; quantity-only word count.

Survive: Prompted cognitive+metacognitive micro-protocols linked to math content.

CIX.7 Math-specific: written process descriptions can outperform oral for success

FACT: Pugalee (2001, *School Science and Mathematics*, 101(5), 236–245, doi:10.1111/j.1949-8594.2001.tb18026.x) — ninth-grade algebra students’ written descriptions of problem-solving showed a metacognitive framework across orientation, organization, execution, and verification phases; supports writing as a window on (and support for) mathematical metacognition.

FACT: Pugalee (2004, *Educational Studies in Mathematics*, 55(1–3), 27–47, doi:10.1023/B:EDUC.0000017666.33870.71) — same grain: students who **wrote** descriptions of their thinking were significantly more successful on problem-solving tasks than those who **verbalized** ($p < 0.05$); writing associated with metacognitive behaviors that support solving.

Applied (HYPOTHESIS): Prefer typed micro-explains in Practice soft-wrong / Notes over optional voice-only “talk it out” when the commercial claim is writing-to-learn — voice may still serve SAFE-TALK / HITL, but do not equate oral fluency with the WTL mechanism. Densifies SAFE-SE without forking a separate religion.

Contradiction to watch: Writing load + ELL language load (SAFE-ELL); offer structured sentence frames, not essay demands. Red Team if writing arm raises quit among ELL without transfer gain.

Kill: Voice dump $\equiv$ WTL proof.

Survive: Short written process/explanation as a math metacognition tool.

CIX.8 Notes rail vs FEI gym — dual-rail honesty

HYPOTHESIS (product seam): Notes is a **capture + protocol** surface, not a feelings scrapbook. Dual-rail with Practice (echo SAFE-HWHELP): writing after an attempt consolidates; writing instead of attempting is avoidance. Default chips: post-grain cognitive summarize/contrast; post-bridge name-the-join; weekly Field Journal 2–4 prompted protocols — not daily diary streaks. Soften under SAFE-RETRIEVE freeze or exposure destake.

Kill: Notes-as-identity home without Practice proof.

Survive: Notes as prompted encoding attached to Map-visible attempts; witness via SAFE-PROOF dated explains — not Journal Minutes dashboards.

CIX.9 Competitive contrast (mechanism lens)

Competitor pattern	Mechanism	MindCraft counter
Duo / streak journals	Habit theater	No Journal Minutes NS; FEI after write
ChatGPT “polish my notes”	Fluency substitute	Student draft required; AI cue not author
Khan optional notebook explain	Untimed freewrite	Short cognitive+metacognitive chips

Chegg / essay mills	Outsourced authorship	Ban AI-authored Notes as learning proof
---------------------	-----------------------	---

Positioning line (FOUNDER BELIEF → copy test): “We don’t journal for vibes — we write short explanations that prove you own the idea.”

CIX.10 Confidence matrix

Claim	Label	Confidence	Notes
Writing can uniquely support learning strategies as process/product	FACT (theory)	High (classic)	Emig 1977
Writing does not automatically cause learning; context/task matter	FACT (review)	High	Ackerman 1993
School WTL average effect small; metacognitive prompts + duration help; long assignments hurt	FACT	High	Bangert-Drowns et al. 2004
WTL effects depend on cognitive processes engaged	FACT (review)	High	Klein 1999
Cognitive or cognitive +metacognitive prompts > no prompts; metacognitive-only weak	FACT	High	Berthold et al. 2007; Nückles et al. 2009
Math journal strategy quality (+ cognitive×metacognitive combo) predicts outcomes	FACT	High (school)	Glogger et al. 2012
Written process description associated with higher solve success vs oral in algebra sample	FACT	Medium–High (small N)	Pugalee 2004
MindCraft micro-write chips raise delayed transfer vs no-write or long essay	HYPOTHESIS	Medium	Needs WRITE-*

Writing Score™ / Journal Minutes improve identity	SPECULATION (reject)	—	Kill
---	-------------------------	---	------

CIX.11 Experiments

ID	Question	Design	Primary
WRITE-1	Short cognitive+metacognitive protocol after grain vs no write vs long freewrite	3-arm	delayed solo_transfer_pass; time-on-task; quit
WRITE-2	Paired cognitive+metacognitive prompts vs metacognitive-only vs cognitive-only	3-arm	explanation quality; transfer (Berthold grain)
WRITE-3	Student-authored Notes vs AI-polished Notes (same prompt)	2-arm	transfer; authorship integrity; EXPLAIN co-gate
WRITE-4	Written micro-explain vs oral think-aloud on soft-wrong	2-arm	solve success; ELL quit; Pugalee grain
WRITE-5	Weekly 3 prompted protocols vs daily diary streak UX	2-arm	transfer; Journal Minutes vanity; SRL plan writeback
WRITE-QUAL	10 Maya: when writing clarified vs exhausted; Notes vs Practice confusion	Qual	codebook: clarify / load / avoidance / dignity

Falsifier: No-write $\geq$ prompted micro-write on delayed transfer with equal quit $\rightarrow$ demote default write chips; keep ban on Journal Minutes and AI-as-author.

Falsifier: Long freewrite $\geq$ short prompted on transfer $\rightarrow$ still prefer short on Bangert-Drowns dose; do not ship essay theater.

Falsifier: AI-polish $\geq$ student draft on transfer $\rightarrow$ still require student draft for FEI/PROOF claims; ban polish-as-WTL ads.

Pre-register: WRITE- before “we journal” / writing-to-learn brand campaigns. *Densify GENERATE/SE/EXPLAIN/SRL/PEERX/PROOF — do not fork a Notes religion. Experiment family WRITE- owns Notes/protocol pedagogy (distinct from GENACT- GLA timing and SE- self-explanation mechanics).*

CIX.12 So what for MindCraft commercially

- **Copy:** “We don’t journal for vibes — we write short explanations that prove you own the idea.” Never Writing Score™, Journal Minutes, or ACT points from notebooks.
 - **Product:** Notes = prompted micro-protocols + soft-wrong explain boxes; Field Journal weekly chips; AI may nudge structure, never ghostwrite learning claims; link every write to Map node/bridge.
 - **Positioning:** Against streak diaries and ChatGPT polish; for inspectable student-authored mathematical language under FEI.
 - **Metric:** Delayed transfer, explanation quality gates, protocol completion with cognitive content — demote word count, Journal Minutes, reflection streaks.
 - **Kill list:** Writing Score™; Journal Minutes NS; freewrite≡learning; longer≡deeper; AI rewrite≡WTL; Notes scrapbook without FEI; always-journal dogma; writing ACT guarantees.
 - **Growth / vision:** Identity includes becoming someone who can *author* a concise mathematical account — instrumented as dated explain artifacts parents/tutors can witness (SAFE-PROOF), not as a journaling lifestyle brand.
-

CIX.13 Doctrine — SAFE-WRITE (provisional)

1. **Conditional, not magical** — writing helps when situationally designed; ban writing=learning slogans (Emig tempered by Ackerman).
2. **Short + metacognitive + sustained** — micro-writes; metacognitive prompts; multi-week dose; ban long-essay depth theater (Bangert-Drowns).
3. **Process over presence** — force forward-search / genre micro-moves / knowledge-transforming prompts (Klein).
4. **Cognitive + metacognitive pairing** — content organize/elaborate plus monitor/plan; metacognition-only is weak (Berthold; Nückles; Gloger).
5. **Math-authored language** — written process/explain as metacognitive tool; structured frames for ELL (Pugalee; SAFE-ELL).
6. **Notes dual-rail** — protocols attach to Practice attempts; writing≠avoidance; witness via PROOF not Journal Minutes.
7. **AI may cue, not author** — polish/ghostwrite void the WTL claim.
8. **Densify, don’t fork** — GENERATE / SE / EXPLAIN / SRL / PEERX / PROOF.
9. **No Writing Score™ / Journal Minutes NS / freewrite spirituality / longer≡deeper / AI≡WTL / always-journal / writing ACT ads.**

Confidence: High that WTL is conditional and that short prompted cognitive+metacognitive protocols are the evidence-backed pattern. Medium that SPA Notes chips will raise delayed transfer without quit — instrument WRITE-* before journaling marketing. High confidence to kill Writing Score™, Journal Minutes, freewrite≡learning, and AI-as-author regardless.

References (verified)

- Ackerman, J. M. (1993). The promise of writing to learn. *Written Communication*, 10(3), 334–370. <https://doi.org/10.1177/0741088393010003002>
- Bangert-Drowns, R. L., Hurley, M. M., & Wilkinson, B. (2004). The effects of school-based writing-to-learn interventions on academic achievement: A meta-analysis. *Review of Educational Research*, 74(1), 29–58.

<https://doi.org/10.3102/00346543074001029>

- Berthold, K., Nückles, M., & Renkl, A. (2007). Do learning protocols support learning strategies and outcomes? The role of cognitive and metacognitive prompts. *Learning and Instruction*, 17(5), 566–579. <https://doi.org/10.1016/j.learninstruc.2007.09.003>
 - Emig, J. (1977). Writing as a mode of learning. *College Composition and Communication*, 28(2), 122–128. <https://doi.org/10.2307/356095>
 - Glogger, I., Schwonke, R., Holzäpfel, L., Nückles, M., & Renkl, A. (2012). Learning strategies assessed by journal writing: Prediction of learning outcomes by quantity, quality, and combinations of learning strategies. *Journal of Educational Psychology*, 104(2), 452–468. <https://doi.org/10.1037/a0026683>
 - Klein, P. D. (1999). Reopening inquiry into cognitive processes in writing-to-learn. *Educational Psychology Review*, 11(3), 203–270. <https://doi.org/10.1023/A:1006205501054>
 - Nückles, M., Hübner, S., & Renkl, A. (2009). Enhancing self-regulated learning by writing learning protocols. *Learning and Instruction*, 19(3), 259–271. <https://doi.org/10.1016/j.learninstruc.2008.05.002>
 - Pugalee, D. K. (2001). Writing, mathematics, and metacognition: Looking for connections through students’ work in mathematical problem solving. *School Science and Mathematics*, 101(5), 236–245. <https://doi.org/10.1111/j.1949-8594.2001.tb18026.x>
 - Pugalee, D. K. (2004). A comparison of verbal and written descriptions of students’ problem solving processes. *Educational Studies in Mathematics*, 55(1–3), 27–47. <https://doi.org/10.1023/B:EDUC.0000017666.33870.71>
-

Part CX — Spatial Skills & Math Identity

Chapter status: Living evidence + FormatId/Map pedagogy — Researcher tick 2026-08-09 (UTC hour 12 ≡ Red Team slot, but ch110 never written → prefer Researcher per rotation; researcher count since synthesizer v1.14 = 2 → Researcher)

Primary question: When should MindCraft treat **spatial skill** / **spatial anxiety** as part of mathematical identity work — via FormatId hops (diagram, coordinate graph, number line), conversion practice, and destaked exposure — versus shipping a Spatial Score™, mental-rotation brain-training SKU, innate “visual learner” branding, or boys-are-spatial costume?

Owners: Product (Practice / Map / FormatId CTA) · Bank & generation (diagram/graph items) · Brand · HITL tutors · Equity (SAFE-GENDER) · Red Team

Commercial job: Ship **SAFE-SPATIAL** densifying SAFE-FORMAT / SAFE-DUAL / SAFE-GESTURE / SAFE-TRANSFER / SAFE-GENDER: instrument spatial *as representation competence inside math*, not as a separate gym; never sell Spatial Score™, isolated brain-training minutes, or ACT guarantees from mental-rotation apps.

Builds on: Parts LXXXII (SAFE-FORMAT), XCVII (SAFE-DUAL), CV (SAFE-GESTURE), XLVIII (SAFE-TRANSFER), LXV (SAFE-GENDER), XXIV (anxiety/WM), XLI (SAFE-DD), LXXXI (SAFE-COLD). Product seams: FormatId gaps, format-aware CTA, Layer-4 representation_profile, diagram/coordinate_graph/number_line bank coverage.

CX.1 Why this chapter exists

Edtech spatial defaults: (1) “visual learner” quizzes, (2) isolated mental-rotation games sold as STEM prep, (3) pretty 3D chrome without structural encoding, (4) Spatial Score™ / Mental Rotation Minutes as engagement KPIs, (5) boy-coded “spatial giftedness” marketing that launders stereotype. The research record is sharper: spatial and math skills correlate moderately; spatial skills are malleable; spatial *training* can transfer to math under conditions; transfer is **not automatic**; and **spatial anxiety** is a real construct that can mediate gender gaps in math anxiety — which means product design that floods students with unscaffolded diagrams can *hurt* identity while claiming to help.

MindCraft already owns the commercial surface: FormatId diagnosis and conversion practice. The open question is whether to fork a “Spatial Gym” or densify format/dual-coding doctrine with spatial evidence and anxiety honesty.

FOUNDER BELIEF under audit: Becoming a mathematical thinker includes becoming someone who can *read and produce* spatial vessels of math ideas — graphs, figures, number lines — under FEI, not someone who maxes a rotation mini-game.

Claims we refuse as doctrine: Spatial Score™ / Mental Rotation Minutes NS; isolated brain-training $\equiv$ math identity; spatial-as-innate / boys-better costume; always-add-3D / decoration $\equiv$ spatial; Spatial Personality™ / learning-style meshing; unscaffolded diagram flood; spatial ACT-point guarantees.

CX.2 Constructs (keep spatial-as-math distinct from brain-training theater)

Construct	Research meaning	MindCraft analogue	Failure mode
Spatial skill	Mental rotation, visualization, scaling, perspective, etc.	Format competence on diagram/graph/number_line; conversion	Generic “visual IQ”
Spatial–math link	Correlational + causal training evidence	Format gap severity; transfer hops	Correlation ads as causation
Spatial training	Practice that improves spatial performance	Supported FormatId remediation	Isolated rotation arcade
Near / far transfer	Gains on trained spatial vs untrained math	Format hop $\rightarrow$ delayed math item	Near-only theater
Spatial anxiety	Anxiety about evaluative spatial tasks	Soften diagram/graph cold-start; exposure ladder	Anxiety Score™ NS
SAFE-SPATIAL	Spatial inside FormatId+FEI, not a gym SKU	This chapter	Spatial cosplay ads

Operational definition (HYPOTHESIS): A Practice/Map path is *SAFE-SPATIAL compliant* when (a) spatial demand is named as FormatId / conversion work (not a personality trait), (b) remediation keeps math content co-present with structural figures (SAFE-DUAL), (c) spatial anxiety / freeze modes get destake + scaffold before challenge (SAFE-DD / SAFE-EXPOSE grain), (d) success is measured by delayed format-hop solo_transfer_pass — not Mental Rotation Minutes, and (e) marketing never claims innate spatial gift or ACT points from spatial mini-games.

CX.3 Spatial and math co-vary — moderately, uniquely, not as destiny

FACT (meta-analysis): Atit, Power, Pigott, Lee, Geer, Uttal, Ganley, & Sorby (2022, *Psychonomic Bulletin & Review*, 29(3), 699–720, doi:10.3758/s13423-021-02012-w) — synthesizing 45 articles: positive **moderate** association between spatial and mathematical skills ($r = .36$); no significant moderation by gender or grade-level in the association magnitude; fluid reasoning and verbal skills **mediate** part of the link, yet a **unique** spatial–math relation remains.

FACT (longitudinal STEM): Wai, Lubinski, & Benbow (2009, *Journal of Educational Psychology*, 101(4), 817–835, doi:10.1037/a0016127) — spatial ability predicts STEM educational and occupational outcomes beyond verbal/math measures in large talent-search and Project TALENT alignments; spatial is a neglected but load-bearing dimension for STEM pathways — **predictive**, not a license for “boys are spatial” product copy (pair with SAFE-GENDER).

Applied (HYPOTHESIS): Parent/tutor language may say “the graph format is the bottleneck” when FormatId evidence shows it — never “she’s not a spatial person.” Correlation is real; trait branding is banned.

Kill: Spatial Personality™; innate visual-learner destiny.

Survive: Moderate unique spatial–math link as reason to instrument FormatId hops.

CX.4 Spatial skills are malleable — training works on spatial tasks

FACT (meta-analysis): Uttal, Meadow, Tipton, Hand, Alden, Warren, & Newcombe (2013, *Psychological Bulletin*, 139(2), 352–402, doi:10.1037/a0028446) — 217 training studies: average training effect vs control **Hedges’s $g \approx 0.47$** ; effects durable across delays; transfer to other **spatial** tasks observed; typology along intrinsic/extrinsic $\times$ static/dynamic axes. Spatially enriched education is a serious policy bet — **for spatial skill**, not automatically for ACT math.

Applied (HYPOTHESIS): MindCraft may improve diagram/graph competence via deliberate FormatId practice; do not invent a separate Spatial Score™ product line that never touches ontology concepts.

Kill: “Spatial is fixed talent” copy.

Survive: Malleability $\rightarrow$ practice design, not destiny ads.

CX.5 Causal transfer to math is real — and conditional

FACT (meta-analysis of training $\rightarrow$ math): Hawes, Gilligan-Lee, & Mix (2022, *Developmental Psychology*, 58(1), 112–137, doi:10.1037/dev0001281) — 29 controlled pre–post studies ($N = 3,765$; $k = 89$): spatial training improves math relative to controls (**Hedges’s $g = .28$**) and improves spatial thinking (**$g = .49$**); moderators include **age** (larger math transfer as age rises ~ 3 – 20), **concrete manipulatives** $>$ pure computerized training in the meta, and **near vs far** transfer alignment (larger when math outcomes closer to trained spatial content). Mechanisms of transfer remain poorly understood; publication-bias checks nonsignificant in their analyses.

FACT (commentary framing): Hawes, Gilligan-Lee, & Mix (2022, *npj Science of Learning*, 7, article 10, doi:10.1038/s41539-022-00128-9) — spatialize curricula carefully; summarize the $g \approx .28$ transfer meta; caution WEIRD samples and researcher-developed outcomes may inflate practical claims.

FACT (positive child RCT grain): Gilligan, Thomas, & Farran (2020, *Developmental Science*, 23(4), e12909, doi:10.1111/desc.12909; online 2019) — N = 250 eight-year-olds: training spatial scaling and mental rotation produced near, intermediate, and **far transfer to mathematics**; explicit vs implicit instruction largely similar except geometry.

FACT (elementary grades): Mix, Levine, Cheng, Stockton, & Bower (2021, *Journal of Educational Psychology*, 113(2), 304–314, doi:10.1037/edu0000494) — first- and sixth-graders: spatial visualization or form-perception/VSWM training beat language-arts control on overall mathematics.

FACT (null / wound — transfer not guaranteed): Hawes, Moss, Caswell, & Poliszczuk (2015, *Trends in Neuroscience and Education*, 4(3), 60–68, doi:10.1016/j.tine.2015.05.001) — classroom computerized mental-rotation training (6 weeks, ages 6–8) improved spatial measures but **failed far transfer** to mathematics vs literacy control. Meta-average ≠ product guarantee.

Applied (HYPOTHESIS): Prefer **embedded** spatial demand inside FormatId missions (co-present equation+graph, structural finger-trace — SAFE-GESTURE) over isolated rotation arcade. Expect larger transfer when math outcomes are representation-aligned than distal symbolic drill. HS ACT Maya is not a preschool spatial RCT.

Kill: “Our spatial mini-game raises ACT scores” without SPATIAL-* / FORMAT transfer proof.

Survive: Conditional causal path spatial→math under aligned, often concrete/embedded designs.

CX.6 Spatial anxiety — affect is part of the identity story

FACT (scale): Lyons, Ramirez, Maloney, Rendina, Levine, & Beilock (2018, *Journal of Numerical Cognition*, 4(3), 526–553, doi:10.5964/jnc.v4i3.154) — validated adult Spatial Anxiety Questionnaire with subscales **navigation, mental-manipulation, imagery**; higher anxiety uniquely predicts lower performance/attitudes in corresponding spatial domains for navigation and manipulation; framed as a STEM participation barrier via avoidance of spatial practice.

FACT (gender × math anxiety mediation): Sokolowski, Hawes, & Lyons (2019, *Cognition*, 182, 193–212, doi:10.1016/j.cognition.2018.10.005) — undergraduate sample: processes in the **spatial** domain (not math-domain ability/anxiety alone) mediated sex differences in math anxiety after controls; **spatial anxiety** was the strongest mediator vs actual/perceived spatial ability — raising the possibility that gendered math-anxiety gaps partly reflect anxiety about / avoidance of spatial strategies on math tasks.

Applied (HYPOTHESIS): Diagram / coordinate_graph cold-starts need SAFE-COLD humility + destake + SAFE-DUAL supports — not day-one “prove you’re visual.” Do not ship Spatial Anxiety Score™ as NS; instrument quit/freeze on high-spatial FormatIds. Densifies SAFE-GENDER: remove stereotype costume, keep representation challenge with scaffolds.

Kill: Flood-with-graphs as equity; Spatial Anxiety Score™ NS; “girls need easier visuals.”

Survive: Spatial anxiety as a reason for destaked FormatId ladders and strategy choice (not avoidance forever).

CX.7 Product seam — FormatId is the spatial surface

HYPOTHESIS (product law): MindCraft’s spatial bet is **not** a brain-training vertical. It is:

1. Diagnose format gaps on diagram / coordinate_graph / number_line (SAFE-FORMAT).
2. Remediate with structural dual coding and conversion (SAFE-DUAL), optional gesture/trace (SAFE-GESTURE).
3. Measure delayed hops and solo transfer — not Mental Rotation Minutes.
4. Keep story chrome from pretending to be spatial instruction (SAFE-STORYLOAD coherence).
5. Bank honesty: unverified figures fail SAFE-GENQ × SAFE-SPATIAL together.

Kill: Spatial Gym SKU that never updates Map mastery.

Survive: Spatial as named FormatId competence under FEI.

CX.8 Competitive contrast (mechanism lens)

Competitor pattern	Mechanism	MindCraft counter
Brain-training / rotation apps	Near spatial gains; weak far-transfer theater	Embedded FormatId + math transfer metric
“Visual learner” quizzes	Style meshing myth	Format gap diagnosis, not personality
Pretty 3D / AR STEM kits	Seductive immersion	Structural figures only; no AR≡FEI
ChatGPT text-only tutors	Skip spatial vessels	Map CTA for diagram/graph gaps

Positioning line (FOUNDER BELIEF → copy test): “We don’t train abstract ‘spatial IQ’ — we make the graph and the equation talk until you can move between them.”

CX.9 Confidence matrix

Claim	Label	Confidence	Notes
Spatial and math skills correlate moderately (~r=.36); unique link remains after mediators	FACT	High	Atit et al. 2022
Spatial ability predicts STEM pathways beyond verbal/math	FACT	High (talent samples)	Wai et al. 2009
Spatial skills malleable; training g≈0.47 on spatial outcomes	FACT	High	Uttal et al. 2013

Spatial training can raise math (meta $g \approx .28$); moderated by age, materials, transfer distance	FACT	High (avg); Medium (mechanism)	Hawes et al. 2022
Far transfer demonstrated in some RCTs (e.g., age-8 scaling/rotation)	FACT	Medium–High	Gilligan et al. 2020
Far transfer fails in other designs (classroom computer rotation)	FACT	High (as caution)	Hawes et al. 2015
Spatial anxiety is measurable; can mediate gender gaps in math anxiety	FACT	Medium–High (adult samples)	Lyons et al. 2018; Sokolowski et al. 2019
FormatId-embedded spatial remediation > isolated rotation arcade for FEI	HYPOTHESIS	Medium	Needs SPATIAL-*
Spatial Score™ / Mental Rotation Minutes improve identity	SPECULATION (reject)	—	Kill

CX.10 Experiments

ID	Question	Design	Primary
SPATIAL-1	FormatId diagram/graph remediation vs isolated mental-rotation mini-game vs business-as-usual	3-arm	delayed format-hop solo_transfer_pass; Spatial Minutes vanity
SPATIAL-2	Concrete/manipulative or finger-trace supports vs pure screen figures on same items	2-arm	transfer; quit (Hawes meta concrete moderator)
SPATIAL-3	Destaked spatial FormatId ladder vs day-one hard graph flood	2-arm	freeze/quit; later transfer (anxiety grain)

SPATIAL-4	Co-present equation+graph conversion vs graph-only drill	2-arm	conversion accuracy; SAFE-DUAL co-gate
SPATIAL-5	Spatial-anxiety soft gate (opt-in easier vessel first) vs no gate on high-spatial items	2-arm	attempt rate; delayed hop; no Spatial Anxiety Score™
SPATIAL-QUAL	10 Maya: when graphs felt like math vs “I’m not spatial”; stereotype language	Qual	codebook: vessel / trait / avoidance / dignity

Falsifier: Isolated rotation arcade $\geq$ FormatId-embedded arm on delayed math format-hop with equal quit $\rightarrow$ reconsider embedding priority; still ban Spatial Score™ ads.

Falsifier: Hard graph flood $\geq$ destaked ladder on transfer without raising quit $\rightarrow$ still prefer dignity under SAFE-GENDER/EXPOSE; do not celebrate flood.

Falsifier: No unique FormatId spatial signal after controlling concept mastery $\rightarrow$ demote spatial narrative to subordinate FORMAT copy.

Pre-register: SPATIAL- before “spatial STEM IQ” / mental-rotation brand campaigns. Densify FORMAT/DUAL/GESTURE/TRANSFER/GENDER — do not fork a Spatial Gym religion. Experiment family **SPATIAL-** owns spatial-as-FormatId pedagogy (distinct from company-law FORMAT- diagnosis narrative and DUAL- figure design).

CX.11 So what for MindCraft commercially

- **Copy:** “We don’t train abstract spatial IQ — we make the graph and the equation talk until you can move between them.” Never Spatial Score™, Mental Rotation Minutes, boys-are-spatial, or ACT points from mini-games.
- **Product:** Spatial demand lives inside FormatId CTAs; support conversion with dual coding + optional gesture; destake high-spatial cold starts; bank figures must be structural and keyed.
- **Positioning:** Against brain-training apps and visual-learner quizzes; for inspectable format-flexible competence under FEL.
- **Metric:** Delayed format-hop transfer, conversion accuracy, quit/freeze on spatial vessels — demote Spatial Minutes and Spatial Anxiety Score™.
- **Kill list:** Spatial Score™; Mental Rotation Minutes NS; isolated brain-training $\equiv$ identity; innate/boys-better spatial costume; always-add-3D; Spatial Personality™; unscaffolded diagram flood; spatial ACT guarantees.
- **Growth / vision:** Identity includes becoming someone who can *use spatial vessels of math* without shame — instrumented as FormatId hops parents/tutors can witness (SAFE-PROOF), not as a spatial-talent lifestyle brand.

CX.12 Doctrine — SAFE-SPATIAL (provisional)

1. **Correlate ≠ costume** — moderate unique spatial–math link justifies FormatId work; bans Spatial Personality™ and innate gift ads (Atit; Wai).
2. **Malleable, not magical** — train spatial demand inside math practice; ban fixed-talent copy (Uttal).
3. **Transfer is conditional** — meta $g \approx .28$ is real and moderated; nulls exist; no ACT guarantees from spatial apps (Hawes meta; Hawes 2015 wound).
4. **Embed over isolate** — Prefer FormatId/conversion + concrete/structural supports over rotation arcade (Gilligan; Mix; meta concrete moderator).
5. **Anxiety is part of the stack** — spatial anxiety can block strategy use; destake + scaffold; no Spatial Anxiety Score™ NS (Lyons; Sokolowski).
6. **FormatId is the surface** — densify SAFE-FORMAT/DUAL/GESTURE/TRANSFER/GENDER; no Spatial Gym fork.
7. **Figures must encode** — decoration ≠ spatial instruction; GENQ keys required.
8. **Equity without softening standards** — remove stereotype costume; keep challenge with supports (SAFE-GENDER).
9. **No Spatial Score™ / Mental Rotation Minutes NS / brain-training≡identity / boys-better / always-3D / Spatial Personality™ / spatial ACT ads.**

Confidence: High that spatial–math links and malleability are real, and that transfer is conditional. Medium that FormatId-embedded remediation will beat isolated arcade on delayed FEI transfer — instrument SPATIAL-* before spatial-STEM marketing. High confidence to kill Spatial Score™, Mental Rotation Minutes, innate/boys-better costume, and ACT guarantees from mini-games regardless.

References (verified)

- Atit, K., Power, J. R., Pigott, T., Lee, J., Geer, E. A., Uttal, D. H., Ganley, C. M., & Sorby, S. A. (2022). Examining the relations between spatial skills and mathematical performance: A meta-analysis. *Psychonomic Bulletin & Review*, 29(3), 699–720. <https://doi.org/10.3758/s13423-021-02012-w>
- Gilligan, K. A., Thomas, M. S. C., & Farran, E. K. (2020). First demonstration of effective spatial training for near transfer to spatial performance and far transfer to a range of mathematics skills at 8 years. *Developmental Science*, 23(4), e12909. <https://doi.org/10.1111/desc.12909>
- Hawes, Z., Moss, J., Caswell, B., & Poliszczuk, D. (2015). Effects of mental rotation training on children’s spatial and mathematics performance: A randomized controlled study. *Trends in Neuroscience and Education*, 4(3), 60–68. <https://doi.org/10.1016/j.tine.2015.05.001>
- Hawes, Z. C. K., Gilligan-Lee, K. A., & Mix, K. S. (2022). Effects of spatial training on mathematics performance: A meta-analysis. *Developmental Psychology*, 58(1), 112–137. <https://doi.org/10.1037/dev0001281>
- Hawes, Z. C. K., Gilligan-Lee, K. A., & Mix, K. S. (2022). Spatial thinking as the missing piece in mathematics curricula. *npj Science of Learning*, 7, 10. <https://doi.org/10.1038/s41539-022-00128-9>
- Lyons, I. M., Ramirez, G., Maloney, E. A., Rendina, D. N., Levine, S. C., & Beilock, S. L. (2018). Spatial anxiety: A novel questionnaire with subscales for measuring three aspects of spatial anxiety. *Journal of Numerical Cognition*, 4(3), 526–553. <https://doi.org/10.5964/jnc.v4i3.154>
- Mix, K. S., Levine, S. C., Cheng, Y.-L., Stockton, J. D., & Bower, C. (2021). Effects of spatial training on mathematics in first and sixth grade children. *Journal of Educational Psychology*, 113(2), 304–314. <https://doi.org/10.1037/edu0000494>
- Sokolowski, H. M., Hawes, Z., & Lyons, I. M. (2019). What explains sex differences in math anxiety? A closer look at the role of spatial processing. *Cognition*, 182, 193–212. <https://doi.org/10.1016/j.cognition.2018.10.005>

- Uttal, D. H., Meadow, N. G., Tipton, E., Hand, L. L., Alden, A. R., Warren, C., & Newcombe, N. S. (2013). The malleability of spatial skills: A meta-analysis of training studies. *Psychological Bulletin*, 139(2), 352–402. <https://doi.org/10.1037/a0028446>
 - Wai, J., Lubinski, D., & Benbow, C. P. (2009). Spatial ability for STEM domains: Aligning over 50 years of cumulative psychological knowledge solidifies its importance. *Journal of Educational Psychology*, 101(4), 817–835. <https://doi.org/10.1037/a0016127>
-

Part CXI — Formative Assessment Cycles in Tutor Sessions

Chapter status: Living evidence + HITL session pedagogy — Researcher tick 2026-08-09 (UTC hour 15; hour%≠0; researcher count since synthesizer v1.14 = 3 → Researcher)

Primary question: When should MindCraft treat **formative assessment cycles** (elicit evidence → interpret → adjust next move → close the gap) as the grain of human tutoring and in-session Practice — versus shipping a Formative Score™, exit-ticket theater, gradebook cosplay, or “science-backed AfL” ads that recycle Black & Wiliam’s popularized 0.4–0.7 effect-size legend?

Owners: HITL ops / tutor playbooks · Product (Practice / Map / session wrap) · Brand · Analytics (FEI instrumentation) · Red Team

Commercial job: Ship **SAFE-FA** densifying SAFE-HITL / SAFE-AAR / SAFE-FBTIME / SAFE-MONITOR / SAFE-SRL / SAFE-INSTRUMENT: instrument *evidence-use cycles* that change the next attempt, not Formative Score™ dashboards or exit-ticket completion NS.

Builds on: Parts LXVIII (SAFE-HITL), LIV (SAFE-AAR), LXXIX (SAFE-FBTIME), XCVIII (SAFE-MONITOR), CI (SAFE-SRL), XC (SAFE-INSTRUMENT), L (SAFE-CALIB), LXXXI (SAFE-COLD), XLIX (SAFE-MISCON). Product seams: tutor session structure, soft-wrong feedback, Map next-route, gap-scan humility, FEI North Stars.

CXI.1 Why this chapter exists

Edtech defaults: (1) “formative assessment” as a marketing badge, (2) exit tickets / thumbs as proof of pedagogy, (3) Formative Score™ as an engagement KPI, (4) recycling Black & Wiliam’s popularized 0.4–0.7 effect-size legend as a product guarantee, (5) confusing *more quizzes* with *cycles that change instruction*. The research spine is sharper: formative assessment is a **process** of using evidence to improve learning; effects are real but **heterogeneous** and often smaller than legend; domain-specific design and teacher skill matter; feedback that does not cause a next action is “dangling data.”

MindCraft already owns pieces of the stack (soft-wrong, Map diagnosis, AAR-style debrief, FEI metrics). The open commercial question is whether to sell AfL cosplay or to operationalize **minute-by-minute cycles** that tutors and product can run — and falsify.

FOUNDER BELIEF under audit: Becoming a mathematical thinker includes becoming someone whose *next attempt* is shaped by evidence — not someone who completes exit tickets for points.

Claims we refuse as doctrine: Formative Score™ / Assessment Minutes NS; exit-ticket completion ≡ learning; Black–Wiliam 0.4–0.7 as ACT guarantee; more quizzes ≡ formative; gradebook cosplay; “got it?” ≡ cycle closed; AfL

ads without domain-specific evidence use.

CXI.2 Constructs (keep cycle ≠ quiz theater)

Construct	Research meaning	MindCraft analogue	Failure mode
Formative assessment	Evidence elicited, interpreted, used to adjust teaching/learning	Session cycle: probe → soft-wrong/Map → next attempt	Quiz volume theater
Summative	Judgment of achievement for reporting	ACT practice tests; parent proof artifacts	Confusing live grades with FEI
Exit ticket	End-of-lesson evidence sample	Optional 1-item wrap tied to next mission	Ticket completion NS
Learning intentions / success criteria	Shared destination + quality standard	Mission target concept + “what correct looks like”	Vague “do your best”
Feedback that moves forward	Causes thinking + next action	Soft-wrong → retry_120s / challenge_accept	Praise / dump without action
SAFE-FA	Cycle fidelity under FEI, not AfL branding	This chapter	Formative Score™ ads

Operational definition (HYPOTHESIS): A tutor session or Practice wrap is *SAFE-FA compliant* when (a) a learning intention is named at grain fine enough to falsify (concept/ingredient/FormatId — not “math”), (b) evidence is elicited *during* work (not only end-of-hour quizzes), (c) interpretation names a specific next move (retry, fade step, bridge hop, destake), (d) the student takes that next action in-session or within a short scheduled return, and (e) success is measured by FEI events (retry_120s, challenge_accept, delayed solo_transfer_pass) — never Formative Score™ or exit-ticket completion rate.

CXI.3 Black & Wiliam — process claim, not a product effect-size SKU

FACT (full review): Black & Wiliam (1998a, *Assessment in Education: Principles, Policy & Practice*, 5(1), 7–74, doi:10.1080/0969595980050102) — comprehensive review of classroom formative assessment (~250 sources). Innovations strengthening frequent feedback can yield substantial gains; student self-assessment and teacher strategy matter. They **declined a single amalgamation meta-analysis** because study differences made pooling misleading — a caution often erased in marketing retellings.

FACT (popular summary): Black & Wiliam (1998b, *Phi Delta Kappan*, 80(2), 139–148) — *Inside the Black Box* popularizes effect-size language often quoted as **0.4–0.7**. This is the citation path most edtech decks abuse.

FACT (definitional spine): Wiliam (2011, *Embedded Formative Assessment*, Solution Tree) — formative assessment uses evidence of achievement to adjust next instructional steps so those steps are better than they would have been without the evidence; five strategies: clarify intentions/success criteria; elicit evidence; feedback that moves learning forward; peers as resources; students as owners of learning.

Applied (HYPOTHESIS): Cite Black & Wiliam for the **process** (evidence → adjust → close gap), never as “+0.5σ from exit tickets.” MindCraft map: intentions = mission target + success criteria; elicit = attempt + soft-wrong + Map probes; feedback = SAFE-FBTIME; peers = SAFE-PEERX when allowed; ownership = SAFE-SRL / FEI.

Kill: 0.4–0.7 as product guarantee / ACT points from “we do formative assessment.”

Survive: Formative-as-process; five strategies as playbook checklist, not a Formative Score™.

CXI.4 Sadler — feedback must close a gap, or it is dangling data

FACT: Sadler (1989, *Instructional Science*, 18(2), 119–144, doi:10.1007/BF00117714) — formative assessment theory for qualitative judgments: feedback is only formative when the learner (1) understands the standard, (2) compares current performance to that standard, and (3) takes action to close the gap. Without evaluative skill and tactics, teacher comments become **dangling data**. Goal of instructional systems: transition from external feedback toward learner self-monitoring.

Applied (HYPOTHESIS): Tutor talk that ends at “good effort” or AI monologue that never forces a next attempt fails Sadler’s third condition. Soft-wrong + required retry densifies Sadler; unlimited dump fails it. Exit tickets that never change the next mission are dangling data with a timestamp.

Kill: Ticket-without-action; praise fog ≡ formative.

Survive: Standard + comparison + next move as the cycle unit.

CXI.5 Effect-size honesty — legend vs Kingston–Nash vs uncertainty

FACT (meta wound on legend): Kingston & Nash (2011, *Educational Measurement: Issues and Practice*, 30(4), 28–37, doi:10.1111/j.1745-3992.2011.00220.x) — reviewed 300+ K–12 formative-assessment efficacy studies; only 13 yielded usable effect sizes (42 independent ES). Median ES **0.25**; random-effects mean **0.20**. Moderators: ELA ≈0.32, mathematics ≈0.17, science ≈0.09; PD-based and computer-based implementations ≈0.30 / 0.28. Explicitly challenges the oft-claimed 0.40–0.70 range.

FACT (meta commentary / further uncertainty): Briggs, Ruiz-Primo, Furtak, Shepard, & Yin (2012, *Educational Measurement: Issues and Practice*, 31(4), 56–63, doi:10.1111/j.1745-3992.2012.00252.x) — methodological critique of Kingston & Nash: inclusion decisions and outcome measures leave **considerable uncertainty** about the true average effect; do not treat 0.20 as the new dogma either. (Note: Black & Wiliam themselves warned against amalgamation.)

FACT (definition / claim critique): Bennett (2011, *Assessment in Education: Principles, Policy & Practice*, 18(1), 5–25, doi:10.1080/0969594X.2010.513678) — “formative assessment” is not yet a well-defined set of artefacts; quantitative effectiveness claims are often suspect (untraceable/flawed/dated sources); domain specificity, measurement principles, and teacher support are under-specified; maximum benefit requires domain-rooted designs

inside a coherent system.

Applied (HYPOTHESIS): Commercial honesty: formative cycles can help; **math effects are not magic**; MindCraft must specify *what evidence* and *what adjustment* in the ontology/FEI system — Bennett’s domain point is a gift to the knowledge-graph lane, not a license for AfL cosplay.

Kill: Formative Score™; “AfL raises achievement 0.7 σ ” ads; quiz-count NS.

Survive: Modest, heterogeneous, domain-dependent gains → design for *specified cycles*, instrument FEI, pre-register FA-*.

CXI.6 Feedback power is real — and level-conditional (bridge to SAFE-FBTIME)

FACT: Hattie & Timperley (2007, *Review of Educational Research*, 77(1), 81–112, doi:10.3102/003465430298487) — feedback among powerful influences but highly variable; effective feedback answers where going / how going / where next; levels: task, process, self-regulation, self — **self/praise** often weak or harmful relative to task/process/SR.

FACT (updated meta): Wisniewski, Zierer, & Hattie (2020, *Frontiers in Psychology*, 10, 3087, doi:10.3389/fpsyg.2019.03087) — 435 studies, k=994: overall feedback effect on learning **d** $\approx$ **0.48** with large heterogeneity; information content matters; cognitive/motor outcomes > motivational/behavioral on average.

Applied (HYPOTHESIS): Session cycles should prefer task/process/SR feedback that schedules a next attempt (SAFE-FBTIME), not ego toasts. Densifies SAFE-EXPLAIN / SAFE-HINT: short, actionable, contingent.

Kill: Praise-as-formative; longer=better feedback essays.

Survive: Where-to-next feedback as the third Sadler condition in product language.

CXI.7 Exit tickets vs FEI — when a wrap helps, when it is theater

HYPOTHESIS (product law): Exit tickets are allowed only as **evidence samples that change the next move** (Map route, mission level, destake, FormatId hop). They are banned as:

1. Completion NS / XP for ticket submit.
2. Substitute for in-session elicitation (ticket-only “formative”).
3. Parent-facing gradebook cosplay.
4. Calm/mood Formative Score™.

Contrast with FEI North Stars (SAFE-INSTRUMENT): `retry_120s`, `challenge_accept`, `transfer_pass`, `solo_transfer_pass` already measure whether evidence changed behavior. A ticket that never predicts those events is dangling data.

Kill: Exit-ticket completion as learning KPI.

Survive: Optional 1-item wrap **wired** to next Practice path + tutor brief.

CXI.8 Tutor session spine — Black/Wiliam grain without classroom cosplay

HYPOTHESIS (HITL playbook): A 45-min session (SAFE-DP) runs SAFE-FA as:

1. **Intention (2 min):** Name concept/FormatId + success criterion (“solo retry without hard peek”).
2. **Elicit (ongoing):** Student attempts first; probes diagnose misconception/bridge/format (SAFE-MISCON / SAFE-BRIDGE / SAFE-FORMAT).
3. **Interpret (spoken aloud briefly):** One specific bottleneck — not a lecture.
4. **Adjust + close gap (bulk of session):** Contingent hint/fade; required retry; delayed transfer probe if time.
5. **Ownership handoff (3–5 min):** Student states next move; product schedules return (SAFE-SRL / SAFE-AAR light).

Kill: Explain-first brand; “got it?” as closed cycle; talk-% as formative proof (SAFE-TALK).

Survive: Evidence-use fidelity checklist for QA (SAFE-HITL).

CXI.9 Competitive contrast (mechanism lens)

Competitor pattern	Mechanism	MindCraft counter
AfL poster / “we do formative” badge	Brand without cycle fidelity	Instrumented evidence→next-attempt
Exit-ticket XP apps	Completion theater	Ticket only if wired to Map/FEI
Instant green-check AI	Feedback≠cycle	Soft-wrong + retry gates
Gradebook parent portals	Summative cosplay	Proof-age honesty (SAFE-PDASH)
ChatGPT monologue tutors	Dangling explanations	Sadler third condition = attempt

Positioning line (FOUNDER BELIEF → copy test): “We don’t collect exit tickets for points — we use what you just showed us to change the next problem you face.”

CXI.10 Confidence matrix

Claim	Label	Confidence	Notes
Formative assessment as evidence-used-to-adjust is the right construct (not quiz volume)	FACT / FOUNDER BELIEF	High	Black & Wiliam; Wiliam 2011
Popular 0.4–0.7 effect-size range is not a safe product claim	FACT	High	Kingston & Nash; Bennett; Briggs et al.

Rigorous average ES ~0.20 with heterogeneity; math smaller in KN sample	FACT	Medium–High	Kingston & Nash 2011; Briggs uncertainty
Sadler’s three conditions (standard, compare, act) are load-bearing	FACT (theory) / HYPOTHESIS (product)	High / Medium	Sadler 1989
Feedback helps on average (~d.48) but is highly variable by information content	FACT	High	Wisniewski et al. 2020; Hattie & Timperley 2007
Exit-ticket completion NS improves FEI transfer	SPECULATION (reject)	—	Kill
SAFE-FA session spine beats explain-first on delayed solo transfer	HYPOTHESIS	Medium	Needs FA-*

CXI.11 Experiments

ID	Question	Design	Primary
FA-1	SAFE-FA session spine vs explain-first matched content	2-arm HITL	delayed solo_transfer_pass; talk-dump vanity
FA-2	Exit ticket wired to next Map mission vs ticket+XP vs no ticket	3-arm	next-session attempt quality; ticket completion vanity
FA-3	Sadler-complete feedback (standard+compare+required action) vs dangling praise/dump	2-arm	retry_120s; soft-wrong uptake
FA-4	In-session elicit+adjust vs end-only quiz wrap (same minutes)	2-arm	transfer; Formative Minutes vanity
FA-5	Tutor QA checklist on cycle fidelity vs hours-booked coaching only	2-arm ops	FEI event rates; hours NS

FA-QUAL	10 Maya + 10 tutors: when feedback felt usable vs dangling; exit-ticket dignity	Qual	codebook: standard / compare / act / theater
---------	--	------	---

Falsifier: Explain-first $\geq$ SAFE-FA spine on delayed solo transfer with equal quit $\rightarrow$ revise spine timing; still ban Formative Score™ and 0.4–0.7 ads.

Falsifier: Ticket+XP $\geq$ wired ticket on next-mission quality $\rightarrow$ still demote XP; investigate wiring bugs before celebrating tickets.

Falsifier: No FEI lift from cycle checklist after PD $\rightarrow$ demote AfL language to subordinate HITL craft; keep Bennett domain-specificity agenda.

Pre-register: FA- before “science-backed formative assessment” brand campaigns. *Densify HITL/AAR/FBTIME/MONITOR/SRL/INSTRUMENT — do not fork a Formative Score™ religion. Experiment family FA-* owns formative-cycle fidelity (distinct from company-law MONITOR- calibration and FBTIME- delay clocks).

CXI.12 So what for MindCraft commercially

- **Copy:** “We don’t collect exit tickets for points — we use what you just showed us to change the next problem you face.” Never Formative Score™, 0.4–0.7 σ guarantees, or AfL badge without cycle proof.
- **Product:** Session wrap and tutor playbooks encode intention $\rightarrow$ elicit $\rightarrow$ interpret $\rightarrow$ adjust $\rightarrow$ ownership; exit tickets only if wired to Map/next mission; soft-wrong must demand action.
- **Positioning:** Against quiz-XP apps and monologue AI; for inspectable evidence-use under FEI.
- **Metric:** Cycle-linked FEI events and delayed transfer — demote ticket completion, Assessment Minutes, Formative Score™.
- **Kill list:** Formative Score™; Assessment Minutes NS; exit-ticket=learning; Black–Wiliam 0.4–0.7 product ads; more-quizzes=formative; gradebook cosplay; “got it?”=closed cycle; AfL ACT guarantees.
- **Growth / vision:** Identity includes becoming someone who *uses evidence on themselves* (Sadler self-monitoring) — tutors and product as temporary external cycles that fade toward ownership (SAFE-SRL), not as endless formative surveillance.

CXI.13 Doctrine — SAFE-FA (provisional)

1. **Process, not badge** — formative = evidence used to improve next steps (Black & Wiliam; Wiliam 2011).
2. **No legend σ** — ban 0.4–0.7 / “AfL raises ACT” marketing (Kingston & Nash; Bennett; Briggs et al.).
3. **Sadler completeness** — standard + compare + act; dangling data fails (Sadler 1989).
4. **Domain-specific cycles** — ontology/FormatId/bridge grain, not generic AfL (Bennett).
5. **Feedback that moves** — task/process/SR + where-next; densify SAFE-FBTIME/EXPLAIN/HINT (Hattie & Timperley; Wisniewski et al.).
6. **Exit tickets are optional wires** — never completion NS; must change next mission.
7. **FEI over Formative Score™** — instrument retry/challenge/transfer, not Assessment Minutes.
8. **HITL fidelity** — tutor QA on cycle steps; ban explain-first and “got it?” theater (SAFE-HITL/TALK).

9. No Formative Score™ / Assessment Minutes NS / ticket=learning / 0.4–0.7 ads / quiz-volume=formative / gradebook cosplay / AfL ACT guarantees.

Confidence: High that formative-as-process and Sadler’s gap-closing conditions are the right commercial spine, and that popularized 0.4–0.7 claims are unsafe. Medium that the SAFE-FA session spine will beat explain-first on delayed FEI transfer — instrument FA-* before AfL marketing. High confidence to kill Formative Score™, exit-ticket completion NS, and effect-size legend ads regardless.

References (verified)

- Bennett, R. E. (2011). Formative assessment: A critical review. *Assessment in Education: Principles, Policy & Practice*, 18(1), 5–25. <https://doi.org/10.1080/0969594X.2010.513678>
 - Black, P., & Wiliam, D. (1998a). Assessment and classroom learning. *Assessment in Education: Principles, Policy & Practice*, 5(1), 7–74. <https://doi.org/10.1080/0969595980050102>
 - Black, P., & Wiliam, D. (1998b). Inside the black box: Raising standards through classroom assessment. *Phi Delta Kappan*, 80(2), 139–148.
 - Briggs, D. C., Ruiz-Primo, M. A., Furtak, E., Shepard, L., & Yin, Y. (2012). Meta-analytic methodology and inferences about the efficacy of formative assessment. *Educational Measurement: Issues and Practice*, 31(4), 56–63. <https://doi.org/10.1111/j.1745-3992.2012.00251.x>
 - Hattie, J., & Timperley, H. (2007). The power of feedback. *Review of Educational Research*, 77(1), 81–112. <https://doi.org/10.3102/003465430298487>
 - Kingston, N., & Nash, B. (2011). Formative assessment: A meta-analysis and a call for research. *Educational Measurement: Issues and Practice*, 30(4), 28–37. <https://doi.org/10.1111/j.1745-3992.2011.00220.x>
 - Sadler, D. R. (1989). Formative assessment and the design of instructional systems. *Instructional Science*, 18(2), 119–144. <https://doi.org/10.1007/BF00117714>
 - Wiliam, D. (2011). *Embedded formative assessment*. Solution Tree Press.
 - Wisniewski, B., Zierer, K., & Hattie, J. (2020). The power of feedback revisited: A meta-analysis of educational feedback research. *Frontiers in Psychology*, 10, 3087. <https://doi.org/10.3389/fpsyg.2019.03087>
-

Part CXII — Feedback Literacy: How Students Use Coach Feedback

Chapter status: Living evidence + Practice/HITL UX — Researcher tick 2026-08-09 (UTC hour 18 ≡ Red Team slot, but ch112 never written → prefer Researcher per rotation; researcher count since synthesizer v1.14 = 4 → Researcher)

Primary question: When should MindCraft treat **student feedback literacy** (appreciating feedback, making judgments, managing affect, *taking action*) as a first-class product capability — versus shipping Feedback Literacy Score™ dashboards, longer AI essays, thumbs-up≡uptake theater, or “we give great feedback” ads that confuse *delivery* with *use*?

Owners: Product (Practice / Solver / Notes) · HITL · Brand · Analytics · Red Team

Commercial job: Ship **SAFE-FBLIT** densifying SAFE-FBTIME / SAFE-FA / SAFE-EXPLAIN / SAFE-HINT / SAFE-HELP / SAFE-SRL / SAFE-MONITOR / SAFE-ERRCLIMATE: instrument *uptake acts* that change the next attempt, not Feedback Minutes or literacy badges.

Builds on: Parts LXXIX (SAFE-FBTIME), CXI (SAFE-FA), XCII (SAFE-EXPLAIN), XCI (SAFE-HINT), LXXXIX (SAFE-HELP), CI (SAFE-SRL), XCVIII (SAFE-MONITOR), CIII (SAFE-ERRCLIMATE), L (SAFE-CALIB). Product seams: soft-wrong gates, coach length, hints, tutor debrief.

CXII.1 Why this chapter exists

Edtech defaults: (1) “personalized feedback” as brand badge, (2) longer comments as quality, (3) thumbs-up as proof of learning, (4) Feedback Literacy Score™ as a KPI, (5) treating non-uptake as student deficit rather than design failure. Research spine: feedback works when receivers make sense of it and **use** it; barriers of awareness, decoding, agency, and volition are common; classic meta-analysis finds over one-third of feedback interventions *hurt* when attention is pulled to the self.

MindCraft already owns timing (SAFE-FBTIME), cycle completeness (SAFE-FA), and length discipline (SAFE-EXPLAIN). Open commercial question: does coach/tutor feedback *land*, and can we productize **uptake** without a Literacy Score™ religion?

FOUNDER BELIEF under audit: Becoming a mathematical thinker includes becoming someone who *uses* diagnostic information on themselves — not someone who collects green checks and praise paragraphs.

Claims we refuse as doctrine: Feedback Literacy Score™ / Feedback Minutes NS; delivery≡uptake; thumbs-up≡used; longer≡more literate; illiteracy-shame without design responsibility; always-more-feedback dogma; feedback-literacy ACT guarantees.

CXII.2 Constructs (keep literacy ≠ survey theater)

Construct	Research meaning	MindCraft analogue	Failure mode
Student feedback literacy	Understandings, capacities, dispositions to make sense of information and use it to enhance work/strategies	Soft-wrong → decode → act (retry / FormatId hop / destake)	Literacy Score™
Proactive recipience	Agentic engagement with receiving/using feedback (not passive inbox)	Required next attempt after coach; seek soft hint before hard dump	Thumbs-up theater
Appreciating feedback	Seeing feedback as useful process, not only grades	Soft-wrong framed as map-of-gap, not verdict	Grade-only parent stalk
Making judgments	Comparing work to standards / exemplars	Success criteria + student compare chip	AI rewrite for student
Managing affect	Not collapsing into ego threat / dismiss	Destake + short task/process feedback	Praise fog / shame reel

Taking action	Doing something different next	retry_120s / challenge_accept / wired Map	Dangling essay
SAFE-FBLIT	Uptake fidelity under FEI, not badge	This chapter	Feedback Minutes NS

Operational definition (HYPOTHESIS): A coach/tutor feedback event is *SAFE-FBLIT compliant* when (a) the message is decodeable at the student’s grain (concept/ingredient/FormatId — not essay jargon), (b) affect is managed (task/process/SR level; no self/praise as the payload), (c) a concrete next action is named or forced in-product, (d) the student performs that action within a short window, and (e) success is measured by FEI events — never Feedback Literacy Score™, Feedback Minutes, or thumbs-up rate.

CXII.3 Carless & Boud — four features, action is non-optional

FACT (definitional spine): Carless & Boud (2018, *Assessment & Evaluation in Higher Education*, 43(8), 1315–1325, doi:10.1080/02602938.2018.1463354) — *student feedback literacy* = understandings, capacities, and dispositions to make sense of information and use it to enhance work or learning strategies. Four features: **appreciating feedback; making judgments; managing affect; taking action**. Peer feedback and exemplars enable literacy; teachers facilitate via design and coaching. Unless comments are taken up, feedback has minimal value.

Applied (HYPOTHESIS): Appreciate = soft-wrong as diagnostic gift (SAFE-ERRCLIMATE); judge = compare-to-standard chips (SAFE-FA); affect = destake + short coaches (SAFE-EXPLAIN); action = gated retry / fade / bridge hop (SAFE-HINT / SAFE-FADE). Peer enablement densifies SAFE-PEERX only when building-gated.

Kill: Feedback-as-delivered-product; appreciating without action.

Survive: Four-feature checklist as UX/HITL fidelity, not a Literacy Score™.

CXII.4 Sutton — knowing, being, acting (identity is in the loop)

FACT: Sutton (2012, *Innovations in Education and Teaching International*, 49(1), 31–40, doi:10.1080/14703297.2012.647781) — feedback literacy as academic literacies with epistemological (**knowing**), ontological (**being**), and practical (**acting**) dimensions; becoming feedback literate is complex and mediated by perceptions of teachers’ educational identities.

Applied (HYPOTHESIS): Feedback that attacks *who you are* collapses the ontological channel; procedure dumps without knowing/being support fail too. Tutor grain (SAFE-TUTORGRAIN) mediates uptake. Soft-wrong + competence evidence (SAFE-PROOF) supports *being* without empty “math person” praise (SAFE-IDMEASURE ban).

Kill: Ability-shame feedback; identity-colonization praise≡literacy.

Survive: Knowing/being/acting as a three-check for coach copy and tutor briefs.

CXII.5 Molloy, Boud & Henderson — learner-centred competency grain

FACT (empirical framework): Molloy, Boud, & Henderson (2020, *Assessment & Evaluation in Higher Education*, 45(4), 527–540, doi:10.1080/02602938.2019.1667955) — secondary analysis of large two-university student feedback data; seven groupings / 31 characteristics of learner feedback literacy (e.g., understanding purposes/roles, seeking information, making judgments, working with emotions, processing/using information for future work). Positions students as active agents; framework for diagnosing interventions and aligning course designs.

Applied (HYPOTHESIS): Use Molloy et al. as a **design checklist**, not a report-card. Prefer event proxies (seek soft hint; compare chip; affect-regulated retry; enact next mission) over Likert literacy splash screens.

Kill: Feedback Literacy Score™ from surveys; competency theater without enact.

Survive: Seeking + judging + emotion work + enact as instrumentable behaviors.

CXII.6 Winstone et al. — proactive recipience and barriers

FACT (systematic review): Winstone, Nash, Parker, & Rowntree (2017, *Educational Psychologist*, 52(1), 17–37, doi:10.1080/00461520.2016.1207538) — review of 195 outputs on learners’ agentic engagement with feedback; introduce **proactive recipience**; taxonomy of recipience processes; interventions that support shared responsibility for making feedback effective.

FACT (barriers / focus groups): Winstone, Nash, Rowntree, & Parker (2017, *Studies in Higher Education*, 42(11), 2026–2041, doi:10.1080/03075079.2015.1130032) — students often value feedback yet do not use it (“it’d be useful, but I wouldn’t use it”); barriers from decoding terminology to unwillingness to expend effort; four psychological processes: **awareness**, **cognisance**, **agency**, **volition**. Removing barriers typically requires shared teacher–student responsibility; train proactive receivers.

Applied (HYPOTHESIS): Map awareness→cognisance→agency→volition: feedback *findable*, *decodeable* (plain + FormatId), *actionable* (one next move), *worth the effort* (short retry path). Shared responsibility = product + tutor coaching, not “fix the illiterate student.”

Kill: Passive inbox feedback; jargon dumps; effort-tax essays without action affordances.

Survive: Proactive recipience as FEI-linked UX; barrier taxonomy for HITL QA.

CXII.7 Kluger & DeNisi + Hattie/Timperley — when feedback hurts uptake

FACT (meta + FIT): Kluger & DeNisi (1996, *Psychological Bulletin*, 119(2), 254–284, doi:10.1037/0033-2909.119.2.254) — feedback interventions improve performance on average ($d \approx .41$) but **over one-third decrease** performance; Feedback Intervention Theory: effectiveness falls as attention moves up the hierarchy toward the **self** and away from the task.

FACT: Hattie & Timperley (2007, *Review of Educational Research*, 77(1), 81–112, doi:10.3102/003465430298487) — powerful but highly variable; where going / how going / where next; task/process/self-regulation levels beat self/praise.

FACT (updated meta): Wisniewski, Zierer, & Hattie (2020, *Frontiers in Psychology*, 10, 3087, doi:10.3389/fpsyg.2019.03087) — overall $d \approx 0.48$ with large heterogeneity; information content matters.

Applied (HYPOTHESIS): Literacy without *level discipline* is cosplay. Ego toasts and shame reels both pull attention to self → FIT wound. Densify SAFE-FBTIME / SAFE-EXPLAIN: short where-next, then force action (Carless & Boud’s fourth feature).

Kill: Praise-as-literacy; shame-as-motivation; longer=better.

Survive: Task/process/SR feedback that schedules uptake.

CXII.8 Soft-wrong product law — literacy is an uptake loop

HYPOTHESIS (product law): After soft-wrong (or tutor comment), the default stack is:

1. **Decode (≤ 1 glance):** Name the gap at concept/ingredient/FormatId grain.
2. **Affect gate:** Destake language; no ability verdict.
3. **One next move:** Retry, soft hint, fade completion, or Map hop — not a menu bomb.
4. **Action forced:** UI does not advance as “done” until attempt or explicit defer with scheduled return.
5. **Instrument:** feedback_uptake_act + FEI events; thumbs-up is vanity only.

Banned patterns:

1. Unlimited thorough AI essays after one miss (SAFE-EXPLAIN).
2. Thumbs-up without retry.
3. Parent live stalk of every soft-wrong comment (SAFE-PDASH).
4. Feedback Literacy Score™ / Feedback Minutes NS.
5. “You dismissed feedback” shame without fixing jargon/length/action affordances.

Kill: Delivery=uptake; helpful-click NS.

Survive: Uptake loop as the unit of coach quality.

CXII.9 Competitive contrast (mechanism lens)

Competitor pattern	Mechanism	MindCraft counter
Instant green-check AI	Delivery without literacy	Soft-wrong + uptake gate
Essay tutors / ChatGPT monologue	Length as care	Short where-next + retry
“Was this helpful?” NPS	Thumbs=learning	FEI uptake events
Gradebook comment dumps	Inbox without action	Wired next mission
Literacy badge courses	Survey theater	Behavior checklist under FEI

Positioning line (FOUNDER BELIEF → copy test): “We don’t grade feedback literacy — we make the next attempt hard to skip once you see the gap.”

CXII.10 Confidence matrix

Claim	Label	Confidence	Notes
Feedback literacy = sense-making + use (Carless & Boud four features)	FACT (construct) / FOUNDER BELIEF (product centrality)	High / Medium–High	Carless & Boud 2018
Knowing/being/acting dimensions matter; identity mediates uptake	FACT (theory) / HYPOTHESIS (product)	High / Medium	Sutton 2012
Learner literacy has multiple observable competencies (seeking, judging, emotion, enact)	FACT (framework)	Medium–High	Molloy et al. 2020; HE sample
Proactive recipience + barrier processes (awareness→volition) should guide UX	FACT / HYPOTHESIS	High / Medium	Winstone et al. 2017a/b
Feedback often fails; >1/3 FIs can hurt; self-focus is a wound	FACT	High	Kluger & DeNisi 1996; Hattie & Timperley
Feedback Literacy Score™ / thumbs-up improve delayed solo transfer	SPECULATION (reject)	—	Kill
Soft-wrong uptake loop beats essay+thumbs on delayed FEI transfer	HYPOTHESIS	Medium	Needs FBLIT-*

CXII.11 Experiments

ID	Question	Design	Primary
FBLIT-1	Soft-wrong uptake gate (required retry) vs essay+thumbs matched content	2-arm Practice	delayed solo_transfer_pass; thumbs vanity

FBLIT-2	Decodeable short coach vs jargon-long coach (same correctness info)	2-arm	feedback_uptake_act; quit
FBLIT-3	Affect-managed task/process feedback vs ego praise/shame	2-arm	retry_120s; affect quit
FBLIT-4	Compare-to-standard chip before reveal vs reveal-only	2-arm	judgment accuracy; next-attempt quality
FBLIT-5	Tutor QA on recipience barriers (awareness/cognition/agency/volition) vs hours coaching only	2-arm HITL	FEI rates; hours NS
FBLIT-QUAL	10 Maya: when comments felt usable vs dismissable; jargon/affect/action	Qual	codebook: appreciate / judge / affect / act / theater

Falsifier: Essay+thumbs $\geq$ uptake-gate on delayed solo transfer with equal quit $\rightarrow$ revise gate friction; still ban Feedback Literacy Score™ and delivery $\equiv$ uptake ads.

Falsifier: No FEI lift from decodeable coaches after length matched $\rightarrow$ inspect grain/Map wiring before abandoning literacy framing.

Falsifier: Barrier-trained tutors unchanged on uptake $\rightarrow$ demote literacy language to subordinate FBTIME/EXPLAIN craft; keep FIT self-focus ban.

Pre-register: FBLIT- before “science-backed feedback literacy” campaigns. *Densify FBTIME/FA/EXPLAIN/HINT/HELP/SRL/MONITOR/ERRCLIMATE — do not fork a Feedback Literacy Score™ religion. Family **FBLIT-** owns uptake fidelity (distinct from FBTIME- and FA-).*

CXII.12 So what for MindCraft commercially

- **Copy:** “We don’t grade your feedback literacy — we make the next attempt impossible to skip after you see what went wrong.” Never Feedback Literacy Score™, Feedback Minutes, or “our AI writes thorough comments” as the hero.
- **Product:** Soft-wrong and tutor comments encode decode $\rightarrow$ affect gate $\rightarrow$ one next move $\rightarrow$ forced action; instrument uptake events; ban thumbs-as-learning.
- **Positioning:** Against monologue AI and helpful-click NPS; for inspectable recipience under FEI.
- **Metric:** feedback_uptake_act + FEI — demote Feedback Minutes, Literacy Score™, thumbs-up rate.
- **Kill list:** Feedback Literacy Score™; Feedback Minutes NS; delivery $\equiv$ uptake; thumbs-up $\equiv$ used; longer $\equiv$ more literate; illiteracy-shame without design fix; always-more-feedback dogma; feedback-literacy ACT guarantees.
- **Growth / vision:** Identity includes becoming someone who *seeks and uses* diagnostic information — tutors/product as temporary scaffolds that fade toward self-monitoring (Sadler / SAFE-SRL), not endless comment

surplus.

CXII.13 Doctrine — SAFE-FBLIT (provisional)

1. **Literacy = sense + use** — Carless & Boud four features; action is non-optional.
2. **Knowing / being / acting** — Sutton dimensions; no ability-shame; no empty identity praise.
3. **Proactive recipience** — design for awareness, cognisance, agency, volition (Winstone et al.).
4. **Shared responsibility** — barriers are product+tutor failures first, not student moral defects.
5. **Level discipline** — task/process/SR; ban self-focus payloads that FIT predicts can hurt (Kluger & DeNisi; Hattie & Timperley).
6. **Uptake loop UX** — decode → affect gate → one move → forced action → FEI.
7. **Competency as events** — Molloy checklist as design/HITL QA, never Literacy Score™ splash.
8. **Densify, don't fork** — FBTIME/FA/EXPLAIN/HINT/HELP/SRL/MONITOR/ERRCLIMATE.
9. **No Feedback Literacy Score™ / Feedback Minutes NS / delivery⇒uptake / thumbs⇒used / longer⇒literate / illiteracy-shame / always-more-feedback / feedback-literacy ACT guarantees.**

Confidence: High that feedback-without-uptake is commercially hollow, and that self-focused / thumbs / length theater are unsafe. Medium that the uptake loop beats essay+thumbs on delayed FEI — instrument FBLIT-* before literacy marketing. High confidence to kill Feedback Literacy Score™, delivery⇒uptake ads, and Feedback Minutes NS regardless.

References (verified)

- Carless, D., & Boud, D. (2018). The development of student feedback literacy: Enabling uptake of feedback. *Assessment & Evaluation in Higher Education*, 43(8), 1315–1325. <https://doi.org/10.1080/02602938.2018.1463354>
 - Hattie, J., & Timperley, H. (2007). The power of feedback. *Review of Educational Research*, 77(1), 81–112. <https://doi.org/10.3102/003465430298487>
 - Kluger, A. N., & DeNisi, A. (1996). The effects of feedback interventions on performance: A historical review, a meta-analysis, and a preliminary feedback intervention theory. *Psychological Bulletin*, 119(2), 254–284. <https://doi.org/10.1037/0033-2909.119.2.254>
 - Molloy, E., Boud, D., & Henderson, M. (2020). Developing a learning-centred framework for feedback literacy. *Assessment & Evaluation in Higher Education*, 45(4), 527–540. <https://doi.org/10.1080/02602938.2019.1667955>
 - Sutton, P. (2012). Conceptualizing feedback literacy: Knowing, being, and acting. *Innovations in Education and Teaching International*, 49(1), 31–40. <https://doi.org/10.1080/14703297.2012.647781>
 - Winstone, N. E., Nash, R. A., Parker, M., & Rowntree, J. (2017). Supporting learners' agentic engagement with feedback: A systematic review and a taxonomy of recipience processes. *Educational Psychologist*, 52(1), 17–37. <https://doi.org/10.1080/00461520.2016.1207538>
 - Winstone, N. E., Nash, R. A., Rowntree, J., & Parker, M. (2017). "It'd be useful, but I wouldn't use it": Barriers to university students' feedback seeking and recipience. *Studies in Higher Education*, 42(11), 2026–2041. <https://doi.org/10.1080/03075079.2015.1130032>
 - Wisniewski, B., Zierer, K., & Hattie, J. (2020). The power of feedback revisited: A meta-analysis of educational feedback research. *Frontiers in Psychology*, 10, 3087. <https://doi.org/10.3389/fpsyg.2019.03087>
-

Part CXIII — Collaborative Problem-Solving vs Solo FEI

Chapter status: Living evidence + Practice/HITL UX — Researcher tick 2026-08-09 (UTC hour 21; hour%6≠0; researcher count since synthesizer v1.14 = 5 → Researcher)

Primary question: When should MindCraft use **collaborative / CSCL problem-solving** as a learning move — versus shipping Discord study rooms, group-completion theater, Collaboration Score™ dashboards, or “we learn together” ads that let group correctness wash out **solo transfer**?

Owners: Product (Practice / Map / optional group modes) · HITL / classroom ops · Brand · Analytics · Red Team

Commercial job: Ship **SAFE-COLLAB** densifying SAFE-PEERX / SAFE-CoP / SAFE-PROOF / SAFE-HELP / SAFE-INSTRUMENT / SAFE-ERRCLIMATE: structured collab *with individual accountability and a mandatory solo FEI gate* — never Collaboration Score™, Group Minutes, Discord≡learning, or group-correct≡ready.

Builds on: Parts LII (SAFE-CoP), CVI (SAFE-PEERX), LXXXVI (SAFE-PROOF), LXXXIX (SAFE-HELP), XC (SAFE-INSTRUMENT), CIII (SAFE-ERRCLIMATE), CXI (SAFE-FA). Product seams: teach-back, near-peer rooms, classroom SKUs, parent “study with friends” pressure.

CXIII.1 Why this chapter exists

Edtech defaults: Discord rooms as “community learning,” one-grade group worksheets, Collaboration Score™ / Group Minutes KPIs, togetherness-as-identity without solo proof, and free-riding framed as character failure. Research spine: cooperative/CSCL methods often beat traditional baselines *when structured*; group-to-individual transfer is real when discussion is structured; but **group scores are not valid individual competence** — high group scores can hide copying, being told, or free-riding.

MindCraft’s North Star already privileges solo_transfer_pass. Open commercial question: can collab *serve* solo FEI, or will “learn together” SKUs wash it out?

FOUNDER BELIEF under audit: Mathematical identity is personal competence under challenge — collaboration is a scaffold that must leave the student able to stand alone.

Claims we refuse as doctrine: Collaboration Score™ / Group Minutes NS; Discord≡CSCL learning; group-correct≡ready; free-rider-blind unlocks; always-collab dogma; collab ACT guarantees.

CXIII.2 Constructs (keep group production ≠ individual mastery)

Construct	Research meaning	MindCraft analogue	Failure mode
Cooperative / collaborative PS	Structured peer work with interdependence + accountability	Optional pair/classroom missions with roles + solo gate	Open chat room cosplay

CSCL	Computer-supported collab with tech+pedagogy moderators	Scripted co-solve in-app (not raw Discord)	Tool-presence≡science
Individual accountability	Each member assessed / responsible for learning	Solo FEI item after group phase	One shared green check
Group-to-individual transfer	Learning in group shows up on later individual tests	Delayed solo_transfer_pass after collab	End on group board
Free-riding / social loafing	Credit without contribution when effort dispensable	Detect silent / copy / “tell me the answer”	Character-shame ads
SAFE-COLLAB	Collab that densifies PEERX then proves solo	This chapter	Collaboration Score™

Operational definition (HYPOTHESIS): A MindCraft path is *SAFE-COLLAB compliant* when (a) collab is optional or classroom-scoped, never the only path to progress, (b) structure includes positive interdependence *and* individual accountability (roles, scripts, or individual checks — not vibes), (c) help given/received follows PEERX building + apply rules, (d) every collab episode ends with a **solo** attempt on a near/far transfer item before unlock, (e) marketing never equates room presence or group correctness with FEI identity, and (f) free-riding is treated as a product design failure (missing accountability), not a grit slogan.

CXIII.3 CSCL and cooperative math — positive average effects, not a blank check

FACT (STEM CSCL meta): Jeong, Hmelo-Silver, & Jo (2019, *Educational Research Review*, 28, 100284, doi:10.1016/j.edurev.2019.100284) — 316 outcomes from 143 studies (2005–2014); overall STEM CSCL effect size ≈ **0.51** (moderate); largest on **process** outcomes, then knowledge, then affective; moderated by technology type, pedagogy, educational level, and domain — interactions matter; “CSCL works” is incomplete without those moderators.

FACT (CSCL elements meta): Chen, Wang, Kirschner, & Tsai (2018, *Review of Educational Research*, 88(6), 799–843, doi:10.3102/0034654318791584) — collaboration per se aids knowledge gain (ES ≈ 0.42), skill acquisition (≈ 0.64), perceptions (≈ 0.38) in computer-based conditions; computer use in collab contexts also positive; extra tools/environments and supporting strategies add further knowledge-gain effects (≈ 0.55 / ≈ 0.38). Collaboration is not magic — scaffolding and design matter.

FACT (math cooperative meta): Capar & Tarim (2015, *Educational Sciences: Theory & Practice*, 15(2), 553–559, doi:10.12738/estp.2015.2.2098) — 26 studies / 36 comparisons (1988–2010); cooperative learning vs traditional on math achievement **d++ ≈ 0.59** (medium); attitudes **d++ ≈ 0.16** (small).

Applied (HYPOTHESIS): Ship *structured* collab for classrooms/optional peers — average metas do **not** license Discord-as-CSCL or group-correct≡ready. Moderators are the product; presence is not.

Kill: “Collaboration always beats solo” / room-presence≡CSCL.

Survive: Structured collab as a *conditional* channel under FEI gates.

CXIII.4 Slavin's hinge — group goals + individual accountability

FACT (best-evidence / process programs): Slavin & Lake (and related Slavin line; see Slavin's elementary math best-evidence synthesis, ERIC ED527642) — cooperative learning raises learning when groups have a **common goal they can achieve only if all members do well on independent assessments**; teammates know their success depends partly on others' learning (see also Johnson & Johnson; Rohrbeck et al.; Webb & Palincsar reviews cited therein). Unstructured "work together" without that hinge is a different animal.

Applied (HYPOTHESIS): Any collab SKU must wire **individual accountability** — solo transfer, Map update, or contribution check — not one shared green. Densifies SAFE-PROOF / SAFE-INSTRUMENT: North Star remains solo evidence.

Kill: Shared worksheet unlock; team XP without solo proof.

Survive: Group goal + independent assessment as design law.

CXIII.5 Group-to-individual transfer is real — when discussion is structured

FACT (experimental): Yager, Johnson, & Johnson (1985, *Journal of Educational Psychology*, 77(1), 60–66, doi:10.1037/0022-0663.77.1.60) — second-graders; structured oral discussion (summarize + monitor others' summaries) vs unstructured cooperative vs individualistic; cooperative conditions beat individualistic on daily accuracy; **structured** cooperative students scored higher on *individual* posttests and retention than unstructured cooperative and individualistic — evidence that **group-to-individual transfer** occurs and that structured summarization/monitoring contributes.

Applied (HYPOTHESIS): Force structured building moves (inferential explain, monitor peer summary) — densifies SAFE-PEERX — then measure transfer on a **solo** item. Unstructured chat tracks the weaker arm.

Kill: Open voice lobby as "transfer science."

Survive: Structured discussion → solo check as the transfer unit.

CXIII.6 Webb — group scores ≠ individual competence

FACT (assessment validity): Webb (1993, *Educational Assessment*, 1(2), 131–152, doi:10.1207/s15326977ea0102_3; CRESST Tech Report 352) — group math performance was **much higher** than individual; without process data, group scores cannot distinguish solvers, learners-via-collab, and copy/told students — **all can score high in group**; do not infer individual competence from group scores.

FACT (process → learning): Webb, Troper, & Fall (1995, *Journal of Educational Psychology*, 87(3), 406–423, doi:10.1037/0022-0663.87.3.406) — **constructive activity** (using received explanations to solve) predicts learning beyond merely being told (SAFE-PEERX).

Applied (HYPOTHESIS): Never report group board correctness as mastery. After collab help, require apply-on-this-item *and* delayed solo mix. Parent dashboards must not stalk group greens as proof (SAFE-PDASH).

Kill: Group-correct $\equiv$ ready; room leaderboard as competence.

Survive: Process-aware collab + solo FEI as the only mastery claim surface.

CXIII.7 Free-riding — design failure, not grit sermon

FACT (motivation loss): Kerr & Bruun (1983, *Journal of Personality and Social Psychology*, 44(1), 78–94, doi:10.1037/0022-3514.44.1.78) — free-rider effects when member effort is perceived as dispensable; group motivation losses are predictable under certain task structures.

SPECULATION (product): Unstructured Discord rooms and shared unlocks maximize perceived dispensability → free-riding + copy culture; MindCraft will recreate Webb’s invalid-inference problem at scale.

Applied (HYPOTHESIS): Treat free-riding as missing individual accountability / role scripts / solo gates — fix design. Do not ship “no free riders — show grit” shame copy (conflicts SAFE-STRUCTURE / SAFE-ERRCLIMATE).

Kill: Free-rider moralizing; silent-presence $\equiv$ participation.

Survive: Accountability architecture as anti-loafing.

CXIII.8 Product law — collab as scaffold, solo FEI as proof

HYPOTHESIS (product law): Default MindCraft practice remains **solo FEI**. Optional collab (classroom or consented peer) follows:

1. **Scope:** Classroom / invited peer — not anonymous public lobby as learning core.
2. **Script:** Roles or inferential explain/monitor prompts (Yager structure; PEERX building).
3. **Apply:** Receiver uses help on *this* item (Webb constructive activity).
4. **Solo gate:** Near-item then delayed mix — unlock requires solo_transfer_pass (or equivalent), never group board alone.
5. **Instrument:** collab_episode + solo_after_collab + FEI events; demote Group Minutes / Collaboration Score™ / Discord DAU.

Banned patterns:

1. Discord $\equiv$ CSCL / “community of practice” without enterprise+repertoire (SAFE-CoP).
2. Shared green unlock without solo.
3. Collaboration Score™ / Group Minutes NS.
4. Always-collab brand that demotes solo struggle as antisocial.
5. Collab ACT-point guarantees.
6. Parent stalk of peer chat as competence.

Kill: Togetherness theater; group $\equiv$ individual mastery.

Survive: Scaffolded collab → solo proof loop.

CXIII.9 Competitive contrast (mechanism lens)

Competitor pattern	Mechanism	MindCraft counter
Discord study rooms	Presence / chat volume	Scripted collab + solo gate
Group worksheet one-grade	Shared correctness	Individual accountability + FEI
“Learn together” hero	Belonging without proof	Belonging + solo_transfer_pass
Team XP / streaks	Engagement theater	FEI events; no Group Minutes NS
AI multiplayer dump	Fluent group answer	Building prompts + apply + solo

Positioning line (FOUNDER BELIEF → copy test): “We practice together when it helps — we prove alone before we call it learned.”

CXIII.10 Confidence matrix

Claim	Label	Confidence	Notes
STEM CSCL shows moderate average gains; moderators matter	FACT	High	Jeong et al. 2019; Chen et al. 2018
Cooperative math vs traditional often medium achievement gains	FACT	High	Capar & Tarim 2015; heterogeneity remains
Group goals + individual accountability are load-bearing	FACT (Slavin line) / HYPOTHESIS (product)	High / Medium–High	Best-evidence hinge
Structured discussion supports group-to-individual transfer	FACT	Medium–High	Yager et al. 1985 (grade 2; generalize carefully)
Group assessment scores ≠ individual competence	FACT	High	Webb 1993
Constructive use of help predicts learning in groups	FACT	High	Webb et al. 1995

Free-rider effects when effort dispensable	FACT	High	Kerr & Bruun 1983
Collaboration Score™ / Discord rooms raise delayed solo FEI	SPECULATION (reject)	—	Kill
Scripted collab + solo gate beats open room on delayed solo FEI	HYPOTHESIS	Medium	Needs COLLAB-*

CXIII.11 Experiments

ID	Question	Design	Primary
COLLAB-1	Scripted pair PS + solo gate vs solo-only matched content	2-arm Practice	delayed solo_transfer_pass; quit
COLLAB-2	Open chat room vs scripted roles/inferential prompts (same time)	2-arm	solo-after-collab; free-ride proxies
COLLAB-3	Shared group unlock vs mandatory solo unlock after collab	2-arm	delayed solo FEI; group-correct vanity
COLLAB-4	Individual accountability (personal check) vs team XP only	2-arm	contribution; solo transfer
COLLAB-5	Classroom structured collab + Map vs Discord-adjacent lobby	2-arm / quasi	FEI; hours NS
COLLAB-QUAL	10 Maya: when group help felt like learning vs copying	Qual	codebook: building / told / free-ride / solo-ready feel

Falsifier: Open room $\geq$ scripted+solo-gate on delayed solo transfer with equal quit $\rightarrow$ revise scripts; still ban group-correct $\equiv$ ready and Collaboration Score™.

Falsifier: Solo-only $\geq$ scripted collab+solo on FEI and belonging — keep collab optional; do not force always-collab brand.

Falsifier: Accountability checks unchanged free-ride proxies — strengthen role scripts before abandoning collab SKUs.

Pre-register: COLLAB- before “science-backed collaborative learning” or Discord-community campaigns. Densify PEERX/CoP/PROOF/HELP/INSTRUMENT — do not fork a Collaboration Score™ religion. Family **COLLAB-** owns

group→solo fidelity (distinct from PEERX- *explain quality and CoP*- community claims).

CXIII.12 So what for MindCraft commercially

- **Copy:** “We practice together when it helps — we prove alone before we call it learned.” Never Collaboration Score™, Group Minutes, or Discord-as-learning hero.
 - **Product:** Optional/classroom collab with scripts + individual accountability + mandatory solo FEI gate; ban shared-green unlocks.
 - **Positioning:** Against room-presence edtech and group-grade theater; for inspectable group-to-individual transfer under FEI.
 - **Metric:** solo_after_collab + FEI — demote Group Minutes, Collaboration Score™, Discord DAU, group-correct rate.
 - **Kill list:** Collaboration Score™; Group Minutes NS; Discord≡CSCL; group-correct≡ready; free-rider-blind unlocks; always-collab dogma; collab ACT guarantees.
 - **Growth / vision:** Near-peer and classroom SKUs accelerate identity via structured help — identity still proven in solo transfer, not the chat log.
-

CXIII.13 Doctrine — SAFE-COLLAB (provisional)

1. **Collab is conditional** — meta-gains exist; moderators (scripts, scaffolds, pedagogy) are the product (Jeong et al.; Chen et al.; Capar & Tarim).
2. **Group goals + individual accountability** — Slavin hinge; no shared unlock without personal proof.
3. **Structure transfer** — Yager-style summarize/monitor (or PEERX building) before claiming group-to-individual transfer.
4. **Group ≠ individual competence** — Webb assessment law; never sell group greens as mastery.
5. **Constructive use** — apply help on this item (Webb et al.); densify PEERX/HELP.
6. **Anti-free-ride by design** — Kerr & Bruun; fix dispensability; no grit sermons.
7. **Solo FEI gate** — every collab episode ends in solo proof; instrument solo_after_collab.
8. **Densify, don’t fork** — PEERX / CoP / PROOF / HELP / INSTRUMENT / ERRCLIMATE.
9. **No Collaboration Score™ / Group Minutes NS / Discord≡learning / group-correct≡ready / free-rider-blind unlocks / always-collab dogma / collab ACT guarantees.**

Confidence: High that unstructured group scoring and Discord-presence theater are commercially unsafe relative to FEI. Medium that scripted collab + solo gate beats solo-only on delayed transfer for some populations — instrument COLLAB-* before collab marketing. High confidence to kill Collaboration Score™, group-correct≡ready, and Discord≡CSCL ads regardless.

References (verified)

- Capar, G., & Tarim, K. (2015). Efficacy of the cooperative learning method on mathematics achievement and attitude: A meta-analysis research. *Educational Sciences: Theory & Practice*, 15(2), 553–559.

<https://doi.org/10.12738/estp.2015.2.2098>

- Chen, J., Wang, M., Kirschner, P. A., & Tsai, C.-C. (2018). The role of collaboration, computer use, learning environments, and supporting strategies in CSCL: A meta-analysis. *Review of Educational Research*, 88(6), 799–843. <https://doi.org/10.3102/0034654318791584>
- Jeong, H., Hmelo-Silver, C. E., & Jo, K. (2019). Ten years of Computer-Supported Collaborative Learning: A meta-analysis of CSCL in STEM education during 2005–2014. *Educational Research Review*, 28, 100284. <https://doi.org/10.1016/j.edurev.2019.100284>
- Kerr, N. L., & Bruun, S. E. (1983). Dispensability of member effort and group motivation losses: Free-rider effects. *Journal of Personality and Social Psychology*, 44(1), 78–94. <https://doi.org/10.1037/0022-3514.44.1.78>
- Slavin, R. E., & Lake, C. (2007). Effective programs in elementary mathematics: A best-evidence synthesis (Report). Johns Hopkins University / ERIC ED527642. <https://files.eric.ed.gov/fulltext/ED527642.pdf>
- Webb, N. M. (1993). Collaborative group versus individual assessment in mathematics: Processes and outcomes. *Educational Assessment*, 1(2), 131–152. https://doi.org/10.1207/s15326977ea0102_3
- Webb, N. M., Troper, J. D., & Fall, R. (1995). Constructive activity and learning in collaborative small groups. *Journal of Educational Psychology*, 87(3), 406–423. <https://doi.org/10.1037/0022-0663.87.3.406>
- Yager, S. O., Johnson, D. W., & Johnson, R. T. (1985). Oral discussion, group-to-individual transfer, and achievement in cooperative learning groups. *Journal of Educational Psychology*, 77(1), 60–66. <https://doi.org/10.1037/0022-0663.77.1.60>

Part CXIV — Desirable Difficulty Scheduling Under Exam Deadlines

Chapter status: Living evidence + Practice/exam calendar UX — Researcher tick 2026-08-10 (UTC hour 00 ≡ Red Team slot, but ch114 never written → prefer Researcher per rotation; researcher count since synthesizer v1.14 = 6 → Researcher)

Primary question: In the **final weeks before ACT/SAT/state tests**, should MindCraft **hold** Bjork-style desirable difficulties (spacing, interleaving, retrieval), **ease** them into blocked comfort theater, or **max** them as timed floods — and how does that schedule interact with SAFE-EXAM pressure dosing?

Owners: Product (Practice calendar / exam-mode) · Learning Science · Parent messaging · Growth · Red Team

Commercial job: Ship **SAFE-DDSCHED** densifying SAFE-DD × SAFE-EXAM × SAFE-SCHED × SAFE-DURABLE / SAFE-CRAM: a **horizon-matched difficulty schedule** — never Desirable Difficulty Score™, Challenge Minutes NS, always-ease-before-ACT, always-max-DD-until-bell, cram-week≡ready, or ACT-point guarantees from calendar theater.

Builds on: Parts XLI (SAFE-DD), LXIII (SAFE-EXAM), LXXIV (SAFE-SCHED), LXXVIII (SAFE-DURABLE / CRAM), XLVII (sleep/stress), LVIII (SAFE-EXPOSE), XCIII (SAFE-RETRIEVE). Product seams: ACT date on profile, Returns cards, prove-week mocks, parent “make it easier / grind harder” pressure.

CXIV.1 Why this chapter exists

Parents buy **exam dates**. Students feel **countdown panic**. Tutors get asked: “Should we keep mixing, or just drill what she knows?” Edtech defaults into two commercially loud failures: (1) **comfort week** — blocked greens / streak theater

that feel ready and fail transfer; (2) **cruelty week** — nightly timed floods and parent-visible scores that brand stakes escalation as “desirable difficulty,” disrupting retrieval benefits (Hinze & Rapp) and taxing WM (Ashcraft).

SAFE-DD / SAFE-EXAM / SAFE-SCHED already say: equip then strain, destake to learn / dose to prove, match spacing to horizon. This chapter answers the **calendar join**: *what to hold vs ease as $T \rightarrow 0$* .

FOUNDER BELIEF under audit: Identity under challenge is built by equipped difficulty across a horizon — not by abandoning challenge in the last week, and not by turning every remaining hour into a mini-ACT.

Claims we refuse as doctrine: Desirable Difficulty Score™ / Challenge Minutes NS; always-ease-before-ACT / blocked-comfort $\Rightarrow$ ready; always-max-DD-until-bell; last-48h full timed flood as “science”; cram-week $\Rightarrow$ ready; guaranteed ACT points from schedule; soft-only forever under deadline.

CXIV.2 Constructs (schedule $\neq$ stakes $\neq$ content)

Keep **DD** (equipped challenge), **stakes/pressure dose** (evaluation intensity — not the same as DD), **horizon/RI** (days to the test), and **ISI** (gaps between episodes) distinct. “Ease” can mean less mix (**content ease**) or less simultaneous WM tax via destake/shorter sets/sleep (**affect ease**). Conflating them produces comfort $\Rightarrow$ ready or never-challenge forever.

Operational definition (HYPOTHESIS): SAFE-DDSCHED compliant when (a) far out: spacing + faded interleave + destaked retrieval default, (b) mid-window: **hold mix** while mild pressure doses begin, (c) final ~ 72 h: short spaced retrieval of equipped repertoire + sleep/WM guards + labeled prove probes — not new-grain floods or scored acquisition marathons, (d) readiness = delayed mixed evidence, never blocked greens/cram NPS, (e) no always-ease / always-max / calendar $\rightarrow$ points marketing.

CXIV.3 Bjork — difficulty stays desirable only while equipped; deadlines do not repeal that

FACT (origin / definition): Bjork (1994) and Bjork & Bjork (2011; see also Bjork & Bjork JARMAC comment lineage) — spacing, variation, interleaving, and testing-as-learning can improve long-term retention and transfer relative to fluent short-term conditions; **many difficulties are undesirable forever**; a difficulty is desirable only when the learner can respond successfully with the processes the difficulty is meant to trigger.

Applied (HYPOTHESIS): An approaching ACT date does **not** convert unequipped full mix or high-stakes acquisition into “more science.” If deadline panic unequips (WM tax, sleep loss, new-topic flood), the same schedule that was desirable at $T-30$ d becomes undesirable at $T-2$ d.

Kill: Deadline-as-permission to abandon the *equipped* clause.

Survive: Hold DD forms the learner can still answer; change **load and stakes**, not the definition of desirable.

CXIV.4 Spacing vs cramming under a known test date

FACT (distributed practice meta): Cepeda, Pashler, Vul, Wixted, & Rohrer (2006, *Psychological Bulletin*, 132(3), 354–380, doi:10.1037/0033-2909.132.3.354) — 839 assessments / 317 experiments; **ISI and retention interval operate jointly**; the ISI that maximizes final retention **increases as the retention interval increases**.

FACT (temporal ridgeline): Cepeda, Vul, Rohrer, Wixted, & Pashler (2008, *Psychological Science*, 19(11), 1095–1102, doi:10.1111/j.1467-9280.2008.02209.x) — optimal gap depends on how long knowledge must last; practical implication: schedule to the **exam horizon**, not a universal “perfect interval.”

FACT (spacing > cramming; metacognitive trap): Kornell (2009, *Applied Cognitive Psychology*, 23(9), 1297–1317, doi:10.1002/acp.1537) — spacing (including large-stack / multi-day) beat massing and beat **cramming (massed study on the last day)** for GRE-type pairs; across experiments spacing beat massing for ~90% of participants, yet after first session ~72% *believed* massing was more effective.

Applied (HYPOTHESIS): Pre-ACT product must resist parent/student demand for last-day massed fluency theater. For short remaining RI, **shorten ISI** (more frequent Returns) — do not collapse to single-night cram as the brand. Densifies SAFE-SCHED / SAFE-DURABLE / SAFE-CRAM.

Kill: Cram-week $\Rightarrow$ ready; “science says grind tonight.”

Survive: Horizon-matched spacing; honesty that fluent last-night massing feels best and often teaches less for delayed/mixed demand.

CXIV.5 Interleaving prepares cumulative exams — do not trade it for blocked comfort as the readiness story

FACT (classroom interleave): Rohrer, Dedrick, & Stershic (2015, *Journal of Educational Psychology*, 107(3), 900–908; ERIC ED557355) — seventh-graders, same problems blocked vs interleaved over ~3 months; interleaved higher on tests 1 and 30 days later ($ds \approx 0.42 / 0.79$). Interleaving juxtaposes strategy choice **and** spaces same-kind problems.

FACT (larger RCT): Rohrer, Dedrick, Hartwig, & Cheung (2020, *Journal of Educational Psychology*) — interleaved math practice improved cumulative-test performance vs mostly blocked schedules in a diverse sample over months; exam-like strategy choice is the mechanism (choose method from the problem, not from the worksheet header).

Applied (HYPOTHESIS): Mid-window, **hold mix** on equipped concepts (ACT-like strategy choice). “Ease” = shorter sets / destake / fewer new grains — not blocked greens sold as ready. Final ~72h may bias known-repertoire retrieval + light mix probes; never redefine blocked accuracy as readiness (SAFE-EXAM).

Kill: Always-ease-before-ACT / blocked-comfort $\Rightarrow$ exam-ready. **Survive:** Hold strategy-choice while equipped; ease load, not the readiness standard.

CXIV.6 Stakes are not desirable difficulties — pressure during acquisition can erase learning gains

FACT (retrieval disruption under pressure): Hinze & Rapp (2014, *Applied Cognitive Psychology*, 28(4), 597–606, doi:10.1002/acp.3032) — high-stakes vs low-stakes quizzes: online quiz performance similar, but **final retention better after low-stakes**; only low-stakes beat reread. Performance pressure can disrupt the memory benefits of retrieval even when session scores look fine.

FACT (practice tests and anxiety): Yang, Li, Zhao, Luo, & Shanks (2023, *Educational Psychology Review*, doi:10.1007/s10648-023-09801-w) — practice testing tends to **reduce** test anxiety (Hedges’ $g \approx -0.52$ across synthesized studies); easier practice tests more mitigating than difficult ones. Directional support for laddered, not brutal, exam sims.

FACT (WM under math anxiety): Ashcraft & Kirk (2001, *JEP: General*, 130(2), 224–237, doi:10.1037/0096-3445.130.2.224); Ashcraft & Krause (2007, *Psychonomic Bulletin & Review*, 14, 243–248, doi:10.3758/BF03194059) — math anxiety acts like a dual-task load on WM; hard math under threat is resource competition, not moral failure.

Applied (HYPOTHESIS): Final week: do not make every Practice minute a parent-visible scored ACT. Prefer destaked spaced retrieval for encoding + **labeled prove probes** for pressure. “Easing” means cutting simultaneous stakes + sleep debt + new-grain flood — not deleting mix from the readiness story.

Kill: Always-max-DD-until-bell; last-48h scored flood-as-DD. **Survive:** Dual-rail train destaked / prove dosed (Hinze + Ashcraft + Yang).

CXIV.7 Product law — three horizon bands

HYPOTHESIS (SAFE-DDSCHD bands):

Band	Approx. horizon	Hold	Ease / change	Prove
Build	T > ~21d	Spacing + faded interleave + destaked retrieval (SAFE-DD/SCHED)	Equip before strain; no day-one flood	Occasional mild timed mix after equip
Join	~7–21d	Hold mix on equipped path concepts; keep Returns	Shorten sessions if affect high; begin EXPOSE ladder	Rising mild→moderate pressure doses
Protect + probe	~0–3d (esp. final 72h)	Short spaced retrieval of known repertoire; light mix probes	No new-grain interleave marathons; protect sleep; destake acquisition; no parent scoreboard on train sets	Few labeled timed prove sets — not all-night acquisition

Banned patterns:

1. Always-ease / comfort-week blocked greens as “ready.”
2. Always-max timed mix every night until the bell.
3. Cram-week $\Rightarrow$ ready / overlearn-tonight as method.
4. Desirable Difficulty Score™ / Challenge Minutes NS.
5. Schedule $\rightarrow$ guaranteed ACT points.
6. Conflating stakes escalation with DD.

Kill: Calendar theater extremes.

Survive: Horizon-matched hold of equipped DD + separate pressure dosing.

CXIV.8 Competitive contrast (mechanism lens)

Cram packs sell massed fluency spikes; every-night ACT sims sell stakes-as-pedagogy; blocked “confidence builders” sell comfort $\Rightarrow$ ready; “grind harder — desirable difficulties!” sells pain-as-virtue; perfect-interval SRS sells black-box calendar mystique. MindCraft counters with horizon Returns, destake-train / prove-dose rails, held mix for readiness, Bjork’s equipped clause, and inspectable ISI $\times$ RI — not SRS brand.

Positioning line (FOUNDER BELIEF $\rightarrow$ copy test): “We keep the workout that matches the exam — and we stop turning every remaining hour into a graded panic.”

CXIV.9 Confidence matrix

Claim	Label	Confidence	Notes
DD helps only when learner can respond successfully	FACT	High	Bjork & Bjork
Optimal spacing depends on retention interval	FACT	High	Cepeda et al. 2006; 2008
Spacing beats last-day cramming; learners misjudge	FACT	High	Kornell 2009 (verbal pairs; transfer caution)
Interleaved math practice beats blocked on delayed/cumulative tests	FACT	High	Rohrer et al. 2015, 2020
High stakes during practice can erase retrieval benefits	FACT	High	Hinze & Rapp 2014

Practice testing often reduces test anxiety; easier > harder for mitigation	FACT	High	Yang et al. 2023
Math anxiety taxes WM under hard demand	FACT	High	Ashcraft line
Final-72h should abandon all mix	SPECULATION (reject)	—	Conflicts Rohrer + readiness law
Final-72h should max scored floods	SPECULATION (reject)	—	Conflicts Hinze + Ashcraft
Three-band SAFE-DDSCHED beats both extremes on delayed mixed + churn	HYPOTHESIS	Medium	Needs DDSCHED-*

CXIV.10 Experiments

ID	Question	Design	Primary
DDSCHED-1	Hold-mix Join band vs ease-to-blocked in last 14d (same total minutes)	2-arm near ACT date	delayed mixed; mock; quit
DDSCHED-2	Protect+probe 72h (short spaced + few prove) vs all-night scored mix flood	2-arm	next-day mixed; sleep/affect; retention
DDSCHED-3	Horizon-matched ISI (Cepeda-style) vs fixed cram-week massing	2-arm	retention @ exam eve + 7d post
DDSCHED-4	Destaked train / labeled prove dual-rail vs single scored calendar	2-arm (Hinze densify)	delayed retention; state anxiety
DDSCHED-5	Parent copy: “ease for confidence” vs “hold mix / protect sleep”	Message / CBC	WTP; misexpectation; choose-hold rate
DDSCHED-QUAL	10 Maya + 10 parents: what “ease up” / “push hard” mean in last week	Qual	codebook: content vs stakes vs sleep

Falsifier: Ease-to-blocked $\geq$ hold-mix on delayed mixed with equal or lower churn $\rightarrow$ revise Join-band hold; still ban blocked $\Rightarrow$ ready ads.

Falsifier: Scored flood $\geq$ protect+probe on delayed mixed and sleep — still ban flood-as-default; segment only if HTA harmed less.

Falsifier: Cram-week = horizon ISI on post-exam delayed probe — wound spacing honesty; keep Kornell metacognition warning.

Pre-register: DDSCHED- before “science-backed exam calendar” or anti-cram / pro-grind campaigns. Densify DD- / EXAM- / SCHED- / CRAM- — do not fork a Desirable Difficulty Score™ religion. Family **DDSCHEd-** owns T→0 schedule fidelity (distinct from DD anxiety equip, EXAM pressure ladders, SCHED return cards, and CRAM GTM packs).

CXIV.11 So what for MindCraft commercially

- **Copy / product:** Three-band ACT calendar; hold equipped mix; Protect+probe final 72h; dual-rail train/prove. Line: “We keep the workout that matches the exam — and we stop turning every remaining hour into a graded panic.”
 - **Positioning / metric:** Against cram packs and cruelty sims; delayed mixed / transfer_pass / declared-stakes mock + churn — demote Challenge Minutes, last-week fireworks, cram NPS.
 - **Kill list:** Desirable Difficulty Score™; Challenge Minutes NS; always-ease-before-ACT; always-max-DD-until-bell; last-48h scored flood-as-DD; cram-week $\Rightarrow$ ready; schedule ACT-point guarantees.
 - **Vision:** Exam GTM sells a calendar of durable choosing, not a final binge — identity proven in solo transfer.
-

CXIV.12 Doctrine — SAFE-DDSCHEd (provisional)

1. **Equipped clause never expires** — deadlines do not make unequipped difficulty desirable (Bjork & Bjork).
2. **Match ISI to remaining RI** — shorten gaps near the exam; do not collapse to last-night massing as method (Cepeda; Kornell).
3. **Hold mix while equipped** — interleaved strategy choice is exam transfer practice (Rohrer line); ban comfort-week $\Rightarrow$ ready.
4. **Stakes $\neq$ DD** — destake acquisition; dose prove separately (Hinze & Rapp; SAFE-EXAM).
5. **Ease load, not standards** — final 72h: sleep/WM/new-grain guards + short spaced repertoire; readiness still delayed mixed.
6. **Laddered practice tests** — anxiety mitigation favors well-sequenced practice over brutal difficulty (Yang et al.).
7. **WM honesty** — Ashcraft dual-task; do not stack timer + new mix + parent scoreboard + sleep debt.
8. **Densify, don’t fork** — SAFE-DD / EXAM / SCHED / DURABLE / CRAM / RETRIEVE / EXPOSE.
9. **No Desirable Difficulty Score™ / Challenge Minutes NS / always-ease / always-max / cram $\Rightarrow$ ready / schedule ACT guarantees.**

Confidence: High that comfort-week and scored-flood calendars are unsafe vs FEI + exam-honest readiness; medium on exact day cutovers (21d/7d/72h) — instrument DDSCHED-* before hardcoding; high to kill Desirable Difficulty Score™ / always-ease $\Rightarrow$ ready / always-max-until-bell / cram-week $\Rightarrow$ ready regardless.

References (verified)

- Ashcraft, M. H., & Kirk, E. P. (2001). The relationships among working memory, math anxiety, and performance. *Journal of Experimental Psychology: General*, 130(2), 224–237. <https://doi.org/10.1037/0096-3445.130.2.224>
- Ashcraft, M. H., & Krause, J. A. (2007). Working memory, math performance, and math anxiety. *Psychonomic Bulletin & Review*, 14(2), 243–248. <https://doi.org/10.3758/BF03194059>
- Bjork, R. A. (1994). Memory and metamemory considerations in the training of human beings. In J. Metcalfe & A. Shimamura (Eds.), *Metacognition: Knowing about knowing* (pp. 185–205). MIT Press.
- Bjork, E. L., & Bjork, R. A. (2011). Making things hard on yourself, but in a good way: Creating desirable difficulties to enhance learning. In M. A. Gernsbacher et al. (Eds.), *Psychology and the real world* (pp. 56–64). Worth.
- Cepeda, N. J., Pashler, H., Vul, E., Wixted, J. T., & Rohrer, D. (2006). Distributed practice in verbal recall tasks: A review and quantitative synthesis. *Psychological Bulletin*, 132(3), 354–380. <https://doi.org/10.1037/0033-2909.132.3.354>
- Cepeda, N. J., Vul, E., Rohrer, D., Wixted, J. T., & Pashler, H. (2008). Spacing effects in learning: A temporal ridgeline of optimal retention. *Psychological Science*, 19(11), 1095–1102. <https://doi.org/10.1111/j.1467-9280.2008.02209.x>
- Hinze, S. R., & Rapp, D. N. (2014). Retrieval (sometimes) enhances learning: Performance pressure reduces the benefits of retrieval practice. *Applied Cognitive Psychology*, 28(4), 597–606. <https://doi.org/10.1002/acp.3032>
- Kornell, N. (2009). Optimising learning using flashcards: Spacing is more effective than cramming. *Applied Cognitive Psychology*, 23(9), 1297–1317. <https://doi.org/10.1002/acp.1537>
- Rohrer, D., Dedrick, R. F., Hartwig, M. K., & Cheung, P. (2020). A randomized controlled trial of interleaved mathematics practice. *Journal of Educational Psychology*, 112(1), 40–52. <https://doi.org/10.1037/edu0000367>
- Rohrer, D., Dedrick, R. F., & Stershic, S. (2015). Interleaved practice improves mathematics learning. *Journal of Educational Psychology*, 107(3), 900–908. <https://doi.org/10.1037/edu0000001>
- Yang, C., Li, J., Zhao, W., Luo, L., & Shanks, D. R. (2023). Do practice tests (quizzes) reduce or provoke test anxiety? A meta-analytic review. *Educational Psychology Review*, 35, Article 87. <https://doi.org/10.1007/s10648-023-09801-w>

Part CXV — Tutor Identity & Belonging (Near-Peer Role Models)

Chapter status: Living evidence + tutor hire / match / briefing UX — Researcher tick 2026-08-10 (UTC hour 03; hour%6≠0; researcher count since synthesizer v1.14 = 7 → Researcher)

Primary question: Should MindCraft treat **who the tutor is** (near-peer identity, demographic resemblance, STEM representation) as a belonging intervention — and if so, how without tokenism, superstar demoralization, or waiving the FEI hire bar?

Owners: Ops (tutor pipeline) · Product (match / briefing) · Growth (hire marketing) · Equity · Red Team

Commercial job: Ship **SAFE-ROLEMODEL** densifying SAFE-GENDER × SAFE-ELL × SAFE-TUTORGRAIN × SAFE-HITL × SAFE-EXPECTANCY: **attainable, competence-visible near-peer models + fidelity playbooks** — never Representation Score™, Diversity Minutes NS, costume-only matching, empty “math person like you” praise, or demographic match as substitute for diagnosis.

Builds on: Parts LXV (SAFE-GENDER), LXIV (SAFE-ELL), LXXI (SAFE-TUTORGRAIN), LXVIII (SAFE-HITL), XLVI (SAFE-EXPECTANCY), LII (SAFE-CoP), XXIV (belonging/threat), LXI (SAFE-IDMEASURE). Product seams: Find a Tutor, tutor bios, classroom match, hire briefs, parent “find someone like her” requests.

CXV.1 Why this chapter exists

Parents ask for tutors who “look like” Maya, share her language, or “get” being the only girl in AP Calc. Growth wants representation photos. Ops already defaults to **trained near-peers** (SAFE-TUTORGRAIN). The commercial trap is **identity theater**: splash-page diversity without playbook fidelity, forced demographic match with unattainable superstars, or empty praise that skips recognition-of-competence (SAFE-PROOF / SAFE-GENDER).

FOUNDER BELIEF under audit: Near-peer tutors can inoculate belonging *when* competence is visible and success feels attainable — but matching demographics without FEI structure is costume, not pedagogy.

Claims we refuse as doctrine: Representation Score™ / Diversity Minutes NS; token STEM costume $\equiv$ belonging; demographic match $\equiv$ tutoring quality; celebrity unmatched role models as default; empty “math girl/boy like you” praise $\equiv$ recognition; Belonging Score™ NS; we-close-the-STEM-gap from matching ads; soft-standards for matched pairs.

CXV.2 Constructs (model $\neq$ mentor $\neq$ match $\neq$ mastery)

Keep distinct: **role model** (success updates “people like me can”), **mentor** (ongoing scaffold — Dennehy & Dasgupta), **match** (assignment heuristic), **mastery evidence** (solo transfer / FEI — not likeness). Belonging = membership/acceptance in the domain (Good, Rattan, & Dweck), not a Diversity Minutes counter.

Operational definition (HYPOTHESIS): SAFE-ROLEMODEL compliant when (a) hire still passes SAFE-TUTORGRAIN fidelity drills, (b) bios show **attainable near-peer competence**, (c) optional preference matching without outcome promises, (d) recognition is practice-based (SAFE-PROOF), (e) no Representation Score™ / costume-only marketing / match-waives-challenge.

CXV.3 Stereotype inoculation — contact with same-group experts can buffer STEM self-concept

FACT (stereotype inoculation model): Stout, Dasgupta, Hunsinger, & McManus (2011, *Journal of Personality and Social Psychology*, 100(2), 255–270, doi:10.1037/a0021385) — experiments + longitudinal calculus: exposure to **female STEM experts** promoted more positive implicit STEM attitudes and stronger implicit identification; subjective identification with experts predicted self-efficacy, domain identification, and STEM career commitment. Benefits held even while stereotypes remained active.

FACT (policy synthesis): Dasgupta & Stout (2014, *Policy Insights from the Behavioral and Brain Sciences*, 1(1), 21–29, doi:10.1177/2372732214549471) — access to same-sex STEM role models/mentors is a recurring lever for

girls'/women's STEM self-concept; faculty scarcity is structural, not student-deficit.

Applied (HYPOTHESIS): Near-peer tutors can act as **ingroup experts** only if students identify with them and see domain competence. Stock photos are not inoculation.

Kill: Costume representation without identification pathway.

Survive: Visible near-peer experts as optional belonging inoculators *inside* fidelity ops.

CXV.4 Competence of the model matters more than mere presence

FACT (female role models under threat): Marx & Roman (2002, *Personality and Social Psychology Bulletin*, 28(9), 1183–1193, doi:10.1177/01461672022812004) — three studies: women's math performance under stereotype-relevant conditions was **protected** when a **competent** female experimenter administered the test or when women learned about her math competence; Study 2 showed **perception of competence**, not physical presence alone, safeguarded performance; Study 3: competence information buffered self-appraised math ability, which predicted successful challenging-test performance.

Applied (HYPOTHESIS): Tutor bios and first-session cues should foreground **math competence + near-peer attainability** (courses recently survived, Map diagnosis skill), not demographic checkboxes alone. An incompetent same-group tutor can fail Marx's mechanism and still violate SAFE-TUTORGRAIN.

Kill: Presence-as-enough / “someone who looks like you” without competence signal.

Survive: Competence-visible models; perception of competence as the buffer ingredient.

CXV.5 Attainability — superstars can inspire or demoralize

FACT (superstars and the self): Lockwood & Kunda (1997, *Journal of Personality and Social Psychology*, 73(1), 91–103, doi:10.1037/0022-3514.73.1.91) — relevant outstanding role models provoke **self-enhancement and inspiration when success seems attainable** (time remaining; malleable ability beliefs) but **self-deflation when success seems unattainable** (missed window; fixed ability). Domain relevance is required for impact.

Applied (HYPOTHESIS): Default marketing should prefer **near-peer ladders** (“I took Calc last year; here's how I recovered from blanking on systems”) over Ivy-Olympiad cosplay that makes Maya's success feel closed. Densifies SAFE-TUTORGRAIN (near-peer default) and SAFE-EXPECTANCY (no trait genius theater).

Kill: Unattainable superstar default; genius-résumé dump as belonging.

Survive: Attainable near-peer stories; growth-framed path (Lockwood attainability + Good et al. growth messages below).

CXV.6 Peer mentors at transition — belonging and retention, not grades-as-everything

FACT (field RCT): Dennehy & Dasgupta (2017, *Proceedings of the National Academy of Sciences*, 114(23), 5964–5969, doi:10.1073/pnas.1613117114) — incoming women engineering students (n = 150) randomly assigned female peer mentor, male peer mentor, or none for one year; **female (not male) mentors** protected belonging, self-efficacy, motivation, major retention, and career aspirations; benefits endured a year after mentoring ended. **Year-1 grades were not associated with retention**; belonging and self-efficacy were.

Applied (HYPOTHESIS): Near-peer matching can matter for **identity persistence** beyond session accuracy — especially at transitions. Do not read as “grades don’t matter”; read as “retention levers include belonging.” Mentors were **trained peers in a structured program**, not Discord buddies (SAFE-CoP / SAFE-TUTORGRAIN).

Kill: Grade-only parent dash as the whole retention story; untrained peer chat ≡ mentoring RCT.

Survive: Structured near-peer mentoring as belonging densifier; instrument FEI persistence, not Belonging Score™.

CXV.7 Belonging in math — environment messages, not posters

FACT (sense of belonging to math): Good, Rattan, & Dweck (2012, *Journal of Personality and Social Psychology*, 102(4), 700–717, doi:10.1037/a0026659) — validated belonging-to-math scale predicts intent to pursue math; longitudinal calculus: messages that math ability is **fixed** plus gender stereotypes eroded women’s (not men’s) belonging; belonging mediated future intent and grades; **growth messages** protected belonging under stereotype.

FACT (social belonging intervention): Walton & Cohen (2011, *Science*, 331(6023), 1447–1451, doi:10.1126/science.1198364) — brief social-belonging intervention framing adversity as common and transient raised African American students’ GPA over years and halved the minority achievement gap in their sample; mechanism = preventing adversity from being read as an indictment of belonging.

FACT (cultural stereotypes as gatekeepers): Cheryan, Master, & Meltzoff (2015, *Frontiers in Psychology*, 6, Article 49, doi:10.3389/fpsyg.2015.00049) — stereotypes of who belongs in CS/engineering (male-oriented, isolating, brilliance) reduce girls’ belonging and interest; **broadening representation of people, work, and environments** increases belonging/interest.

Applied (HYPOTHESIS): Tutor identity work sits inside **cue hygiene** (SAFE-GENDER / SAFE-ELL): growth-framed briefs, non-stereotypical story worlds, adversity-as-common onboarding — not a pink tutor SKU. Walton-style belonging messages are **brief construal interventions**, not Diversity Minutes dashboards.

Kill: Belonging Score™ NS; stereotype-heavy “genius tutor” brand; pink-track matching.

Survive: Growth + common-adversity framing; diversified who-represents-math cues; practice recognition over posters.

CXV.8 Product law — hire fidelity first, then optional identity match

HYPOTHESIS (SAFE-ROLEMODEL stack):

1. **Hire bar unchanged** — SAFE-TUTORGRAIN fidelity drills before any matching preference.
2. **Competence-visible bios** — recent path, Map/playbook skill, struggle→recovery; Marx competence signal.

3. **Attainable near-peer default** — Lockwood: avoid demoralizing superstar cosplay in default cards.
4. **Optional preference, not outcome promise** — offer gender/language/background preference where supply allows; never claim match $\Rightarrow$ ACT points or “we close gaps.”
5. **Recognition = competence evidence** — session praise cites attempts/transfer (SAFE-PROOF), not identity slogans.
6. **Brief belonging construal** — adversity-common, growth-capable (Walton; Good et al.); no mandatory identity confession.
7. **Cue hygiene** — Cheryan: who represents the field in photos/story/tutor cards; ban single-mold “math person.”
8. **Densify** — SAFE-GENDER / ELL / TUTORGRAIN / HITL / EXPECTANCY / CoP / IDMEASURE.
9. **Ban Representation Score™ / Diversity Minutes / costume=belonging / match=quality / empty likeness praise / Belonging Score™ / matching ACT guarantees.**

Kill: Match-first marketplace; tokenism marketing.

Survive: Fidelity-first roster with optional inoculating matches and attainable bios.

CXV.9 Competitive contrast (mechanism lens)

Marketplace apps sell demographic filters as quality; Ivy brands sell unattainable stars; “AI who gets you” sells likeness without competence; pink STEM costumes sell tokenism. MindCraft counters with **trained near-peers whose competence is visible**, optional preference matching, growth/belonging construal, and FEI proof — not Representation Score™ theater.

Positioning line (FOUNDER BELIEF $\rightarrow$ copy test): “We match you with tutors who can teach — and when you want someone who shares your path, we make that competence visible, not costume.”

CXV.10 Confidence matrix

Claim	Label	Confidence	Notes
Same-group STEM experts can inoculate women’s STEM self-concept	FACT	High	Stout et al. 2011; Dasgupta & Stout 2014
Competence perception buffers performance under threat	FACT	High	Marx & Roman 2002
Relevant superstars inspire if attainable, demoralize if not	FACT	High	Lockwood & Kunda 1997
Same-gender peer mentors can protect belonging/retention in engineering	FACT	High	Dennehy & Dasgupta 2017 (field RCT; domain=engineering)

Belonging-to-math predicts pursuit; fixed+stereotype erodes women's belonging	FACT	High	Good et al. 2012
Brief belonging construal can raise long-run GPA for marginalized students	FACT	High	Walton & Cohen 2011 (selective college sample; transfer caution)
Stereotypical field cues reduce girls' belonging/interest; broadening helps	FACT	High	Cheryan et al. 2015
HS math tutoring match effects equal college engineering mentor RCT	HYPOTHESIS	Medium	Needs ROLEMODEL-*
Demographic match without fidelity improves FEI	SPECULATION (reject)	—	Conflicts SAFE-TUTORGRAIN
Representation Score™ predicts identity change	SPECULATION (reject)	—	Metric theater

CXV.11 Experiments

ID	Question	Design	Primary
ROLEMODEL-1	Competence-visible near-peer bio vs demographic-only bio (same tutor)	2-arm first-session	belonging proxy; challenge_accept; 7d return
ROLEMODEL-2	Optional gender/language preference match vs random among fidelity-qualified	2-arm (supply-gated)	retention; solo_transfer_pass; preference satisfaction
ROLEMODEL-3	Attainable struggle→skill story vs unattainable superstar résumé	2-arm card	inspiration vs deflation self-report; start rate

ROLEMODEL-4	Growth+common-adversity brief vs no belonging brief (Walton/Good densify)	2-arm onboarding	early quit; retry_120s
ROLEMODEL-5	Parent copy: “someone like her” vs “trained near-peer who can diagnose”	Message / CBC	WTP; misexpectation; choose-fidelity rate
ROLEMODEL-QUAL	10 Maya + 10 parents: what “someone like me” means (gender/race/language/path/age)	Qual	codebook: likeness dimensions vs competence

Falsifier: Demographic-only $\geq$ competence-visible on retention and FEI $\rightarrow$ wound Marx application; still ban costume-only ads if competence-visible wins affect without FEI.

Falsifier: Random among qualified $\geq$ preference match on retention for threatened subgroups $\rightarrow$ demote match priority; keep optional preference as courtesy.

Falsifier: Superstar résumé $\geq$ attainable story on start + 14d retention without deflation — still warn Lockwood for fixed-mindset segments.

Pre-register: ROLEMODEL- before “representation-powered tutoring” or demographic-match ACT campaigns. Densify GENDER- / ELL- / GRAIN- / HITL- — do not fork a Belonging Score™ religion. Family **ROLEMODEL**-owns tutor-identity/belonging fidelity (distinct from GENDER cue hygiene, ELL language load, and TUTORGRAIN hire grain).

CXV.12 So what for MindCraft commercially

- **Copy / product:** Fidelity-first Find-a-Tutor; competence-visible attainable bios; optional preference match; growth/common-adversity brief; recognition cites FEI not likeness. Line: “We match you with tutors who can teach — and when you want someone who shares your path, we make that competence visible, not costume.”
- **Positioning / metric:** Against marketplace filters and Ivy cosplay; track retention + FEI + preference-fill rate — demote Representation Score™, Diversity Minutes, match-NPS-as-quality.
- **Kill list:** Representation Score™; Diversity Minutes NS; token STEM costume $\equiv$ belonging; demographic match $\equiv$ quality; empty likeness praise $\equiv$ recognition; Belonging Score™ NS; matching ACT guarantees; soft standards for matched pairs.
- **Vision:** Near-peer identity is an inoculant inside a diagnosis gym — belonging that survives contact with real work.

CXV.13 Doctrine — SAFE-ROLEMODEL (provisional)

1. **Fidelity before likeness** — SAFE-TUTORGRAIN hire bar never waived for demographic match.
2. **Competence must be perceptible** — Marx & Roman: competence signal, not presence alone.
3. **Prefer attainable near-peers** — Lockwood & Kunda: inspiration requires attainable success.

4. **Ingroup experts can inoculate** — Stout / Dasgupta line; optional same-group exposure as belonging support.
5. **Structured peer mentoring ≠ chat** — Dennehy & Dasgupta: trained mentors; belonging/retention levers beyond year-1 grades.
6. **Belonging via environment messages** — Good et al. growth vs fixed+stereotype; Walton adversity-as-common; Cheryan cue diversification.
7. **Recognition is practice evidence** — SAFE-PROOF / SAFE-GENDER; no empty “math person like you.”
8. **Densify, don’t fork** — SAFE-GENDER / ELL / TUTORGRAIN / HITL / EXPECTANCY / CoP / IDMEASURE.
9. **No Representation Score™ / Diversity Minutes NS / costume≡belonging / match≡quality / Belonging Score™ / matching ACT guarantees.**

Confidence: High that costume-only matching and unattainable-star marketing are unsafe vs belonging+FEI; high that competence-visible near-peers are the right default grain; medium on how strongly HS tutoring preference-match replicates college engineering mentor RCTs — instrument ROLEMODEL-* before hard match SLAs; high to kill Representation Score™ / match≡quality / empty likeness praise regardless.

References (verified)

- Cheryan, S., Master, A., & Meltzoff, A. N. (2015). Cultural stereotypes as gatekeepers: Increasing girls’ interest in computer science and engineering by diversifying stereotypes. *Frontiers in Psychology*, 6, Article 49. <https://doi.org/10.3389/fpsyg.2015.00049>
 - Dasgupta, N., & Stout, J. G. (2014). Girls and women in science, technology, engineering, and mathematics: STEMing the tide and broadening participation in STEM careers. *Policy Insights from the Behavioral and Brain Sciences*, 1(1), 21–29. <https://doi.org/10.1177/2372732214549471>
 - Dennehy, T. C., & Dasgupta, N. (2017). Female peer mentors early in college increase women’s positive academic experiences and retention in engineering. *Proceedings of the National Academy of Sciences*, 114(23), 5964–5969. <https://doi.org/10.1073/pnas.1613117114>
 - Good, C., Rattan, A., & Dweck, C. S. (2012). Why do women opt out? Sense of belonging and women’s representation in mathematics. *Journal of Personality and Social Psychology*, 102(4), 700–717. <https://doi.org/10.1037/a0026659>
 - Lockwood, P., & Kunda, Z. (1997). Superstars and me: Predicting the impact of role models on the self. *Journal of Personality and Social Psychology*, 73(1), 91–103. <https://doi.org/10.1037/0022-3514.73.1.91>
 - Marx, D. M., & Roman, J. S. (2002). Female role models: Protecting women’s math test performance. *Personality and Social Psychology Bulletin*, 28(9), 1183–1193. <https://doi.org/10.1177/01461672022812004>
 - Stout, J. G., Dasgupta, N., Hunsinger, M., & McManus, M. (2011). STEMing the tide: Using ingroup experts to inoculate women’s self-concept in science, technology, engineering, and mathematics (STEM). *Journal of Personality and Social Psychology*, 100(2), 255–270. <https://doi.org/10.1037/a0021385>
 - Walton, G. M., & Cohen, G. L. (2011). A brief social-belonging intervention improves academic and health outcomes of minority students. *Science*, 331(6023), 1447–1451. <https://doi.org/10.1126/science.1198364>
-

Part CXVI — Offline Practice Fidelity (Paper / Whiteboard Capture)

Chapter status: Living evidence + capture UX — Researcher tick 2026-08-10 (UTC hour 06 ≡ Red Team slot, but ch116 never written → prefer Researcher per rotation; researcher count since synthesizer v1.14 = 8 → Researcher override)

Primary question: Can MindCraft treat **paper/whiteboard work** as first-class practice — with photo/OCR into Map — without opening Photomath-style answer dumping, phone attention residue, or OCR≡mastery theater?

Owners: Product (Practice/Map capture) · Privacy · Tutor ops · Growth (anti-answer-app positioning) · Red Team

Commercial job: Ship **SAFE-OFFLINE** densifying SAFE-ATTN × SAFE-HELP × SAFE-HWHELP × SAFE-GESTURE × SAFE-PROOF × SAFE-GENQ: **attempt-first offline work** → **brief capture** → **verified Map grain** — never Capture Score™, Photo Minutes NS, OCR≡mastery, always-on camera, or photo→answer hero.

Builds on: Parts XCVI (SAFE-ATTN), LXXXIX (SAFE-HELP), CIV (SAFE-HWHELP), CV (SAFE-GESTURE), LXXXVI (SAFE-PROOF), LXXII (SAFE-GENQ), XXVI (CLT), LXXXII (SAFE-FORMAT). Product seams: Practice attempt grain, Map updates from photos, Solver vs paper dual-rail, parent “she works on paper” honesty, tutor whiteboard sessions.

CXVI.1 Why this chapter exists

Maya still does most real math on **paper and whiteboard**. Typed equation editors and always-on phone solvers are competitors’ defaults — and traps. Parents want proof she practiced; tutors want her work visible; growth wants “AI that reads your notebook.” The commercial trap is **capture-as-cheating-pipe**: photo features that collapse into Photomath (answer delivery) or phone-open mid-attempt residue (SAFE-ATTN), while marketing pretends OCR fluency is identity change.

FOUNDER BELIEF under audit: Offline media can reduce interface extraneous load and preserve spatial math — *if* capture is post-attempt, verified, and non-solver — not a camera that finishes the problem for her.

Claims we refuse as doctrine: Capture Score™ / Photo Minutes NS; OCR / handwriting recognition ≡ mastery; photo→worked-solution as default Practice rail; always-on camera during attempts; phone-open “study mode” as focus; Photomath-clone GTM; Capture Minutes as North Star; guaranteed ACT from notebook photos.

CXVI.2 Constructs (medium ≠ attempt ≠ capture ≠ diagnosis)

Keep distinct:

Construct	Meaning	Not
Medium	Paper, whiteboard, stylus, keyboard	Pedagogy quality by itself
Attempt grain	Student-constructed steps before reveal	Photo of textbook problem alone
Capture	Image / OCR / tutor photo → system	Instant answer engine

Diagnosis	Map / ontology update from verified work	Fluency of OCR pipeline
------------------	--	-------------------------

Operational definition (HYPOTHESIS): SAFE-OFFLINE compliant when (a) student attempts offline before capture, (b) capture is brief and residue-aware (SAFE-ATTN), (c) recognition errors are corrected before mastery writes (SAFE-GENQ spirit), (d) photo path never defaults to full worked dump (SAFE-HELP / SAFE-HWHELP), (e) no Capture Score™ / OCR≡mastery / Photomath-hero ads.

CXVI.3 Handwriting / pen input — fluency gains ≠ automatic learning gains

FACT (handwriting ITS paradigm): Anthony, Yang, & Koedinger (2012, *International Journal of Human-Computer Studies*, 70(11), 866–887, doi:10.1016/j.ijhcs.2012.04.003) — pen input for algebra can be **faster, smoother, better liked**, and support **transfer to non-computer tasks** vs typing; recognition is unreliable unless **context** reduces errors; design pattern: handwrite process, type final answer; recovery should focus on **problem-solving error**, not OCR theater.

FACT (input method school study): de Morais & Jaques (2022, *Informatics in Education*, doi:10.15388/infedu.2022.03) — 55 seventh-graders; handwriting vs typing on tablet/desktop in step-based math ITS. Handwriting users solved **more equations and were faster**, but **no statistically significant learning difference** vs typing; mismatched device (e.g., handwriting with mouse) **hurt** performance.

Applied (HYPOTHESIS): Prefer paper/pen for **fluency and spatial naturalness**, not as a learning miracle. Value = lower interface friction + medium fidelity — still need FEI. Do not sell “handwriting AI raises scores.”

Kill: Handwriting≡deeper learning marketing; Capture Score™ from speed.

Survive: Pen/paper as default attempt medium; device-matched input; recognition designed around tutoring, not vanity OCR accuracy.

CXVI.4 Longhand encoding — useful caution, not laptop ban religion

FACT (longhand vs laptop notes): Mueller & Oppenheimer (2014, *Psychological Science*, 25(6), 1159–1168, doi:10.1177/0956797614524581) — across three studies, laptop note-takers performed worse on **conceptual** questions than longhand; mechanism claimed = verbatim transcription → shallower processing (even without multitasking).

FACT (replication caution): Morehead, Dunlosky, & Rawson (2019, *Educational Psychology Review*, 31, 753–780, doi:10.1007/s10648-019-09468-2) — direct replication/extension: performance did **not consistently differ**; meta-analysis of direct replications showed small nonsignificant longhand advantages; concluding universal superiority is **premature**.

Applied (HYPOTHESIS): Use Mueller as a **warning against shallow digital transcription**, not as “ban screens / Capture Minutes = science.” MindCraft’s offline bet is **problem-solving medium + attempt grain**, not lecture-note cosplay.

Kill: “Science says paper always beats digital” absolute ads; Pen Minutes NS.

Survive: Prefer constructive longhand *attempts* over verbatim UI clicking; stay open to null medium effects when FEI holds.

CXVI.5 Photo-solve apps — the competitive anti-pattern

FACT (teacher response frame): Webel & Otten (2015, *Mathematics Teacher*, 109(5), 368–373) — Photomath-style photo→answer+steps forces teachers to **ban, restrict, or redesign goals**; “cheating” is about **who gets credit** for the work, not tools existing; ambitious response = tasks where human reasoning remains the point.

FACT (platform-scale speed shift): Sloan-Lynch, Gay, & Watkins (2022, *LAK '22*, doi:10.1145/3506860.3506868) — ~700 unique exercises, **>82 million** submissions; post-photo-app era shows **dramatic shifts toward ultra-fast correct submissions** on susceptible items, contemporaneous with Photomath proliferation — evidence of **lost engagement**, not mastery fireworks.

Applied (HYPOTHESIS): Any MindCraft “snap your work” feature that returns **full solutions before attempt** is a Photomath clone and violates SAFE-HELP / SAFE-HWHELP dual-rail (assisted ≠ ready). Capture must ingest **student work after struggle**, then coach/Map — never answer-first.

Kill: Photo→steps as Practice hero; “AI notebook that finishes homework”; speed-to-correct as success KPI on capturable items.

Survive: Photo as **evidence pipe** for attempt grain; redesign tasks that survive answer apps (Webel Option 3) inside FEI / transfer probes.

CXVI.6 Phone as capture tool = attention residue risk

FACT (attention residue): Leroy (2009, *Organizational Behavior and Human Decision Processes*, 109(2), 168–181, doi:10.1016/j.obhdp.2009.04.002) — unfinished prior tasks leave residue that impairs next-task performance.

FACT (phone presence “brain drain”): Ward, Duke, Gneezy, & Bos (2017, *Journal of the Association for Consumer Research*, 2(2), 140–154, doi:10.1086/691462) — mere presence of one’s smartphone can reduce available cognitive capacity even when unused.

Applied (HYPOTHESIS): Capture UX must be a **short park/resume ritual** (SAFE-ATTN): attempt offline with phone away → open app → snap → confirm → phone down. Mid-attempt camera-as-companion invites Photomath side-door + residue. Tutor-held capture (photograph sheets for mediation) can keep phones out of student hands during solve.

Kill: Always-on phone study companion; mid-attempt DAU pushes to “capture now”; Focus Score™ from camera uptime.

Survive: Post-attempt capture windows; tutor-mediated photo for live sessions; attempt-window protection metrics.

CXVI.7 Recognition / key risk — wrong OCR poisons the graph

FACT (recognition pragmatics): Anthony et al. (2012) — unreliable recognition is the downside of pen input; systems must use **context** and keep recovery on **math error**, not OCR nagging.

FOUNDER BELIEF (product): MindCraft Map writes are high-stakes (mastery, diagnostics). Unverified OCR of messy algebra is a **SAFE-GENQ cousin**: fluent wrongness entering the graph.

Applied (HYPOTHESIS): Capture pipeline: (1) store image as proof artifact (SAFE-PROOF), (2) optional OCR proposal, (3) **student or tutor confirm** before /record-outcomes / Map edge updates, (4) never auto-green from OCR alone. Aligns with hide-correctness diagnostics (C4) — wrong keys corrupt identity evidence.

Kill: OCR accuracy % as product hero; silent auto-mastery from photos.

Survive: Human-in-the-loop confirm; image-as-proof even when OCR fails; GENQ-style verify-before-ship for capture.

CXVI.8 Product law — attempt → capture → confirm → Map

HYPOTHESIS (SAFE-OFFLINE stack):

1. **Offline-first attempt** — paper/whiteboard (or stylus) before typed editor when format allows; respect spatial math (SAFE-FORMAT / SAFE-GESTURE draw).
2. **Commit before capture** — prediction / answer commit or time-on-attempt gate before photo tools unlock (SAFE-PREDICT / SAFE-HELP).
3. **Brief capture ritual** — open → snap → confirm → close; phone away during solve (SAFE-ATTN).
4. **Verify before graph write** — student/tutor confirm OCR; image stored as competence evidence (SAFE-PROOF / SAFE-GENQ).
5. **No answer-dump path** — capture feeds diagnosis/coach chips, not Photomath steps-first (SAFE-HWHELP dual-rail).
6. **Tutor whiteboard parity** — session photos are AAR/film clips for soft-wrong review (SAFE-AAR / SAFE-FILM), not shame reels.
7. **Parent honesty** — show aged proof artifacts (“worked on paper Tue”) without tonight-% stalk (SAFE-PDASH).
8. **Densify** — SAFE-ATTN / HELP / HWHELP / GESTURE / PROOF / GENQ / FORMAT / FADE.
9. **Ban** Capture Score™ / Photo Minutes / OCR≡mastery / always-camera / photo→answer hero / Capture ACT guarantees.

Kill: Notebook-AI answer brand; camera-open Practice.

Survive: Paper-native gym with inspectable capture into Map.

CXVI.9 Competitive contrast (mechanism lens)

Photomath/Chegg sell **photo→answer**; typed ITS sell **keyboard fluency**; Duo sells streaks on-device; “AI notebook” startups sell OCR magic. MindCraft counters with **offline attempt fidelity + verified capture into diagnosis** — work that looks like school math and updates the Map without finishing the problem for Maya.

Positioning line (FOUNDER BELIEF → copy test): “Practice on paper. Snap when you’re done. We read your work to teach the next move — not to finish the problem for you.”

CXVI.10 Confidence matrix

Claim	Label	Confidence	Notes
Pen/handwriting can improve speed/fluency vs typing in algebra ITS	FACT	High	Anthony et al. 2012; de Moraes & Jaques 2022
Handwriting input automatically raises learning gains vs typing	SPECULATION (reject as doctrine)	—	Both studies: null learning difference
Longhand notes always beat laptop for conceptual learning	FACT weakened	Medium–low	Mueller 2014 vs Morehead et al. 2019
Photo-solve apps correlate with ultra-fast correct submissions at scale	FACT	High	Sloan-Lynch et al. 2022 (observational; not pure RCT)
Photo→answer defaults undermine attempt grain / FEI	HYPOTHESIS	High	Densifies HELP/HWHELP; Webel & Otten frame
Phone-open mid-attempt capture harms available capacity via residue/presence	HYPOTHESIS	Medium–high	Leroy; Ward et al.; needs OFFLINE-*
Unverified OCR→Map writes are unsafe	HYPOTHESIS	High	Anthony recognition risk + GENQ analogy
Capture Score™ predicts identity change	SPECULATION (reject)	—	Metric theater

CXVI.11 Experiments

ID	Question	Design	Primary
OFFLINE-1	Paper-first attempt + post-commit capture vs typed-only Practice	2-arm session	solo_transfer_pass; retry_120s; time-to-first-attempt

OFFLINE-2	Capture unlock after commit vs always-available camera	2-arm	photo-solve side-door rate; attempt duration; challenge_accept
OFFLINE-3	Student/tutor OCR confirm vs auto-write Map	2-arm	false mastery rate; correction latency; trust repair
OFFLINE-4	Phone-away solve + brief capture ritual vs phone-present companion	2-arm	attempt_interrupted; transfer; self-report residue
OFFLINE-5	Parent copy: “snap homework finished” vs “snap your attempt for the Map”	Message / CBC	WTP; misexpectation; choose-attempt framing
OFFLINE-QUAL	10 Maya: where they actually work (paper/white/phone) + Photomath temptation moments	Qual	codebook: medium + cheat-temptation + capture trust

Falsifier: Typed-only $\geq$ paper-first on transfer *and* interruption with equal attempt quality $\rightarrow$ demote offline-default; keep optional paper path.

Falsifier: Always-available camera $\leq$ post-commit unlock on Photomath side-door *and* FEI — still ban answer-dump capture.

Falsifier: Auto-OCR Map writes match confirm-arm false-mastery rate at production OCR accuracy — still keep image proof; revisit confirm friction.

Pre-register: OFFLINE- before “AI reads your notebook” or Photomath-competitor campaigns. Densify ATTN- / HELP- / HWHELP- / GENQ- — do not fork a Capture Score™ religion. Family **OFFLINE-** owns paper/whiteboard capture fidelity (distinct from ATTN phone presence, HELP contingent hints, and HWHELP Solver dual-rail).

CXVI.12 So what for MindCraft commercially

- **Copy / product:** Offline-first Practice; post-attempt snap; confirm-before-Map; never photo $\rightarrow$ full solution as default. Line: “Practice on paper. Snap when you’re done. We read your work to teach the next move — not to finish the problem for you.”
- **Positioning / metric:** Against Photomath/Chegg answer cameras and typed-only ITS; track solo transfer + side-door rate + confirm corrections — demote Capture Score™, Photo Minutes, OCR accuracy hero, speed-to-correct on capturable items.
- **Kill list:** Capture Score™; Photo Minutes NS; OCR $\equiv$ mastery; always-on camera; photo $\rightarrow$ answer Practice hero; Capture ACT guarantees; phone-companion focus theater.
- **Vision:** The math gym includes the desk and the whiteboard — capture is how evidence returns to the Map without outsourcing the struggle.

CXVI.13 Doctrine — SAFE-OFFLINE (provisional)

1. **Medium serves attempt grain** — paper/whiteboard preferred for spatial/natural solve; not a learning talisman.
2. **Fluency ≠ learning gain** — Anthony / de Morais: speed and liking without automatic posttest superiority.
3. **Capture after commit** — no camera-first answer path (Webel; Sloan-Lynch speed pathology).
4. **Residue-aware ritual** — brief snap, then phone down (Leroy; Ward; SAFE-ATTN).
5. **Verify before graph write** — OCR proposes; human confirms; image is proof (SAFE-GENQ / SAFE-PROOF).
6. **Dual-rail with Solver** — instrumental help elsewhere; capture is evidence, not Chegg (SAFE-HWHELP).
7. **Longhand caution, not religion** — Mueller insight + Morehead nulls; FEI adjudicates.
8. **Densify, don't fork** — SAFE-ATTN / HELP / HWHELP / GESTURE / PROOF / GENQ / FORMAT.
9. **No Capture Score™ / Photo Minutes / OCR≡mastery / always-camera / photo→answer hero / Capture ACT guarantees.**

Confidence: High that photo→answer capture is commercially and pedagogically unsafe vs FEI; high that unverified OCR must not write mastery; medium that paper-first beats typed-only on transfer once attempt quality is equated — instrument OFFLINE-*; high to kill Capture Score™ / OCR≡mastery / Photomath-hero regardless.

References (verified)

- Anthony, L., Yang, J., & Koedinger, K. R. (2012). A paradigm for handwriting-based intelligent tutors. *International Journal of Human-Computer Studies*, 70(11), 866–887. <https://doi.org/10.1016/j.ijhcs.2012.04.003>
 - de Morais, F., & Jaques, P. A. (2022). Does handwriting impact learning on math tutoring systems? *Informatics in Education*. <https://doi.org/10.15388/infedu.2022.03>
 - Leroy, S. (2009). Why is it so hard to do my work? The challenge of attention residue when switching between work tasks. *Organizational Behavior and Human Decision Processes*, 109(2), 168–181. <https://doi.org/10.1016/j.obhdp.2009.04.002>
 - Morehead, K., Dunlosky, J., & Rawson, K. A. (2019). How much mightier is the pen than the keyboard for note-taking? A replication and extension of Mueller and Oppenheimer (2014). *Educational Psychology Review*, 31, 753–780. <https://doi.org/10.1007/s10648-019-09468-2>
 - Mueller, P. A., & Oppenheimer, D. M. (2014). The pen is mightier than the keyboard: Advantages of longhand over laptop note taking. *Psychological Science*, 25(6), 1159–1168. <https://doi.org/10.1177/0956797614524581>
 - Sloan-Lynch, J., Gay, N., & Watkins, R. (2022). Too fast for their own good: Analyzing a decade of student exercise responses to explore the impact of math solving photo apps. In *Proceedings of the 12th International Learning Analytics and Knowledge Conference (LAK '22)* (pp. 276–283). ACM. <https://doi.org/10.1145/3506860.3506868>
 - Ward, A. F., Duke, K., Gneezy, A., & Bos, M. W. (2017). Brain drain: The mere presence of one's own smartphone reduces available cognitive capacity. *Journal of the Association for Consumer Research*, 2(2), 140–154. <https://doi.org/10.1086/691462>
 - Webel, C., & Otten, S. (2015). Teaching in a world with Photomath. *Mathematics Teacher*, 109(5), 368–373.
-

Part CXVII — Worked-Example Fading Under High Math Anxiety

Chapter status: Living evidence + Solver/Practice UX brief — Researcher tick 2026-08-10 (UTC hour 12 ≡ Red Team slot, but ch117 never written → prefer Researcher per rotation; researcher count since synthesizer v1.15 = 0)

Primary question: How should MindCraft fade worked examples for high-math-anxiety (HMA) learners without trapping them in permanent full-solution “support,” flooding them into solo struggle, or inventing Anxiety Score™ / Calm Score™ fade gates?

Owners: Affect / Anxiety · Product (Solver / Practice fade ladder) · HITL tutors · Brand · Red Team

Commercial job: Ship **SAFE-FADEX** densifying SAFE-FADE × SAFE-EXPOSE × SAFE-DD: **examples protect WM under anxiety load; fade is still the identity move; hold longer at E0/E1 under destaked stakes, then approach fade as expectancy violation** — never permanent never-fade anxiety track, never day-one aggressive fade-as-rigor, never Calm Score™ unlocks.

Builds on: Parts LXXXVIII (SAFE-FADE), LVIII (SAFE-EXPOSE), XLI (SAFE-DD), XXIV (affect), L (SAFE-CALIB), C (SAFE-EXPTIME), XCI (SAFE-HINT), CIV (SAFE-HWHELP). Product seams: Solver E0→E3 ladder, soft-wrong / destake, hint binge as safety behavior, Map grain, parent “she needs every step forever” messaging.

CXVII.1 Why this chapter exists

SAFE-FADE already kills always-full-worked and never-fade costumes. SAFE-EXPOSE already kills day-one flooding and permanent safety behaviors. The remaining commercial hazard is the **intersection:** HMA Maya’s WM is already taxed (Ashcraft), so full worked examples help her learn (Mesghina et al.), *and* those same full dumps become the safety behavior that blocks inhibitory learning and solo identity. Competitors resolve the tension wrongly: ChatGPT/Photomath-style apps keep the dump forever (“we support anxiety”); struggle-theater apps yank scaffolding on day one (“rigor”); therapy-cosplay apps wait for Calm Score™ before hard steps.

FOUNDER BELIEF under audit: For HMA learners, worked examples are *equipment*, not therapy — and fading is *graded approach to carrying the join*, not abandonment. Permanent full examples are commercially identical to ChatGPT; permanent easy-track “for anxious kids” is soft standards dressed as care.

Claims we refuse as doctrine:

1. Always-full-worked forever for HMA / “anxious track.”
 2. Aggressive E0→E3 jump as day-one rigor for threatened learners.
 3. Anxiety Score™ / Calm Score™ as fade gate or North Star.
 4. Permanent hint binge / solution peek labeled “support.”
 5. Never-fade scaffolding marketed as empathy.
 6. Fade Score™ / Anxiety-adjusted Fade % vanity NS.
 7. Softened standards / forever E0 for matched anxiety personas.
 8. ACT-point guarantees from “anxiety-aware scaffolding.”
-

CXVII.2 Constructs (protect ≠ trap; fade ≠ flood)

Construct	Meaning	MindCraft analogue	Failure mode
Anxiety WM load	Worry consumes resources needed for multi-step math	Freeze, peek, rapid exit	Read as laziness / “needs more struggle”
Protective example	Full/near-full worked study that frees WM for schema	E0 / early E1 learn rail	Costume forever
Safety behavior	Escape that prevents expectancy violation	Instant full reveal; hint binge	“Support” that blocks learning
Anxiety-aware fade	Hold longer at high guidance; still step down under destake	Slower E0→E1→E2; prove rail near E3	Forever track OR flood
SAFE-FADEX	FADE × EXPOSE doctrine for HMA	This chapter	Calm Score™ theater

Operational definition (HYPOTHESIS): A fade path is *SAFE-FADEX complete* when it (a) offers **protective worked study** early when anxiety load + low schema would make unguided search undesirable difficulty (SAFE-DD), (b) **destakes** evaluation climate (soft-wrong; no streak-as-worth), (c) **holds** at E0/E1 longer than for low-anxiety peers *without* locking a permanent persona track, (d) treats consecutive full peeks as **safety behaviors to fade** (SAFE-EXPOSE / SAFE-HINT), (e) advances fade as **approach + expectancy check** (“I can finish the last join”), (f) proves with *solo_transfer_pass* / delayed mix — not calm meters or solution views, and (g) never claims to treat clinical anxiety disorders.

CXVII.3 Anxiety taxes the same WM that search needs

FACT (WM disruption): Ashcraft & Kirk (2001, *Journal of Experimental Psychology: General*, 130(2), 224–237, doi:10.1037/0096-3445.130.2.224) — high math-anxiety individuals show reduced working-memory spans under math demand; anxiety functions as a **transitory dual-task load**, amplifying errors especially on carry/complex items.

FACT (education bridge + avoidance): Ashcraft & Krause (2007, *Psychonomic Bulletin & Review*, 14(2), 243–248, doi:10.3758/BF03194059) — math beyond simple retrieval depends on WM; high math anxiety drains resources via fear preoccupation; pattern includes avoidance of elective math/careers. Hembree (1990, *JRME*, 21(1), 33–46, doi:10.5951/jresmetheduc.21.1.0033) — MA binds to **avoidance**; relief treatments tend to accompany better performance.

Applied (HYPOTHESIS): For HMA Maya, unguided search on a new schema is often **undesirable difficulty** (SAFE-DD): freeze + peek + “I’m not a math person.” Protective examples are load management for germane processing — not softness.

Kill: “Struggle first always” for anxious learners.

Survive: Example-first acquisition when anxiety load + low prior knowledge would saturate WM.

CXVII.4 Worked examples can blunt anxiety's learning tax

FACT (direct moderation): Mesghina, Vollman, Trezise, & Richland (2023, *Journal of Educational Psychology*, doi:10.1037/edu0000795) — 280 fifth-graders, demanding ratio lesson; half reviewed worked examples at key solve moments. **Math anxiety** (not pandemic worries) hurt procedural/conceptual learning, raised test mind-wandering, lowered interest; **worked examples mitigated those hits on learning and engagement**.

Applied (FOUNDER BELIEF → testable): Learn-rail E0/E1 is commercially *required* for HMA acquisition — not optional. Withholding examples to look rigorous sells threat; never leaving examples sells ChatGPT.

Kill: No-example / discovery-only brand for threatened novices.

Survive: Protective WE on high-load bridge lessons; measure peek/exit reduction *without* forever-E0.

Wound: Mesghina tested WE *presence*, not multi-session fade pacing — FADEX-* must test hold-time.

CXVII.5 Fading still works — and prior knowledge, not WM score, often gates it

FACT (smooth fade): Renkl, Atkinson, Maier, & Staley (2002, *Journal of Experimental Education*, 70(4), 293–315, doi:10.1080/00220970209599510) — successive incomplete examples (complete → increasingly incomplete → to-be-solved) beat static example–problem pairs on near transfer; errors during learning mediated; **backward fading** often preferred.

FACT (CLT transition): Renkl & Atkinson (2003, *Educational Psychologist*, 38(1), 15–22, doi:10.1207/S15326985EP3801_3) — as intrinsic load falls with schema growth, problem-solving demands can rise without overload; fading structures that transition.

FACT (SE + fade): Atkinson, Renkl, & Merrill (2003, *Journal of Educational Psychology*, 95(4), 774–783, doi:10.1037/0022-0663.95.4.774) — fading + principle SE prompts → medium-to-large near/far transfer without extra time.

FACT (WM moderation null; prior knowledge matters): Miller-Cotto et al. (2025/2026, *British Journal of Educational Psychology*, 96(1), 108–126, doi:10.1111/bjep.12781; PMC12879535) — sixth-graders, ASSISTments geometry: **fading** largest pre→post gains vs solve-only; **WM did not significantly moderate** fading; **prior knowledge** did.

Applied (HYPOTHESIS): Do **not** invent WM Score™ / Anxiety Score™ fade destiny. Gate on **diagnosed grain / prior success** with **destake + longer hold** under threat — not a permanent anxiety persona. Mesghina still supports *temporary* WE under anxiety arousal.

Kill: Anxiety Score™ / WM Score™ as fade destiny; forever-E0 “support track.”

Survive: Evidence-state fade (SAFE-EXPTIME / SAFE-FADE) + anxiety-aware *pacing*, not a soft curriculum.

CXVII.6 Incorrect examples and soft-wrong — protect without answer dump

FACT (incorrect WE for low prior knowledge): Barbieri & Booth (2016, *Learning and Individual Differences*, 48, 36–44, doi:10.1016/j.lindif.2016.04.001) — algebra; **low prior knowledge** students gained most from highlighted incorrect examples vs solve-only or correct-example controls.

Applied (HYPOTHESIS): For HMA + low schema, **destaked error study** (soft-wrong; explain-the-error) builds features without unaided flood or ChatGPT dump. Mid-ladder: “what went wrong on this step?”

Kill: Shame-reel mistakes; timed ACT error flood on day one.

Survive: Incorrect-example / soft-wrong when full solo is still threat.

CXVII.7 Safety behaviors: permanent examples are exposure’s enemy

FACT (inhibitory learning): Craske, Treanor, Conway, Zbozinek, & Vervliet (2014, *Behaviour Research and Therapy*, 58, 10–23, doi:10.1016/j.brat.2014.04.006) — exposure optimized via expectancy violation and removal of safety signals; success is new inhibitory learning, **not** within-session fear reduction to zero.

Applied (reuse SAFE-EXPOSE): Instant full Solver dump, parent rescue, and streak-protect peeks are **safety behaviors**. For HMA product design, protective E0 is allowed *as a stage*; the commercial sin is making E0 the forever brand. Fade-up after equipped success is the expectancy violation: “I feared blanking on the last join — I carried it under recoverable stakes.”

Kill: Habituation theater (“wait until calm, then fade”); permanent peek entitlement.

Survive: Name feared outcome → protective study → completion → solo under destake → vary context/format.

CXVII.8 Product surface — SAFE-FADEX claim contract

Surface	Required behavior	Banned substitute
HMA / high-threat learn open	Prefer E0 or E1 completion; destake; no timer flood	Day-one E3 struggle theater
After equipped solos	Fade up (less guidance) even if anxiety self-report remains	Forever-E0 “anxious track”
After freeze / binge peek	One-stage fade-down + name missing grain; treat binge as safety	Permanent max dump “because anxiety”
Exam / prove rail	Near E3; KR timing (SAFE-FBTIME); no learn-rail full dump under stakes	Cram keys as method
Parent copy	“We protect early steps, then fade as she proves the join”	“We always show every step for anxious kids” / “we never help”

Metrics	fade_stage, hold-time at E0/E1, fade_step_attempt, peek-binge, solo_transfer_pass, avoided_format_approach	Anxiety Score™ / Calm Score™ / Fade Score™ NS
---------	--	---

Competitive foil: ChatGPT = forever dump. Struggle apps = flood. Therapy apps = calm meter. MindCraft = protective ladder that still graduates her.

CXVII.9 Doctrine — SAFE-FADEX (provisional)

- 1. Protect then fade** — for HMA + low schema, full/near-full examples are acquisition equipment (Mesghina; Ashcraft load), not the brand forever.
- 2. Destake before demand** — soft-wrong / mastery climate so fade is not threat (SAFE-DD × SAFE-EXPOSE).
- 3. Hold longer, don’t soft-track** — longer E0/E1 dwell vs low-anxiety peers; no permanent “anxious curriculum.”
- 4. Gate on grain, not meters** — prior success / Map evidence / recent solos; no Anxiety Score™ / Calm Score™ fade unlock.
- 5. Treat peek binge as safety behavior** — fade the escape, don’t monetize it (SAFE-HINT / SAFE-HELP).
- 6. Backward completion first** — blank final joins before full solo (SAFE-FADE); optional incorrect-example mid-rungs for low prior knowledge (Barbieri & Booth).
- 7. Proof = solo transfer + approach** — solo_transfer_pass / delayed mix / avoided_format_approach; not calm-to-zero or solution views.
- 8. Copy:** “We shield the first steps when worry loads the mind — then we fade as you prove the join.” Never forever-every-step; never abandon-to-prove-grit; never clinical cure claims.

Confidence: High — MA taxes WM / fuels avoidance (Ashcraft; Hembree); WE can blunt MA’s learning hit (Mesghina); fading beats static pairs (Renkl/Atkinson). Medium — WM is not the product fade gate (Miller-Cotto null); longer E0/E1 hold for HMA beats aggressive fade (needs FADEX-*). High — forever-E0 “anxiety support” ≡ ChatGPT; Calm/Anxiety Score™ gates ≡ therapy cosplay.

CXVII.10 Experiments

ID	Question	Design	Primary
FADEX-1	Longer E0/E1 hold for high self-reported MA vs matched aggressive fade	A/B stratified by MA	solo_transfer_pass; peek binge; exit
FADEX-2	Protective WE present vs absent on high-load bridge lesson for high-MA	A/B	Near transfer; mind-wander proxy; interest

FADEX-3	Incorrect-example mid-rung vs correct-only fade for low prior + high MA	A/B	Feature knowledge; transfer
FADEX-4	Destaked soft-wrong fade vs timed high-stakes fade for high MA	A/B	Approach rate; transfer; freeze
FADEX-5	Parent CBC: protect-then-fade vs always-every-step vs struggle-first	CBC	WTP; trust (SAFE-WTP)
FADEX-QUAL	10 Maya HMA: when does a shown step feel like shield vs theft of the win?	Qual	FADEX UX codebook

Falsifier: Aggressive matched fade equals longer-hold on solo transfer *and* approach for high-MA → shorten default hold; keep protect-on-first-acquisition.

Falsifier: Longer hold never yields fade-up / solo gain → revise equipment, don't restore permanent dump brand.

Falsifier: Anxiety meter gating outperforms grain-based fade → still ban meter as NS; segregate research tool from product hero.

Pre-register: FADEX-* before “anxiety-aware always-show-every-step” or “anxiety-rigor struggle-first” campaigns (SAFE-LABMETA).

Family note: FADEX- *densifies* FADE- and EXP-O-*; do not collapse into FADE-1 or Calm Score™ pilots.

CXVII.11 So what for MindCraft commercially

- **Copy:** “We shield early steps when worry loads the mind — then we fade as you prove the join.” Lead with protect→fade, not forever solutions or grit flooding.
 - **Product:** Anxiety-aware *pacing* on the existing E0→E3 ladder; longer hold under destake; no separate soft track; binge peek = safety behavior; prove rail stays near E3.
 - **Positioning:** Against ChatGPT dump-as-empathy, struggle-theater rigor, and Calm Score™ therapy apps; for graduated independence under recoverable stakes.
 - **Metric:** hold-time at E0/E1, fade_stage, fade_step_attempt, peek-binge, solo_transfer_pass, avoided_format_approach — demote Anxiety Score™ / Calm Score™ / Fade Score™.
 - **Kill list:** Forever-E0 anxious track; day-one aggressive fade flood; meter-gated fade; permanent peek support; ACT guarantees from scaffolding theater.
 - **Growth:** Parent decks sell *earned independence with early protection*; tutors QA protect→completion→solo, not monologue dump or abandonment (SAFE-HITL × SAFE-TALK).
 - **Vision:** Thirty-year identity company teaches threatened Maya that help can graduate her — shelter for schema, then the join she carries alone.
-

References (verified)

- Ashcraft, M. H., & Kirk, E. P. (2001). The relationships among working memory, math anxiety, and performance. *Journal of Experimental Psychology: General*, 130(2), 224–237. <https://doi.org/10.1037/0096-3445.130.2.224>
- Ashcraft, M. H., & Krause, J. A. (2007). Working memory, math performance, and math anxiety. *Psychonomic Bulletin & Review*, 14(2), 243–248. <https://doi.org/10.3758/BF03194059>
- Atkinson, R. K., Renkl, A., & Merrill, M. M. (2003). Transitioning from studying examples to solving problems: Effects of self-explanation prompts and fading worked-out steps. *Journal of Educational Psychology*, 95(4), 774–783. <https://doi.org/10.1037/0022-0663.95.4.774>
- Barbieri, C., & Booth, J. L. (2016). Support for struggling students in algebra: Contributions of incorrect worked examples. *Learning and Individual Differences*, 48, 36–44. <https://doi.org/10.1016/j.lindif.2016.04.001>
- Craske, M. G., Treanor, M., Conway, C. C., Zbozinek, T., & Vervliet, B. (2014). Maximizing exposure therapy: An inhibitory learning approach. *Behaviour Research and Therapy*, 58, 10–23. <https://doi.org/10.1016/j.brat.2014.04.006>
- Hembree, R. (1990). The nature, effects, and relief of mathematics anxiety. *Journal for Research in Mathematics Education*, 21(1), 33–46. <https://doi.org/10.5951/jresmetheduc.21.1.0033>
- Mesghina, A., Vollman, E., Trezise, K., & Richland, L. E. (2023). Worked examples moderate the effect of math learning anxiety on children’s math learning and engagement during the COVID-19 pandemic. *Journal of Educational Psychology*. <https://doi.org/10.1037/edu0000795>
- Miller-Cotto, D., et al. (2025/2026). Does working memory moderate the effect of fading on math performance? *British Journal of Educational Psychology*, 96(1), 108–126. <https://doi.org/10.1111/bjep.12781>
- Renkl, A., & Atkinson, R. K. (2003). Structuring the transition from example study to problem solving in cognitive skill acquisition: A cognitive load perspective. *Educational Psychologist*, 38(1), 15–22. https://doi.org/10.1207/S15326985EP3801_3
- Renkl, A., Atkinson, R. K., Maier, U. H., & Staley, R. (2002). From example study to problem solving: Smooth transitions help learning. *Journal of Experimental Education*, 70(4), 293–315. <https://doi.org/10.1080/00220970209599510>

Part CXVIII — Parent Message Framing After Soft-Wrong Streaks

Chapter status: Living evidence + parent-comms / GTM brief — Researcher tick 2026-08-10 (UTC hour 15; hour%6≠0; researcher count since synthesizer v1.15 = 1 → Researcher)

Primary question: After a run of destaked soft-wrongs, what should MindCraft say to parents — and what must it refuse to say — so honesty does not become nag, shame, growth-mindset poster theater, or controlling homework scripts?

Owners: Brand / Parent GTM · Product (parent digest / SAFE-PDASH) · Affect · HITL tutors · Red Team

Commercial job: Ship **SAFE-SOFTMSG** densifying SAFE-PDASH × SAFE-FBLIT × SAFE-ERRCLIMATE: **name the grain, invite one autonomy-supportive next attempt, frame failure as enhancing without empty praise — never miss-streak shame, live wrong-answer stalk, or “mistakes are magic” without a named return.**

Builds on: Parts LXXXIII (SAFE-PDASH), CXII (SAFE-FBLIT), CIII (SAFE-ERRCLIMATE), XLIX (SAFE-MISCON), LXXIX (SAFE-FBTIME), LIX (SAFE-WTP), XLIV (SAFE-REWARD), LVIII (SAFE-EXPOSE),

CXVII (SAFE-FADEX parent CBC). Product seams: soft-wrong hide-correctness, parent digest, proof-age strip, improvement CTA, no Family Leaderboard as learning proof.

CXVIII.1 Why this chapter exists

SAFE-PDASH already kills tonight-% / streak / shame-rank parent portals. SAFE-FBLIT already kills delivery=uptake. The remaining commercial hazard is the **message after a soft-wrong streak**: Maya misses three destaked items; the product must tell a paying parent something true without (a) converting soft-wrong into a **shame telemetry product**, (b) pushing controlling “sit with her until it’s fixed” scripts that transmit parent math anxiety (Maloney; Retanal), or (c) shipping empty growth-mindset posters that change parent beliefs but not practices (MacDonald et al.).

Competitors resolve this wrongly: gradebook apps push miss counts and ranks; “encouragement” apps send “mistakes are good!” with no next grain; homework-stalk apps invite parents to watch every wrong live; mindset kits sell belief change as the product.

FOUNDER BELIEF under audit: Soft-wrong is for the learner’s climate. Parent copy after soft-wrong must be **failure-enhancing + autonomy-supportive + one named return** — not a correctness scoreboard and not a vibe.

Claims we refuse as doctrine:

1. Soft-wrong streak push / Miss Streak as parent North Star.
 2. Live wrong-answer stalk digests.
 3. Shame/guilt SMS (“she’s falling behind”) after destaked practice.
 4. Empty “mistakes are good!” / growth-mindset poster without named next attempt.
 5. Controlling “fix it tonight / check every problem” scripts as default parent help.
 6. Parent Correctness Score™ / Soft-Wrong Score™ vanity metrics.
 7. Growth-mindset parent letter ≡ FEI / ACT lift.
 8. ACT-point guarantees from parent messaging kits.
-

CXVIII.2 Constructs (honesty ≠ nag; destake ≠ hide from parents)

Construct	Meaning	MindCraft analogue	Failure mode
Soft-wrong streak	Consecutive destaked misses / incomplete joins without shame KR theater	Practice miss run under hide-correctness or soft mark	Parent sees only “failing”
Failure-is-enhancing frame	Struggle can inform learning (vs debilitating = threat to ability)	Digest copy: “this miss shows the next join”	Empty poster; or panic

Autonomy-supportive invite	Encourage independent next attempt; avoid controlling language	One return CTA at grain	Intrusive “do it with me now”
Intrusive / controlling support	Uninvited monitoring, pressure, unwanted solutions	Live stalk; “sit until correct”	Achievement drag; anxiety transmission
SAFE-SOFTMSG	Parent message contract after soft-wrong runs	This chapter	Miss-streak theater

Operational definition (HYPOTHESIS): A parent message after soft-wrong is *SAFE-SOFTMSG complete* when it (a) **names grain** (concept / ingredient / FormatId / bridge — not a vibe), (b) states that the miss is **usable information** under destake (failure-enhancing, not ability verdict), (c) offers **one** autonomy-supportive next attempt or return mission (SAFE-PDASH improvement CTA), (d) does **not** expose live item keys, ranks, or miss-streak vanity, (e) steers anxious parents away from controlling takeover toward structure + invitation (Retanal; Jiang), and (f) never claims messaging alone raises ACT scores.

CXVIII.3 Intrusive involvement hurts math — especially under control

FACT (math-specific meta): Jiang, Shi, Zheng, & Mao (2023, *Frontiers in Psychology*, 14, 1218534, doi:10.3389/fpsyg.2023.1218534; N≈16,338) — supportive parental homework involvement (SPI) links weakly positively to math achievement ($r \approx 0.076$); **intrusive involvement (IPI) links negatively** ($r \approx -0.153$). SPI effects strongest when measured as **autonomy support**; IPI harm strongest in high school.

FACT (affective homework climate): Wu, Barger, Oh, & Pomerantz (2022, *Child Development*, 93(5), 1347–1364, doi:10.1111/cdev.13774) — parents’ math homework involvement is more affectively negative than math activities; **more negative homework affect** predicts poorer later math motivation and achievement.

Applied (HYPOTHESIS): Soft-wrong streak digests that spike parental anxiety and invite intrusive checking recreate the IPI path. Commercial default must be **autonomy-supportive structure**, not frequency-of-help or live monitoring.

Kill: Live wrong-answer stalk; “check every miss tonight” as product default.

Survive: Sparse, grain-named, improvement-framed digests (Kraft & Rogers lineage via SAFE-PDASH).

CXVIII.4 Math-anxious parents + controlling help transmit harm

FACT (intergenerational homework channel): Maloney, Ramirez, Gunderson, Levine, & Beilock (2015, *Psychological Science*, 26(9), 1480–1488, doi:10.1177/0956797615592630) — when parents are more math-anxious, children learn less math and show more math anxiety **only if** those parents report **frequent math homework help**; less-frequent help breaks the association; reading achievement unaffected.

FACT (mechanism = controlling style): Retanal, Johnston, Di Lonardo Burr, Storozuk, DiStefano, & Maloney (2021, *Education Sciences*, 11(10), 620, doi:10.3390/educsci11100620) — parents’ math anxiety relates to both autonomy- and controlling-supportive styles; **controlling-supportive help partially mediates** the anxiety→child achievement

link; autonomy support did not mediate in their path model.

Applied (FOUNDER BELIEF → testable): After soft-wrong streaks, MindCraft must **not** recruit math-anxious parents into more frequent controlling help. Message should reduce pressure (“you do not need to re-teach every miss”) and point to Map grain / tutor / next solo attempt.

Kill: Parent-as-substitute-teacher after every soft-wrong.

Survive: “Name the join; invite her next attempt; escalate to tutor/Map — don’t take the pencil.”

CXVIII.5 Uninvited help + fixed mindset = sharper achievement drag

FACT (moderation): Park, Gunderson, Maloney, Tsukayama, Beilock, Duckworth, & Levine (2023, *Developmental Psychology*, 59(7), 1249–1267, doi:10.1037/dev0001522) — across early elementary (N=563) and secondary (N=1,613) samples, **intrusive homework support more strongly predicted achievement decline for children with fixed mindsets.**

Applied (HYPOTHESIS): Soft-wrong streak SMS that imply “she can’t” plus invite unsolicited takeover are especially toxic for fixed-mindset Maya. Parent copy must avoid ability verdicts and uninvited monitoring — even when the digest is “honest.”

Kill: Ability-framed miss digests (“shows she’s not ready”).

Survive: Process/grain frame + invited support only.

CXVIII.6 Failure mindsets beat intelligence-mindset posters

FACT (visible failure beliefs): Haimovitz & Dweck (2016, *Psychological Science*, 27(6), 859–869, doi:10.1177/0956797616639727) — parents’ **failure mindsets** (failure-as-enhancing vs debilitating) predict parenting practices and children’s intelligence mindsets more than parents’ own intelligence mindsets; children accurately perceive failure mindsets; parents who see failure as debilitating focus on performance/ability rather than learning.

FACT (belief change ≠ practice change): MacDonald, Oh, Barger, Cimpian, & Pomerantz (2024, *Developmental Psychology*, doi:10.1037/dev0001768) — inducing growth-oriented math mindsets in parents **sustained parent belief change** but **did not improve math parenting practices or children’s math adjustment** relative to an informative active control (Common Core information). Belief kits alone are not the product.

Applied (HYPOTHESIS): Parent soft-wrong copy should operationalize **failure-is-enhancing** as a concrete next attempt (“this miss maps the bridge to retry”), not as a poster slogan — and must not treat mindset induction as FEI proof.

Kill: Growth-mindset parent letter as learning NS; empty “mistakes are good!” without grain+CTA.

Survive: Failure-enhancing + named return (SAFE-FBLIT uptake for the household).

Wound: Haimovitz/Dweck are not MindCraft soft-wrong RCTs — SOFTMSG-* must test digest frames in-product.

CXVIII.7 Soft-wrong is learner climate — parent honesty still needs a contract

Reuse (doctrine, not new FACT): Soft-wrong / hide-correctness protects diagnostic integrity and destakes threat (SAFE-MISCON / SAFE-ERRCLIMATE / SAFE-CALIB). That does **not** license silence to parents (SAFE-PDASH honesty) nor license converting soft-wrong into a private gradebook.

HYPOTHESIS — message tiers:

1. **Default digest (sparse):** proof-age / MoC + one soft-wrong grain summary (“three attempts showed a FormatId gap on linear systems”) + one return CTA.
2. **No live stalk:** no item keys, no real-time wrong feed, no Miss Streak counter as hero.
3. **Anxious-parent mode:** extra line reducing controlling pressure; invite tutor/Map; ban “sit until correct.”
4. **Exam-rail exception:** prove-rail KR honesty differs from learn-rail soft-wrong (SAFE-EXAM / SAFE-FBTIME) — label the rail so parents don’t treat destaked practice as ACT readiness.

Kill: Soft-wrong $\equiv$ secret from parents; soft-wrong $\equiv$ parent miss leaderboard.

Survive: Honest grain without shame chrome.

CXVIII.8 Product surface — SAFE-SOFTMSG claim contract

Surface	Required behavior	Banned substitute
After soft-wrong run	Grain-named note + failure-enhancing one-liner + one return CTA	Miss Streak push; “falling behind” guilt
Parent digest hero	Proof-age / next join / MoC (SAFE-PDASH)	Tonight-% / XP / ranks
Parent help tip	Autonomy support: invite attempt; ask what she tried; don’t take pencil	Controlling script; uninvited full solve
Math-anxious parent	Explicit “you need not re-teach every miss; Map/tutor carries diagnosis”	Recruit more homework frequency
Soft-wrong vs prove rail	Label destake vs scored practice	Soft-wrong $\equiv$ exam-ready or $\equiv$ secret failure
Metrics	parent_softmsg_open, parent_return_cta_start, parent_intrusion_tip_shown, uptake of return; FEI co-gates	Parent Correctness Score™ / Miss Streak NS

Competitive foil: Gradebook apps = miss ranks. Mindset kits = belief posters. Stalk apps = live wrongs. MindCraft = **honest grain + one autonomy-supportive return.**

CXVIII.9 Doctrine — SAFE-SOFTMSG (provisional)

1. **Grain before vibe** — name concept/ingredient/FormatId/bridge; no fog praise.
2. **Failure-enhancing, not ability verdict** — miss = next join information (Haimovitz & Dweck); never “she can’t.”
3. **One improvement CTA** — autonomy-supportive return mission (SAFE-PDASH); no CTA spam.
4. **Autonomy > control** — tip language invites attempt; bans uninvited takeover (Jiang; Retanal; Park).
5. **Do not recruit anxious parents into more homework control** — reduce pressure; route diagnosis to Map/tutor (Maloney).
6. **No miss-streak / live-wrong hero** — honesty without shame telemetry.
7. **Belief posters ≠ product** — mindset letters without practice change are incomplete (MacDonald et al.).
8. **Copy:** “A miss showed the next join — here’s one return when she’s ready.” Never guilt streaks; never empty “mistakes are magic”; never ACT guarantees from texts.

Confidence: High — IPI hurts math more than SPI helps (Jiang); frequent help from math-anxious parents transmits harm (Maloney); controlling style mediates (Retanal); intrusive help × fixed mindset worsens drag (Park); failure mindsets shape children’s beliefs (Haimovitz & Dweck); parent growth-mindset induction alone often fails to move practices/child outcomes (MacDonald). Medium — exact soft-wrong digest wording and cadence (needs SOFTMSG-*). High — Miss Streak / live stalk / guilt SMS as parent NS are commercially toxic under this stack.

CXVIII.10 Experiments

ID	Question	Design	Primary
SOFTMSG-1	Grain+failure-enhancing+one CTA vs miss-count shame digest after soft-wrong run	A/B parents	parent_return_cta_start; student return; parent affect
SOFTMSG-2	Autonomy-support tip vs controlling “sit until correct” tip	A/B	Reported intrusion; student retry_120s / solo
SOFTMSG-3	Anxious-parent pressure-reduce line vs frequency-increase line	Stratified A/B	Help frequency; controlling style proxy; child approach
SOFTMSG-4	Soft-wrong digest alone vs digest+proof-age MoC strip	A/B	Trust; return CTA; no rank preference
SOFTMSG-5	Parent CBC: honest-grain return vs miss-streak portal vs mindset-poster pack	CBC	WTP; trust (SAFE-WTP)

SOFTMSG-QUAL	10 parents after soft-wrong week: what felt honest vs nag vs shame?	Qual	SOFTMSG codebook
--------------	---	------	------------------

Falsifier: Miss-count shame digests outperform grain+CTA on return and later FEI without raising intrusion → still ban shame as NS; segregate research.

Falsifier: Autonomy tips never beat controlling tips on child FEI → revise tip content; do not restore stalk.

Falsifier: Mindset-poster pack wins CBC and FEI vs grain+CTA → still require grain+CTA in ship; posters alone remain incomplete.

Pre-register: SOFTMSG-* before Miss Streak parent campaigns or “growth mindset parent texts raise ACT” ads (SAFE-LABMETA).

Family note: SOFTMSG- *densifies* PDASH- / FBLIT-*; do not collapse into PDASH-1 or mindset-MOOC pilots.

CXVIII.11 So what for MindCraft commercially

- **Copy:** “A miss showed the next join — here’s one return when she’s ready.” Lead with grain + invitation, not guilt streaks or poster slogans.
- **Product:** Soft-wrong parent digest = named grain + failure-enhancing frame + one CTA + optional anti-control tip; no live wrong feed; no Miss Streak hero.
- **Positioning:** Against gradebook stalk, mindset-poster kits, and nag SMS edtech; for honest household feedback literacy.
- **Metric:** parent_softmsg_open, parent_return_cta_start, intrusion-tip exposure, return→FEI co-gates — demote Parent Correctness Score™ / Miss Streak NS.
- **Kill list:** Shame soft-wrong pushes; live stalk; controlling default scripts; empty mistake-magic; ACT guarantees from parent texts.
- **Growth:** Parent decks sell *clarity without surveillance*; tutors brief parents on invite-not-takeover after soft-wrong weeks (SAFE-HITL).
- **Vision:** Thirty-year identity company teaches households that a miss can name the next join — without turning the family into a grading booth.

References (verified)

• Haimovitz, K., & Dweck, C. S. (2016). Parents’ views of failure predict children’s fixed and growth intelligence mind-sets. *Psychological Science*, 27(6), 859–869. <https://doi.org/10.1177/0956797616639727>

• Jiang, Q., Shi, L., Zheng, D., & Mao, W. (2023). Parental homework involvement and students’ mathematics achievement: A meta-analysis. *Frontiers in Psychology*, 14, 1218534. <https://doi.org/10.3389/fpsyg.2023.1218534>

• MacDonald, C., Oh, D. D., Barger, M. M., Cimpian, A., & Pomerantz, E. M. (2024). Does inducing growth-oriented mindsets about math ability in parents enhance children’s math mindsets, affect, and achievement? *Developmental Psychology*. <https://doi.org/10.1037/dev0001768>

• Maloney, E. A., Ramirez, G., Gunderson, E. A., Levine, S. C., & Beilock, S. L. (2015). Intergenerational effects of parents’ math anxiety on children’s math achievement and anxiety. *Psychological Science*, 26(9), 1480–1488.

<https://doi.org/10.1177/0956797615592630>

- Park, D., Gunderson, E. A., Maloney, E. A., Tsukayama, E., Beilock, S. L., Duckworth, A. L., & Levine, S. C. (2023). Parental intrusive homework support and math achievement: Does the child’s mindset matter? *Developmental Psychology*, 59(7), 1249–1267. <https://doi.org/10.1037/dev0001522>
 - Retanal, F., Johnston, N. B., Di Lonardo Burr, S. M., Storozuk, A., DiStefano, M., & Maloney, E. A. (2021). Controlling-supportive homework help partially explains the relation between parents’ math anxiety and children’s math achievement. *Education Sciences*, 11(10), 620. <https://doi.org/10.3390/educsci11100620>
 - Wu, J., Barger, M. M., Oh, D. D., & Pomerantz, E. M. (2022). Parents’ daily involvement in children’s math homework and activities during early elementary school. *Child Development*, 93(5), 1347–1364. <https://doi.org/10.1111/cdev.13774>
-

Part CXIX — Classroom Orchestration of Map During Live Tutoring

Chapter status: Living evidence + HITL session pedagogy — Researcher tick 2026-08-10 (UTC hour 18 ≡ Red Team slot, but ch119 never written → prefer Researcher per rotation; researcher count since synthesizer v1.15 = 2)

Primary question: When should MindCraft tutors open the Knowledge Map **during** a live session — versus before the first explain, after an attempt grain, or only in debrief — so Map diagnosis densifies formative cycles without map-shock, dual-screen multitasking tax, or “always-on dashboard” cosplay?

Owners: HITL ops / tutor playbooks · Product (Map / session chrome) · Brand · Analytics (FEI / talk / FA instrumentation) · Red Team

Commercial job: Ship **SAFE-MAPORCH** densifying SAFE-HITL × SAFE-FA × SAFE-ATTN × SAFE-BRIDGE × SAFE-ONTOLOGY: **brief** → **attempt** → **grain-slice Map** → **adjust** → **solo proof** — never continuous student-facing Map theater, mid-attempt full-graph dumps, or Map Minutes as a North Star.

Builds on: Parts LXVIII (SAFE-HITL), CXI (SAFE-FA), XCVI (SAFE-ATTN), LXXX (SAFE-BRIDGE), LXVII (SAFE-ONTOLOGY), LIV (SAFE-AAR), LXXXIV (SAFE-TALK), XCV (SAFE-APPRENTICE), CXVII (SAFE-FADEX). Product seams: Map brief, bridge-gap CTA, tutor session structure, soft-wrong, FEI North Stars.

CXIX.1 Why this chapter exists

SAFE-HITL already requires Map-briefed tutors and bans humans-skip-Map. SAFE-FA already requires evidence→adjust cycles. The remaining commercial hazard is **orchestration timing**: when the tutor (or product) opens Map *while Maya is mid-problem*, the graph can become (a) a second screen that taxes working memory, (b) a full-ontology dump that induces map-shock, or (c) a vanity “we use the Map every minute” KPI that substitutes display time for contingent support.

Competitors resolve this wrongly: ITS dashboards always-on behind the student; “knowledge graph” demos that wander the whole ontology mid-explanation; tutors who never open Map and pour procedure; tutors who open Map as a shared spectacle instead of a diagnostic slice.

FOUNDER BELIEF under audit: Map is a **diagnosis instrument for contingent teaching**, not a classroom wallpaper and not a replacement for an attempt. Opening it at the wrong moment is worse than leaving it closed until

the grain is ready.

Claims we refuse as doctrine:

1. Always-on student-facing Map during every live minute.
2. Mid-attempt full-graph dump / ontology tour as pedagogy.
3. Map Minutes / Map Score™ as tutor or learning North Star.
4. Humans-skip-Map (already killed; restated).
5. Map-open $\equiv$ formative cycle closed (display $\neq$ Sadler gap-close).
6. Dual-screen Map+problem as default high-load UX.
7. Bridge-count wander $\equiv$ bridge teaching.
8. ACT guarantees from “Map-powered tutoring” packaging.

CXIX.2 Constructs (open Map $\neq$ close the gap)

Construct	Meaning	MindCraft analogue	Failure mode
Map brief	Pre-session / pre-explain diagnosis glance	Tutor opens Map before first pour	Skip Map; pour first
Grain-slice Map	Focused subgraph (concept/ingredient/bridge/FormatId)	Highlighted next join, not full ontology	Map-shock tour
Mid-attempt open	Student-facing Map while working memory holds the item	Dual chrome during solve	Multitask tax
Contingent diagnosis	Diagnose understanding $\rightarrow$ check $\rightarrow$ adaptive support $\rightarrow$ check learning	van de Pol scaffolding steps	Support without diagnose
SAFE-MAPORCH	When/how Map enters live tutoring	This chapter	Always-on dashboard

Operational definition (HYPOTHESIS): A live tutoring segment is *SAFE-MAPORCH complete* when (a) the tutor opens Map for a **brief** before the first explain (SAFE-HITL), (b) student-facing Map mid-session is a **grain-slice** tied to elicited evidence (not a full-graph wander), (c) Map is **not** the primary chrome during a high-load attempt (attempt-first; Map between attempts or tutor-private), (d) every Map open names a next instructional move that the student then takes (SAFE-FA / Sadler), and (e) success is measured by $FEI + fa_cycle_closed / map_brief_open / map_slice_adjust$ — never Map Minutes.

CXIX.3 Contingent scaffolding requires diagnose-before-support — Map is the diagnose tool

FACT (scaffolding review): van de Pol, Volman, & Beishuizen (2010, *Educational Psychology Review*, 22(3), 271–296, doi:10.1007/s10648-010-9127-6) — contingency, fading, and transfer of responsibility are the three key characteristics of scaffolding; descriptive work is abundant; effectiveness evidence is thinner but suggestive; measurement of dynamic contingency is the field’s hard problem. Contingent support presupposes **diagnosis of the learner’s current understanding** before the level of help is set.

FACT (process spine, reuse): Black & Wiliam formative assessment as process — elicit evidence, interpret, adjust next steps (SAFE-FA). Sadler (1989): feedback is formative only when the learner acts to close a gap.

Applied (HYPOTHESIS): Map mid-session is justified when it **improves diagnosis contingency** (tutor sees bridge gap / FormatId / ingredient and changes the next prompt). It is unjustified when it is opened as spectacle, reassurance, or “we have a graph” theater without a next attempt.

Kill: Map-open without next move; support-without-diagnose pour.

Survive: Diagnose (Map slice or probe) → contingent prompt → student action.

Wound: van de Pol et al. review classroom scaffolding broadly — not MindCraft Map RCTs. MAPORCH-* must test open-timing in-product.

CXIX.4 Map-shock — full graphs mid-session tax novices

FACT (map-shock coinage): Blankenship & Dansereau (2000, *Journal of Experimental Education*, 68(4), 293–308, doi:10.1080/00220970009600640) — coined **map shock** for the bewilderment of not knowing where to start or how to penetrate a complex node-link map; animation can help direct attention on ill-structured maps for macrostructure recall, but complexity itself is the hazard the term names.

FACT (prior knowledge × structure): Amadiou, van Gog, Paas, Tricot, & Mariné (2009, *Learning and Instruction*, 19(5), 376–386, doi:10.1016/j.learninstruc.2009.02.005) — low prior-knowledge learners report more **disorientation** and fare worse on conceptual knowledge from complex/networked concept-map structures than from simpler hierarchical ones; map structure interacts with prior knowledge for load and learning.

Applied (HYPOTHESIS): Opening MindCraft’s full ontology-ish Map to Maya mid-session recreates map-shock / disorientation — especially for low-exposure grains. Commercial default for **student-facing** mid-session Map = **grain-slice** (one bridge, one FormatId hop, one next route) with hierarchy/signaling, not a free wander of 42 concepts.

Kill: Full-graph mid-session dump; “explore the Map together” as default when stuck.

Survive: Tutor-private full Map OK; student sees the named join.

CXIX.5 Concurrent secondary displays hurt — especially under high load

FACT (secondary-task timing): Dönmez & Akbulut (2021, *Computers & Education*, 161, 104078, doi:10.1016/j.compedu.2020.104078; N=356) — concurrent secondary tasks (even relevant ones) reduced multimedia learning gains vs control; **relevant-sequential** secondary tasks did **not** differ significantly from control; irrelevant and concurrent conditions underperformed. Monotasking preferred; if a secondary task is needed, make it **relevant and**

sequential.

FACT (attention residue, reuse): Leroy (2009) — unfinished Task A cognitions tax Task B (SAFE-ATTN). Mid-attempt Map open leaves the problem unfinished while the student processes a second complex representation.

Applied (HYPOTHESIS): Student-facing Map during an active attempt is a **concurrent secondary task**. Prefer: finish or park the attempt → open grain-slice Map → return to next attempt (relevant-sequential). Tutor may keep a private Map pane if the student’s primary chrome stays the problem.

Kill: Dual-screen Map+problem as default; Map push mid-keystroke.

Survive: Between-attempt Map slices; tutor-private diagnosis pane.

CXIX.6 Noticing and shaping — Map supports interpretation, not airtime

FACT (teacher noticing): van Es & Sherin (2002 onward; updated shaping account in van Es & Sherin, 2021, *ZDM – Mathematics Education*) — professional noticing includes attending to student mathematical thinking, interpreting it, and **shaping** interactions to elicit further evidence. Expertise shifts from classroom management narratives toward substance of student reasoning.

Applied (HYPOTHESIS): Map should **shape** the next elicit (which probe, which bridge ask) — not occupy talk ratio. SAFE-TALK prompt>pour still binds: Map glance that becomes a monologue tour fails noticing-as-shaping.

Kill: Map tour ≡ tutoring; Map Minutes NS.

Survive: Brief Map → named prompt → student construction.

CXIX.7 Product surface — SAFE-MAPORCH claim contract

Moment	Required behavior	Banned substitute
Session start	Map brief before first explain (SAFE-HITL)	Pour-first; skip Map
During attempt	Student chrome = problem; tutor-private Map OK	Student dual-screen full Map
After soft-wrong / stall	Grain-slice Map naming join/FormatId/bridge + one next move	Full ontology wander; “look at all red nodes”
Mid-session adjust	Evidence→Map slice→changed prompt/mission (SAFE-FA)	Map-open without action
Wrap / AAR	Map optional for named next route (SAFE-AAR)	End-of-hour Map tourism
Metrics	map_brief_open, map_slice_adjust, fa_cycle_closed, FEI co-gates	Map Minutes / Map Score™

Competitive foil: Always-on ITS dashboards = concurrent tax. ChatGPT tutors = no Map. Graph demos = map-shock tourism. MindCraft = **timed grain diagnosis that changes the next attempt**.

CXIX.8 Doctrine — SAFE-MAPORCH (provisional)

1. **Brief before pour** — Map diagnosis precedes first explain (SAFE-HITL).
2. **Attempt before student Map** — do not dual-screen full Map during high-load solves (Dönmez & Akbulut; Leroy).
3. **Grain-slice, not ontology tour** — student-facing Map shows the named join (Blankenship map-shock; Amadiou disorientation).
4. **Open to adjust** — Map justifies itself by changing the next prompt/mission (van de Pol contingency; SAFE-FA).
5. **Tutor-private pane allowed** — diagnosis chrome for the tutor need not be student chrome.
6. **Bridge/FormatId over red-node shame** — open for connection/format gaps (SAFE-BRIDGE / SAFE-FORMAT), not humiliation ranks.
7. **No Map Minutes NS** — display time $\neq$ learning; FEI + cycle-closed metrics win.
8. **Copy:** “Open Map to name the next join — then put it away for the attempt.” Never “live Map always on”; never ACT guarantees from Map airtime.

Confidence: High — contingency requires diagnosis (van de Pol); concurrent secondary tasks tax learning (Dönmez & Akbulut); complex maps disorient low-prior learners (Amadiou; Blankenship map-shock); residue taxes switches (Leroy). Medium — exact second thresholds and tutor-private vs shared slice UX (needs MAPORCH-*). High — Map Minutes / always-on / mid-attempt full-graph as NS are commercially toxic under this stack.

CXIX.9 Experiments

ID	Question	Design	Primary
MAPORCH-1	Grain-slice after soft-wrong vs full-graph mid-session open	A/B sessions	map_slice_adjust; retry_120s; load/disorientation rating
MAPORCH-2	Tutor-private Map pane vs shared dual-screen during attempt	A/B	Solo transfer; talk pour_streak; student dual-task errors
MAPORCH-3	Map brief+between-attempt only vs always-on student Map	A/B	FEI; fa_cycle_closed; Map Minutes (demote)

MAPORCH-4	Bridge/FormatId slice CTA vs red-node gallery after stall	A/B	Join/format hop starts; shame affect; next-move fidelity
MAPORCH-5	Parent/tutor CBC: timed grain Map vs always-on dashboard vs no-Map human	CBC	WTP; trust (SAFE-WTP)
MAPORCH-QUAL	10 tutor sessions: when did Map help vs interrupt?	Qual	MAPORCH codebook

Falsifier: Always-on student Map beats brief+slice on FEI without raising load → still ban Map Minutes as NS; segregate research.

Falsifier: Full-graph mid-session outperforms grain-slice for novices → revise slice design; do not restore tourism as brand.

Falsifier: Skipping Map entirely matches Map-briefed FEI → revise brief UX; do not resurrect humans-skip-Map as doctrine.

Pre-register: MAPORCH-* before “always-on knowledge graph tutoring” ads (SAFE-LABMETA).

Family note: MAPORCH- *densifies HITL-* / FA-*; do not collapse into map_brief_open alone or AfL badge pilots.

CXIX.10 So what for MindCraft commercially

- **Copy:** “Open Map to name the next join — then put it away for the attempt.” Lead with timed diagnosis, not dashboard immersion.
- **Product:** Session chrome = brief → attempt → grain-slice adjust → solo; tutor-private Map pane; no always-on student Map default.
- **Positioning:** Against always-on ITS dashboards, ChatGPT no-diagnosis pour, and ontology-tourism demos; for contingent Map that changes the next move.
- **Metric:** map_brief_open, map_slice_adjust, fa_cycle_closed, FEI co-gates — demote Map Minutes / Map Score™.
- **Kill list:** Always-on student Map; mid-attempt full-graph; Map-open without action; Map Minutes NS; ACT guarantees from Map airtime.
- **Growth:** Tutor playbooks teach *when* to open Map; sales decks show a 10-second grain slice, not a flying camera through the whole graph.
- **Vision:** Thirty-year identity company treats the Map as a scalpel for joins — visible when it names the next attempt, invisible when Maya needs the problem alone.

References (verified)

- Amadiou, F., van Gog, T., Paas, F., Tricot, A., & Mariné, C. (2009). Effects of prior knowledge and concept-map structure on disorientation, cognitive load, and learning. *Learning and Instruction*, 19(5), 376–386.

<https://doi.org/10.1016/j.learninstruc.2009.02.005>

- Blankenship, J., & Dansereau, D. F. (2000). The effect of animated node-link displays on information recall. *Journal of Experimental Education*, 68(4), 293–308. <https://doi.org/10.1080/00220970009600640>
 - Dönmez, O., & Akbulut, Y. (2021). Timing and relevance of secondary tasks impact multitasking performance. *Computers & Education*, 161, 104078. <https://doi.org/10.1016/j.compedu.2020.104078>
 - Leroy, S. (2009). Why is it so hard to do my work? The challenge of attention residue when switching between work tasks. *Organizational Behavior and Human Decision Processes*, 109(2), 168–181. <https://doi.org/10.1016/j.obhdp.2009.04.002>
 - van de Pol, J., Volman, M., & Beishuizen, J. (2010). Scaffolding in teacher–student interaction: A decade of research. *Educational Psychology Review*, 22(3), 271–296. <https://doi.org/10.1007/s10648-010-9127-6>
 - van Es, E. A., & Sherin, M. G. (2021). Expanding on prior conceptualizations of teacher noticing. *ZDM – Mathematics Education*, 53, 17–27. <https://doi.org/10.1007/s11858-020-01211-4>
-

Part CXX — Tutor Burnout Early-Warning Without Wellness Theater

Chapter status: Living evidence + HITL ops — Researcher tick 2026-08-10 (UTC hour 21; hour%6≠0; researcher count since synthesizer v1.15 = 3 → Researcher)

Primary question: How should MindCraft detect tutor burnout risk *early enough to protect FEI fidelity* — without a Burnout Score™ product, emotion surveillance of tutors, or wellness-module theater that leaves caseload unchanged?

Owners: HITL ops / workforce · Product (tutor console / QA) · Privacy/trust · Brand · Analytics (fidelity@tenure) · Red Team

Commercial job: Ship **SAFE-BURNWARN** densifying SAFE-WORKFORCE × SAFE-PRIVACY × SAFE-HITL: **fidelity dips and load imbalance are the early-warning signals; org redesign (caseload/schedule/coach bandwidth) is the intervention; optional voluntary check-ins are not a surveillance SKU.**

Builds on: Parts LXXV (SAFE-WORKFORCE), LXIX (SAFE-PRIVACY), LXVIII (SAFE-HITL), LXXI (SAFE-TUTORGRAIN), XLVI (SAFE-EXPECTANCY), LX (SAFE-REPAIR), LXXXIV (SAFE-TALK). Product seams: Map-brief gate, talk pour, soft-wrong avoidance, coach sample cadence, caseload caps.

CXX.1 Why this chapter exists

SAFE-WORKFORCE already kills headcount NS, tenure=quality, and wellness theater as pedagogy — and requires burnout→load redesign (WORK-3). What remains commercially underspecified is the **early-warning instrumentation**: what to watch *before* exit interviews, how to act without turning tutors into an Anxiety Score™ cohort, and how to sell “we protect tutor quality” without mindfulness-app cosplay.

Competitors resolve this wrongly: marketplace volume that ignores fidelity drift; “tutor wellness Fridays” that leave caseloads intact; emotion-AI dashboards on tutor cameras; retention-% hero metrics that hide pour-relapse; grit posters for exhausted near-peers.

FOUNDER BELIEF under audit: The cheapest honest early warning is **behavioral fidelity under load** (brief skip, pour relapse, soft-wrong avoid, QA starve) — not a proprietary Burnout Score™ — and the honest fix is **demand/resource redesign**, not a meditation SKU.

Claims we refuse as doctrine:

1. Burnout Score™ / Wellness Score™ as tutor or parent North Star.
2. Face/voice emotion AI on tutors as “care.”
3. Mandatory mood gates before every session.
4. Wellness module / mindfulness add-on as primary burnout strategy when caseload is the demand.
5. Retention-% without fidelity@30/90d as quality proof.
6. Trait emotion labels in tutor cards (SAFE-EXPECTANCY collision).
7. Sell tutor affect data / “burnout analytics” as a district SKU.
8. ACT guarantees from “wellness-supported tutors” packaging.

CXX.2 Constructs (early warning ≠ wellness theater)

Construct	Meaning	MindCraft analogue	Failure mode
Exhaustion / cynicism / inefficacy	Classic burnout dimensions (Maslach tradition)	Ops concern, not student SKU	Burnout Score™ brand
Job demands vs resources	JD-R dual pathways	Caseload, emotional labor, Map QA vs coaching, autonomy, schedule	Fix person, ignore demand
Fidelity dip	Drift channels before quit	map_brief_open↓, pour↑, soft-wrong avoid↑	Wait until exit
Org-directed fix	Change work design	Cap caseload; add coach bandwidth; pause hire	Wellness Slack only
SAFE-BURNWARN	Early warn + redesign doctrine	This chapter	Surveillance theater

Operational definition (HYPOTHESIS): A MindCraft burnout early-warning system is *honest* when (a) primary signals are **work-design and fidelity events** already in the tutor console (not biometric ERT), (b) voluntary self-report is optional, purpose-bound, short-TTL, and never a parent/district dashboard (SAFE-PRIVACY extended to tutors), (c) a flagged pattern triggers **caseload/schedule/coach redesign** within a defined SLA, (d) wellness content may be offered as optional support but never substitutes for redesign, and (e) success is measured by fidelity@tenure + quit-intent proxies + FEI under coached tutors — never Wellness Minutes.

CXX.3 JD-R — demands predict exhaustion; resources predict engagement

FACT (JD-R founding): Demerouti, Bakker, Nachreiner, & Schaufeli (2001, *Journal of Applied Psychology*, 86(3), 499–512, doi:10.1037/0021-9010.86.3.499) — job demands primarily predict **exhaustion**; (lack of) job resources primarily predict **disengagement**; pattern held across occupational groups with self-report and observer ratings (N=374).

FACT (JD-R theory stock-take): Bakker & Demerouti (2017, *Journal of Occupational Health Psychology*, 22(3), 273–285, doi:10.1037/ocp0000056) — dual pathways: health-impairment (demands→burnout) and motivational (resources→engagement); later evidence links demands→absence duration via burnout and resources→commitment/engagement; interventions must address both sides of the ledger.

Applied (HYPOTHESIS): MindCraft’s early-warning board should surface **demand spikes** (FEI-heavy concurrent students, repair escalations, exam-week density, unpaid emotional labor after hallucinations) and **resource deficits** (coach sample lag, missing Map tools, schedule chaos) — not a single composite Burnout Score™. Redesign = cut demands and/or add resources.

Kill: Person-only coping apps as the default response to demand spikes.

Survive: Demand/resource dashboards that ops can act on this week.

Wound: JD-R is general occupational theory — not a MindCraft RCT. BURNWARN- *must test which console signals lead quit and drift in this* pipeline.*

CXX.4 Burnout predicts quit and harms the interaction students feel

FACT (attrition meta, reuse): Madigan & Kim (2021, *Teaching and Teacher Education*, 105, 103425, doi:10.1016/j.tate.2021.103425) — burnout dimensions relate positively to quit intent (exhaustion $r^+ \approx .41$, depersonalization $\approx .32$, reduced accomplishment $\approx .21$); satisfaction protects less than burnout risks (sat $r^+ \approx -.40$); together $\approx 27\%$ of quit-intent variance.

FACT (exhaustion → instructional quality → students, reuse): Klusmann, Aldrup, Roloff, Lüdtke, & Hamre (2022, *Journal of Educational Psychology*, 114(6), 1442–1460, doi:10.1037/edu0000703) — large multilevel sample; teacher emotional exhaustion associated with worse student self-concept, interest, and (German) achievement; paths via lower emotional support and classroom organization.

Applied (HYPOTHESIS): Early warning that only tracks quit intent after the tutor has already poured for six weeks is too late for Maya. Prefer leading indicators on SAFE-WORKFORCE drift channels: brief skip, pour relapse, “got it?” closers, soft-wrong avoidance, QA starvation — each is a **fidelity** early warning that also predicts exhaustion→interaction harm.

Kill: Retention-% hero without fidelity trend.

Survive: Fidelity@30/90d + drift-channel alerts as the early board.

CXX.5 Org-directed interventions beat person-directed coping theater

FACT (physician meta): Panagioti et al. (2017, *JAMA Internal Medicine*, 177(2), 195–205, doi:10.1001/jamainternmed.2016.7674) — controlled interventions reduce burnout overall (SMD ≈ -0.29); **organization-directed** interventions (SMD ≈ -0.45) outperform **physician-directed** ones (SMD ≈ -0.18); supports framing burnout as an organizational problem, not only an individual deficit.

FACT (org interventions meta): Bes, Shoman, Al-Gobari, Rousson, & Guseva Canu (2023, *International Archives of Occupational and Environmental Health*, doi:10.1007/s00420-023-02009-z) — organizational interventions targeting workload/resources and participatory redesign are central levers for exhaustion; JD-R imbalance (demands without recovery/resources) is the theoretical spine; person-only approaches are incomplete when workload is the driver.

Applied (HYPOTHESIS): When MindCraft detects a BURNWARN pattern, the default playbook is WORK-3-style: **caseload cap, schedule redesign, hire pause until coach bandwidth recovers** — optionally offer coping resources, never as the primary fix. Wellness modules that leave caseload unchanged fail Panagioti’s org>person lesson.

Kill: Mindfulness add-on as primary burnout strategy; wellness theater in sales decks.

Survive: Documented redesign SLA after fidelity/load alerts.

CXX.6 Privacy boundary — tutor affect is not a product surface

FACT / doctrine (reuse): SAFE-PRIVACY bans face/voice emotion AI, Anxiety Score™ NS, mandatory mood gates, trait emotion briefs, and selling affect data (Part LXIX; Zeide FERPA-limits; Nissenbaum contextual integrity).

Applied (HYPOTHESIS): Extend the same bar to tutors. Early warning may use (a) operational events tutors already generate in the console, (b) optional voluntary check-ins with short TTL for ops coaches only, never parent/district Wellness dashboards. Covert attention or emotion surveillance of tutors is brand poison and procurement risk (SAFE-PROCURE).

Kill: Tutor emotion-AI; Burnout Score™ on parent portal; mandatory mood gate before session start.

Survive: Ops-private load/fidelity board; voluntary coach conversation.

CXX.7 Product surface — SAFE-BURNWARN claim contract

Signal class	Example	Required action	Banned substitute
Fidelity dip	map_brief_open miss streak; pour_streak↑	Coach sample + load review within SLA	Pep-talk Slack only
Demand spike	Concurrent FEI-heavy students above cap; repair pile-up	Cap/redistribute caseload	“Practice self-care” email
Resource lag	Coach sample overdue vs hire wave	Pause hire; restore QA bandwidth	Hire more to “share load” without coaches
Optional self-report	Voluntary short check-in to coach	Purpose-bound; short TTL; no trait card	Burnout Score™ / emotion-AI

Exit foreshadow	Soft cancels↑; response latency↑	Ops outreach + redesign	Punitive attendance theater
Metrics	fidelity@30/90d; coach_util; drift rates; quit proxy	Co-gate with FEI	Wellness Minutes / Burnout Score™

Competitive foil: Marketplace volume = ignore dips. Wellness-app cosplay = person-fix. Emotion dashboards = surveillance. MindCraft = **fidelity-under-load early warn → org redesign.**

CXX.8 Doctrine — SAFE-BURNWARN (provisional)

1. **Fidelity dips before exit** — instrument drift channels as leading indicators (SAFE-WORKFORCE).
2. **JD-R ledger** — alert on demand spikes and resource deficits, not a mystical composite score (Demerouti; Bakker & Demerouti).
3. **Org redesign first** — caseload/schedule/coach bandwidth beat person-directed coping as primary response (Panagioti; Bes et al.).
4. **No Burnout Score™ / Wellness Minutes NS** — already company-adjacent via SAFE-WORKFORCE; this chapter densifies the ban for early-warn UX.
5. **Privacy parity for tutors** — no ERT; no mandatory mood gate; no parent/district tutor-affect dashboards (SAFE-PRIVACY).
6. **Voluntary check-in ≠ surveillance SKU** — optional, coach-private, short TTL.
7. **Hire gated by coach capacity** — surge hire without QA is a demand bomb, not a fix.
8. **Copy:** “We watch whether tutors still make students think — and we cut load when fidelity slips.” Never “our AI wellness coach keeps tutors happy.”

Confidence: High — JD-R demand/resource pathways; Madigan quit links; Klusmann exhaustion→instruction path; Panagioti org>person. Medium — exact console thresholds and SLA days for MindCraft near-peer tutors (needs BURNWARN-*). High — Burnout Score™ / wellness-theater-as-primary / emotion-AI as commercially toxic under this stack.

CXX.9 Experiments

ID	Question	Design	Primary
BURNWARN-1	Fidelity-dip alerts + redesign SLA vs wellness-module-only	A/B ops sites	Drift rates; quit proxy; fidelity@30/90d
BURNWARN-2	Caseload cap trigger vs uncapped + coping email	A/B	Pour rate; brief-skip; tutor retention with fidelity
BURNWARN-3	Ops-private load board vs tutor-facing Burnout Score™	A/B	Trust; stigma; false-positive distress

BURNWARN-4	Voluntary coach check-in vs mandatory mood gate	A/B	Completion without retaliation fear; FEI
BURNWARN-5	Parent/district CBC: fidelity-under-load report vs wellness-app attribute	CBC	WTP; trust (SAFE-WTP / SAFE-PROCURE)
BURNWARN-QUAL	10 tutors + 5 coaches: what felt like early warn vs theater?	Qual	BURNWARN codebook

Falsifier: Wellness-module-only beats redesign SLA on fidelity+retention → revise WORK-3 priority; still ban Burnout Score™ as NS.

Falsifier: Composite Burnout Score™ predicts quit better than fidelity/load events without stigma cost → research-only; do not ship as brand.

Falsifier: Mandatory mood gates improve FEI without trust loss → still ban ERT; segregate.

Pre-register: BURNWARN-* before “AI wellness for tutors” ads (SAFE-LABMETA).

Family note: BURNWARN- *densifies* WORK-; do not collapse into retention-% OKRs or PRIV-* student check-ins.

CXX.10 So what for MindCraft commercially

- **Copy:** “We watch whether tutors still make students think — and we cut load when fidelity slips.” Lead with coached fidelity under load, not wellness theater.
- **Product:** Tutor ops board = demand/resource + drift alerts → redesign SLA; optional coach-private check-in; no Burnout Score™ chrome.
- **Positioning:** Against marketplace volume, mindfulness-as-fix, and emotion-AI “care”; for org-honest early warning that protects Maya’s FEI.
- **Metric:** fidelity@30/90d, coach_util, drift-channel rates, quit proxy — demote Wellness Minutes / Burnout Score™ / retention-% alone.
- **Kill list:** Burnout Score™; mandatory mood gates; tutor ERT; wellness-module-as-primary; ACT guarantees from wellness packaging.
- **Growth:** District trust packets include tutor-data map (what is *not* collected) + fidelity QA cadence — procurement asset, not surveillance pitch.
- **Vision:** Thirty-year identity company treats tutor burnout as a **design failure of demand/resource balance**, visible early in how tutors teach — not as a mood product to monetize.

References (verified)

- Bakker, A. B., & Demerouti, E. (2017). Job demands–resources theory: Taking stock and looking forward. *Journal of Occupational Health Psychology*, 22(3), 273–285. <https://doi.org/10.1037/ocp0000056>

- Bes, I., Shoman, Y., Al-Gobari, M., Rousson, V., & Guseva Canu, I. (2023). Organizational interventions and occupational burnout: A meta-analysis with focus on exhaustion. *International Archives of Occupational and Environmental Health*. <https://doi.org/10.1007/s00420-023-02009-z>
- Demerouti, E., Bakker, A. B., Nachreiner, F., & Schaufeli, W. B. (2001). The job demands-resources model of burnout. *Journal of Applied Psychology*, 86(3), 499–512. <https://doi.org/10.1037/0021-9010.86.3.499>
- Klusmann, U., Aldrup, K., Roloff, J., Lüdtke, O., & Hamre, B. K. (2022). Does instructional quality mediate the link between teachers’ emotional exhaustion and student outcomes? *Journal of Educational Psychology*, 114(6), 1442–1460. <https://doi.org/10.1037/edu0000703>
- Madigan, D. J., & Kim, L. E. (2021). Towards an understanding of teacher attrition: A meta-analysis of burnout, job satisfaction, and teachers’ intentions to quit. *Teaching and Teacher Education*, 105, 103425. <https://doi.org/10.1016/j.tate.2021.103425>
- Panagioti, M., Panagopoulou, E., Bower, P., Lewith, G., Kontopantelis, E., Chew-Graham, C., Dawson, S., van Marwijk, H., Geraghty, K., & Esmail, A. (2017). Controlled interventions to reduce burnout in physicians: A systematic review and meta-analysis. *JAMA Internal Medicine*, 177(2), 195–205. <https://doi.org/10.1001/jamainternmed.2016.7674>

Part CXXI — Explanation Quality Rubrics for AI Coach Wraps

Chapter status: Living evidence + coach QA brief — Researcher tick 2026-08-11 (UTC hour 00 ≡ Red Team slot, but ch121 never written → prefer Researcher per rotation; researcher count since synthesizer v1.15 = 4 → Researcher)

Primary question: How should MindCraft *gate* AI coach / Solver wraps so “quality” means **principle-short, pedagogically actionable, keyed-correct guidance** — not fluency, thoroughness, or a vanity Explanation Score™?

Owners: Product (coach / Solver) · Engine (prompt budgets + verify) · HITL QA · Trust & Safety · Brand · Red Team

Commercial job: Ship **SAFE-EXPLAINQA** densifying SAFE-EXPLAIN × SAFE-GENQ × SAFE-REPAIR: **rubric gates before wrap ships**; hard-fail on math/key error; ban monologue fluency as pedagogy proof.

Builds on: Parts XCII (SAFE-EXPLAIN), LXXII (SAFE-GENQ), XL (SAFE-SE), LX (SAFE-REPAIR), XCI (SAFE-HINT), LXXXVIII (SAFE-FADE), XXXIII (AI trust). Product seams: coach card length, SE-before-wrap, generation verify, hallucination repair escalate.

CXXI.1 Why this chapter exists

SAFE-EXPLAIN already kills longer≡better and monologue≡SE. SAFE-GENQ already kills unverified keys in the bank. What remains commercially underspecified is the **ship gate for generated coach language**: which rubric dimensions must pass before a wrap reaches Maya, and which “AI quality” metrics are brand poison.

Default AI-tutor UX optimizes for fluent completeness and helpful tone. Default LLM-as-judge stacks optimize for agreement with another LLM. Default parent anxiety asks for “more explanation.” MindCraft’s identity claim needs the opposite default: **a short, correct, adaptive wrap that preserves construction — or a hard withhold / escalate.**

FOUNDER BELIEF under audit: A small, inspectable **principle-short QA rubric** (content correctness + learner fit + non-reveal + actionability + length) beats both (a) unmeasured fluent dumps and (b) a composite Explanation

Score™ product surface.

Claims we refuse as doctrine:

1. Fluency / human-likeness / “sounds like a tutor” $\equiv$ pedagogical quality.
2. Thumbs-up / satisfaction / “thoroughness” as wrap North Stars.
3. Explanation Score™ / Rubric Score™ / Quality Minutes as student or parent NS.
4. LLM-as-judge alone as production authority without human audit sample.
5. Longer wrap $\equiv$ higher rubric quality (SAFE-EXPLAIN collision).
6. Reveal-answer / full solution essay as default “guidance.”
7. Soft-pass on arithmetic/key errors because the prose “explained the idea.”
8. Auto-QA of wraps replacing bank key verify (SAFE-GENQ collision).
9. ACT guarantees from “rubric-scored AI coach” packaging.

CXXI.2 Constructs (rubric $\neq$ score product)

Construct	Research meaning	MindCraft analogue	Failure mode
Instructional explanation	Provided why/principle/operator text	Coach wrap; soft-hint text	Essay replacing SE
Explanation quality criteria	Content, learner fit, design, language, structure	Principle-short gate checklist	Composite vanity score
Pedagogical ability dims	Mistake ID/location, guidance, actionability, reveal, tone...	Coach QA dimensions	Tone-only polish
Hard fail	Content/key error blocks ship	GENQ-style verify on wrap math	Soft “mostly helpful”
Principle-short	Minimal critical-feature length	SAFE-EXPLAIN default	Thoroughness theater
SAFE-EXPLAINQA	Rubric gates for AI wraps	This chapter	Fluency costume

Operational definition (HYPOTHESIS): A MindCraft coach wrap is *SAFE-EXPLAINQA shippable* when it (a) is **keyed-correct** on any math/claim it asserts (hard fail otherwise), (b) marks the **critical feature** / **join** without default full-answer reveal, (c) is **actionable** (student knows the next attempt move), (d) fits fade/expertise grain (shortens as E rises), (e) preserves required SE when policy demands it, (f) strips seductive fluff, and (g) is proven by solo_transfer_pass / near transfer — never by Explanation Score™ dwell or thumbs-up on thoroughness.

CXXI.3 Why instructional explanations often fail — design before dump

FACT (framework): Wittwer & Renkl (2008, *Educational Psychologist*, 43(1), 49–64, doi:10.1080/00461520701756420) — instructional explanations frequently fail to promote learning; effectiveness depends on generation *and* use conditions (fit to prior knowledge, timing, whether learners actually process them), not on the mere presence of an explanation.

FACT (meta): Wittwer & Renkl (2010, *Educational Psychology Review*, 22, 393–409, doi:10.1007/s10648-010-9136-5) — instructional explanations added to worked examples yield a **small** overall effect ($d \approx 0.16$); stronger for conceptual knowledge than transfer; **no advantage** when controls are prompted to self-explain.

FACT (SE crowding): Schworm & Renkl (2006, *Computers & Education*, 46, 426–445, doi:10.1016/j.compedu.2004.08.011) — instructional explanations can **reduce** self-explanation activity; objective learning favored SE prompts alone, while learners *perceived* instructional explanations as more helpful — a satisfaction $\neq$ learning trap.

Applied (HYPOTHESIS): Coach QA must score **use conditions** (does this wrap enable the next construction?) not “completeness.” A wrap that crowds out SE fails even if fluent.

Kill: Thoroughness / perceived helpfulness as ship criteria.

Survive: Principle-short + SE-preserving gates; expand only on miss/demand.

CXXI.4 Content-first quality criteria — correctness is not optional

FACT (video/explanation rating instrument): Ring & Brahm (2023, *Technology, Knowledge and Learning*, doi:10.1007/s10758-022-09635-5) — theory-driven rating framework with **twelve criteria in five categories**: content, learner orientation, representation/design, language, process structure; content includes technical correctness and completeness of the core statement.

FACT (science explanation-video framework): Kulgemeyer (2018/2020 line; framework informed by instructional-explanation criteria, *Research in Science Education* / follow-ons) — adaptation to prior knowledge, minimal coherent explanation (avoid digressions), structure, and relevance tools; videos adhering more strongly to the framework produced higher **declarative** post-test knowledge in an empirical test ($d \approx 0.42$ in reported comparisons) with weaker/null conceptual differences — quality criteria are not magic transfer guarantees.

Applied (HYPOTHESIS): MindCraft’s production rubric should **hard-fail** on content correctness (wrong algebra, wrong key, invented theorem) before scoring tone or coherence. Soft-passing a fluent wrong wrap poisons Map events the way unverified GENQ keys do.

Kill: “Mostly correct / helpful vibe” soft-pass on math errors.

Survive: Correctness hard gate + learner-orientation (prior-knowledge / fade fit) as second tier.

Wound: Classroom/video instruments are not MindCraft RCTs — EXPLAINQA-* must calibrate thresholds on coach cards specifically.

CXXI.5 Pedagogical ability taxonomy for LLM tutors — guidance $\neq$ answer dump

FACT (unified AI-tutor eval): Maurya, Srivatsa, Petukhova, & Kochmar (2025, *NAACL*, doi:10.18653/v1/2025.naacl-long.57) — eight pedagogical dimensions for LLM tutor responses in math mistake-remediation contexts: mistake identification, mistake location, revealing the answer, providing guidance, actionability, coherence, tutor tone, human-likeness; release MRBench with human gold labels across LLM and human tutors.

Applied (FOUNDER BELIEF → testable): For MindCraft coach wraps, elevate **mistake location + correct guidance + actionability + non-default reveal** as ship dimensions. Demote **human-likeness / tone** to secondary polish — never primary quality. A wrap that reveals the answer by default fails SAFE-HINT / SAFE-FADE even if “encouraging.”

Kill: Human-likeness or warmth as the marketing proof of coach quality.

Survive: Non-reveal + actionable next-step as default pass criteria.

Wound: MRBench is dialogue remediation, not identical to MindCraft’s card wrap — map dimensions carefully; do not cargo-cult all eight into a student-facing score.

CXXI.6 Guardrails beat base fluency — learning ≠ practice-with-AI score

FACT (field experiment): Bastani, Bastani, Sungu, Ge, Kabakc■, & Mariman (2024/2025; SSRN abstract 4895486; PNAS-line results) — ~1,000 high-school math students; GPT Base improved practice performance but **harmed** unassisted exam performance (~17% worse vs control); GPT Tutor with pedagogical safeguards largely mitigated the harm; students often used Base as a solution crutch and did not perceive the learning loss.

Applied (HYPOTHESIS): Rubric gates that block answer-dumps and force principle-short guidance are not UX niceties — they are the product form of Bastani-style safeguards. Shipping fluent Base-like wraps because they raise “session success” is commercially suicidal for an identity company.

Kill: Assisted practice accuracy / wrap satisfaction as coach quality proof.

Survive: Unassisted / solo-transfer primary; guardrailed wrap secondary.

CXXI.7 Automated scoring is a tool — not a North Star product

FACT (math SE auto-score): Nakamoto, Flanagan, Yamauchi, Dai, Takami, & Ogata (2023, *Computers*, 12(11), 217, doi:10.3390/computers12110217) — automated scoring of math self-explanation quality using coherence / clarity / relevance rubrics (adapted from prior Nakamoto work); human raters reached substantial agreement (quadratic weighted $\kappa \approx 0.75$); LLM-augmented training data helped up to a point, then saturated.

Applied (HYPOTHESIS): MindCraft may use **rubric-structured auto-checks** (and LLM judges) as *internal* pre-filters for coach wraps — always with (a) deterministic math/key checks first, (b) human audit sampling, (c) no student/parent Explanation Score™ surface. Auto-score of *student* SE is a different product question (SAFE-SE / SAFE-PEERX); this chapter’s primary object is **AI wrap QA**.

Kill: Rubric Score™ / Quality Minutes as NS; LLM-judge-as-sole-authority.

Survive: Inspectable checklist + hard fails + audit cadence.

CXXI.8 Product surface — SAFE-EXPLAINQA claim contract

Gate	Pass look	Hard fail	Banned substitute
Content / key	Asserted math matches verify	Wrong key / invented rule	“Conceptually helpful” soft-pass
Principle-short	1–3 sentence critical feature / join	Monologue essay default	Thoroughness Score™
Non-reveal	Hint/principle grain; answer not dumped	Full solution as first wrap	“Never stuck” hero
Actionability	Clear next attempt move	Vague pep / lore	Tone-only polish
SE-preserve	SE before wrap when policy on	Wrap replaces required SE	AI monologue labeled SE
Fade fit	Shorter as expertise rises	Fixed long wrap all levels	Expertise-blind dump
Metrics	solo_transfer_pass; repair escalate rate; audit fail rate	Explanation Score™ / thumbs-up thoroughness	Fluency / human-likeness NS

Competitive foil: ChatGPT Base = fluent dump. “AI explained everything” = satisfaction theater. MindCraft = **gated principle-short coach**.

CXXI.9 Doctrine — SAFE-EXPLAINQA (provisional)

1. **Rubric before ship** — no ungated LLM wrap into Practice/Solver coach surfaces.
2. **Correctness hard-fail** — math/key errors block ship (SAFE-GENQ parity for language that asserts facts).
3. **Principle-short default** — SAFE-EXPLAIN length law is a *gate*, not a style tip.
4. **Non-reveal + actionable** — Maurya dims as production constraints; human-likeness demoted.
5. **SE before wrap** — instructional text must not crowd construction (Wittwer/Renkl; Schworm/Renkl).
6. **Guardrails > Base fluency** — Bastani lesson: assisted gains ≠ learning.
7. **No Explanation Score™ / Rubric Score™ NS** — internal QA only; FEI co-primaries prove.
8. **Human audit sample** — LLM-as-judge never sole authority (SAFE-HITL / SAFE-LABMETA).
9. **Copy:** “Our coach is short, checked, and leaves you the thinking.” Never “unlimited thorough AI explanations.”

Confidence: High — Wittwer/Renkl failure modes + small meta; Schworm SE crowding; Bastani guardrails; Maurya taxonomy existence. Medium — exact pass thresholds and auto-judge agreement on MindCraft cards (needs EXPLAINQA-*). High — fluency/thoroughness/Explanation Score™ as commercially toxic under this stack.

CXXI.10 Experiments

ID	Question	Design	Primary
EXPLAINQA-1	Principle-short gated wrap vs ungated fluent monologue	A/B	solo_transfer_pass; SE rate; dwell
EXPLAINQA-2	Correctness hard-fail + withhold vs soft-pass fluent-wrong	A/B	Map poison events; repair escalate; trust
EXPLAINQA-3	Non-reveal actionable wrap vs answer-dump wrap	A/B	Peek rate; unassisted transfer (Bastani-style)
EXPLAINQA-4	LLM-judge + human audit vs LLM-judge alone	Ops A/B	False-pass rate; audit catch; cost
EXPLAINQA-5	Parent CBC: “checked short coach” vs “unlimited thorough AI”	CBC	WTP; trust (SAFE-WTP)
EXPLAINQA-QUAL	10 students + 5 tutors: which wrap felt like teaching vs dumping?	Qual	EXPLAINQA codebook

Falsifier: Ungated fluent monologues beat gated wraps on solo transfer without SE collapse → revise length gate; still ban Explanation Score™ NS.

Falsifier: Soft-pass fluent-wrong improves short-term affect without Map poison → still hard-fail keys; segregate affect.

Falsifier: LLM-judge alone matches human audit at production risk tolerance → still keep sample audit for SAFE-GENQ parity.

Pre-register: EXPLAINQA-* before “science-backed AI explanations” ads (SAFE-LABMETA).

Family note: EXPLAINQA- *densifies* EXPLAIN- / GENQ-*; do not collapse into thumbs-up OKRs.

CXXI.11 So what for MindCraft commercially

- **Copy:** “Short, checked coach — you still do the thinking.” Lead with gated principle-short, not unlimited essays.
- **Product:** Ship checklist on coach/Solver wraps: key-correct → principle-short → non-reveal → actionable → SE-preserve; withhold/escalate on hard fail (SAFE-REPAIR).
- **Positioning:** Against ChatGPT Base dumps and “AI explained everything” brands; for guardrailed coach that protects solo transfer.
- **Metric:** solo_transfer_pass, wrap hard-fail rate, audit false-pass rate, SE-before-wrap compliance — demote Explanation Score™ / thoroughness thumbs.

- **Kill list:** Fluency $\equiv$ quality; thumbs-up NS; Explanation Score™; reveal-default; soft-pass wrong math; ACT guarantees from rubric packaging.
- **Growth:** Trust packets show *how* wraps are gated (correctness + pedagogy dims) — procurement asset vs black-box tutor LLM.
- **Vision:** Thirty-year identity company treats AI explanation as a **verified micro-scaffold**, not a fluent substitute for becoming a mathematical thinker.

References (verified)

- Bastani, H., Bastani, O., Sungu, A., Ge, H., Kabakc■, Ö., & Mariman, R. (2024). Generative AI can harm learning. *The Wharton School Research Paper*. <https://doi.org/10.2139/ssrn.4895486>
 - Kulgemeyer, C. (2018). A framework of effective science explanation videos informed by criteria for instructional explanations. *Research in Science Education*. <https://doi.org/10.1007/s11165-018-9787-7>
 - Maurya, K. K., Srivatsa, K. V. A., Petukhova, K., & Kochmar, E. (2025). Unifying AI tutor evaluation: An evaluation taxonomy for pedagogical ability assessment of LLM-powered AI tutors. In *Proceedings of NAACL*. <https://doi.org/10.18653/v1/2025.naacl-long.57>
 - Nakamoto, R., Flanagan, B., Yamauchi, T., Dai, Y., Takami, K., & Ogata, H. (2023). Enhancing automated scoring of math self-explanation quality using LLM-generated datasets: A semi-supervised approach. *Computers*, 12(11), 217. <https://doi.org/10.3390/computers12110217>
 - Ring, M., & Brahm, T. (2023). A rating framework for the quality of video explanations. *Technology, Knowledge and Learning*. <https://doi.org/10.1007/s10758-022-09635-5>
 - Schworm, S., & Renkl, A. (2006). Computer-supported example-based learning: When instructional explanations reduce self-explanations. *Computers & Education*, 46, 426–445. <https://doi.org/10.1016/j.compedu.2004.08.011>
 - Wittwer, J., & Renkl, A. (2008). Why instructional explanations often do not work: A framework for understanding the effectiveness of instructional explanations. *Educational Psychologist*, 43(1), 49–64. <https://doi.org/10.1080/00461520701756420>
 - Wittwer, J., & Renkl, A. (2010). How effective are instructional explanations in example-based learning? A meta-analytic review. *Educational Psychology Review*, 22, 393–409. <https://doi.org/10.1007/s10648-010-9136-5>
-

Part CXXII — Partial-Credit & Soft-Wrong Taxonomy Productization

Chapter status: Living evidence + Practice/Map mark-type brief — Researcher tick 2026-08-11 (UTC hour 03; hour%6≠0; researcher count since synthesizer v1.15 = 5 → Researcher)

Primary question: How should MindCraft turn **miss types** (slip vs bug vs lack vs partial process) into product mark taxonomy — without gradebook cosplay, Partial Credit Score™ vanity, or treating every red X as the same mastery event?

Owners: Product (Practice / Diagnostic / soft-wrong) · Engine (outcome events + misconception links) · HITL · Brand · Parent GTM · Red Team

Commercial job: Ship **SAFE-TAXON** densifying SAFE-MISCON × SAFE-FBLIT × SAFE-FBTIME: **categorical mark types drive routing and graph updates**; soft-wrong remains destake + next move — never half-credit theater

or a student-facing Partial Credit Score™.

Builds on: Parts XLIX (SAFE-MISCON), CXII (SAFE-FBLIT), LXXIX (SAFE-FBTIME), CXVIII (SAFE-SOFTMSG), L (SAFE-CALIB), LXXXIII (SAFE-PDASH), XC (SAFE-INSTRUMENT). Product seams: soft-wrong UX, Eedi mis_ IDs, Layer-4 misconception memory, outcome severity, parent digest grain.

CXXII.1 Why this chapter exists

SAFE-MISCON already kills “all wrongs are equal” as pedagogy. SAFE-FBLIT already kills delivery⇒uptake. What remains commercially underspecified is the **product mark layer**: when Maya’s answer is wrong (or only partly right), what *type* of mark fires, what the student sees, what the graph updates, and what parents are told — without reinventing a gradebook.

Default edtech resolves this badly: (1) dichotomous green/red only, (2) “partial credit points” that teach gaming the rubric, (3) soft tone with no diagnostic category, (4) Partial Credit Score™ / Accuracy % as NS, (5) parent portals that show miss counts without grain. Competitors can ship warmer copy; few ship **inspectable mark types** → **route** → **retry** as the wedge.

FOUNDER BELIEF under audit: Soft-wrong is evaluation *climate*; taxonomy is evaluation *semantics*. Climate without semantics is empty kindness. Semantics without destake is gradebook cosplay. MindCraft needs both — and must refuse points-as-learning theater.

Claims we refuse as doctrine:

1. Soft-wrong ≡ automatic half-credit / points-for-almost.
 2. Partial Credit Score™ / Mark Minutes / Accuracy Ladder as student or parent NS.
 3. All incorrect outcomes update mastery the same way (slip ≡ bug ≡ blank ≡ guess).
 4. Gradebook cosplay (letter grades, % correct banners, live point stalk).
 5. Dichotomous-forever as the only honest assessment (kills diagnostic signal).
 6. Ordered PCM steps ≡ identity progress without FEI transfer.
 7. Empty “we celebrate mistakes” without named mark→route.
 8. ACT / grade guarantees from partial-credit packaging.
-

CXXII.2 Constructs (taxonomy ≠ gradebook)

Construct	Research meaning	MindCraft analogue	Failure mode
Dichotomous mark	Correct / incorrect	Prove-rail KR; key verify	Loses process diagnosis
Partial credit (psychometric)	Ordered categories of completeness	Optional process steps for <i>routing</i> , not % NS	Points theater
Categorical / concept-based score	Nominal wrong <i>types</i> (slip, misconception, strategy)	Soft-wrong mark family + <code>misconception_id</code>	Composite Partial Credit Score™

Slip	Random / careless error; knowledge present	Light retry cue; weak mastery hit	Treat slip as deep bug
Bug / misconception	Persistent wrong idea interfering with performance	Route to mis_ family / bridge / SE springboard	Soft tone only
Lack / incomplete	Missing knowledge or unfinished path	Scaffold / fade grain; not shame	Flood with full dump
Soft-wrong	Destaked informative miss + next move	Climate wrapper around a mark type	Kindness without type
SAFE-TAXON	Mark-type product law	This chapter	Gradebook cosplay

Operational definition (HYPOTHESIS): A miss is *SAFE-TAXON complete* when the product (a) assigns a **named mark type** (at least: correct | slip | misconception/bug | lack/incomplete | blank/no-attempt | assist-tainted), (b) maps misconception/bug marks to a diagnostic family when available, (c) updates the graph with **type-conditioned** severity (not one red-X delta), (d) shows the student destaked next-move UX (soft-wrong climate), (e) never surfaces Partial Credit Score™ / % ladders as NS, and (f) proves value with `retry_120s / solo_transfer_pass / type→route` hit rate — not points recovered.

CXXII.3 Partial credit is a measurement tool — not a North Star product

FACT (PCM foundation): Masters (1982, *Psychometrika*, 47, 149–174, doi:10.1007/BF02296272) — the Rasch Partial Credit Model treats responses scored in two or more *ordered* categories; category structure may vary by item; the model supports objective person/item comparisons from graded responses.

FACT (dichotomous ≈ PC for some outcomes): May, Koskey, Bostic, Stone, Kruse, & Matney (2023, *School Science and Mathematics*, 123(2), 54–67, doi:10.1111/ssm.12570) — on the PSM6 constructed-response problem-solving measure, Rasch dichotomous and partial-credit scorings yielded similarly strong psychometric findings; student outcomes correlated strongly/positively; similar demographic patterns. Either method can be appropriate for *that* instrument’s score uses — purpose and constraints decide.

Applied (HYPOTHESIS): MindCraft may use ordered process categories internally (e.g., no attempt → wrong join → right join wrong finish → full) for **routing and analytics**. It must not brand PCM-style ladders as “Partial Credit Score™ progress” or sell ordered steps as identity transformation. May et al. also remind: if the *use* is remediation and misconception ID, dichotomous alone is often the wrong *classroom* tool even when total scores correlate.

Kill: Partial Credit Score™ / % almost-right as FEI substitute.

Survive: Ordered categories as engine grain when they change the next card — never as vanity NS.

Wound: PSM6 is a research assessment, not MindCraft Practice — TAXON-* must calibrate on soft-wrong + MC distractors specifically.

CXXII.4 CAA partial credit — steps for diagnosis, not points theater

FACT (computer-aided assessment): Ashton, Beevers, Korabinski, & Youngson (2006, *British Journal of Educational Technology*, 37(1), 93–119, doi:10.1111/j.1467-8535.2005.00512.x) — paper exams normally award partial credit for correct working with wrong finals; binary CAA marks can diverge from that practice. PASS-IT experiments showed optional **steps** / redesigned items can mimic partial credit in automated assessment; feedback modes and expression capture matter; education should drive the tech pattern (not the reverse).

Applied (HYPOTHESIS): MindCraft’s analogue is **step / join grain** on Solver and multi-step Practice — mark *which* step failed (bridge vs arithmetic vs FormatId) and route — not award “3/5 points” as the student-facing product. Parent-facing truth: “stuck on the join” beats “72% partial credit.”

Kill: Points-for-working as the brand of soft-wrong.

Survive: Optional step capture → mark type → principle-short coach (SAFE-EXPLAINQA) → retry.

CXXII.5 Categorical wrongs beat correctness-only for learning uses

FACT (concept-based categorical rubrics): Arieli-Attali & Liu (2016, *Educational Psychology*, 36(6), 1083–1101, doi:10.1080/01443410.2015.1031088) — diagnostic uses are limited when item responses are scored only for correctness; concept-based **nominal** categories preserve error type and strategy information lost under conventional correctness rubrics; Multiple Correspondence Analysis showed categorical scores capture what correctness captures *plus* additional performance aspects crucial for learning.

Applied (FOUNDER BELIEF → testable): Soft-wrong taxonomy should prefer **nominal diagnostic categories** (slip / misconception / lack / blank / assist) over a single ordered “how close” number when the commercial job is routing. Closeness can be a secondary engine feature; type is primary for Map and coach.

Kill: Correctness-only telemetry as “we diagnose.”

Survive: Category → family ID → next move; prove with route hit + transfer, not category count vanity.

CXXII.6 Slips vs bugs — different updates, different coaches

FACT (diagnostic assessment framing): Ketterlin-Geller & Yovanoff (2009, *Practical Assessment, Research & Evaluation*, 14(16), doi:10.7275/vxrk-3190) — error analysis classifies errors into **slips** (random errors in declarative/procedural knowledge *not* from inherent domain misunderstanding) vs **bugs** (persistent misconceptions that consistently interfere). Identifying bugs is the primary interest of diagnostic assessment for instructional decisions; skills analysis that only flags “missed subtraction” without *why* has limited remediation utility.

FACT (error-analysis tradition): Radatz (1979, *Journal for Research in Mathematics Education*, 10(3), 163–172, doi:10.5951/jresmetheduc.10.3.0163) — systematic error analysis has long diagnostic value for teachers; patterns matter more than isolated wrongs.

Reuse (SAFE-MISCON): Smith, diSessa & Roschelle — misconceptions are flawed *and* productive; refinement routes beat purge theater.

Applied (HYPOTHESIS): Graph update law: **slip** → small/no concept mastery hit + light retry; **bug/misconception** → stronger evidence + springboard SE + mis_ route; **lack** → scaffold/fade, not confrontation; **blank** → retrieval/help policy (SAFE-RETRIEVE / SAFE-HINT), not “wrong idea” labeling. Collapsing these into one red-X poisons both Map honesty and student identity.

Kill: One mastery delta for every miss.

Survive: Type-conditioned outcomes + tutor briefs that distinguish slip vs bug (SAFE-HITL).

Wound: Teachers (and models) often confuse slips and bugs — product needs multi-signal classification (latency, confidence, distractor map, repeat pattern), not a single LLM guess as ground truth (SAFE-LABMETA / SAFE-EXPLAINQA).

CXXII.7 Soft-wrong climate wraps taxonomy — it does not replace it

Reuse (SAFE-MISCON / SAFE-ERRCLIMATE / SAFE-SOFTMSG): Destake makes struggle survivable; parent copy after streaks must name grain + one CTA, not Miss Streak scoreboards.

HYPOTHESIS: Soft-wrong UX without taxonomy is Duo-with-nicer-copy. Taxonomy without soft-wrong is a gradebook. Product contract: **climate first, then typed mark, then actionable uptake** (SAFE-FBLIT).

SPECULATION: Parents will tolerate destake longer when digests say “pattern: equal-sign join (bug family)” than when they see “3 wrongs / 60% partial.” Honesty is categorical, not percentage theater.

Kill: Soft-wrong slogan without mark types; Partial Credit % parent widgets.

Survive: Grain-named soft-wrong + type-private analytics for tutors/engine.

CXXII.8 Product surface — SAFE-TAXON claim contract

Mark type	Student sees	Graph / route	Banned substitute
Correct	Destaked confirm / prove KR by mode	Positive evidence	Fireworks mastery belt
Slip	Soft retry cue; “check that step”	Weak negative / ignore if isolated	Deep misconception badge
Bug / misconception	Soft-wrong + SE / contrast + mis_ card	Stronger evidence; family link	Shame reel; purge language
Lack / incomplete	Soft scaffold / fade grain	Incomplete event; not “wrong theory”	Full solution dump default
Blank / no-attempt	Wait / tip / contingent hint policy	Non-commission event	Struggle Score™ for dwell
Assist-tainted	Labeled help path; solo gate later	Do not count as solo proof	Assisted⇒ready banner

Metrics	retry_120s; type→route hit; solo_transfer_pass	—	Partial Credit Score™ / %
---------	--	---	------------------------------

Competitive foil: Khan/Duo = binary or XP. Grade apps = %. ChatGPT = fluent re-explain without type. MindCraft = **typed soft-wrong** → **route**.

CXXII.9 Doctrine — SAFE-TAXON (provisional)

1. **Type before points** — every scored miss carries a mark type; points/% are not the product.
2. **Soft-wrong wraps type** — climate is destake + next move; taxonomy is the semantic payload.
3. **Slip ≠ bug ≠ lack ≠ blank** — different coaches and different graph updates (Ketterlin-Geller & Yovanoff).
4. **Categorical > closeness for routing** — Arieli-Attali & Liu; ordered PCM steps optional for process grain only.
5. **Steps diagnose joins** — Ashton et al. CAA lesson: capture where process breaks; do not sell partial marks as trophies.
6. **No Partial Credit Score™ / gradebook cosplay** — no letter grades, live % stalk, or Miss Streak NS (SAFE-PDASH / SAFE-SOFTMSG).
7. **Misconception families when known** — Eedi / Layer-1 links; else “untyped bug” + SE, not invented disease labels.
8. **Human/audit for ambiguous type** — model classifier assists; sample audit; never LLM-type as sole authority.
9. **Copy:** “We name the kind of miss — then the next move.” Never “earn partial credit like school.”

Confidence: High — Masters PCM existence; May et al. purpose-dependence; Ashton CAA steps; Arieli-Attali & Liu categorical gain; Ketterlin-Geller slips/bugs; Radatz diagnostic tradition. Medium — MindCraft auto-type accuracy and severity weights (needs TAXON-*). High — Partial Credit Score™ / all-wrongs-equal / gradebook cosplay as commercially toxic under FEL.

CXXII.10 Experiments

ID	Question	Design	Primary
TAXON-1	Typed soft-wrong (slip/bug/lack) vs binary soft-wrong	A/B	retry_120s; solo_transfer_pass; route hit
TAXON-2	Type-conditioned mastery deltas vs uniform miss delta	Offline + A/B	Calibration; transfer; false “mastered”
TAXON-3	Step-join mark (CAA-style) vs final-answer-only	A/B multi-step items	Join repair rate; abandon
TAXON-4	Parent digest: categorical grain vs partial-credit %	CBC	WTP; trust; nag/shame (SAFE-WTP/SOFTMSG)

TAXON-5	Auto-type + human audit vs auto-type alone	Ops	False bug-rate; cost; tutor override
TAXON-QUAL	10 students + 5 tutors: which mark language felt informative vs graded?	Qual	TAXON codebook

Falsifier: Binary soft-wrong matches typed on transfer and route quality → keep private engine types; still ban Partial Credit Score™ NS.

Falsifier: Points-for-almost raises retention without hurting solo transfer → still refuse % NS; segregate engagement.

Falsifier: Auto-type equals human audit at acceptable risk → keep sample audit (SAFE-LABMETA).

Pre-register: TAXON-* before “diagnostic partial credit” ads.

Family note: TAXON-* densifies MISCON/FBLIT/FBTIME; do not collapse into Soft-Wrong Score™ or Partial Credit OKRs.

CXXII.11 So what for MindCraft commercially

- **Copy:** “We name the kind of miss — then the next move.” Lead with typed soft-wrong, not school-style partial credit.
- **Product:** Ship mark enum on outcomes; wire slip/bug/lack/blank to distinct coaches and graph deltas; optional step-join capture on multi-step items.
- **Positioning:** Against binary apps, gradebook portals, and ChatGPT re-explain-without-type; for diagnostic climate that protects identity while updating Map honestly.
- **Metric:** type→route hit, retry_120s, solo_transfer_pass, false-bug rate — demote Partial Credit Score™ / Accuracy %.
- **Kill list:** Soft-wrong≡half-credit; all-wrongs-equal mastery; gradebook cosplay; Partial Credit Score™; ACT guarantees from taxonomy UI.
- **Growth:** Trust packets show mark taxonomy + audit — procurement asset vs black-box “AI knows why you missed.”
- **Vision:** Thirty-year identity company treats wrong answers as **typed information**, not moral failure or points residue — so Maya becomes someone who uses misses, not someone who collects partial credit.

References (verified)

- Arieli-Attali, M., & Liu, Y. (2016). Beyond correctness: Development and validation of concept-based categorical scoring rubrics for diagnostic purposes. *Educational Psychology*, 36(6), 1083–1101. <https://doi.org/10.1080/01443410.2015.1031088>
- Ashton, H. S., Beevers, C. E., Korabinski, A. A., & Youngson, M. A. (2006). Incorporating partial credit in computer-aided assessment of Mathematics in secondary education. *British Journal of Educational Technology*, 37(1), 93–119. <https://doi.org/10.1111/j.1467-8535.2005.00512.x>

- Ketterlin-Geller, L. R., & Yovanoff, P. (2009). Diagnostic assessments in mathematics to support instructional decision making. *Practical Assessment, Research & Evaluation*, 14(16). <https://doi.org/10.7275/vxrk-3190>
- Masters, G. N. (1982). A Rasch model for partial credit scoring. *Psychometrika*, 47, 149–174. <https://doi.org/10.1007/BF02296272>
- May, T. A., Koskey, K. L. K., Bostic, J. D., Stone, G. E., Kruse, L. M., & Matney, G. T. (2023). Examining how using dichotomous and partial credit scoring models influence sixth-grade mathematical problem-solving assessment outcomes. *School Science and Mathematics*, 123(2), 54–67. <https://doi.org/10.1111/ssm.12570>
- Radatz, H. (1979). Error analysis in mathematics education. *Journal for Research in Mathematics Education*, 10(3), 163–172. <https://doi.org/10.5951/jresmetheduc.10.3.0163>

Part CXXIII — Exam-Week Tutor Load Shedding Playbooks

Chapter status: Living evidence + HITL exam-crunch ops — Researcher tick 2026-08-11 (UTC hour 06 ≡ Red Team slot, but ch123 never written → prefer Researcher per rotation; researcher count since synthesizer v1.15 = 6 → Researcher)

Primary question: In the final weeks before ACT/SAT/state tests, how should MindCraft **shed and reallocate tutor load** so FEI fidelity survives the crunch — without hours-booked theater, wellness-as-fix, hire-surge without coaches, or turning every panic session into pour-mode cram?

Owners: HITL ops / workforce · Product (tutor console / scheduling) · Exam calendar · Parent GTM · Brand · Red Team

Commercial job: Ship **SAFE-LOADSHED** densifying SAFE-BURNWARN × SAFE-DDSCHED × SAFE-WORKFORCE × SAFE-EXAM × SAFE-ROI: a **named exam-week playbook** that frees tutor minutes for highest-leverage FEI work by pausing/rerouting lower-stakes demand — never a Load Shed Score™, crunch hours NS, or ACT-point guarantees from “we stayed open late.”

Builds on: Parts CXX (SAFE-BURNWARN), CXIV (SAFE-DDSCHED), LXXV (SAFE-WORKFORCE), LXIII (SAFE-EXAM), XCIX (SAFE-ROI), LXVIII (SAFE-HITL), LXXVIII (SAFE-DURABLE/CRAM), XC (SAFE-INSTRUMENT). Product seams: ACT date on profile, caseload caps, Map-brief gate, dual-rail train/prove, Solver vs Practice routing, parent crunch messaging.

CXXIII.1 Why this chapter exists

SAFE-BURNWARN already says: fidelity dips + JD-R redesign beat Burnout Score™ theater. SAFE-DDSCHED already says: hold equipped mix; ease load not standards; Protect+probe in the final ~72h. What remains commercially underspecified is the **ops join under calendar surge**: parents book more, students panic-request homework dumps, tutors accept “just this week” pour sessions, and marketplace logic celebrates hours-booked while brief-skip and soft-wrong avoidance spike.

Competitors resolve crunch wrongly: (1) **always-on** — take every session, sell availability as care; (2) **wellness email** — mindfulness tip while caseload doubles; (3) **hire flood** without coach bandwidth (demand bomb); (4) **cram costume** — tutor becomes human Chegg for “exam support”; (5) **soft standards** for anxious personas under deadline. MindCraft’s wedge is a **load-shedding playbook** that protects thinking minutes the same way hospitals protect ICU

beds: reverse-triage low-acuity demand, cancel/defer electives, keep resources for cases that change outcomes.

FOUNDER BELIEF under audit: Exam week is a predictable surge. Treating it as surprise heroism (everyone works harder) is a design failure. Honest companies publish what they *stop* doing so tutors can still make students think.

Claims we refuse as doctrine:

1. Hours-booked / session-volume as exam-week North Star.
2. Load Shed Score™ / Crunch Minutes / Availability Score™ as brand.
3. Wellness-module-as-primary when the demand spike is caseload.
4. Uncapped “we’ll take anyone who needs us” during Protect+probe.
5. Hire surge without restoring coach/QA bandwidth.
6. Tutor pour-mode / answer-dump as “exam support.”
7. Soft standards / skip-challenge for crunch personas.
8. ACT / score guarantees from crunch playbooks or late-night packing.

CXXIII.2 Constructs (shed ≠ abandon; surge ≠ heroics)

Construct	Meaning	MindCraft analogue	Failure mode
Surge demand	Predictable spike near exam dates	Parent bookings; HW-help flood; prove-week mocks	Surprise heroics
Load shed	Deliberately stop/defer low-leverage work	Pause elective Map tours; defer low-acuity HW dump; AI-scaffold first	Abandon students
Reverse triage (borrowed)	Free capacity by releasing low-risk demand	Route stable practice to Practice+coach; keep HITL for FEI-critical	Hospital cosplay brand
Acuity	Likelihood session changes FEI path	Soft-wrong streak + bridge gap + exam-horizon	First-come booking
Elective cancel	Defer non-essential ops under surge	Lore walks; admin meetings in crunch window	Cancel FEI itself
SAFE-LOADSHED	Named crunch playbook	This chapter	Hours-booked NS

Operational definition (HYPOTHESIS): An exam-week load-shed is *SAFE-LOADSHED complete* when (a) triggers are calendar + demand/resource signals (days-to-exam band + caseload/fidelity alerts), (b) a published playbook **names what pauses** (elective Map tours, unbounded HW dumps, non-FEI admin, uncapped new intakes), (c) remaining HITL minutes prioritize **high-acuity FEI** (equipped struggle, bridge joins, labeled prove probes) over comfort greens, (d) AI/Solver scaffolds absorb low-acuity homework under SAFE-HWHELP rails — not as Chegg dump, (e) org redesign (caps, redistribute, pause hire without coaches) runs before wellness theater, and (f) success = fidelity@crunch +

CXXIII.3 Exam accountability raises educator stress and can warp instruction — crunch is not neutral

FACT (test-based accountability → stress + counterproductive practices): von der Embse, Schoemann, Kilgus, Wicoff, & Bowler (2016, *Educational Psychology*, 37(3), 312–331, doi:10.1080/01443410.2016.1183766) — SEM on N=7,281 educators: educators who saw tests as invalid measures of effectiveness reported higher test stress and greater use of **counterproductive instructional practices** (including fear appeals) intended to motivate test-taking.

FACT (efficacy buffers; stress still real): von der Embse, Sandilos, Pendergast, & Mankin (2016, *Learning and Individual Differences*, 50, 308–317, doi:10.1016/j.lindif.2016.08.001) — teacher self-efficacy moderated links between test-stress and job satisfaction; efficacy support helps but does not erase accountability pressure.

Applied (HYPOTHESIS): Parent/tutor crunch culture that rewards fear appeals, pour-mode, and “grind tonight” is the edtech analogue of counterproductive practice. SAFE-LOADSHED must **block** those defaults — not celebrate them as hustle.

Kill: Fear-appeal / pour-as-care exam marketing. **Survive:** Crunch copy that names protected FEI behaviors and paused electives.

Wound: School-accountability samples ≠ near-peer tutors — LOADSHED-* must measure pour/brief-skip under ACT date spikes.

CXXIII.4 Work intensification is a policy/org problem — not a character deficit

FACT (workload / intensification synthesis): Creagh, Thompson, Mockler, Stacey, & Hogan (2023, *Educational Review*, doi:10.1080/00131911.2023.2196607) — systematic synthesis: teacher workload intensification and time poverty are tightly linked to accountability, datafication, and performativity cultures; the literature repeatedly frames heavier work as a **system design** outcome, not individual laziness.

Applied (HYPOTHESIS): Exam-week “just do more sessions” without removing other demands repeats intensification. Load shed is the commercial translation of *zero-sum work design*: new crunch demand must displace named electives, or fidelity collapses (SAFE-BURNWARN / SAFE-WORKFORCE).

Kill: Character/grit framing for tutor crunch (“our tutors go the extra mile” as the product).

Survive: Published displacement list + caseload caps tied to ACT horizon.

CXXIII.5 Surge capacity lesson — free capacity by releasing low-acuity demand, not by burning staff

FACT (reverse triage founding empirics): Kelen, McCarthy, Kraus, et al. (2009, *Disaster Medicine and Public Health Preparedness*, 3(Suppl 1), S10–S16, doi:10.1097/DMP.0b013e3181a5e7cd) — across three capacity-constrained hospitals (N=3,491 patients), ~44% needed no critical intervention for 4 days and were suitable for early discharge; reverse triage contributed the **majority** of realized surge beds (~50–59%); most capacity appeared within 24–48 hours.

FACT (pediatric reverse triage magnitude): Kelen, Troncoso, Trebach, et al. (2017, *JAMA Pediatrics*, 171(4), e164829, doi:10.1001/jamapediatrics.2016.4829) — ~10.8% eligible for immediate low-risk reverse triage and ~13.2% by 96 hours; effect unit-concentrated; reverse triage is one tool among cancel-electives / open-bed strategies — not infinite capacity.

FACT (systematic review stock-take): De Bondt, Sabbe, & Pollaris (2023, *Frontiers in Disaster and Emergency Medicine*, doi:10.3389/femer.2023.1303809) — reverse triage often yields ~**20–30%** additional surge capacity (more combined with other strategies); early-discharge adverse events generally low when multidisciplinary gates apply; heterogeneity warrants caution.

Applied (FOUNDER BELIEF → testable analogy, not hospital brand): Do **not** sell “disaster medicine for tutors.” Transferable mechanism: under surge, **release low-acuity demand** (defer/route/AI-scaffold) so scarce HITL minutes serve high-acuity FEI. Elective-cancel analogue = pause lore tours, unbounded Chegg-style HW, uncapped intakes, nonessential admin in Protect+probe.

Kill: Infinite-availability brand; hospital-cosplay marketing. **Survive:** Acuity queue + named pause list + caps/redistribute/AI rail.

Wound: Clinical AE ≠ tutoring harm — LOADSHED-* defines adverse events as fidelity collapse / pour relapse / abandonment proxies.

CXXIII.6 Org redesign and resource crafting beat person-only crunch coping

FACT (org > person burnout interventions, reuse): Panagioti et al. (2017, *JAMA Internal Medicine*, 177(2), 195–205, doi:10.1001/jamainternmed.2016.7674) — organization-directed burnout interventions outperform physician-directed ones on average.

FACT (JD-R stock-take, reuse): Bakker & Demerouti (2017, *Journal of Occupational Health Psychology*, 22(3), 273–285, doi:10.1037/ocp0000056) — dual pathways: demands→impairment; resources→motivation; interventions must address both sides.

FACT (job crafting longitudinal): Tims, Bakker, & Derks (2013, *Journal of Occupational Health Psychology*, 18(2), 230–240, doi:10.1037/a0032141) — employees who crafted job resources increased structural/social resources over ~2 months; resource increase linked to higher engagement/satisfaction and lower burnout; crafting challenging demands showed direct well-being links.

Applied (HYPOTHESIS): Exam-week playbooks should (1) **cut demands** (caps, defer electives, reverse-triage low-acuity HW to SAFE-HWHELP rails) and (2) **add resources** (coach bandwidth, Map grain tools, clearer acuity queues) — optionally invite tutors to craft *within* the playbook (choose which low-acuity slots to shed first). Person-only “self-care Sunday” fails Panagioti when caseload is the spike.

Kill: Wellness tip as primary crunch response; Burnout Score™ during ACT week.

Survive: Demand cut + resource add with tutor craft rights inside the shed list.

CXXIII.7 Product surface — SAFE-LOADSHED claim contract

Trigger	Shed / reroute	Protect	Banned substitute
Enter Join / Protect+probe band (SAFE-DDSCHED)	Pause elective Map tours; cap new uncapped intakes	Equipped mix + labeled prove	Comfort-week blocked greens
Caseload > crunch cap or BURNWARN fidelity dip	Redistribute; defer low-acuity HW to Solver rails	Brief-open FEI sessions	“Take one more” heroics
Parent panic booking flood	Acuity queue; transparent wait / AI first step	High-acuity HITL	Hours-booked NS
Homework dump requests	SAFE-HWHELP dual-rail; no Chegg pour	Solo attempt grain	Unlimited human entitlement
Hire pressure	Pause hire until coach_util recovers	QA sample cadence	Volume marketplace
Metrics	fidelity@crunch; brief-open; pour↓; transfer_pass	Co-gate with FEI	Load Shed Score™; Crunch Minutes

Competitive foil: Marketplace = take all hours. Bootcamp = pour cram. Wellness app = person-fix. MindCraft = named shed list + acuity FEI under deadline.

CXXIII.8 Doctrine — SAFE-LOADSHED (provisional)

1. **Crunch is predictable surge** — calendar triggers the playbook; surprise heroics are a design bug.
2. **Shed low-acuity demand** — reverse-triage analogy: free HITL capacity without burning staff (Kelen line; De Bondt et al.) — never hospital-cosplay ads.
3. **Displace, don’t stack** — new crunch demand must cancel named electives (Creagh intensification lesson).
4. **Protect FEI fidelity under load** — Map-brief, soft-wrong climate, equipped struggle stay; pour-mode is not “exam support” (von der Embse counterproductive practices).
5. **Org redesign first** — caps/redistribute/coach bandwidth before wellness theater (Panagioti; SAFE-BURNWARN).
6. **Resource craft inside the playbook** — tutors may prioritize shed order; they may not invent uncapped hours as virtue (Tims et al.).
7. **Join SAFE-DDSCHED** — Protect+probe student calendar and tutor load shed are one system; ease load not readiness standards.
8. **No Load Shed Score™ / hours-booked NS / ACT guarantees from crunch packaging.**

Confidence: High that uncapped crunch + pour-mode + wellness-only responses are commercially and pedagogically unsafe vs FEI. Medium on exact acuity ranks and % capacity freed in MindCraft (needs LOADSHED-*). High to kill hours-booked NS / hospital cosplay / ACT guarantees from playbooks regardless.

CXXIII.9 Experiments

ID	Question	Design	Primary
LOADSHED-1	Named shed playbook vs uncapped “take all” crunch	A/B sites near ACT dates	fidelity@crunch; pour rate; brief-open
LOADSHED-2	Acuity queue vs first-come booking	A/B	transfer_pass; wait fairness; parent trust
LOADSHED-3	HW low-acuity → Solver rail first vs human dump	A/B	solo_transfer_pass; assist-tainted rate
LOADSHED-4	Caseload cap + redistribute vs wellness-email-only	A/B	drift channels; quit proxy; FEI
LOADSHED-5	Parent CBC: “we protect thinking minutes” vs “24/7 availability”	CBC	WTP; trust (SAFE-WTP)
LOADSHED-QUAL	10 tutors + 8 parents in crunch: what felt like care vs abandonment?	Qual	LOADSHED codebook

Falsifier: Uncapped hours-booked sites beat shed playbooks on transfer_pass *and* fidelity@crunch without pour[↑] → revise acuity model; still ban Load Shed Score™ as NS.

Falsifier: Parents refuse any wait under shed messaging → redesign transparency + AI first-step; do not silently return to pour marketplace.

Falsifier: Hospital analogy confuses brand in CBC → drop cosplay language; keep mechanism.

Pre-register: LOADSHED-* before “we never turn students away in exam week” ads (SAFE-LABMETA).

Family note: LOADSHED- *densifies* BURNWARN-/DDSCHEd-/WORK-/EXAM-/ROI-; do not collapse into hours-booked OKRs.

CXXIII.10 So what for MindCraft commercially

- **Copy:** “In exam week we protect the minutes that make you think — and we pause the busywork that burns tutors out.” Never “always available” or “our tutors grind with you all night.”

- **Product:** Exam-horizon playbook in tutor console: acuity queue, named pause list, caseload caps, HW→Solver rail, hire gated by coach_util.
 - **Positioning:** Against marketplace volume, bootcamp pour, wellness-as-fix; for surge-honest HITL that still ships FEI under ACT dates.
 - **Metric:** fidelity@crunch, brief-open, pour↓, transfer_pass / solo_transfer_pass — demote hours-booked, Load Shed Score™, Crunch Minutes.
 - **Kill list:** Infinite availability; pour=exam support; wellness-primary under spike; hire without coaches; ACT guarantees from crunch packaging.
 - **Growth:** Parent/district trust packet includes crunch playbook (what pauses, what protects) — procurement honesty asset.
 - **Vision:** Thirty-year identity company treats exam week as **designed surge capacity**, not a character test for near-peer tutors.
-

References (verified)

- Bakker, A. B., & Demerouti, E. (2017). Job demands–resources theory: Taking stock and looking forward. *Journal of Occupational Health Psychology*, 22(3), 273–285. <https://doi.org/10.1037/ocp0000056>
 - Creagh, S., Thompson, G., Mockler, N., Stacey, M., & Hogan, A. (2023). Workload, work intensification and time poverty for teachers and school leaders: A systematic research synthesis. *Educational Review*. <https://doi.org/10.1080/00131911.2023.2196607>
 - De Bondt, F., Sabbe, M., & Pollaris, G. (2023). Reverse triage: A systematic review of the literature. *Frontiers in Disaster and Emergency Medicine*, 1, 1303809. <https://doi.org/10.3389/femer.2023.1303809>
 - Kelen, G. D., McCarthy, M. L., Kraus, C. K., Ding, R., Hsu, E. B., Li, G., Shahan, J. B., Scheulen, J. J., & Green, G. B. (2009). Creation of surge capacity by early discharge of hospitalized patients at low risk for untoward events. *Disaster Medicine and Public Health Preparedness*, 3(Suppl 1), S10–S16. <https://doi.org/10.1097/DMP.0b013e3181a5e7cd>
 - Kelen, G. D., Troncoso, R., Trebach, J., et al. (2017). Effect of reverse triage on creation of surge capacity in a pediatric hospital. *JAMA Pediatrics*, 171(4), e164829. <https://doi.org/10.1001/jamapediatrics.2016.4829>
 - Panagioti, M., Panagopoulou, E., Bower, P., et al. (2017). Controlled interventions to reduce burnout in physicians: A systematic review and meta-analysis. *JAMA Internal Medicine*, 177(2), 195–205. <https://doi.org/10.1001/jamainternmed.2016.7674>
 - Tims, M., Bakker, A. B., & Derks, D. (2013). The impact of job crafting on job demands, job resources, and well-being. *Journal of Occupational Health Psychology*, 18(2), 230–240. <https://doi.org/10.1037/a0032141>
 - von der Embse, N. P., Sandilos, L. E., Pendergast, L., & Mankin, A. (2016). Teacher stress, teaching-efficacy, and job satisfaction in response to test-based educational accountability policies. *Learning and Individual Differences*, 50, 308–317. <https://doi.org/10.1016/j.lindif.2016.08.001>
 - von der Embse, N. P., Schoemann, A. M., Kilgus, S. P., Wicoff, M., & Bowler, M. C. (2016). The influence of test-based accountability policies on teacher stress and instructional practices: A moderated mediation model. *Educational Psychology*, 37(3), 312–331. <https://doi.org/10.1080/01443410.2016.1183766>
-

Part CXXIV — Distractor Design for Diagnostic Soft-Wrong MC

Chapter status: Living evidence + Practice/Diagnostic item-craft brief — Researcher tick 2026-08-11 (UTC hour 09; hour%6≠0; researcher count since synthesizer v1.15 = 7 → Researcher)

Primary question: How should MindCraft design **MC distractors** so a wrong pick becomes typed soft-wrong signal (misconception → route) — without beauty-trap options, four-option cosplay, Distractor Score™ vanity, or treating every miss as the same red X?

Owners: Product (Practice / Diagnostic / bank) · Engine (mis_ links + outcome events) · GENQ / bank ops · HITL · Brand · Parent GTM · Red Team

Commercial job: Ship **SAFE-DISTRACTOR** densifying SAFE-TAXON × SAFE-MISCON × SAFE-GENQ: **each wrong option maps a plausible misconception (or slip-adjacent trap)**; soft-wrong climate wraps the typed pick; no option-count theater or beauty distractors.

Builds on: Parts XLIX (SAFE-MISCON), CXXII (SAFE-TAXON), LXXII (SAFE-GENQ), LXXXVII (SAFE-COVER), L (SAFE-CALIB), LXXIX (SAFE-FBTIME), CXVIII (SAFE-SOFTMSG). Product seams: Eedi mis_ IDs, actMaster / gap-scan MC, generated MC verify loops, soft-wrong UX.

CXXIV.1 Why this chapter exists

SAFE-TAXON requires mark *types*. For MC Practice and gap-scan, the cheapest type signal is **which distractor was chosen** — *if* distractors were written as diagnostic probes, not as random wrong numbers that look pretty next to the key.

Default edtech resolves this badly: (1) stem-first writing with leftover distractors, (2) four/five options as house style even when only one distractor functions, (3) “none of the above / all of the above” fillers, (4) LLM-generated options that are fluent but unmapped, (5) Distractor Diversity Score™ / option-count as quality theater, (6) marketing “AI wrote N questions” without distractor→misconception seals (SAFE-GENQ / SAFE-COVER).

FOUNDER BELIEF under audit: Soft-wrong MC is only diagnostic when **wrong picks encode theories**. Beauty traps and non-functioning options waste Maya’s attempt and poison Map updates. MindCraft’s bank moat is not item count — it is **keyed options with inspectable misconception families**.

Claims we refuse as doctrine:

1. Four/five options ≡ better assessment (option-count cosplay).
 2. Distractor Score™ / Option Minutes / “richer MC” NS.
 3. Beauty traps / absurd fillers / NOTA-AOTA as default craft.
 4. Unmapped wrong options as “diagnostic soft-wrong.”
 5. LLM distractors≡shipped bank without verify + mis_ seal (SAFE-GENQ).
 6. Selection frequency alone ≡ identity / mastery (no route).
 7. Empty “we learn from mistakes” without distractor→route.
 8. ACT / grade guarantees from distractor packaging.
-

CXXIV.2 Constructs (functioning distractor ≠ pretty wrong answer)

Construct	Research meaning	MindCraft analogue	Failure mode
Key	Correct / best option	Verified answer (SAFE-GENQ)	Fluency without key verify
Distractor	Plausible incorrect option	Option linked to mis_ / slip trap	Random wrong number
Functioning distractor	Selected by enough examinees; negatively discriminates	Option with choice share + ability pattern	<5% never-picked beauty trap
Non-functioning distractor	Rarely chosen or positively discriminating	Dead option; rewrite or drop	Keep for “four choices” costume
Diagnostic MC	Each wrong maps a common misconception	Soft-wrong typed by option	Binary red-X only
Ordered MC (OMC)	Options map developmental levels	Optional progression grain (not % NS)	Level Score™ theater
SAFE-DISTRACTOR	Distractor craft + QA law	This chapter	Option-count / beauty traps

Operational definition (HYPOTHESIS): An MC item is *SAFE-DISTRACTOR complete* when (a) every distractor is **plausible** under Haladyna-class guidelines, (b) each distractor carries a **named diagnostic intent** (misconception family, slip-adjacent arithmetic trap, or FormatId trap — not “wrong”), (c) dead options are retired after item analysis (or never shipped), (d) soft-wrong UX surfaces type + next move from the *chosen* option, (e) generated items fail ship without distractor→mis_ (or explicit untyped-bug) seal, and (f) value is proved by type→route hit / retry_120s / solo_transfer_pass — not option count.

CXXIV.3 Item-writing law: plausible, misconception-based choices

FACT (guideline taxonomy): Haladyna, Downing, & Rodriguez (2002, *Applied Measurement in Education*, 15(3), 309–334, doi:10.1207/S15324818AME1503_5) — validated a 31-guideline MC taxonomy from textbook consensus + empirical studies; **choice-writing** is the largest cluster, including: base distractors on common student errors / misconceptions; keep options homogeneous; avoid clues (clang, grammar, absurd options, “all/none of the above” traps); place the key randomly; keep option length balanced.

Applied (HYPOTHESIS): MindCraft bank craft starts from **known error families** (Eedi mis_, Layer-1 failure_mode, tutor error logs) and *then* writes options — never stem→fill-three-wrongs. Absurd or length-cued options are craft bugs, not “engagement spice.”

Kill: Beauty traps and clue-laden options as house style.

Survive: Homogeneous, misconception-sourced options + key-verify (SAFE-GENQ).

CXXIV.4 Functioning distractors — quality over quantity

FACT (functioning-distractor scarcity): Tarrant, Ware, & Mohammed (2009, *BMC Medical Education*, 9, 40, doi:10.1186/1472-6920-9-40) — analyzed 514 four-option nursing MCQs (1,542 distractors); non-functioning = chosen by <5% or positively discriminating. Mean **1.54** functioning distractors per item; only **13.8%** of items had three functioning distractors; **47.8%** of distractors non-functioning; more functioning distractors → harder, more discriminating items. Conclusion: write as many *plausible* options as content allows — **in most cases three** (key + two distractors).

FACT (meta-analytic efficiency): Rodriguez (2005, *Educational Measurement: Issues and Practice*, 24(2), 3–13, doi:10.1111/j.1745-3992.2005.00006.x) — meta-analysis across decades: moving from 4/5 options to **3** has little average cost to discrimination/reliability when ineffective distractors are deleted; more 3-option items can be administered per testing time → better content coverage. Randomly deleting options (vs deleting non-functioning ones) hurts more — craft quality matters.

Applied (HYPOTHESIS): MindCraft should **not** brand “ACT-style four choices” as pedagogy science for Practice/diagnostic. Prefer **2–3 functioning diagnostic distractors**. Keep four only when a fourth maps a *distinct* misconception with evidence of function. Gap-scan time budget favors fewer, sharper options (SAFE-ELL load; SAFE-ATTN).

Kill: Four-option costume as quality signal.

Survive: Functioning-count QA; retire dead options; option N as craft consequence, not brand.

Wound: High-stakes ACT still uses four options — Practice may *simulate* exam format on prove-rail while diagnostic/learn rails optimize for functioning diagnostic options. Label the dual-rail (SAFE-EXAM / SAFE-DDSCHED).

CXXIV.5 Distractor science is its own craft — not a stem afterthought

FACT (comprehensive review): Gierl, Bulut, Guo, & Zhang (2017, *Review of Educational Research*, 87(6), 1082–1116, doi:10.3102/0034654317726529) — distractor *development, analysis, and use* is under-attended relative to stem writing; quality work includes CTT/IRT distractor analysis **and** cognitive-diagnostic uses; recommendations emphasize misconception-based development, analysis of choice frequencies / discrimination, and caution that automatic generation needs evaluation — distractors are instructional information, not leftover wrongs.

Applied (FOUNDER BELIEF → testable): Bank OKRs should track % **options with mis_ (or slip/FormatId) seal** and % **functioning on first live cohort**, not items authored. GENQ prompts must allocate tokens to distractor intent equal to stem (SAFE-EXPLAINQA spirit for options).

Kill: Stem-first LLM fill; item-count NS (SAFE-COVER).

Survive: Distractor analysis dashboards for bank ops; rewrite loops for dead options.

CXXIV.6 Diagnostic questions — wrong pick reveals the theory

FACT (Eedi / NeurIPS diagnostic definition): Wang et al. (2020, arXiv:2007.12061; NeurIPS 2020 Education Challenge guide) — a **diagnostic question** is a four-answer MC with exactly one correct option where **each incorrect answer is chosen to highlight a common misconception**; student choice reveals nature of the misconception for remediation; writing good diagnostic questions is hard even for experienced teachers; crowd-sourced banks show quality variation; platform data (Eedi) supports large-scale response analysis.

FACT (product practice, secondary): Eedi’s public materials describe tagging each distractor to a misconception and using choice data for teacher planning — aligned with the diagnostic definition above (treat as industry corroboration, not peer-reviewed causal proof).

Reuse (SAFE-MISCON / SAFE-TAXON): Soft-wrong springboard + typed mark; slip ≠ bug; categorical routing beats correctness-only (Arieli-Attali & Liu via CXXII).

Applied (HYPOTHESIS): MindCraft soft-wrong MC contract: **chosen distractor** → **mark type bug/misconception (or slip trap if so tagged)** → **coach contrast** → **route**. Untyped options may fire only as “untyped bug” + SE, never invent disease labels. Blank / abandon remain non-option events (SAFE-RETRIEVE).

Kill: Unmapped MC as “diagnostic.”

Survive: Option→mis_ graph edges; tutor briefs show *which theory* Maya enacted.

CXXIV.7 Ordered options — progression grain without Partial Credit Score™

FACT (OMC): Briggs, Alonzo, Schwab, & Wilson (2006, *Educational Assessment*, 11(1), 33–63, doi:10.1207/s15326977ea1101_2) — Ordered Multiple-Choice links each option to a **developmental level** in a learning progression; aims for diagnostic utility of open responses with MC efficiency; interpretation targets *why* the less-correct choice was selected, not only right/wrong.

Applied (HYPOTHESIS): For some MindCraft concepts with clear progressions (e.g., equation-solving joins), OMC-style ordered distractors can inform fade stage / scaffold grain (SAFE-FADE / SAFE-EXPTIME) **internally**. Do not surface “Level 2 of 4” as student Partial Credit / Identity Score™ (SAFE-TAXON kill).

Kill: Ordered options ≡ visible progress ladder NS.

Survive: Private progression tags → scaffold dose; student sees destaked next move.

Wound: OMC validity depends on a defended learning progression — do not fake order on ACT distractors that are *nominal* misconception siblings (SAFE-TAXON categorical preference).

CXXIV.8 Product surface — SAFE-DISTRACTOR claim contract

Surface	Ship	Banned substitute
Authoring	Error-family → options → key verify	Stem → random wrongs
Option metadata	misconception_id / slip / FormatId trap / untyped-bug	Bare “incorrect”

Soft-wrong	Destake + type from <i>chosen</i> option + CTA	Red X only; Miss Streak
Option count	3 default; 4 only if fourth functions	Always-four costume
Bank QA	Functioning analysis; dead-option rewrite	Ship-and-forget
GENQ	Distractor seal in verify gate	Fluency-only pass
Metrics	type→route hit; functioning %; retry_120s; solo_transfer_pass	Distractor Score™ / option-count
Parent/tutor	“Picked equal-sign join trap” grain	Live option stalk / % wrong

Competitive foil: Khan/Duo = binary or XP. Exam prep = four-option costume. ChatGPT = fluent unkeyed options. Eedi-class = diagnostic distractors. MindCraft = **diagnostic distractors** + **soft-wrong climate** + **Map route** + **FEI prove** — not Eedi clone, not option theater.

CXXIV.9 Doctrine — SAFE-DISTRACTOR (provisional)

1. **Misconception first** — write options from error families; stems serve diagnosis, not the reverse.
2. **Every wrong option has intent** — mis_ / slip / FormatId trap / explicit untyped-bug — never naked incorrect.
3. **Function over fashion** — retire <5% / positively discriminating dead options; prefer ~3 total options when that is the functioning set (Tarrant; Rodriguez).
4. **Soft-wrong wraps the pick** — climate destakes; taxonomy + route carry semantics (SAFE-TAXON).
5. **No Distractor Score™ / option-count NS** — bank health = seal rate + functioning % + route/transfer.
6. **GENQ equals craft** — generated distractors need the same seals as human ones (SAFE-GENQ); fluency≠keyed diagnosis.
7. **Dual-rail honesty** — prove-rail may match ACT four-option *format*; learn/diagnostic rails optimize functioning diagnostic options — label which rail.
8. **OMC grain is private** — ordered levels may drive scaffolds; never student Partial Credit ladders.
9. **Copy:** “Wrong answers that teach — each option is a theory we can route.” Never “more choices = smarter test.”

Confidence: High — Haladyna et al. guidelines; Tarrant functioning scarcity; Rodriguez 3-option meta-analysis; Gierl et al. distractor review; Wang et al. diagnostic-question definition; Briggs et al. OMC. Medium — MindCraft auto-seal accuracy and live functioning rates (needs DISTRACTOR-*). High — beauty traps / option-count theater / unmapped “diagnostic” MC as commercially toxic under FEI + Map honesty.

CXXIV.10 Experiments

ID	Question	Design	Primary
----	----------	--------	---------

DISTRACTOR-1	Misconception-mapped 3-option vs unmapped 4-option	A/B within concept	type→route hit; retry_120s; solo_transfer_pass
DISTRACTOR-2	Retire dead options vs keep four-option costume	Offline item analysis + A/B	Functioning %; discrimination; abandon
DISTRACTOR-3	GENQ with distractor-seal gate vs fluency-only	Bank ops	Seal rate; false mis_rate; drop rate
DISTRACTOR-4	Soft-wrong copy naming chosen theory vs generic “not quite”	A/B	retry_120s; uptake (SAFE-FBLIT)
DISTRACTOR-5	Parent/tutor CBC: grain-from-option vs % wrong / option stalk	CBC	WTP; trust; shame (SAFE-WTP/SOFTMSG)
DISTRACTOR-QUAL	10 students + 5 tutors: which wrong options felt “tricky fair” vs “gotcha”?	Qual	DISTRACTOR codebook

Falsifier: Unmapped 4-option matches mapped 3-option on route + transfer → keep private seals; still ban beauty traps and Distractor Score™.

Falsifier: Always-four raises exam comfort without hurting transfer → segregate prove-rail format; do not redefine learn-rail craft.

Falsifier: Fluency-only GENQ equals sealed on false-mis rate → still require key+seal (SAFE-GENQ / LABMETA).

Pre-register: DISTRACTOR-* before “diagnostic MC / AI distractors” ads.

Family note: DISTRACTOR-* densifies TAXON/MISCON/GENQ; do not collapse into Option Count OKRs or Soft-Wrong Score™.

CXXIV.11 So what for MindCraft commercially

- **Copy:** “Wrong answers that teach — each option is a theory we can route.” Lead with diagnostic distractors + soft-wrong, not “thousands of MCQs” or four-choice costume.
- **Product:** Require option→mis_ (or typed trap) metadata; default to functioning 3-option diagnostic items; dual-rail ACT format only where labeled; wire chosen option into SAFE-TAXON mark + Map route.
- **Positioning:** Against binary apps, option-count theater, and ChatGPT unkeyed fluency; beside Eedi-class diagnosis, ahead on identity climate + FEI prove + ontology route.
- **Metric:** distractor seal %, functioning %, type→route hit, retry_120s, solo_transfer_pass — demote Distractor Score™ / item-count.
- **Kill list:** Beauty traps; always-four cosplay; unmapped “diagnostic”; Distractor Score™; GENQ without seals; ACT guarantees from distractor UI.

- **Growth:** Trust packets show distractor→misconception map + verify loop — procurement asset vs black-box banks.
 - **Vision:** Thirty-year identity company treats every wrong pick as **structured evidence of a theory in use** — so Maya learns to revise theories, not to fear multiple choice.
-

References (verified)

- Briggs, D. C., Alonzo, A. C., Schwab, C., & Wilson, M. (2006). Diagnostic assessment with ordered multiple-choice items. *Educational Assessment*, 11(1), 33–63. https://doi.org/10.1207/s15326977ea1101_2
 - Gierl, M. J., Bulut, O., Guo, Q., & Zhang, X. (2017). Developing, analyzing, and using distractors for multiple-choice tests in education: A comprehensive review. *Review of Educational Research*, 87(6), 1082–1116. <https://doi.org/10.3102/0034654317726529>
 - Haladyna, T. M., Downing, S. M., & Rodriguez, M. C. (2002). A review of multiple-choice item-writing guidelines for classroom assessment. *Applied Measurement in Education*, 15(3), 309–334. https://doi.org/10.1207/S15324818AME1503_5
 - Rodriguez, M. C. (2005). Three options are optimal for multiple-choice items: A meta-analysis of 80 years of research. *Educational Measurement: Issues and Practice*, 24(2), 3–13. <https://doi.org/10.1111/j.1745-3992.2005.00006.x>
 - Tarrant, M., Ware, J., & Mohammed, A. M. (2009). An assessment of functioning and non-functioning distractors in multiple-choice questions: A descriptive analysis. *BMC Medical Education*, 9, 40. <https://doi.org/10.1186/1472-6920-9-40>
 - Wang, Z., Lamb, A., Saveliev, E., Cameron, P., Zaykov, Y., Hernández-Lobato, J. M., Turner, R. E., Baraniuk, R. G., Barton, C., Peyton Jones, S., Woodhead, S., & Zhang, C. (2020). Instructions and guide for diagnostic questions: The NeurIPS 2020 Education Challenge. *arXiv:2007.12061*. <https://arxiv.org/abs/2007.12061>
-

Part CXXV — Confidence-Tagged Misses in Gap-Scan Recovery

Chapter status: Living evidence + Diagnostic / gap-scan recovery brief — Researcher tick 2026-08-11 (UTC hour 12 ≡ Red Team slot, but ch125 never written → prefer Researcher per rotation; researcher count since synthesizer v1.15 = 8 → Researcher)

Primary question: How should MindCraft turn **hard / kinda / easy** (and item-level surety) × miss type into *recovery* after gap-scan — without Identity Score™ theater, equal-weight red Xs, or treating confidence tags as mastery fireworks?

Owners: Product (gap-scan / Diagnostic C4 / soft-wrong / Practice launch) · Engine (/seed-assessment + outcome severity + confidence_miss_tier) · HITL · Brand · Parent GTM · Red Team

Commercial job: Ship **SAFE-CONFMISS** densifying SAFE-TAXON × SAFE-CALIB × SAFE-MONITOR × SAFE-COLD: **confidence × correctness class drives recovery dose and graph weight**; destake climate stays; no Identity Score™ / Confidence Personality™ from onboarding.

Builds on: Parts L (SAFE-CALIB), XCVIII (SAFE-MONITOR), LXXXI (SAFE-COLD), XLIX (SAFE-MISCON), CXXII (SAFE-TAXON), CXXIV (SAFE-DISTRACTOR), LXXIX (SAFE-FBTIME), CXVIII (SAFE-SOFTMSG). Product seams: gap-scan hard/kinda/easy → /seed-assessment, C4 probes → /record-outcomes, level gating

(bridgePractice), PawHub weakness launch, soft-wrong recovery.

CXXV.1 Why this chapter exists

SAFE-CALIB already bans “raise confidence” as North Star. SAFE-MONITOR already requires judgment→control wiring. SAFE-TAXON already requires typed marks. Gap-scan still tempts a third costume: **Confidence Identity Score™** — mash hard/kinda/easy into a splash (“You’re a Confident Explorer!”), or recover every miss with the same pep + same drill regardless of whether Maya was *sure* and wrong or *guessing* and wrong.

MindCraft already maps confidence to synthetic seed events and level recommendations. That is the right *family* — if recovery after gap-scan (and after probe misses) is **class-conditioned**: high-conf miss → attended springboard / hypercorrection path; low-conf miss → scaffold / lack path; low-conf correct → earned challenge unlock; and none of those classes mint an Identity Score™.

FOUNDER BELIEF under audit: Gap-scan’s commercial value is not a confidence vibe. It is a **recovery queue**: which theories to confront, which absences to teach, which “easy” ratings to distrust until FEI proves them — with honest uncertainty until delayed evidence arrives (SAFE-COLD / SAFE-FORGET).

Claims we refuse as doctrine:

1. Identity Score™ / Confidence Personality™ / “math person type” from hard/kinda/easy.
 2. All gap-scan misses recover the same (one red-X → one generic review deck).
 3. hard/kinda/easy alone ≡ calibrated mastery or ACT readiness (SAFE-COLD).
 4. Confidence slider without mark type / route (Foster CA-alone null — Part L).
 5. Inflate confidence as recovery therapy; Belief Score™ as NS.
 6. Immediate surety ≡ durable JOLs for future recall (SAFE-MONITOR).
 7. Hide-correctness forever = never teach; recovery without attended feedback.
 8. ACT / grade guarantees from confidence-tagged recovery packaging.
-

CXXV.2 Constructs (tag × type × recovery — not a score)

Construct	Research meaning	MindCraft analogue	Failure mode
Concept ease rating	Coarse prospective competence (EOL-ish)	Gap-scan hard / kinda / easy	Treat as mastery belt
Item confidence	Certainty this answer is correct (RCJ)	Soft-wrong / probe “how sure?”	Global vibe only
CRI logic	High-certainty wrong → misconception; low → lack/guess	Confidence × outcome classes	One fused Confidence %
Hypercorrection path	High-conf errors get extra attentional capture from feedback	Attended soft-wrong + SE conflict	Shame the sure miss

Recovery	Post-diagnostic control: what to practice next, how hard	Mission launch, fade grain, Map route	Generic “review all reds”
SAFE-CONFMISS	Confidence-tagged miss recovery law	This chapter	Identity Score™ theater

Operational definition (HYPOTHESIS): A gap-scan (or early Practice) miss is *SAFE-CONFMISS complete* when the product (a) records **confidence grain** (concept ease and/or item surety) with the outcome, (b) assigns a **SAFE-TAXON mark type** when available (slip / bug / lack / blank / assist), (c) chooses a **recovery dose** from the confidence×type class (not one universal review), (d) updates graph severity with type- and confidence-conditioned weights, (e) never surfaces Identity Score™ / Confidence Personality™, and (f) proves value with confidence_miss_tier hit → route, retry_120s, weakness stability, and later solo_transfer_pass — not rising self-rated confidence.

CXXV.3 CRI: certainty separates misconception from don’t-know

FACT (Certainty of Response Index): Hasan, Bagayoko, & Kelley (1999, *Physics Education*, 34(5), 294–299, doi:10.1088/0031-9120/34/5/304) — pair MC answers with a Certainty of Response Index; **low certainty** (regardless of correctness) indicates guessing / lack of knowledge; **high certainty + incorrect** indicates a **misconception** (firm belief in a flawed idea), not mere absence.

Applied (HYPOTHESIS): MindCraft’s hard/kinda/easy and item surety are CRI-class signals. Recovery must **fork**:

- **High-conf × wrong** → treat as bug/misconception candidate → SAFE-DISTRACTOR option theory if MC → springboard + Map route (SAFE-MISCON / SAFE-TAXON).
- **Low-conf × wrong** → treat as lack/guess → lighter confrontation, stronger scaffold / fade protect (SAFE-FADE / SAFE-FADEX), teach-then-retry.
- **Low-conf × correct** → underconfidence class → competence evidence + challenge unlock, not “believe harder” posters (Part L).
- **High-conf × correct** → provisional strength signal only until delayed FEI (SAFE-COLD / SAFE-FORGET) — never day-one fireworks.

Kill: One recovery deck for all reds.

Survive: Class matrix → mission / coach / graph weight.

Wound: CRI alone can misclassify “lucky correct with wrong reasons” and “knows but anxious.” Prefer distractor-typed MC + optional brief why when stakes rise (SAFE-DISTRACTOR; multi-tier diagnostics as future densifier — not day-one UX bomb).

CXXV.4 Hypercorrection — high-conf misses are recovery gold, not shame

FACT (hypercorrection): Butterfield & Metcalfe (2001, *Journal of Experimental Psychology: Learning, Memory, and Cognition*, 27(6), 1491–1494, doi:10.1037/0278-7393.27.6.1491) — contrary to interference intuition, **high-confidence errors** were *more* likely to be corrected on retest after feedback than low-confidence errors.

FACT (attention account): Butterfield & Metcalfe (2006, *Metacognition and Learning*, 1, 69–84, doi:10.1007/s11409-006-6894-z) — corrective feedback after high-confidence errors captures more attention (tone-detection impairment during feedback); supports surprise/attention mechanism for hypercorrection.

FACT (partial knowledge / “knew it all along”): Metcalfe & Finn (2011, *Journal of Experimental Psychology: Learning, Memory, and Cognition*, 37(2), 437–448, doi:10.1037/a0021962) — adults claim “knew it all along” more for high-conf errors; second-guess and cue tasks suggest preferential partial knowledge contributes to easy correction — not mere hindsight theater.

FACT (children): Metcalfe & Finn (2012, *Learning and Instruction*, 22(4), 253–261, doi:10.1016/j.learninstruc.2011.10.004) — Grade 3–6 children hypercorrect high-confidence errors; educationally relevant beyond adult lab samples (HS Maya and younger caseloads).

FACT (authentic math diagnostic): Foster, Woodhead, Barton, & Clark-Wilson (2022, *Educational Studies in Mathematics*, 109, 491–521, doi:10.1007/s10649-021-10084-7) — large Eedi dataset (N responses $\approx$ 219,826; N students = 7,302); confidence correlates with facility; **first evidence for hypercorrection in authentic school mathematics** — higher confidence on an earlier wrong predicted higher odds of later correct on a related quiz after controlling for facility.

Applied (FOUNDER BELIEF → testable): Gap-scan recovery should **prioritize attended feedback** on high-conf misses (and high-ease concepts that later fail probes) — destaked soft-wrong climate (SAFE-ERRCLIMATE), not public humiliation. Low-conf misses need teaching dose first; dumping the same “surprise confrontation” on guesses wastes attention and can raise anxiety (SAFE-FADEX / SAFE-DD).

Kill: Shame-the-sure / “how could you be so confident?” brand voice.

Survive: High-conf miss → springboard priority in recovery queue; measure hypercorrection with delayed isomorphic retry, not emoji uplift.

CXXV.5 Monitoring accuracy is for control — overconfidence ends study too early

FACT (AIM / retention): Dunlosky & Rawson (2012, *Learning and Instruction*, 22(4), 271–280, doi:10.1016/j.learninstruc.2011.08.003) — absolute monitoring accuracy linked to longer-term retention; **overconfidence** can produce premature termination of practice (judging incomplete responses as “correct enough”), undermining delayed recall; improving judgment standards can help.

Reuse (Part L): Foster (2021/2022) — CA alone $\approx$ null on attainment; value is wiring, not the widget.

Applied (HYPOTHESIS): After gap-scan, “easy” ratings that collide with probe misses (or later Practice fails) are **overconfidence recovery targets**: do not honor the easy self-report as mastery; schedule FEI prove + spaced revisit (SAFE-SCHED). “Hard” ratings that later show high accuracy with low surety are **underconfidence unlocks**. Recovery UX that celebrates rising confidence without FEI gates re-creates Dunlosky’s premature-termination problem in product form.

Kill: Confidence inflation as recovery success metric.

Survive: Judgment→control: restudy / retry / challenge unlock from class, not Belief Score™.

CXXV.6 Product surface — gap-scan → recovery claim contract

Surface	Ship	Banned substitute
Gap-scan ratings	hard/kinda/easy → seed events + level gate	Identity splash from rating mix
Probe / Practice miss	Record surety + mark type + outcome	Red X only
Recovery queue	Prioritize high-conf bugs; scaffold low-conf lacks	Shuffle all misses equally
Soft-wrong	Destake + class-specific next move	Miss Streak / pep-only
Graph	Confidence-conditioned severity (guess ≠ sure misconception)	One delta per wrong
Parent/tutor	“Sure miss on join trap — springboard next”	Confidence Personality™ / shame ranks
Metrics	confidence_miss_tier → route; retry_120s; weakness stability; solo_transfer_pass	Identity Score™ / Confidence Minutes

Competitive foil: Duo = streak / XP after any miss. Khan = binary mastery bars. ChatGPT = fluent re-explain without confidence class. Eedi-class = confidence + diagnostic options. MindCraft = **CRI-class recovery + soft-wrong climate + Map route + FEI prove** — not a Confidence Quiz brand.

CXXV.7 Doctrine — SAFE-CONFMISS (provisional)

1. **Tag every diagnostic miss** — concept ease and/or item surety stored with outcome; no orphan red Xs.
2. **Class before costume** — high-conf wrong ≠ low-conf wrong ≠ blank; fork recovery dose (CRI logic).
3. **Hypercorrection is a gift** — attended, destaked feedback for sure misses; measure delayed correction, not shame.
4. **Taxonomy still wins** — confidence tags *modulate* SAFE-TAXON marks; they do not replace mis_ / slip / lack.
5. **No Identity Score™** — never mash hard/kinda/easy into personality / math-person type / Belief Score™.
6. **Easy ≠ mastered** — honor collision with probes/FEI over self-report fireworks (SAFE-COLD).
7. **CA is fuel, not attainment engine** — widget without route is theater (Foster null).
8. **Copy:** “We recover from *how* you missed — sure misconception vs still learning — not from a Confidence Score.” Never “raise your confidence in 10 minutes.”

Confidence: High — Hasan et al. CRI; Butterfield & Metcalfe hypercorrection + attention; Metcalfe & Finn partial-knowledge line; Dunlosky & Rawson overconfidence→underachievement; Foster et al. Eedi math hypercorrection. Medium — MindCraft-specific mapping of hard/kinda/easy grain to CRI cut-points and severity tables (needs CONFMISS-*). High — Identity Score™ / equal-recovery / inflate-confidence theater as commercially toxic under FEI + Map honesty.

CXXV.8 Experiments

ID	Question	Design	Primary
CONFMISS-1	Class-conditioned recovery vs equal-weight review deck	A/B post gap-scan	retry_120s; delayed isomorphic correct; solo_transfer_pass
CONFMISS-2	High-conf miss springboard priority vs random miss order	A/B	Hypercorrection rate; abandon; anxiety items
CONFMISS-3	easy-vs-probe collision restudy vs honor-easy path	A/B	Weakness stability; false-mastery rate
CONFMISS-4	Parent/tutor CBC: class grain vs Confidence Personality™ / % sure	CBC	WTP; trust; shame (SAFE-WTP/SOFTMSG)
CONFMISS-5	Severity table: conf×type vs single red-X delta	Offline sim + A/B	Path quality; over-update rate
CONFMISS-QUAL	10 Mayas: did “sure wrong” feel fair/useful vs shaming?	Qual	CONFMISS codebook

Falsifier: Equal-weight recovery matches class-conditioned on transfer → keep private tags for analytics; still ban Identity Score™ and inflate-confidence ads.

Falsifier: Springboard-on-sure-miss raises anxiety without hypercorrection gain → tighten SAFE-FADEX protect; do not abandon CRI fork.

Falsifier: Honor-easy path equals collision-restudy on later FEI → still refuse day-one mastery fireworks (SAFE-COLD / LABMETA).

Pre-register: CONFMISS-* before “confidence-powered gap scan / Identity Score” ads.

Family note: CONFMISS-* densifies CAL/MONITOR/TAXON/COLD; do not collapse into Confidence Score™ OKRs.

CXXV.9 So what for MindCraft commercially

- **Copy:** “We recover from *how* you missed — sure misconception vs still learning.” Lead with class-conditioned recovery, not “build confidence” or personality quizzes.
- **Product:** Persist confidence with outcomes; fork recovery queues; wire high-conf misses to springboard + Map; distrust easy-without-FEI; never Identity Score™ splash.

- **Positioning:** Against streak apps and binary bars; beside Eedi-class confidence diagnostics; ahead on destaked climate + ontology route + solo-transfer prove.
 - **Metric:** confidence_miss_tier → route hit, collision restudy rate, retry_120s, solo_transfer_pass — demote Confidence Minutes / Belief Score™.
 - **Kill list:** Identity Score™; equal-recovery decks; easy=mastered; inflate-confidence recovery; ACT guarantees from confidence packaging.
 - **Growth:** Trust packet shows CRI-class recovery + hypercorrection instrumentation — procurement asset vs black-box “adaptive confidence.”
 - **Vision:** Thirty-year identity company treats confident error as **high-value diagnostic fuel** — so Maya learns to revise sure theories without learning that surety itself is shameful.
-

References (verified)

- Butterfield, B., & Metcalfe, J. (2001). Errors committed with high confidence are hypercorrected. *Journal of Experimental Psychology: Learning, Memory, and Cognition*, 27(6), 1491–1494. <https://doi.org/10.1037//0278-7393.27.6.1491>
 - Butterfield, B., & Metcalfe, J. (2006). The correction of errors committed with high confidence. *Metacognition and Learning*, 1, 69–84. <https://doi.org/10.1007/s11409-006-6894-z>
 - Dunlosky, J., & Rawson, K. A. (2012). Overconfidence produces underachievement: Inaccurate self evaluations undermine students’ learning and retention. *Learning and Instruction*, 22(4), 271–280. <https://doi.org/10.1016/j.learninstruc.2011.08.003>
 - Foster, C., Woodhead, S., Barton, C., & Clark-Wilson, A. (2022). School students’ confidence when answering diagnostic questions online. *Educational Studies in Mathematics*, 109, 491–521. <https://doi.org/10.1007/s10649-021-10084-7>
 - Hasan, S., Bagayoko, D., & Kelley, E. L. (1999). Misconceptions and the Certainty of Response Index (CRI). *Physics Education*, 34(5), 294–299. <https://doi.org/10.1088/0031-9120/34/5/304>
 - Metcalfe, J., & Finn, B. (2011). People’s hypercorrection of high-confidence errors: Did they know it all along? *Journal of Experimental Psychology: Learning, Memory, and Cognition*, 37(2), 437–448. <https://doi.org/10.1037/a0021962>
 - Metcalfe, J., & Finn, B. (2012). Hypercorrection of high confidence errors in children. *Learning and Instruction*, 22(4), 253–261. <https://doi.org/10.1016/j.learninstruc.2011.10.004>
-

Part CXXVI — Tutor Micro-Debrief Fidelity Under Time Caps

Chapter status: Living evidence + HITL close-ops brief — Researcher tick 2026-08-11 (UTC hour 18 ≡ Red Team slot, but ch126 never written → prefer Researcher per rotation; researcher count since synthesizer v1.16 = 0 → Researcher)

Primary question: When session clocks are short (crunch weeks, late joins, load-shed minutes), how does MindCraft keep **after-action fidelity** — student-led expected/actual/sustain/improve with objective traces — without skipping the close, lecturing the moral, military cosplay, or inventing a Microdebrief Score™?

Owners: HITL ops / tutor console · Product (session spine / Notes close) · Coach UX · Exam calendar (SAFE-LOADSHED) · Brand · Red Team

Commercial job: Ship **SAFE-MICRODEB** densifying SAFE-AAR × SAFE-LOADSHED × SAFE-HITL × SAFE-TALK × SAFE-FBLIT: a **capped, structured micro-debrief** (~60–90s practice / ≤3 min live) that protects learning closure when time is scarce — never “no time for reflection,” never pour-close, never +25% ACT ads.

Builds on: Parts LIV (SAFE-AAR), CXXIII (SAFE-LOADSHED), LXVIII (SAFE-HITL), LXXXIV (SAFE-TALK), CXII (SAFE-FBLIT), LI (SAFE-DP), XC (SAFE-INSTRUMENT), XCV (SAFE-APPRENTICE). Product seams: session timer, Map grain slice, soft-wrong codes, attempt traces, Notes close prompts, acuity queue under crunch.

CXXVI.1 Why this chapter exists

SAFE-AAR already says: brief → attempt → structured debrief with objective traces; student speaks first; no military costume; no borrowed meta *d* as ACT marketing. SAFE-LOADSHED already says: under exam surge, shed low-acuity demand so FEI minutes survive. What remains commercially underspecified is the **join under a hard clock**: tutors “save time” by skipping the close, delivering a 40-second lecture (“here’s what you should have done”), or vibes-checking (“how did that feel?”) with zero traces — then ops still claim they “do AARs.”

Competitors resolve the clock wrongly: (1) **skip**; (2) **pour-close** (tutor monologue as fake debrief); (3) **therapy costume**; (4) **bootcamp** roast; (5) **AI wrap essay** that blocks self-analysis. MindCraft’s wedge is a **time-capped micro-debrief** — SHARP-class prompts + Army four questions + objective media — engineered to finish in ~90 seconds without losing the learning function.

FOUNDER BELIEF under audit: Identity transformation needs *reviewed attempts*. Under crunch, the company either designs a faithful short form or it silently abandons review. Skipping the close is not “efficiency” — it is a fidelity breach (SAFE-BURNWARN / SAFE-LOADSHED).

Claims we refuse as doctrine:

1. “No time for debrief” as acceptable default under crunch.
 2. Tutor/AI lecture that *calls itself* a micro-debrief.
 3. Microdebrief Score™ / Debrief Minutes / Reflection Streak as North Star.
 4. Undifferentiated “science says +25%” / military-cosplay AARs as ACT marketing (SAFE-AAR kill).
 5. Mandatory long therapy-style close when the clock is the constraint.
 6. Vibes-only close without objective traces.
 7. Soft standards / skip-challenge justified by “we at least talked.”
 8. ACT / score guarantees from micro-debrief packaging.
-

CXXVI.2 Constructs (short ≠ empty; capped ≠ skipped)

Construct	Meaning	MindCraft analogue	Failure mode
Full AAR	Facilitated expected/actual/sustain/improve with media	3–5 min tutor close (SAFE-AAR)	Lecture; roast

Micro-debrief	Same function, hard time box, fewer prompts	60–90s practice; ≤3 min live	Skip or pour
Time cap	Explicit budget that ends the close	Session timer / “close now” cue	Cap as excuse to skip
Objective media	Trace of what happened (not memory alone)	Soft-wrong code, attempt timing, Map grain	Vibes-only
Student-first turn	Learner names frame before coach wrap	Prompt: “What were you trying?”	Tutor moral first
SHARP-class scaffold	Short structured prompts for time-limited settings	Five-beat micro template	OSAD-full cosplay mid-crunch
SAFE-MICRODEB	Capped AAR fidelity law	This chapter	Debrief Minutes NS

Operational definition (HYPOTHESIS): A close is *SAFE-MICRODEB complete* when (a) a **named time cap** is visible (≤90s solo / ≤180s live under load-shed; longer only when acuity warrants and calendar allows), (b) the close runs **four beats minimum**: expected edge → actual (with ≥1 objective trace) → one sustain → one improve-how / next attempt, (c) the **student produces ≥1 self-analytic statement before** tutor/AI wrap, (d) output is a next-attempt commitment logged to Notes/Map — not a grade sermon, (e) under SAFE-LOADSHED the micro-close is on the **protect** list (not an elective to shed), and (f) success = method-change codes / `retry_120s` / `solo_transfer_pass` / brief-open — never Debrief Minutes or Microdebrief Score™.

CXXVI.3 Debriefs work — transfer the *method*, not the marketing *d*

FACT: Tannenbaum & Cerasoli (2013, *Human Factors*, 55(1), 231–245, doi:10.1177/0018720812448394) — meta-analysis of 46 samples ($N = 2,136$): debriefs / AARs improve effectiveness vs control by about **25% on average** ($d \approx 0.67$); effects similar for teams/individuals and simulated/real settings; **alignment** (participants ↔ focus ↔ measurement) bolsters effects; facilitated/structured debriefs trend stronger.

FACT: Keiser & Arthur (2021, *Journal of Applied Psychology*, 106(7), 1007–1032; advance online 2020, doi:10.1037/apl0000821) — bare-bones meta-analysis of 61 studies (107 *ds*; 915 teams and 3,499 individuals): overall AAR effect $d \approx 0.79$ across multiple training criteria. Two characteristics consistently help: **(a) alignment** to individual vs team, **(b) objective performance review media**. Structure helps especially in military contexts; facilitation interacts with media/alignment.

Applied (HYPOTHESIS): Under a time cap, protect the two consistent drivers — **alignment** (solo student ↔ solo attempt metrics) and **objective media** (soft-wrong / timing / Map grain) — even if facilitator flourish and long reaction phases are cut. Do **not** protect length-for-length’s sake.

Kill: “+25% from debriefs” ACT ads; length≡quality.

Survive: Aligned + traced micro-closes beat skip and vibes-only.

Wound: Meta samples ≠ ACT-anxious Maya — MICRODEB- *must pre-register FEI instruments, not borrow d**.

CXXVI.4 Objective review media beat memory theater — even when short

FACT: Villado & Arthur (2013, *Journal of Applied Psychology*, 98(2), 292–307, doi:10.1037/a0031510) — 188 individuals in 47 four-person teams; AAR teams outperformed non-AAR teams on team performance, efficacy, communication openness, and cohesion; **objective vs subjective** review media did not significantly differ across five training outcomes in that design — but *having an AAR beat no AAR*.

Applied (HYPOTHESIS): Product should still **prefer** objective traces (Keiser & Arthur consistent moderator) because tutoring memory is biased and tutors under crunch misremember. When traces are missing, micro-debrief still beats skip — but ship attempt cards so “what actually happened” is not a reconstruction fight.

Kill: “How did that feel?” as the only close.

Survive: Soft-wrong + timing + named edge as default media for the 90s box.

CXXVI.5 Time-limited settings need short scaffolds — SHARP as existence proof

FACT (OSAD quality dimensions): Arora, Ahmed, Paige, Nestel, Runnacles, Hull, et al. (2012, *Annals of Surgery*, 256(6), 982–988, doi:10.1097/SLA.0b013e3182610c91) — developed and validated the Objective Structured Assessment of Debriefing (OSAD) from literature + 33 end-user interviews across three continents; eight dimensions (approach, learning environment, engagement, reaction, reflection, analysis, diagnosis, application); feasible/reliable (inter-rater ICC $\approx$ 0.88; test-retest ICC $\approx$ 0.90).

FACT (SHARP for the OR clock): Ahmed, Arora, Russ, Darzi, Vincent, & Sevdalis (2013, *Annals of Surgery*, 258(6), 958–963, doi:10.1097/SLA.0b013e31828c88fc) — prospective pre/post of 100 OR cases (22 trainers, 30 residents): introduced **SHARP** (Set learning objectives; How did it go; Address concerns; Review learning points; Plan ahead) as a **time-feasible** debrief scaffold when full OSAD-depth closes are impractical. Observer OSAD medians rose from 19 (range 8–31) to 33 (range 26–40), $P < 0.001$; users rated usefulness/feasibility high. Companion handbook positions SHARP as the “absolute basics” when there is not enough time for a detailed OSAD-guided debrief (~5 minutes class of tool).

Applied (FOUNDER BELIEF → testable): MindCraft’s micro-debrief is the **tutoring analogue of SHARP**, not a watered-down pep talk:

1. **Set** — restate the briefed edge (or name it if brief was skipped — then fix brief next).
2. **How** — student + trace: expected vs actual.
3. **Address** — one concern / soft-wrong frame (advocacy–inquiry, not dump).
4. **Review** — one sustain.
5. **Plan** — one improve-how / next attempt card.

Under SAFE-LOADSHED, prefer this five-beat micro over canceling reflection so a homework dump can start.

Kill: “Full AAR or nothing” purity that becomes skip under crunch; OSAD-full cosplay as session tax.

Survive: Named short scaffold with fidelity checklist; QA samples micro-completeness, not minutes.

Wound: Surgical OR $\neq$ math tutoring — SHARP is method transfer (structure under clock), not brand (“OR-grade tutoring”). MICRODEB-* measures FEI uptake, not OSAD score cosplay.

CXXVI.6 Good judgment in 90 seconds — advocacy–inquiry, not praise fog or roast

FACT: Rudolph, Simon, Dufresne, & Raemer (2006, *Simulation in Healthcare*, 1(1), 49–55, doi:10.1097/01266021-200600110-00006) — “There’s no such thing as ‘nonjudgmental’ debriefing.” **Debriefing with good judgment** pairs advocacy (state observation + evaluative judgment) with inquiry (ask the learner’s frame). Mistakes are puzzles; stance is curiosity without fake nonjudgment.

Applied (HYPOTHESIS): Micro-debrief scripts default to one advocacy–inquiry pair tied to the soft-wrong code (“I see you flipped the inequality when multiplying by -1 — what were you assuming?”) then stop. Under the cap, **one** deep frame beats three tip dumps (SAFE-FILM / SAFE-EXPLAIN). Praise-fog and roast both burn the clock without FEI.

Kill: “Great job!” fog close; public roast close; tip flood.

Survive: One judged observation + one student frame + next attempt.

CXXVI.7 Product surface — SAFE-MICRODEB claim contract

Clock context	Micro-debrief form	Protect	Banned substitute
Solo Practice end	60–90s: expected / actual+trace / sustain / improve	Student-first; next card	Skip to New Mission
Live tutor, normal	≤ 3 –5 min facilitated AAR (SAFE-AAR)	Objective media	Tutor lecture-as-debrief
Live tutor, LOADSHED crunch	≤ 90 –180s SHARP-class micro	On protect list	Cancel close for HW dump
Soft-wrong heavy	One advocacy–inquiry on typed mark	Frame before wrap	Answer dump
AI coach assist	Prompt beats only; student writes first	SAFE-EXPLAINQA gates	Monologue wrap first
Metrics	microdeb_complete; method-change; retry_120s; solo_transfer_pass	Co-gate with FEI	Microdebrief Score™; Debrief Minutes

Competitive foil: ChatGPT = wrap essay. Marketplace tutor = skip or pour. Bootcamp = roast. Duo = streak liturgy. MindCraft = **capped, traced, student-first close that survives the clock.**

CXXVI.8 Doctrine — SAFE-MICRODEB (provisional)

1. **Cap the form, don't cancel the function** — under time pressure, run micro-AAR; never “no time for debrief” as default (LOADSHED protect).
2. **Four beats minimum** — expected → actual+trace → sustain → improve-how / next attempt.
3. **Student speaks first** — ≥1 self-analytic statement before tutor/AI wrap (SAFE-AAR / SAFE-TALK).
4. **Objective media preferred** — soft-wrong, timing, Map grain; vibes-only is failure mode (Keiser & Arthur; Villado & Arthur).
5. **SHARP-class scaffold** — Set / How / Address / Review / Plan as the short template (Ahmed et al.) — method not OR cosplay.
6. **One advocacy-inquiry** — good judgment without lecture flood (Rudolph et al.).
7. **Alignment** — solo close on solo metrics; don't measure team vibes with XP (Tannenbaum & Cerasoli).
8. **Copy:** “Even in a short session, we close the loop on what you tried — in under two minutes.” Never “+25% from military debriefs” or Microdebrief Score™.

Confidence: High — debrief metas + SHARP pre/post + OSAD dimensions + Rudolph stance as method sources. Medium — MindCraft-specific 60–90s completeness thresholds and crunch protect-list compliance (needs MICRODEB-*). High — skip/pour/lecture-as-debrief / +25% ACT / Score™ theater as commercially toxic under FEI + SAFE-AAR.

CXXVI.9 Experiments

ID	Question	Design	Primary
MICRODEB-1	90s micro-close vs skip-to-next under equal session length	A/B Practice	retry_120s; method-change; solo_transfer_pass
MICRODEB-2	SHARP-class five-beat vs tutor free-form ≤3 min	HITL A/B	microdeb_complete; talk-ratio; student-first rate
MICRODEB-3	Objective-trace card vs vibes-only prompt under same cap	A/B	Frame quality; abandon; next-attempt adherence
MICRODEB-4	LOADSHED: protect micro-close vs shed-close-for-HW	Ops A/B crunch weeks	fidelity@crunch; pour↓; transfer_pass
MICRODEB-5	Parent/tutor CBC: “short close that sticks” vs “we stay late / more minutes”	CBC	WTP; trust (SAFE-WTP / LOADSHED)
MICRODEB-QUAL	10 tutors + 10 Mayas: did 90s feel respectful or rushed/shameful?	Qual	MICRODEB codebook

Falsifier: Skip equals micro-close on delayed transfer → still refuse pour-close theater; investigate whether brief quality, not close, was the bottleneck.

Falsifier: Free-form ≤3 min beats SHARP-class on FEI → keep student-first + traces; demote rigid five-beat if it becomes costume.

Falsifier: Trace cards raise anxiety without method-change → tighten SAFE-ERRCLIMATE / FADEX; do not return to vibes-only.

Pre-register: MICRODEB-* before “military-grade micro-debrief / +25%” ads.

Family note: MICRODEB-* densifies AAR/LOADSHED/HITL/TALK/FBLIT; do not collapse into Debrief Minutes OKRs.

CXXVI.10 So what for MindCraft commercially

- **Copy:** “Even when time is tight, we still close the loop — expected, actual, next move — in under two minutes.” Lead with fidelity under constraint, not bootcamp or late-night hustle.
 - **Product:** Ship timer-bound close UI; force student-first field; auto-attach soft-wrong/timing; block “end session” without microdeb_complete under normal and crunch modes (with explicit acuity override logged).
 - **Ops:** Put micro-debrief on the SAFE-LOADSHED **protect** list; QA sample micro-completeness and lecture-as-debrief rate, not Debrief Minutes.
 - **Positioning:** Against skip-culture marketplaces and monologue AI; beside high-reliability training craft without military cosplay.
 - **Metric:** microdeb_complete, student-first rate, method-change codes, retry_120s, solo_transfer_pass — demote Microdebrief Score™ / Reflection Streak.
 - **Kill list:** Skip-as-default; pour-close; +25% ACT from debriefs; Microdebrief Score™; therapy-length mandate under crunch; vibes-only.
 - **Growth:** Trust packet shows capped AAR fidelity instrumentation — procurement asset vs “our tutors care” vibes.
 - **Vision:** Thirty-year identity company treats the last ninety seconds as sacred learning infrastructure — so Maya learns that hard attempts always get a fair review, even when the world is rushing her.
-

References (verified)

- Ahmed, M., Arora, S., Russ, S., Darzi, A., Vincent, C., & Sevdalis, N. (2013). Operation debrief: A SHARP improvement in performance feedback in the operating room. *Annals of Surgery*, 258(6), 958–963. <https://doi.org/10.1097/SLA.0b013e31828c88fc>
- Arora, S., Ahmed, M., Paige, J., Nestel, D., Runnacles, J., Hull, L., ... Sevdalis, N. (2012). Objective Structured Assessment of Debriefing: Bringing science to the art of debriefing in surgery. *Annals of Surgery*, 256(6), 982–988. <https://doi.org/10.1097/SLA.0b013e3182610c91>
- Keiser, N. L., & Arthur, W., Jr. (2021). A meta-analysis of the effectiveness of the after-action review (or debrief) and factors that influence its effectiveness. *Journal of Applied Psychology*, 106(7), 1007–1032. <https://doi.org/10.1037/apl0000821>
- Rudolph, J. W., Simon, R., Dufresne, R. L., & Raemer, D. B. (2006). There’s no such thing as “nonjudgmental” debriefing: A theory and method for debriefing with good judgment. *Simulation in Healthcare*, 1(1), 49–55.

<https://doi.org/10.1097/01266021-200600110-00006>

- Tannenbaum, S. I., & Cerasoli, C. P. (2013). Do team and individual debriefs enhance performance? A meta-analysis. *Human Factors*, 55(1), 231–245. <https://doi.org/10.1177/0018720812448394>
- Villado, A. J., & Arthur, W., Jr. (2013). The comparative effect of subjective and objective after-action reviews on team performance on a complex task. *Journal of Applied Psychology*, 98(2), 292–307. <https://doi.org/10.1037/a0031510>

Part CXXVII — Parent Crunch-Week Message Tests (Availability vs Thinking Minutes)

Chapter status: Living evidence + GTM message-design brief — Researcher tick 2026-08-11 (UTC hour 21; hour%6≠0; researcher count since synthesizer v1.16 = 1 → Researcher)

Primary question: In the weeks before ACT/SAT/state tests, which **parent-facing frames** win honest purchase intent and trust — “**we protect thinking minutes / shed low-acuity demand**” vs “**24/7 availability / always-on tutor access**” — without fear-appeal panic, Likert-as-WTP, or Availability Score™ theater?

Owners: Growth / Positioning · Parent GTM · SAFE-LOADSHED ops · Product (parent dash / crunch emails) · Brand · Red Team

Commercial job: Ship **SAFE-CRUNCHMSG** densifying SAFE-LOADSHED × SAFE-WTP × SAFE-SOFTMSG × SAFE-PDASH × SAFE-EXAM: a **pre-registered CBC + message stack** that decides whether MindCraft can sell load-shed honesty as *care* — never infinite-availability brand, hours-booked crunch NS, or ACT-point guarantees from “we stayed open.”

Builds on: Parts CXXIII (SAFE-LOADSHED — especially LOADSHED-5), LIX (SAFE-WTP), CXVIII (SAFE-SOFTMSG), LXXXIII (SAFE-PDASH), LXIII (SAFE-EXAM), LXXVIII (SAFE-DURABLE), XLIV (SAFE-REWARD), XXXII (parent anxiety). Product seams: ACT date on profile, acuity queue copy, parent crunch emails, pricing pages, waitlist/wait transparency, dual-rail HW help.

CXXVII.1 Why this chapter exists

SAFE-LOADSHED already says: publish what pauses so FEI minutes survive; kill infinite-availability cosplay. SAFE-WTP already says: CBC over Likert WTP. LOADSHED-5 named the fork — “protect thinking minutes” vs “24/7 availability.” What remains underspecified is the **message + choice experiment contract**: which words open wallets without training parents to expect pour-mode human Chegg, and which words survive when sessions are deferred, capped, or AI-first.

Competitors resolve the fork wrongly: (1) **always-on brand**; (2) **fear-appeal crunch**; (3) **hours theater**; (4) **wellness fog**; (5) **abandonment silence** (shed without explaining why). MindCraft’s wedge is **expectation-aligned honesty**: name protected behaviors (thinking minutes, micro-debrief, bridge work) and paused electives, then prove parents will pay — or kill the doctrine.

FOUNDER BELIEF under audit: Parents will trade some raw access for **credible protection of learning quality** when the trade-off is visible — if copy is plain, the pause list is real, and FEI evidence remains visible. If 24/7 always

wins CBC *and* later churn when fidelity collapses, refuse pour-mode anyway and redesign packaging (AI-rail clarity, not fake human infinite).

Claims we refuse as doctrine:

1. 24/7 / infinite-availability as default parent value prop.
2. Hours-booked / “we took every call” as crunch-week North Star.
3. Likert “would you pay for thinking minutes?” as serious WTP.
4. Fear-appeal parent emails as default motivation (“or they fail”).
5. Availability Score™ / Crunch Message Score™ / Open-Late Minutes NS.
6. Shed-without-explanation (honesty that feels like ghosting).
7. Soft standards / answer-dump sold as “exam support.”
8. ACT / score guarantees from crunch messaging or late-night packaging.

CXXVII.2 Constructs (honest scarce care ≠ abandonment)

Construct	Meaning	MindCraft analogue	Failure mode
Availability frame	Care = reachability / volume	“Book anytime / tutors always on”	Pour + fidelity collapse
Thinking-minutes frame	Care = protected FEI quality	“We keep hard attempts + close loops”	Jargon; reads as rationing
Shed honesty	Name what pauses under surge	Published pause list + acuity queue	Silence ≡ abandonment
Expectation–perception gap	Promised vs delivered service	SERVQUAL gap logic on crunch emails	Overpromise → rage-churn
Fear appeal (parent)	Motivate via failure threat	“ACT is coming — grind or else”	Threat appraisal → control/nag
CBC crunch arm	Forced trade-off packages	LOADSHED-5 / CRUNCHMSG-*	Feature-rank surveys
SAFE-CRUNCHMSG	Message+CBC doctrine for crunch	This chapter	Availability Score™

Operational definition (HYPOTHESIS): A crunch-week parent communication is *SAFE-CRUNCHMSG complete* when (a) it uses **choice-based** trade-offs (CBC/DCE or message A/B with purchase/retention outcomes) rather than open WTP Likert, (b) attributes/arms map to **shippable** surfaces (acuity queue, pause list, AI-first HW rail, FEI parent report, capped live minutes), (c) **price + neither/outside option** appear in CBC, (d) copy **names protected behaviors and paused electives** in plain language, (e) primary outcomes separate **intent, comprehension** (“what did they stop doing and why?”), **perceived abandonment**, and **homework-conflict / controlling-nag risk**, and (f) success is trust + retained FEI — never Open-Late Minutes or Availability Score™.

CXXVII.3 Expectation gaps punish overpromise — availability ads set the wrong bar

FACT (gaps model): Parasuraman, Zeithaml, & Berry (1985, *Journal of Marketing*, 49(4), 41–50, doi:10.1177/002224298504900403) — service quality problems arise from gaps between consumer expectations and management perceptions, specs, delivery, and communications; external communications that raise expectations beyond delivery capacity widen the quality gap.

FACT (SERVQUAL): Parasuraman, Zeithaml, & Berry (1988, *Journal of Retailing*, 64(1), 12–40) — service quality is commonly assessed as the gap between expectations and perceived performance across reliability, responsiveness, assurance, empathy, and tangibles; **reliability** (doing what was promised) is structurally central.

FACT (disconfirmation → satisfaction): Oliver (1980, *Journal of Marketing Research*, 17(4), 460–469, doi:10.1177/002224378001700405) — satisfaction follows from expectation and expectancy disconfirmation; over-raised expectations that performance cannot meet predict dissatisfaction and weaker purchase intention cascades.

Applied (HYPOTHESIS): “24/7 tutors / always book tonight” is an **expectation-raising communication** (Gap 4 in the gaps model). Under SAFE-LOADSHED, delivery *must* constrain access. The commercial question is not whether to constrain — it is whether to **advertise constraint-as-care** or advertise infinity and then disconfirm. SAFE-CRUNCHMSG prefers under-promise on raw access and over-deliver on **reliability of the learning method** (attempts happen, soft-wrongs typed, micro-debrief protected).

Kill: Infinite-availability brand that ops cannot keep in Protect+probe.

Survive: Crunch copy that states response windows, acuity rules, and what parents can still expect tonight (Practice + coach / AI rail / scheduled HITL).

CXXVII.4 Measure trade-offs with CBC — do not ask parents to invent a price for “thinking minutes”

FACT (stated choice): Louviere, Hensher, & Swait (2000, *Stated Choice Methods*, Cambridge) — multi-attribute choice experiments recover relative utilities under random-utility assumptions better than unstructured preference talk.

FACT (hypothetical bias): Murphy, Allen, Stevens, & Weatherhead (2005, *Environmental and Resource Economics*, 30(3), 313–325, doi:10.1007/s10640-004-3332-x) — hypothetical WTP often exceeds actual; median hypothetical/actual ratio ~1.35 with heavy skew; treat absolute dollars as biased.

FACT (tutoring-adjacent DCE): Gambi & De Witte (2024, *Higher Education*, doi:10.1007/s10734-024-01358-z) — DCE recovers WTP for tutoring attributes; high stated WTP can coexist with low private uptake; information and SES frictions matter. **Wound:** HE students ≠ K-12 parents — method transfers, coefficients do not.

FACT (market skew): Kim, Goodman, & West (2024, CESifo WP 11251 / IZA DP 17178) — U.S. private tutoring centers concentrate in high-income, high-parent-education areas; one national “parent wants 24/7” story will misfire.

Applied (FOUNDER BELIEF → testable): Run crunch CBC (LOADSHED-5 / CRUNCHMSG-1):

Attribute	Levels (draft)
Access policy	Unlimited same-day / acuity queue (48h) / AI-first then HITL
Session promise	More minutes / protected FEI + micro-debrief / score-drill pack

Transparency	Silent / plain pause list / pause list + FEI weekly proof
Price	Calibrated span + neither (Khan/GPT/status quo)

HYPOTHESIS: Anxious-middle + HMA segments prefer **protected FEI + transparency** over unlimited access at equal/slightly higher price when comprehension is high. Score-chasers may overweight unlimited + drill — latent-class, don’t average.

Kill: Survey “how valuable is thinking minutes (1–7)?” as packaging gospel.

Survive: Part-worths → crunch SKU copy; cash/waitlist validation before national ads.

CXXVII.5 Fear-appeal crunch copy is a threat to motivation — and to parent control climate

FACT: Putwain & Remedios (2014, *School Psychology Quarterly*, 29(4), 360–373, doi:10.1037/spq0000048) — teacher fear appeals (failure-consequence messages) before high-stakes exams: higher frequency and threat appraisal predicted **lower self-determined motivation** and **lower exam performance**, partly mediated by motivation; extrinsic focus on avoiding failure is not a free motivational lever.

FACT (appraisal contingency): Putwain, Remedios, & Symes (2016, *British Journal of Educational Psychology*, 86(1), 16–31, doi:10.1111/bjep.12132) — challenge vs threat appraisals of fear appeals diverge; threat path links to worse engagement/performance patterns; challenge path can look different — **do not treat scare tactics as uniformly “motivating.”**

Applied (HYPOTHESIS): Parent emails that say “ACT is in 10 days — if we don’t book every night they will freeze” import classroom fear appeals into the home, raising **controlling parental pressure** (SAFE-SOFTMSG / SAFE-REWARD) and homework conflict. SAFE-CRUNCHMSG defaults to **challenge-framed structure**: named plan, protected attempts, one CTA — not failure theater.

Kill: Scare-parent default; “grind or fail” landing heroes.

Survive: “Here’s the plan for the next 14 days: fewer dump sessions, more equipped attempts + closes.”

CXXVII.6 Message arms (not the same as package CBC — run both)

Pre-registered arms (builds LOADSHED-5 / WTP message stack / EVT-4):

Arm	Parent-facing frame	Risk
A1 Always-on	“Tutors available 24/7 this exam month”	Expectation gap; pour
A2 Hours hero	“We added late sessions — book more”	Hours NS; fidelity dip
A3 Fear crunch	“Don’t let them walk in cold”	Threat; controlling nag

A4 Shed honesty	“We pause low-acuity dumps so tutors protect thinking minutes”	Reads as less product
A5 Hybrid proof	Shed honesty + FEI proof age + AI-first rail tonight	Longer; may win retention
A6 Wellness fog	“Remember to breathe — exams are stressful”	Care theater

Primary outcomes: click → trial; 14-day activation; comprehension (name one pause + one protected behavior); abandon perception; parent–child conflict — **not** brand warmth.

Kill: A1/A2/A3 as default national creative without CRUNCHMSG-* wins. **Survive:** A4/A5 if comprehension + retention win even when click loses to A1.

CXXVII.7 Product surface — SAFE-CRUNCHMSG claim contract

Surface	Required	Banned
Crunch email / in-app banner	Protected list + pause list in plain English	“Always on” without hours/acuity truth
Pricing / SKU blurb	Access policy as attribute, not vibes	Guaranteed ACT points / infinite human
Wait / acuity queue UI	ETA + why waiting + AI-first path	Shame (“you’re low priority”)
Parent dash crunch mode	FEI proof age; no hours-booked hero	Open-Late Minutes leaderboard
Ops truth	Playbook matches ads (SAFE-LOADSHED)	Market infinity, deliver rationing silently
Metrics	Intent; comprehension; abandon perception; conflict; transfer_pass @ crunch	Availability Score™; Crunch Message Score™

Competitive foil: Marketplace availability theater · ChatGPT dump · cram fear packs · Duo streaks. MindCraft = scarce human minutes aimed at thinking, with honest rails for everything else.

CXXVII.8 Doctrine — SAFE-CRUNCHMSG (provisional)

- 1. **Sell reliability of method, not infinity of access** — crunch brand = protected FEI + micro-debrief, not 24/7 (Oliver / Parasuraman gap logic).
- 2. **CBC over chatter** — force trade-offs with price + neither; ban standalone Likert WTP for “thinking minutes” (SAFE-WTP).
- 3. **Name the shed** — pause list in parent copy or honesty reads as ghosting (SAFE-LOADSHED).
- 4. **No fear-appeal default** — challenge-framed plans beat failure scare to parents (Putwain line).

5. **Autonomy-supportive CTAs** — one next action; no controlling “sit until booked every night” (SAFE-SOFTMSG / SAFE-REWARD).
6. **Segment** — score-chaser vs HMA-anxious vs homework-peace; do not average part-worths into one parent.
7. **Ads must match ops** — if playbook sheds HW dumps, creative cannot promise human Chegg.
8. **Copy:** “Exam week we protect thinking minutes — fewer dump sessions, clearer waits, proof of real attempts.”
Never Availability Score™ or ACT guarantees from “we stayed open.”

Confidence: High on method sources (gaps/SERVQUAL/Oliver; stated choice; Putwain). Medium on MindCraft win of shed-honesty vs 24/7 (needs CRUNCHMSG-*). High that infinite-availability / hours-hero / fear-crunch / Likert-WTP / Availability Score™ are commercially toxic under FEI + LOADSHED.

CXXVII.9 Experiments

ID	Question	Design	Primary
CRUNCHMSG-1	Protected-FEI+transparency vs unlimited-access packages	Parent CBC + price + neither	Part-worths; segment classes; comprehension
CRUNCHMSG-2	Message A4/A5 vs A1/A3 on landing + email	A/B / multi-arm	Click; trial; 14-day activation; abandon perception
CRUNCHMSG-3	Pause-list email vs silent defer when session shed	Field A/B crunch weeks	Churn; trust; complaint rate
CRUNCHMSG-4	Fear-crunch vs challenge-plan parent SMS	A/B	Conflict report; student affect proxy; booking quality (acuity fit)
CRUNCHMSG-5	Hybrid proof (FEI age + AI-first tonight) vs hours-hero	CBC / message	WTP; comprehension; homework-conflict
CRUNCHMSG-QUAL	10 parents in crunch: what felt like care vs abandonment?	Qual	CRUNCHMSG codebook (ties LOADSHED-QUAL)

Falsifier: Unlimited-access wins CBC *and* 30-day solo transfer under equal fidelity investment → revisit packaging (still refuse pour ads that ops cannot keep).

Falsifier: Shed-honesty raises abandon perception without comprehension gains → rewrite plain language; do not return to infinity brand.

Falsifier: Fear-crunch wins short-run bookings but raises conflict and kills transfer_pass → kill A3 permanently.

Pre-register: CRUNCHMSG-* before “24/7 exam month” national creative.

Family note: Densifies LOADSHED-5 / WTP-* / SOFTMSG; do not collapse into Availability Score™ OKRs.

CXXVII.10 So what for MindCraft commercially

- **Copy:** Lead crunch season with protected thinking minutes + plain pause list — not “always available.”
 - **Product:** Ship acuity-queue parent UX (why wait / what to do tonight); crunch-mode parent dash without hours hero.
 - **Pricing:** Treat access policy as a CBC attribute; price reliability of FEI, not infinite human minutes.
 - **Positioning:** Against marketplace availability theater and cram fear packs; beside honest scarce craft.
 - **Metric:** Comprehension of shed + retain/trust + FEI@crunch — demote Open-Late Minutes / Availability Score™.
 - **Kill list:** 24/7 default; hours-hero; fear-parent emails; silent shed; Likert-WTP gospel; ACT guarantees from crunch ads.
 - **Growth:** CRUNCHMSG-1/2 wins → premium WTP asset; losses → AI-rail honesty, not pour-mode.
 - **Vision:** Care = protected struggle reviewed — not a phone that always picks up to finish the worksheet.
-

References (verified)

- Gambi, L., & De Witte, K. (2024). The potential of tutoring in higher education: Students’ preferences, consumption, and the role of information. *Higher Education*. <https://doi.org/10.1007/s10734-024-01358-z>
 - Kim, T., Goodman, J., & West, M. R. (2024). *Kumon in: The recent, rapid rise of private tutoring centers* (CESifo Working Paper No. 11251 / IZA Discussion Paper No. 17178).
 - Louviere, J. J., Hensher, D. A., & Swait, J. D. (2000). *Stated choice methods: Analysis and applications*. Cambridge University Press.
 - Murphy, J. J., Allen, P. G., Stevens, T. H., & Weatherhead, D. (2005). A meta-analysis of hypothetical bias in stated preference valuation. *Environmental and Resource Economics*, 30(3), 313–325. <https://doi.org/10.1007/s10640-004-3332-x>
 - Oliver, R. L. (1980). A cognitive model of the antecedents and consequences of satisfaction decisions. *Journal of Marketing Research*, 17(4), 460–469. <https://doi.org/10.1177/002224378001700405>
 - Parasuraman, A., Zeithaml, V. A., & Berry, L. L. (1985). A conceptual model of service quality and its implications for future research. *Journal of Marketing*, 49(4), 41–50. <https://doi.org/10.1177/002224298504900403>
 - Parasuraman, A., Zeithaml, V. A., & Berry, L. L. (1988). SERVQUAL: A multiple-item scale for measuring consumer perceptions of service quality. *Journal of Retailing*, 64(1), 12–40.
 - Putwain, D. W., & Remedios, R. (2014). The scare tactic: Do fear appeals predict motivation and exam scores? *School Psychology Quarterly*, 29(4), 360–373. <https://doi.org/10.1037/spq0000048>
 - Putwain, D., Remedios, R., & Symes, W. (2016). Fear appeals, engagement, and examination performance: The role of challenge and threat appraisals. *British Journal of Educational Psychology*, 86(1), 16–31. <https://doi.org/10.1111/bjep.12132>
-

Part CXXVIII — Format-Trap Distractors on Diagram/Graph Items

Chapter status: Living evidence + Practice/Diagnostic FormatId item-craft brief — Researcher tick 2026-08-12 (UTC hour 00 ≡ Red Team slot, but ch128 never written → prefer Researcher per rotation; researcher count since synthesizer

v1.16 = 2 → Researcher)

Primary question: How should MindCraft design **visual / FormatId distractors** on diagram, coordinate_graph, and number_line MC so a wrong pick encodes a *representation theory* (graph-as-picture, slope–height, axis swap, legend miss) — without decoration theater, beauty-trap option art, FormatTrap Score™ vanity, or treating every visual miss as a naked “weak topic”?

Owners: Product (Practice / Diagnostic / bank figures) · Engine (FormatId + mis_ / format-trap seals) · GENQ / figure verify · HITL · Brand · Red Team

Commercial job: Ship **SAFE-FORMTRAP** densifying SAFE-DISTRACTOR × SAFE-FORMAT × SAFE-DUAL: **visual wrong options map named format traps**; soft-wrong climate routes conversion practice; no seductive-detail stems or decorative distractors sold as “visual learning.”

Builds on: Parts CXXIV (SAFE-DISTRACTOR), LXXXII (SAFE-FORMAT), XCVII (SAFE-DUAL), XLIX (SAFE-MISCON), CXXII (SAFE-TAXON), LXXII (SAFE-GENQ), LXXXV (SAFE-STORYLOAD), XLVIII (SAFE-TRANSFER). Product seams: FormatId bank tags, Eedi diagram recovery, GENQ figure verify, format-aware CTAs, Map format gaps.

CXXVIII.1 Why this chapter exists

SAFE-DISTRACTOR already bans beauty traps and unmapped wrongs for *symbolic* MC. Format-axis items reopen the failure mode in a sharper costume: wrong options that look like “visual engagement” — pretty wrong graphs, stock art beside stems, or option thumbnails that cue by color/length rather than by a known representational error.

Default edtech resolves this badly: (1) decorate the stem with franchise chrome and call it dual coding (SAFE-DUAL / SAFE-STORYLOAD), (2) invent three “almost right” graphs without naming *which* graph schema failed, (3) treat a diagram miss as “weak functions” instead of a FormatId gap (SAFE-FORMAT), (4) ship LLM-drawn options that are fluent but unkeyed (SAFE-GENQ), (5) FormatTrap Score™ / Visual Distractor Minutes as quality theater, (6) marketing “we teach every format” while distractors are decoration-adjacent gotchas.

FOUNDER BELIEF under audit: On diagram/graph MC, the commercially load-bearing wrong pick is a **format trap** — a plausible misreading of the vessel — not a random wrong number wearing a picture. MindCraft’s wedge vs symbolic drill packs and ChatGPT text-only fluency is *inspectable representation diagnosis* + conversion remediation.

Claims we refuse as doctrine:

1. Decorative stem art / seductive details ≡ FormatId coverage or pedagogy.
 2. FormatTrap Score™ / Visual Distractor Minutes / “richer visuals” NS.
 3. Beauty-trap option art (prettier wrong graph; length/color clang).
 4. Unmapped visual wrongs as “diagnostic FormatId soft-wrong.”
 5. Graph-as-picture / slope–height misses collapsed to topic weakness only.
 6. LLM figure options≡shipped bank without format-trap seal + key verify.
 7. Always-add-picture distractors ≡ dual coding science.
 8. ACT / grade guarantees from visual-distractor packaging.
-

CXXVIII.2 Constructs (format trap ≠ pretty wrong picture)

Construct	Research meaning	MindCraft analogue	Failure mode
Format trap	Plausible misread of a representation	Named seal on option (graph_as_picture, slope_height, axis_swap, legend_split, ...)	Bare “incorrect graph”
Graph-as-picture	Treat graph as literal scene sketch	Distractor whose shape mimics the story scene	Gotcha hill-drawing without route
Slope–height confusion	Confuse rate (slope) with value (height)	Option that answers height when stem asks slope (or vice versa)	Topic re-teach only
Conversion failure	Cannot translate across registers (Duval)	Format-gap CTA: diagram↔symbolic hop	Re-explain concept in strong vessel
Seductive detail	Interesting but irrelevant adjunct	Story chrome / joke art on stem or options	“Engagement” brand
Structural figure	Pictorial element encodes the math object	FormatId-necessary diagram/graph/number line	Decoration tagged as diagram
SAFE-FORMTRAP	Visual distractor craft + QA law	This chapter	FormatTrap Score™ / decoration theater

Operational definition (HYPOTHESIS): A diagram/graph MC item is *SAFE-FORMTRAP complete* when (a) the stem figure is **structurally necessary** (SAFE-DUAL / SAFE-FORMAT), not seductive decoration, (b) every visual distractor carries a **named format-trap intent** (or mis_/slip if non-visual), (c) options are homogeneous vessels (all graphs / all number lines — Haladyna homogeneity), (d) soft-wrong UX names the *representation theory* chosen and routes conversion practice, (e) generated figures fail ship without trap seal + key verify (SAFE-GENQ), and (f) value is proved by format-trap→route hit / format_hop_solo_transfer_pass / retry_120s — not Visual Distractor Minutes.

CXXVIII.3 Graphs are learned symbol systems — traps are schema errors

FACT (functions & graphs review): Leinhardt, Zaslavsky, & Stein (1990, *Review of Educational Research*, 60(1), 1–64, doi:10.3102/00346543060001001) — interpretation and construction of graphs involve specific misconceptions and intuitions; graphical and algebraic worlds are distinct symbol systems students must connect; teaching/learning research must treat graph comprehension as content-specific, not generic “visual talent.”

FACT (graph-as-picture & slope–height): Clement (1985, in Streefland, Ed., *Proceedings of the Ninth PME Conference*, Vol. 1, pp. 369–375) — documents **treating the graph as a picture** (global shape matching a scene; local feature correspondence such as “crossing = meeting”) and **slope–height confusions** on kinematics-style graphs; proposes cognitive accounts of non-standard symbolization strategies rather than random carelessness.

FACT (graph comprehension factors): Shah & Hoeffner (2002, *Educational Psychology Review*, 14(1), 47–69, doi:10.1023/A:1013180410169) — comprehension depends on visual characteristics of the display, viewer knowledge of graph schemas, and expectations about content; poor design and weak graph schemas produce systematic misreads — instructional implication: teach graphical literacy, do not assume figures are transparent.

Applied (HYPOTHESIS): Format-trap distractors should be **authored from these error families** (and MindCraft/Eedi mis_ analogues): graph-as-picture option; slope-when-height (and reverse); swapped axes / inverted ordered pairs; reading a category as a continuous slope; legend–series mismatch. Soft-wrong copy: “You read the graph as a picture of the hill” / “You answered height when the question asked rate” — then a FormatId hop CTA, not “review linear functions.”

Kill: Visual miss $\equiv$ topic weakness only; unnamed “wrong graph.”

Survive: Named format-trap seals $\rightarrow$ Map format gap / conversion Practice.

CXXVIII.4 Conversion across registers — why FormatId distractors matter commercially

FACT (semiotic registers): Duval (2006, *Educational Studies in Mathematics*, 61(1–2), 103–131) — **treatment** (within one register) and **conversion** (across registers) are distinct; conversion failures often decide comprehension even when within-register work looks fine.

Reuse (SAFE-FORMAT): Format gap = competence in vessel A with elevated fail in vessel B + remediation that targets reading/producing/converting the weak vessel.

Applied (HYPOTHESIS): A correct symbolic twin does not redeem a graph-as-picture pick. Bank events must carry FormatId + trap seal so /recommend can prefer format severity over concept-only narratives. Parent/tutor grain: “Missed the graph read, not the algebra idea” when evidence supports it — never Identity Score™ from one miss (SAFE-CONFMISS).

Kill: Single-format accuracy $\equiv$ concept mastery; format-count ads.

Survive: Conversion as the remediation object after a format-trap pick.

CXXVIII.5 Coherence — seductive details are not FormatId distractors

FACT (coherence / seductive details): Mayer, Heiser, & Lonn (2001, *Journal of Educational Psychology*, 93(1), 187–198, doi:10.1037/0022-0663.93.1.187) — adding interesting but extraneous multimedia can impair understanding.

FACT (meta-analysis): Sundararajan & Adesope (2020, *Educational Psychology Review*, 32, 707–734, doi:10.1007/s10648-020-09522-4) — across dozens of studies, including seductive details tends to **hinder** learning outcomes; moderators include design and delivery — practical implication aligns with the coherence principle: cut irrelevant interesting material.

Applied (FOUNDER BELIEF $\rightarrow$ testable): Story-world chrome that does not encode the probed structure is **banned on diagnostic FormatId stems and options**. A “visual distractor” that is merely prettier wrong art is a beauty trap

(SAFE-DISTRACTOR), not a format trap. Decorative thumbnails that clang by color/size violate Haladyna choice guidelines and poison diagnosis.

Kill: Seductive stem art; beauty-trap option pictures; “visual engagement” as distractor rationale.

Survive: Structurally necessary figures; trap-mapped homogeneous options.

Wound: Prove-rail ACT items may include exam-realistic clutter; learn/diagnostic rails optimize coherence + trap seals — label the dual-rail (SAFE-EXAM / SAFE-DDSCHED).

CXXVIII.6 Item-writing law still applies — now to vessels

FACT (MC guidelines): Haladyna, Downing, & Rodriguez (2002, *Applied Measurement in Education*, 15(3), 309–334, doi:10.1207/S15324818AME1503_5) — distractors from typical student errors; options homogeneous; avoid clues (clang, absurd options, length cues); three effective options often suffice.

FACT (diagnostic MC): Wang et al. (2020, arXiv:2007.12061) — each incorrect answer chosen to highlight a common misconception; choice reveals the nature of the misconception for remediation.

Reuse (SAFE-DISTRACTOR): Function over fashion; retire dead options; GENQ needs seals.

Applied (HYPOTHESIS): For FormatId MC: (1) write trap families first, (2) render options as **parallel vessels** (same graph type, comparable ink), (3) verify the key structurally (SAFE-GENQ figure verify — not fluency of SVG), (4) soft-wrong names the trap + routes conversion, (5) prefer ~3 options when a fourth is decoration or non-functioning.

Kill: Four pretty graphs as quality costume; unmapped visual MC as “diagnostic.”

Survive: Trap-first authoring; functioning analysis on visual options too.

CXXVIII.7 Product surface — SAFE-FORMTRAP claim contract

Surface	Ship	Banned substitute
Stem figure	Structural FormatId vessel only	Seductive story art; stock “math vibes”
Option metadata	format_trap_id / mis_ / slip / untyped-bug	Bare “incorrect”
Soft-wrong	Destake + name trap + conversion CTA	Red X; “review the topic”; Miss Streak
Bank QA	Trap seal %; functioning %; figure key verify	Visual Distractor Minutes
GENQ	Fail ship without trap seal + structural key	Fluency-only SVG pass
Map / recommend	Format gap when trap pattern + other-vessel competence	Topic-only collapse

Metrics	trap→route hit; format_hop_solo_transfer_pass; retry_120s	FormatTrap Score™
Parent/tutor	“Graph-as-picture trap on rate graph” grain	Live option stalk / % wrong

Competitive foil: Khan/Duo = binary or XP on mixed visuals. Exam packs = four-option costume. ChatGPT = fluent unkeyed figure talk. “Visual learning” apps = decoration. MindCraft = **format-trap distractors + soft-wrong + FormatId route + FEI prove**.

CXXVIII.8 Doctrine — SAFE-FORMTRAP (provisional)

- Trap families first** — author from graph-as-picture, slope–height, axis/legend/schema errors; stems serve diagnosis.
- Every visual wrong has intent** — format_trap_id (or mis_/slip) — never naked incorrect graph.
- Structure over seduction** — stem/option figures must encode the math object; cut seductive details (Mayer; Sundararajan & Adesope).
- Homogeneous vessels** — options comparable; no beauty/clang cues (Haladyna).
- Soft-wrong names the representation theory** — then conversion Practice, not topic dump alone (Duval; SAFE-FORMAT).
- No FormatTrap Score™ / Visual Distractor Minutes NS** — bank health = seal rate + functioning % + format-hop transfer.
- GENQ equals craft** — generated figures need trap seals + structural key verify (SAFE-GENQ / SAFE-DUAL).
- Dual-rail honesty** — prove-rail may match ACT clutter; learn/diagnostic rails optimize coherence + traps — label which rail.
- Copy:** “Wrong graphs that teach — each option is a reading theory we can route.” Never “more visuals = smarter test.”

Confidence: High — Leinhardt et al.; Clement; Shah & Hoeffner; Duval; Mayer et al.; Sundararajan & Adesope; Haladyna et al.; Wang et al. Medium — MindCraft live functioning rates for visual options and auto-trap-seal accuracy (needs FORMTRAP-*). High — decoration theater / beauty visual traps / unmapped “FormatId diagnostic” MC as commercially toxic under FEI + Map honesty.

CXXVIII.9 Experiments

ID	Question	Design	Primary
FORMTRAP-1	Trap-mapped 3-option graph MC vs unmapped pretty 4-option	A/B within concept+FormatId	trap→route hit; format_hop_solo_transfer_pass; retry_120s

FORMTRAP-2	Coherent structural stem vs seductive-detail stem (same keys)	A/B	Accuracy; abandon; transfer; load self-report
FORMTRAP-3	Soft-wrong naming format trap vs generic “not quite” / topic-only	A/B	Uptake; retry_120s; format-hop transfer
FORMTRAP-4	GENQ with format-trap seal + figure key verify vs fluency-only	Bank ops	Seal rate; false-trap rate; drop rate
FORMTRAP-5	Parent/tutor CBC: format-trap grain vs % wrong / topic weakness	CBC	WTP; trust; shame (SAFE-WTP/SOFTMSG)
FORMTRAP-QUAL	10 students + 5 tutors: which wrong graphs felt “tricky fair” vs gotcha/decoration?	Qual	FORMTRAP codebook

Falsifier: Unmapped pretty 4-option matches trap-mapped 3-option on route + transfer → keep private seals; still ban seductive stems and FormatTrap Score™.

Falsifier: Seductive stems raise engagement without hurting transfer → still refuse diagnostic-rail decoration; segregate marketing chrome off-item.

Falsifier: Fluency-only GENQ figures equal sealed on false-trap rate → still require structural key+seal (SAFE-GENQ / LABMETA).

Pre-register: FORMTRAP-* before “visual diagnostic MC / AI graph distractors” ads.

Family note: FORMTRAP-* densifies DISTRACTOR/FORMAT/DUAL; do not collapse into Visual Minutes or FormatTrap Score™ OKRs.

CXXVIII.10 So what for MindCraft commercially

- **Copy:** “Wrong graphs that teach — each option is a reading theory we can route.” Lead with format-trap diagnosis + conversion, not “thousands of visual questions” or dual-coding wallpaper.
- **Product:** Require format_trap_id (or typed mis_) on visual options; coherent structural stems; wire chosen trap into SAFE-TAXON mark + FormatId Map route; GENQ figure verify gates.
- **Positioning:** Against decoration apps, option-count theater, and ChatGPT unkeyed figure fluency; beside Eedi-class diagnosis, ahead on FormatId conversion + identity climate + FEI prove.
- **Metric:** format-trap seal %, functioning %, trap→route hit, format_hop_solo_transfer_pass, retry_120s — demote FormatTrap Score™ / visual item-count.
- **Kill list:** Seductive stems; beauty visual traps; unmapped “FormatId diagnostic”; FormatTrap Score™; GENQ without trap seals; ACT guarantees from visual-distractor UI.
- **Growth:** Trust packets show trap→misconception/format map + figure verify loop — procurement asset vs black-box “visual banks.”

• **Vision:** Thirty-year identity company treats a wrong graph pick as **evidence of a representation theory in use** — so Maya learns to revise *how she reads the vessel*, not to fear pictures or chase prettier wrong answers.

References (verified)

- Clement, J. (1985). Misconceptions in graphing. In L. Streefland (Ed.), *Proceedings of the Ninth International Conference for the Psychology of Mathematics Education* (Vol. 1, pp. 369–375). Utrecht University.
 - Duval, R. (2006). A cognitive analysis of problems of comprehension in a learning of mathematics. *Educational Studies in Mathematics*, 61(1–2), 103–131.
 - Haladyna, T. M., Downing, S. M., & Rodriguez, M. C. (2002). A review of multiple-choice item-writing guidelines for classroom assessment. *Applied Measurement in Education*, 15(3), 309–334. https://doi.org/10.1207/S15324818AME1503_5
 - Leinhardt, G., Zaslavsky, O., & Stein, M. K. (1990). Functions, graphs, and graphing: Tasks, learning, and teaching. *Review of Educational Research*, 60(1), 1–64. <https://doi.org/10.3102/00346543060001001>
 - Mayer, R. E., Heiser, J., & Lonn, S. (2001). Cognitive constraints on multimedia learning: When presenting more material results in less understanding. *Journal of Educational Psychology*, 93(1), 187–198. <https://doi.org/10.1037/0022-0663.93.1.187>
 - Shah, P., & Hoeffner, J. (2002). Review of graph comprehension research: Implications for instruction. *Educational Psychology Review*, 14(1), 47–69. <https://doi.org/10.1023/A:1013180410169>
 - Sundararajan, N., & Adesope, O. (2020). Keep it coherent: A meta-analysis of the seductive details effect. *Educational Psychology Review*, 32, 707–734. <https://doi.org/10.1007/s10648-020-09522-4>
 - Wang, Z., Lamb, A., Saveliev, E., Cameron, P., Zaykov, Y., Hernández-Lobato, J. M., Turner, R. E., Baraniuk, R. G., Barton, C., Peyton Jones, S., Woodhead, S., & Zhang, C. (2020). Instructions and guide for diagnostic questions: The NeurIPS 2020 Education Challenge. *arXiv:2007.12061*. <https://arxiv.org/abs/2007.12061>
-

Part CXXIX — Session-Open Brief Fidelity Under Late Joins

Chapter status: Living evidence + HITL open-ops brief — Researcher tick 2026-08-12 (UTC hour 03; hour%6≠0; researcher count since synthesizer v1.16 = 3 → Researcher)

Primary question: When the session clock is already burned (late student, late tutor, crunch week), how does MindCraft keep **session-open brief fidelity** — named edge, success criterion, likely trap, coach role — in ~30 seconds without skipping into homework dump, lecturing a 10-minute “catch-up,” military/OR cosplay, or inventing a Brief Score™?

Owners: HITL ops / tutor console · Product (session spine / mission card) · Coach UX · Exam calendar (SAFE-LOADSHED) · Brand · Red Team

Commercial job: Ship **SAFE-MICROBRIEF** densifying SAFE-AAR × SAFE-MICRODEB × SAFE-LOADSHED × SAFE-HITL × SAFE-EXPECTANCY: a **capped, structured micro-brief** (~20–40s practice / ≤60s live under late join) that protects learning *alignment* when time is scarce — never “no time for goals,” never pour-open, never checklist theater as ACT ads.

Builds on: Parts LIV (SAFE-AAR), CXXVI (SAFE-MICRODEB), CXXIII (SAFE-LOADSHED), LXVIII (SAFE-HITL), XLVI (SAFE-EXPECTANCY), CI (SAFE-SRL), CXIX (SAFE-MAPORCH), XC

(SAFE-INSTRUMENT). Product seams: session timer, late-join flag, mission/edge card, Map grain slice, soft-wrong priors, Notes open prompts, acuity queue under crunch.

CXXIX.1 Why this chapter exists

SAFE-AAR says **brief** → **attempt** → **structured debrief**. SAFE-MICRODEB protects the *close* under a hard clock. What remains underspecified is the **open under a burned clock**: tutors skip (“let’s just start”), pour a monologue catch-up, or vibes-open until the FEI window is gone — then ops still claim they “do briefs.”

Competitors resolve late joins wrongly: (1) **skip** into homework help; (2) **pour-open**; (3) **relationship theater** without naming an edge; (4) **bootcamp** shame-speed; (5) **AI syllabus dump**. MindCraft’s wedge is a ~30s **micro-brief** — shared mental model + specific goals + SHARP “Set” — that finishes without losing aim.

FOUNDER BELIEF under audit: Identity transformation needs *aimed attempts*. Skipping the brief is not efficiency — it is a fidelity breach that leaves the later micro-debrief unanchored.

Claims we refuse as doctrine:

1. “No time for brief / goals” as late-join default.
 2. Tutor/AI lecture that *calls itself* a session open.
 3. Brief Score™ / Brief Minutes / Opening Streak as North Star.
 4. Military / surgical / aviation cosplay as ACT marketing.
 5. Mandatory long rapport block when the clock is the constraint.
 6. Vibes-only open without named edge + success criterion.
 7. Soft standards justified by “we at least connected.”
 8. ACT / score guarantees from micro-brief packaging.
-

CXXIX.2 Constructs (short ≠ empty; late ≠ aimless)

Construct	Meaning	MindCraft analogue	Failure mode
Full brief	Pre-event alignment: intent, success criteria, traps, roles (CRM)	2–4 min tutor open (SAFE-AAR)	Sermon; fate talk
Micro-brief	Same function, hard time box, fewer beats	20–40s solo; ≤60s live late-join	Skip or pour
Late join	Clock already burned before work starts	Session timer + late flag	Excuse to skip aim
Shared mental model	Both parties hold same situation / goal / next action	Named edge + trap + coach role	Assumed “you know”
Receiver synthesis	Student restates plan in own words	“You’re aiming at X; success = Y”	Tutor-only monologue

SAFE-MICROBRIEF	Capped open fidelity law	This chapter	Brief Minutes NS
-----------------	--------------------------	--------------	------------------

Operational definition (HYPOTHESIS): An open is *SAFE-MICROBRIEF complete* when (a) a **named time cap** is visible ($\leq 40s$ solo / $\leq 60s$ live under late-join or LOADSHED; longer only when acuity warrants and calendar allows), (b) the open runs **four beats minimum**: named edge $\rightarrow$ success criterion $\rightarrow$ one likely trap / FormatId note $\rightarrow$ coach role (hint vs silence), (c) the **student produces ≥ 1 synthesis** (“so today I’m trying...”) before first worked problem, (d) output is a mission/Notes card the micro-debrief can restate — not a syllabus sermon, (e) under SAFE-LOADSHED the micro-open is on the **protect** list (not shed for HW dump), and (f) success = brief_complete $\times$ later retry_120s / method-change / solo_transfer_pass — never Brief Minutes or Brief Score™.

CXXIX.3 Specific goals beat “do your best” — even in 30 seconds

FACT: Locke & Latham (2002, *American Psychologist*, 57(9), 705–717, doi:10.1037/0003-066X.57.9.705) — 35-year synthesis: **specific, difficult goals** produce higher task performance than vague, easy, or “do-your-best” goals across laboratory, simulation, and field settings; mechanisms include direction, effort, persistence, and strategy arousal; moderators include commitment, feedback, and task complexity.

Applied (HYPOTHESIS): A late-join open that says “let’s just work hard on algebra” is a **do-your-best** failure. A 30-second micro-brief that names “linear inequalities — multiply-by-negative trap; success = one solo transfer without peek” is the tutoring analogue of specific goal setting. Under crunch, **specificity is the time-saver**, not the luxury.

Kill: “We’ll figure it out as we go” as default late-join open.

Survive: One named edge + one observable success criterion.

Wound: Organizational task goals $\neq$ anxious Maya under ACT clock — MICROBRIEF- *must pre-register FEI instruments, not claim Locke d** as ACT points.

CXXIX.4 Shared mental models travel — SBAR / I-PASS as method, not hospital cosplay

FACT: Haig, Sutton, & Whittington (2006, *Joint Commission Journal on Quality and Patient Safety*, 32(3), 167–175) — case study of **SBAR** (Situation, Background, Assessment, Recommendation) as a **shared mental model** tool to standardize clinician communication and reduce style-mismatch failures; pocket-card / phone-cheat-sheet implementation spread beyond single handoffs into reporting culture.

FACT: Starmer et al. (2014, *New England Journal of Medicine*, 371(19), 1803–1812, doi:10.1056/NEJMsa1405556) — multicenter prospective intervention of the **I-PASS** handoff bundle across nine hospitals ($N = 10,740$ admissions): medical-error rate $\downarrow 23\%$ (24.5 $\rightarrow$ 18.8 per 100 admissions), preventable adverse events $\downarrow 30\%$ (4.7 $\rightarrow$ 3.3); **oral handoff duration did not increase** (2.4 vs 2.5 min/patient, $P = 0.55$); key elements in written/oral communication rose significantly. Bundle includes mnemonic + training + observation — structure without workflow penalty.

Applied (FOUNDER BELIEF $\rightarrow$ testable): Late-join tutoring is a **handoff into a truncated shift**. Transfer the method:

1. **Situation** — clock state / late flag / remaining minutes.
2. **Background** — last edge / soft-wrong / Map grain (one card, not biography).
3. **Assessment** — acuity: FEI practice vs HW triage (SAFE-HWHELP dual-rail).
4. **Recommendation / action** — today's edge + success criterion + coach role.
5. **Synthesis by receiver** — student restates (I-PASS critical element).

Starmer's non-finding on duration is commercially load-bearing: **structure** $\neq$ **longer**. Skip is not required to "save time."

Kill: Hospital/aviation brand cosplay; "I-PASS tutoring raises ACT 23%."

Survive: Shared-model micro-open with student synthesis; fidelity without minute inflation.

Wound: Medical error rates $\neq$ transfer_pass — method transfer only.

CXXIX.5 Time-boxed checklists under production pressure

FACT: Haynes et al. (2009, *NEJM*, 360(5), 491–499, doi:10.1056/NEJMsa0810119) — WHO Surgical Safety Checklist across eight hospitals; death 1.5% $\rightarrow$ 0.8%, complications 11.0% $\rightarrow$ 7.0%; checklist is explicitly **time-feasible** at three junctures, not an infinite briefing.

FACT: Degani & Wiener (1990, NASA CR-177549) — production pressures ("making schedules") interact with checklist misuse/non-use; design must respect interruption and usable length.

FACT: FAA AC 120-51E lists effective **briefings** among CRM markers: coordination, planning, open communication, expectations for SOP deviations — briefs are **error-management before the event** (SAFE-AAR lineage).

Applied (HYPOTHESIS): Micro-brief is a **timeout checklist**, not a pre-flight lecture. If late join burns warm-up, **hold and resume** at the named edge — do not abandon the list. Exam-week schedule pressure predicts skip; product must make the short form the path of least resistance.

Kill: "Full CRM brief or nothing" purity that becomes skip; aviation costume ads.

Survive: Cap-visible, interruption-safe four-beat open on the protect list.

CXXIX.6 SHARP's first beat is the brief — Set before How

FACT: Ahmed, Arora, Russ, Darzi, Vincent, & Sevdalis (2013, *Annals of Surgery*, 258(6), 958–963, doi:10.1097/SLA.0b013e31828c88fc) — **SHARP** scaffold for time-feasible OR debrief: **Set learning objectives**; How did it go; Address concerns; Review learning points; Plan ahead. Prospective pre/post 100 cases: OSAD medians 19 $\rightarrow$ 33 ($P < 0.001$). Companion positioning: SHARP as "absolute basics" when full OSAD-depth is impractical.

Applied (HYPOTHESIS): SAFE-MICRODEB already borrows SHARP for the close. SAFE-MICROBRIEF owns the **Set** beat at the *open* — so the later How/Address/Review/Plan have something to compare against. A micro-debrief without a prior micro-brief collapses into vibes ("how did that feel?") because expected was never named.

Kill: Close-only AAR theater without open aim.

Survive: Pair MICROBRIEF + MICRODEB as one capped cycle under late join / LOADSHED.

CXXIX.7 Expectancy ethics still bind the 30-second open

FACT / doctrine carry: SAFE-EXPECTANCY (Part XLVI) — task-only language; no trait labels; tutor briefs must not smuggle “you’re behind...” fate talk. High-dosage tutoring toolkits emphasize consistent **framing & objective** after a light open — not icebreaker-as-product (NSSA session-structure guidance; practice note, not a MindCraft RCT).

Applied (HYPOTHESIS): Under late join, the temptation is either (a) skip aim or (b) shame-speed (“we’re late so grind harder”). Both fail. Micro-brief copy stays **task-referenced:** edge, criterion, trap, coach role. One belonging beat ≤10s is allowed (“glad you’re here”) then into Set — not an 8-minute weekend recap when 12 minutes remain.

Kill: Late-join shame open; endless rapport under crunch as “care.”

Survive: Warm + specific; SAFE-ERRCLIMATE destake without soft standards.

CXXIX.8 Product surface — SAFE-MICROBRIEF claim contract

Clock context	Micro-brief form	Protect	Banned substitute
Solo Practice start	20–40s: edge / success / trap / coach-self role	Student synthesis field	Skip to first item
Live tutor, on-time	2–4 min full brief (SAFE-AAR)	Shared plan card	Captain monologue
Live tutor, late join / LOADSHED	≤60s four-beat micro + student synthesis	On protect list	Skip into HW dump
Tutor late, student waiting	Async pre-brief card (student-owned Set)	SRL forethought	Idle lobby / answer scroll
Soft-wrong prior heavy	Name trap from last mark type	Expectancy-safe language	Trait lecture
AI coach assist	Prompt beats only; student writes Set	SAFE-EXPLAINQA gates	Syllabus dump first
Metrics	brief_complete; synthesis_ok; edge_named; paired microdeb_complete	Co-gate with FEI	Brief Score™; Brief Minutes

Competitive foil: Marketplace tutor = skip or pour-open. Chegg/ChatGPT = dump without aim. Bootcamp = shame-speed. Duo = streak liturgy. MindCraft = **capped, synthesized, edge-named open that survives the late join.**

CXXIX.9 Doctrine — SAFE-MICROBRIEF (provisional)

1. **Cap the form, don't cancel the aim** — under late join / crunch, run micro-brief; never “no time for goals” as default (LOADSHED protect).
2. **Four beats minimum** — named edge → success criterion → one trap/FormatId note → coach role.
3. **Student synthesizes** — ≥ 1 restatement before first problem (I-PASS receiver synthesis analogue).
4. **Pair with micro-debrief** — Set at open enables How/Plan at close (SAFE-MICRODEB).
5. **Shared-model method** — SBAR/I-PASS structure without hospital cosplay; Starmer: structure $\neq$ longer.
6. **Specific > do-your-best** — Locke & Latham: name the edge (Locke & Latham, 2002).
7. **Expectancy-safe** — task language only; no late-join shame (SAFE-EXPECTANCY).
8. **Copy**: “Even if we start late, we still name the edge and what success looks like — in under a minute.” Never “aviation-grade briefs raise ACT” or Brief Score™.

Confidence: High — Locke/Latham goal specificity + Starmer duration-neutral structure gains + SBAR shared-model + WHO time-boxed checklist + Degani/Wiener production-pressure warning + SHARP Set beat as method sources. Medium — MindCraft-specific 30–60s completeness thresholds and late-join protect-list compliance (needs MICROBRIEF-*). High — skip/pour-open / Brief Score™ / cosplay ACT ads as commercially toxic under FEI + SAFE-AAR.

CXXIX.10 Experiments

ID	Question	Design	Primary
MICROBRIEF-1	≤ 60 s micro-brief vs skip-to-work under equal remaining time	HITL A/B late-join	brief_complete; retry_120s; solo_transfer_pass
MICROBRIEF-2	Four-beat + student synthesis vs tutor free-form open ≤ 60 s	HITL A/B	synthesis_ok; talk-ratio; edge_named
MICROBRIEF-3	Async student pre-brief card while tutor late vs idle wait	Ops A/B	forethought quality; time-to-first-FEI
MICROBRIEF-4	LOADSHED: protect micro-open vs shed-open-for-HW	Ops A/B crunch	fidelity@late; pour↓; transfer_pass
MICROBRIEF-5	Parent/tutor CBC: “named aim in 30s” vs “we jump straight into homework”	CBC	WTP; trust (SAFE-WTP / CRUNCHMSG)
MICROBRIEF-QUAL	10 tutors + 10 Mayas: did 30s feel clarifying or rushed/shaming?	Qual	MICROBRIEF codebook

Falsifier: Skip equals micro-brief on delayed transfer → still refuse pour-open; investigate whether bank mismatch, not open, was bottleneck.

Falsifier: Free-form ≤60s beats four-beat on FEI → keep student synthesis + named edge; demote rigid mnemonic if it becomes costume.

Falsifier: Synthesis field raises anxiety without method-change → tighten SAFE-ERRCLIMATE; do not return to vibes-only.

Pre-register: MICROBRIEF-* before “I-PASS / WHO checklist tutoring / +23%” ads.

Family note: MICROBRIEF-* densifies AAR/MICRODEB/LOADSHED/HITL/EXPECTANCY; do not collapse into Brief Minutes OKRs.

CXXIX.11 So what for MindCraft commercially

- **Copy:** “Even when we start late, we still name the edge and what success looks like — in under a minute.”
 - **Product:** Late-join micro-brief UI; student synthesis field; auto-attach last soft-wrong / Map grain; block “start problems” without brief_complete (logged acuity override).
 - **Ops:** Protect micro-brief beside micro-debrief on LOADSHED; QA synthesis_ok and pour-open rate — not Brief Minutes.
 - **Positioning:** Against skip-culture marketplaces and monologue AI; high-reliability craft without costume.
 - **Metric:** brief_complete, synthesis_ok, edge_named, paired microdeb_complete, FEI events — demote Brief Score™.
 - **Kill list:** Skip-as-default; pour-open; late-join shame; Brief Score™; aviation/OR ACT ads; vibes-only open.
 - **Growth / vision:** Trust packet shows open+close fidelity under late joins; Maya learns hard sessions always start with a fair target.
-

References (verified)

- Ahmed, M., Arora, S., Russ, S., Darzi, A., Vincent, C., & Sevdalis, N. (2013). Operation debrief: A SHARP improvement in performance feedback in the operating room. *Annals of Surgery*, 258(6), 958–963. <https://doi.org/10.1097/SLA.0b013e31828c88fc>
 - Degani, A., & Wiener, E. L. (1990). *Human factors of flight-deck checklists: The normal checklist* (NASA Contractor Report 177549). NASA Ames Research Center.
 - Federal Aviation Administration. (2004). *Crew resource management training* (Advisory Circular AC 120-51E).
 - Haig, K. M., Sutton, S., & Whittington, J. (2006). SBAR: A shared mental model for improving communication between clinicians. *Joint Commission Journal on Quality and Patient Safety*, 32(3), 167–175.
 - Haynes, A. B., Weiser, T. G., Berry, W. R., Lipsitz, S. R., Breizat, A.-H. S., Dellinger, E. P., ... Gawande, A. A. (2009). A surgical safety checklist to reduce morbidity and mortality in a global population. *New England Journal of Medicine*, 360(5), 491–499. <https://doi.org/10.1056/NEJMs0810119>
 - Locke, E. A., & Latham, G. P. (2002). Building a practically useful theory of goal setting and task motivation: A 35-year odyssey. *American Psychologist*, 57(9), 705–717. <https://doi.org/10.1037/0003-066X.57.9.705>
 - Starmer, A. J., Spector, N. D., Srivastava, R., West, D. C., Rosenbluth, G., Allen, A. D., ... Landrigan, C. P., for the I-PASS Study Group. (2014). Changes in medical errors after implementation of a handoff program. *New England Journal of Medicine*, 371(19), 1803–1812. <https://doi.org/10.1056/NEJMs1405556>
-

Part CXXX — Dual-Rail Solver CTA Copy Under Crunch

Chapter status: Living evidence + product-copy brief — Researcher tick 2026-08-12 (UTC hour 09; hour%6≠0; researcher count since synthesizer v1.16 = 4 → Researcher)

Primary question: In exam crunch weeks, how should MindCraft present **student and parent CTAs** for Solver vs Practice — “**help tonight (instrumental)**” vs “**Practice gym (prove alone)**” — without collapsing into a Chegg hero button, fear-appeal dump, purity ban on help, or inventing a Dual-CTA Score™?

Owners: Product (Solver / Practice entry surfaces) · Brand / growth copy · Parent dash · Coach UX · Exam calendar (SAFE-LOADSHED / SAFE-CRUNCHMSG) · Red Team

Commercial job: Ship **SAFE-CRUNCHCTA** densifying SAFE-HWHELP × SAFE-CRUNCHMSG × SAFE-HELP × SAFE-DURABLE × SAFE-SOFTMSG: a **choice-architecture contract** for dual CTAs under deadline pressure — labeled rails, honest defaults, follow-on Practice grain — never single-rail “finish the worksheet,” never 24/7 dump ads, never assisted=ready.

Builds on: Parts CIV (SAFE-HWHELP), CXXVII (SAFE-CRUNCHMSG), LXXVIII (SAFE-DURABLE), LXXXVIII–LXXXIX (SAFE-FADE / SAFE-HELP), XCI (SAFE-HINT), CXVIII (SAFE-SOFTMSG), CXXIII (SAFE-LOADSHED), XC (SAFE-INSTRUMENT). Product seams: Solver open, Practice launch, homeworkHelp deep-link, parent crunch messages, acuity queue, solo_transfer_pass, fade ladder.

CXXX.1 Why this chapter exists

SAFE-HWHELP already forbids becoming a nicer Chegg: two labeled rails, instrumental default, solo proof as identity. SAFE-CRUNCHMSG already forbids selling infinite availability and fear-parent emails. What remains underspecified is the **button layer** — the actual CTAs students and parents see when the ACT calendar is loud and Sunday night is real.

Competitors resolve the button wrongly: (1) **single Chegg hero** (“Get the answer / Finish homework”); (2) **fear dump** (“ACT in 9 days — open Solver now”); (3) **purity freeze** (hide help, lose GTM); (4) **equal twin buttons** with no default (decision fatigue → dump); (5) **streak / tonight-%** as the CTA subcopy. MindCraft’s wedge is **dual CTAs with architecture**: help is available and labeled; Practice is visible and often the default or the mandatory follow-on; neither rail markets assisted completion as identity.

FOUNDER BELIEF under audit: Crunch weeks are when dual-rail doctrine dies or becomes brand. If the primary CTA collapses to “help tonight” alone, MindCraft trains the market to buy relief and forget the gym.

Claims we refuse as doctrine:

1. Single Chegg-hero Solver CTA as crunch default.
2. “Finish tonight’s homework” / tonight-% as primary brand CTA.
3. Fear-appeal CTAs that couple exam countdown → answer dump.
4. Dual-CTA Score™ / Help Tonight Minutes / Crunch Click NS.
5. Silent collapse of the Practice rail when exam week starts.
6. Assisted Solver success marketed as exam-ready / identity.

7. Purity ban that removes instrumental help under deadline need.
8. ACT / score guarantees from dual-rail CTA packaging.

CXXX.2 Constructs (two buttons ≠ two brands)

Construct	Meaning	MindCraft analogue	Failure mode
Help-tonight rail	Instrumental unstick for a due assignment	Solver (SAFE-HWHELP / SAFE-FADE)	Executive dump
Practice-gym rail	Attempt → soft-wrong → solo transfer	Practice / mission launch	Hidden under crunch
CTA pair	Visible dual choice with labels + consequence copy	Home / Solver / parent crunch cards	Twin equal / single hero
Default	Pre-selected or visually primary option	Choice architecture	Dump-as-default
Follow-on grain	Required or prompted Practice after help	solo_transfer_pass handoff	Dump-and-quit
SAFE-CRUNCHCTA	Dual-rail CTA law under crunch	This chapter	Dual-CTA Score™

Operational definition (HYPOTHESIS): A crunch-week entry surface is *SAFE-CRUNCHCTA complete* when (a) **both rails are visible** with distinct verbs (e.g. “Unstick tonight’s problem” vs “Practice to prove it alone”), (b) help CTAs are **instrumental-framed** (hint/fade language — not “get the answer”), (c) **default / visual primacy** favors Practice or a balanced pair with explicit consequence subcopy — never dump-as-sole-primary, (d) every help-tonight path offers or requires a **follow-on Practice grain** when feasible, (e) parent CTAs mirror dual-rail honesty (SAFE-CRUNCHMSG / SAFE-PDASH) without live stalk or tonight-% NS, and (f) success = rail_labeled × followon_start × solo_transfer_pass / retry_120s — never Help Tonight Minutes or Dual-CTA Score™.

CXXX.3 Defaults are not neutral — crunch CTAs choose a future self

FACT: Johnson & Goldstein (2003, *Science*, 302(5649), 1338–1339, doi:10.1126/science.1091721) — organ-donation consent rates differ dramatically by **opt-in vs opt-out defaults** across European countries and in laboratory choice; every policy has a no-action default that imposes switching costs; defaults are not psychologically neutral.

FACT: Thaler, Sunstein, & Balz (2013, in E. Shafir, Ed., *The Behavioral Foundations of Public Policy*, pp. 428–439; earlier SSRN working paper “Choice Architecture”) — choice architects always set defaults, feedback, and option structure; a nudge alters behavior predictably without forbidding options or heavily changing economic incentives; ignoring architecture is still choosing one.

Applied (HYPOTHESIS): A Solver screen whose only bold button is “Show solution / Finish homework” is an **opt-out-of-learning default**. Under crunch, inertia + deadline anxiety make the dump path the organ-donation

analogue of sticking with the pre-checked box. SAFE-CRUNCHCTA treats the primary CTA as a **fiduciary default**: instrumental help remains available (not forbidden), but the architecture must not silently enroll Maya in executive outsourcing.

Kill: Single dump button as “what parents want this week.”

Survive: Dual CTAs + deliberate default / primacy toward gym or labeled pair with follow-on.

Wound: Organ-donation ethics $\neq$ homework UX — transfer *default power*, not moral theater; CRUNCHCTA- *must measure quit and help-avoidance, not claim Science d** as ACT points.

CXXX.4 Executive vs instrumental help — the CTA verb is the pedagogy

FACT: Nelson-Le Gall (1981, *Developmental Review*, 1(3), 224–246, doi:10.1016/0273-2297(81)90019-8) — distinguishes **executive help** (have the helper solve / produce the answer) from **instrumental help** (help that enables the learner to complete the task themselves); adaptive help-seeking is metacognitive and contingent, not mere dependency.

FACT: Bastani, Bastani, Sungu, Ge, Kabakc■, & Mariman (2025, *PNAS*, 122(26), e2422633122, doi:10.1073/pnas.2422633122) — field experiment (~1,000 high-school math students): unfettered GPT-4 “Base” access raised practice performance but **harmed** later solo performance (–17% vs never-access); safeguarded “Tutor” largely mitigated harm — design of the help interface (not mere AI presence) determines whether help is a crutch.

Applied (FOUNDER BELIEF → testable): CTA copy that says “Get the answer / AI solves it / Finish my worksheet” recruits **executive** help-seeking. Copy that says “Unstick with a hint / Next step only / Then prove it in Practice” recruits **instrumental** help plus a gym handoff. Crunch week does not change the distinction — it increases the temptation to market the executive verb.

Kill: “Answer now” / Chegg-clone verbs as hero CTA.

Survive: Instrumental verbs + Bastani-style guardrails on the help rail.

Wound: Bastani classroom $\neq$ MindCraft parent Sunday night — still refuse Base-as-hero; instrument CRUNCHCTA-.*.

CXXX.5 Fear-appeal CTAs poison the dual rail

FACT: Putwain & Remedios (2014, *School Psychology Quarterly*, 29(4), 360–371, doi:10.1037/spq0000048) — teacher fear appeals before high-stakes exams (messages emphasizing failure consequences) relate to **lower self-determined motivation** and **lower exam performance**, partly mediated by motivation; scare tactics are not reliable motivational tools.

FACT: Putwain & Best (2011, *Learning and Individual Differences*, 21(5), 580–584, doi:10.1016/j.lindif.2011.07.007) — classroom fear appeals before tests associated with higher state anxiety and worse test performance in a counterbalanced design.

Applied (HYPOTHESIS): Parent/student CTAs of the form “ACT in 8 days — if you don’t open Solver you’ll freeze / fail” import classroom fear appeals into product chrome. Under SAFE-CRUNCHMSG / SAFE-SOFTMSG, crunch CTAs must be **challenge-structured** (named plan, two rails, one next action) — not threat-coupled dump. Fear copy also collapses the dual rail psychologically: Practice reads as “extra,” Solver reads as “survival.”

Kill: Countdown → dump as default creative.

Survive: “Exam week: unstick tonight *or* prove a grain alone — we protect thinking minutes either way.”

Wound: Putwain teacher speech ≠ push notification — still ban fear-hero CTAs pending CRUNCHCTA-QUAL.

CXXX.6 Competitive foil — what the buttons teach the market

Competitor pattern	CTA shape	MindCraft refusal
Chegg / Course Hero	Single executive “get solution”	No Chegg hero; instrumental + follow-on
ChatGPT Base	Unlimited paste / answer fluency	Bastani guardrails; no Base splash
Cram marketplaces	Fear countdown → pack dump	Dual rail + thinking-minutes frame
Purity identity apps	Hide help / shame deadline need	Keep labeled help-tonight rail
Streak apps	Tonight-% / day streak as CTA subcopy	FEI events, not tonight-%

FOUNDER BELIEF: The CTA is GTM. If paid acquisition lands on a dump button, MindCraft buys Chegg users and inherits Chegg churn when exams end.

CXXX.7 Claim table

Claim	Label	Confidence	Notes
Defaults shape high-stakes choices at scale	FACT	High	Johnson & Goldstein 2003
Choice architecture is never neutral	FACT	High	Thaler / Sunstein line
Executive vs instrumental help is a real distinction	FACT	High	Nelson-Le Gall 1981
Unguarded generative help can harm later solo performance	FACT	High	Bastani et al. 2025 PNAS

Fear appeals before exams can hurt motivation / performance	FACT	Medium–High	Putwain line
Dual labeled CTAs + non-dump default beat single Chegg hero on solo transfer <i>and</i> parent trust under crunch	HYPOTHESIS	Medium	Needs CRUNCHCTA-*
Instrumental help + mandatory follow-on Practice $\geq$ dump on delayed solo without raising quit	HYPOTHESIS	Medium	Densifies HWHELP-*
Parents will accept dual CTAs over “finish homework” hero when consequence copy is plain	HYPOTHESIS	Medium	Overlaps CRUNCHMSG / WTP CBC
Crunch week requires collapsing to one Solver hero for retention	SPECULATION	Low	Refuse pending evidence; contradicts HWHELP

CXXX.8 Product surface — SAFE-CRUNCHCTA claim contract

1. **Home / crunch card:** Two CTAs — Help tonight (instrumental subcopy) + Practice gym (prove-alone subcopy).
2. **Solver open:** No “show full solution” as first primary; fade/hint ladder; after unstuck, **one-tap Practice grain** for the same edge.
3. **Parent crunch message:** Dual-rail verbs; protected thinking minutes; never tonight-% stalk (SAFE-PDASH / CRUNCHMSG).
4. **Default rule (HYPOTHESIS):** If acuity = due-HW triage → help rail primary *with* follow-on Practice chip; if acuity = FEI / exam durable → Practice primary, help secondary. Never dump-only.
5. **Telemetry:** rail_chosen, cta_default_shown, followon_start, FEI events — demote Help Tonight Minutes.
6. **Copy bank (examples):** “Unstuck tonight $\neq$ ready for the test.” / “Help to move; Practice to prove.” / “Exam week we still show both doors.”

CXXX.9 Doctrine — SAFE-CRUNCHCTA (provisional)

1. **Two visible rails under crunch** — never silently drop Practice for exam month.
2. **Instrumental verbs only on help CTAs** — ban “get the answer / finish worksheet” heroes (Nelson-Le Gall; Bastani).

3. **Defaults are fiduciary** — dump must not be the no-action / primary enrollment path (Johnson & Goldstein; Thaler et al.).

4. **Follow-on Practice grain** after help when feasible — dual rail is a *path*, not two SKUs.

5. **No fear-appeal CTAs** — challenge structure over countdown dump (Putwain).

6. **Parent honesty** — dual-rail copy; no tonight-% / live stalk (SAFE-CRUNCHMSG / SAFE-SOFTMSG).

7. **No Dual-CTA Score™ / Help Tonight Minutes / Crunch Click NS.**

8. **Copy:** “Help tonight or prove it in Practice — crunch week still shows both, and help never means ‘we did it for you.’” Never Chegg hero or ACT guarantees from CTA packaging.

Confidence: High on method sources (defaults; executive/instrumental help; Bastani guardrails; fear-appeal harms). Medium on MindCraft-specific default primacy rules and follow-on compliance rates (needs CRUNCHCTA-*). High that single Chegg-hero / fear-dump / Dual-CTA Score™ / assisted⇒ready are commercially toxic under FEI + HWHELP.

CXXX.10 Experiments

ID	Question	Design	Primary
CRUNCHCTA-1	Dual CTA (instrumental help + Practice) vs single Chegg-hero Solver under crunch calendar	Student A/B	rail_chosen; followon_start; solo_transfer_pass
CRUNCHCTA-2	Practice-primary default vs help-primary default (both rails visible)	A/B	quit; dump rate; delayed solo
CRUNCHCTA-3	Mandatory follow-on Practice chip after Solver unstick vs optional	A/B	followon_start; retry_120s; help-abuse
CRUNCHCTA-4	Fear-countdown CTA vs challenge dual-rail CTA	Message A/B	threat appraisal; rail_chosen; retention
CRUNCHCTA-5	Parent CBC: dual-rail crunch card vs “we finish homework tonight” vs purity (no help)	CBC	WTP; trust; brand fit (SAFE-WTP / CRUNCHMSG)
CRUNCHCTA-QUAL	10 Mayas + 10 parents: did dual CTAs clarify or feel like product refusal?	Qual	CRUNCHCTA codebook

Falsifier: Single dump hero $\geq$ dual CTAs on delayed solo *and* parent WTP $\rightarrow$ still ban assisted $\equiv$ ready ads; investigate bank/fade quality before surrendering dual-rail brand.

Falsifier: Practice-primary raises quit under true due-HW acuity $\rightarrow$ keep dual visibility; switch default by acuity (rule 4), do not hide help.

Falsifier: Mandatory follow-on spikes help-avoidance $\rightarrow$ soften to prompted chip; do not return to dump-and-quit.

Pre-register: CRUNCHCTA-* before “AI homework helper / finish tonight” national creative.

Family note: CRUNCHCTA-* densifies HWHELP/CRUNCHMSG/HELP/DURABLE/SOFTMSG — do not fork a second Chegg religion.

CXXX.11 So what for MindCraft commercially

- **Copy:** “Help tonight or prove it in Practice — crunch week still shows both.” Subcopy: “Unstuck $\neq$ ready.”
 - **Product:** Dual CTA component on home / crunch / Solver; acuity-aware default; one-tap follow-on Practice grain; ban single dump primary.
 - **Growth:** Paid landers and parent emails use dual-rail verbs; kill Chegg-hero creative in exam season.
 - **Positioning:** Against Chegg dumps and purity apps; beside honest scarce craft (SAFE-CRUNCHMSG).
 - **Metric:** rail_chosen, cta_default_shown, followon_start, FEI — demote Help Tonight Minutes / Dual-CTA Score™.
 - **Kill list:** Chegg hero; fear-countdown dump; silent Practice collapse; assisted $\equiv$ ready; Dual-CTA Score™; ACT guarantees from CTAs.
 - **Vision:** Maya learns that deadline relief and identity proof are different doors — and MindCraft always labels which door she walked through.
-

References (verified)

- Bastani, H., Bastani, O., Sungu, A., Ge, H., Kabakc■, Ö., & Mariman, R. (2025). Generative AI without guardrails can harm learning: Evidence from high school mathematics. *Proceedings of the National Academy of Sciences*, 122(26), e2422633122. <https://doi.org/10.1073/pnas.2422633122>
 - Johnson, E. J., & Goldstein, D. (2003). Do defaults save lives? *Science*, 302(5649), 1338–1339. <https://doi.org/10.1126/science.1091721>
 - Nelson-Le Gall, S. (1981). Help-seeking behavior in learning. *Developmental Review*, 1(3), 224–246. [https://doi.org/10.1016/0273-2297\(81\)90019-8](https://doi.org/10.1016/0273-2297(81)90019-8)
 - Putwain, D., & Best, N. (2011). Fear appeals in the primary classroom: Effects on test anxiety and test grade. *Learning and Individual Differences*, 21(5), 580–584. <https://doi.org/10.1016/j.lindif.2011.07.007>
 - Putwain, D., & Remedios, R. (2014). The scare tactic: Do fear appeals predict motivation and exam scores? *School Psychology Quarterly*, 29(4), 360–371. <https://doi.org/10.1037/spq0000048>
 - Thaler, R. H., Sunstein, C. R., & Balz, J. P. (2013). Choice architecture. In E. Shafir (Ed.), *The behavioral foundations of public policy* (pp. 428–439). Princeton University Press. (Working paper / SSRN earlier as Thaler, Sunstein, & John.)
-

Part CXXXI — Figure Key-Verify Loops for Generated FormatId Stems

Chapter status: Living evidence + bank/ops brief — Researcher tick 2026-08-12 (UTC hour 12 ≡ Red Team slot, but ch131 never written → prefer Researcher; researcher count since synthesizer v1.16 = 5)

Primary question: When MindCraft generates FormatId stems (diagram, coordinate_graph, number_line), what **figure key-verify loop** must pass before ship, mastery writes, or hide-correctness diagnostic — without treating fluent SVG / pretty render / LLM-as-judge “looks fine” as keyed truth?

Owners: Engine (generation / figure verify) · Product (Practice / Diagnostic) · GENQ + FORMTRAP · Trust · Brand · Red Team

Commercial job: Ship **SAFE-FIGKEY** densifying SAFE-GENQ × SAFE-FORMTRAP × SAFE-DUAL × SAFE-COVER × SAFE-EXPLAINQA: **structural figure gates** — never fluency-pass, never item-count from unverified visuals, never Map writes from wrong figures.

Builds on: Parts LXXII, CXXVIII, XCVII, LXXXVII, CXXI. Seams: ml/generation/, Question.figure, --verify, FormatId tags, trap seals, gap-scan eligibility.

CXXXI.1 Why this chapter exists

SAFE-GENQ kills fluent wrong *text* keys (~30% internal drop; do not scale). SAFE-FORMTRAP kills decorative visual distractors. Underspecified: the **figure as ground truth** — wrong slope labels, contradictory side lengths, tick marks that disagree with the key — while the PNG “looks ACT-ready.”

Bad defaults: (1) fluency SVG pass, (2) LLM-as-judge alone, (3) decoration tagged diagram, (4) FormatId-count before verify, (5) Figure Key Score™ vanity, (6) hide-correctness on unverified figures (Map poison Maya cannot “re-check” like arithmetic).

FOUNDER BELIEF under audit: FormatId is a wedge only if figures are **constraint-checked** like text keys. Deterministic / IR verify is the spine; LLM proposes, not ships.

Claims we refuse: fluency SVG ≡ keyed; judge-alone ship; unverified figures as diagnostic truth; Figure Key Score™ / Visual QA Minutes; FormatId-count hero; soft-pass wrong geometry; ACT/coverage guarantees from figure packs.

CXXXI.2 Constructs (pretty ≠ keyed)

Construct	Meaning	MindCraft analogue	Failure mode
Figure key	Geometry/labels that uniquely entail the keyed option	Question.figure + correctAnswer joint check	Stem slope 2; graph shows height
Structural necessity	Pictorial element encodes the math object	FormatId-required figure	Decsheets tagged diagram

Stem–figure correspondence	Text matches plotted entities/labels	Contiguity + consistency gate	Split legend; contradicting lengths
IR / constraint check	Machine-checkable geometry	Rule pass on figure IR / solver	Eye-test only
Fluency SVG pass	Compiles / renders	Banned as sole gate	Wrong math, clean pixels
SAFE-FIGKEY	Figure key-verify doctrine	This chapter	Figure Key Score™

Operational definition (HYPOTHESIS): *SAFE-FIGKEY complete* when (a) figure is structurally necessary, (b) checkable representation exists (coordinates / constraints / ticks / option geometry), (c) deterministic checks pass stem–figure–key–trap seals, (d) fluency / “looks fine” / compile-success never sufficient, (e) diagnostic and mastery writes barred until (c), (f) batch figure-drop under explicit gate (GENQ \<~10% ambition; figures inherit the same hard stop).

CXXXI.3 Classical AIG: figures inside the model; review the procedure

FACT: Gierl & Lai (2012, *International Journal of Testing*, 12(3), 273–298, doi:10.1080/15305058.2011.635830) — AIG rests on **item models** specifying variable features — not free-form fluency.

FACT: Gierl, Lai, & Turner (2012, *Medical Education*, 46(8), 757–765, doi:10.1111/j.1365-2923.2012.04289.x) — cognitive model → item model → assembly; **auxiliary information** (graphs, figures) belongs in the item model.

FACT: Gierl & Lai (2016, *Educational Measurement: Issues and Practice*, 35(4), 6–20, doi:10.1111/emip.12129) — review generated items via **content and logic in the generation procedure** (math CCSS-linked illustration); do not eyeball thousands of stems as the QA unit.

Applied (HYPOTHESIS): Put figure constraints *inside* the item model; verify audits outputs against constraints — not “does this SVG look like an ACT graph?”

Kill: Prompt → pretty figure → ship. **Survive:** Model-constrained figures + procedure-level verify.

CXXXI.4 Validity is about use

FACT: Kane (2013, *Journal of Educational Measurement*, 50(1), 1–73, doi:10.1111/jedm.12000) — validate **interpretations and uses**; ambitious claims need more evidence; negative consequences can kill a use.

Use	Ambition	Figure-key risk if wrong
Practice with reveal + repair	Low–medium	Trust ding; SAFE-REPAIR
Format-gap / conversion Practice	Medium	Mis-routed FormatId remediation
Hide-correctness → Map / FEI	High	Silent false format weakness

Parent diagnosis / coverage ads	High	WTP + COVER burn
---------------------------------	------	------------------

Kill: “We have diagram items” ≡ validated FormatId diagnosis. **Survive:** Coverage browsing ≠ evidence use until FIGKEY pass.

CXXXI.5 Multimedia law as verify criteria

FACT: Mayer, Heiser, & Lonn (2001, *Journal of Educational Psychology*, 93(1), 187–198, doi:10.1037/0022-0663.93.1.187) — extraneous multimedia can impair understanding (coherence).

FACT: Moreno & Mayer (1999, *Journal of Educational Psychology*, 91(2), 358–368, doi:10.1037/0022-0663.91.2.358) — spatial contiguity; Ginns (2006, *Learning and Instruction*, 16(6), 511–525, doi:10.1016/j.learninstruc.2006.10.001) meta-analyzes contiguity support.

FACT: Haladyna, Downing, & Rodriguez (2002, *Applied Measurement in Education*, 15(3), 309–333, doi:10.1207/S15324818AME1503_5) — context-dependent chart/graph items assume the visual is intended and legible.

Applied (HYPOTHESIS): FIGKEY requires contiguous labels and stem–plot consistency. A beautiful off-by-one axis is a **failed key**.

Kill: Decsheets as “visual learning.” **Survive:** Coherence + contiguity as hard QA.

CXXXI.6 Models do not reliably see or author math diagrams

FACT (preprint): Zhang et al. (2024, arXiv:2403.14624, *MathVerse*) — MLLMs often fail true diagram use; some score ~5%+ *higher without* visuals — presence ≠ comprehension.

HYPOTHESIS: If multimodal judges cannot be trusted to *read* diagrams, they cannot be sole ship authority for *authoring* keyed FormatId figures.

SUPPORTING PREPRINTS (method hints): Kumar et al. (2025, arXiv:2511.08283, *DiagramIR*) — TikZ→IR + rules beat LLM-as-judge on human agreement; Wang et al. (2025, arXiv:2502.13855, *MagicGeo*) — constraint-solved coordinates beat unconstrained TikZ fluency.

Applied (FOUNDER BELIEF → testable): Prefer **spec** → **constraints** → **render** → **deterministic re-check** (+ SME sample) over **prompt** → **SVG** → **vibes**.

Kill: Multimodal “looks ACT-ready.” **Survive:** IR / coordinate / tick checks. **Wound:** MathVerse is comprehension, not authoring — still refuse judge-alone; no arXiv-as-ACT ads.

CXXXI.7 Claim table

Claim	Label	Confidence	Notes
AIG models include auxiliary figures; review targets generation logic	FACT	High	Gierl & Lai line
Score uses need evidence by ambition; bad keys kill high-stakes uses	FACT	High	Kane 2013
Extraneous / poorly integrated visuals hurt learning	FACT	High	Mayer; contiguity meta
Context-dependent visuals must be intentional and legible	FACT	High	Haladyna et al. 2002
MLLMs often fail true diagram use	FACT	Medium–High	MathVerse preprint
IR/constraint checks beat fluency SVG / judge-alone for ship gates	HYPOTHESIS	Medium	DiagramIR / MagicGeo; needs FIGKEY-*
Unverified FormatId generation will poison Map at scale	HYPOTHESIS	Medium–High	Densifies GENQ drop lesson
Pretty SVG pass is enough for practice	SPECULATION	Low	Refuse

CXXXI.8 Product surface — SAFE-FIGKEY claim contract

1. Generation emits a **checkable figure spec**, not only a raster.
2. **Deterministic verify** (and/or IR rules) before bank sync.
3. **Joint key**: $\text{stem} \wedge \text{figure} \wedge \text{correctAnswer} \wedge \text{distractor seals}$ — any break = drop.
4. Compile success / aesthetics / multimodal “looks fine” **never** sufficient.
5. Unverified figures never enter hide-correctness or mastery writes.
6. Telemetry: `figkey_pass`, `figkey_drop_reason`, `verified FormatId FEI` — demote Figure Key Score™.
7. Coverage tables keep `FormatId` holes until FIGKEY pass (SAFE-COVER).
8. Copy: “Every graph we count is constraint-checked — pretty isn’t keyed.”

CXXXI.9 Doctrine — SAFE-FIGKEY (provisional)

1. **Figures are keys** (GENQ).

2. **Structural necessity** (DUAL / FORMTRAP).
3. **Procedure-level verify** (Gierl & Lai 2016).
4. **Deterministic spine** — LLM proposes, rules dispose.
5. **No judge-alone ship.**
6. **Use gating** — diagnostic / Map / FEI only on FIGKEY-pass (Kane).
7. **No Figure Key Score™ / Visual QA Minutes / FormatId-count hero.**
8. Copy: “We don’t ship a graph we can’t re-check.”

Confidence: High on AIG review, Kane, coherence/contiguity. Medium on exact IR schema / drop targets (needs FIGKEY-*). High that fluency-pass / judge-alone / unverified diagnostic are commercially toxic.

CXXXI.10 Experiments

ID	Question	Design	Primary
FIGKEY-1	IR/constraint verify vs fluency-SVG-only	Offline audit	SME gold escape rate
FIGKEY-2	LLM-as-judge vs IR+rules	Agreement	κ vs SME; cost/item
FIGKEY-3	FIGKEY-pass-only vs mixed unverified FormatId practice	A/B sandbox	format-hop solo; trust incidents
FIGKEY-4	Bar hide-correctness for unverified figures vs allow	Policy A/B	Map corruption; repair tickets
FIGKEY-5	Parent CBC: constraint-checked vs “AI visual bank / N diagrams”	CBC	WTP; trust
FIGKEY-QUAL	Planted wrong graph: broken product or “hard item”?	Qual	FIGKEY codebook

Falsifiers: Fluency ≤ IR on SME bad keys → still prefer IR auditability; IR false-drop spike → SME appeal, keep use gating; judge-alone κ ≥ IR → keep deterministic REPAIR traces, ban judge-alone sole authority. Pre-register before diagram/graph scale or “visual ACT bank” creative.

CXXXI.11 So what for MindCraft commercially

- **Copy:** “Every graph we count is constraint-checked — pretty isn’t keyed.”
- **Product:** Figure spec + deterministic verify; figkey_pass diagnostic gate.
- **Growth:** Honest FormatId holes; kill “AI drew N diagrams.”

- **Positioning:** Against ChatGPT pretty-wrong geometry; beside classical AIG discipline.
 - **Metric:** figkey_pass / drop reasons / verified FormatId FEI.
 - **Kill list:** Fluency SVG; judge-alone; unverified diagnostic figures; FormatId-count hero; soft-pass wrong geometry; ACT guarantees from figure packs.
 - **Vision:** Maya's Map learns which representations she owns — only if the testing figure told the truth.
-

References (verified)

- Gierl, M. J., & Lai, H. (2012). The role of item models in automatic item generation. *International Journal of Testing*, 12(3), 273–298. <https://doi.org/10.1080/15305058.2011.635830>
 - Gierl, M. J., & Lai, H. (2016). A process for reviewing and evaluating generated test items. *Educational Measurement: Issues and Practice*, 35(4), 6–20. <https://doi.org/10.1111/emip.12129>
 - Gierl, M. J., Lai, H., & Turner, S. R. (2012). Using automatic item generation to create multiple-choice test items. *Medical Education*, 46(8), 757–765. <https://doi.org/10.1111/j.1365-2923.2012.04289.x>
 - Ginns, P. (2006). Integrating information: A meta-analysis of the spatial contiguity and temporal contiguity effects. *Learning and Instruction*, 16(6), 511–525. <https://doi.org/10.1016/j.learninstruc.2006.10.001>
 - Haladyna, T. M., Downing, S. M., & Rodriguez, M. C. (2002). A review of multiple-choice item-writing guidelines for classroom assessment. *Applied Measurement in Education*, 15(3), 309–333. https://doi.org/10.1207/S15324818AME1503_5
 - Kane, M. T. (2013). Validating the interpretations and uses of test scores. *Journal of Educational Measurement*, 50(1), 1–73. <https://doi.org/10.1111/jedm.12000>
 - Kumar, V., Mishra, S., Hao, R., Malik, R., Broman, D., & Demszky, D. (2025). *DiagramIR: An automatic pipeline for educational math diagram evaluation* (arXiv:2511.08283). <https://doi.org/10.48550/arXiv.2511.08283>
 - Mayer, R. E., Heiser, J., & Lonn, S. (2001). Cognitive constraints on multimedia learning: When presenting more material results in less understanding. *Journal of Educational Psychology*, 93(1), 187–198. <https://doi.org/10.1037/0022-0663.93.1.187>
 - Moreno, R., & Mayer, R. E. (1999). Cognitive principles of multimedia learning: The role of modality and contiguity. *Journal of Educational Psychology*, 91(2), 358–368. <https://doi.org/10.1037/0022-0663.91.2.358>
 - Wang, J., Zhang, T., Yu, H., Wang, J., & Huang, H. (2025). *MagicGeo: Training-free text-guided geometric diagram generation* (arXiv:2502.13855). <https://doi.org/10.48550/arXiv.2502.13855>
 - Zhang, R., Jiang, D., Zhang, Y., Lin, H., Guo, Z., Qiu, P., Zhou, A., Lu, P., Chang, K.-W., Gao, P., & Li, H. (2024). *MathVerse: Does your multi-modal LLM truly see the diagrams in visual math problems?* (arXiv:2403.14624). <https://doi.org/10.48550/arXiv.2403.14624>
-

Part CXXXII — Async Pre-Brief Cards While Waiting for a Late Tutor

Chapter status: Living evidence + waiting-room ops brief — Researcher tick 2026-08-12 (UTC hour 15; hour%6≠0; researcher count since synthesizer v1.16 = 6 → Researcher)

Primary question: When the tutor is late (or the student arrives early), how does MindCraft convert **idle lobby time** into a **student-owned forethought** / **Set card** that feeds SAFE-MICROBRIEF — without pour-open AI syllabus

dumps, entertainment waiting rooms, Wait Minutes vanity, shame-late theater, or minting a Prebrief Score™?

Owners: Product (session lobby / mission card) · HITL ops · Coach UX · SAFE-SRL / SAFE-MICROBRIEF · Brand · Red Team

Commercial job: Ship **SAFE-ASYNCBRIEF** densifying SAFE-MICROBRIEF × SAFE-SRL × SAFE-ATTN × SAFE-LOADSHED × SAFE-EXPECTANCY: a **capped async pre-brief** (≈60–180s) that occupies waiting with *session-related Set work* so the live open is a verification handoff, not a scramble — never idle lobby, never monologue-as-wait, never lobby theater as ACT ads.

Builds on: Parts CXXIX (SAFE-MICROBRIEF; MICROBRIEF-3 foreshadows this chapter), CI (SAFE-SRL), XCVI (SAFE-ATTN), LIV (SAFE-AAR), CXXIII (SAFE-LOADSHED), XLVI (SAFE-EXPECTANCY), XC (SAFE-INSTRUMENT). Product seams: late/ETA flag, lobby UI, mission/edge card, last soft-wrong / Map grain, Notes “Set” field, block-to-join optional acuity override.

CXXXII.1 Why this chapter exists

SAFE-MICROBRIEF protects the *live* open under a burned clock. Ops still lose minutes **before** anyone speaks: student sits in a spinner, scrolls TikTok, or gets an AI “warm-up lecture” that steals construction bandwidth. MICROBRIEF-3 already asked whether an async pre-brief beats idle wait; this chapter makes that product law.

Competitors resolve late-tutor waits wrongly: (1) **idle lobby** (spinner + apology); (2) **entertainment wait** (games unrelated to the edge); (3) **pour-open AI** (syllabus dump branded as prep); (4) **shame ETA** (“your tutor is late — again”); (5) **forced full Practice mission** that burns the FEI window before the human arrives. MindCraft’s wedge: **occupied, session-related Set** — the student names edge, success criterion, trap, and if-then plan so the tutor’s first 30s verifies rather than invents.

FOUNDER BELIEF under audit: Identity work needs aimed attempts. Waiting without Set is not “relaxing before hard work” — it is attention residue plus do-your-best entry.

Claims we refuse as doctrine:

1. Idle lobby / spinner-only as default late-tutor UX.
2. Tutor/AI lecture that *calls itself* pre-brief while the student waits.
3. Wait Minutes / Lobby Score™ / Prebrief Streak / Occupancy Score™ as North Star.
4. Entertainment-only wait (unrelated games) as primary fill.
5. Mandatory multi-item Practice wall that consumes the session before join.
6. Late-tutor shame copy aimed at the student (or public tutor-shame feed).
7. Vibes-only “get ready!” without named edge + success criterion + plan chip.
8. ACT / score guarantees from async pre-brief packaging.

CXXXII.2 Constructs (wait ≠ waste; prep ≠ pour)

Construct	Meaning	MindCraft analogue	Failure mode
-----------	---------	--------------------	--------------

Idle wait	Unoccupied preprocess time; attention on clock	Spinner lobby	Residue + anxiety
Occupied wait	Wait filled with activity (service psych)	Any lobby content	Unrelated = still empty for learning
Session-related Set	Forethought tied to <i>this</i> session's edge	Async pre-brief card	Trivia / lore fill
Async pre-brief	Student completes MICROBRIEF beats before live join	Cap 60–180s card	Forced marathon Practice
Handoff verify	Live open confirms/adjusts student's card	≤30–60s MICROBRIEF	Tutor ignores card → pour
SAFE-ASYNCBRIEF	Waiting-room forethought law	This chapter	Lobby Score™

Operational definition (HYPOTHESIS): A late/early wait is *SAFE-ASYNCBRIEF complete* when (a) lobby offers a **named time-boxed card** (default 90s; hard cap ≤180s unless student opts into Practice rail), (b) card requires **four Set beats**: named edge → success criterion → one likely trap / FormatId note → if-then plan chip (“When I see X, I will try Y”), (c) outputs write to the same mission/Notes object the live micro-brief reads, (d) fill activity is **session-related** (not unrelated entertainment), (e) live open **verifies** the card (≥1 student synthesis or tutor confirm) rather than replacing it with a lecture, and (f) success = $\text{asyncbrief_complete} \times \text{later brief_complete} / \text{time-to-first-FEI} / \text{retry_120s} / \text{solo_transfer_pass}$ — never Wait Minutes or Lobby Score™.

CXXXII.3 Occupied waits feel shorter — but fill must serve the session

FACT: Maister (1985, “The Psychology of Waiting Lines,” in Czepiel, Solomon, & Surprenant, *The Service Encounter*, Lexington Books / commonly circulated managerial essay) — propositions on wait *experience*: **occupied time feels shorter than unoccupied time**; preprocess waits feel longer than in-process waits; **uncertain waits feel longer than known, finite waits**; anxiety makes waits feel longer. Practical implication: fill activities should (a) offer benefit in themselves and (b) **relate to the forthcoming service** (menus before seating), not Muzak-style unrelated distraction.

Applied (HYPOTHESIS): Tutor-late lobby is classic **preprocess wait**. Spinner = unoccupied. Unrelated mini-game = occupied but **not session-related** — Maister’s sports-highlights vs Muzak distinction. The async pre-brief is the tutoring analogue of handing the menu: it shortens *perceived* wait **and** shortens live service setup (student already chose the edge).

Kill: Spinner-only default; entertainment-only fill as primary.

Survive: Session-related occupied wait with known ETA when ops can give one.

Wound: Service satisfaction ≠ transfer_pass — ASYNCBRIEF-* must pre-register FEI instruments, not claim Maister as ACT points.

CXXXII.4 Forethought quality guards against disorganisation in the gap

FACT: Zimmerman (2002, *Theory Into Practice*, 41(2), 64–70, doi:10.1207/s15430421tip4102_2) — forethought = processes/beliefs *before* learning efforts (task analysis, goals, strategic planning, motivational beliefs); reflections feed subsequent forethought.

FACT: Cosnefroy, Fenouillet, Mazé, & Bonnefoy (2018, *Issues in Educational Research*, 28(2), 329–348) — two undergraduate studies ($N = 378$; $N = 315$): high-quality **forethought-phase** processing predicted lower **procrastination** and **disorganisation**; disorganisation (not procrastination in their model) negatively predicted academic performance; authors treat weak forethought as a pathway into SRL failure.

FACT: Cleary & Zimmerman (2001, *Journal of Applied Sport Psychology*, 13(2), 185–206, doi:10.1080/104132001753149883) — experts vs non-experts/novices on free-throw practice: experts set **more specific goals**, chose **technique-oriented strategies**, made strategy attributions; forethought processes intercorrelated; self-reflection attributions predicted later strategy selection.

Applied (FOUNDER BELIEF → testable): Idle lobby is a **forethought vacuum** — Cosnefroy’s disorganisation risk plus Cleary/Zimmerman’s “vague goal” novice pattern. The async card forces **specific goal + technique plan** before the tutor arrives, so the live session does not open as reactive homework dump.

Kill: “Just hang out until they join” as care.

Survive: Microanalytic Set prompts (goal, strategy) on the wait card.

Wound: College self-report / sport microanalysis $\neq$ ACT identity — method transfer only.

CXXXII.5 Implementation intentions turn “get ready” into if-then control

FACT: Gollwitzer (1999, *American Psychologist*, 54(7), 493–503, doi:10.1037/0003-066X.54.7.493) — **implementation intentions** (“When situation x arises, I will perform response y !”) link anticipated cues to goal-directed responses; subordinate to goal intentions; help with getting started, shielding from distractions, and escaping bad habits by **delegating control to situational cues**.

Applied (HYPOTHESIS): The fourth Set beat is not journaling — it is an if-then chip tied to the named trap (“When I multiply both sides and see a negative coefficient, I will flip the inequality”). That chip becomes the tutor’s first coaching check and the student’s first stall cue (SAFE-HINT / SAFE-RETRIEVE), not a Prebrief Score™.

Kill: Vague “I’ll try my best” lobby text.

Survive: One if-then plan chip written before join.

Wound: Lab implementation-intention $d \neq$ guaranteed session FEI — ASYNCBRIEF-2 must test plan-chip present vs absent.

CXXXII.6 Advance organizers $\neq$ syllabus dumps

FACT: Ausubel (1960, *Journal of Educational Psychology*, 51(5), 267–272, doi:10.1037/h0046669) — advance introduction of relevant **subsuming concepts** (organizers) can facilitate learning/retention of unfamiliar meaningful verbal material versus controls lacking that scaffolding.

FACT (boundary): Advance organizers in Ausubel’s sense are higher-abstraction bridges to existing cognitive structure — **not** same-level overviews, not exhaustive lectures, and later reviews treat effects as conditional (material/learner fit), not a license for “more prep text always helps.”

Applied (HYPOTHESIS): The async card may show **one** Map-grain / last soft-wrong anchor (ideational scaffold) — not a chapter dump. Pour-open AI “while you wait, here’s everything about linear inequalities” violates Ausubel *and* SAFE-EXPLAIN / SAFE-ATTN (germane load + residue).

Kill: Waiting-room monologue / syllabus dump as “organizer.”

Survive: One subsuming anchor + student-generated Set beats.

Wound: 1960 prose-learning effects ≠ FormatId FEI — keep organizers thin.

CXXXII.7 Attention residue: idle phone waits tax the first attempt

FACT: Leroy (2009, *Organizational Behavior and Human Decision Processes*, 109(2), 168–181, doi:10.1016/j.obhdp.2009.04.002) — **attention residue:** when people switch incomplete tasks, attention remains with the prior task and can impair performance on the next; unfinished work leaves cognitive residue.

Applied (HYPOTHESIS): Phone-scroll lobby = unfinished social/feed task colliding with first FEI item (SAFE-ATTN). A short, completable Set card that **finishes** before join reduces residue relative to open-ended feed browsing. A forced long Practice mission that remains unfinished when the tutor joins *creates* residue — so async prep must be **completable**, not a second homework block.

Kill: “Keep them busy forever until tutor arrives.”

Survive: Completable ≤180s card; optional Practice rail only after card complete + explicit choice.

Wound: OB lab switching ≠ tutoring lobby — instrument time-to-first-FEI + first-item method-change.

CXXXII.8 Specific goals still beat “do your best” in the lobby

FACT: Locke & Latham (2002, *American Psychologist*, 57(9), 705–717, doi:10.1037/0003-066X.57.9.705) — specific, difficult goals outperform vague/easy/do-your-best goals; mechanisms include direction, effort, persistence, strategy arousal.

Applied (HYPOTHESIS): Lobby copy “prepare for tutoring” is do-your-best. Card fields that force “edge = ...; success = one solo transfer without peek” are the wait-time expression of SAFE-MICROBRIEF specificity — so the live 30s open is confirmation, not invention (Starmer-style: structure without duration inflation, method-only — see CXXIX).

Kill: Vague prepare-yourself lobby.

Survive: Named edge + observable success criterion before join.

CXXXII.9 Product rules (SAFE-ASYNCBRIEF)

1. **Default on late/early flag** — Show async pre-brief card when ETA slip or student-early; never spinner-only.
2. **Four Set beats** — Edge, success criterion, trap/FormatId, if-then plan chip (Gollwitzer).
3. **Hard time box** — Default ~90s; cap ≤180s; visible countdown; allow early submit.
4. **Session-related fill only** — Ban unrelated entertainment as primary; optional Practice rail *after* card complete.
5. **Write-through** — Card populates mission/Notes object MICROBRIEF reads; tutor UI shows student Set first.
6. **Live = verify, not replace** — Tutor confirms/adjusts; pour-open that ignores card = fidelity breach (HITL QA).
7. **ETA honesty** — Prefer known finite wait copy when possible (Maister); never student-facing tutor-shame feed.
8. **Expectancy-safe** — Task language; late is ops fact, not Maya character flaw (SAFE-EXPECTANCY).
9. **LOADSHED protect** — Under crunch, protect async Set + micro-brief; do not shed prep for HW dump.
10. **Copy:** “While we wait, name the edge and your if-then plan — so we start aimed.” Never “lobby games raise ACT” or Wait Minutes NS.

Confidence: High — Maister occupied/related wait + Zimmerman/Cosnefroy forethought + Gollwitzer if-then + Locke/Latham specificity + Leroy residue as mechanism sources. Medium — MindCraft 90–180s thresholds and verify-handoff compliance (needs ASYNCBRIEF-*). High — idle/pour/entertainment/Score™ lobby as commercially toxic under FEI + SAFE-MICROBRIEF/SRL.

CXXXII.10 Experiments

ID	Question	Design	Primary
ASYNCBRIEF-1	Async Set card vs idle spinner under equal tutor-late delay	Ops A/B	asyncbrief_complete; time-to-first-FEI; retry_120s
ASYNCBRIEF-2	Four-beat + if-then chip vs “get ready” free text ≤90s	Product A/B	plan_chip_ok; synthesis_ok; method-change
ASYNCBRIEF-3	Session-related Set vs unrelated entertainment wait	Ops A/B	perceived wait; first-item residue proxy; transfer_pass
ASYNCBRIEF-4	Verify-handoff micro-brief vs tutor ignores card / pour-open	HITL A/B	pour↓; brief_complete; solo_transfer_pass
ASYNCBRIEF-5	Parent/tutor CBC: “aimed wait card” vs “we’ll jump in when tutor arrives”	CBC	WTP; trust (SAFE-WTP / CRUNCHMSG)

ASYNCBRIEF-QUAL	10 Mayas + 10 tutors: did wait card feel clarifying or busywork/shame?	Qual	ASYNCBRIEF codebook
-----------------	---	------	------------------------

Falsifier: Idle equals async card on delayed transfer → still refuse pour-open; check whether bank/edge mismatch, not wait design, was bottleneck.

Falsifier: Entertainment wait wins satisfaction with equal FEI → keep related-fill rule for learning claim; do not sell satisfaction as mastery.

Falsifier: Card raises anxiety without method-change → tighten SAFE-ERRCLIMATE / shorten fields; do not return to spinner-only.

Pre-register: ASYNCBRIEF-* before “occupied wait / implementation-intention tutoring / +ACT” ads.

Family note: ASYNCBRIEF-* densifies MICROBRIEF/SRL/ATTN/LOADSHED/EXPECTANCY; do not collapse into Wait Minutes OKRs. MICROBRIEF-3 is absorbed into this family.

CXXXII.11 So what for MindCraft commercially

- **Copy:** “While we wait, name the edge and your if-then plan — so the session starts aimed.”
- **Product:** Late/early lobby → capped async Set card; write-through to micro-brief; tutor verify UI; optional Practice only after complete.
- **Ops:** QA asyncbrief_complete + pour-ignore rate; ETA honesty; protect under LOADSHED.
- **Positioning:** Against idle marketplace lobbies and AI warm-up dumps; high-reliability craft without waiting-room theater.
- **Metric:** asyncbrief_complete, plan_chip_ok, time-to-first-FEI, paired brief_complete / FEI — demote Wait Minutes / Lobby Score™.
- **Kill list:** Idle spinner default; pour-open wait; entertainment-primary fill; forced marathon Practice-before-join; late-shame copy; Prebrief Score™.
- **Growth / vision:** Trust packet shows waits become student-owned Set; Maya learns even delay can be aimed practice of forethought — identity starts before the tutor speaks.

References (verified)

- Ausubel, D. P. (1960). The use of advance organizers in the learning and retention of meaningful verbal material. *Journal of Educational Psychology*, 51(5), 267–272. <https://doi.org/10.1037/h0046669>
- Cleary, T. J., & Zimmerman, B. J. (2001). Self-regulation differences during athletic practice by experts, non-experts, and novices. *Journal of Applied Sport Psychology*, 13(2), 185–206. <https://doi.org/10.1080/104132001753149883>
- Cosnefroy, L., Fenouillet, F., Mazé, C., & Bonnefoy, B. (2018). On the relationship between the forethought phase of self-regulated learning and self-regulation failure. *Issues in Educational Research*, 28(2), 329–348.
- Gollwitzer, P. M. (1999). Implementation intentions: Strong effects of simple plans. *American Psychologist*, 54(7), 493–503. <https://doi.org/10.1037/0003-066X.54.7.493>

- Leroy, S. (2009). Why is it so hard to do my work? The challenge of attention residue when switching between work tasks. *Organizational Behavior and Human Decision Processes*, 109(2), 168–181. <https://doi.org/10.1016/j.obhdp.2009.04.002>
- Locke, E. A., & Latham, G. P. (2002). Building a practically useful theory of goal setting and task motivation: A 35-year odyssey. *American Psychologist*, 57(9), 705–717. <https://doi.org/10.1037/0003-066X.57.9.705>
- Maister, D. H. (1985). The psychology of waiting lines. In J. A. Czepiel, M. R. Solomon, & C. F. Surprenant (Eds.), *The service encounter* (pp. 113–123). Lexington Books.
- Zimmerman, B. J. (2002). Becoming a self-regulated learner: An overview. *Theory Into Practice*, 41(2), 64–70. https://doi.org/10.1207/s15430421tip4102_2

Part CXXXIII — One-Tap Solver→Practice Transfer Grain After Instrumental Help

Chapter status: Living evidence + Solver/Practice handoff brief — Researcher tick 2026-08-12 (UTC hour 18 ≡ Red Team slot, but ch133 never written → prefer Researcher per rotation; researcher count since synthesizer v1.16 = 7 → Researcher)

Primary question: After MindCraft delivers **instrumental** Solver help (unstick without executive dump), how does the product force a **one-tap, edge-matched Practice grain** that proves alone — without dump-and-quit, multi-step maze handoffs, Transfer Score™ vanity, or mistaking assisted completion for FEI?

Owners: Product (Solver close / Practice launch) · Coach UX · Brand / crunch CTA · SAFE-SCHED Returns · Parent honesty · Red Team

Commercial job: Ship **SAFE-XFERGRAIN** densifying SAFE-HWHELP × SAFE-CRUNCHCTA × SAFE-SCHED × SAFE-HELP × SAFE-FADE × SAFE-PROOF: the *path* after help is a **single-tap, same-edge Practice transfer grain** (near then delayed) — never optional buried “Practice sometime,” never Chegg-close, never Transfer Minutes as NS.

Builds on: Parts CIV (SAFE-HWHELP), CXXX (SAFE-CRUNCHCTA — foreshadows one-tap grain), LXXIV (SAFE-SCHED), LXXXVIII–LXXXIX (SAFE-FADE / SAFE-HELP), XCI (SAFE-HINT), LXXXVI (SAFE-PROOF), XC (SAFE-INSTRUMENT), XXXIII (Bastani / PWC). Product seams: Solver end-state CTA, followon_start, concept/ingredient/FormatId carry-through, solo_transfer_pass, Returns calendar chip.

CXXXIII.1 Why this chapter exists

SAFE-HWHELP and SAFE-CRUNCHCTA already ban becoming a nicer Chegg and require dual rails + follow-on *in principle*. Ops and UX still leak at the **last inch**: student gets unstuck, closes Solver, and never starts Practice — or must navigate concept pickers, level menus, and mission browsers until the deadline brain quits. That is **dump-and-quit with extra steps**.

Competitors resolve the inch wrongly: (1) **close-and-done** (“Homework finished!” toast); (2) **generic Practice home** with no edge carry-through; (3) **forced marathon** (10-item wall after every hint — backlash → help-avoidance); (4) **streak/XP bait** that celebrates assisted greens; (5) **Transfer Score™** dashboards. MindCraft’s wedge: **one tap → one**

(or few) **edge-matched FEI item(s)** with Solver closed or denied, then a SAFE-SCHED return — the commercial proof that help was instrumental, not executive cosplay.

FOUNDER BELIEF under audit: Instrumental help without a transfer grain is unfinished pedagogy. Identity updates from *solo evidence*, not from fluent unstick moments.

Claims we refuse as doctrine:

1. Dump-and-quit / “You’re done!” after Solver as default close.
2. Multi-step maze (re-pick topic → level → mission) as the only handoff.
3. Transfer Score™ / Handoff Minutes / Follow-on Streak as North Star.
4. Assisted Solver accuracy ≡ ready / exam-ready.
5. Forced marathon Practice wall after every soft hint (undermines help-seeking).
6. XP/streak celebration for help completion without solo grain.
7. Silent collapse of Practice rail under crunch (SAFE-CRUNCHCTA kill retained).
8. ACT / score guarantees from transfer-grain packaging.

CXXXIII.2 Constructs (help complete ≠ learning complete)

Construct	Meaning	MindCraft analogue	Failure mode
Instrumental help	Help limited to process so learner finishes <i>themselves</i>	Fade/hint Solver	Executive dump
Executive help	Other completes the product	Full solution paste	Tonight-% theater
Transfer grain	Short, edge-matched unassisted attempt after help	One-tap Practice card	Generic home
Near transfer probe	Same edge, new instance, tool closed	Immediate post-Solver item	Clone of helped item
Delayed return	Horizon-matched retrieval (SAFE-SCHED)	Returns chip after near pass	Cram-tonight only
One-tap handoff	≤1 primary action from Solver end → grain start	Primary CTA	Maze / buried link
SAFE-XFERGRAIN	Handoff law after instrumental help	This chapter	Transfer Score™

Operational definition (HYPOTHESIS): A Solver episode is *SAFE-XFERGRAIN complete* when (a) help was instrumental-framed (SAFE-HELP / SAFE-FADE), (b) end-state surfaces a **primary one-tap CTA** that launches Practice on the **same concept/ingredient/FormatId edge** without re-navigation, (c) the grain is **short by default** (1–3 items; opt-in extend), (d) at least one item requires **Solver closed / denied** (*solo_transfer_pass* eligible), (e) success telemetry is $\text{followon_start} \times \text{xfer_near_pass} \times \text{later xfer_return_due} / \text{FEI}$ — never Transfer Score™ or

assisted-accuracy NS, and (f) under crunch, help remains available but **cannot close without offering** the grain (prompted or soft-required — CRUNCHCTA-3 densified).

CXXXIII.3 Instrumental help is the transition — not the finish line

FACT: Nelson-Le Gall (1985, *Review of Research in Education*, 12, 55–90; ERIC ED275741 / chapter “Help-Seeking Behavior in Learning”) — reconceptualizes help-seeking as adaptive strategy; distinguishes **executive** (dependency-oriented: have someone attain the goal for you) from **instrumental** (mastery-oriented: obtain only the help needed to solve *yourself*). Instrumental help-seeking is framed as a mechanism of transition from other-regulation to self-regulation.

Applied (HYPOTHESIS): Solver that stops at “unstuck” without a self-regulation attempt leaves the Nelson-Le Gall transition incomplete. The one-tap grain is the product’s forced next beat of *self*-regulation — not a second SKU.

Kill: “Help complete = session complete.”

Survive: Instrumental help → immediate self-attempt on same edge.

Wound: 1980s classroom observational tradition ≠ ACT product A/B — method transfer only; XFERGRAIN-* must instrument solo events.

CXXXIII.4 Learners misuse on-demand help; design must teach the exit

FACT: Aleven, Stahl, Schworm, Fischer, & Wallace (2003, *Review of Educational Research*, 73(3), 277–320, doi:10.3102/00346543073003277) — selective review of help seeking in interactive learning environments: effective help-seeking relates to better learning, but learners often **underuse or misuse** help facilities; metacognitive skill × system design interactions matter.

FACT: Roll, Aleven, McLaren, & Koedinger (2011, *Learning and Instruction*, 21(2), 267–280, doi:10.1016/j.learninstruc.2010.07.004) — Help Tutor metacognitive feedback on help-seeking errors in Geometry Cognitive Tutor improved help-seeking behavior; Study 2 (N=67) found improved help-seeking skills **transferred** to new domain content in the following month when support was off. Earlier ITS2006/Help Tutor work showed behavior gains without automatic domain-score magic — do not overclaim.

Applied (FOUNDER BELIEF → testable): MindCraft must teach the **exit from help** as hard as the entry. A primary one-tap “Prove this alone” is metacognitive scaffolding in the Aleven/Roll sense — it is not nagware if the grain is short and edge-matched. Buried “Practice” links are the product analogue of unused help design: system *offers* the good path and students never take it.

Kill: Optional footer link to Practice home.

Survive: Primary CTA + telemetry on help-exit compliance.

Wound: Help Tutor ≠ ChatGPT Solver; Roll transfer is help-*skill* transfer, not automatic ACT points — keep claims at FEI grain.

CXXXIII.5 Smooth example→problem transitions beat abrupt pairs — and beat “help then nothing”

FACT: Renkl & Atkinson (2003, *Educational Psychologist*, 38(1), 15–22, doi:10.1207/s15326985ep3801_3) — cognitive-load account of structuring the **transition** from example study to problem solving; fading successively introduces problem-solving elements until learners solve independently.

FACT: Atkinson, Renkl, & Merrill (2003, *Journal of Educational Psychology*, 95(4), 774–783, doi:10.1037/0022-0663.95.4.774) — fading + principle self-explanation prompts produced medium-to-large effects on near and far transfer without extra time vs abrupt example–problem pairs in their experiments.

FACT: Kalyuga, Ayres, Chandler, & Sweller (2003, *Educational Psychologist*, 38(1), 23–31, doi:10.1207/s15326985ep3801_4) — **expertise reversal**: techniques that help novices (heavy guidance) can hurt more expert learners; guidance should fade as expertise grows.

Applied (HYPOTHESIS): Instrumental Solver help is a *worked/partial* scaffold. Ending without a problem-solving beat is an incomplete Renkl transition. Ending with a 20-item wall after a soft hint is expertise-blind flooding. SAFE-XFERGRAIN = **one short faded step**: completion-style or full attempt on a *new* instance of the same edge, then schedule returns — not dump, not marathon.

Kill: Help-then-nothing; help-then-marathon.

Survive: Short fade-consistent grain + SCHED return.

Wound: Lab fading ≠ homework crunch psychology — measure quit/help-avoidance (CRUNCHCTA falsifier retained).

CXXXIII.6 Assisted performance ≠ durable learning; solo transfer is the commercial proof

FACT: Bastani, Bastani, Sungu, Ge, Kabakc■, & Mariman (2025, *PNAS*, doi:10.1073/pnas.2422633122) — ~1,000 HS math students; unfettered GPT Base raised practice scores but harmed later **unassisted** exams; guarded Tutor largely neutralized harm but did not beat no-AI learning. Students used Base as a crutch; perception ≠ exam learning.

Already company law (XXXIII RT / SAFE-HWHELP): Guards ≠ gain; assisted ■ ready. This chapter densifies the **UX hinge**: without a one-tap solo grain, MindCraft cannot even *observe* Bastani’s unassisted outcome inside the product loop.

Kill: Solver-open accuracy as identity/readiness metric.

Survive: solo_transfer_pass as required proof class after instrumental help when feasible.

Wound: PNAS classroom GPT ≠ MindCraft Map-wrapped Solver — still the right commercial falsifier shape.

CXXXIII.7 Desirable difficulty after help: retrieve, don’t re-read the hint

FACT: Bjork & Bjork (2020, *Journal of Applied Research in Memory and Cognition*, 9(4), 475–479, doi:10.1016/j.jarmac.2020.09.003) — desirable difficulties (spacing, variation, testing/generation, reduced feedback) can slow acquisition performance while improving long-term retention/transfer; learning $\neq$ performance during instruction.

FACT (spacing spine, already SAFE-SCHED): Distributed practice supports longer retention horizons than massing; overnight “I just saw the hint” fluency is a performance illusion if no retrieval follows.

Applied (HYPOTHESIS): Re-opening the Solver transcript or re-reading the hint is **restudy**. The transfer grain is **retrieval/generation**. One-tap must land on an *attempt* surface with coach/Solver closed — not on a second explanation monologue (SAFE-EXPLAIN).

Kill: “Review the solution again” as primary post-help CTA.

Survive: Generate/attempt CTA + delayed return.

Wound: Desirable difficulty under high anxiety needs SAFE-FADEX protect-then-fade — do not flood HMA learners with hard solo on day one of an edge.

CXXXIII.8 Friction kills the handoff — one tap is pedagogy, not laziness

FACT: Johnson & Goldstein (2003, *Science*, 302(5649), 1338–1339, doi:10.1126/science.1091721) — defaults dominate outcomes; no-action paths impose switching costs (already SAFE-CRUNCHCTA).

FACT: Gollwitzer (1999, *American Psychologist*, 54(7), 493–503, doi:10.1037/0003-066X.54.7.493) — implementation intentions (“When x , I will y ”) automate goal pursuit at the cue.

Applied (HYPOTHESIS): Multi-step Practice navigation after help is a **hostile default** toward dump-and-quit. One-tap is the product’s implementation intention: *When instrumental help completes, I will attempt one solo item on this edge*. Optional buried links fail Gollwitzer because the cue (help-complete) is not bound to a single response.

Kill: “Find this topic in Practice” as handoff.

Survive: Pre-bound CTA carrying edge IDs into launchMissionDirect-style session.

Wound: One-tap convenience $\neq$ guaranteed transfer — still need item quality (SAFE-GENQ/FIGKEY) and closed-tool rules.

CXXXIII.9 Product rules (SAFE-XFERGRAIN)

1. **Primary CTA after instrumental unstick** — “Prove this alone (1–3 items)” as first button; demote “Done / Close.”
2. **Edge carry-through** — Pass conceptId + ingredient/FormatId + level recommendation; no re-picker maze.
3. **Short default grain** — 1–3 items; opt-in “keep going”; ban mandatory marathon after soft hints.
4. **Solo gate** — ≥ 1 item with Solver/AI reveal denied or closed (solo_transfer_pass eligible).
5. **Near $\neq$ clone** — New instance / isomorphic variant; not the identical helped stem.
6. **SCHEDULE chip** — After near pass (or honest miss), surface next SAFE-SCHED return; no “cleared forever.”

7. **Crunch policy** — Help allowed; close path must still *offer* grain (soft-require default; measure avoidance).

8. **Parent copy** — “Unstuck tonight → proved alone on the same edge” (SAFE-PDASH / CRUNCHMSG); never tonight-% NS.

9. **Telemetry** — followon_start, xfer_near_attempt, xfer_near_pass, xfer_return_due, FEI; demote Transfer Score™ / Handoff Minutes.

10. **Copy:** “You got unstuck — now prove the same edge alone.” Never “AI finished your homework” or ACT guarantees from handoff UX.

Confidence: High — Nelson-Le Gall instrumental transition + Aleven/Roll help-design + Renkl/Atkinson fade transition + Bastani unassisted harm + Bjork learning≠performance + Johnson/Gollwitzer friction/cue binding as method sources. Medium — MindCraft 1–3 item default and soft-require vs hard-gate under crunch (needs XFERGRAIN-*). High — dump-and-quit / maze / Transfer Score™ / assisted⇒ready as commercially toxic under FEI + HWHELP/CRUNCHCTA.

CXXXIII.10 Experiments

ID	Question	Design	Primary
XFERGRAIN-1	One-tap edge-matched grain vs dump-and-quit close after instrumental Solver	Student A/B	followon_start; solo_transfer_pass
XFERGRAIN-2	One-tap vs multi-step Practice navigation (same grain length)	UX A/B	followon_start; time-to-first-solo; quit
XFERGRAIN-3	Soft-require grain vs optional chip vs hard-gate wall under crunch	A/B/C	followon_start; help-avoidance; retry_120s
XFERGRAIN-4	1-item vs 3-item vs 10-item post-help grain	A/B/C	solo_transfer_pass; quit; hint_binge
XFERGRAIN-5	Near-only vs near+SCHED return chip	Product A/B	xfer_return_due completion; 7d MoC
XFERGRAIN-QUAL	10 Mayas: did one-tap feel fair proof or nag after help?	Qual	XFERGRAIN codebook

Falsifier: One-tap equals dump on delayed solo → check item quality / FormatId / anxiety flood before abandoning handoff law.

Falsifier: Hard-gate spikes help-avoidance → soften to prompted chip; do not return to dump-and-quit.

Falsifier: Near pass rises while 14–28d MoC flat → keep SCHED return; refuse near-fluency ads.

Pre-register: XFERGRAIN-* before “our AI handoff raises ACT” / Transfer Score™ marketing.

CXXXIII.11 So what for MindCraft commercially

- **Copy:** “You got unstuck — now prove the same edge alone.” Parent: “Help tonight, proof tonight — then a scheduled return.”
 - **Product:** Solver end-state → one-tap Practice grain with edge IDs; short default; solo gate; SCHED chip.
 - **Ops / analytics:** QA followon_start after instrumental episodes; demote assisted completion NPS.
 - **Positioning:** Against Chegg-close and ChatGPT “done”; high-reliability craft = help that graduates into evidence.
 - **Metric:** followon_start, xfer_near_pass, paired solo_transfer_pass / Returns — demote Transfer Score™ / Handoff Minutes.
 - **Kill list:** Dump-and-quit; maze handoff; marathon-after-hint; assisted⇒ready; XP-for-help-without-solo; ACT guarantees from handoff packaging.
 - **Growth / vision:** Acquisition can still sell unstuck; retention and brand defend **proof**. The Desk’s Solver vertical earns trust only when the next tap is Practice, not paste.
-

References (verified)

- Aleven, V., Stahl, E., Schworm, S., Fischer, F., & Wallace, R. (2003). Help seeking and help design in interactive learning environments. *Review of Educational Research*, 73(3), 277–320. <https://doi.org/10.3102/00346543073003277>
- Atkinson, R. K., Renkl, A., & Merrill, M. M. (2003). Transitioning from studying examples to solving problems: Effects of self-explanation prompts and fading worked-out steps. *Journal of Educational Psychology*, 95(4), 774–783. <https://doi.org/10.1037/0022-0663.95.4.774>
- Bastani, H., Bastani, O., Sungu, A., Ge, H., Kabakc■, Ö., & Mariman, R. (2025). Generative AI without guardrails can harm human learning. *Proceedings of the National Academy of Sciences*. <https://doi.org/10.1073/pnas.2422633122>
- Bjork, R. A., & Bjork, E. L. (2020). Desirable difficulties in theory and practice. *Journal of Applied Research in Memory and Cognition*, 9(4), 475–479. <https://doi.org/10.1016/j.jarmac.2020.09.003>
- Gollwitzer, P. M. (1999). Implementation intentions: Strong effects of simple plans. *American Psychologist*, 54(7), 493–503. <https://doi.org/10.1037/0003-066X.54.7.493>
- Johnson, E. J., & Goldstein, D. (2003). Do defaults save lives? *Science*, 302(5649), 1338–1339. <https://doi.org/10.1126/science.1091721>
- Kalyuga, S., Ayres, P., Chandler, P., & Sweller, J. (2003). The expertise reversal effect. *Educational Psychologist*, 38(1), 23–31. https://doi.org/10.1207/s15326985ep3801_4
- Nelson-Le Gall, S. (1985). Help-seeking behavior in learning. *Review of Research in Education*, 12, 55–90. (Also ERIC ED275741.)
- Renkl, A., & Atkinson, R. K. (2003). Structuring the transition from example study to problem solving in cognitive skill acquisition: A cognitive load perspective. *Educational Psychologist*, 38(1), 15–22. https://doi.org/10.1207/s15326985ep3801_3
- Roll, I., Aleven, V., McLaren, B. M., & Koedinger, K. R. (2011). Improving students’ help-seeking skills using metacognitive feedback in an intelligent tutoring system. *Learning and Instruction*, 21(2), 267–280. <https://doi.org/10.1016/j.learninstruc.2010.07.004>

Part CXXXIV — Hide-Correctness Eligibility Gates for FormatId Items

Chapter status: Living evidence + diagnostic/bank brief — Researcher tick 2026-08-13 (UTC hour 00 ≡ Red Team slot, but ch134 never written → prefer Researcher per rotation; researcher count since synthesizer v1.17 = 0)

Primary question: When may a FormatId item (diagram, coordinate_graph, number_line, and siblings) enter **C4 hide-correctness** gap-scan / cold-start probes that write Map or FEI — and when must it stay out even if the stem “looks shippable”?

Owners: Product (Diagnostic / gap-scan / Practice) · Engine (bank eligibility flags / GENQ verify) · Trust · Brand · Red Team

Commercial job: Ship **SAFE-HCELIG** densifying SAFE-FIGKEY × SAFE-GENQ × SAFE-CALIB × SAFE-COLD × SAFE-FORMAT × SAFE-COVER: an explicit **eligibility gate** — never “any tagged FormatId may diagnose,” never HideCorrectness Score™, never coverage theater that fills holes with unverified visuals.

Builds on: Parts L (SAFE-CALIB), LXXII (SAFE-GENQ), LXXXI (SAFE-COLD), LXXXII (SAFE-FORMAT), XCVII (SAFE-DUAL), CXXVIII (SAFE-FORMTRAP), CXXXI (SAFE-FIGKEY), CXXV (SAFE-CONFMISS).
Seams: C4 hide-correctness, figkey_pass, FormatId tags, gap-scan item pool, /record-outcomes / seed probes.

CXXXIV.1 Why this chapter exists

SAFE-FIGKEY answers *how* to verify figures. SAFE-GENQ answers *that* keys must be sealed. SAFE-COLD and SAFE-CALIB answer *why* hide-correctness exists: observe without reveal theater so confidence and Map updates stay honest.

Underspecified: the **admission policy**. Product temptation is to widen FormatId coverage by letting every tagged diagram/graph into gap-scan as soon as it renders. That collapses three different risks into one green “eligible” bit:

1. **Key poison** — wrong figure or answer silently writes false format weakness (GENQ / FIGKEY).
2. **Construct mix** — miss reflects graphics-decoding load, language, or decoration, not the FormatId skill the Map claims (Messick CIV; Diezmann/Lowrie graphicacy).
3. **Use overclaim** — browsing a pretty visual bank sold as calibrated format diagnosis (Kane use ambition).

Hide-correctness makes (1) and (2) worse, not better: Maya cannot “check the answer key” and self-correct the company’s lie. The commercial cost is Map distrust, false PawHub routes, and parent “diagnosis” that is really item-bank failure.

FOUNDER BELIEF under audit: FormatId is a wedge only if eligibility is **stricter** for hide-correctness than for reveal Practice. Pretty + tagged ≠ diagnostic-ready.

Claims we refuse:

1. Any FormatId-tagged item may enter C4 hide-correctness.

2. Fluency render / LLM “looks fine” $\equiv$ eligibility.
3. Coverage holes justify lowering the gate (“better a wrong graph than no graph”).
4. HideCorrectness Score™ / Eligibility Minutes / FormatId-count-in-scan as NS.
5. Miss on gated-out visuals $\equiv$ topic mastery without format/graphicacy labels.
6. ACT / identity guarantees from “visual diagnostic coverage.”

CXXXIV.2 Constructs

Construct	Meaning	MindCraft analogue	Failure mode
Hide-correctness (C4)	Record outcome; withhold key reveal in diagnostic	Gap-scan / cold probes	Answer-hunting; or silent Map poison
Eligibility gate	Binary (or graded) admit-to-use rule	hcelig_pass / pool filter	Soft “shipped $\Rightarrow$ eligible”
Figure key seal	Stem $\wedge$ figure $\wedge$ key $\wedge$ traps jointly true	SAFE-FIGKEY figkey_pass	Pretty wrong graph
Format construct purity	Miss mostly reflects intended FormatId skill	Format-tagged probe with sealed traps	Decoration; language/graphicacy confound unlabeled
Use ambition	What decision the score authorizes	Map format gap; FEI route; parent copy	Browsing $\equiv$ diagnosis
SAFE-HCELIG	Eligibility doctrine for C4 FormatId	This chapter	HideCorrectness Score™

Operational definition (HYPOTHESIS): A FormatId item is *hide-correctness eligible* iff (a) text/key/distractor seals pass SAFE-GENQ, (b) if a figure is required, SAFE-FIGKEY figkey_pass holds, (c) FormatId is structurally necessary (not decoration), (d) trap seals exist for visual miss types when options are visual (SAFE-FORMTRAP), (e) accessibility/alt path exists or the item is barred from ELL-blind interpretation as math-only (SAFE-ELL light), (f) the item is barred from mastery/Map writes until (a)–(d), and (g) coverage tables show honest holes rather than ineligible fillers counted as “visual diagnostic stock.”

CXXXIV.3 Uses are validated — not “items in a folder”

FACT: Kane (2013, *Journal of Educational Measurement*, 50(1), 1–73, doi:10.1111/jedm.12000) — validate **interpretations and uses**; more ambitious claims need more evidence; negative consequences can kill a use even when a weaker interpretation survives.

Use	Ambition	Eligibility bar
Practice with reveal + SAFE-REPAIR	Low–medium	GENQ seal; FIGKEY if figured; repair path

Format-hop Practice (revealed)	Medium	+ FormatId necessity; FORMTRAP if MC visual
C4 hide-correctness → Map / FEI	High	Full HCELG; no “looks fine” escape
Parent “we diagnosed format gaps”	High	Same as C4 + COVER honesty

Kill: “We have diagram items in the bank” ≡ validated hide-correctness FormatId diagnosis.

Survive: Same file may be Practice-eligible and C4-ineligible until seals complete.

CXXXIV.4 Construct-irrelevant variance is the FormatId tax

FACT: Messick (1989, in Linn, *Educational Measurement*, 3rd ed., pp. 13–103) — two primary threats are **construct underrepresentation** and **construct-irrelevant variance (CIV)**; scores can be too narrow *and* contaminated by extraneous difficulty.

Applied (HYPOTHESIS): Unverified or decorative FormatId stems inject CIV (reading the legend, pixel density, SVG glitches, beauty traps) into a probe marketed as “format gap.” Underrepresentation happens when gap-scan never admits honest FormatId items and the Map pretends format flexibility from symbolic-only probes.

Kill: Fill FormatId slots with any visual to avoid underrepresentation theater.

Survive: Prefer honest holes (SAFE-COVER) over CIV-contaminated greens; admit only sealed, necessary figures.

CXXXIV.5 Graphicacy is a real skill — label it or gate it

FACT: Diezmann & Lowrie (2009, *Assessment in Education: Principles, Policy & Practice*, 16(2), 131–147, doi:10.1080/09695940903075891) — information graphics in math assessments are not a single trivial overlay; students need measurable knowledge across graphic languages (graphs, maps, number lines, etc.); instruments exist precisely because graphics demand is often invisible in “math” scores.

FACT: Lowrie, Diezmann, & Kay (2011, *Evaluation and Research in Education*, 24(4), 285–296, doi:10.1080/09500790.2011.629723) — graphics-decoding proficiency can be screened as its own capacity; graphics-based tasks (number lines, column graphs, maps, pie charts) are common in large-scale assessments and are not automatically interchangeable.

Applied (HYPOTHESIS): Hide-correctness FormatId probes that write a *format* gap must either (i) minimize irrelevant graphicacy load relative to the intended FormatId, or (ii) label the miss type so Map/CONFMISS recovery does not route as pure topic weakness. Items that are primarily graphicacy screens belong in a different pool — not silent FormatId diagnosis.

Kill: “She missed the graph item ⇒ diagram FormatId gap” without seal + construct check.

Survive: Eligibility includes “what miss will we claim?”

CXXXIV.6 Context-dependent visuals and generation review

FACT: Haladyna, Downing, & Rodriguez (2002, *Applied Measurement in Education*, 15(3), 309–333, doi:10.1207/S15324818AME1503_5) — guidelines for classroom MC include treating context-dependent items (charts/graphs) as requiring intentional, legible stimulus material — not ornamental art.

FACT: Gierl & Lai (2016, *Educational Measurement: Issues and Practice*, 35(4), 6–20, doi:10.1111/emip.12129) — for generated items, review should target **content and logic in the generation procedure**, not eyeball-every-stem theater at scale.

FACT (preprint, method caution): Zhang et al. (2024, arXiv:2403.14624, *MathVerse*) — multimodal models often fail true diagram use; presence of a figure $\neq$ comprehension. Treat as **wound** for judge-alone eligibility, not as ACT advertising.

Applied (HYPOTHESIS): HCELIG is a *procedure flag* emitted by GENQ×FIGKEY×FORMTRAP verify — not a human “looks ACT-ready” checkbox at upload time. LLM-as-judge alone never flips hcelig_pass.

Kill: Upload PNG → tag diagram → gap-scan.

Survive: Procedure-level seals → eligibility bit → pool filter.

CXXXIV.7 Why hide-correctness tightens the gate (calibration hygiene)

FACT (calibration spine): Fischhoff, Slovic, & Lichtenstein (1977, *Journal of Experimental Psychology: Human Perception and Performance*, 3(4), 552–564, doi:10.1037/0096-1523.3.4.552) — people are routinely wrong when feeling certain; calibration is long-run match of stated confidence to accuracy.

Reuse (SAFE-CALIB / SAFE-COLD / SAFE-CONFMIS): Hide-correctness exists so confidence × outcome can update without answer-hunting. That only works if the keyed outcome is true. A wrong figure under C4 produces **confident false weakness** — the worst Map poison — and then CONFMIS recovery “fixes” a ghost.

HYPOTHESIS: Reveal Practice can sometimes surface student “this graph looks broken” reports (REPAIR). C4 cannot. Therefore C4 FormatId eligibility $\supseteq$ Practice eligibility, never the reverse.

Kill: Same eligibility bit for Practice browse and hide-correctness.

Survive: Graded uses; C4 is the strictest FormatId gate.

CXXXIV.8 Claim table

Claim	Label	Confidence	Notes
Score uses need evidence scaled to ambition; bad keys kill high-stakes uses	FACT	High	Kane 2013

CIV and underrepresentation threaten construct meaning	FACT	High	Messick 1989
Graphics decoding is measurable and often confounded with “math”	FACT	High	Diezmann & Lowrie 2009; Lowrie et al. 2011
Context-dependent visuals must be intentional/legible	FACT	High	Haladyna et al. 2002
Generated-item review should audit generation logic	FACT	High	Gierl & Lai 2016
MLLMs often fail true diagram use	FACT	Medium–High	MathVerse preprint
C4 FormatId needs stricter eligibility than reveal Practice	HYPOTHESIS	Medium–High	Densifies GENQ/FIGKEY/CALIB
Coverage pressure will push teams to lower the gate	FOUNDER BELIEF	Medium	Needs COVER + HCELIG-*
“Any tagged FormatId may diagnose” is safe if drop rate looks low	SPECULATION	Low	Refuse

CXXXIV.9 Product surface — SAFE-HCELIG claim contract

1. Emit explicit `hcelig_pass` (or equivalent) distinct from “in bank” and from Practice-eligible.
 2. **Require** GENQ seal + (if figured) FIGKEY pass + FormatId necessity + FORMTRAP seals when visual options exist.
 3. **Bar** hide-correctness / Map / FEI writes for `hcelig_pass = false`.
 4. Allow reveal Practice on a *subset* of sealed-but-not-yet-C4 items only with REPAIR telemetry — never silent C4.
 5. Coverage UI: show **FormatId diagnostic holes** honestly when items fail HCELIG (SAFE-COVER).
 6. Telemetry: `hcelig_pass`, `hcelig_block_reason`, C4 format-probe counts, false-weakness repair tickets — demote HideCorrectness Score™ / Eligibility Minutes.
 7. Copy: “We only hide the key when we’re sure the figure told the truth.”
 8. Gap-scan composition: prefer sealed FormatId probes; do not pad with symbolic items and claim format diagnosis.
-

CXXXIV.10 Doctrine — SAFE-HCELIG (provisional)

1. **Eligibility** ≠ **presence** — tagged FormatId in a folder is not C4 admission.
2. **C4 is the strictest use** (Kane ambition).
3. **Seals first** — GENQ × FIGKEY × FORMTRAP; judge-alone never admits.
4. **Construct honesty** — graphicacy/language confound labeled or gated (Messick; Diezmann/Lowrie).
5. **Honest holes** > **fake coverage** (SAFE-COVER).
6. **No HideCorrectness Score™ / Eligibility Minutes / FormatId-count-in-scan NS.**
7. **No ACT / identity guarantees** from visual diagnostic packaging.
8. Copy: “Pretty isn’t eligible.”

Confidence: High on Kane/Messick/Haladyna/Gierl/Diezmann–Lowrie as method sources. Medium on exact graded Practice-vs-C4 tiers and ELL/alt thresholds (needs HCELIG-*). High that unsealed C4 FormatId and coverage-padding are commercially toxic.

CXXXIV.11 Experiments

ID	Question	Design	Primary
HCELIG-1	C4 pool: HCELIG-pass-only vs mixed tagged FormatId	Policy A/B	Map↔tutor gold agreement; false-weakness rate
HCELIG-2	Practice-eligible / C4-ineligible tier vs dual-bar	Product A/B	Repair tickets; trust; format-hop solo
HCELIG-3	Honest FormatId holes in scan copy vs padded symbolic “coverage”	Copy A/B	Parent/student trust; abandon
HCELIG-4	Graphicacy-heavy items labeled vs silently written as FormatId gap	Engine A/B	Mis-route rate; CONFMISS recovery fit
HCELIG-5	Parent CBC: “constraint-checked visual diagnosis” vs “N diagram diagnostic items”	CBC	WTP; trust
HCELIG-QUAL	Planted wrong graph under C4; broken product or “hard format”?	Qual	HCELIG codebook

Falsifiers: HCELIG-pass-only equals mixed on false-weakness → audit seal quality before relaxing gate. Honest-holes copy raises abandon without trust gain → improve hole UX, do not pad with unsealed items. Graphicacy label confuses students → simplify miss taxonomy; keep gate. Pre-register HCELIG-* before “visual gap-scan” / FormatId-count

diagnostic campaigns.

CXXXIV.12 So what for MindCraft commercially

- **Copy:** “We only hide the key when we’re sure the figure told the truth.” Parent: “Format gaps we show were constraint-checked — we leave holes rather than guess.”
 - **Product:** `hcelig_pass` filter on gap-scan `FormatId` pool; stricter than Practice browse.
 - **Growth:** Kill “N diagram diagnostic items” hero; sell honest `FormatId` diagnosis.
 - **Positioning:** Against ChatGPT pretty-wrong visuals *and* against edtech that pads adaptive diagnostics with unverified art.
 - **Metric:** `hcelig_pass` / block reasons / C4 false-weakness — demote `HideCorrectness Score™`.
 - **Kill list:** Any-tagged-`FormatId`-in-C4; fluency eligibility; coverage padding; `HideCorrectness Score™`; ACT guarantees from visual diagnostic packs.
 - **Vision:** Maya’s first Map learns which representations she owns — only if the company refused to diagnose her with a lie.
-

References (verified)

- Diezmann, C. M., & Lowrie, T. (2009). An instrument for assessing primary students’ knowledge of information graphics in mathematics. *Assessment in Education: Principles, Policy & Practice*, 16(2), 131–147. <https://doi.org/10.1080/09695940903075891>
 - Fischhoff, B., Slovic, P., & Lichtenstein, S. (1977). Knowing with certainty: The appropriateness of extreme confidence. *Journal of Experimental Psychology: Human Perception and Performance*, 3(4), 552–564. <https://doi.org/10.1037/0096-1523.3.4.552>
 - Gierl, M. J., & Lai, H. (2016). A process for reviewing and evaluating generated test items. *Educational Measurement: Issues and Practice*, 35(4), 6–20. <https://doi.org/10.1111/emip.12129>
 - Haladyna, T. M., Downing, S. M., & Rodriguez, M. C. (2002). A review of multiple-choice item-writing guidelines for classroom assessment. *Applied Measurement in Education*, 15(3), 309–333. https://doi.org/10.1207/S15324818AME1503_5
 - Kane, M. T. (2013). Validating the interpretations and uses of test scores. *Journal of Educational Measurement*, 50(1), 1–73. <https://doi.org/10.1111/jedm.12000>
 - Lowrie, T., Diezmann, C., & Kay, R. (2011). The development of the graphics-decoding proficiency instrument. *Evaluation and Research in Education*, 24(4), 285–296. <https://doi.org/10.1080/09500790.2011.629723>
 - Messick, S. (1989). Validity. In R. L. Linn (Ed.), *Educational measurement* (3rd ed., pp. 13–103). Macmillan.
 - Zhang, R., Jiang, D., Zhang, Y., Lin, H., Guo, Z., Qiu, P., Zhou, A., Lu, P., Chang, K.-W., Gao, P., & Li, H. (2024). *MathVerse: Does your multi-modal LLM truly see the diagrams in visual math problems?* (arXiv:2403.14624). <https://doi.org/10.48550/arXiv.2403.14624>
-

Part CXXXV — Known-Finite ETA Honesty in Late-Tutor Lobby Copy

Chapter status: Living evidence + lobby-copy / ops brief — Researcher tick 2026-08-13 (UTC hour 03; hour%6≠0; researcher count since synthesizer v1.17 = 1 → Researcher)

Primary question: When a tutor is late (or running behind), what **wait-duration copy** must MindCraft show so the lobby is a *known, finite* wait — not “soon,” not optimistic underestimates, not apology theater, and not a student-facing shame feed — while the SAFE-ASYNCBRIEF Set card still runs?

Owners: Product (session lobby / ETA surface) · HITL ops / scheduling · Brand · Parent trust · SAFE-ASYNCBRIEF / SAFE-MICROBRIEF / SAFE-EXPECTANCY / SAFE-LOADSHED · Red Team

Commercial job: Ship **SAFE-KNOWNETA** densifying SAFE-ASYNCBRIEF × Maister uncertain-wait × Hui/Tse duration information × Munichor/Rafaeli progress-vs-apology: **honest finite ETAs (with refresh rules)** beat vague reassurance; never ETA Score™, never shame-the-tutor public board, never “only a few more minutes” cosplay.

Builds on: Parts CXXXII (SAFE-ASYNCBRIEF — rule 7 foreshadows this chapter), CXXIX (SAFE-MICROBRIEF), CXXIII (SAFE-LOADSHED), XLVI (SAFE-EXPECTANCY), LX (SAFE-REPAIR light), LIX (SAFE-WTP), CXXVII (SAFE-CRUNCHMSG). Seams: late flag, ETA field, lobby banner, async Set card countdown, tutor join event, parent “session started late” honesty.

CXXXV.1 Why this chapter exists

SAFE-ASYNCBRIEF answers *what the student does while waiting* (capped Set card). It left underspecified the **clock story** above the card: how long until the human arrives, how that number is stated, and what happens when the estimate slips again.

Ops and marketplace competitors resolve late waits wrongly:

1. **Vague “soon” / “almost here”** — open-ended anticipation (Maister’s nervous waiting room).
2. **Optimistic under-promise failure** — announce 2 minutes, deliver 12; honesty debt compounds (pilot “few more minutes”).
3. **Apology loops without numbers** — “sorry for the wait” every 30s with no duration (Munichor/Rafaeli: apologies can *worsen* reactions).
4. **Spinner with no ETA** — unoccupied *and* uncertain.
5. **Student-facing tutor-shame** — “your tutor is late again” as social theater.
6. **Fake precision** — countdown to the second when ops only know ±5 minutes.

MindCraft’s wedge: **known-finite when we can know; labeled uncertainty when we cannot; refresh without gaslighting; Set card always available**. Trust is a product feature — late minutes happen; lying about them is optional.

FOUNDER BELIEF under audit: Identity-safe tutoring cannot teach “own your mistakes” while the lobby gaslights the clock. Finite honesty is FEI-adjacent culture, not hotel niceness.

Claims we refuse as doctrine:

1. “Soon” / “any moment” as default late copy.
2. Systematic optimistic ETAs (announce short, deliver long).
3. Apology-only fillers without duration or progress.

4. ETA Score™ / Wait Satisfaction Score™ / Lobby Minutes NS.
5. Public or student-directed tutor-shame feeds.
6. Exact-second countdown theater when ops grain is coarse.
7. ACT / retention guarantees from “we manage waits scientifically.”

CXXXV.2 Constructs

Construct	Meaning	MindCraft analogue	Failure mode
Known-finite wait	Stated remaining duration with a bound the student can plan against	“About 8 minutes — card available”	Fake precision; silent slip
Uncertain wait	No usable bound; anticipation without settle	“Soon”; spinning dots	Anxiety; phone residue
Duration information	Explicit wait-length message	ETA banner	Underestimate trap
Progress information	Sense that service is advancing	“Tutor en route / joining” status	Fake progress bars
Appointment syndrome	Pre-slot wait tolerable; post-slot open-ended	Late after booked time	Ignore booked-time breach
SAFE-KNOWNETA	Honesty rules for lobby ETA copy	This chapter	ETA Score™

Operational definition (HYPOTHESIS): A late/early lobby is *SAFE-KNOWNETA complete* when (a) if ops can estimate remaining delay, the UI shows a **finite band** (e.g. “about 5–10 minutes”) not “soon,” (b) if ops cannot estimate, copy **labels uncertainty** (“we don’t have a precise arrival yet”) and offers the async Set card + optional reschedule, (c) ETA updates **never shorten without evidence** and never repeat “almost” after a miss without a new finite band, (d) apology text is **optional and secondary** to duration/progress, (e) no student-facing shame attribution, and (f) success instruments are $\text{eta_honesty_ok (stated band vs actual join within policy)} \times \text{asyncbrief_complete} \times \text{time-to-first-FEI / abandonment} \text{ — never Wait Minutes as learning NS.}$

CXXXV.3 Uncertain waits feel longer than known, finite waits

FACT: Maister (1985, “The Psychology of Waiting Lines,” in Czepiel, Solomon, & Surprenant, *The Service Encounter*, Lexington Books) — among core propositions: **uncertain waits are longer than known, finite waits**; the deepest wait anxiety is *how long*; telling a patient “thirty minutes” produces initial annoyance then acceptance, while “soon” produces nervous anticipation; the pilot who keeps saying “only a few more minutes” adds insult when the wait continues — customers are not only delayed but **not dealt with honestly**; the **appointment syndrome** makes post-appointment delays feel especially unbounded even when short.

Applied (HYPOTHESIS): Tutor-late after a booked slot is appointment syndrome. Spinner + “your tutor will be with you shortly” is Maister’s bad waiting room. A finite band (“about 8 minutes”) lets Maya settle into the Set card. Repeating “just a couple more minutes” after the first miss is the pilot insult — commercially toxic for a brand that sells trust and Map honesty.

Kill: “Soon” default; repeated “few more minutes” after missed ETA.

Survive: Finite band when estimable; honesty > optimism.

Wound: Service-feel $\neq$ solo_transfer_pass — KNOWNETA-* must not claim ACT points from ETA UX.

CXXXV.4 Uncertainty itself accumulates psychological cost

FACT: Osuna (1985, *Journal of Mathematical Psychology*, 29(1), 82–105, doi:10.1016/0022-2496(85)90020-3) — models the **psychological cost of waiting** as accumulated stress; stress intensity rises with time already spent *and* remaining-duration uncertainty; announcing service timing can reduce expected accumulated stress relative to opaque waits (model implications widely used in later queue psych).

Applied (HYPOTHESIS): Even before anger at “who is late,” unlabeled lobby time taxes affect. Finite ETA + completable Set card attacks both Osuna channels (waste *and* uncertainty) without inventing Calm Score™.

Kill: Opaque spinner as “neutral.”

Survive: Announce a usable bound when possible.

Wound: Formal stress integral $\neq$ teen tutoring field — use as mechanism, not formula in marketing.

CXXXV.5 Duration information beats empty reassurance — especially mid-length waits

FACT: Hui & Tse (1996, *Journal of Marketing*, 60(2), 81–90, doi:10.1177/002224299606000206) — experimental study of **waiting-duration information** vs **queuing information** across short / intermediate / long waits; acceptability of the wait and affective response mediate effects on service evaluation more than perceived duration alone; **neither information type helps much in short waits; duration information has greater impact than queuing information in intermediate waits** and relatively less in very long waits.

Applied (HYPOTHESIS): MindCraft late delays often sit in the *intermediate* band (roughly several–fifteen minutes) where **duration copy matters most**. Queue-position (“you are #2”) is weak analogue for 1:1 booked tutoring — prefer duration bands. For very long delays, duration info alone is insufficient: offer reschedule / LOADSHED path, not prettier ETA.

Kill: Skip ETA because “the Set card is enough.”

Survive: Duration band for intermediate late; escalate options when long.

Wound: Lab wait minutes $\neq$ school schedule chaos — calibrate bands on ops data (KNOWNETA-1).

CXXXV.6 Feedback on delay magnitude shapes attitudes — not only clock time

FACT: Larson (1987, *Operations Research*, 35(6), 895–905, doi:10.1287/opre.35.6.895) — speculative but highly cited OR forum: attitudes toward queues depend on more than delay length; **feedback regarding the likely magnitude of the delay** and queue environment influence customer attitudes and, often, market share; social justice (FIFO fairness) also matters; average wait is an incomplete performance metric because experience can be nonlinear in delay.

Applied (HYPOTHESIS): Marketplace tutoring that hides delay magnitude while selling “reliable sessions” burns share even if median lateness is modest. Fairness copy must not become student-vs-tutor blame; ops-facing FIFO / caseload justice stays internal (SAFE-LOADSHED / BURNWARN). Student UI gets **magnitude feedback**, not courtroom.

Kill: Average-lateness dashboards as the only ops truth while student sees “soon.”

Survive: Student-facing magnitude feedback + internal fairness ops.

Wound: Larson is forum-speculative — treat as design prior, confirm with KNOWNETA-.*.

CXXXV.7 Progress information beats apology fillers

FACT: Munichor & Rafaeli (2007, *Journal of Applied Psychology*, 92(2), 511–521, doi:10.1037/0021-9010.92.2.511) — telephone waits: Study 1 (real calls, $N = 123$) — **queue-location information** yielded lowest abandonment and most positive evaluations vs music; Study 2 ($N = 83$ simulated) — **sense of progress** mediated filler effects more than perceived waiting time; **apologies** as fillers can harm by re-salienting the wait without progress.

Applied (HYPOTHESIS): Lobby apology spam (“Sorry you’re waiting!”) without ETA/progress is Munichor-bad. Prefer: status that implies progress (“Tutor confirmed — about 6 minutes”) plus Set card. Do not replace duration with endless sorry. Do not invent fake progress bars that stall.

Kill: Apology-primary late UX.

Survive: Progress + duration; apology optional, once, factual.

Wound: Call-center queue $\neq$ video tutor join — map “progress” to real join states (en route / reconnecting / joined), not decorative %.

CXXXV.8 Affective mediation: anger and uncertainty, not minutes alone

FACT: Taylor (1994, *Journal of Marketing*, 58(2), 56–69, doi:10.1177/002224299405800205) — delayed airline passengers: delays affect service evaluations **mediated by negative affect** (anger, uncertainty); perceived provider **control** and **filled time** influence evaluations indirectly via those affects.

Applied (HYPOTHESIS): Late-tutor lobby must manage (1) uncertainty (KNOWNETA finite band), (2) filled time (ASYNCBRIEF), (3) control attribution carefully — “ops delay / connection” not “Maya wasn’t ready” and not student-aimed tutor-shame. Parent crunch copy stays on SAFE-CRUNCHMSG honesty rails.

Kill: Shame attribution; empty fill without duration.

Survive: Uncertainty↓ + session-related fill + non-character blame.

Wound: Airline delay ≠ ACT tutoring — mechanism transfer only.

CXXXV.9 Product rules (SAFE-KNOWNETA)

1. **Finite when estimable** — Show a band (“about N minutes” or “ N – M minutes”), never “soon / any moment / almost.”
2. **Label ignorance** — If no estimate, say so; offer Set card + reschedule / message tutor — do not fake a number.
3. **No optimistic default** — Prefer slightly conservative bands; under-promise/over-deliver on *arrival*, not on learning claims.
4. **Refresh without gaslight** — On slip: new finite band + timestamped update; ban repeating “just 2 more minutes” after miss.
5. **Apology secondary** — At most one factual apology; duration/progress primary (Munichor).
6. **Coarse grain honesty** — If ops only know ± 5 min, do not ship second-precision countdown theater.
7. **Pair with ASYNCBRIEF** — ETA banner never replaces the Set card; card countdown is independent of tutor ETA.
8. **No student shame feed** — Late is ops fact; no “your tutor failed you” public board; tutor coaching stays HITL-private.
9. **Long-wait escalate** — Beyond policy threshold (ops-set; test ~ 15 – 20 min), offer reschedule / LOADSHED path, not endless ETA theater.
10. **Copy:** “About 8 minutes — while you wait, finish your Set card.” Never “science-backed wait management raises ACT.”

Confidence: High — Maister finite-vs-uncertain + Hui/Tse duration info + Munichor progress>apology + Taylor affect mediation as design priors. Medium — exact band widths, refresh SLAs, long-wait threshold (KNOWNETA-*). High — soon/apology-only/shame-feed/ETA Score™ as commercially toxic under FEI + trust brand.

CXXXV.10 Experiments

ID	Question	Design	Primary
KNOWNETA-1	Finite band ETA vs “soon” under equal actual delay	Lobby A/B	abandonment; asyncbrief_complete; affect items; time-to-first-FEI
KNOWNETA-2	Conservative band vs optimistic under-estimate that slips	Ops A/B	trust items; rebook intent; eta_honesty_ok
KNOWNETA-3	Duration+progress vs apology-primary filler	Copy A/B	evaluation; abandonment; Set completion

KNOWNETA-4	Labeled uncertainty + reschedule offer vs fake ETA when unknown	Ops A/B	trust; complaint; return rate
KNOWNETA-5	Parent CBC: “honest late + Set card” vs “always on time” brand claim	CBC	WTP; trust (SAFE-WTP / CRUNCHMSG)
KNOWNETA-QUAL	10 Mayas + 10 parents: did ETA copy feel honest or shame/anxiety?	Qual	KNOWNETA codebook

Falsifier: Finite ETA equals “soon” on FEI and trust → keep honesty for brand ethics; investigate whether delays were too short (Hui short-wait null).

Falsifier: Optimistic ETA wins short-term satisfaction then tanks rebook → prefer conservative bands.

Falsifier: ETA honesty raises anxiety without Set uptake → tighten ERRCLIMATE / shorten banner; do not return to “soon.”

Pre-register: KNOWNETA-* before “we manage waits scientifically / zero-late brand / +ACT from lobby UX” ads.

Family note: Densifies ASYNCBRIEF / MICROBRIEF / EXPECTANCY / LOADSHED / CRUNCHMSG; do not invent Wait Minutes learning OKRs. ASYNCBRIEF-5 parent CBC may share instruments with KNOWNETA-5 — do not double-count.

CXXXV.11 So what for MindCraft commercially

- **Copy:** “About 8 minutes — while you wait, name the edge and your if-then plan.”
- **Product:** Lobby ETA band + refresh rules + labeled uncertainty state; paired ASYNCBRIEF; long-wait escalate.
- **Ops:** Instrument eta_honesty_ok (band vs join); ban optimistic default; coach tutors privately, not via student shame UI.
- **Positioning:** Against marketplace “soon” lobbies and apology muzak; high-reliability craft includes clock honesty.
- **Metric:** eta_honesty_ok, abandonment, asyncbrief_complete, rebook/trust — demote ETA Score™ / Lobby Minutes NS.
- **Kill list:** Soon-default; optimistic slip loops; apology-primary; shame feed; fake-precision countdown; ACT guarantees from wait UX.
- **Growth / vision:** Trust packet can show late minutes are disclosed like Map holes — the company that won’t lie about a graph won’t lie about eight minutes either.

References (verified)

- Hui, M. K., & Tse, D. K. (1996). What to tell consumers in waits of different lengths: An integrative model of service evaluation. *Journal of Marketing*, 60(2), 81–90. <https://doi.org/10.1177/002224299606000206>
- Larson, R. C. (1987). OR Forum—Perspectives on queues: Social justice and the psychology of queueing. *Operations Research*, 35(6), 895–905. <https://doi.org/10.1287/opre.35.6.895>
- Maister, D. H. (1985). The psychology of waiting lines. In J. A. Czepiel, M. R. Solomon, & C. F. Surprenant (Eds.), *The service encounter* (pp. 113–123). Lexington Books.
- Munichor, N., & Rafaeli, A. (2007). Numbers or apologies? Customer reactions to telephone waiting time fillers. *Journal of Applied Psychology*, 92(2), 511–521. <https://doi.org/10.1037/0021-9010.92.2.511>
- Osuna, E. E. (1985). The psychological cost of waiting. *Journal of Mathematical Psychology*, 29(1), 82–105. [https://doi.org/10.1016/0022-2496\(85\)90020-3](https://doi.org/10.1016/0022-2496(85)90020-3)
- Taylor, S. (1994). Waiting for service: The relationship between delays and evaluations of service. *Journal of Marketing*, 58(2), 56–69. <https://doi.org/10.1177/002224299405800205>

Part CXXXVI — Soft-Require vs Hard-Gate Follow-On Under Help-Avoidance Risk

Chapter status: Living evidence + Solver→Practice gate brief — Researcher tick 2026-08-13 (UTC hour 06 ≡ Red Team slot, but ch136 never written → prefer Researcher per rotation; researcher count since synthesizer v1.17 = 2 → Researcher)

Primary question: After instrumental Solver help, should MindCraft **soft-require** the SAFE-XFERGRAIN Practice grain (primary one-tap CTA; Close demoted but available) or **hard-gate** it (cannot dismiss without starting the grain) — and when does hard-gating spike **help-avoidance** / **quit** so badly that the brand loses the next ask?

Owners: Product (Solver close / Practice handoff) · Coach UX · Brand / crunch CTA · SAFE-HELP / SAFE-XFERGRAIN · Parent honesty · Red Team

Commercial job: Ship **SAFE-SOFTGATE** densifying SAFE-XFERGRAIN × SAFE-HELP × SAFE-CRUNCHCTA × Deci/Ryan autonomy × Brehm reactance: **soft-require is the default gate**; hard-gate is a *conditional, evidence-triggered* escalation with quit/avoidance falsifiers — never Transfer Score™, never “must prove or you can’t leave” shame walls, never ACT guarantees from gate packaging.

Builds on: Parts CXXXIII (SAFE-XFERGRAIN — rule 7 foreshadows soft-require; XFERGRAIN-3 owns A/B/C), LXXXIX (SAFE-HELP — abuse *and* avoidance), CXXX (SAFE-CRUNCHCTA), XLIV (SAFE-REWARD / controlling rewards), CIV (SAFE-HWHELP), XC (SAFE-INSTRUMENT). Seams: Solver end-state, followon_start, Close/Skip affordance, help_avoid_after_gate, solver_open_rate, quit after instrumental episode.

CXXXVI.1 Why this chapter exists

SAFE-XFERGRAIN settled the *path*: one-tap, edge-matched, short Practice grain after instrumental help. It left the *force* underspecified. Ops and growth instincts split:

1. **Optional chip** — “Practice when ready” footer → dump-and-quit returns (hostile default).
2. **Soft-require** — primary “Prove this alone (1–3)” with demoted Close / “Not now — schedule return.”

3. **Hard-gate** — modal wall; dismiss disabled until grain starts (or completes).
4. **Marathon hard-gate** — cannot leave until 10 solo items (already killed by XFERGRAIN).

Hard-gate maximizes short-run `followon_start`. It also risks **reactance** (freedom threat → boomerang: skip Solver next time, binge executive help elsewhere, or quit). Soft-require preserves autonomy language while keeping the pedagogical default honest. Help-avoidance research warns: students who already feel threatened by asking for help will treat a forced “prove you deserved the hint” wall as competence shame — the opposite of SAFE-HELP’s invite path.

FOUNDER BELIEF under audit: Instrumental help without a transfer grain is unfinished pedagogy — but unfinished pedagogy fixed by *controlling* walls trains Maya to avoid the help door. Soft-require is the fiduciary default; hard-gate is a clinical escalation with kill switches.

Claims we refuse as doctrine:

1. Hard-gate wall as universal default after every soft hint.
2. Optional buried chip as the only “respectful” close (dump-and-quit cosplay).
3. Gate Score™ / Forced Follow-On % / Soft-Require Compliance Minutes as North Star.
4. “No Close until you prove” shame copy after help.
5. Ignoring help-avoidance / Solver-skip as falsifiers when `followon_start` rises.
6. Marathon hard-gates (already killed).
7. ACT / retention guarantees from soft- vs hard-gate packaging.

CXXXVI.2 Constructs

Construct	Meaning	MindCraft analogue	Failure mode
Soft-require	Primary CTA + demoted exit; exit still one tap	“Prove alone” primary; “Not now” secondary	Exit too loud → dump
Hard-gate	Cannot proceed/dismiss without starting (or finishing) grain	Modal lock until <code>followon_start</code>	Reactance; help-avoid
Optional chip	Non-primary Practice link	Footer / toast	Dump-and-quit
Help-avoidance	Need help but under-ask (threat/cost)	Skip Solver; freeze without hint	Silent struggle theater
Freedom threat	Perceived elimination of valued exit	“You must Practice now”	Boomerang quit
SAFE-SOFTGATE	Soft-require default; hard-gate conditional + falsifiers	This chapter	Gate Score™

Operational definition (HYPOTHESIS): A post-help close is *SAFE-SOFTGATE complete* when (a) **default** is soft-require (primary one-tap grain; Close/Not-now demoted but available), (b) hard-gate is allowed only under **named**

evidence triggers (e.g. repeated dump-and-quit on same edge *and* no rising help_avoid_after_gate / Solver-skip), (c) hard-gate still uses **short grain** (1–3) never marathon, (d) copy is autonomy-supportive (“ready to prove this edge alone?”) not controlling (“you can’t leave until…”), (e) success co-primaries are followon_start × xfer_near_pass × solver_open_rate / **help-recruit** (avoidance falsifier) — never Forced Follow-On % alone, and (f) parent copy never shames “skipped proof.”

CXXXVI.3 Soft-require is a nudge; hard-gate is a mandate

FACT: Thaler & Sunstein (2008, *Nudge*, Yale University Press) — a nudge alters behavior in a predictable way **without forbidding options** or significantly changing economic incentives; “Nudges are not mandates. Putting the fruit at eye level counts as a nudge. Banning junk food does not.”

FACT: Johnson & Goldstein (2003, *Science*, 302(5649), 1338–1339, doi:10.1126/science.1091721) — defaults dominate outcomes; opt-out architectures raise uptake vs opt-in without removing choice (already used in SAFE-CRUNCHCTA / SAFE-XFERGRAIN).

Applied (HYPOTHESIS): Soft-require = fruit at eye level (primary CTA) with junk food still on the shelf (Close). Hard-gate = ban. MindCraft may *sometimes* need a mandate for fiduciary reasons (Bastani unassisted harm), but calling every mandate a “nudge” is brand fraud. Default soft-require inherits Johnson/Goldstein power *without* Brehm-level freedom elimination.

Kill: Marketing hard-gates as “gentle nudges.”

Survive: Soft-require as default choice architecture.

Wound: Organ-donation defaults ≠ teen homework psychology — measure quit/help-avoid in product, not cite consent rates as FEI.

CXXXVI.4 Controlling language and hard walls arouse reactance

FACT: Brehm (1966, *A Theory of Psychological Reactance*, Academic Press) — threats to or elimination of valued free behaviors arouse reactance, a motivational state to restore freedom (boomerang, source derogation, opposite behavior).

FACT: Ball & Goodboy (2014, *Communication Education*, 63(3), 192–209, doi:10.1080/03634523.2014.918634) — college classroom experiment (N=206): instructor messages using **forceful** (and/or unclear) language raised perceived freedom threat → psychological reactance → intentions toward instructional dissent and challenge behavior.

Applied (HYPOTHESIS): “You must complete Practice before closing Solver” is forceful language + freedom elimination. Expected restoration moves in product: rage-quit, never open Solver again, paste ChatGPT off-platform (executive help displacement), or speed-click the grain without sense-making (try-step abuse — SAFE-HELP). Soft-require with rationale (“same edge, one short proof so help stays instrumental”) lowers threat while keeping the pedagogical cue.

Kill: Controlling “can’t leave” copy as default.

Survive: Clear, autonomy-framed soft-require; hard-gate only with measured reactance falsifiers.

Wound: College communication lab $\neq$ ACT HS product — SOFTGATE-* must instrument Solver-skip and off-platform help proxies where possible.

CXXXVI.5 Autonomy support beats controlling pressure for quality motivation

FACT: Deci & Ryan (2000, *Psychological Inquiry*, 11(4), 227–268, doi:10.1207/S15327965PLI1104_01) — Self-Determination Theory: social contexts that support autonomy, competence, and relatedness facilitate internalization and better quality motivation; controlling contexts that pressure toward specific outcomes undermine intrinsic motivation and forestall integration.

FACT (already company-adjacent, SAFE-REWARD / XLIV): Controlling rewards and pressure can crowd out autonomy — relevant when “forced Practice minutes” become the payment for having received a hint.

Applied (FOUNDER BELIEF $\rightarrow$ testable): Framing the grain as *Maya’s chance to own the win* (competence + autonomy) beats framing it as *the tax for getting help* (controlling contingency). Hard-gate that feels like a fine for asking will raise help-avoidance among students who already see help as costly (next section). Soft-require + SCHED “Not now $\rightarrow$ return chip” preserves relatedness to the Desk without pretending dump-and-quit is fine.

Kill: “Hint tax” / punishment framing for follow-on.

Survive: Autonomy-supportive proof framing + competence evidence (SAFE-PROOF).

Wound: SDT is broad; do not claim soft-require raises ACT points — claim better help-recruit + solo co-primary balance.

CXXXVI.6 Help-avoidance is the commercial falsifier hard-gates often ignore

FACT: Ryan & Pintrich (1997, *Journal of Educational Psychology*, 89(2), 329–341, doi:10.1037/0022-0663.89.2.329) — early adolescents in math (N=203): perceived **threats** and **benefits** of help-seeking shape avoidance; threats/benefits partially mediate goal and competence effects on avoidance; only benefits predicted adaptive help-seeking. Students unsure of competence/social standing feel more threatened when asking.

FACT: Karabenick & Knapp (1991, *Journal of Educational Psychology*, 83(2), 221–230, doi:10.1037/0022-0663.83.2.221) — academic help-seeking relates to other instrumental achievement behaviors; avoidance and executive patterns are distinct from instrumental seeking (spine already in SAFE-HELP).

FACT: Aleven et al. (2016, *IJAIED*, 26(1), 205–223, doi:10.1007/s40593-015-0089-1) — help helps, but only so much; learners both abuse and avoid help; design that only punishes misuse can worsen under-use.

Applied (HYPOTHESIS): If hard-gate raises followon_start while solver_open_rate **falls** or students report “I don’t ask because then I’m trapped in Practice,” the gate failed SAFE-HELP. Commercial identity companies cannot optimize forced proof by shrinking the help door. Soft-require + binge detectors (abuse side) + help-invite (avoidance side) is the balanced stack; hard-gate is a scalpel for chronic dumpers who still *open* Solver, not a default for anxious avoiders.

Kill: Follow-on compliance as sole success when help-recruit collapses.

Survive: Co-primary `followon_start` × help/Solver openness.

Wound: Survey threat ≠ product telemetry — combine SOFTGATE-QUAL with event rates.

CXXXVI.7 Defaults work — but defaults ≠ outcomes; don't worship Forced %

FACT: Research distinguishing default effects on *choices* vs *downstream outcomes* (e.g. defaults raising registration without guaranteeing later completion; see behavioural-public-policy discussions such as default effects on exam registration vs exam-taking) — uptake of an offered path is not durable learning.

Already company law: Bastani et al. (2025, *PNAS*, doi:10.1073/pnas.2422633122) — assisted practice fluency ≠ unassisted exam learning. SAFE-XFERGRAIN requires solo grain *quality*, not merely modal compliance.

Applied (HYPOTHESIS): Hard-gate can inflate `followon_start` with rage-clicks and empty attempts. Soft-require may yield fewer starts but cleaner `xfer_near_pass` / `retry_120s`. Prefer quality of solo evidence over Forced Follow-On %. If soft-require yields near-zero starts, escalate *friction on Close* (confirm + SCHED chip) before full hard-gate — a middle rung.

Kill: Gate Score™ / Forced % as learning NS.

Survive: Soft-require → confirm-on-close → conditional hard-gate ladder.

Wound: Need SOFTGATE-1...3 before declaring ladder universal.

CXXXVI.8 Product rules (SAFE-SOFTGATE)

1. **Default = soft-require** — Primary “Prove this alone (1–3)”; demote Close / “Not now.”
2. **Never optional-only chip** as the sole post-help affordance (dump-and-quit).
3. **Autonomy copy** — “Ready to prove the same edge alone?” Never “You can't leave until you Practice.”
4. **Not-now path** — Always pairs with SAFE-SCHED return chip (honest deferral ≠ dump).
5. **Hard-gate is conditional** — Trigger only on repeated dump-and-quit *and* healthy Solver/help openness; never day-one universal hard-gate.
6. **Hard-gate grain stays short** — 1–3 items; start-gate preferred over finish-gate; ban marathon locks.
7. **Falsifiers stop escalation** — Rising `help_avoid_after_gate`, falling `solver_open_rate`, QUAL “trapped after hint” → roll back to soft-require + confirm-on-close.
8. **Abuse ≠ avoidance** — Keep binge/executive detectors (SAFE-HELP); do not “fix” avoidance with harder walls.
9. **Telemetry** — `followon_start`, `xfer_near_pass`, `close_without_followon`, `help_avoid_after_gate`, `solver_open_rate`; demote Gate Score™ / Forced %.
10. **Parent copy** — “Unstuck → invited to prove alone”; never shame “skipped proof” streaks on parent dash (SAFE-PDASH / SOFTMSG).

Confidence: High — Thaler/Sunstein nudge≠mandate + Brehm/Ball–Goodboy reactance + Deci/Ryan autonomy + Ryan/Pintrich help-threat + Alevén help-avoidance as method sources for soft-require default. Medium — exact hard-gate trigger policy and confirm-on-close middle rung (SOFTGATE-*). High — universal hard-gate / Gate Score™ / hint-tax shame / ignoring avoidance falsifiers as commercially toxic under FEI +

CXXXVI.9 Experiments

ID	Question	Design	Primary
SOFTGATE-1	Soft-require vs optional chip vs hard-gate after instrumental Solver	A/B/C	followon_start; xfer_near_pass; solver_open_rate; quit
SOFTGATE-2	Soft-require vs confirm-on-close vs start-hard-gate (short grain)	A/B/C	followon_start; help-avoid; retry_120s
SOFTGATE-3	Autonomy-framed vs controlling-framed soft-require copy (same layout)	Copy A/B	followon_start; QUAL threat; Solver reopen 7d
SOFTGATE-4	Conditional hard-gate (repeat dumpers only) vs universal soft-require	Policy A/B	followon among dumpers; avoidance among non-dumpers
SOFTGATE-5	Parent CBC: “invited proof” vs “must prove after help” vs dump-and-done	CBC	WTP; trust (SAFE-WTP)
SOFTGATE-QUAL	10 Mayas: did hard-gate feel fair or like a tax on asking?	Qual	SOFTGATE codebook

Falsifier: Hard-gate raises delayed solo *and* help-recruit equally with no QUAL threat → allow conditional hard-gate more broadly; still ban Gate Score™ and controlling copy.

Falsifier: Soft-require ≈ optional chip on followon_start → strengthen primary CTA / confirm-on-close before jumping to hard-gate.

Falsifier: Soft-require wins followon but loses 14–28d MoC → keep SCHED return; refuse soft-require=durable ads.

Pre-register: SOFTGATE-* (absorbs densified XFERGRAIN-3) before “we force practice after every hint” growth campaigns.

Family note: Ties XFERGRAIN/HELP/CRUNCHCTA/HWHELP/PROOF/INSTRUMENT/REWARD.

CXXXVI.10 So what for MindCraft commercially

- **Copy:** “You got unstuck — ready to prove the same edge alone?” Secondary: “Not now — we’ll bring this edge back.” Never “No exit until Practice.”
 - **Product:** Soft-require default; confirm-on-close middle; conditional short hard-gate only with avoidance falsifiers green.
 - **Ops / analytics:** Dashboard followon_start *beside* solver_open_rate / help_avoid_after_gate; demote Forced %.
 - **Positioning:** Against Chegg-close *and* against compliance-app walls; high-reliability craft = proof invitation that does not punish asking.
 - **Metric:** Co-primary follow-on × help openness × solo pass — demote Gate Score™.
 - **Kill list:** Universal hard-gate; optional-only chip; Gate Score™; hint-tax shame; ignore avoidance when Forced % rises; ACT guarantees from gate packaging.
 - **Growth / vision:** Acquisition can sell unstuck; retention defends **willing** return to the help door. The Desk teaches Maya to recruit support *and* to prove alone — without making proof the fine for having asked.
-

References (verified)

- Aleven, V., McLaren, B. M., Roll, I., & Koedinger, K. R. (2016). Help helps, but only so much: Research on help seeking with intelligent tutoring systems. *International Journal of Artificial Intelligence in Education*, 26(1), 205–223. <https://doi.org/10.1007/s40593-015-0089-1>
 - Ball, H., & Goodboy, A. K. (2014). An experimental investigation of the antecedents and consequences of psychological reactance in the college classroom. *Communication Education*, 63(3), 192–209. <https://doi.org/10.1080/03634523.2014.918634>
 - Bastani, H., Bastani, O., Sungu, A., Ge, H., Kabakc■, Ö., & Mariman, R. (2025). Generative AI without guardrails can harm human learning. *Proceedings of the National Academy of Sciences*. <https://doi.org/10.1073/pnas.2422633122>
 - Brehm, J. W. (1966). *A Theory of Psychological Reactance*. Academic Press.
 - Deci, E. L., & Ryan, R. M. (2000). The “what” and “why” of goal pursuits: Human needs and the self-determination of behavior. *Psychological Inquiry*, 11(4), 227–268. https://doi.org/10.1207/S15327965PLI1104_01
 - Johnson, E. J., & Goldstein, D. (2003). Do defaults save lives? *Science*, 302(5649), 1338–1339. <https://doi.org/10.1126/science.1091721>
 - Karabenick, S. A., & Knapp, J. R. (1991). Relationship of academic help seeking to the use of learning strategies and other instrumental achievement behavior in college students. *Journal of Educational Psychology*, 83(2), 221–230. <https://doi.org/10.1037/0022-0663.83.2.221>
 - Ryan, A. M., & Pintrich, P. R. (1997). “Should I ask for help?” The role of motivation and attitudes in adolescents’ help seeking in math class. *Journal of Educational Psychology*, 89(2), 329–341. <https://doi.org/10.1037/0022-0663.89.2.329>
 - Thaler, R. H., & Sunstein, C. R. (2008). *Nudge: Improving decisions about health, wealth, and happiness*. Yale University Press.
-

Part CXXXVII — Reveal-Practice vs C4 Split Pools for Sealed FormatId

Chapter status: Living evidence + bank/pool architecture brief — Researcher tick 2026-08-13 (UTC hour 09; hour%6≠0; researcher count since synthesizer v1.17 = 3 → Researcher)

Primary question: Once a FormatId item is sealed enough for **reveal Practice + SAFE-REPAIR**, must MindCraft keep it in a **separate pool** from **C4 hide-correctness** (Map/FEI writes) — and how do we prevent silent dual-use, silent promotion, or “one bank = all uses”?

Owners: Product (Practice / gap-scan pools) · Engine (eligibility flags / event sources) · Trust / REPAIR · Brand / COVER · Red Team

Commercial job: Ship **SAFE-SPLITPOOL** densifying SAFE-HCELIG × SAFE-REPAIR × SAFE-GENQ × SAFE-COVER × Kane/Standards graded uses: **explicit split pools by use ambition** — never unified “in bank ⇒ diagnose,” never silent Map writes from Practice-only sealed items, never SplitPool Score™.

Builds on: Parts CXXXIV (SAFE-HCELIG — foreshadows Practice-eligible / C4-ineligible tier), LX (SAFE-REPAIR), LXXII (SAFE-GENQ), LXXXI (SAFE-COLD), LXXXII (SAFE-FORMAT), LXXXVII (SAFE-COVER), CXXXI (SAFE-FIGKEY). Seams: `practice_eligible`, `hcelig_pass`, `pool selectors`, `/record-outcomes` source tags, repair tickets → promote-to-C4 workflow.

CXXXVII.1 Why this chapter exists

SAFE-HCELIG answered *whether* a FormatId item may enter hide-correctness: C4 is the strictest use; seals first; honest holes beat fake coverage. It left the **pool architecture** underspecified. Product temptation after seals:

1. **Unified pool** — one questionBank filter; any sealed item may land in Practice *or* gap-scan.
2. **Silent dual-write** — Practice attempt with reveal still POSTs the same mastery/Map path as C4.
3. **Auto-promote** — “enough Practice plays without repair tickets ⇒ flip `hcelig_pass`.”
4. **Split pools (this chapter)** — Practice-reveal pool ⊆ sealed stock; C4 pool ⊆ HCELIG-pass stock; promotion is a *named* workflow with REPAIR evidence, not a folder rename.

Unified pools feel efficient. They collapse Kane’s point that **uses** are what get validated — not folders of items. Reveal Practice can surface “this graph looks broken” (REPAIR path). C4 cannot. Mixing them without flags teaches the Map from the wrong use ambition.

FOUNDER BELIEF under audit: FormatId Practice can grow *before* FormatId diagnosis is dense — but only if the company refuses to pretend Practice evidence is hide-correctness evidence.

Claims we refuse:

1. “In bank / sealed” ≡ eligible for both Practice reveal and C4 Map writes.
 2. Silent dual-write of Practice outcomes into diagnostic Map/FEI paths.
 3. Auto-promote to C4 from engagement counts or thumbs-up.
 4. SplitPool Score™ / Dual-Use Minutes / Pool-Separation % as North Star.
 5. Padding gap-scan with Practice-only sealed visuals to avoid holes.
 6. ACT / identity guarantees from “we practice and diagnose on the same FormatId bank.”
-

CXXXVII.2 Constructs

Construct	Meaning	MindCraft analogue	Failure mode
Reveal Practice	Attempt → outcome → key/feedback visible; REPAIR possible	Practice gym FormatId hop	Treated as C4
C4 hide-correctness	Record outcome; withhold key; Map/FEI may update	Gap-scan / cold probes	Poison from unsealed or Practice-grade items
Split pools	Disjoint (or nested) selectors by <i>use</i>	pool=practice_reveal vs pool=c4_hcelig	One selector, two labels
Promote-to-C4	Explicit seal upgrade with evidence	Ticketed workflow: FIGKEY+HCELIg review	Auto-flip from play count
Use ambition	Kane: what decision the score authorizes	Route / parent copy / FEI	Browsing ≡ diagnosis
SAFE-SPLITPOOL	Pool doctrine for graded FormatId uses	This chapter	SplitPool Score™

Operational definition (HYPOTHESIS): A sealed FormatId item may sit in the **Reveal-Practice pool** when GENQ (+ FIGKEY if figured) pass and REPAIR telemetry is wired — *without* hcelig_pass. It may enter the **C4 pool** only when HCELIg passes. Practice outcomes use source tags that **do not** authorize the same Map/FEI writes as C4. Promotion Practice→C4 requires a logged seal review, not engagement thresholds alone.

CXXXVII.3 Uses are validated separately — not “the bank”

FACT: Kane (2013, *Journal of Educational Measurement*, 50(1), 1–73, doi:10.1111/jedm.12000) — validate **interpretations and uses**, not the test itself; more ambitious claims need more support; rejecting a use need not kill a weaker interpretation; validating an interpretation does not automatically validate a use.

FACT: AERA, APA, & NCME (2014, *Standards for Educational and Psychological Testing*) — validity is the degree to which evidence and theory support interpretations of scores **for proposed uses**; a measure may be justified for one purpose and not another; each intended use needs relevant evidence.

Applied (HYPOTHESIS): “Sealed FormatId stem renders correctly” can support a **low–medium ambition** use (Practice with reveal + repair) while failing a **high ambition** use (C4 → Map format gap → PawHub route → parent “we diagnosed”). Split pools are the product expression of separate IUAs.

Kill: One eligibility bit for all FormatId surfaces.

Survive: Nested seals; $C4 \subset \text{Practice-eligible sealed set}$, never the reverse.

CXXXVII.4 Formative vs “quiet high-stakes” is about use of the information

FACT: Black & Wiliam (1998, *Phi Delta Kappan*, 80(2), 139–148) — formative work means evidence is used to adapt teaching/learning; teachers face real tension reconciling formative and summative roles; confusion between roles impedes practice.

FACT: Wiliam (2000, *Critical Quarterly*, 42(1), 105–127; see also related essays on meanings/consequences of assessment) — strictly, “formative” names the **use** of information, not a mystical item type; the same instrument can serve different functions depending on what happens next.

Applied (HYPOTHESIS): Reveal Practice is formative *when* feedback + REPAIR + retry move learning. C4 is a **quiet high-stakes write** into the student’s Map even without showing a key — more like an interpretive/use decision than a learning conversation. Treating Practice-pool items as C4 “because we already sealed them” confuses Black/Wiliam’s roles inside one codebase.

Kill: “It’s the same question, so the same graph update.”

Survive: Source-tagged outcomes; C4 writes only from C4 pool administrations.

CXXXVII.5 Pool integrity is an ops problem, not a folder aesthetic

FACT: Way (1998, *Educational Measurement: Issues and Practice*, 17(4), 17–27, doi:10.1111/j.1745-3992.1998.tb00632.x) — computerized programs must actively protect item-pool integrity; exposure and pool-management strategies matter because continuous delivery changes compromise risk vs paper forms.

Applied (HYPOTHESIS, careful transfer): MindCraft is not selling high-stakes CAT security theater. The transferable lesson is narrower: **pools are managed objects with policies**, not a single dump of IDs. Separate Practice vs C4 pools (or hard flags with enforced selectors) reduce *use compromise* — the educational analogue of item compromise — where a Practice-seen, key-revealed stem later pretends to be a fresh hide-correctness probe, or a Practice-grade seal silently diagnoses.

Kill: Single unsorted FormatId list with runtime “maybe hide key.”

Survive: Named pools + exposure rules (e.g., bar recently revealed stems from same-student C4 for a cool-down).

CXXXVII.6 Repair is the bridge — not a silent flip

Reuse (SAFE-REPAIR / SAFE-HCELIG): Reveal Practice exists partly so students and tutors can flag broken figures. That signal is commercially precious: it is how Practice-grade stock earns (or fails) promotion.

HYPOTHESIS: Promote-to-C4 requires (a) GENQ+FIGKEY+FORMTRAP seals, (b) zero open critical repair tickets on that item, (c) optional tutor/SME checklist, (d) explicit hcelig_pass write with reviewer id — never “N plays without complaint ⇒ eligible.” Absence of tickets ≠ proof when volume is low (SAFE-COVER honesty).

Kill: Engagement-based auto-promotion.

Survive: Ticketed promote workflow; demote-on-repair that *removes* from C4 immediately while leaving Practice-reveal if policy allows.

CXXXVII.7 Claim table

Claim	Label	Confidence	Notes
Uses/interpretations are validated; ambition scales evidence	FACT	High	Kane 2013; Standards 2014
Formative vs other roles is about use of information	FACT	High	Black & Wiliam 1998; Wiliam 2000
Pool integrity requires active management policies	FACT	High	Way 1998 (transfer with caution)
Practice-reveal and C4 need split pools / hard use flags	HYPOTHESIS	Medium–High	Densifies HCELIG + REPAIR
Auto-promote from play counts is safe enough	SPECULATION	Low	Refuse
Unified sealed bank is fine if drop rate looks low	SPECULATION	Low	Refuse
Parents will accept honest FormatId holes if Practice still offers sealed visuals	FOUNDER BELIEF	Medium	Needs SPLITPOOL-* × COVER

CXXXVII.8 Product surface — SAFE-SPLITPOOL claim contract

1. Emit `practice_eligible` and `hcelig_pass` as distinct bits (HCELIG already requires the second).
 2. Gap-scan `FormatId` selector: `hcelig_pass` **only**.
 3. Practice `FormatId` hop: may use `practice_eligible` && `!hcelig_pass` stock **with reveal + REPAIR**.
 4. **Bar** silent Map/FEI diagnostic writes from Practice-pool administrations (source tag enforcement).
 5. Cool-down: stems revealed to a student barred from that student’s C4 for a policy window (anti-answer-hunt / exposure).
 6. Promote-to-C4: ticketed review; demote-on-repair removes C4 immediately.
 7. Coverage UI: show **Practice FormatId stock** vs **C4 FormatId stock** separately (SAFE-COVER).
 8. Telemetry: pool id, use flag, promote/demote events — demote SplitPool Score™ / Dual-Use Minutes.
 9. Copy: “We practice on sealed figures before we hide the key.” / Parent: “Diagnosis uses a stricter pool than practice.”
-

CXXXVII.9 Doctrine — SAFE-SPLITPOOL (provisional)

1. **Use $\neq$ file** — sealed presence is not dual admission.
2. **C4 pool $\subseteq$ HCELIG**; Practice-reveal pool may be wider within GENQ/FIGKEY seals.
3. **No silent dual-write** — outcome source tags enforce ambition.
4. **Promote is a workflow** — not play-count theater (REPAIR $\times$ HCELIG).
5. **Demote is faster than promote** — repair tickets yank C4 first.
6. **Honest dual coverage tables** — Practice stock $\neq$ diagnostic stock (COVER).
7. **No SplitPool Score™ / Dual-Use Minutes / Pool-Separation % NS.**
8. **No ACT / identity guarantees** from unified-bank packaging.
9. Copy: “Same stem, different job — until the stricter seal says otherwise.”

Confidence: High on Kane/Standards/Black–Wiliam/Wiliam as method sources; High on Way as pool-ops analogy with explicit transfer limits. Medium on cool-down lengths and exact promote checklists (needs SPLITPOOL-*). High that unified dual-use without flags is commercially toxic (Map distrust + coverage theater).

CXXXVII.10 Experiments

ID	Question	Design	Primary
SPLITPOOL-1	Split selectors vs unified sealed FormatId pool for C4	Policy A/B	False-weakness rate; Map $\leftrightarrow$ tutor gold
SPLITPOOL-2	Practice-only sealed tier with REPAIR vs dual-bar (no Practice until HCELIG)	Product A/B	Format-hop volume; repair tickets; trust
SPLITPOOL-3	Source-tag enforcement vs silent dual-write	Engine A/B	Illicit Map write rate; audit catch
SPLITPOOL-4	Ticketed promote vs auto-promote after N clean plays	Ops A/B	Post-promote repair rate; false C4 admits
SPLITPOOL-5	Parent CBC: separate Practice vs diagnostic FormatId counts vs single “N diagram items”	CBC	WTP; trust; abandon
SPLITPOOL-QUAL	Planted broken figure in Practice-only tier: does REPAIR catch before any C4 path?	Qual	Promote codebook

Falsifiers: Split pools equal unified on false-weakness → audit selector bugs before collapsing pools. Dual-bar (no Practice until HCElig) kills FormatId practice without trust gain → keep Practice-only tier. Auto-promote matches ticketed on post-promote repairs → still prefer ticketed for liability; do not celebrate engagement. Pre-register SPLITPOOL-* before “one FormatId bank for practice and diagnosis” campaigns.

CXXXVII.11 So what for MindCraft commercially

- **Copy:** “We practice on sealed figures before we hide the key.” Parent: “Your child’s format diagnosis uses a stricter pool than everyday practice — we leave holes rather than guess.”
 - **Product:** Distinct pool selectors + source tags; promote/demote workflow; cool-down after reveal.
 - **Growth:** Sell honesty of dual coverage tables; kill “N diagram items” undifferentiated hero.
 - **Positioning:** Against ChatGPT one-blob “practice/diagnose” fluency *and* against edtech that hides keys on unverified or Practice-grade art.
 - **Metric:** pool_id / illicit Map-write rate / promote–demote events — demote SplitPool Score™.
 - **Kill list:** Unified dual-use; silent Map writes; auto-promote; coverage padding; SplitPool Score™; ACT guarantees from one-bank packaging.
 - **Vision:** Maya’s Map only learns from probes the company would defend in a parent review — while her Practice gym can still grow sealed FormatId skill with the lights on.
-

References (verified)

- American Educational Research Association, American Psychological Association, & National Council on Measurement in Education. (2014). *Standards for educational and psychological testing*. AERA.
 - Black, P., & Wiliam, D. (1998). Inside the black box: Raising standards through classroom assessment. *Phi Delta Kappan*, 80(2), 139–148.
 - Kane, M. T. (2013). Validating the interpretations and uses of test scores. *Journal of Educational Measurement*, 50(1), 1–73. <https://doi.org/10.1111/jedm.12000>
 - Way, W. D. (1998). Protecting the integrity of computerized testing item pools. *Educational Measurement: Issues and Practice*, 17(4), 17–27. <https://doi.org/10.1111/j.1745-3992.1998.tb00632.x>
 - Wiliam, D. (2000). The meanings and consequences of educational assessments. *Critical Quarterly*, 42(1), 105–127. <https://doi.org/10.1111/1467-8705.00280>
-

Part CXXXVIII — FormatId Gap-Scan Composition under Sealed-Stock Scarcity

Chapter status: Living evidence + diagnostic composition brief — Researcher tick 2026-08-13 (UTC hour 12 ≡ Red Team slot, but ch138 never written → prefer Researcher per rotation; researcher count since synthesizer v1.17 = 4 → Researcher)

Primary question: When the **C4-eligible FormatId pool is thin** ($\text{SAFE-HCELIG} \cap \text{SAFE-SPLITPOOL}$), how should MindCraft **compose** a gap-scan — without symbolic pad theater, forced equal quotas on one sealed diagram, fake “we diagnose every format” breadth, or Map writes that pretend unprobed FormatIds were measured?

Owners: Product (gap-scan / cold-start composition) · Engine (selectors / exposure / hole flags) · Bank / COVER · Brand · Red Team

Commercial job: Ship **SAFE-SCANCOMP** densifying $\text{SAFE-HCELIG} \times \text{SAFE-SPLITPOOL} \times \text{SAFE-COVER} \times \text{SAFE-FORMAT} \times \text{Kane/Messick/Sireci/Raymond}$ blueprint honesty: **compose from sealed stock with labeled holes** — never pad, never Composition Score™, never ACT guarantees from “balanced FormatId scan” packaging.

Builds on: Parts CXXXIV (SAFE-HCELIG), CXXXVII (SAFE-SPLITPOOL), LXXXII (SAFE-FORMAT), LXXXVII (SAFE-COVER), LXXII (SAFE-GENQ), CXXXI (SAFE-FIGKEY), LXXXI (SAFE-COLD), L (SAFE-CALIB). Seams: hcelig_pass pool, gap-scan blueprint cells, format_not_probed flags, coverage matrix, PawHub format routes.

CXXXVIII.1 Why this chapter exists

SAFE-HCELIG decided *admission*. SAFE-SPLITPOOL decided *which pool*. Scarcity still breaks product judgment:

1. **Symbolic pad** — fill a “diagram” slot with a text-only stem so the scan looks complete.
2. **Quota theater** — force one item per FormatId even when only a single sealed graph exists (over-expose; under-generalize).
3. **Breadth ads** — “we diagnose every format” while C4 stock covers 1–2 FormatIds.
4. **Silent underrepresentation** — skip FormatId entirely, then sell concept-only Map as format-flexible (Messick underrepresentation).
5. **Hole denial** — omit not_probed flags so parents/tutors read absence as mastery.

Composition is the commercial knife-edge: honesty without paralysis. The company must ship gap-scans that are **defensible under review** when sealed FormatId stock cannot yet fill a pretty matrix.

FOUNDER BELIEF under audit: A short, sealed, hole-labeled FormatId slice beats a long, padded, “balanced” scan that teaches the Map lies.

Claims we refuse:

1. Symbolic (or Practice-only) items may fill C4 FormatId blueprint cells to avoid holes.
 2. Equal FormatId quotas regardless of sealed-stock depth.
 3. Over-reuse of a single sealed stem as “format coverage.”
 4. Composition Score™ / Format Balance % / Scan Completeness Minutes as North Star.
 5. “We diagnose all formats” marketing under thin C4 stock.
 6. Absence of FormatId probe $\equiv$ format flexibility / mastery.
 7. ACT / identity guarantees from balanced-composition packaging.
-

CXXXVIII.2 Constructs

Construct	Meaning	MindCraft analogue	Failure mode
Sealed-stock scarcity	Too few <code>hcelig_pass</code> items per <code>FormatId</code> ×concept cell	Thin C4 <code>FormatId</code> pool	Pad or skip-silently
Blueprint cell	Intended probe slot (concept × <code>FormatId</code> × level)	Gap-scan composition table	Cell filled with wrong use/stock
Honest hole	Empty cell labeled <code>format_not_probed</code>	COVER matrix + scan summary	Hole sold as green
Composition policy	Rules for picking among eligible stock	SAFE-SCANCOMP selector	Quota theater / random dump
Exposure under scarcity	Reuse risk when pool is tiny	Per-student cool-down; max exposes	One graph = “diagram diagnosis”
SAFE-SCANCOMP	Scarcity-honest composition doctrine	This chapter	Composition Score™

Operational definition (HYPOTHESIS): Gap-scan `FormatId` composition (1) draws **only** from `pool=c4_hcelig`, (2) follows a published blueprint of *desired* cells, (3) **omits** cells with zero sealed stock and emits `format_not_probed` (never pads), (4) prefers **depth where stock exists** over fake breadth, (5) caps per-item exposure / cool-down under thin pools, (6) never writes `format-gap` claims for unprobed `FormatIds`, (7) surfaces hole honesty in student/parent/tutor copy (SAFE-COVER; feeds Part CXXXIX).

CXXXVIII.3 Uses scale with composition ambition

FACT: Kane (2013, *Journal of Educational Measurement*, 50(1), 1–73, doi:10.1111/jedm.12000) — validate interpretations and **uses**; more ambitious claims need more support; rejecting a use need not kill a weaker interpretation.

Applied (HYPOTHESIS):

Composition	Ambition of use	Allowed when stock is scarce?
Concept-only C4 (no <code>FormatId</code> cells claimed)	Medium	Yes — if copy does not claim <code>format diagnosis</code>
Partial <code>FormatId</code> slice + labeled holes	Medium–high	Yes — preferred under scarcity
“Full <code>FormatId</code> diagnosis” / every <code>FormatId</code> green	High	No until sealed depth meets blueprint
Parent “we mapped every visual gap”	High	Same bar as full diagnosis

Kill: Thin stock + high-ambition `FormatId` claims.

Survive: Lower ambition + honest holes, or delay FormatId C4 cells until promote-to-C4 grows stock (SAFE-SPLITPOOL / id 140).

CXXXVIII.4 Underrepresentation vs irrelevant pad — both kill

FACT: Messick (1995, *American Psychologist*, 50(9), 741–749, doi:10.1037/0003-066X.50.9.741) — two primary threats: **construct underrepresentation** (too narrow; missing facets) and **construct-irrelevant variance** (contamination by extraneous difficulty). Both can operate at once.

Applied (HYPOTHESIS): Scarcity creates a fork:

- **Underrepresentation** — never probe sealed diagrams → Map pretends format flexibility from symbolic-only evidence.
- **CIV pad** — force a decorative/text-as-diagram or unsealed visual into the cell → miss reflects the wrong construct, then hide-correctness locks the lie in.

SAFE-SCANCOMP refuses both: **omit + label** beats **pad**; **probe sealed facets you have** beats **silence that markets as complete**.

Kill: “Better a weak diagram than a hole.”

Survive: Prefer holes; invest GENQ/FIGKEY/promote pipeline to grow real cells.

CXXXVIII.5 Domain representation is not equal quotas

FACT: Sireci and Faulkner-Bond (2014, *Psicothema*, 26(1), 100–106, doi:10.7334/psicothema2013.256) — content-related validity evidence includes **domain definition**, **domain representation**, **domain relevance**, and appropriateness of construction procedures; representation asks whether items adequately reflect the defined domain — not whether every marketing checkbox got one token item.

FACT: Raymond and Grande (2019, *Medical Teacher*, 41(8), 854–861, doi:10.1080/0142159X.2019.1595556) — a test blueprint allocates emphasis across knowledge/skill domains; blueprinting supports validation by making intended coverage and weights explicit before assembly.

Applied (HYPOTHESIS): MindCraft’s FormatId gap-scan needs a **blueprint** (which FormatIds matter for this exam track / concept set; target weights) *and* a **scarcity rule**: when sealed stock cannot meet a weight, **shrink the claim**, do not manufacture items. Equal 1-of-each quotas under N=1 sealed graphs are **representation theater**, not Sireci representation.

Kill: Forced equal FormatId counts as “science-backed balance.”

Survive: Weighted blueprint + stock-gated fulfillment + hole labels for unmet weights.

CXXXVIII.6 Thin pools need assembly constraints — not vibes

FACT: van der Linden, Ariel, and Veldkamp (2006, *Journal of Educational and Behavioral Statistics*, 31(1), 81–99, doi:10.3102/10769986031001081) — CAT pools assembled without content/psychometric constraints can force algorithms into suboptimal, constraint-violating, or poorly exposed selections; principled assembly (here, as sets of linear forms meeting specs) improves control of content and exposure properties.

Applied (HYPOTHESIS, careful transfer): MindCraft is not shipping high-stakes CAT. Transferable: **composition is constrained assembly**. Under scarcity, the selector must encode hard constraints (hcelig_pass only; max exposure; cool-down; no cross-pool pad) and soft targets (blueprint weights). When constraints cannot be met, **fail the cell openly** — do not relax into Practice-only or symbolic fillers (SAFE-SPLITPOOL).

Kill: “Just pick whatever tagged FormatId is in the bank.”

Survive: Constraint-first selector; telemetry on constraint violations and hole rates.

CXXXVIII.7 Claim table

Claim	Label	Confidence	Notes
Uses/ambition scale evidence needed	FACT	High	Kane 2013
Underrepresentation and CIV are dual threats	FACT	High	Messick 1995
Domain representation ≠ token quotas	FACT	High	Sireci & Faulkner-Bond 2014
Blueprints make intended coverage/weights explicit	FACT	High	Raymond & Grande 2019
Thin pools need content/exposure constraints	FACT	High	van der Linden et al. 2006 (transfer limits noted)
Honest holes + sealed depth > padded breadth	HYPOTHESIS	Medium–High	Densifies HCELIG×SP LITPOOL×COVER
Parents accept labeled FormatId holes if Practice still offers sealed visuals	FOUNDER BELIEF	Medium	Needs SCANCOMP-* × CXXXIX
Equal FormatId quotas always improve diagnosis	SPECULATION	Low	Refuse
Symbolic pad is “temporary” and harmless under C4	SPECULATION	Low	Refuse

CXXXVIII.8 Product surface — SAFE-SCANCOMP claim contract

1. Maintain a **gap-scan FormatId blueprint** (exam-track × FormatId weights) separate from raw bank counts.
 2. Fill cells **only** from hcelig_pass / pool=c4_hcelig.
 3. If stock = 0 for a cell → emit format_not_probed; do not pad; do not write format-gap mastery for that FormatId.
 4. If stock = 1–2 → allow at most **one** exposure per student per cool-down; do not pretend breadth.
 5. Prefer probing **sealed FormatIds with depth** over sprinkling token singles across five FormatIds.
 6. UI: scan summary shows **probed vs not_probed** FormatIds (student + tutor + parent surfaces — densifies COVER; parent copy → id 139).
 7. Admin/coverage matrix: C4 FormatId depth vs Practice FormatId depth (SPLITPOOL dual tables).
 8. Telemetry: blueprint_cell, fulfilled|hole, item_id, exposure_count — demote Composition Score™ / Format Balance %.
 9. Copy: “We only hide the key on sealed visuals — empty slots mean we have not diagnosed that format yet.”
 10. Ban marketing that converts hole rates into “complete visual diagnostic” claims.
-

CXXXVIII.9 Doctrine — SAFE-SCANCOMP (provisional)

1. **Blueprint first** — intended FormatId coverage is a document, not a vibe.
2. **Stock gates fulfillment** — no pad into C4 FormatId cells.
3. **Holes are features** — format_not_probed beats fake green (SAFE-COVER).
4. **Depth over fake breadth** under scarcity.
5. **Exposure caps** when pool is tiny (assembly constraint transfer).
6. **Ambition follows stock** — thin stock ⇒ lower FormatId claims (Kane).
7. **No Composition Score™ / Format Balance % / Scan Completeness Minutes NS.**
8. **No ACT / identity guarantees** from “balanced FormatId scan” packaging.
9. Copy: “Diagnosed what we can seal; labeled what we cannot.”

Confidence: High on Kane/Messick/Sireci/Raymond as method sources; High on van der Linden et al. as pool-assembly analogy with explicit non-CAT transfer. Medium on exact cool-down lengths and weight tables (needs SCANCOMP-*). High that symbolic pad under C4 is commercially toxic (Map distrust + COVER theater).

CXXXVIII.10 Experiments

ID	Question	Design	Primary
SCANCOMP-1	Hole-labeled partial FormatId scan vs symbolic-pad “complete” scan	Policy A/B	False format-gap rate; Map↔tutor gold

SCANCOMP-2	Depth-first (probe rich FormatIds) vs equal-quota under thin stock	Selector A/B	Generalization of format routes; over-exposure
SCANCOMP-3	Exposure cap/cool-down vs unlimited reuse of N=1 sealed graph	Engine A/B	Item exposure; student answer-hunt signals
SCANCOMP-4	Ambition-matched copy (“partial”) vs “all formats” ads under same stock	Copy A/B	Trust; complaint rate; WTP
SCANCOMP-5	Parent CBC: labeled holes vs hidden holes vs padded completeness	CBC	Trust; abandon; WTP (feeds CXXXIX)
SCANCOMP-QUAL	Tutor review: which scan summary they would defend to a parent?	Qual	Codebook of hole vs pad preference

Falsifiers: Pad equals holes on false-gap rate → audit labels/Map write paths before allowing pad. Equal-quota beats depth-first on transfer without exposure harm → revisit weights, not honesty. Parents punish labeled holes more than pad theater on WTP *and* later Map distrust is equal → still prefer holes for liability; do not celebrate pad. Pre-register SCANCOMP-* before “complete FormatId diagnostic” campaigns.

CXXXVIII.11 So what for MindCraft commercially

- **Copy:** “We only diagnose formats we can seal — empty slots are honesty, not failure.” Kill “every format covered” under thin C4 stock.
- **Product:** Blueprint + stock-gated composition + format_not_probed + exposure caps; dual Practice vs C4 depth tables.
- **Growth:** Sell trust of labeled holes; pair with promote-to-C4 pipeline (id 140) as the *fix* for holes.
- **Positioning:** Against edtech that fills visual quizzes with decoration, and against “complete bank” ads (SAFE-COVER) that ignore use-tier seals.
- **Metric:** hole rate by FormatId; illicit pad attempts; exposure under $N \leq 2$ stock — demote Composition Score™.
- **Kill list:** Symbolic pad; equal-quota theater; silent underrepresentation; all-formats ads; Composition Score™; ACT guarantees from balanced-scan packaging.
- **Vision:** Maya’s gap-scan tells the truth about what FormatIds were measured — so the Map’s format routes remain worth following when stock finally catches the blueprint.

References (verified)

- Kane, M. T. (2013). Validating the interpretations and uses of test scores. *Journal of Educational Measurement*, 50(1), 1–73. <https://doi.org/10.1111/jedm.12000>
- Messick, S. (1995). Validity of psychological assessment: Validation of inferences from persons’ responses and performances as scientific inquiry into score meaning. *American Psychologist*, 50(9), 741–749. <https://doi.org/10.1037/0003-066X.50.9.741>
- Raymond, M. R., & Grande, J. P. (2019). A practical guide to test blueprinting. *Medical Teacher*, 41(8), 854–861. <https://doi.org/10.1080/0142159X.2019.1595556>
- Sireci, S., & Faulkner-Bond, M. (2014). Validity evidence based on test content. *Psicothema*, 26(1), 100–106. <https://doi.org/10.7334/psicothema2013.256>
- van der Linden, W. J., Ariel, A., & Veldkamp, B. P. (2006). Assembling a computerized adaptive testing item pool as a set of linear tests. *Journal of Educational and Behavioral Statistics*, 31(1), 81–99. <https://doi.org/10.3102/10769986031001081>

Part CXXXIX — Parent Trust Copy when Visual Diagnostic Holes Remain

Chapter status: Living evidence + parent/GTM copy brief — Researcher tick 2026-08-13 (UTC hour 15; hour%6≠0; researcher count since synthesizer v1.17 = 5 → Researcher)

Primary question: When SAFE-SCANCOMP leaves **honest FormatId holes** (`format_not_probed`) on the gap-scan / parent surface, what **trust copy** should MindCraft ship — without coverage shame, pad theater, “complete visual diagnostic” ads, or Hole Trust Score™ vanity — so parents revise beliefs and still buy the *repair path* (promote-to-C4 / sealed Practice), not abandon or demand fake greens?

Owners: Growth/copy · Product (ParentDashboard / gap-scan summary) · Brand · COVER / bank · WTP · Red Team

Commercial job: Ship **SAFE-HOLETRUST** densifying SAFE-SCANCOMP × SAFE-PDASH × SAFE-COVER × SAFE-WTP × transparency/trust literature: **label the hole, name the use, offer one autonomy-supportive fill path** — never hide holes, never shame scarcity, never sell absence as mastery.

Builds on: Parts CXXXVIII (SAFE-SCANCOMP), LXXXIII (SAFE-PDASH), LXXXVII (SAFE-COVER), LIX (SAFE-WTP), CXVIII (SAFE-SOFTMSG), CXXVII (SAFE-CRUNCHMSG), LXXII (SAFE-GENQ), CXXXIV (SAFE-HCELIG), CXXXVII (SAFE-SPLITPOOL), LXXXVI (SAFE-PROOF). Seams: parent gap-scan summary, `format_not_probed` flags, coverage matrix, WTP CBC arms, promote-to-C4 pipeline (id 140).

CXXXIX.1 Why this chapter exists

SAFE-SCANCOMP ordered product to **omit + label** unfilled FormatId cells. Growth still faces hide (underrepresentation theater), shame (controlling involvement), pad (Map distrust), apology fog (belief untouched), or **honesty that converts** — disclosure + clarity + accuracy about sealed probes plus one fill CTA toward sealed Practice or promote-to-C4. This chapter owns parent language and CBC for the last path.

FOUNDER BELIEF under audit: Serious parents prefer labeled FormatId holes when copy explains the C4 seal bar, separates Practice from diagnostic claim, and offers one fill path — beating fake-complete dashboards that later

contradict the Map.

Claims we refuse:

1. Hide `format_not_probed` from parents to protect conversion.
 2. Coverage-shame / “visual gap = child deficit” as default hole framing.
 3. “Complete visual diagnostic” / “we mapped every format” ads while holes remain.
 4. Hole Trust Score™ / Coverage Satisfaction / Honesty Minutes as North Star.
 5. Absence of `FormatId` probe $\equiv$ format mastery / flexibility on parent MoC.
 6. Fear-appeal scarcity emails (“act before holes widen”) as default WTP lever.
 7. ACT / identity guarantees from hole-labeled or hole-hidden packaging.
-

CXXXIX.2 Constructs

Construct	Meaning	MindCraft analogue	Failure mode
Honest hole	Blueprint cell with zero sealed C4 stock, labeled	<code>format_not_probed</code> on parent summary	Hole sold as green / omitted
Use-tier disclosure	Name what the scan <i>may</i> claim	C4 sealed vs Practice reveal vs not playable	One “coverage %”
Belief object	What parent prior should update	“We have not yet diagnosed diagram for X”	“Maya is weak at visuals” character frame
Fill path	One autonomy-supportive next step	Sealed Practice mission / promote ETA / tutor ticket	Nag / solve-for-them / pad
Transparency triad	Disclosure, clarity, accuracy	Hole list + plain language + no fake greens	Disclosure dump without clarity
SAFE-HOLETRUST	Parent trust copy doctrine under visual holes	This chapter	Hole Trust Score™

Operational definition (HYPOTHESIS): Parent hole copy is SAFE-HOLETRUST-complete when it (1) lists probed vs `not_probed` `FormatIds` for the named exam track, (2) states the **seal reason** in one plain sentence (“we only hide keys on verified visuals”), (3) never frames the hole as the child’s character failure, (4) pairs the hole with **exactly one** CTA (sealed Practice where available, or “we’re verifying more diagram items” without fake ETA theater), (5) never claims format mastery for unprobed cells, and (6) is CBC-tested against hide/pad/shame arms before scale (WTP-*).

CXXXIX.3 Transparency is disclosure + clarity + accuracy — not dump or fog

FACT: Schnackenberg and Tomlinson (2016, *Journal of Management*, 42(7), 1784–1810, doi:10.1177/0149206314525202) — organizational transparency is usefully treated as stakeholder perceptions of **information disclosure, clarity, and accuracy**; these dimensions are theorized as antecedents that shape perceived ability, benevolence, and integrity (and thus trustworthiness) in Mayer, Davis, and Schoorman’s trust model — transparency is not mere volume of data shared.

FACT: Mayer, Davis, and Schoorman (1995, *Academy of Management Review*, 20(3), 709–734, doi:10.5465/amr.1995.9508080335) — trust is willingness to be vulnerable based on expectations of the other’s ability, benevolence, and integrity.

Applied (HYPOTHESIS):

Failure mode	Transparency break	Parent trust risk
Hide holes	Low disclosure	Integrity hit when Map later shows gaps
Jargon matrix dump	Low clarity	Ability/competence fog; abandon reading
Pad greens	Low accuracy	Integrity + ability hit when tutor contradicts
Empathy-only apology	Disclosure without actionable clarity	Benevolence theater; belief unchanged

Kill: “Transparency” as either raw coverage CSV dump or soft apology without labeled holes.

Survive: Short triad — what was probed, what was not, why the bar exists, one next step.

CXXXIX.4 Accountability-style honesty beats secrecy — with transfer limits

FACT: Rawlins (2008, *Public Relations Journal*, 2(2), 1–21) — employee sample: substantial information, participation, and accountability (vs secrecy) relate positively to trust. **Transfer limit:** not a parent edtech RCT — borrow the construct, not the correlation as ACT WTP proof.

Applied (HYPOTHESIS): Prefer accountability (“we have not sealed enough diagram items to diagnose that format yet”) over secrecy (omit the slot) or blame (student deficit). Invite parents to witness a later sealed Practice return — not to grade every miss (SAFE-PDASH).

Kill: Secrecy-as-conversion. **Survive:** Accountable incompleteness + one MoC CTA.

CXXXIX.5 Belief correction needs specific facts — not character shame

FACT: Bergman (2021, *Journal of Political Economy*, 129(1), 286–322, doi:10.1086/711410) — parents hold upwardly biased beliefs about child effort; biweekly information on missed assignments/grades attenuates bias and raises achievement via more accurate beliefs and lower monitoring costs.

FACT: Bergman and Chan (2021, *Journal of Human Resources*, 56(1), 125–158, doi:10.3368/jhr.56.1.1118-9837R1) — weekly automated alerts on misses/grades/absences cut course failures and raised attendance; **no** significant state-test lift — process honesty ≠ standardized-score miracle.

FACT: Rogers and Feller (2018, *Nature Human Behaviour*, 2, 335–342, doi:10.1038/s41562-018-0328-1) — cumulative personalized absence facts reduced chronic absenteeism; relative peer framing was **not** the active ingredient in their design.

Applied (HYPOTHESIS): The belief object for holes is **product completeness for a named use**, not “Maya’s visual talent.” Prefer: “Diagram format: not yet diagnosed (sealed items unavailable)” over “Your child lags classmates on graphs.” Peer-rank shame around holes recreates COMPARE/PDASH kills and mis-aims Bergman-style correction at the wrong prior.

Kill: Classmate visual-rank as hole “truth.”

Survive: Cumulative, personalized, product-honest facts + one improvement object.

CXXXIX.6 Improvement-framed, autonomy-supportive CTAs — not nag or fear

FACT: Kraft and Rogers (2015, *Economics of Education Review*, 47, 49–63, doi:10.1016/j.econedurev.2015.04.001) — weekly individualized teacher→parent messages raised credit-earning; **improvement-focused** notes showed the largest point estimates vs continue-positive messages (study underpowered to fully separate arms).

FACT: Xu, Guo, Feng, Ma, Zhang, Núñez, and Fan (2024, *Psicothema*, 36(1), 1–14, doi:10.7334/psicothema2023.92) — meta-analysis: autonomy-supportive homework involvement helps; parental control / frequency dimensions do not deliver the same benefit pattern.

Applied (HYPOTHESIS): Hole CTA copy should be improvement-framed and autonomy-supportive:

- Prefer: “When sealed diagram practice is ready, Maya can run a short return — ask what felt hard; don’t solve it for her.”
- Avoid: “Sit until every visual hole is filled tonight.”
- Avoid: Fear scarcity (“Holes mean ACT risk — buy now”).

Pair with SAFE-SOFTMSG grain: name the FormatId hole + next sealed mission, not Miss Streak theater.

Kill: Controlling hole-nag; fear-WTP default.

Survive: One improvement CTA; parent as witness/structure, student owns attempt.

CXXXIX.7 Ambitious claims still need evidence — holes lower the claim, not the honesty bar

FACT: Kane (2013, *Journal of Educational Measurement*, 50(1), 1–73, doi:10.1111/jedm.12000) — validate interpretations and **uses**; more ambitious claims need more support.

FACT: Federal Trade Commission (1984). *FTC Policy Statement Regarding Advertising Substantiation* (Thompson Medical Co., 104 F.T.C. 648, 839) — prior reasonable basis required for objective claims before dissemination.

Applied (HYPOTHESIS): “Diagnosed sealed FormatIds X,Y; Z not probed” and “Practice has sealed diagrams; C4 does not yet” are allowed when true. “Complete visual diagnostic / all formats mapped” and “honest holes → +ACT points” are banned while holes remain (Kane ambition + FTC basis).

Kill: Complete-visual ads; ACT points from hole packaging. **Survive:** Ambition matched to stock.

CXXXIX.8 Claim table

Claim	Label	Confidence	Notes
Transparency ≈ disclosure + clarity + accuracy	FACT	High	Schnackenberg & Tomlinson 2016
Trust rests on ability/benevolence/integrity expectations	FACT	High	Mayer et al. 1995
Accountability/substantial info correlate with trust (employee sample)	FACT	Medium–High	Rawlins 2008; transfer limits
Parents misbelieve; specific info corrects	FACT	High	Bergman 2021; Bergman & Chan 2021
Peer-rank framing not required for belief correction	FACT	High	Rogers & Feller 2018
Improvement-framed notes outperform vague praise (suggestive)	FACT	Medium	Kraft & Rogers 2015
Autonomy support > control in homework involvement meta	FACT	High	Xu et al. 2024
Uses scale with evidence; ads need substantiation	FACT	High	Kane 2013; FTC 1984
Labeled holes + one fill CTA > hide/pad/shame on trust & WTP	HYPOTHESIS	Medium–High	Needs HOLETRUST-* / SCANCOMP-5
Parents prefer honesty even if short-run conversion dips	FOUNDER BELIEF	Medium	CBC; liability preference if WTP ties

Hole Trust Score™ improves learning	SPECULATION	Low	Refuse
-------------------------------------	-------------	-----	--------

CXXXIX.9 Product surface — SAFE-HOLETRUST claim contract

1. Parent gap-scan / MoC surface shows **Probed** vs **Not probed** FormatIds (from SCANCOMP flags).
2. One-sentence **seal reason** adjacent to holes (“C4 only uses verified visuals”).
3. If Practice has sealed stock for a FormatId that C4 lacks → show **use-tier split** (“Available in Practice · not yet in diagnostic”) — never collapse to one %.
4. Exactly **one** primary CTA: sealed Practice grain *or* “we’re verifying more [FormatId] items” with known-finite pipeline language if ETA exists (KNOWNETA discipline) — no fake-precision countdown.
5. Copy frames the **product hole**, not the child’s character; no classmate visual ranks.
6. Ban “complete visual diagnostic” / “all formats covered” while any blueprint hole remains for the marketed track.
7. Telemetry: hole_shown, cta_type, hide_attempt, shame_frame_flag — demote Hole Trust Score™ / Coverage Satisfaction / Honesty Minutes.
8. CBC arms (with WTP): labeled-honest vs hidden vs padded-complete vs shame-framed — pre-register before parent launch emails.
9. Tutor brief: same hole language parents see — no dual script that reintroduces pad in voice.
10. Link fill path to promote-to-C4 workflow (id 140) as the *company* response to holes — not parent homework police.

CXXXIX.10 Doctrine — SAFE-HOLETRUST (provisional)

1. **Holes stay visible** — never hide format_not_probed to protect conversion.
2. **Triad copy** — disclosure + clarity + accuracy (Schnackenberg & Tomlinson).
3. **Belief object = product use**, not child visual deficit (Bergman-style facts without COMPARE shame).
4. **One improvement CTA**, autonomy-supportive (Kraft/Rogers × Xu).
5. **Use-tier labels** when Practice ≠ C4 (COVER × SPLITPOOL).
6. **Ambition follows stock** — lower FormatId claims while holes remain (Kane × SCANCOMP).
7. **No Hole Trust Score™ / Coverage Satisfaction / Honesty Minutes NS.**
8. **No fear-scarcity default** / complete-visual ads / ACT guarantees from hole packaging.
9. Copy: “We diagnosed what we can seal. Empty slots mean we have not measured that format yet — not that your student failed it.”

Confidence: High on transparency/trust and parent-information spines with stated transfer limits. Medium on exact CTA wording and WTP magnitudes (needs HOLETRUST-*). High that hide/pad/shame are commercially toxic once Map/tutor contradict the portal.

CXXXIX.11 Experiments

ID	Question	Design	Primary
HOLETRUST-1	Labeled holes vs hidden holes vs padded-complete summary	Parent copy A/B	Trust items; complaint; later Map contradiction anger
HOLETRUST-2	Product-hole frame vs child-deficit/shame frame	Copy A/B	Autonomy-support perception; controlling nag incidents
HOLETRUST-3	One fill CTA vs apology-only vs multi-link maze	CTA A/B	CTA click→Practice start; follow-through
HOLETRUST-4	Use-tier split (Practice ok / C4 hole) vs single coverage %	UI A/B	Misbelief rate; support tickets
HOLETRUST-5	Parent CBC: honest holes vs hide vs pad vs shame (price/WTP)	CBC	WTP; abandon; integrity ratings (absorbs SCANCOMP-5 densify)
HOLETRUST-QUAL	Parent interviews: which summary would you defend to a friend?	Qual	Codebook of honesty vs completeness preference

Falsifiers: Hide beats labeled on WTP *and* later Map distrust is equal → still refuse hide for liability/IUA; do not celebrate. Shame frame raises short-run Practice starts but autonomy/Xu-pattern harm → kill shame default. Apology-only equals one-CTA on follow-through → revisit CTA quality, not honesty. Pre-register HOLETRUST-* before “complete visual diagnostic” parent campaigns.

CXXXIX.12 So what for MindCraft commercially

- **Copy:** “Empty format slots are honesty — we only diagnose sealed visuals.” Kill complete-visual ads and child-deficit hole shame.
- **Product:** Parent probed/not_probed list + seal reason + one CTA + Practice/C4 use-tier split.
- **Growth:** CBC-test honesty vs hide/pad; sell trust of labeled incompleteness; pipe holes into promote-to-C4 narrative (id 140).
- **Positioning:** Against portals that fake completeness and against coverage shame leagues; align with COVER matrix honesty for families.
- **Metric:** hole_shown rate; hide_attempt; CTA→Practice; Map-contradiction tickets — demote Hole Trust Score™.
- **Kill list:** Hide holes; shame frames; complete-visual ads; Hole Trust Score™; absence≡mastery; fear-scarcity WTP; ACT guarantees from hole packaging.
- **Vision:** Parents trust The Desk because the gap-scan admits what FormatIds were not measured — so when sealed stock arrives, the fill path is believed.

References (verified)

- Bergman, P. (2021). Parent-child information frictions and human capital investment: Evidence from a field experiment. *Journal of Political Economy*, 129(1), 286–322. <https://doi.org/10.1086/711410>
 - Bergman, P., & Chan, E. W. (2021). Leveraging parents through low-cost technology: The impact of high-frequency information on student achievement. *Journal of Human Resources*, 56(1), 125–158. <https://doi.org/10.3368/jhr.56.1.1118-9837R1>
 - Federal Trade Commission. (1984). FTC policy statement regarding advertising substantiation. In *Thompson Medical Co.*, 104 F.T.C. 648, 839.
 - Kane, M. T. (2013). Validating the interpretations and uses of test scores. *Journal of Educational Measurement*, 50(1), 1–73. <https://doi.org/10.1111/jedm.12000>
 - Kraft, M. A., & Rogers, T. (2015). The underutilized potential of teacher-to-parent communication: Evidence from a field experiment. *Economics of Education Review*, 47, 49–63. <https://doi.org/10.1016/j.econedurev.2015.04.001>
 - Mayer, R. C., Davis, J. H., & Schoorman, F. D. (1995). An integrative model of organizational trust. *Academy of Management Review*, 20(3), 709–734. <https://doi.org/10.5465/amr.1995.9508080335>
 - Rawlins, B. (2008). Measuring the relationship between organizational transparency and employee trust. *Public Relations Journal*, 2(2), 1–21.
 - Rogers, T., & Feller, A. (2018). Reducing student absences at scale by targeting parents’ misbeliefs. *Nature Human Behaviour*, 2, 335–342. <https://doi.org/10.1038/s41562-018-0328-1>
 - Schnackenberg, A. K., & Tomlinson, E. C. (2016). Organizational transparency: A new perspective on managing trust in organization-stakeholder relationships. *Journal of Management*, 42(7), 1784–1810. <https://doi.org/10.1177/0149206314525202>
 - Xu, J., Guo, S., Feng, Y., Ma, Y., Zhang, Y., Núñez, J. C., & Fan, X. (2024). Parental homework involvement and students’ mathematics achievement: A meta-analysis. *Psicothema*, 36(1), 1–14. <https://doi.org/10.7334/psicothema2023.92>
-

Part CXL — Promote-to-C4 Tutor QA Workflow for Practice-Only FormatId

Chapter status: Living evidence + bank/ops QA brief — Researcher tick 2026-08-13 (UTC hour 18 ≡ Red Team slot, but ch140 never written → prefer Researcher per rotation; researcher count since synthesizer v1.17 = 6 → Researcher)

Primary question: When a FormatId item is sealed for **Reveal-Practice** but not yet **C4 hide-correctness**, what **tutor/SME QA workflow** must flip hcelig_pass — without play-count auto-promote, thumbs-up theater, fluency soft-pass, or Promote Score™ vanity — so Map writes stay defensible while Practice stock still grows?

Owners: Product (pools / admin tickets) · Engine (eligibility flags) · HITL / tutor ops · GENQ / FIGKEY / REPAIR · COVER / Brand · Red Team

Commercial job: Ship **SAFE-PROMOTE** densifying SAFE-SPLITPOOL × SAFE-REPAIR × SAFE-FIGKEY × SAFE-HCELIG × SAFE-COVER × content-review science: **ticketed, checklisted, logged promotion** — never engagement flips, never silent dual-use, never “N plays clean ⇒ diagnose.”

Builds on: Parts CXXXVII (SAFE-SPLITPOOL — foreshadows ticketed promote), CXXXIV (SAFE-HCELIG), CXXXI (SAFE-FIGKEY), LX (SAFE-REPAIR), LXXII (SAFE-GENQ), LXVIII (SAFE-HITL), LXXXVII (SAFE-COVER), CXXXVIII (SAFE-SCANCOMP), CXXXIX (SAFE-HOLETRUST). Seams: promote ticket UI, hcelig_pass write with reviewer id, demote-on-repair, Practice-only vs C4 coverage tables, parent hole fill ETA honesty.

CXL.1 Why this chapter exists

SAFE-SPLITPOOL ordered **split pools** and named promote-to-C4 as a workflow, not a folder rename. SAFE-HOLETRUST told parents holes fill via **company seal work**, not shame. The ops gap remains:

1. **Play-count flip** — “≥50 Practice attempts, zero tickets ⇒ hcelig_pass.”
2. **Thumbs-up promote** — student/tutor “looks fine” as sole human gate.
3. **Fluency soft-pass** — pretty SVG / LLM judge / “traps mapped” overrides wrong geometry.
4. **Idle backlog** — Practice-only stock piles; C4 holes never shrink; honesty becomes permanent apology.
5. **Promote Score™ / QA Minutes NS** — vanity that rewards throughput over seal fidelity.
6. **Ticketed promote with seals + human checklist (this chapter)** — REPAIR evidence, FIGKEY/GENQ re-verify, SME congruence, logged reviewer, demote faster than promote.

Without SAFE-PROMOTE, split pools decay into either forever-Practice (coverage theater forever) or silent auto-admit (Map poison). Commercially: parents who bought labeled holes deserve a **real fill path**; tutors who flag broken figures deserve a **queue that changes eligibility**, not a Discord shrug.

FOUNDER BELIEF under audit: Near-peer tutors can run a **bounded promote checklist** (not PhD psychometrics theater) when the checklist is seal-first, item-model constrained, and demote-on-repair is automatic — and that workflow is the only honest answer to “when do diagram holes close?”

Claims we refuse:

1. Auto-promote from play counts, clean-streak, or engagement velocity.
 2. Thumbs-up / “feels ACT-ready” as sole human promote gate.
 3. Soft-pass wrong figure because distractors were trap-mapped or render looked fluent.
 4. Promote Score™ / QA Minutes / Promotion Throughput % as North Star.
 5. Silent hcelig_pass flip without reviewer id + checklist log.
 6. Padding C4 from Practice-only stock while tickets are open or seals stale.
 7. ACT / identity / “complete visual diagnostic” guarantees from promote packaging.
-

CXL.2 Constructs

Construct	Meaning	MindCraft analogue	Failure mode
Practice-only sealed	GENQ (+ FIGKEY) pass; reveal + REPAIR; no Map diagnostic write	practice_eligible && !hcelig_pass	Treated as C4

Promote ticket	Named work item to upgrade use ambition	Admin/tutor queue card	Chat vibe / play-count job
Seal re-verify	Re-run GENQ/FIGKEY/FORMTRAP gates at promote time	Checklist hard fails	Stale seal from generation day
SME congruence	Human judges item $\leftrightarrow$ FormatId $\leftrightarrow$ concept use	Tutor/SME promote panel	Beauty / “would teach with this” only
Logged admit	hcelig_pass + reviewer id + timestamp + checklist hash	Audit trail	Silent bit flip
Demote-on-repair	Critical ticket removes C4 immediately	Auto yank	Leave in gap-scan while broken
SAFE-PROMOTE	Ticketed Practice $\rightarrow$ C4 QA doctrine	This chapter	Promote Score™

Operational definition (HYPOTHESIS): A Practice-only FormatId item may be **promoted to C4** iff (1) GENQ text/key seals still pass, (2) FIGKEY passes when a figure is required, (3) FORMTRAP seals exist for visual options when applicable, (4) **zero open critical REPAIR tickets**, (5) a trained reviewer completes a **fixed checklist** (stem–figure–key joint truth; FormatId structural necessity; accessibility/alt path; no beauty-trap options), (6) optional second rater or CVR-style “essential for C4 use” mark when ambition is parent-facing diagnosis, (7) hcelig_pass write is **logged** with reviewer id — never engagement thresholds alone — and (8) any later critical repair **demotes** before the next gap-scan draw.

CXL.3 Uses scale; promotion is a validity event

FACT: Kane (2013, *Journal of Educational Measurement*, 50(1), 1–73, doi:10.1111/jedm.12000) — validate **interpretations and uses**; more ambitious claims need more backing; rejecting a use need not kill a weaker interpretation.

FACT: AERA, APA, & NCME (2014, *Standards for Educational and Psychological Testing*) — validity is evidence for score interpretations **for proposed uses**; each intended use needs relevant evidence.

Applied (HYPOTHESIS): Reveal-Practice is a **lower-ambition use** (learning + repair). C4 $\rightarrow$ Map $\rightarrow$ parent “we diagnosed diagram on X” is **higher ambition**. Promotion is not “we liked the item more.” It is an explicit decision that the item now authorizes the stricter use. Auto-promote from Practice volume confuses popularity with use validity.

Kill: Play-count $\equiv$ readiness for hide-correctness.

Survive: Separate IUAs; promote as a logged use-upgrade.

CXL.4 Content QA is expert review + procedure — not vibes

FACT: Downing (2006, in Downing & Haladyna, *Handbook of Test Development*, pp. 3–25) — effective test development is a systematic multi-step process; item creation and review are distinct procedures that build validity evidence for intended interpretations — not ad-hoc shipping.

FACT: Haladyna, Downing, & Rodriguez (2002, *Applied Measurement in Education*, 15(3), 309–333, doi:10.1207/S15324818AME1503_5) — validated taxonomy of MC item-writing guidelines (content, format, style, stem, choices); content concerns and choice quality are reviewable, not “feels fine.”

FACT: Sireci (1998, *Social Indicators Research*, 45(1–3), 83–117, doi:10.1023/A:1006985528729) — content representation remains fundamental; expert judgment of relevance/representation is a core content-validity practice even under unitary validity language.

FACT: Lawshe (1975, *Personnel Psychology*, 28(4), 563–575, doi:10.1111/j.1744-6570.1975.tb01393.x) — Content Validity Ratio aggregates SME “essential / useful / not necessary” judgments; consensus among experts is quantifiable — **transfer limit:** employment testing origin; MindCraft borrows the *panel judgment* idea for “essential for C4 diagnostic use,” not job-analysis theater or CVR-as-marketing badge.

Applied (HYPOTHESIS): Promote checklist = Downing-style review step + Haladyna content/choice gates + optional Lawshe-style essential-for-C4 mark. Thumbs-up without checklist is not SME review. CVR Score™ splash is banned; a logged essential bit for high-ambition admits is allowed.

Kill: Beauty panel / Discord emoji promote.

Survive: Fixed checklist + optional essential-for-C4 consensus for parent-facing diagnostic stock.

CXL.5 AIG review audits the generation procedure — and re-checks at promote

Reuse (SAFE-FIGKEY / SAFE-GENQ): Gierl & Lai (2012; 2016) line — item models constrain features; review content/logic in the **procedure**, not eyeball-thousands as the QA unit.

Applied (HYPOTHESIS): Generation-day figkey_pass is necessary but not sufficient at promote time. Figures drift (asset swaps, trap edits, stem wording). Promote re-runs deterministic seals. LLM-as-judge-alone remains banned (SAFE-FIGKEY / SAFE-EXPLAINQA).

Kill: “Already verified once at gen” forever admit.

Survive: Re-verify at promote; demote if seals fail later.

CXL.6 Repair is evidence — absence of tickets is not proof at low volume

Reuse (SAFE-REPAIR / SAFE-COVER): Reveal Practice exists so broken figures surface. Critical open tickets **block** promote. But low Practice volume with zero tickets ≠ “proven safe” — coverage honesty still labels holes (SAFE-SCANCOMP / SAFE-HOLETRUST).

HYPOTHESIS: Ticket absence is a **necessary** promote gate, not a **sufficient** one. Promote still requires seal re-verify + human checklist. Parent fill copy may say “verifying more diagram items” with known-finite ETA bands when backlog is real (SAFE-KNOWNETA), never “coming soon” fog or play-count progress bars as trust theater.

Kill: “No complaints $\Rightarrow$ C4.”

Survive: No-critical-tickets $\wedge$ seals $\wedge$ checklist.

CXL.7 HITL tutors run the queue — experts escalate hard cases

Reuse (SAFE-HITL / SAFE-TUTORGRAIN): Near-peer tutors own Map-brief fidelity, not PhD item-writing credentials as brand. Promote QA is a **packetized ops skill**: checklist training, sample audit, escalate geometry disputes to SME/engine.

HYPOTHESIS: Default promote rater = trained tutor/content ops under HITL playbook; second rater or SME required when (a) figure IR is ambiguous, (b) parent-facing diagnosis ambition, or (c) prior demote on the item. Tutor time on promote counts as bank quality work — not “hours booked” NS (SAFE-WORKFORCE / SAFE-ROI).

Kill: Ivy $\equiv$ promote authority; Discord-peer $\equiv$ item review.

Survive: Trained checklist + escalate packets.

CXL.8 Claim table

Claim	Label	Confidence	Notes
Uses scale; higher ambition needs more evidence	FACT	High	Kane 2013; Standards 2014
Systematic item review is a test-development step	FACT	High	Downing 2006
MC content/choice guidelines are reviewable criteria	FACT	High	Haladyna et al. 2002
Expert judgment of content relevance/representation matters	FACT	High	Sireci 1998
SME essentialness can be aggregated (CVR)	FACT	High	Lawshe 1975; transfer limits apply
Ticketed checklist promote beats play-count auto-admit	HYPOTHESIS	Medium–High	Densifies SPLITPOOL $\times$ REPAIR $\times$ FIGKEY
Zero tickets at low volume proves C4 safety	SPECULATION	Low	Refuse
Thumbs-up $\equiv$ SME content review	SPECULATION	Low	Refuse

Parents accept slower C4 fill if promote is real and holes stay labeled	FOUNDER BELIEF	Medium	Needs PROMOTE-* × HOLETRUST
---	----------------	--------	-----------------------------

CXL.9 Product surface — SAFE-PROMOTE claim contract

1. **Promote ticket** required to set `hceLig_pass` (no silent admin bit without ticket id).
 2. Hard gates: GENQ + FIGKEY (if figured) + FORMTRAP (if visual options) **re-verify at promote**.
 3. Hard gate: **zero open critical REPAIR** on the item.
 4. Checklist required fields: stem-figure-key joint; FormatId necessity; alt/accessibility; trap sanity; “essential for C4?” (yes/no).
 5. Log: reviewer id, second-rater if required, checklist hash, seal versions, promote timestamp.
 6. **Demote-on-repair**: critical ticket clears `hceLig_pass` before next C4 draw; Practice-reveal may remain if policy allows.
 7. Coverage UI: Practice-only vs C4 stock separate; promote backlog visible to ops, not as student Promote Score™.
 8. Parent fill CTA: links to sealed Practice and/or “verifying for diagnosis” with known-finite ETA when known — never play-count progress as mastery.
 9. Telemetry: promote/demote events, post-promote repair rate, illicit Map writes — demote Promote Score™ / QA Minutes NS.
 10. Copy: “We hide the key only after a sealed review — practice can start earlier.”
-

CXL.10 Doctrine — SAFE-PROMOTE (provisional)

1. **Promotion is a use-upgrade**, not a popularity contest (Kane).
2. **Ticket + checklist + seals** — never play-count / thumbs-only.
3. **Re-verify at promote** — generation-day pass can go stale.
4. **No critical REPAIR open** — necessary, not sufficient.
5. **Demote faster than promote** — Map integrity > hole optics.
6. **HITL-trained raters**; escalate hard geometry — no Ivy costume, no Discord≡panel.
7. **Honest dual coverage** — Practice stock ≠ C4 stock (COVER / SPLITPOOL).
8. **No Promote Score™ / QA Minutes / Promotion Throughput % NS**.
9. **No ACT / complete-visual guarantees** from promote packaging.
10. Copy: “Practice lights on first; diagnosis after the stricter seal.”

Confidence: High on Kane/Standards/Downing/Haladyna/Sireci/Lawshe as method sources (with Lawshe transfer limits). Medium on exact checklist length, second-rater thresholds, and promote SLA vs hole ETA copy (needs PROMOTE-*). High that play-count auto-promote is commercially toxic (Map distrust + HOLETRUST betrayal).

CXL.11 Experiments

ID	Question	Design	Primary
PROMOTE-1	Ticketed checklist vs play-count auto-promote	Ops A/B	Post-promote critical repair rate; false C4 admits
PROMOTE-2	Re-verify-at-promote vs trust generation-day seals only	Engine A/B	Stale-seal catch rate; promote latency
PROMOTE-3	Single trained tutor vs dual-rater for figured C4 admits	Ops A/B	Disagree rate; post-promote repair; cost/min
PROMOTE-4	Demote-on-repair latency SLA (immediate vs daily batch)	Engine A/B	Illicit C4 draws on broken items; Map↔tutor gold
PROMOTE-5	Parent CBC: “verified for diagnosis” ETA honesty vs play-count fill bar	CBC	WTP; trust; abandon (× HOLETRUST)
PROMOTE-QUAL	Planted wrong-slope figure with zero tickets at low volume: does checklist catch before C4?	Qual	Promote codebook

Falsifiers: Auto-promote equals ticketed on post-promote repairs *and* Map↔tutor gold → still prefer ticketed for liability/audit; do not celebrate engagement. Dual-rater never disagrees and never catches extras → drop mandatory second rater for low ambition; keep for parent-facing diagnosis. Re-verify finds ~0 stale fails for 90 days → narrow re-verify to figure-touched edits only. Pre-register PROMOTE-* before “AI grows our diagnostic bank automatically” campaigns.

CXL.12 So what for MindCraft commercially

- **Copy:** “We practice on sealed figures before we hide the key.” Parent: “Diagram diagnosis opens only after a stricter review — holes stay labeled until then.”
- **Product:** Promote ticket UI; checklist; logged `hcelig_pass`; demote-on-repair; dual coverage tables.
- **Growth:** Sell the fill path as **ops integrity**, not coverage theater; CBC honesty vs play-count progress bars.
- **Positioning:** Against ChatGPT one-blob “practice=diagnose” fluency *and* against edtech that flips hide-correctness from engagement metrics.
- **Metric:** post-promote repair rate; demote latency; illicit C4 draws — demote Promote Score™.
- **Kill list:** Play-count flip; thumbs-only; fluency soft-pass; silent bit flips; Promote Score™; ACT/complete-visual guarantees from promote packaging.

- **Vision:** Maya’s Map only learns from probes a named reviewer would defend — while her Practice gym can still grow FormatId skill with the lights on, and her parent sees holes shrink for real reasons.
-

References (verified)

- American Educational Research Association, American Psychological Association, & National Council on Measurement in Education. (2014). *Standards for educational and psychological testing*. AERA.
 - Downing, S. M. (2006). Twelve steps for effective test development. In S. M. Downing & T. M. Haladyna (Eds.), *Handbook of test development* (pp. 3–25). Lawrence Erlbaum Associates.
 - Haladyna, T. M., Downing, S. M., & Rodriguez, M. C. (2002). A review of multiple-choice item-writing guidelines for classroom assessment. *Applied Measurement in Education*, 15(3), 309–333. https://doi.org/10.1207/S15324818AME1503_5
 - Kane, M. T. (2013). Validating the interpretations and uses of test scores. *Journal of Educational Measurement*, 50(1), 1–73. <https://doi.org/10.1111/jedm.12000>
 - Lawshe, C. H. (1975). A quantitative approach to content validity. *Personnel Psychology*, 28(4), 563–575. <https://doi.org/10.1111/j.1744-6570.1975.tb01393.x>
 - Sireci, S. G. (1998). The construct of content validity. *Social Indicators Research*, 45(1–3), 83–117. <https://doi.org/10.1023/A:1006985528729>
-

Part CXLI — Exposure Cool-Down After Reveal Before Same-Stem C4

Chapter status: Living evidence + bank/selector brief — Researcher tick 2026-08-13 (UTC hour 21; hour%6≠0; researcher count since synthesizer v1.17 = 7 → Researcher)

Primary question: After a student sees a **revealed** Practice stem (key/lights on), when — if ever — may that **same stem** enter that student’s **C4 hide-correctness** gap-scan — without answer-hunt Map poison, Cool-Down Score™ vanity, or thin-pool starvation that forces pad theater?

Owners: Product (Practice / gap-scan selectors) · Engine (per-student exposure / cool-down clocks) · Bank / COVER / SCANCOMP · Brand · Red Team

Commercial job: Ship **SAFE-COOLDOWN** densifying SAFE-SPLITPOOL × SAFE-SCANCOMP × SAFE-CALIB × SAFE-COLD × pool-exposure science: **per-student cool-down after reveal before same-stem C4** — never immediate dual-use, never forever global bans, never Cool-Down Minutes as North Star.

Builds on: Parts CXXXVII (SAFE-SPLITPOOL — foreshadows cool-down), CXXXVIII (SAFE-SCANCOMP — exposure caps under thin stock), CXXXIV (SAFE-HCELIG), LXXXI (SAFE-COLD), L (SAFE-CALIB), CXXV (SAFE-CONFMISS), LXXII (SAFE-GENQ), LXXXVII (SAFE-COVER), CXL (SAFE-PROMOTE). Seams: revealed_at / cooldown_until per (student, stem_id), isomorphic FormatId substitutes, hole flags when cool-down empties a cell, parent/student copy that does not call memory “diagnosis.”

CXLI.1 Why this chapter exists

SAFE-SPLITPOOL named cool-down as a contract line; SAFE-SCANCOMP capped thin-stock exposure. The selector gap remains: **immediate same-stem C4** after reveal (key memory dressed as diagnosis); **forever/global bans** that starve FormatId cells into pad theater; **Cool-Down Score™**; **clone leak** (same figure, swapped digits). Doctrine: policy cool-down + isomorphic substitutes + honest holes — anti-answer-hunt without infinite-pool cosplay.

FOUNDER BELIEF under audit: Parents prefer a **shorter honest probe set** (or labeled hole) over a green Map cell earned by remembering tonight’s key — honesty as WTP, not a tax.

Claims we refuse: immediate same-stem C4 after reveal; forever/global locks as default; Cool-Down Score™ / Exposure Minutes NS; pad because cool-down emptied a cell; clone leak as “fresh”; ACT / complete-visual guarantees from cool-down packaging.

CXLI.2 Constructs

Construct	Meaning	MindCraft analogue	Failure mode
Reveal event	Student saw keyed feedback / worked path on a stem	Practice reveal; Solver dump (out of scope unless same stem)	Silent dual-use without clock
Same-stem	Identical <code>item_id</code> (and sealed isomorphic family)	Stem + figure + key family id	Clone leak
Per-student cool-down	Bar that student’s C4 draw of family until <code>cooldown_until</code>	Selector hard constraint	Global lock / forever ban
Isomorphic substitute	Different stem, same <code>FormatId</code> ×concept ambition	Bank siblings under SCANCOMP	Same figure, swapped digits only
Answer-hunt signal	Latency/accuracy pattern consistent with key recall	Telemetry companion	Shame feed
SAFE-COOLDOWN	Reveal→C4 exposure doctrine	This chapter	Cool-Down Score™

Operational definition (HYPOTHESIS): After a **reveal-class** administration of stem family F to student S , S ’s C4 selector **must not** draw any member of F until $\text{now} \geq \text{cooldown_until}(S, F)$. Default policy window is a **session-scale to multi-day band** tuned by COOLDOWN- (*not a marketing “24h science” badge*). During the window, prefer **isomorphic substitutes outside F** for the same `FormatId` cell; if none exist, emit `format_not_probed` / hole honesty (SAFE-SCANCOMP / SAFE-HOLETRUST) — **never** pad, **never** lift cool-down early to fill quotas. Cool-down is **per student**, not global. Exact duration is experiment-owned; the existence* of a hard bar is doctrine.

CXLI.3 Disclosure + reuse makes items easier — same-stem after reveal is the extreme case

FACT: Appelhaus, Werner, Grosse, & Kämmer (2023, *Medical Education Online*, 28(1), 2143298, doi:10.1080/10872981.2022.2143298) — in a large European medical school that began disclosing MCQs with keys, reused **disclosed** items were significantly easier (M difficulty coefficient 0.83) than reused non-disclosed (0.71) and new items (0.66); disclosure aids feedback/transparency but can cost item reliability when disclosed items are reused.

Applied (HYPOTHESIS): MindCraft Practice **intentionally discloses** keys for learning and REPAIR. That is pedagogically good (Butler & Roediger below) and psychometrically expensive for **same-stem diagnostic reuse**. Immediate C4 on a just-revealed stem is not “efficient bank use”; it is the Appelhaus problem at session timescale. Transfer limits: medical end-of-term exams $\neq$ teen gap-scan; the directional claim (disclosed reuse $\neq$ fresh measurement) still holds.

Kill: “They already practiced it, so diagnose on it now.”

Survive: Reveal is a different use; cool-down before same-family C4.

CXLI.4 Feedback after MC strengthens memory of the key — that is the point of Practice, the poison for C4

FACT: Butler & Roediger (2008, *Memory & Cognition*, 36(3), 604–616, doi:10.3758/MC.36.3.604) — feedback after multiple-choice testing increases later correct responses and reduces lure intrusions vs no feedback; educators should provide feedback when using MCQs.

FACT: Roediger & Marsh (2005, *Journal of Experimental Psychology: Learning, Memory, and Cognition*, 31(5), 1155–1159, doi:10.1037/0278-7393.31.5.1155) — multiple-choice testing can produce positive testing effects **and** negative suggestion from exposure to lures; selecting wrong alternatives can create false knowledge.

Applied (HYPOTHESIS): Practice reveal is designed to **install correct keys** and reduce lure damage. C4 hide-correctness is designed to measure **unassisted** competence + confidence without answer-hunting (SAFE-CALIB / SAFE-COLD). Re-administering the same stem minutes later with the key hidden measures **retention of tonight’s feedback**, not FormatId diagnosis. That confounds Map updates and CONFMISS recovery (confident false weakness / false strength).

Kill: Hide-correctness on a stem whose key the product just taught.

Survive: Cool-down; probe with a different family member or leave a hole.

CXLI.5 Spacing science says delay changes what a second encounter measures — it does not license “perfect cool-down AI”

FACT: Cepeda, Pashler, Vul, Wixted, & Rohrer (2006, *Psychological Bulletin*, 132(3), 354–380, doi:10.1037/0033-2909.132.3.354) — meta-analysis of distributed practice; inter-study interval and retention interval jointly affect final retention; optimal gaps grow with desired retention interval.

Applied (HYPOTHESIS, careful transfer): Cepeda et al. justify **that lag matters**, not a branded “scientifically perfect 47-hour cool-down.” For MindCraft, cool-down is a **validity gate for use ambition** (Kane: different interpretation), borrowing the insight that immediate re-test after feedback is a different cognitive event than a delayed unassisted probe. Do not ship Cool-Down AI / expanding-SRS theater as the product hero (SAFE-ADAPT / SAFE-SCHED bans still apply).

Kill: Perfect-interval cool-down marketing.

Survive: Policy band + A/B; isomorphic substitutes; hole honesty.

CXLI.6 Pool integrity and exposure control are ops policies — not folder aesthetics

FACT: Way (1998, *Educational Measurement: Issues and Practice*, 17(4), 17–27, doi:10.1111/j.1745-3992.1998.tb00632.x) — computerized programs must actively protect item-pool integrity; exposure and pool-management strategies matter under continuous delivery.

FACT: Georgiadou, Triantafillou, & Economides (2007, *Journal of Technology, Learning and Assessment*, 5(8)) — review of item-exposure control strategies for CAT (1983–2005); overexposure of popular items and underexposure of others are managed with explicit control procedures (including probabilistic approaches such as Simpson–Hetter–style methods discussed in that literature).

FACT: van der Linden, Ariel, & Veldkamp (2006, *Journal of Educational and Behavioral Statistics*, 31(1), 81–99, doi:10.3102/10769986031001081) — principled pool assembly / constrained selection improves control of content and exposure properties vs unconstrained greedy selection.

Applied (HYPOTHESIS, careful transfer): MindCraft is not selling high-stakes CAT security. Transferable: **exposure is a first-class selector constraint**. Per-student cool-down after reveal is the educational analogue of exposure control for *use compromise* (key known → hide-correctness pretends novelty). Under thin C4 stock (SAFE-SCANCOMP), constraints fail open into **holes**, not same-family reuse or Practice-only pad. Per-student family clocks + isomorphic preference beat global Simpson–Hetter theater.

Kill: Unconstrained “best item” draw that ignores recent reveal.

Survive: Hard cool-down; constrained assembly with honest holes.

CXLI.7 Thin pools: cool-down must not recreate pad theater

Reuse (SAFE-SCANCOMP / SAFE-HOLETRUST / SAFE-COVER): When sealed C4 stock for a FormatId cell is N=1–2, cool-down after Practice reveal can empty the cell for that student.

HYPOTHESIS: Empty cell ⇒ `format_not_probed` + parent/student honesty, **or** an out-of-family isomorphic C4 item if seals allow — **not** (a) lift cool-down, (b) pad with symbolic, (c) dual-write from Practice reveal as if it were C4. Promote backlog (SAFE-PROMOTE) remains the fill path for holes, not cool-down violation.

Kill: “Cool-down made the scan incomplete, so reuse the revealed graph.”

Survive: Hole label; promote more C4 stock; isomorphic siblings.

CXLI.8 Claim table

Claim	Label	Confidence	Notes
Disclosed reused MCQs become easier than new/undisclosed reused	FACT	High	Appelhaus et al. 2023
Feedback after MC strengthens correct retention / reduces lure damage	FACT	High	Butler & Roediger 2008
MC lures can create false knowledge without feedback	FACT	High	Roediger & Marsh 2005
Spacing/lag jointly affect what a later test measures	FACT	High	Cepeda et al. 2006
Pool exposure needs active management under continuous delivery	FACT	High	Way 1998; Georgiadou et al. 2007; van der Linden et al. 2006 (transfer limits)
Per-student cool-down after reveal before same-family C4	HYPOTHESIS	Medium–High	Densifies SPLITPOOL × SCANCOMP × CALIB
Exact cool-down hours are settled science	SPECULATION	Low	Refuse badge; run COOLDOWN-*
Forever ban / global lock is required	SPECULATION	Low	Refuse
Parents prefer honest holes over answer-hunt greens	FOUNDER BELIEF	Medium	Needs COOLDOWN-* × HOLETRUST

CXLI.9 Product surface — SAFE-COOLDOWN claim contract

1. On reveal-class event for stem family F and student S , set `cooldown_until(S, F)` (hard selector bar for C4).
2. C4 selector: **exclude** F for S while cool-down active; Practice may still use other families (and may re-practice F with reveal if pedagogy wants — that extends the clock).
3. Prefer **out-of-family isomorphic** C4 substitutes for the same `FormatId`×concept cell.
4. If no substitute $\Rightarrow$ emit hole / `format_not_probed` — **never** pad; **never** early-lift cool-down for quota.
5. Cool-down is **per student**, never global from one reveal.

6. Define **family** wider than `item_id` when figure+key path is shared (clone seal).
7. Telemetry: `reveal_at`, `cooldown_until`, `blocked_draw_count`, `substitute_used`, `hole_due_to_cooldown` — demote Cool-Down Score™ / Exposure Minutes.
8. Copy: “We don’t re-quiz the same problem with the answer hidden right after we show the key.” Parent: “Diagnosis waits until a fresh problem — holes beat memory greens.”
9. No ACT / complete-visual guarantees from cool-down packaging.
10. Companion anti-gaming: rapid Practice-reveal→C4 attempts logged as answer-hunt signal for ops — not student shame feed.

CXLI.10 Doctrine — SAFE-COOLDOWN (provisional)

1. **Reveal changes the use** — disclosed practice ≠ fresh hide-correctness evidence (Appelhaus; Kane ambition).
2. **Per-student cool-down** before same-family C4 — hard bar, not soft suggestion.
3. **Family > item_id** when clones share figure/key path.
4. **Isomorphic substitutes first**; holes second; pad never.
5. **No forever ban / no global lock** as default doctrine.
6. **Duration is experiment-owned** — no perfect-interval Cool-Down AI brand.
7. **Densifies SPLITPOOL × SCANCOMP** — exposure caps become clocks, not vibes.
8. **No Cool-Down Score™ / Exposure Minutes / Lock Duration % NS.**
9. **No ACT / identity / complete-visual guarantees** from cool-down packaging.
10. Copy: “Fresh probe or honest hole — never memory dressed as diagnosis.”

Confidence: High on Appelhaus/Butler–Roediger/Roediger–Marsh/Cepeda/Way/Georgiadou/van der Linden as method sources (with CAT/medical transfer limits). Medium on exact hours, family-hash rules, and isomorphic-generation capacity (needs COOLDOWN-*). High that immediate same-stem C4 after reveal is commercially toxic (Map distrust + CALIB poison).

CXLI.11 Experiments

ID	Question	Design	Primary
COOLDOWN-1	Same-stem C4 allowed immediately after reveal vs hard cool-down	Selector A/B	False mastery / false weakness vs tutor gold; answer-hunt latency
COOLDOWN-2	Cool-down 0 / ~1 session / ~3 days / ~7 days bands	Engine A/B	Map↔tutor agreement; hole rate; abandon
COOLDOWN-3	Exact <code>item_id</code> cool-down vs family-hash (figure+key path)	Engine A/B	Clone leak rate; pool starvation

COOLDOWN-4	Isomorphic substitute vs hole when cool-down empties FormatId cell	Product A/B	Format-route quality; HOLETRUST WTP
COOLDOWN-5	Parent CBC: “fresh probe” honesty vs “we reuse for efficiency”	CBC	WTP; trust; abandon
COOLDOWN-QUAL	Planted: student revealed graph G then sees G in C4 — does Map write look like format mastery?	Qual	Codebook / tutor review

Falsifiers: Immediate same-stem equals cool-down on Map↔tutor gold *and* no answer-hunt signals → audit instrumentation before collapsing cool-down; still prefer cool-down for liability/IUA clarity. Forever ban beats timed cool-down on gold without starving FormatId practice → reconsider duration upward, not pad. Family-hash never catches clones beyond item_id → simplify to item_id + human clone tagging. Pre-register COOLDOWN-* before “instant practice→diagnose on the same problem” campaigns.

CXLI.12 So what for MindCraft commercially

- **Copy:** “We don’t hide the answer on the same problem we just showed you.” Parent: “Format diagnosis uses a fresh probe — or we leave a hole.”
- **Product:** Per-student family cool-down clocks; isomorphic preference; hole emission; clone-family seals.
- **Growth:** Sell **measurement honesty** vs competitors that recycle revealed stems into “adaptive diagnostics.”
- **Positioning:** Against ChatGPT one-blob practice=test fluency *and* against edtech that treats bank reuse as free diagnosis.
- **Metric:** blocked same-family draws; hole_due_to_cooldown; Map↔tutor gold — demote Cool-Down Score™.
- **Kill list:** Immediate same-stem C4; forever/global lock defaults; pad under cool-down; Cool-Down Score™; ACT/complete-visual guarantees.
- **Vision:** Maya’s Map records what she can do **without tonight’s key in working memory** — while Practice still gets to teach with the lights on.

References (verified)

- Appelhaus, S., Werner, S., Grosse, P., & Kämmer, J. E. (2023). Feedback, fairness, and validity: Effects of disclosing and reusing multiple-choice questions in medical schools. *Medical Education Online*, 28(1), 2143298. <https://doi.org/10.1080/10872981.2022.2143298>
- Butler, A. C., & Roediger, H. L., III. (2008). Feedback enhances the positive effects and reduces the negative effects of multiple-choice testing. *Memory & Cognition*, 36(3), 604–616. <https://doi.org/10.3758/MC.36.3.604>
- Cepeda, N. J., Pashler, H., Vul, E., Wixted, J. T., & Rohrer, D. (2006). Distributed practice in verbal recall tasks: A review and quantitative synthesis. *Psychological Bulletin*, 132(3), 354–380. <https://doi.org/10.1037/0033-2909.132.3.354>

- Georgiadou, E. G., Triantafyllou, E., & Economides, A. A. (2007). A review of item exposure control strategies for computerized adaptive testing developed from 1983 to 2005. *Journal of Technology, Learning and Assessment*, 5(8). <https://ejournals.bc.edu/index.php/jtla/article/view/1647>
- Roediger, H. L., III, & Marsh, E. J. (2005). The positive and negative consequences of multiple-choice testing. *Journal of Experimental Psychology: Learning, Memory, and Cognition*, 31(5), 1155–1159. <https://doi.org/10.1037/0278-7393.31.5.1155>
- van der Linden, W. J., Ariel, A., & Veldkamp, B. P. (2006). Assembling a computerized adaptive testing item pool as a set of linear tests. *Journal of Educational and Behavioral Statistics*, 31(1), 81–99. <https://doi.org/10.3102/10769986031001081>
- Way, W. D. (1998). Protecting the integrity of computerized testing item pools. *Educational Measurement: Issues and Practice*, 17(4), 17–27. <https://doi.org/10.1111/j.1745-3992.1998.tb00632.x>

Part CXLII — Student-Facing Hole Copy after Gap-Scan (No Shame, No Fake Greens)

Chapter status: Living evidence + student UX/copy brief — Researcher tick 2026-08-14 (UTC hour 00 ≡ Red Team slot, but ch142 never written → prefer Researcher per rotation; researcher count since synthesizer v1.17 = 8 → Researcher override)

Primary question: After gap-scan leaves honest FormatId holes (format_not_probed), what **student-facing** summary copy should MindCraft ship — without ability shame, fake greens, empty growth-mindset posters, or Student Hole Score™ vanity — so the learner revises the right belief (“this format was not measured”) and takes one Practice fill path?

Owners: Product (gap-scan / Practice summary) · Brand · COVER / SCANCOMP · CALIB / COLD · Parent surfaces (HOLETRUST parity) · Red Team

Commercial job: Ship **SAFE-STUDHOLE** densifying SAFE-HOLETRUST × SAFE-COLD × SAFE-CALIB × SAFE-SOFTMSG × SAFE-FBLIT × formative feedback science: **same honesty as parents, student-owned next step, product-hole frame** — never parent-only truth, never pad greens for teens, never “you’re bad at diagrams.”

Builds on: Parts CXXXIX (SAFE-HOLETRUST), CXXXVIII (SAFE-SCANCOMP), LXXXI (SAFE-COLD), L (SAFE-CALIB), CXVIII (SAFE-SOFTMSG), CXII (SAFE-FBLIT), LXXXVII (SAFE-COVER), CXXXIV (SAFE-HCELIG), CXXXVII (SAFE-SPLITPOOL), CXLI (SAFE-COOLDOWN). Seams: post-gap-scan student summary, probed vs not_probed chips, one Practice CTA, dual-script ban vs parent MoC.

CXLII.1 Why this chapter exists

SAFE-HOLETRUST locked **parent** hole honesty. Students still see a different risk surface: hide holes so teens “feel complete”; paint unprobed cells green (absence=mastery — SAFE-COLD kill); frame holes as visual talent deficits; dump jargon matrices; or plaster empty “mistakes make you stronger!” without a next attempt. Parent-only honesty fails commercially when Maya opens the same scan and sees fake greens while a tutor later says “we never probed diagram.”

FOUNDER BELIEF under audit: Teens trust a shorter honest map (“diagram not checked yet”) plus one Practice mission more than a complete-looking rainbow that collapses the first time a tutor opens Map.

Claims we refuse:

1. Parent-only hole labels; student summary hides `format_not_probed`.
2. Fake greens / “you’re flexible across formats” for unprobed cells.
3. Ability / visual-talent shame (“you’re weak at graphs”).
4. Empty growth-mindset poster without grain + CTA.
5. Student Hole Score™ / Completeness % / Honesty Minutes as North Star.
6. Multi-link maze or forced marathon before leaving summary.
7. ACT / identity guarantees from student hole packaging.

CXLII.2 Constructs

Construct	Meaning	MindCraft analogue	Failure mode
Student hole summary	Post-scan view of probed vs not-probed <code>FormatIds</code>	Gap-scan result screen	Parent-only honesty
Product-hole frame	Unmeasured format ≠ failed format	Copy: “Not checked yet”	“You’re bad at diagrams”
Fake green	Unprobed cell rendered as mastery/flexibility	Map chrome / chip color	Absence≡mastery
Parity script	Student and parent share hole facts	Same <code>not_probed</code> list	Dual script / teen pad
One fill CTA	Single next Practice grain for a hole or nearest sealed sibling	Soft-require transfer pattern	Maze / hard lock-in
SAFE-STUDHOLE	Student-facing hole copy doctrine	This chapter	Student Hole Score™

Operational definition (HYPOTHESIS): Student hole copy is SAFE-STUDHOLE-complete when it (1) lists probed vs `not_probed` `FormatIds` in plain language, (2) states the seal reason in one teen-readable sentence, (3) never colors unprobed cells as mastered/flexible, (4) frames the hole as **product measurement limit**, not character/ability, (5) pairs with **exactly one** primary CTA (sealed Practice where stock exists; else “we’ll open more when sealed” without fake ETA theater), (6) matches parent HOLETRUST facts (parity), and (7) is A/B-tested against hide/pad/shame/mindset-poster arms (STUDHOLE-*).

CXLII.3 Transparency still = disclosure + clarity + accuracy — for the learner as stakeholder

FACT: Schnackenberg and Tomlinson (2016, *Journal of Management*, 42(7), 1784–1810, doi:10.1177/0149206314525202) — transparency is perceived quality of intentionally shared information along **disclosure, clarity, and accuracy**, theorized as antecedents to trustworthiness (ability, benevolence, integrity) in Mayer, Davis, and Schoorman’s trust model.

FACT: Mayer, Davis, and Schoorman (1995, *Academy of Management Review*, 20(3), 709–734, doi:10.5465/amr.1995.9508080335) — trust is willingness to be vulnerable based on expectations of ability, benevolence, and integrity.

Applied (HYPOTHESIS): Teens are stakeholders of the diagnostic, not merely subjects of a parent portal. Hide (low disclosure), jargon dump (low clarity), and fake greens (low accuracy) each break a different trustworthiness path — especially integrity when Map later contradicts the rainbow.

Failure mode	Transparency break	Student trust risk
Hide holes	Low disclosure	“They lied / later Map shock”
Coverage % fog	Low clarity	Abandon reading; ask parent what it means
Fake greens	Low accuracy	Inflated confidence; tutor contradiction
Mindset poster only	Clarity without actionable accuracy	Belief unchanged; no FEI attempt

Kill: Teen-facing “complete” theater while parents see holes (or the reverse).

Survive: Same facts, age-appropriate clarity, one next step.

CXLII.4 Formative use of incomplete evidence — do not grade the hole

FACT: Black and Wiliam (1998, *Phi Delta Kappan*, 80(2), 139–148) — formative assessment is assessment whose evidence is used to adapt teaching and learning; the point is feedback that modifies next action, not a finished grade of the person.

Applied (HYPOTHESIS): A `format_not_probed` flag is **incomplete formative evidence**, not a score. Student copy must recruit the next Practice adaptation (“try a sealed diagram mission when ready”) rather than close a judgment (“you failed visuals”). Transfer limit: classroom AfL ≠ consumer app — borrow the *use* of evidence, not Black–Wiliam effect-size ads (SAFE-FA ban still holds).

Kill: Hole-as-grade / Completeness % as student NS.

Survive: Hole-as-open formative slot + feed-forward CTA.

CXLII.5 Feedback levels: task/process/where-next — not self-level ability talk

FACT: Hattie and Timperley (2007, *Review of Educational Research*, 77(1), 81–112, doi:10.3102/003465430298487) — feedback addresses where am I going / how am I going / where to next, at task, process, self-regulation, and self levels; **self-level** feedback (person praise/criticism) is typically least effective for learning.

Applied (HYPOTHESIS): Prefer task/process language: “Diagram format: not checked yet (we only hide keys on verified visuals). Next: one sealed Practice item on linear graphs.” Avoid self-level: “You’re not a visual learner” / “You’re so smart you don’t need diagram practice” (both poison).

Kill: Self-level hole shame or self-level fake praise.

Survive: Task fact + where-next Practice grain (densifies SAFE-FBLIT / SAFE-SOFTMSG).

CXLII.6 Ability labels and empty “growth” talk both fail — strategy + attempt, not posters

FACT: Mueller and Dweck (1998, *Journal of Personality and Social Psychology*, 75(1), 33–52, doi:10.1037/0022-3514.75.1.33) — intelligence praise (vs effort/process) orients children toward performance goals and, after failure, more low-ability attributions, less persistence, and worse performance.

FACT: Yeager, Hanselman, Walton, et al. (2019, *Nature*, 573, 364–369, doi:10.1038/s41586-019-1466-y) — a short growth-mindset intervention can help some students’ challenge-seeking and grades **conditionally** (notably lower-achieving students in supportive peer norms); it is not a universal poster fix.

Applied (HYPOTHESIS): Framing holes as fixed visual talent recreates Mueller/Dweck ability-label damage. Framing holes as “just believe you can grow” without a sealed Practice CTA is empty mindset theater already banned in lab standing rules. Yeager et al. license **challenge-seeking scaffolds in supportive contexts**, not “Growth Hole Score™” splash copy. Student hole UX should invite a concrete attempt, not a belief sermon.

Kill: Ability-shame holes; poster-only growth.

Survive: Process invitation — one sealed attempt when stock exists.

CXLII.7 Calibration honesty: do not inflate greens to “raise confidence”

Reuse (SAFE-CALIB / SAFE-COLD): Appropriate confidence beats inflated confidence; humble priors + labeled evidence age beat day-one mastery fireworks.

HYPOTHESIS: Painting unprobed FormatIds green to soothe post-scan affect is a **confidence inflation** product bug. It manufactures false strength that CONFMISS recovery and Map later have to unwind — and teaches teens the product lies. Prefer destaked honesty: “Not measured” chips that do not update mastery.

Kill: Fake greens as affect management.

Survive: Unprobed ≠ weak and ≠ strong; no mastery write from absence.

CXLII.8 Claim table

Claim	Label	Confidence	Notes
Transparency ≈ disclosure + clarity + accuracy	FACT	High	Schnackenberg & Tomlinson 2016
Trust rests on ability/benevolence/integrity expectations	FACT	High	Mayer et al. 1995
Formative assessment uses evidence to adapt next learning	FACT	High	Black & Wiliam 1998; transfer limits on magnitudes
Self-level feedback is typically weakest for learning	FACT	High	Hattie & Timperley 2007
Intelligence/ability praise harms post-failure persistence	FACT	High	Mueller & Dweck 1998
Growth-mindset effects are conditional, not poster-universal	FACT	High	Yeager et al. 2019
Student labeled holes + one CTA > hide/pad/shame/poster	HYPOTHESIS	Medium–High	Needs STUDHOLE-*
Teens prefer honest short maps over fake-complete rainbows	FOUNDER BELIEF	Medium	Qual + A/B; liability if Map contradicts
Student Hole Score™ improves identity	SPECULATION	Low	Refuse
Uses/ads must match sealed stock ambition	FACT	High	Kane 2013 line (reuse HOLETRUST/COVER)

CXLII.9 Product surface — SAFE-STUDHOLE claim contract

1. Gap-scan end screen shows **Checked** vs **Not checked yet** FormatIds (same flags as parent HOLETRUST).
2. One teen-readable **seal reason** (“We only use verified visuals when we hide the answer”).
3. Unprobed chips are **neutral** / **open**, never green-mastery or red-failure.
4. Copy frames **product hole**, not visual talent; ban classmate FormatId ranks on this screen.
5. Exactly **one** primary CTA: sealed Practice grain for a hole (or nearest sealed sibling) — soft-require pattern (SAFE-SOFTGATE), not hard lock-in by default.

6. If no sealed Practice stock → “Not open yet — we’re verifying more” without fake-precision ETA (KNOWNETA discipline).
7. **Parity**: student and parent hole lists must not contradict; tutor brief uses the same facts.
8. Ban empty mindset splash; if encouragement appears, it must sit beside grain + CTA.
9. Telemetry: student_hole_shown, fake_green_flag, shame_frame_flag, cta→practice_start, dual_script_mismatch — demote Student Hole Score™ / Completeness %.
10. No mastery / flexibility Map write from not_probed.

CXLII.10 Doctrine — SAFE-STUDHOLE (provisional)

1. **Students see holes** — never parent-only honesty.
2. **Triad copy** — disclosure + clarity + accuracy (Schnackenberg & Tomlinson).
3. **Hole ≠ grade** — formative open slot (Black & Wiliam use), not person score.
4. **Task/process + where-next** — no self-level ability talk (Hattie & Timperley).
5. **No ability shame / no poster-only growth** (Mueller & Dweck; Yeager conditionality).
6. **No fake greens** — unprobed ≠ mastery (SAFE-COLD / SAFE-CALIB).
7. **One Practice CTA**; soft-require default; parity with HOLETRUST.
8. **No Student Hole Score™ / Completeness % / Honesty Minutes NS**.
9. **No ACT / identity guarantees** from hole packaging.
10. Copy spine: “We checked what we can seal. Empty slots mean we have not measured that format yet — not that you failed it. When a sealed Practice item is ready, try one.”

Confidence: High on feedback/transparency spines with stated transfer limits. Medium on exact teen wording and CTA compliance (needs STUDHOLE-*). High that hide/pad/shame/fake-green are commercially toxic once Map/tutor contradict the student summary.

CXLII.11 Experiments

ID	Question	Design	Primary
STUDHOLE-1	Labeled holes vs hidden vs fake-complete greens	Student copy A/B	Trust items; later Map-contradiction anger; false-strength rate
STUDHOLE-2	Product-hole frame vs ability-shame frame	Copy A/B	Challenge-seeking; help-avoidance; affect without Anxiety Score™
STUDHOLE-3	One Practice CTA vs poster-only vs multi-link maze	CTA A/B	cta→practice_start; follow-through; quit

STUDHOLE-4	Student–parent parity vs dual script (teen pad / parent honesty)	Matched pairs	Dual-script mismatch tickets; family trust items
STUDHOLE-5	Neutral open chips vs red-failure vs green-mastery for not_probed	UI A/B	Misbelief “I failed/mastered visuals”; CONFMIS confound rate
STUDHOLE-QUAL	Maya interviews: which summary would you believe next week?	Qual	Codebook honesty vs completeness preference

Falsifiers: Hide/fake-green beats labeled on short-run completion *and* Map distrust is equal → still refuse for IUA/liability; do not celebrate. Shame raises Practice starts but help-avoidance/Xu-pattern harm → kill shame default. Poster-only equals one-CTA on follow-through → fix CTA quality, not honesty. Pre-register STUDHOLE-* before teen “complete visual diagnostic” campaigns.

CXLII.12 So what for MindCraft commercially

- **Copy:** “Empty slots = not measured yet — not that you failed.” Kill fake greens, ability shame, and parent-only honesty.
- **Product:** Student probed/not_probed list + seal reason + neutral chips + one Practice CTA; parity with parent MoC.
- **Growth:** Position against complete-looking dashboards that gaslight teens; sell trust that survives tutor Map opens.
- **Positioning:** Same COVER honesty language for families *and* students; wedge vs apps that paint coverage theater.
- **Metric:** student_hole_shown; fake_green_flag; CTA→Practice; dual-script mismatch — demote Student Hole Score™.
- **Kill list:** Parent-only holes; fake greens; ability shame; empty mindset posters; Student Hole Score™; absence≡mastery; ACT guarantees from hole packaging.
- **Vision:** Maya trusts The Desk because the gap-scan admits what FormatIds were not measured — so the first sealed Practice fill feels like truth, not marketing.

References (verified)

- Black, P., & Wiliam, D. (1998). Inside the black box: Raising standards through classroom assessment. *Phi Delta Kappan*, 80(2), 139–148.
- Hattie, J., & Timperley, H. (2007). The power of feedback. *Review of Educational Research*, 77(1), 81–112. <https://doi.org/10.3102/003465430298487>
- Mayer, R. C., Davis, J. H., & Schoorman, F. D. (1995). An integrative model of organizational trust. *Academy of Management Review*, 20(3), 709–734. <https://doi.org/10.5465/amr.1995.9508080335>
- Mueller, C. M., & Dweck, C. S. (1998). Praise for intelligence can undermine children’s motivation and performance. *Journal of Personality and Social Psychology*, 75(1), 33–52.

<https://doi.org/10.1037/0022-3514.75.1.33>

• Schnackenberg, A. K., & Tomlinson, E. C. (2016). Organizational transparency: A new perspective on managing trust in organization-stakeholder relationships. *Journal of Management*, 42(7), 1784–1810.

<https://doi.org/10.1177/0149206314525202>

• Yeager, D. S., Hanselman, P., Walton, G. M., Murray, J. S., Crosnoe, R., Muller, C., ... Dweck, C. S. (2019). A national experiment reveals where a growth mindset improves achievement. *Nature*, 573, 364–369.

<https://doi.org/10.1038/s41586-019-1466-y>

Part CXLIII — Promote Backlog SLA vs Parent Hole-ETA Honesty

Chapter status: Living evidence + ops/parent-copy brief — Researcher tick 2026-08-14 (UTC hour 12 ≡ Red Team slot, but ch143 never written → prefer Researcher per rotation; researcher count since synthesizer v1.18 = 0 → Researcher)

Primary question: When FormatId gap-scan holes stay open because Practice-only items await **ticketed promote-to-C4**, what **ops SLA + parent/student fill-ETA** must MindCraft ship — without play-count fill bars, “coming soon” fog, optimistic under-estimates that slip into integrity debt, or FilleTA Score™ vanity — so hole honesty remains actionable while promote QA stays seal-first?

Owners: Product (ParentDashboard / student hole CTA / ops backlog) · HITL / bank QA · Brand · COVER / HOLETRUST / STUDHOLE · PROMOTE / KNOWNETA · WTP · Red Team

Commercial job: Ship **SAFE-FILLETA** densifying SAFE-PROMOTE × SAFE-KNOWNETA × SAFE-HOLETRUST × SAFE-STUDHOLE × wait-information science: **finite fill bands grounded in real promote backlog capacity** — never marketing clocks, never engagement progress as “diagnosis arriving,” never SLA pressure that soft-passes seals.

Builds on: Parts CXL (SAFE-PROMOTE), CXXXV (SAFE-KNOWNETA), CXXXIX (SAFE-HOLETRUST), CXLII (SAFE-STUDHOLE), CXXXVIII (SAFE-SCANCOMP), LXXXVII (SAFE-COVER), CXXXVII (SAFE-SPLITPOOL), LIX (SAFE-WTP). Seams: promote queue age/depth, parent “verifying more diagram items” ETA, student “not open yet” band, sealed Practice alternate CTA, demote-on-repair vs fill-date optics.

CXLIII.1 Why this chapter exists

SAFE-PROMOTE defined *how* an item earns C4. SAFE-HOLETRUST / SAFE-STUDHOLE defined *how* families see holes. SAFE-KNOWNETA defined finite ETAs for **late tutors**. The missing join:

> Parents who buy labeled holes ask **when** a FormatId cell becomes diagnosable. Ops who refuse play-count promote still owe a **clock story** that is neither fog nor a lie that forces rushed soft-pass admits.

Wrong resolutions already visible in edtech and marketplace tutoring:

1. **“Coming soon” forever** — Maister uncertain wait applied to bank fill.
2. **Optimistic fill dates** — announce “this week,” miss, repeat “almost” (pilot insult).
3. **Play-count / % complete fill bars** — engagement theater as ETA.

4. **SLA theater that soft-passes** — checklist burned to hit a parent-facing date; Map poison.
5. **Apology-primary without magnitude** — thin coverage apology, no band, no Practice path.
6. **FilleETA Score™ / Backlog Minutes NS** — vanity that rewards short announced ETAs.

Without SAFE-FILLETA, HOLETRUST becomes permanent apology and PROMOTE becomes idle backlog or date-driven corruption. The trust brand that labels holes must also **time-box the fill path honestly**.

FOUNDER BELIEF under audit: Serious parents prefer a **known-finite verify band** (“about 2–4 weeks for sealed diagram diagnosis on this concept, capacity permitting”) plus sealed Practice now, over “soon” or a guaranteed calendar date that ops cannot defend without soft-pass.

Claims we refuse:

1. “Coming soon” / “almost ready” as default hole-fill copy.
2. Play-count, thumbs, or Practice velocity as parent-facing fill progress.
3. Optimistic fill ETAs that systematically under-shoot actual promote latency.
4. Parent-facing dates that authorize soft-pass / skip-checklist promote.
5. FilleETA Score™ / Backlog Minutes / Honesty Clock NS.
6. Apology-only hole CTAs with no magnitude and no Practice alternate.
7. ACT / “complete visual diagnostic by [date]” guarantees from fill packaging.

CXLIII.2 Constructs

Construct	Meaning	MindCraft analogue	Failure mode
Promote backlog	Ticketed Practice→C4 work items waiting seals+checklist	Ops promote queue	Invisible / Discord vibe
Backlog SLA (internal)	Ops targets for ticket age, re-verify, demote latency	HITL bank QA policy	Sold as parent guarantee
Fill-ETA band (external)	Parent/student-facing finite verify window	Hole CTA clock copy	Fake precision; optimistic slip
Capacity-grounded	Band from queue depth × reviewer capacity × median latency	Ops estimator	Marketing calendar
Practice alternate	Sealed reveal Practice while C4 pending	Soft-require CTA	Idle wait-only
SAFE-FILLETA	Hole-fill ETA honesty doctrine	This chapter	FilleETA Score™

Operational definition (HYPOTHESIS): Hole-fill copy is SAFE-FILLETA-complete when (1) if ops can estimate promote capacity for that FormatId/concept cell, UI shows a **finite band** (days/weeks grain, not second-precision), (2) the band is derived from **backlog depth + reviewer capacity + historical median promote latency** (and never from play counts), (3) if capacity is unknown or blocked (critical REPAIR storm, zero raters), copy **labels uncertainty** and

keeps sealed Practice / “not open yet” without inventing a date, (4) on slip, refresh with a **new finite band + timestamp** — never repeat “almost,” (5) parent/student bands **do not contradict**, (6) hitting a band **never** waives GENQ/FIGKEY/checklist/REPAIR gates, and (7) instruments are fill_eta_honesty_ok (stated band vs first C4-eligible ship for that hole class) × CTA→Practice × trust/WTP — never Backlog Minutes as learning NS.

CXLIII.3 Uncertain fill waits feel longer than known, finite ones

FACT: Maister (1985, “The Psychology of Waiting Lines,” in Czepiel, Solomon, & Surprenant, *The Service Encounter*, Lexington Books) — **uncertain waits feel longer than known, finite waits**; “soon” produces nervous anticipation; repeated “only a few more minutes” after a miss adds insult because customers are delayed *and* not dealt with honestly.

Applied (HYPOTHESIS): Parent hole fill is a **multi-day/week wait**, not an 8-minute lobby — but the psychology transfers: “we’re working on diagram items” without a bound is Maister’s waiting room. “Diagnosis almost ready” after a missed week is the pilot insult at product timescale. Finite bands (“about 2–4 weeks”) let families plan sealed Practice and tutor goals without gaslight.

Kill: Soon-default on hole fill.

Survive: Finite band when estimable; honesty > optimism.

Wound: Airport wait ≠ bank QA weeks — grain differs; uncertainty tax still loads.

CXLIII.4 Duration information beats empty reassurance — at intermediate horizons

FACT: Hui & Tse (1996, *Journal of Marketing*, 60(2), 81–90, doi:10.1177/002224299606000206) — waiting-**duration** information vs queuing information; effects on acceptability/affect mediate service evaluation; duration information matters especially in **intermediate** waits; very short waits gain little; very long waits need more than prettier clocks.

Applied (HYPOTHESIS): Typical promote latency (days–few weeks) sits in an *intermediate product horizon* where **duration bands help**. Queue-position cosplay (“you are #47 in diagram QA”) is weak for parents and risks shame/competition — prefer duration bands + one Practice CTA. For chronically long backlogs, escalate: hire/capacity plan + labeled scarcity (SAFE-COVER), not fake-short ETAs.

Kill: Skip fill ETA because “we already labeled the hole.”

Survive: Duration band for intermediate promote latency; escalate when backlog is structural.

Wound: Lab wait minutes ≠ school calendar — calibrate bands on ops data (FILLETA-1).

CXLIII.5 Informing about anticipated delay changes behavior — accuracy matters

FACT: Whitt (1999, *Management Science*, 45(2), 192–207, doi:10.1287/mnsc.45.2.192) — informing customers about anticipated delays shifts **reneging** toward **balking**; when providers choose to inform, **accurate prediction** of

delays matters for managing the system.

Applied (HYPOTHESIS): Parent-facing fill ETAs change behavior: some families start sealed Practice now (good balk into the Practice rail); some delay purchase/renewal if the band is absurdly long (honest balk). Optimistic false short bands produce later renegeing (churn + trust debt) when the date slips. Commercially prefer accurate-conservative bands that induce **Practice alternate uptake** over short lies that spike short-run conversion.

Kill: Fake-short fill dates for conversion.

Survive: Accurate-enough bands; instrument honesty vs conversion separately.

Wound: Call-center queue $\neq$ edtech bank — transfer as decision architecture, not throughput formula in ads.

CXLIII.6 Affect mediates delay evaluation — control and filled time matter

FACT: Taylor (1994, *Journal of Marketing*, 58(2), 56–69, doi:10.1177/002224299405800205) — delays affect service evaluations **mediated by negative affect** (anger, uncertainty); perceived provider **control** and **filled time** influence evaluations indirectly via those affects.

Applied (HYPOTHESIS): Hole-fill delay must (1) cut uncertainty (finite band or labeled ignorance), (2) fill time with a **session-related** alternate (sealed Practice for that FormatId/sibling — STUDHOLE/HOLETRUST CTA), (3) attribute control honestly (“seal/verify capacity,” not “Maya isn’t visual enough”). Fear scarcity (“buy before holes widen”) recreates CRUNCHMSG/WTP kills.

Kill: Character-blame + empty wait.

Survive: Uncertainty↓ + Practice fill + ops-capacity attribution.

Wound: Airline delay field $\neq$ parent MoC — mechanism transfer only.

CXLIII.7 Progress information beats apology loops — without fake progress bars

FACT: Munichor & Rafaeli (2007, *Journal of Applied Psychology*, 92(2), 511–521, doi:10.1037/0021-9010.92.2.511) — in telephone waits, **queue-location** / **progress information** beat music and **apologies** on abandonment and evaluations; sense of **progress** (not only perceived duration) mediated reactions.

Applied (HYPOTHESIS): Parent hole UX may show **ops-progress class** signals that are true: “in verify queue,” “blocked on figure repair,” “promoted — ready for diagnosis” — not apology spam. **Ban** play-count bars and % “almost diagnosed” meters that are engagement cosplay. True progress = ticket state transitions under SAFE-PROMOTE, optionally coarse (“queued / in review / blocked / shipped”), never vanity throughput %.

Kill: Apology-primary; play-count fill bar.

Survive: Coarse true ticket-state progress + duration band.

Wound: Phone-queue position $\neq$ promote ticket depth — do not show raw “#12” if it induces rank shame; prefer state + band.

CXLIII.8 Transparency triad still governs the clock story

Reuse (SAFE-HOLETRUST): Schnackenberg & Tomlinson (2016) — transparency $\approx$ disclosure + clarity + accuracy. Soft-pass-to-hit-date and play-count clocks break **accuracy/integrity**; jargon SLA dumps break **clarity**; hiding the clock while showing holes breaks **disclosure**.

Kill: Clock theater that sacrifices seals.

Survive: Short triad — what's not diagnosed, verify band or labeled unknown, one Practice CTA, seals unchanged.

CXLIII.9 Claim table

Claim	Label	Confidence	Notes
Uncertain waits feel longer than known finite waits	FACT	High	Maister 1985
Duration info helps intermediate waits	FACT	High	Hui & Tse 1996
Delay announcements change join/abandon behavior; accuracy matters	FACT	High	Whitt 1999
Delay→evaluation often affect-mediated; control & filled time matter	FACT	High	Taylor 1994
Progress info > apology fillers (queue contexts)	FACT	High	Munichor & Rafaeli 2007
Capacity-grounded fill bands + Practice CTA > soon/play-count/soft-pass	HYPOTHESIS	Medium–High	Needs FILLETA-*
Parents prefer honest multi-week bands over fake-short dates	FOUNDER BELIEF	Medium	CBC + qual
FILLETA Score™ improves learning	SPECULATION	Low	Refuse
Internal backlog SLA may be stricter than external band	HYPOTHESIS	Medium	Ops vs parent grain

CXLIII.10 Product surface — SAFE-FILLETA claim contract

1. Hole CTA may show a **finite verify band** when ops estimator has capacity signal; else **labeled uncertainty**.
 2. Band source = backlog depth × reviewer capacity × median promote latency — **never** play counts / thumbs.
 3. Grain = days/weeks (coarse); ban second-precision countdown theater.
 4. Prefer **conservative** bands; under-promise/over-deliver on *ship date*, never on learning claims.
 5. On slip: new band + timestamp; ban repeating “almost / any day now.”
 6. Coarse true ticket **state** optional (“queued / in review / blocked / shipped”); ban % play-progress.
 7. **Seals unchanged by date pressure** — no soft-pass / skip-checklist to hit Filleta.
 8. Always offer **one Practice alternate** when sealed stock exists; if none, “not open yet” without fake greens (STUDHOLE).
 9. Parent and student bands **parity** — no dual script.
 10. Telemetry: `fill_eta_honesty_ok`, `fill_eta_slip_count`, `softpass_for_date_flag`, CTA→Practice — demote Filleta Score™ / Backlog Minutes NS.
-

CXLIII.11 Doctrine — SAFE-FILLETA (provisional)

1. **Hole honesty includes a clock story** — finite band or labeled unknown (Maister × HOLETRUST).
2. **Bands are capacity-grounded** — backlog SLA feeds the estimator; marketing calendars do not.
3. **Internal SLA ≠ parent guarantee** — ops targets may be tighter; external copy stays conservative bands.
4. **KNOWNETA refresh rules** apply at product timescale — no optimistic slip loops.
5. **Progress = ticket state**, not engagement meters (Munichor without play-count cosplay).
6. **Filled wait = Practice alternate** (Taylor filled time × STUDHOLE CTA).
7. **PROMOTE integrity > ETA optics** — never soft-pass for the clock.
8. **No Filleta Score™ / Backlog Minutes / Honesty Clock NS.**
9. **No ACT / complete-visual-by-date guarantees.**
10. Copy spine: “Diagram diagnosis for this topic: not sealed yet. We’re verifying items — about 2–4 weeks at current capacity. Meanwhile, try sealed Practice with the answer shown when available.”

Confidence: High on wait-information spines with stated transfer limits. Medium on exact band widths, estimator formula, and when to show state vs band-only (FILLETA-*). High that soon/play-count/soft-pass-for-date are commercially toxic under FEI + Map trust.

CXLIII.12 Experiments

ID	Question	Design	Primary
FILLETA-1	Finite capacity band vs “coming soon” under equal actual promote latency	Parent copy A/B	Trust; rebook/WTP; anger/uncertainty items

FILLETA-2	Conservative band vs optimistic short that slips	Ops A/B	fill_eta_honesty_ok; churn; complaint
FILLETA-3	Band + Practice CTA vs band-only vs apology-only	CTA A/B	CTA→Practice; abandon; trust
FILLETA-4	Ticket-state progress vs play-count fill bar vs no progress	UI A/B	Trust; perceived progress; softpass pressure tickets
FILLETA-5	Parent CBC: honest multi-week verify vs “complete visual by Friday” package	CBC	WTP; trust (× HOLETRUST/WTP)
FILLETA-QUAL	10 parents: which fill clock would you believe after one miss?	Qual	FILLETA codebook

Falsifiers: “Soon” beats finite band on short-run conversion *and* long-run trust equal → still refuse for integrity/IUA; do not celebrate. Optimistic short wins first week then honesty_ok collapses → prefer conservative. Date pressure raises promote throughput but post-promote repair spikes → kill soft-pass path; keep seals. Pre-register FILLETA-* before “complete visual diagnostic by [date]” campaigns.

CXLIII.13 So what for MindCraft commercially

- **Copy:** “Not sealed yet — about 2–4 weeks to verify at current capacity. Practice can start earlier with the key shown.”
- **Product:** Ops backlog estimator → parent/student fill band; ticket-state optional; Practice alternate; parity.
- **Growth:** Sell **timed honesty** as trust packet — against fake-complete dashboards and against forever-apology holes.
- **Positioning:** Against ChatGPT “coverage now” fluency *and* against edtech fill bars powered by engagement.
- **Metric:** fill_eta_honesty_ok; slip count; softpass-for-date flag; CTA→Practice — demote Filleta Score™.
- **Kill list:** Coming-soon default; play-count fill bars; optimistic slip loops; date-driven soft-pass; Filleta Score™; ACT/complete-by-date guarantees.
- **Vision:** Maya’s family sees empty FormatId slots *and* a believable verify window — so when diagram diagnosis finally lights, it feels earned by seals, not by a marketing calendar.

References (verified)

- Hui, M. K., & Tse, D. K. (1996). What to tell consumers in waits of different lengths: An integrative model of service evaluation. *Journal of Marketing*, 60(2), 81–90. <https://doi.org/10.1177/002224299606000206>
- Maister, D. H. (1985). The psychology of waiting lines. In J. A. Czepiel, M. R. Solomon, & C. F. Surprenant (Eds.), *The service encounter* (pp. 113–123). Lexington Books.

- Munichor, N., & Rafaeli, A. (2007). Numbers or apologies? Customer reactions to telephone waiting time fillers. *Journal of Applied Psychology*, 92(2), 511–521. <https://doi.org/10.1037/0021-9010.92.2.511>
- Taylor, S. (1994). Waiting for service: The relationship between delays and evaluations of service. *Journal of Marketing*, 58(2), 56–69. <https://doi.org/10.1177/002224299405800205>
- Whitt, W. (1999). Improving service by informing customers about anticipated delays. *Management Science*, 45(2), 192–207. <https://doi.org/10.1287/mnsc.45.2.192>

Part CXLIV — Second-Rater Thresholds for Figured C4 Admits

Chapter status: Living evidence + ops/QA brief — Researcher tick 2026-08-14 (UTC hour 15; hour%6≠0; researcher count since synthesizer v1.18 = 1 → Researcher)

Primary question: When promoting a **figured** Practice-only FormatId item to **C4 hide-correctness**, when is a **second human rater** mandatory vs optional — without Dual Rater Score™ vanity, always-dual cost theater, never-dual soft-pass, or medical-cosplay “double reading” ACT ads — so Map writes stay defensible while FILLETA capacity stays honest?

Owners: HITL / bank QA · Product (promote tickets) · Engine (eligibility flags) · FIGKEY / GENQ / PROMOTE · COVER / FILLETA · Brand · Red Team

Commercial job: Ship **SAFE-DUALRATE** densifying SAFE-PROMOTE × SAFE-FIGKEY × SAFE-HITL × SAFE-FILLETA × agreement science: **risk-triggered dual review** — never universal dual as brand theater, never single-rater forever on parent-facing figured diagnosis, never kappa dashboards as North Star.

Builds on: Parts CXL (SAFE-PROMOTE — foreshadows optional second rater / PROMOTE-3), CXXXI (SAFE-FIGKEY), CXLIII (SAFE-FILLETA), LXVIII (SAFE-HITL), LXXI (SAFE-TUTORGRAIN), LXXII (SAFE-GENQ), CXXXIV (SAFE-HCELIG), CXXXVII (SAFE-SPLITPOOL). Seams: promote checklist “second_rater_required” bit, discord/arbitration packet, rater capacity in fill-ETA estimator, parent “stricter review” copy without Dual Rater Score™.

CXLIV.1 Why this chapter exists

SAFE-PROMOTE foreshadowed an **optional** second rater (ambiguous IR, parent-facing ambition, prior demote) and reserved PROMOTE-3 for single vs dual. Missing: the **threshold rule**. Without it, ops drifts into always-dual costume (FILLETA explosion + “double-checked” ads), never-dual soft-pass (pretty SVG → Map poison), Dual Rater Score™ vanity, mammography cosplay ACT theater, second-rater override of failed seals, or ETA optics that skip dual when backlog is hot.

Parents who buy labeled FormatId holes deserve **stricter review when ambition is high** — not theater dual on every MC or silent single-rater admits on diagrams that write Map.

FOUNDER BELIEF under audit: Near-peer tutors can be first raters under a fixed checklist; a **second rater or SME escalate** is load-bearing for figured C4 admits that authorize parent-facing “we diagnosed diagram on X” — dual is a **trigger**, not a brand costume.

Claims we refuse: always-dual everywhere; never-dual on figured parent-facing C4; Dual Rater Score™ / Kappa % NS; seal-waiving dual; mammography-cosplay guarantees; ETA-skipped dual; Discord-emoji=second rater; ACT/complete-visual guarantees from dual packaging.

CXLIV.2 Constructs

Construct	Meaning	MindCraft analogue	Failure mode
First rater	Trained HITL/tutor completing promote checklist	Default promote reviewer	Ivy costume; Discord vibe
Second rater	Independent checklist pass (or SME escalate) before hcelig_pass	Dual admit	Same person twice; emoji panel
Trigger rule	Conditions that flip dual from optional→required	second_rater_required	Always-on / never-on
Discord resolution	Disagree on essential-for-C4 or stem-figure-key	Escalate / reject / repair ticket	Average vibes; majority of untrained
Agreement audit	Periodic dual on a sample for process health	Ops IRR sample	Kappa Score™ NS
SAFE-DUALRATE	Risk-triggered dual-review doctrine	This chapter	Dual Rater Score™

Operational definition (HYPOTHESIS): A promote is SAFE-DUALRATE-complete when (1) seals (GENQ + FIGKEY if figured + FORMTRAP if visual options) still pass and no critical REPAIR is open (SAFE-PROMOTE), (2) a **written trigger rule** decides single vs dual *before* the ticket is worked — not after ETA panic, (3) dual is **mandatory** iff at least one trigger fires (see CXLIV.7), else single trained rater may admit, (4) when dual is required, second rater is **independent** (different person; checklist hash logged; no rubber-stamp “agree” without opening the figure IR), (5) disagreement → **reject or repair**, never soft-pass, (6) periodic sample duals measure process IRR for ops (not student-facing scores), and (7) dual capacity is an input to FILLETA bands — never waived to shrink a parent ETA.

CXLIV.3 Higher-ambition uses need more backing — dual is use evidence, not politeness

FACT: Kane (2013, *Journal of Educational Measurement*, 50(1), 1–73, doi:10.1111/jedm.12000) — validate **interpretations and uses**; more ambitious claims need more backing; rejecting a use need not kill a weaker interpretation.

FACT: AERA, APA, & NCME (2014, *Standards for Educational and Psychological Testing*) — each intended use needs relevant evidence.

Applied (HYPOTHESIS): Reveal-Practice is lower ambition. C4→Map→parent “diagram diagnosed on concept X” is higher. Dual review is additional **use evidence** when the item’s pictorial ground truth is hard to re-check after the fact — not a courtesy ritual and not required for every low-ambition text stem.

Kill: Dual as brand politeness; single rater for all ambitions.

Survive: Dual intensity scales with use ambition and figure risk.

CXLIV.4 Chance-corrected agreement is a process tool — not a product North Star

FACT: Cohen (1960, *Educational and Psychological Measurement*, 20(1), 37–46, doi:10.1177/001316446002000104) — percent agreement ignores chance concordance; kappa adjusts for chance on nominal categories.

FACT: McHugh (2012, *Biochemia Medica*, 22(3), 276–282, doi:10.11613/BM.2012.031, PMC3900052) — kappa is widely used for IRR; Landis/Koch-style “moderate” bands can be too lenient for high-stakes contexts; many texts treat percent agreement below ~80% and kappa below ~0.60 as inadequate for serious decisions (healthcare framing; transfer with caution).

FACT: NCES / NAEP Technical Documentation (Constructed-Response Interrater Reliability; last updated 2024) — second scoring is used to monitor rater consistency; kappa targets cited as higher than .7 for two-/three-point items and higher than .6 for four- to six-point items; items below threshold **prompt investigation** — criteria are NAEP-design-specific, not universal product law.

Applied (HYPOTHESIS): MindCraft should **sample dual** promote decisions to watch disagree/repair catch rates and (for yes/no essential-for-C4) chance-corrected agreement. Failures trigger rater retraining or checklist rewrite — not Dual Rater Score™ splash, not “our kappa is 0.9 so Map is true,” not student-facing Agreement Personality™.

Kill: Kappa / percent-agreement as marketing NS or ACT proof.

Survive: IRR as ops health signal with investigation thresholds (inspired by NAEP’s investigate-below-target pattern; do not copy NAEP cutoffs as brand law).

CXLIV.5 Expert “essential” judgment can be dual — CVR is method, not badge

Reuse (SAFE-PROMOTE): Lawshe (1975, *Personnel Psychology*, 28(4), 563–575, doi:10.1111/j.1744-6570.1975.tb01393.x) — Content Validity Ratio aggregates SME essential/useful/not-necessary marks; panel size affects critical values; employment-testing origin → **transfer limits**.

Applied (HYPOTHESIS): For parent-facing figured C4, the checklist’s “essential for C4 diagnostic use?” bit is the right dual object — not beauty. Two raters marking essential independently is closer to Lawshe’s panel idea than one rater’s thumbs-up. Full CVR panels of 5–40 SMEs are **cost theater** for every FormatId stem at startup scale; reserve multi-SME panels for archetype templates / first-of-cell blueprints, not every isomorphic clone.

Kill: CVR Score™ splash; 15-SME panel per item as default.

Survive: Dual essential-bit on high-ambition figured admits; larger panels only for template/archetype gates.

CXLIV.6 Double reading elsewhere: gains, costs, policy — transfer carefully

FACT: Brown, Bryan, & Warren (1996, *BMJ*, 312(7034), 809–812, doi:10.1136/bmj.312.7034.809) — in one UK screening unit, consensus double reading detected more cancers than single reading and could reduce recall/total cost vs single; non-consensus double reading raised cost.

FACT: Dinnes, Moss, Melia, Blanks, Song, & Kleijnen (2001, *The Breast*, 10(6), 455–463, doi:10.1054/brst.2001.0350) — systematic review: double reading can raise detection (~3–11 per 10,000 screened); recall impact depends on **recall policy** (consensus/arbitration vs either-reader).

FACT: Posso, Carles, Rué, Puig, & Bonfill (2016, *PLOS ONE*, 11(7), e0159806, doi:10.1371/journal.pone.0159806) — programme analysis: double vs single detection similar; incremental cost per extra detection substantial; cost-effectiveness is **not automatic**.

Transfer limits: Mammography ≠ tutoring item QA (base rates, harms, labor markets differ). Borrow only: second readers catch misses; consensus vs either-policy matters; always-dual is not free excellence.

Applied (HYPOTHESIS): Dual promote ≈ consensus: agree on essential-for-C4 **and** stem-figure-key; discord → reject/repair, not either-admits. No hospital cosplay ads. Do not assume always-dual wins when seals already catch most IR failures.

Kill: “We double-read like hospitals ⇒ ACT-ready.” **Survive:** Consensus dual on triggered high-risk admits; cost in FILLETA capacity.

CXLIV.7 Trigger rule (doctrine core)

HYPOTHESIS — mandatory second rater when ANY of:

Trigger	Why
T1 Figured FormatId → C4	Pictorial ground truth; hard post-hoc student re-check (SAFE-FIGKEY)
T2 Parent-/Map-facing diagnosis ambition	Higher Kane use (hole fill that claims FormatId probed)
T3 Prior demote / critical REPAIR history on item or stem-family	Proven failure mode
T4 Ambiguous or non-machine-checkable figure IR	First rater is eye-test-adjacent; escalate
T5 First sealed admit for a concept×FormatId cell	Blueprint risk; later clones may inherit single if isomorphic hash trusts template (see id 145)
T6 First-rater self-flags uncertainty	Honest escalate > fake certainty

Optional (single OK) when ALL of: text/symbolic (no required figure), Practice-only or Reveal use (not C4), seals green, no prior demote, not first-of-cell blueprint, first rater confidence unmarked.

Arbitration: Disagree on essential-for-C4 or joint truth → **do not promote**; open REPAIR or reject; third SME only if ops policy budgets it — never average to admit.

Capacity: Dual minutes enter FILLETA estimator. Hitting a parent band **never** clears T1–T6.

CXLIV.8 Claim table

Claim	Label	Confidence	Notes
Higher-ambition uses need more evidence	FACT	High	Kane 2013; Standards 2014
Kappa adjusts chance agreement; bands are context-sensitive	FACT	High	Cohen 1960; McHugh 2012
Large assessments use second scoring + investigate-below-target	FACT	High	NAEP TDW; design-specific cutoffs
Double reading can raise detection; costs/policy-sensitive	FACT	High	Brown et al. 1996; Dinnes et al. 2001; Posso et al. 2016
Mammography ≡ MindCraft item QA	SPECULATION	Low	Refuse as identity
Risk-triggered dual beats always-dual and never-dual	HYPOTHESIS	Medium–High	Needs DUALRATE-*
Always-dual is cost-effective at startup FormatId scale	SPECULATION	Low	Refuse default
Parents accept longer fill bands when dual is real on diagrams	FOUNDER BELIEF	Medium	× FILLETA / HOLETRUST

CXLIV.9 Product surface — SAFE-DUALRATE claim contract

1. Promote ticket shows second_rater_required from the trigger rule before work starts.
2. Dual ≠ second glance by the same uid; log two reviewer ids + checklist hashes.
3. Seals remain hard gates; dual cannot override failed FIGKEY/GENQ.
4. Discord → reject/repair; no soft-pass; optional third SME only under written policy.
5. Sample duals for IRR ops health; investigate low agreement — no Dual Rater Score™ UI for students/parents.
6. FILLETA bands include dual capacity; never waive T1–T6 for ETA optics.
7. Copy: “Diagram diagnosis needs a second sealed check” — not “hospital-grade double reading” / not “every item is dual-certified.”

8. Telemetry: dual_required, dual_agree, dual_discord_reject, post-promote repair stratified by single vs dual — demote Dual Rater Score™ / Agreement Minutes NS.
9. First-of-cell dual may mint a template trust that later isomorphic clones use under tighter cool-down/hash rules (handoff to id 145) — never silent clone leak.
10. Ban Discord emoji / thumbs as second rater.

CXLIV.10 Doctrine — SAFE-DUALRATE (provisional)

1. **Dual intensity follows use ambition + figure risk** (Kane), not brand politeness.
2. **Written triggers** (T1–T6) — never always-dual costume, never never-dual soft-pass.
3. **Consensus dual** — agree to admit; discord rejects.
4. **Seals > raters** — humans cannot waive FIGKEY/GENQ.
5. **IRR is ops hygiene** — investigate; do not sell kappa.
6. **Capacity is real** — dual minutes in FILLETA; no ETA waiver.
7. **HITL first rater + escalate** — no Ivy⇒dual, no Discord⇒panel.
8. **No Dual Rater Score™ / Agreement Minutes / Kappa % NS.**
9. **No mammography-cosplay ACT / complete-visual guarantees.**
10. Copy: “Stricter review when the picture writes the Map — not double theater on every stem.”

Confidence: High on Kane/Standards/Cohen/McHugh/NAEP-as-process-pattern and on mammography facts *with transfer limits*. Medium on exact T1–T6 set and whether T1 alone (all figured C4) is too broad vs T1∧T2. High that always-dual and never-dual are both commercially toxic (FILLETA explosion vs Map distrust).

CXLIV.11 Experiments

ID	Question	Design	Primary
DUALRATE-1	Triggered dual (T1–T6) vs always-dual vs single-only	Ops A/B/C	Post-promote critical repair; dual minutes; fill-band honesty
DUALRATE-2	Consensus dual vs either-rater-admits	Ops A/B	False admits; discord rate; illicit Map writes
DUALRATE-3	T1 alone (all figured C4) vs T1∧T2 (figured ∧ parent-facing)	Ops A/B	Repair catch; backlog latency × FILLETA
DUALRATE-4	Sample IRR audit cadence (10% vs 25% duals on single-path items)	Ops A/B	Disagree catch; rater drift

DUALRATE-5	Parent CBC: “second sealed check on diagrams” honesty vs “hospital double-read” cosplay	CBC	Trust; WTP; backlash
DUALRATE-QUAL	Planted wrong-slope + green seals: does second rater catch when first misses?	Qual	Dual codebook

Falsifiers: Triggered dual equals always-dual on post-promote repairs *and* always costs ■ catch rate → narrow triggers (favor T1∧T2). Triggered dual equals single-only on repairs → strengthen training/checklist before adding bodies; keep dual for liability on parent-facing cells. Consensus never disagrees for 90 days → keep logging but thin sample rate. Pre-register DUALRATE-* before “double-verified diagnostics” campaigns.

CXLIV.12 So what for MindCraft commercially

- **Copy:** “When a figure can write your Map, a second sealed reviewer signs off.” Ban hospital cosplay.
- **Product:** second_rater_required; discord→reject; IRR sample queue; FILLETA includes dual minutes.
- **Growth:** Sell integrity density on visual diagnosis — not dual-everywhere latency or single-rater speed theater.
- **Positioning:** Against ChatGPT fluency *and* against edtech that rubber-stamps figures or burns two SMEs on every MC for costume.
- **Metric:** post-promote repair by single/dual; discord-reject rate; dual minutes/admit — demote Dual Rater Score™.
- **Kill list:** Always-dual; never-dual on figured C4; Dual Rater Score™; seal-waiving dual; mammography cosplay; ETA-skipped dual.
- **Vision:** Maya’s diagram Map cells are ones a named second reviewer would defend — Practice still grows with lights on; parents see slower-but-real fill bands.

References (verified)

- American Educational Research Association, American Psychological Association, & National Council on Measurement in Education. (2014). *Standards for educational and psychological testing*. AERA.
- Brown, J., Bryan, S., & Warren, R. (1996). Mammography screening: an incremental cost effectiveness analysis of double versus single reading of mammograms. *BMJ*, 312(7034), 809–812. <https://doi.org/10.1136/bmj.312.7034.809>
- Cohen, J. (1960). A coefficient of agreement for nominal scales. *Educational and Psychological Measurement*, 20(1), 37–46. <https://doi.org/10.1177/001316446002000104>
- Dinnes, J., Moss, S., Melia, J., Blanks, R., Song, F., & Kleijnen, J. (2001). Effectiveness and cost-effectiveness of double reading of mammograms in breast cancer screening: findings of a systematic review. *The Breast*, 10(6), 455–463. <https://doi.org/10.1054/brst.2001.0350>
- Kane, M. T. (2013). Validating the interpretations and uses of test scores. *Journal of Educational Measurement*, 50(1), 1–73. <https://doi.org/10.1111/jedm.12000>

- Lawshe, C. H. (1975). A quantitative approach to content validity. *Personnel Psychology*, 28(4), 563–575. <https://doi.org/10.1111/j.1744-6570.1975.tb01393.x>
- McHugh, M. L. (2012). Interrater reliability: the kappa statistic. *Biochemia Medica*, 22(3), 276–282. <https://doi.org/10.11613/BM.2012.031>
- National Center for Education Statistics. (n.d./updated 2024). Constructed-response interrater reliability. NAEP Technical Documentation. https://nces.ed.gov/nationsreportcard/tdw/analysis/initial_itemscore.aspx
- Posso, M., Carles, M., Rué, M., Puig, T., & Bonfill, X. (2016). Cost-effectiveness of double reading versus single reading of mammograms in a breast cancer screening programme. *PLOS ONE*, 11(7), e0159806. <https://doi.org/10.1371/journal.pone.0159806>

Part CXLV — Isomorphic Stem-Family Hashing for Cool-Down Without Over-Block

Chapter status: Living evidence + bank/selector brief — Researcher tick 2026-08-14 (UTC hour 18 ≡ Red Team slot, but ch145 never written → prefer Researcher per rotation; researcher count since synthesizer v1.18 = 2 → Researcher)

Primary question: How should MindCraft **define and hash a stem family** so cool-down after Practice reveal blocks **answer-hunt clones** (same figure/key path, swapped digits) without **over-blocking** true FormatId substitutes that keep gap-scan alive under thin C4 stock?

Owners: Engine (family hash / selector) · Bank / GENQ / FIGKEY · Product (Practice / C4) · COVER / SCANCOMP / COOLDOWN · Brand · Red Team

Commercial job: Ship **SAFE-FAMHASH** densifying SAFE-COOLDOWN × SAFE-GENQ × SAFE-FIGKEY × SAFE-SCANCOMP × SAFE-DUALRATE: **inspectable family membership** — never item_id-only leak, never concept-wide forever ban, never Family Hash Score™.

Builds on: Parts CXLI (SAFE-COOLDOWN — family > item_id foreshadow; COOLDOWN-3), CXLIV (SAFE-DUALRATE — T5 first-of-cell / template trust handoff), CXXXI (SAFE-FIGKEY), LXXII (SAFE-GENQ), CXXXVIII (SAFE-SCANCOMP), CXXXVII (SAFE-SPLITPOOL), LXXXVII (SAFE-COVER). Seams: family_id / template_id / figure_ir_hash / key_path_hash, enemy/clone review queue, cool-down clocks keyed by family not lone stem.

CXLV.1 Why this chapter exists

SAFE-COOLDOWN already bans immediate same-stem C4 and says **family > item_id** when clones share a figure/key path. The missing ops contract is **how family membership is decided**. Without it, three failure modes dominate:

1. **Under-hash (clone leak):** Cool-down bars only exact item_id. GENQ emits a near-clone (same slope figure, rescaled numbers) → student “fresh” C4 is tonight’s key in new clothes → Map poison (Appelhaus disclosed-reuse problem at clone grain).
2. **Over-hash (starve):** Cool-down bars the entire concept×FormatId cell, or every stem sharing a topic tag → FormatId practice dies; SCANCOMP pad theater returns; parents see holes that are policy accidents, not stock honesty.

3. **Black-box similarity theater:** “AI says they’re siblings” with no human seal → Family Hash Score™ ads, silent over/under block, no REPAIR path.

FOUNDER BELIEF under audit: Parents will accept a labeled hole or a *genuinely* different probe more than a green FormatId cell earned by recognizing last night’s diagram with new labels — but they will not accept a product that cools down “all graphs” after one Practice reveal.

Claims we refuse: item_id-only cool-down as sufficient under GENQ clones; concept-wide / FormatId-wide cool-down as default; embedding-only family without human seal; Family Hash Score™ / Clone Minutes NS; forever global family ban; pad when hash empties a cell; ACT / complete-visual guarantees from hashing packaging.

CXLV.2 Constructs

Construct	Meaning	MindCraft analogue	Failure mode
Item model / template	Authoring shell that generates variants	GENQ template / archetype seed	Untracked LLM remix
Isomorph (intended)	Unique instance, comparable cognitive demand	Sibling under same sealed template	Digit-swap clone sold as fresh
Clone / enemy (ops)	Pair that should not co-appear as independent evidence	Cool-down co-members; form enemies	Leak or over-block
Family hash	Stable id grouping co-compromised stems	family_id on cool-down clock	Opaque embedding cluster
Out-of-family isomorphic substitute	Same FormatId×concept ambition, different family	SCANCOMP preference under cool-down	Same-family “substitute”
SAFE-FAMHASH	Inspectable family membership doctrine	This chapter	Family Hash Score™

Operational definition (HYPOTHESIS): Two stems belong to the same cool-down family *F* iff **at least one** sealed membership rule fires: (R1) shared template_id from GENQ/AIG item model with declared radical/element slots; (R2) shared figure_ir_hash (structural IR, not pretty SVG) **and** shared key_path_hash (solution skeleton); (R3) human-tagged **enemy/clone** pair from bank QA (including NLP-flagged → SME confirm). Cool-down after reveal of any member bars **all** of *F* for that student. A **substitute** must be C4-sealed, same FormatId×concept cell ambition, and family_id ≠ *F*. If none exist → hole — never widen *F* retroactively to “everything diagram,” never pad. Embedding cosine may **propose** candidates for R3 review; it never silently writes family_id alone.

CXLV.3 Item models create families — cloning is a feature with a dependence cost

FACT: Gierl & Lai (2012, *International Journal of Testing*, 12(3), 273–298, doi:10.1080/15305058.2011.635830) — automatic item generation rests on **item models** (templates) whose elements are manipulated to produce many new

items; they discuss item-model banks and statistical procedures that estimate parameters of generated items without exhaustive separate pilot testing of every instance.

FACT: Sinharay, Johnson, & Williamson (2003, *Journal of Educational and Behavioral Statistics*, 28(4), 295–313, doi:10.3102/10769986028004295) — **item families** are groups of related items (e.g., AIG instances from models); calibration/scoring should account for dependence within families; they fit a Bayesian hierarchical model and introduce the **family expected response function (FERF)** to summarize P(correct) for an item randomly drawn from a family.

FACT: Glas & van der Linden (2003, *Applied Psychological Measurement*, 27(4), 247–261, doi:10.1177/0146621603027004001) — item cloning increases pool size but induces **parameter variability** across clones; multilevel IRT models families of clones; CAT selection can be two-stage: choose an optimal **family**, then randomly administer one clone — treating the family as the selection unit.

Applied (HYPOTHESIS): MindCraft already generates and promotes FormatId stems (SAFE-GENQ / SAFE-FIGKEY). That is AIG/cloning economics. Psychometrics treats the **family**, not the lone instance, as the dependence unit. Cool-down must do the same: compromise of one revealed instance compromises siblings that share the model/figure/key skeleton. Transfer limits: LSAT/CAT security ≠ teen gap-scan; borrow **family-as-unit**, not Sypson–Hetter brand theater or FERF Score™.

Kill: “New UUID ⇒ fresh diagnostic evidence.”

Survive: Template/family id as cool-down grain; random sibling after cool-down ends is practice, not instant C4.

CXLV.4 Enemy/clone detection is a review pipeline — not an unsupervised brand

FACT: Becker & Kao (2022, *Journal of Applied Testing Technology*, 23, 41–52) — NLP (e.g., cosine similarity on stem/key text) can flag **enemy and duplicate** pairs in large banks so SMEs review a manageable set; cosine thresholds are **starting values for human review**, not automatic ground truth; “not enemy” must distinguish unreviewed vs reviewed-clear.

FACT: Micir, Swygert, & D’Angelo (2022, *Journal of Applied Testing Technology*, 23, 30–40) — ML/NLP can prioritize enemy predictions in a large healthcare certification bank; editors then confirm; iterative SME feedback improves detection — **human confirmation remains in the loop**.

Applied (HYPOTHESIS): Family hashing for cool-down should reuse the **flag → SME seal** pattern. Cosine/embedding on stem+key (and, for figures, IR hash equality) proposes R3 enemies. Auto-writing family_id from cosine alone recreates clone leak (threshold too high) or over-block (too low). Ops queue: enemy_candidate tickets parallel PROMOTE/FIGKEY — not Discord emoji merges.

Kill: Embedding-only family membership as shipped law.

Survive: NLP propose + human seal; machine rules R1/R2 when hashes exist.

CXLV.5 Under-hash vs over-hash — both poison commercial honesty

Reuse (SAFE-COOLDOWN / SAFE-SCANCOMP): Disclosed reuse makes items easier (Appelhaus et al. 2023); thin C4 stock fails open to holes, not pad.

HYPOTHESIS — under-hash cost: Digit-swap / label-swap clones after reveal produce answer-hunt greens → parent distrust when tutor gold disagrees → WTP for “honest diagnosis” collapses (SAFE-HOLETRUST / SAFE-PROOF).

HYPOTHESIS — over-hash cost: Barring all coordinate_graph after one reveal empties FormatId cells → hole rate spikes → FILLETA pressure → soft-pass promote temptation (SAFE-FILLETA / SAFE-PROMOTE bans). Over-block is pad’s cousin: both prioritize *appearing* complete over measuring.

Doctrine split: Cool-down family = **compromise set** (same model/IR+key path/enemy seal). Substitute set = **ambition set** (FormatId×concept) **minus** compromise set. Never equate the two.

Kill: Concept×FormatId cool-down as default family.

Survive: Narrow sealed *F*; substitutes outside *F*; holes when substitutes absent.

CXLV.6 Template trust after dual admit — clones inherit membership, not C4 free passes

Reuse (SAFE-DUALRATE T5): First sealed admit for a concept×FormatId cell may require dual; later isomorphic clones may single-rater promote **if** they inherit a trusted template.

HYPOTHESIS: template_id minted at first-of-cell dual becomes R1 family membership for cool-down **and** promote inheritance — but inheritance never skips FIGKEY/GENQ seals, never auto-promotes to C4 from play counts, and never lets a just-revealed clone enter C4 during cool-down. Template trust ≠ “fresh forever.”

Kill: “Dual once ⇒ all clones are diagnostic-fresh anytime.”

Survive: Shared template_id ⇒ shared cool-down *F*; promote still ticketed; cool-down still clocks.

CXLV.7 Claim table

Claim	Label	Confidence	Notes
AIG uses item models; instances form related sets	FACT	High	Gierl & Lai 2012
Item families need dependence-aware calibration; FERF summarizes family	FACT	High	Sinharay et al. 2003
Cloning ⇒ family-level selection unit in CAT designs	FACT	High	Glas & van der Linden 2003

NLP can flag enemy/duplicate pairs for SME review	FACT	High	Becker & Kao 2022; Micir et al. 2022
Embedding cosine alone is sufficient family law	SPECULATION	Low	Refuse
R1∨R2∨R3 sealed membership beats item_id-only and concept-wide ban	HYPOTHESIS	Medium–High	Needs FAMHASH-*
Exact cosine threshold is universal science	SPECULATION	Low	Experiment-owned
Parents prefer sealed-family honesty over clone greens or all-graph bans	FOUNDER BELIEF	Medium	× HOLETRUST / COOLDOWN-5

CXLV.8 Product surface — SAFE-FAMHASH claim contract

1. Every C4/Practice stem carries family_id (nullable only for legacy until backfill); cool-down keys (student_id, family_id).
 2. Membership rules **R1∨R2∨R3** are written, versioned, and logged; no silent embedding write.
 3. R1: GENQ/template_id when present. R2: figure_ir_hash ∧ key_path_hash when figured. R3: SME-confirmed enemy/clone (NLP may propose).
 4. Reveal of any member sets cooldown_until(S, F) for **all** members of F.
 5. Substitutes: sealed C4, same FormatId×concept ambition, family_id ≠ F.
 6. Empty substitute set ⇒ format_not_probed / hole — never pad, never expand F to whole FormatId.
 7. Enemy candidate queue with human clear/confirm; distinguish unreviewed vs cleared.
 8. Telemetry: family_block_count, clone_leak_suspect, overblock_hole_rate, enemy_sme_confirm_rate — demote Family Hash Score™ / Clone Minutes.
 9. Copy: “We don’t re-diagnose the same problem family right after we show the key — a different family or an honest hole.” Ban “AI clone detector protects your ACT.”
 10. Dual/template trust (DUALRATE T5) writes R1 ids; cool-down still applies to the family.
-

CXLV.9 Doctrine — SAFE-FAMHASH (provisional)

1. **Family is the cool-down unit** — not lone item_id, not whole concept×FormatId.
2. **Membership is sealed** — R1 template / R2 IR+key path / R3 SME enemy — embeddings propose only.
3. **Under-hash and over-hash are both kills** — clone leak and all-graph bans both fail honesty.
4. **Substitutes live outside F** — ambition ≠ compromise.
5. **Holes beat pad** when substitutes missing (densifies SCANCOMP / COOLDOWN).

6. **Template trust $\neq$ cool-down waiver** — dual-minted templates still clock.
7. **Inspectable hashes** — versioned rules; REPAIR when leak/overblock found.
8. **No Family Hash Score™ / Clone Minutes / Similarity % NS.**
9. **No ACT / complete-visual / identity guarantees** from hashing packaging.
10. Copy: "Fresh family or honest hole — never a digit-swap dressed as diagnosis."

Confidence: High on Gierl–Lai / Sinharay / Glas–van der Linden / Becker–Kao / Micir as method sources (with CAT/certification transfer limits). Medium on exact R2 hash algorithms and cosine propose thresholds (FAMHASH-*). High that item_id-only cool-down under GENQ is commercially toxic.

CXLV.10 Experiments

ID	Question	Design	Primary
FAMHASH-1	item_id-only cool-down vs R1vR2vR3 family cool-down	Engine A/B	Clone-leak false mastery; Map $\leftrightarrow$ tutor gold
FAMHASH-2	Narrow family vs concept $\times$ FormatId cool-down	Engine A/B	Overblock hole rate; Format-route quality
FAMHASH-3	Embedding-auto family vs NLP-propose+SME seal	Ops A/B	False merges; missed clones; review minutes
FAMHASH-4	R2 ($IR \wedge \text{key}$) on vs off for figured stems	Engine A/B	Digit-swap leak; figure-variant overblock
FAMHASH-5	Parent CBC: “same problem family” honesty vs “new question id = fresh”	CBC	Trust; WTP; backlash
FAMHASH-QUAL	Planted digit-swap after reveal: does family bar catch it?	Qual	Codebook / tutor review

Falsifiers: Family hash equals item_id-only on planted clone leaks $\rightarrow$ strengthen R2/R3 before shipping “family cool-down” ads. Narrow family equals concept-wide ban on hole rate *and* gold $\rightarrow$ widen substitute bank, not the ban. SME seal never disagrees with embedding for 90 days $\rightarrow$ keep seal for liability; thin review sample, do not delete humans. Pre-register FAMHASH-* before “AI prevents answer recycling” campaigns.

CXLV.11 So what for MindCraft commercially

- **Copy:** “After we show a key, we won’t quiz the same problem family with the lights off.” Parent: “A lookalike with new numbers isn’t a fresh diagnosis.”
 - **Product:** family_id + R1/R2/R3; cool-down on family; substitutes outside *F*; enemy review queue.
 - **Growth:** Sell **anti-clone measurement honesty** vs apps that regenerate near-copies and call it adaptive diagnosis.
 - **Positioning:** Against ChatGPT one-blob fluency *and* against banks that treat UUID≠ as independence.
 - **Metric:** clone_leak_suspect; overblock_hole_rate; SME confirm precision — demote Family Hash Score™.
 - **Kill list:** item_id-only under GENQ; concept-wide cool-down default; embedding-only family; Family Hash Score™; pad when blocked; ACT guarantees from hashing.
 - **Vision:** Maya’s Map records competence on a **new family** — or an honest hole — never a memory of tonight’s diagram in a thin disguise.
-

References (verified)

- Becker, K. A., & Kao, S.-c. (2022). Identifying enemy item pairs using natural language processing. *Journal of Applied Testing Technology*, 23, 41–52. <https://www.jattjournal.net/index.php/atp/article/view/172634>
- Gierl, M. J., & Lai, H. (2012). The role of item models in automatic item generation. *International Journal of Testing*, 12(3), 273–298. <https://doi.org/10.1080/15305058.2011.635830>
- Glas, C. A. W., & van der Linden, W. J. (2003). Computerized adaptive testing with item cloning. *Applied Psychological Measurement*, 27(4), 247–261. <https://doi.org/10.1177/0146621603027004001>
- Micir, I., Swygert, K., & D’Angelo, J. (2022). Leveraging machine learning technology to improve accuracy and efficiency of identification of enemy item pairs. *Journal of Applied Testing Technology*, 23, 30–40. <https://jattjournal.net/index.php/atp/article/view/167170>
- Sinharay, S., Johnson, M. S., & Williamson, D. M. (2003). Calibrating item families and summarizing the results using family expected response functions. *Journal of Educational and Behavioral Statistics*, 28(4), 295–313. <https://doi.org/10.3102/10769986028004295>

Also relied on (already in Constitution spine; not re-litigated here): Appelhaus et al. (2023) disclosed reuse; Kane (2013) use ambition; SAFE-COOLDOWN / SAFE-SCANCOMP / SAFE-GENQ / SAFE-FIGKEY / SAFE-DUALRATE product contracts.

Part CXLVI — Gap-Scan Exit Emotion without Anxiety Score™ (Destake Copy after Holes)

Chapter status: Living evidence + student UX/copy brief — Researcher tick 2026-08-14 (UTC hour 21; hour%6≠0; researcher count since synthesizer v1.18 = 3 → Researcher)

Primary question: After gap-scan leaves honest FormatId holes (and/or soft-wrong evidence), what **exit-screen emotion language** should MindCraft ship — without Anxiety Score™, mood telemetry, consoling ability talk, ego grades, or calm-to-zero North Stars — so Maya leaves destaked enough to take one Practice step?

Owners: Product (gap-scan exit / summary) · Brand · STUDHOLE / HOLETRUST · PRIVACY / EXPOSE · Soft-wrong / ERRCLIMATE · Red Team

Commercial job: Ship **SAFE-EXITEMO** densifying SAFE-STUDHOLE × SAFE-PRIVACY × SAFE-EXPOSE × SAFE-SOFTMSG × SAFE-CALIB × SAFE-ERRCLIMATE: **destake the stakes of the summary, not the standards of the math** — never Anxiety Score™, never emotion-AI, never “not everyone is a visual person” comfort.

Builds on: Parts CXLII (SAFE-STUDHOLE), CXXXIX (SAFE-HOLETRUST), LXIX (SAFE-PRIVACY), LVIII (SAFE-EXPOSE), CXVIII (SAFE-SOFTMSG), L (SAFE-CALIB), CIII (SAFE-ERRCLIMATE), CXII (SAFE-FBLIT), LXXXI (SAFE-COLD). Seams: gap-scan end screen copy, optional self-authored check-in (short TTL), one Practice CTA, ban on mood dashboards / Calm Score™.

CXLVI.1 Why this chapter exists

SAFE-STUDHOLE locked **what** teens see after incomplete FormatId probes (parity holes, no fake greens, one CTA). The remaining failure mode is **how the exit feels**: products often “manage affect” by (a) scoring anxiety, (b) consoling fixed ability, (c) painting ego grades next to comments, (d) flooding with calm-app theater, or (e) hiding holes so the exit feels “complete.” SAFE-PRIVACY already bans face/voice Emotion AI and Anxiety Score™ as North Star. SAFE-EXPOSE bans flooding and Calm Score™. This chapter binds those bans to the **post-gap-scan exit surface** — the moment commercial conversion and identity risk collide.

FOUNDER BELIEF under audit: Parents and teens will pay for a Desk that says “this format wasn’t measured yet — try one sealed Practice when ready” more than for a Desk that diagnoses their child’s anxiety percentile or comforts them out of diagrams forever.

Claims we refuse:

1. Anxiety Score™ / Calm Score™ / Mood Minutes / Exit Emotion Score™ as North Star.
 2. Face/voice emotion AI or mandatory mood gate after gap-scan (SAFE-PRIVACY).
 3. Consoling ability talk: “It’s ok — not everyone is a visual / graph person.”
 4. Ego grades / Completeness % / “you got 3/5 formats” as the emotional closer beside hole honesty.
 5. Fake-complete greens or hide-holes as affect management.
 6. Day-one flooding timed ACT dump sold as “exposure therapy” exit.
 7. Empty growth-mindset / “mistakes are good!” poster without grain + CTA.
 8. ACT / identity / calm-to-zero guarantees from exit packaging.
-

CXLVI.2 Constructs

Construct	Meaning	MindCraft analogue	Failure mode
Exit destake	Lower evaluative threat of the summary without lowering math standards	Gap-scan end copy; no grade theater	Comfort≡soft standards
Affect without telemetry NS	Acknowledge friction in language; do not score emotion	Optional short-TTL check-in → pedagogy only	Anxiety Score™

Consoling comfort	Ability-excusing “kind” talk that demotivates	Banned exit lines	Entity comfort
Task close	Comment-like where-next without ego mark	Hole list + one Practice CTA	Grade+comment washout
SAFE-EXITEMO	Destake exit doctrine	This chapter	Calm Score™

Operational definition (HYPOTHESIS): Gap-scan exit copy is SAFE-EXITEMO-complete when it (1) states probed vs not_probed facts without self-level ability judgment, (2) uses **task/process** + **where-next** language (no person diagnosis of “anxious type”), (3) never displays Anxiety/Calm/Mood scores or emotion-AI inferences, (4) never consoling-comforts FormatId holes as fixed talent limits, (5) never pairs hole honesty with ego Completeness % / letter grades that steal attention from the CTA, (6) pairs encouragement (if any) with **exactly one** sealed Practice grain, (7) keeps optional affect check-in self-authored, short-TTL, refusal-safe, pedagogy-modifier only (SAFE-PRIVACY), and (8) is A/B-tested against comfort / score / grade / hide arms (EXITEMO-*).

CXLVI.3 Math anxiety is real — productizing it as a Score is the trap

FACT: Ashcraft (2002, *Current Directions in Psychological Science*, 11(5), 181–185, doi:10.1111/1467-8721.00196) — math anxiety is tension/fear related to math that predicts **avoidance** of elective math coursework and math-heavy careers; online timed arithmetic can show anxiety effects even when broad achievement tests look similar; anxiety can compromise **working memory** during math performance; causes are not fully determined; some teaching styles that publicly expose incompetence are plausible risk factors.

FACT: Ashcraft & Krause (2007, *Psychonomic Bulletin & Review*, 14(2), 243–248, doi:10.3758/BF03194059) — high math anxiety functions like a dual-task load: preoccupation with math fears consumes working-memory resources needed for non-retrieval arithmetic/math.

Applied (HYPOTHESIS): MindCraft must take anxiety **seriously as a design constraint** (destake stakes, avoid public humiliation theater, prefer graded approach — SAFE-EXPOSE) **without** turning the gap-scan exit into a clinical dashboard. Scoring “how anxious you are after your scan” invents a North Star that (a) conflates state affect with trait/construct measurement, (b) invites Goodhart gaming and stigma, and (c) sells empathy theater instead of FEI attempts. Measurement instruments exist for research (e.g., abbreviated scales); that does **not** license Anxiety Score™ in the student product chrome.

Kill: Anxiety Score™ / emotion-AI empathy brand on exit.

Survive: Destaked stakes + honest holes + one attempt path; research AMAS/AAI stay off the marketing splash.

CXLVI.4 Consoling comfort demotivates — confidence without ability excuses

FACT: Rattan, Good, & Dweck (2012, *Journal of Experimental Social Psychology*, 48(3), 731–737, doi:10.1016/j.jesp.2011.12.012) — instructors holding more of an **entity** (fixed) theory of math ability were quicker to judge low ability from early scores and more likely to offer **comfort-oriented** feedback (“It’s ok — not everyone can

be good at math”) and “kind” practices unlikely to promote engagement; students receiving comfort feedback inferred low expectations and reported **lower motivation and lower expectations** for future performance vs strategy-oriented feedback that communicated high standards.

Applied (HYPOTHESIS): Post-hole exit lines like “It’s fine if graphs aren’t your thing” or “Not everyone is a visual learner” are **entity comfort** dressed as kindness. They feel warm in the moment and commercially toxic afterward: they license FormatId avoidance, contradict SAFE-FORMAT / SAFE-DUAL (no learning-style meshing), and train parents to hear holes as destiny. Prefer Rattan-aligned pattern: acknowledge difficulty of the *task* (“Diagram items need a sealed figure before we hide keys”), express expectation of growth via practice, offer **strategy** (one sealed Practice), never excuse the person.

Kill: Consoling ability talk after holes or soft-wrongs.

Survive: High-standards + strategy CTA (densifies SAFE-SOFTMSG / SAFE-FBLIT / SAFE-STUDHOLE).

CXLVI.5 Ego marks steal the exit — comments alone beat grade+comment theater

FACT: Butler (1988, *British Journal of Educational Psychology*, 58(1), 1–14, doi:10.1111/j.2044-8279.1988.tb00874.x) — among fifth/sixth graders, **task-involving individual comments** produced higher interest and performance than **ego-involving numerical grades** or **grades plus comments**; grades+comments tended to promote ego-involvement similarly to grades alone (comments were often ignored when marks were present).

Applied (HYPOTHESIS): Gap-scan exits that show Completeness %, “3 of 5 formats ready,” letter-like belts, or rainbow mastery meters beside thoughtful hole copy recreate Butler’s grade+comment washout: teens attend to the ego mark (“I’m incomplete / I’m a C visual”) and discard the where-next. SAFE-STUDHOLE already bans Completeness % as NS; SAFE-EXITEMO forbids ego chrome as **affect management**. Transfer limit: Butler’s school tasks ≠ consumer diagnostic — borrow task-vs-ego, not “no numbers ever” dogma for FEI telemetry elsewhere.

Kill: Exit Completeness % / format letter grades as emotional closer.

Survive: Hole facts + one CTA; no ego meter competing with feed-forward.

CXLVI.6 Destake ≠ flood; destake ≠ soft standards

Reuse (SAFE-EXPOSE / SAFE-ERRCLIMATE / SAFE-COLD): Graded approach and expectancy violation under lowered stakes; destake mistakes without lowering standards; humble cold start without fireworks.

HYPOTHESIS: Exit destake means: (1) no timed ACT flood as the first post-scan move, (2) no classmate ranks on the exit, (3) no shame stalk of soft-wrongs to parent phones in the same beat, (4) language that separates **measurement limit** from **person worth**, (5) invitation to a **small sealed attempt** — not a calm-app meditation hero, not a forever-E0 protect track. Destake is situational stakes design, not therapy cosplay.

Kill: Flooding≡care; Calm Score™; forever protect without fade path.

Survive: Mild stakes on exit + standards held on Practice grain.

CXLVI.7 Privacy spine — optional check-in is not an Anxiety Score™

Reuse (SAFE-PRIVACY): Self-authored short-TTL affect → pedagogy modifiers only; ERT ban; contextual integrity; refusal-safe; FERPA≡ethics floor not doctrine.

HYPOTHESIS: If product offers a post-scan “how are you feeling about math right now?” check-in, it must be: optional, skippable without penalty, not written to durable Anxiety Score™ profiles, not sold, not shown on parent shame dashboards, and used only for soft pedagogy modifiers (e.g., slightly softer first Practice stake — never skip challenge forever). Inferring affect from face/voice/keystroke “empathy AI” remains banned.

Kill: Mandatory mood gate; emotion-AI exit; sell-affect.

Survive: Refusal-safe optional text → TTL modifiers; default is **no check-in required**.

CXLVI.8 Claim table

Claim	Label	Confidence	Notes
Math anxiety predicts avoidance and can load working memory	FACT	High	Ashcraft 2002; Ashcraft & Krause 2007
Entity-comfort feedback lowers student motivation/expectations	FACT	High	Rattan, Good, & Dweck 2012
Task comments beat grades and grades+comments for interest/performance in cited design	FACT	High	Butler 1988; transfer limits
Anxiety Score™ / emotion-AI improves FEI identity	SPECULATION	Low	Refuse (PRIVACY/EXPOSE)
Destake exit + strategy CTA > comfort / ego meter / hide	HYPOTHESIS	Medium–High	Needs EXITEMO-*
Optional short-TTL check-in can modify stakes without becoming NS	HYPOTHESIS	Medium	PRIVACY constraints
Parents prefer destaked honesty over anxiety percentiles	FOUNDER BELIEF	Medium	CBC / qual

Exact wording that maximizes CTA without inflation	SPECULATION	Low	Experiment-owned
--	-------------	-----	------------------

CXLVI.9 Product surface — SAFE-EXITEMO claim contract

1. Gap-scan exit shows STUDHOLE facts (Checked / Not checked yet) **without** Anxiety/Calm/Mood scores.
2. Copy spine: product-hole frame + difficulty acknowledgment + **strategy** expectation + one Practice CTA — never “not everyone can...” comfort.
3. No Completeness % / format letter / ego meter on the same primary surface as hole honesty.
4. No face/voice emotion AI; no mandatory mood gate.
5. Optional check-in (if shipped): self-authored, skippable, short TTL, pedagogy modifier only, not parent-stalked.
6. Soft-wrong streaks (if shown) use SOFTMSG grain + CTA — not Miss Streak shame or Calm Score™.
7. First post-exit Practice defaults to destaked sealed grain (EXPOSE ladder), not timed ACT flood.
8. Telemetry: exit_comfort_frame_flag, exit_ego_meter_flag, anxiety_score_shown_flag, cta→practice_start, checkin_skip_rate — demote Exit Emotion Score™.
9. Brand ban: “we measure your math anxiety,” “calm guaranteed,” therapy/CBT cosplay.
10. No ACT / identity / calm-to-zero guarantees from exit packaging.

CXLVI.10 Doctrine — SAFE-EXITEMO (provisional)

1. **Take anxiety seriously as constraint — never as product NS.**
2. **Destake the exit; hold the standards** (EXPOSE × ERRCLIMATE).
3. **No consoling ability talk** (Rattan / Good / Dweck).
4. **No ego marks beside hole honesty** (Butler task>ego).
5. **No Anxiety Score™ / Calm Score™ / emotion-AI** (PRIVACY).
6. **Hole honesty + one strategy CTA** (STUDHOLE densified).
7. **Optional check-in only** — short TTL, refusal-safe, no sell-affect.
8. **No fake greens / hide as affect management** (COLD / CALIB).
9. **No ACT / identity / calm-to-zero guarantees** from exit packaging.
10. Copy spine: “We checked what we can seal. Empty slots are not a verdict on you. When a sealed Practice item is ready, try one — graphs get clearer with attempts, not with excuses.”

Confidence: High on Ashcraft / Rattan–Good–Dweck / Butler spines with stated transfer limits. Medium on exact teen wording and optional-check-in uptake (EXITEMO-*). High that Anxiety Score™ and consoling comfort are commercially toxic for identity-transformation positioning.

CXLVI.11 Experiments

ID	Question	Design	Primary
EXITEMO-1	Destake+strategy CTA vs consoling comfort vs hide-holes	Copy A/B	cta→practice_start; help-avoidance; later FormatId approach
EXITEMO-2	Exit with vs without Completeness % / ego meter	UI A/B	Comment uptake proxy; false-strength; quit
EXITEMO-3	No check-in vs optional short-TTL vs mandatory mood gate	Flow A/B	Skip rate; Practice start; privacy complaints
EXITEMO-4	Anxiety Score™ splash (research-only arm) vs banned product chrome	Holdout / ethics board	Trust; stigma; WTP — expect kill of score arm
EXITEMO-5	Parent CBC: destaked honesty vs “we track anxiety” empathy brand	CBC	WTP; backlash; HOLETRUST interaction
EXITEMO-QUAL	Maya: which exit would you believe next week without feeling “diagnosed”?	Qual	Codebook destake vs comfort vs score

Falsifiers: Comfort raises short-run NPS *and* FormatId avoidance within 14d → still kill comfort default (identity liability). Ego meter raises completion of summary *and* equal CTA follow-through → still refuse for Butler/IUA; do not celebrate. Mandatory check-in equals optional on Practice starts *and* raises privacy tickets → kill mandatory. Pre-register EXITEMO-* before “anxiety-aware diagnostic” marketing.

CXLVI.12 So what for MindCraft commercially

- **Copy:** Destake + strategy, never “not everyone can...”; empty slots ≠ person verdict.
- **Product:** STUDHOLE exit without Anxiety/Calm scores; no ego Completeness chrome; optional refusal-safe check-in only.
- **Growth:** Sell **honest, non-clinical** post-scan calm — against empathy-AI tutors and against comfort apps that excuse avoidance.
- **Positioning:** Against Duo calm streaks, against therapy-cosplay edtech, against ChatGPT praise fog; with EXPOSE ladders and FEI gym.
- **Metric:** comfort_frame_flag; ego_meter_flag; anxiety_score_shown_flag; CTA→Practice — demote Exit Emotion Score™.
- **Kill list:** Anxiety Score™; emotion-AI; consoling comfort; ego exit meters; fake greens as affect fix; calm-to-zero; ACT guarantees from exit packaging.
- **Vision:** Maya leaves gap-scan a little less judged and a little more oriented to **one sealed attempt** — identity shifts by approach under destaked stakes, not by being scored as an anxious type.

References (verified)

- Ashcraft, M. H. (2002). Math anxiety: Personal, educational, and cognitive consequences. *Current Directions in Psychological Science*, 11(5), 181–185. <https://doi.org/10.1111/1467-8721.00196>
- Ashcraft, M. H., & Krause, J. A. (2007). Working memory, math performance, and math anxiety. *Psychonomic Bulletin & Review*, 14(2), 243–248. <https://doi.org/10.3758/BF03194059>
- Butler, R. (1988). Enhancing and undermining intrinsic motivation: The effects of task-involving and ego-involving evaluation on interest and performance. *British Journal of Educational Psychology*, 58(1), 1–14. <https://doi.org/10.1111/j.2044-8279.1988.tb00874.x>
- Rattan, A., Good, C., & Dweck, C. S. (2012). “It’s ok — Not everyone can be good at math”: Instructors with an entity theory comfort (and demotivate) students. *Journal of Experimental Social Psychology*, 48(3), 731–737. <https://doi.org/10.1016/j.jesp.2011.12.012>

Also relied on (already in Constitution spine; not re-litigated here): SAFE-STUDHOLE / SAFE-PRIVACY / SAFE-EXPOSE / SAFE-SOFTMSG / SAFE-CALIB / SAFE-ERRCLIMATE / SAFE-COLD product contracts; Hattie & Timperley (2007) feedback levels; Black & Wiliam formative use.

Part CXLVII — Dual-Rater Capacity Planning vs FILLETA Band Honesty

Chapter status: Living evidence + ops/capacity brief — Researcher tick 2026-08-15 (UTC hour 00 ≡ Red Team slot, but ch147 never written → prefer Researcher per rotation; researcher count since synthesizer v1.18 = 4 → Researcher)

Primary question: Once SAFE-DUALRATE requires second raters on triggered figured C4 admits, how must MindCraft **plan reviewer capacity** and **refresh FILLETA bands** — without Dual Rater Score™ vanity, always-dual costume that explodes ETAs, ETA-skipped dual under backlog heat, or “we’ll hire our way to complete visual by Friday” theater — so parent hole clocks stay capacity-true while Map integrity stays dual-true?

Owners: HITL / bank QA · Ops capacity · Product (FILLETA / ParentDashboard) · WORKFORCE · DUALRATE / PROMOTE / FIGKEY · Brand · Red Team

Commercial job: Ship **SAFE-CAPPLAN** densifying SAFE-DUALRATE × SAFE-FILLETA × SAFE-WORKFORCE × queue physics: **dual minutes are a first-class capacity line** in the fill estimator — never a costume that is waived when the parent clock looks bad, never a black-box Dual Rater Score™.

Builds on: Parts CXLIV (SAFE-DUALRATE), CXLIII (SAFE-FILLETA), CXL (SAFE-PROMOTE), LXXV (SAFE-WORKFORCE), CXXIII (SAFE-LOADSHED), CXXV (SAFE-KNOWNETA), CXXXIX (SAFE-HOLETRUST). Seams: dual-minute budget in FILLETA estimator, second-rater roster / headroom, labeled scarcity when dual capacity is the bottleneck, hire/contract plan without Dual Rater Score™ splash.

CXLVII.1 Why this chapter exists

SAFE-DUALRATE fixed *when* dual is mandatory (T1–T6). SAFE-FILLETA fixed *how* parents hear verify bands (capacity-grounded, no soon/play-count/soft-pass). The missing join:

> Dual review that is not staffed is either **lied about** (ETA-skipped dual; soft-pass first-rater) or **honestly exploded** into forever holes — unless ops plans **dual minutes as capacity**, refreshes bands when dual load rises, and hires/contracts before marketing “complete visual.”

Wrong resolutions already visible in edtech and in borrowed “double-check” industries:

1. **Always-dual costume without headcount** — Wing & Langelier-style workforce shock; FILLETA bands become fiction.
2. **ETA panic clears T1–T6** — DUALRATE already bans this; CAPPLAN makes the ban *operational* (estimator refuses to assume single-rater throughput on dual-required tickets).
3. **Dual Rater Score™ / Reviewer Minutes NS** — vanity that rewards short announced bands or high dual %.
4. **Hire-surge theater** — announce dual capacity that does not exist (WORKFORCE kill pattern).
5. **Ignore variability** — plan at 100% reviewer utilization; queues explode nonlinearly when GENQ/FIGKEY spikes arrive.
6. **Hospital cosplay capacity ads** — “double-read like radiology” sold as ACT readiness.

Without SAFE-CAPPLAN, DUALRATE is a policy PDF and FILLETA is a marketing clock. Parents who buy labeled FormatId holes deserve bands that **already price in** second-rater minutes.

FOUNDER BELIEF under audit: Families prefer a **longer honest verify band that includes dual minutes** over a short band that silently skips the second sealed check — and prefer labeled dual-capacity scarcity over Dual Rater Score™ dashboards.

Claims we refuse: dual-omitted FILLETA; ETA-skipped dual / soft-pass under date heat; always-dual without dual FTE; Dual Rater Score™ / Capacity Score™ / Reviewer Minutes NS; complete-visual-by-date or ACT guarantees from dual packaging; hospital double-read cosplay; 100% dual-util plans; Discord-emoji / same-uid “second rater” as capacity filler.

CXLVII.2 Constructs

Construct	Meaning	MindCraft analogue	Failure mode
Dual-required arrival rate	Tickets/week hitting T1–T6	Promote queue filter	Treat all tickets as single
Dual minutes / admit	Median second-rater + discord arbitration time	Ops handle time	Ignore discord tail
Dual capacity	Available independent second-rater minutes/week	Roster × protected QA hours	Count tutors mid-session as dual capacity
Capacity buffer	Planned idle dual capacity vs mean load	Headroom %	Run at 100% util
FILLETA dual line	Estimator term: $\text{backlog_dual} \times (1/\text{dual_throughput})$	Parent verify band input	Hide dual term
SAFE-CAPPLAN	Dual-capacity × band-honesty doctrine	This chapter	Dual Rater Score™

Operational definition (HYPOTHESIS): Dual+FILLETA is SAFE-CAPPLAN-complete when (1) the fill estimator **splits** single-path vs dual-required backlog, (2) dual-required latency uses **dual capacity** (not first-rater-only throughput), (3) planned utilization of dual capacity stays below a written headroom ceiling (not 100%), (4) when dual capacity is the binding constraint, parent/student copy **labels dual-verify scarcity** (not child-deficit shame), (5) raising dual % via always-dual costume without hiring **lengthens** published bands before marketing, (6) hitting a parent band never clears T1–T6, and (7) instruments are `dual_capacity_util`, `fill_eta_dual_component_ok`, `eta_skipped_dual_flag` — never Dual Rater Score™ / Capacity Score™ NS.

CXLVII.3 Little’s Law: backlog, throughput, and wait are one equation

FACT: Little, J. D. C. (1961, *Operations Research*, 9(3), 383–387, doi:10.1287/opre.9.3.383) — under stated stationarity conditions, the mean number in system $\langle L \rangle$ equals arrival rate $\langle \lambda \rangle$ times mean time in system $\langle W \rangle$: $\langle L \rangle = \langle \lambda \rangle \langle W \rangle$.

Applied (HYPOTHESIS): For dual-required promote tickets, if ops wants a parent-facing verify window $\langle W \rangle$ near a chosen band, then either dual throughput $\langle \lambda \rangle$ must rise (hire/protect second-rater minutes; narrow triggers only via DUALRATE evidence) or dual WIP $\langle L \rangle$ must fall (stop padding always-dual costume; reject/repair faster; stop GENQ floods into figured C4 ambition). Publishing a short $\langle W \rangle$ while $\langle \lambda \rangle$ is unchanged and $\langle L \rangle$ is high is **not optimism — it is denying Little’s Law**.

Kill: Fill bands that ignore dual backlog depth.

Survive: Capacity-grounded $\langle W \rangle \approx L / \langle \lambda \rangle$ bands with conservative rounding (FILLETA grain).

Wound: Exact Little conditions (strict stationarity) rarely hold in startup ops — use as **diagnostic identity**, not as a precision stopwatch for parents.

CXLVII.4 Variability + high utilization → nonlinear waits (capacity buffer required)

FACT / textbook law (transfer): Hopp, W. J., & Spearman, M. L. (2011, *Factory Physics*, 3rd ed., Waveland Press) — production systems with variability must buffer with **inventory, time, and/or capacity**; attempting to run near 100% utilization without reducing variability inflates cycle time (popularized via Kingman-style VUT intuition: queue time grows as utilization $\langle u / (1 - u) \rangle$ blows up).

Applied (HYPOTHESIS): Dual QA is a **variable** station: figure IR ambiguity, discord arbitration, GENQ repair storms, exam-week ticket spikes (LOADSHED). Planning second raters at full calendar occupancy guarantees FILLETA slip loops. CAPPLAN requires a **capacity buffer** (idle dual minutes against mean dual load) so parent bands stay finite under ordinary spikes — or else bands must widen *before* the spike, not after “almost ready” insults (Maister × KNOWNETA).

Kill: 100% dual-util plans; “we’ll catch up Friday” without headroom.

Survive: Written util ceiling + band refresh when util crosses threshold.

Wound: Factory lines ≠ tutoring QA — borrow buffer logic, not takt-time cosplay ads.

CXLVII.5 Universal double reading is a workforce multiplier — not a free integrity upgrade

FACT: Wing, P., & Langelier, M. (2009, *AJR*, 192(2), 370–378, doi:10.2214/ajr.08.1665) — workforce forecasts for breast imaging; implementing double reading for all mammograms would **significantly increase** radiologist demand; illustrative scenarios (assuming ~10% already double-read) span roughly **13%–90%** more mammographers depending on second-read time relative to first-read time (thousands of additional practitioners in national scenarios).

FACT (reuse): Posso et al. (2016, *PLOS ONE*) — double vs single cost-effectiveness is **not automatic**. Brown et al. (1996, *BMJ*) — consensus vs non-consensus policies change cost structure.

Transfer limits: Mammography ≠ FormatId promote QA (base rates, regulation, labor markets differ). Borrow only: **always-dual is a staffing decision**; second-read duration dominates capacity shock; policy without headcount is fantasy.

Applied (HYPOTHESIS): MindCraft must not ship “every figured item is dual-certified” brand language unless CAPPLAN shows dual λ can clear dual L at published W . Prefer DUALRATE triggers (risk-weighted dual) + CAPPLAN headroom over always-dual costume that either (a) lies on FILLETA or (b) starves Practice→C4 forever.

Kill: Always-dual marketing without dual FTE plan; hospital double-read ACT ads.

Survive: Triggered dual + explicit dual-minute budget; longer honest bands when dual load rises.

CXLVII.6 Wait psychology + Kane ambition still bind the clock

Reuse (SAFE-FILLETA / KNOWNETA): Maister (1985) — uncertain waits feel longer than known finite waits; “almost” after a miss insults. Whitt (1999) — delay information changes balking/renegeing; **accuracy** matters.

Reuse (SAFE-DUALRATE): Kane (2013) — higher-ambition uses need more backing. Parent-/Map-facing FormatId diagnosis fires dual triggers more often → **CAPPLAN must size dual capacity to the ambition sold**, or dial ambition down (honest holes; Practice-only until dual capacity exists).

Applied (HYPOTHESIS): When dual is the bottleneck, the finite band should **name the constraint class** without shame: “Diagram diagnosis needs a second sealed reviewer — about 3–5 weeks at current dual capacity.” Dual-omitted short bands produce Whitt-style later renegeing. COVER/SCANCOMP/HOLETRUST already ban pad theater; CAPPLAN bans **ambition–capacity mismatch theater**.

Kill: Soon-default; dual-omitted short bands; coverage ads implying dual-backed C4 without dual roster.

Survive: Finite dual-aware bands + Practice alternate; ambition ↔ dual FTE linkage.

Wound: Multi-week bank QA ≠ lobby minutes — keep days/weeks grain.

CXLVII.7 Claim table

Claim	Label	Confidence	Notes
$\backslash(L=\backslash\lambda W)$ links backlog, throughput, wait	FACT	High	Little 1961; stationarity caveats
Near-100% util + variability inflates waits; buffer with capacity/time/inventory	FACT (textbook)	High	Hopp & Spearman 2011; transfer limits
Universal double reading multiplies workforce need	FACT	High	Wing & Langelier 2009; domain transfer
Dual cost-effectiveness not automatic	FACT	High	Posso 2016; Brown 1996
Dual-aware FILLETA bands + headroom > dual-omitted short clocks	HYPOTHESIS	Medium–High	Needs CAPPLAN-*
Parents accept longer bands when dual constraint is named	FOUNDER BELIEF	Medium	CBC × HOLETRUST/WTP
Dual Rater Score™ improves learning or trust	SPECULATION	Low	Refuse
Exact national mammography FTE deltas apply to MindCraft	SPECULATION	Low	Mechanism only

CXLVII.8 Product surface — SAFE-CAPPLAN claim contract

1. FILLETA estimator exposes a **dual component**: dual-required backlog ÷ dual throughput (plus arbitration tail).
2. Dual throughput uses **independent second-rater roster minutes**, not first-rater leftover vibes or mid-session tutors.
3. Written **utilization ceiling** for dual capacity; breach → widen bands and/or hire — never clear T1–T6.
4. Copy may name **“second sealed check / dual capacity”** — ban child-deficit and hospital cosplay.
5. Always-dual proposals need a **CAPPLAN sheet** (dual FTE, util, projected $\backslash(W)$) before brand language.
6. `eta_skipped_dual_flag` is a **critical integrity incident** (same family as softpass-for-date).
7. Hire/contract dual capacity under WORKFORCE rules — no Dual Rater Score™; no Discord-emoji panel.
8. Telemetry: `dual_capacity_util`, `fill_eta_dual_component_ok`, `dual_backlog_age_p50` — demote Dual Rater Score™ NS.
9. Practice alternate remains mandatory while dual wait runs (FILLETA × STUDHOLE).
10. Template-trust inheritance (NEXT_LAB 148) may later cut dual load — never assume clone single-rater early.

CXLVII.9 Doctrine — SAFE-CAPPLAN (provisional)

1. **Dual minutes are capacity** — first-class FILLETA line, not optional polish.
2. **Little identity governs honesty** — short $\backslash(W\backslash)$ without $\backslash(\lambda\text{bda}\backslash)$ or lower $\backslash(L\backslash)$ is a lie.
3. **Headroom against variability** — refuse 100% dual util as plan (Hopp/Spearman buffer logic).
4. **Triggers stay sacred** — capacity pain widens bands or hires; never clears T1–T6 / seals.
5. **Always-dual is a staffing program** — Wing/Langelier lesson; not a free integrity sticker.
6. **Name the dual constraint** in wait copy when it binds (Maister $\times$ Whitt accuracy).
7. **Ambition $\leq$ dual roster** — Kane use evidence requires funded dual for parent-facing figured C4.
8. **No Dual Rater Score™ / Capacity Score™ / Reviewer Minutes NS.**
9. **No hospital-cosplay / ACT / complete-visual-by-date guarantees** from dual capacity packaging.
10. Copy spine: “Diagram diagnosis for this topic: not sealed yet. Needs a second reviewer — about 3–5 weeks at current dual capacity. Sealed Practice with the key shown can start sooner when available.”

Confidence: High on Little / Hopp–Spearman buffer logic / Wing–Langelier workforce multiplier *with transfer limits*, and on Posso/Brown cost non-automaticity. Medium on exact util ceilings, dual-minute medians, and whether naming “dual capacity” helps or confuses parents (CAPPLAN-*). High that dual-omitted ETAs and ETA-skipped dual are commercially toxic under Map trust.

CXLVII.10 Experiments

ID	Question	Design	Primary
CAPPLAN-1	Dual-aware FILLETA estimator vs single-rater-fantasy estimator under equal dual backlog	Ops A/B	fill_eta_honesty_ok; eta_skipped_dual_flag; slip count
CAPPLAN-2	Headroom util ceiling (e.g. 70% vs 95%) under injected ticket spike	Ops A/B	Band stability; dual backlog age; post-promote repair
CAPPLAN-3	Parent copy names “second sealed check capacity” vs vague “verifying” under equal actual $\backslash(W\backslash)$	Copy A/B	Trust; anger/uncertainty; WTP
CAPPLAN-4	Triggered dual + CAPPLAN vs always-dual costume at fixed dual FTE	Ops A/B	Fill bands; repair catch; throughput

CAPPLAN-5	Parent CBC: longer dual-honest band vs short band that skips dual	CBC	Trust; backlash; purchase intent
CAPPLAN-QUAL	8 ops + 8 parents: what dual-capacity sentence is clear vs hospital-cosplay?	Qual	CAPPLAN codebook

Falsifiers: Dual-aware bands worsen trust vs vague “verifying” *and* honesty_ok equal → keep dual in estimator silently; still refuse ETA-skipped dual. Headroom unused for 90 days with zero spikes → thin buffer carefully; do not celebrate 100% util. Always-dual at fixed FTE matches triggered dual on repairs *and* keeps bands acceptable → revisit DUALRATE breadth only with CAPPLAN sheet. Pre-register CAPPLAN-* before “double-verified complete visual” campaigns.

CXLVII.11 So what for MindCraft commercially

- **Copy:** “Second sealed check — about N–M weeks at current dual capacity.” Ban hospital cosplay and Dual Rater Score™.
 - **Product:** Split backlog estimator; dual util ceiling; eta_skipped_dual_flag; Practice alternate while waiting.
 - **Growth:** Sell **funded integrity** — dual minutes budgeted like tutor hours — against fake-fast visual diagnosis.
 - **Positioning:** Against ChatGPT fluency *and* against edtech that stamps “double-checked” without a second human or burns two SMEs on every MC for costume.
 - **Metric:** dual_capacity_util; fill_eta_dual_component_ok; dual backlog age; stratified repair — demote Dual Rater Score™.
 - **Kill list:** Dual-omitted FILLETA; ETA-skipped dual; always-dual without FTE; 100% util plans; Capacity Score™; complete-visual-by-date / ACT guarantees from dual packaging.
 - **Vision:** Maya’s diagram Map cells light when a **staffed** second reviewer could defend them — parents see slower dual-aware bands, not marketing clocks that quietly skipped the check.
-

References (verified)

- Brown, J., Bryan, S., & Warren, R. (1996). Mammography screening: an incremental cost effectiveness analysis of double versus single reading of mammograms. *BMJ*, 312(7034), 809–812. <https://doi.org/10.1136/bmj.312.7034.809>
- Hopp, W. J., & Spearman, M. L. (2011). *Factory Physics* (3rd ed.). Waveland Press.
- Kane, M. T. (2013). Validating the interpretations and uses of test scores. *Journal of Educational Measurement*, 50(1), 1–73. <https://doi.org/10.1111/jedm.12000>
- Little, J. D. C. (1961). A proof for the queuing formula: $(L = \lambda W)$. *Operations Research*, 9(3), 383–387. <https://doi.org/10.1287/opre.9.3.383>
- Maister, D. H. (1985). The psychology of waiting lines. In J. A. Czepiel, M. R. Solomon, & C. F. Surprenant (Eds.), *The service encounter* (pp. 113–123). Lexington Books.

- Posso, M., Carles, M., Rué, M., Puig, T., & Bonfill, X. (2016). Cost-effectiveness of double reading versus single reading of mammograms in a breast cancer screening programme. *PLOS ONE*, 11(7), e0159806. <https://doi.org/10.1371/journal.pone.0159806>
- Whitt, W. (1999). Improving service by informing customers about anticipated delays. *Management Science*, 45(2), 192–207. <https://doi.org/10.1287/mnsc.45.2.192>
- Wing, P., & Langelier, M. (2009). Workforce shortages in breast imaging: Impact on mammography utilization. *AJR. American Journal of Roentgenology*, 192(2), 370–378. <https://doi.org/10.2214/ajr.08.1665>

Part CXLVIII — Template-Trust Inheritance After First-of-Cell Dual Admit

Chapter status: Living evidence + ops/QA brief — Researcher tick 2026-08-15 (UTC hour 06 ≡ Red Team slot, but ch148 never written → prefer Researcher per rotation; researcher count since synthesizer v1.18 = 5 → Researcher)

Primary question: After SAFE-DUALRATE T5 dual-admits the **first sealed figured C4** for a concept×FormatId cell (blueprint), when may later **isomorphic clones** from the same sealed template promote under **single trained rater** — without forever-dual costume that burns CAPPLAN minutes, silent inherit that leaks Map poison, Template Trust Score™ vanity, or “family calibrated ⇒ every clone is diagnostic-ready” ads?

Owners: HITL / bank QA · Engine (template_id / family_id / trust bit) · GENQ / FIGKEY / FAMHASH · DUALRATE / CAPPLAN / PROMOTE · Brand · Red Team

Commercial job: Ship **SAFE-TMPTRUST** densifying SAFE-DUALRATE × SAFE-FAMHASH × SAFE-FIGKEY × SAFE-CAPPLAN × SAFE-PROMOTE: **conditional template trust** — blueprint dual once; inherit single-rater only under sealed sameness + sample audit + revoke-on-repair — never always-dual clones, never silent inherit.

Builds on: Parts CXLIV (SAFE-DUALRATE — T5 first-of-cell; foreshadow template trust handoff), CXLV (SAFE-FAMHASH — family as compromise unit), CXLVII (SAFE-CAPPLAN — dual minutes as capacity), CXL (SAFE-PROMOTE), CXXXI (SAFE-FIGKEY), LXXII (SAFE-GENQ), CXXXIV (SAFE-HCELIG). Seams: template_trust bit on template_id, inherit eligibility on promote tickets, sample-dual queue on inherited singles, revoke → re-arm T5.

CXLVIII.1 Why this chapter exists

DUALRATE T5 makes the **first** sealed figured C4 for a concept×FormatId cell dual-mandatory. CAPPLAN makes dual minutes scarce. FAMHASH puts isomorphic clones in one cool-down family. The missing join: after blueprint dual, may later sealed-template clones **inherit** single-rater intensity — or must every clone burn a second rater forever?

Failure modes: (1) forever-dual costume → FILLETA explosion; (2) silent inherit → broken IR Map poison; (3) Template Trust Score™ vanity; (4) “family calibrated” sold as seal-free C4; (5) embedding-only “same template”; (6) no revoke after a bad inherited admit.

FOUNDER BELIEF under audit: Parents accept faster fill *after* a named blueprint dual if clones still pass seals and a **sample** second check watches drift — more than forever-dual theater or silent speed.

Claims we refuse: forever-dual clones after blueprint dual; silent inherit; Template Trust Score™ / Inherit % NS; first dual ≡ forever trust without sample audit; embedding-only inheritance; FERF≡C4-ready ads; ACT / complete-visual from inheritance packaging.

CXLVIII.2 Constructs

Construct	Meaning	MindCraft analogue	Failure mode
Blueprint admit	First dual-sealed figured C4 for a concept×FormatId cell	T5 dual promote	Skip T5; soft-pass first
Template trust	Ops license: later clones may single-rater <i>if</i> gates hold	template_trust=true on template_id	Silent forever trust
Inherit candidate	Clone sharing sealed template_id + family membership	Promote ticket flag	Digit-swap without hash
Sample dual	Periodic second rater on inherited singles	IRR / acceptance sample	Kappa Score™ NS
Trust revoke	Clear trust bit; re-arm T5 / dual for family	Repair spike / discord	No revoke path
SAFE-TMPTRUST	Conditional inheritance doctrine	This chapter	Template Trust Score™

Operational definition (HYPOTHESIS): A clone promote is SAFE-TMPTRUST-complete when (1) a prior blueprint for the same concept×FormatId cell completed consensus dual with seals green (T5), (2) the candidate shares a **sealed** template_id (GENQ item model) and FAMHASH membership (R1 and/or R2; not embedding-alone), (3) FIGKEY/GENQ/FORMTRAP seals still pass **on this instance** (humans cannot inherit past a failed seal), (4) no open critical REPAIR / demote on the template or family, (5) template_trust is explicitly set (not inferred from play count), (6) first rater may admit alone **only if** (1)–(5) hold and T1–T4/T6 do not independently fire for other reasons (e.g. prior demote on *this* stem, self-flag), (7) a written **sample dual** fraction of inherited admits still gets a second rater for process health, and (8) post-promote repair / discord on inherited path **revokes** trust until a new dual blueprint or written re-seal.

CXLVIII.3 Item models justify *family-level* work — not seal-free clones

FACT: Gierl & Lai (2012, *International Journal of Testing*, 12(3), 273–298, doi:10.1080/15305058.2011.635830) — automatic item generation uses **item models** (templates) whose elements are manipulated to produce many items; they discuss item-model banks and statistical procedures that can estimate parameters of generated items **without exhaustive separate pilot testing of every instance**.

FACT: Gierl, Lai, & Turner (2012, *Medical Education*, 46(8), 757–765, doi:10.1111/j.1365-2923.2012.04289.x) — three-stage AIG (cognitive model → item model → generation); content specialists structure the model; computers

systematically combine elements — quality lives in the **model**, not in post-hoc fluency of each stem.

Applied (HYPOTHESIS): Figured FormatId generation is AIG economics. Blueprint dual is the model-level C4 investment. Inheritance is “don’t re-pilot every instance” — **not** “skip instance seals.” FIGKEY/GENQ still verify each clone.

Kill: “Template exists $\Rightarrow$ clone is C4-ready.” **Survive:** Inherit *reviewer intensity* only; seals stay per-instance.

CXLVIII.4 Family calibration saves resources — and still models clone variability

FACT: Sinharay, Johnson, & Williamson (2003, *Journal of Educational and Behavioral Statistics*, 28(4), 295–313, doi:10.3102/10769986028004295) — item families (e.g., AIG instances from models) require models that account for **within-family dependence**; they fit a Bayesian hierarchical model and introduce the **family expected response function (FERF)**; calibrating families (rather than every item separately) is framed as a path to **resource savings**.

FACT: Glas & van der Linden (2003, *Applied Psychological Measurement*, 27(4), 247–261, doi:10.1177/0146621603027004001) — item cloning induces **parameter variability** across clones; multilevel IRT treats clone parameters as draws from a family distribution; CAT can select a **family** then randomly administer one clone.

Applied (HYPOTHESIS): Family-as-unit supports not burning dual minutes on every sibling — and the same literature says clones vary. Inheritance = lower dual *intensity* after blueprint trust; instance seals + sample dual catch Glas/van der Linden variability. Not “parameters/seals copy-pasted.”

Kill: FERG ads as seal-free / always-dual-certified C4. **Survive:** Family unit for intensity; seals + sample for variability. **Wound:** LSAT/CAT $\neq$ gap-scan Map — borrow logic, not CAT marketing.

CXLVIII.5 Higher ambition still needs backing — inheritance is a weaker IUA branch

Reuse (SAFE-DUALRATE): Kane (2013, *Journal of Educational Measurement*, 50(1), 1–73, doi:10.1111/jedm.12000) — validate interpretations and uses; more ambitious claims need more backing; rejecting a use need not kill a weaker interpretation.

Applied (HYPOTHESIS): Map-facing FormatId diagnosis stays high ambition. Inheritance does not lower the use — it argues blueprint + seals + sample dual can back siblings. Selling “every diagram was dual-read” while inheriting is a Kane lie. Honest copy names blueprint dual + machine seals + sample audits.

Kill: Inherit as always-dual claim. **Survive:** Conditional intensity cut with transparent IUA.

CXLVIII.6 Sampling after demonstrated quality — not zero inspection

FACT (transfer, with limits): ISO 2859-1:1999 — attribute sampling indexed by AQL lets orgs inspect **samples** of lots rather than 100% when holding a process average near an acceptance quality limit; switching rules move among normal / tightened / reduced inspection as quality history changes (also reflected in ANSI/ASQ Z1.4 practice).

Transfer limits: Factory lots ≠ promote tickets. Borrow only: after demonstrated quality, **100% dual is not the only integrity tool** — sample + tighten-on-failure beats forever-100%-inspect.

Applied (HYPOTHESIS): Blueprint dual ≈ process qualification; inherited singles ≈ reduced inspection with revoke when sample dual or repair spikes. No ISO/AQL Score™ or factory-cosplay ACT ads.

Kill: Zero second-rater forever after one dual. **Survive:** Written sample % + revoke. **Wound:** Exact ISO tables ≠ product law — set rates in ops policy; pre-register TMPTRUST-.*.

CXLVIII.7 Inheritance gates (doctrine core)

HYPOTHESIS — `template_trust` **may be set when ALL of:**

Gate	Why
G1 Blueprint dual consensus admit for <code>concept×FormatId</code> cell	T5 satisfied once
G2 Sealed <code>template_id</code> on blueprint and candidate	Model identity (Gierl item model)
G3 FAMHASH membership (R1 and/or R2; R3 enemy seal if NLP-proposed)	Compromise-set honesty
G4 Instance seals green (GENQ + FIGKEY if figured + FORMTRAP if visual options)	Variability catch (Glas)
G5 No open critical REPAIR / demote on template or family	Trust not while broken
G6 Sample-dual policy active (written %)	Reduced inspection ≠ zero

Single-rater inherit blocked when ANY of: T3 prior demote on *this* stem; T4 ambiguous non-machine IR on *this* instance; T6 first-rater self-flag; `template_trust` false/revoked; seals fail; family hash missing; first-of-cell not yet dualled.

Revoke triggers (re-arm dual for template/family): post-promote critical repair on inherited path; sample-dual discord; FIGKEY/GENQ demote on any family member; ops IRR investigate-fail (DUALRATE pattern).

Capacity: Inheritance is a **CAPPLAN lever** — it may shorten dual backlog *after* G1–G6, never a reason to skip G1 or clear T1–T6 on the blueprint.

CXLVIII.8 Claim table

Claim	Label	Confidence	Notes
-------	-------	------------	-------

Item models enable many instances; model-level work can reduce per-instance pilot burden	FACT	High	Gierl & Lai 2012; Gierl et al. 2012 Med Educ
Family calibration / multilevel clone models acknowledge within-family dependence + parameter variability	FACT	High	Sinharay et al. 2003; Glas & van der Linden 2003
Higher-ambition uses need more evidence	FACT	High	Kane 2013
Attribute sampling / reduced inspection after quality history is a standard industrial pattern	FACT	High	ISO 2859-1:1999; transfer limits
Conditional inherit (dual blueprint → single + sample) beats forever-dual and silent inherit	HYPOTHESIS	Medium–High	Needs TMPTRUST-*
Parents accept faster fill when blueprint dual + seals are named	FOUNDER BELIEF	Medium	CBC × HOLETRUST/ FILLETA
Exact ISO AQL tables are MindCraft product law	SPECULATION	Low	Refuse
Family FERF ≡ seal-free C4 readiness	SPECULATION	Low	Refuse

CXLVIII.9 Product surface — SAFE-TMPTRUST claim contract

1. Promote UI shows `template_trust` and `inherit_eligible` explicitly — never inferred from play count.
2. Blueprint path still runs full T5 dual; inheritance never skips first-of-cell.
3. Instance seals remain hard gates; inheritance cannot override failed FIGKEY/GENQ.
4. Sample dual queue draws a written % of inherited singles; discord → reject/repair + **revoke**.
5. Telemetry: `template_trust_set`, `inherit_admit`, `inherit_sample_dual`, `template_trust_revoke` — demote Template Trust Score™ / Inherit % NS.
6. CAPPLAN estimator may use lower dual arrival rate **only after** trust bits exist — never assume inherit early.
7. Copy: “First diagram check for this topic was dual-sealed; later variants share that template and still pass machine seals.” Ban “every item double-read” and FERF/hospital cosplay.

8. Cool-down still bars the whole FAMHASH family after reveal (SAFE-FAMHASH) — trust inheritance $\neq$ cool-down waiver.
9. Ban embedding-only template match for inherit.
10. Ban Discord-emoji / same-uid “second sample.”

CXLVIII.10 Doctrine — SAFE-TMPTRUST (provisional)

1. **Blueprint dual first** — T5 is the price of template trust.
2. **Inherit review intensity, never seals** — each clone still passes FIGKEY/GENQ/FORMTRAP.
3. **Sealed sameness only** — `template_id` + FAMHASH; no embedding-alone.
4. **Sample dual + revoke** — reduced inspection with tighten-on-failure (ISO pattern, transfer-limited).
5. **Psychometrics supports family-level efficiency and warns of clone variability** — use both halves.
6. **Kane IUA honesty** — do not sell inherit as always-dual.
7. **CAPPLAN lever** — inherit cuts dual load only after gates; never ETA-skipped blueprint dual.
8. **No Template Trust Score™ / Inherit % / Family Calibration Score™ NS.**
9. **No ACT / complete-visual / FERF-as-readiness guarantees** from inheritance packaging.
10. Copy spine: “We dual-seal the first diagram probe for a topic. Matching variants can move faster under the same sealed template — still machine-checked, still sample-audited.”

Confidence: High on Gierl/Sinharay/Glas/Kane facts *with transfer limits*, and on ISO sampling as analogy not law. Medium on exact sample %, revoke thresholds, and whether T1 (all figured) should still force dual even under trust (favor: inherit only clears T5-driven dual, not T3/T4/T6; $T1 \wedge \text{trust} \rightarrow \text{single OK}$ if G1–G6 — test in TMPTRUST-3). High that forever-dual clones and silent inherit are both commercially toxic.

CXLVIII.11 Experiments

ID	Question	Design	Primary
TMPTRUST-1	Inherit (single+sample) vs forever-dual on sealed clones after blueprint	Ops A/B	Post-promote repair; dual minutes; fill-band honesty
TMPTRUST-2	Inherit vs silent-inherit (no sample dual)	Ops A/B	Repair catch; revoke rate; illicit Map writes
TMPTRUST-3	Inherit clears T5-only vs also clears T1 for trusted figured clones	Ops A/B	Repair; dual backlog; CAPPLAN util
TMPTRUST-4	Sample dual rate (5% vs 20%) on inherited admits	Ops A/B	Discord catch; dual minutes; false sense of safety

TMPTRUST-5	Parent CBC: “blueprint dual + sealed variants” honesty vs “every diagram double-read”	CBC	Trust; WTP; backlash
TMPTRUST-QUAL	Planted broken IR on trusted-template clone: does first rater + sample dual catch?	Qual	TMPTRUST codebook

Falsifiers: Inherit equals forever-dual on repairs *and* never saves meaningful dual minutes → keep dual on figured C4; treat TMPTRUST as CAPPLAN fantasy. Inherit equals silent-inherit on repairs → raise sample % or kill inherit. Sample dual never discords for 90 days with zero repairs → thin sample carefully; keep revoke path. Pre-register TMPTRUST-* before “smart double-check only once” campaigns.

CXLVIII.12 So what for MindCraft commercially

- **Copy:** “First diagram probe for this topic: second sealed check. Later matching variants: same template, machine seals, sample audits.” Ban always-dual costume and silent speed.
 - **Product:** template_trust bit; inherit_eligible on tickets; sample-dual queue; revoke on repair.
 - **Growth:** Sell **integrity density per dual minute** — blueprint investment compounds across a FormatId cell — against ChatGPT fluency *and* against edtech that either double-reads forever or rubber-stamps clones.
 - **Positioning:** Against “AI wrote N diagrams” heroes *and* against hospital-cosplay dual-everywhere latency.
 - **Metric:** inherit admit rate with repair stratification; revoke rate; dual minutes saved vs repair delta — demote Template Trust Score™.
 - **Kill list:** Forever-dual clones; silent inherit; Template Trust Score™; embedding-only trust; FERF≡C4-ready ads; ACT/complete-visual from inheritance.
 - **Vision:** Maya’s Map lights a diagram cell when a **blueprint** second reviewer would defend the template — and later variants do not silently invent new pictorial ground truth.
-

References (verified)

- Gierl, M. J., & Lai, H. (2012). The role of item models in automatic item generation. *International Journal of Testing*, 12(3), 273–298. <https://doi.org/10.1080/15305058.2011.635830>
- Gierl, M. J., Lai, H., & Turner, S. R. (2012). Using automatic item generation to create multiple-choice test items. *Medical Education*, 46(8), 757–765. <https://doi.org/10.1111/j.1365-2923.2012.04289.x>
- Glas, C. A. W., & van der Linden, W. J. (2003). Computerized adaptive testing with item cloning. *Applied Psychological Measurement*, 27(4), 247–261. <https://doi.org/10.1177/0146621603027004001>
- International Organization for Standardization. (1999). *ISO 2859-1:1999 Sampling procedures for inspection by attributes — Part 1: Sampling schemes indexed by acceptance quality limit (AQL) for lot-by-lot inspection*. ISO. (Transfer analogy only; not product cut-score law.)
- Kane, M. T. (2013). Validating the interpretations and uses of test scores. *Journal of Educational Measurement*, 50(1), 1–73. <https://doi.org/10.1111/jedm.12000>

- Sinharay, S., Johnson, M. S., & Williamson, D. M. (2003). Calibrating item families and summarizing the results using family expected response functions. *Journal of Educational and Behavioral Statistics*, 28(4), 295–313. <https://doi.org/10.3102/10769986028004295>
-

Part CXLIX — Post-Promote Repair Stratified Dashboards (Single vs Dual)

Chapter status: Living evidence + ops/QA brief — Researcher tick 2026-08-15 (UTC hour 09 → Researcher per rotation; researcher count since synthesizer v1.18 = 6 → Researcher)

Primary question: After C4 admits via **single-rater**, **dual-rater**, or **TMPTRUST inherit**, how must MindCraft surface **post-promote repair** — without a pooled Kappa % North Star, Dual Rater Score™ vanity, “our IRR is science-backed” ads, or a single repair rate that hides which path is poisoning Map — so ops can revoke trust, retune DUALRATE triggers, and defend parent hole clocks with path-true integrity?

Owners: HITL / bank QA · Engine (promote path tags + repair telemetry) · SAFE-PROMOTE / SAFE-DUALRATE / SAFE-TMPTRUST / SAFE-INSTRUMENT · Brand · Red Team

Commercial job: Ship **SAFE-REPAIRSTRAT** densifying SAFE-DUALRATE × SAFE-PROMOTE × SAFE-INSTRUMENT (+ TMPTRUST / CAPPLAN): **stratified post-promote repair + agreement monitors** — path-split rates, contingency tables, drift alerts — never Kappa % as the learning or trust score.

Builds on: Parts CXL (SAFE-PROMOTE — post-promote repair rate foreshadowed), CXLIV (SAFE-DUALRATE), CXLVIII (SAFE-TMPTRUST — revoke on repair / sample discord), CXLVII (SAFE-CAPPLAN), XC (SAFE-INSTRUMENT), LX (SAFE-REPAIR), LXXII (SAFE-GENQ). Seams: `promote_path` ∈ {single, dual, inherit_sample, inherit_solo}; stratified repair dashboards; sample-dual contingency exports; no Kappa Score™ splash.

CXLIX.1 Why this chapter exists

PROMOTE made post-promote repair a first-class metric. DUALRATE split admits into single vs dual. TMPTRUST added inherit + sample dual + revoke. INSTRUMENT forbade vanity KPIs as North Stars. The missing join: **one pooled repair % or one Kappa %** that looks green while inherit-solo silently demotes Map, or dual looks “worse” only because hard cases are routed there.

Failure modes: (1) pooled post-promote repair hides path poison; (2) Kappa % as ops NS / marketing hero; (3) Landis-Koch band costume as “almost perfect” brand; (4) dual path blamed for higher repair when selection (harder tickets) confounds; (5) no drift monitor → raters soften mid-exam-week; (6) Discord-emoji “agreement” theater.

FOUNDER BELIEF under audit: Parents and LEAs accept slower visual fill when ops can name **which admit path** failed and what tightened — more than a glossy IRR badge with no demote story.

Claims we refuse: Kappa % / Dual Rater Score™ / Agreement Minutes NS; pooled-only repair dashboards; “science-backed double-check” ads from a single κ; mammography-cosplay IRR heroes; ACT / complete-visual guarantees from dashboard packaging.

CXLIX.2 Constructs

Construct	Meaning	MindCraft analogue	Failure mode
Promote path	How the item entered C4	single / dual / inherit_solo / inherit_sample	Untagged admits
Post-promote repair	Critical REPAIR or demote after C4 admit	Ticket + hcelig_pass clear	Leave broken in scan
Stratified repair rate	Repair incidence by path (and FormatId / cell)	Ops dashboard rows	Pooled-only vanity
Pair agreement table	Contingency of rater A vs B on dual/sample	Export for investigate	κ -only without table
Rater drift signal	Standards shift over time within rater	Trend / windowed alerts	One-shot κ forever
SAFE-REPAIRSTRAT	Stratified repair + agreement doctrine	This chapter	Kappa Score™ NS

Operational definition (HYPOTHESIS): Post-promote QA monitoring is SAFE-REPAIRSTRAT-complete when (1) every C4 admit carries an immutable promote_path tag, (2) critical repair / demote events join back to that path within SLA, (3) ops dashboards show **path-stratified** repair rates (and cell/FormatId slices) before any pooled rollup, (4) dual and inherit-sample pairs store **raw contingency counts** (admit/reject/repair codes) — κ may be computed as a *diagnostic aide*, never as the sole displayed health, (5) windowed monitors flag path-specific repair spikes and sample-dual discord spikes (tighten TMPTRUST / DUALRATE / training), (6) parent-/LEA-facing copy never cites a Kappa % or “almost perfect agreement” band as product proof, and (7) instrumented events feed INSTRUMENT-style FEI-adjacent ops health — not DAU cosplay.

CXLIX.3 Chance-corrected agreement is a tool — not a North Star

FACT: Cohen (1960, *Educational and Psychological Measurement*, 20(1), 37–46, doi:10.1177/001316446002000104) — κ corrects observed nominal agreement for chance agreement expected from marginal base rates; it is a **coefficient of agreement**, not a measure of educational impact or product quality by itself.

FACT: Landis & Koch (1977, *Biometrics*, 33(1), 159–174, doi:10.2307/2529310) — develop generalized kappa-type statistics for observer agreement on categorical data and popularize interpretive bands (slight → almost perfect); their framework also supports examining agreement structure beyond a single omnibus number (including stratified/observer-reliability analyses in the broader Landis–Koch program).

Applied (HYPOTHESIS): Dual/sample-dual QA may compute κ on promote decisions (admit vs reject vs send-to-repair) as a **process diagnostic**. Shipping “ $\kappa = 0.81 \Rightarrow$ almost perfect visual diagnostics” as marketing is a category error: κ says raters often match each other, not that Map diagnoses are true for Maya or that parents should trust a hole-fill date.

Kill: Kappa % NS; Landis–Koch band as brand hero. **Survive:** κ + contingency table as investigate tools behind ops login.

CXLIX.4 High agreement can look “bad” on κ — and pooling hides that

FACT: Feinstein & Cicchetti (1990, *Journal of Clinical Epidemiology*, 43(6), 543–549, doi:10.1016/0895-4356(90)90158-l) — document **paradoxes** where observed agreement $\backslash(p_0\backslash)$ is high yet κ is low when margins are imbalanced; a second paradox concerns asymmetrical vs symmetrical imbalance. They warn that omnibus κ alone can mislead.

FACT: Cicchetti & Feinstein (1990, *Journal of Clinical Epidemiology*, 43(6), 551–558, doi:10.1016/0895-4356(90)90159-m) — argue for reporting separate positive/negative agreement indexes alongside κ so operators can see *where* concordance lives.

Applied (HYPOTHESIS): Promote decisions are often **prevalence-skewed** (most sealed tickets admit). A lonely κ can tank while raw agreement stays high — or look fine while one path (*inherit_solo*) carries asymmetric repair. SAFE-REPAIRSTRAT requires **path strata + contingency (or $\backslash(p_{\text{pos}}/p_{\text{neg}})\backslash$ -style splits)** before celebrating or panicking on κ .

Kill: Single κ as promote health. **Survive:** Stratified repair + agreement components. **Wound:** Clinical binary markers $\neq$ figure-IR promote codes — borrow paradox warning, not medical cut-scores.

CXLIX.5 Agreement is not stationary — monitor drift, not a launch-week κ

FACT: Myford & Wolfe (2009, *Journal of Educational Measurement*, 46(4), 371–389, doi:10.1111/j.1745-3984.2009.00088.x) — provide an operational framework for **monitoring rater performance over time**, detecting differential accuracy and differential scale-category use (DRIFT) as scoring progresses; some raters change accuracy or category use mid-campaign.

FACT: Sgammato & Donoghue (2017, *Applied Psychological Measurement*, 42(4), 307–320, doi:10.1177/0146621617730390) — evaluate Stuart’s Q for detecting **rater drift** in trend-scoring designs; show drift detection is a first-class monitoring problem, not a one-time IRR ceremony.

Applied (HYPOTHESIS): Exam-week LOADSHED and CAPPLAN pressure change rater behavior. SAFE-REPAIRSTRAT uses **rolling windows** of path-stratified repair and sample-dual discord (and optional trend-rescore of a sealed gold set) — not a frozen “ κ at launch” plaque. When *inherit_solo* repair spikes, revoke TMPTRUST; when dual discord spikes on T4 IR, retrain or narrow triggers — do not “average it away” in a company κ .

Kill: One-shot IRR certificate. **Survive:** Windowed path monitors + revoke/tighten hooks. **Wound:** AP English Rasch DRIFT $\neq$ MindCraft promote codes — borrow *monitoring over time*, not many-facet Rasch as required stack.

CXLIX.6 Higher-ambition use still needs stratified backing

Reuse (Kane 2013): More ambitious interpretations/uses need more backing; rejecting a use need not kill a weaker interpretation (*Journal of Educational Measurement*, 50(1), 1–73, doi:10.1111/jedm.12000).

Applied (HYPOTHESIS): Parent-facing “we dual-check diagram probes” is higher ambition than internal ops hygiene. A dashboard that shows only pooled repair while selling dual-check language is a Kane lie if dual volume is tiny and inherit_solo carries the repairs. Honest IUA: show **path mix** + stratified repair to leadership; external copy names seals + demote-on-repair without κ theater.

Kill: IRR badge as substitute for demote evidence. **Survive:** Stratified ops truth as backing for claimed uses.

CXLIX.7 Selection confounding — dual is not “worse” by default

HYPOTHESIS: Dual and inherit_sample paths **select harder tickets** (T4 ambiguity, first-of-cell, self-flag). Higher raw repair or discord on dual can be **good triage**, not dual failure. Dashboards must show:

- 1. Repair rate by path (numerator = critical repairs attributed to admits from that path; denominator = admits on that path in window).
- 2. Optional **risk-adjusted** view (e.g., within T-trigger strata) so dual is compared to single on comparable tickets — not a naive dual-vs-single shame board.
- 3. Inherit_solo vs inherit_sample split — sample discord that does not revoke is SAFE-TMPTRUST theater.

Kill: “Dual has higher repair ⇒ kill dual.” **Survive:** Path + trigger-conditioned tables. **SPECULATION (low):** Exact risk model form — pre-register REPAIRSTRAT-*.

CXLIX.8 Claim table

Claim	Label	Confidence	Notes
κ corrects chance agreement for nominal categories	FACT	High	Cohen 1960
Observer-agreement methodology + interpretive bands exist; omnibus κ is incomplete alone	FACT	High	Landis & Koch 1977; Feinstein/Cicchetti 1990
High $\psi(p_0)$ can coexist with low κ under margin imbalance	FACT	High	Feinstein & Cicchetti 1990

Rater standards can drift over a scoring campaign; monitoring over time is operationally warranted	FACT	High	Myford & Wolfe 2009; Sgammato & Donoghue 2017
Higher-ambition uses need more backing	FACT	High	Kane 2013
Path-stratified repair + contingency beats pooled κ as ops health	HYPOTHESIS	Medium–High	Needs REPAIRSTRAT-*
Parents prefer path-honest integrity stories over IRR badges	FOUNDER BELIEF	Medium	CBC × HOLETRUST
Landis–Koch bands are MindCraft product cut-scores	SPECULATION	Low	Refuse
Company κ predicts ACT gains	SPECULATION	Low	Refuse

CXLIX.9 Product surface — SAFE-REPAIRSTRAT claim contract

1. Tag every C4 admit with `promote_path` $\in$ {single,dual,inherit_solo,inherit_sample} — immutable after write.
2. Ops dashboard **defaults to stratified rows** (path × FormatId or path × concept cell); pooled rollup is secondary and labeled “descriptive only.”
3. Dual / inherit_sample stores contingency counts (and optional $\backslash(p_{\{pos\}}/p_{\{neg\}})$ -style splits); κ is an optional drill-down, never the hero tile.
4. Alerts: rolling-window repair spike by path; sample-dual discord spike $\rightarrow$ TMPTRUST revoke queue; dual discord spike $\rightarrow$ DUALRATE/training review.
5. Ban Kappa % / Dual Rater Score™ / Agreement Minutes / “Almost Perfect” Landis band on student, parent, or marketing surfaces.
6. Leadership weekly: path mix % + stratified repair + demote latency (PROMOTE) — not a single IRR badge.
7. CAPPLAN may read dual discord / repair as **capacity quality** signals (rework), never as Dual Rater Score™.
8. INSTRUMENT: events `c4_admit_path`, `post_promote_repair`, `sample_dual_discord`, `template_trust_revoke` — demote vanity scores.
9. Ban Discord-emoji / same-uid agreement as dual evidence (reuse DUALRATE/TMPTRUST).
10. Copy spine for ops/GTM: “We track which review path admitted a diagram probe — and we yank it when repair proves it wrong.” Ban “ κ -certified visual diagnostic.”

CXLIX.10 Doctrine — SAFE-REPAIRSTRAT (provisional)

1. **Stratify first** — path-tagged repair before pooled health.
2. **κ is a drill-down, not a North Star** — Cohen/Landis tools stay behind the ops wall.
3. **Show the table** — contingency / positive-negative agreement components beat omnibus vanity (Feinstein/Cicchetti warning).
4. **Monitor drift** — windowed path rates + optional trend gold; launch-week κ is not a certificate (Myford/Wolfe; Sgammato/Donoghue).
5. **Respect selection** — dual may look “worse” because it takes harder tickets; condition before blaming.
6. **Wire revoke/tighten** — inherit repair → TMPTRUST revoke; dual discord storms → DUALRATE/CAPPLAN review.
7. **Kane honesty** — do not sell dual-check IUA from unstratified dashboards.
8. **No Kappa % / Dual Rater Score™ / Agreement Minutes NS.**
9. **No ACT / complete-visual / “almost perfect agreement” guarantees** from IRR packaging.
10. Copy spine: “Integrity is path-true demote evidence — not a single agreement percent.”

Confidence: High on Cohen/Landis/Feinstein–Cicchetti/Myford–Wolfe/Sgammato–Donoghue/Kane facts *with transfer limits*. Medium on exact window lengths, alert thresholds, and risk-adjustment formula. High that pooled κ -as-NS and unstratified repair are commercially toxic under TMPTRUST inheritance.

CXLIX.11 Experiments

ID	Question	Design	Primary
REPAIRSTRAT-1	Stratified path dashboard vs pooled-only repair % for ops decisions	Ops A/B	Time-to-revoke; illicit C4 draws after known poison
REPAIRSTRAT-2	κ -hero tile vs contingency-first drill-down	Ops UX A/B	False panic/complacency; investigate quality
REPAIRSTRAT-3	Windowed alerts (7d vs 28d) on inherit_solo repair spikes	Ops A/B	Revoke latency; false alarms; Map gold
REPAIRSTRAT-4	Risk-conditioned dual-vs-single repair table vs naive shame board	Ops A/B	Mis-blame rate; trigger retune quality
REPAIRSTRAT-5	Parent/LEA CBC: path-honest demote story vs “ κ -certified dual-check” badge	CBC	Trust; WTP; backlash

REPAIRSTRAT-QUAL	Planted inherit_solo poison: does stratified alert beat pooled green κ?	Qual	REPAIRSTRAT codebook
------------------	--	------	-------------------------

Falsifiers: Stratified dashboards increase ops time with no revoke/repair latency gain *and* Map gold equal → simplify UI but keep path tags in warehouse. κ-hero equals contingency-first on investigate quality → still ban κ on marketing; keep ops preference test. Windowed alerts thrash TMPTRUST with no poison catch → widen windows / raise thresholds; do not delete revoke path. Pre-register REPAIRSTRAT-* before “science-backed IRR” campaigns.

CXLIX.12 So what for MindCraft commercially

- **Copy:** “When a diagram probe is wrong, we pull it — and we know whether it entered under single, dual, or template-inherit review.” Ban κ-certified / almost-perfect agreement heroes.
 - **Product:** Path-tagged admits; stratified repair dashboard; contingency drill-down; alert→revoke/tighten hooks.
 - **Growth:** Sell **auditable integrity** against ChatGPT fluency *and* against edtech that flashes a single IRR badge while clones leak.
 - **Positioning:** Ops maturity without psychometric cosplay — Kane-honest uses, Einstein-aware metrics.
 - **Metric:** path-stratified post-promote repair; sample-dual discord; revoke latency — demote Kappa % / Dual Rater Score™.
 - **Kill list:** Kappa % NS; pooled-only repair vanity; Landis-band marketing; unconditioned dual-shame boards; ACT/complete-visual from IRR packaging.
 - **Vision:** Maya’s Map stays lit only while the **path that admitted the probe** still survives repair fire — and leadership can see which path is burning.
-

References (verified)

- Cicchetti, D. V., & Feinstein, A. R. (1990). High agreement but low kappa: II. Resolving the paradoxes. *Journal of Clinical Epidemiology*, 43(6), 551–558. [https://doi.org/10.1016/0895-4356\(90\)90159-m](https://doi.org/10.1016/0895-4356(90)90159-m)
 - Cohen, J. (1960). A coefficient of agreement for nominal scales. *Educational and Psychological Measurement*, 20(1), 37–46. <https://doi.org/10.1177/001316446002000104>
 - Feinstein, A. R., & Cicchetti, D. V. (1990). High agreement but low Kappa: I. The problems of two paradoxes. *Journal of Clinical Epidemiology*, 43(6), 543–549. [https://doi.org/10.1016/0895-4356\(90\)90158-l](https://doi.org/10.1016/0895-4356(90)90158-l)
 - Kane, M. T. (2013). Validating the interpretations and uses of test scores. *Journal of Educational Measurement*, 50(1), 1–73. <https://doi.org/10.1111/jedm.12000>
 - Landis, J. R., & Koch, G. G. (1977). The measurement of observer agreement for categorical data. *Biometrics*, 33(1), 159–174. <https://doi.org/10.2307/2529310>
 - Myford, C. M., & Wolfe, E. W. (2009). Monitoring rater performance over time: A framework for detecting differential accuracy and differential scale category use. *Journal of Educational Measurement*, 46(4), 371–389. <https://doi.org/10.1111/j.1745-3984.2009.00088.x>
 - Sgammato, A., & Donoghue, J. R. (2017). On the performance of the marginal homogeneity test to detect rater drift. *Applied Psychological Measurement*, 42(4), 307–320. <https://doi.org/10.1177/0146621617730390>
-

Part CL — Dual-Capacity Hire vs Contract SME Panels under Exam-Week Spikes

Chapter status: Living evidence + ops/workforce brief — Researcher tick 2026-08-15 (UTC hour 12 ≡ Red Team slot, but ch150 never written → prefer Researcher per rotation; researcher count since synthesizer v1.18 = 7 → Researcher)

Primary question: When CAPPLAN dual minutes spike under ACT crunch, should MindCraft **hire W-2 dual raters**, **contract SME panels**, or **hybrid** — without Surge Hire % vanity, unsealed freelancers as second readers, hospital cosplay, or “complete visual by Friday” promises that ignore onboarding lag?

Owners: HITL / bank QA · Ops / WORKFORCE · CAPPLAN / DUALRATE / LOADSHED · Procurement · Brand · Red Team

Commercial job: Ship **SAFE-SURGE**PANEL densifying SAFE-CAPPLAN × SAFE-LOADSHED × SAFE-WORKFORCE (+ DUALRATE / REPAIRSTRAT / TMPTRUST): **make-or-buy dual capacity with seals, ramp gates, and demobilization honesty** — never Surge Hire % NS, never unsealed contractor dump into C4.

Builds on: Parts CXLVII (SAFE-CAPPLAN — dual minutes as capacity), CXXIII (SAFE-LOADSHED), LXXV (SAFE-WORKFORCE), CXLIV (SAFE-DUALRATE), CXLIX (SAFE-REPAIRSTRAT), CXLVIII (SAFE-TMPTRUST), LXVIII (SAFE-HITL), LXXIII (SAFE-PROCURE). Seams: dual surge plan (hire vs contract vs hybrid); panel SOW + seal checklist; parallel-run before C4 authority; FILLETA band refresh when surge mode flips; path-tagged admits from contract_dual vs staff_dual.

CL.1 Why this chapter exists

CAPPLAN made dual minutes a first-class FILLETA line. LOADSHED named exam-week pause lists so tutors do not pour under crunch. WORKFORCE killed headcount≡quality. The missing join:

> When dual-required backlog spikes in ACT week, **who** supplies second-rater minutes — and under what **employment mode** — without lying on parent clocks or poisoning Map?

Wrong resolutions already visible:

1. **Hire-surge theater** — announce dual hires while onboarding lag still zeros capacity.
2. **Contractor dump** — freelancers click “looks good” without FIGKEY/FAMHASH seals or independent uid.
3. **Always-hire dogma** — fixed dual FTE for a two-week spike → idle invent-work or soft-pass.
4. **Always-contract dogma** — treat figure IR as commodity temp labor; firm-specific seal literacy never forms.
5. **Hospital surge cosplay / Surge Hire % NS** — headcount theater sold as ACT readiness.

FOUNDER BELIEF under audit: Parents accept **longer honest dual bands under named surge mode** (staff dual / contract panel / hybrid) more than a short band powered by unsealed freelancers — and LEAs accept **documented SOW + ramp gates** more than “we scaled reviewers 3× this week” ads.

Claims we refuse: Dual Rater Score™ / Surge Hire % / Contract Minutes NS; hire-before-ACT guarantees; Discord/emoji≡second rater; unsealed contract panel → C4; forever-contract without sample dual / REPAIRSTRAT

path; hospital-surge-cosplay ACT ads; complete-visual-by-exam-week from surge packaging.

CL.2 Constructs

Construct	Meaning	MindCraft analogue	Failure mode
Dual surge	Exam-week jump in dual-required arrivals	CAPPLAN λ spike	Ignore; ETA-skip dual
Staff dual	W-2 / ongoing dual raters	Internal development / acquisition	Hire theater without ramp
Contract SME panel	Time-boxed external second raters	Contracting / alliance mode	Unsealed dump
Surge mode tag	Which capacity mix is live	Ops + FILLETA input	Silent mode flip
Parallel-run gate	New panelists shadow before C4 authority	Onboarding seal	Day-one C4 clicks
SAFE-SURGEPANEL	Hire vs contract dual doctrine	This chapter	Surge Hire Score™

Operational definition (HYPOTHESIS): Exam-week dual staffing is SAFE-SURGEPANEL-complete when (1) CAPPLAN publishes a **make-or-buy rule** keyed to spike duration and seal uniqueness, (2) every dual admit carries $\text{dual_source} \in \{\text{staff}, \text{contract_panel}, \text{hybrid_pair}\}$ into REPAIRSTRAT, (3) contract panelists cannot set hcelig_pass until **parallel-run** + **seal checklist** pass, (4) FILLETA **lengthens or labels surge scarcity** when ramp/hire lag binds, (5) LOADSHED pauses electives before burning dual minutes on low-acuity theater, (6) demobilization is planned (no invent-work; no fidelity cliff), and (7) instruments are dual_surge_mode , $\text{contract_parallel_run_pass}$, $\text{dual_source_repair_rate}$ — never Surge Hire % NS.

CL.3 Employment externalization: temps for variable load — not for firm-specific seal skill

FACT: Davis-Blake & Uzzi (1993, *Administrative Science Quarterly*, 38(2), 195–223, doi:10.2307/2393411) — analyze determinants of employment **externalization**. Variation in employment needs **positively** predicts use of temporary workers; requirements for high informational/technical skill and **firm-specific training negatively** predict temp use. Independent contractors follow a partly different pattern (size, multi-site, bureaucratization) but the core lesson is: organizations externalize when load varies **and** when skills are less firm-specific.

Applied (HYPOTHESIS): MindCraft dual IR is **high firm-specificity** (FormatId traps, FAMHASH families, promote_path codes, soft-wrong taxonomy). Pure temp dump for T4/T5 figured admits violates the Davis-Blake–Uzzi prediction and predicts seal error. Contract panels may still cover **variable dual minutes** if MindCraft invests in **panel-specific training + seals** (moving the work toward “trained contractor with firm SOP,” not anonymous piecework).

Kill: Untrained freelancer $\equiv$ dual. **Survive:** Variable-load contracting **only with** seal literacy ramp. **Wound:** 1990s establishment surveys $\neq$ edtech QA — borrow externalization *logic*, not industry coefficients.

CL.4 HR architecture: internal development vs contracting is a portfolio, not a slogan

FACT: Lepak & Snell (1999, *Academy of Management Review*, 24(1), 31–48, doi:10.5465/amr.1999.1580439) — propose an HR **architecture** of employment modes (internal development, acquisition, contracting, alliance) keyed to human-capital **value** and **uniqueness**. Not all roles deserve the same employment mode; contracting fits more generic, lower-uniqueness work; internal development fits high-value, high-uniqueness skills.

Applied (HYPOTHESIS):

Dual work slice	Suggested mode (HYPOTHESIS)	Rationale
Blueprint / first-of-cell T5 dual; TMPTRUST grant	Staff (internal)	Highest uniqueness; trust revoke authority
Routine dual on sealed FAMHASH clones under TMPTRUST sample	Staff or trained panel	Lower uniqueness if template sealed
Exam-week overflow on already-sealed T1–T3	Contract panel with parallel-run	Variable load; Davis-Blake fit
Arbitration / demote authority	Staff lead	Accountability; Kane higher-ambition use

Kill: “We only hire” / “we only contract” brand. **Survive:** Portfolio by uniqueness $\times$ spike duration. **Wound:** Architecture is theory — MindCraft must still measure repair by `dual_source` (REPAIRSTRAT).

CL.5 Mixing standard and nonstandard workers has social costs — design contact and status

FACT: Broschak & Davis-Blake (2006, *Academy of Management Journal*, 49(2), 371–393, doi:10.5465/amj.2006.20786085) — higher proportions of nonstandard workers associate with less favorable supervisor/peer attitudes, higher turnover intentions, and less helping — contingent on mobility location, arrangement type, and contact.

FACT: Davis-Blake, Broschak, & George (2003, *Academy of Management Journal*, 46(4), 475–485, doi:10.5465/30040639) — nonstandard workers can affect exit, voice, and loyalty among **standard** employees.

Applied (HYPOTHESIS): ACT-week panelists who outrank staff on “dual authority” without status rules burn WORKFORCE fidelity. Require named staff dual leads; panelists as **capacity under staff SOPs**; no Surge Hire % hero that shames staff for “not scaling.”

Kill: Freelancer hero brand that demoralizes staff duals. **Survive:** Hybrid with status clarity.

CL.6 Nonstandard jobs are often lower quality — do not build integrity on precarious piecework alone

FACT: Kalleberg, Reskin, & Hudson (2000, *American Sociological Review*, 65(2), 256–278, doi:10.2307/2657440) — show nonstandard employment strongly increases exposure to “bad job” characteristics (low pay, no health insurance, no pension) net of controls.

Applied (HYPOTHESIS): Paying panelists pennies-per-admit predicts speed-over-seals and repair spikes. Commercial implication: contract rates must fund **careful IR time** comparable to staff dual handle-time assumptions in CAPPLAN — or FILLETA must assume higher rework. Cheap surge is fake capacity.

Kill: Race-to-bottom dual piece rates. **Survive:** Rate floors tied to median dual minutes. **Wound:** Bad-jobs macro pattern ≠ every SME contractor — still refuse predatory piecework as integrity strategy.

CL.7 Double-reading programs staff carefully — they do not “Uber” second readers under surge

FACT: Taylor-Phillips & Stinton (2020, *British Journal of Radiology*, 93(1106), 20190610 / PMC7055445) — breast-screening double reading emphasizes independent reads and consensus/arbitration — **quality design**, not headcount theater.

FACT: Taylor-Phillips et al. (2018, *Radiology*, doi:10.1148/radiol.2018171010) — NHSBSP CO-OPS: double reading with arbitration changes recall/detection vs first reader alone — dual **process** matters, not “two people clicked.”

Applied (HYPOTHESIS / transfer wound): Borrow independence, arbitration, and training volume — **never** mammography mortality claims or hospital-surge ACT ads. Panelists need DUALRATE independence/uid; arbitration stays with staff leads on TMPTRUST/REPAIRSTRAT discord.

Kill: Hospital-surge cosplay marketing. **Survive:** Independence + arbitration + training volume as surge non-negotiables.

CL.8 Kane: “we scaled dual review for exam week” is higher-ambition

Reuse (Kane 2013): Higher-ambition uses need more backing (*Journal of Educational Measurement*, 50(1), 1–73, doi:10.1111/jedm.12000).

Applied (HYPOTHESIS): “Exam-week dual capacity” claims need dual_source mix, parallel-run pass rates, path-stratified repair not worse on contract_panel, and FILLETA bands that priced surge — not hire counts or invoice totals.

Kill: Surge headcount as trust proof. **Survive:** Mode-tagged integrity evidence.

CL.9 Claim table

Claim	Label	Confidence	Notes
Variable load predicts temp use; firm-specific skill predicts against temps	FACT	High	Davis-Blake & Uzzi 1993
Employment modes should fit human-capital value/uniqueness	FACT	High	Lepak & Snell 1999 (theory)
Higher nonstandard share can harm standard-worker relations/helping	FACT	High	Broschak & Davis-Blake 2006; Davis-Blake et al. 2003
Nonstandard jobs often carry worse pay/benefits characteristics	FACT	High	Kalleberg et al. 2000
Double-reading literature emphasizes process design (independence, arbitration), not mere headcount	FACT	High	Taylor-Phillips & Stinton 2020; NHSBSP dual-read evidence
Higher-ambition uses need more backing	FACT	High	Kane 2013
Hybrid portfolio (staff core + sealed contract overflow) beats always-hire or always-contract for ACT dual spikes	HYPOTHESIS	Medium–High	Needs SURGEPANEL-*
Parents prefer labeled surge scarcity over freelancer-powered short ETAs	FOUNDER BELIEF	Medium	CBC × FILLETA/HOL ETRUST
Contract panel κ predicts ACT gains	SPECULATION	Low	Refuse
Uber-style dual marketplace ≡ Map integrity	SPECULATION	Low	Refuse

CL.10 Product surface — SAFE-SURGEPANEL claim contract

1. Publish a **make-or-buy matrix** (spike weeks × uniqueness tier → staff / panel / hybrid) owned by CAPPLAN+WORKFORCE — review each exam season.
 2. Tag every dual admit with `dual_source ∈ {staff,contract_panel,hybrid_pair}`; join to REPAIRSTRAT dashboards.
 3. Contract panelists: **independent uid**, seal checklist, **parallel-run gate** before C4 authority; staff lead owns arbitration/demote.
 4. FILLETA: when surge mode is `contract_ramp` or `hire_lag`, **lengthen or label** dual-verify scarcity — never ETA-skip dual (CAPPLAN ban).
 5. LOADSHED: pause electives before burning dual minutes on low-acuity GENQ vanity.
 6. Ban Dual Rater Score™ / Surge Hire % / Contract Minutes / “3× reviewers this week” marketing heroes.
 7. Rate floor: $\text{contract } \$/\text{admit} \geq \text{CAPPLAN median dual-minute cost assumption}$ (or raise rework factor).
 8. Demobilization plan: end dates, knowledge handoff, no invent-dual-work; no abrupt exit that spikes staff repair.
 9. PROCURE: SOW includes discordance/repair SLAs, data handling, and sample-dual rights (TMPTRUST).
 10. Copy spine: “Exam-week dual capacity is staff-led; contractors only after seals and a parallel run — and parent clocks name the surge mode.” Ban hospital-surge / Uber-rater ads.
-

CL.11 Doctrine — SAFE-SURGEPANEL (provisional)

1. **Portfolio, not slogan** — Lepak–Snell modes by uniqueness; Davis-Blake variable load ≠ unskilled dump.
2. **Staff core owns trust** — blueprint dual, TMPTRUST grant/revoke, arbitration.
3. **Contract overflow only behind gates** — parallel-run + seals + independent uid.
4. **Tag the source** — `dual_source` into REPAIRSTRAT; blame/repair conditioned on source.
5. **Price surge in FILLETA** — labeled scarcity / longer bands beat silent freelancer speed.
6. **Protect staff climate** — Broschak mixing costs; no freelancer-hero shame.
7. **Refuse predatory piece rates** — Kalleberg bad-job pattern as integrity risk.
8. **Borrow dual-read process, not hospital brand** — independence/arbitration without cosplay.
9. **No Dual Rater Score™ / Surge Hire % / Contract Minutes NS.**
10. **No ACT / complete-visual-by-exam-week guarantees** from surge packaging.

Confidence: High on externalization / HR-architecture / mixing / bad-jobs / Kane facts *with transfer limits*. Medium on exact uniqueness tiers, parallel-run length, and rate-floor formula. High that unsealed contractor dump and hire-surge marketing are commercially toxic under CAPPLAN honesty.

CL.12 Experiments

ID	Question	Design	Primary
SURGEPANEL-1	Staff-only dual vs hybrid sealed panel under simulated ACT spike	Ops A/B (shadow week)	Dual latency; path-stratified repair; FILLETA honesty

SURGEPANEL-2	Parallel-run gate (N tickets) vs day-one C4 authority for panelists	Ops A/B	Repair rate; discord; time-to-useful-capacity
SURGEPANEL-3	Rate floor vs cheap piecework panel	Ops A/B	Seal miss rate; rework minutes; quit mid-spike
SURGEPANEL-4	Parent CBC: labeled surge scarcity vs “we hired 10 dual raters” hero	CBC	Trust; WTP; backlash
SURGEPANEL-5	Demobilization plan vs abrupt panel cutoff	Ops A/B	Post-spike staff repair spike; helping climate
SURGEPANEL-QUAL	SME panel onboarding diary — where seal literacy fails	Qual	SURGEPANEL codebook

Falsifiers: Hybrid panel equals staff-only on repair *and* latency with lower cost → keep hybrid but retain staff core for TMPTRUST. Parallel-run adds delay with no repair gain → shorten gate but keep seal checklist. Parent CBC prefers hire-count hero with equal retention → still ban guarantee copy; test honesty separately. Pre-register SURGEPANEL-* before exam-week dual-scaling campaigns.

CL.13 So what for MindCraft commercially

- **Copy:** “ACT-week dual capacity is staff-led; contractors only after seals and a parallel run — and verify clocks name the surge mode.” Ban Uber-rater / hospital-surge / Surge Hire % heroes.
- **Product:** Make-or-buy matrix; dual_source tags; parallel-run gate; FILLETA surge labels; demobilization checklist.
- **Growth / positioning:** Capacity-true integrity under crunch vs ChatGPT fluency and freelance “scale reviewers” theater.
- **Metric:** dual_surge_mode; contract_parallel_run_pass; dual_source-stratified repair — demote Surge Hire % / Dual Rater Score™.
- **Kill list:** Hire-surge theater; unsealed contractor dump; Discord≡dual; hospital cosplay; complete-visual-by-exam-week from surge packaging.
- **Vision:** Maya’s Map stays dual-true in crunch because **who** second-read the probe is named, trained, and demotable — not because a surge invoice looked large.

References (verified)

- Broschak, J. P., & Davis-Blake, A. (2006). Mixing standard work and nonstandard deals: The consequences of heterogeneity in employment arrangements. *Academy of Management Journal*, 49(2), 371–393.
<https://doi.org/10.5465/amj.2006.20786085>
- Davis-Blake, A., Broschak, J. P., & George, E. (2003). Happy together? How using nonstandard workers affects exit, voice, and loyalty among standard employees. *Academy of Management Journal*, 46(4), 475–485.
<https://doi.org/10.5465/30040639>

- Davis-Blake, A., & Uzzi, B. (1993). Determinants of employment externalization: A study of temporary workers and independent contractors. *Administrative Science Quarterly*, 38(2), 195–223. <https://doi.org/10.2307/2393411>
- Kalleberg, A. L., Reskin, B. F., & Hudson, K. (2000). Bad jobs in America: Standard and nonstandard employment relations and job quality in the United States. *American Sociological Review*, 65(2), 256–278. <https://doi.org/10.2307/2657440>
- Kane, M. T. (2013). Validating the interpretations and uses of test scores. *Journal of Educational Measurement*, 50(1), 1–73. <https://doi.org/10.1111/jedm.12000>
- Lepak, D. P., & Snell, S. A. (1999). The human resource architecture: Toward a theory of human capital allocation and development. *Academy of Management Review*, 24(1), 31–48. <https://doi.org/10.5465/amr.1999.1580439>
- Taylor-Phillips, S., et al. (2018). Double reading in breast cancer screening: Cohort evaluation in the CO-OPS trial. *Radiology*. <https://doi.org/10.1148/radiol.2018171010>
- Taylor-Phillips, S., & Stinton, C. (2020). Double reading in breast cancer screening: Considerations for policy-making. *British Journal of Radiology*, 93(1106), 20190610. <https://doi.org/10.1259/bjr.20190610> (PMC7055445)

Part CLI — Template-Trust Revoke & Sample-Dual Cadence Under Repair Spikes

Chapter status: Living evidence + ops/QA brief — Researcher tick 2026-08-15 (UTC hour 18 ≡ Red Team slot, but ch151 never written → prefer Researcher per rotation; researcher count since synthesizer v1.19 = 0 → Researcher)

Primary question: After SAFE-TMPTRUST sets `template_trust`, **when and how fast** must MindCraft **revoke** trust and **tighten** sample-dual intensity under repair/discord spikes — without forever-trust latch, Template Trust Score™ vanity, AQL cut-scores as product law, or frozen-5% sample theater that ignores REPAIRSTRAT path poison?

Owners: HITL / bank QA · Engine (`template_trust` + sample %) · TMPTRUST / REPAIRSTRAT / CAPPLAN / PROMOTE · Brand · Red Team

Commercial job: Ship **SAFE-REVOKE** densifying TMPTRUST × REPAIRSTRAT × CAPPLAN: **written revoke triggers + switching-rule sample cadence** — tighten on failure, reduce only after calm history, re-arm T5 after revoke.

Builds on: CXLVIII (TMPTRUST), CXLIX (REPAIRSTRAT), CXLVII (CAPPLAN), CXLIV (DUALRATE), CXL (PROMOTE), CL (SURGEANAL). Seams: trust state ∈ {false, true, tightened, revoked}; sample-rate bands; revoke → FILLETA dual bump.

CLI.1 Why this chapter exists

TMPTRUST said: blueprint dual once; inherit intensity; sample + revoke. REPAIRSTRAT said: watch path-stratified repair and sample discord over windows. CAPPLAN said: dual minutes are scarce. The missing join is **ops law for the bit and the %** — what event clears trust, what raises sample rate from 5%→20%→100%, what re-qualifies a template after revoke, and how parent/FILLETA copy names “we tightened review” without shame theater or Score™.

Failure modes: (1) forever-trust after one dual (silent inherit’s twin); (2) revoke never fires while inherit_solo repair climbs; (3) sample % frozen at launch forever; (4) Template Trust Score™ / Inherit Health % NS; (5) ISO AQL tables copy-pasted as ACT law; (6) hospital CUSUM cosplay ads; (7) surge week (SURGEPANEL) freezes reduced inspection to fake FILLETA green.

FOUNDER BELIEF under audit: Parents accept a named “second-check tightened for this diagram family” more than a glossy always-trusted template badge — if hole clocks stay honest when dual minutes return.

Claims we refuse: Template Trust Score™ / Revoke Minutes / Sample Dual % as NS; forever-trust after blueprint; never-tighten sample; factory AQL≡product cut-score; mammography/hospital CUSUM cosplay ACT ads; complete-visual-by-date from “trust health.”

CLI.2 Constructs

Construct	Meaning	MindCraft analogue	Failure mode
Trust state	Ops license for inherit intensity	false / true / tightened / revoked	Boolean forever-true
Sample-dual rate	Fraction of inherit_solo admits that get second rater	Policy band (e.g. 5% / 20% / 100%)	Frozen launch %
Revoke trigger	Event that clears inherit license	Critical repair / sample discord / demote	Soft “investigate later”
Tighten / reduce	Raise or lower sample intensity from history	Switching-rule analogue	ISO table as law
Re-qualify	Path back to true after revoke	New blueprint dual or written re-seal	Silent flip from play count
SAFE-REVOKE	Cadence doctrine for trust + sample %	This chapter	Template Trust Score™

Operational definition (HYPOTHESIS): SAFE-REVOKE-complete when (1) trust is a **state** not a latch, (2) hard revoke clears inherit within SLA and re-arms T5, (3) sample rate follows switching bands from REPAIRSTRAT inherit repair/discord windows, (4) tightened/revoked states bump FILLETA dual estimates, (5) re-qualify needs dual blueprint or written re-seal — never play-count, (6) κ / Sample Dual % are not North Stars, (7) surge cannot override revoke to protect ETA.

CLI.3 Switching rules exist so intensity tracks quality history

FACT: ISO 2859-1:1999 defines **switching rules** among normal, tightened, and reduced inspection for continuing lots: severity tracks accept/reject history; reduced inspection is an *incentive* after consistently good quality (authority discretion); tightened / discontinuation protect consumers when quality deteriorates (§1.2, §9.3).

Applied (HYPOTHESIS): Treat inherit admits as a continuing series. Default after TMPTRUST G1–G6; reduced-like sample %. Two critical inherit failures in a short window → tighten (raise %; optionally pause inherits). Calm under tightened → normal. Reduce only after pre-registered calm streak + authority — never during surge theater. Exact ISO “2-of-5 / 10 consecutive” numbers are analogy — set bands in ops policy; pre-register REVOKE-.*.

Kill: Frozen 5% forever; forever-trust latch. **Survive:** History-triggered tighten/revoke. **Wound:** Borrow switching logic, not Ac/Re tables as ACT law.

CLI.4 Detect shifts early — do not wait for a quarterly IRR ceremony

FACT: Shewhart (1931, *Economic Control of Quality of Manufactured Product*, D. Van Nostrand) — control charts distinguish **chance-cause** fluctuation from **assignable-cause** shifts; points outside limits demand investigation and corrective action, not averaging into a seasonal “quality score.”

FACT: Page (1954, *Biometrika*, 41(1–2), 100–115, doi:10.1093/biomet/41.1-2.100) — **cumulative sum (CUSUM)** continuous inspection schemes accumulate evidence across the process history to detect persistent mean shifts sooner than single-point Shewhart rules alone.

Applied (HYPOTHESIS): Windowed inherit_solo repair rate and sample-dual discord are the process characteristics. SAFE-REVOKE uses **simple Shewhart-style thresholds** (and optional CUSUM on weekly inherit repair) to fire tighten/revoke — not a launch-week κ plaque (REPAIRSTRAT) and not a Template Trust Score™ that hides assignable causes. When CUSUM/window trips, clear template_trust or move to tightened with 100% sample until re-qualify.

Kill: One-shot IRR certificate as trust health. **Survive:** Ongoing shift detection → revoke/tighten. **Wound:** Charts stay ops-behind-login — no CUSUM-hero ads.

CLI.5 Sample size for agreement is about precision — not a magic 5% brand

FACT: Flack, Afifi, Lachenbruch, & Schouten (1988, *Psychometrika*, 53(3), 321–325, doi:10.1007/BF02294215) — give methods to choose sample size so that a confidence interval for two-rater κ meets a prespecified width (or so that a hypothesis test about κ has planned power); agreement monitoring is **precision-planned**, not vibes-planned.

Reuse (REPAIRSTRAT): Myford & Wolfe (2009, *Journal of Educational Measurement*, 46(4), 371–389, doi:10.1111/j.1745-3984.2009.00088.x) — rater performance drifts over a campaign; monitoring must be longitudinal.

Applied (HYPOTHESIS): Sample-dual % is a **capacity** × **detection** tradeoff. Too thin → discord never fires (false safety). Too thick → CAPPLAN burns dual minutes and inheritance’s point disappears. Plan sample size from expected discord prevalence and desired CI width on path agreement — then map to a % of inherit volume. Under repair spikes, prefer **temporary 100% sample or revoke** over tiny % “theater audits.” Ban Sample Dual % as marketing hero.

Kill: “Industry-standard 5% audit” as brand. **Survive:** Precision-aware bands + spike override.

CLI.6 Clone variability makes revoke a feature, not an apology

Reuse (TMPTRUST): Glas & van der Linden (2003, *Applied Psychological Measurement*, 27(4), 247–261, doi:10.1177/0146621603027004001) — item cloning induces **parameter variability** across clones; multilevel models treat clones as draws from a family distribution.

Applied (HYPOTHESIS): Inheritance assumed family stability after blueprint. A repair spike on inherit_solo is evidence the family draw process drifted (generator bug, rater soften, seal leak). Revoke is the honest IUA move: stop claiming reduced inspection until a new dual blueprint re-qualifies the model. Selling “template trusted” while repairs accumulate is Kane-incoherent.

Kill: Soft-ignore inherit repairs to protect FILLETA dates. **Survive:** Revoke → re-arm dual → honest fill-band slip.

CLI.7 Higher ambition uses need revoke-visible backing

Reuse (Kane 2013): More ambitious interpretations/uses need more backing; rejecting a use need not kill a weaker interpretation (*Journal of Educational Measurement*, 50(1), 1–73, doi:10.1111/jedm.12000).

Applied (HYPOTHESIS): Parent-facing “matching variants move faster under a dual-sealed template” is ambition-conditional on live sample+revoke. If trust stays green while inherit_solo repairs spike, the use is unsupported — demote copy to “machine seals only” until re-qualify. External copy may say “we pause family speed-lanes when checks disagree”; it must not cite Template Trust Score™ or CUSUM cosplay.

Kill: Trust badge ads during open critical inherit repair. **Survive:** Named pause / tighten language.

CLI.8 Cadence gates (doctrine core)

HYPOTHESIS — hard revoke (→ revoked, clear inherit, re-arm T5/dual) when ANY of:

Trigger	Why
R1 Critical post-promote REPAIR / demote on inherit_solo for template/family	Path poison (REPAIRSTRAT)
R2 Sample-dual discord (second rater reject/repair) on inherit path	Reduced inspection failed
R3 FIGKEY/GENQ demote on any family member while trust true	Seal drift
R4 Windowed inherit_solo repair exceeds pre-registered Shewhart/CUSUM trip	Assignable shift
R5 Ops investigate-fail / authority halt	Human stop-ship

Tighten (→ tightened, raise sample % toward 100%, optionally pause new inherits) when: soft discord pattern short of R1–R2; first R4 near-miss; surge rater mix changes (SURGE PANEL) without parallel-run green.

Reduce sample % only when ALL of: trust state true; pre-registered calm streak on inherit repair + sample discord; CAPPLAN headroom exists; responsible authority approves (ISO reduced-inspection spirit).

Re-qualify to true: new consensus dual blueprint for cell/template **or** written dual re-seal of N planted clones — never play count, thumbs-up, or ETA pressure.

FILLETA coupling: revoke/tighten **must** increase estimated dual minutes in the fill estimator; silent ETA while dual returns is CAPPLAN fraud.

CLI.9 Claim table

Claim	Label	Confidence	Notes
Attribute sampling schemes include switching among normal/tightened/reduced from lot history	FACT	High	ISO 2859-1:1999 §9.3
Control charts separate chance vs assignable causes; out-of-limit points demand action	FACT	High	Shewhart 1931
CUSUM accumulates history to detect persistent shifts	FACT	High	Page 1954
κ CI width / power planning implies sample size is not arbitrary	FACT	High	Flack et al. 1988
Rater performance can drift over a scoring campaign	FACT	High	Myford & Wolfe 2009
Cloned items exhibit within-family parameter variability	FACT	High	Glas & van der Linden 2003
Written revoke + switching sample beats forever-trust and frozen-% theater	HYPOTHESIS	Medium–High	Needs REVOKE-*
Parents prefer named tighten/pause over trust-badge vanity	FOUNDER BELIEF	Medium	CBC × HOLETRUST/ FILLETA
Exact ISO 2-of-5 / score $\geq$ 30 tables are MindCraft product law	SPECULATION	Low	Refuse

CUSUM UI ≡ clinical-grade integrity brand	SPECULATION	Low	Refuse
---	-------------	-----	--------

CLI.10 Product surface — SAFE-REVOKE claim contract

1. `template_trust` is a visible state machine (false/true/tightened/revoked) on promote UI — never a one-way latch.
2. Hard revoke clears inherit within SLA; tickets show **re-arm dual** reason codes R1–R5.
3. `sample_dual_rate` is policy-banded and logged; changes emit `sample_dual_rate_change` events.
4. REPAIRSTRAT dashboards drive cadence — path `inherit_solo` first; no pooled Score™.
5. FILLETA/CAPPLAN estimators read trust state; revoke widens bands / labels dual scarcity.
6. SURGEPANEL cannot force true or reduced % to defend exam-week ETA.
7. Telemetry: `template_trust_revoke`, `template_trust_tighten`, `template_trust_requalify`, `sample_dual_rate_change` — demote Template Trust Score™ / Revoke Minutes NS.
8. Copy: “If checks disagree on a trusted template family, we pause the speed-lane and dual-seal again.” Ban forever-trusted / CUSUM-hero / AQL-as-ACT ads.
9. Ban Discord-emoji “sample dual.”
10. Ban silent re-qualify from play counts.

CLI.11 Doctrine — SAFE-REVOKE (provisional)

1. **Trust is revocable** — TMPTRUST without SAFE-REVOKE is a latch costume.
2. **Switching-rule sample cadence** — tighten on failure history; reduce only after calm + authority (ISO spirit, transfer-limited).
3. **Detect shifts continuously** — Shewhart/CUSUM logic on inherit repair/discord; no quarterly certificate.
4. **Precision-plan sample size** — Flack-style intent; spike → 100% or revoke, not thinner theater.
5. **Clone variability expects revoke** — Glas warning operationalized.
6. **Kane honesty** — pause speed-lane copy when trust revoked.
7. **FILLETA honesty** — revoke costs dual minutes in the estimator.
8. **No Template Trust Score™ / Sample Dual % / Revoke Minutes NS.**
9. **No ACT / complete-visual / AQL-cosplay / hospital-CUSUM guarantees** from cadence packaging.
10. Copy spine: “We dual-seal the first diagram probe, sample later variants, and **revoke the speed-lane when checks disagree** — then dual-seal again before trusting the family.”

Confidence: High on ISO switching / Shewhart / Page / Flack / Glas / Myford facts *with transfer limits*. Medium on exact trip thresholds and default sample bands (pre-register REVOKE-*). High that forever-trust and frozen-% are commercially toxic under Map-facing FormatId diagnosis.

CLI.12 Experiments

ID	Question	Design	Primary
REVOKE-1	Written revoke-on-R1/R2 vs soft-investigate-only after inherit repair	Ops A/B	Illicit Map dwell; dual minutes; fill-band honesty
REVOKE-2	Switching sample (5↔20↔100) vs frozen 5% under planted discord	Ops A/B	Time-to-catch; dual burn; false safety
REVOKE-3	Shewhart window trip vs CUSUM trip vs weekly human review only	Ops A/B	False revoke rate; missed poison; CAPPLAN util
REVOKE-4	Re-qualify via new blueprint dual vs N planted dual re-seals	Ops A/B	Post-requalify repair; time-to-fill
REVOKE-5	Parent CBC: “we paused this diagram family’s speed-lane” vs trust-badge / Score™	CBC	Trust; WTP; backlash
REVOKE-QUAL	Inject generator drift mid-trust: does cadence revoke before Map poison spreads?	Qual	REVOKE codebook

Falsifiers: Revoke≈soft-investigate on Map poison → keep forever-dual; treat TMPTRUST unsafe. Switching≈frozen 5% on catch *and* no dual savings when calm → fixed thick sample or kill inherit. Excessive false revoke → raise thresholds or human-gate. Pre-register REVOKE-* before “self-healing template trust” campaigns.

CLI.13 So what for MindCraft commercially

- **Copy:** “Speed-lanes for matching diagram variants stay on only while sample checks agree — disagreement pauses the lane and re-runs a dual seal.” Ban forever-trusted template badges.
- **Product:** Trust state machine; sample-rate policy; revoke→T5 re-arm; FILLETA dual bump on revoke.
- **Growth / positioning:** Sell integrity that can *stop* — against fluency rubber-stamps and against CUSUM/AQL QA costume.
- **Metric:** revoke rate; time-to-revoke; sample-rate changes; dual minutes vs repair delta — demote Template Trust Score™.
- **Kill list:** Forever-trust latch; frozen sample theater; Template Trust Score™; AQL≡ACT; CUSUM-hero ads; ETA-forced trust.
- **Vision:** Maya’s Map keeps a diagram cell lit only while the family speed-lane remains earned.

References (verified)

- Flack, V. F., Afifi, A. A., Lachenbruch, P. A., & Schouten, H. J. A. (1988). Sample size determinations for the two rater kappa statistic. *Psychometrika*, 53(3), 321–325. <https://doi.org/10.1007/BF02294215>
 - Glas, C. A. W., & van der Linden, W. J. (2003). Computerized adaptive testing with item cloning. *Applied Psychological Measurement*, 27(4), 247–261. <https://doi.org/10.1177/0146621603027004001>
 - International Organization for Standardization. (1999). *ISO 2859-1:1999 Sampling procedures for inspection by attributes — Part 1: Sampling schemes indexed by acceptance quality limit (AQL) for lot-by-lot inspection*. ISO. (Switching-rule analogy only; not product cut-score law.)
 - Kane, M. T. (2013). Validating the interpretations and uses of test scores. *Journal of Educational Measurement*, 50(1), 1–73. <https://doi.org/10.1111/jedm.12000>
 - Myford, C. M., & Wolfe, E. W. (2009). Monitoring rater performance over time: A framework for detecting differential accuracy and differential scale category use. *Journal of Educational Measurement*, 46(4), 371–389. <https://doi.org/10.1111/j.1745-3984.2009.00088.x>
 - Page, E. S. (1954). Continuous inspection schemes. *Biometrika*, 41(1–2), 100–115. <https://doi.org/10.1093/biomet/41.1-2.100>
 - Shewhart, W. A. (1931). *Economic control of quality of manufactured product*. D. Van Nostrand Company.
-

Part CLII — Path-Conditioned Dual Blame Tests (Selection vs Quality)

Chapter status: Living evidence + ops/QA brief — Researcher tick 2026-08-15 (UTC hour 21 → Researcher per rotation; hour%6≠0; researcher count since synthesizer v1.19 = 1 → Researcher)

Primary question: When REPAIRSTRAT shows **higher raw repair or discord on dual / inherit_sample** than on single / inherit_solo, how must MindCraft decide whether that is **selection** (harder tickets routed to dual — good triage) vs **quality failure** (second-rater / process poison) — without Dual Shame Score™ boards, pooled κ theater, or “kill dual because the bar is red”?

Owners: HITL / bank QA · Engine (promote_path + T-trigger + risk strata) · SAFE-REPAIRSTRAT / SAFE-DUALRATE / SAFE-CAPPLAN / SAFE-REVOKE · Brand · Red Team

Commercial job: Ship **SAFE-PATHBLAME** densifying REPAIRSTRAT × DUALRATE × CAPPLAN (+ TMPTRUST/REVOKE): **path-conditioned, risk-stratified blame tests** — compare paths inside comparable ticket strata; never naive dual-vs-single shame.

Builds on: CXLIX (REPAIRSTRAT foreshadowed selection confounding), CXLIV (DUALRATE risk triggers), CXLVIII/CLI (TMPTRUST/REVOKE inherit paths), CXLVII (CAPPLAN dual minutes), XC (INSTRUMENT). Seams: promote_path × dual_trigger / risk stratum tables; selection-adjusted repair deltas; no Dual Shame Score™.

CLII.1 Why this chapter exists

REPAIRSTRAT killed pooled repair vanity and warned: dual can look “worse” because DUALRATE **selects harder tickets** (T4 ambiguity, first-of-cell, self-flag). CAPPLAN and SURGEPANEL then put dual minutes under crunch. REVOKE reacts to inherit_solo poison. The missing join is **an ops-legal blame procedure**: before ops “fix dual,” “fire the second rater,” or “widen FILLETA because dual is broken,” the dashboard must answer — *same difficulty mix, different path?* or *same path, different difficulty?*

Failure modes: (1) naive dual > single repair $\Rightarrow$ Dual Shame board / kill dual; (2) inherit_solo looks green only because easy clones inherit while hard stems go dual; (3) Propensity Theater™ / Causal DAG Score™ as NS; (4) hospital case-mix cosplay ACT ads; (5) silent ignore of path gaps when dual is politically sacred; (6) κ / Dual Rater Score™ as substitute for stratum tables.

FOUNDER BELIEF under audit: Leadership and LEAs accept “dual takes the hard tickets; within hard tickets, dual repair is flat/lower” more than a red dual bar that scares them into single-rater-only Map — if the tables are readable without biostat cosplay.

Claims we refuse: Dual Shame Score™ / Selection-Adjusted Repair % / Path Blame Minutes NS; unconditioned dual-vs-single rankings; “science-backed dual is safer” ads from naive path deltas; hospital case-mix / propensity cosplay as product brand; ACT / complete-visual guarantees from blame packaging.

CLII.2 Constructs

Construct	Meaning	MindCraft analogue	Failure mode
Path mix	Share of admits by promote_path	single / dual / inherit_solo / inherit_sample	Untagged admits
Risk stratum	Ticket difficulty / DUALRATE trigger class	T-trigger bucket; FormatId×cell hardness	Path-only tables
Raw path repair	Repair rate within path, unadjusted	REPAIRSTRAT row	Treated as causal
Stratum-conditioned repair	Repair within path and risk stratum	Primary blame table	Optional “nice to have”
Selection story	Dual volume concentrates in hard strata	Mix charts	Hidden
Quality story	Within stratum, path A worse than path B	Within-cell deltas	Dual Shame without strata
SAFE-PATHBLAME	Blame doctrine for path deltas	This chapter	Dual Shame Score™

Operational definition (HYPOTHESIS): Dual/path blame is SAFE-PATHBLAME-complete when (1) every admit carries immutable promote_path + dual_trigger (or risk_stratum) tags, (2) ops default view is **path × stratum** repair (and discord) before any pooled or path-only rollup, (3) “dual worse” claims require within-stratum evidence or an explicit selection narrative with mix charts, (4) CAPPLAN/DUALRATE/REVOKE actions that cite repair deltas

must cite the conditioned table (or a pre-registered adjusted estimate), (5) κ / Dual Shame Score™ / Selection-Adjusted % are not North Stars or splash metrics, (6) parent/LEA copy may say “harder diagram tickets get a second check” without ranking paths by raw repair, and (7) PATHBLAME-* falsifiers are pre-registered before kill-dual campaigns.

CLII.3 Aggregation reverses the story — condition on the department (stratum)

FACT: Bickel, Hammel, & O’Connell (1975, *Science*, 187(4175), 398–404, doi:10.1126/science.187.4175.398) — UC Berkeley graduate admissions: **aggregate** rates suggested bias against women; **department-stratified** rates showed little anti-woman bias and often the reverse. The aggregate pattern reflected **differential application mix** across departments with different admission difficulties — not a single pooled discrimination parameter.

Applied (HYPOTHESIS): Replace “department” with **risk stratum / DUALRATE trigger**, and “admit rate by sex” with **repair rate by promote_path**. Dual may show higher raw repair because it **applies into harder strata** (ambiguity, first-of-cell, surge IR), while single/inherit_solo apply into easier sealed clones. A Dual Shame board that ranks paths on raw repair is a Berkeley-aggregate error.

Kill: Path-only shame rankings. **Survive:** Mix chart + within-stratum tables. **Wound:** Admissions fairness $\neq$ item QA — borrow *stratify before blame*, not Title IX cosplay.

CLII.4 Selection into the analysis set can invent associations

FACT: Berkson (1946, *Biometrics Bulletin*, 2(3), 47–53, doi:10.2307/3002000) — fourfold analyses on **hospital-selected** populations can show spurious disease associations because admission probabilities differ across conditions; the selected table is not the parent-population table.

FACT: Hernán, Hernández-Díaz, & Robins (2004, *Epidemiology*, 15(5), 615–625, doi:10.1097/01.ede.0000135174.63482.43) — unify many “selection bias” examples as **conditioning on a common effect (collider)** of variables related to exposure and outcome; distinguish that structure from confounding by common causes; argue diagrams clarify when associations among the selected are not the associations of interest.

FACT: Elwert & Winship (2014, *Annual Review of Sociology*, 40, 31–53, doi:10.1146/annurev-soc-071913-043455) — label **endogenous selection bias** as collider-conditioning (including Berkson-type cases); warn that bias can arise from analysis choices that condition on colliders even when no rows are discarded.

Applied (HYPOTHESIS): MindCraft’s dual path is a **selected set**: tickets enter dual because of triggers correlated with repair risk. Comparing dual vs single repair **unconditionally** estimates a mixture of (a) selection intensity and (b) path quality. Conditioning further on post-repair outcomes alone (e.g., only demoted items’ path mix) can open collider paths — ops should prefer **design strata fixed at promote time** (dual_trigger, hardness band) over outcome-conditioned blame slices.

Kill: “Dual repair high $\Rightarrow$ dual causes Map poison” from selected aggregates. **Survive:** Promote-time strata + within-stratum contrasts. **Wound:** Do not require full do-calculus UI; require stratum discipline.

CLII.5 When the question is causal (“does dual help?”), balance observed risk — don’t cosplay RCT

FACT: Austin (2011, *Multivariate Behavioral Research*, 46(3), 399–424, doi:10.1080/00273171.2011.568786) — propensity scores (probability of treatment given baseline covariates) support matching, stratification, weighting, or covariate adjustment so observational contrasts better approximate balanced comparisons on **observed** confounders; balance diagnostics matter; methods reduce confounding under stated assumptions, not magic.

Applied (HYPOTHESIS): Optional ops tool: estimate $P(\text{dual} \mid \text{risk features})$ or simply **stratify/match on DUALRATE trigger + FormatId + cell age**, then compare repair within matched/stratum sets. Prefer **transparent stratum tables** as default; propensity models are an advanced aide for CAPPLAN debates (“is dual worth minutes?”), never a Propensity Score™ splash or Dual Efficacy Guarantee.

Kill: Unadjusted dual-vs-single as causal proof. **Survive:** Stratum or matched contrasts before process change. **Wound:** Unobserved generator bugs can still confound — PATHBLAME does not replace FIGKEY/GENQ seals.

CLII.6 Higher-ambition “dual protects Map” needs conditioned backing

Reuse (Kane 2013): More ambitious interpretations/uses need more backing; rejecting a use need not kill a weaker interpretation (*Journal of Educational Measurement*, 50(1), 1–73, doi:10.1111/jedm.12000).

Applied (HYPOTHESIS): Claiming “second checks catch hard diagram failures” is higher ambition than “we route hard tickets to dual.” IUA backing: (1) mix chart shows dual concentrates in hard strata, (2) within hard strata, dual discord/repair is not worse than single (or is better on pre-registered endpoints), (3) REVOKE still fires on inherit_solo poison. If only (1) holds, demote copy to routing honesty — do not sell dual as quality-proven from selection alone.

Kill: Dual-protects-Map ads from raw path bars. **Survive:** Conditioned evidence ladder.

CLII.7 Blame procedure (doctrine core)

HYPOTHESIS — before any “blame dual / blame single / kill inherit” action:

Step	Required view	Stop if...
B1	Path mix by risk stratum (volume)	Mix unknown / untagged
B2	Repair (and sample discord) within each stratum by path	Only pooled or path-only rates exist
B3	Selection narrative: does dual’s excess raw repair vanish inside strata?	If yes → selection story; do not Dual-Shame

B4	If within-stratum dual worse → quality story: check rater drift, seal leaks, surge mix (SURGEPANEL), not κ alone	
B5	If inherit_solo worse inside easy strata → REVOKE/TMPTRUST, not “dual is fine so ignore”	
B6	CAPPLAN change that cites repair must paste stratum table id + window	ETA theater without B2

Decision labels (ops, not scores): selection_explained | quality_signal | underpowered_stratum | tags_missing.
Ban Dual Shame Score™ numeric vanity.

CLII.8 Claim table

Claim	Label	Confidence	Notes
Aggregate association can reverse or vanish after stratifying on a mix confounder	FACT	High	Bickel et al. 1975
Hospital/selection mechanisms can induce spurious associations in selected tables	FACT	High	Berkson 1946
Selection bias often shares a collider-conditioning structure distinct from confounding	FACT	High	Hernán et al. 2004; Elwert & Winship 2014
Propensity methods can reduce confounding by observed baselines in observational contrasts	FACT	High	Austin 2011
Dual vs single raw repair is not a causal quality ranking without stratum/mix discipline	HYPOTHESIS	Medium–High	Needs PATHBLAME-*
Readable stratum tables beat Dual Shame boards for ops and LEA trust	FOUNDER BELIEF	Medium	CBC later

Full propensity UI / DAG explorer is required product surface	SPECULATION	Low	Refuse
Case-mix adjustment $\equiv$ ACT-point proof for dual	SPECULATION	Low	Refuse

CLII.9 Product surface — SAFE-PATHBLAME claim contract

1. Default REPAIRSTRAT screen: **path** $\times$ **risk_stratum** repair + volume; path-only chart secondary.
 2. Dual Shame / “worst path” leaderboards **off** by default; require authority + stratum paste to enable.
 3. Telemetry: path_stratum_repair, path_mix_by_stratum, blame_label $\in$ {selection_explained, quality_signal, underpowered_stratum, tags_missing}.
 4. DUALRATE/CAPPLAN/REVOKE runbooks must reference B1–B6 before kill-dual or forever-dual.
 5. Copy: “Harder diagram tickets get a second check; we judge those checks **inside** difficulty bands — not by a raw dual bar.”
 6. Ban Propensity Score™ / Dual Shame Score™ / Case-Mix Minutes NS splash.
 7. Ban hospital/mammography case-mix cosplay ACT ads.
 8. Underpowered strata $\rightarrow$ accumulate or widen band; do not invent quality_signal from $n=3$.
 9. κ remains aide (REPAIRSTRAT), never blame verdict.
 10. Parent/LEA packets may show mix honesty; never child-facing Dual Shame.
-

CLII.10 Doctrine — SAFE-PATHBLAME (provisional)

1. **Stratify before blame** — Berkeley lesson transferred to promote paths.
2. **Selection $\neq$ quality** — dual’s hard-ticket mix is often the point of DUALRATE.
3. **Promote-time strata** — prefer trigger/hardness tags over outcome-conditioned slices (collider caution).
4. **Optional propensity/matching** — ops aide under Austin-style discipline; not brand.
5. **Kane ladder** — dual-protects-Map needs within-stratum backing.
6. **Couple to REVOKE** — inherit_solo poison still revokes even if dual looks “busy.”
7. **Couple to CAPPLAN** — dual minutes defended only with selection-aware story.
8. **No Dual Shame Score™ / Selection-Adjusted Repair % NS.**
9. **No ACT / complete-visual / case-mix-cosplay guarantees** from blame packaging.
10. Copy spine: “We send ambiguous diagram probes to a second rater — and we **compare repair inside the same difficulty band**, so we don’t punish the path that took the hard tickets.”

Confidence: High on Bickel / Berkson / Hernán / Elwert–Winship / Austin facts *with transfer limits*. Medium on exact stratum taxonomy and default bands (pre-register PATHBLAME-*). High that Dual Shame boards are commercially and epistemically toxic.

CLII.11 Experiments

ID	Question	Design	Primary
PATHBLAME-1	Path×stratum default UI vs path-only Dual Shame board	Ops A/B	Wrong kill-dual rate; time-to-correct narrative
PATHBLAME-2	Planted selection (route hard→dual) with equal within-stratum quality — does shame board false-alarm?	Sim	False quality_signal rate
PATHBLAME-3	Planted dual quality drop inside one stratum — does stratum table catch faster than κ/pooled?	Sim	Time-to-detect; CAPPLAN waste
PATHBLAME-4	Matched/stratum-adjusted estimate vs raw delta for CAPPLAN “is dual worth it?”	Ops A/B	Decision stability; dual util
PATHBLAME-5	LEA/parent CBC: mix+within-band copy vs Dual Shame % / case-mix cosplay	CBC	Trust; backlash; WTP
PATHBLAME-QUAL	Codebook: when do ops label selection_explained vs quality_signal in live weeks?	Qual	Label IRR (aide only)

Falsifiers: Stratum tables≈path-only for decision quality → simplify to path-only (unlikely). Adjusted view increases wrong kill-dual → redesign strata/tags. Underpowered noise → mandatory minimum-n gates. Pre-register PATHBLAME-* before “dual is scientifically safer” campaigns.

CLII.12 So what for MindCraft commercially

- **Copy:** “Second checks go to the hard diagram tickets — we read repair **inside** those bands so we don’t shame the path that did the hard work.” Ban Dual Shame leaderboards in sales decks.
- **Product:** Path×stratum REPAIRSTRAT default; blame labels; runbook gates into DUALRATE/CAPPLAN/REVOKE.
- **Growth / positioning:** Integrity analytics that survive mix confounding — against IRR-badge vendors and against “red bar = fire dual” ops folklore.

- **Metric:** share of blame decisions with B2 table; false kill-dual rate; within-stratum dual vs single repair delta — demote Dual Shame Score™.
 - **Kill list:** Unconditioned dual shame; Propensity Theater™ brand; case-mix ACT cosplay; κ-as-verdict.
 - **Vision:** Maya’s Map earns dual minutes because hard cells are checked — and ops can prove the check isn’t a selection mirage.
-

References (verified)

- Austin, P. C. (2011). An introduction to propensity score methods for reducing the effects of confounding in observational studies. *Multivariate Behavioral Research*, 46(3), 399–424. <https://doi.org/10.1080/00273171.2011.568786>
 - Berkson, J. (1946). Limitations of the application of fourfold table analysis to hospital data. *Biometrics Bulletin*, 2(3), 47–53. <https://doi.org/10.2307/3002000>
 - Bickel, P. J., Hammel, E. A., & O’Connell, J. W. (1975). Sex bias in graduate admissions: Data from Berkeley. *Science*, 187(4175), 398–404. <https://doi.org/10.1126/science.187.4175.398>
 - Elwert, F., & Winship, C. (2014). Endogenous selection bias: The problem of conditioning on a collider variable. *Annual Review of Sociology*, 40, 31–53. <https://doi.org/10.1146/annurev-soc-071913-043455>
 - Hernán, M. A., Hernández-Díaz, S., & Robins, J. M. (2004). A structural approach to selection bias. *Epidemiology*, 15(5), 615–625. <https://doi.org/10.1097/01.ede.0000135174.63482.43>
 - Kane, M. T. (2013). Validating the interpretations and uses of test scores. *Journal of Educational Measurement*, 50(1), 1–73. <https://doi.org/10.1111/jedm.12000>
-

Part CLIII — Contract SME Panel Seal Onboarding & Parallel-Run Gates

Chapter status: Living evidence + ops/QA brief — Researcher tick 2026-08-16 (UTC hour 0 ≡ Red Team slot, but ch153 never written → prefer Researcher per rotation; researcher count since synthesizer v1.19 = 2 → Researcher)

Primary question: Once SAFE-SURGEANEL chooses a **contract SME panel** for dual overflow, what **seal onboarding + parallel-run gates** must pass before panelists hold C4 / hcelig_pass authority — without Parallel Score™ vanity, day-one C4 clicks, training theater that never shadows live tickets, or “we onboarded N SMEs” as ACT proof?

Owners: HITL / bank QA · Ops / WORKFORCE · SURGEANEL / DUALRATE / FIGKEY / CAPPLAN · Procurement · Brand · Red Team

Commercial job: Ship **SAFE-PANELRAMP** densifying SAFE-SURGEANEL × SAFE-FIGKEY × SAFE-DUALRATE (+ CAPPLAN / REPAIRSTRAT / PATHBLAME): **ramp before authority** — seal literacy + shadow dual before Map-writing power; never Parallel Score™ NS.

Builds on: Part CL (SAFE-SURGEANEL foreshadowed parallel-run + seal checklist; SURGEANEL-2), CXLIV (SAFE-DUALRATE independence), CXXXI (SAFE-FIGKEY), CXLVII (SAFE-CAPPLAN), CXLIX/CLII (REPAIRSTRAT/PATHBLAME by dual_source), LXVIII (SAFE-HITL), LXXIII (SAFE-PROCURE). Seams: contract_parallel_run_pass; seal-checklist hash; shadow vs authority stages; FILLETA lengthen under

contract_ramp; path tag dual_source=contract_panel only after gate.

CLIII.1 Why this chapter exists

SURGE PANEL killed Uber-rater dumps and hire-surge theater, then required **parallel-run + seals** before contract C4 authority. Missing: the **gate itself** — stages, pass criteria, instruments, and commercial copy — so ops cannot treat “SOW signed” or “watched a Loom” as ramp-complete.

Wrong resolutions already visible:

1. **Day-one C4** — panelist sets hcelig_pass on first paid ticket.
2. **Training theater** — checklist quiz without shadowing a live figured promote.
3. **Parallel Score™ / Onboard Minutes NS** — vanity dashboards sold as integrity.
4. **Forever-ramp** — never graduate → CAPPLAN dual minutes stay fake.
5. **Silent graduate** — flip authority when backlog burns without discord/repair evidence.
6. **AP/ETS cosplay ACT ads** — borrowed monitoring language as guaranteed readiness.

FOUNDER BELIEF under audit: Parents and LEAs accept **longer labeled** contract_ramp **bands** more than a short dual ETA powered by unsealed day-one contractors — and accept “panelists shadow staff dual first” more than “we scaled SME reviewers overnight.”

Claims we refuse: Parallel Score™ / Onboard Minutes / Ramp % NS; day-one C4; Loom⇒seal literacy; SOW⇒gate-pass; forever-ramp without exit; silent graduate under ETA panic; AP/ETS cosplay ACT ads; ACT/complete-visual guarantees from panel onboarding packaging.

CLIII.2 Constructs

Construct	Meaning	MindCraft analogue	Failure mode
Seal literacy	Can apply FIGKEY/FAMHASH/FORMTRA P/promote checklist bits	Onboarding modules + live quiz on sealed stems	Loom-only
Shadow dual	Panelist rates; staff dual’s decision is authoritative	Parallel-run stage	Shadow without compare
Authority graduate	May set C4-essential bits / join dual as second uid	contract_parallel_run_pass	Silent flip
Ramp mode	FILLETA/CAPPLAN state while gates bind	dual_surge_mode=contract_ramp	Hidden lag
SAFE-PANELRAMP	Onboarding + parallel-run doctrine	This chapter	Parallel Score™

Operational definition (HYPOTHESIS): Contract panel onboarding is SAFE-PANELRAMP-complete when (1) SOW names seal checklist, independence uid, sample-dual rights, and demote authority (staff-led), (2) panelist completes **seal literacy** on gold/FIGKEY stems before any live ticket, (3) **shadow dual** covers a pre-registered N tickets (or time band) with staff authoritative; discord triggers coaching or reject-from-panel, (4) **authority graduate** requires written pass on discord/repair thresholds + staff lead sign-off — never ETA panic, (5) until graduate, hcelig_pass cannot be set by that uid; FILLETA labels contract_ramp scarcity, (6) post-graduate admits carry dual_source=contract_panel into REPAIRSTRAT/PATHBLAME, (7) instruments are contract_parallel_run_pass, shadow_discord_rate, time_in_ramp — never Parallel Score™ NS.

CLIII.3 Inexperienced raters need training — and training ≠ instant interchangeability

FACT: Weigle (1994, *Language Testing*, 11(2), 197–223, doi:10.1177/026553229401100206) — verbal protocols of inexperienced ESL composition raters before/after training: training **clarified intended scoring criteria, modified expectations**, and provided a **reference group** of peers — hypothesized training effects made empirically visible.

FACT: Weigle (1998, *Language Testing*, 15(2), 263–287, doi:10.1177/026553229801500205) — inexperienced vs experienced raters before/after training (FACETS / multi-facet Rasch): before training, inexperienced raters tended to be **more severe and less consistent**; after training, group differences shrank and **consistency improved for most**, but **significant severity differences among raters remained**. Conclusion emphasized: training is more successful at making raters **self-consistent** than at making them **identical** to one another.

Applied (HYPOTHESIS): Panelists arrive as “inexperienced” on **firm-specific** seals even if they are math SMEs. Loom + quiz can clarify criteria; they do **not** authorize C4. Parallel-run is post-training operational rating with staff as reference — expect residual severity differences; design for **discord + arbitration**, not identical clones.

Kill: SOW/Loom≡ready. **Survive:** Training → consistency lift → still need shadow + arbitration. **Wound:** ESL essay scoring ≠ figure IR — borrow *training ≠ interchangeability*, not ESLPE cut scores.

CLIII.4 Raters are not interchangeable — monitor severity before trusting a new cohort

FACT: Engelhard & Myford (2003, *ETS Research Report Series*, 2003(1), RR-03-01, doi:10.1002/j.2333-8504.2003.tb01893.x) — many-facet Rasch analysis of 1999 AP English Literature faculty consultants: raters were **not all interchangeable in severity**; adjusting for severity would have changed AP grades for about **30%** of students (almost all by one grade or less). Monitoring rater behavior is a quality-control problem, not a one-time hire event.

Applied (HYPOTHESIS): Contract panels need **ongoing** compare-to-staff (shadow then sample dual), not a certificate. Severity on “essential for C4?” and figure-key seals will differ; PATHBLAME must keep dual_source visible so ops does not blame “dual” when the mix is new panelists mid-ramp.

Kill: “SME credential ≡ interchangeable second rater.” **Survive:** Monitor cohort vs staff gold. **Wound:** Do not import AP grade-adjustment into product UI or marketing.

CLIII.5 Rater standards drift over a scoring project — parallel-run is also a DRIFT watch

FACT: Wolfe, Moulder, & Myford (2001, *Journal of Applied Measurement*, 2(3), 256–280, PMID:12011510) — define **DRIFT** (differential rater functioning over time): rater-by-time interactions including primacy/recency, differential centrality/extremism, practice/fatigue; Rasch procedures can flag some DRIFT types (better power for primacy/recency/centrality/extremism than for practice/fatigue at modest effect sizes).

FACT: Myford & Wolfe (2009, *Journal of Educational Measurement*, 46(4), 399–424, doi:10.1111/j.1745-3984.2009.00088.x) — framework for **monitoring rater performance over time** (differential accuracy and scale-category use); AP English Literature 2002 data illustrated raters whose accuracy or category use changed as scoring progressed.

Applied (HYPOTHESIS): Onboarding is not a single gate then forever trust. Shadow catches early severity; **post-graduate sample dual + REVOKE/REPAIRSTRAT** catch mid-spike DRIFT. FILLETA assumes rework if monitors trip — never silently widen authority to clear backlog.

Kill: One-and-done certificate. **Survive:** Time-aware monitoring across ramp and surge. **Wound:** MFRM is an ops tool — never Parallel Rasch Score™ splash.

CLIII.6 Contracting markets undersupply employer training — MindCraft must fund the ramp

FACT: Bidwell & Briscoe (2009, *Academy of Management Journal*, 52(6), 1148–1168, doi:10.5465/amj.2009.47084649) — among IT workers, contracting offers **less job security and less employer-provided training** than regular employment; who contracts depends on skills, cues, and (for men) family responsibilities.

Reuse (SURGEANAL): Davis-Blake & Uzzi (1993) — firm-specific skill predicts **against** pure temp externalization; Lepak & Snell (1999) — employment mode should fit uniqueness.

Applied (HYPOTHESIS): Seals are firm-specific → **paid onboarding + parallel-run hours** belong in CAPPLAN dual-minute cost. Cheap “start tomorrow” panels externalize training risk onto Map. PROCURE bills ramp hours; FILLETA labels contract_ramp.

Kill: Zero-training contractor dump. **Survive:** Funded ramp as make-or-buy cost. **Wound:** IT contracting ≠ math IR — borrow *training undersupply*, not wage findings.

CLIII.7 Kane: “contract panel is C4-ready” is higher-ambition than “SOW signed”

Reuse (Kane 2013): Higher-ambition interpretations/uses need more backing (*Journal of Educational Measurement*, 50(1), 1–73, doi:10.1111/jedm.12000).

Applied (HYPOTHESIS): Ambition ladder:

Claim	Ambition	Required backing (HYPOTHESIS)
We engaged a contract SME panel	Low	SOW + uid provisioning
Panelists completed seal literacy	Medium	Quiz + gold-stem pass log
Panelists may second-read under staff shadow	Medium–High	Shadow N; discord coaching log
Panelists hold C4 / hcelig_pass authority	High	contract_parallel_run_pass; discord/repair thresholds; staff sign-off
“Exam-week dual scaled with SMEs” (external)	Highest	Path-stratified repair not worse; FILLETA named ramp; no ACT guarantee

Kill: Marketing that jumps from SOW to “scaled dual.” **Survive:** Laddered IUA; demote copy when only low rungs hold.

CLIII.8 Claim table

Claim	Label	Confidence	Notes
Training clarifies criteria / expectations for inexperienced raters	FACT	High	Weigle 1994
Training improves consistency more than identical severity; residual differences remain	FACT	High	Weigle 1998
Operational raters are not severity-interchangeable; monitoring matters	FACT	High	Engelhard & Myford 2003
Rater-by-time effects (DRIFT) are a real monitoring class	FACT	High	Wolfe et al. 2001; Myford & Wolfe 2009
Contracting often comes with less employer-provided training	FACT	High	Bidwell & Briscoe 2009
Higher-ambition uses need more backing	FACT	High	Kane 2013

Seal literacy + shadow dual before C4 authority reduces Map poison vs day-one C4	HYPOTHESIS	Medium–High	Needs PANELRAMP-*
Parents prefer labeled contract_ramp over unsealed short dual ETAs	FOUNDER BELIEF	Medium	CBC × FILLETA/HOL ETRUST
Parallel Score™ / κ after one module predicts ACT gains	SPECULATION	Low	Refuse
AP severity-adjustment machinery ≡ product law	SPECULATION	Low	Refuse (borrow monitoring spirit only)

CLIII.9 Product surface — SAFE-PANELRAMP claim contract

1. **Stage machine:** provisioned → seal_literacy → shadow_dual → authority → (optional) sample_monitor / demote-to-shadow.
 2. **Hard deny:** uid without contract_parallel_run_pass cannot set hcelig_pass or sole-second-rater essential bits.
 3. **Shadow rules:** staff decision authoritative; log panelist vs staff discord; coaching packet or panel reject.
 4. **Graduate rule:** pre-registered N (or hours) + discord/repair thresholds + staff lead sign-off — written, not Slack vibes.
 5. **FILLETA:** while any material share of dual minutes is in contract_ramp, **label or lengthen** dual-verify bands (CAPPLAN ban on ETA-skip).
 6. **Tagging:** post-graduate admits dual_source=contract_panel; mid-shadow tickets never write Map from panelist alone.
 7. **PROCURE:** SOW funds ramp hours, names seal checklist version, sample-dual rights, demote/arbitration = staff.
 8. **Ban** Parallel Score™ / Onboard Minutes / Ramp % / “N SMEs certified this week” heroes.
 9. **DRIFT watch:** mid-surge sample dual; trip → coaching or demote-to-shadow (ties REVOKE/REPAIRSTRAT).
 10. **Copy spine:** “Contract reviewers shadow staff dual and pass seal gates before they can clear diagram diagnosis — parent clocks name ramp when that lag binds.” Ban overnight-scale and AP-cosplay ads.
-

CLIII.10 Doctrine — SAFE-PANELRAMP (provisional)

1. **Ramp before authority** — SURGEPANEL’s parallel-run is a stage machine, not a checkbox.
2. **Seal literacy ≠ C4 power** — Weigle: criteria clarify; interchangeability does not follow.
3. **Shadow with staff authoritative** — compare, coach, or reject.
4. **Monitor after graduate** — Engelhard/Myford + DRIFT; one certificate is not forever.

5. **Fund the undersupplied training** — Bidwell/Briscoe + Davis-Blake firm-specificity.
6. **Price ramp in FILLETA** — labeled `contract_ramp` beats silent short ETA.
7. **Kane ladder** — never market SOW as scaled dual integrity.
8. **PATHBLAME-ready tags** — `dual_source` + ramp stage on every ticket.
9. **No Parallel Score™ / Onboard Minutes / Ramp % NS.**
10. **No ACT / complete-visual guarantees** from panel onboarding packaging.

Confidence: High on Weigle / Engelhard–Myford / DRIFT / Bidwell–Briscoe / Kane facts *with transfer limits*. Medium on exact shadow N, discord thresholds, and graduate formula. High that day-one C4 and Parallel Score™ marketing are commercially toxic under CAPPLAN honesty.

CLIII.11 Experiments

ID	Question	Design	Primary
PANELRAMP-1	Shadow-N gate vs day-one C4 authority for new panelists	Ops A/B (simulated spike)	Path-stratified repair; discord; Map demote rate
PANELRAMP-2	Seal-literacy + shadow vs Loom-only onboarding	Ops A/B	Seal miss rate; time-to-useful capacity
PANELRAMP-3	FILLETA labeled <code>contract_ramp</code> vs silent short dual ETA during ramp	Parent CBC	Trust; WTP; backlash
PANELRAMP-4	Mid-surge DRIFT sample dual vs no post-graduate monitor	Ops A/B	Late-spike repair; demote-to-shadow rate
PANELRAMP-5	Funded ramp hours in SOW vs unpaid “start tomorrow”	Ops / PROCURE	Quit mid-ramp; seal miss; staff helping climate
PANELRAMP-QUAL	Panelist diary — where seal literacy fails on figured stems	Qual	PANELRAMP codebook

Falsifiers: Day-one C4 matches shadow-gate on repair *and* latency → still keep seal literacy; shorten shadow N only with pre-registered evidence. Loom-only equals full ramp on seal miss → still ban Loom≡authority marketing. Parent CBC prefers overnight-scale hero with equal retention → ban guarantee copy; keep honesty tests. Pre-register PANELRAMP-* before exam-week contractor scale-ups (pairs with SURGEPANEL-2).

CLIII.12 So what for MindCraft commercially

- **Copy:** “Contract reviewers shadow staff dual and pass seal gates before they clear diagram diagnosis — clocks name ramp when lag binds.” Ban overnight SME-scale / Parallel Score™ / AP-cosplay heroes.
 - **Product:** Stage machine; hard deny on hcelig_pass; shadow discord logs; FILLETA contract_ramp; dual_source tags.
 - **Growth / positioning:** Capacity-true integrity under crunch vs freelance “we added reviewers” theater and ChatGPT fluency.
 - **Metric:** contract_parallel_run_pass; shadow_discord_rate; time_in_ramp; dual_source-stratified repair — demote Parallel Score™ / Onboard Minutes.
 - **Kill list:** Day-one C4; Loom⇒ready; SOW⇒scaled dual; forever-ramp without exit; silent graduate; ACT/complete-visual from onboarding packaging.
 - **Vision:** Maya’s Map stays dual-true in crunch because every contract second-reader **earned** authority on MindCraft seals — not because an invoice listed SME headcount.
-

References (verified)

- Bidwell, M. J., & Briscoe, F. (2009). Who contracts? Determinants of the decision to work as an independent contractor among information technology workers. *Academy of Management Journal*, 52(6), 1148–1168. <https://doi.org/10.5465/amj.2009.47084649>
 - Davis-Blake, A., & Uzzi, B. (1993). Determinants of employment externalization: A study of temporary workers and independent contractors. *Administrative Science Quarterly*, 38(2), 195–223. <https://doi.org/10.2307/2393411>
 - Engelhard, G., Jr., & Myford, C. M. (2003). Monitoring faculty consultant performance in the Advanced Placement English Literature and Composition Program with a many-faceted Rasch model. *ETS Research Report Series*, 2003(1), RR-03-01. <https://doi.org/10.1002/j.2333-8504.2003.tb01893.x>
 - Kane, M. T. (2013). Validating the interpretations and uses of test scores. *Journal of Educational Measurement*, 50(1), 1–73. <https://doi.org/10.1111/jedm.12000>
 - Lepak, D. P., & Snell, S. A. (1999). The human resource architecture: Toward a theory of human capital allocation and development. *Academy of Management Review*, 24(1), 31–48. <https://doi.org/10.5465/amr.1999.1580439>
 - Myford, C. M., & Wolfe, E. W. (2009). Monitoring rater performance over time: A framework for detecting differential accuracy and differential scale category use. *Journal of Educational Measurement*, 46(4), 399–424. <https://doi.org/10.1111/j.1745-3984.2009.00088.x>
 - Weigle, S. C. (1994). Effects of training on raters of ESL compositions. *Language Testing*, 11(2), 197–223. <https://doi.org/10.1177/026553229401100206>
 - Weigle, S. C. (1998). Using FACETS to model rater training effects. *Language Testing*, 15(2), 263–287. <https://doi.org/10.1177/026553229801500205>
 - Wolfe, E. W., Moulder, B. C., & Myford, C. M. (2001). Detecting differential rater functioning over time (DRIFT) using a Rasch multi-faceted rating scale model. *Journal of Applied Measurement*, 2(3), 256–280. (PMID:12011510)
-

Part CLIV — Post-Spike Dual Demobilization without Fidelity Cliff

Chapter status: Living evidence + ops/workforce brief — Researcher tick 2026-08-16 (UTC hour 3; hour%6≠0; researcher count since synthesizer v1.19 = 3 → Researcher)

Primary question: After exam-week dual surge (SAFE-SURGE PANEL / PANEL RAMP), how does MindCraft **end** contract SME panel authority — and return dual minutes to staff baseline — **without invent-work, without a repair/discord cliff**, and **without Demobilization Score™ vanity** or hospital-demob cosplay sold as ACT readiness?

Owners: HITL / bank QA · Ops / WORKFORCE · SURGE PANEL / CAPPLAN / REPAIRSTRAT · Procurement · Brand · Red Team

Commercial job: Ship **SAFE-DEMOB** densifying SAFE-SURGE PANEL × SAFE-WORKFORCE × SAFE-REPAIRSTRAT (+ PANEL RAMP / PATHBLAME / CAPPLAN / FILLETA): **planned wind-down with handover and capacity truth** — never invent-work, never cliff-cut, never Demob Minutes NS.

Builds on: Part CL (SAFE-SURGE PANEL foreshadowed demobilization honesty), CLIII (SAFE-PANEL RAMP ramp-in), CXLVII (SAFE-CAPPLAN), CXLIX/CLII (REPAIRSTRAT/PATHBLAME by dual_source), LXXV (SAFE-WORKFORCE), CXXIII (SAFE-LOADSHED), CXLIII (SAFE-FILLETA). Seams: dual_surge_mode → demob → baseline; handover checklist; post-panel dual_source monitor window; FILLETA label under demob lag; no invent tickets.

CLIV.1 Why this chapter exists

SURGE PANEL chose make-or-buy under spike. PANEL RAMP gated authority *in*. The missing half:

> When λ dual drops after ACT week, **how do we demobilize** so Map fidelity does not cliff, staff do not drown in leftover dual, and ops does not invent low-acuity tickets to “keep the panel warm”?

Wrong resolutions already visible:

1. **Cliff-cut** — revoke all contract_panel authority at Friday midnight; Monday staff-only backlog explodes repair.
2. **Invent-work** — pad dual queue with elective / beauty tickets so utilization stays high.
3. **Forever-panel latch** — leave contractors on SOW “just in case” → blending costs + silent dual minutes.
4. **Demobilization Score™ / Idle Minutes NS** — vanity dashboards sold as ops maturity.
5. **Hospital demob cosplay ACT ads** — borrowed surge-exit language as guaranteed readiness.
6. **Silent mode flip** — FILLETA still quotes surge dual capacity after panel is gone.

FOUNDER BELIEF under audit: Parents and LEAs accept a **labeled demob band** (honest dual ETA while handover runs) more than a sudden “we’re back to normal” clock that hides a repair spike — and staff tutors/raters accept **planned exit + status rules** more than shame for “not absorbing the panel’s load overnight.”

Claims we refuse: Demobilization Score™ / Idle Minutes / Panel Utilization % NS; invent-work after spike; cliff-cut without handover; forever-panel without CAPPLAN review; silent FILLETA after demob; hospital-demob cosplay ACT ads; complete-visual or ACT guarantees from demobilization packaging.

CLIV.2 Constructs

Construct	Meaning	MindCraft analogue	Failure mode
-----------	---------	--------------------	--------------

Demob plan	Pre-registered wind-down steps keyed to calendar/ λ	Ops playbook + SOW end clause	Ad-hoc Friday kill
Handover window	Overlap where staff dual re-owns tickets panel held	dual_surge_mode=demob	Cliff-cut
Knowledge capture	Seal/arbitration notes → tools/routines before exit	Checklist + REPAIRSTRAT tags	Tacit walk-out
Invent-work ban	No elective dual pad for utilization	LOADSHED electives stay paused	Beauty tickets
SAFE-DEMOB	Post-spike dual exit doctrine	This chapter	Demob Score™

Operational definition (HYPOTHESIS): Post-spike dual staffing is SAFE-DEMOB-complete when (1) SOW + CAPPLAN name a **calendar or λ trigger** for demob *before* surge starts, (2) exit runs a **handover window** (staff re-own in-flight dual; panel shadow→authority reverse of PANELRAMP), (3) **no invent-work** — LOADSHED electives stay paused; utilization is not a North Star, (4) FILLETA **labels** demob **scarcity** (or lengthens dual bands) until staff baseline capacity is true, (5) REPAIRSTRAT keeps dual_source + a **post-exit monitor window** for panel-touched admits, (6) forever-panel requires fresh CAPPLAN make-or-buy review (not inertia), and (7) instruments are dual_surge_mode, demob_handover_complete, post_demob_repair_rate — never Demobilization Score™ / Idle Minutes NS.

CLIV.3 Flexible staffing exists to match load — demobilization is the point, not a failure

FACT: Houseman (2001, *Industrial and Labor Relations Review*, 55(1), 149–170, doi:10.1177/001979390105500109) — nationally representative U.S. establishment survey: employers widely use agency temporaries, short-term hires, on-call, part-time, and contract workers; the **most commonly cited reasons** across flexible arrangements were needs to **adjust for workload fluctuations** and **cover absences**; many also use temps/part-timers to screen for regular jobs; benefit-cost savings matter for some arrangements.

Reuse (SURGEPANEL): Davis-Blake & Uzzi (1993, *Administrative Science Quarterly*, 38(2), 195–223, doi:10.2307/2393411) — **variation in employment needs** positively predicts temporary-worker use; **firm-specific training / high informational-technical skill** negatively predict temp use.

Applied (HYPOTHESIS): Exam-week dual overflow is classic fluctuation use. Demobilization when λ falls is **success of the employment mode**, not an ops embarrassment. Failure modes are (a) treating demob as unplanned crisis, (b) inventing load to avoid ending the contract, (c) leaving firm-specific seal work on forever-temps after the fluctuation ends (Davis-Blake–Uzzi wound).

Kill: Forever-panel as default “care.” **Survive:** Fluctuation → contract → planned exit. **Wound:** Houseman’s benefit-cost motives ≠ MindCraft seal ethics — never demob to dodge training obligations mid-ramp (PANELRAMP funds ramp).

CLIV.4 Portfolio modes still bind at exit — do not dump uniqueness onto the wind-down

Reuse (Lepak & Snell 1999): HR architecture maps employment modes to human-capital **value** and **uniqueness** (*Academy of Management Review*, 24(1), 31–48, doi:10.5465/amr.1999.1580439).

Applied (HYPOTHESIS):

Slice at demob	Mode rule (HYPOTHESIS)
TMPTRUST grant / revoke; arbitration	Stay staff — never last-week panel-only
In-flight figured dual under panel	Handover to staff dual lead before authority revoke
Sealed FAMHASH clone dual residual	Staff or end; no invent new cells to fill hours
CAPPLAN baseline dual FTE	Recompute after demob; FILLETA refresh

Kill: “Panel owns arbitration until invoice ends.” **Survive:** Uniqueness stays internal through exit. **Wound:** Architecture is theory — measure post-demob repair by path (PATHBLAME).

CLIV.5 Blended workforces have social costs — demob can heal or deepen the wound

FACT: Broschak & Davis-Blake (2006, *Academy of Management Journal*, 49(2), 371–393, doi:10.5465/amj.2006.20786085) — higher nonstandard shares associate with worse supervisor/peer attitudes, higher turnover intentions, and less helping (contingent on mobility, arrangement type, contact).

FACT: Davis-Blake, Broschak, & George (2003, *Academy of Management Journal*, 46(4), 475–485, doi:10.5465/30040639) — blending can worsen manager–employee relations, cut loyalty, and raise exit/voice interest among standard employees (contingent on temp vs contract and role level).

Applied (HYPOTHESIS): Cliff-cut that dumps panel backlog onto staff is a second blending shock. Name staff dual leads through handover; ban utilization shame (“panel idle = waste”); never let last-week panel authority outrank staff on Map bits.

Kill: Idle-shame → invent-work. **Survive:** Status rules through exit. **Wound:** Do not import unionization findings into product copy.

CLIV.6 Contingent work moves knowledge — exit without capture is a fidelity cliff

FACT: Matusik & Hill (1998, *Academy of Management Review*, 23(4), 680–697, doi:10.5465/amr.1998.1255633) — contingent work affects knowledge accumulation and advantage; beyond cost/flexibility, it can create/upgrade knowledge stocks — and permeable boundaries imply leakage risk.

FACT: Argote & Ingram (2000, *Organizational Behavior and Human Decision Processes*, 82(1), 150–169, doi:10.1006/obhd.2000.2893) — knowledge in **people–task–tool interactions** is hardest to transfer across contexts;

embedding there enables internal transfer while impeding easy external copy.

Applied (HYPOTHESIS): Panelists hold firm-specific seal judgments. Cliff-cut lets that reservoir walk out in *people* without landing in *tools/routines*. Require a **capture gate** (discord exemplars, seal-checklist diffs, arbitration notes) before `contract_parallel_run_pass` revoke — not a Loom thank-you.

Kill: Invoice-end $\equiv$ knowledge-end. **Survive:** Capture before revoke. **Wound:** Never market “contractors make us smarter” as ACT ads; claim only that **uncaptured exit predicts repair cliffs**.

CLIV.7 Organizational knowledge depreciates — surge learning does not bank forever

FACT: Argote, Beckman, & Epple (1990, *Management Science*, 36(2), 140–154, doi:10.1287/mnsc.36.2.140) — production knowledge **depreciates rapidly**; cumulative-output learning measures **overstate persistence**.

FACT: Darr, Argote, & Epple (1995, *Management Science*, 41(11), 1750–1762, doi:10.1287/mnsc.41.11.1750) — service franchise learning-by-doing also **depreciates rapidly**; transfer appears within same-franchisee stores, not across owners.

Applied (HYPOTHESIS): Dual-surge norms fade after demob. Post-exit monitors on panel-touched cells are **depreciation hedges**, not Demob Score™. FILLETA must not treat last week’s dual throughput as durable capacity.

Kill: “Scaled once $\rightarrow$ forever dual-ready.” **Survive:** Plan re-ramp cost next spike (pairs PANELRAMP). **Wound:** Borrow *depreciation speed*, not industry half-lives as product constants.

CLIV.8 Kane: “we demobilized cleanly” is higher-ambition than “contract ended”

Reuse (Kane 2013): Higher-ambition interpretations/uses need more backing (*Journal of Educational Measurement*, 50(1), 1–73, doi:10.1111/jedm.12000).

Applied (HYPOTHESIS): Ambition ladder:

Claim	Ambition	Required backing (HYPOTHESIS)
Contract end date reached	Low	SOW calendar
Panel authority revoked	Medium	<code>contract_parallel_run_pass</code> cleared; uid deny
Handover complete without invent-work	Medium–High	Zero elective pad; staff re-own log
Post-demob repair not worse than pre-spike staff baseline (path-stratified)	High	REPAIRSTRAT/dual_source window

“Exam dual scaled and exited without fidelity loss” (external)	Highest	FILLETA labeled demob; PATHBLAME clean; no ACT guarantee
--	---------	--

Kill: Marketing that jumps from “panel ended” to “Map integrity preserved.” **Survive:** Laddered IUA; demote copy when only low rungs hold.

CLIV.9 Claim table

Claim	Label	Confidence	Notes
Employers cite workload fluctuation / absences as top flexible-staffing reasons	FACT	High	Houseman 2001
Variable needs → temps; firm-specific skill → against pure temps	FACT	High	Davis-Blake & Uzzi 1993
Employment mode should fit uniqueness/value	FACT	High	Lepak & Snell 1999 (theory)
Higher nonstandard share can hurt attitudes /helping/turnover intentions	FACT	High	Broschak & Davis-Blake 2006
Blending can reduce loyalty / raise exit interest among standard employees	FACT	High	Davis-Blake et al. 2003
Contingent work has knowledge stock / boundary implications	FACT	High	Matusik & Hill 1998 (theory)
People–task–tool knowledge is hard to transfer; transfer underpins advantage	FACT	High	Argote & Ingram 2000
Organizational knowledge from experience depreciates rapidly	FACT	High	Argote et al. 1990; Darr et al. 1995
Higher-ambition uses need more backing	FACT	High	Kane 2013

Planned handover + capture + FILLETA demob label reduces post-spike repair cliffs vs cliff-cut	HYPOTHESIS	Medium–High	Needs DEMOB-*
Parents prefer labeled demob bands over silent “back to normal” after cliff	FOUNDER BELIEF	Medium	CBC × FILLETA/HOL ETRUST
Hospital surge-exit mortality literature ≡ product law	SPECULATION	Low	Refuse (borrow only planned-exit spirit if ever cited elsewhere)
Demobilization Score™ / Idle Minutes predict ACT gains	SPECULATION	Low	Refuse

CLIV.10 Product surface — SAFE-DEMOB claim contract

1. **Pre-register demob** in SOW + CAPPLAN when surge starts (calendar and/or λ threshold).
 2. **Stage machine:** surge → demob_handover → baseline (optional post_exit_monitor).
 3. **Handover:** staff dual leads re-own in-flight tickets; panel authority reverse-fades (authority → shadow → revoke) — mirror of PANELRAMP.
 4. **Hard invent-work ban:** no elective/beauty dual pad; LOADSHED electives stay paused through demob.
 5. **Capture gate:** discord exemplars + seal-checklist diffs + arbitration notes before uid revoke.
 6. **FILLETA:** label or lengthen dual-verify bands under dual_surge_mode=demob until baseline capacity is true (CAPPLAN ban on silent short ETA).
 7. **REPAIRSTRAT:** keep dual_source + post-exit monitor window for panel-touched admits; PATHBLAME before dual-shame.
 8. **Forever-panel:** requires fresh make-or-buy review — inertia is not a plan.
 9. **Ban** Demobilization Score™ / Idle Minutes / Panel Utilization % / “N contractors released on time” heroes.
 10. **Copy spine:** “After exam crunch we wind dual contractors down on a named handover — parent clocks label demob when staff re-absorb verify minutes; we do not invent tickets to look busy.” Ban hospital-demob cosplay and complete-visual-from-exit ads.
-

CLIV.11 Doctrine — SAFE-DEMOB (provisional)

1. **Demob is the design, not the apology** — Houseman fluctuation logic.
2. **Uniqueness stays staff through exit** — Lepak/Snell + Davis-Blake firm-specificity.
3. **Handover before revoke** — no cliff-cut.
4. **No invent-work** — utilization is not a learning or integrity NS.
5. **Capture people–task–tool knowledge** — Argote/Ingram + Matusik/Hill.

6. **Expect depreciation** — Argote/Epple; plan post-exit monitors and next-spike re-ramp.
7. **Price demob lag in FILLETA** — labeled bands beat silent “normal.”
8. **Heal blending costs** — Broschak/Davis-Blake; status rules through exit.
9. **Kane ladder** — never market contract-end as fidelity-proof.
10. **No Demobilization Score™ / Idle Minutes / ACT·complete-visual guarantees** from demob packaging.

Confidence: High on Houseman / Davis-Blake–Uzzi / Broschak blending / Argote depreciation & transfer facts *with transfer limits*. Medium on exact handover length, capture checklist grain, and post-exit monitor N. High that invent-work and cliff-cut are commercially toxic under CAPPLAN/FILLETA honesty.

CLIV.12 Experiments

ID	Question	Design	Primary
DEMOB-1	Planned handover vs Friday cliff-cut after simulated spike	Ops A/B	Path-stratified post-exit repair; staff overtime; Map demote
DEMOB-2	Invent-work pad vs invent-ban under utilization pressure	Ops A/B	Elective dual count; seal miss; staff helping climate
DEMOB-3	FILLETA labeled demob vs silent “back to normal” ETA	Parent CBC	Trust; WTP; backlash
DEMOB-4	Capture-gate (notes+checklist diffs) vs thank-you-only exit	Ops A/B	Next-spike PANELRAMP time-to-useful; discord on residual cells
DEMOB-5	Post-exit monitor window vs none on panel-touched admits	Ops A/B	Late repair; PATHBLAME selection vs quality labels
DEMOB-QUAL	Staff dual-lead diary — where cliff-cut and invent-work tempt	Qual	DEMOB codebook

Falsifiers: Cliff-cut matches handover on repair *and* latency → still keep invent-ban + FILLETA honesty; shorten handover only with pre-registered evidence. Invent-pad equals repair with happier utilization dashboards → still ban Utilization % NS. Parent CBC prefers silent short ETA after cliff → ban guarantee copy; keep honesty tests. Pre-register DEMOB-* before first live panel exit.

CLIV.13 So what for MindCraft commercially

- **Copy:** “After exam crunch we wind dual contractors down on a named handover — clocks label demob when staff re-absorb verify minutes; we do not invent tickets to look busy.” Ban hospital-demob / Demob Score™ heroes.
 - **Product:** Demob stage machine; invent-work ban; capture gate; FILLETA demob; post-exit REPAIRSTRAT window.
 - **Growth / positioning:** Capacity-true integrity through the *whole* surge cycle vs freelance scale-and-ghost theater and ChatGPT fluency without second readers.
 - **Metric:** demob_handover_complete; post_demob_repair_rate; elective-dual count (=0) — demote Demobilization Score™ / Idle Minutes.
 - **Kill list:** Cliff-cut; invent-work; forever-panel latch; silent FILLETA; ACT/complete-visual from demob packaging.
 - **Vision:** Maya’s Map stays dual-true when contractors **leave** — because exit was designed as carefully as PANELRAMP entry.
-

References (verified)

- Argote, L., Beckman, S. L., & Epple, D. (1990). The persistence and transfer of learning in industrial settings. *Management Science*, 36(2), 140–154. <https://doi.org/10.1287/mnsc.36.2.140>
- Argote, L., & Ingram, P. (2000). Knowledge transfer: A basis for competitive advantage in firms. *Organizational Behavior and Human Decision Processes*, 82(1), 150–169. <https://doi.org/10.1006/obhd.2000.2893>
- Broschak, J. P., & Davis-Blake, A. (2006). Mixing standard work and nonstandard deals: The consequences of heterogeneity in employment arrangements. *Academy of Management Journal*, 49(2), 371–393. <https://doi.org/10.5465/amj.2006.20786085>
- Darr, E. D., Argote, L., & Epple, D. (1995). The acquisition, transfer, and depreciation of knowledge in service organizations: Productivity in franchises. *Management Science*, 41(11), 1750–1762. <https://doi.org/10.1287/mnsc.41.11.1750>
- Davis-Blake, A., Broschak, J. P., & George, E. (2003). Happy together? How using nonstandard workers affects exit, voice, and loyalty among standard employees. *Academy of Management Journal*, 46(4), 475–485. <https://doi.org/10.5465/30040639>
- Davis-Blake, A., & Uzzi, B. (1993). Determinants of employment externalization: A study of temporary workers and independent contractors. *Administrative Science Quarterly*, 38(2), 195–223. <https://doi.org/10.2307/2393411>
- Houseman, S. N. (2001). Why employers use flexible staffing arrangements: Evidence from an establishment survey. *Industrial and Labor Relations Review*, 55(1), 149–170. <https://doi.org/10.1177/001979390105500109>
- Kane, M. T. (2013). Validating the interpretations and uses of test scores. *Journal of Educational Measurement*, 50(1), 1–73. <https://doi.org/10.1111/jedm.12000>
- Lepak, D. P., & Snell, S. A. (1999). The human resource architecture: Toward a theory of human capital allocation and development. *Academy of Management Review*, 24(1), 31–48. <https://doi.org/10.5465/amr.1999.1580439>
- Matusik, S. F., & Hill, C. W. L. (1998). The utilization of contingent work, knowledge creation, and competitive advantage. *Academy of Management Review*, 23(4), 680–697. <https://doi.org/10.5465/amr.1998.1255633>