

MINDCRAFT

Brand & Business Update

2026-08-11 · For co-founders

Never work alone. Office hours from your room.
The operating system for student work — with Solver, the math engine and Katha's story world, as
its deepest-built vertical.

Part One: Brand Book v2.0 — positioning, voice, Maya, Jordan, vocabulary
Part Two: Business Model v1.0 — segments, revenue, unit economics, GTM, risks

Working draft — every unvalidated number in Part Two is tagged [ASSUMPTION].

joinmindcraft.com · confidential, internal use

PART ONE

Brand Book

Version 2.0 · governs every design, copy, product, and marketing decision at MindCraft

Product name on the web: **The Desk by MindCraft**. App chrome / home screen: **The Desk**.

Pedagogy translation of research: docs/canon/PEDAGOGY.md · Research OS: agent_work/research/ (Constitution v1.16).

Shareable PDF: regenerate with `python3 docs/canon/generate_brand_book.py` → BRAND_BOOK.pdf.

This document governs every design, copy, product, and marketing decision at MindCraft. When two options conflict, this book breaks the tie. When the book itself is unclear, the Brand Essence breaks the tie. Read it before you write a headline, pick a color, name a feature, or ship a screen. Agents: start at docs/canon/README.md.

v2.0 pivot note: MindCraft's positioning broadened from "ACT-math tutoring platform" to "the collaborative workspace for learning." This is a layering, not a replacement: everything this book says about Katha, Maya, Jordan, and math-specific voice is still true and still binding — it now describes the **Solver vertical** (the math/ACT-tutoring surface) rather than the entire company. Sections below are rewritten where the company-level frame changed (Essence, Manifesto, Mission, Vision, personas, personality, voice scope, vocabulary scope, anti-positioning, brand story, application examples) and left as-is where the Solver-specific system still applies directly (naming, visual identity, Katha mechanics, most of vocabulary). See BUSINESS_MODEL.md for the business case built on this positioning.

0. Naming (binding)

Surface	Say
Company / publisher	MindCraft
Marketing product	The Desk by MindCraft
App home screen + in-app chrome	The Desk (text only — no MindCraft / raccoon logo mark in chrome)
Narrative layer	Katha
Student sections	Notes · Solver · Map

Never call the student app “MindCraftNotes” in user-facing copy. Internal Xcode target names may lag; display name is **The Desk**.

1. Brand Essence

MindCraft is the collaborative workspace for learning: a place where students can work, create, and get live help from the right person without leaving their room.

MindCraft — Never work alone.

Office hours from your room.

That is the tagline and the explainer line, in that order, everywhere the company introduces itself. Lead with the entry point — office hours from anywhere — before revealing the larger platform underneath it. Nobody meets MindCraft

for the first time with "operating system for all students." They meet it stuck on something, at 11pm, wondering who could help.

Every decision either removes a reason a student is working alone, or gets out of the way of the person who's about to help them stop. The Desk is the place the work happens; MindCraft is the company that built it. Math is one thing MindCraft helps with, not the definition of what MindCraft is — inside the Solver vertical specifically, the older essence still governs word for word (§10).

Positioning architecture — the order in which the brand reveals itself, entry point to horizon:

Layer	Message
Entry point	Office hours from your room
Core product	A collaborative workspace with on-demand college mentors
Differentiation	Help is attached to the student's actual work — not a fixed tutoring curriculum
Broader platform	Meetings, notes, presentations, projects, content, and AI
Long-term vision	The operating system for student work
Emotional promise	Never work alone

Copy, decks, and product onboarding should generally move down this table in order. Skipping straight to "operating system" without first landing the entry point is the single most common way to lose someone in the first ten seconds.

2. Manifesto

You are alone with it. The blank document. The problem set. The deck due in the morning. The code that throws the same error for the third time. Somewhere there is a person who has done this exact thing before and could tell you in five minutes what is taking you three hours. You do not have that person's number.

So you keep going alone. You Google it. You ask an AI and it hands back an answer to a slightly different question than the one you asked. You wait for office hours that are two days and one deadline away. You nod along in a group project because admitting you're stuck costs more than staying stuck.

Somebody, somewhere back, decided you were not the kind of person who is good at this. Maybe a teacher said it with a red pen. Maybe a test said it with a number. Maybe you said it yourself, quietly, and nobody argued. Maya heard hers in seventh grade, mid-fractions, from a teacher who meant it kindly: "some people are just more verbal." She has been living inside that sentence for four years. Math is only the room she got stuck in first — these days it is also the essay due Monday and the presentation she has to give alone.

They described a moment. You heard a verdict.

Here is what actually happened: you were working alone on something hard, and nobody who understood it was in the room. That is not a broken brain. That is a broken workflow.

MindCraft rebuilds the workflow. Whatever you are stuck on — a proof, an essay, a pitch deck, a function that will not compile — you open the platform, show your actual work, and someone who has done this before is suddenly in the room with you. Not a ticket in a queue. Not a bot guessing at your intent. A person, live, looking at the same page you're looking at.

And then it happens. The pattern surfaces. The fog burns off. You feel the click, the specific relief of suddenly seeing it, and you know you could have found it alone eventually. You didn't have to.

You were never bad at the work. You were working alone.

Never again.

3. Mission

Turn every student working alone into a student who never has to be. Build the operating system where students go to think, create, collaborate, and get unstuck.

4. Vision

The Desk becomes the operating system for everything a student does outside class: the workspace where they take notes, join meetings, build presentations, manage projects, create content, and pull in a mentor or an AI the moment they get stuck — on any subject, any kind of work. `WORLD_VISION.md`'s own Horizon 1 already names this shape, in its own words: "The Desk as the student OS ... Binder, Intel, Connect, Ask, workflows, gap scan, Map, story-framed practice." That horizon was written before the company said the positioning out loud; this book is catching up to what the codebase already started building.

Inside that OS, the **Solver vertical** carries its own fuller vision, unchanged: a living world where the student is the character, every math problem is a mission, mastery visibly changes the map, and a mentor walks in already knowing the gap. That is no longer the whole company vision — it is what the vision looks like specifically for math. Full horizon map, including Katha's math-as-a-world roadmap: `WORLD_VISION.md`. Research thesis under audit (developed within Solver): identity transformation, not mathematics delivery (`agent_work/research/`).

5. The Student

Her name is Maya. She is sixteen, a junior, and she takes the ACT in April. Math is where she got stuck first. It is no longer the only room she gets stuck in.

She is not failing math. That is the part everyone gets wrong. She holds a B minus by memorizing procedures the night before tests and forgetting them by Friday. She writes strong history essays. Her English teacher reads her paragraphs out loud to the class. In seventh grade she switched schools mid-year, landed in a class three weeks deep into fractions, and never caught up. A teacher, meaning to comfort her, said "some people are just more verbal." Maya heard it as a diagnosis. She has been living inside it for four years.

Here is Maya at 2am, six days before the ACT. A practice test glows on her laptop. She has read question 14 five times. Each read, the symbols dissolve a little more. She switches tabs, watches a video about a topic she already understands, because understanding something, anything, quiets the chest-tightness for a minute. She is not lazy. She is rationing shame. Every problem she attempts and misses confirms the verdict, so she has learned to stop attempting. Giving up is not her weakness. Giving up is her armor.

Here is Maya on a different night, three weeks earlier, on the thing she supposedly *is* good at. A slides tab has been open since 4pm: six blank frames and a title she has rewritten four times. Her psych group project is due at 9am and her two partners went quiet in the group chat around dinner. She knows the content — she did the reading, she took the notes — she does not know how to turn it into eight minutes that won't put the class to sleep. She opens a template, deletes it, opens another. At 1am she gives up and writes plain bullet points she knows are flat, which she will read off tomorrow in a voice that says *this is fine* while nothing about it feels fine. Nobody sees this stuck moment. It never shows up on a report card. It just costs her a night, and whatever confidence she walked in with.

What she tells people: "I don't really care about math. I'm more of a words person." What she doesn't say out loud: being a words person doesn't make the blank page easy. It just makes the shame quieter, and better hidden.

What she never tells anyone: she watches the front-row kids finish early and wants, more than almost anything, to know what that feels like from the inside. Not the grade. The feeling. She wants to be the one somebody turns to and asks, "wait, how did you do that?" She has rehearsed explaining things she does not yet understand — the math, and lately, how to actually structure a deck.

What earns her trust: being taken seriously. Real stakes. Plain language about exactly where she's stuck, delivered without a verdict attached, whether the subject is a quadratic or a slide deck. A hard problem given to her like she can handle it, because she can.

What loses her in three seconds: cartoon mascots. "Math is fun!" Confetti for trivial things. Anything that smells remedial. Being told to try again with no new information, as if her attempt were a coin flip that might land differently this time.

Maya does not need motivation. She needs one honest click — on a concept, on a deck, on whatever is actually in front of her — and then she will supply her own.

6. The Tutor

His name is Jordan. He is twenty, a sophomore studying engineering, and he tutors math on MindCraft. He is one mentor among many the platform runs on — a design senior who walks a stranger through her first real slide deck at 10pm, a CS junior who debugs someone else's takehome between classes, a founder-track junior who workshops a pitch on a Tuesday afternoon. MindCraft mentors are not one job description; Jordan is the version whose story we tell in full, because it is one real, specific path through what every mentor on this platform goes through, not a definition that flattens the rest of them.

He remembers his own click with photographic precision: eleventh grade, a physics TA, a spring problem that suddenly became a picture instead of a formula. He has been chasing the chance to hand that moment to someone else ever since.

He tutors partly for money and mostly because explaining things is the best version of himself. His frustration is structural, not personal: he spends the first forty minutes of every session playing detective, poking at problems, guessing where the real gap hides, and by the time he finds it the hour is gone. Worse, his students have learned to

perform understanding. They nod at the right moments to end the session politely. He can feel it happening and cannot always stop it. Every mentor on MindCraft — whatever they actually help with — describes some version of this same tax: too much of the hour spent finding the real problem, not enough spent solving it.

What MindCraft gives him: the map, already drawn. For Jordan that's literal — the engine walks in knowing that Maya's quadratics problem is actually a slopes problem from three concepts back. For a mentor helping with a deck or a pitch, it's the equivalent: what the student already tried, where they got stuck, what they actually need from this session. Either way, the mentor's hour goes to the human part — the analogy, the patience, the eye contact when the click lands — not the detective work that used to eat it.

How the brand treats him, and every mentor like him: as a co-author, never as gig labor. He is the person who carries the flame across the table. Copy aimed at mentors runs peer-level and technical. It never says "learners." It never says "leverage outcomes." It says: here is exactly where your student is stuck, here is why, go do the thing you are great at.

7. Brand Personality

Five adjectives. Every piece of MindCraft output should score on at least three.

1. Cinematic.

MindCraft frames everything as a scene with stakes, never as a curriculum unit. If a screen could open a film, it belongs; if it could open a textbook chapter, it does not.

Sounds like: "The flood is three days out and the fields are still unmarked."

Never sounds like: "Welcome to Chapter 4: How to Format Your Notes."

2. Electric.

The brand chases the click, so its energy is sharp, sudden, and earned rather than constant and cheap. Excitement that fires at everything means nothing; MindCraft saves the voltage for the moment it matters.

Sounds like: "There it is."

Never sounds like: "Great job!!! Keep up the amazing work!"

3. Certain.

MindCraft speaks in declaratives because the engine actually knows things and the story actually goes somewhere. Hedging tells an already-doubtful student that even the software doubts them.

Sounds like: "You can learn this. The map says so."

Never sounds like: "You might want to consider possibly reviewing some earlier topics."

4. Human.

The brand knows exactly what 2am feels like and never pretends otherwise, but warmth here means respect, not babying. It talks to Maya like the capable person she is about to discover she is.

Sounds like: "That gap has been sitting there since seventh grade. It ends today."

Never sounds like: "Don't worry! This is hard for everyone!"

5. Unflinching.

MindCraft names gaps plainly, keeps the stakes real, and refuses both sugarcoating and shame. Honesty without a verdict is the whole trick: the gap is a fact about the map or the draft, never a fact about the person.

Sounds like: "Your thesis statement is the weak point in this draft right now. Here's the way through."

Never sounds like: "Oopsie! Let's give that another whirl!"

8. Voice and Tone

One voice, six dials. The voice never changes; the tone adjusts to what the student is feeling in that moment.

Universal voice rules:

- Second person. Present tense. Active voice.
- Verbs first. Short sentences. Declaratives.
- Stakes before instructions.
- Sentence case everywhere. No exclamation marks in UI chrome. No emoji in product copy, ever.
- The word "yet" is sacred. "You can't do this" never ships. "You can't do this yet" barely ships. "You haven't done this yet" ships.
- Never blame the student. Never baby the student. The distance between those two is where MindCraft lives.

Scope note — two registers, one voice. Everything below splits into **Solver-specific tone** (gap scan, mastery moments, "the click," story splash) and **general workspace-surface tone** (notes, meetings, mentor booking, project work). Solver keeps full cinematic voltage — that system is unchanged. The general workspace surfaces pull from the same five adjectives in §7 but dial Cinematic and Electric down and lean on Certain, Human, and Unflinching instead: calmer, more practical, closer to the elevator pitch's register in §14. A student booking a mentor for a deck at 11pm wants a straight answer fast, not a film title card. That's a different dial setting, not a different brand voice.

UI microcopy

Compressed, kinetic, six words or fewer. Buttons are actions in a story, not functions in software.

- Start button: **Run the mission**
- Diagnostic entry: **Start the gap scan**
- Continue: **Next move**
- Knowledge graph: **Open the map**
- Session end: **That's the run. See what changed.**

Story splash text

This is Katha's territory (see Section 10). Present tense, sensory, a human in a crisis whose shape is the math. Roughly 200 words, ending on the problem itself. The math is never announced; it surfaces.

Example opening: "Kai has been awake for thirty hours. The playtest is at noon, and the character still jumps wrong: too high off the ledge, too flat off the ramp, like the world's gravity changes its mind. Somewhere in the jump code, one relationship is lying."

Error states (technical failures)

Calm, honest, ownership taken by the system. The student never wonders if they broke something, because they didn't.

- Connection loss: **"The connection dropped. We saved your progress. Reconnecting now."**
- Load failure: **"That didn't load. It's on us. Give it a second and pull the map again."**
- Never: "Oops!", "Uh-oh!", "Something went wrong :(", or any error that performs cuteness.

Success states

Understated and electric. The click needs no cheerleading; it needs acknowledgment. The bigger the moment, the fewer the words.

- Correct in practice: **"There it is."**
- Concept mastered: **"Mastered. The map just changed."**
- Gap closed: **"That gap held you for two years. It's gone."**
- Never: confetti language, "Awesome!", "You're a rockstar," streak-guilt framing.

Marketing headlines

These are **Solver-specific** — the verdict-reversal register for the math/ACT surface. For the company-level tagline and elevator pitch, see §1 and §14. Big claims stated flatly, aimed at the verdict Maya carries. Every headline should make the kid who gave up stop scrolling. Product line when naming the app: **The Desk by MindCraft**.

- **You were never bad at math.**
- **Feel what it's like to be good.**
- **Math didn't lose you. The story did.**
- **The kids in the front row aren't smarter. They just clicked earlier.**
- **The Desk by MindCraft — learning that learns you.**

Claims we will not make (research-aligned)

Defendable claims only — see Pedagogy Canon + Constitution kills:

- No ACT / score **point guarantees**.
- No streak / XP / leaderboard as proof of mastery.
- No “AI replaces tutors” or “tutoring is free.”
- No item-count or “complete ACT bank” hero without coverage honesty.
- No Identity / Belief / Anxiety Score™ theater in parent or student copy.

Email subject lines

Sentence case, specific, story-forward. No fake urgency, no "Don't miss out," no discount energy.

- **Your map changed overnight**
- **One gap left in linear equations**
- **Kai's game still crashes. The function is waiting.**
- **Six days to the ACT. Here's your actual weak point.**

Tone dials by context

Context	What the student feels	Tone dial
Gap scan	Exposed, braced for judgment	Steady, factual, zero verdicts
Story splash	Curious, guard lowered	Full cinematic, immersive
Mid-practice miss	Old shame flickering	Direct, forward-pointing, brief
Mastery moment	The click	Electric, spare, let it land
Marketing	Skeptical, burned before	Certain, unflinching, bold
Tutor-facing	Professional, invested	Peer-level, precise, warm
Notes / meetings / project work	Focused, task-forward, not in crisis	Calm, plain, practical — Certain and Human, low Cinematic
Mentor booking	Deciding who to trust with this	Certain, warm, no hype, no cinematic buildup
Company-level marketing	Skeptical, evaluating a new tool	Certain, direct, leads with the entry point (§1), not the whole platform at once

9. Visual Identity

Two stages, one brand — and the split already scales past math: Deep Field carries any story or marketing moment, cream Desk OS carries any workspace surface, whether the student is inside a Solver mission or writing meeting notes.

1. **Story / marketing / Katha** — film title sequence energy: Deep Field, enormous type, high contrast. Never looks like a school LMS.
2. **The Desk OS** (Jesse's Kitchen — Solver's onboarding world — plus the general Work and Create surfaces) — cream paper workspace, soft ink (#143a2e), lime accents. Feels like a real desk a student owns, not a dark theater. Do not force Deep Field onto The Desk surfaces.

Color system

Color	Hex	Name	Role
Near-black	#080e14	Deep Field	Marketing, story, Katha, cinematic moments.
Desk paper	#f7f5f0 / #f8faf7	Field Desk cream	The Desk OS canvas (Work / Create / kitchen chrome).
Ink green	#143a2e	Desk ink	Primary text on cream Desk surfaces.
White / chalk	#f5f5f5	Chalk	Primary text on Deep Field.

Lime	#c4f547	The Click	Mastery, primary CTAs, earned signal.
Red	#c1121f	Stakes	Narrative tension, gap severity, Katha's flame — never student shame.
Navy	#1d3a8a	Depth	Graph edges / secondary structure on dark stages.

Usage rules:

- **Stage-appropriate canvas.** Deep Field for story and marketing. Cream for The Desk OS. Do not mix randomly on one screen.
- **The Click is earned.** Lime means something just happened or is about to. Never two competing lime CTAs.
- **Stakes belongs to the story, never to the student.** Red never marks a miss as “you failed.” A red X on student work is banned.
- **Desk chrome stays calm.** Top-left product mark is the words **The Desk** — no raccoon / Minecraft logo in app chrome (logo stays for marketing site / company mark).

Typography direction

No specific typefaces mandated. The feelings are mandated.

- **Display (Minecraft voice):** a confident geometric grotesque with tight tracking and real presence at poster scale. It should feel like a title card, an athletic brand, a launch countdown. Set it huge: the type scale jumps hard, display sizes at four times body or more, so every screen has one loud statement and quiet everything else.
- **UI text:** a clean humanist sans built for legibility at small sizes, with tabular numerals wherever the engine reports data (mastery, timing, counts). The engine's numbers should feel machined.
- **Story text (Katha voice):** an editorial serif with genuine literary character, the kind of face you'd find in a novel, not a textbook. When the serif appears, the reader knows without being told: the storyteller is speaking now.
- **Never:** rounded "friendly" fonts, handwriting fonts, anything that signals children's education, chalkboard textures, comic-adjacent anything.

Imagery direction

Belongs:

- Cinematic scenes from the stories: the game studio at 2am lit by one monitor, the Nile at flood under a red dawn, rendered like film stills with a single dominant light source.
- The knowledge graph as an object of beauty: constellation-like node fields on Deep Field, navy edges, lime nodes where mastery lives, red where gaps burn.
- Bold geometric abstraction with narrative weight: a parabola as the arc of a jump, a slope as a rising floodline. Math as scenery, not decoration.
- Darkness with a point of light. This is the brand's core image: the click, visualized.

Never belongs:

- Stock photos of any kind. Especially smiling students at laptops.

- Chat bubbles, robot mascots, anthropomorphized anything.
- Clip-art math: floating pi symbols, raining numbers, cartoon calculators.
- Trophies, gold stars, badge-cabinet iconography.
- Anything a school would laminate.

Motion principles

1. **Motion means state change.** Things move because the world changed: a node lights, an edge strengthens, a path redraws. Nothing animates to seem lively.
 2. **Weight over bounce.** Elements move with mass and settle with intent. Nothing springs, jiggles, or celebrates cutely.
 3. **The click gets the budget.** Mastery moments earn the largest, slowest, most considered animation in the product: the node ignites, the map ripples, up to 800ms. Routine UI stays fast and invisible, 150 to 250ms.
 4. **Nothing begs.** No idle loops, no wiggling buttons, no notification-dot anxiety mechanics. The world waits with dignity.
 5. **Dark stays dark.** Transitions move through Deep Field, never through white flashes.
-

10. The Katha Sub-brand

Katha (Sanskrit: story) is the narrative layer of MindCraft. It is a sub-brand with its own voice, its own typography, and its own mark, living entirely inside MindCraft.

Scope. Katha is the storytelling voice of the **Solver vertical** specifically — the math-as-a-world system described in `WORLD_VISION.md`. It is not the company's only storytelling voice. MindCraft's broader workspace surfaces (notes, meetings, mentor booking, project work) tell their own story in the calmer register defined in §8; they do not borrow Katha's flame, serif, or red. If a mentor-booking screen or a meetings feature wants narrative flourish, it does not become a Katha frame by default — Katha's rules below (never mentions the interface, ends at the threshold, etc.) are binding specifically for Solver missions and story splashes.

The division of labor

MindCraft is the system. Katha is the fire.

MindCraft measures, maps, diagnoses, and builds: the engine, the knowledge graph, the gap scan, the mastery model. It speaks in the geometric sans, signals in lime and navy, and deals in facts about the map.

Katha tells the story: every splash, every scene, every mission frame. It speaks in the editorial serif, burns in red and chalk, and deals in humans in trouble. Kai's broken jump. The scribe watching the flood erase the field lines.

The handoff is a designed moment. The engine decides *what* Maya needs (that is MindCraft's job, deterministic and auditable) and then the flame appears, the serif takes over, and Katha gives her a *reason* (that is Katha's job, and the engine never writes a word of it). One brand does the math about the math. The other makes it matter.

How Katha sounds

- Present tense, sensory, specific. Always a person, always a crisis, always stakes.

- Katha never mentions the interface. It does not say "concept," "practice," "mastery," "level," "question," or "click here." The moment story text acknowledges the software, the spell breaks.
- Katha never explains the math. It arranges the crisis so the math is the only door out, then steps aside.
- Katha ends every story at the threshold of the problem, never past it. The student, not the narrator, opens the door.

How Katha looks

- Editorial serif on Deep Field. Generous margins and negative space, like the page of a novel printed on darkness.
- Red and chalk are Katha's colors. Lime almost never appears in Katha's frame: lime belongs to the engine, and the click happens after the story, not inside it.
- Illustrated or rendered scenes composed like film stills: one light source, one figure, one problem in the air.
- Every story signs itself: a small flame mark and the words **A Katha story**.

The flame

The flame mark carries three meanings, in order:

1. **Stories began around fire.** Katha is the oldest technology for making knowledge survivable: wrap it in a narrative and pass it around the flame. The mark claims that lineage.
2. **The click is a strike of light in the dark.** The whole visual system is darkness waiting for ignition. The flame is the promise that ignition comes.
3. **Fire is carried.** Tutor to student, story to student, student to the next student who gave up. Nobody keeps the flame. Everybody hands it on.

Flame rules: the flame appears whenever Katha speaks and nowhere else. Red on Deep Field. Never a mascot, never given a face, never animated in idle loops. It may flare once, briefly, at a mastery moment: the story's stakes resolving into the engine's light.

11. Vocabulary

Words build the world. These pairs are binding for all product copy, marketing, and support.

Scope note. Some rows below are **Solver-specific** — they name pieces of the math/ACT system precisely and should not be borrowed loosely into unrelated surfaces: **mission, gap, the click, gap scan, master, tutor** (as the human title inside a Solver session). Others are **general-workspace** terms that apply company-wide: **you/player, story-driven, earn/unlock**, and the new rows below (mentor, office hours, workspace, bring what you're working on, get unstuck). When in doubt: if the copy is happening inside a Solver mission or story splash, reach for the Solver-specific word; everywhere else, reach for the general one.

We say	We never say	Why
challenge	quiz	A quiz judges you at school. A challenge invites you into a story. "Quiz" is banned outright.

weakness (with care)	deficiency, deficit	"Weakness" is honest and useful when it describes the map, spoken privately, attached to a concept. It never describes the person.
not yet / not this one	wrong	"Wrong" is a verdict on the person. "Not yet" is a fact about time.
not yet / not this one	incorrect	"Incorrect" is "wrong" wearing a lab coat. Same verdict, colder delivery.
look at [the specific thing], then take the next one	try again	"Try again" asks the student to repeat a failure with no new information. MindCraft always points at what to see differently first.
mission	assignment, homework	Assignments come from authority. Missions come from stakes.
story	content	"Content" is what gets poured into feeds. Nobody ever changed their life because of content.
the click	aha moment, lightbulb moment	The click is the brand's own word for its own promise. Greeting-card synonyms dilute it.
gap	falling behind	A gap is a location on a map: findable, closeable, gone when closed. "Behind" is a race the student already believes they lost.
practice	drill	Drilling is what happens to teeth and soldiers. Practice is what musicians and athletes do on purpose.
master	complete, finish	You complete a form. You master a skill. Concepts are never "done"; they are owned.
you / player	user	Users consume software. Players inhabit worlds. In student-facing copy, just say "you."
world / map	app, platform	Student-facing copy never names the software category. The product is a place, not a program.

gap scan	diagnostic test, placement test	Both banned terms contain "test," and "test" is the exact machine that issued Maya's verdict. A scan finds things; a test judges people.
yet	can't (unqualified)	"Yet" is the single most important word in the vocabulary. Every statement of inability carries it or dies in review.
sharp, locked in	smart, dumb	Intelligence-as-identity language is banned in both directions. Praise the seeing, never the ceiling.
quick / short	easy	Calling a problem "easy" sets a trap: miss it and the verdict returns doubled. Describe length, never difficulty-as-judgment.
tutor (Solver-specific copy)	instructor, coach	Inside a Solver mission or math-specific copy, tutor is still the one name for the human across the table.
story-driven	gamified	"Gamified" means points bolted onto boredom. The story is load-bearing. Banned even in internal decks.
earn / unlock	reward	Rewards are pellets from a dispenser. Earning and unlocking are things that happen inside a world because you changed it.
mentor (general workspace copy)	tutor (outside Solver-specific copy)	Outside Solver, "mentor" is the umbrella term for who's on the other side of a session — a writer, designer, coder, or founder, not only a math tutor. Reserve "tutor" for the Solver-specific relationship.
office hours	tutoring session (in company-level marketing)	"Tutoring session" narrows the whole platform back to one subject. "Office hours" is the entry-point promise (§1) — it describes the shape of the help, not its subject.

bring what you're working on	submit your assignment	"Submit" implies grading and one right answer waiting on the other side. MindCraft mentors show up on the student's actual, unfinished work, not a completed hand-in.
workspace	app, tool, platform (in company-level copy)	Same instinct as "world / map" in Solver copy: naming the software category flattens the thing into generic SaaS. "Workspace" names the place, not the product category.
get unstuck	get help, get answers	"Get answers" sounds like a vending machine. "Get unstuck" names the actual outcome: momentum resumes. That is a different promise than being handed a solution.

12. Anti-positioning

What MindCraft is explicitly not, and why each comparison fails on contact.

Not Khan Academy. A library assumes you walked in wanting a book; MindCraft exists for the student who stopped walking in. Khan organizes the world's math superbly for the already-motivated, while MindCraft manufactures the motivation itself, which is the actual missing ingredient for the kid who gave up.

Not Duolingo. Duolingo made the mechanic the point: streaks, guilt, an owl trained to make you open an app, with the content as filler between dopamine hits. MindCraft inverts it: the story is the point and math is the mechanism, because streak anxiety teaches you to launch software, and it has never once taught anyone to see a pattern.

Not a tutoring app. Tutoring marketplaces sell hours and pray that chemistry happens inside them; MindCraft sells the click and uses the engine to make the human hour count, because the tutor should spend forty minutes teaching, not forty minutes guessing where the gap is.

Not a homework helper. Homework helpers optimize for tonight: answer delivered, dependence deepened, gap untouched and compounding. MindCraft optimizes for the version of you that stops needing help, and any feature that hands out answers without building the seeing has failed the brand test by definition.

Not a productivity suite. Notion and Google Workspace hand you an infinite set of blank tools — a doc, a deck, a board — and zero help actually using them; they optimize for organizing work you already know how to do alone. MindCraft inverts it: every tool in the workspace sits one tap from a live person who can actually help, because a better blank page has never once gotten anyone unstuck, and a folder structure has never taught anyone anything.

13. Brand Story

MindCraft started at a tutoring table, watching the same scene repeat. A student flinches at a worksheet. The same student, twenty minutes later, leans all the way in because the problem became a person: a game developer whose character jumps at the wrong height, a scribe watching the flood erase every boundary in the valley. Same math. Different story. Completely different kid.

The founders kept pulling on that thread and found two separable problems. Finding the gap is a systems problem: it takes a knowledge graph, real evidence, and a deterministic engine that never guesses. Caring about the gap is a story problem: it takes stakes, a character, a world that answers back. Every product on the market solved one and ignored the other.

So MindCraft builds both, welded together: an engine that knows exactly where math lost you, and a storyteller that gives you a reason to go back for it.

That was the origin of Solver — and for a while, it was the whole company: a platform for booking college tutors, one subject, one session at a time. But the tutoring table kept teaching the founders things they hadn't asked to learn. A tutor would open a shared doc to sketch a proof, and the student would ask to keep using that same doc for an English paper due the next morning. A tutor would stay online after the booked hour ended because the student was now stuck on something else entirely — a presentation, a coding assignment, a scholarship essay. Students weren't booking a fixed subject with a fixed person. They were showing up with whatever was actually in front of them, hoping the person on the other end could help with that too.

So the platform grew the way real use grows a product: not by a founder's roadmap, but by watching what people did with it anyway. What started as a tutor-connection tool for one subject became a place students opened for office hours, for meetings, for notes, for building a deck at midnight, for pulling in a mentor who happened to know exactly the thing they were stuck on that day — math or not. MindCraft stopped being the company that booked tutoring hours and became the company that builds the room where the work happens, with the right person a tap away whenever it stalls.

Solver — Maya's engine, Katha's stories, the click — is what that room looks like for math specifically, and it remains the deepest, most built-out corner of the workspace. The rest of the workspace is the same bet, applied everywhere else a student gets stuck alone.

14. Application Examples

Ten real contexts, final copy quality. This is what the brand sounds like in the wild. The first six below are **Solver-specific** — math/ACT copy in the engine's and Katha's voice, unchanged from v1.1. The four after that are **company-level** — the workspace's own voice: calmer, broader, and new in v2.0.

Push notification (re-engagement, student mid-story on functions):

Kai's game still crashes at the jump. The function is waiting.

Story splash opening line (fractions, the Nile scribe):

The flood took the field markers in the night, and by morning forty farmers stand at the water's edge, each one certain his land was bigger than this.

Wrong-answer feedback message (practice, linear equations):

Not this one. The slope tells you how fast it changes, not where it starts. Take the next one with that in your pocket.

Dashboard headline (returning student, one gap from a milestone):

One gap stands between you and quadratics. It's the same one from seventh grade. Today's the day.

Marketing site hero headline:

You were never bad at math.

Subhead: Somebody told you that, and you built a life around it. The Desk by MindCraft finds the exact moment math lost you, then tells the story that brings you back.

App Store description, opening lines:

The Desk by MindCraft is where you find out you were never bad at math.

>

A learning engine maps exactly where math lost you: not the chapter, the concept, the precise gap. Then a story world gives you a reason to go back for it. Every concept is a scene. Every problem is a mission. And somewhere in your first week, it happens: the pattern surfaces, the fog burns off, and you feel the click the front-row kids have been feeling all along.

Company tagline (primary, everywhere the company introduces itself):

MindCraft — Never work alone.

Explainer line: Office hours from your room.

Elevator pitch (two-minute, spoken or written — the canonical company pitch):

College students spend hours every day working alone — writing papers, debugging code, studying for exams, creating presentations, or trying to turn an idea into something real. When they get stuck, their options are limited: wait for office hours, search through disconnected online content, ask AI and hope it understands, or schedule an expensive tutoring session days later.

>

MindCraft gives students another option: open the platform, show what they're working on, and instantly enter a collaborative workspace with a college mentor who can help them move forward.

>

This isn't traditional tutoring, where you're assigned one tutor for one subject and locked into recurring sessions. On MindCraft, students can work with different people depending on what they need — a mathematics student for calculus, a designer for a presentation, a founder for a startup idea, or a writer for an essay. They can return to mentors they trust or discover someone new.

>

But MindCraft has grown beyond mentor discovery. It is becoming the workspace where the work itself happens. Students can join virtual office hours, collaborate in live rooms, take notes, create presentations and content,

organize projects, meet with classmates, and use AI alongside real people — all without switching between five different platforms.

>

AI can generate an answer, but learning often requires conversation, judgment, encouragement, and another person who understands where you're stuck. MindCraft combines the scalability of AI with the intelligence and connection of a real academic community.

>

Our vision is to build the operating system for student work: one place to think, create, collaborate, and get unstuck.

>

MindCraft brings the people, tools, and workspace of a university directly into every student's room.

Elevator pitch, 30-second version (same voice, compressed — use where the two-minute version doesn't fit):

MindCraft is office hours from your room. Students bring anything they're working on — an essay, a math problem, a presentation, code, or even a startup idea — and collaborate live with a college mentor who can help them move forward. Unlike traditional tutoring, students aren't locked into one tutor or one subject. MindCraft combines on-demand human guidance, AI, meetings, notes, and creation tools in one collaborative workspace. We're building the operating system where students think, create, and get unstuck.

Company page copy (LinkedIn or equivalent company-profile surface):

Tagline: Office hours from your room. Think, create, collaborate, and get unstuck with AI and real college mentors.

>

About: Being a student means constantly encountering things you have never done before. Solving a difficult problem, writing a paper, creating a presentation, learning a new domain, or applying for your first internship. But the right support is rarely available when you actually need it. MindCraft brings office hours directly into every student's room. Students can bring whatever they are working on, collaborate with college mentors who have done it before, and use AI-powered tools to think, create, and move forward. They can work with different mentors for different needs or return to someone they trust. AI can generate an answer, but learning often requires conversation, judgment, encouragement, and another person who understands where you are stuck. MindCraft combines the scalability of AI with the intelligence and connection of a real academic community. Our vision is to build the operating system for student work. One place to think, create, collaborate, and get unstuck. MindCraft brings the people, tools, and workspace of a university directly into every student's room. As the students we support gain experience of their own, we hope they return as mentors for the students following in their footsteps.

15. Research handshake

Brand copy must stay inside what pedagogy can defend. When research kills a Score™ or packaging claim, Brand Book follows — we do not keep the pretty lie.

Topic	Look first
What to build / measure	docs/canon/PEDAGOGY.md
Evidence + Red Team	agent_work/research/MINDCRAFT_RESEARCH_CONSTITUTION_v1.md
Why / horizons	WORLD_VISION.md
Agent index	docs/canon/README.md

End of Brand Book v2.0. When in doubt, reread the essence: never work alone, office hours from your room, The Desk is the place — and inside Solver specifically, math is still the mechanism, the story is still the point, and everything there still serves the click.

PART TWO

Business Model

Version 1.0 · companion to the Brand Book · every unvalidated figure tagged [ASSUMPTION]

Companion to BRAND_BOOK.md v2.0. Validation plan referenced throughout: MindCraft_Viability_Research_Strategy.md.

How to read this document. This is a business-model draft, not a set of settled facts. Every number that is not directly sourced from an existing repo document is tagged [ASSUMPTION] inline. Tagged numbers are internally consistent with each other (if a take rate and a session price are both assumed, the downstream math ties out to both) but none of them have been validated against real pricing research, real mentor payout expectations, or real conversion data. Before this document is used in an investor deck, a pricing page, or any external commitment, every [ASSUMPTION] needs to be replaced with a real figure or an explicit "still unknown."

1. Positioning recap

MindCraft is the collaborative workspace for learning: students open the platform, show what they're working on, and get live help from a college mentor without leaving their room — entry point "office hours from your room," broadening into a full workspace for notes, meetings, presentations, projects, content, and AI, with the long-term vision of becoming the operating system for student work. Math tutoring (the Solver vertical — Katha, Maya, "the click") is the deepest-built vertical inside that workspace, not the definition of the company. Full positioning, voice, and the tagline system ("MindCraft — Never work alone. / Office hours from your room.") live in BRAND_BOOK.md v2.0; this document assumes that positioning and builds the business case on top of it.

2. Target segments

Who uses the product (not necessarily who pays):

- **Students** — the core user on both sides of the interaction: as a person seeking help (any subject, any kind of work) and, over time, as a mentor once they have experience of their own (the Brand Book's "as the students we support gain experience of their own, we hope they return as mentors" line — this is a real supply-growth mechanic, not just sentiment).
- **Mentors** — currently college students and recent grads; the supply side of the marketplace.
- **Parents** — mostly relevant for the Solver/high-school-ACT slice of the student base; pay for or co-decide on paid tiers, consume the reporting add-on (§3).
- **Institutions** (eventually) — college academic-support offices, high schools, possibly employers/bootcamps — a licensing buyer, not a near-term segment.

Who pays:

- Students/families directly (subscription + marketplace session fees).
- Parents specifically for the reporting/oversight add-on.
- Institutions, once a licensing motion exists (not near-term — see §6).

TAM / SAM / SOM — methodology, not just a number

The right way to size this is bottom-up from campuses, not top-down from a national headline number (a top-down "N million students" figure invites exactly the kind of unsupported scale claim Akshat has already ruled out — see the note

on the killed "1.6 billion" line below). Still, a rough top-down frame is useful as a ceiling check, so both are shown.

Top-down ceiling (methodology: national enrollment counts, not usage data):

- TAM — US students who regularly do "work that could stall and need live help": roughly all US undergraduates plus college-bound high schoolers, on the order of **24–27 million** people [ASSUMPTION – not sourced to a specific NCES/IPEDS pull; needs a real citation before use externally].
- SAM — the slice at institutions/communities where a peer-mentor culture already exists or is plausible (residential colleges, ACT/SAT-prep-active high schools), and where students have some ability to pay directly or through parents: roughly **8–12 million** [ASSUMPTION].
- SOM (3-year horizon) — reachable via the campus-grassroots motion in §6, not paid acquisition: **~50 campuses/communities × ~3,000 addressable students each ≈ 150,000** [ASSUMPTION, both the campus count and per-campus student count].

Bottom-up build (methodology: start from the actual GTM unit — one campus — and multiply):

1. One campus (e.g., Macalester, ~2,100 undergrads) → assume **40–60% of a campus is "addressable"** in year one (students who do enough independent academic/project work to plausibly want this) [ASSUMPTION] → ~900–1,250 addressable students per small liberal-arts campus.
2. Mentor supply per campus: assume **3–5% of addressable students** are willing to mentor for pay [ASSUMPTION], i.e., 30–60 mentors at Macalester's scale — enough to test whether "any subject, on demand" is a real supply story or a thin one (see §8 risk).
3. Multiply the per-campus addressable count by the number of campuses reached via the grassroots playbook in a given year (§6 lays out the sequencing; no specific campus count beyond year one is assumed here) to get SOM at any point in time.

This bottom-up build is the one to actually track — it is falsifiable per campus, which the top-down TAM is not.

Killed claim: the old "1.6 billion [students worldwide]" scale line does not appear in this document and should not appear anywhere else. Akshat explicitly killed it as a scale claim the company cannot defend.

3. Revenue streams

Grounded in two things already in the repo rather than invented from scratch: CODEX_BRIEF.md's existing pricing note ("pricing is invitation-based, not transactional... 'A few seats remain'... 'Apply for a seat'") and backup.md's tutor/parent monetization idea ("auto weekly reports, exact weak concepts, suggested sessions, and receipts of improvement — this sells the business better than 'AI tutor' alone").

Reconciling "invitation-based" with a real revenue model. "Invitation-based, not transactional" describes the *marketing funnel and scarcity framing* at the front door (controlled rollout, "apply for a seat," no self-serve checkout page) — it does not mean the backend is free. It's compatible with every stream below: access is gated by an application/invite, and once in, the streams below are how the gated relationship is monetized.

Stream	Mechanism	Rate [ASSUMPTION]
Mentor-session marketplace take rate	Platform fee on paid mentor sessions booked through MindCraft	20% platform / 80% mentor payout [ASSUMPTION]

Desk OS subscription — free tier	Workspace, notes, rate-limited AI tools, browsing the mentor marketplace	\$0
Desk OS subscription — paid individual tier ("Plus")	Unlimited AI tools, priority mentor matching, a monthly mentor-session credit, unlimited notes/presentation storage	\$12–15/month [ASSUMPTION]
Family/parent reporting add-on	The backup.md monetization loop: automated weekly reports, exact weak concepts (pulled from the same mastery graph that powers Solver), suggested sessions, receipts of improvement	\$10/month per student, stacks on top of a student plan [ASSUMPTION]
Institutional/school licensing (eventual, not near-term)	Per-seat annual license for a college academic-support office or a high school; bulk seats, admin dashboard, aggregate (non-identifying) insight	\$8–15/seat/year [ASSUMPTION]

The family/parent add-on is worth calling out specifically because it is the one revenue line already partially validated by product reality: the mastery graph, weakness detection, and evidence tracking described in `CLAUDE.md` already exist for Solver. Turning that into a parent-facing weekly digest is closer to a packaging decision than a new build — which is exactly why backup.md flags it as selling the business better than "AI tutor" alone. It is still [ASSUMPTION]-tagged here because the price point and attach rate are unvalidated.

4. Unit economics [ASSUMPTION – all figures, kept internally consistent]

Baseline assumptions carried through this section:

- Average mentor session price: **\$25/hour** [ASSUMPTION].
- Marketplace take rate: **20%** → **\$5/hour platform revenue, \$20/hour mentor payout.**
- Average paying-student marketplace usage: **2 hours/month** → **\$10/month** marketplace revenue per active marketplace user [ASSUMPTION].
- Desk OS Plus subscription: **\$12/month** [ASSUMPTION].
- Blended ARPU for a paying, marketplace-active student (subscription + marketplace, not additive for every user — some pay only one or the other): **~\$18–22/month** [ASSUMPTION].

Gross margin:

- Marketplace revenue is close to pure margin minus payment processing (~3%) and support/trust-and-safety overhead — assume **~85% gross margin** on the take-rate line [ASSUMPTION].
- Subscription revenue carries AI inference cost (LLM calls for notes/AI tools, ~\$1–2/active user/month at current model pricing) plus hosting — assume **~75–80% gross margin** on the subscription line [ASSUMPTION], in line with typical SaaS.

- Family add-on is close to zero marginal cost (mostly automated report generation off data already collected) — assume **~90% gross margin** [ASSUMPTION].

CAC:

- Campus-grassroots CAC (the GTM model in §6 — direct recruiting, no paid ads): **\$8–15 per acquired active student** [ASSUMPTION], materially below typical tutoring-marketplace paid-acquisition CAC (which runs \$40–80+ for comparable products going through paid channels — cited as an industry ballpark, not a MindCraft-specific figure).
- Mentor-side acquisition cost via the same grassroots channel (flyers, tutoring-center outreach, a small referral incentive): **\$5–10 per onboarded mentor** [ASSUMPTION].

LTV:

- Assume average active tenure of **18 months** (spans multiple semesters, including summer drop-off) [ASSUMPTION].
- $LTV \approx ARPU \times \text{blended gross margin} \times \text{months retained} \approx \$20 \times 0.80 \times 18 \approx \mathbf{\$288}$ [ASSUMPTION].
- Against a **\$8–15 grassroots CAC**, that implies an LTV:CAC ratio in the **19:1–36:1** range — this is almost certainly too good, and the honest read is that campus-grassroots CAC is not a number that survives past the first one or two founder-adjacent campuses (see §8). Treat this ratio as a best-case floor-tester, not a plan input.

Mentor lifetime value to the platform (a second-order number worth tracking, not just student LTV): a mentor doing ~5 sessions/week for two semesters (~28 weeks) at \$5/hour platform take $\approx \mathbf{\$700}$ in platform revenue per mentor over an academic year [ASSUMPTION]. Mentor supply is therefore not just a cost center (payouts) — a retained, high-volume mentor is a real revenue asset, which argues for investing in mentor retention (the "co-author, never gig labor" treatment in Brand Book §6) as a unit-economics lever, not just a brand nicety.

5. Pricing tiers (draft)

Tier	Price [ASSUMPTION]	What's included
Free	\$0	Workspace (notes, project boards), rate-limited AI tools, browse/apply to mentors, one intro session or async Q&A [ASSUMPTION on the intro-session mechanic]
Individual — Plus	\$12–15/month [ASSUMPTION]	Unlimited AI tools, priority mentor matching, one included mentor-session credit/month, unlimited storage for notes/presentations/projects
Family add-on	+\$10/month [ASSUMPTION], stacks on a student plan	Weekly report, weak-concept summary (Solver-linked where applicable), suggested next sessions, receipts of improvement — the backup.md loop

Institutional / school (eventual)	\$8–15/seat/year [ASSUMPTION], custom above a seat threshold	Bulk seats, admin dashboard, aggregate progress insight, integration with the institution's existing tutoring-center mentor supply where one exists
--	---	---

6. Go-to-market

This is not a new plan — it builds directly on `MindCraft_Viability_Research_Strategy.md`, which already lays out a bottoms-up, campus-grassroots research and validation methodology centered on Macalester and the surrounding St. Paul community. That document should stay the operating reference for the research/validation phase; this section only sequences it into a broader GTM arc.

Phase 0 — Validate before scaling (already planned, not yet run at time of writing). Run the 15-conversation protocol from `MindCraft_Viability_Research_Strategy.md`: 5 students, 5 parents, 5 tutors/mentors, sourced from the locations that document already names (Macalester campus green, library, Dunn Bros/Black Coffee & Waffle Bar on Grand Ave, Summit Ave/Crocus Hill door-to-door, high school pickup zones, the Macalester tutoring center/TA offices). Output: affinity-mapped jobs-to-be-done per stakeholder, and a kill/keep/build list — feeds directly back into §3's revenue-stream prioritization and §8's risk list.

Phase 1 — One-campus bootstrap (Macalester). Recruit mentor supply first (tutoring center, TA office bulletin boards, the \$5 coffee-gift-card flyer approach already specified in the research doc), because a workspace with no one on the other end of it has nothing to sell. Recruit students through the same low-friction, in-person channels. Parents reached via the door-to-door/pickup-zone channels already mapped for the reporting add-on specifically — that's the segment most likely to pay for §3's family stream first.

Phase 2 — Replicate the playbook at 2–4 similar campuses. Same channel mix (campus-center recruiting, in-person conversations, no paid acquisition), chosen for similarity to Macalester (residential, small-to-mid liberal-arts, walkable off-campus student neighborhood) so the playbook transfers with minimal adaptation. This phase is where the CAC assumption in §4 gets its first real test outside a founder's personal network.

Phase 3 — Broaden campus type + open the institutional conversation. Once mentor-supply liquidity and retention are proven at 4–5 campuses, expand to larger state schools and to high schools (where the Solver vertical already has product-market fit signal) and begin institutional-licensing conversations (§3, §5) with academic-support offices at campuses where organic usage is already meaningful — sell the license into the community that already exists, not into a cold institution.

7. Competitive landscape & moat

Competitor type	Example	What it optimizes for	Why MindCraft inverts it (Brand Book §12)
-----------------	---------	-----------------------	---

Tutoring marketplace	Wyzant-style hour marketplaces	Selling hours, matching on availability/price	Sells the click, not the hour — the human session sits inside a persistent workspace with context, not a transaction that ends when the timer does
Productivity suite	Notion, Google Workspace	Infinite blank tools, self-serve, no human in the loop	Every tool is one tap from a live mentor; a better blank page has never gotten anyone unstuck
AI-only tutor	Khanmigo-style AI tutoring	Scaling a single AI interlocutor across every student	Combines AI's scalability with a real human's judgment, encouragement, and read on where a specific student is stuck — the elevator pitch's explicit argument for why AI alone isn't enough
Content library	Khan Academy	Organizing the world's material for the already-motivated	Manufactures the motivation itself via live human connection, for the student who wasn't going to open a library on their own
Gamified habit app	Duolingo	The mechanic (streaks) as the product	The work itself, and the person helping with it, are the product — no engagement-bolted-on boredom mechanic

Moat, in order of durability:

1. **The mastery/evidence graph.** Solver's knowledge graph (per-student mastery, misconception tracking, gap detection — real, shipped, described in `CLAUDE.md`) is the deepest technical asset and the hardest to replicate quickly; it's also what makes the family-reporting stream (§3) close to a packaging exercise rather than new engineering. Extending an analogous "context object" to non-math mentoring sessions (what a student already tried, where they're stuck) is the workspace-wide version of the same moat.
2. **Workspace lock-in.** Once a student's notes, project history, and mentor relationships live on MindCraft, switching cost rises the way it does for any workspace tool — this is the productivity-suite competitors' own moat, turned against a category (tutoring marketplaces) that currently has none.
3. **Trust layer / mentor community.** Verified mentors, session history, and the "co-author, not gig labor" mentor relationship (Brand Book §6) compound over time in a way a pure marketplace's anonymous hour-sellers don't.
4. **Campus-native distribution.** The grassroots GTM in §6 is slow but close to unpaid-CAC; a VC-funded competitor optimizing for paid-channel growth has a structurally worse cost basis in the same early market.

None of these are validated moats yet — they're the intended shape of one, and #1 is the only one with real evidence behind it today (a working knowledge graph). #2–#4 are theses that Phase 0–2 of §6 should be testing.

8. Risks & what needs validation

Ordered roughly by how much this business model depends on each one, most load-bearing first. `MindCraft_Viability_Research_Strategy.md` is the existing plan to de-risk the first three.

1. **Supply-side liquidity outside math.** Solver has real mentor supply and product-market signal. Does "any mentor, any subject, on demand" hold once you leave math — or is coverage thin outside a few popular subjects, undermining the "different mentor for different needs" pitch in §14 of the Brand Book? Needs direct validation via the tutor-interview protocol in the research doc, expanded to ask specifically about subject breadth, not just willingness to tutor.
2. **Willingness to pay for workspace features.** Notion and Google Workspace have trained students to expect notes/docs/presentation tools for free. The model here leans on the marketplace take rate and family add-on to carry revenue rather than the workspace tools themselves — that bet needs the research doc's parent and student questions ("what would make you switch," "how much do you currently spend") answered before it's trusted.
3. **The family/parent reporting add-on's actual demand and price sensitivity.** `backup.md` asserts this "sells the business better than AI tutor alone," and the research doc already has the right parent-facing questions queued up (what would need to be in a session summary, what earns 30-day trust) — but nobody has run those conversations yet as of this document.
4. **CAC durability past founder-adjacent campuses.** The [ASSUMPTION]-tagged \$8–15 grassroots CAC in §4 is very likely a first-campus number, not a steady-state one. The LTV:CAC ratio in §4 is flagged there as too good to plan around — Phase 2 of §6 is the actual test.
5. **On-demand supply/demand matching at the moment a student is stuck.** The core promise ("show what you're working on, instantly enter a workspace with a mentor") requires enough mentors online at the right moment. Cold-start liquidity risk is highest late at night and in low-supply subjects — exactly when the emotional need (per Maya's 2am scene) is highest.
6. **Every rate in §3–§5 is invented.** Session price, take rate, subscription price, and licensing price are placeholders chosen to be internally consistent, not researched. Real comparable pricing (existing tutoring marketplace rates, what parents currently pay per `MindCraft_Viability_Research_Strategy.md`'s own question "how much do you currently spend on tutoring per month?") should replace all of them before any external use.
7. **Institutional sales cycle.** Untested and likely slow (higher-ed and K-12 procurement cycles run long) — §6 deliberately sequences this last and treats it as opportunistic (sell into an already-organic community) rather than a near-term revenue plan.